

ADVANCES IN
COMPUTER
VISION

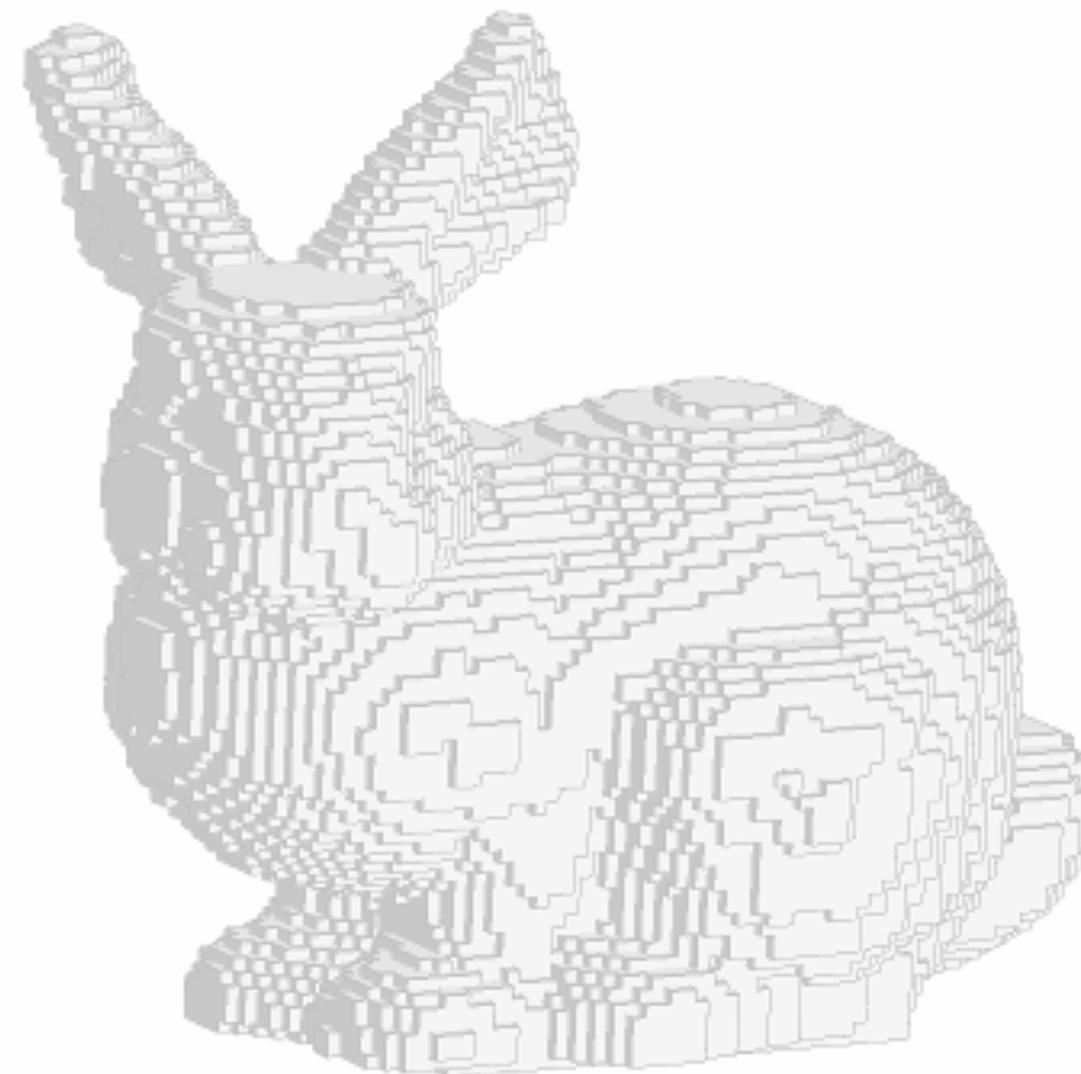
MODULE 1:
Geometry

Differentiable Rendering I

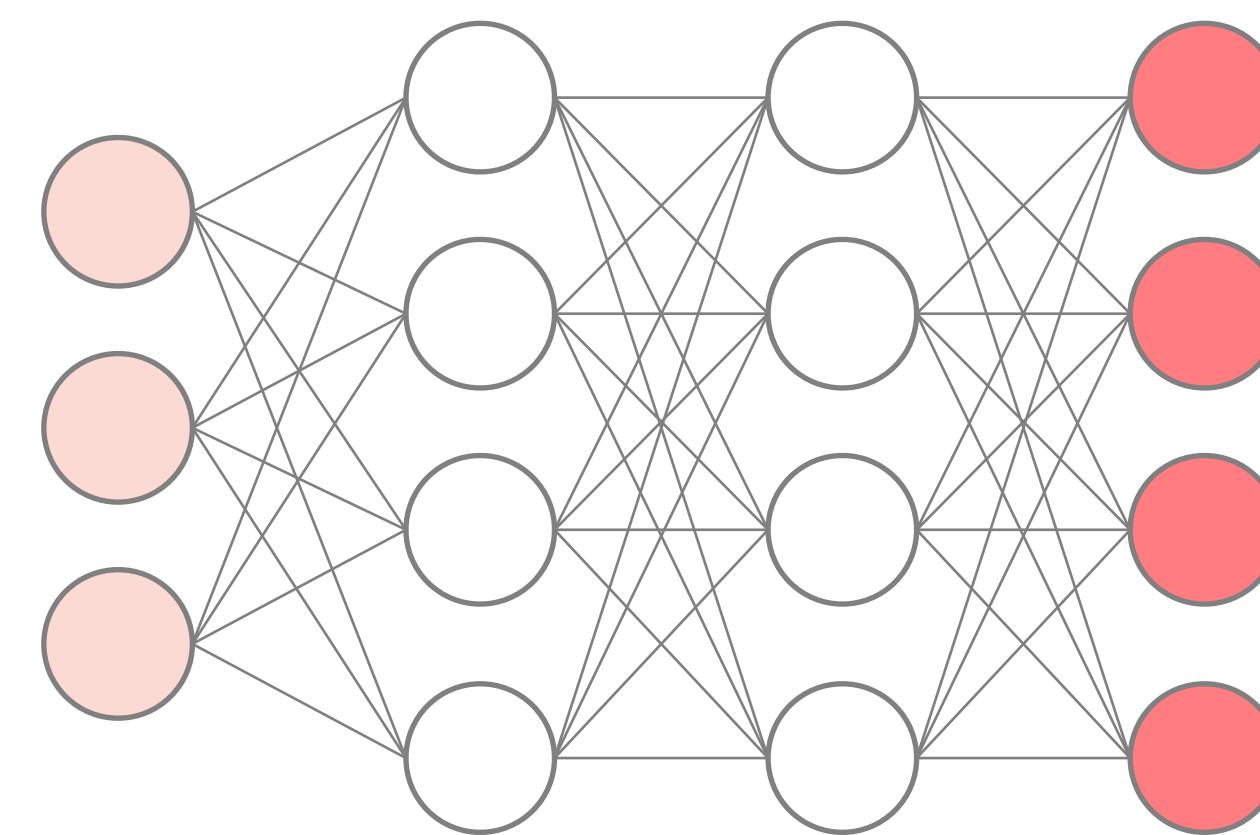


Prof. Vincent Sitzmann

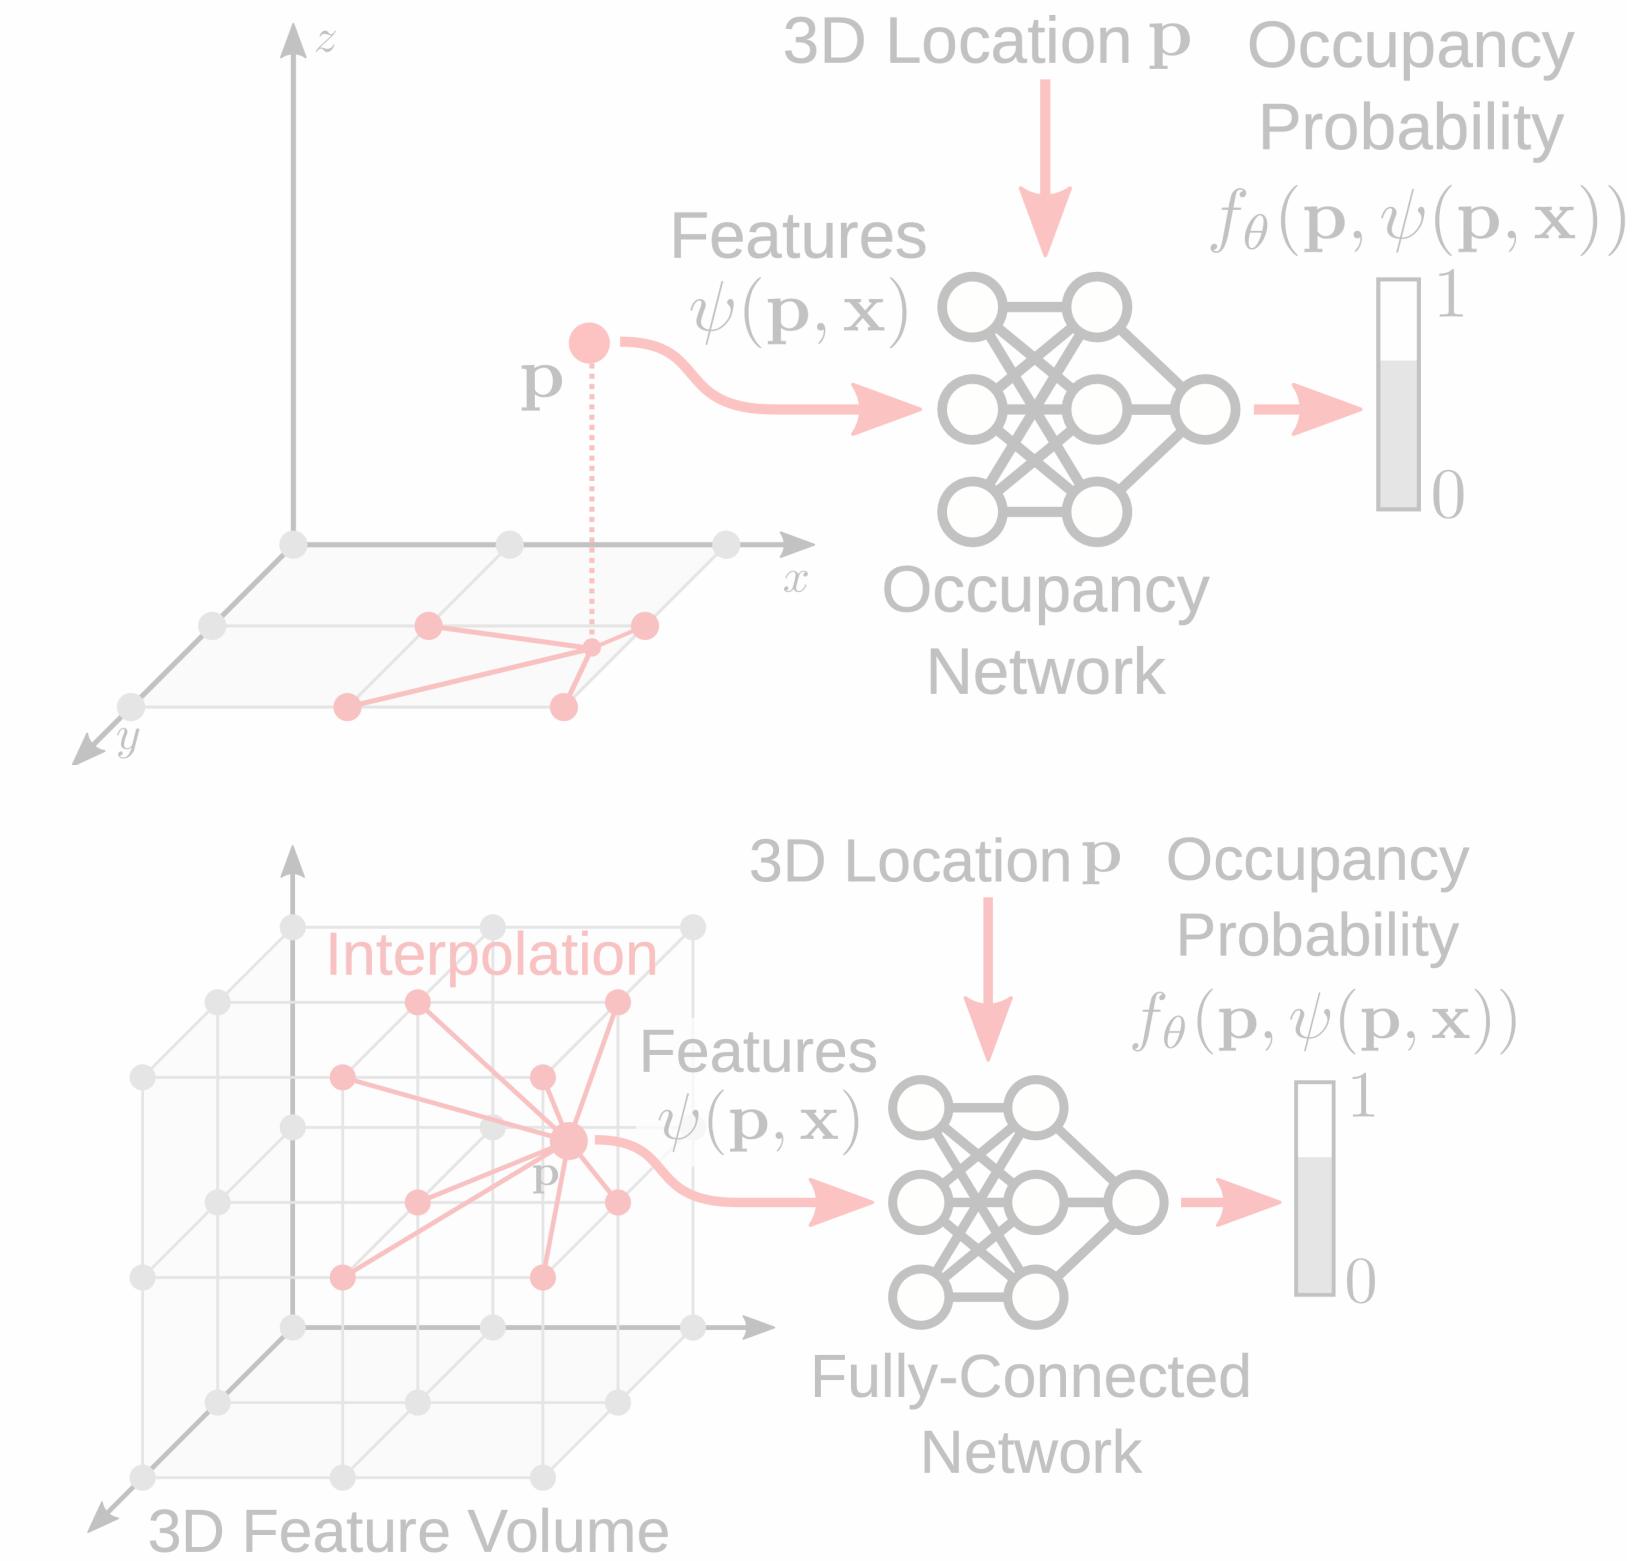
Neural Fields



Discrete Parameterizations

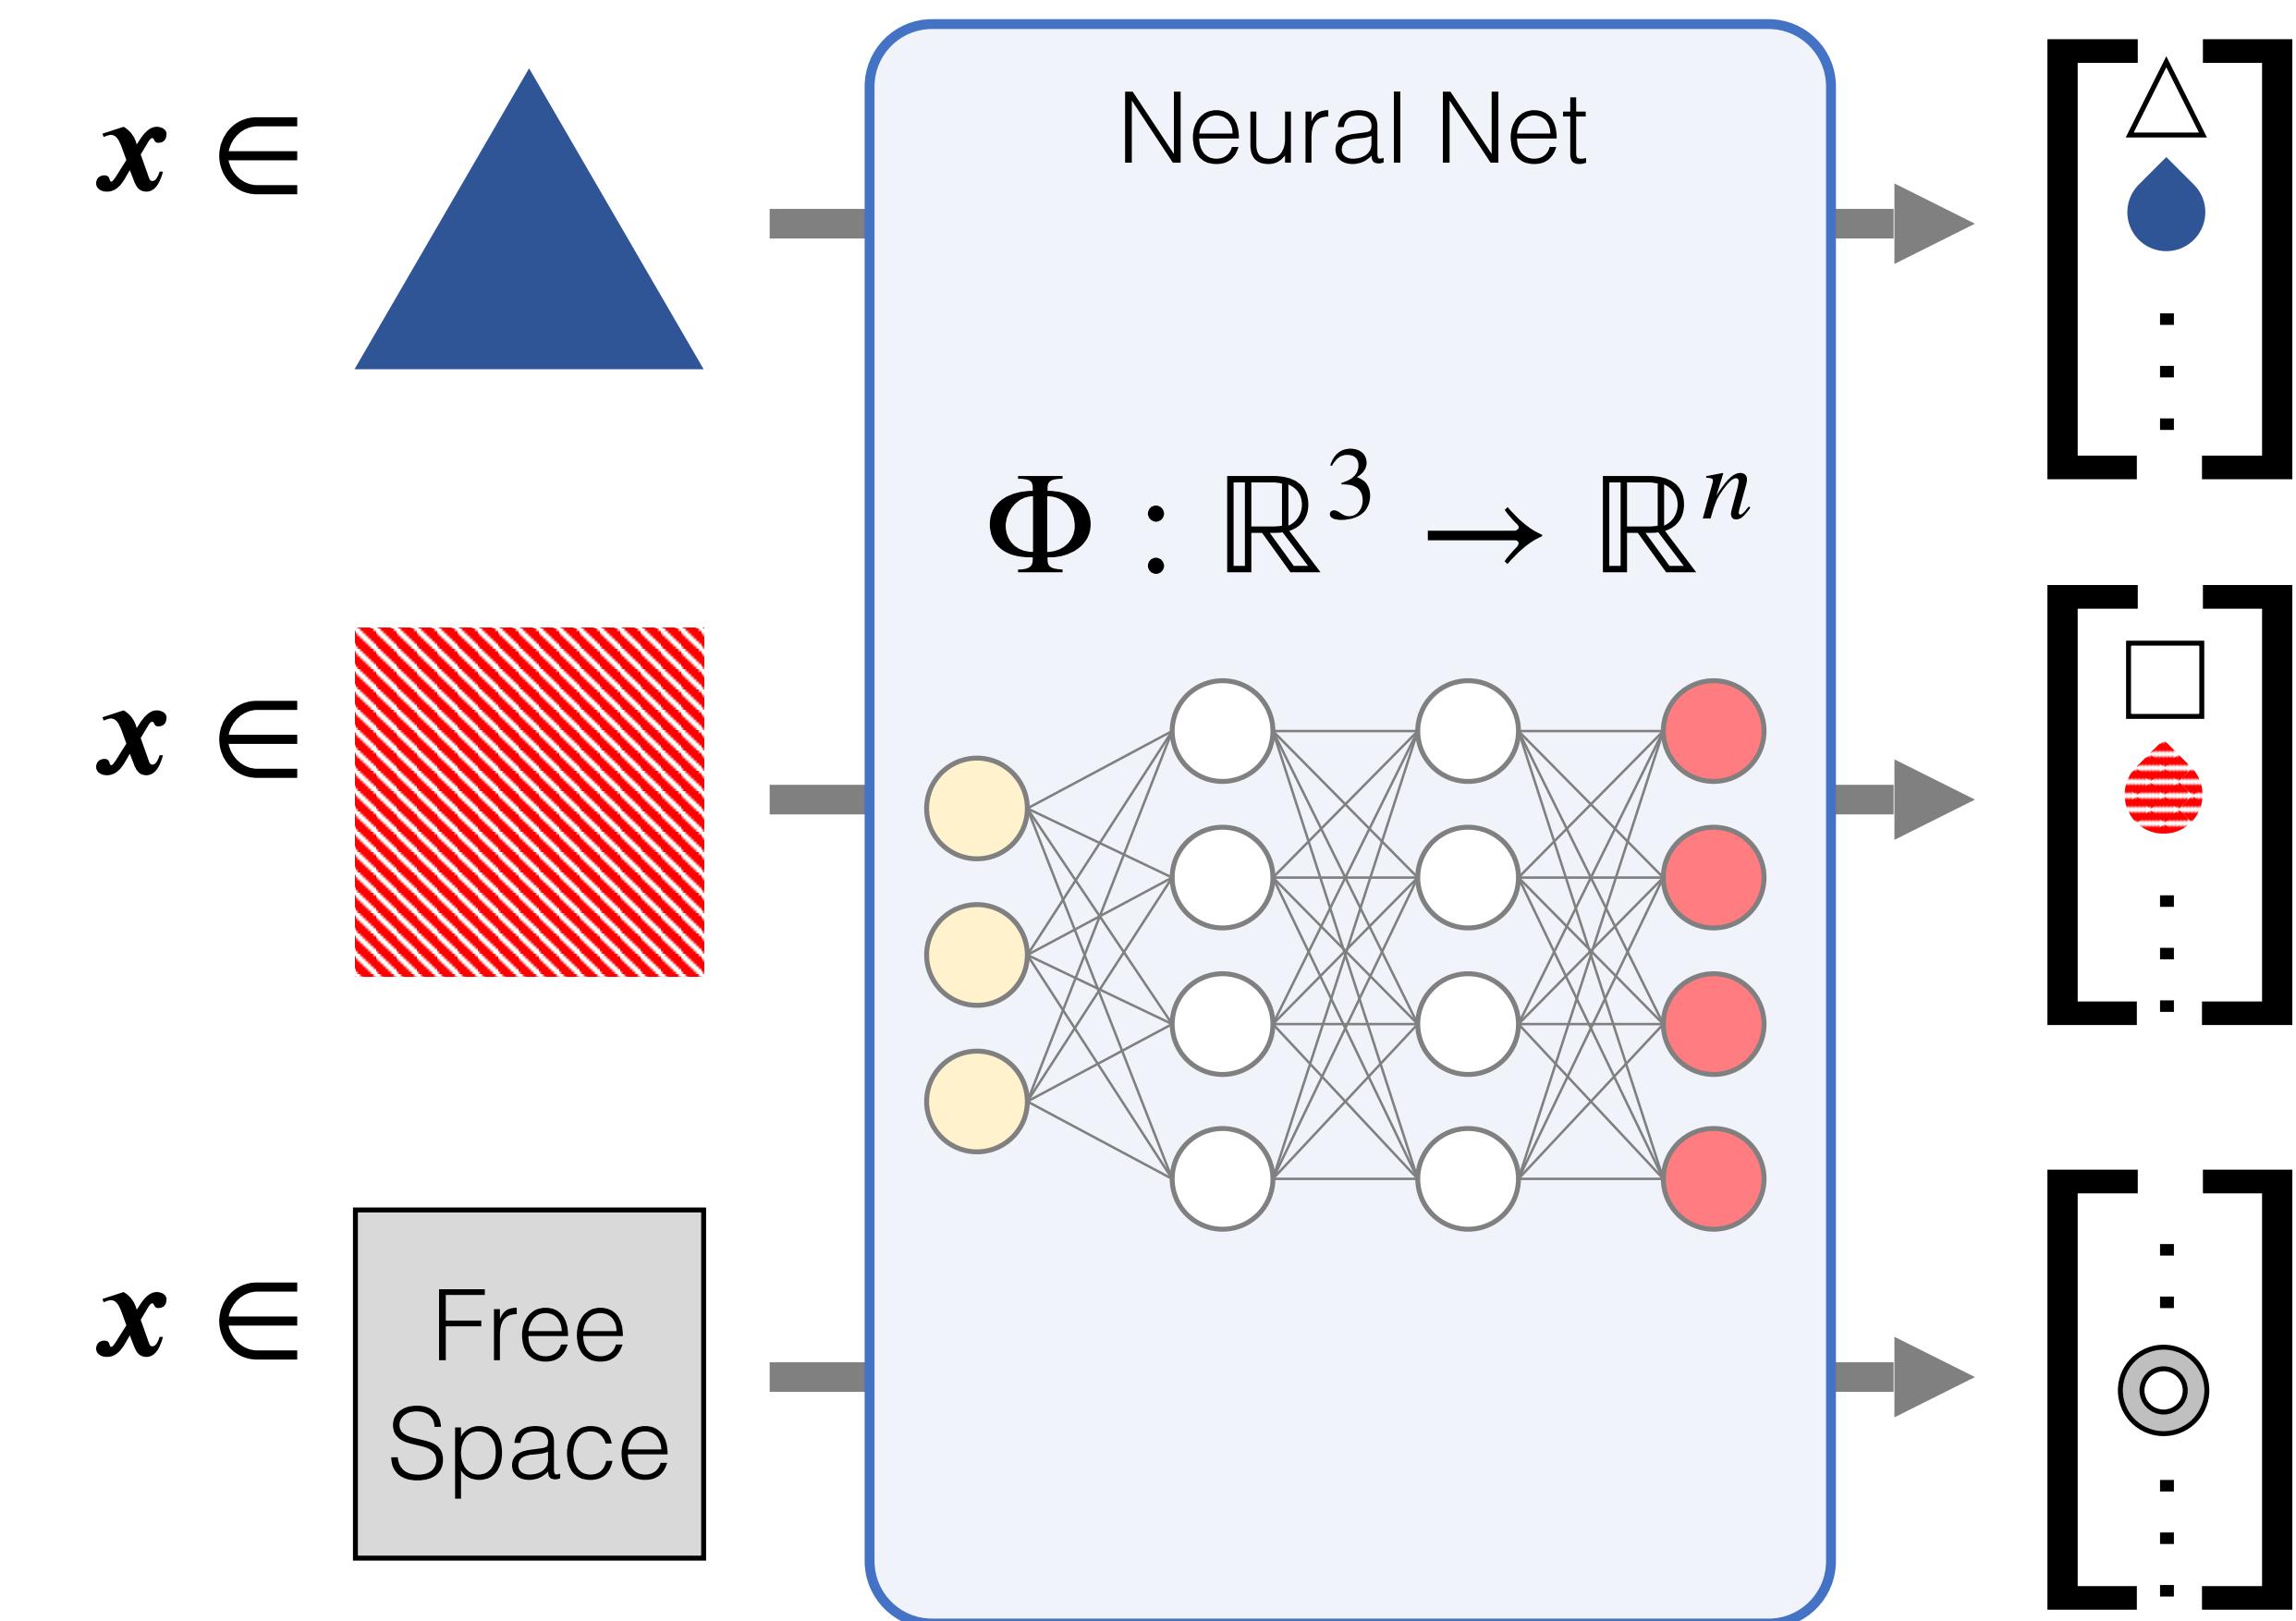
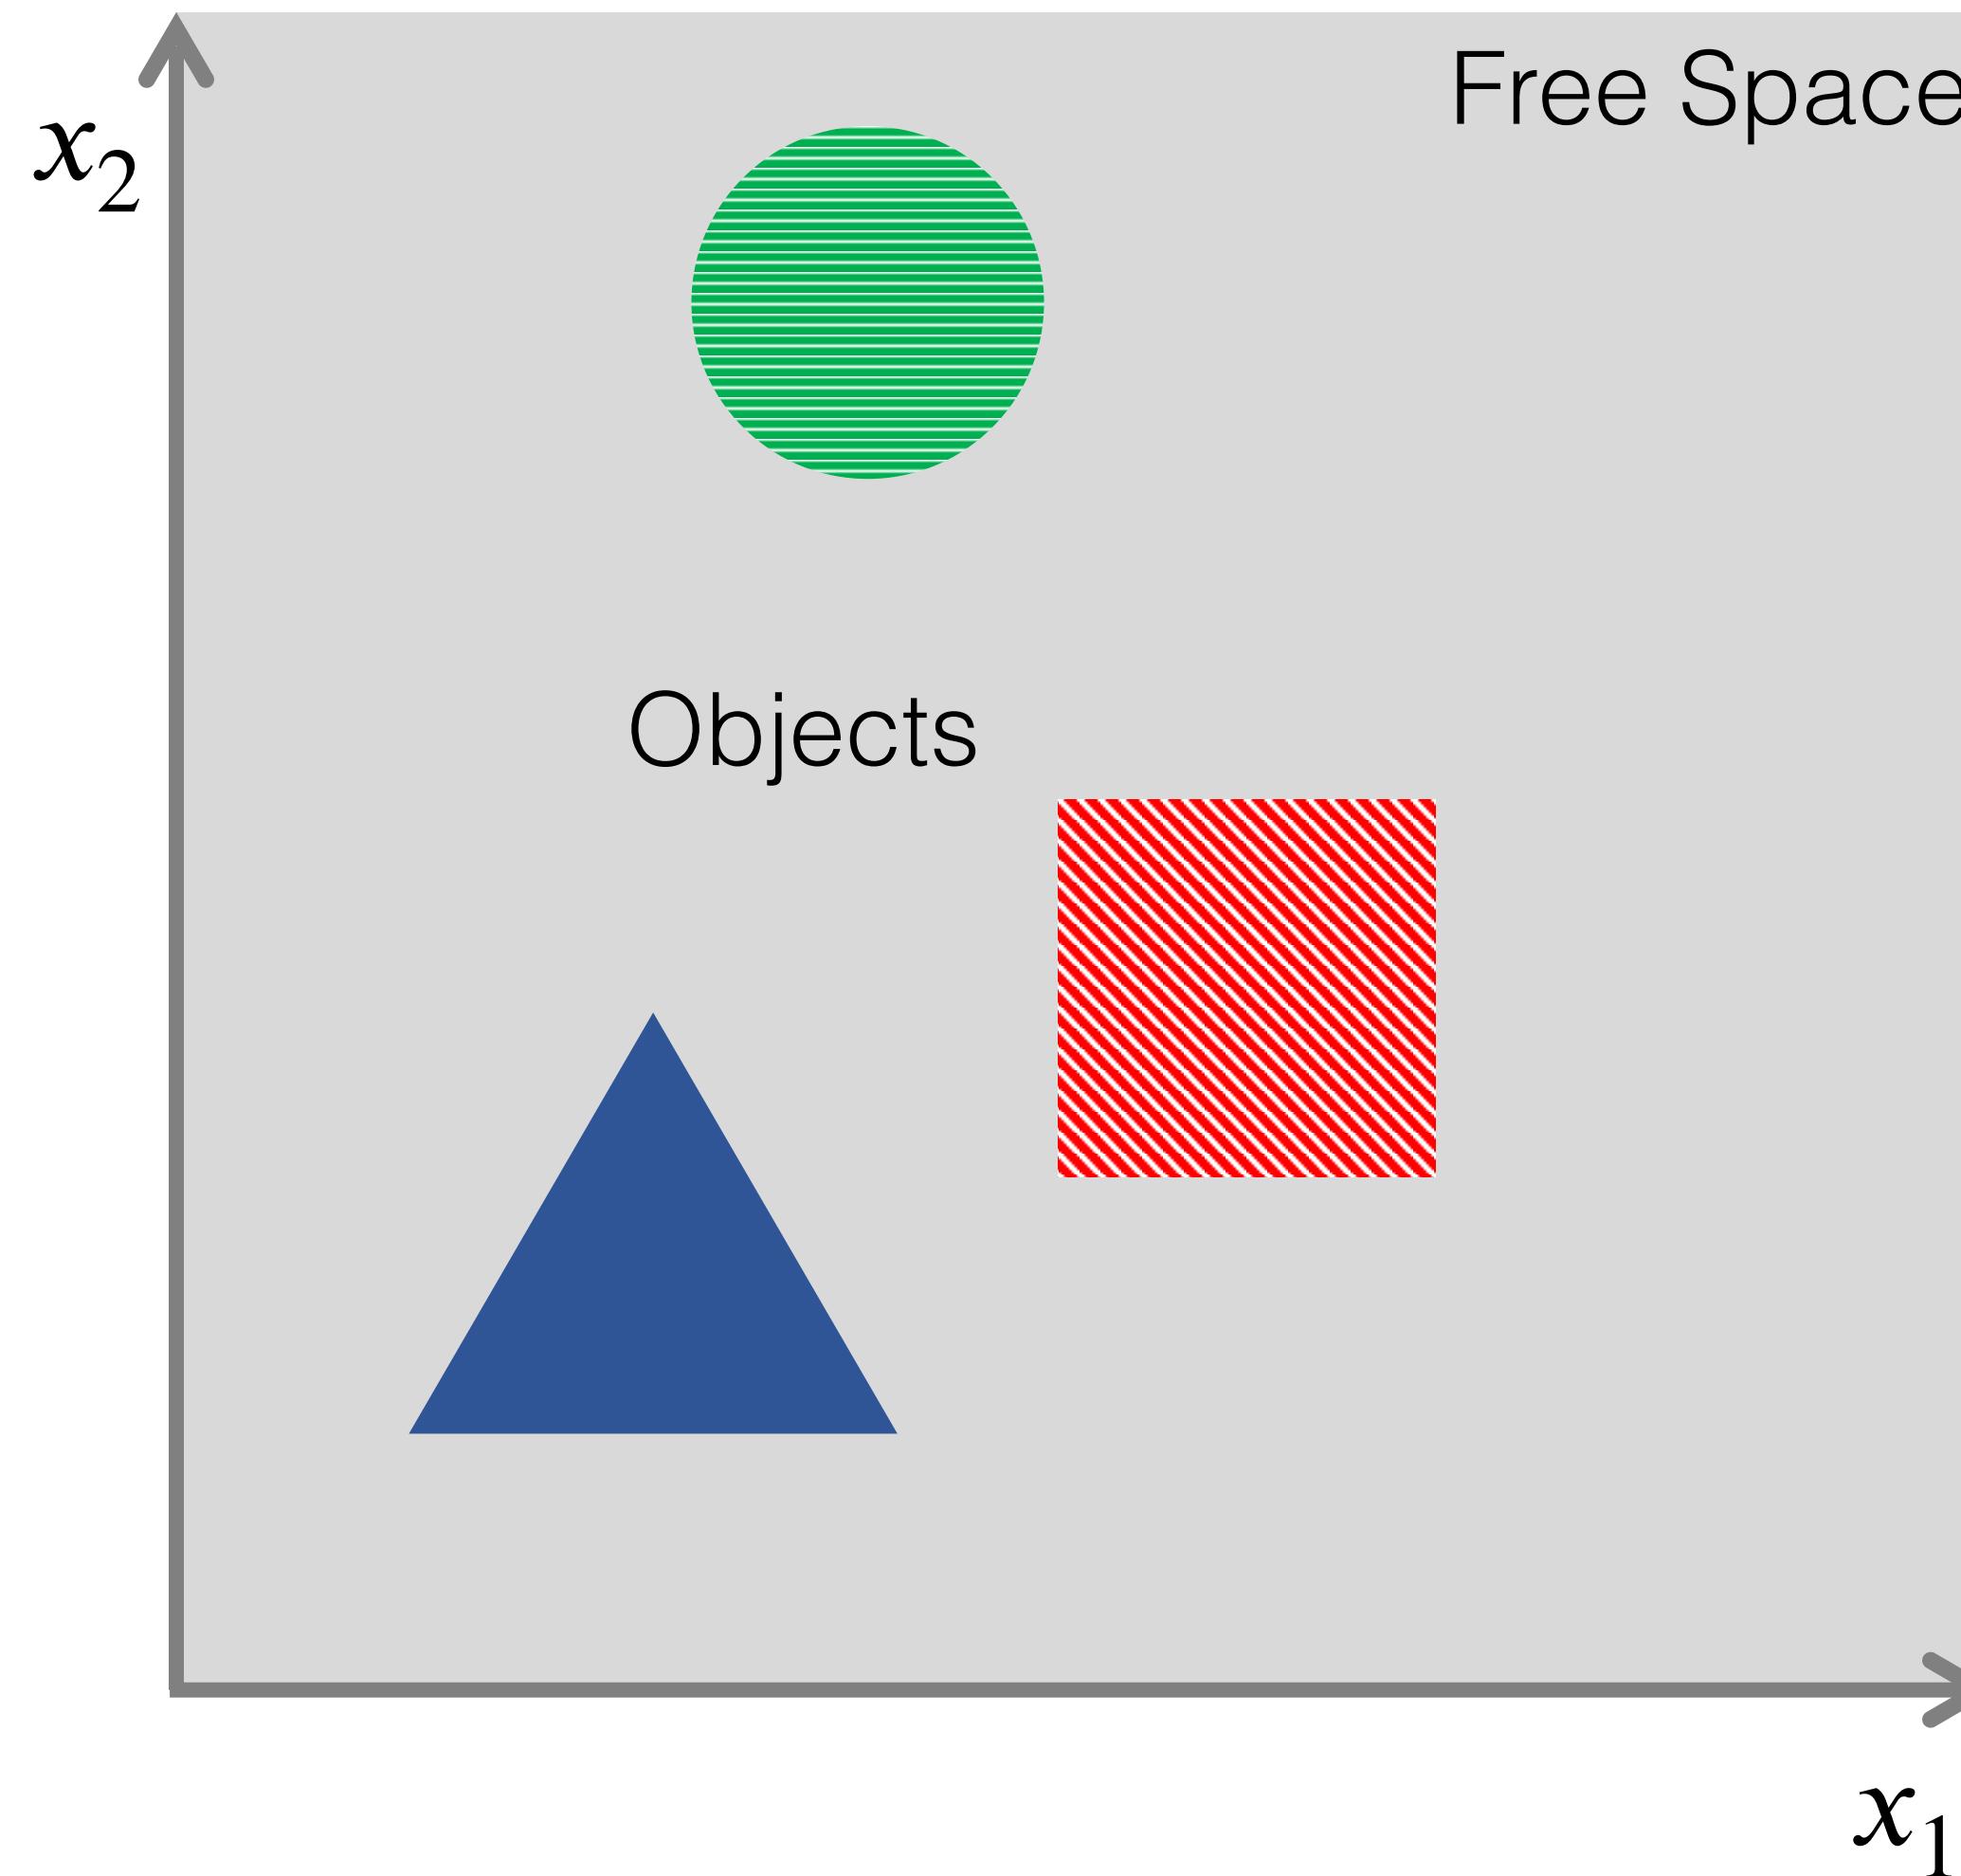


Neural Fields



Hybrid Parameterizations

Neural fields: Parameterize Field as Neural Network!



Occupancy Networks: Learning 3D Reconstruction in Function Space, Mescheder et al.

IM-Net: Learning Implicit Fields for Generative Shape Modeling, Chen et al.

DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation, Park et al. 2019

Sitzmann et al: Scene Representation Networks: Continuous 3D-Structure-Aware Neural Scene Representations, NeurIPS 2019.

Background: Kernel Machines

Given “training” set $\{(x_i, y_i)\}_i$, a kernel machine makes predictions on a point \mathbf{x} by interpolating labels y_i in the training set according to pairwise weights between x_i to \mathbf{x} as measured by a kernel function:

$$f(\mathbf{x}) = \sum_{i=1}^n (\mathbf{K}^{-1}\mathbf{y})_i k(\mathbf{x}_i, \mathbf{x})$$

Where $\mathbf{K}_{ij} = k(\mathbf{x}_i, \mathbf{x}_j)$, k is kernel function with $k \geq 0$.

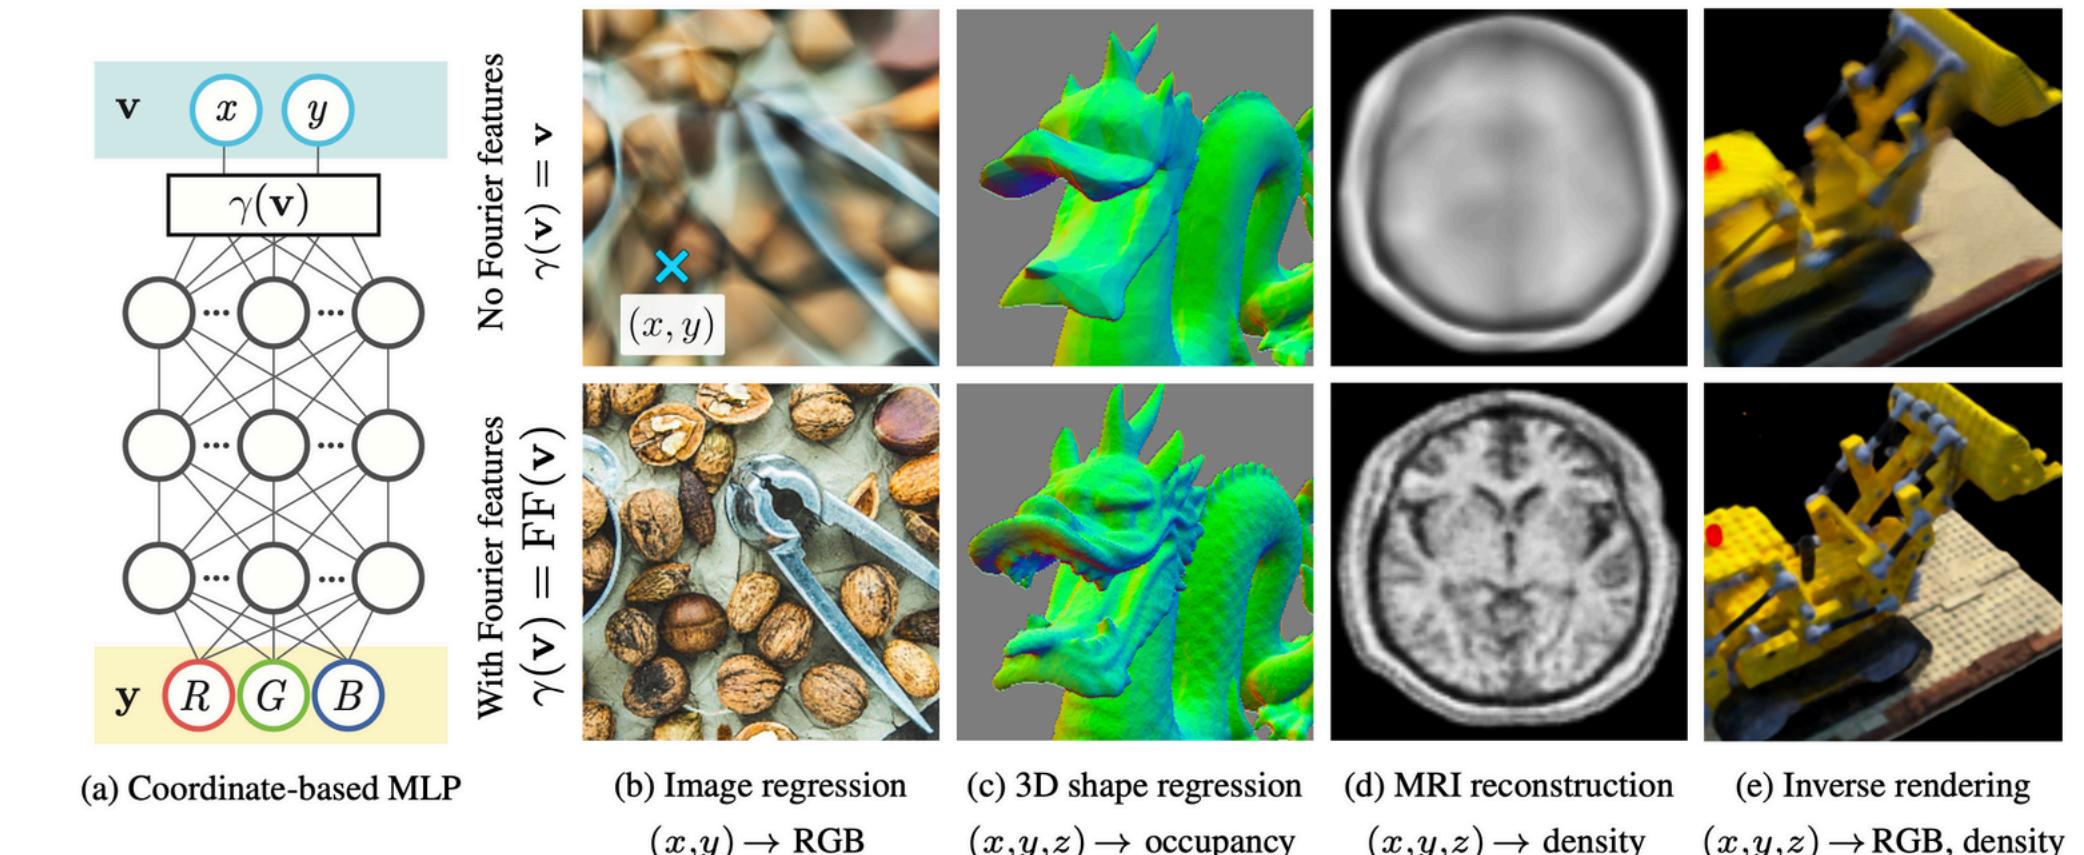
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



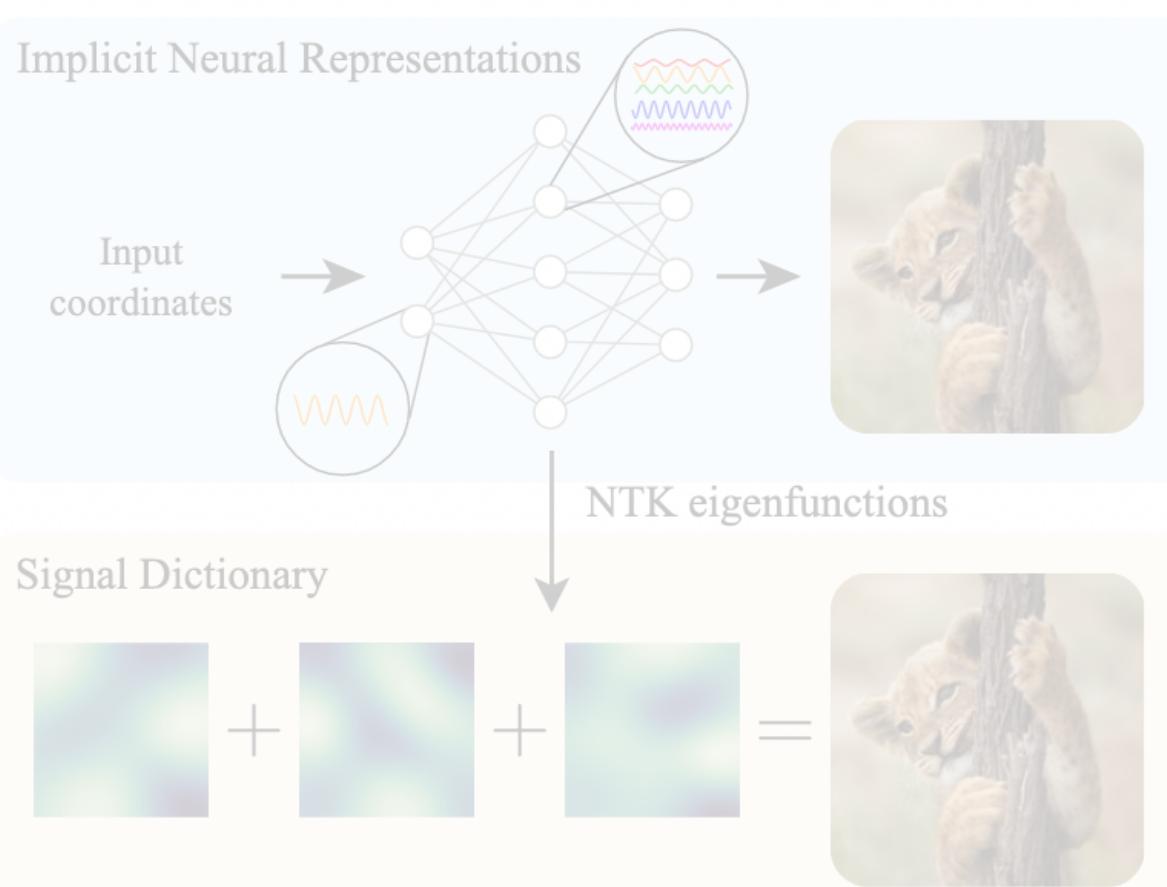
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall et al, NeurIPS 2020



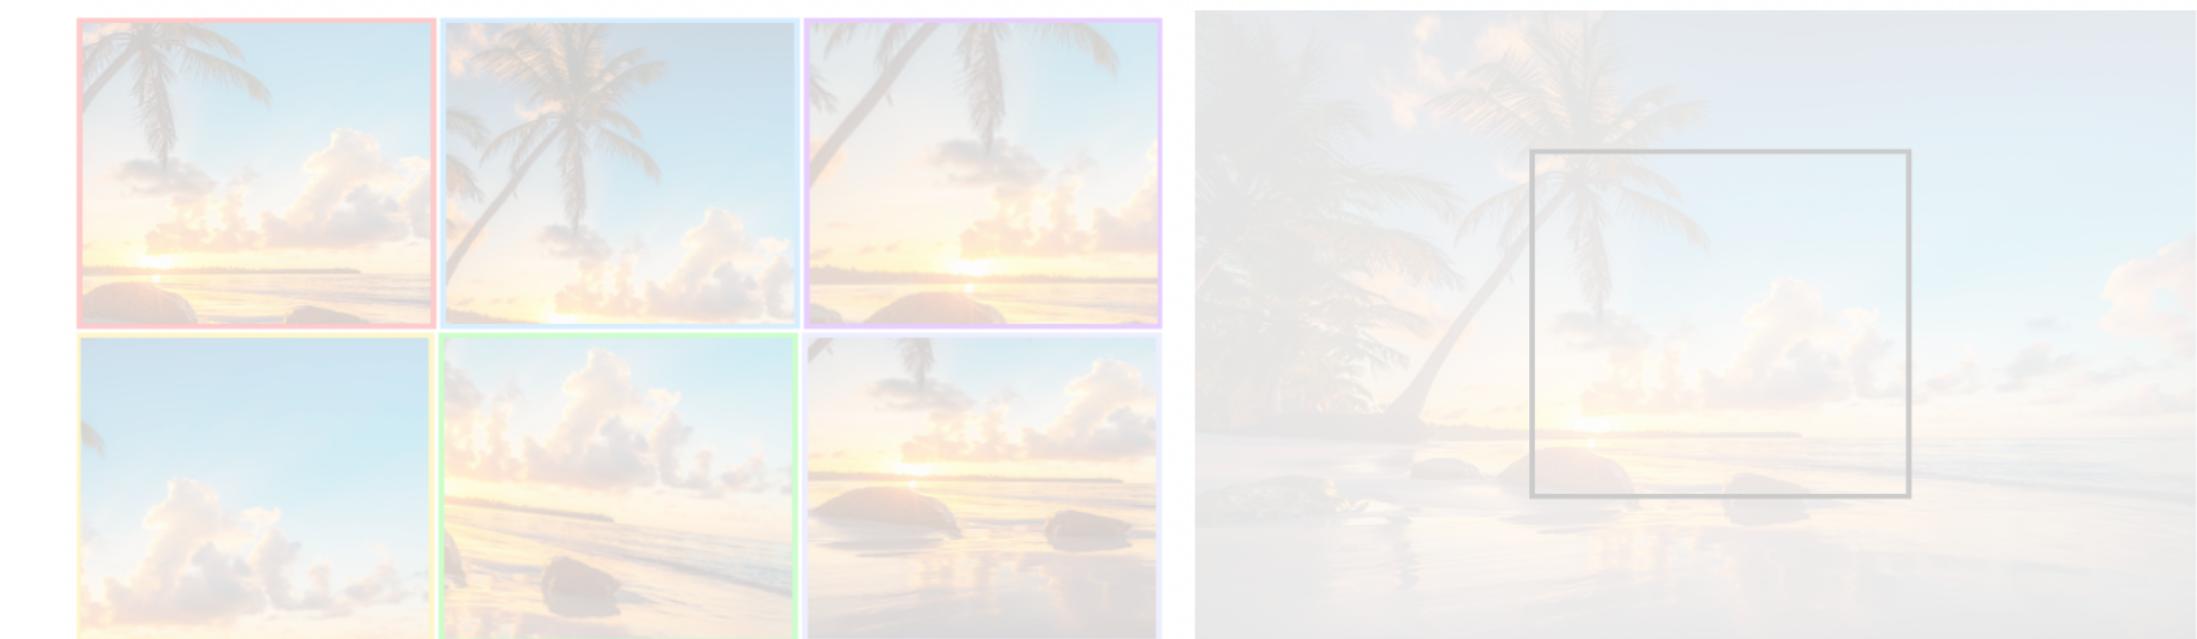
A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



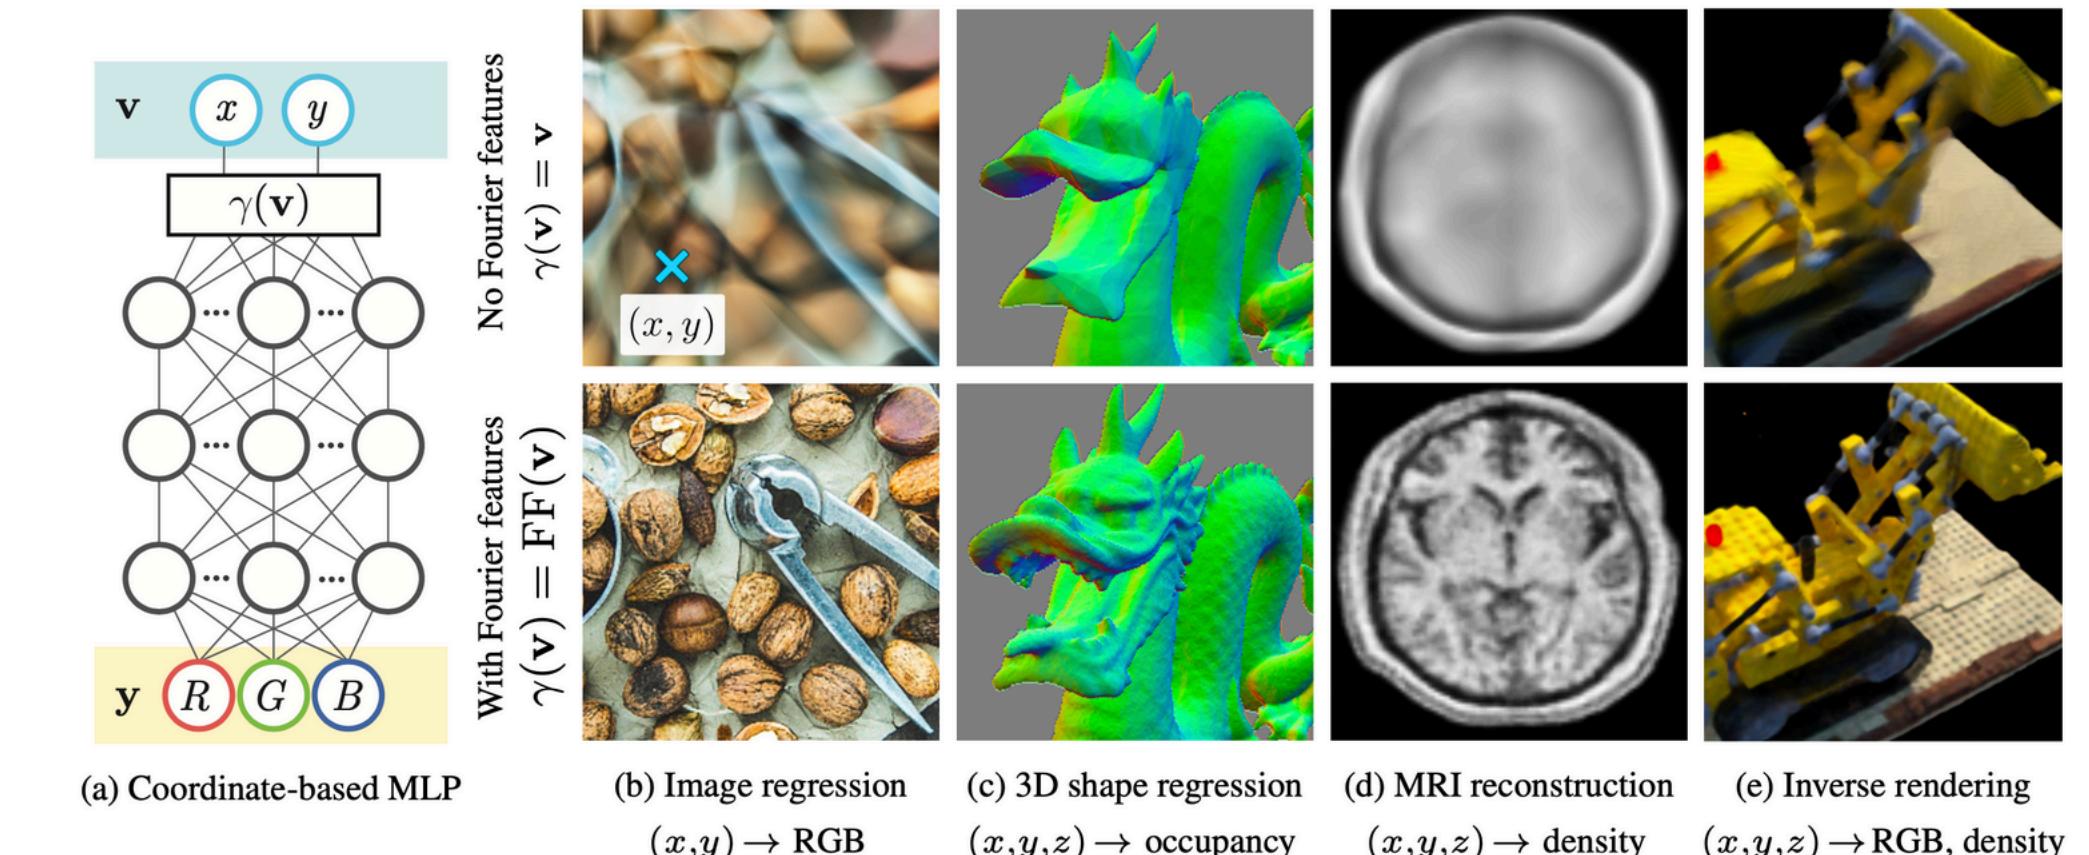
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



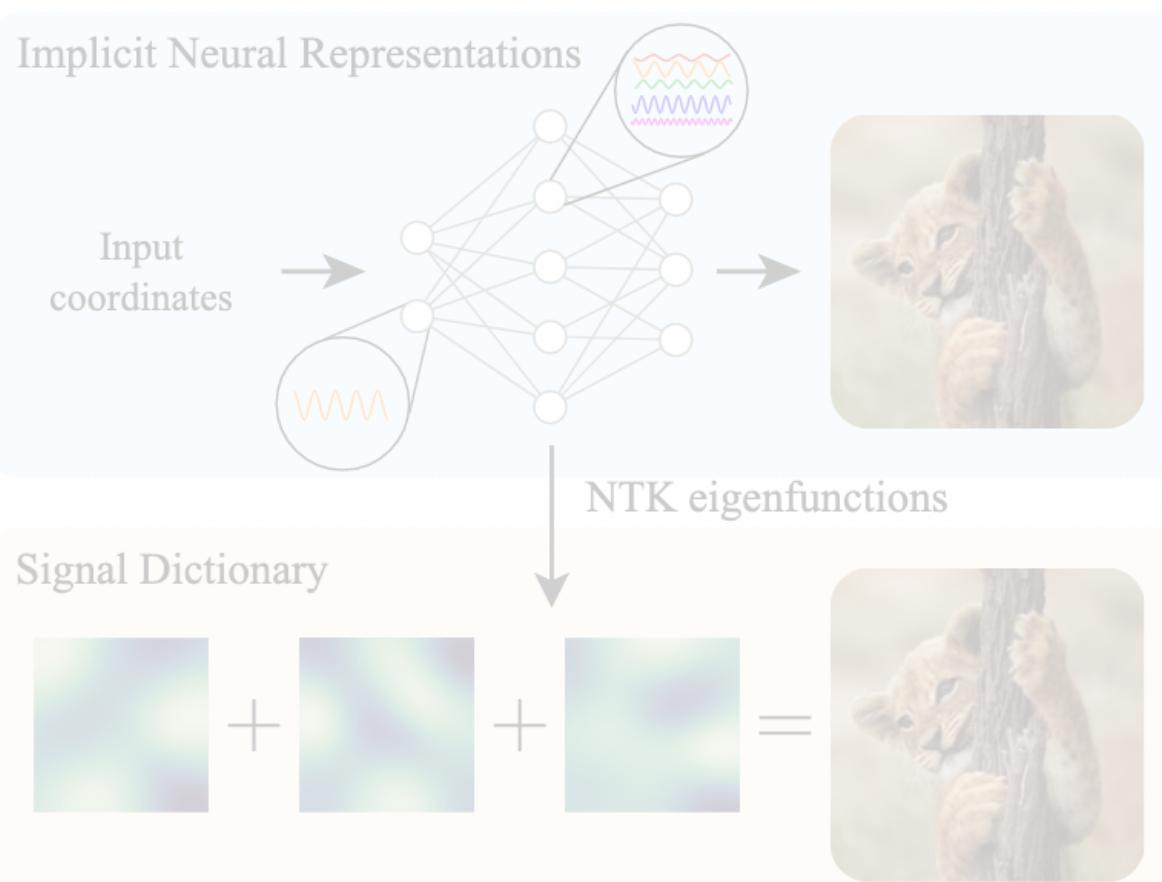
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall et al, NeurIPS 2020



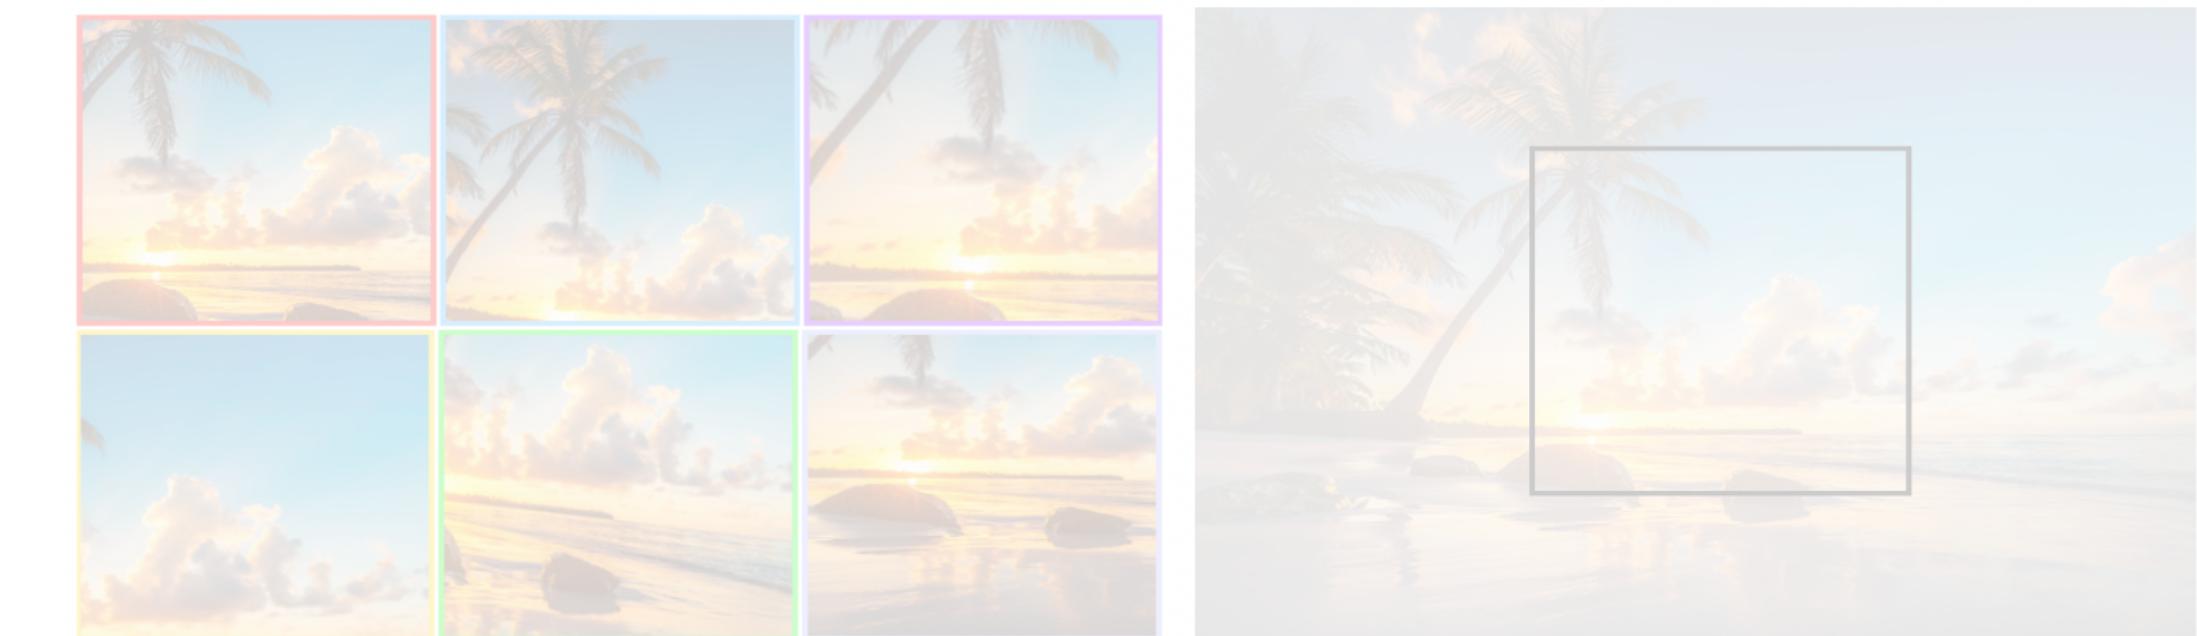
A Structured Dictionary Perspective on Implicit Neural Representations

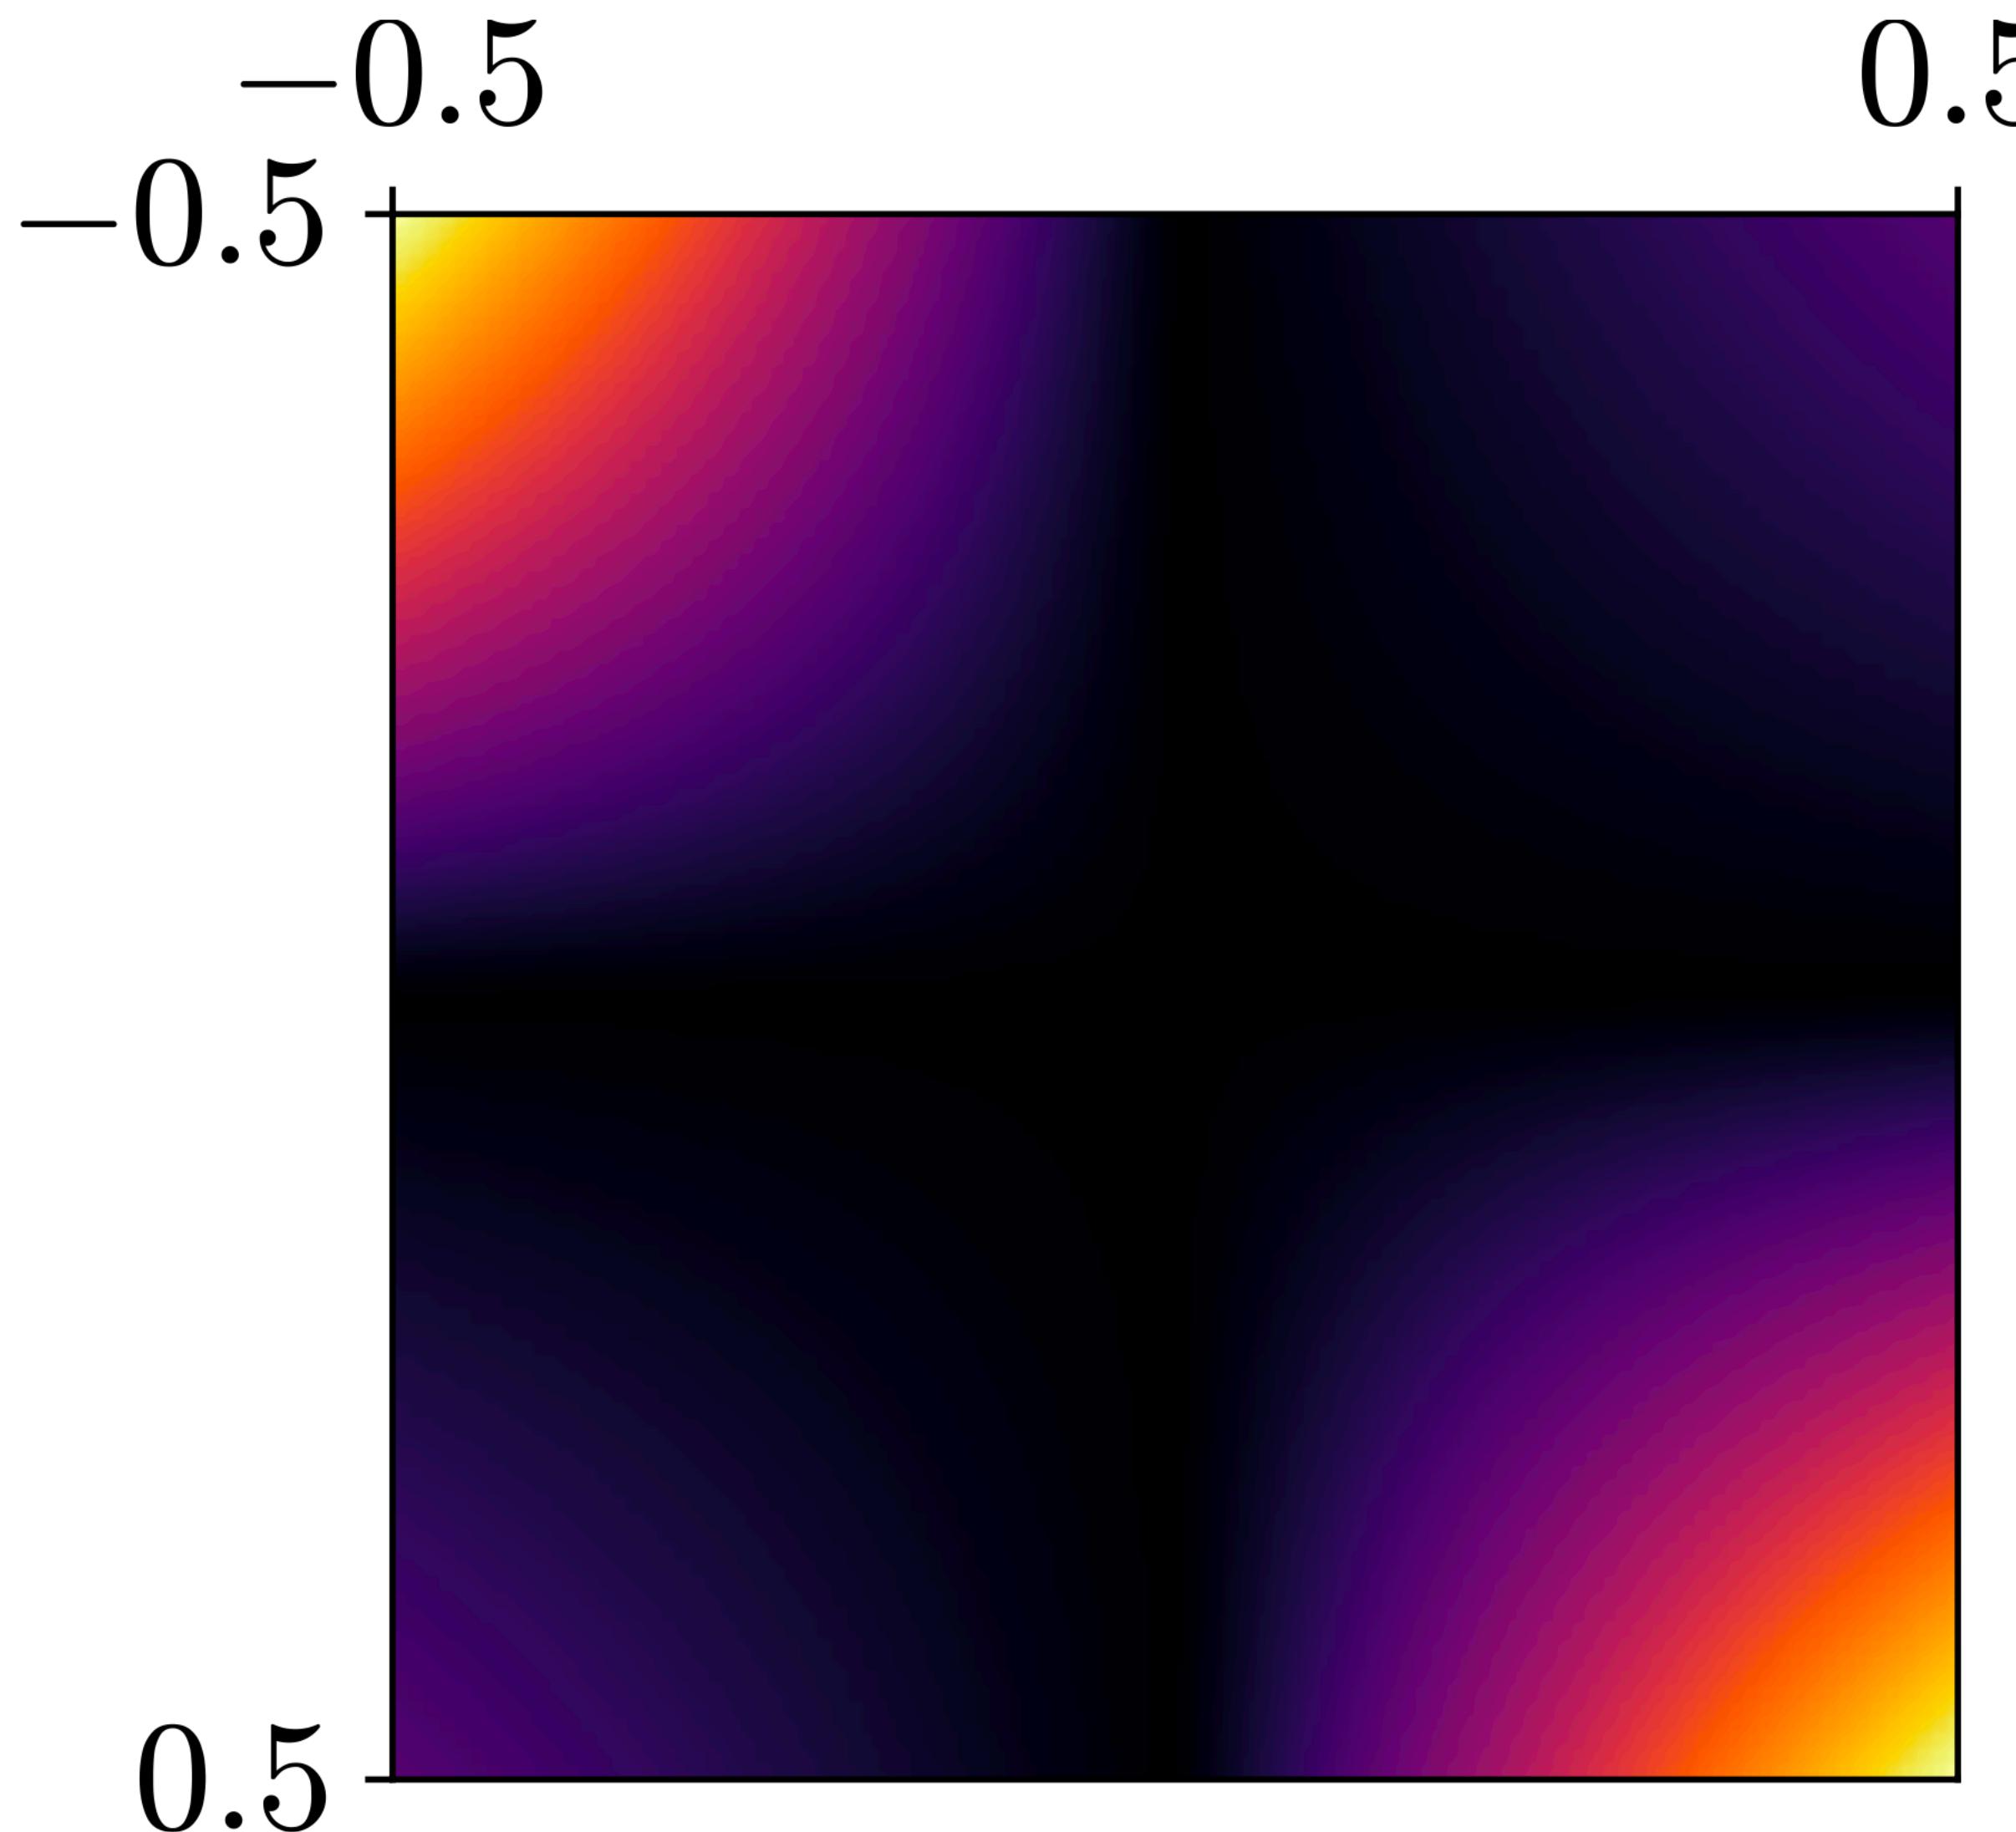
Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



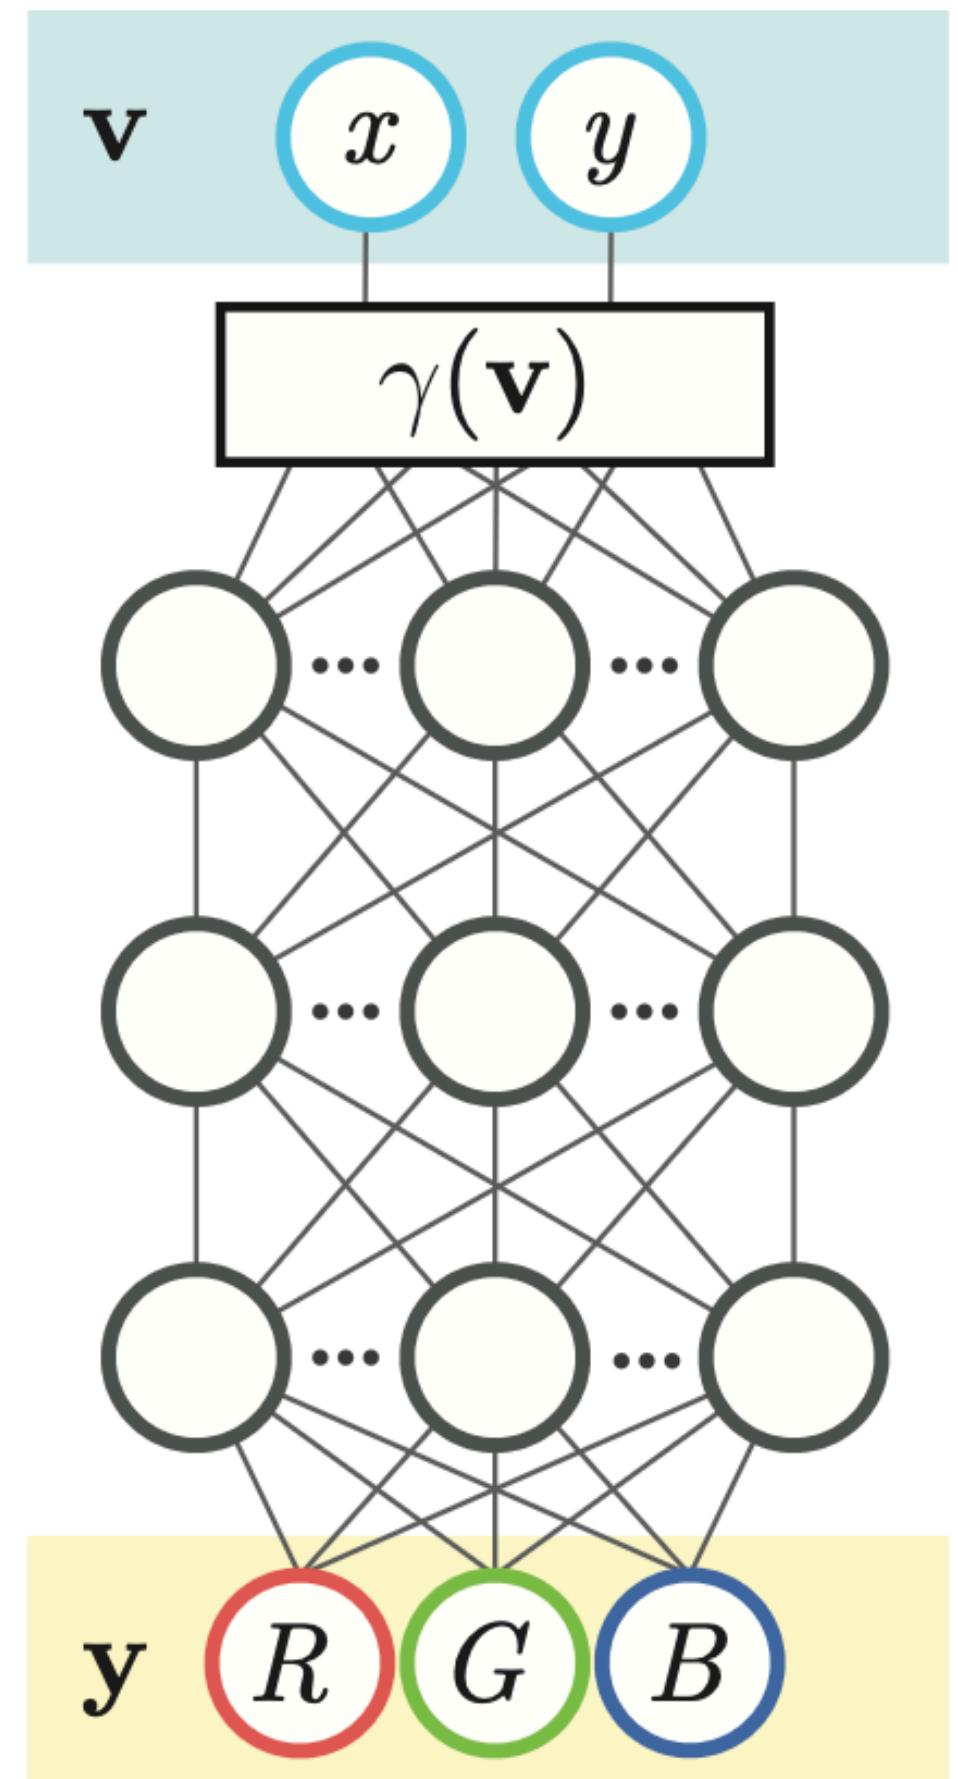


NTK kernel $k_{NTK}(x_i, x_j)$ for a 4-layer
ReLU MLP with one scalar input.

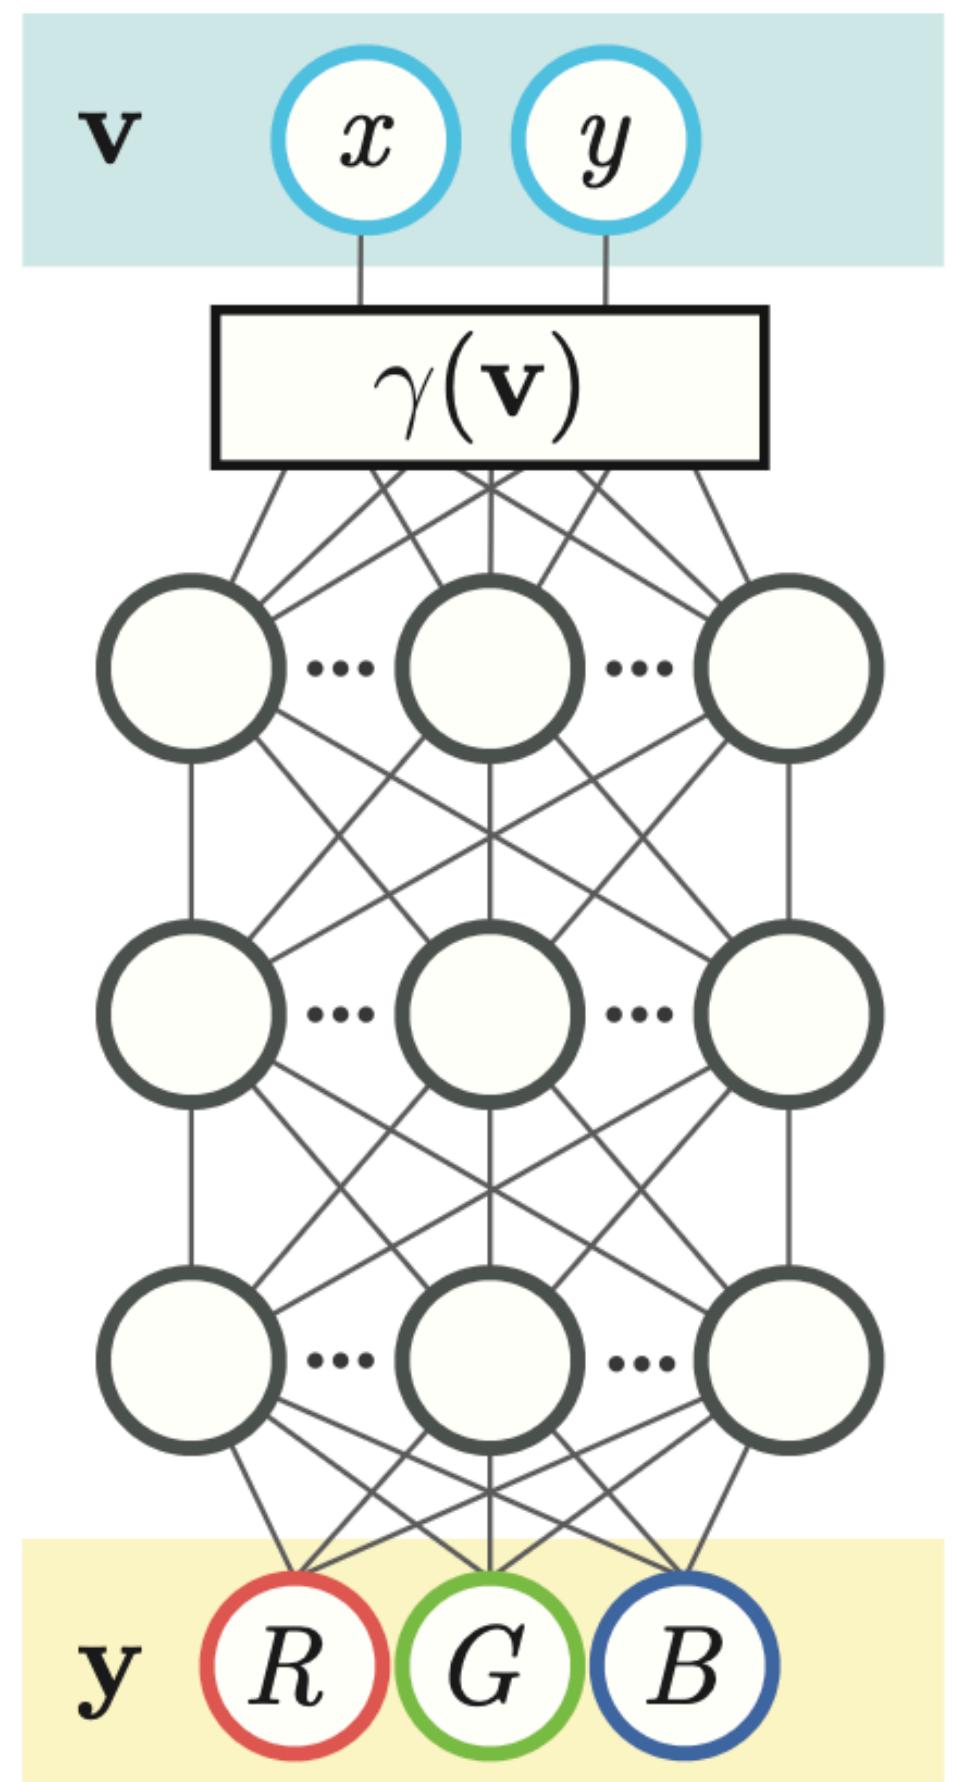
From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

NTK of ReLU neural net is quite non-local.

Feature Embeddings

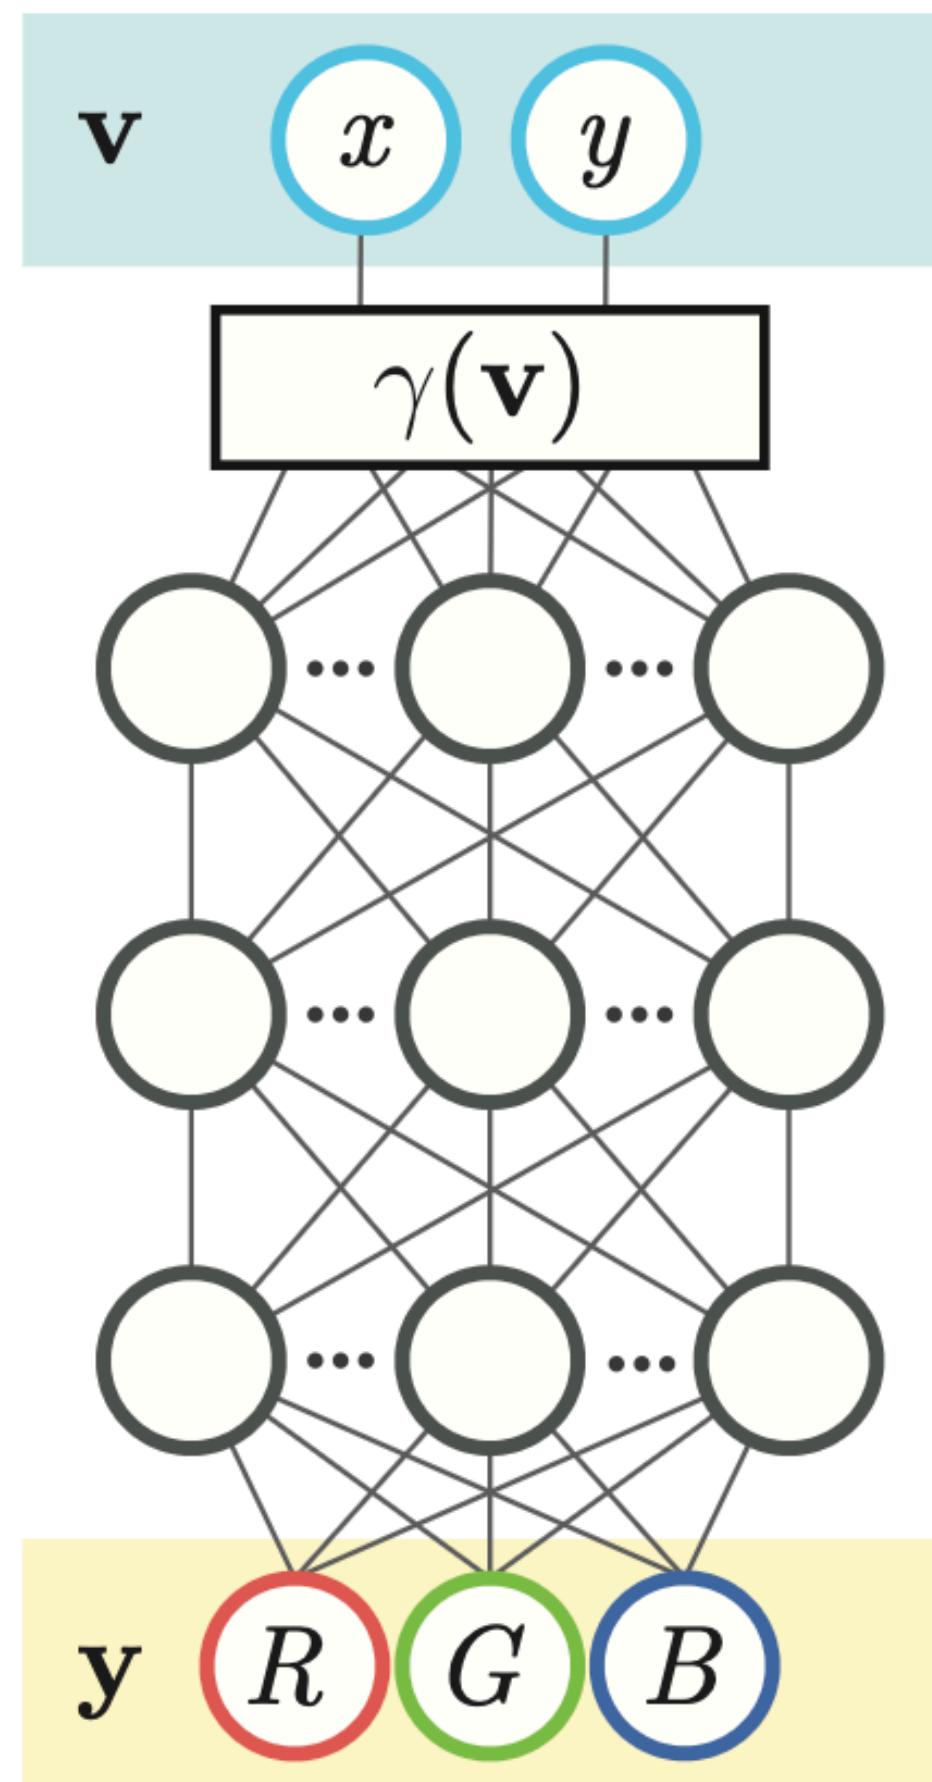


Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

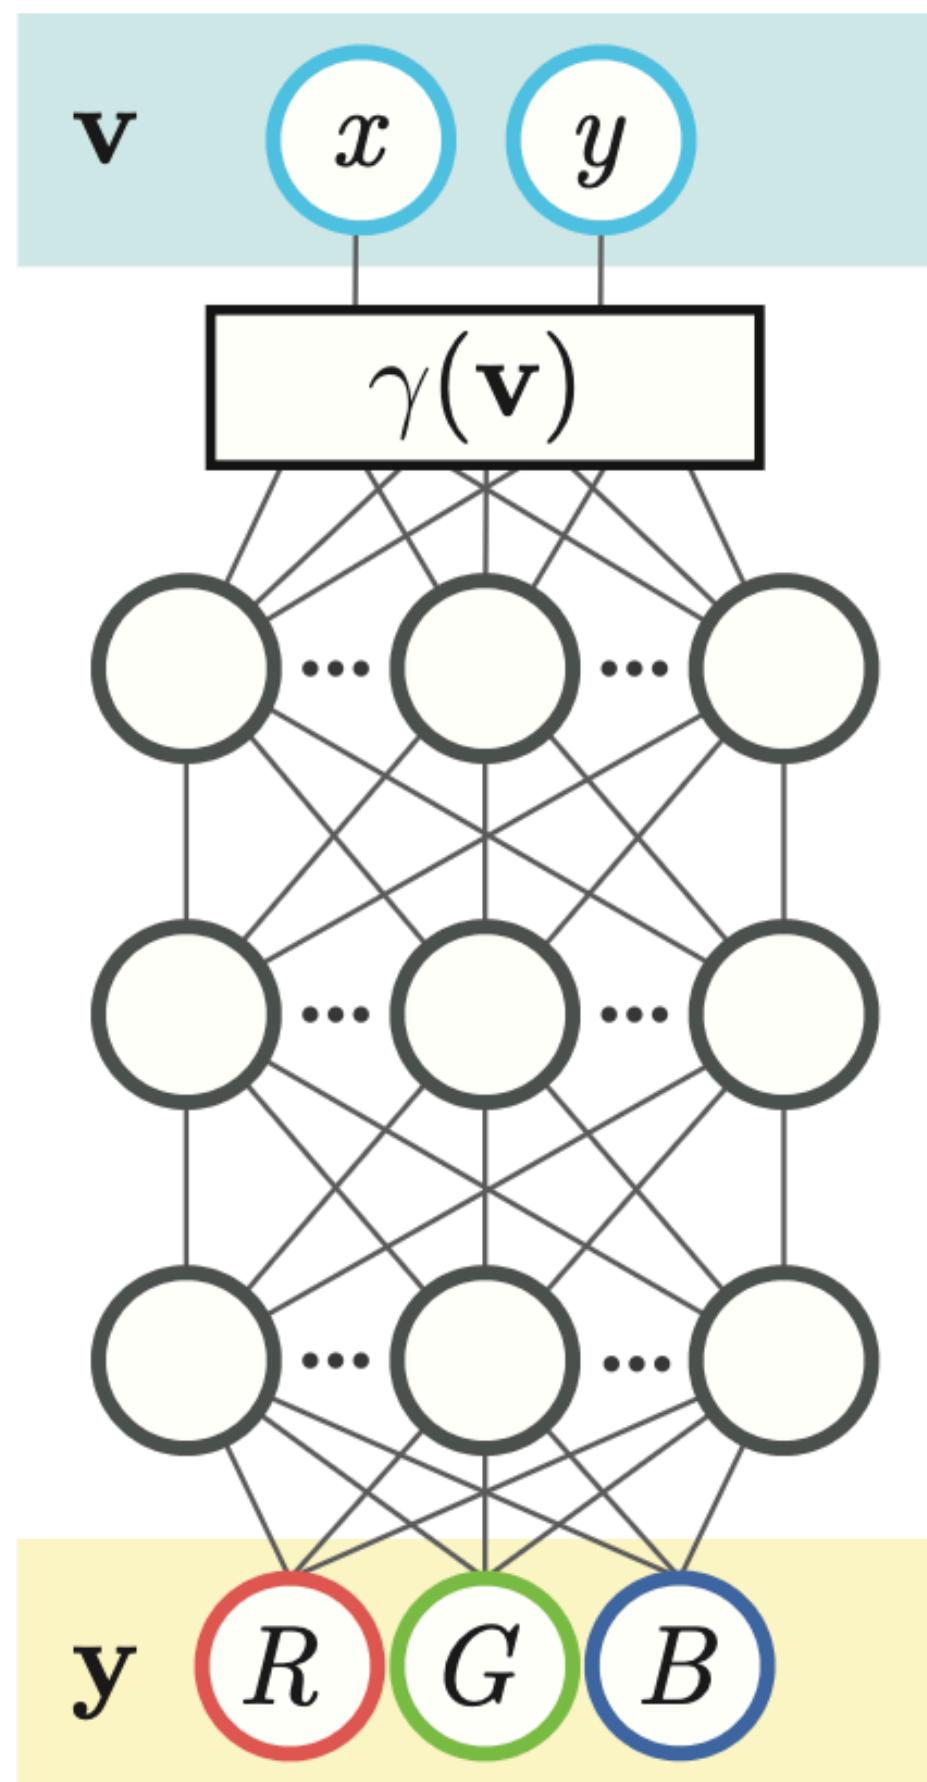
$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

Sinusoidal Embeddings

Zhong et al. ICLR 2020

Mildenhall et al., ECCV 2020

Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

Sinusoidal Embeddings

Zhong et al. ICLR 2020

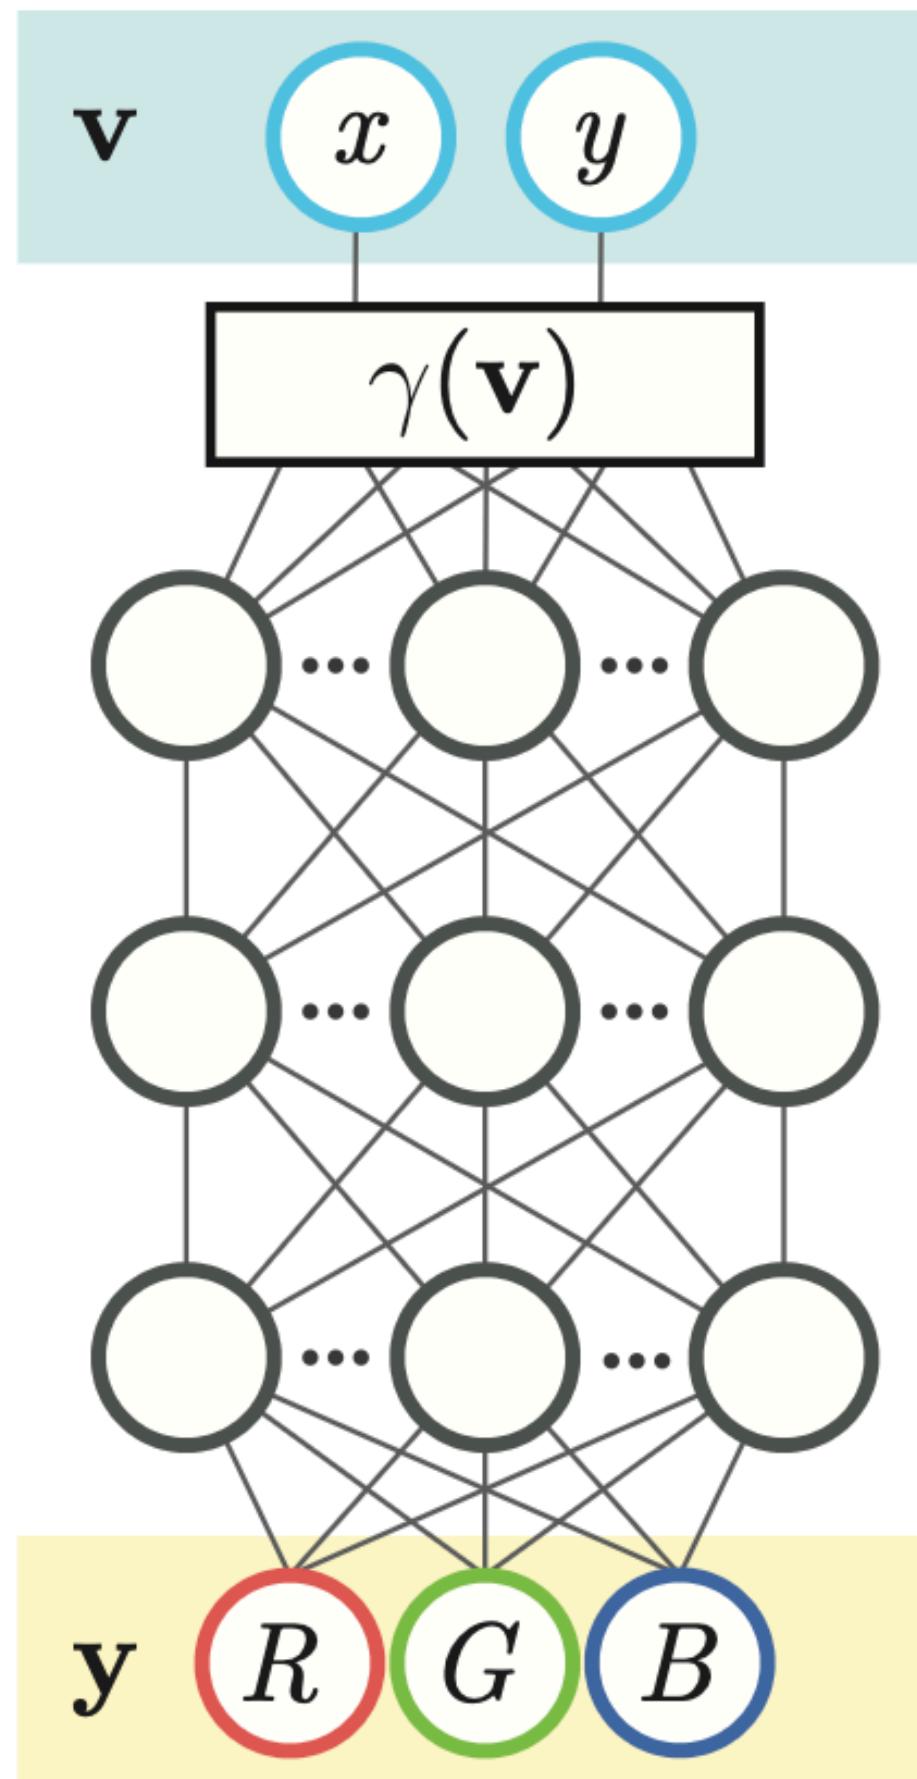
Mildenhall et al., ECCV 2020

$$\gamma(\mathbf{x}) = \exp\left(-\frac{\|t - x\|^2}{2\sigma^2}\right)$$

Gaussian Embeddings

Zheng et al., arXiv 2021

Feature Embeddings



$$\gamma(\mathbf{x}) = [\gamma_1(\mathbf{x}), \dots, \gamma_n(\mathbf{x})]$$

$$\gamma_{2i}(\mathbf{x}) = \sin(2^{i-1}\pi x_i)$$

$$\gamma_{(2i+1)}(\mathbf{x}) = \cos(2^{i-1}\pi x_i)$$

Sinusoidal Embeddings

Zhong et al. ICLR 2020

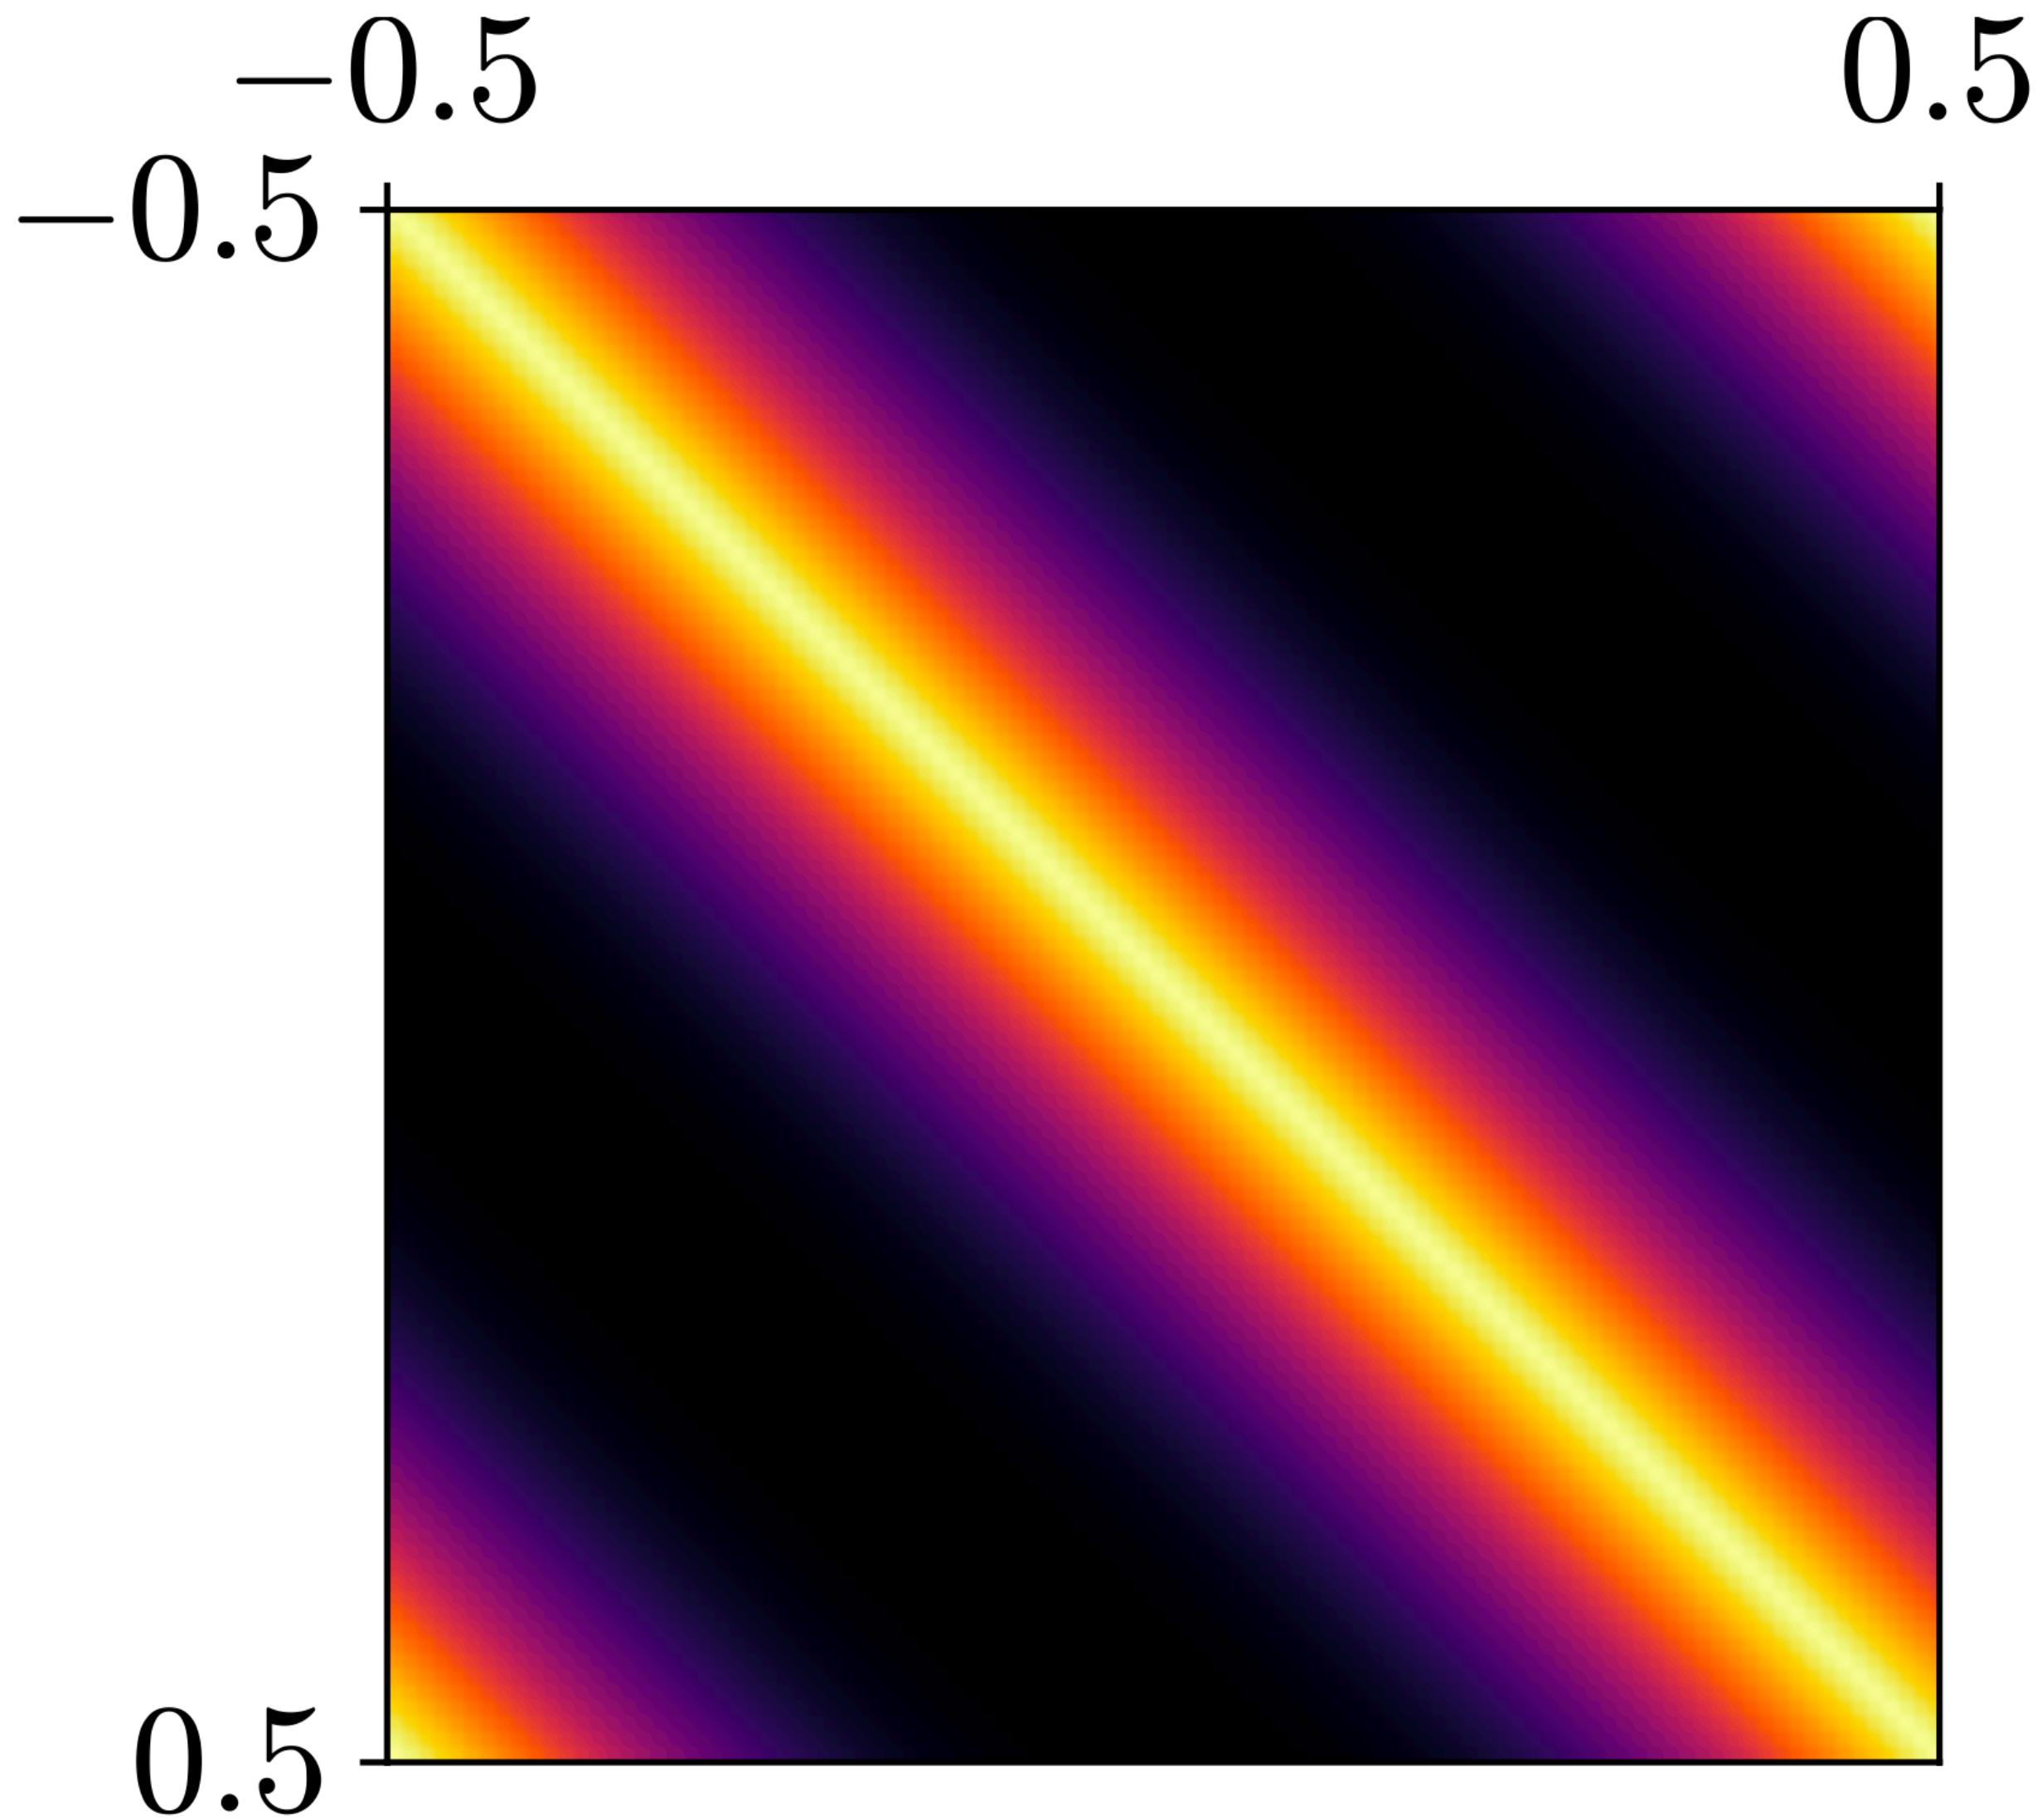
Mildenhall et al., ECCV 2020

$$\gamma(\mathbf{x}) = \exp\left(-\frac{\|t - x\|^2}{2\sigma^2}\right)$$

Gaussian Embeddings

Zheng et al., arXiv 2021

• • •



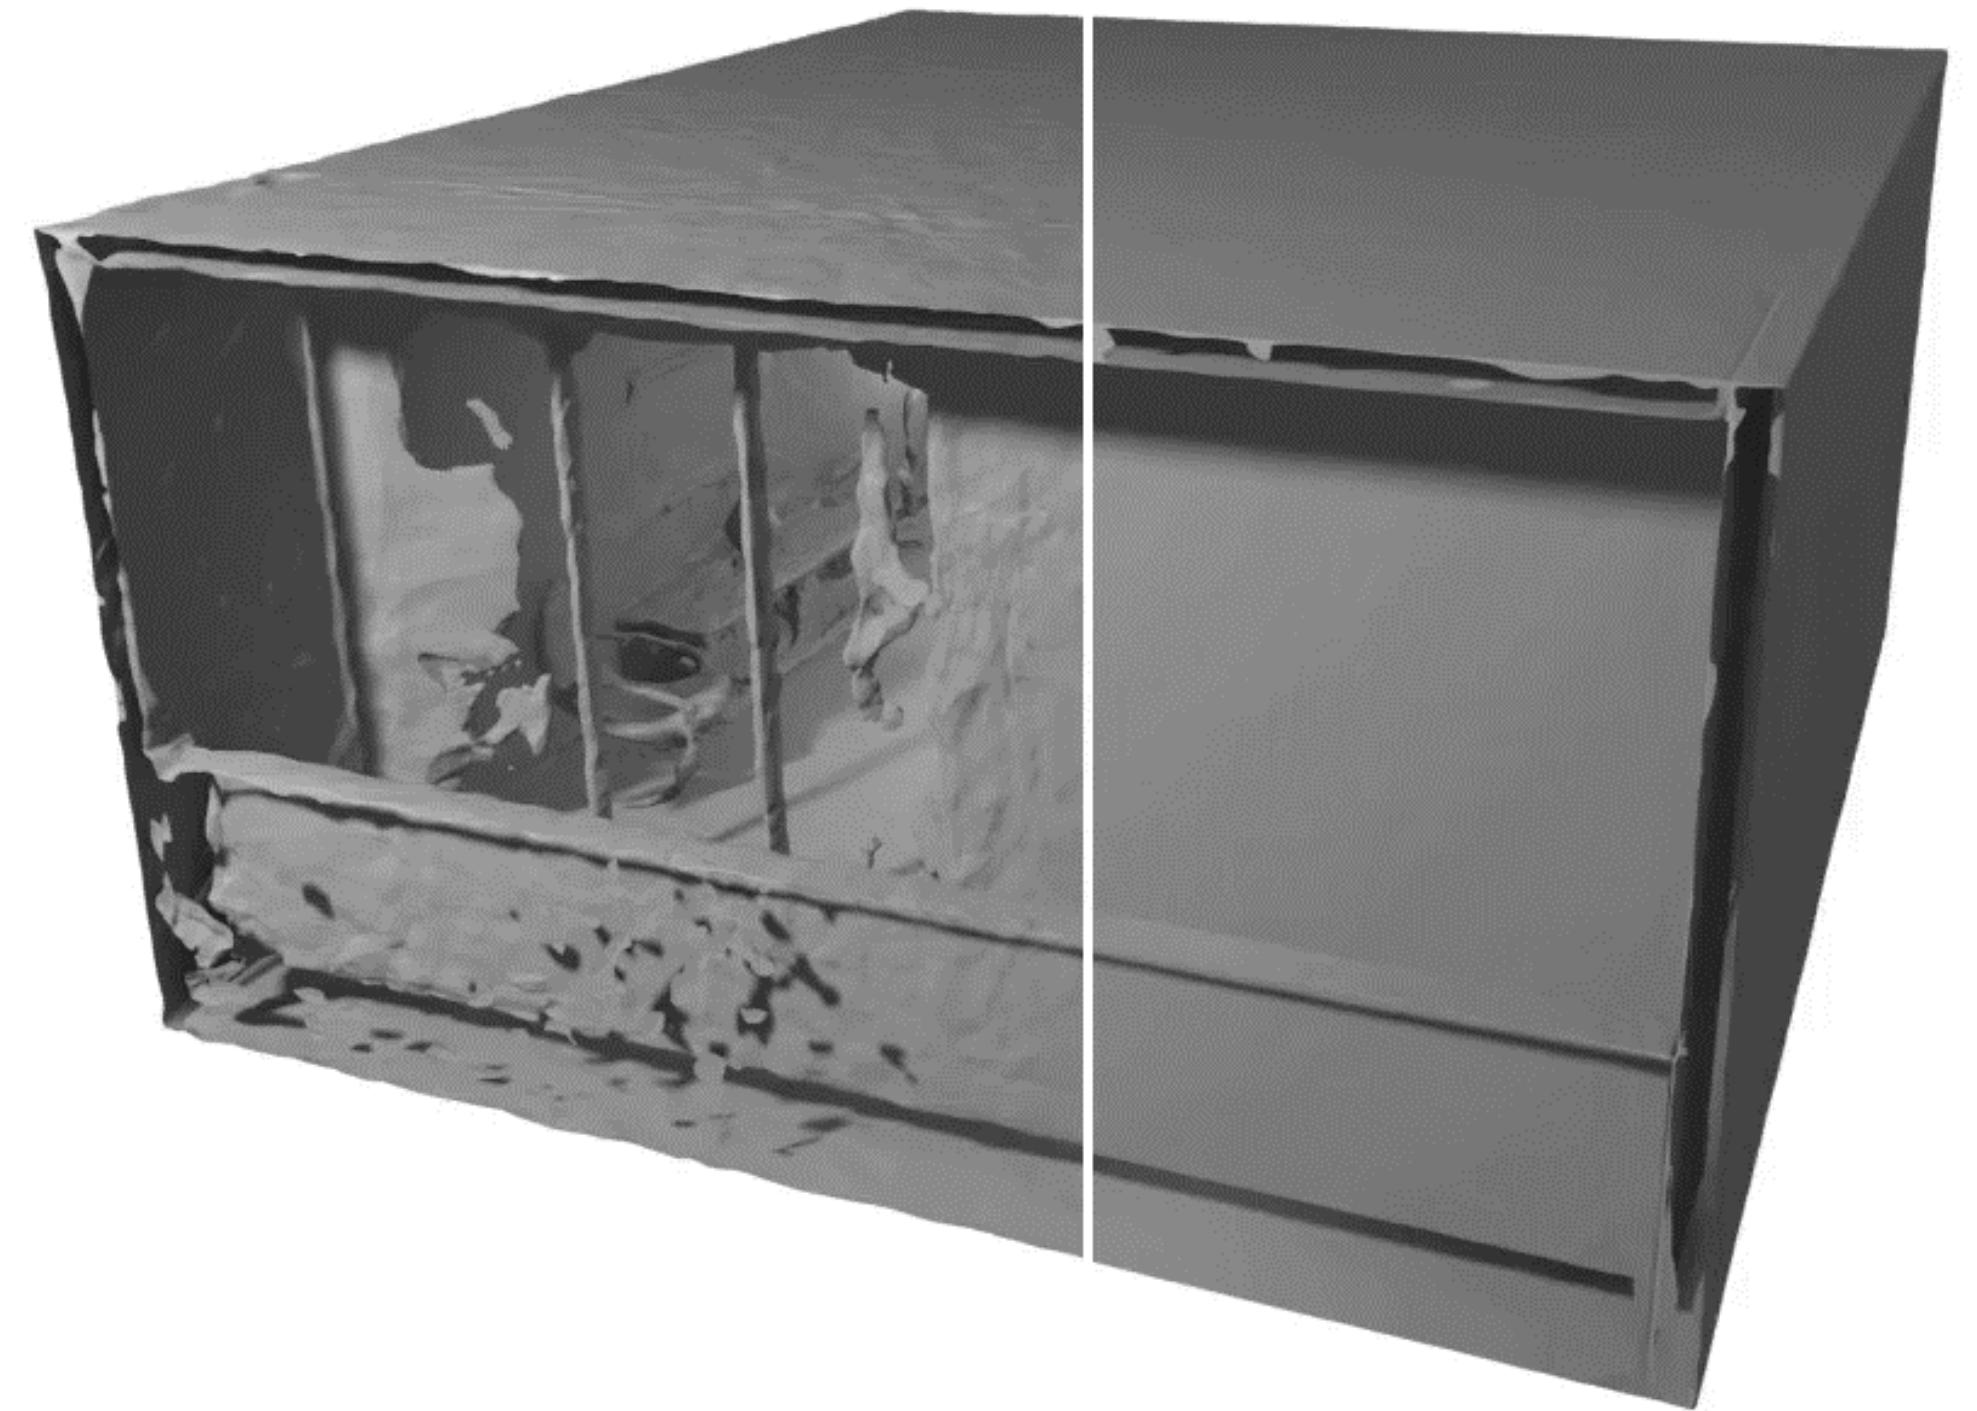
NTK kernel $k_{NTK}(\gamma(x_i), \gamma(x_j))$ for a 4-layer ReLU MLP with one scalar input and **Sinusoidal Feature Mapping**

$$\gamma(x) = [\cos(2\pi x), \sin(2\pi x)].$$

From “Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains”, Tancik, Srinivasan, Mildenhall et al.

NTK is local (note diagonal).

SIREN



ReLU

SIREN⁹

SIREN



ReLU

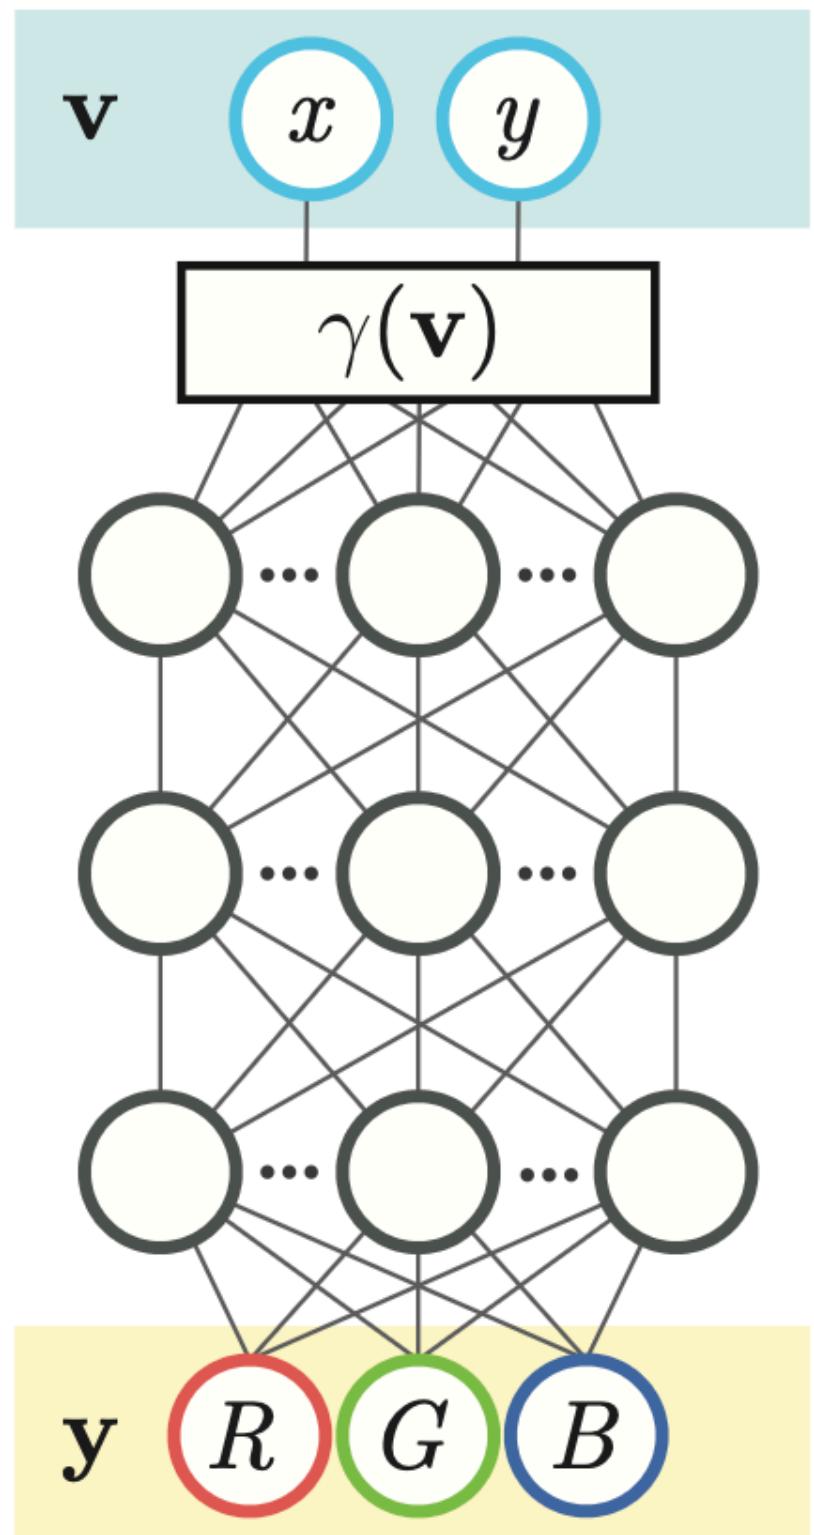
SIREN⁹

SIREN



~1MB compared to
110MB of full mesh!

Fourier Features



(a) Coordinate-based MLP

No Fourier features
 $\gamma(\mathbf{v}) = \mathbf{v}$

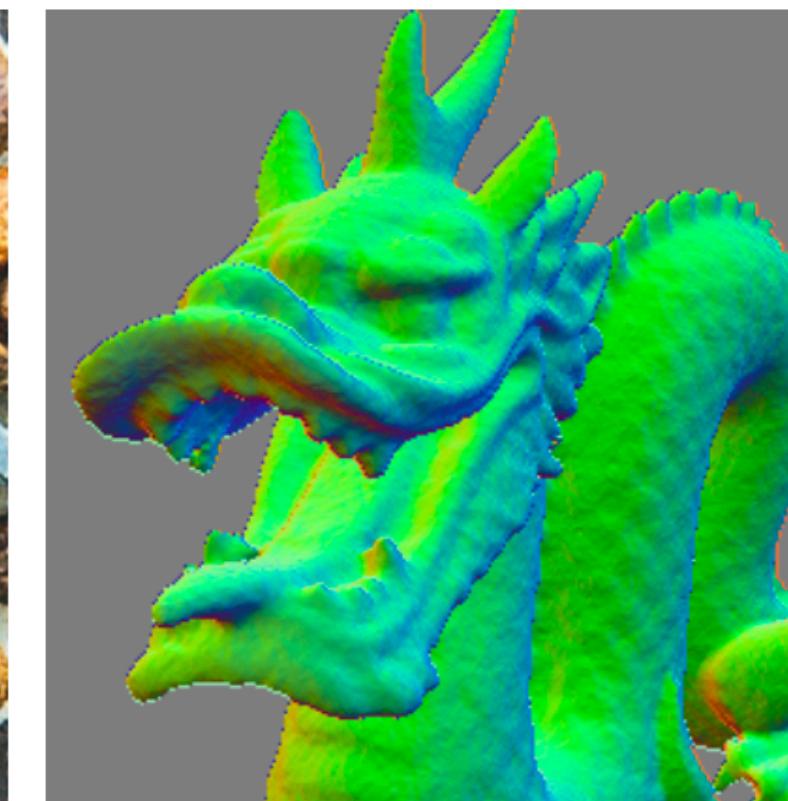
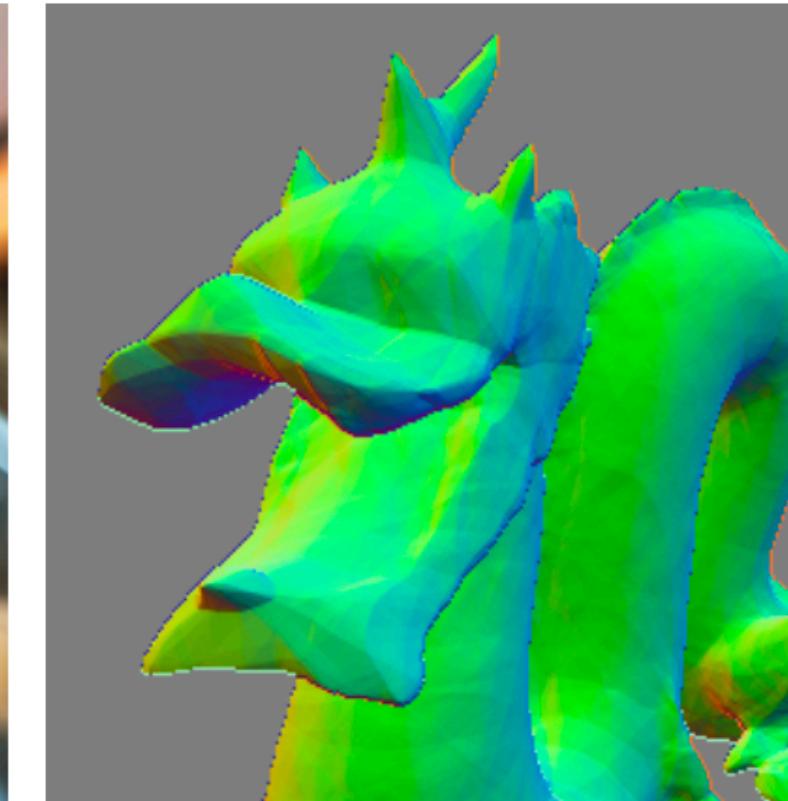


With Fourier features
 $\gamma(\mathbf{v}) = \text{FF}(\mathbf{v})$



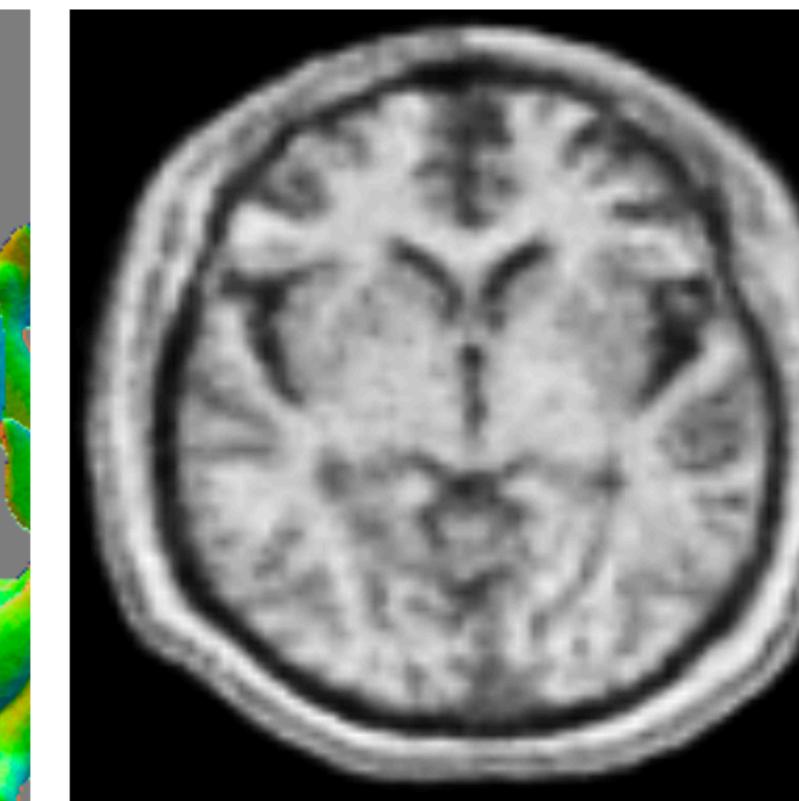
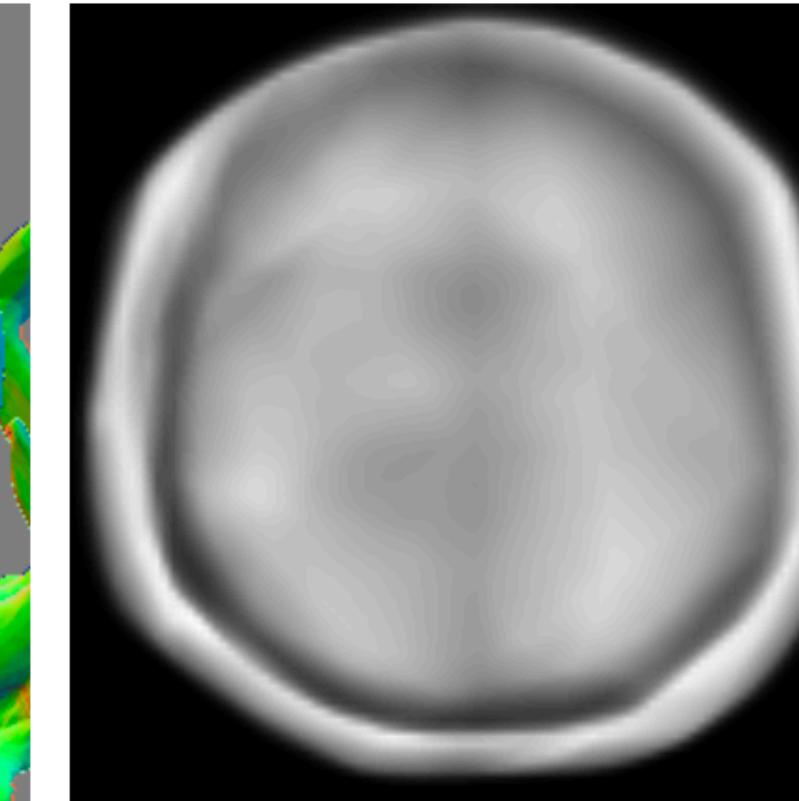
(b) Image regression

$$(x, y) \rightarrow \text{RGB}$$



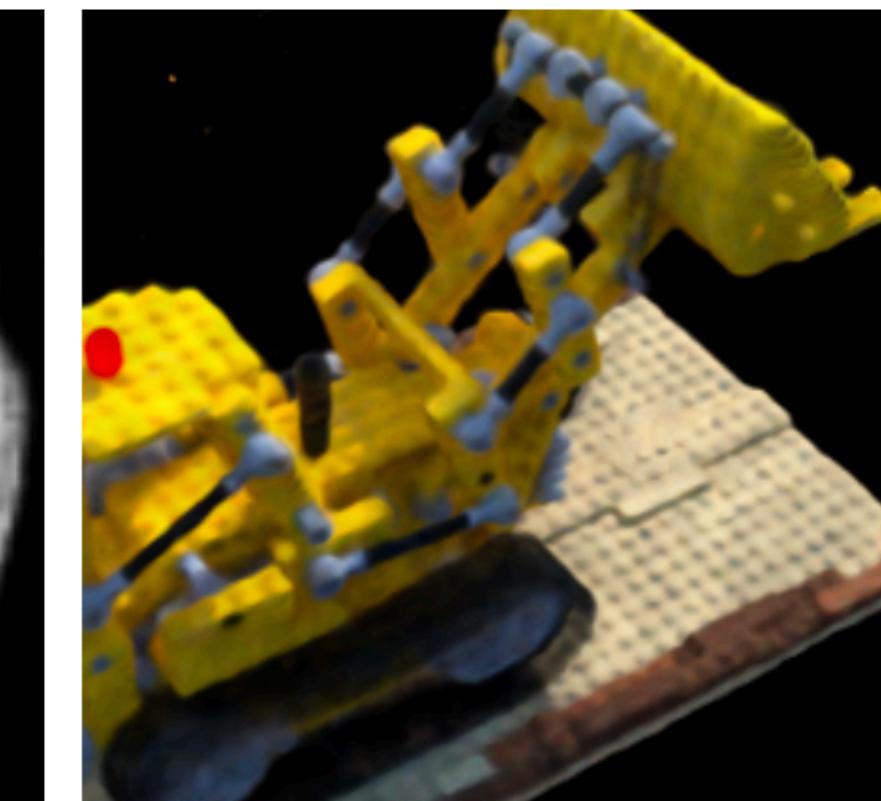
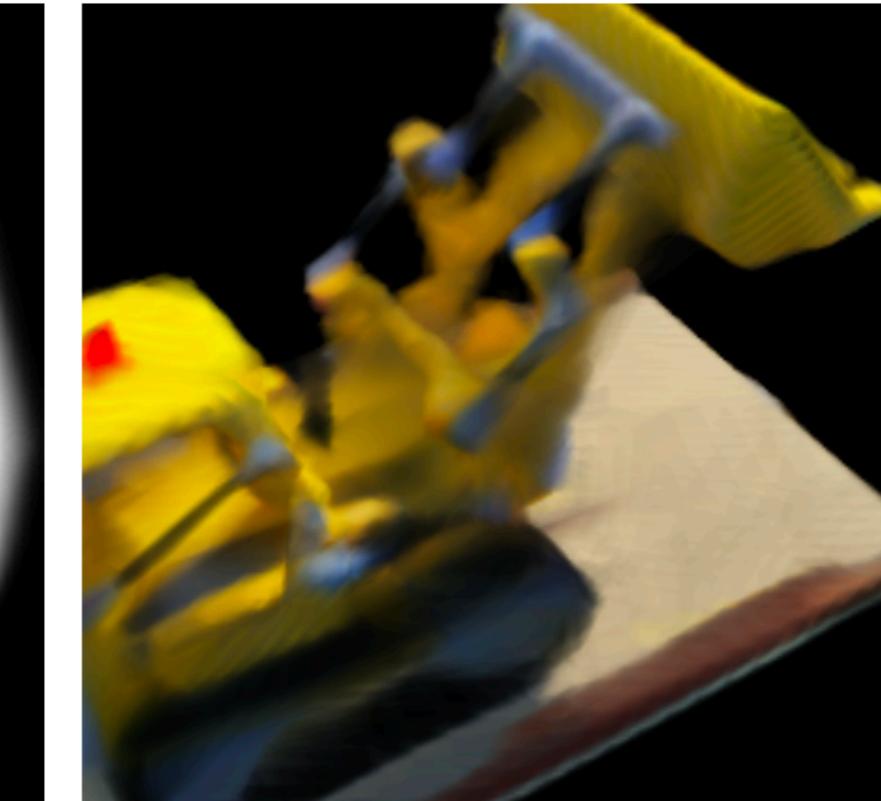
(c) 3D shape regression

$$(x, y, z) \rightarrow \text{occupancy}$$



(d) MRI reconstruction

$$(x, y, z) \rightarrow \text{density}$$



(e) Inverse rendering

$$(x, y, z) \rightarrow \text{RGB, density}$$

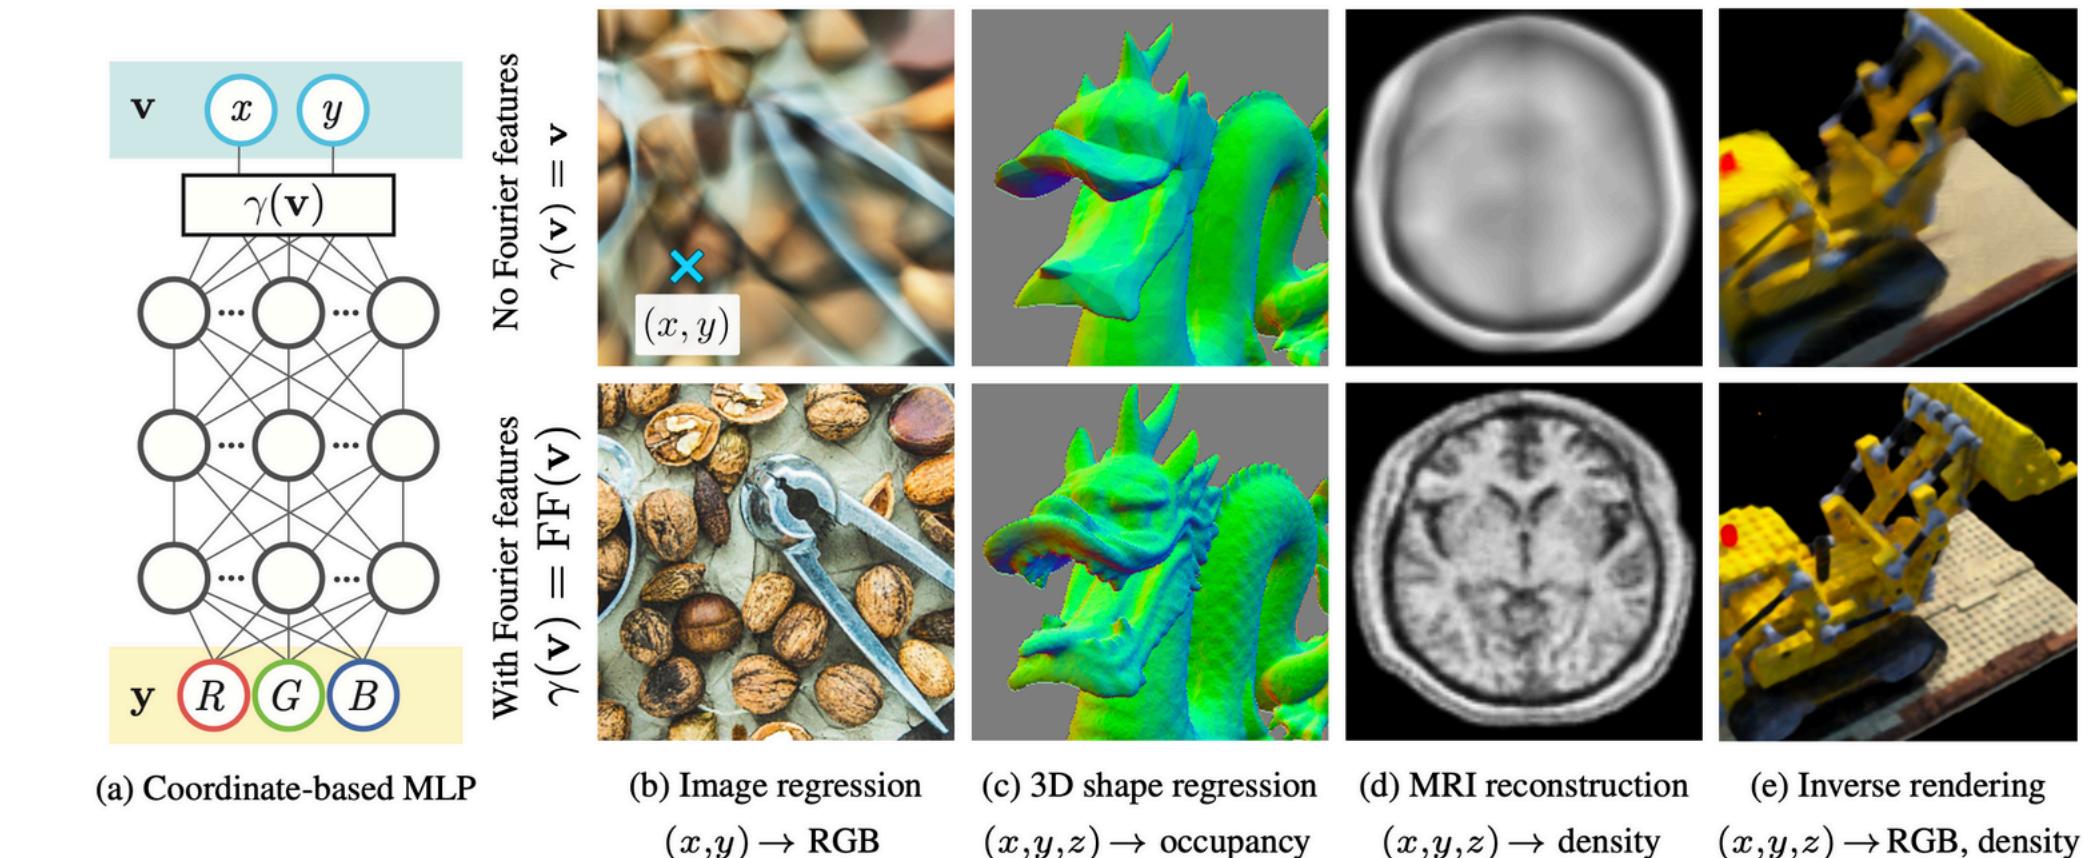
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



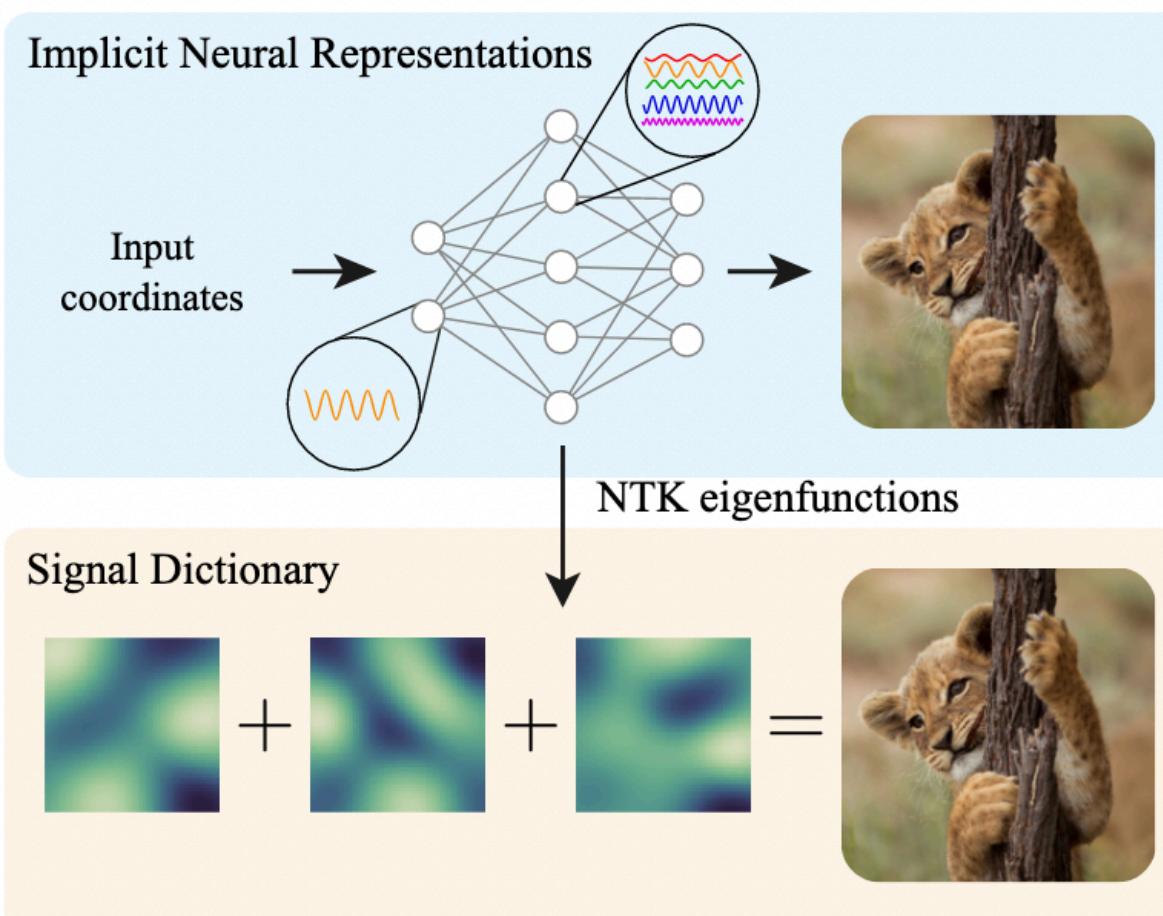
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



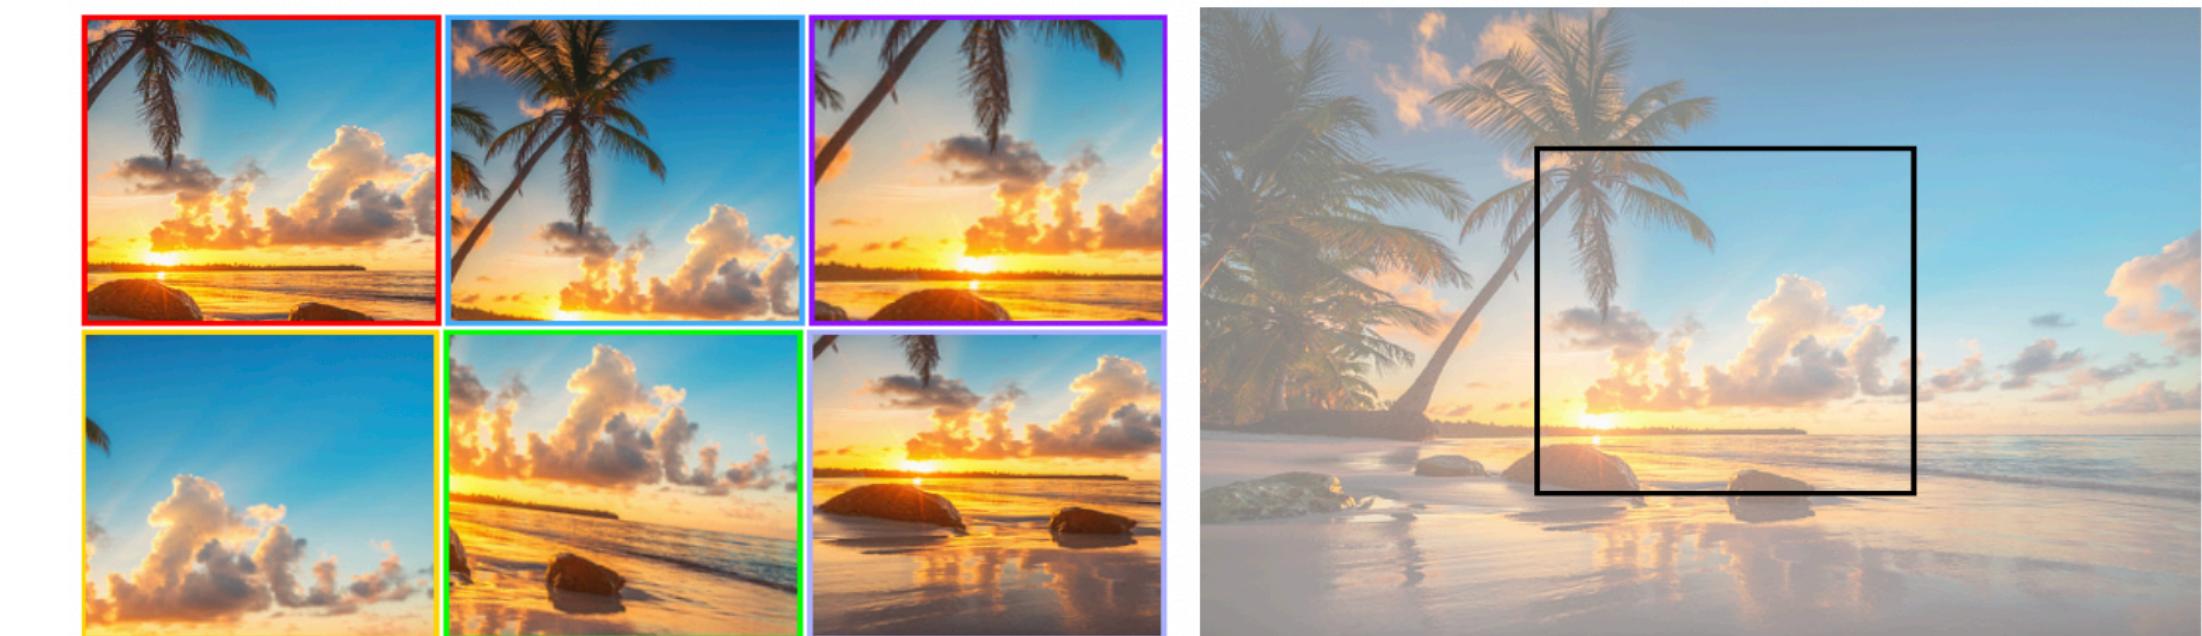
A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

Chng et al. ECCV 2022



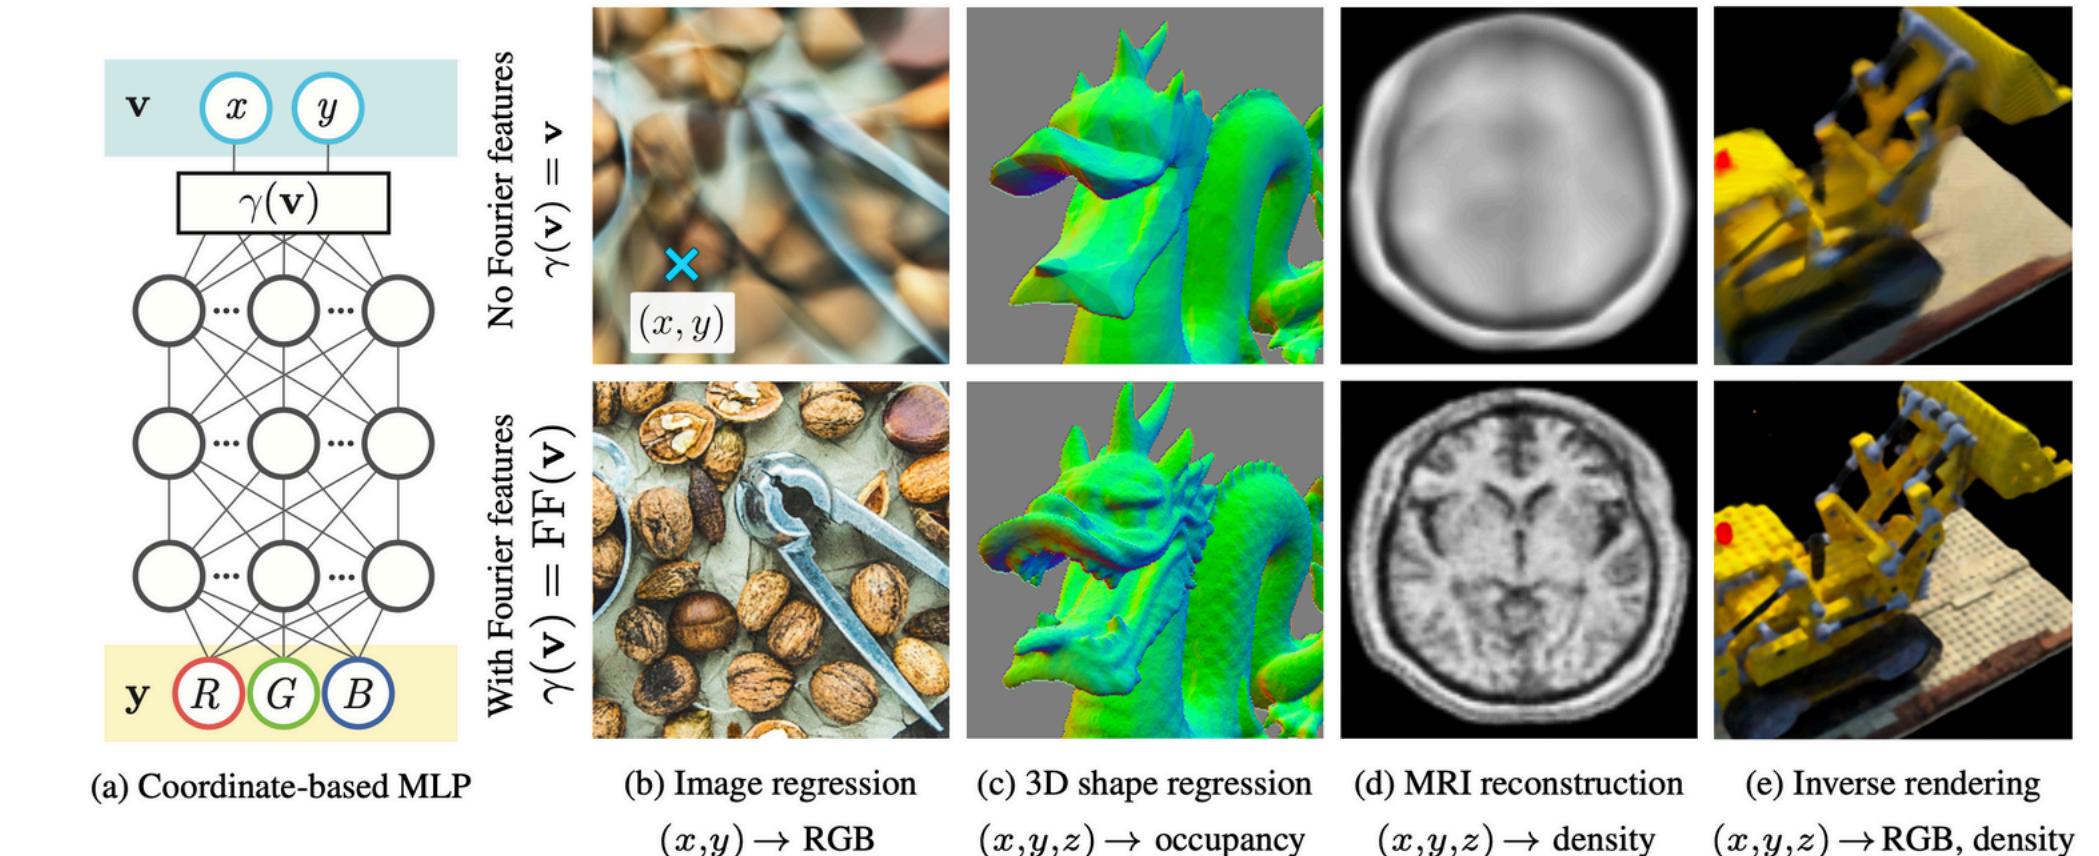
SIREN: Neural Implicit Representations With Periodic Activation Functions

Sitzmann & Martel et al. NeurIPS 2020



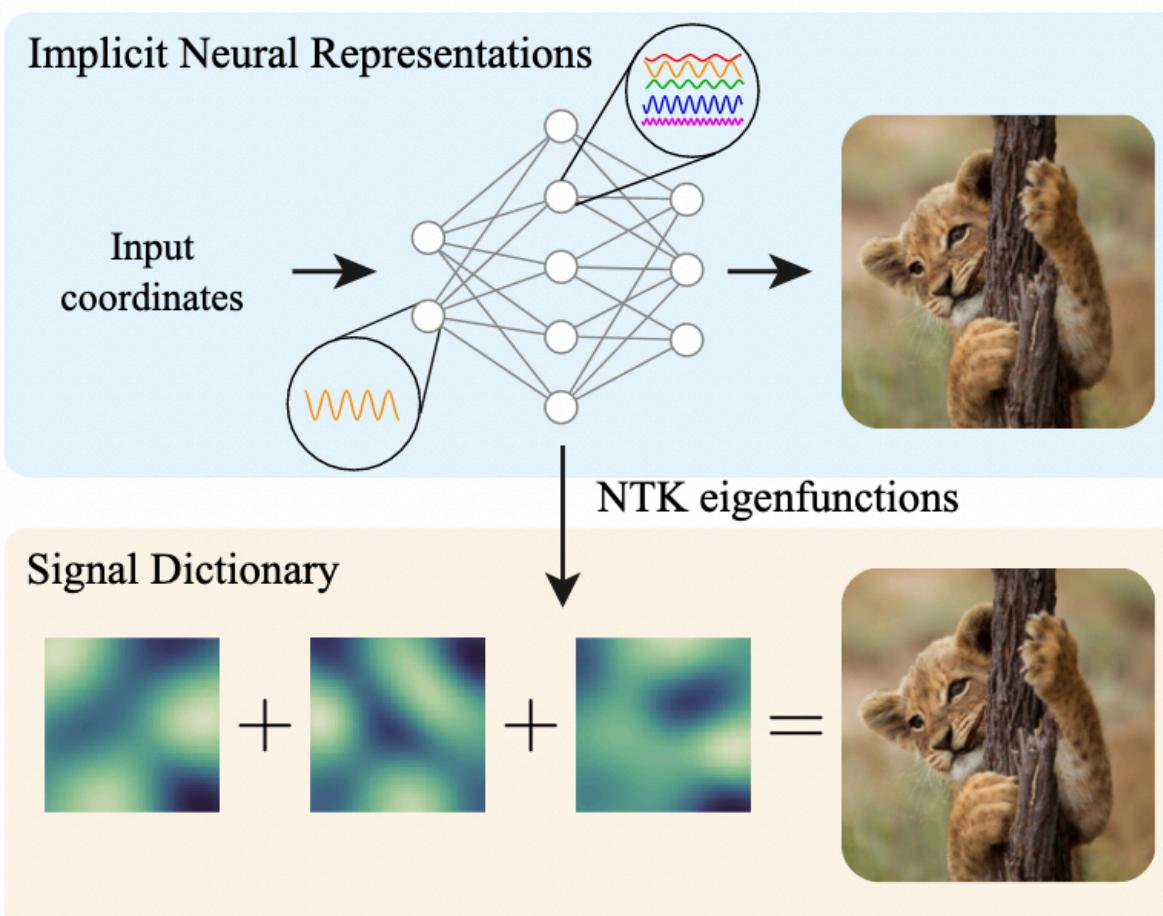
Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains

Tancik, Srinivasan, Mildenhall NeurIPS 2020



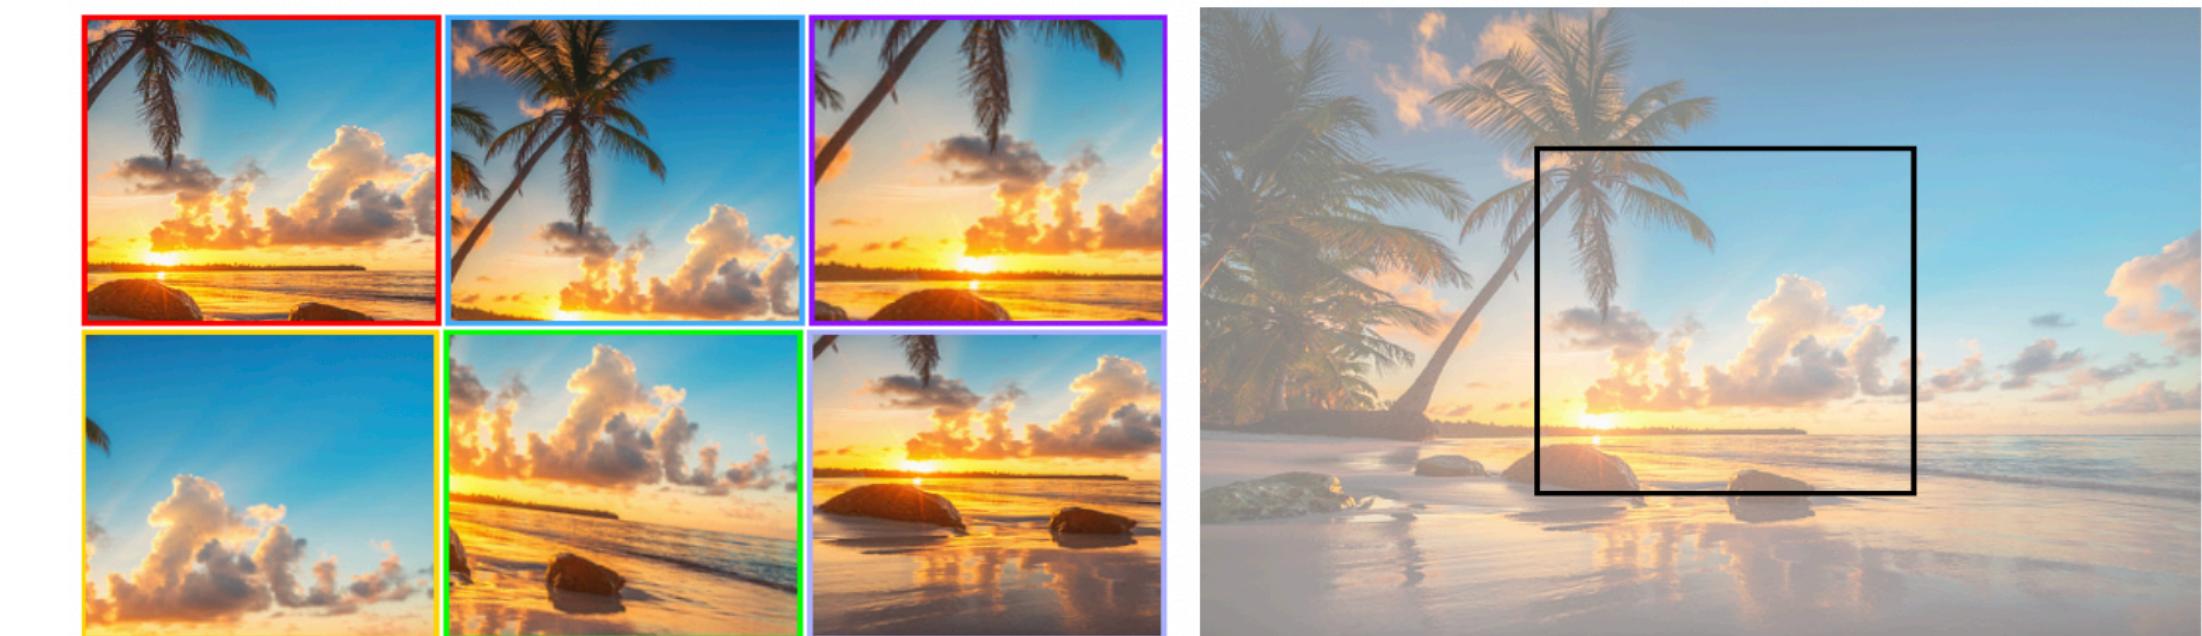
A Structured Dictionary Perspective on Implicit Neural Representations

Yüce & Ortiz-Jiménez et al. CVPR 2022



Gaussian Activated Neural Radiance Fields for High Fidelity Reconstruction & Pose Estimation

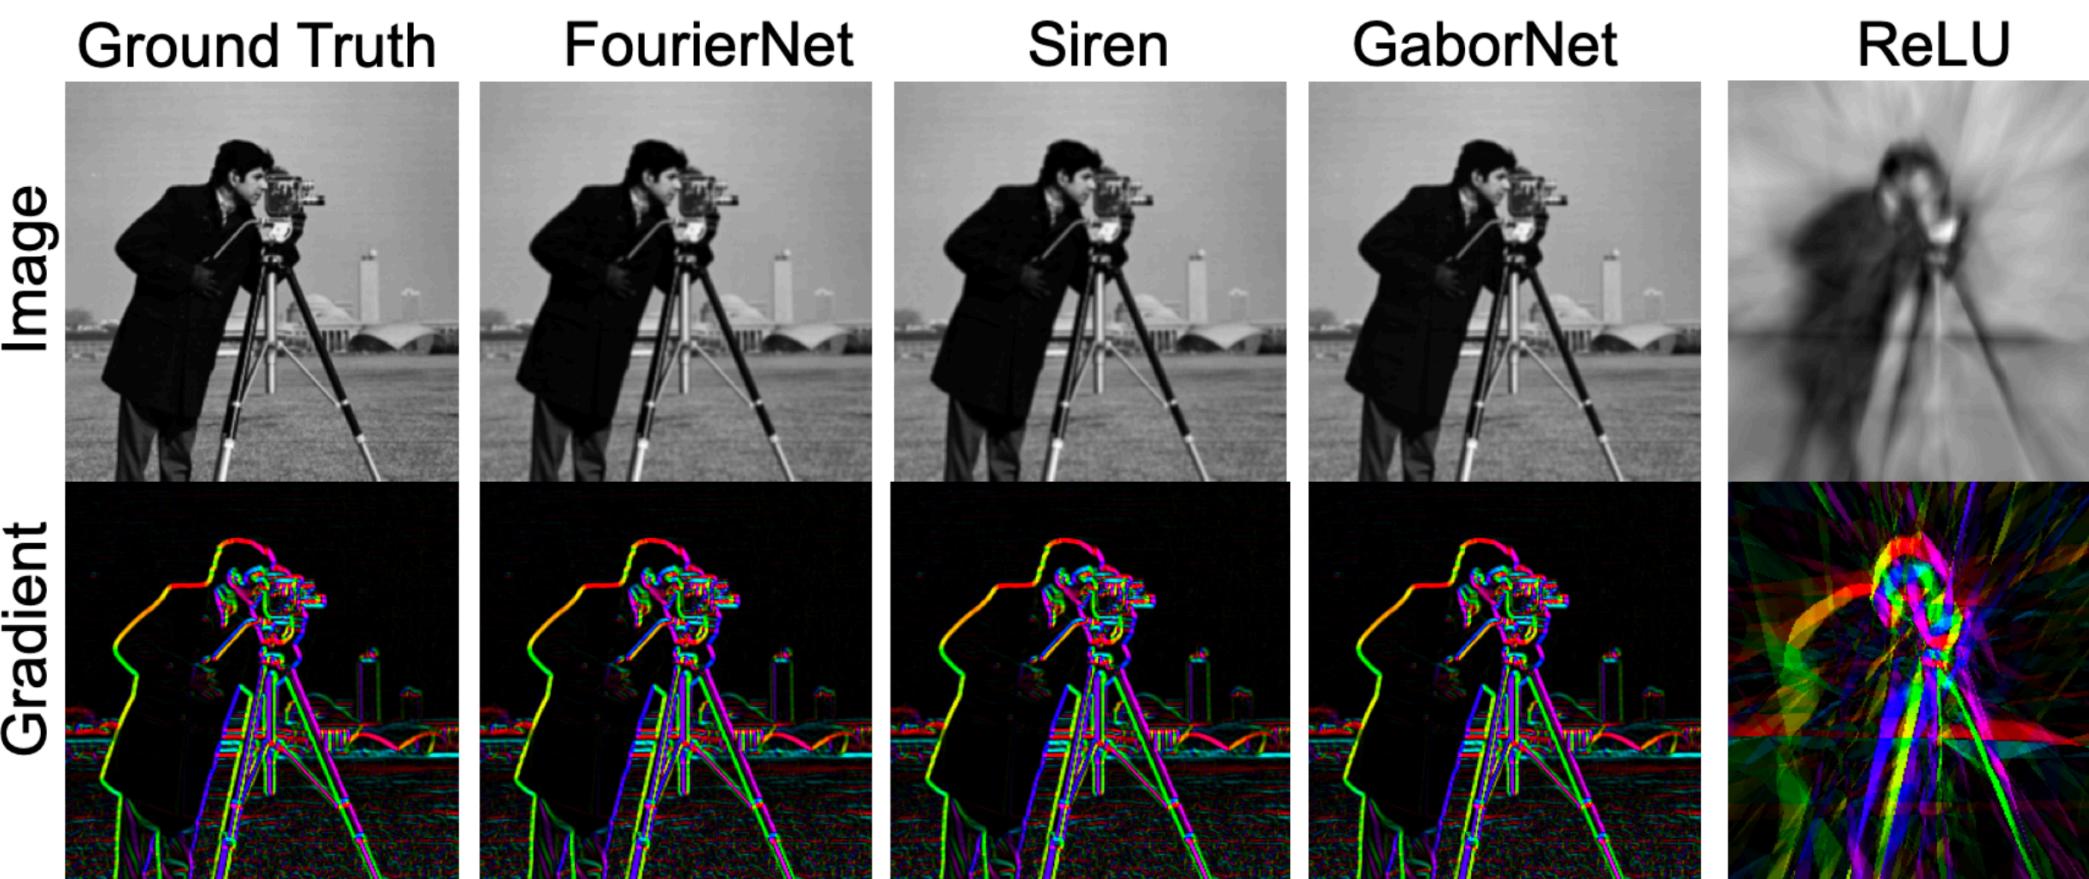
Chng et al. ECCV 2022



Neural (?) Fields with analytical (but not tractably computable) Fourier spectrum!

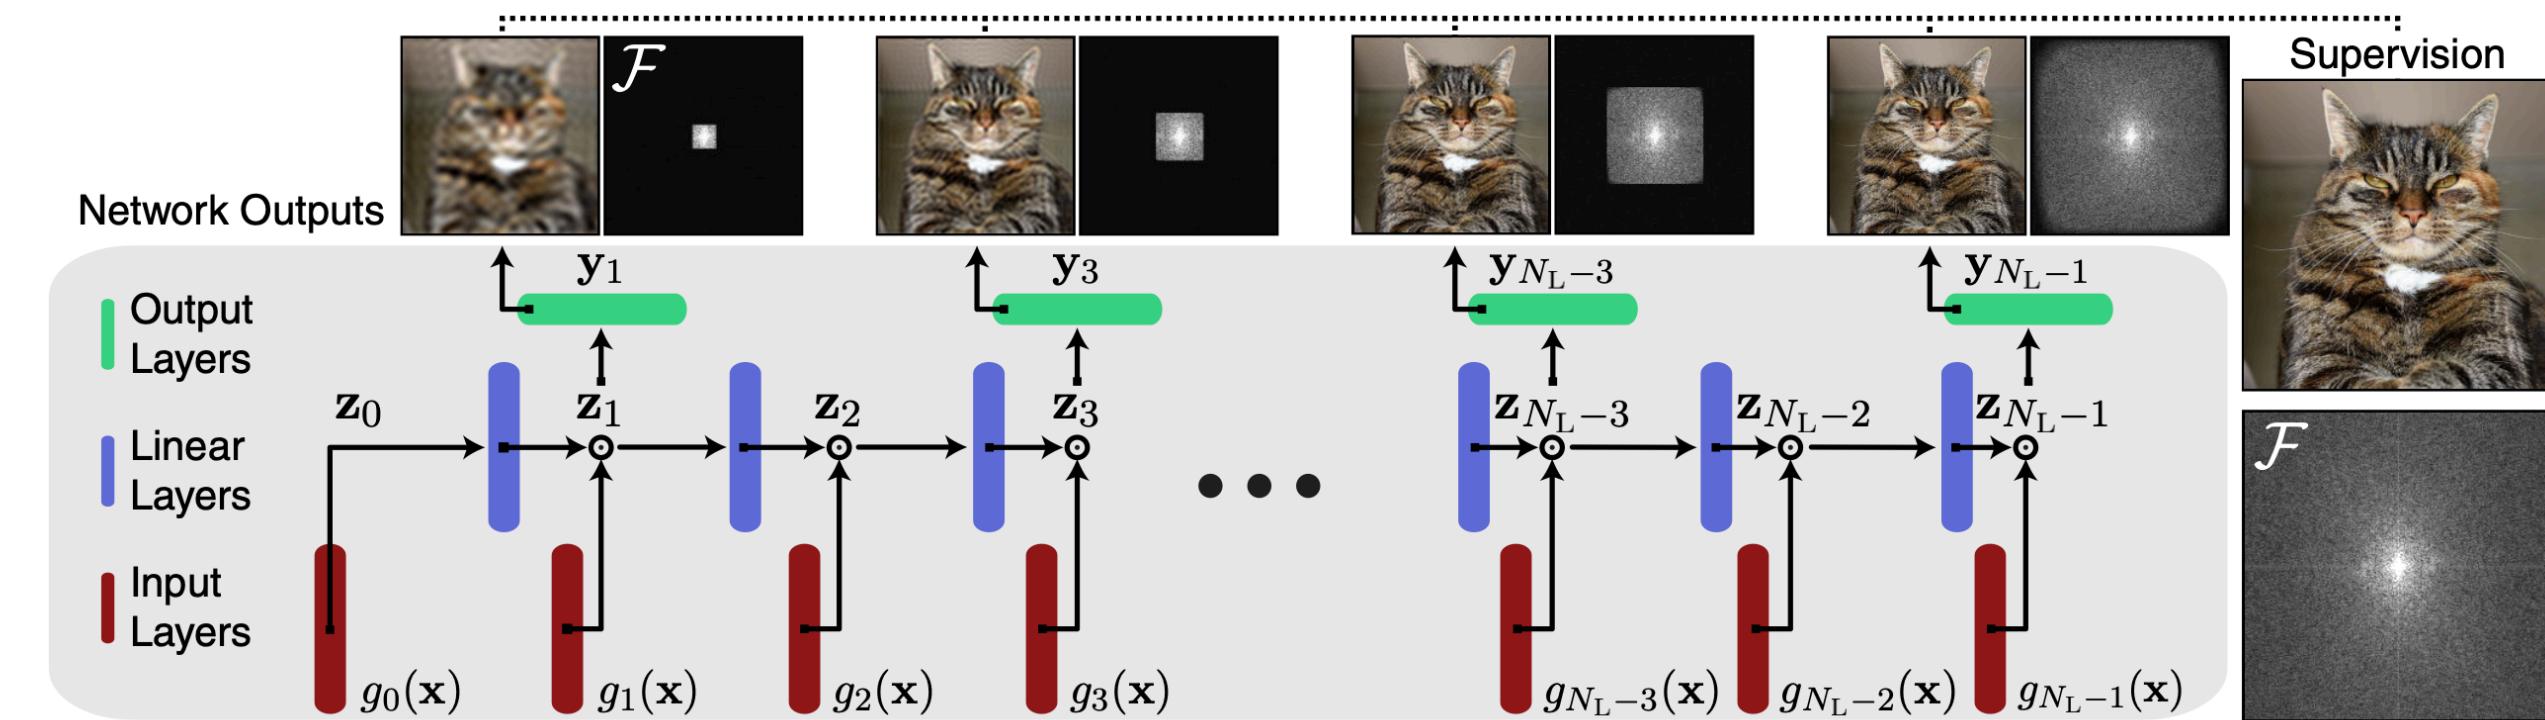
Multiplicative Filter Networks

Fathony et al. ICLR 2021



BACON: Band-limited coordinate networks for multiscale scene representation

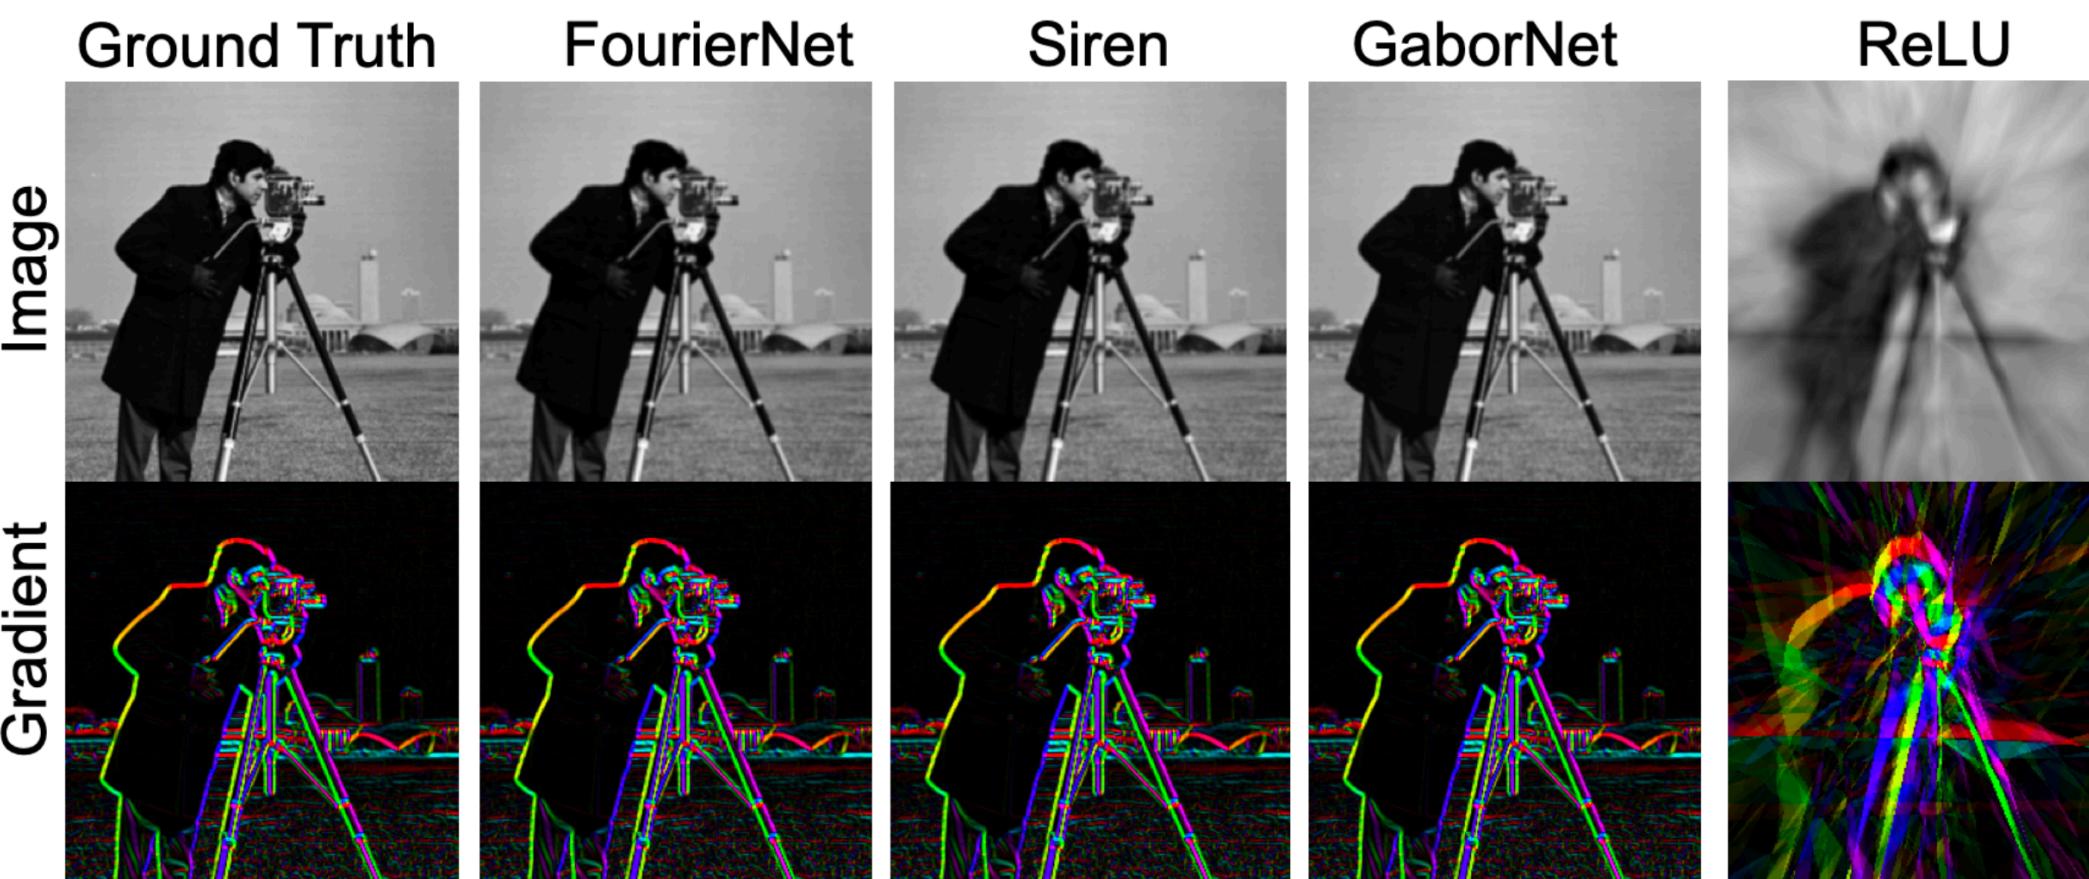
Lindell et al. CVPR 2022



Neural (?) Fields with analytical (but not tractably computable) Fourier spectrum!

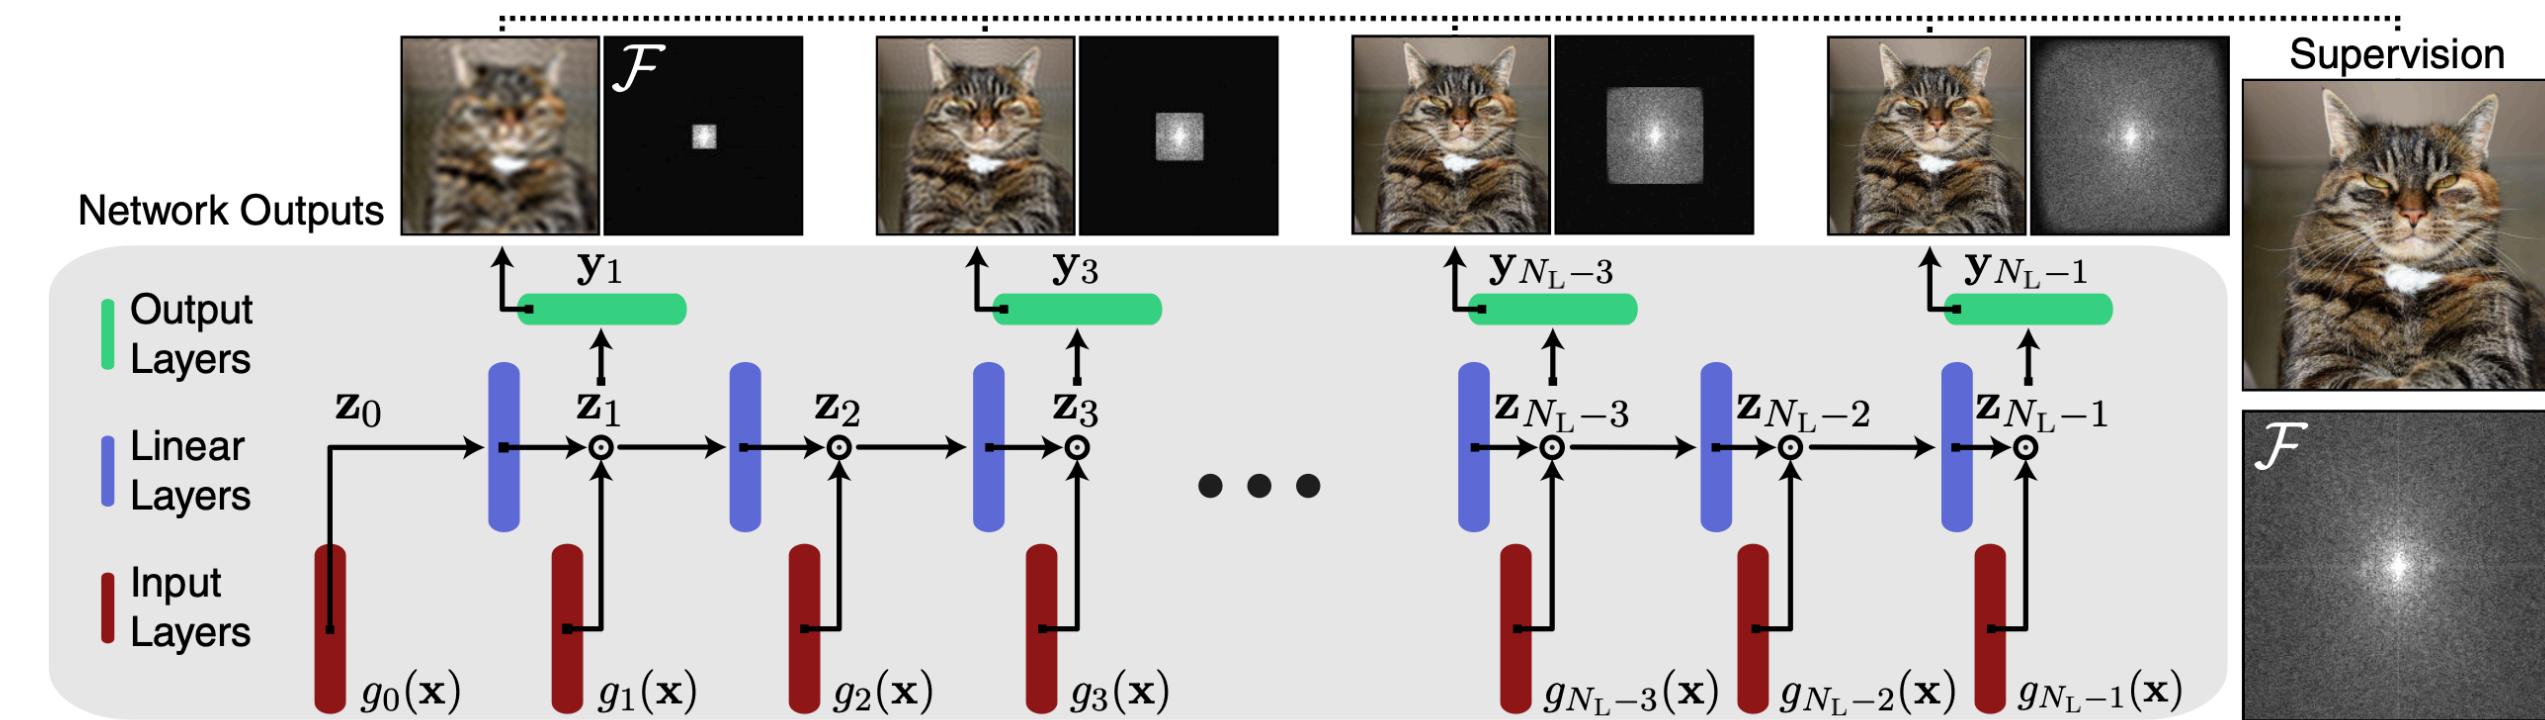
Multiplicative Filter Networks

Fathony et al. ICLR 2021



BACON: Band-limited coordinate networks for multiscale scene representation

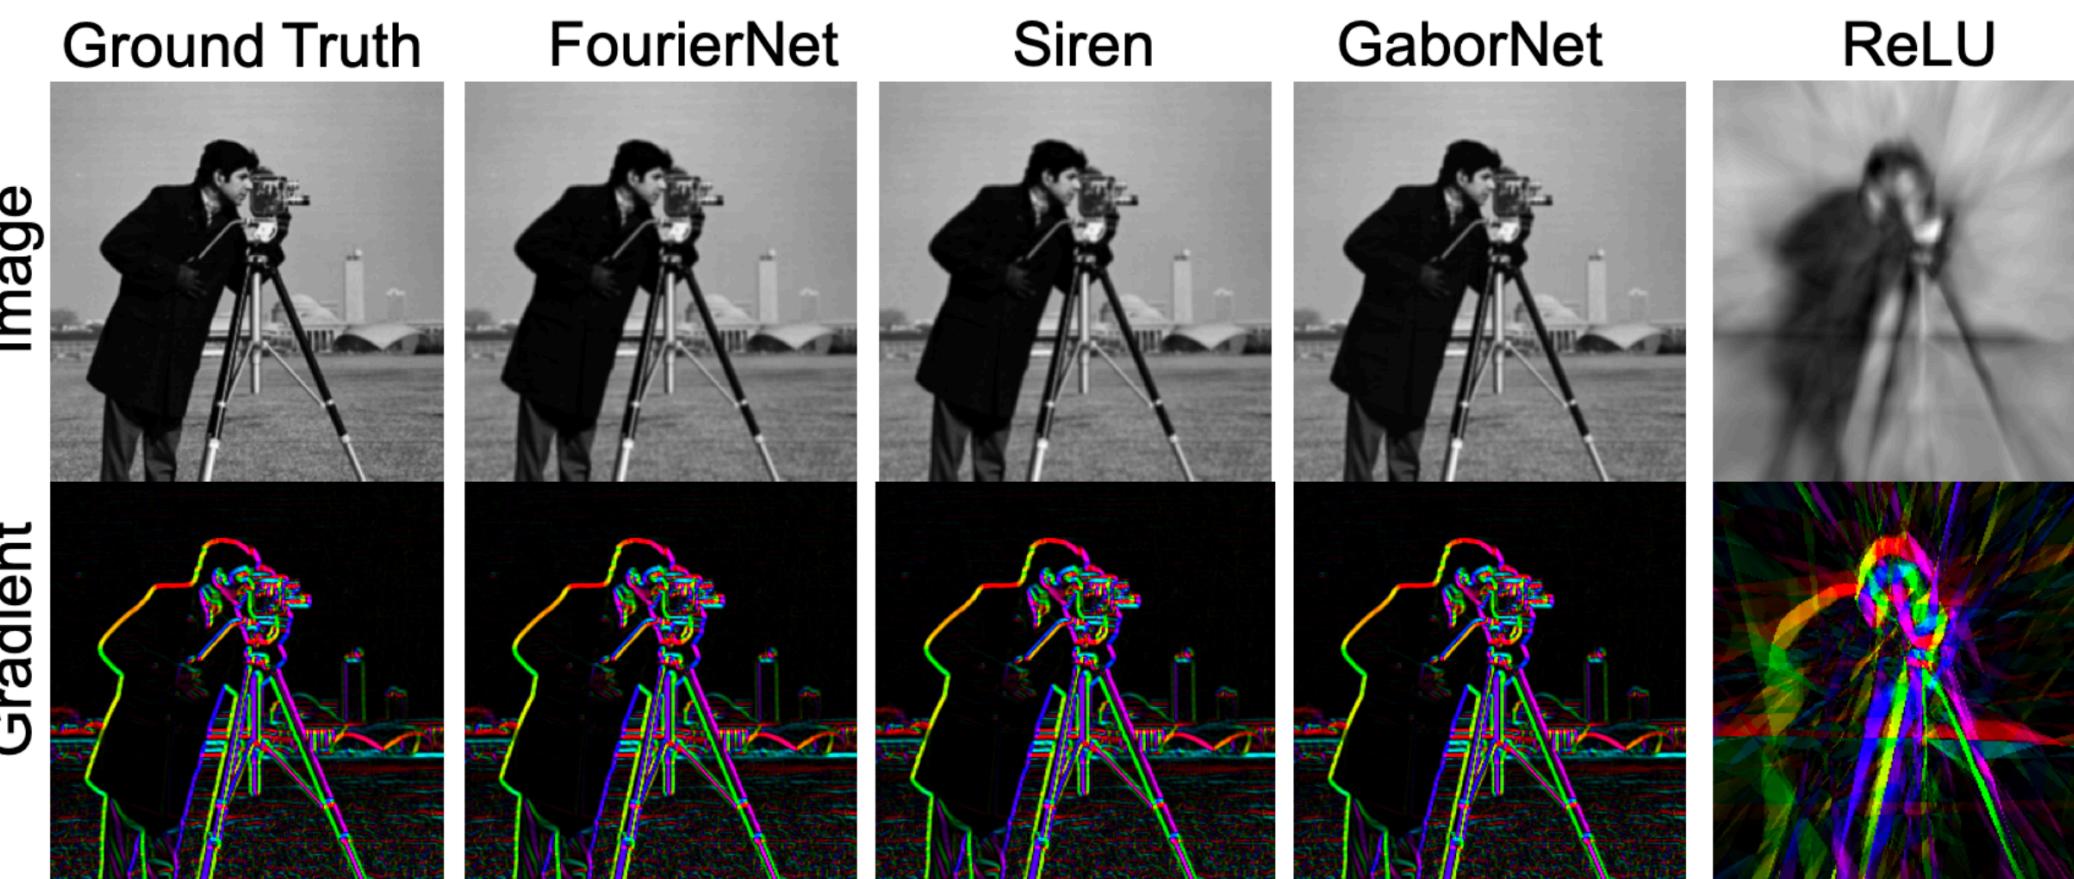
Lindell et al. CVPR 2022



Neural (?) Fields with analytical (but not tractably computable) spectrum!

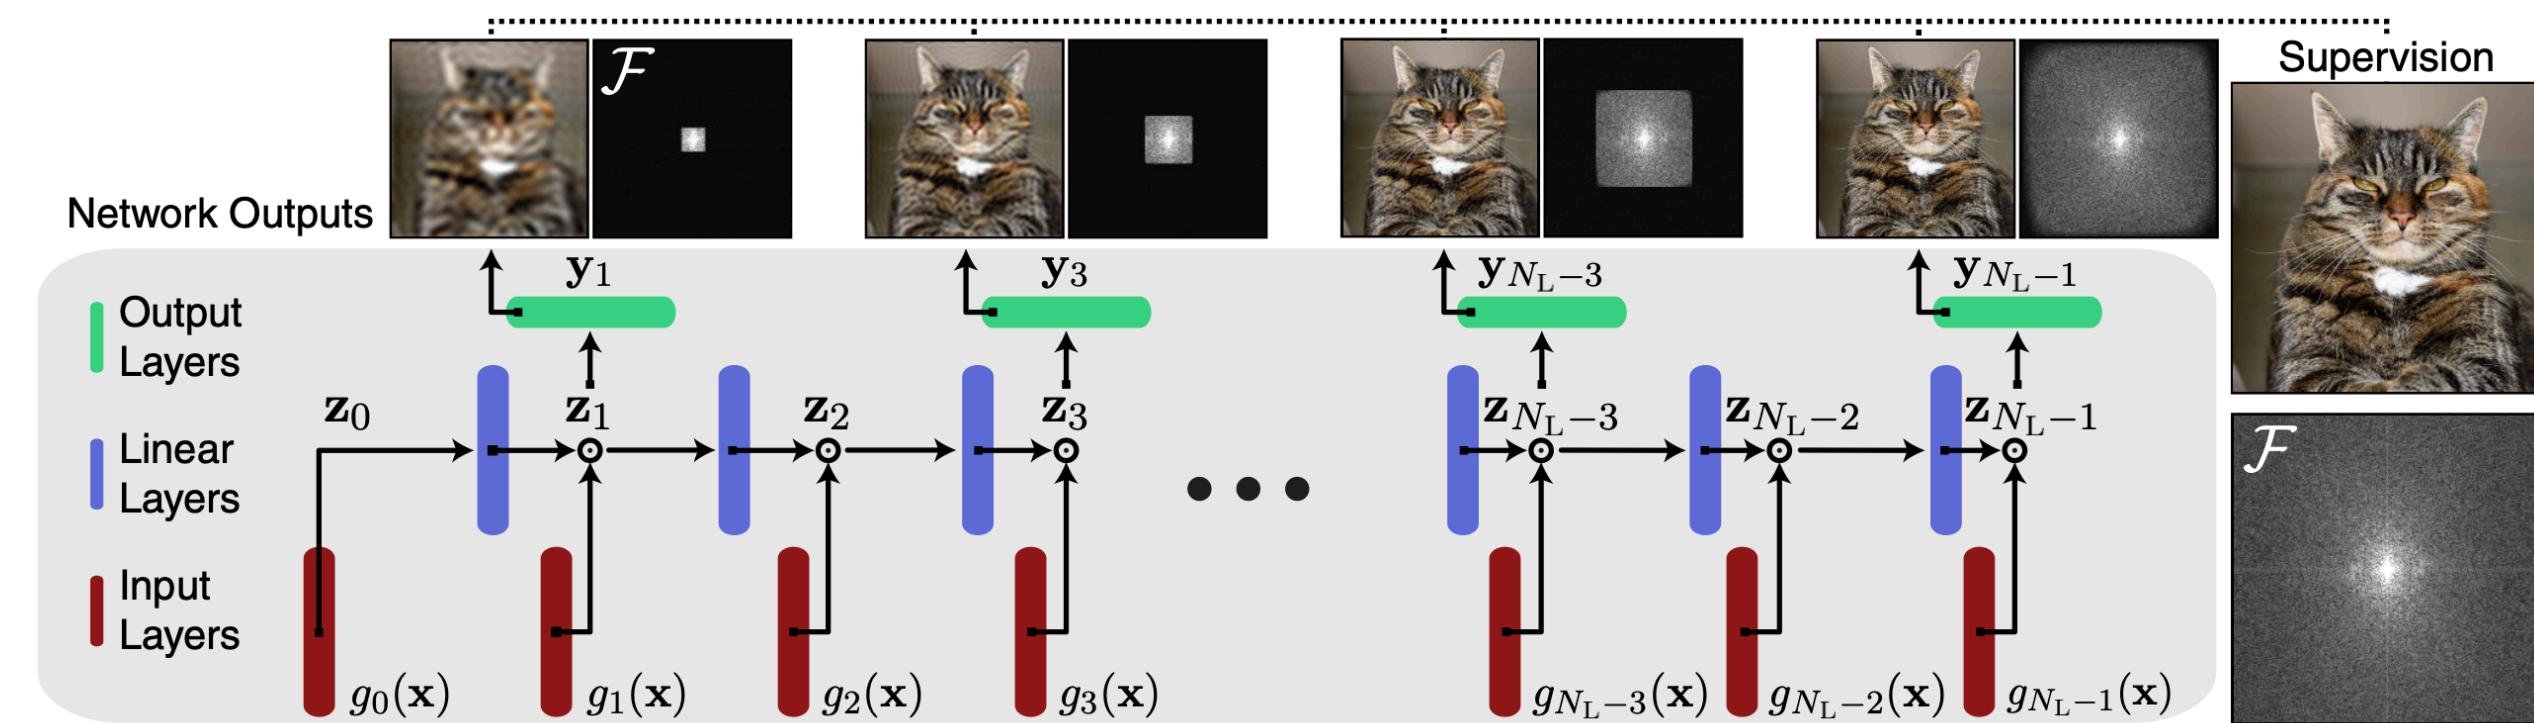
Multiplicative Filter Networks

Fathony et al. ICLR 2021



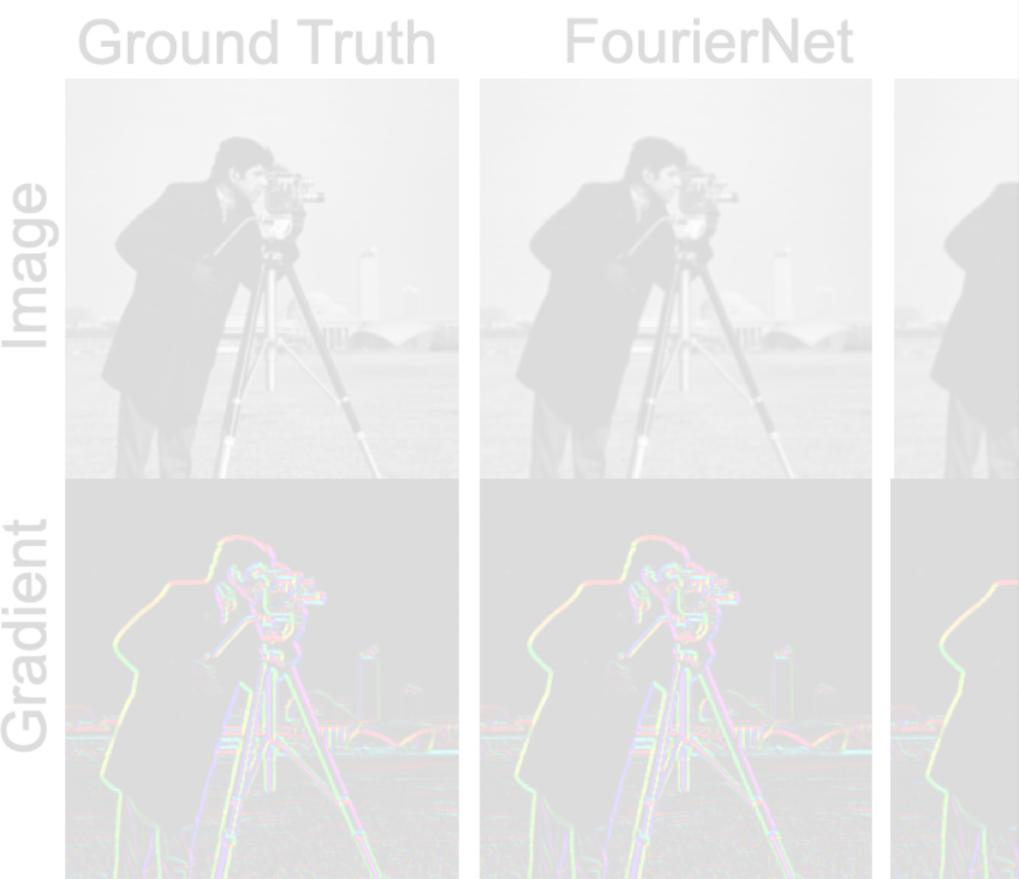
BACON: Band-limited coordinate networks for multiscale scene representation

Lindell et al. CVPR 2022



(but

Multiplicative Fathy et al.



arXiv:2111.11426v4 [cs.CV] 5 Apr 2022

EUROGRAPHICS 2022
D. Meneveaux and G. Patanè
(Guest Editors)

Volume 41 (2022), Number 2
STAR – State of The Art Report

rum!

Neural Fields in Visual Computing and Beyond

Yiheng Xie^{1,2} Towaki Takikawa^{3,4} Shunsuke Saito⁵ Or Litany⁴ Shiqin Yan¹ Numair Khan¹ Federico Tombari^{6,7}
James Tompkin¹ Vincent Sitzmann^{8†} Srinath Sridhar^{1†}

¹Brown University ²Unity Technologies ³University of Toronto ⁴NVIDIA ⁵Meta Reality Labs Research ⁶Google ⁷Technical University of Munich
⁸Massachusetts Institute of Technology [†]Equal advising

<https://neuralfields.cs.brown.edu/>

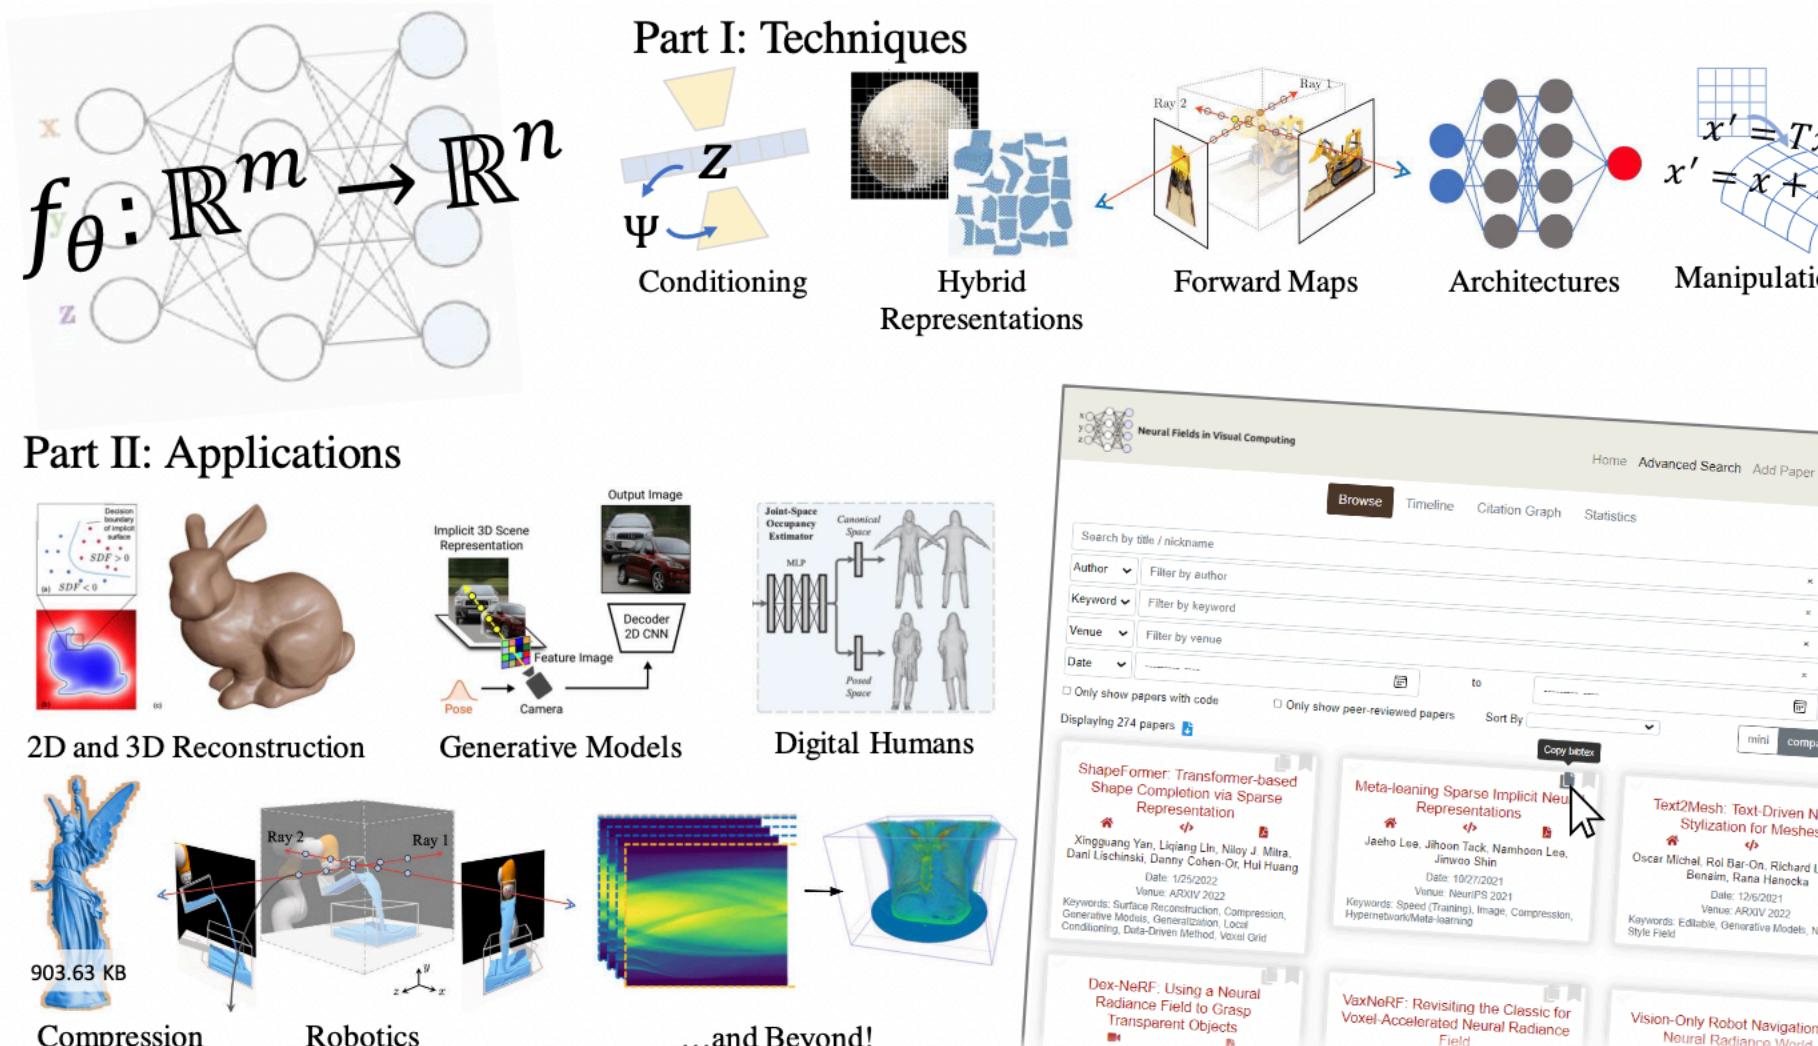
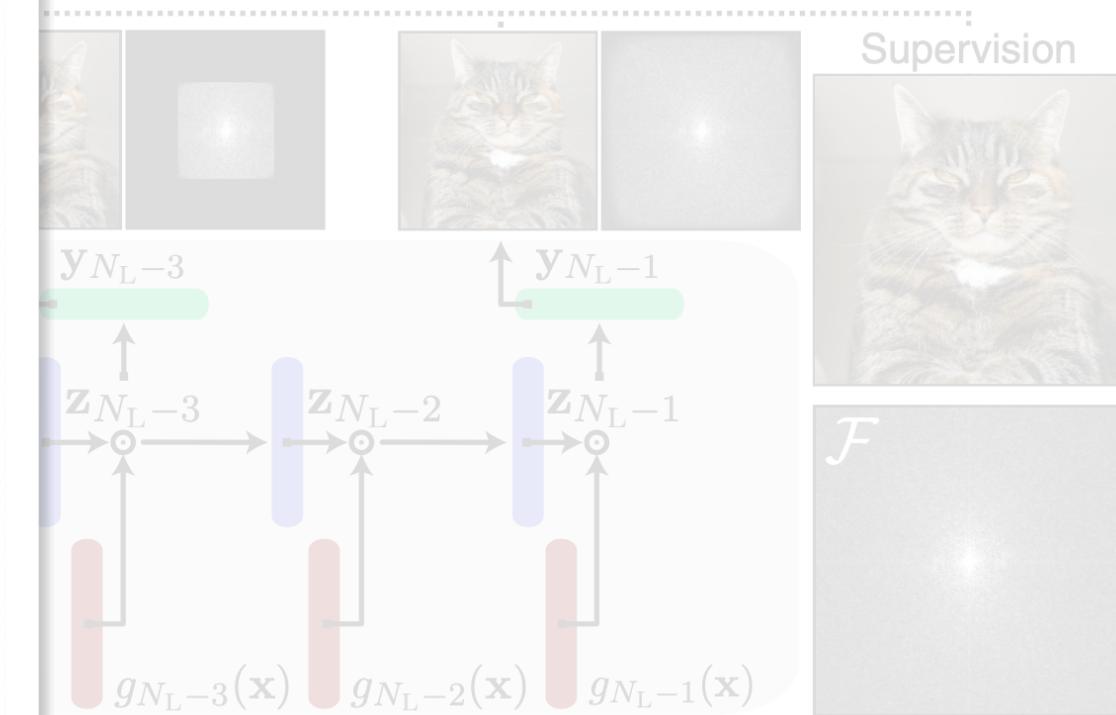


Figure 1: Contribution of this report. Following a survey of over 250 papers, we provide a review of (**Part I**) techniques in neural fields such as prior learning and conditioning, representations, forward maps, architectures, and manipulation, and of (**Part II**) applications in visual computing including 2D image processing, 3D scene reconstruction, generative modeling, digital humans, compression, robotics, and beyond. This report is complemented by a [community-driven website](https://neuralfields.cs.brown.edu/) with search, filtering, bibliographic, and visualization features.

ordinate networks for
representation
CVPR 2022

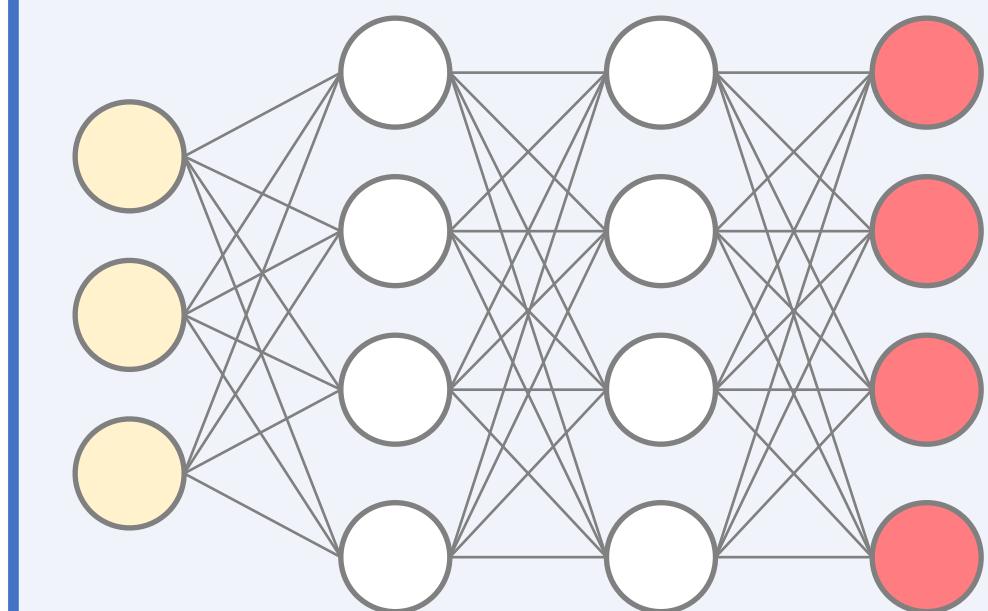


Neural Fields

- Storage memory **does not grow with spatial resolution or number of spatial dimensions.**
- **Slow** sampling: Each query takes a forward pass through the neural net. Means GPU-memory intensive forward passes.
- Does **not** expose locality: Can't identify set of parameters / direction in parameter space that encodes particular spatial location.
- **Inconvenient processing:** cannot run convolutions.
- But: automatic adaptive resolution. Over optimization, will converge to assign more compute to higher-frequency areas.

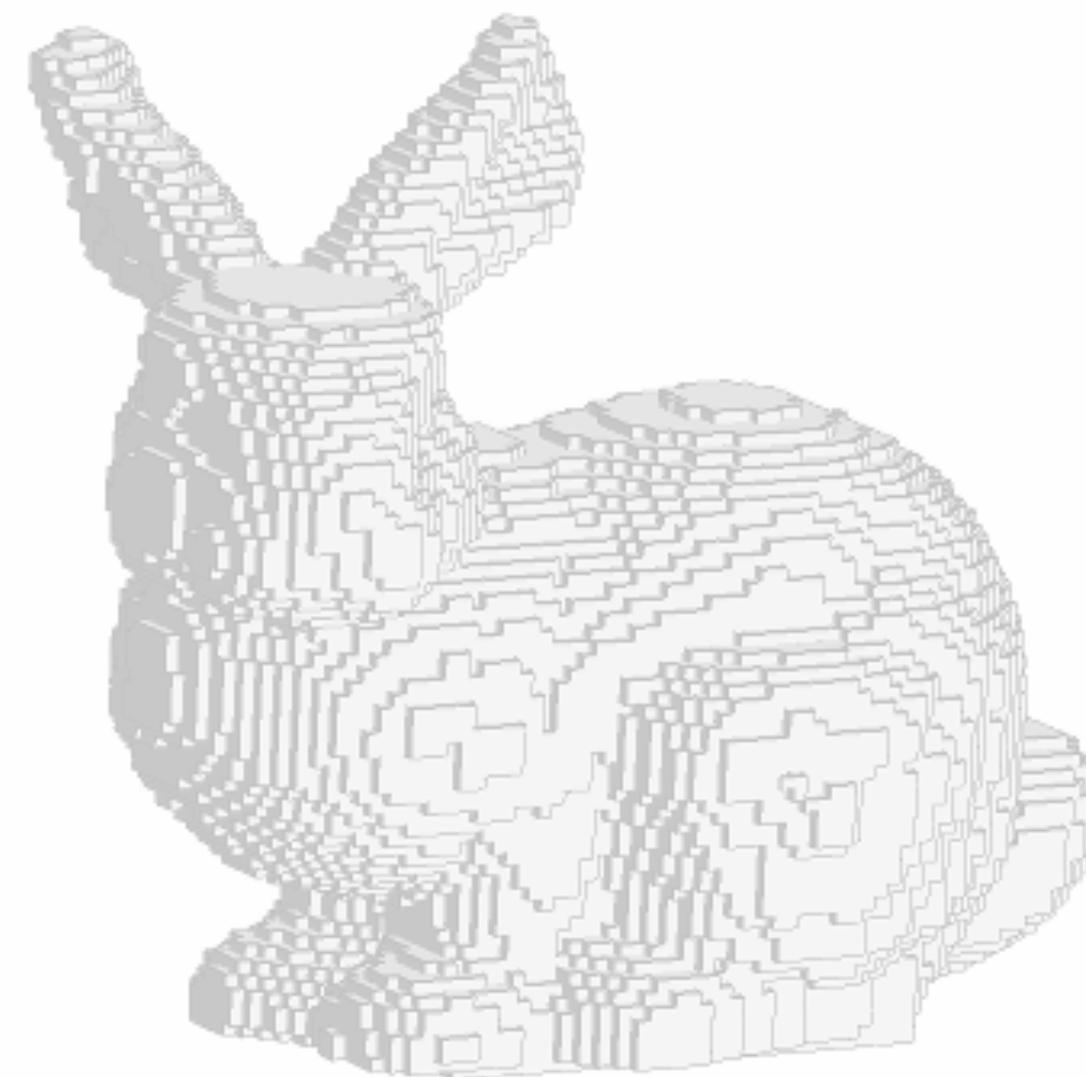
Neural Fields

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$

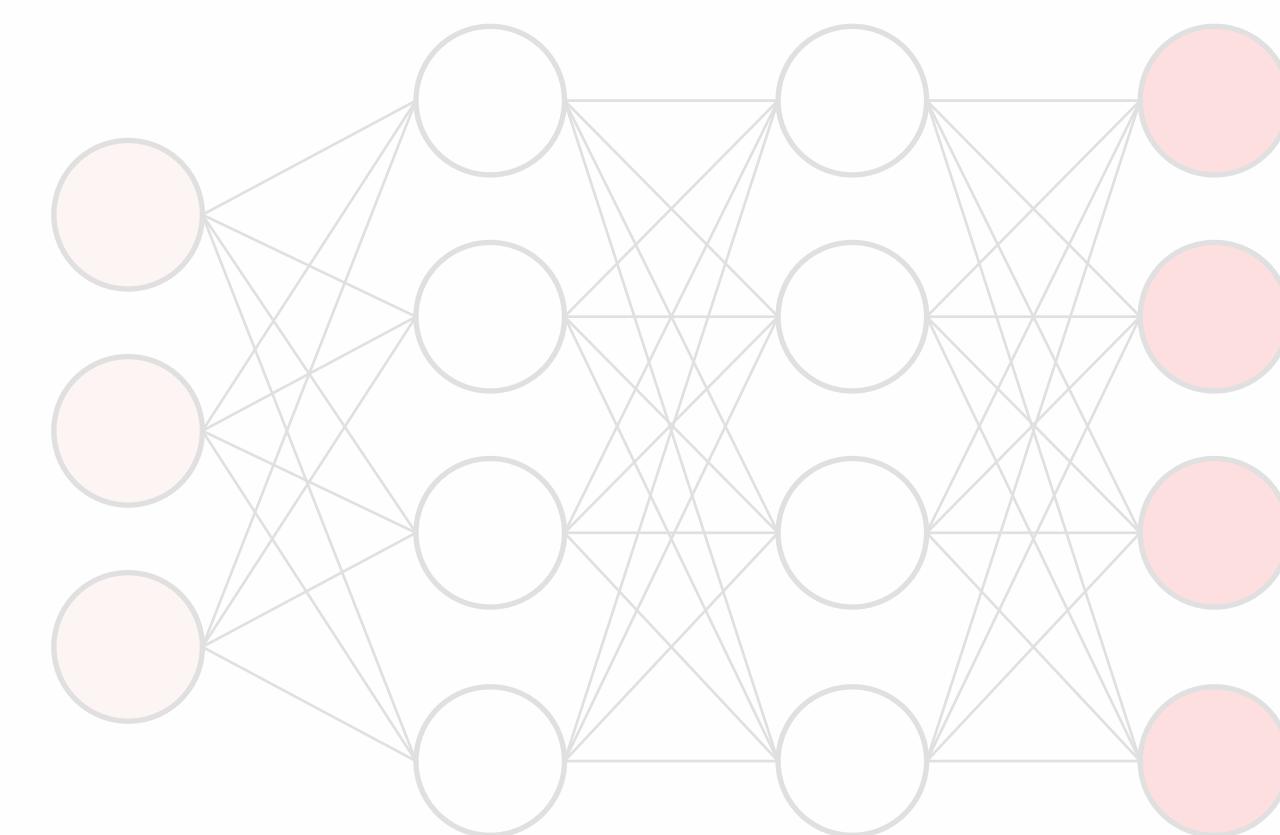


Note: Despite their name, has nothing to do with Machine Learning itself.
It's as “smart” as a voxelgrid.

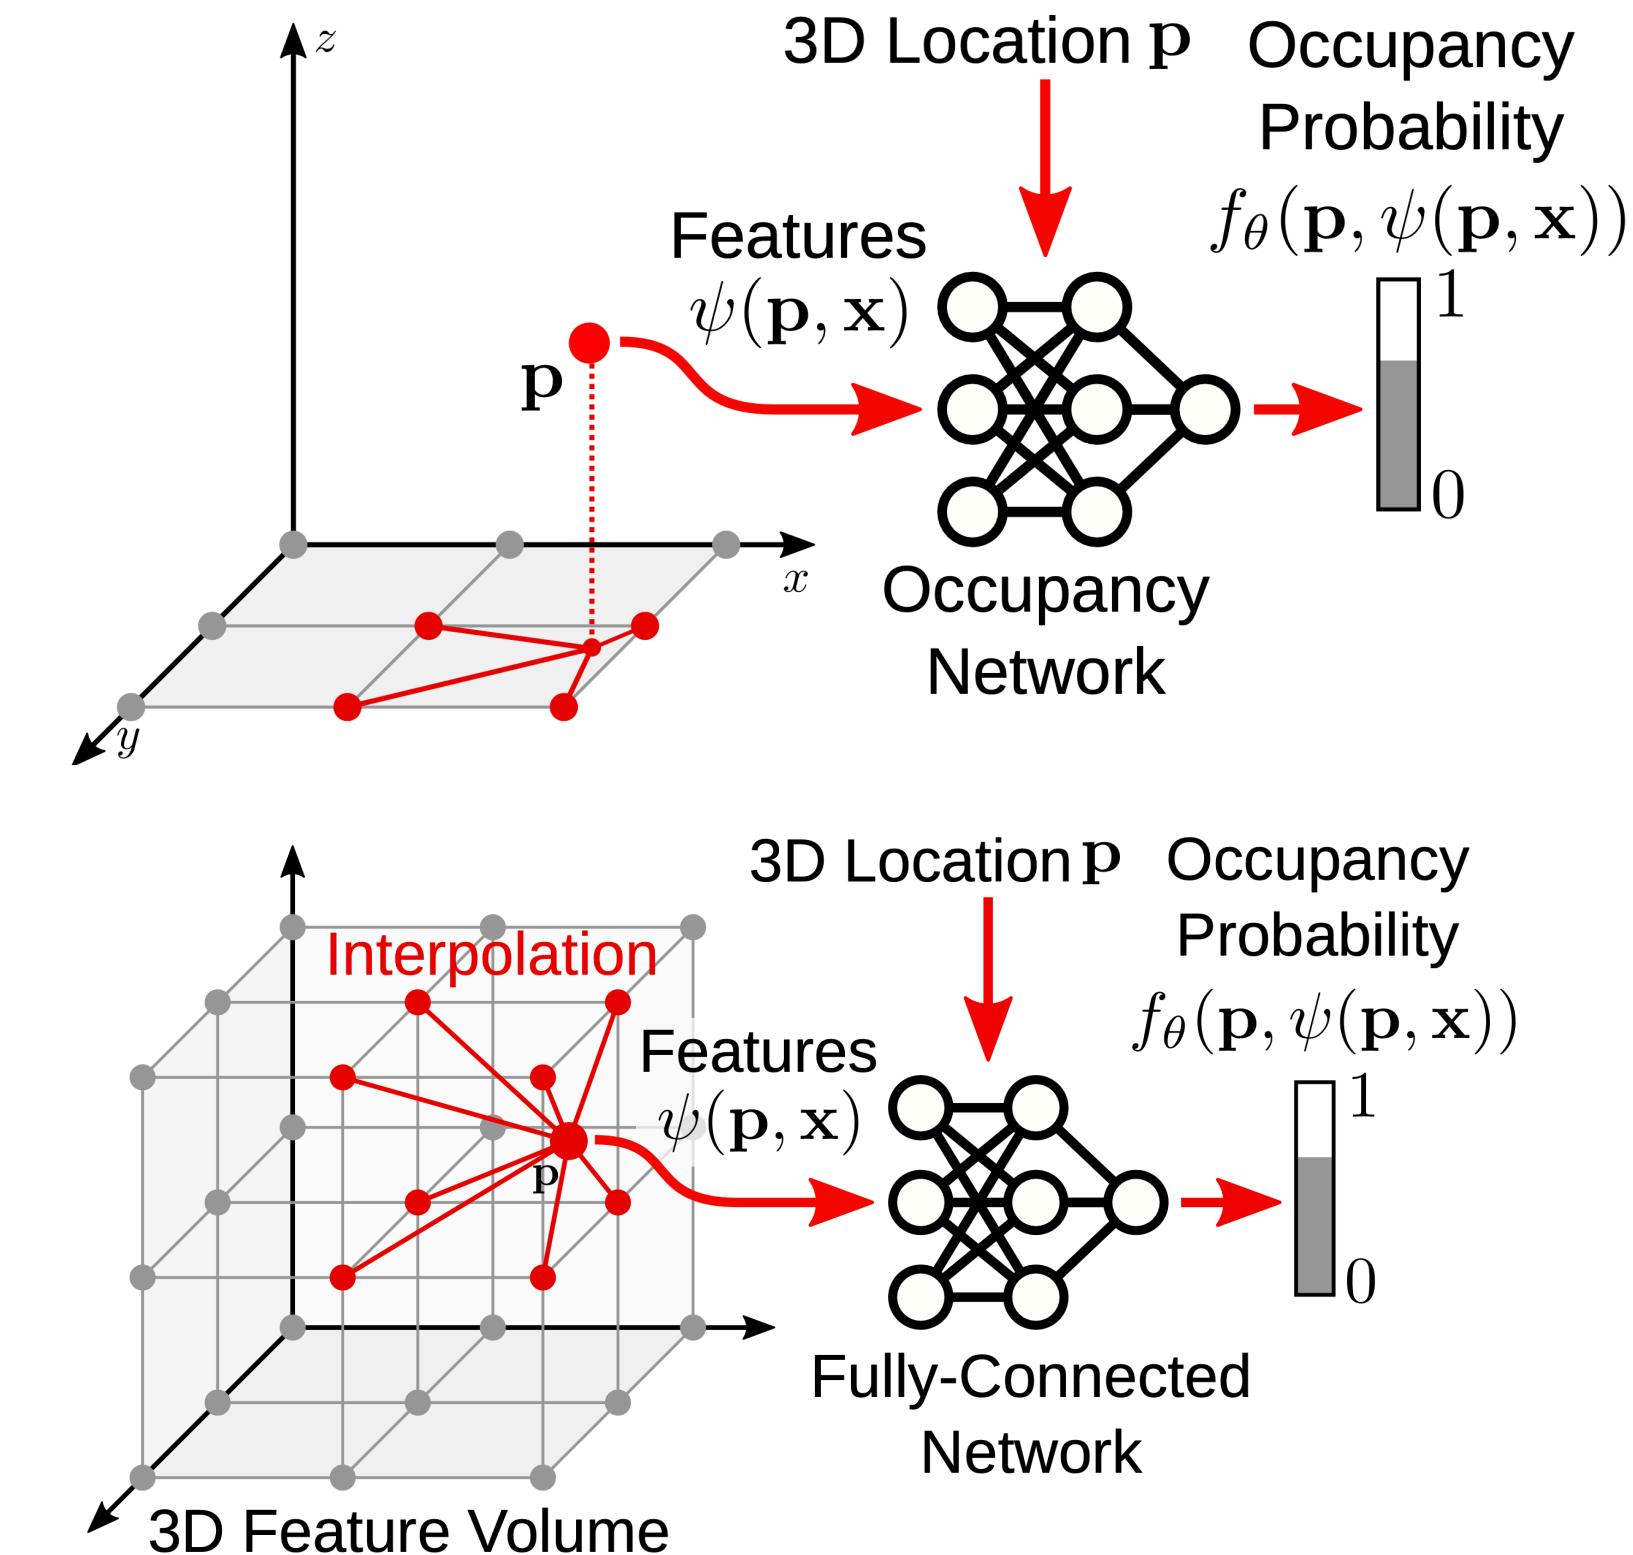
Hybrid discrete / continuous reps



Discrete Parameterizations

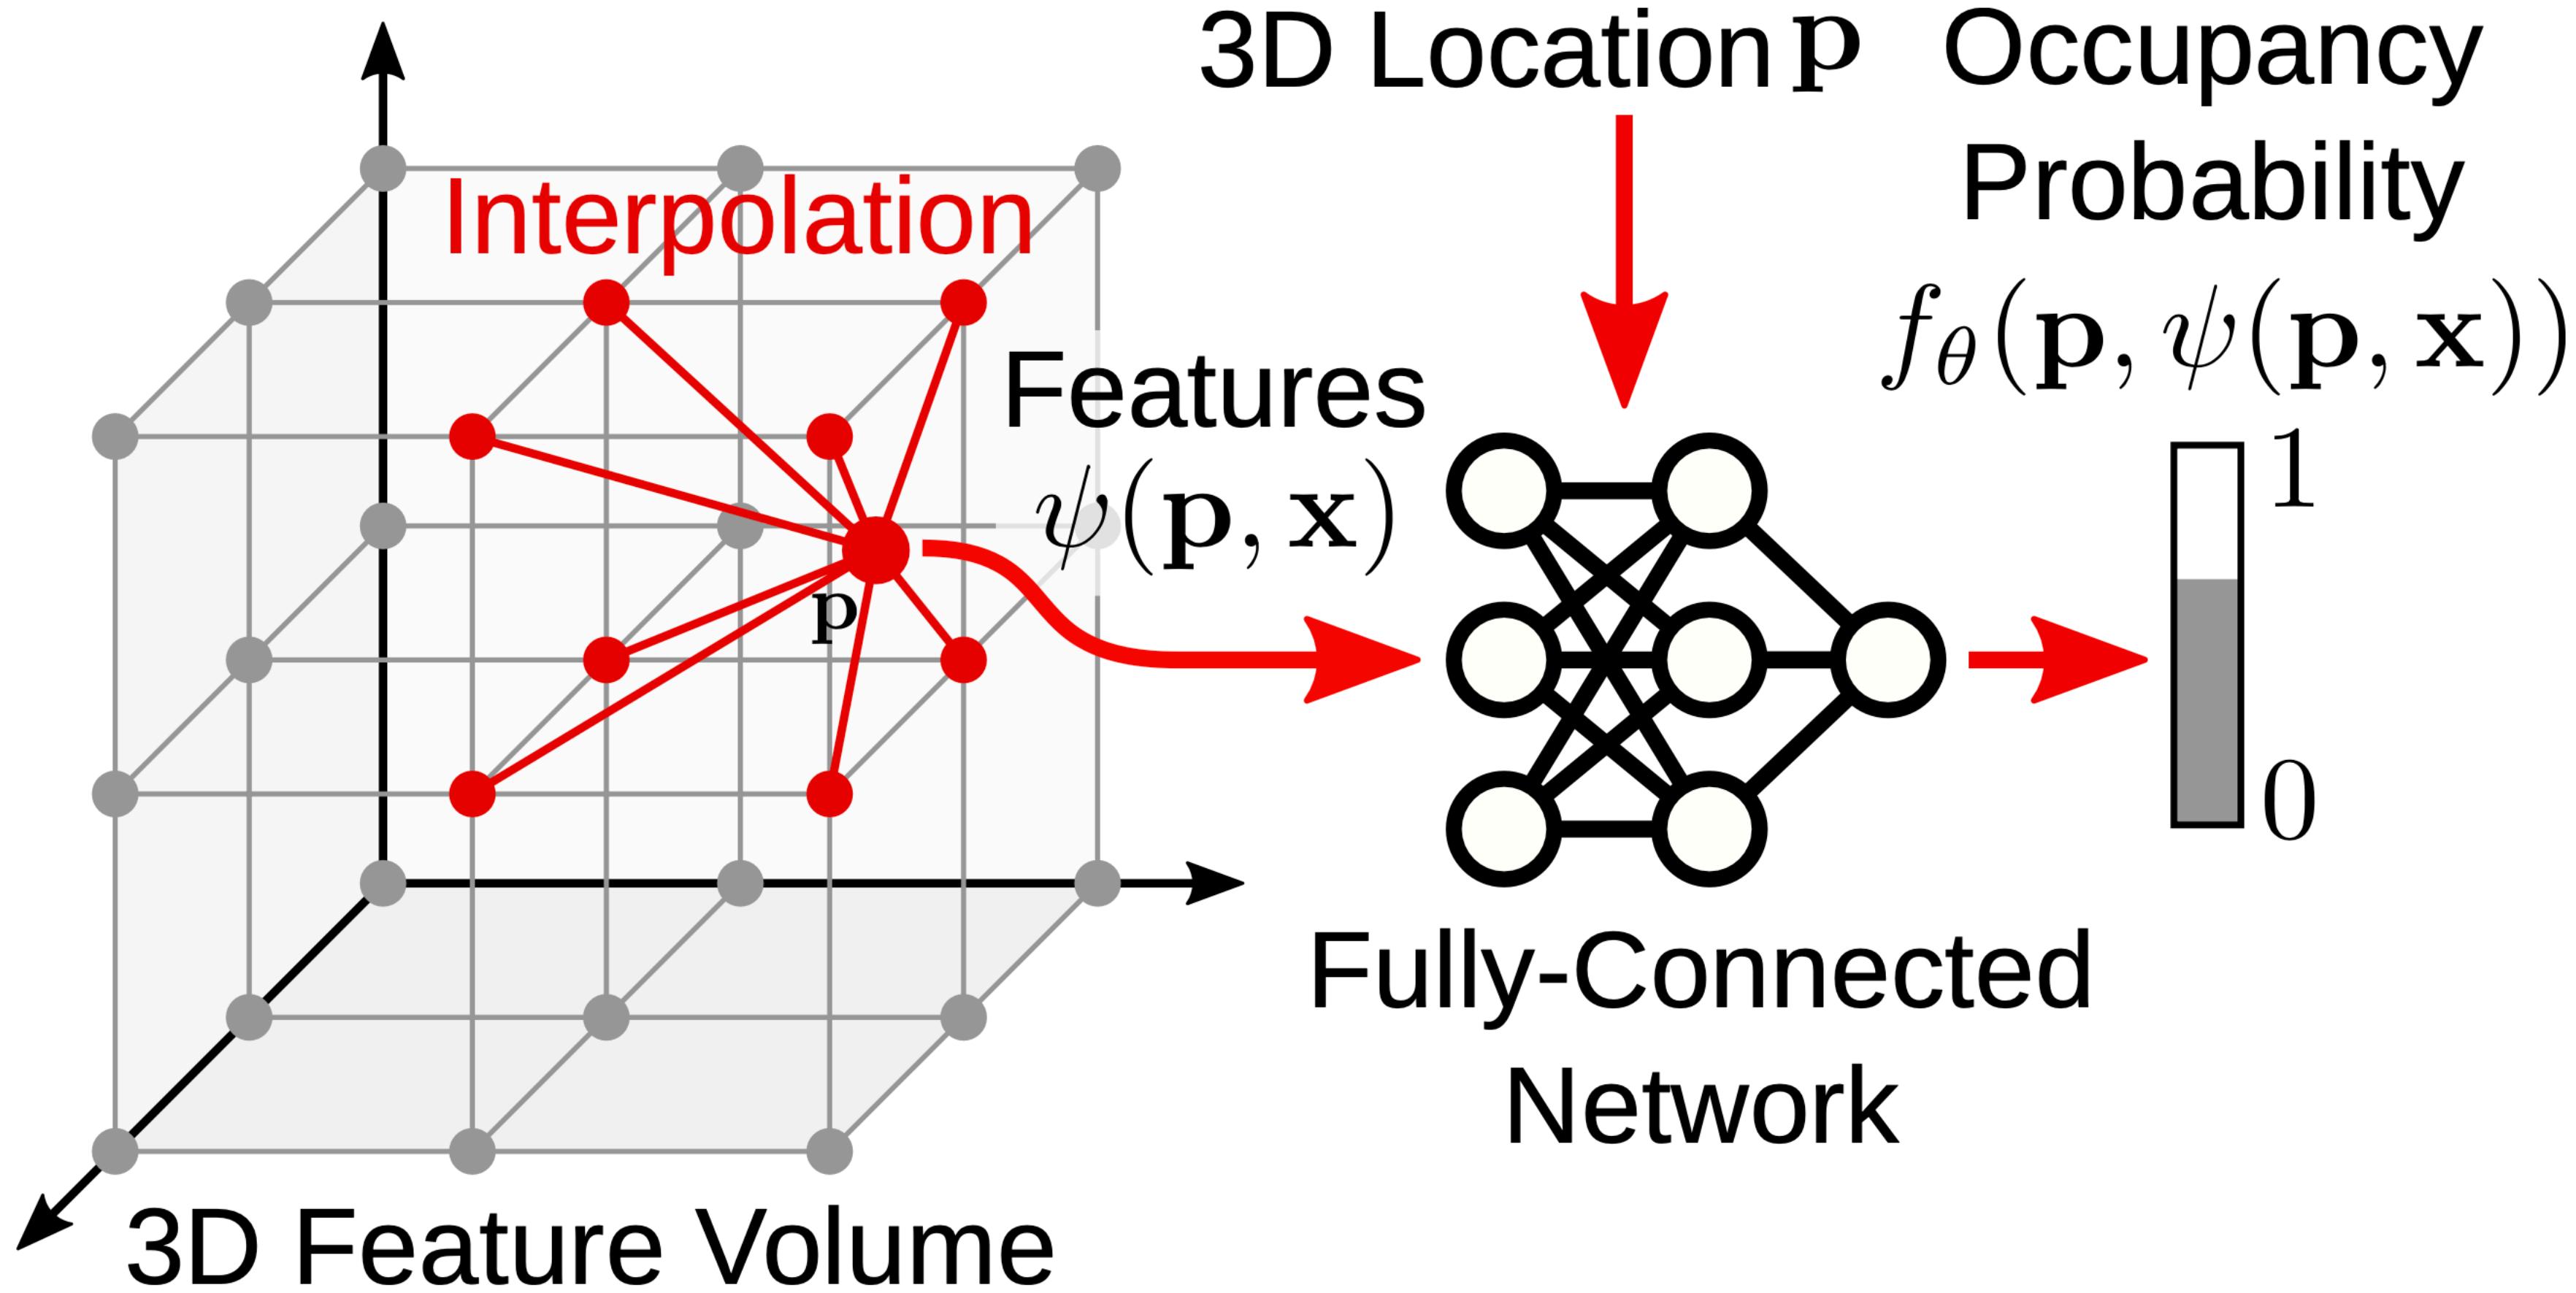


Continuous Parameterizations



Hybrid Parameterizations

Basic idea: Use Neural Net as Interpolation Kernel



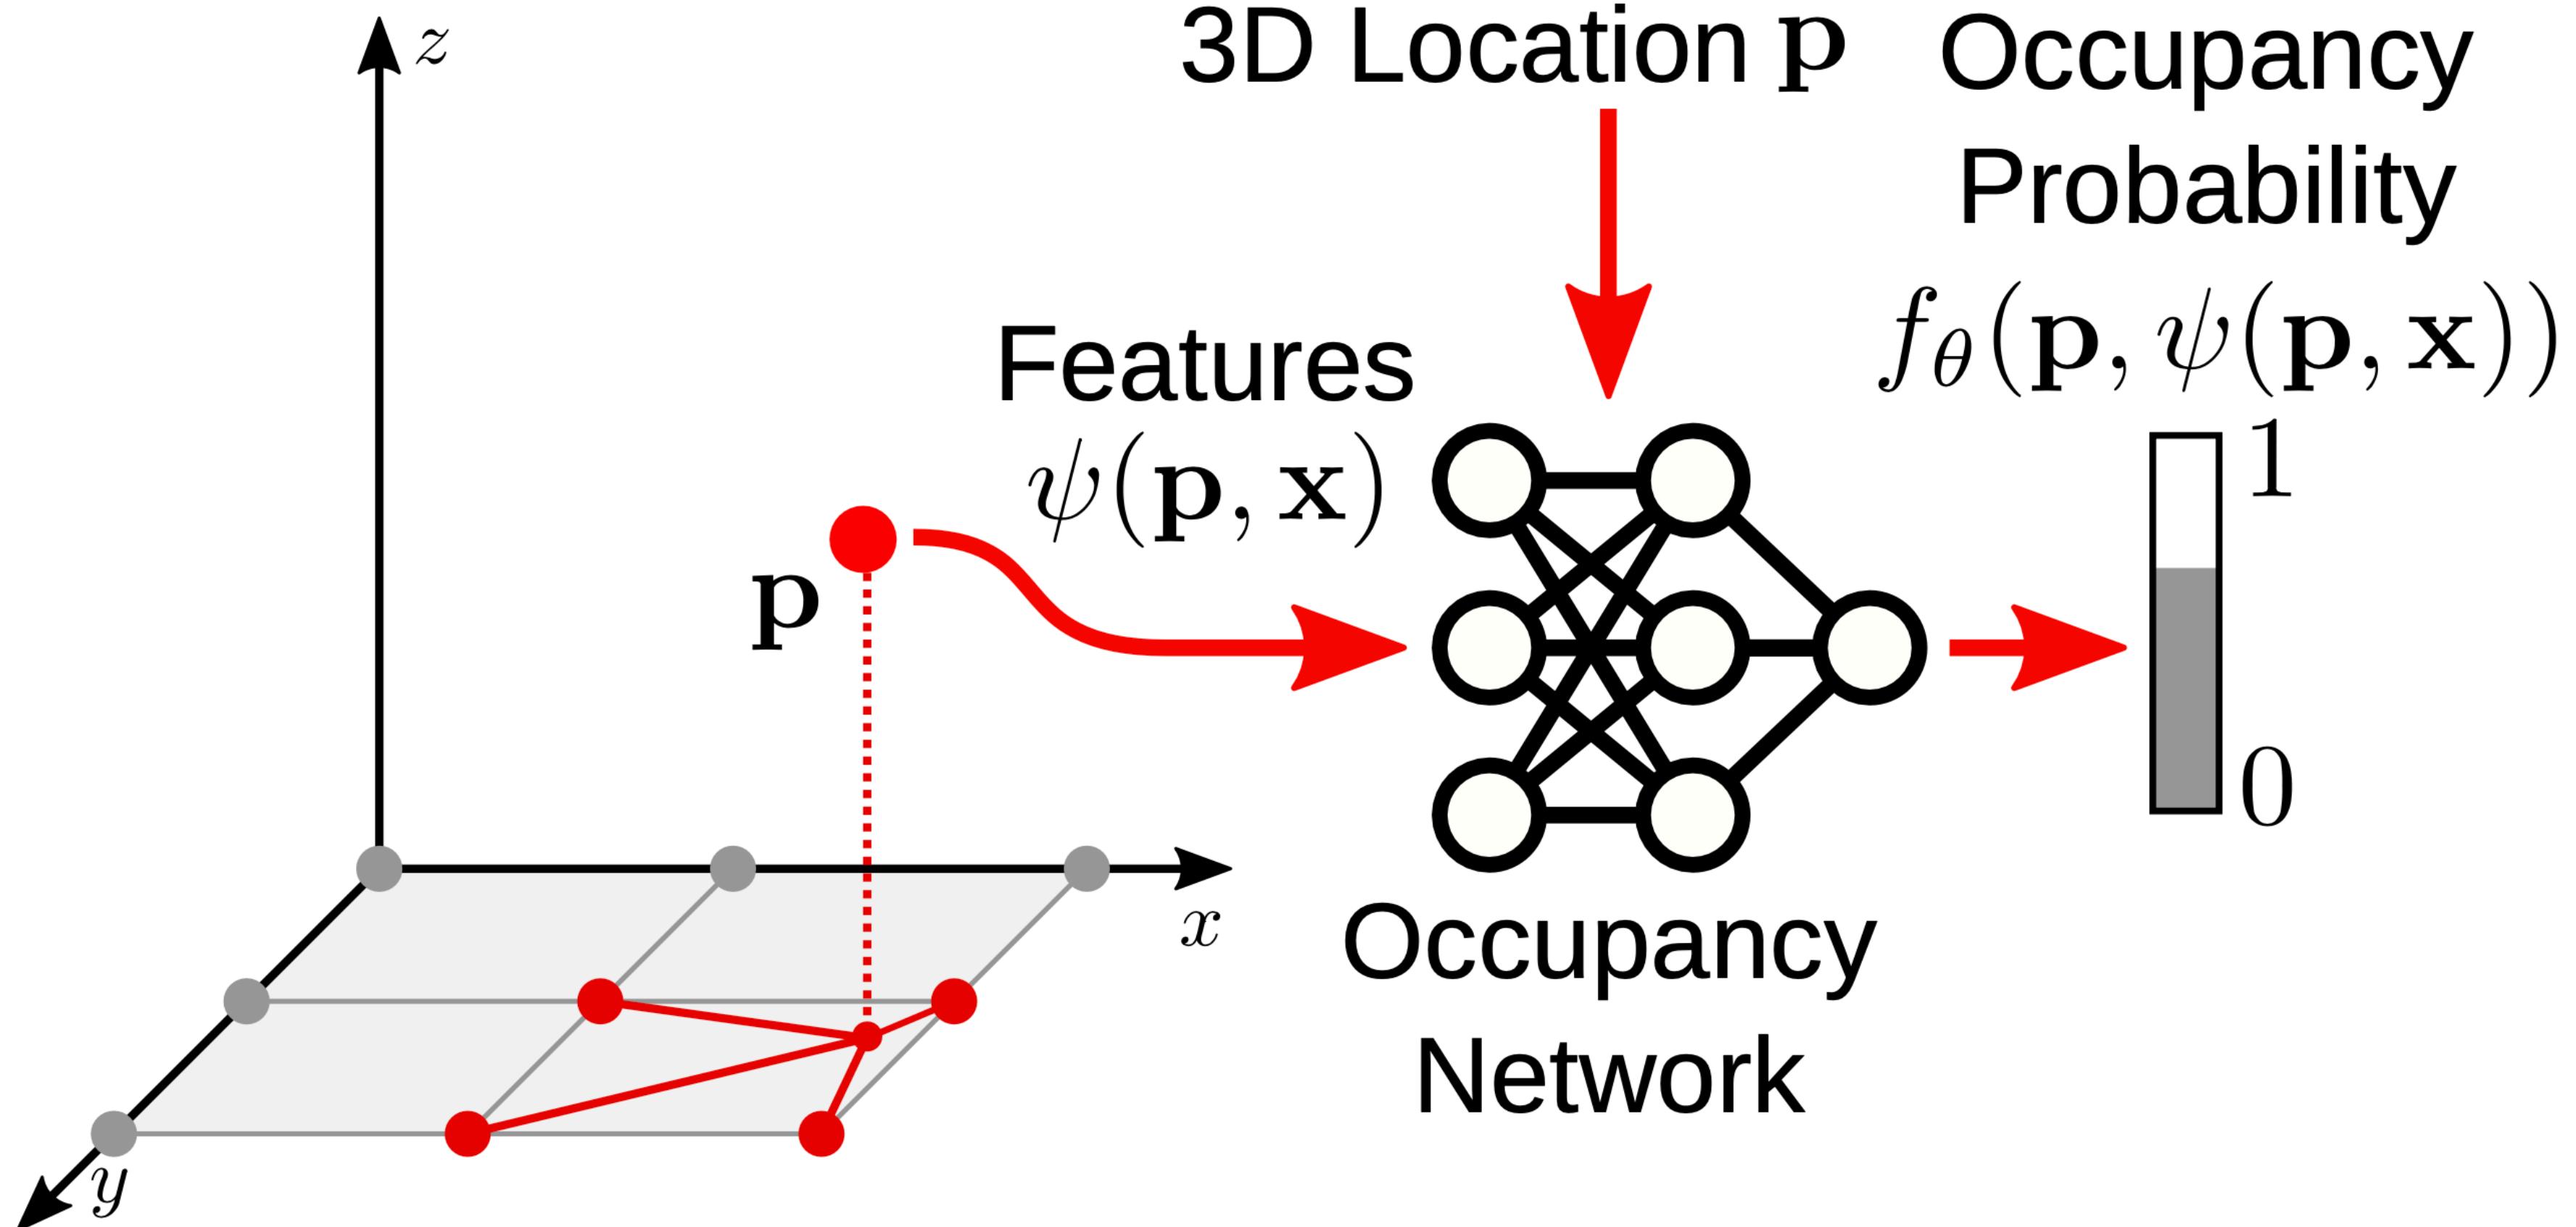
Convolutional Occupancy Networks [Peng et al. 2020]

Local Implicit Grid Representations for 3D Scenes [Jiang et al. 2020]

Implicit Functions in Feature Space for 3D Shape Reconstruction and Completion [Chibane et al. 2020]

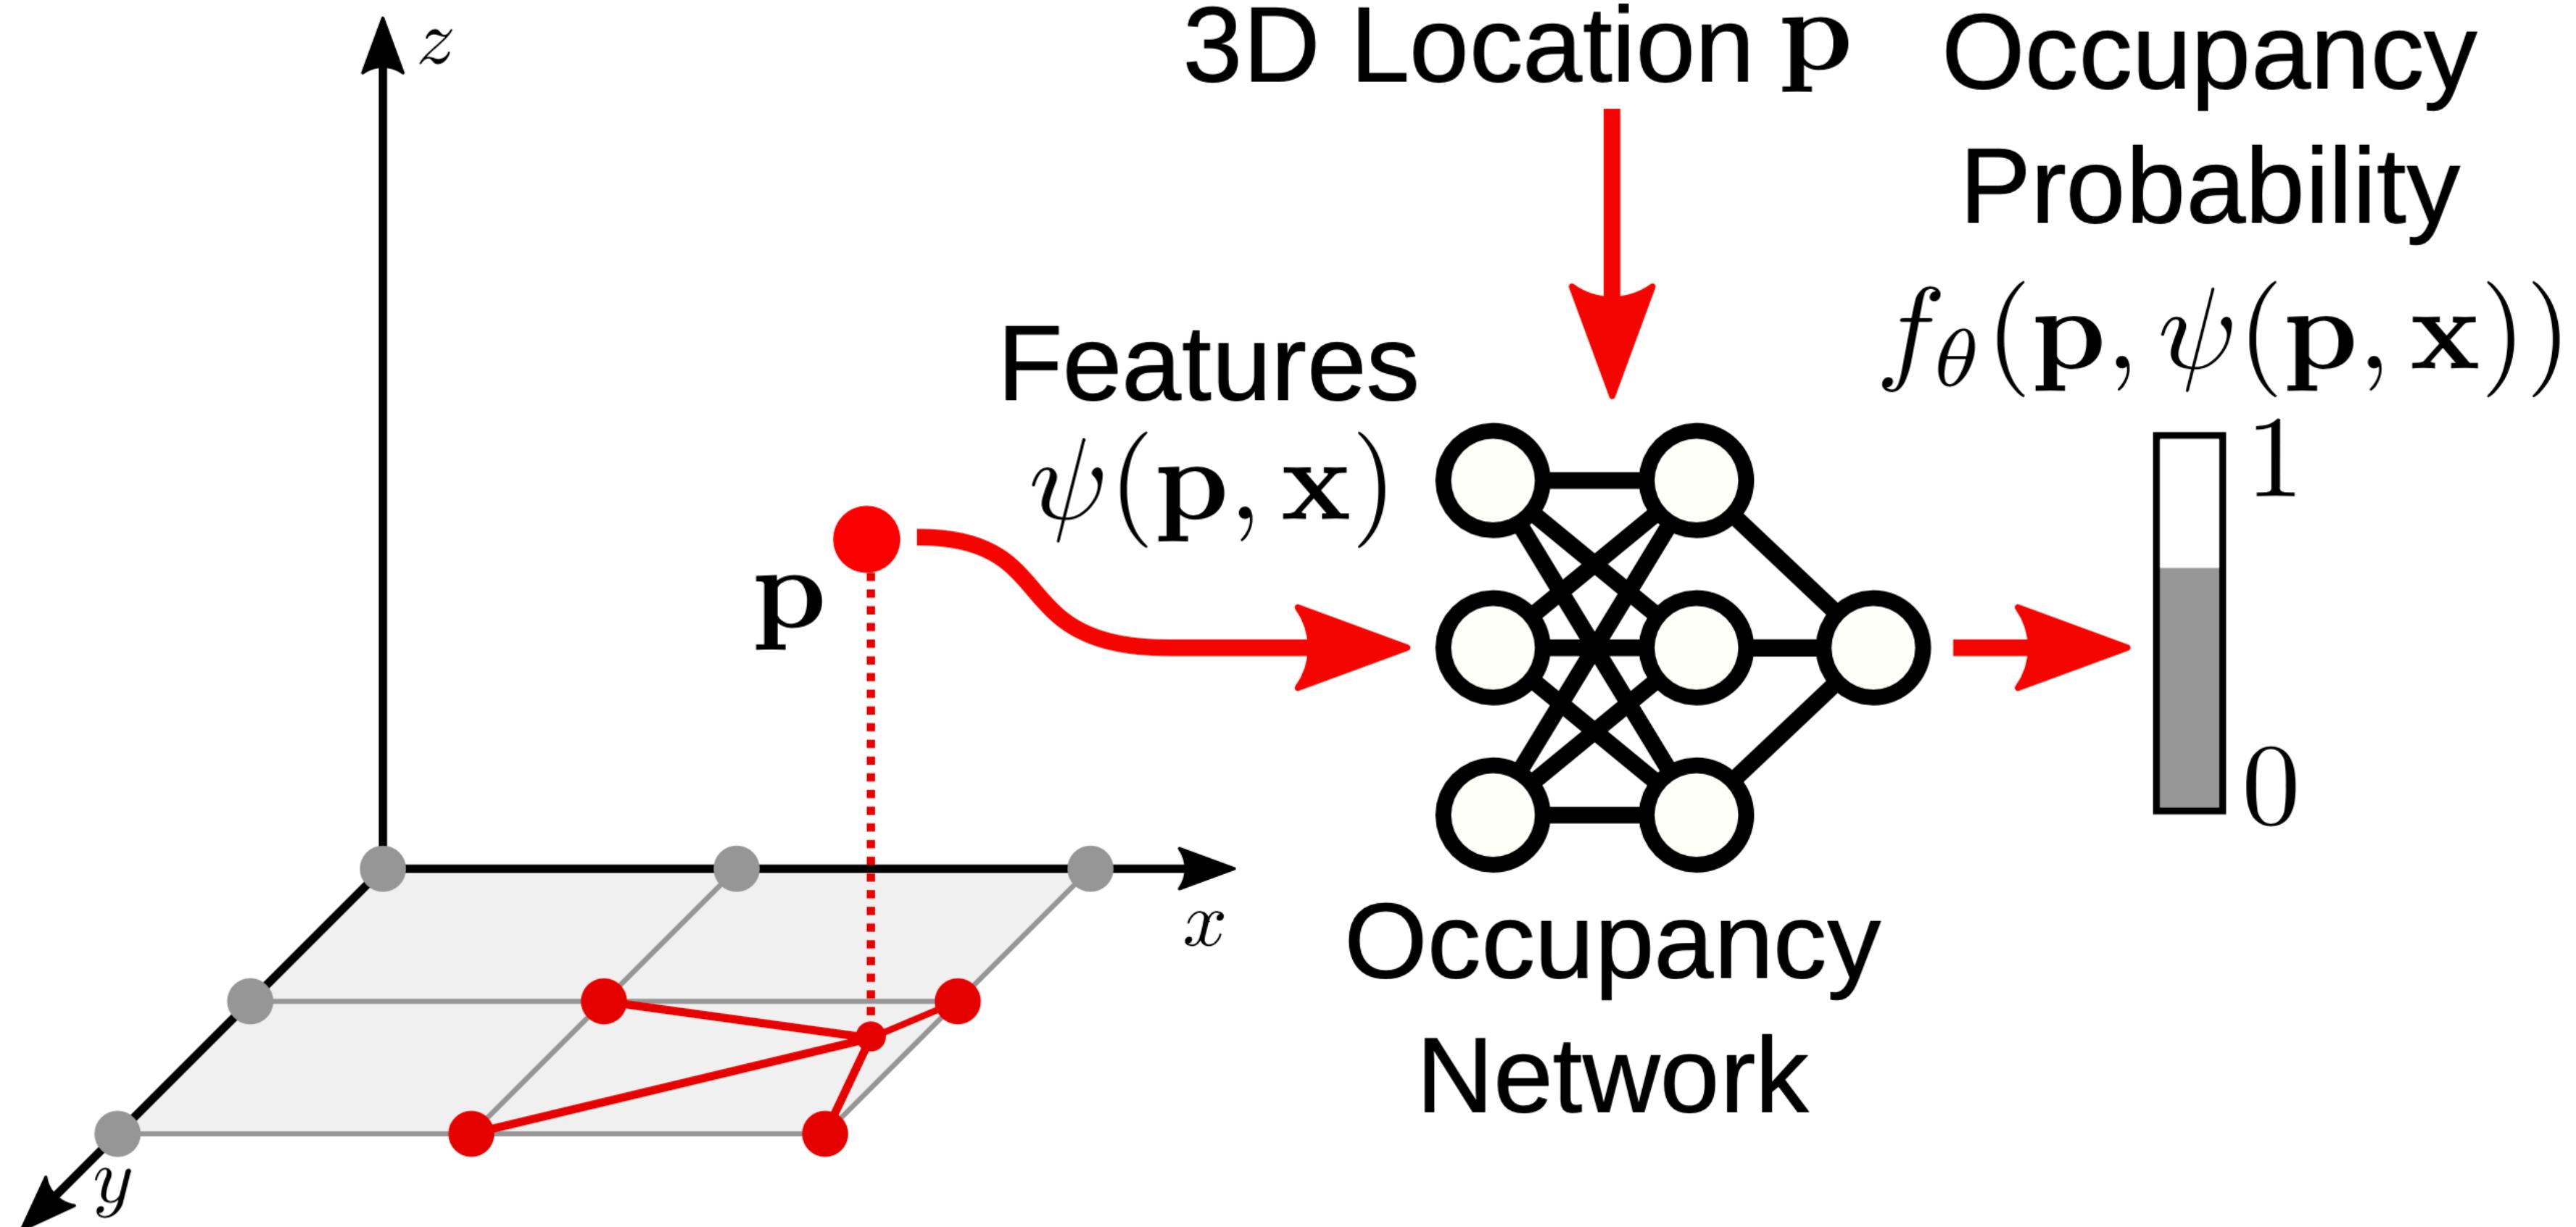
Deep Local Shapes: Learning Local SDF Priors for Detailed 3D Reconstruction [Chabra et al. 2020]

Ground plan & Orthographic Projection



Any assumptions made on scene? Any limitations?

Ground plan & Orthographic Projection



No. You can still represent *any* scene - some info stored in NN, some in grid.

Tri-Planes and Orthographic Projection

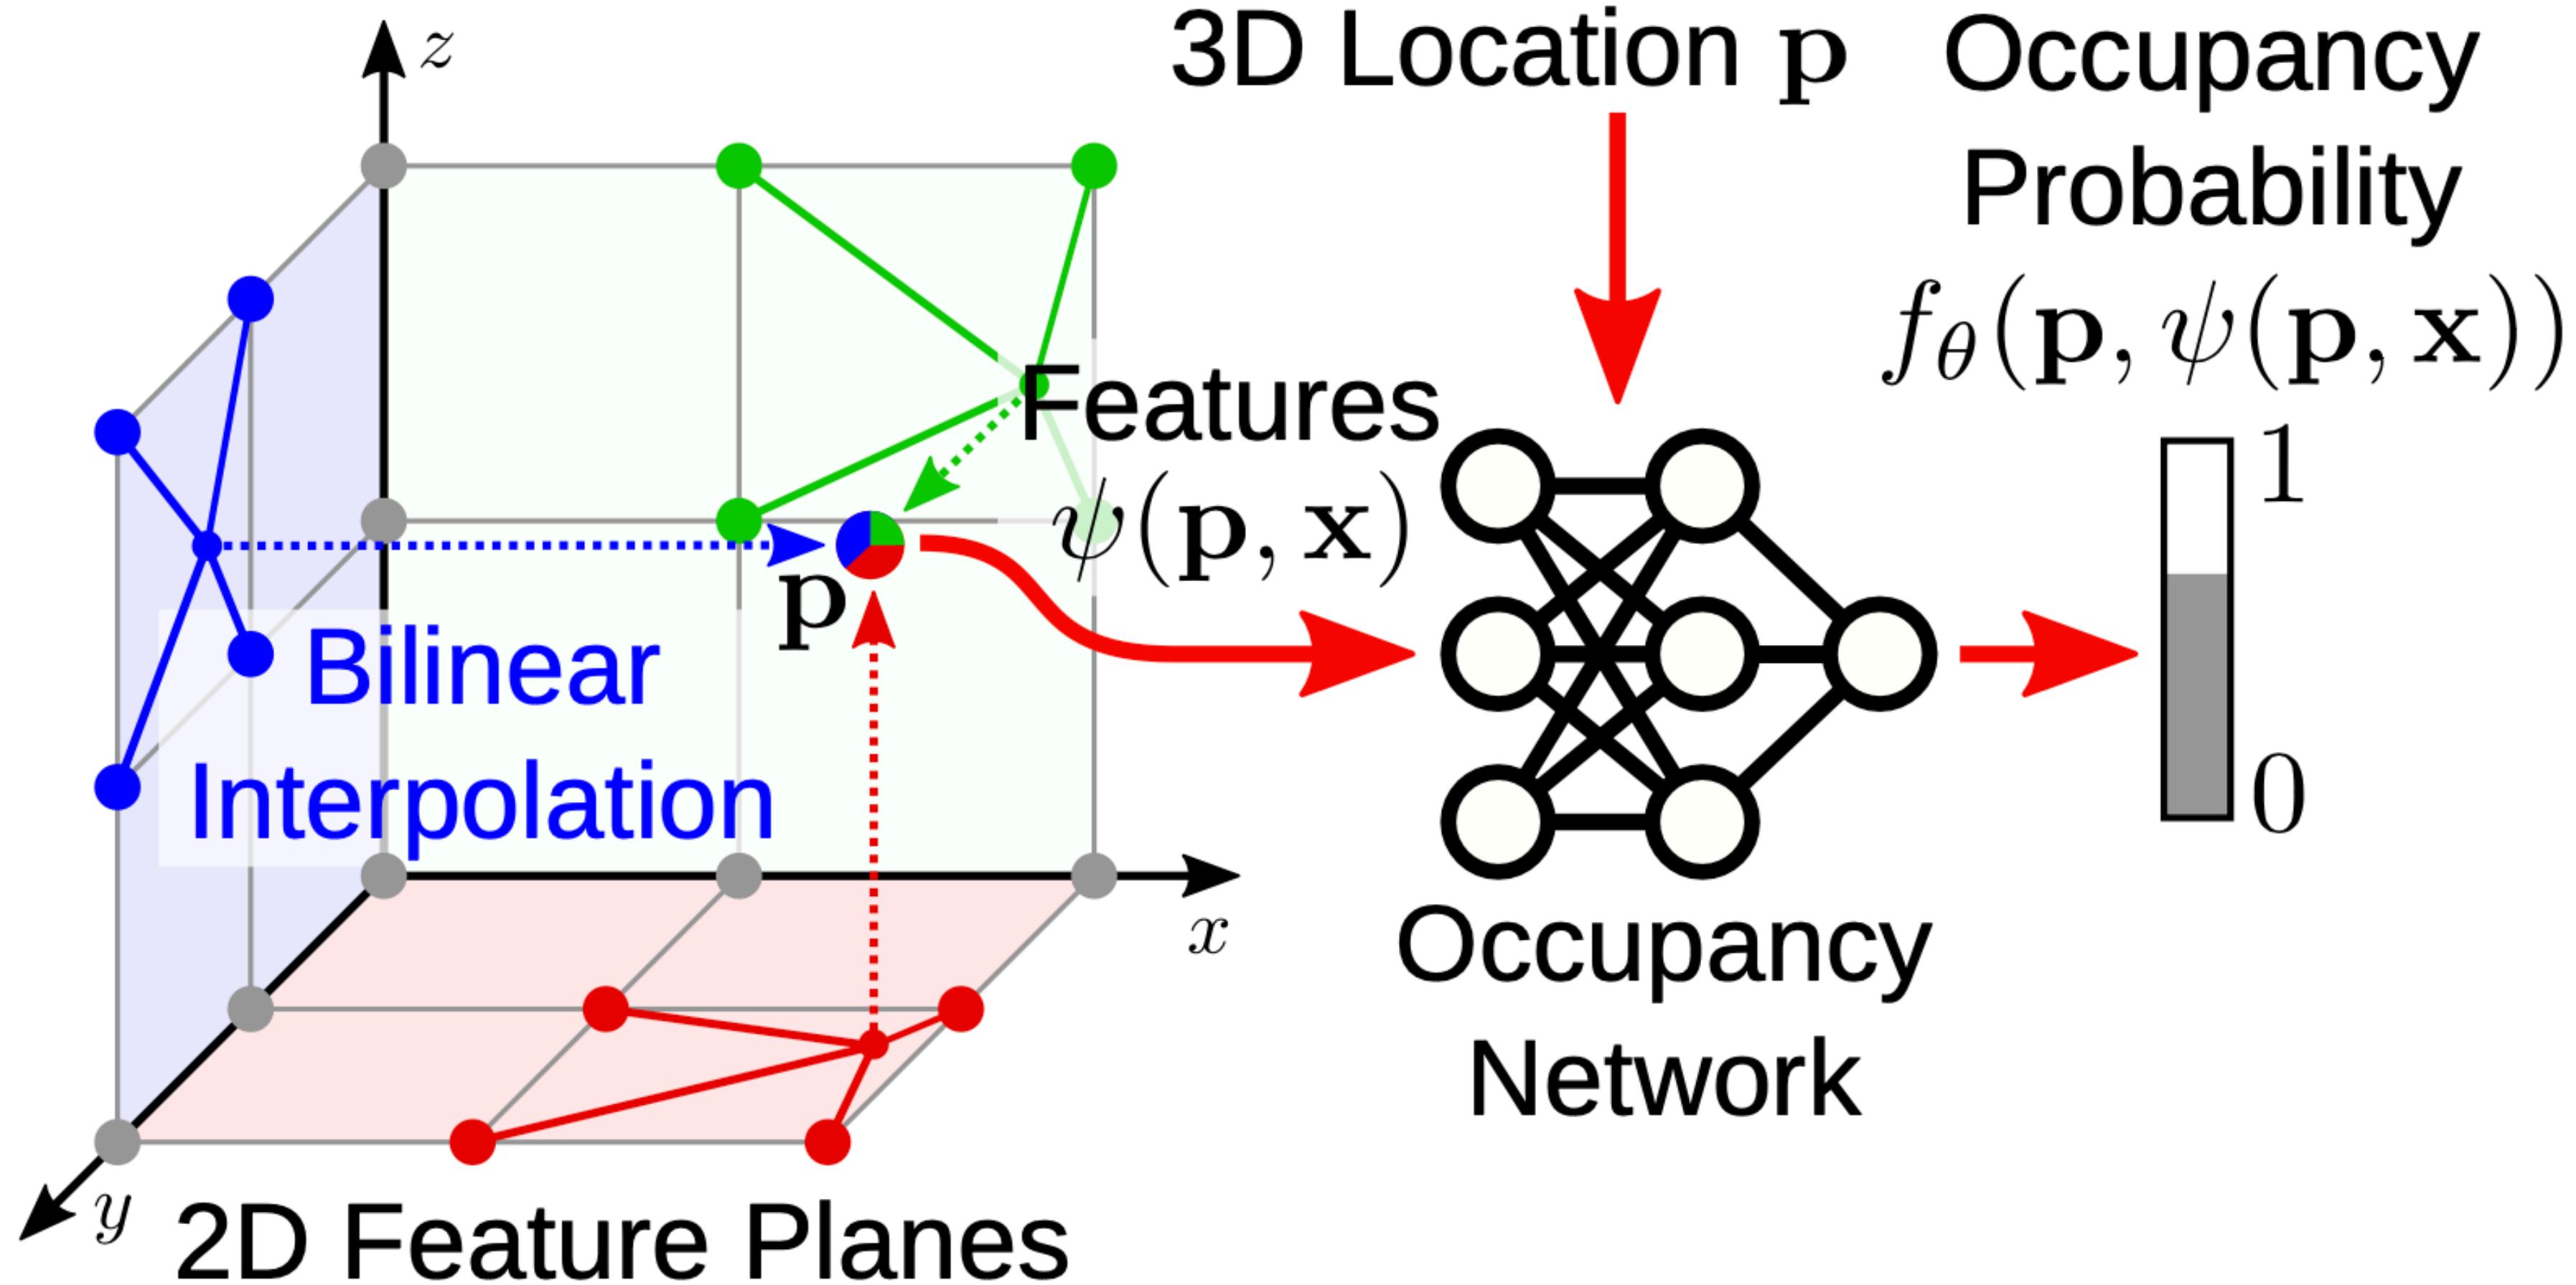
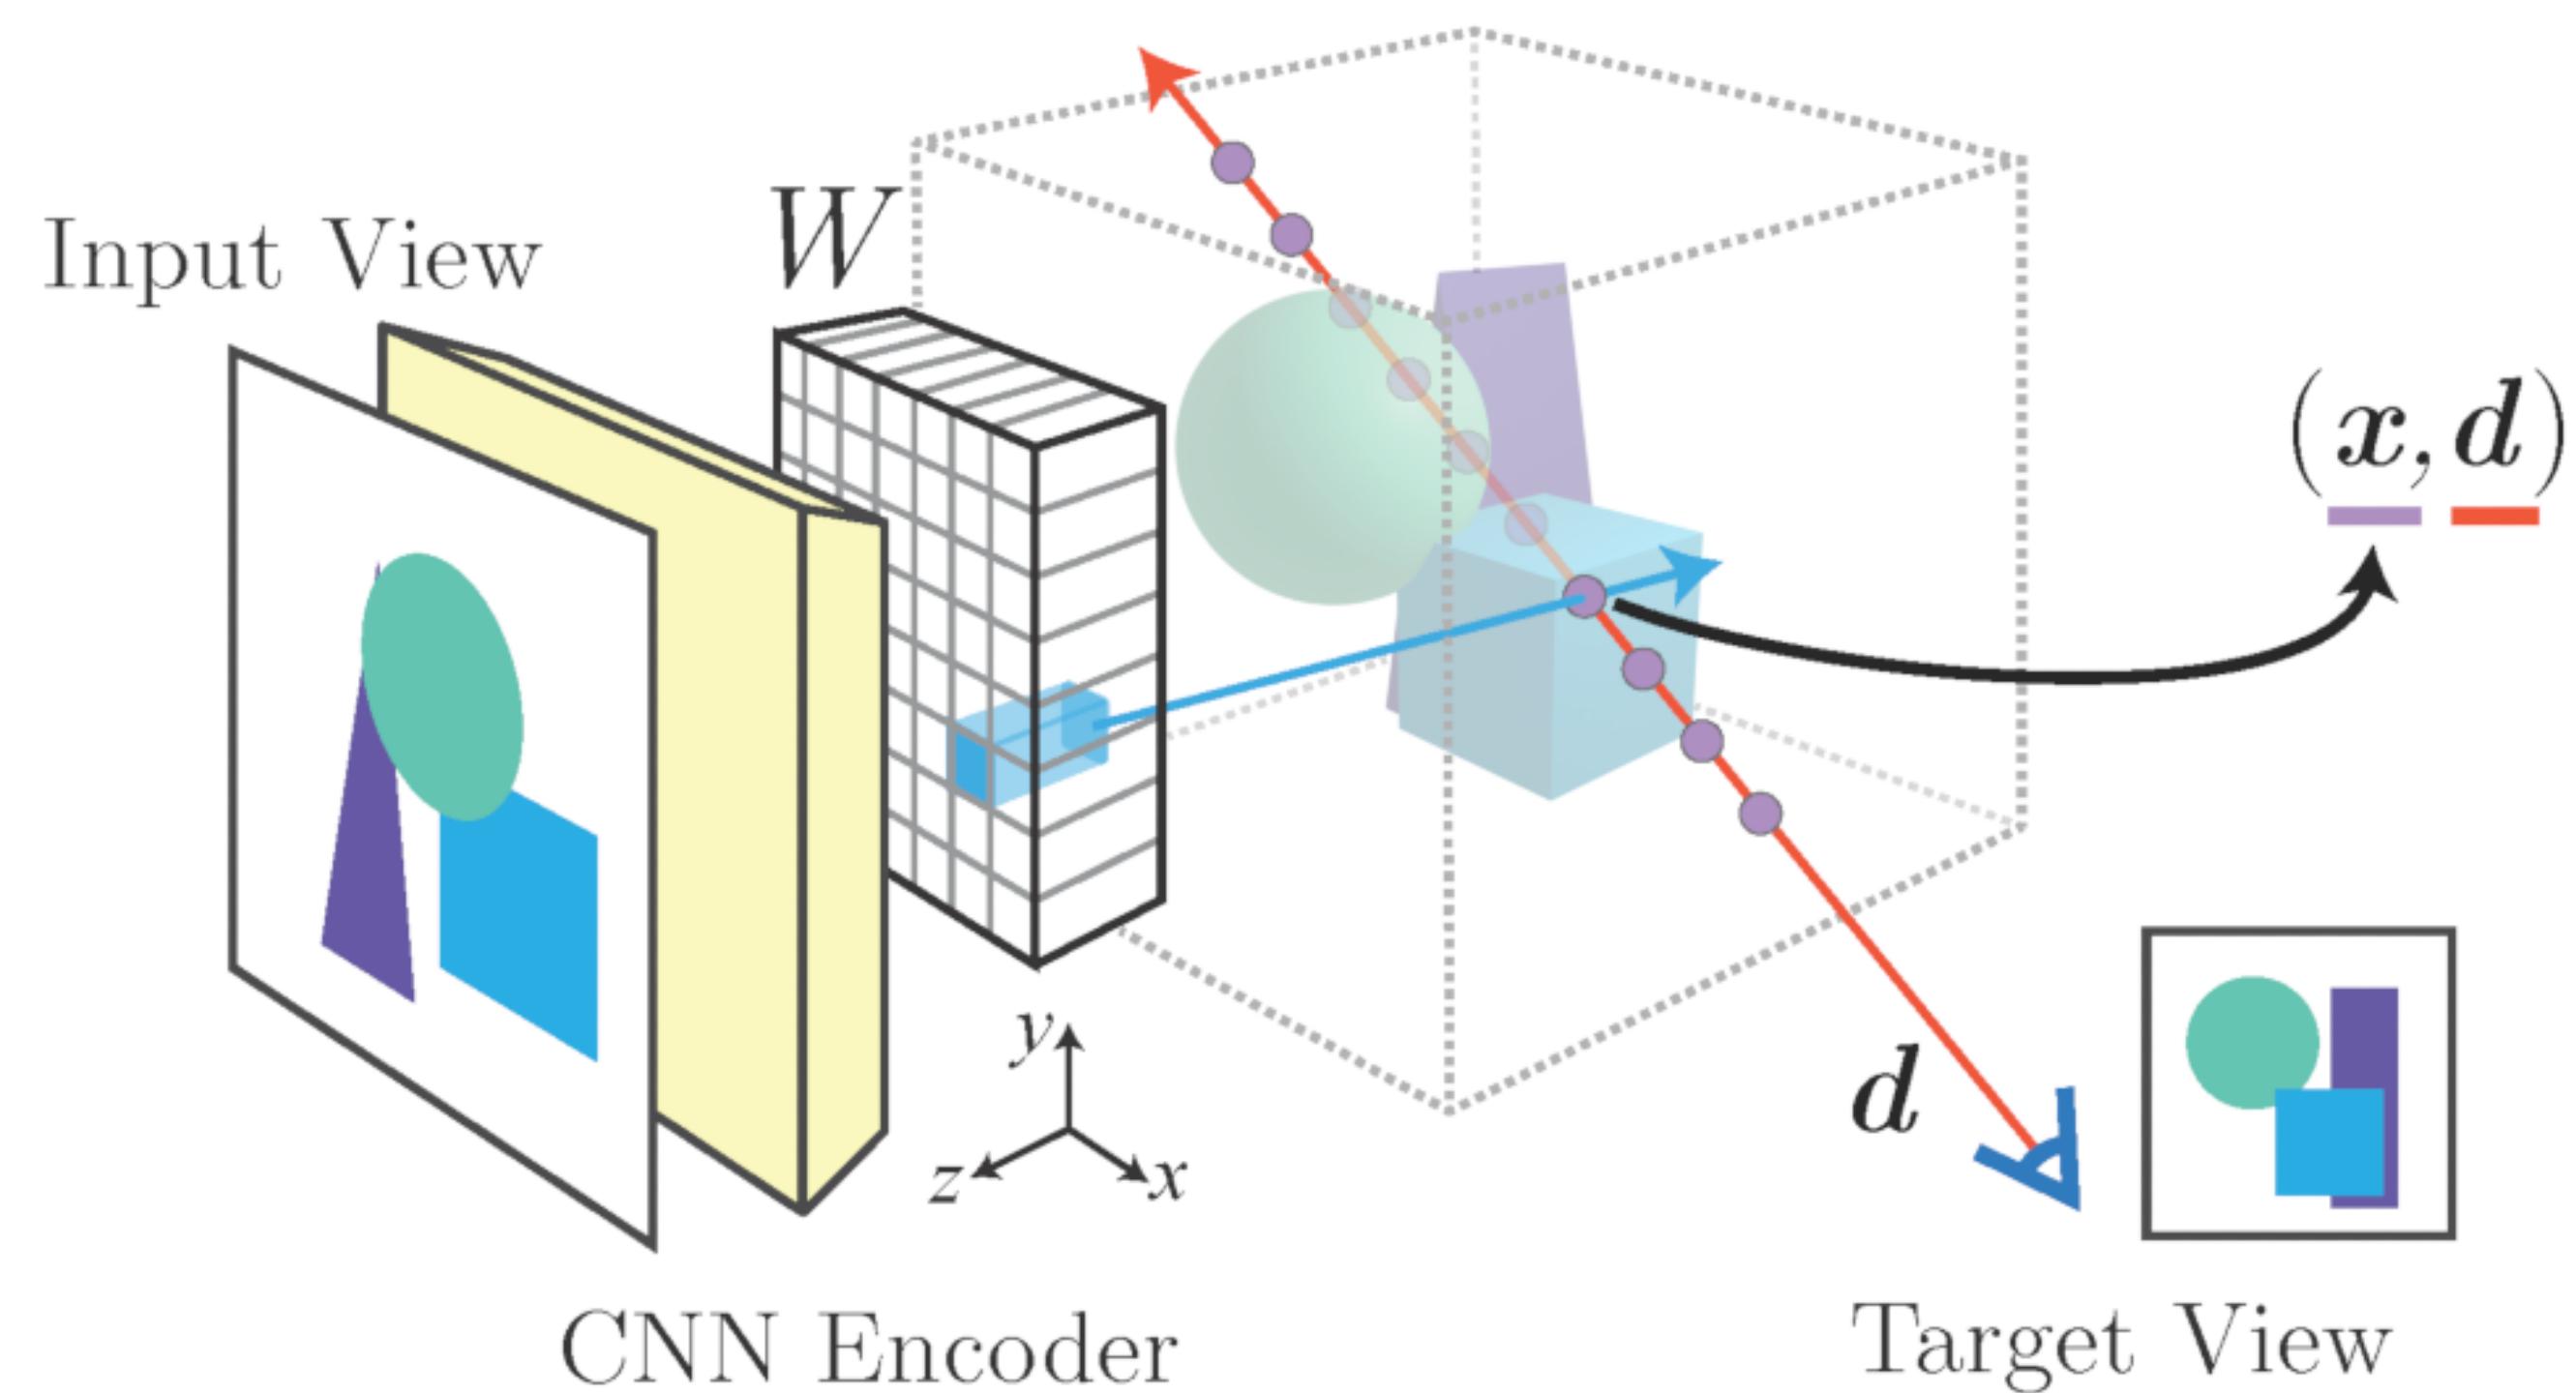
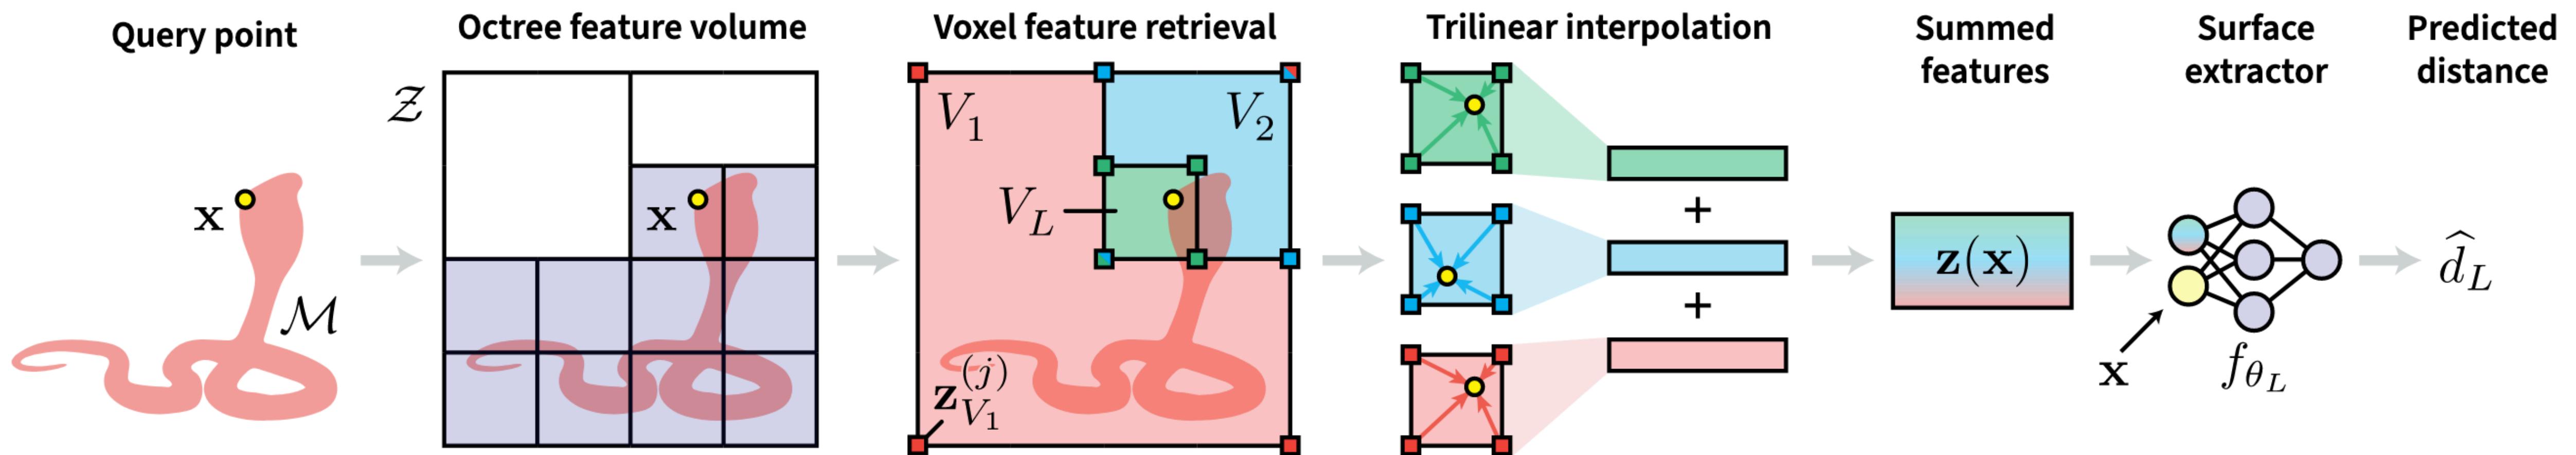


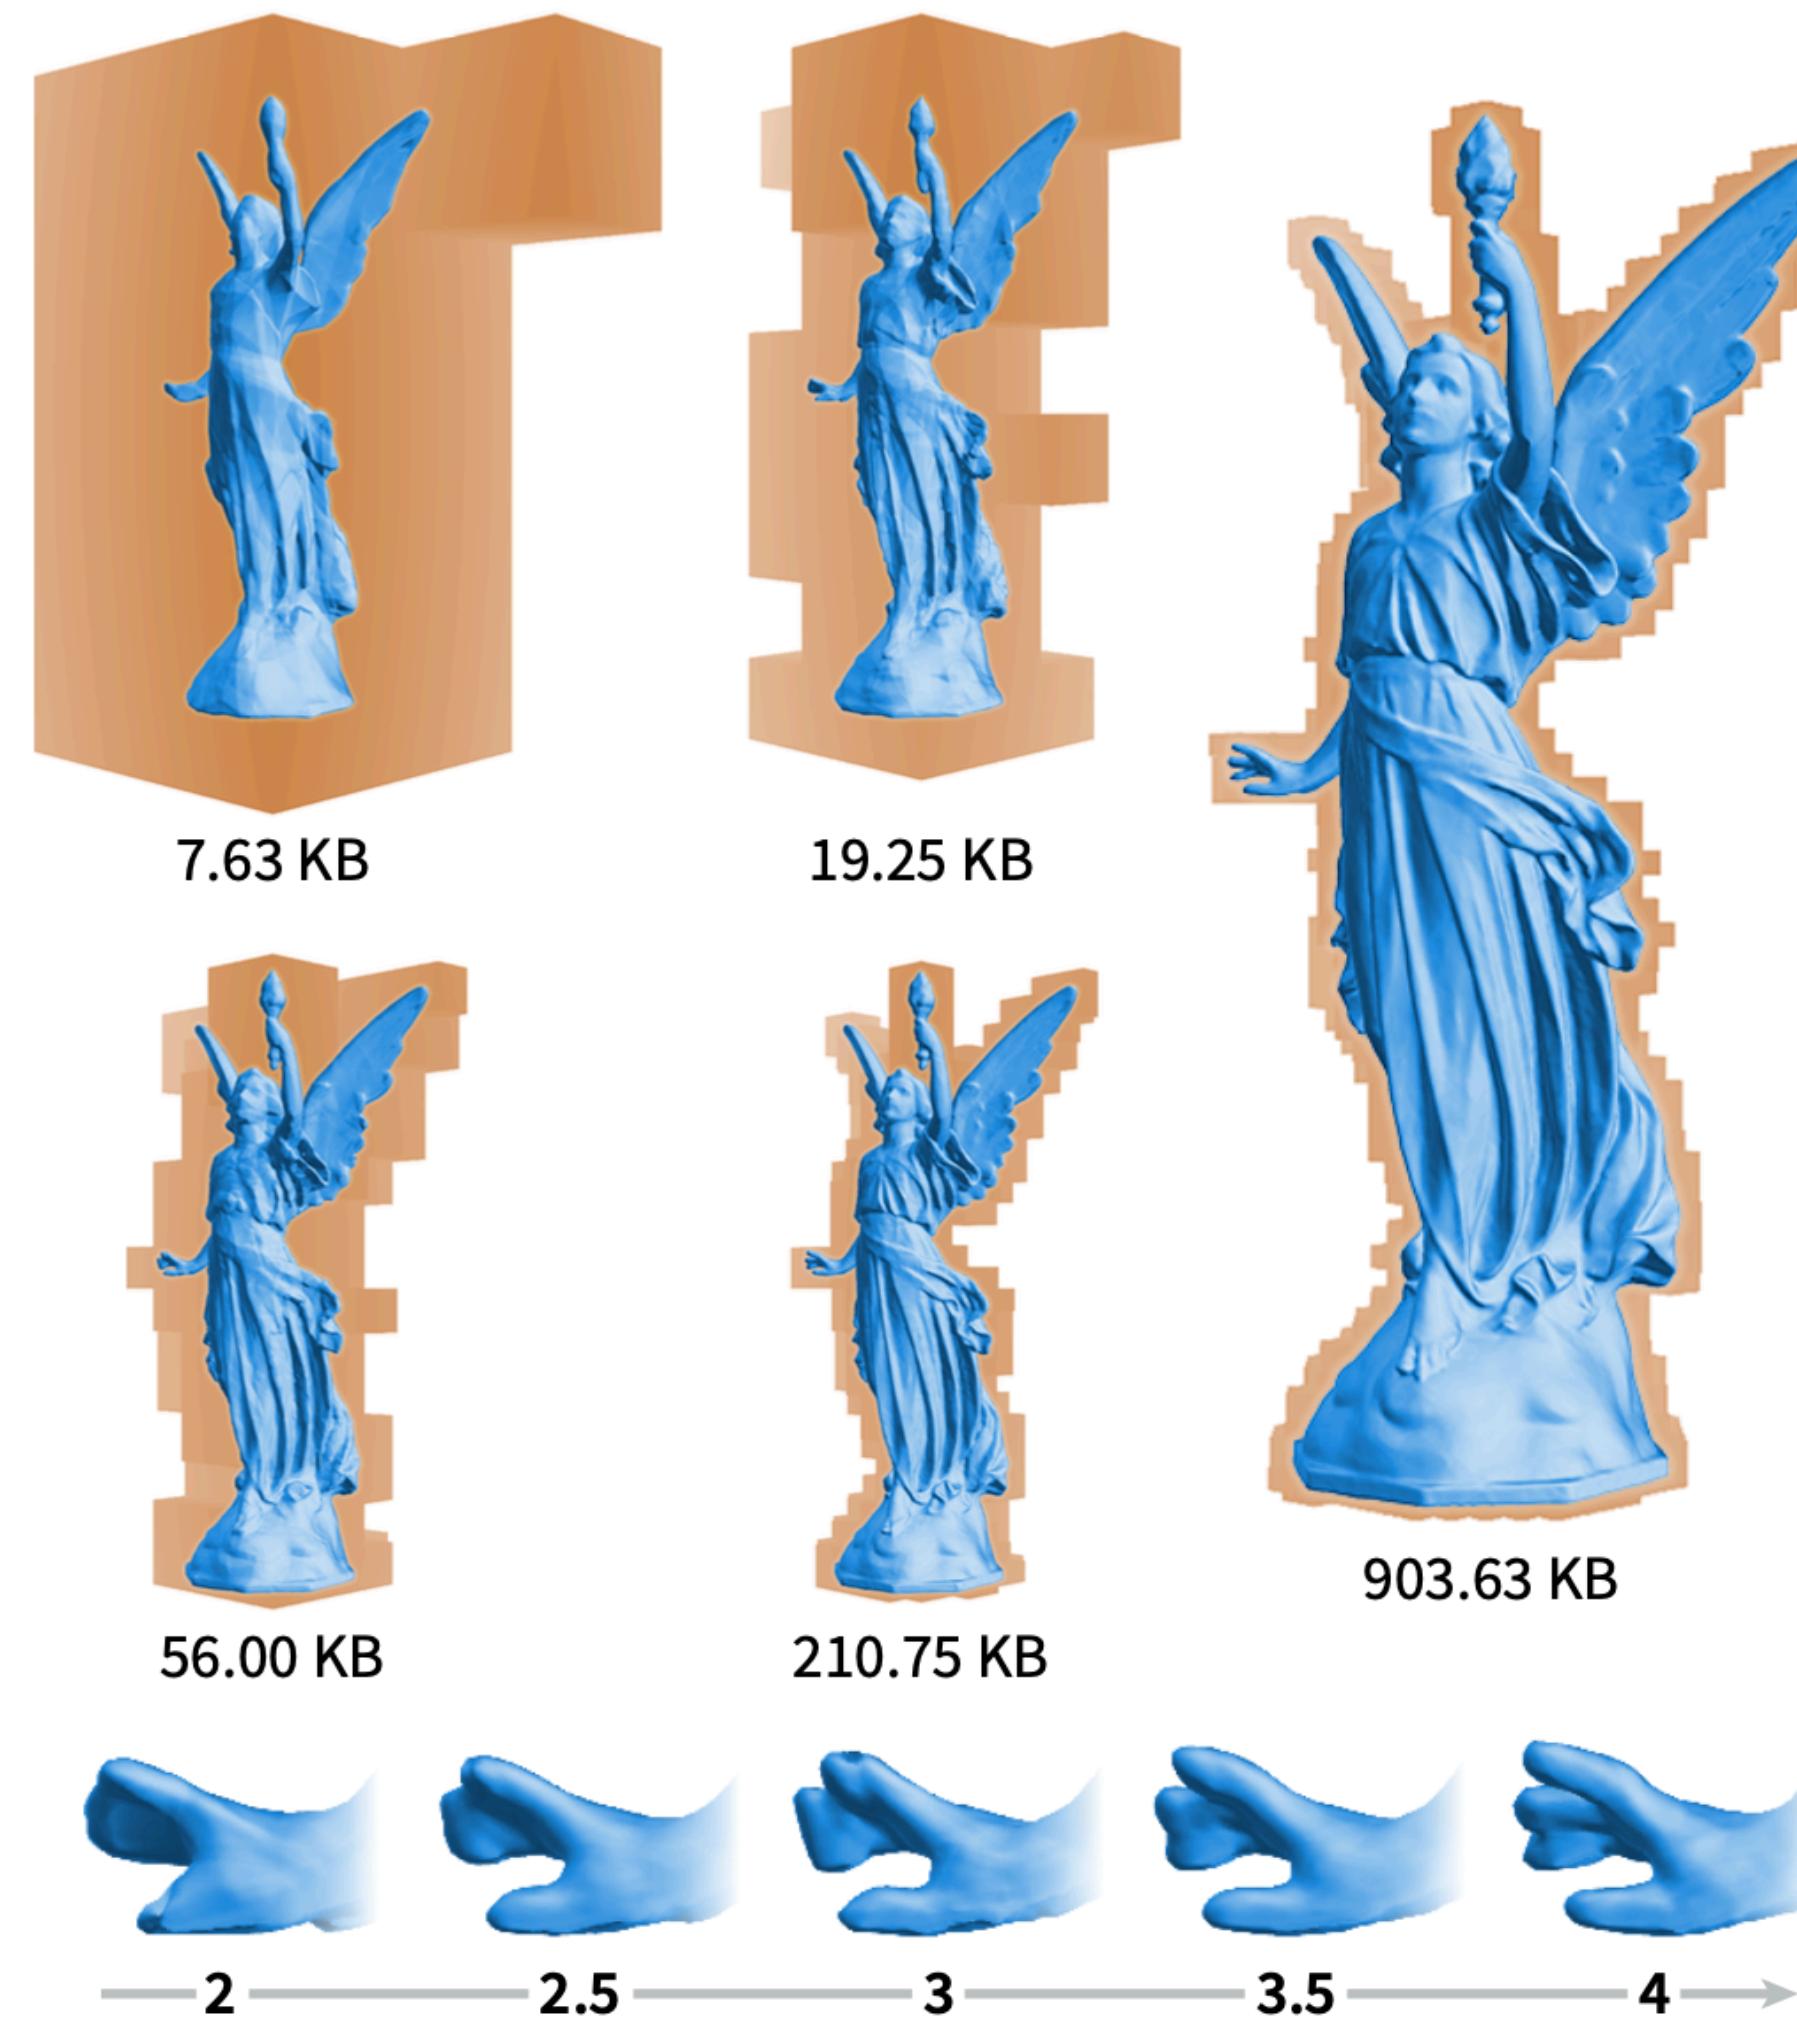
Image Grids & Perspective Projection



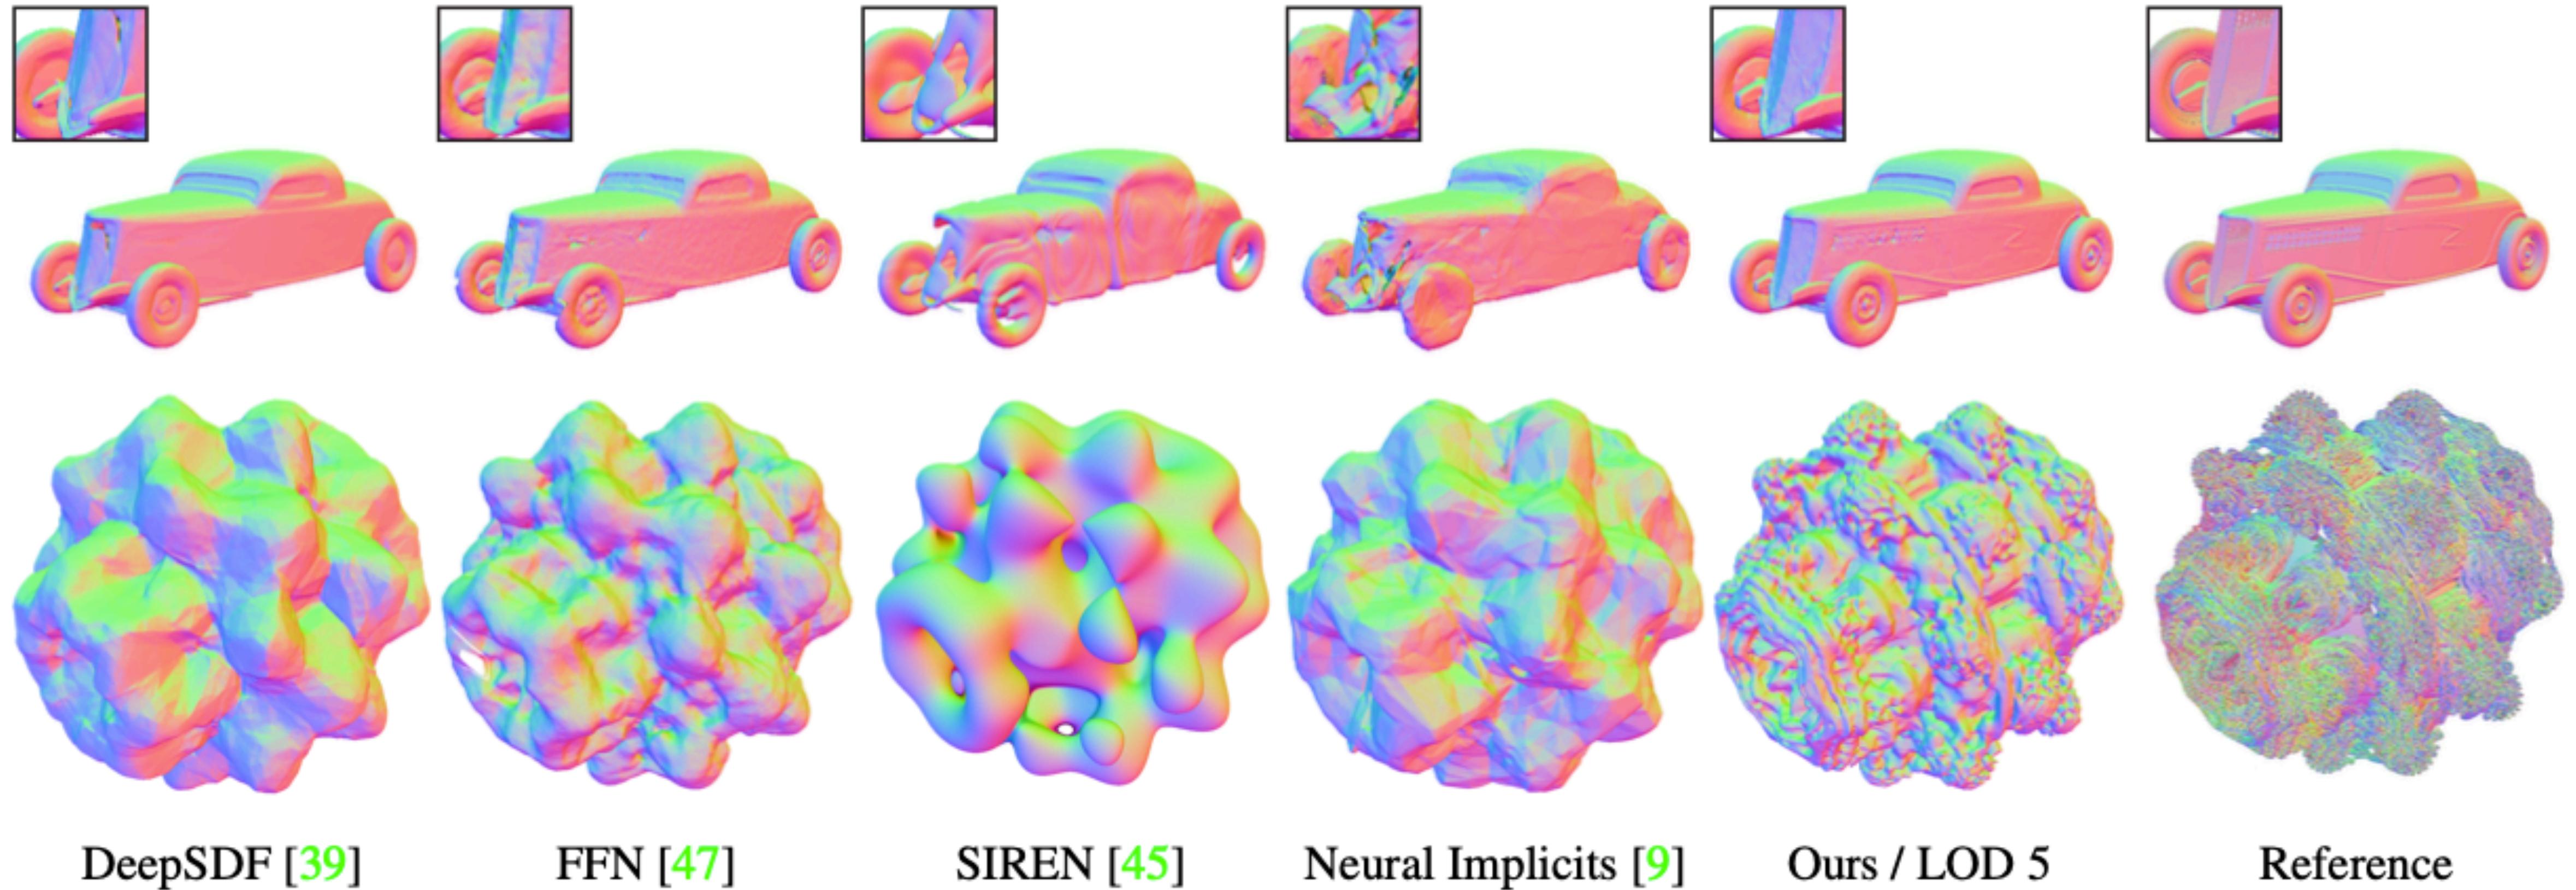
Multi-scale Voxel Grids



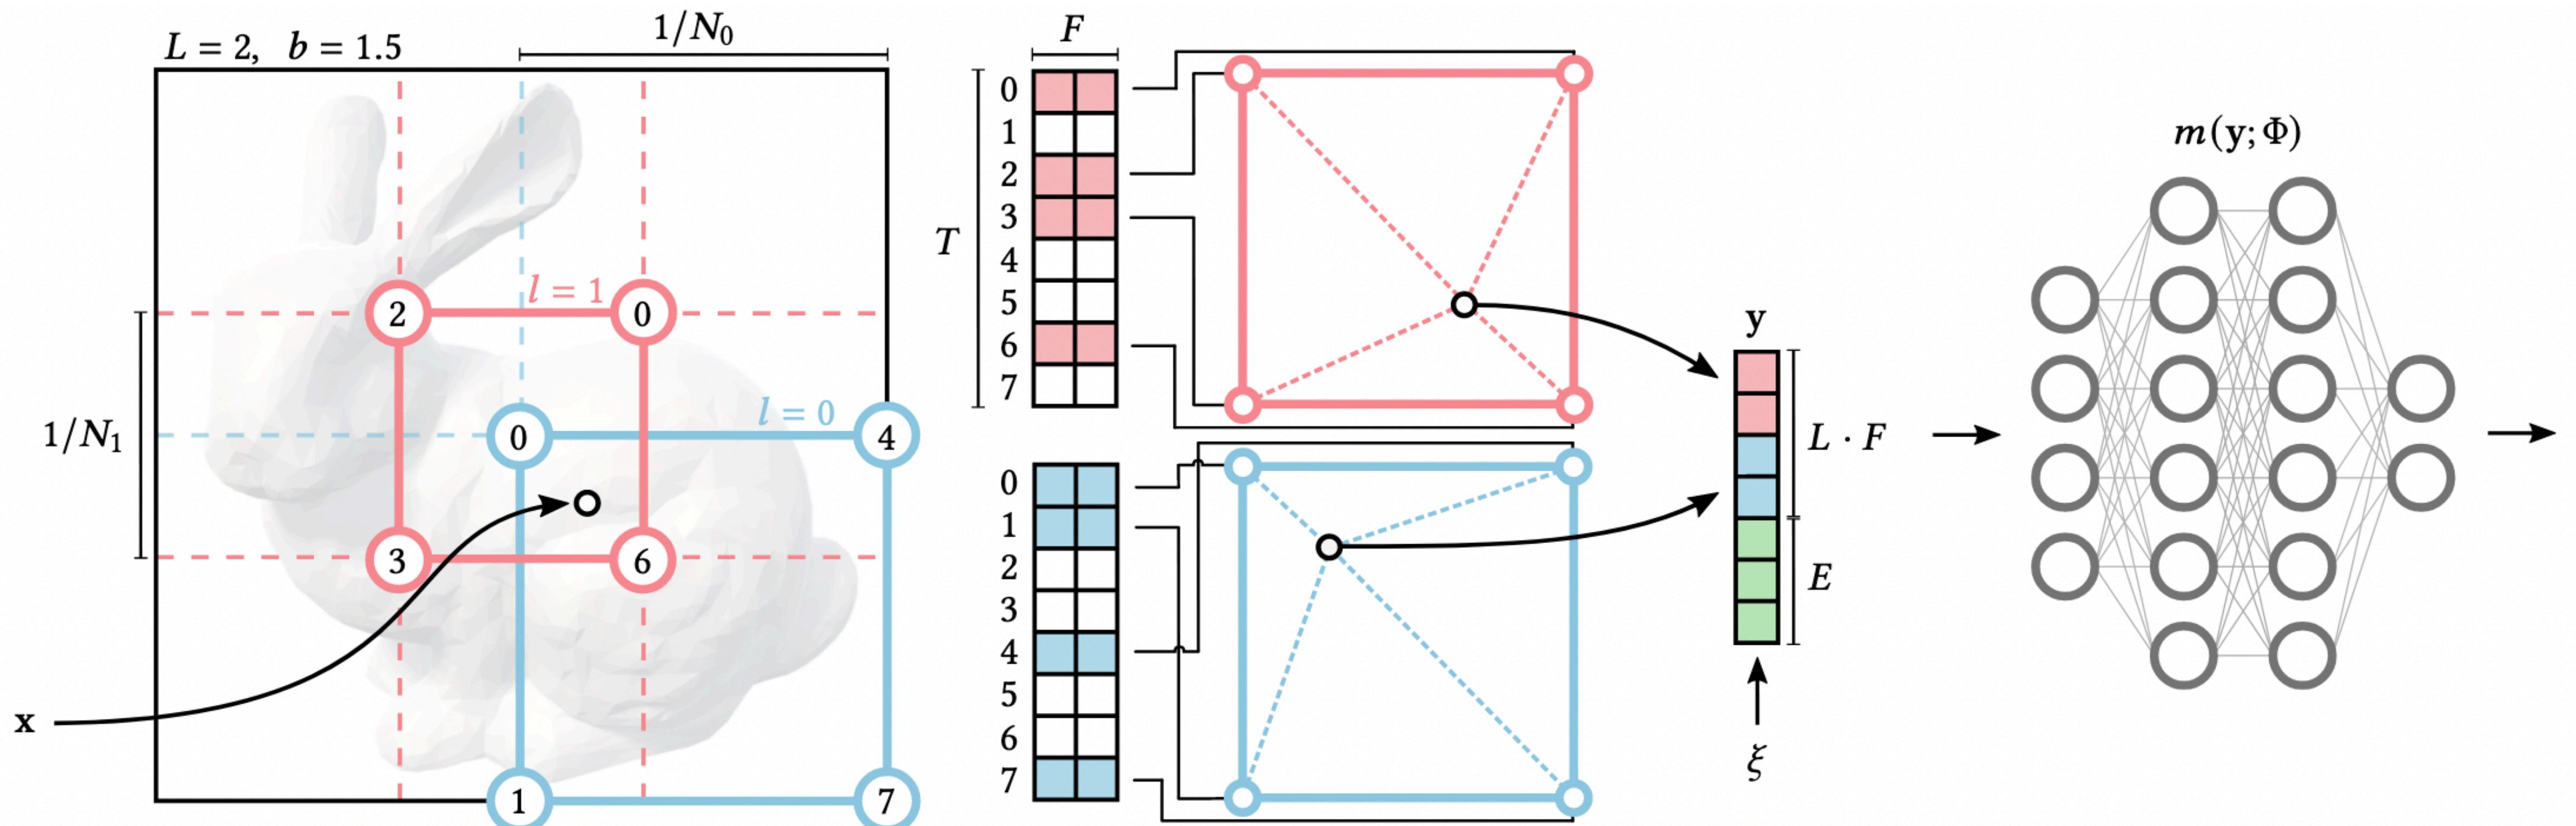
Multi-scale Voxel Grids



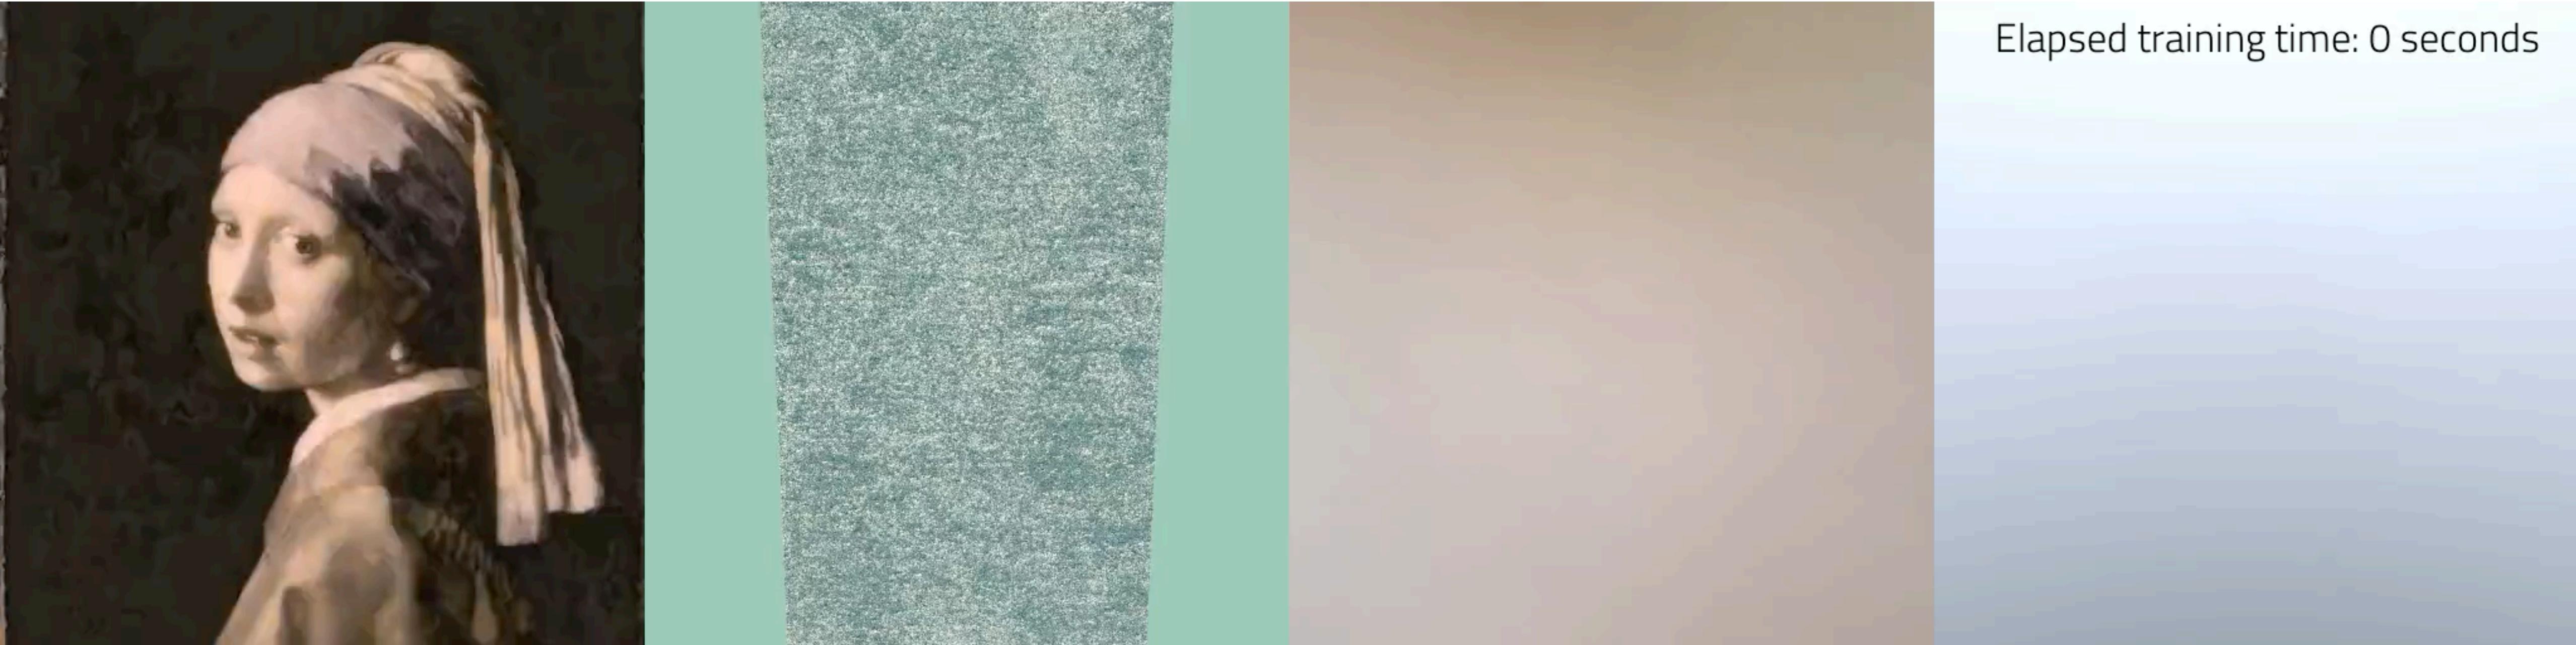
Multi-scale Voxel Grids



Multi-scale Voxel Grids + Hash Table



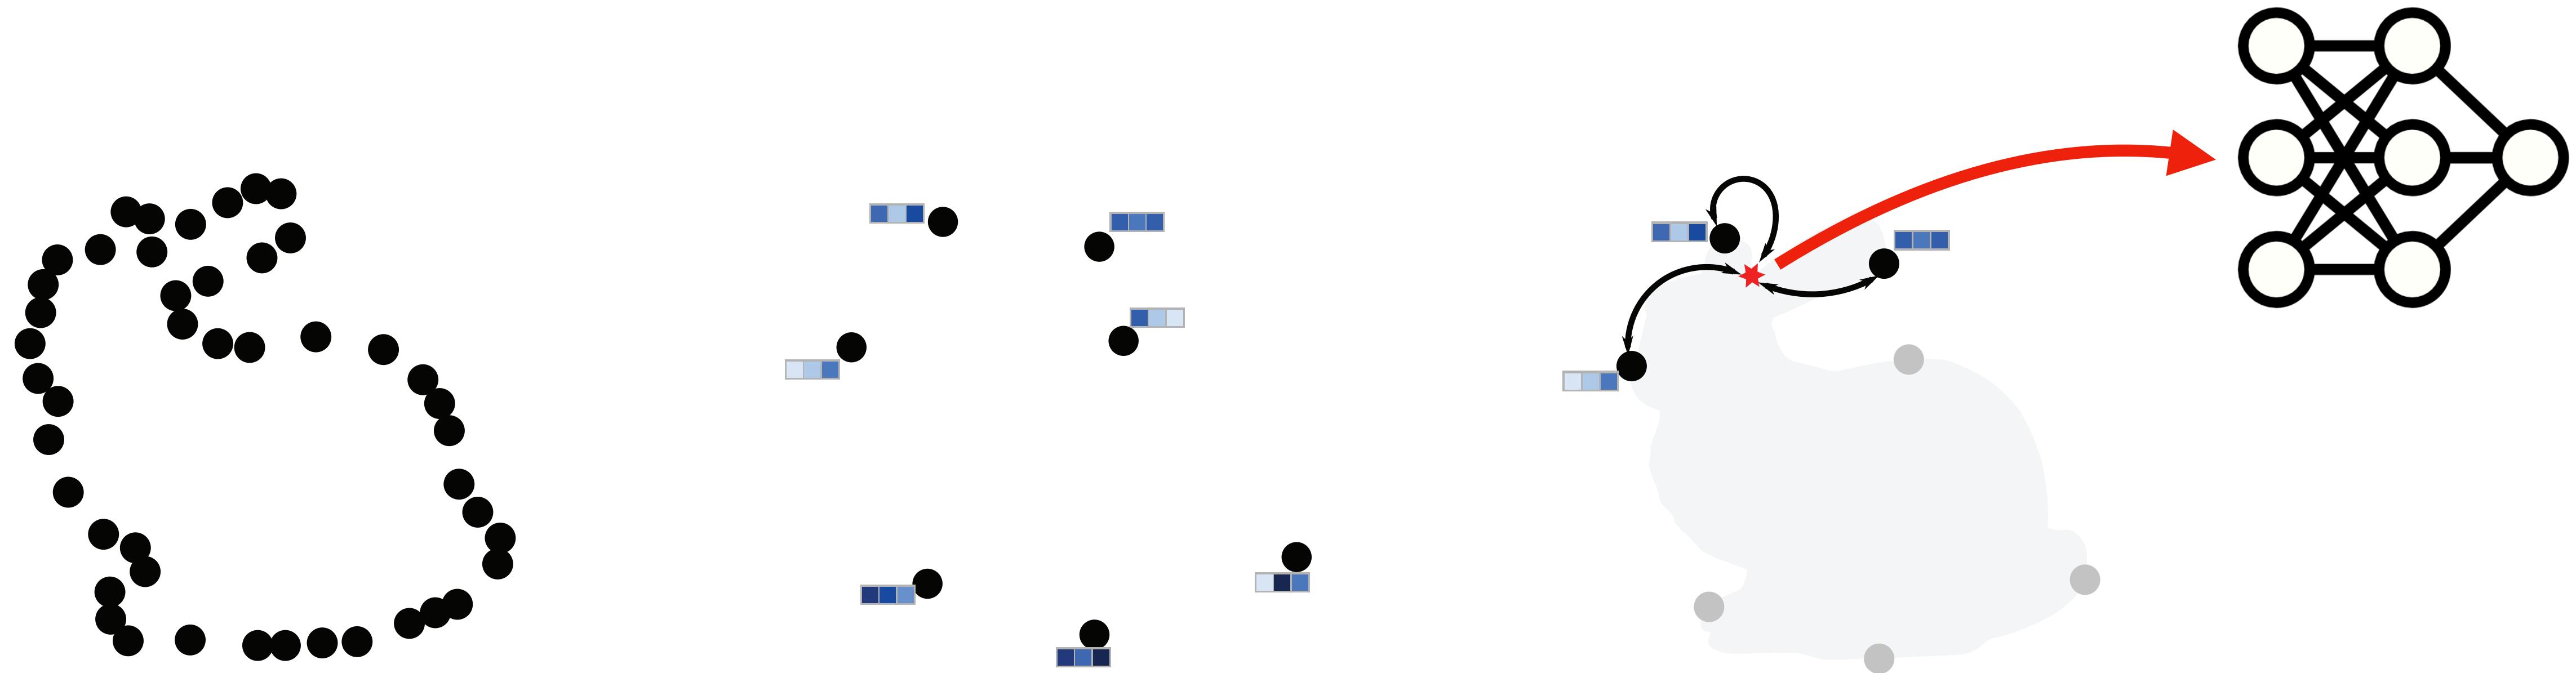
Multi-scale Voxel Grids + Hash Table



Multi-scale Voxel Grids + Hash Table

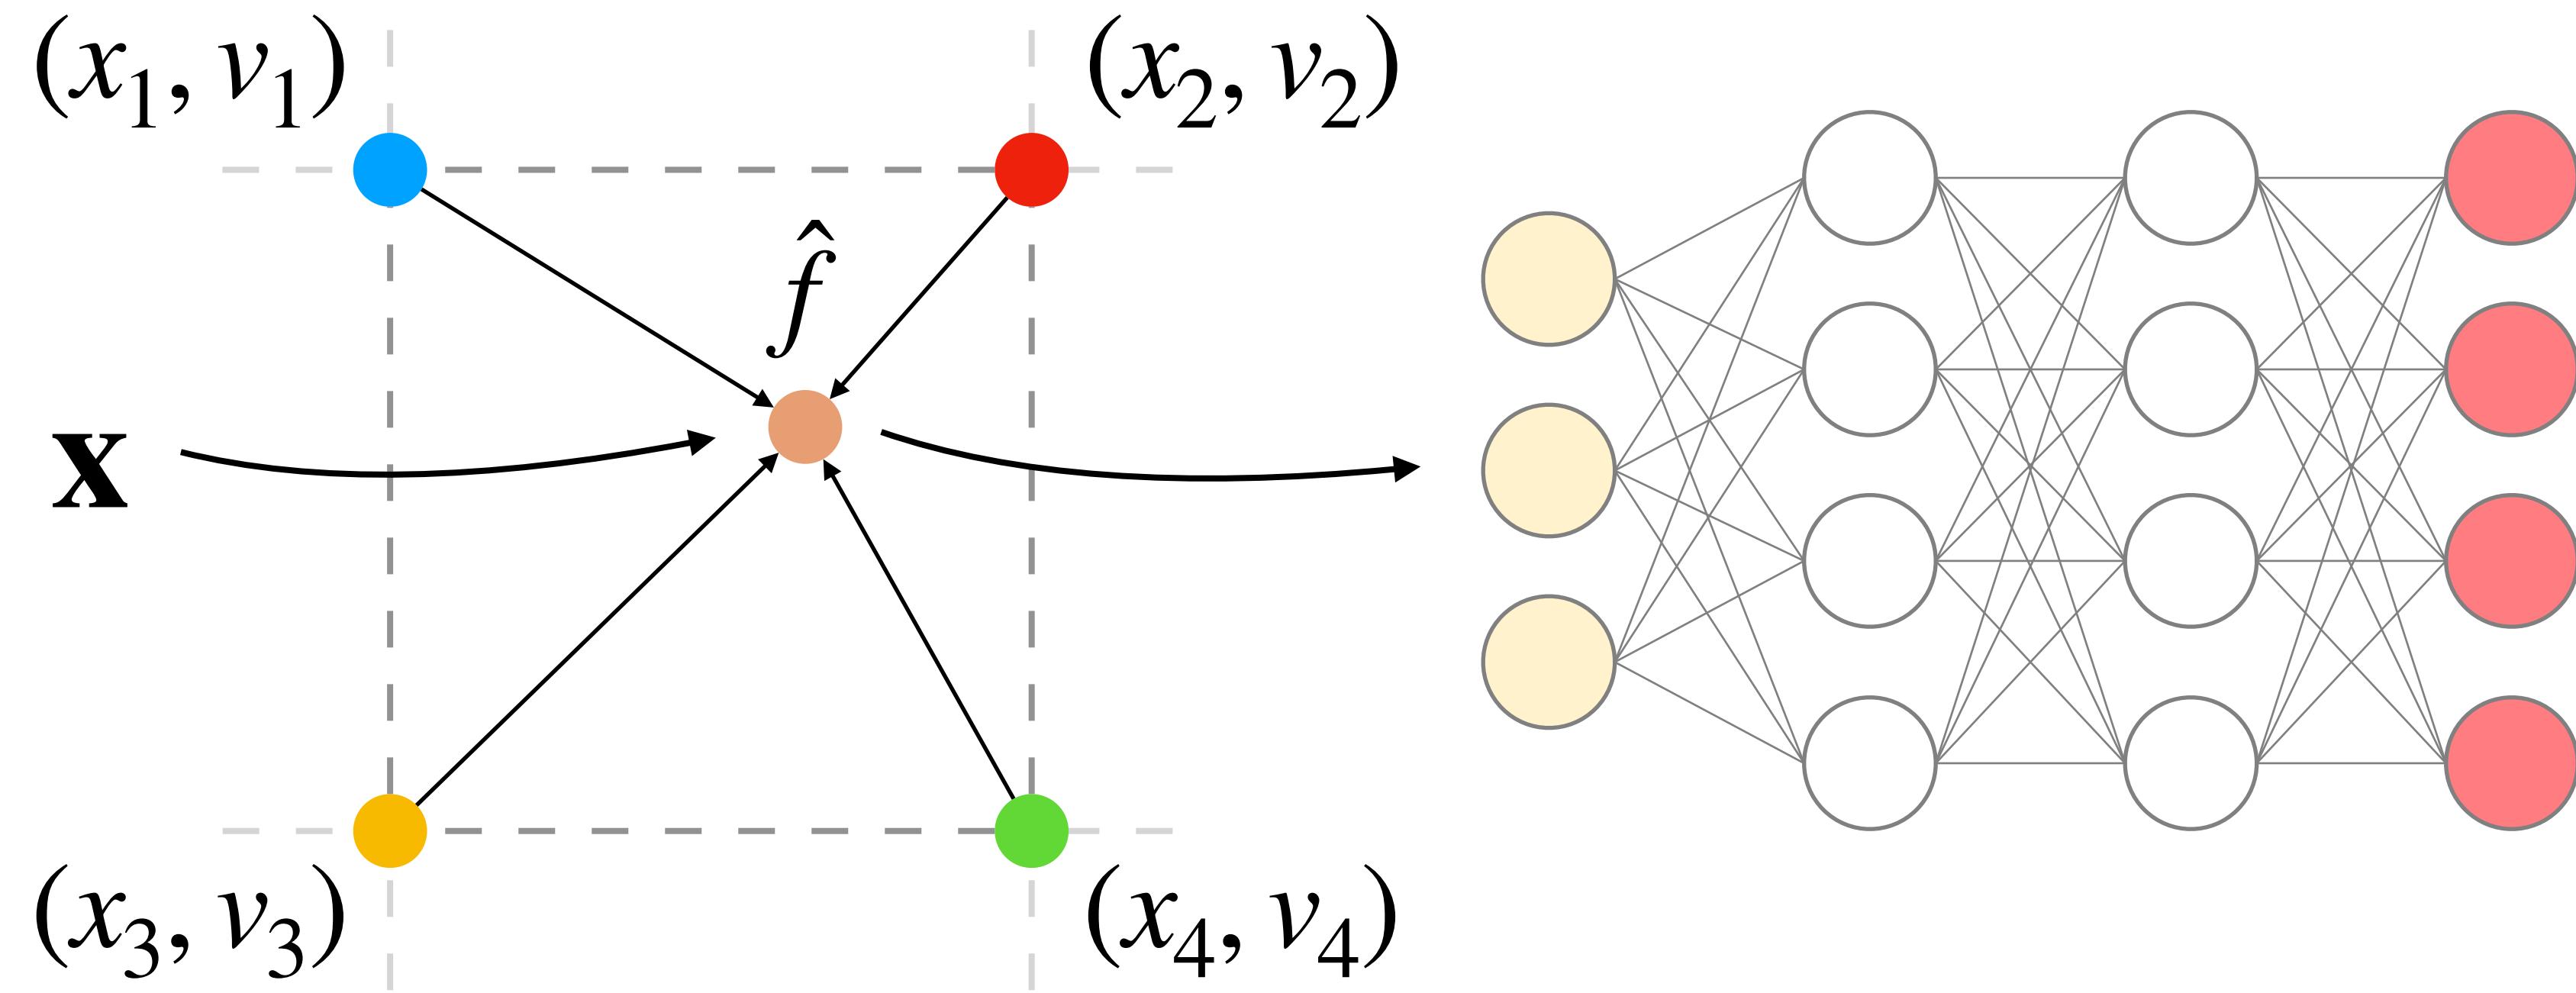


Point Cloud & Nearest Neighbor

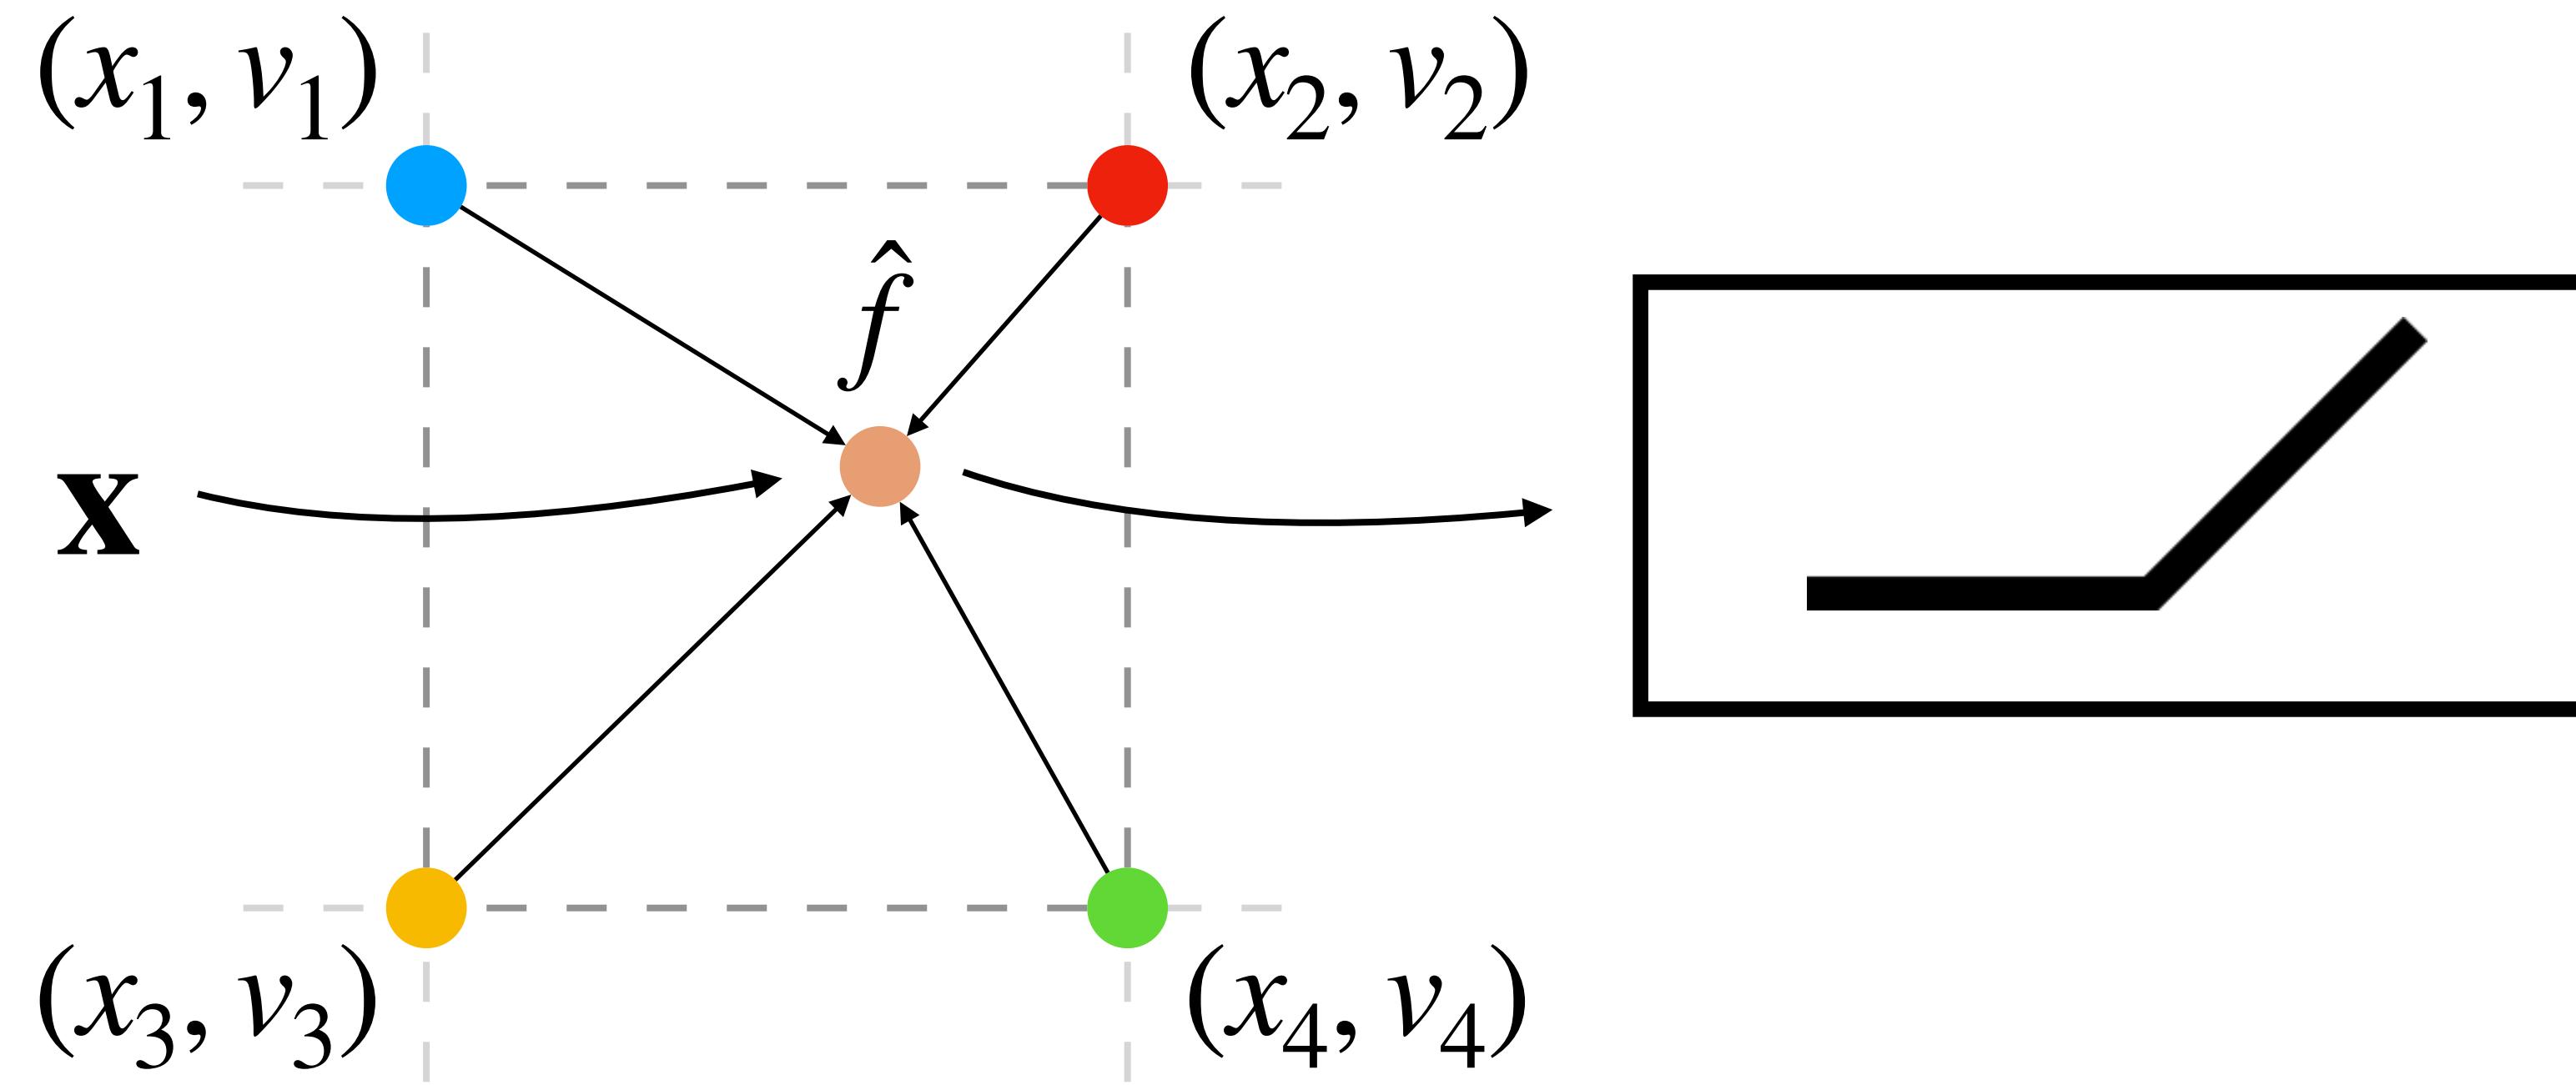


Hybrid representations can be used to decode a surface representation into a volume representation!

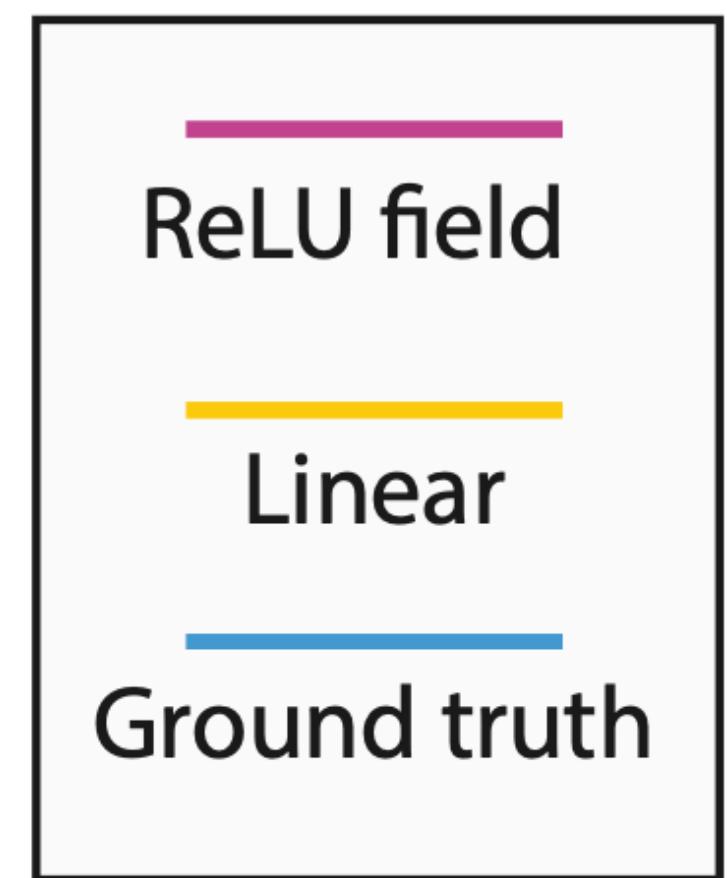
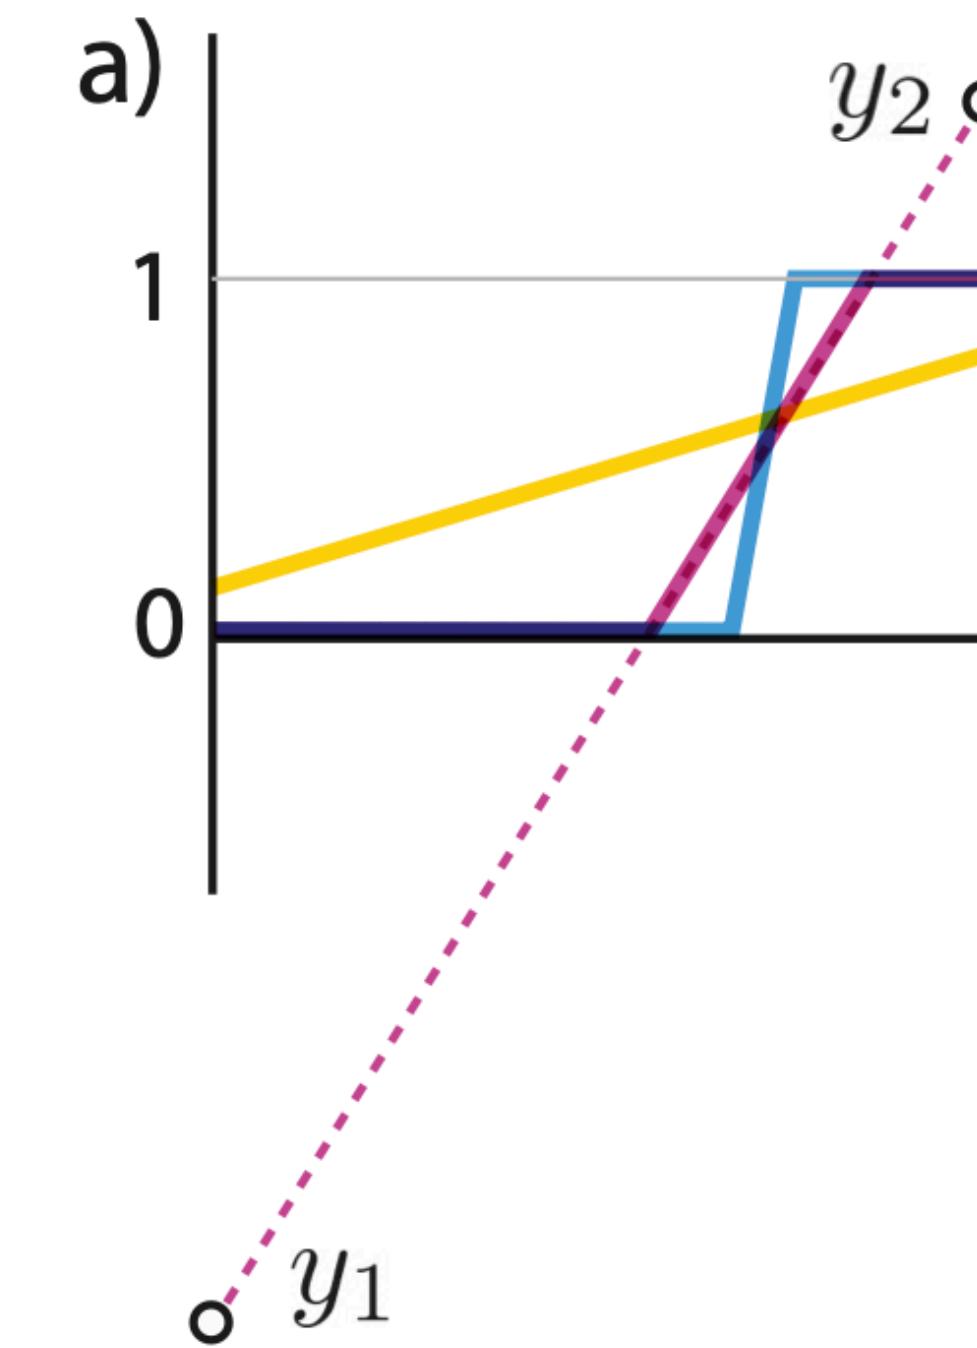
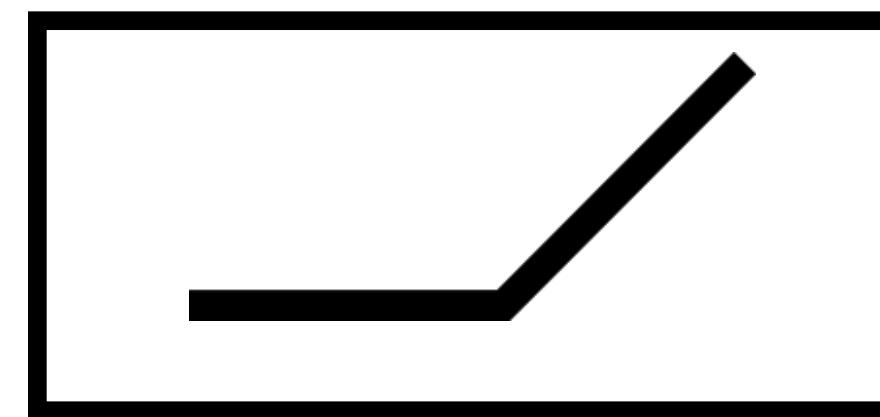
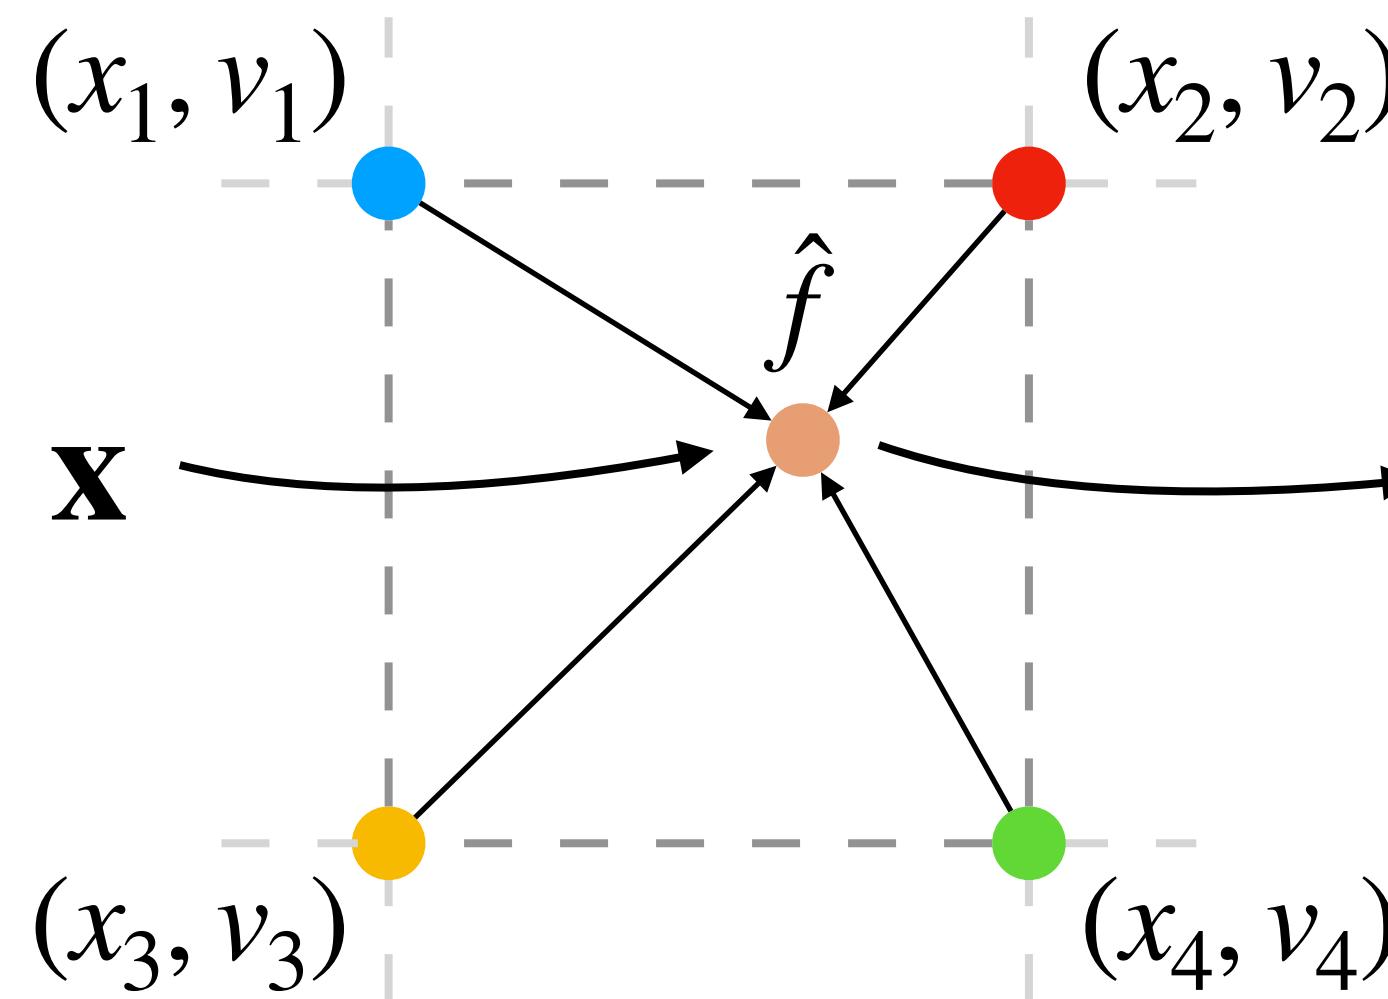
ReLU Fields: How much MLP do you need?



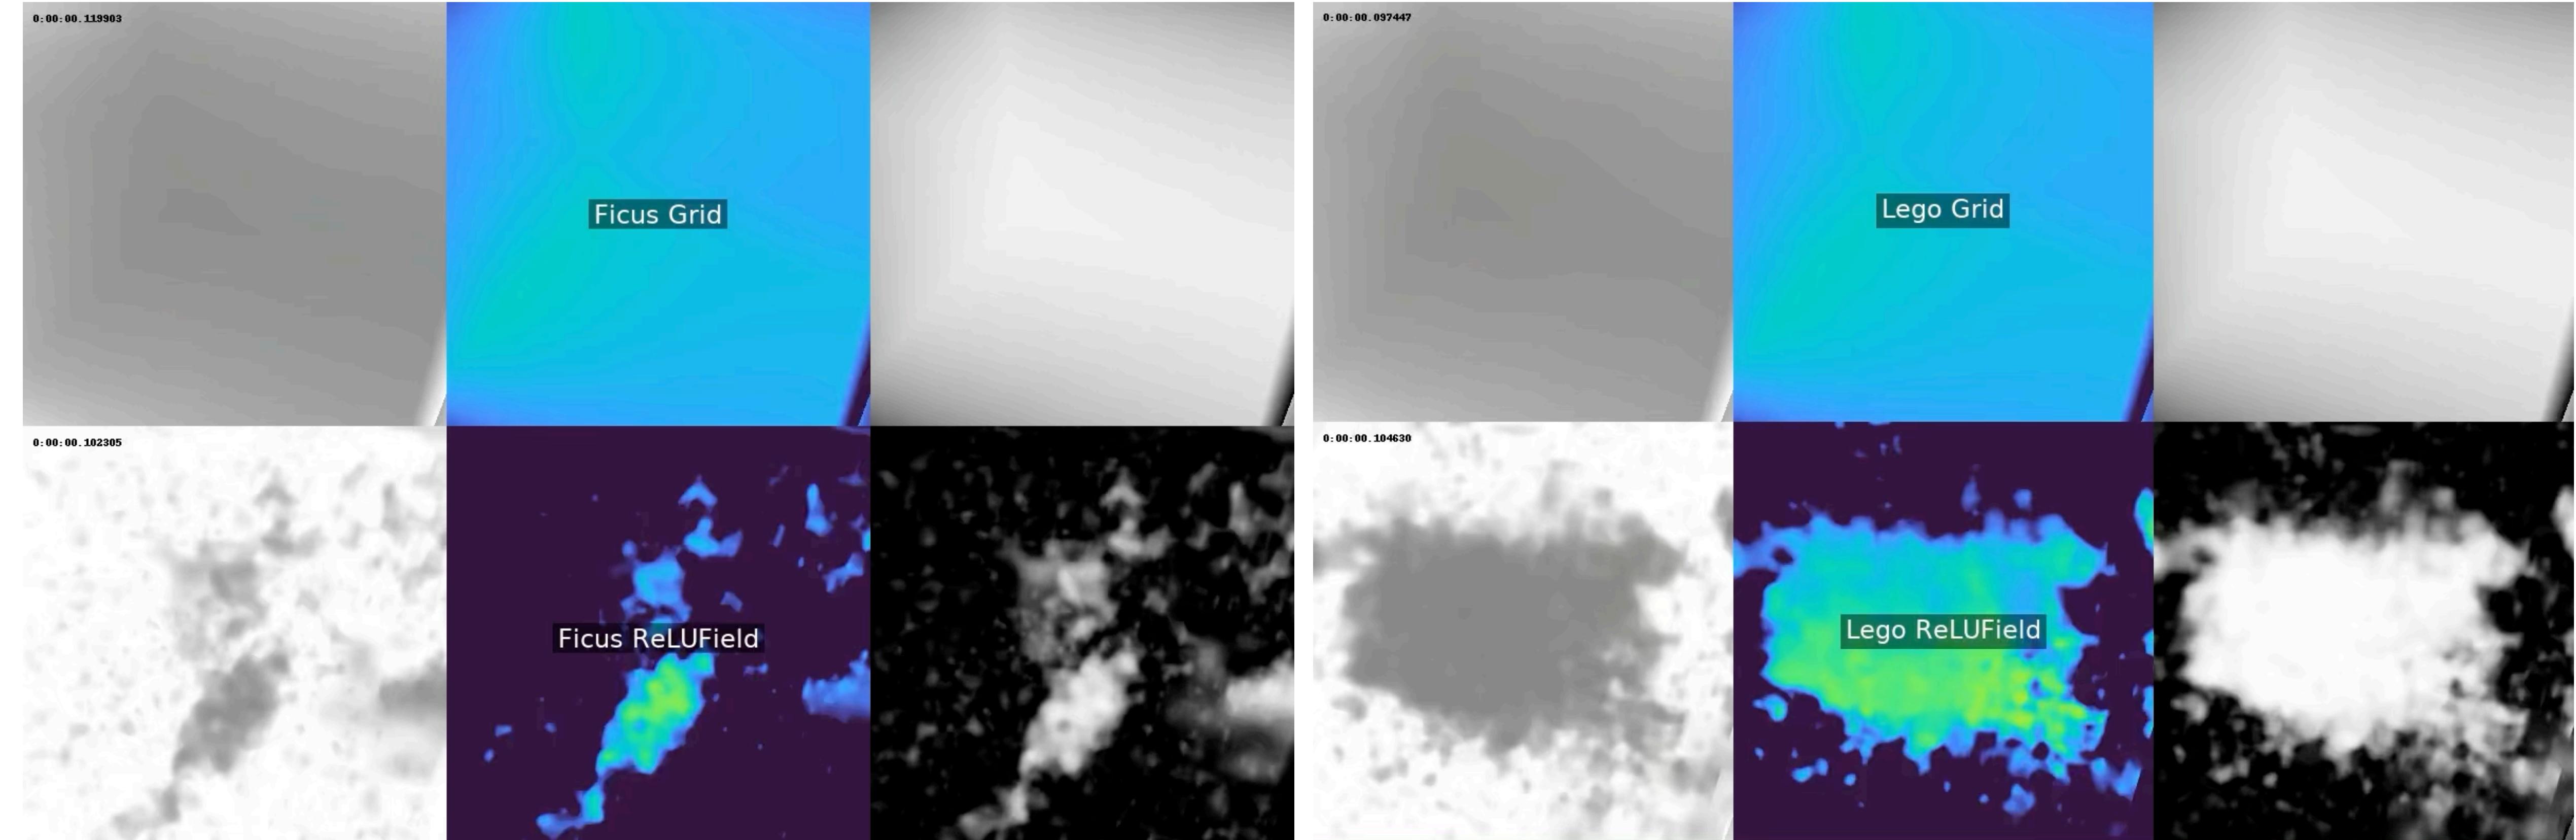
ReLU Fields: How much MLP do you need?



ReLU Fields: How much MLP do you need?

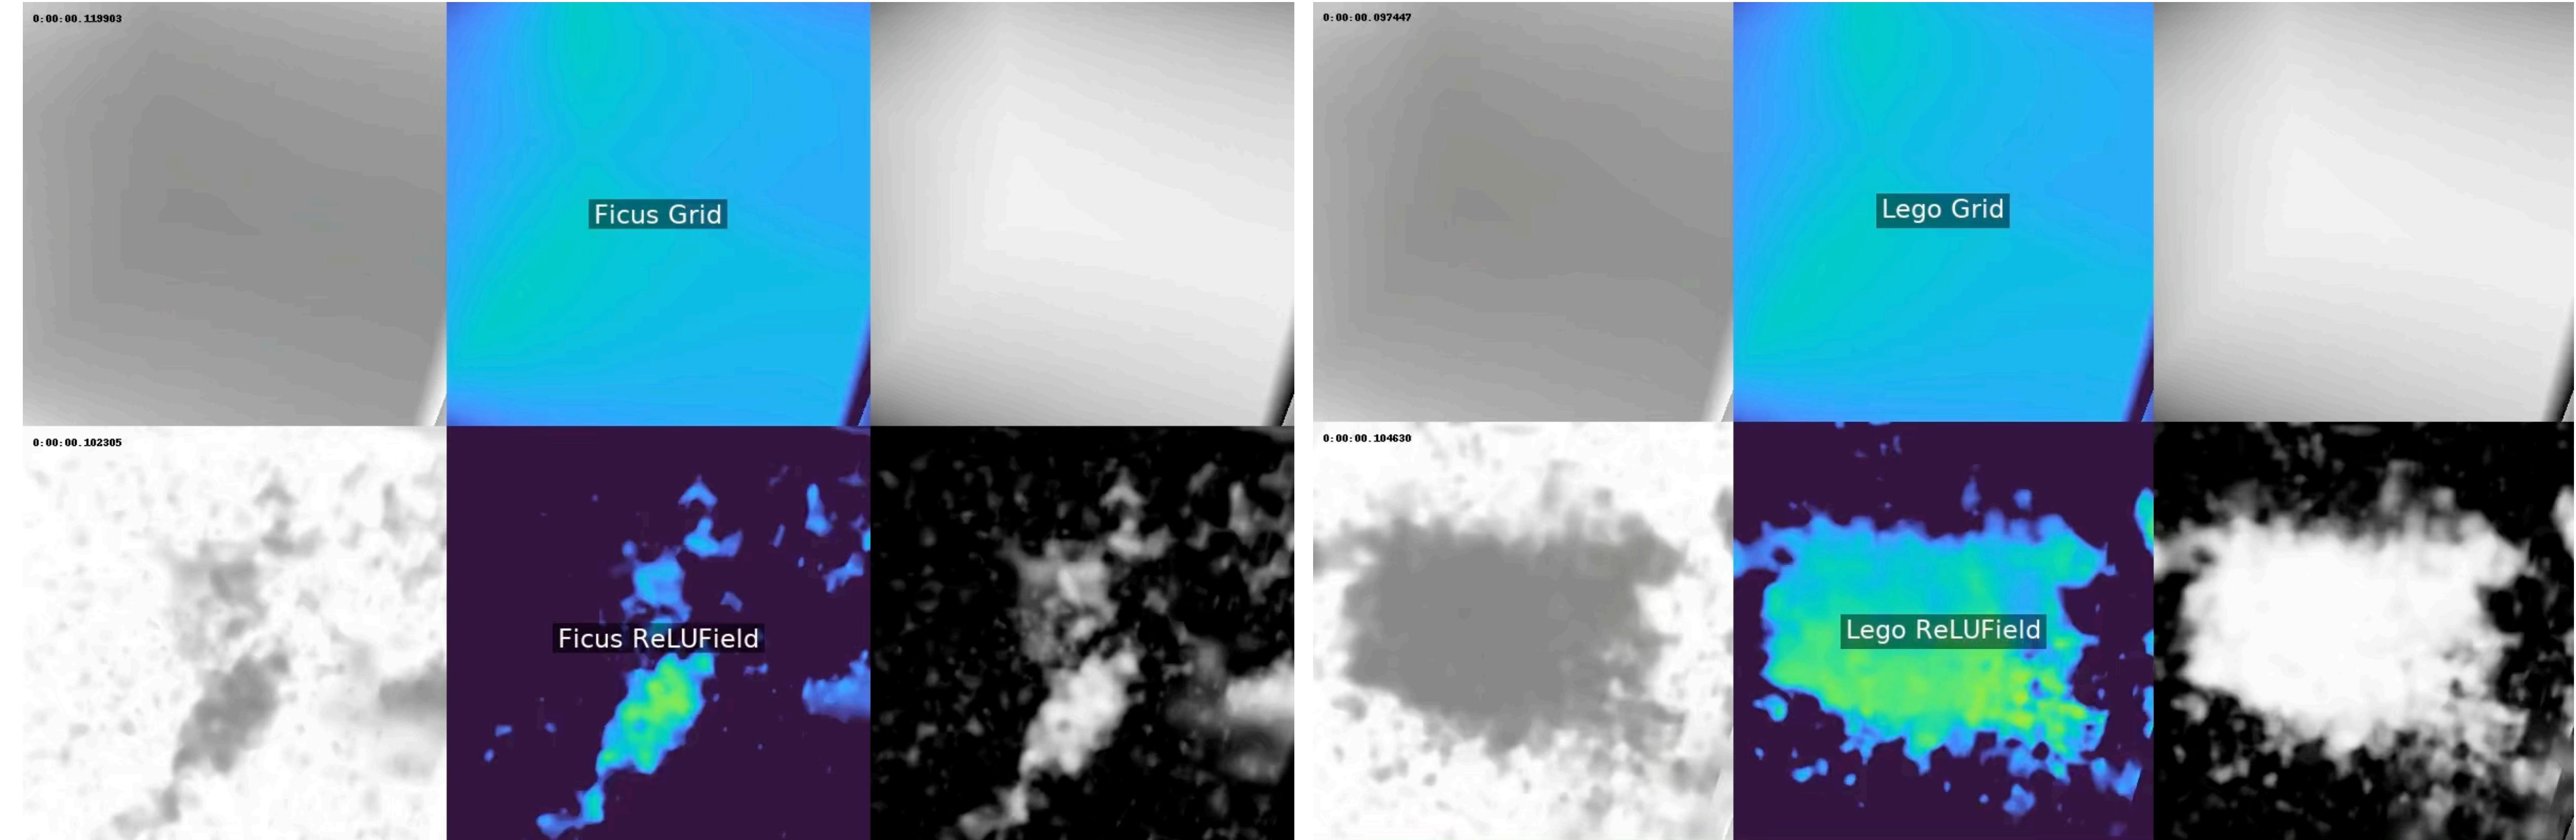


ReLU Fields: How much MLP do you need?



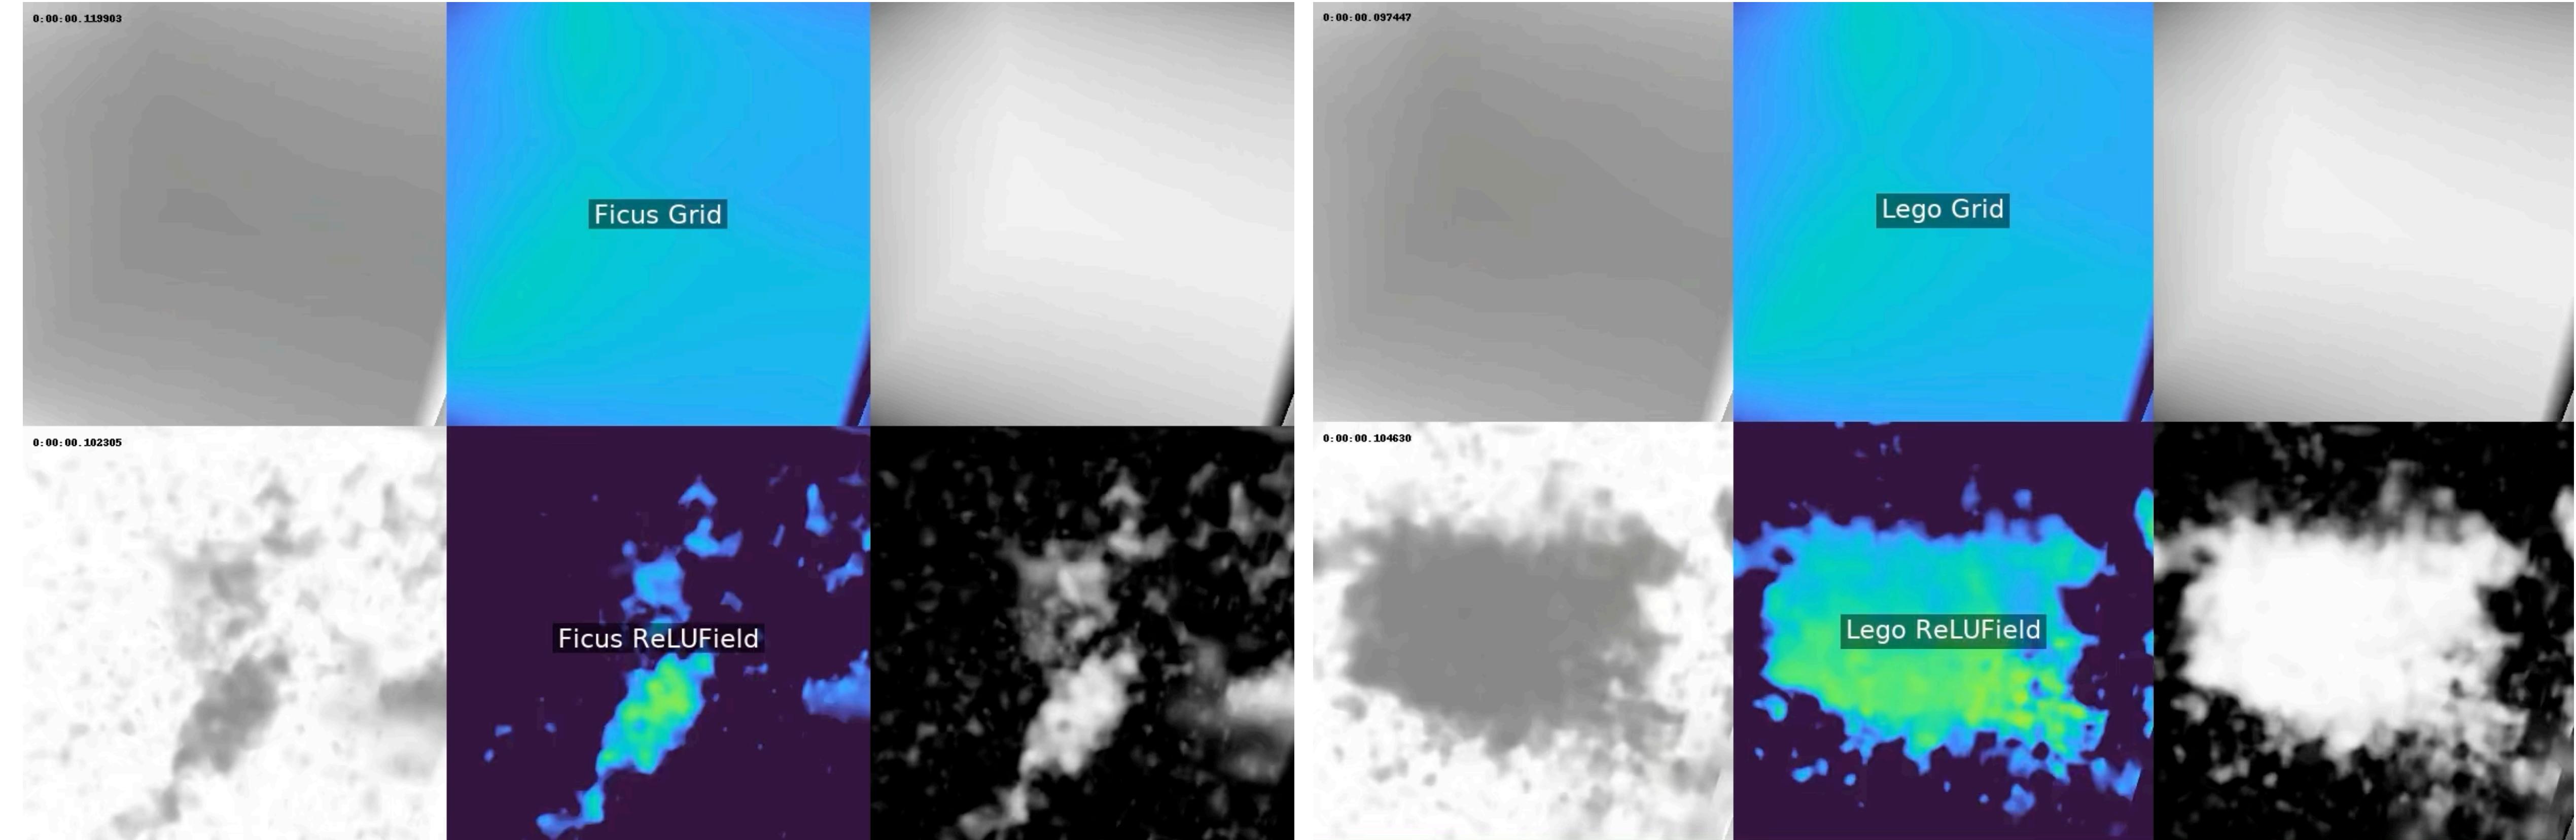
ReLU Fields: The Little Non-linearity That Could (Karnewar et al. 2022)

ReLU Fields: How much MLP do you need?



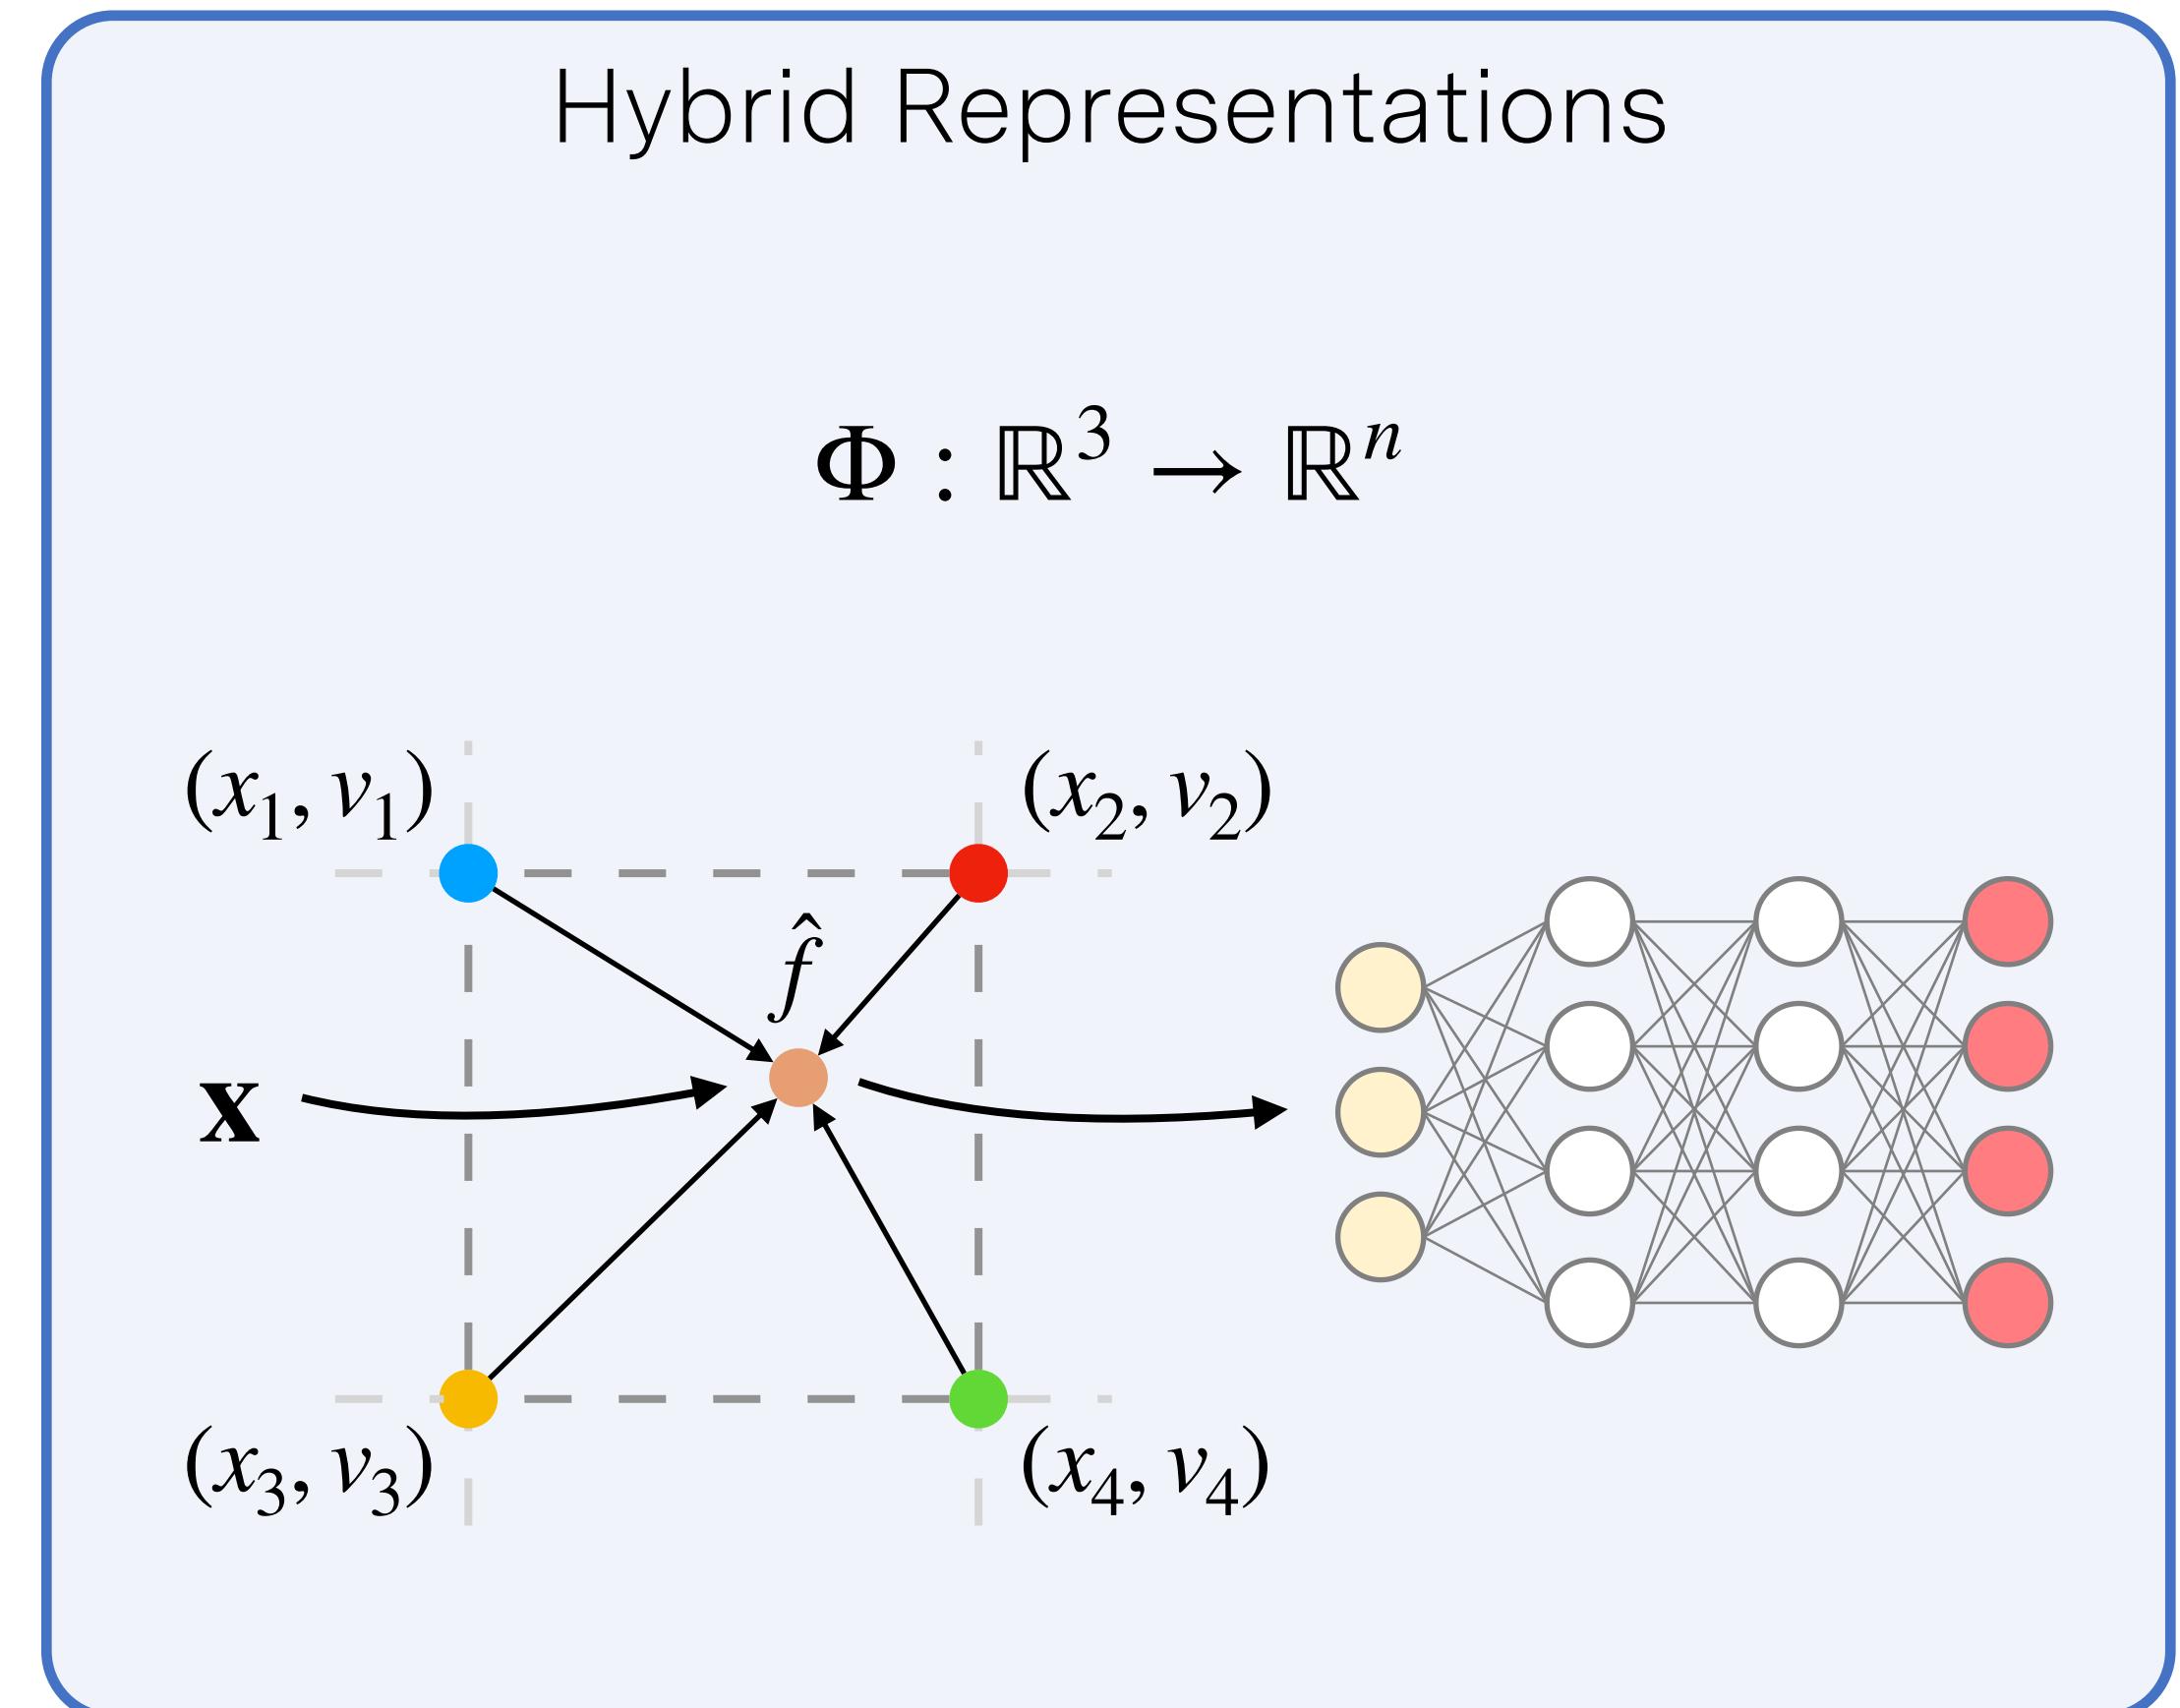
ReLU Fields: The Little Non-linearity That Could (Karnewar et al. 2022)

ReLU Fields: How much MLP do you need?

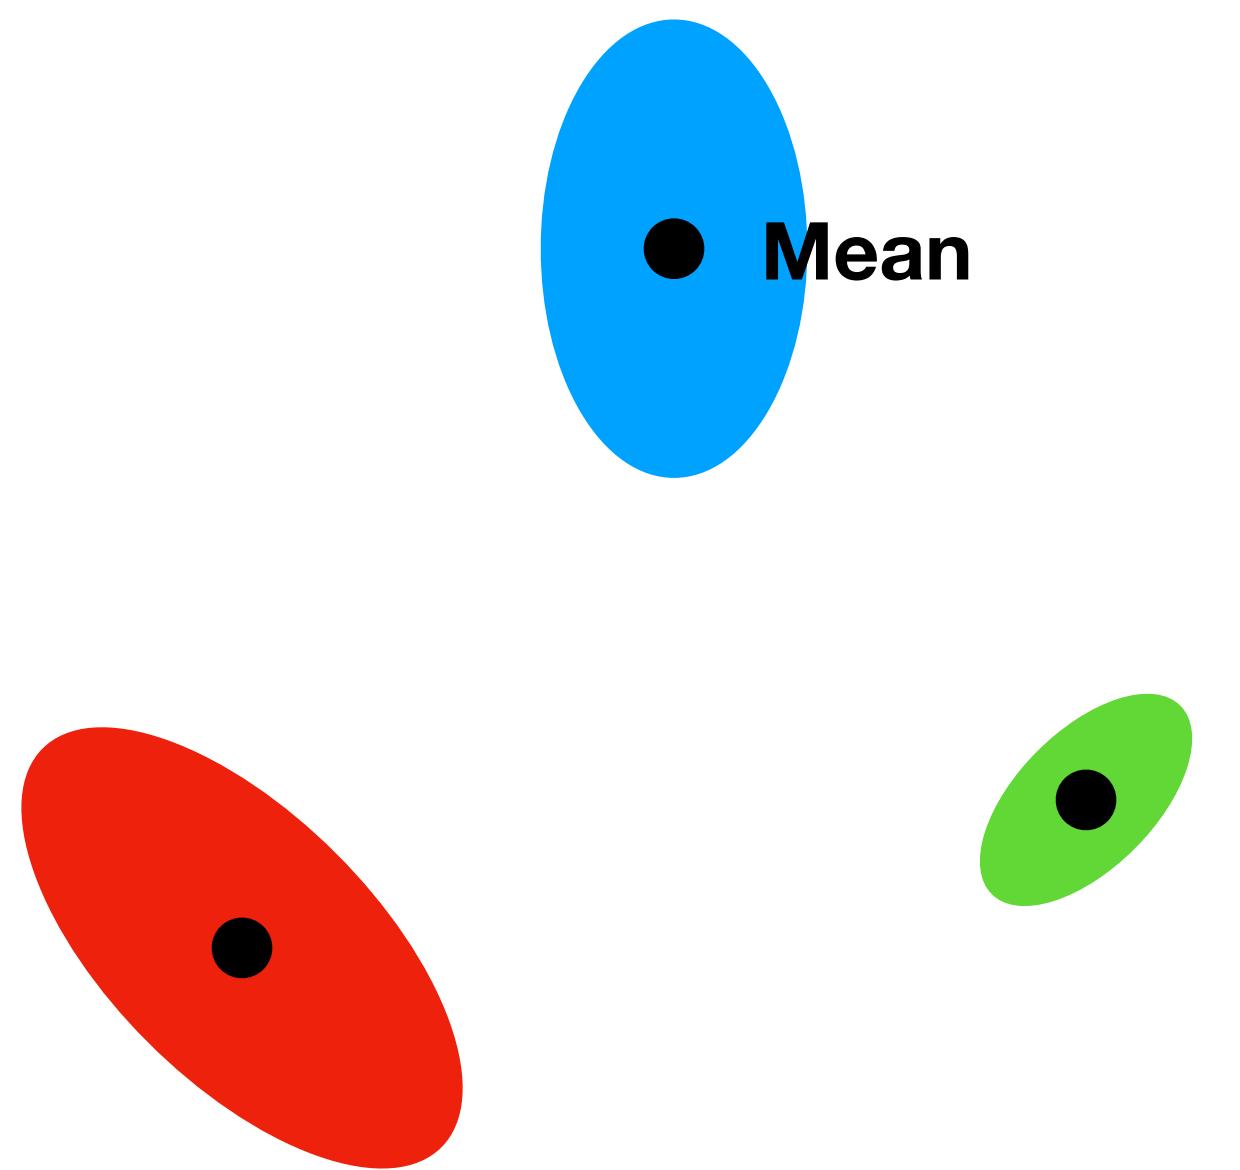


Hybrid Representations

- Very natural to trade off memory with compute.
- Neural network can locally “super resolve” discrete representation: Basically acts as interpolation kernel.
- Can be the best of both worlds: Fast and “continuous”.

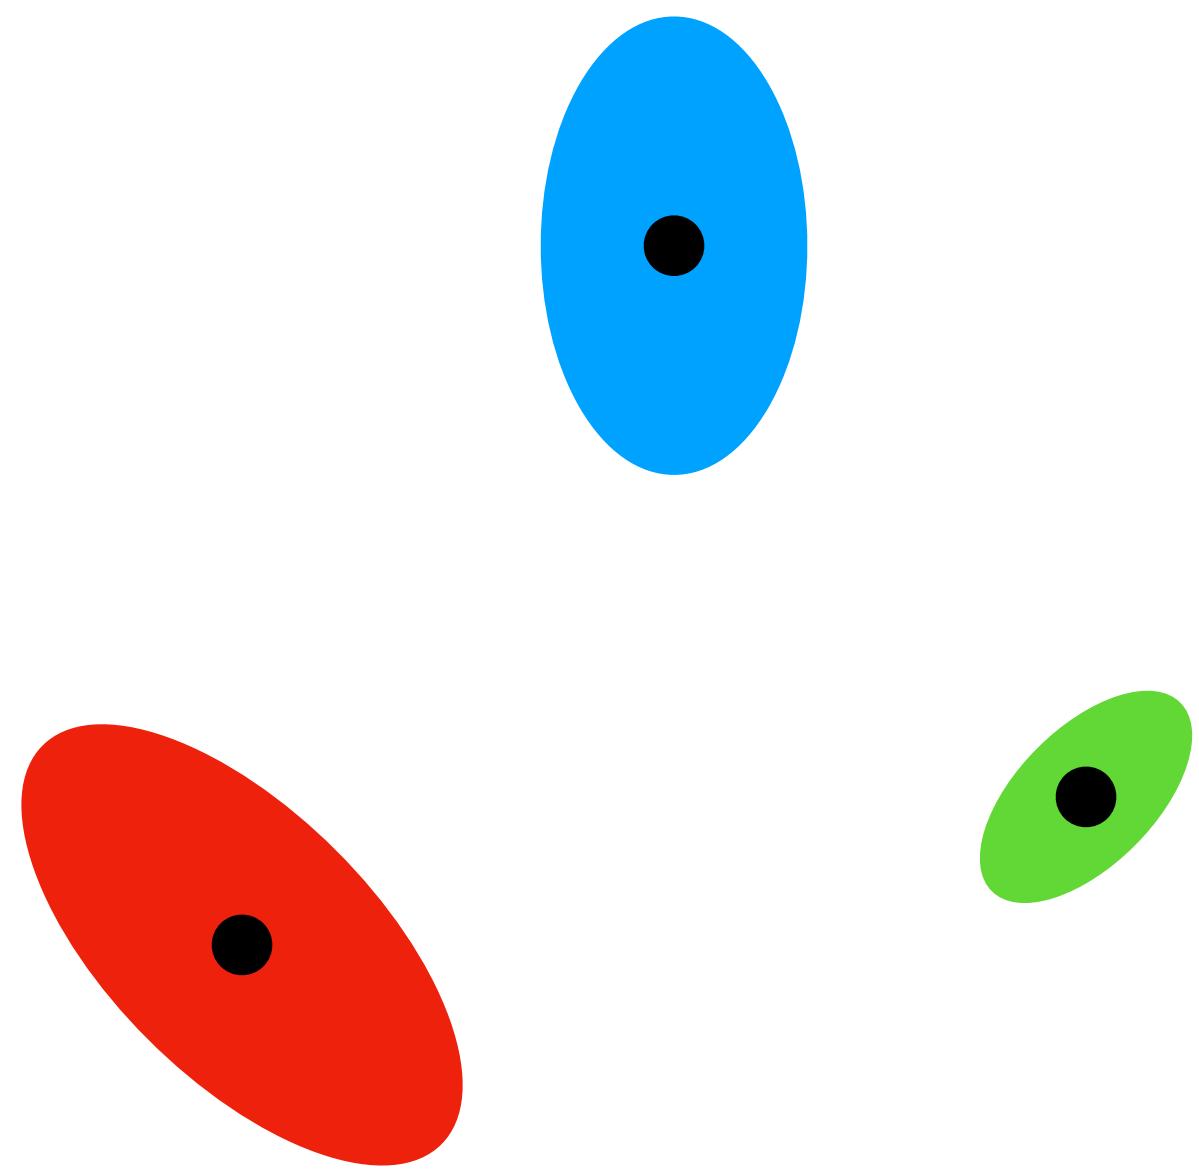


Primitive-based representations



**3D Gaussian Blobs
floating in Space**

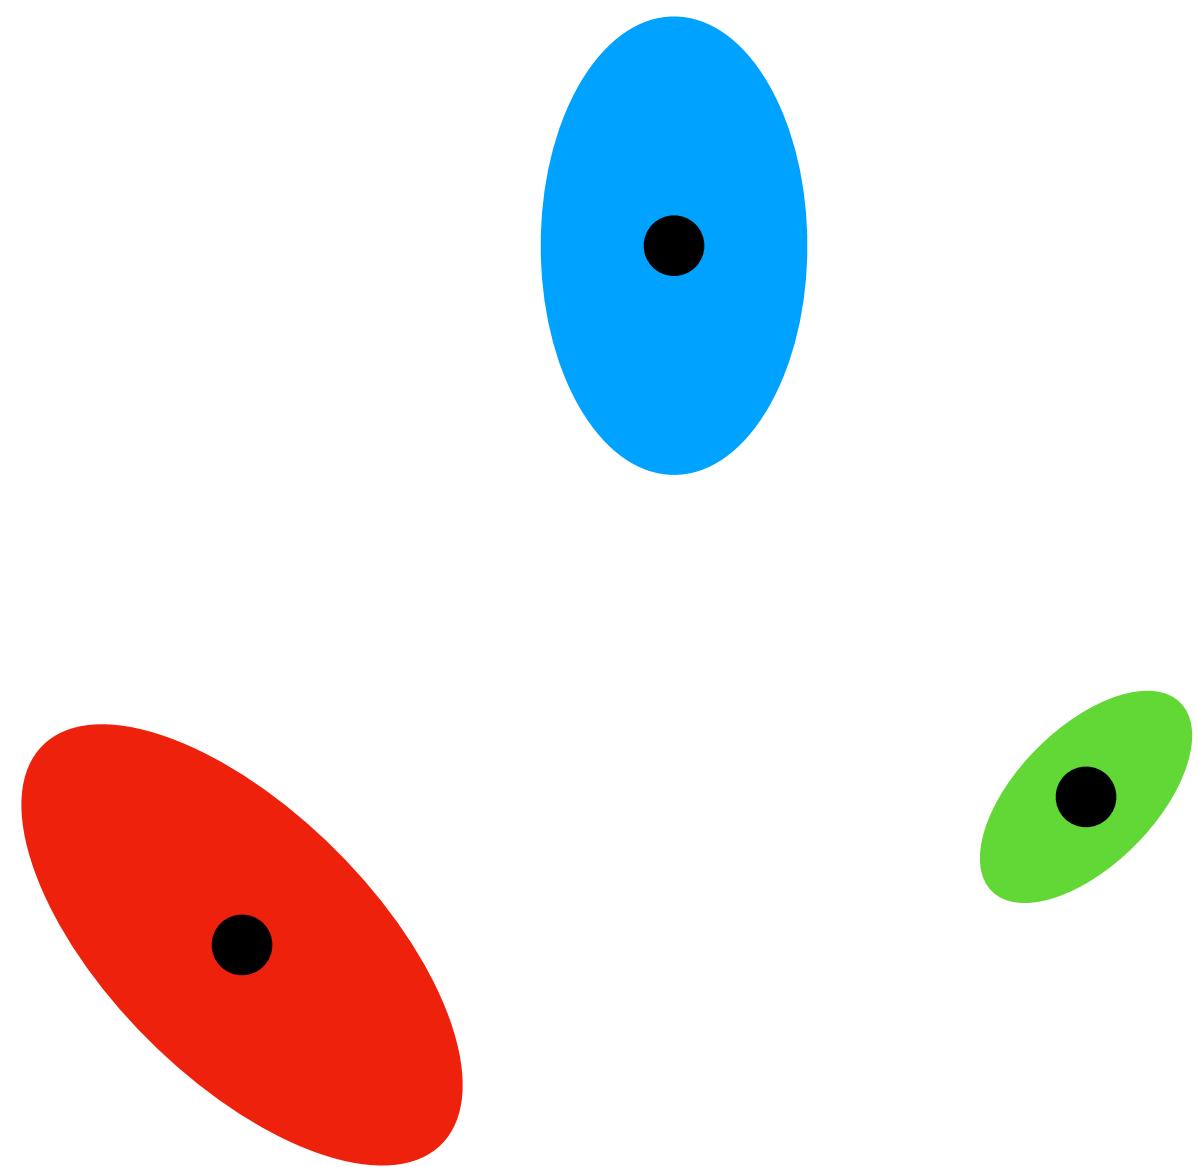
Primitive-based representations



**3D Gaussian Blobs
floating in Space**

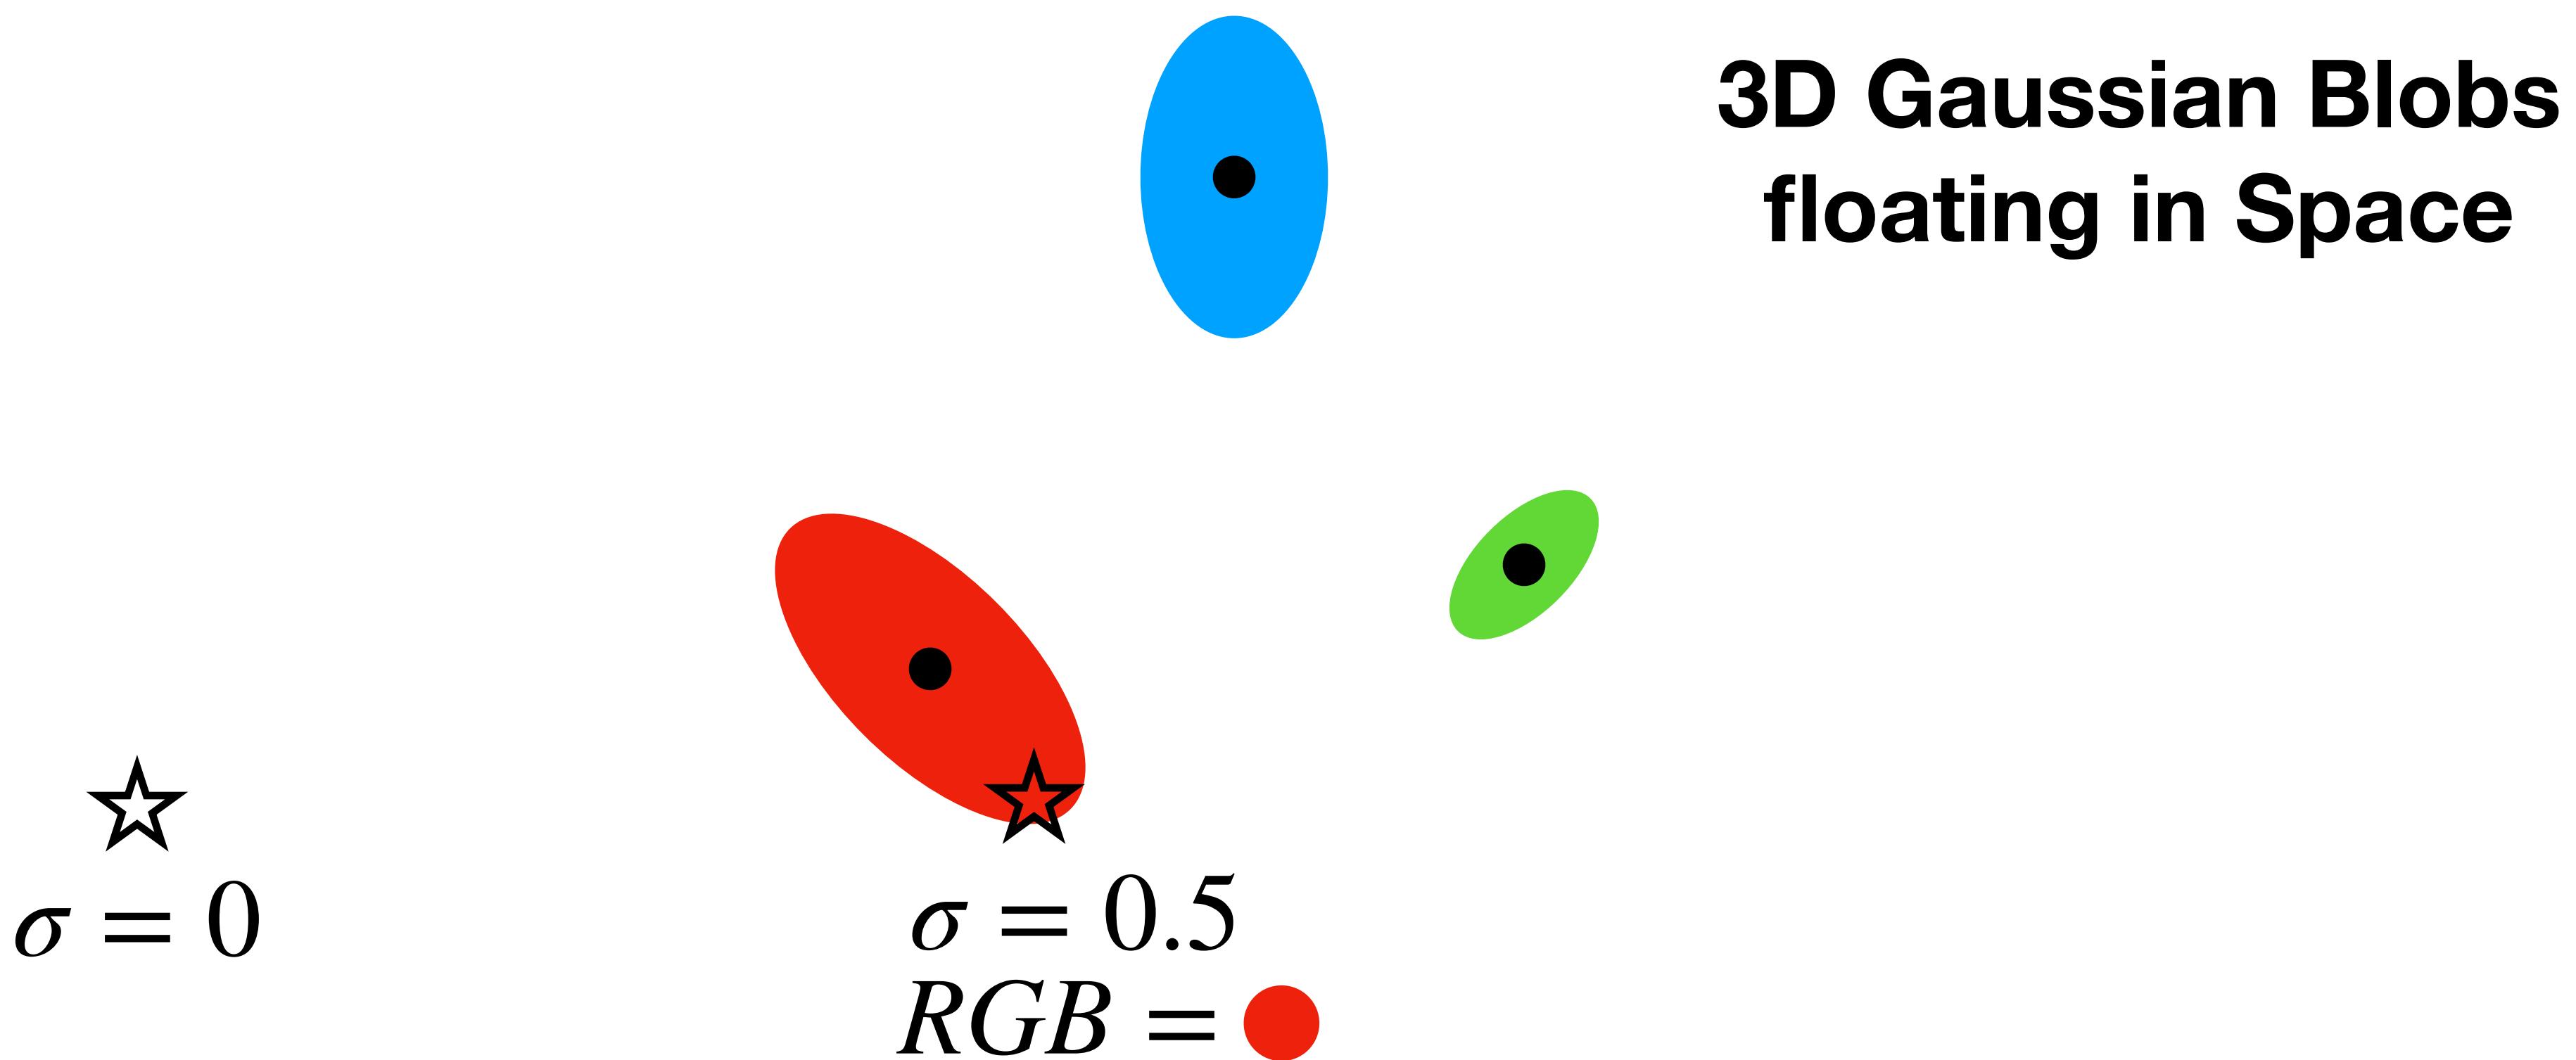
Primitive-based representations

★
 $\sigma = 0$

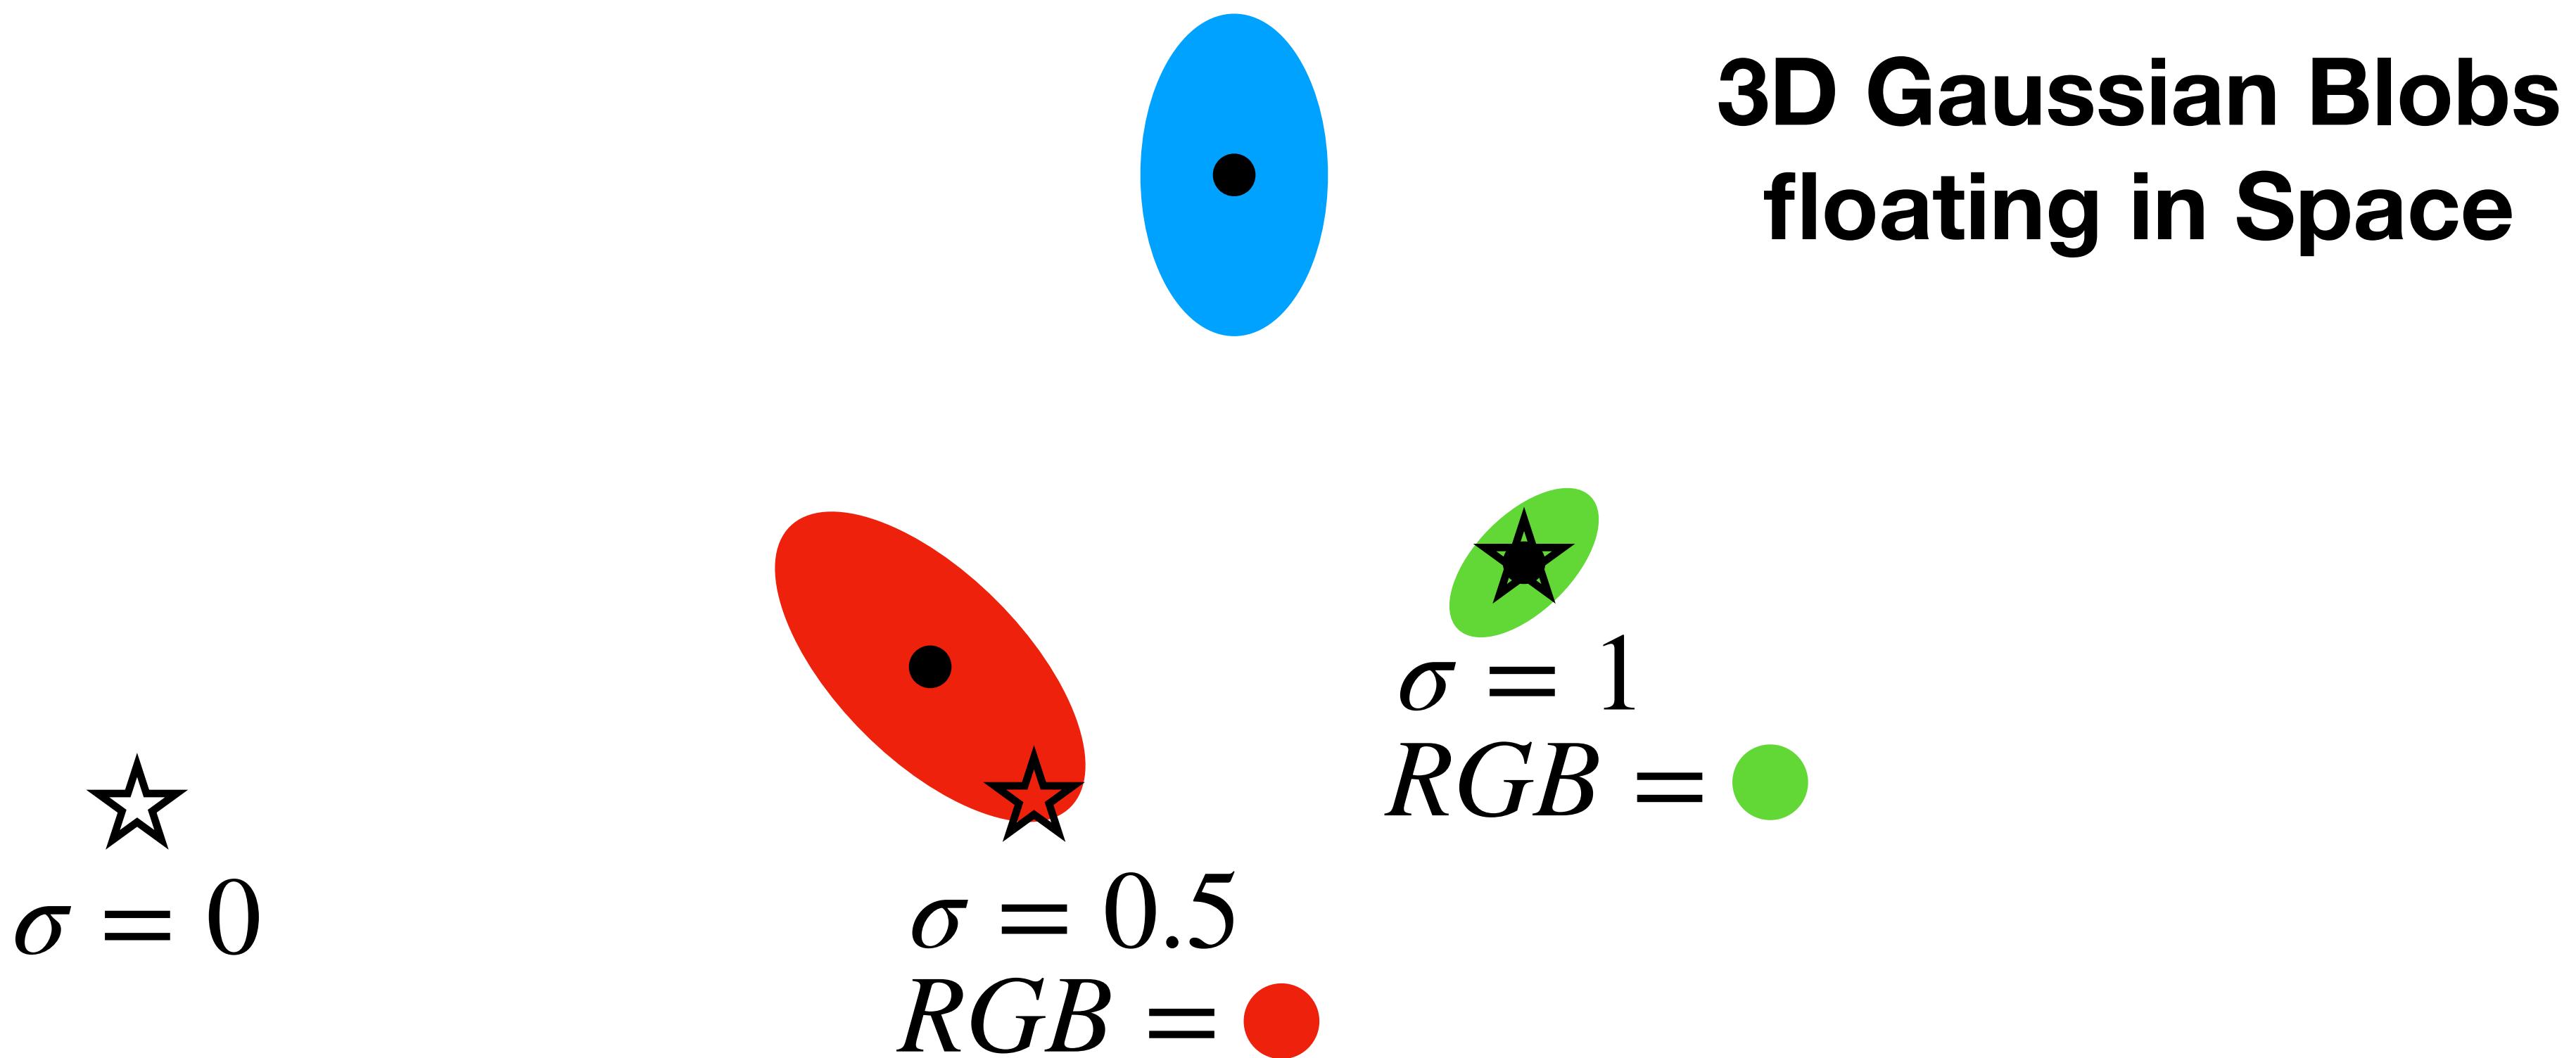


**3D Gaussian Blobs
floating in Space**

Primitive-based representations



Primitive-based representations



Anisotropic Volumetric 3D Gaussians



Final Rendering

3D Gaussian Visualization

Anisotropic Volumetric 3D Gaussians



Final Rendering

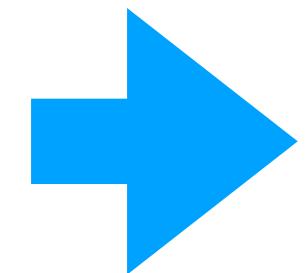
3D Gaussian Visualization

How can we represent unbounded scenes?



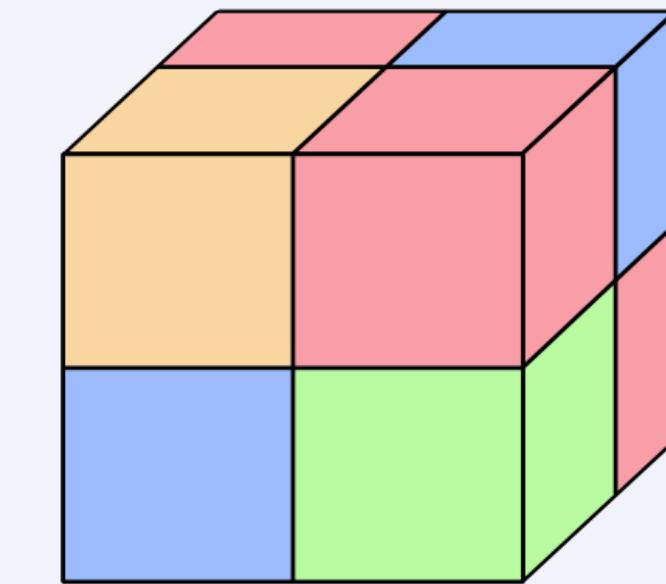
Image Source: Wikipedia

How can we represent unbounded scenes?

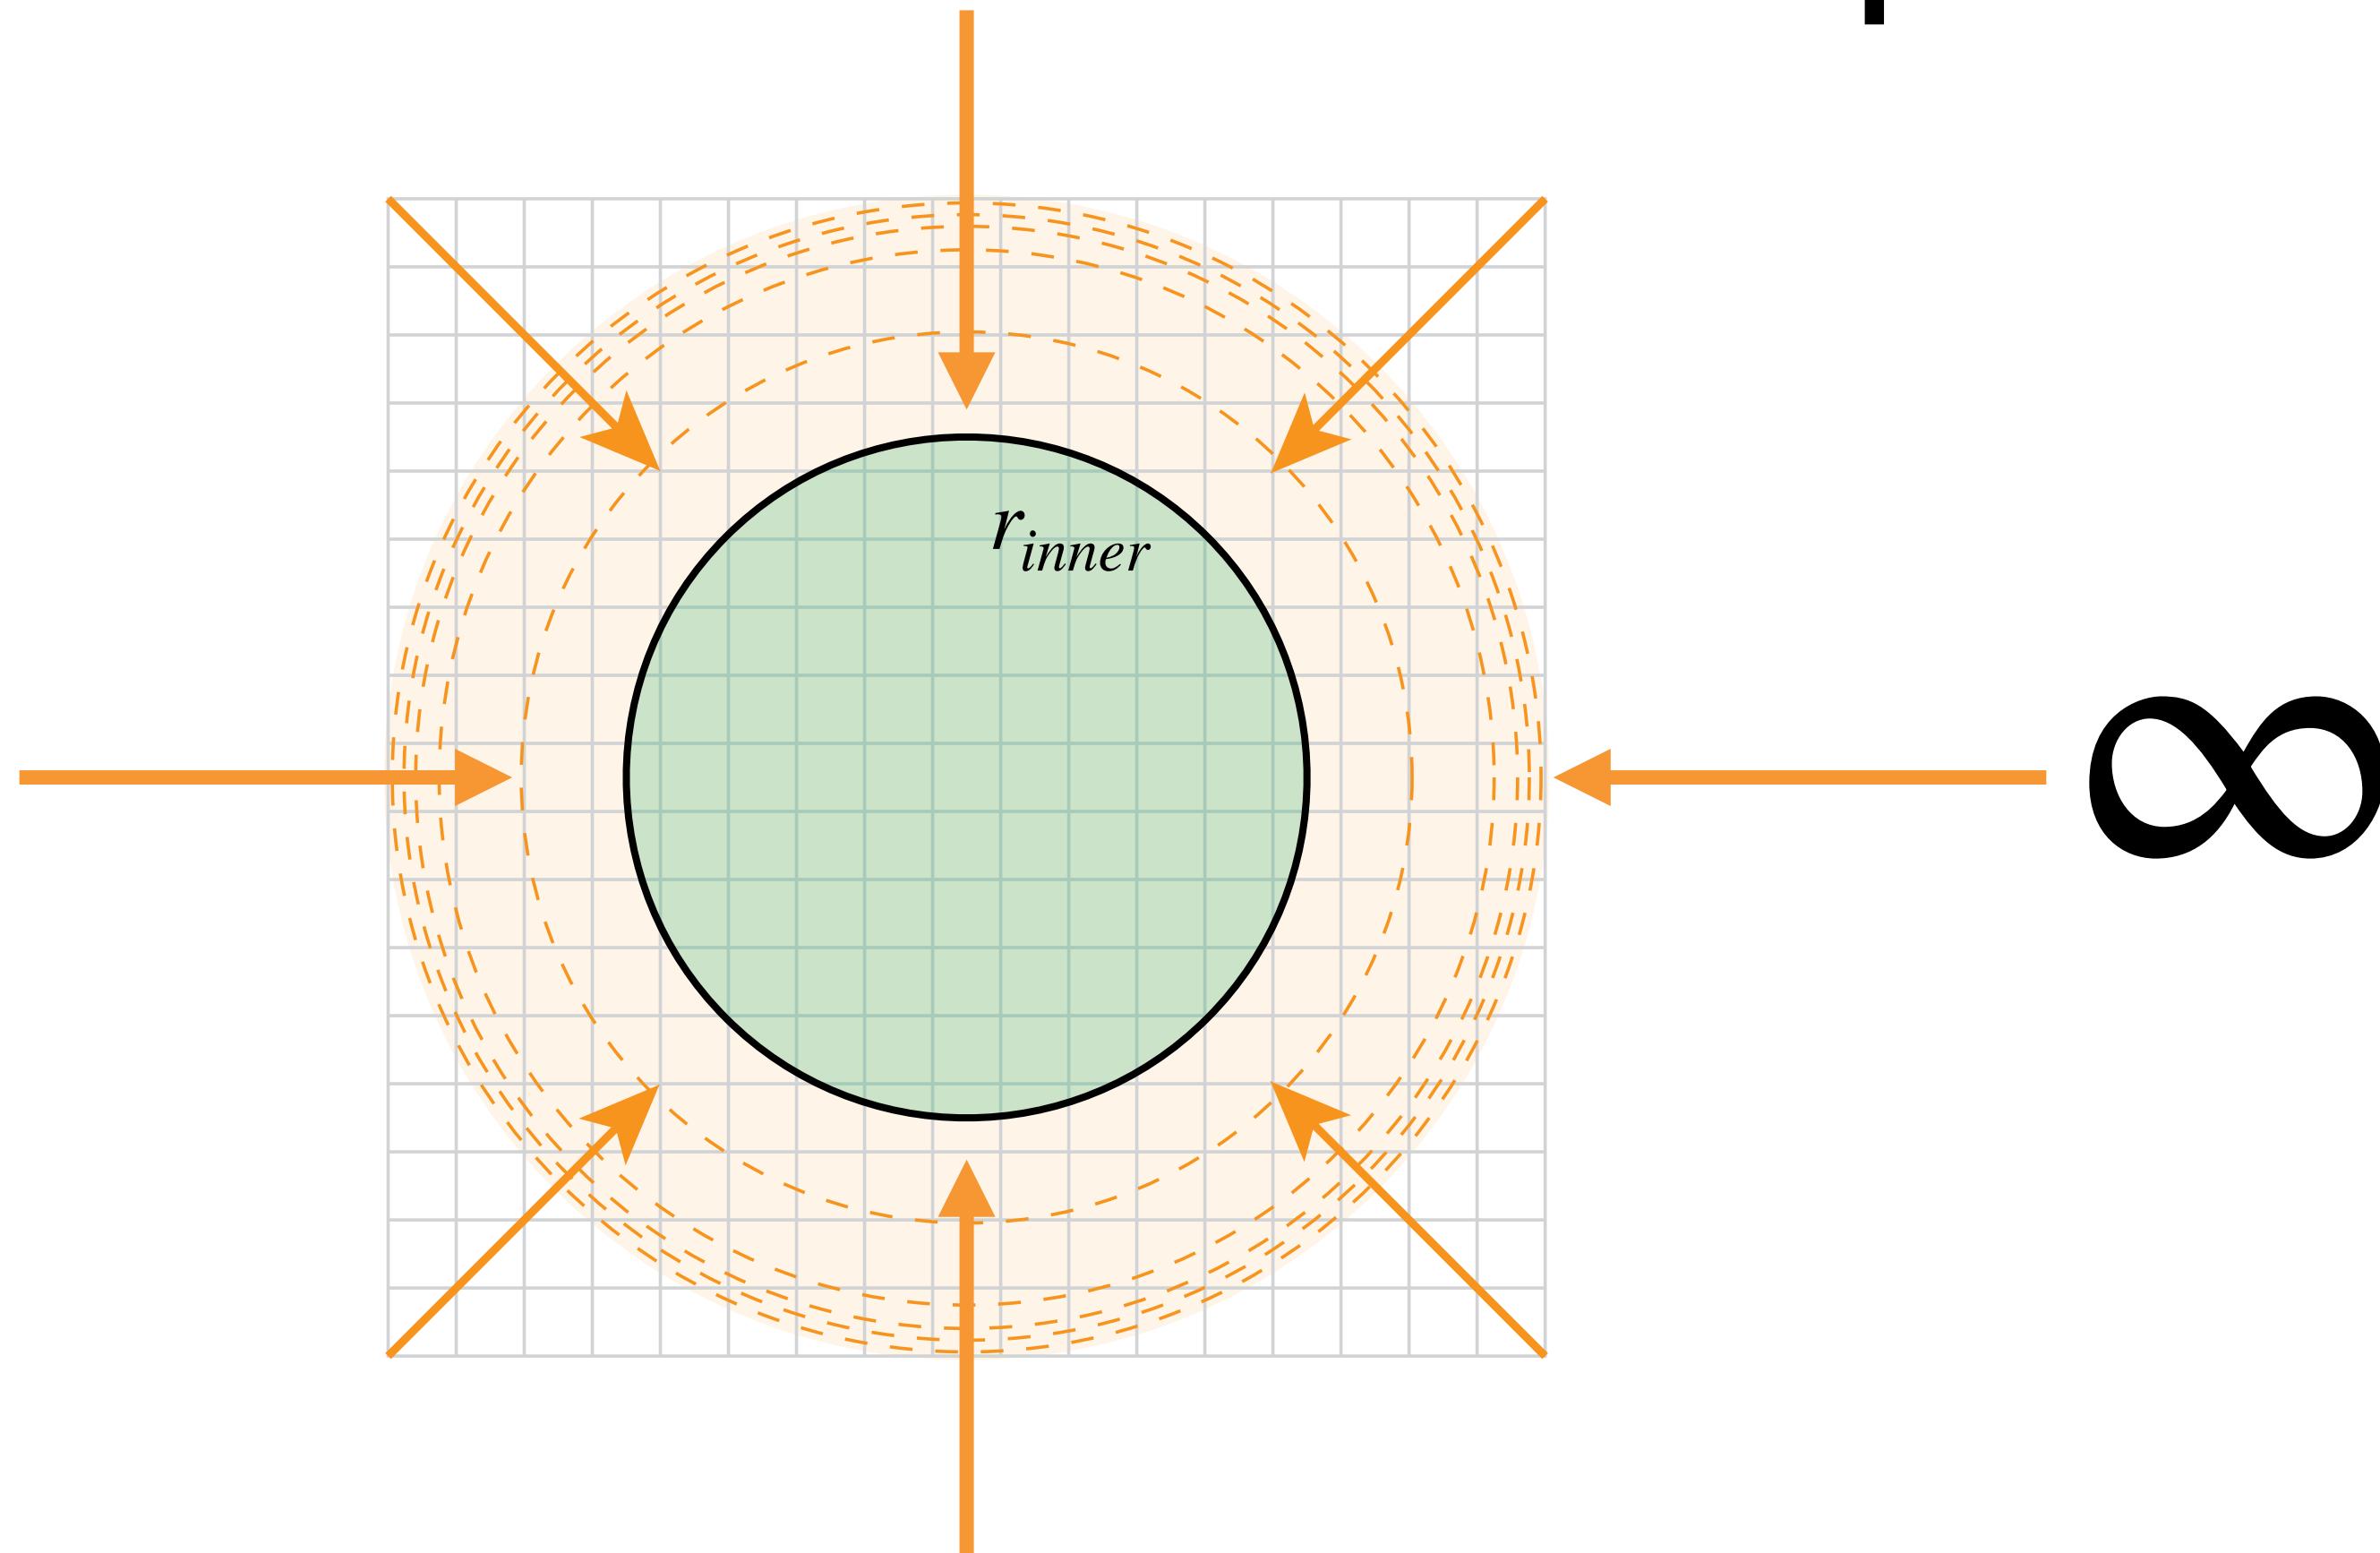


Scene
Representation

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



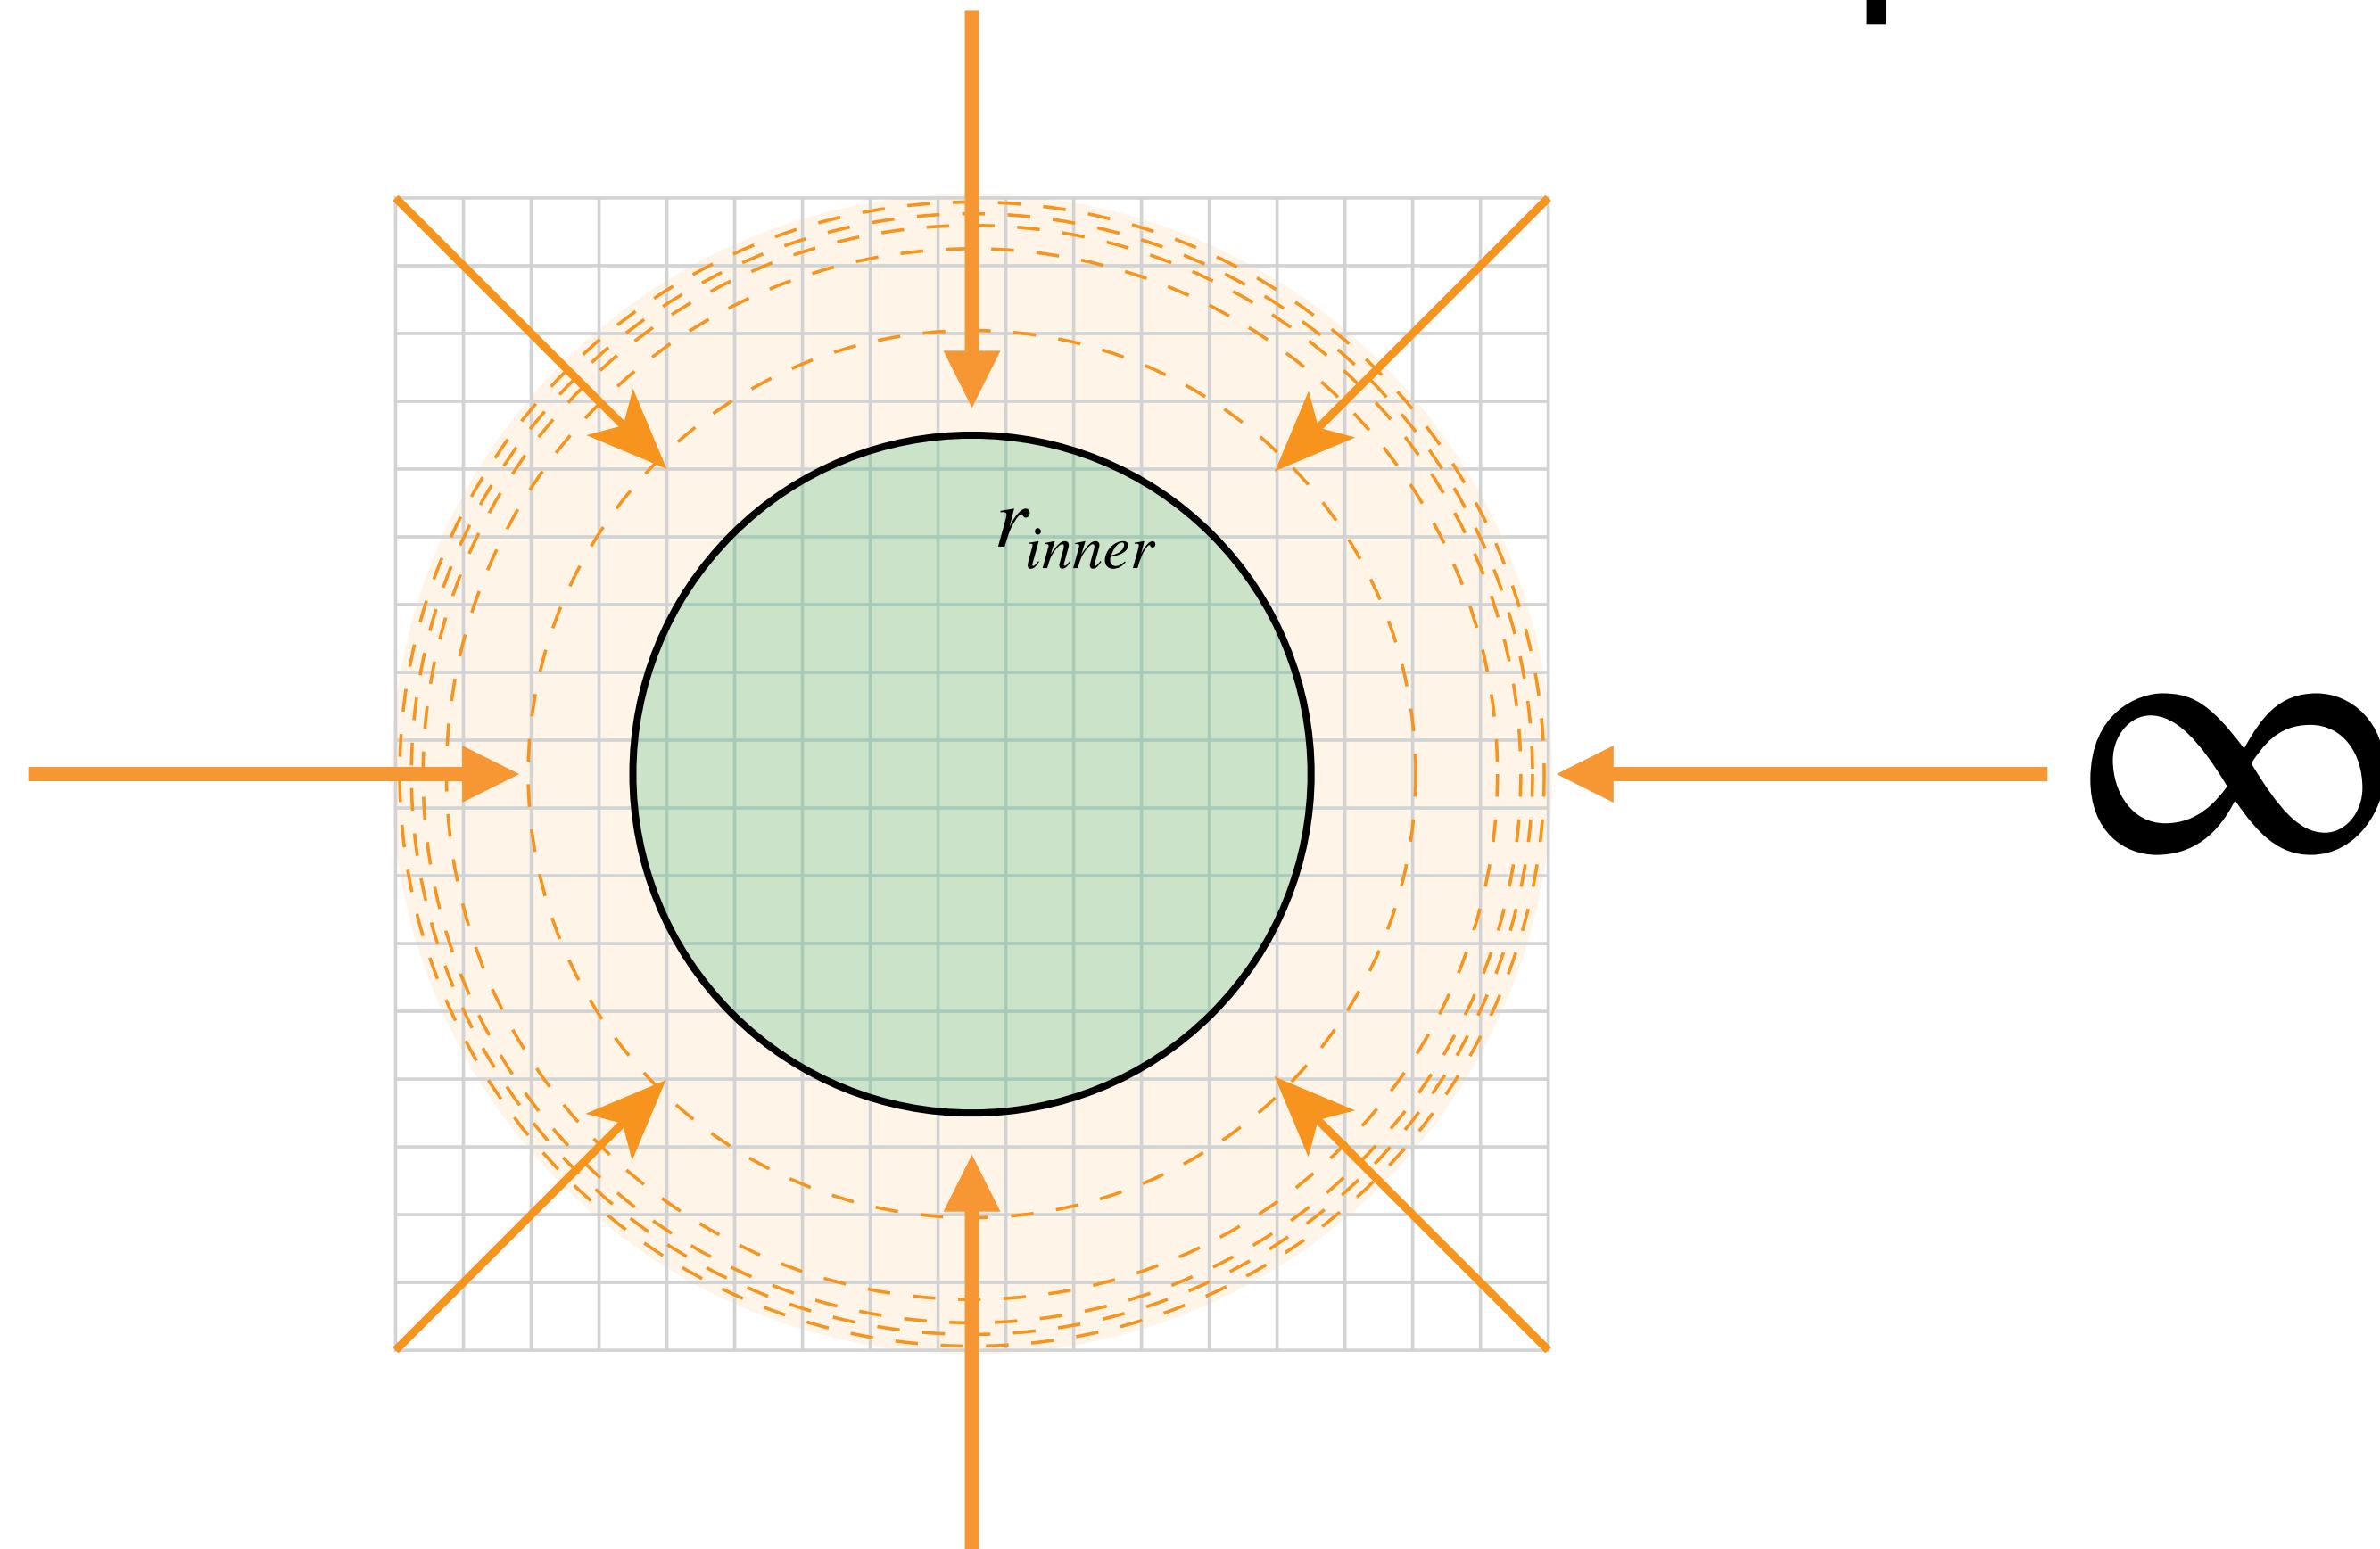
Parameterizing Unbounded Scenes: Compactification



Source: "Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields", Barron et al., CVPR 2022

Come up with nonlinear mapping that “squashes” space, such that infinity is mapped to some **finite** value. Aligns with intuition “further away, less resolution”.

Parameterizing Unbounded Scenes: Compactification



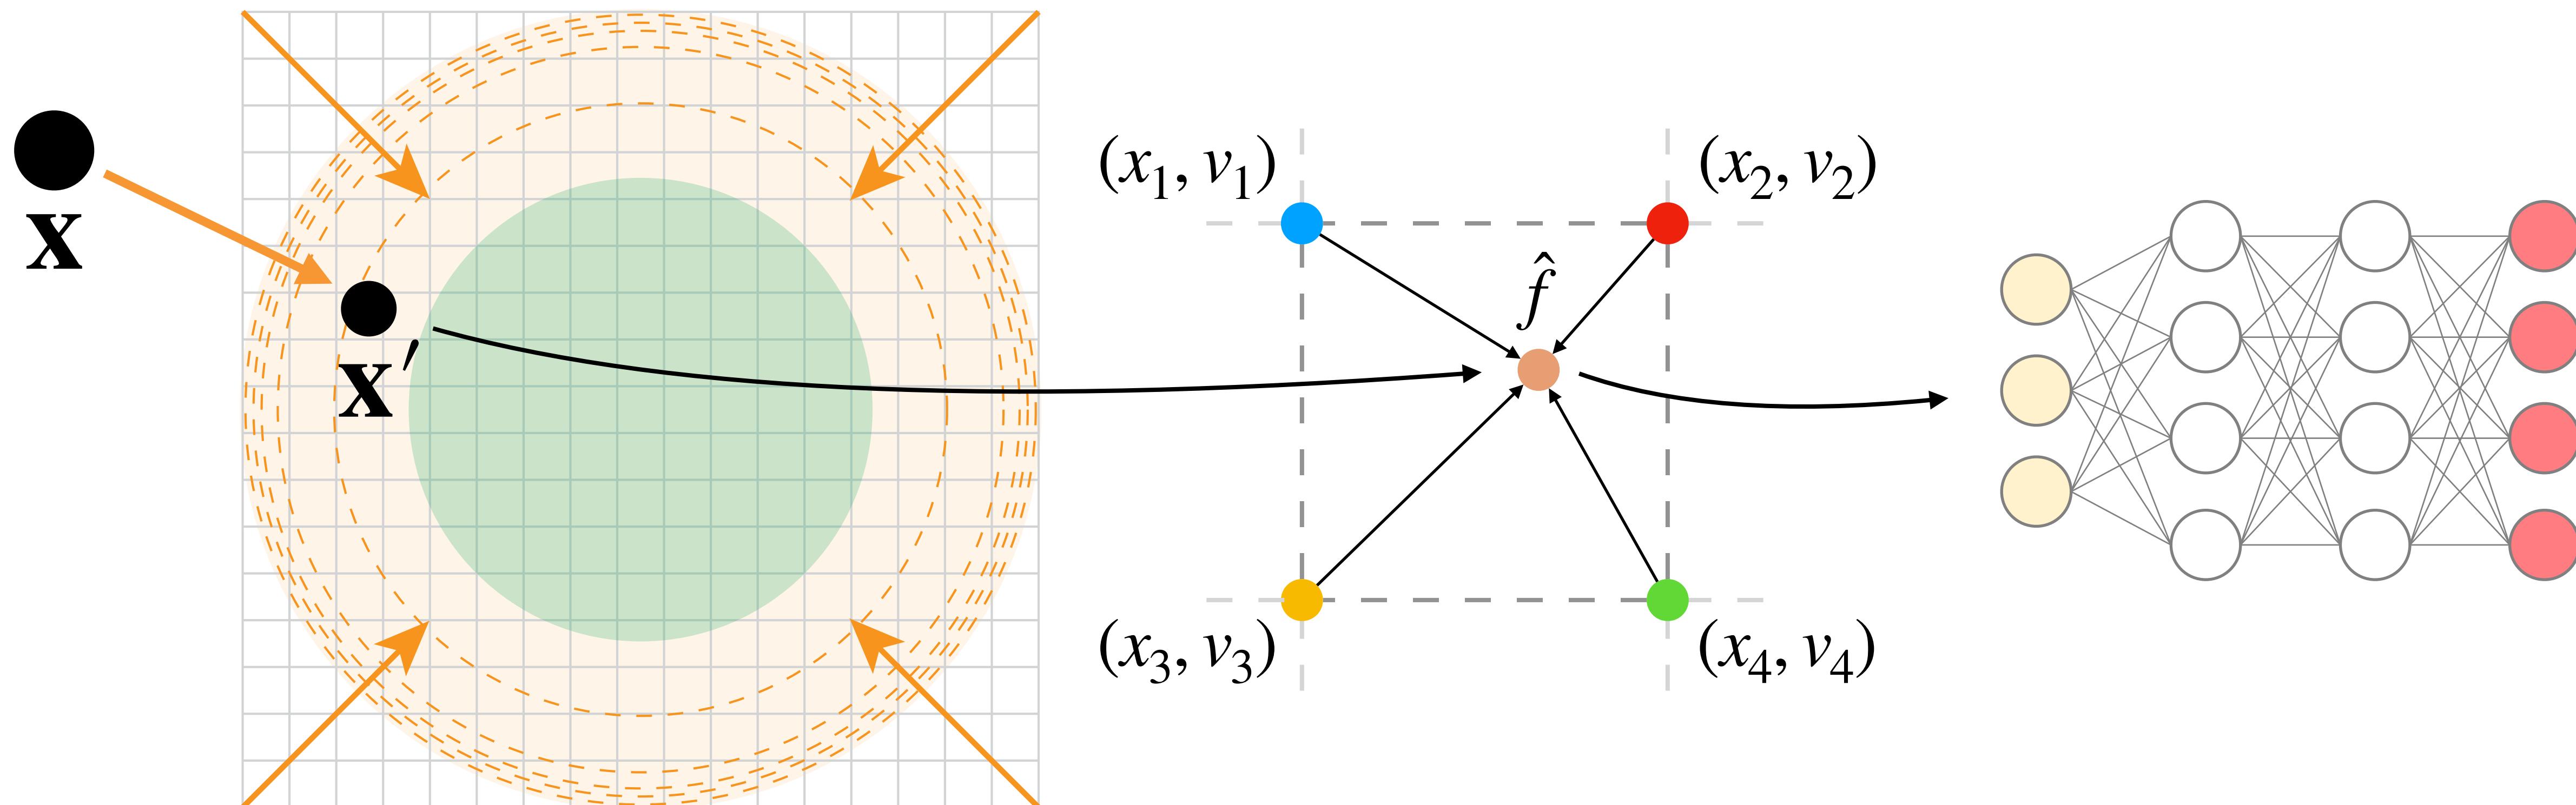
Source: "Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields", Barron et al., CVPR 2022

$$\mathbf{x}' = \frac{(1 + k) - k/\|\mathbf{u}\|}{\mathbf{u}/\|\mathbf{u}\|} r_{inner}$$
$$\mathbf{u} = \frac{\mathbf{x}}{r_{inner}}$$

Example Compactification function from
“Mip-NeRF 360: Unbounded Anti-Aliased Neural
Radiance Fields”, Barron et al., CVPR 2022

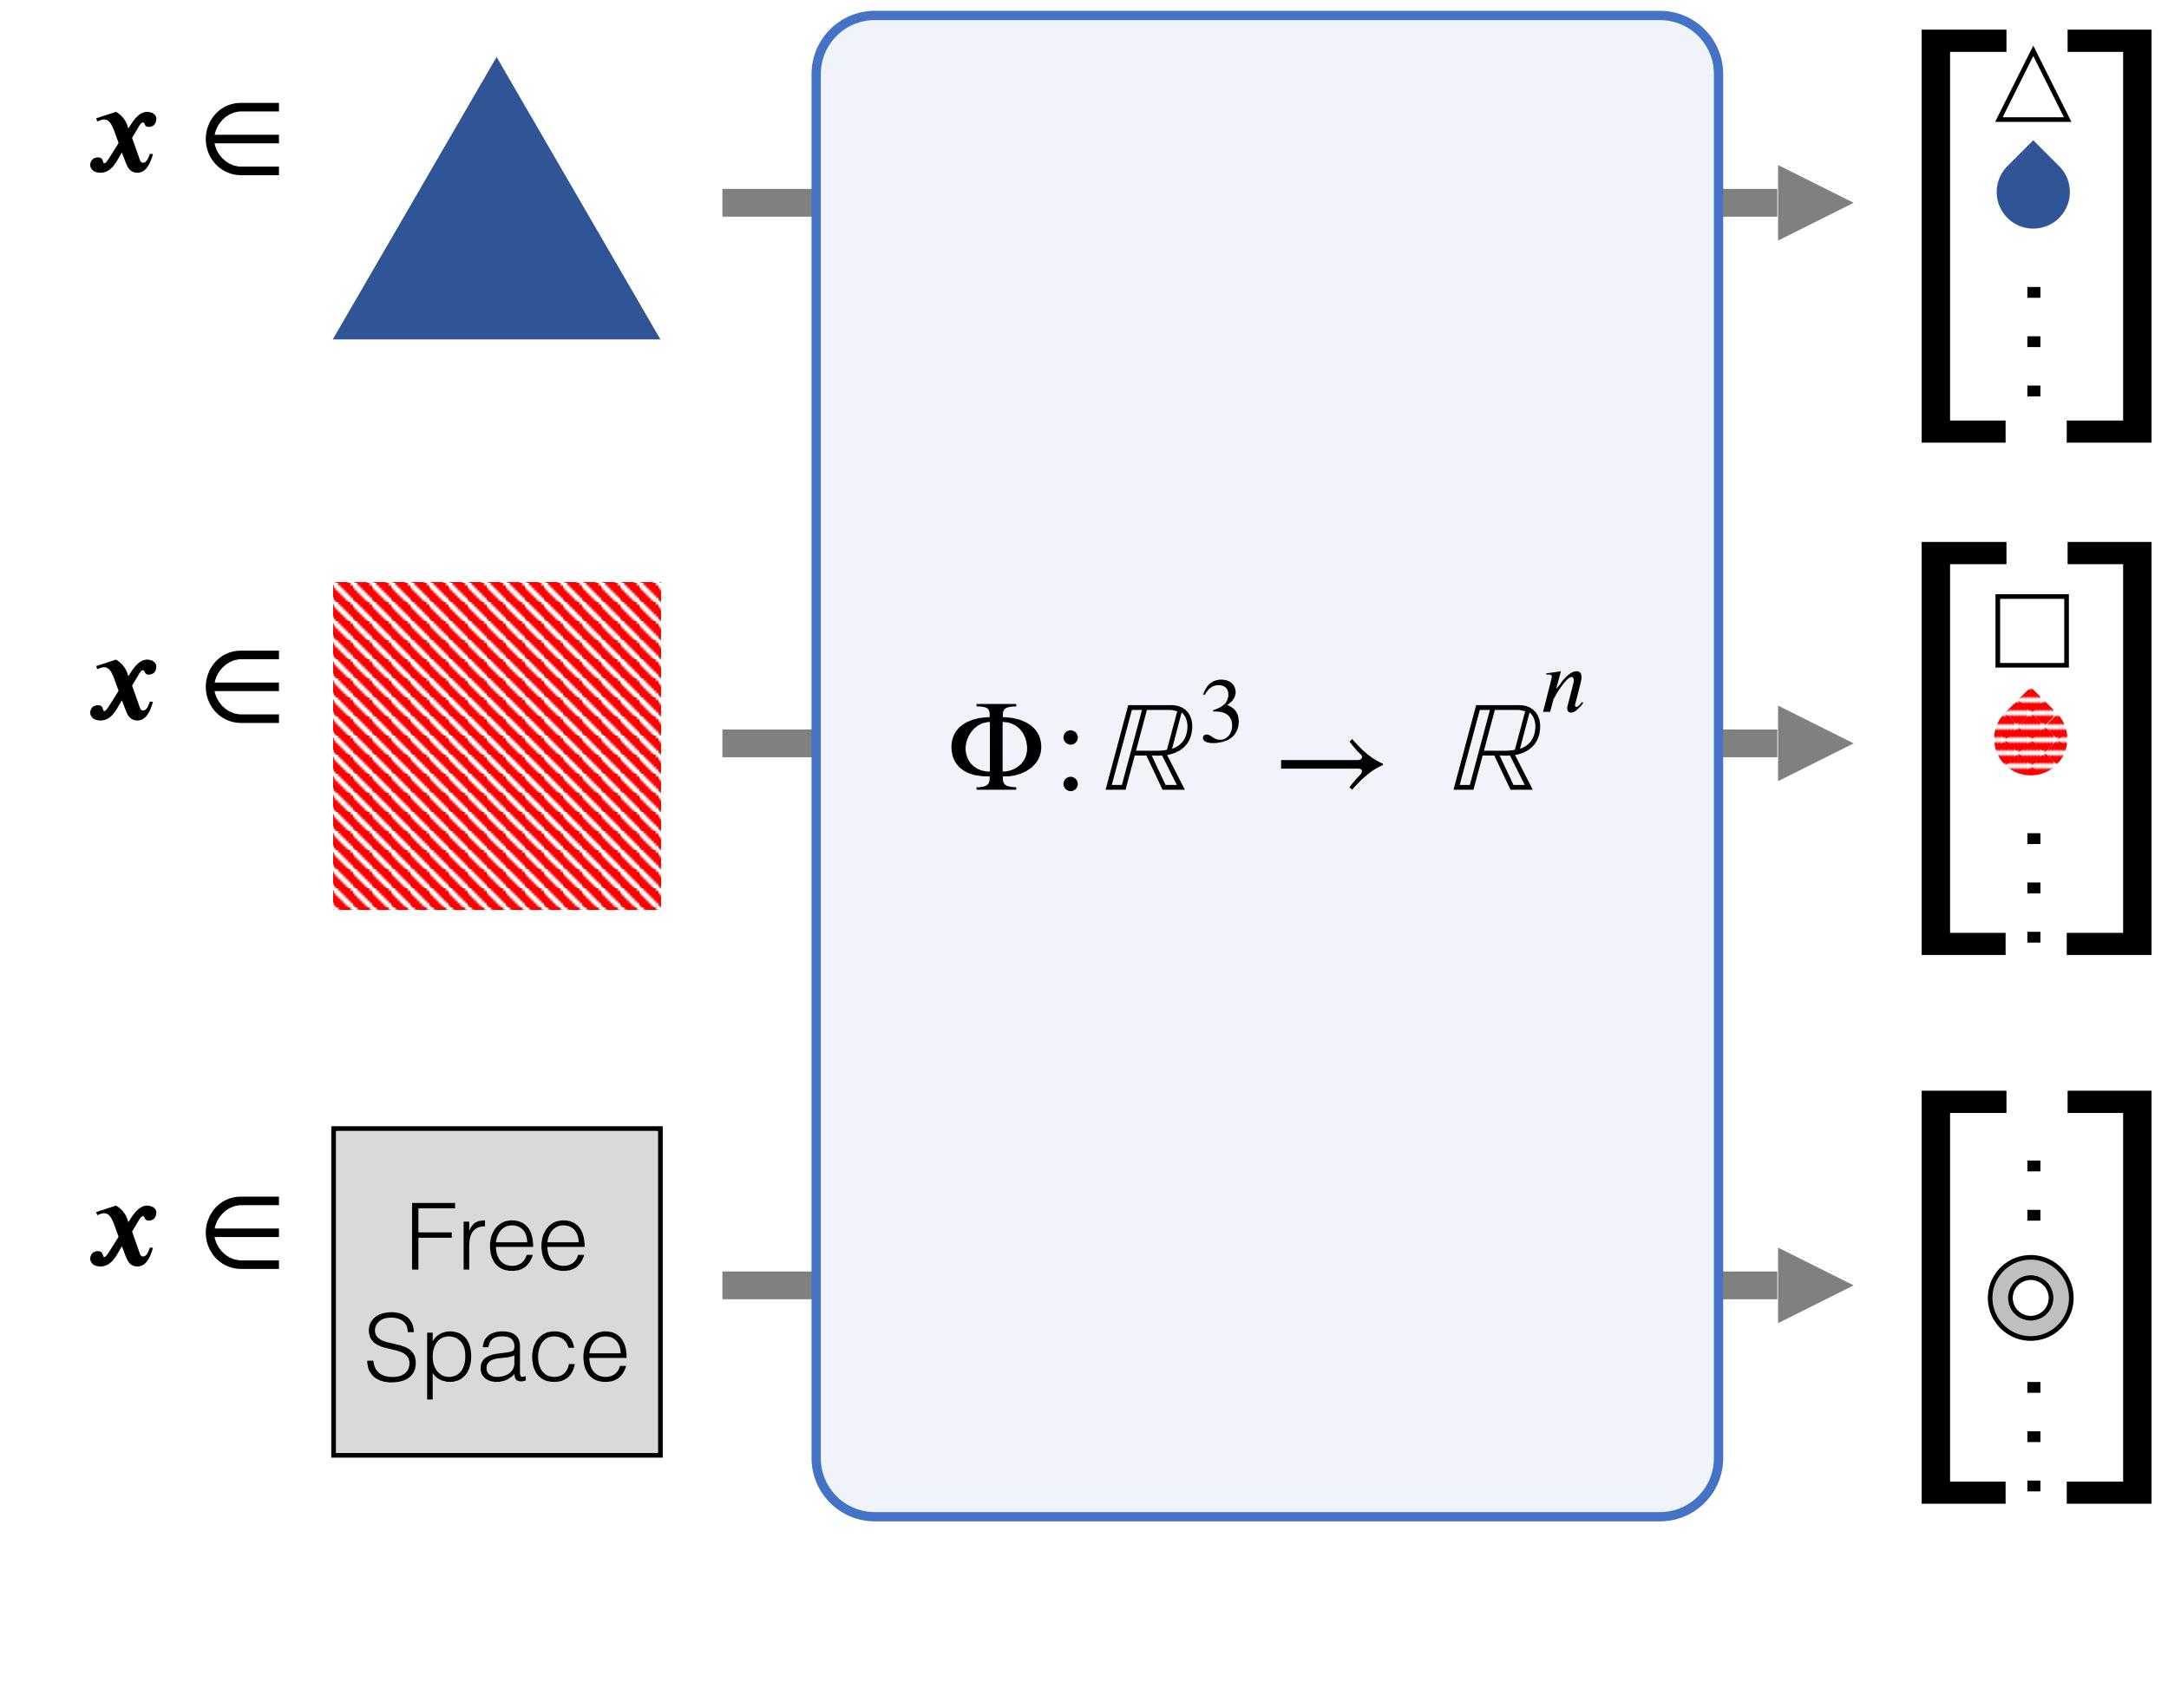
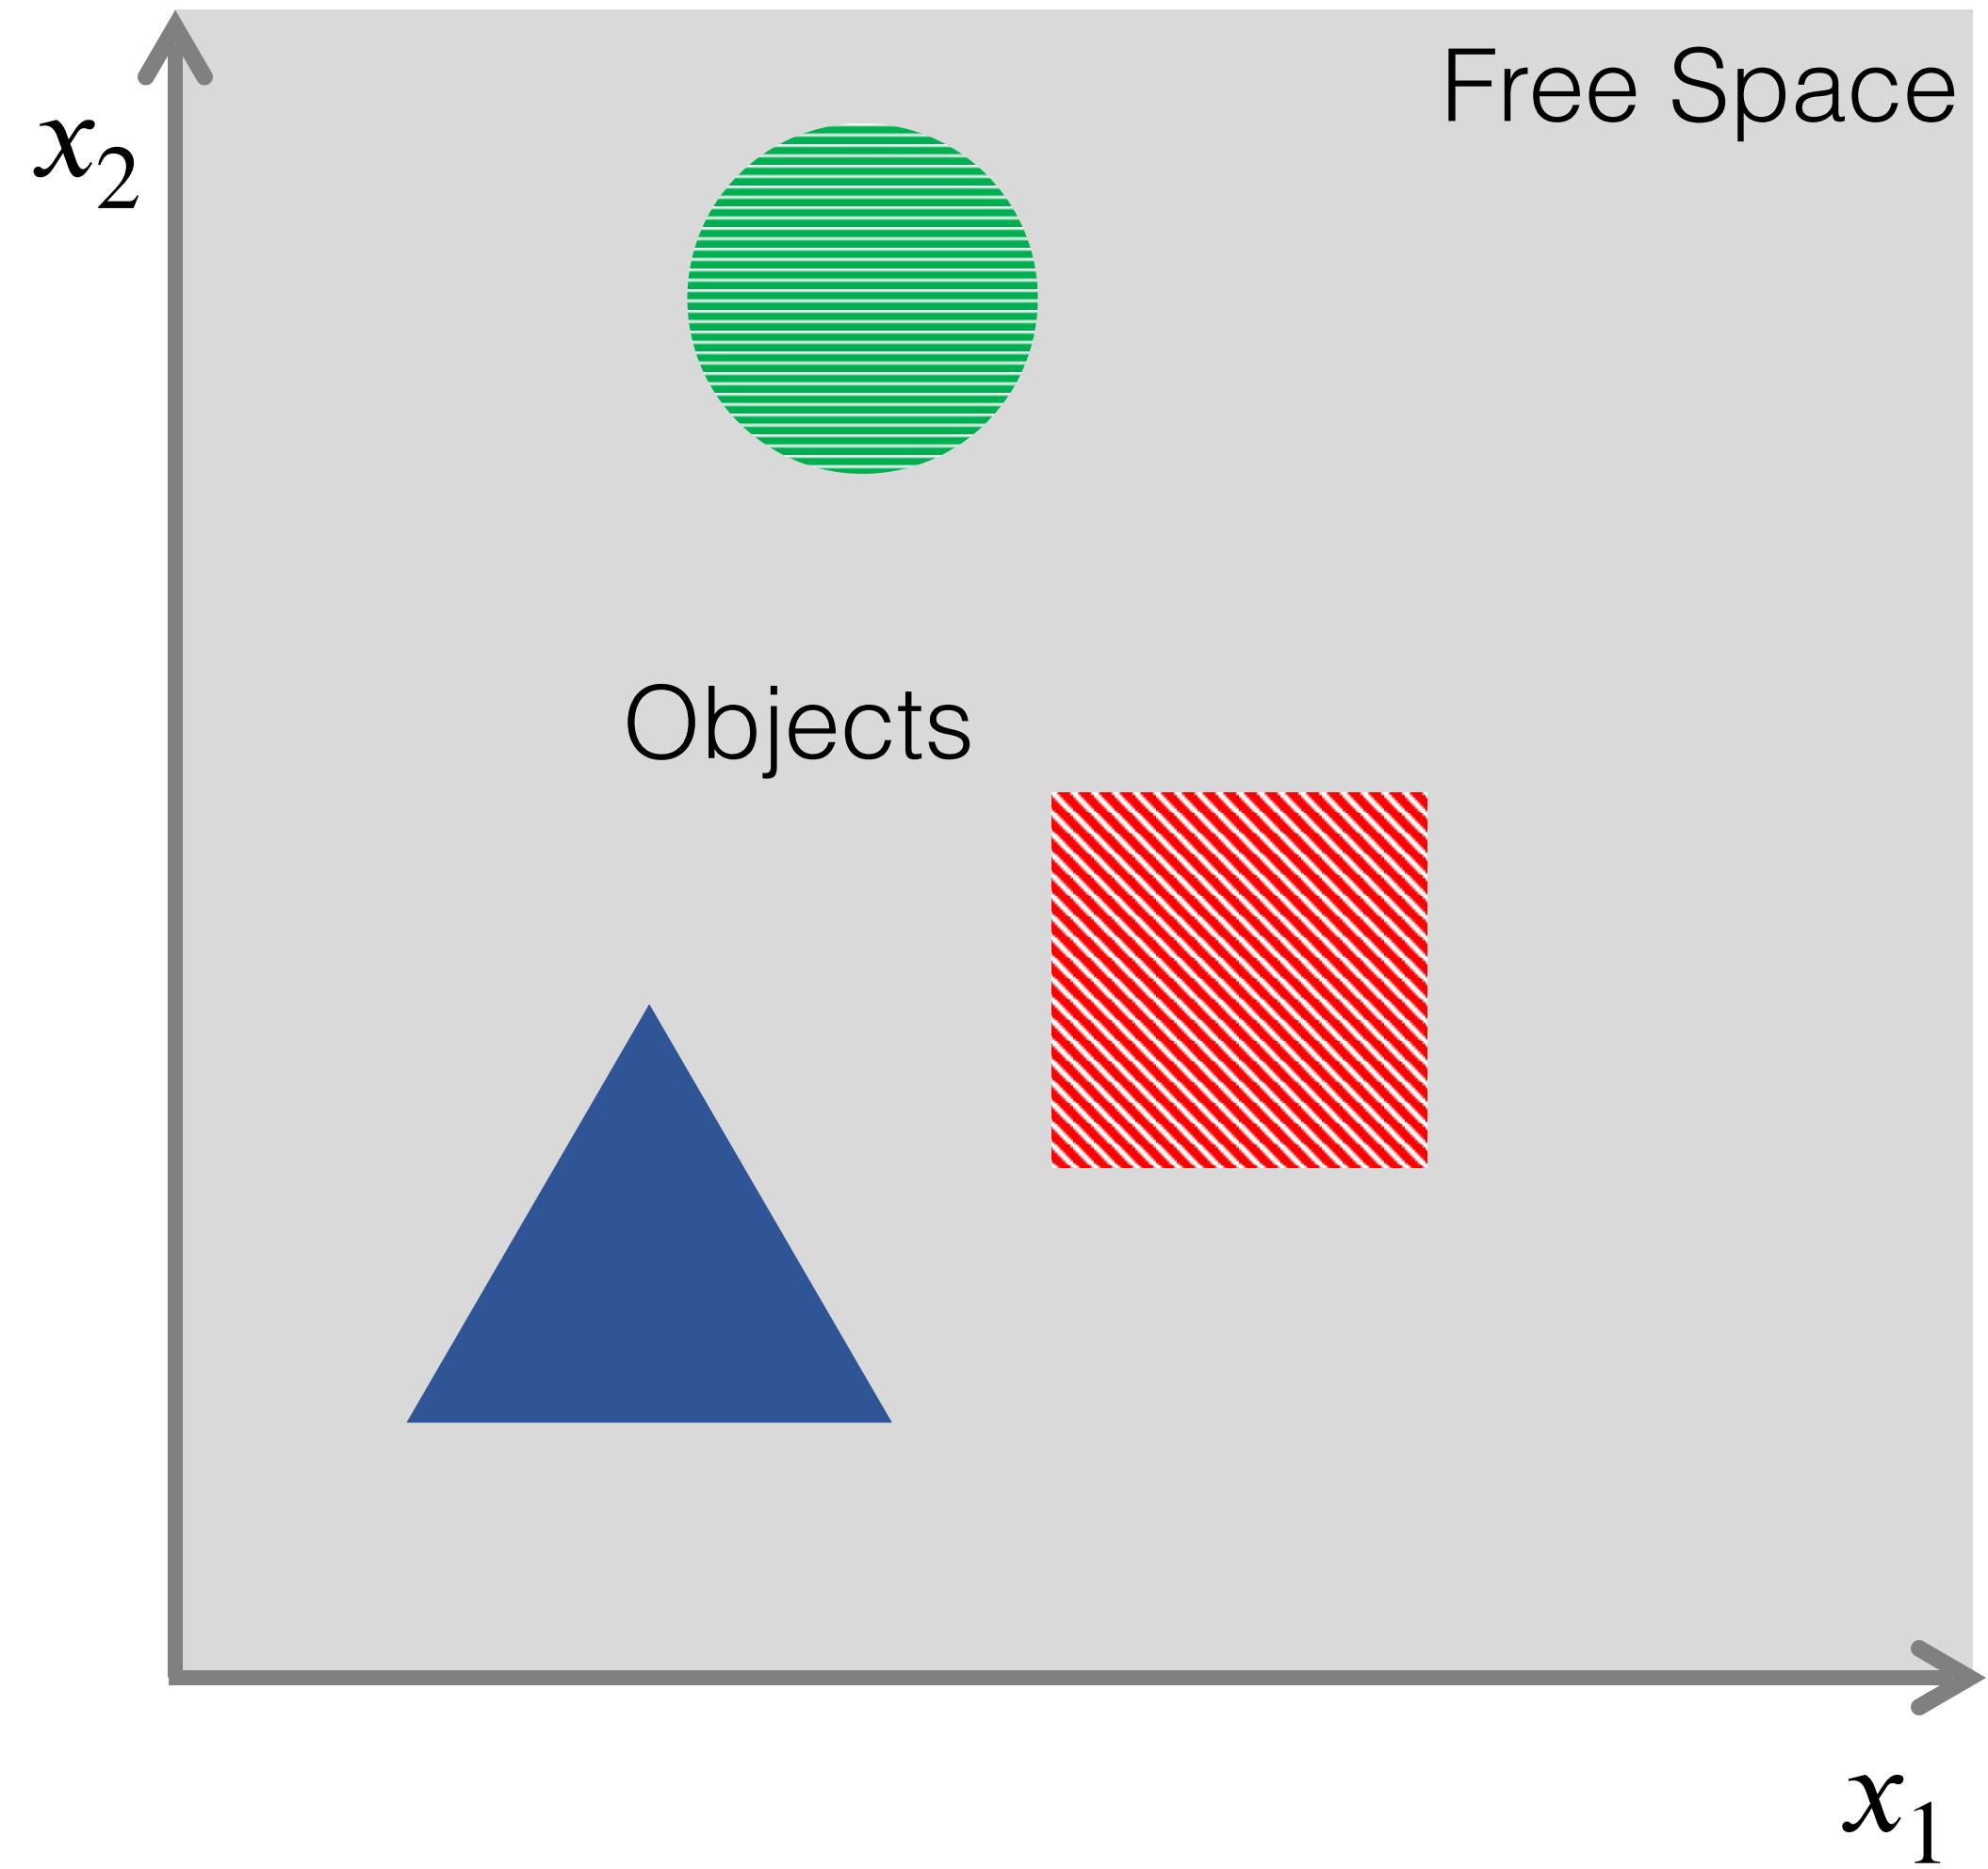
Come up with nonlinear mapping that “squashes” space, such that infinity is mapped to some **finite** value. Aligns with intuition “further away, less resolution”.

Parameterizing Unbounded Scenes: Compactification



Source: "Mip-NeRF 360: Unbounded Anti-Aliased Neural Radiance Fields", Barron et al., CVPR 2022

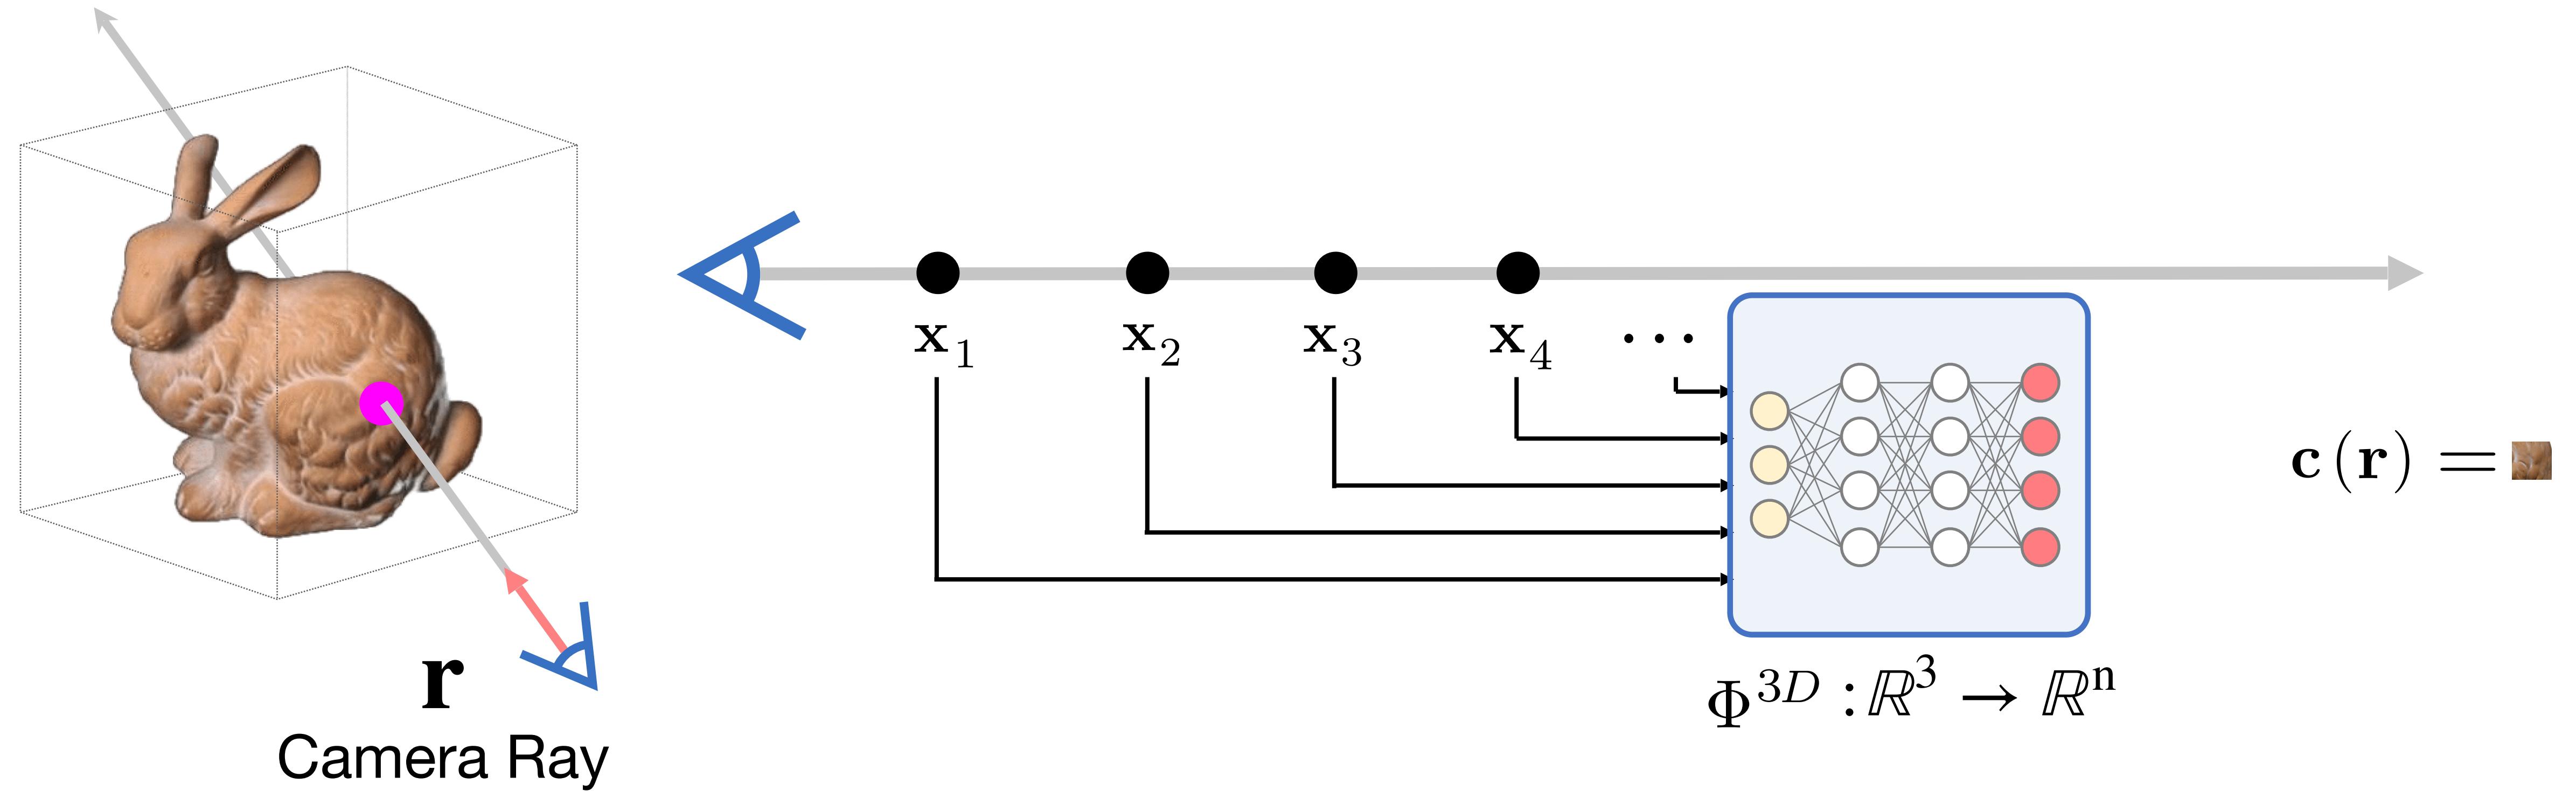
Do we have to parameterize 3D Scene as 3D Function?



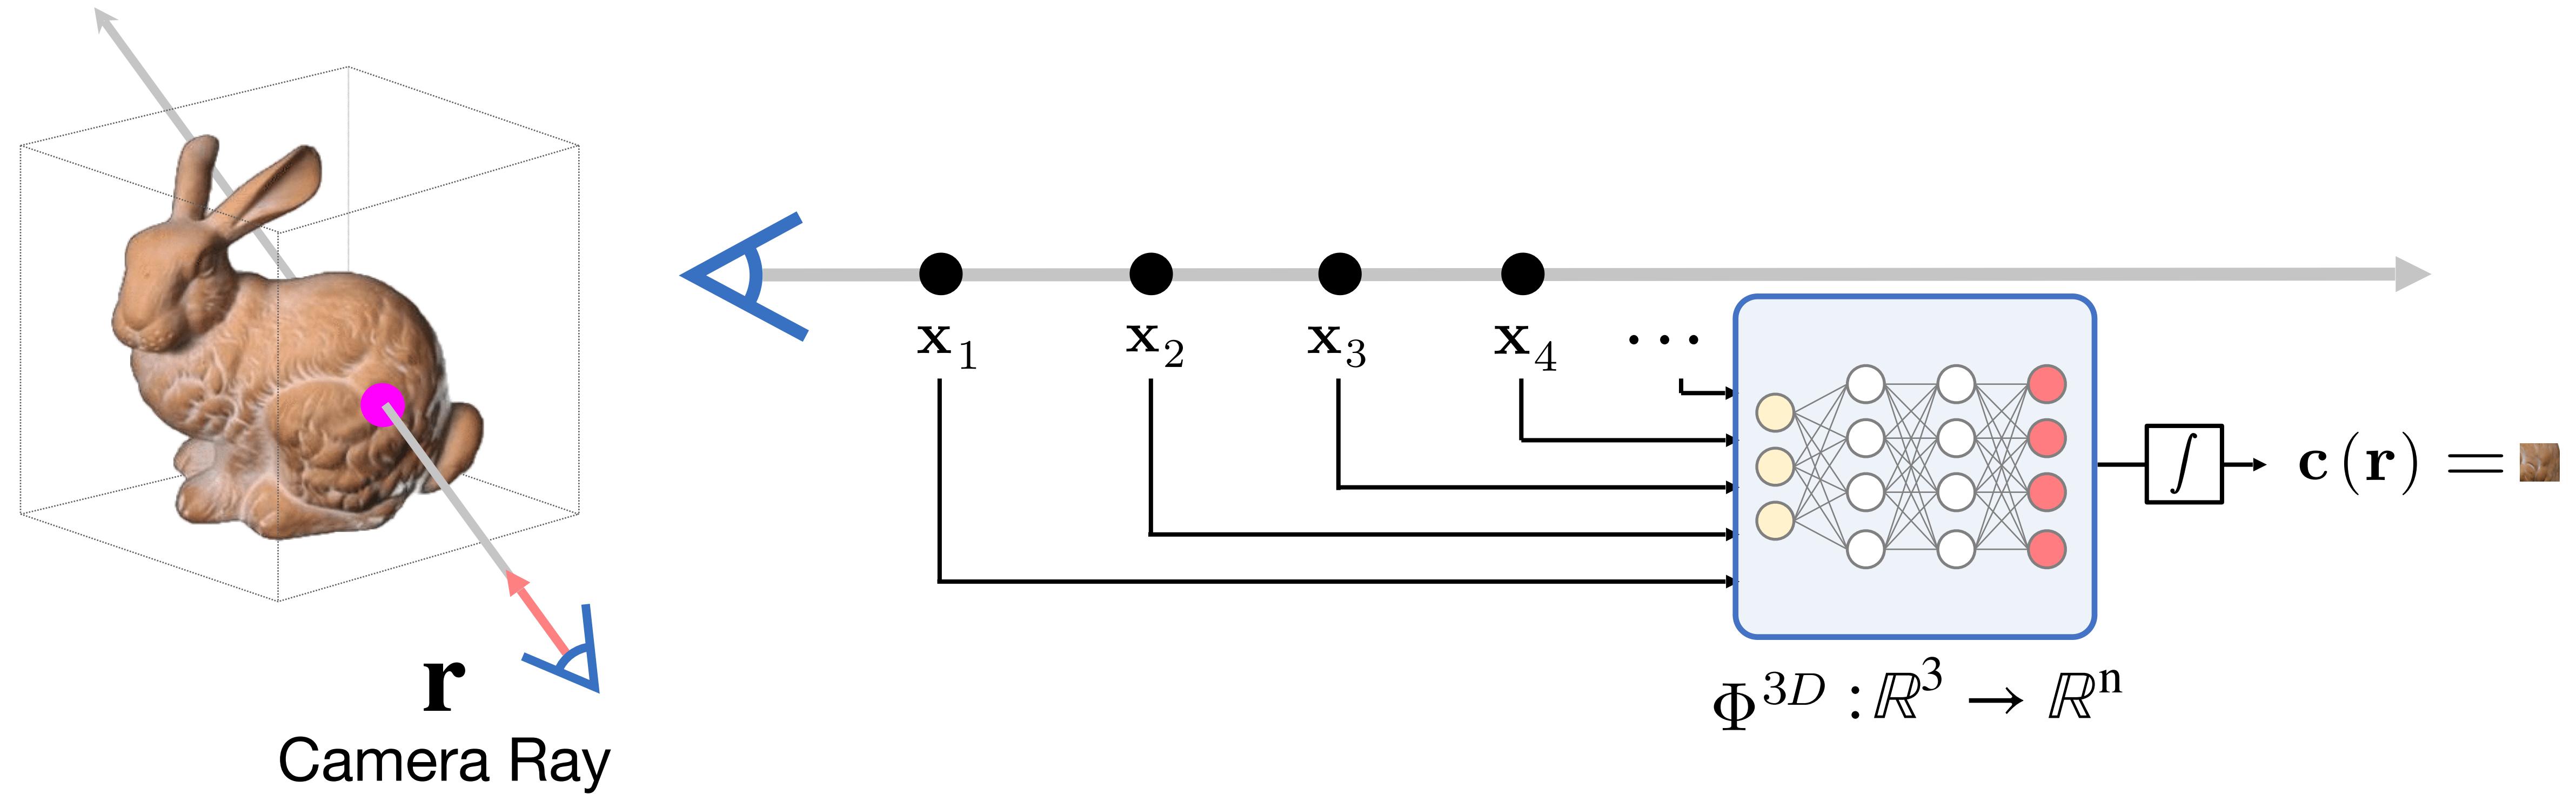
3D-structured Representation

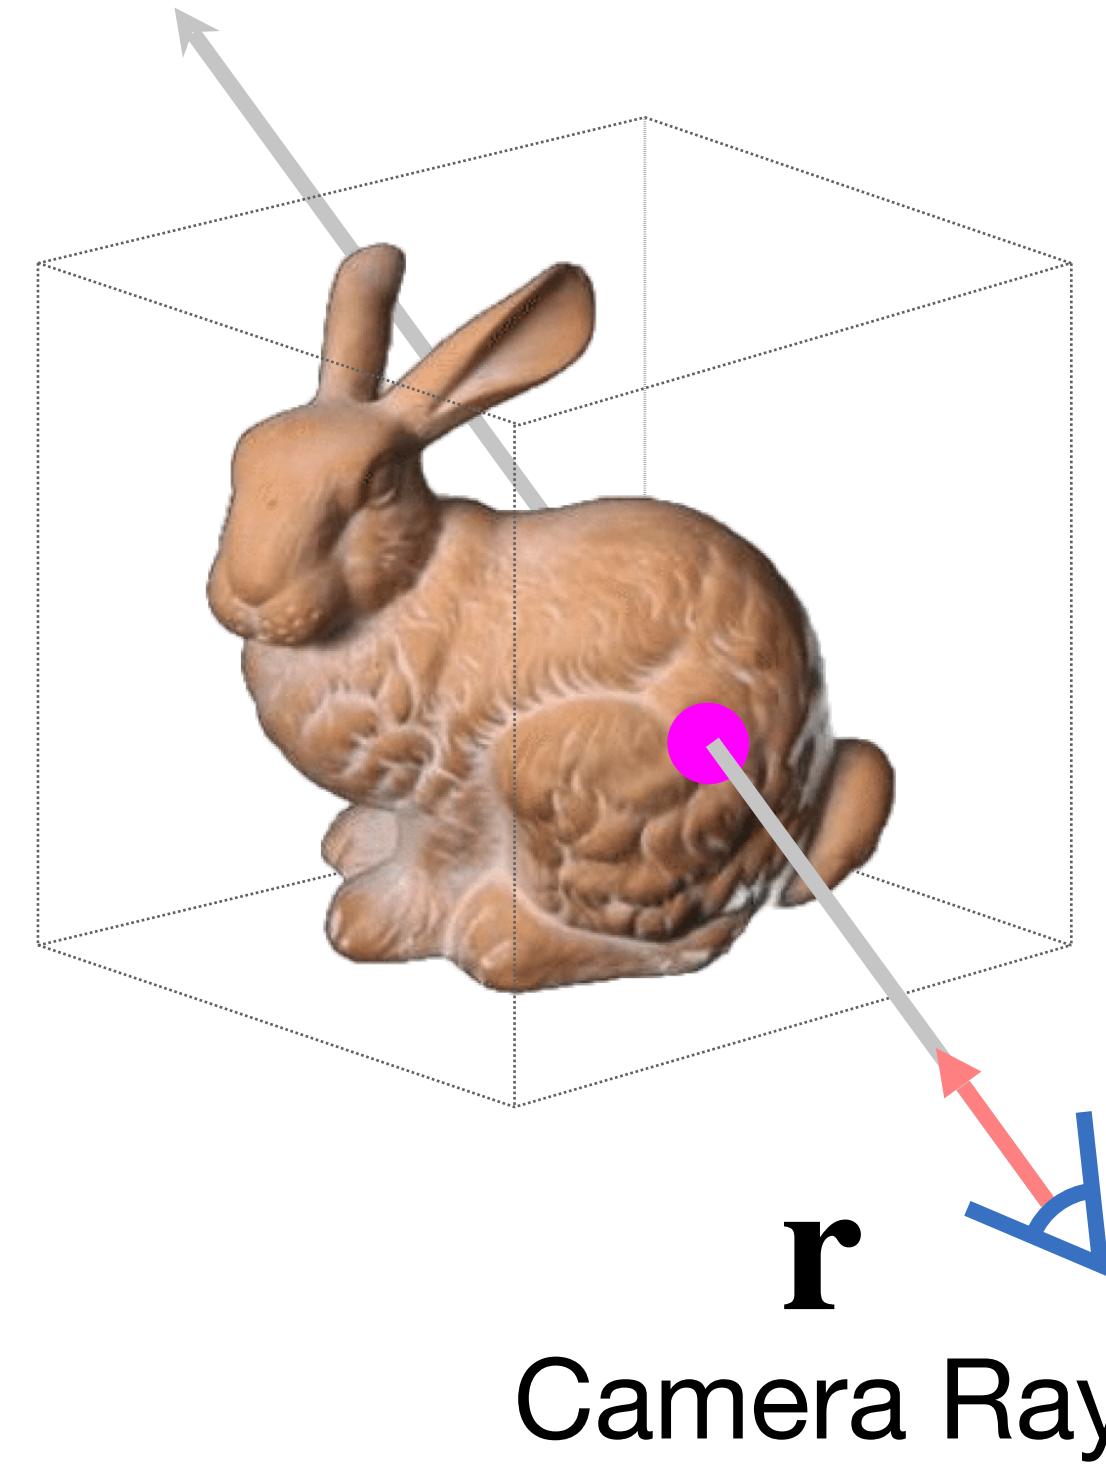


3D-structured Representation



3D-structured Representation





A

x_1

x_2

x_3

x_4

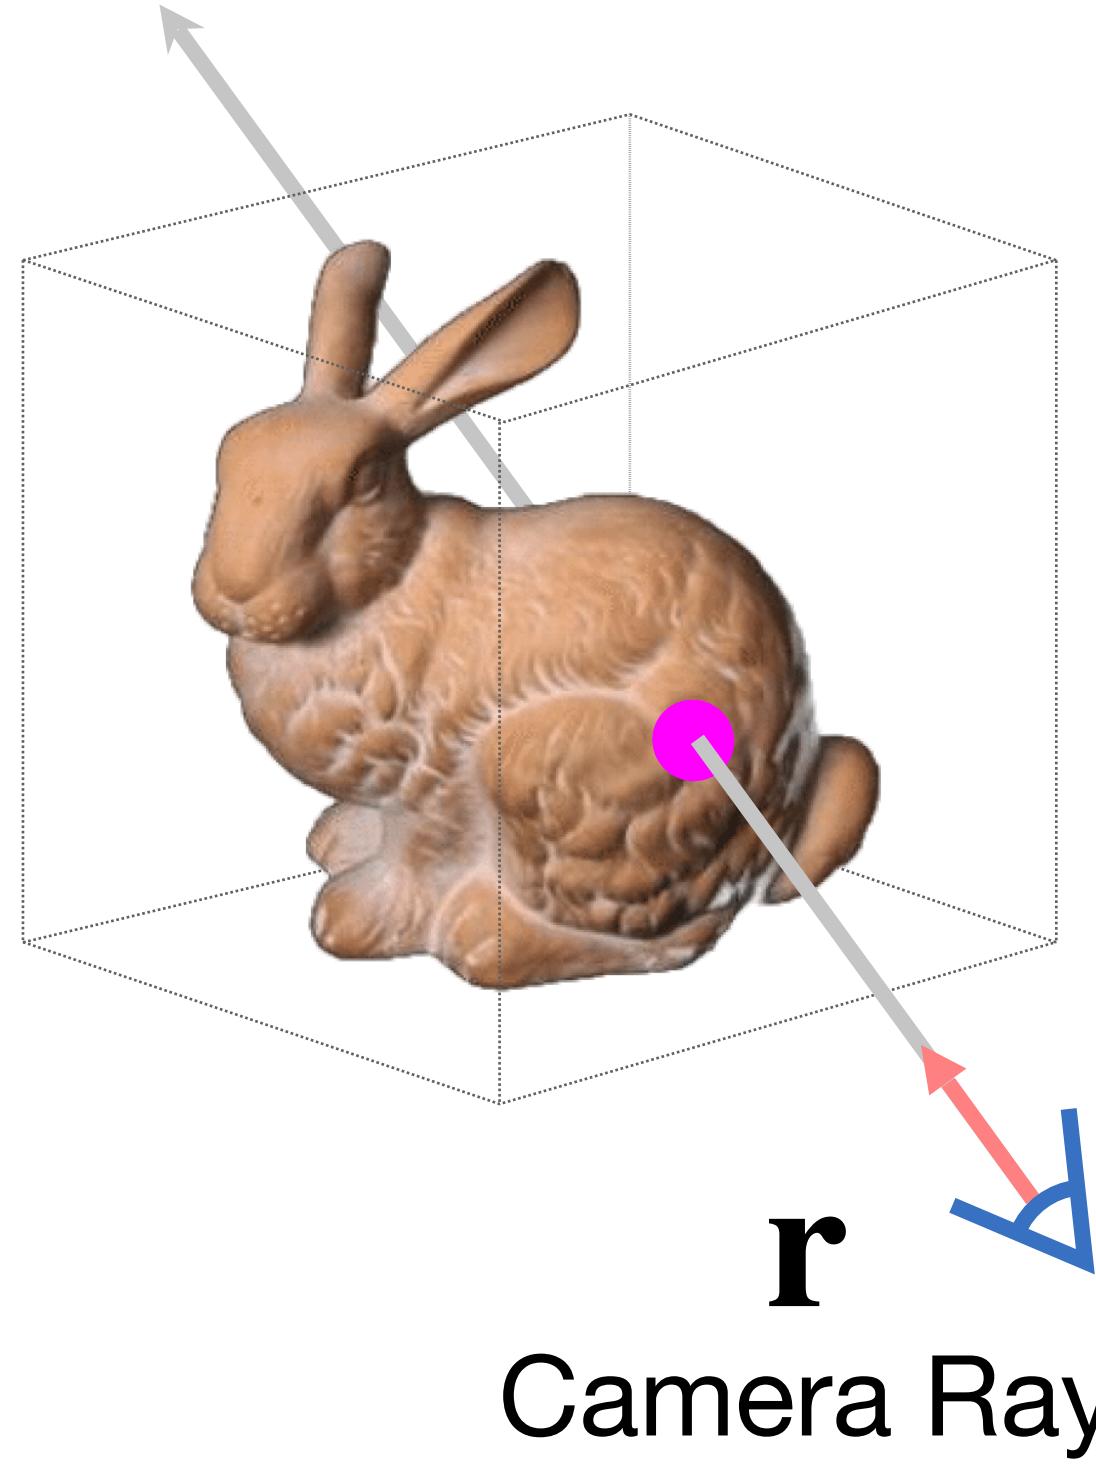
...

Light Field

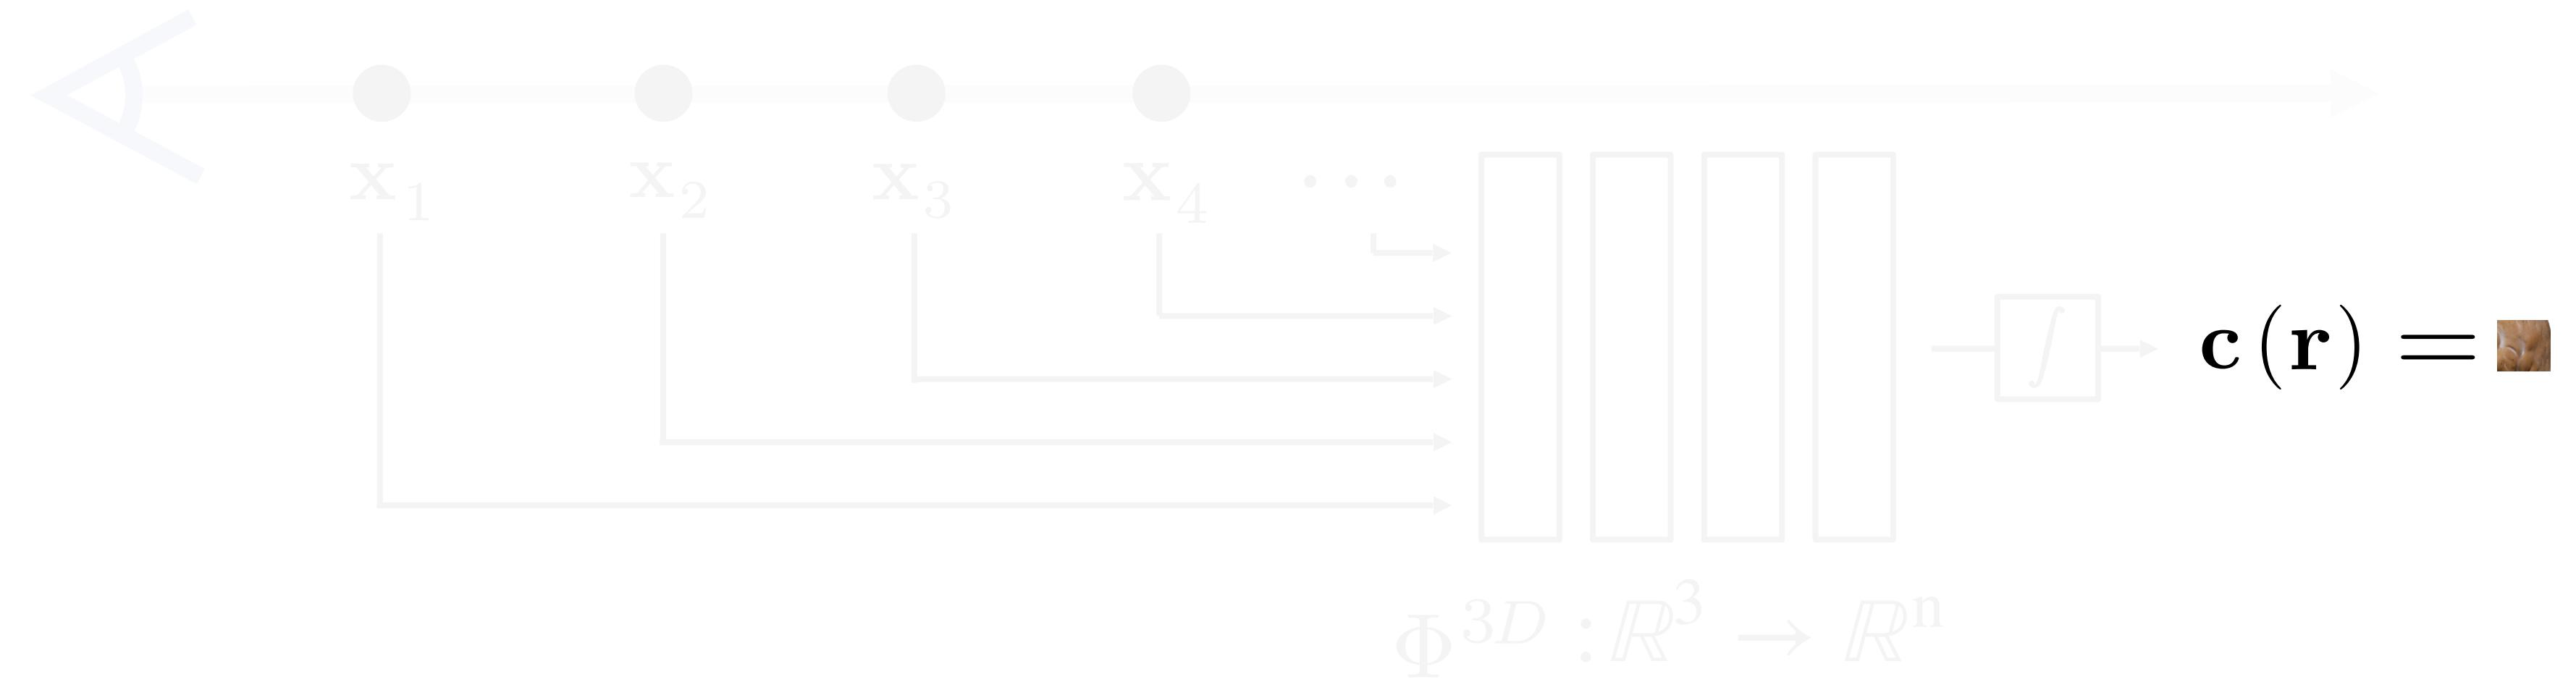
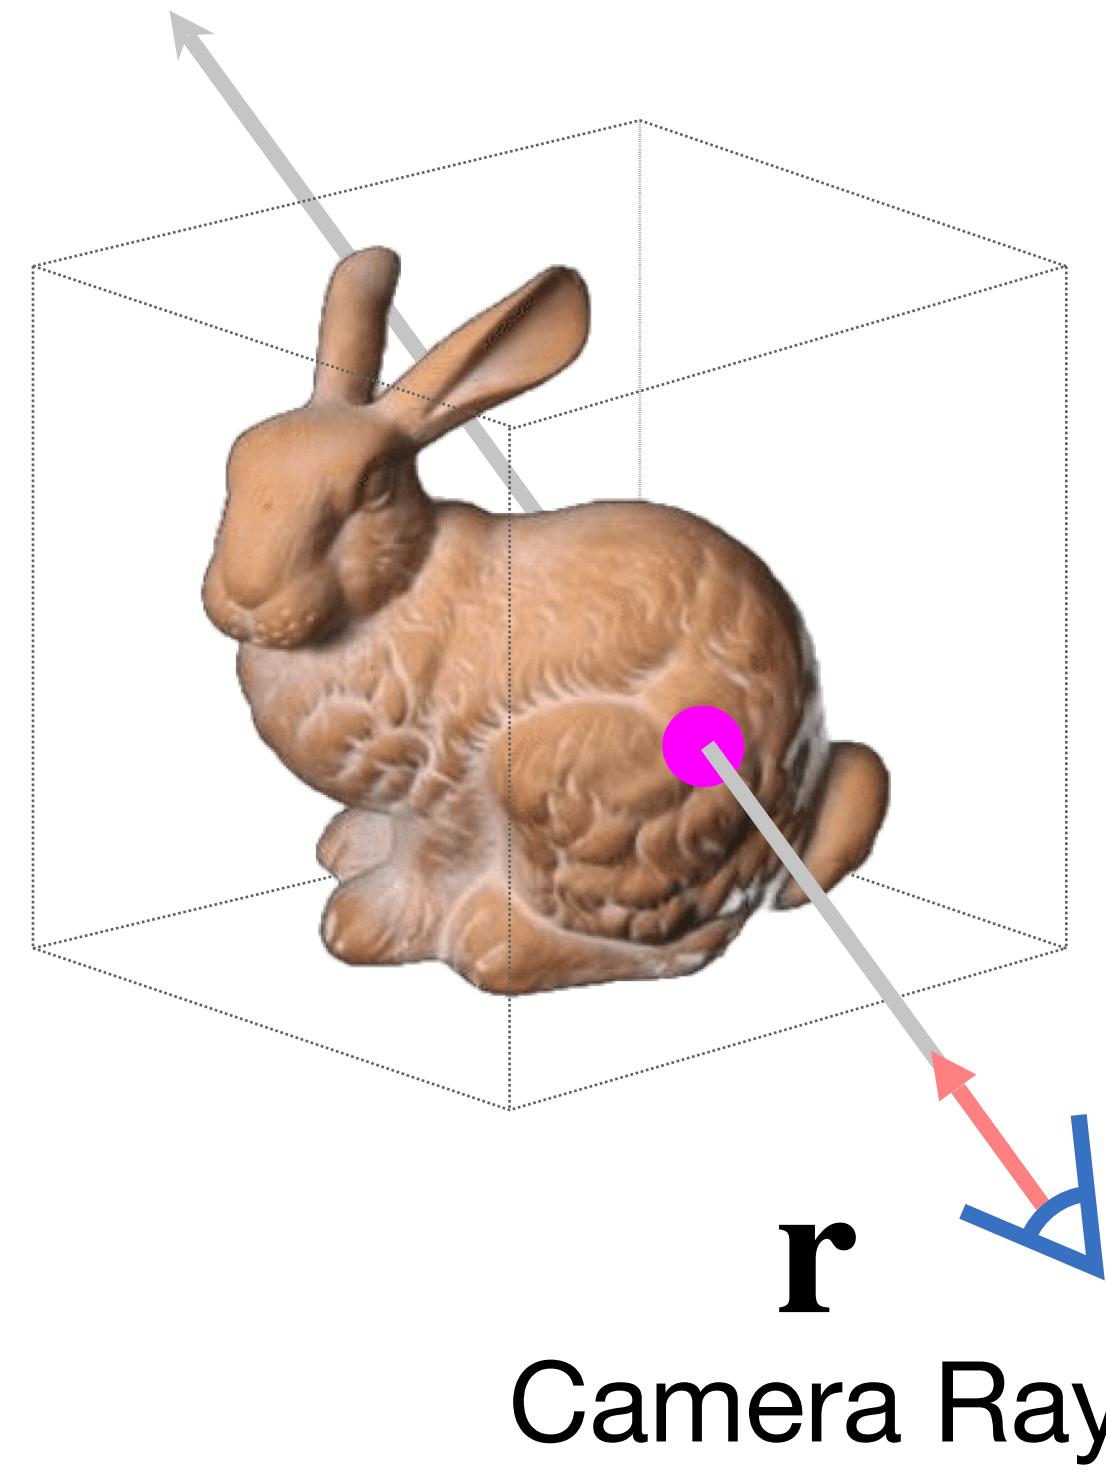
$$\Phi^{3D} : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



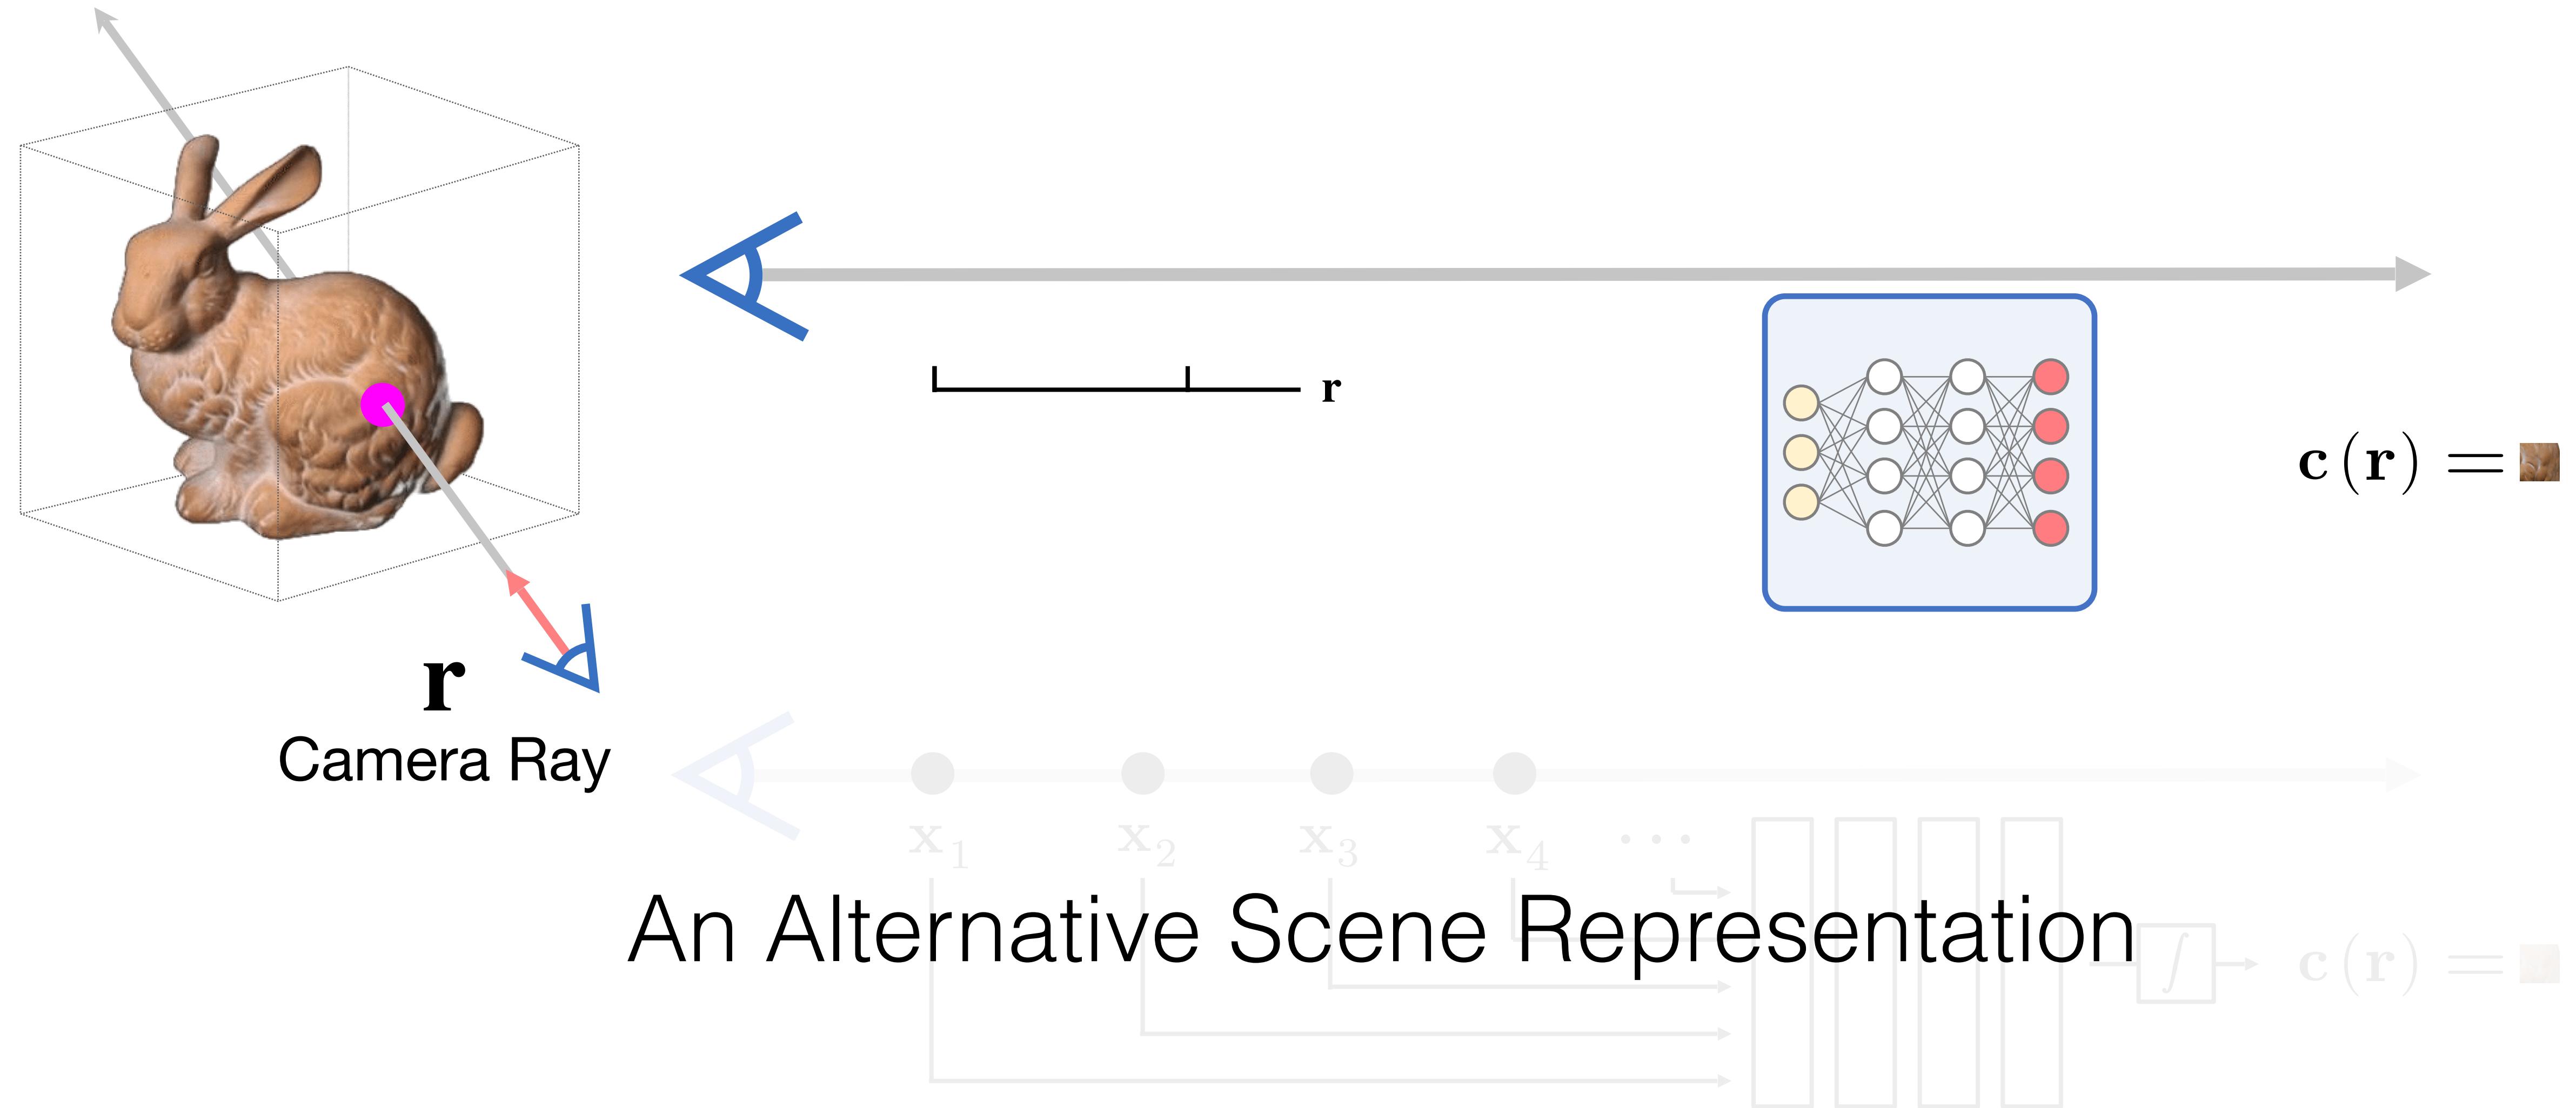
$$c(r) =$$



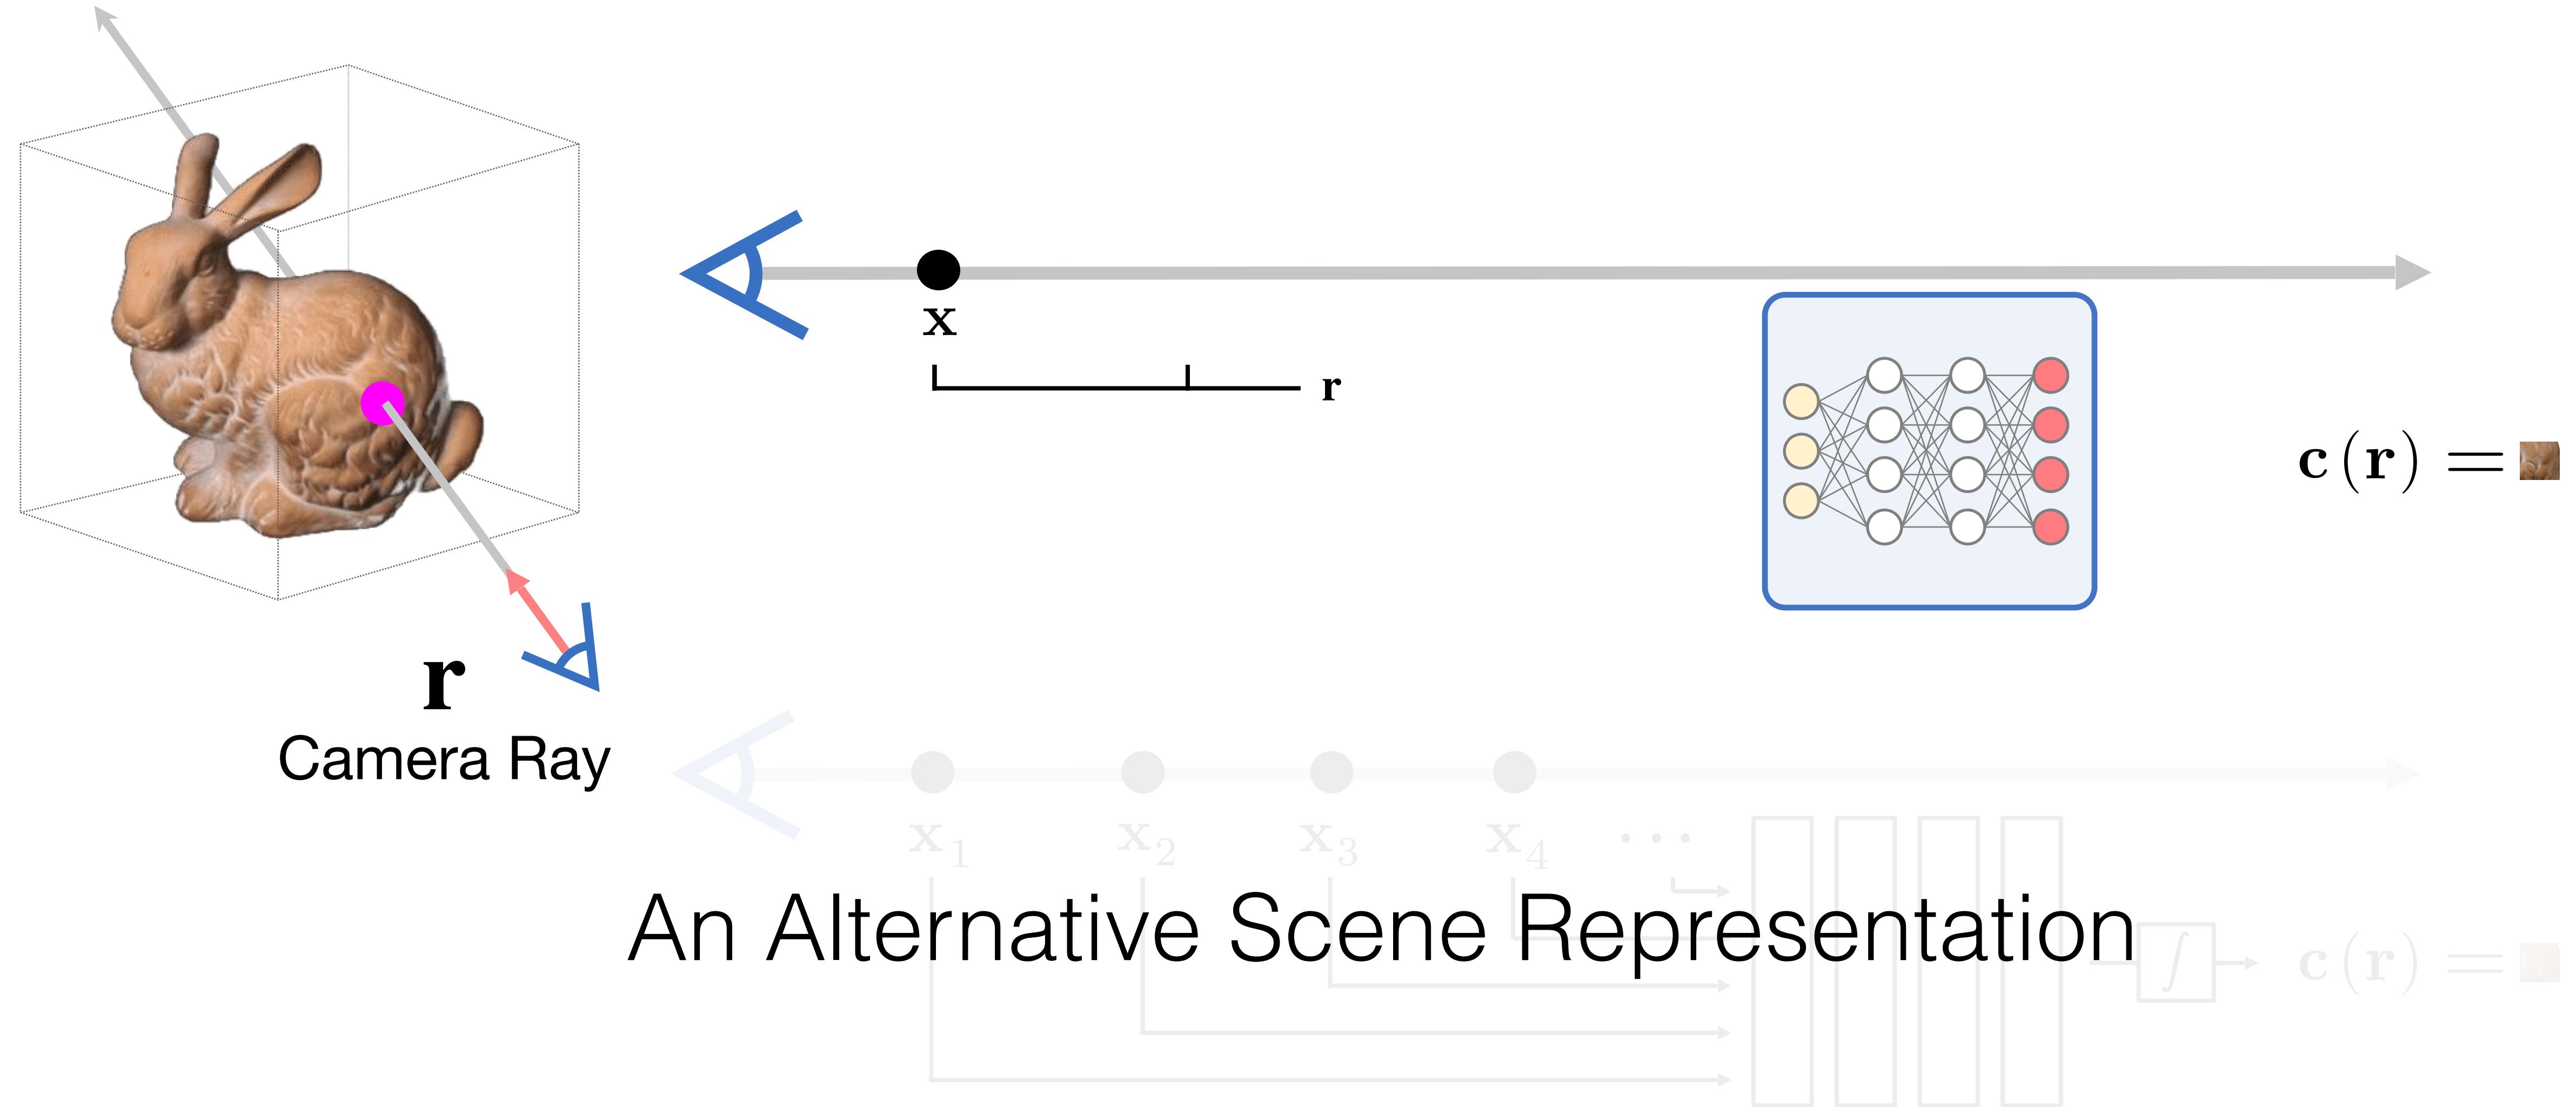
Light Field Networks



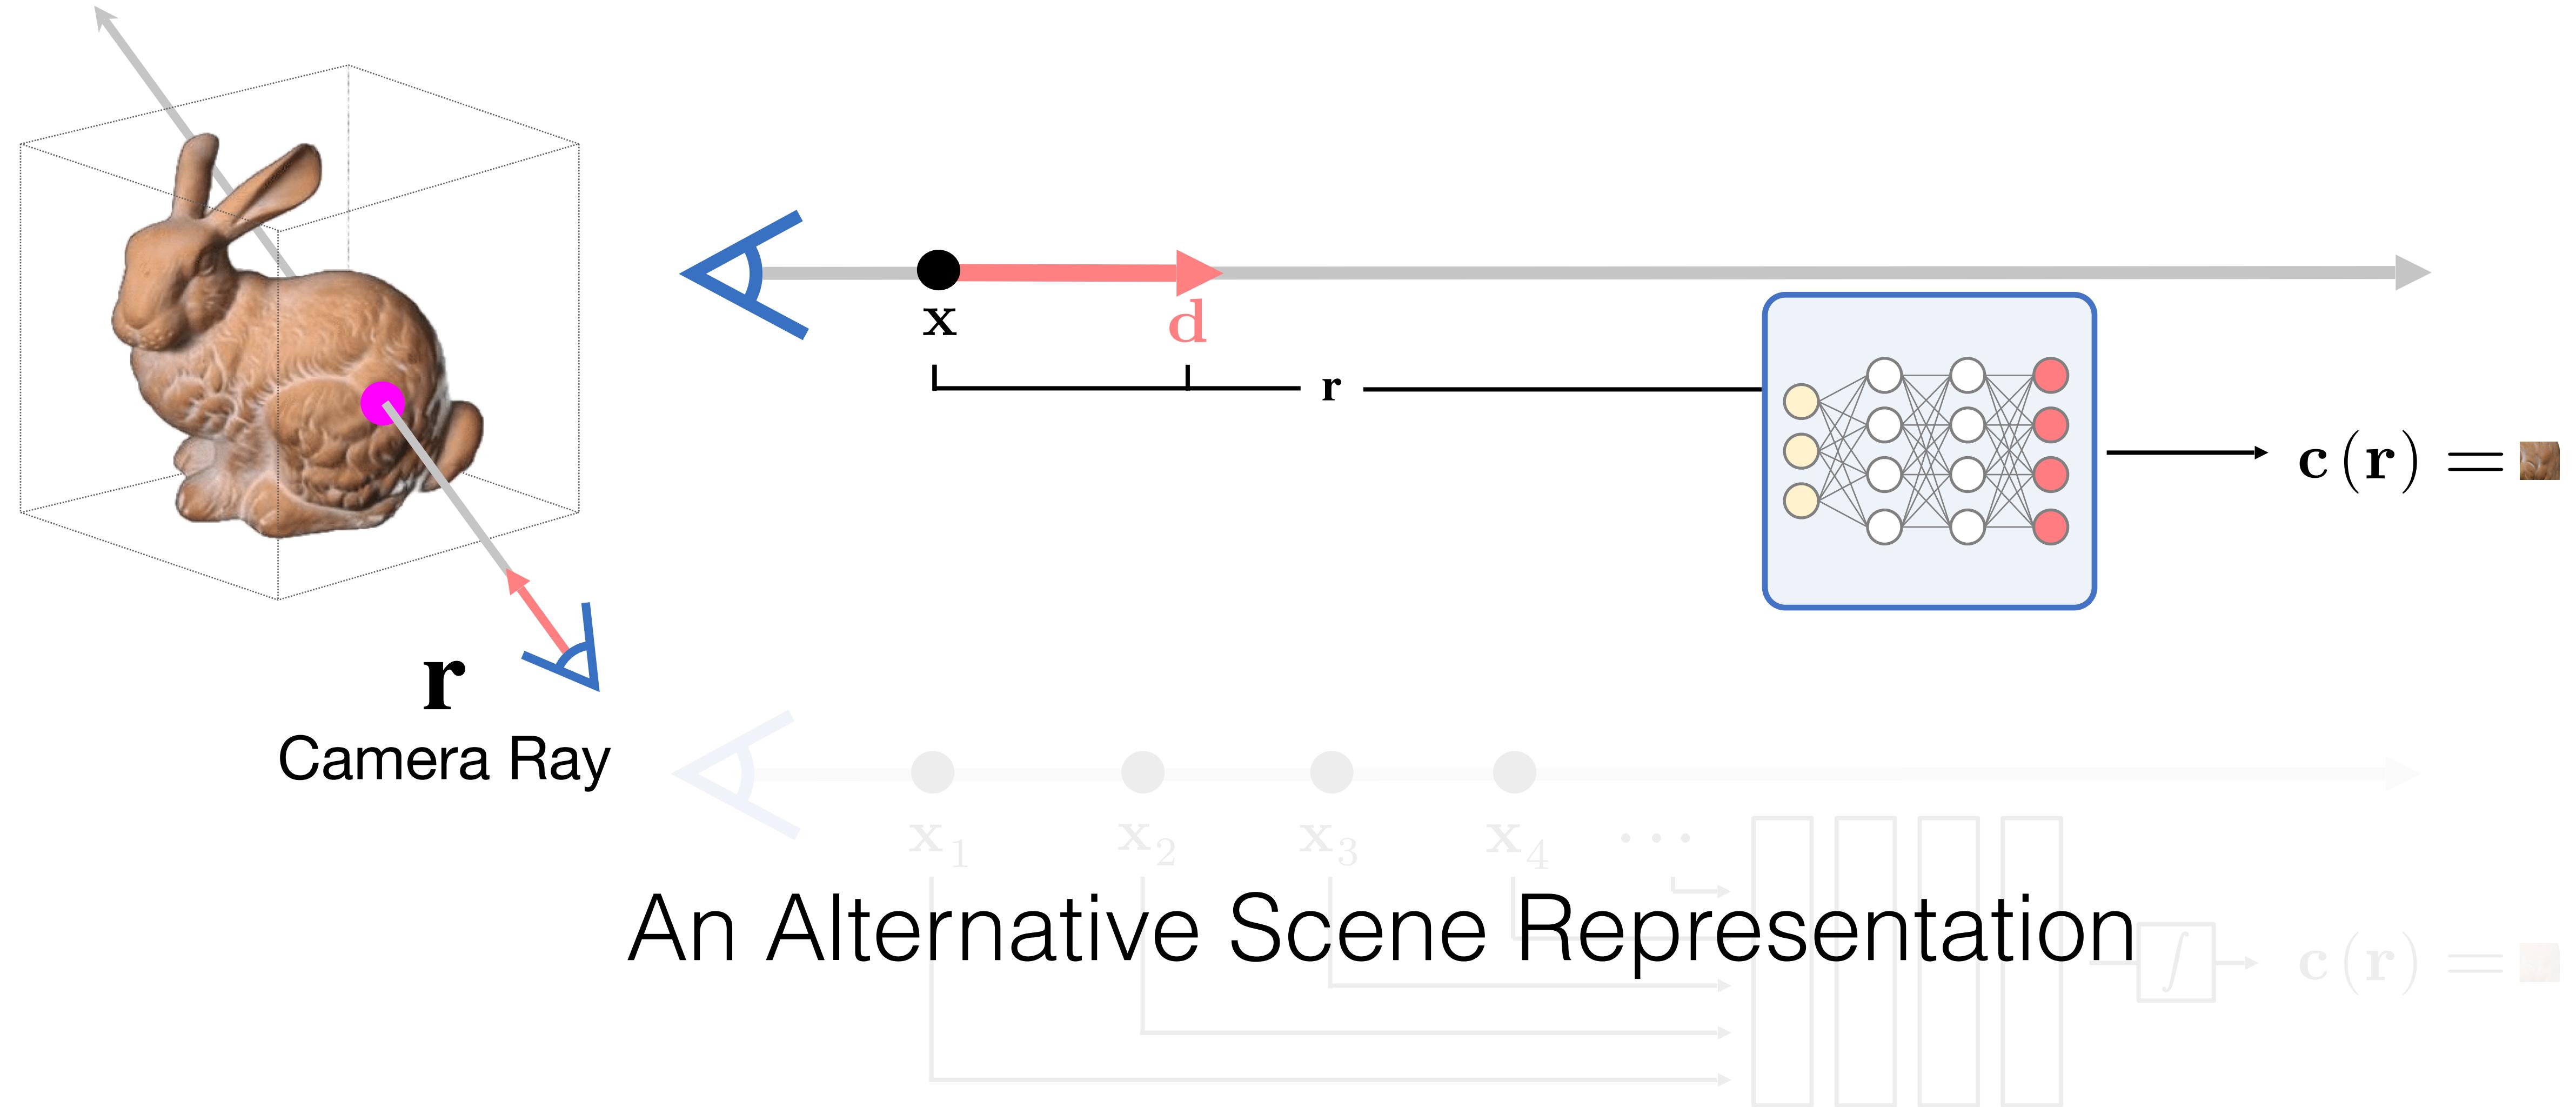
Light Field Networks



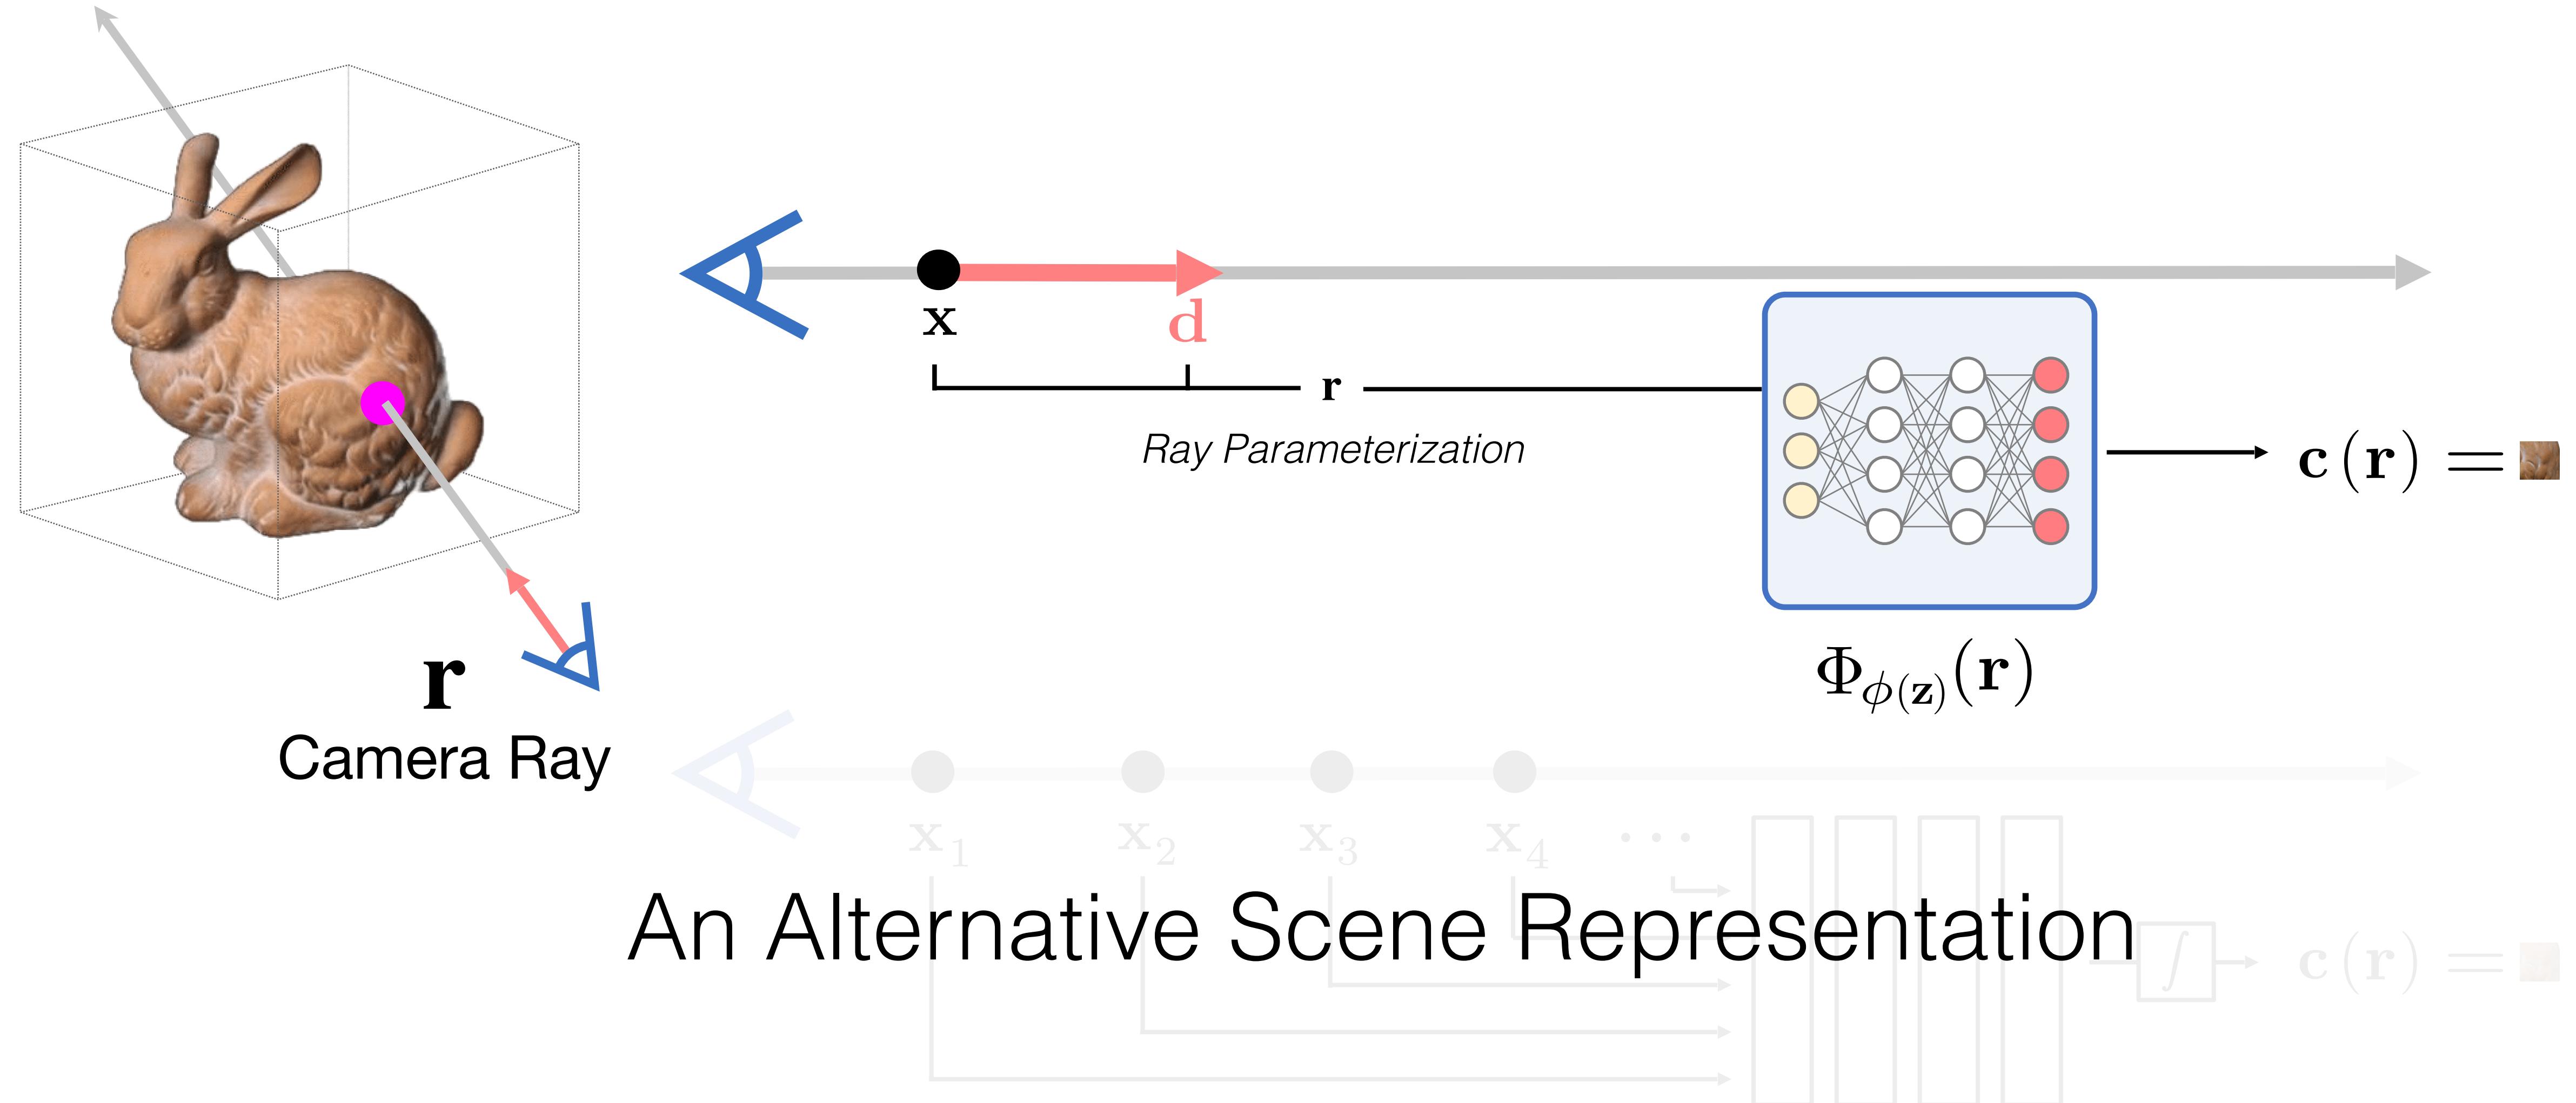
Light Field Networks



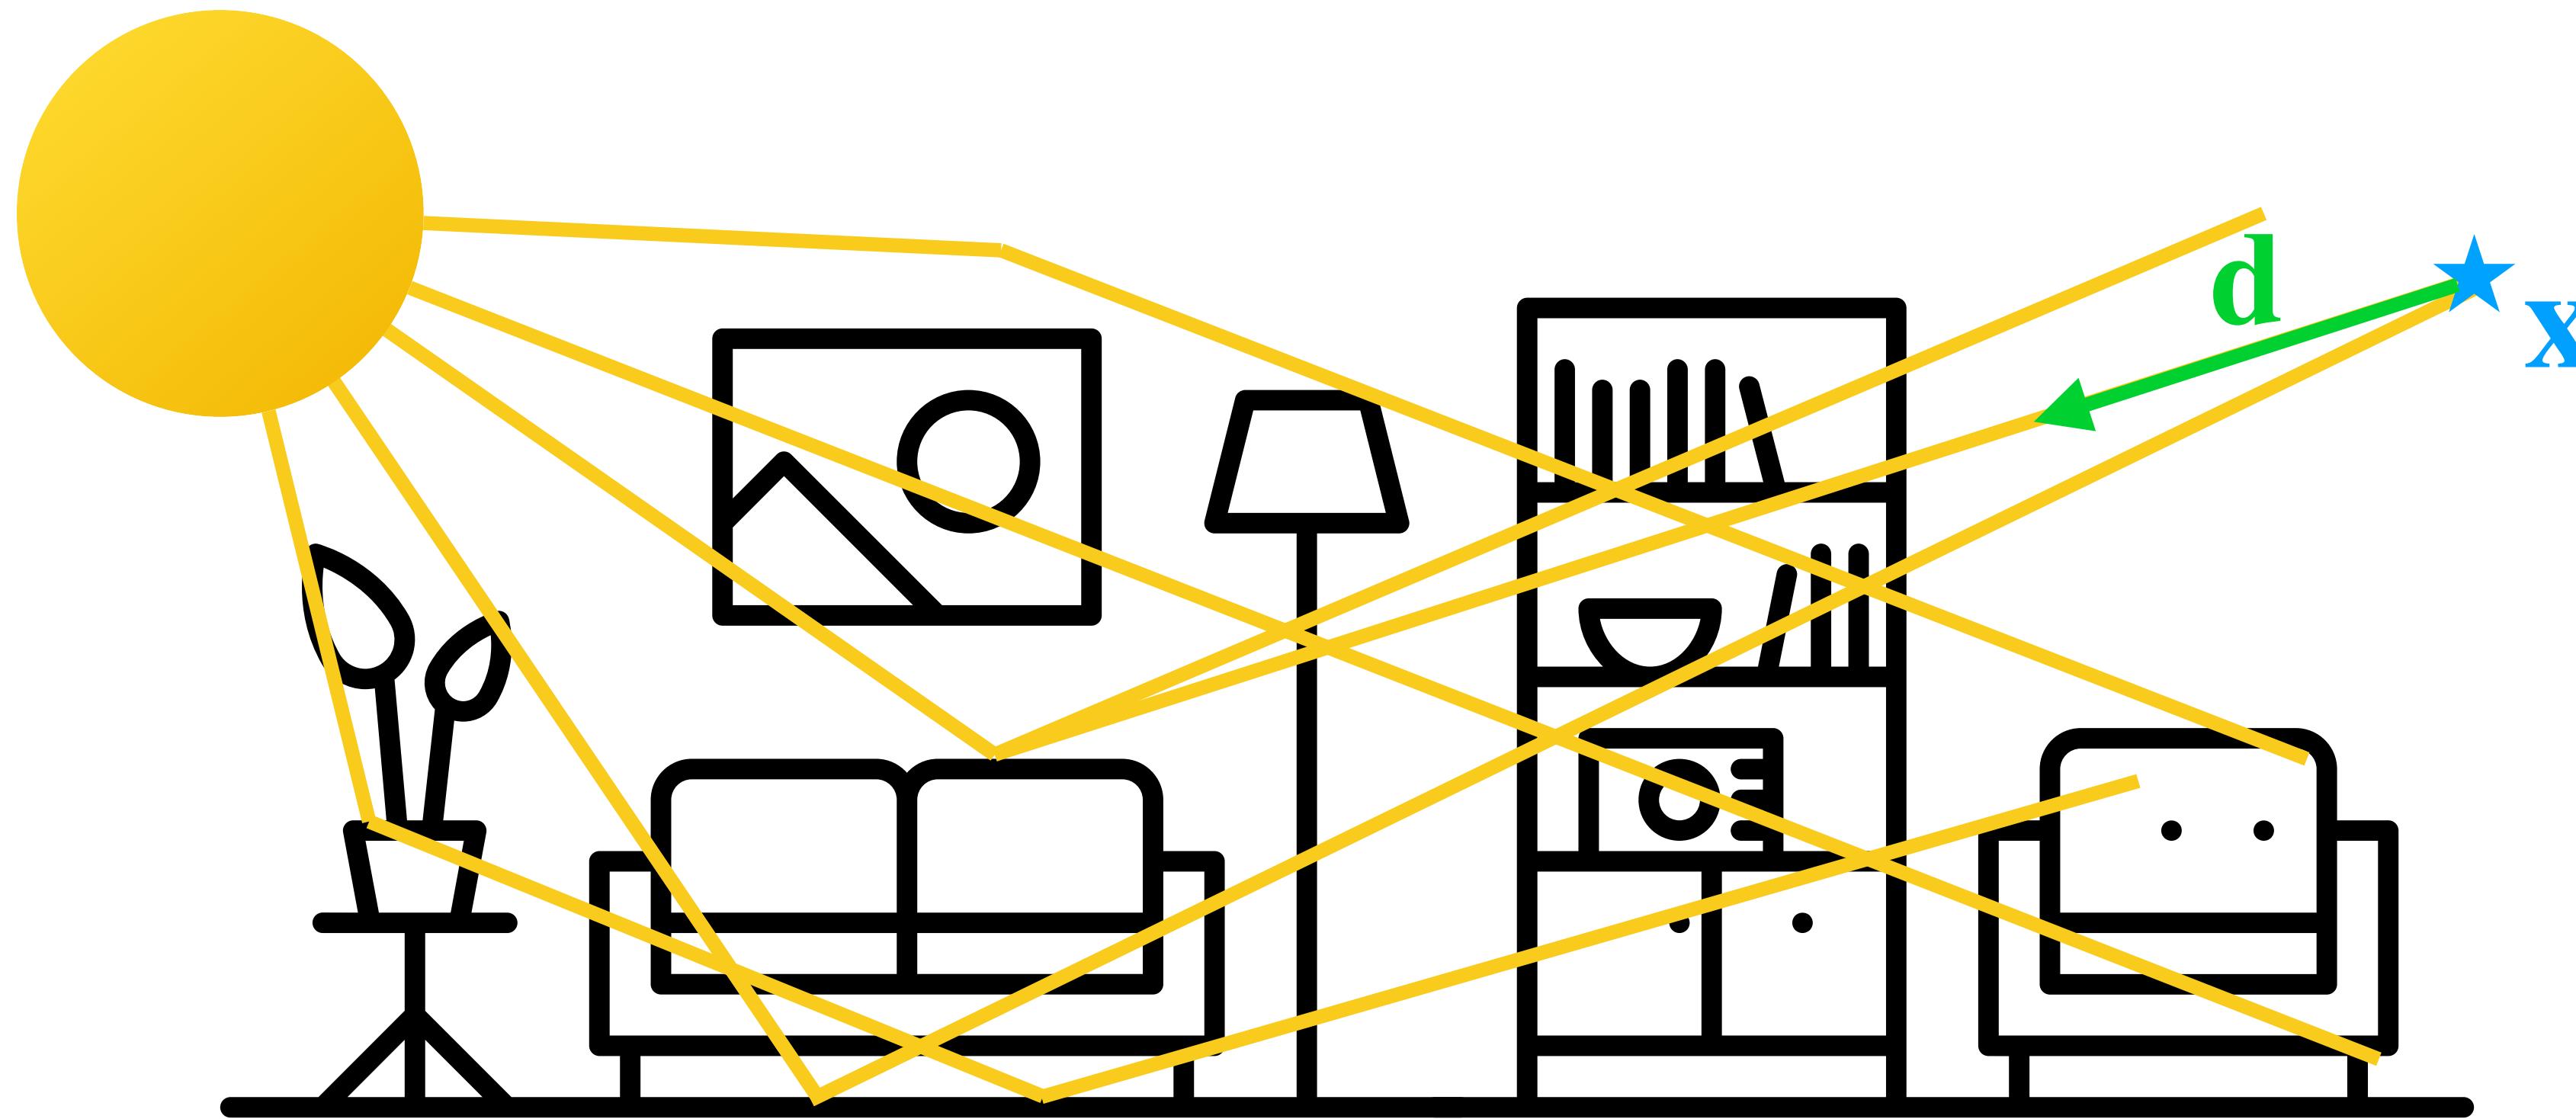
Light Field Networks



Light Field Networks



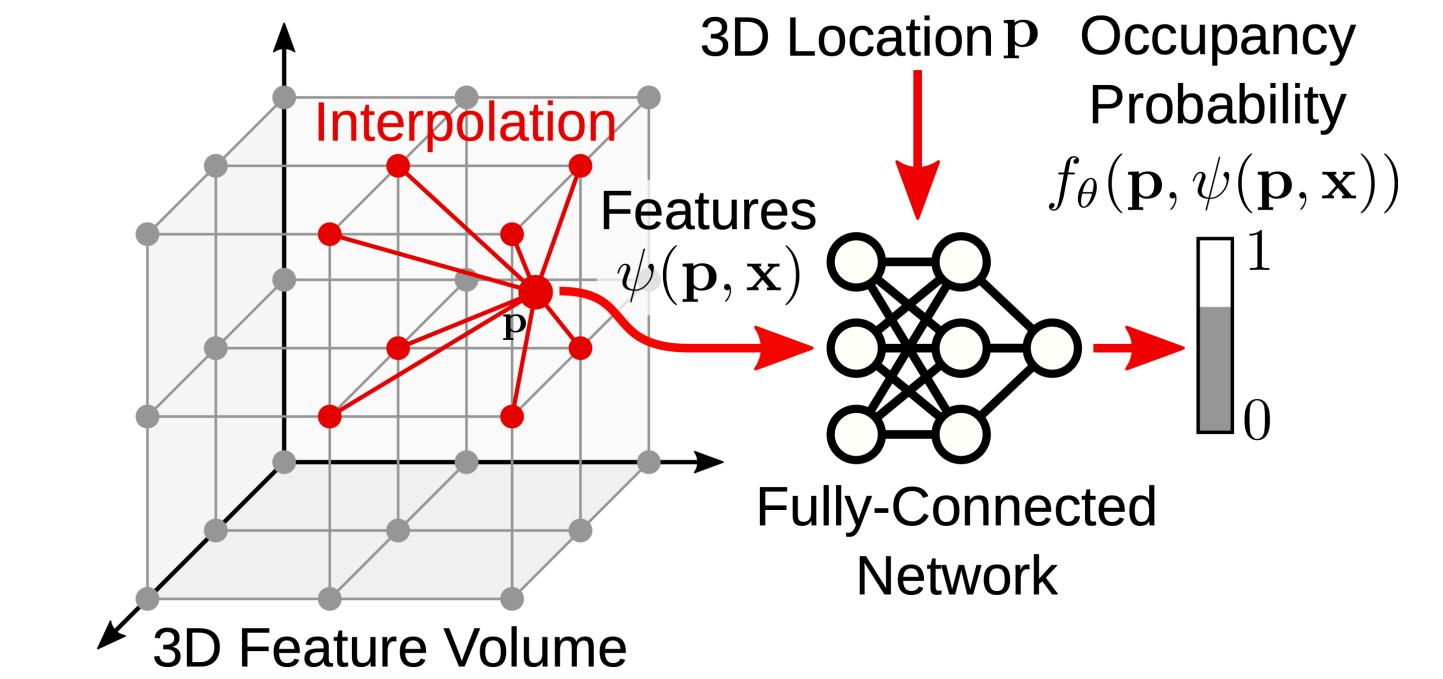
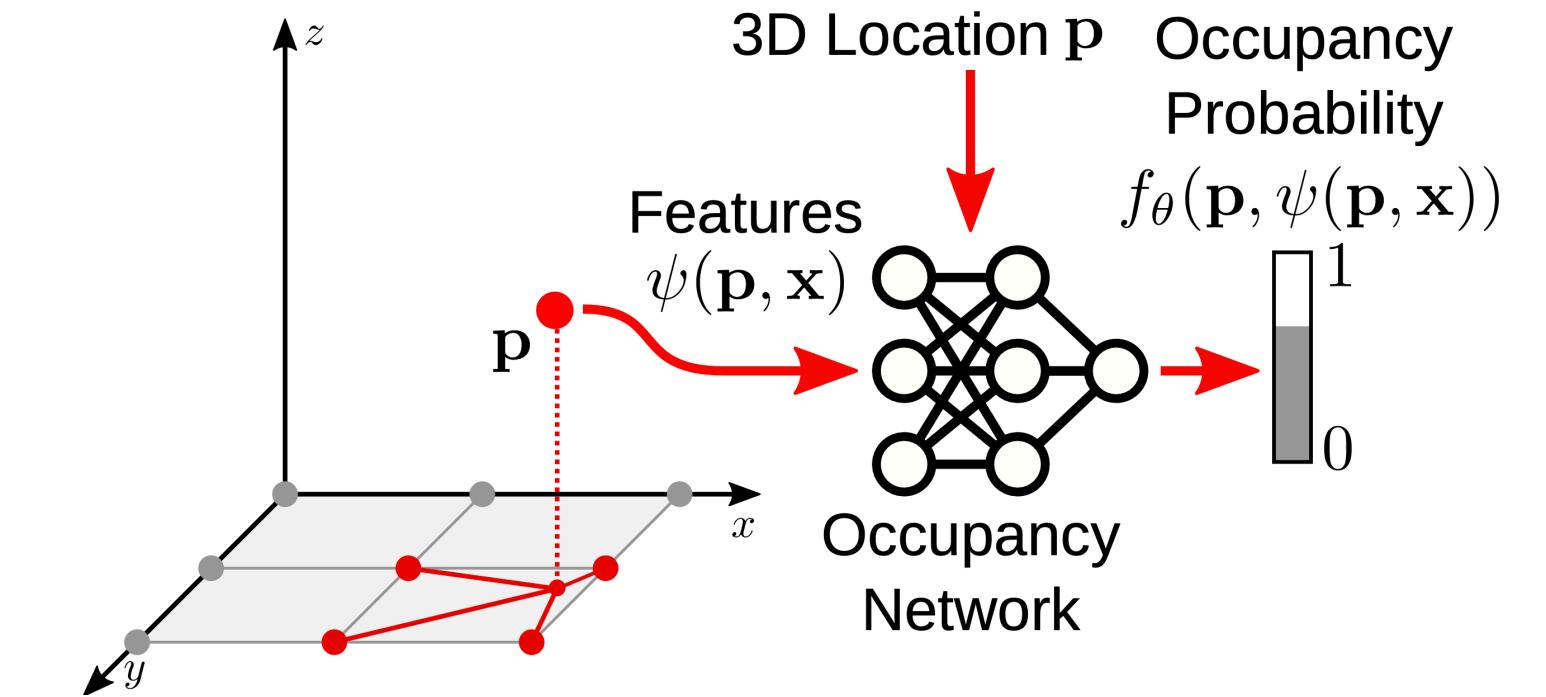
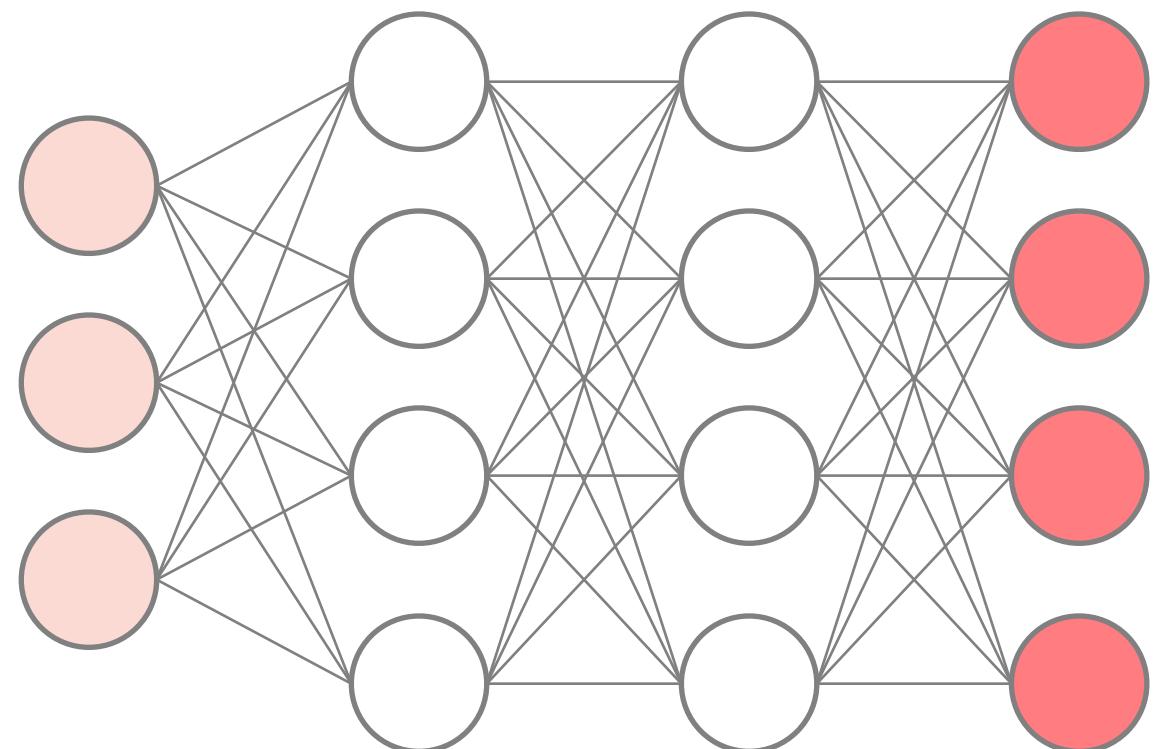
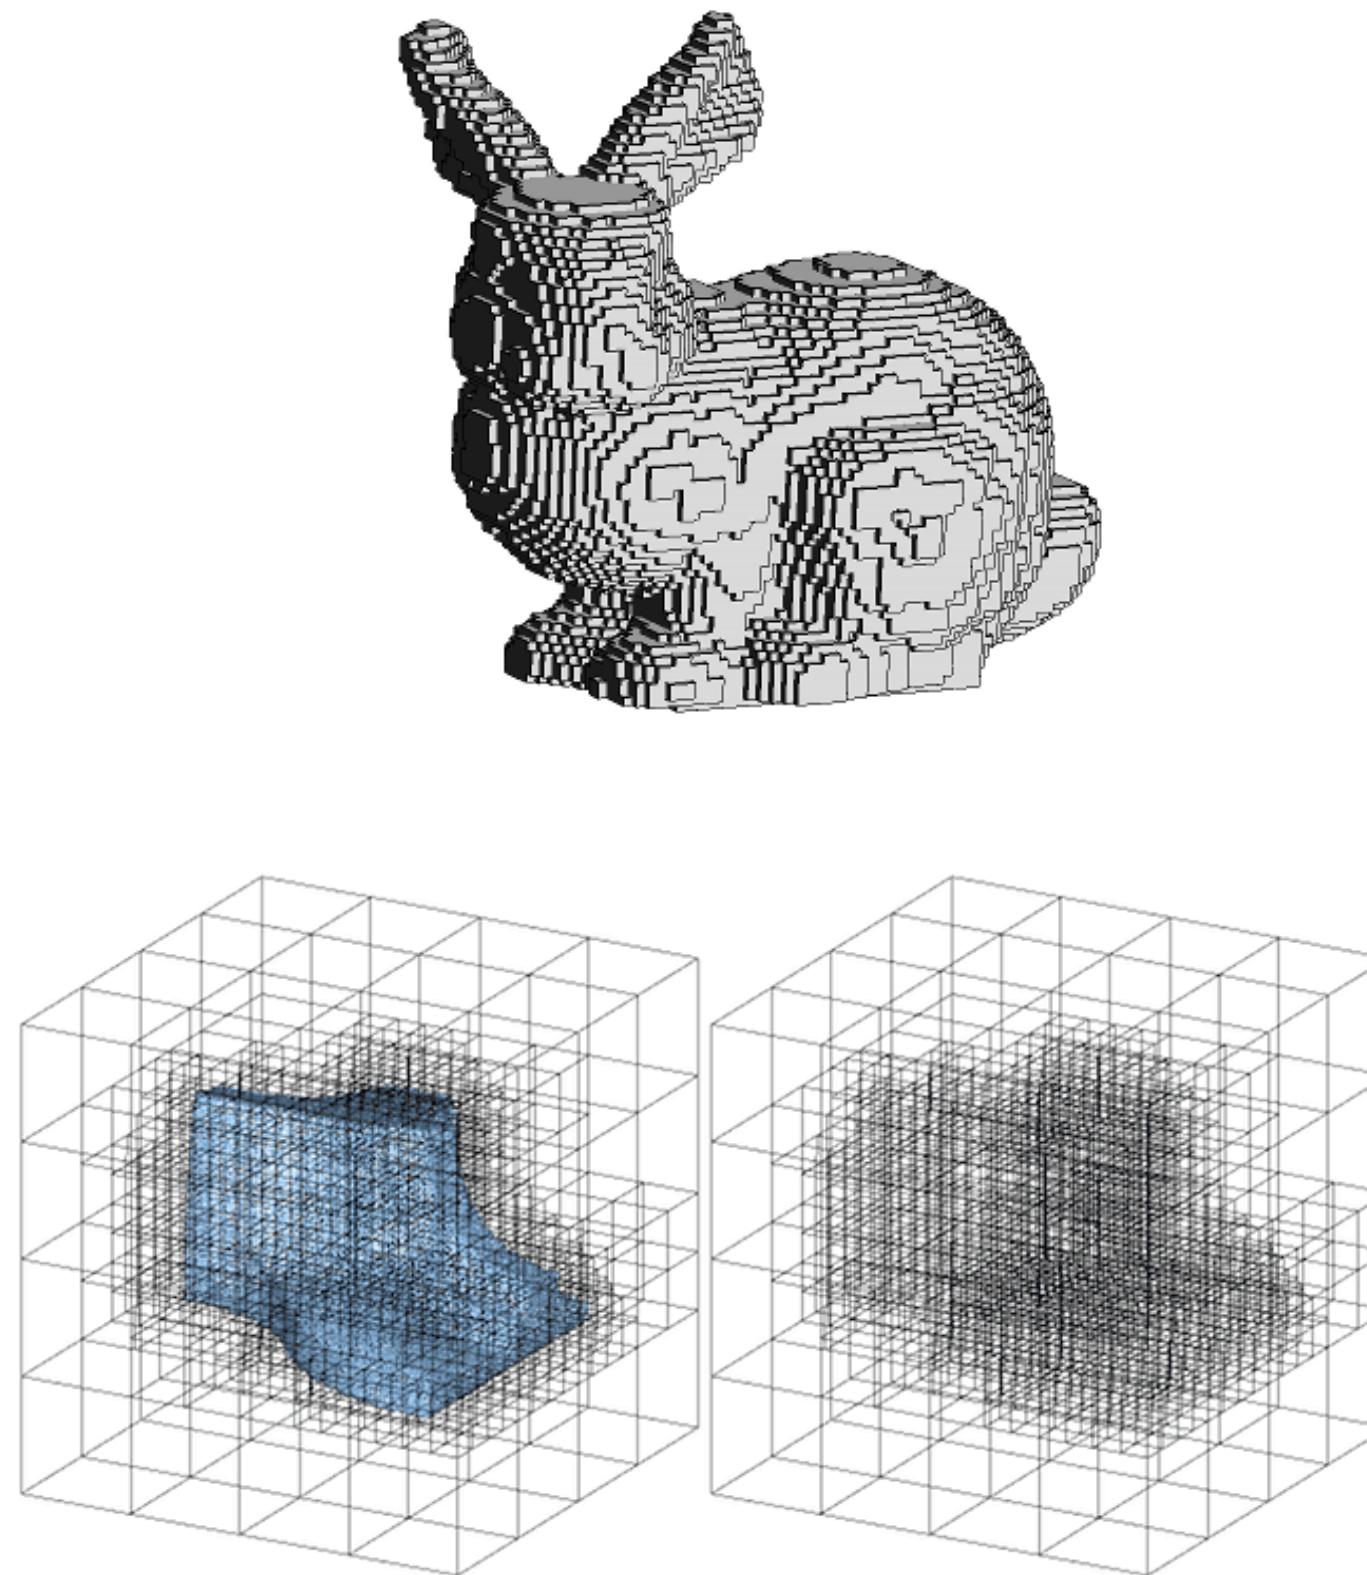
Just because Scene is 3D, don't have to store information in 3D data structure...



$$LF : \mathbb{R}^3 \times \mathbb{S}^2 \rightarrow \mathbb{R}^3, \quad LF(\mathbf{x}, \mathbf{d}) = \mathbf{c}$$

Light Field is 5-dimensional! What parameterizations are possible...?

Summary



- No “Artificial Intelligence”, “understanding” or “Machine Learning” when fitting a single representation to a single 3D thing, even if there are neural networks in them.
- It becomes *a lot* more interesting when considering *inference*, i.e., we want representation to be output of an encoder! More on that later in “Prior-based reconstruction”.

3D Scene

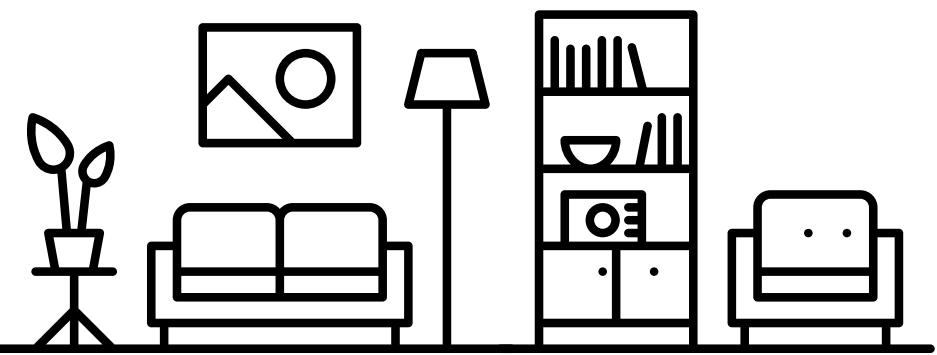


Image
Formation

Graphics

3D Scene

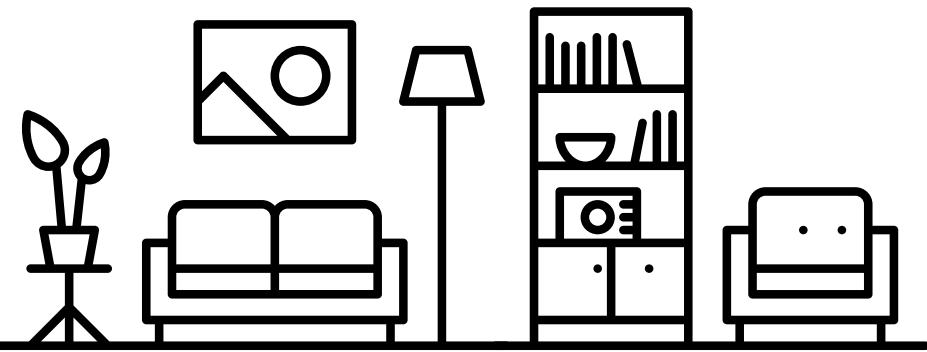


Image
Formation

Neural Scene
Representation

Graphics

3D Scene

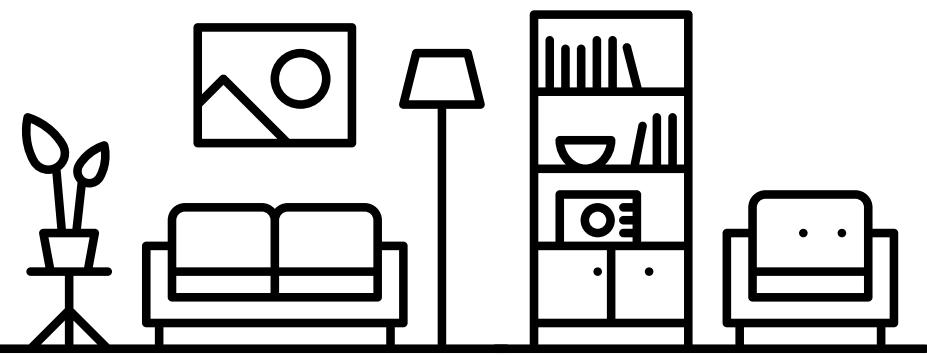


Image
Formation

Inference

Neural Scene
Representation

Graphics

3D Scene

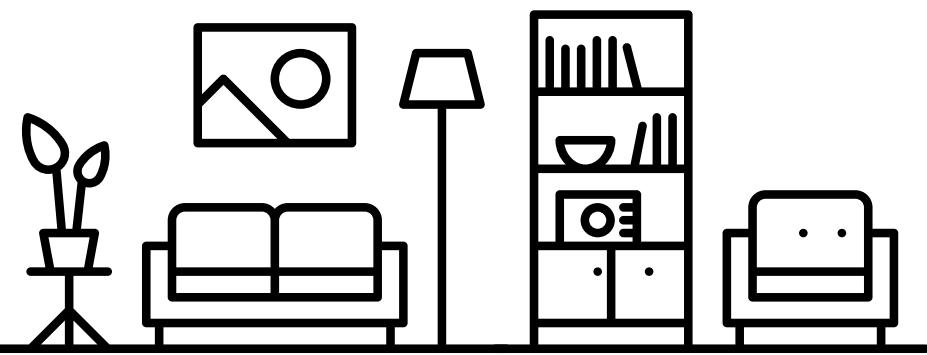


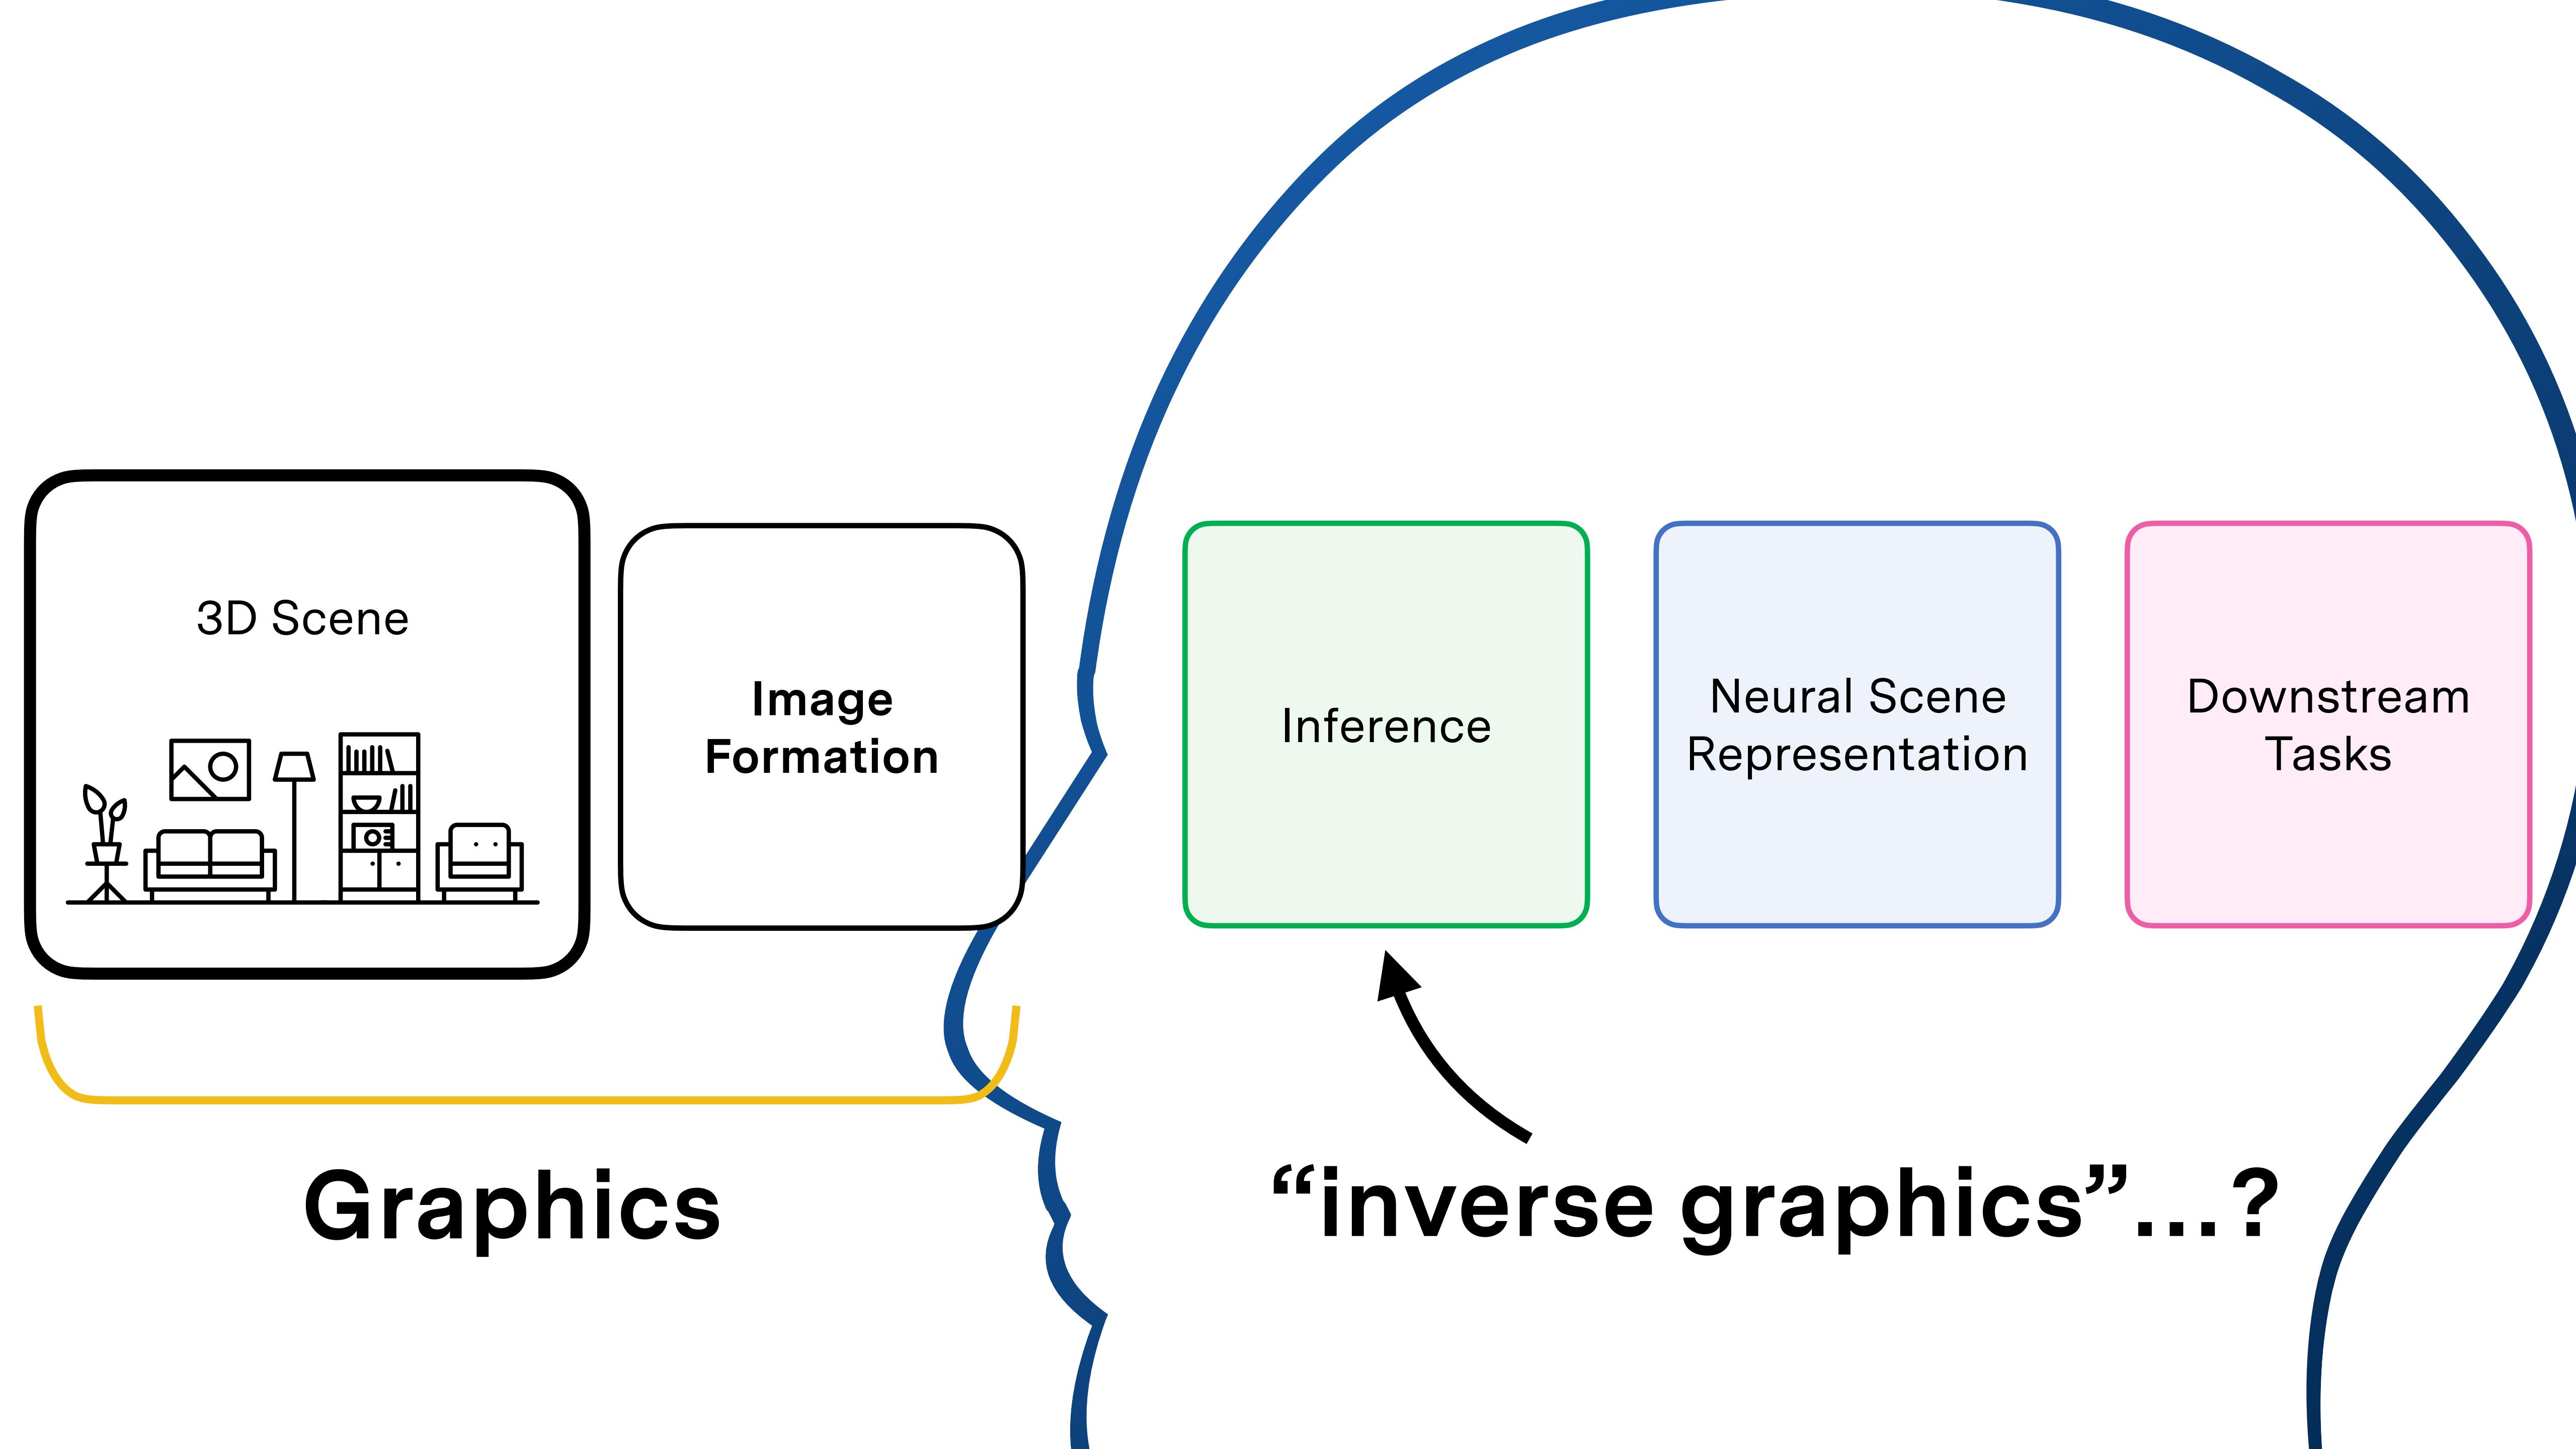
Image
Formation

Inference

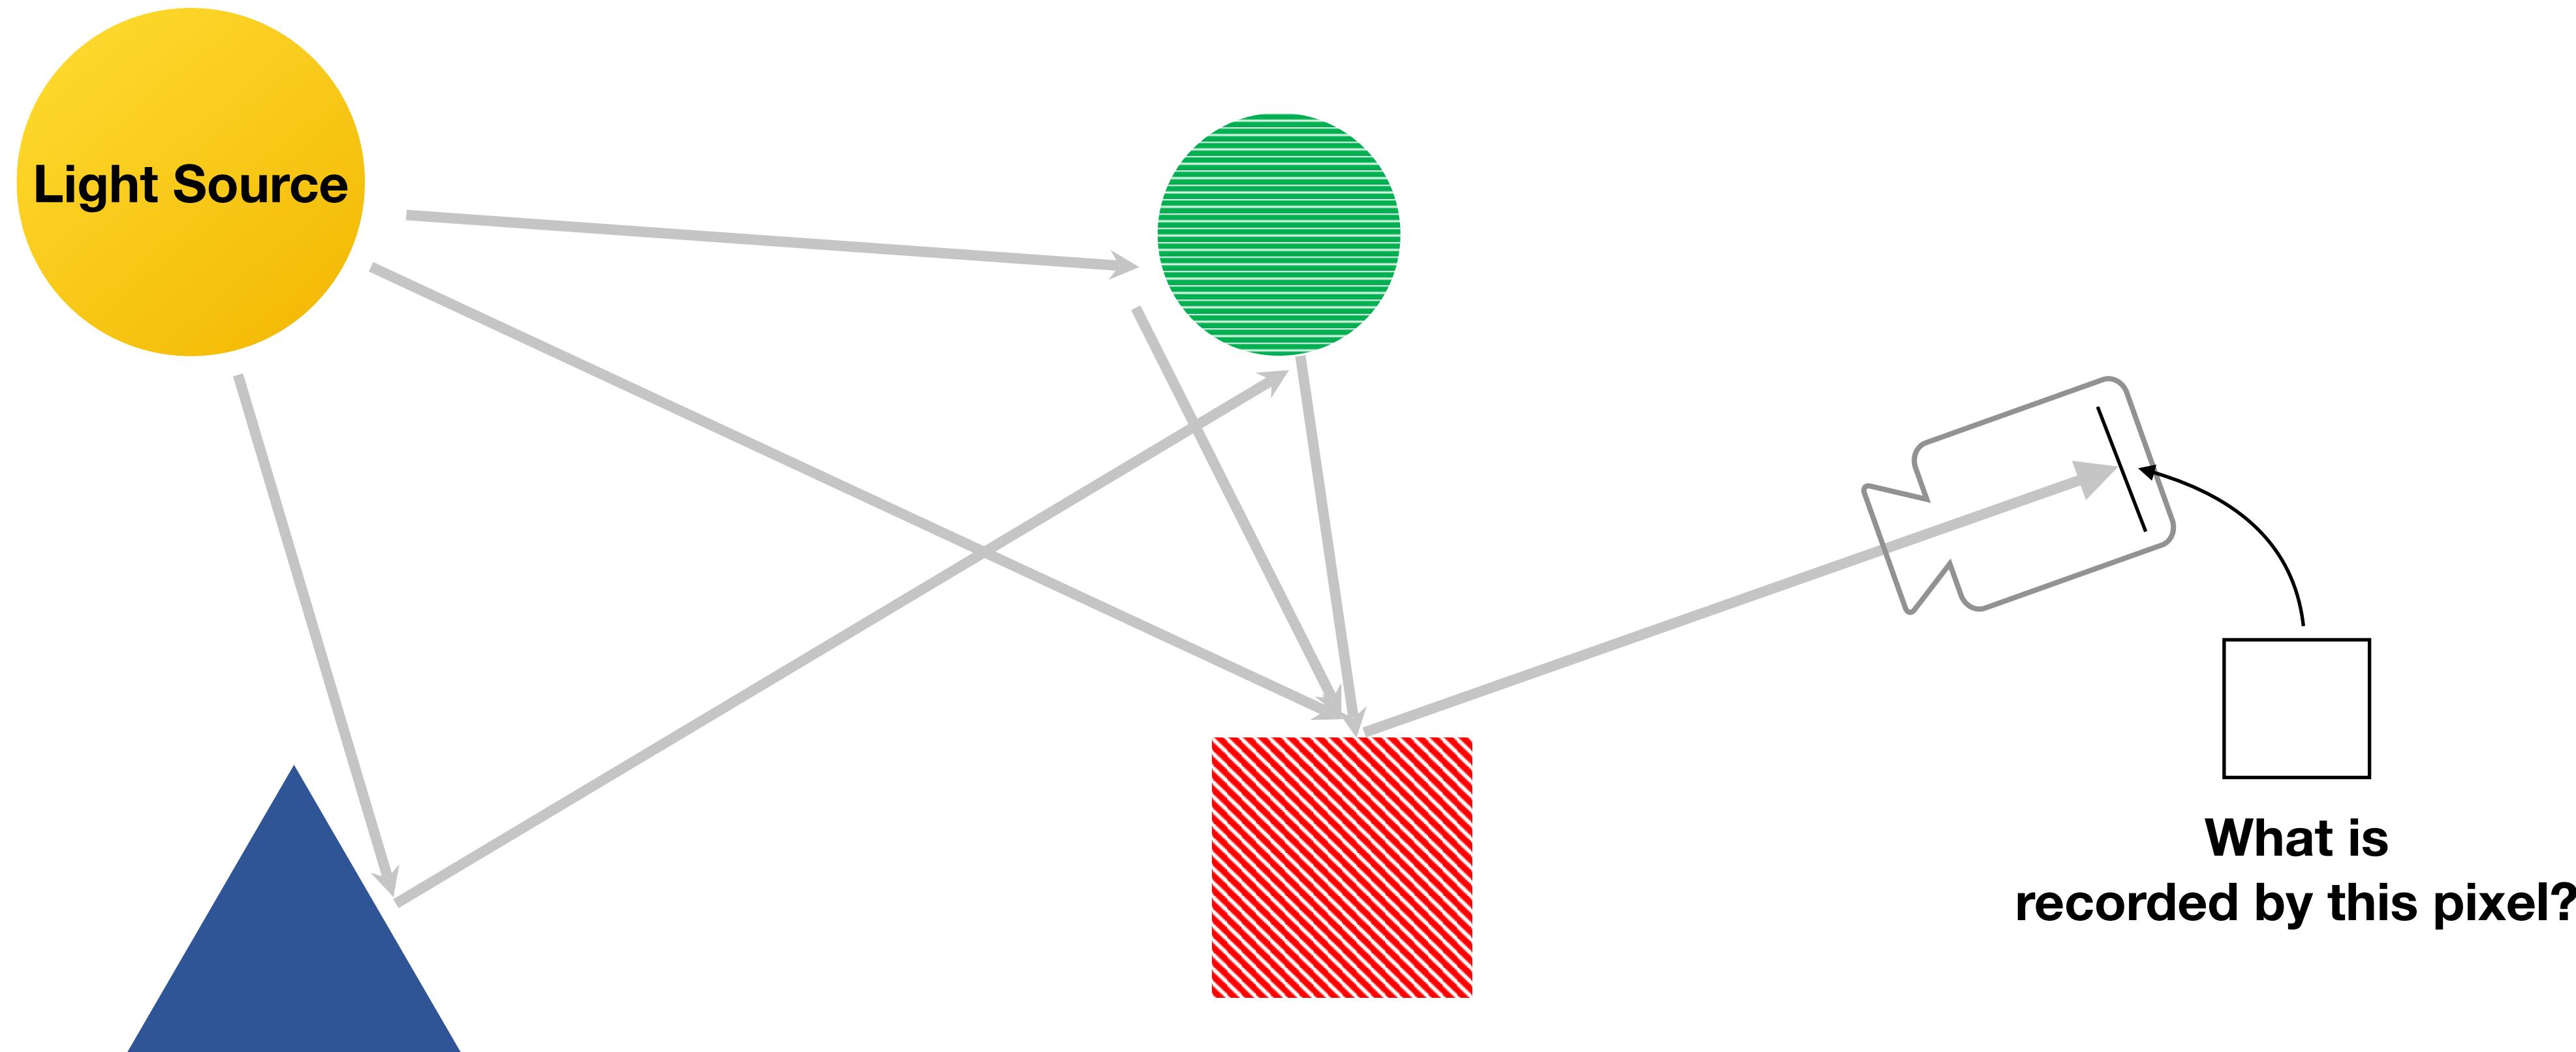
Neural Scene
Representation

Graphics

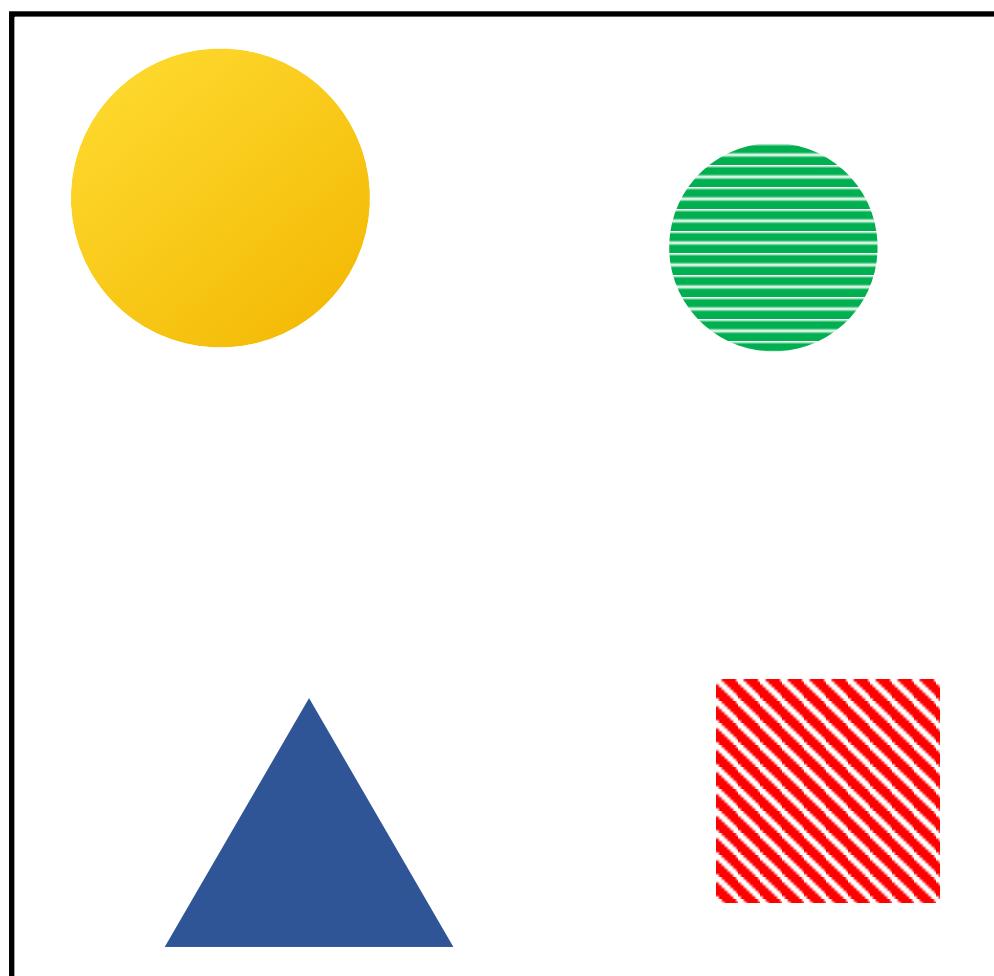
“inverse graphics”...?



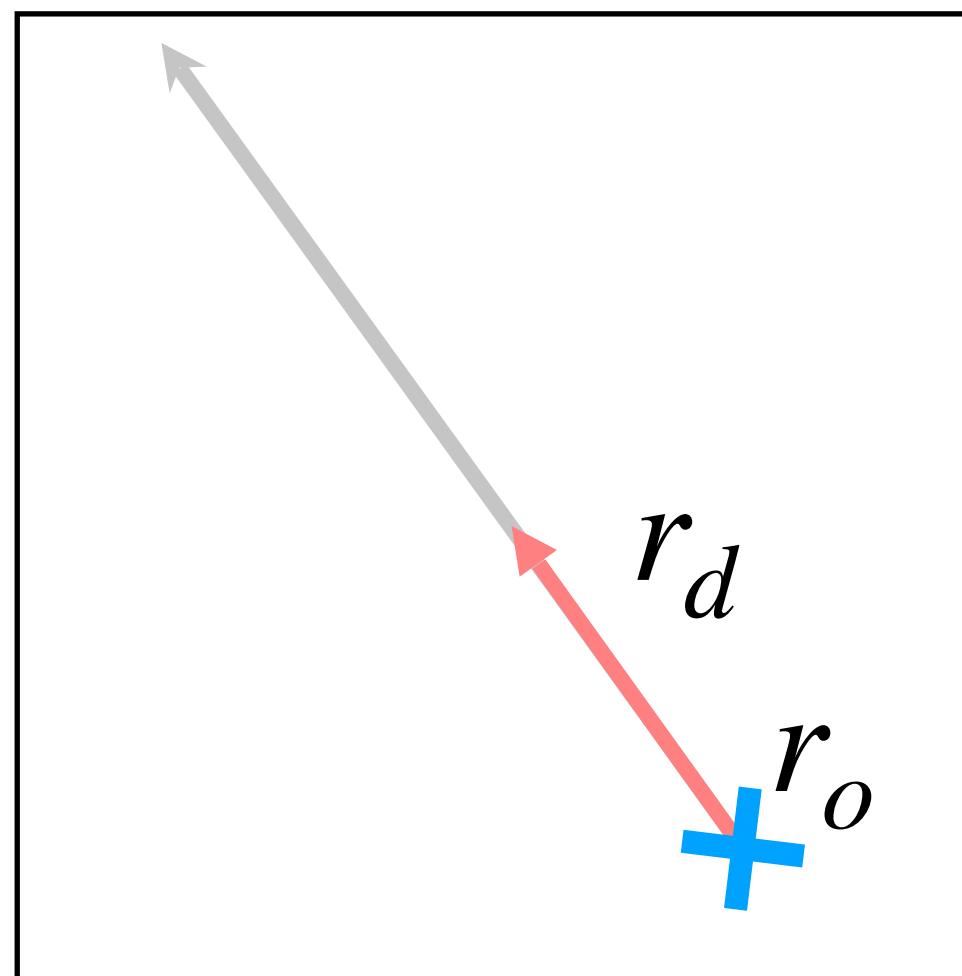
Light Transport & Cameras



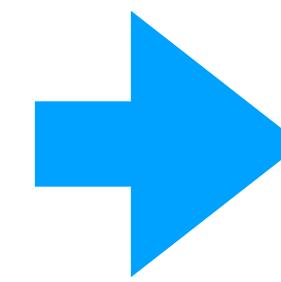
Function Signature of Rendering



Scene : Light sources,
materials, shapes, ...



Camera Ray



Radiance of the ray

The Rendering equation

The moral of the story:

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Radiance

Radiance

- Radiance is a fundamental quantity of light.

Radiance

- Radiance is a fundamental quantity of light.
- Radiance is density of photons at a point, traveling in the same direction, at given wavelength.

Radiance

- Radiance is a fundamental quantity of light.
- Radiance is density of photons at a point, traveling in the same direction, at given wavelength.
- Radiance is energy per time, per unit area, per unit solid angle, per wavelength.

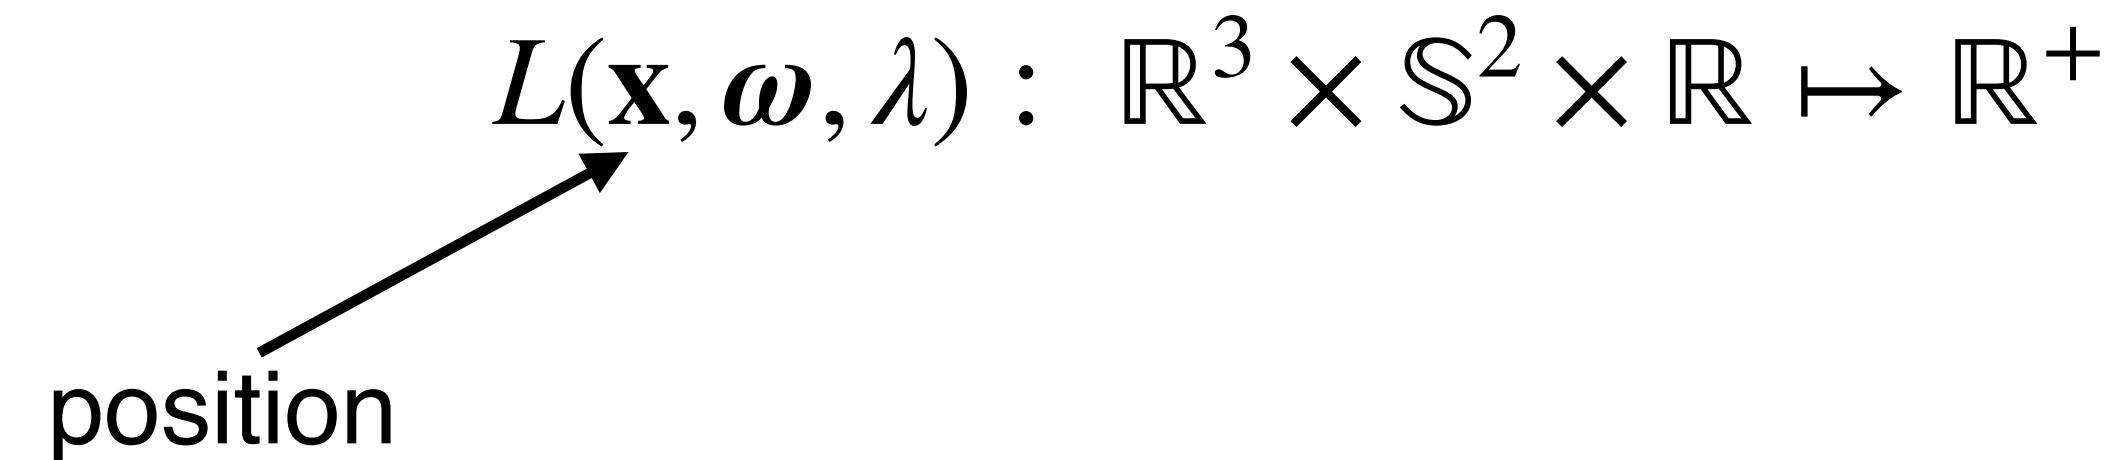
$$L(\mathbf{x}, \boldsymbol{\omega}, \lambda) : \mathbb{R}^3 \times \mathbb{S}^2 \times \mathbb{R} \mapsto \mathbb{R}^+$$

Radiance

- Radiance is a fundamental quantity of light.
- Radiance is density of photons at a point, traveling in the same direction, at given wavelength.
- Radiance is energy per time, per unit area, per unit solid angle, per wavelength.

$$L(\mathbf{x}, \omega, \lambda) : \mathbb{R}^3 \times \mathbb{S}^2 \times \mathbb{R} \mapsto \mathbb{R}^+$$

position



Radiance

- Radiance is a fundamental quantity of light.
- Radiance is density of photons at a point, traveling in the same direction, at given wavelength.
- Radiance is energy per time, per unit area, per unit solid angle, per wavelength.

$$L(\mathbf{x}, \omega, \lambda) : \mathbb{R}^3 \times \mathbb{S}^2 \times \mathbb{R} \mapsto \mathbb{R}^+$$

The diagram shows the function $L(\mathbf{x}, \omega, \lambda)$ with two arrows pointing upwards from the words "position" and "direction" to the first two arguments of the function respectively. The word "position" is positioned below the first argument \mathbf{x} , and the word "direction" is positioned below the second argument ω .

Radiance

- Radiance is a fundamental quantity of light.
- Radiance is density of photons at a point, traveling in the same direction, at given wavelength.
- Radiance is energy per time, per unit area, per unit solid angle, per wavelength.

$$L(\mathbf{x}, \omega, \lambda) : \mathbb{R}^3 \times \mathbb{S}^2 \times \mathbb{R} \mapsto \mathbb{R}^+$$

position direction wavelength

The diagram shows a mathematical function $L(\mathbf{x}, \omega, \lambda)$ with three input variables. The first variable, "position", is represented by a horizontal arrow pointing towards the function. The second variable, "direction", is represented by a vertical arrow pointing upwards towards the function. The third variable, "wavelength", is represented by a diagonal arrow pointing towards the function from the right. The function itself is written as $L(\mathbf{x}, \omega, \lambda) : \mathbb{R}^3 \times \mathbb{S}^2 \times \mathbb{R} \mapsto \mathbb{R}^+$.

The “Light Field”: Radiance for every possible ray

Radiance intuition

Radiance intuition



Radiance intuition



Point Light
Source

Radiance intuition



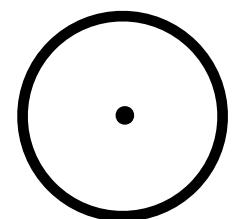
Point Light
Source



Radiance intuition



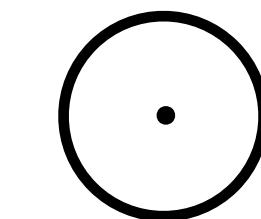
Point Light
Source



2D radiance meter



Radiance intuition

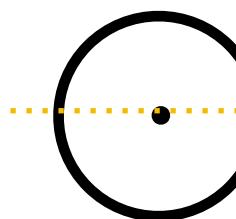


2D radiance meter (not real)

Point Light
Source



Radiance intuition

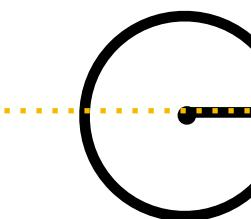


2D radiance meter (not real)

Point Light
Source



Radiance intuition



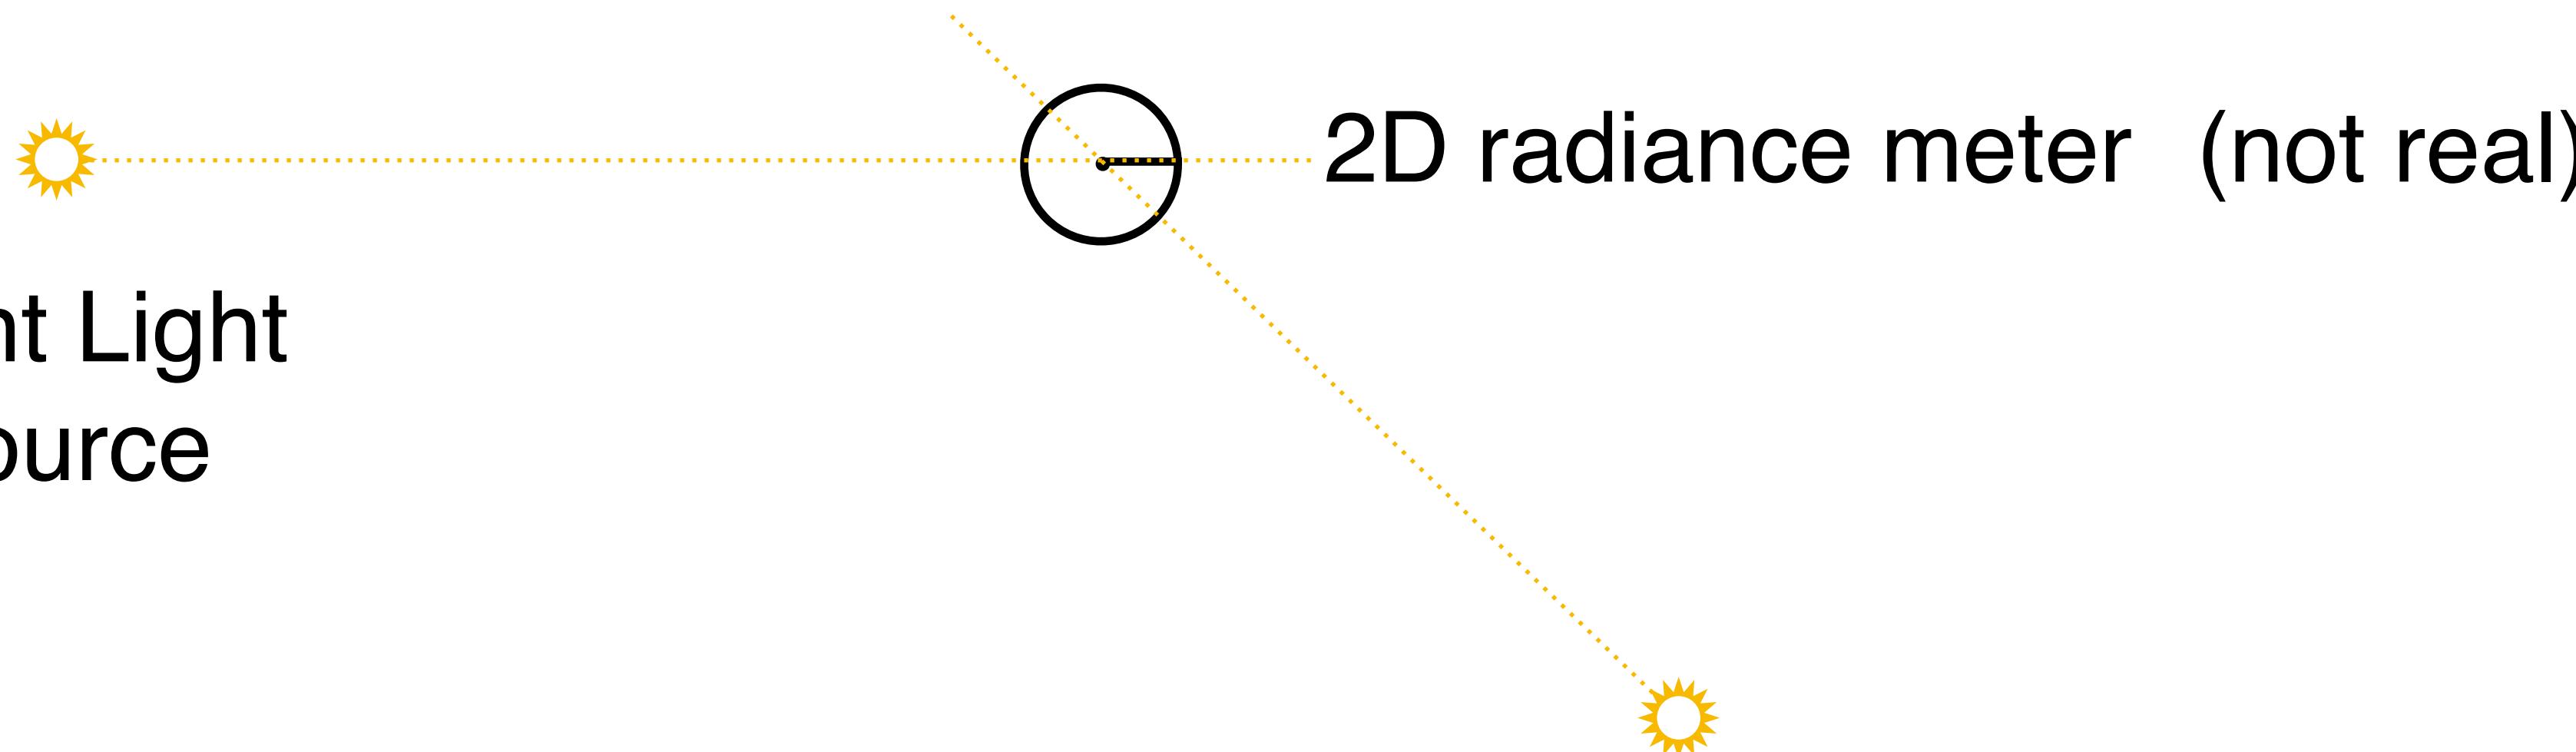
2D radiance meter (not real)

Point Light
Source



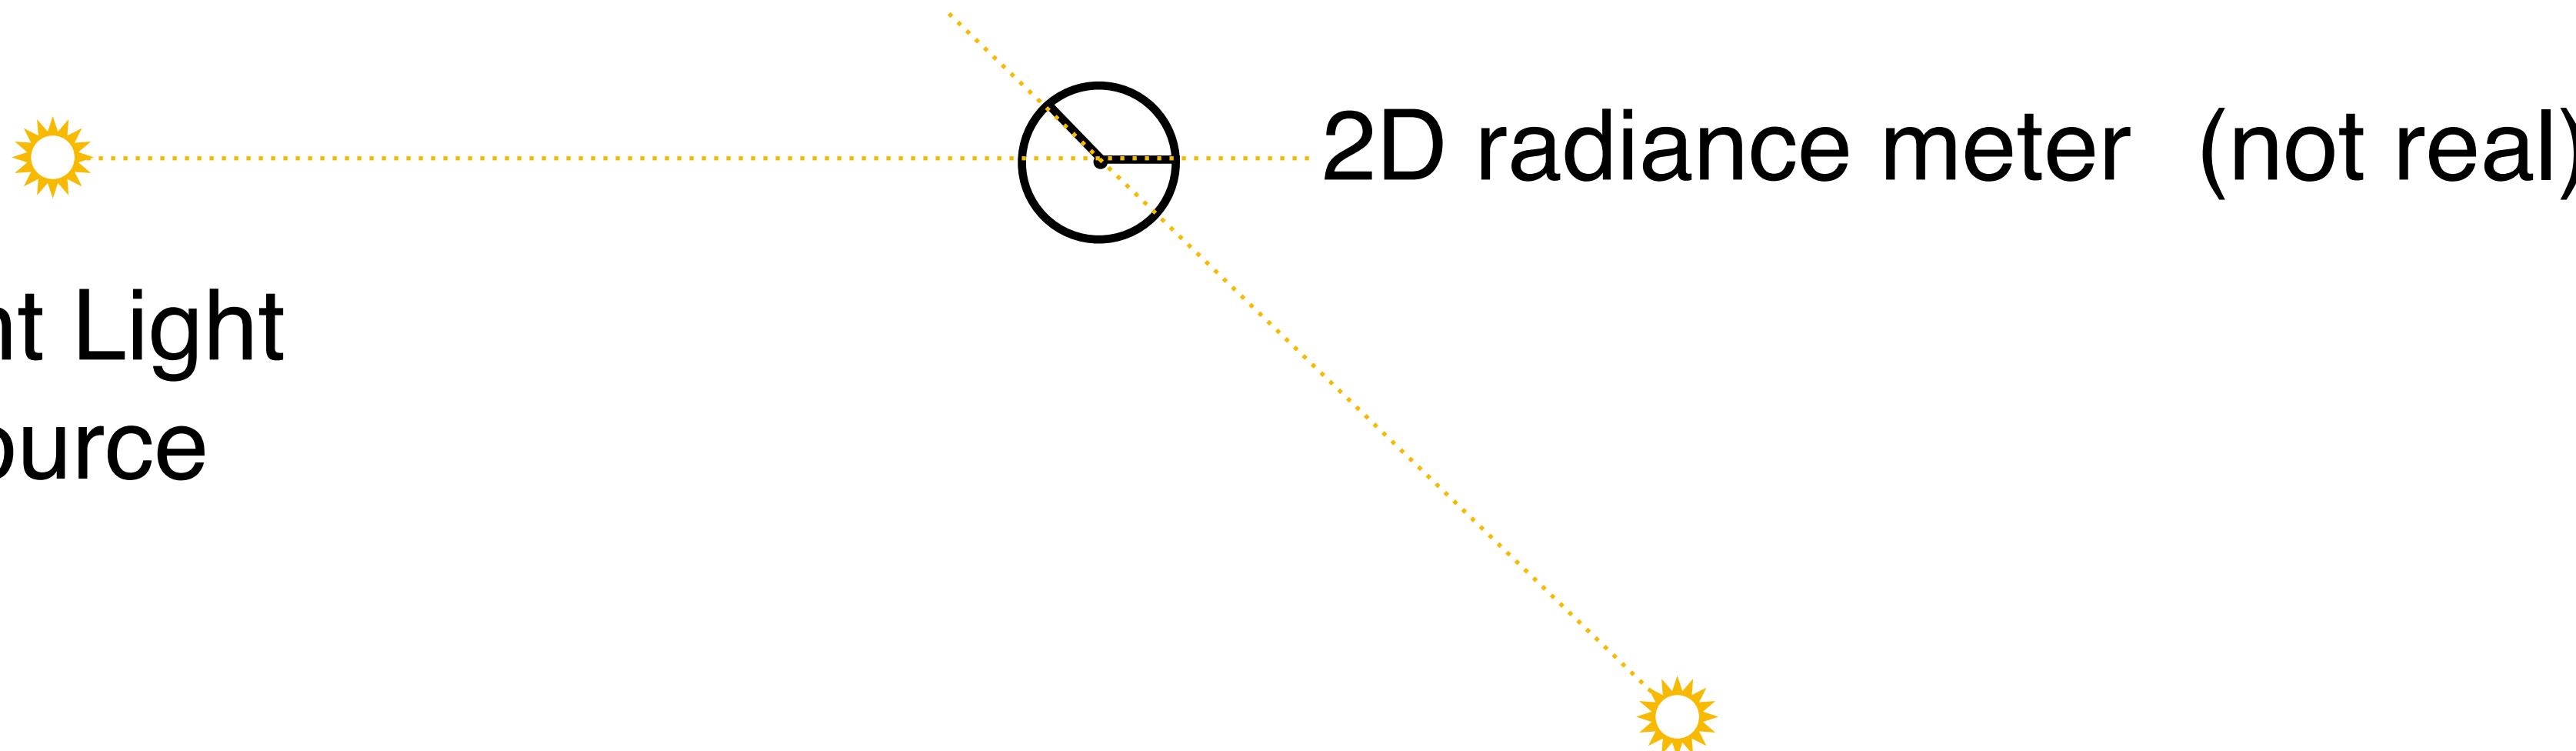
Radiance intuition

Point Light
Source



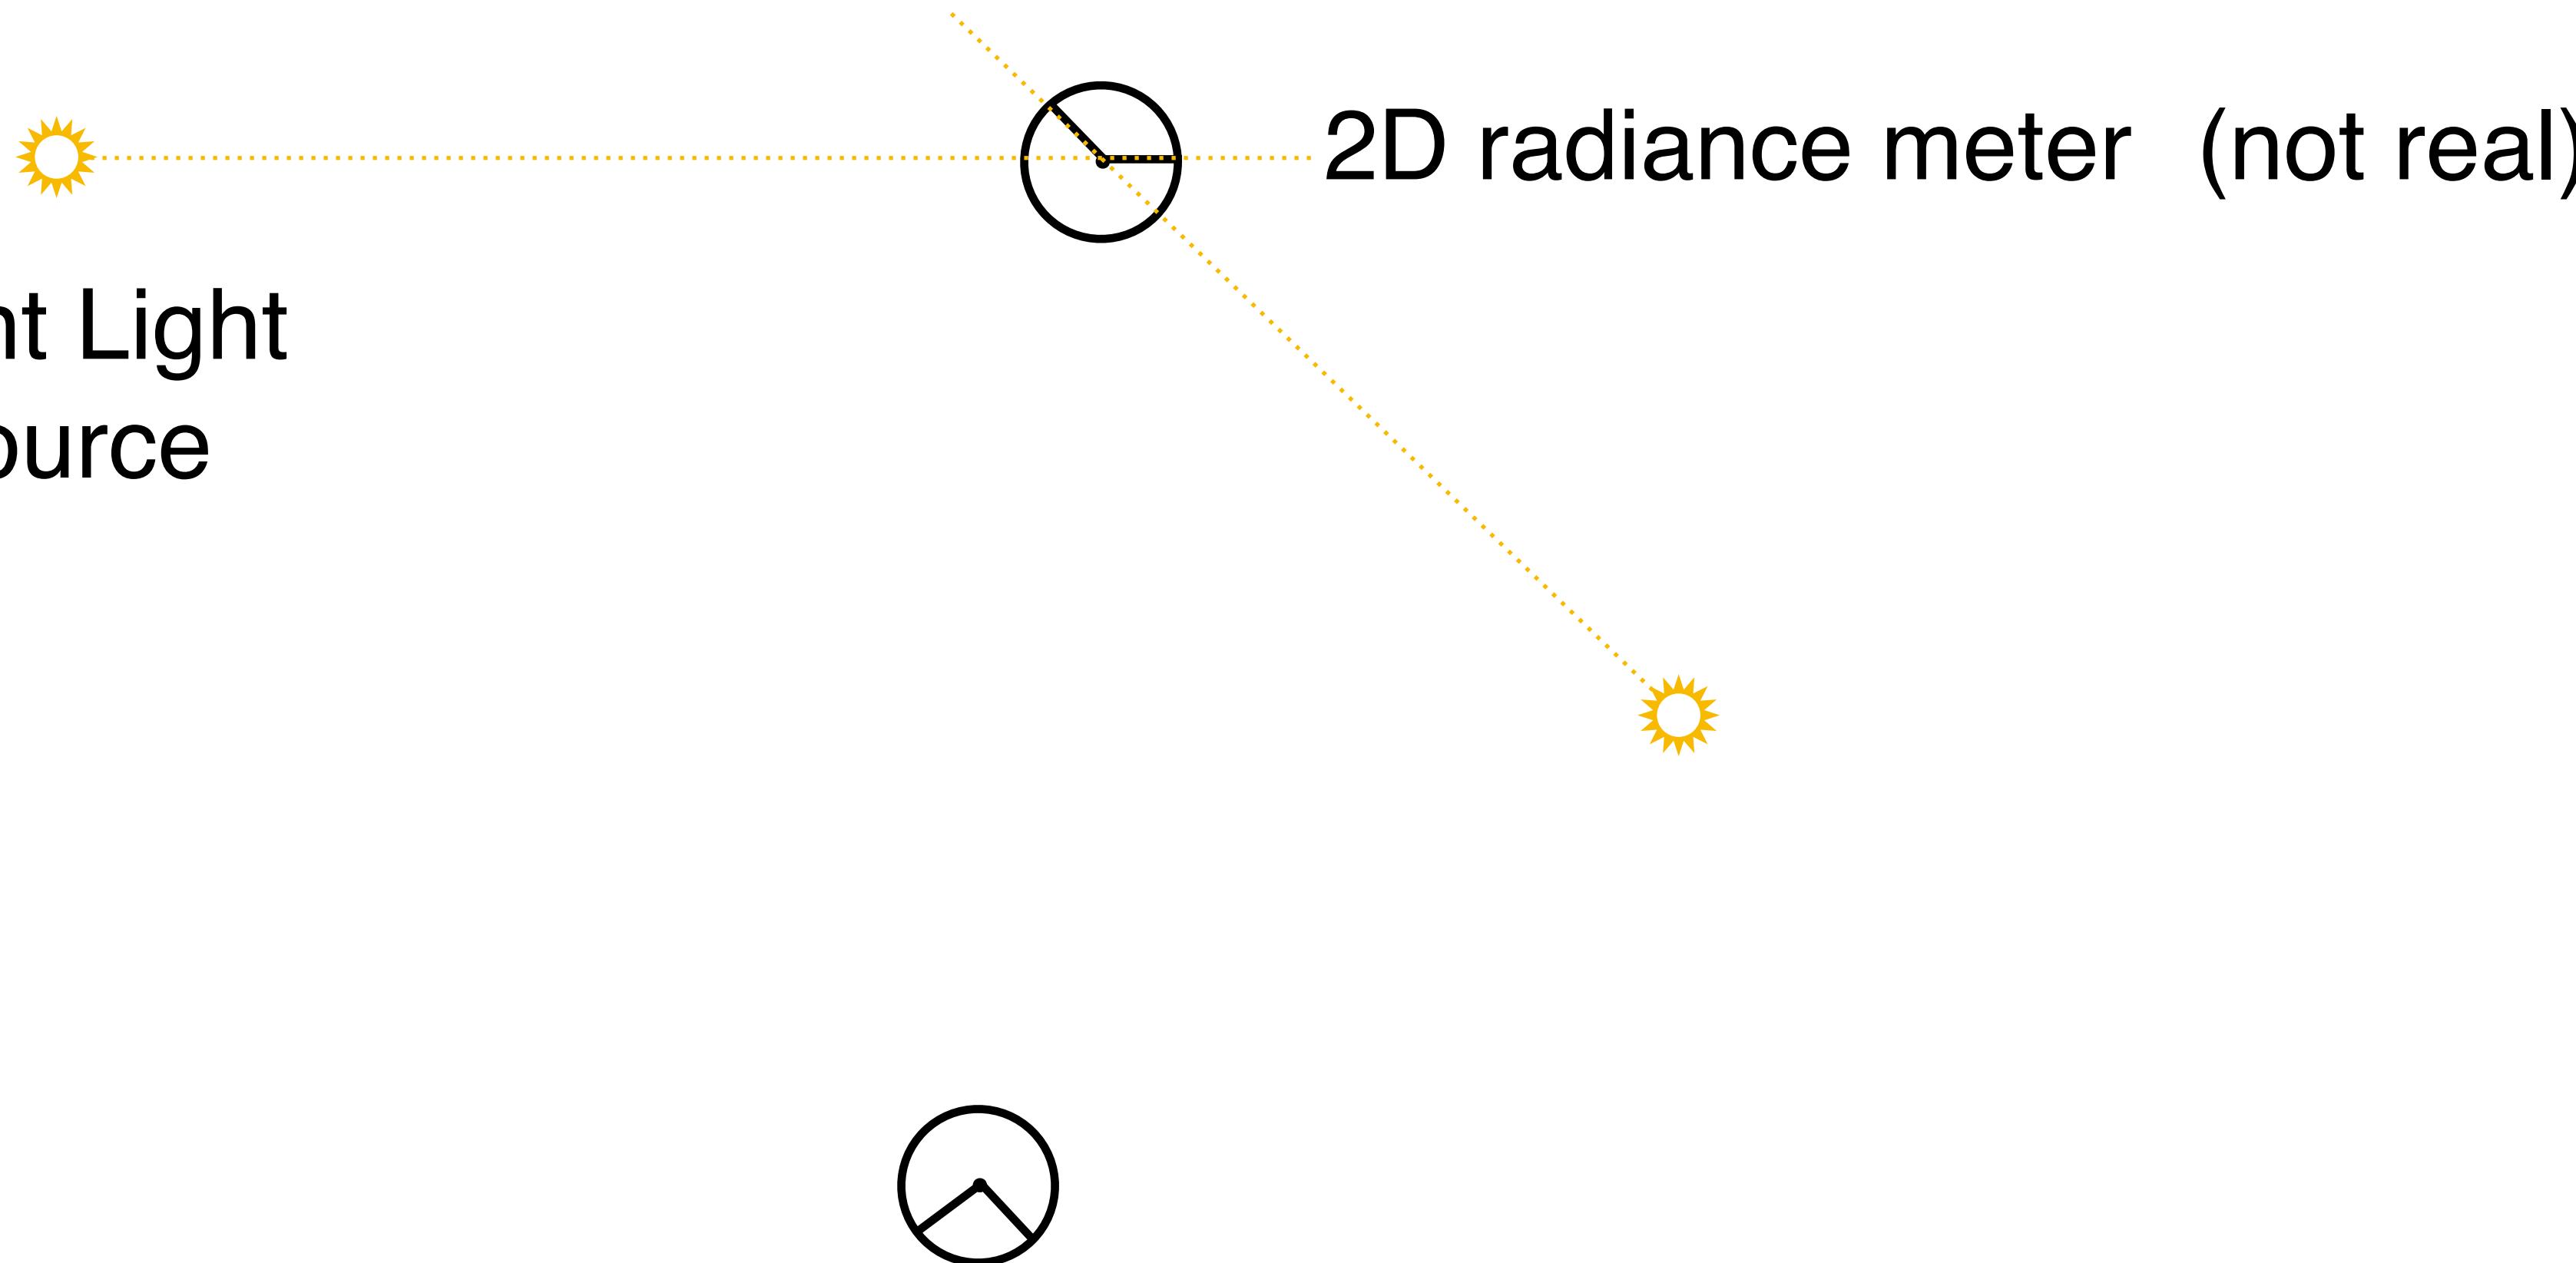
Radiance intuition

Point Light
Source



Radiance intuition

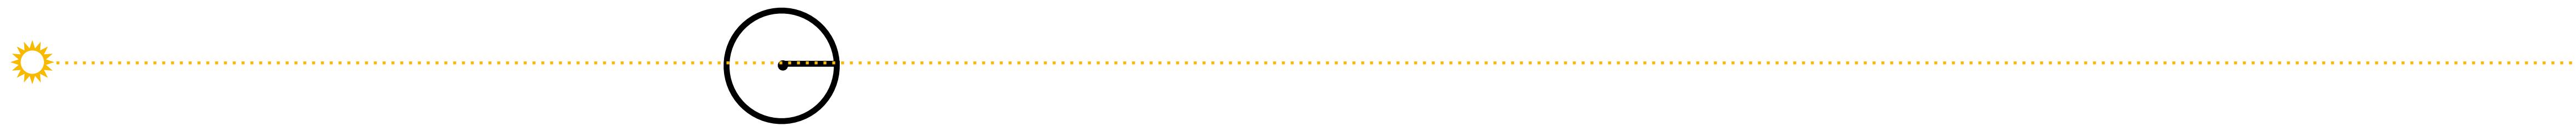
Point Light
Source



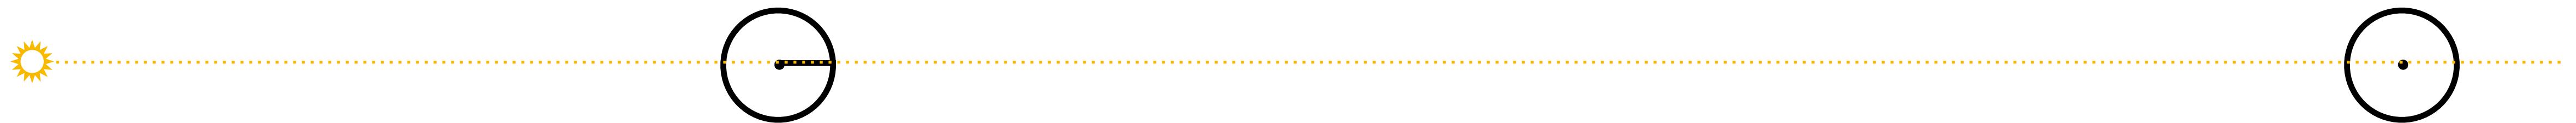
Radiance intuition



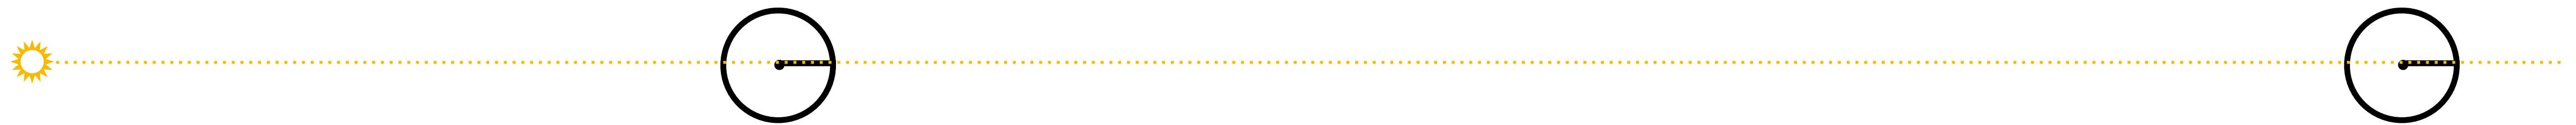
Radiance intuition



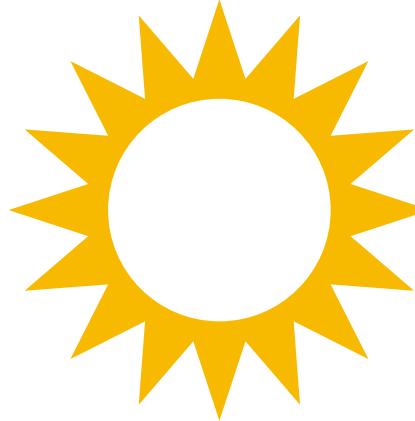
Radiance intuition



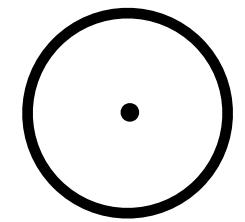
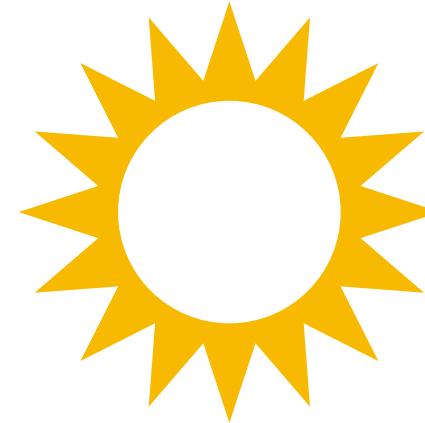
Radiance intuition



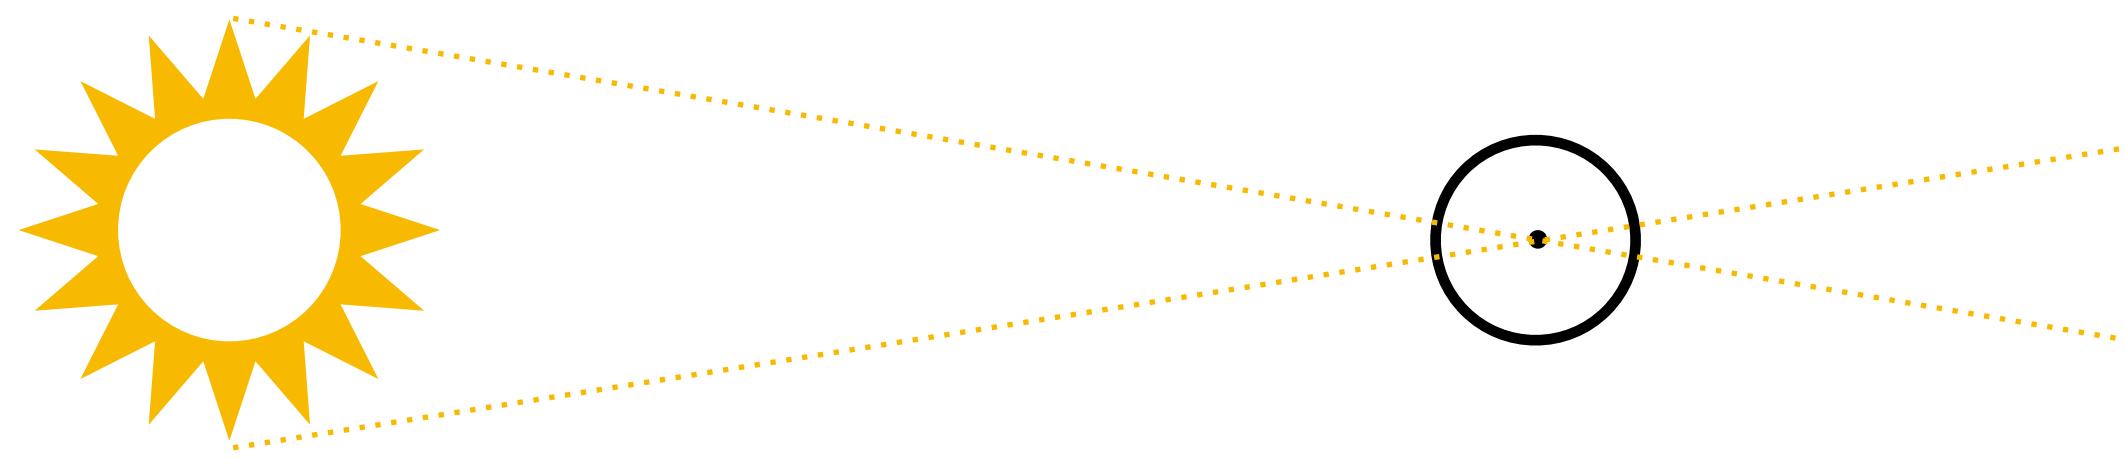
Radiance intuition



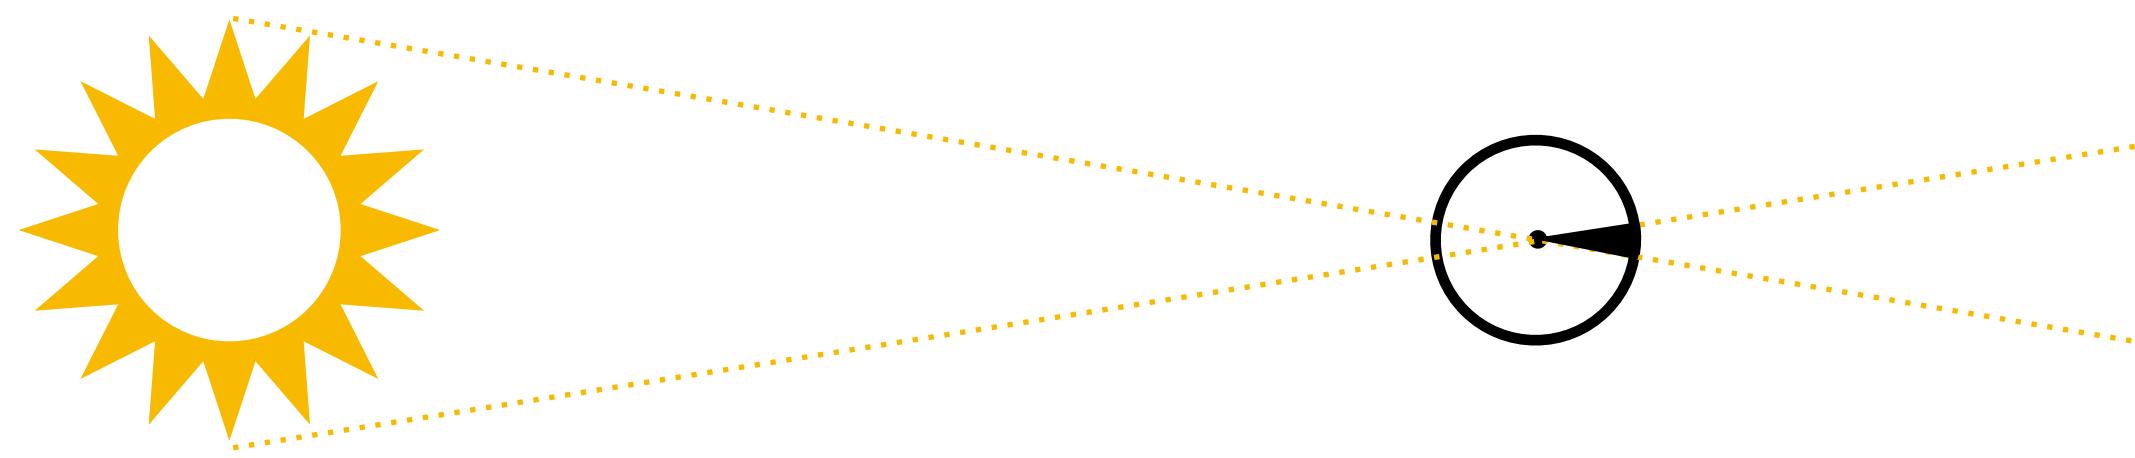
Radiance intuition



Radiance intuition



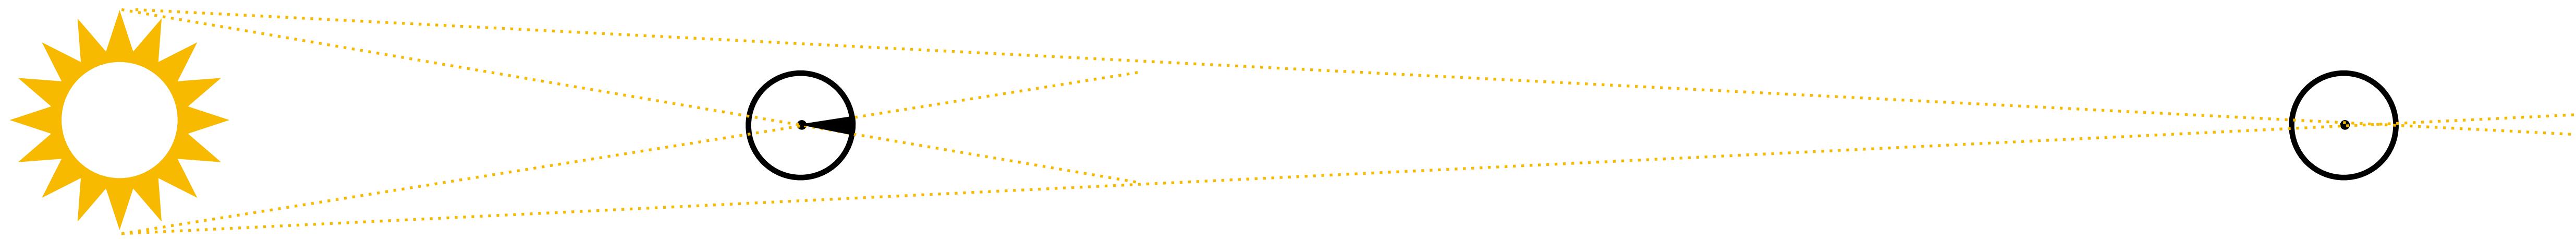
Radiance intuition



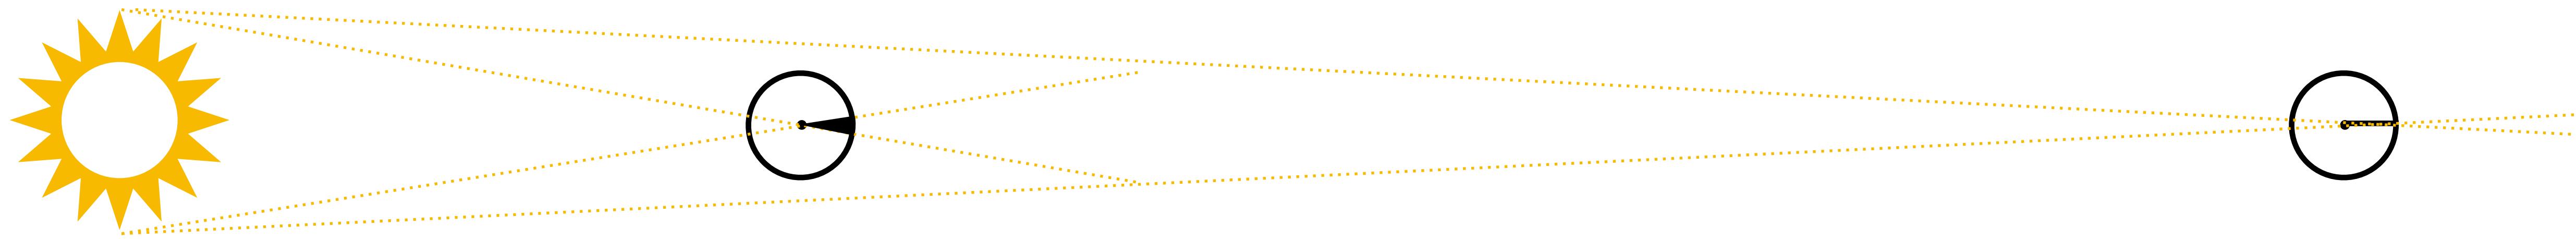
Radiance intuition



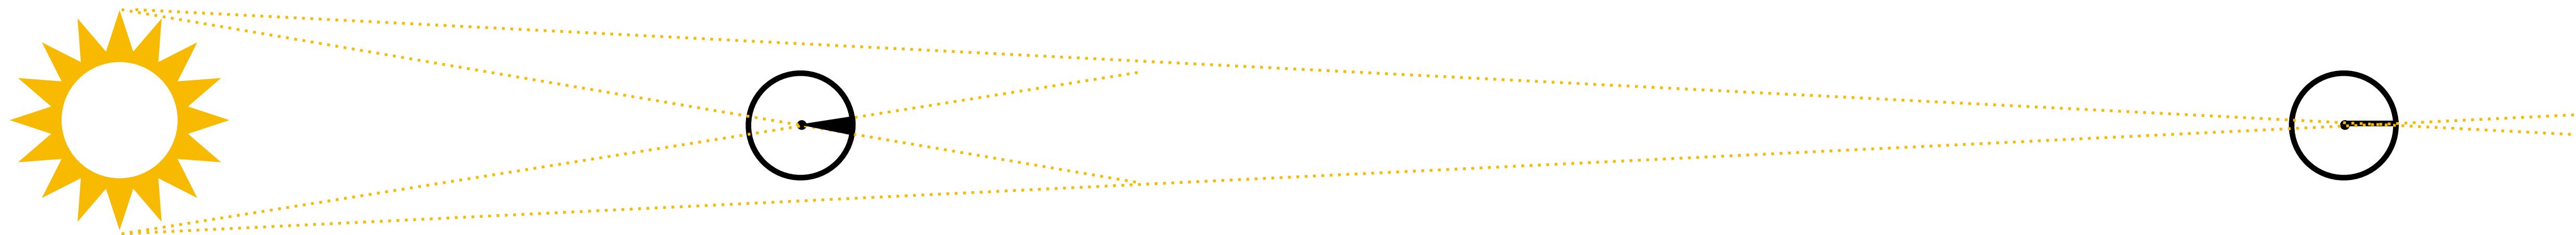
Radiance intuition



Radiance intuition

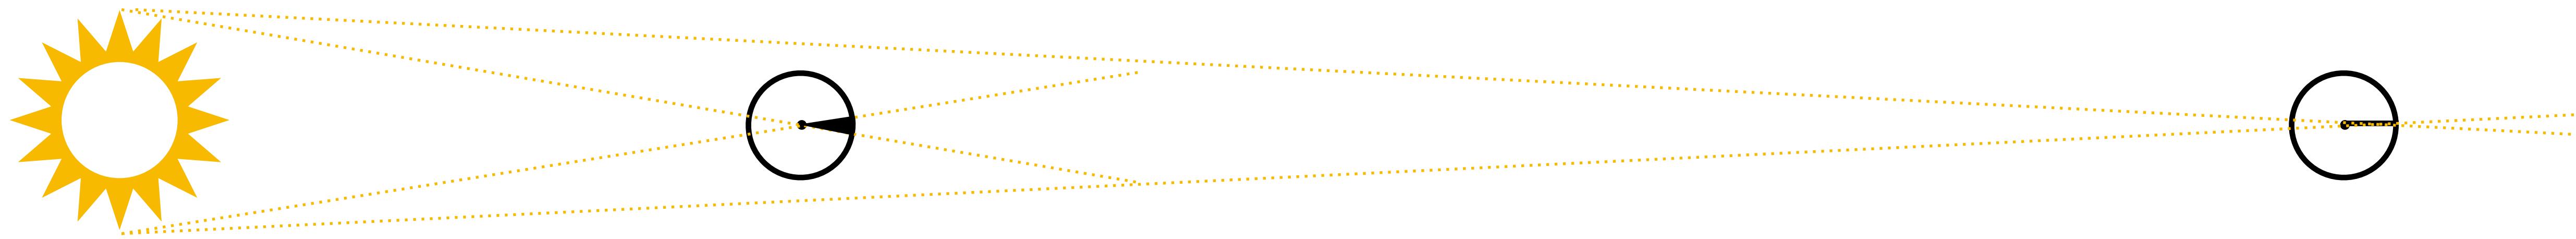


Radiance intuition



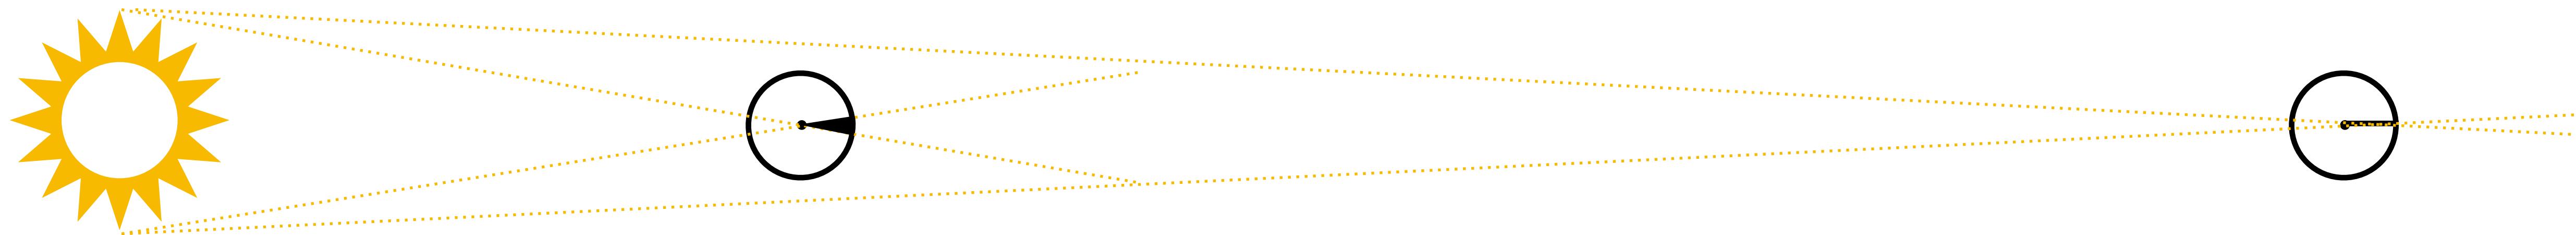
- Radiance properties

Radiance intuition



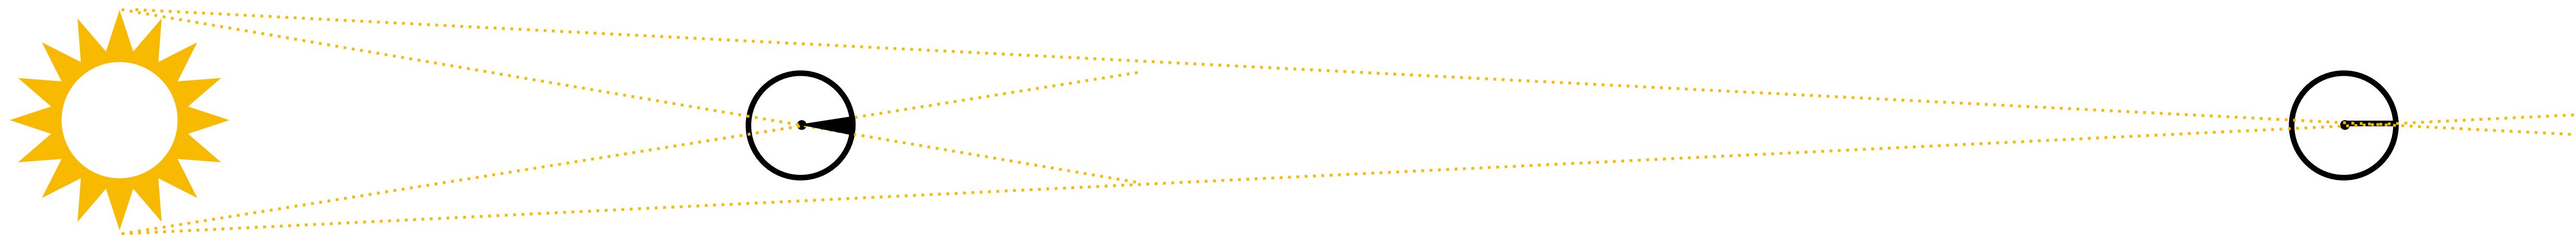
- Radiance properties
 - Radiance along an unblocked ray is constant.

Radiance intuition



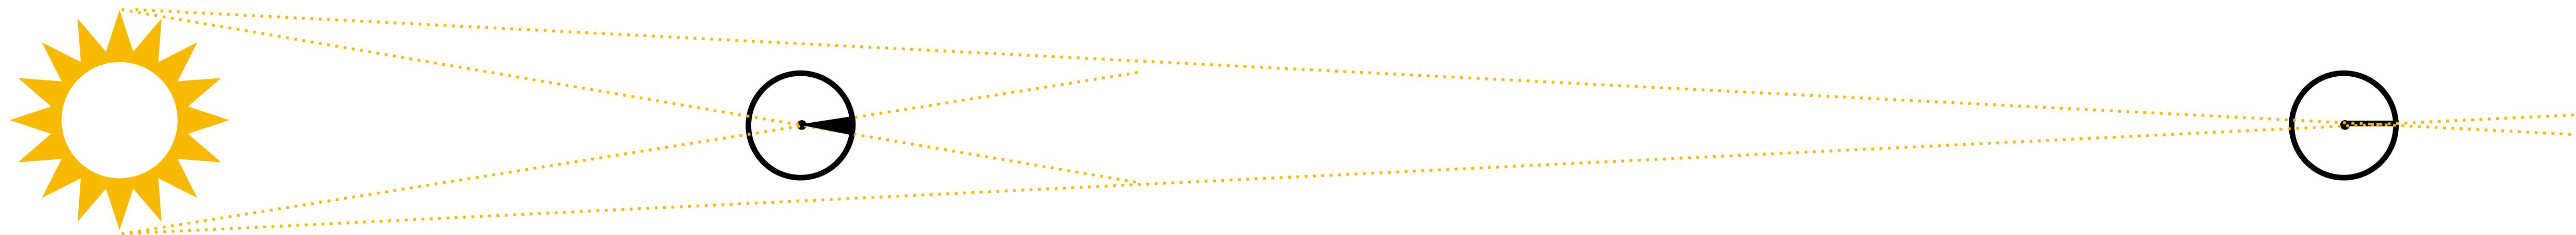
- Radiance properties
 - Radiance along an unblocked ray is constant.
 - Response of sensors is proportional to radiance of visible surface.

Radiance intuition



- Radiance properties
 - Radiance along an unblocked ray is constant.
 - Response of sensors is proportional to radiance of visible surface.
 - Radiance of reflected light is proportional to that of incoming light:

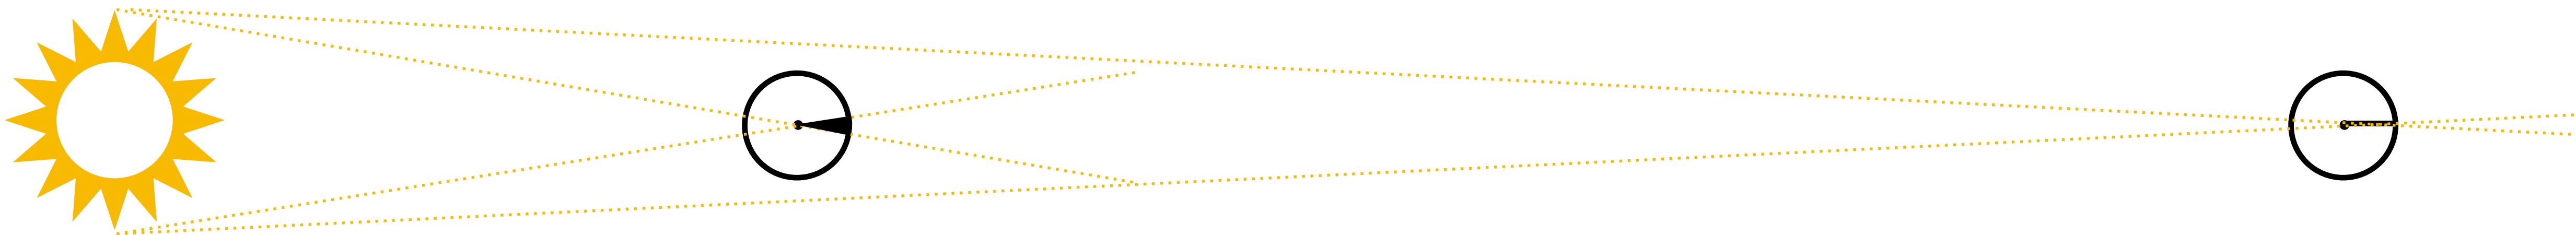
Radiance intuition



- Radiance properties
 - Radiance along an unblocked ray is constant.
 - Response of sensors is proportional to radiance of visible surface.
 - Radiance of reflected light is proportional to that of incoming light:

$$L(\mathbf{x}, \omega_o, \lambda) \propto L(\mathbf{x}, \omega_i, \lambda)$$

Radiance intuition



- Radiance properties
 - Radiance along an unblocked ray is constant.
 - Response of sensors is proportional to radiance of visible surface.
 - Radiance of reflected light is proportional to that of incoming light:

$$L(\mathbf{x}, \omega_o, \lambda) \propto L(\mathbf{x}, \omega_i, \lambda)$$

↑
reflected

Radiance intuition



- Radiance properties
 - Radiance along an unblocked ray is constant.
 - Response of sensors is proportional to radiance of visible surface.
 - Radiance of reflected light is proportional to that of incoming light:

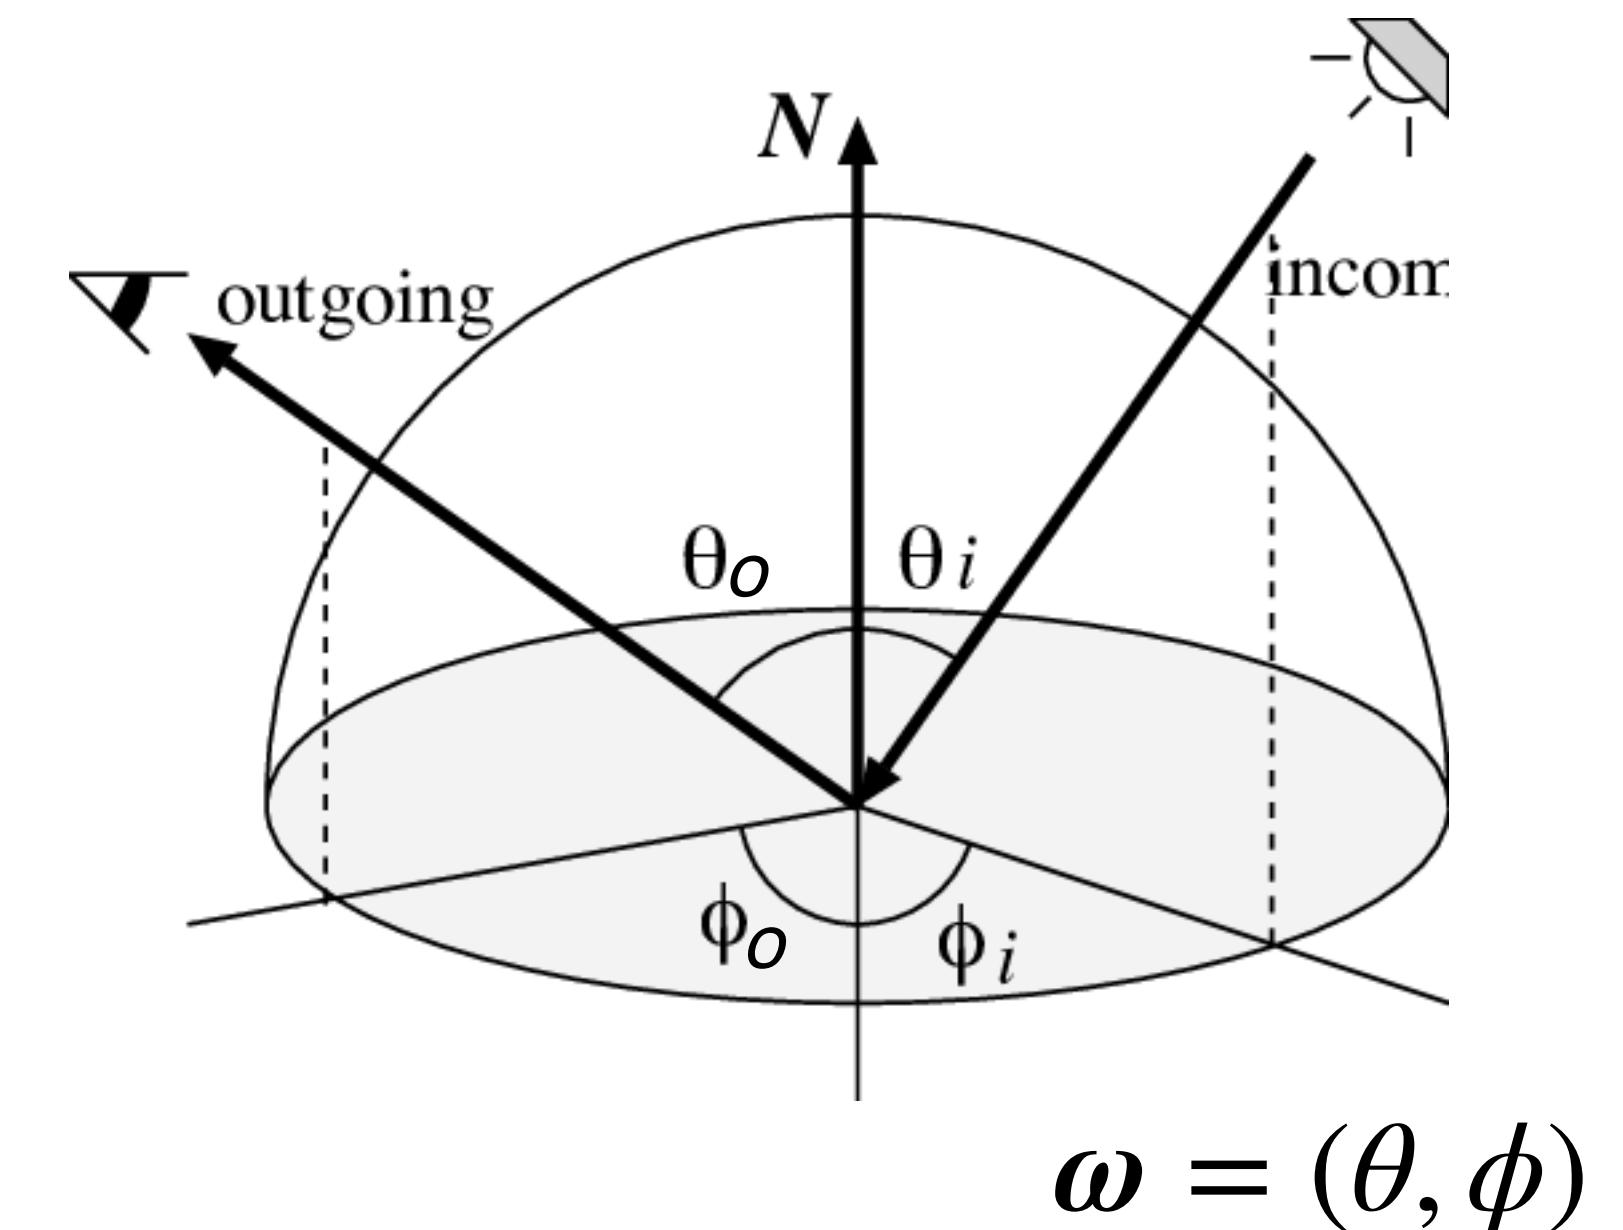
BRDF

- Constant of proportionality defines Bidirectional Reflectance Distribution Function (BRDF).

BRDF

- Constant of proportionality defines Bidirectional Reflectance Distribution Function (BRDF).

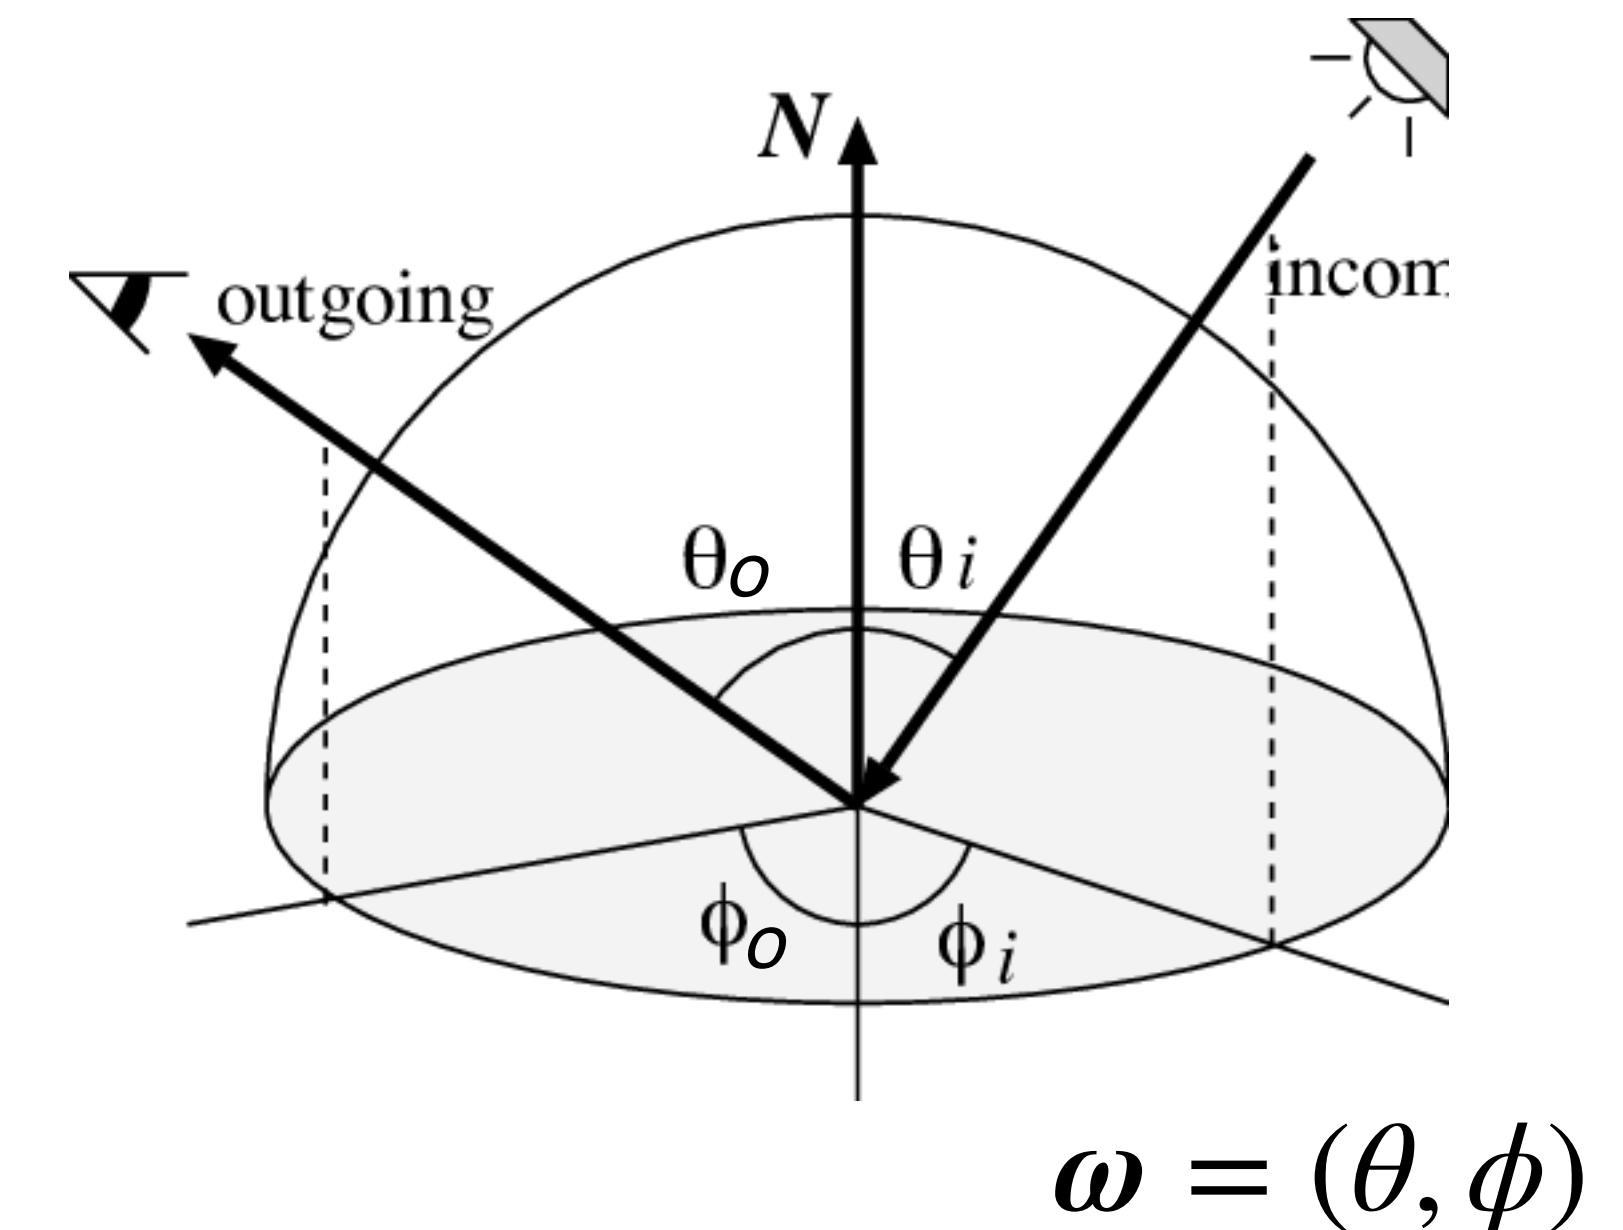
$$f(\omega_i, \omega_o) = \frac{\text{outgoing light in direction } \omega_o}{\text{incoming light in direction } \omega_i}$$



BRDF

- Constant of proportionality defines Bidirectional Reflectance Distribution Function (BRDF).

$$f(\omega_i, \omega_o) = \frac{\text{outgoing light in direction } \omega_o}{\text{incoming light in direction } \omega_i}$$



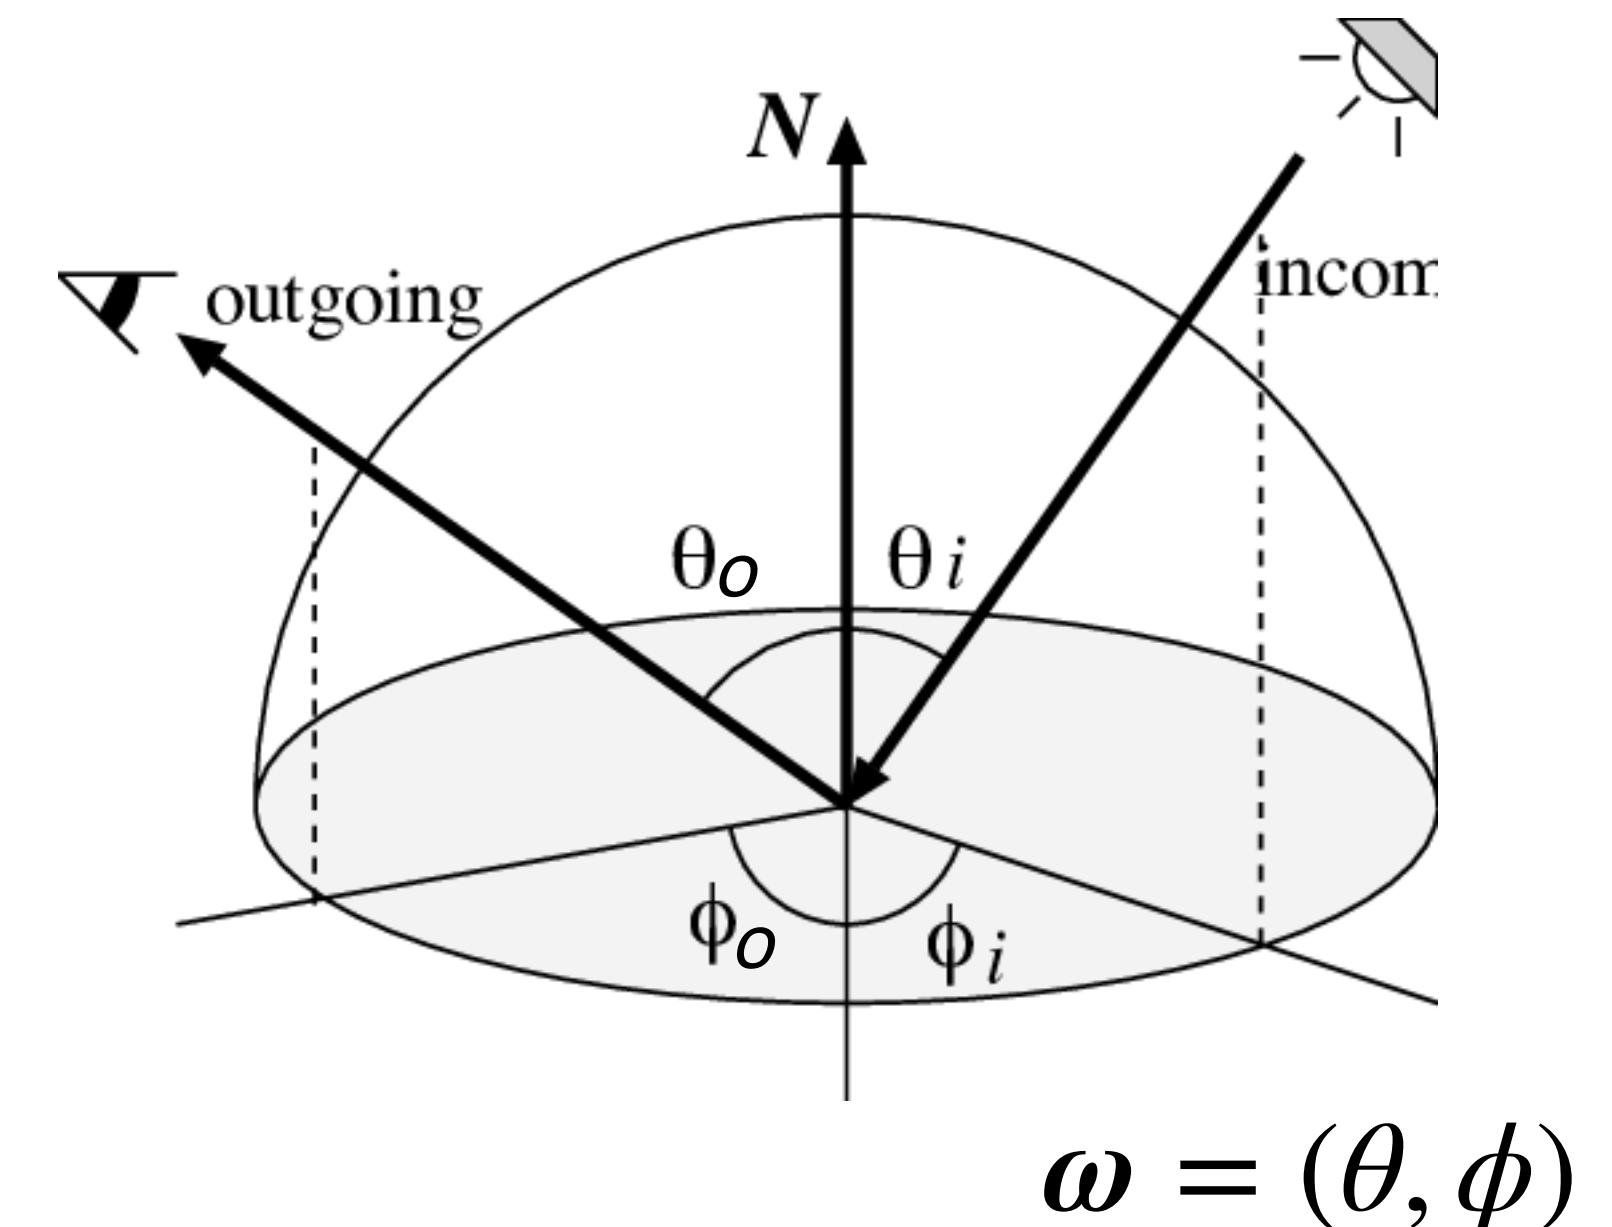
- BRDF describes how a material reflects light.

BRDF

- Constant of proportionality defines Bidirectional Reflectance Distribution Function (BRDF).

$$f(\omega_i, \omega_o) = \frac{\text{outgoing light in direction } \omega_o}{\text{incoming light in direction } \omega_i}$$

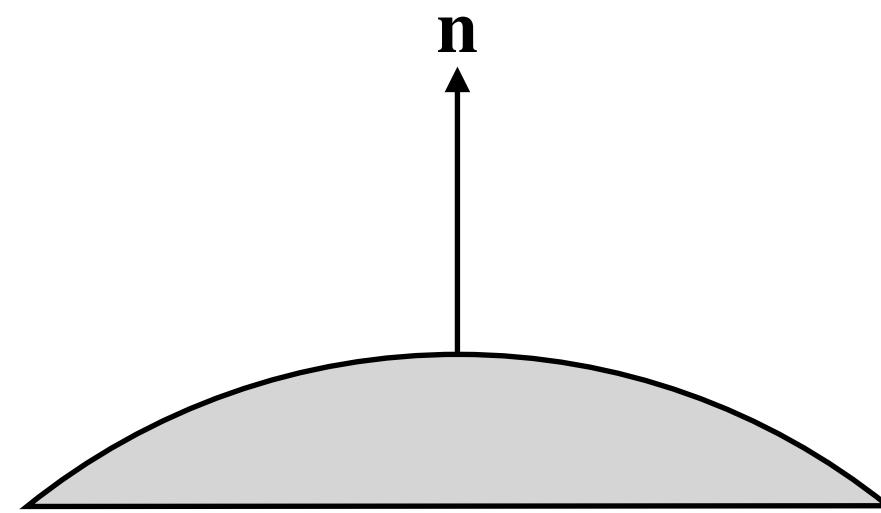
- BRDF describes how a material reflects light.
- Helmholtz reciprocity: $f(\omega_1, \omega_2) = f(\omega_2, \omega_1)$



$$\omega = (\theta, \phi)$$

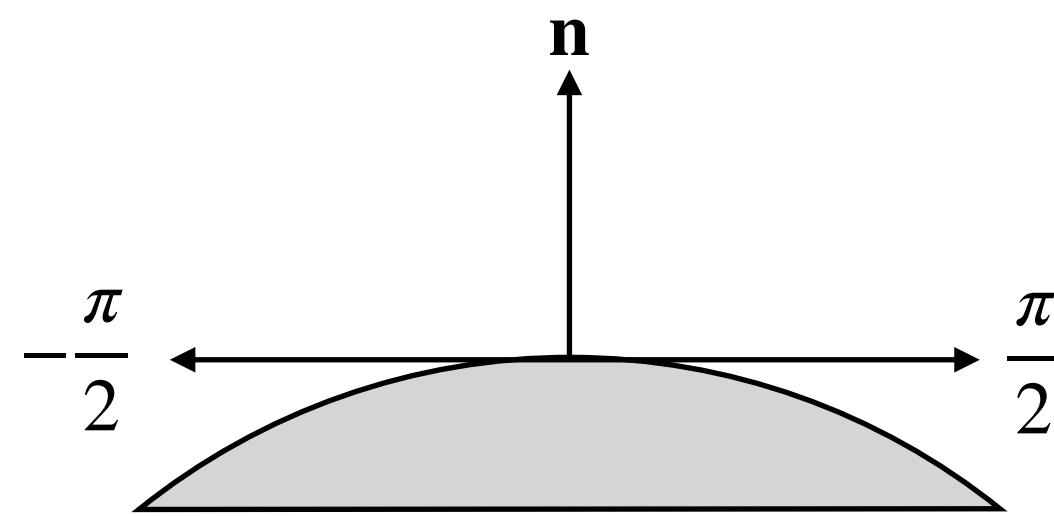
BRDF intuition

perfect mirror



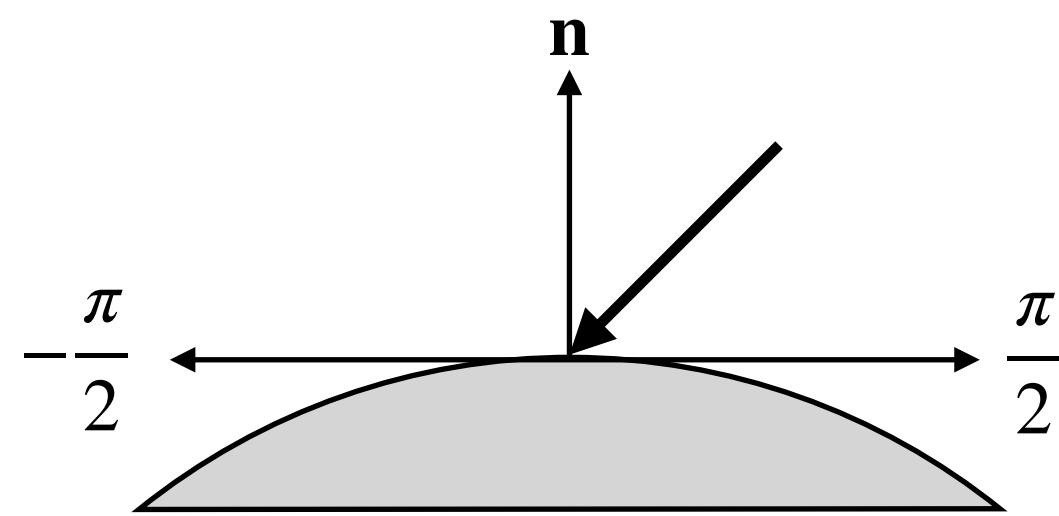
BRDF intuition

perfect mirror



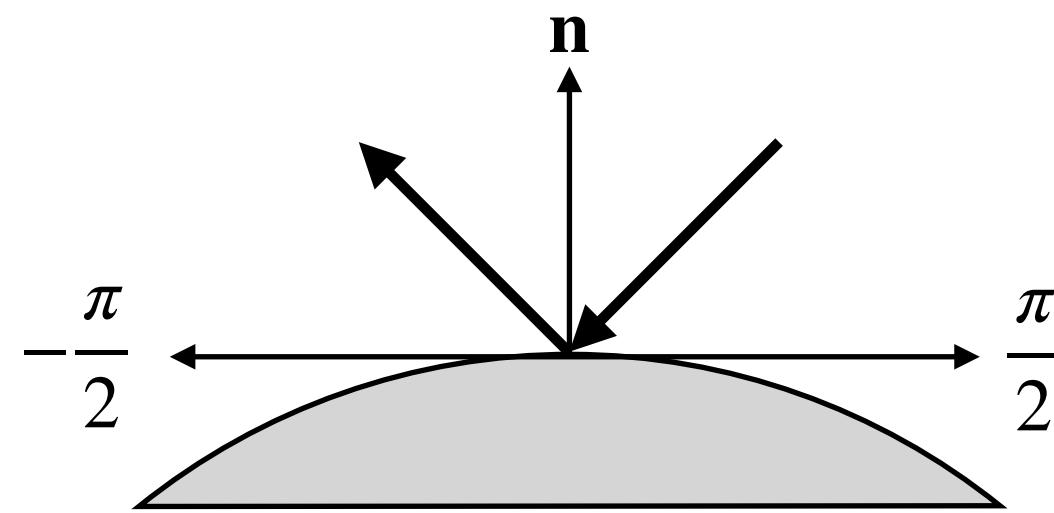
BRDF intuition

perfect mirror



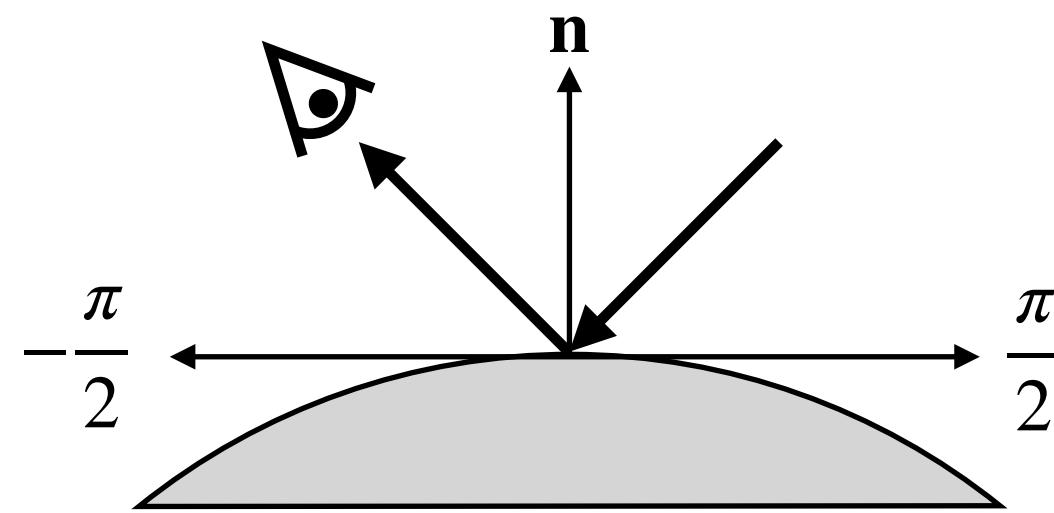
BRDF intuition

perfect mirror



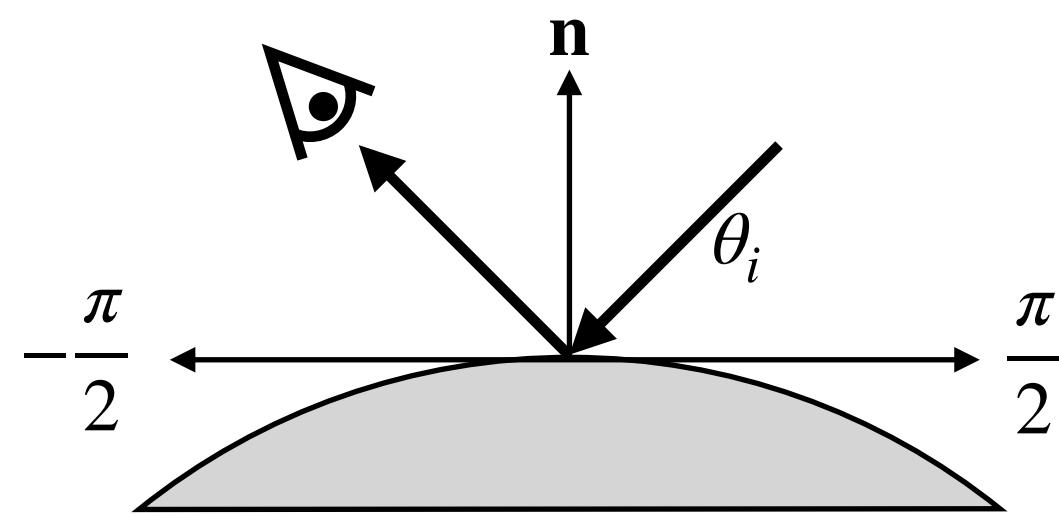
BRDF intuition

perfect mirror



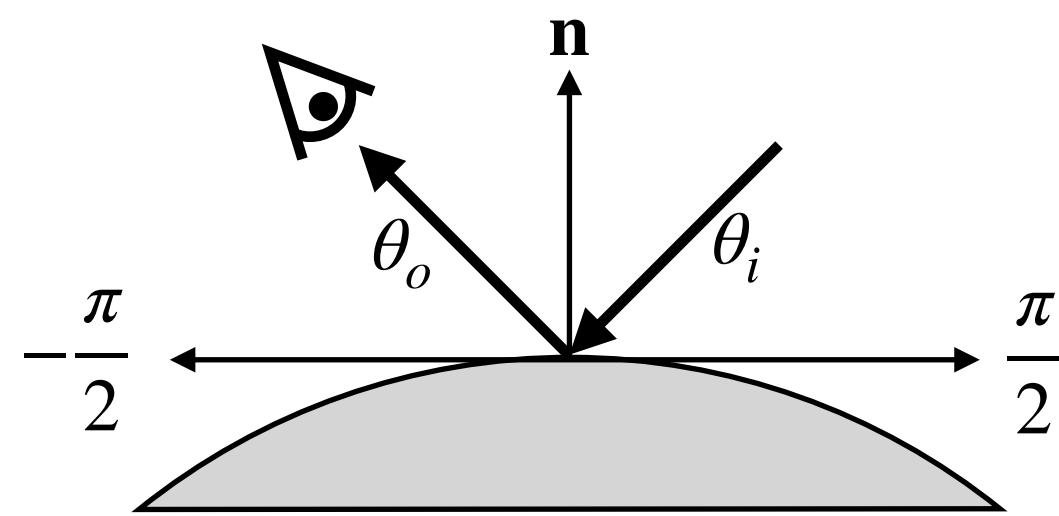
BRDF intuition

perfect mirror



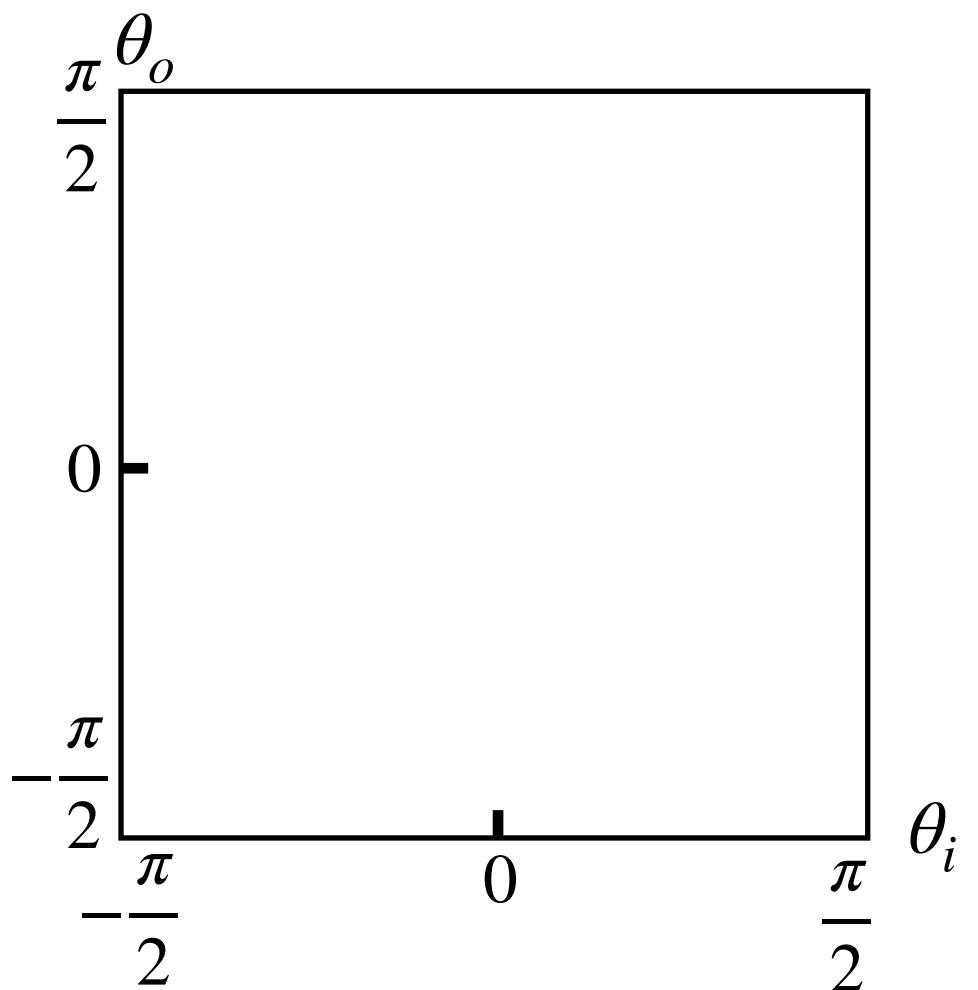
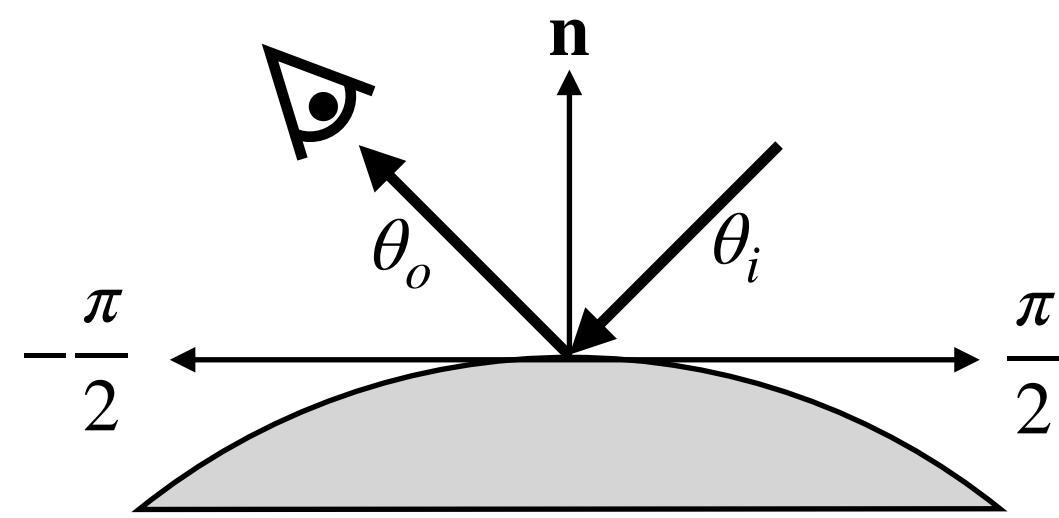
BRDF intuition

perfect mirror



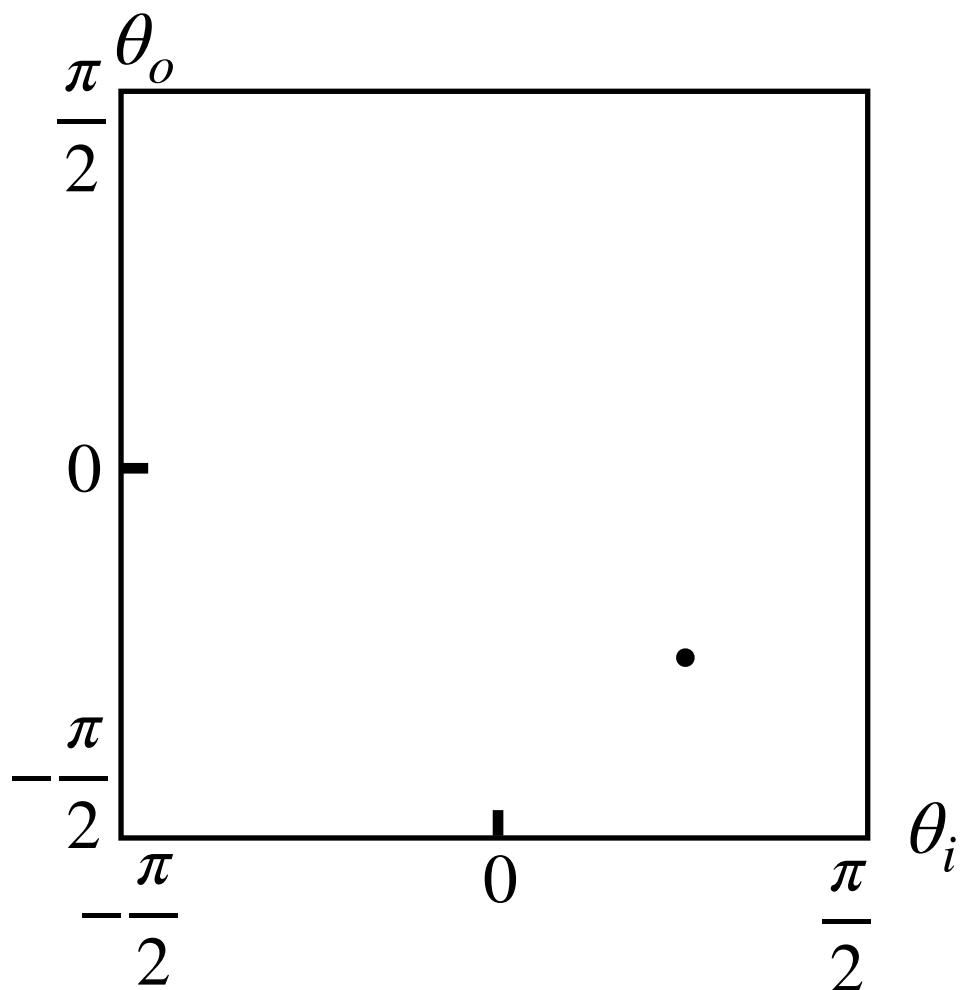
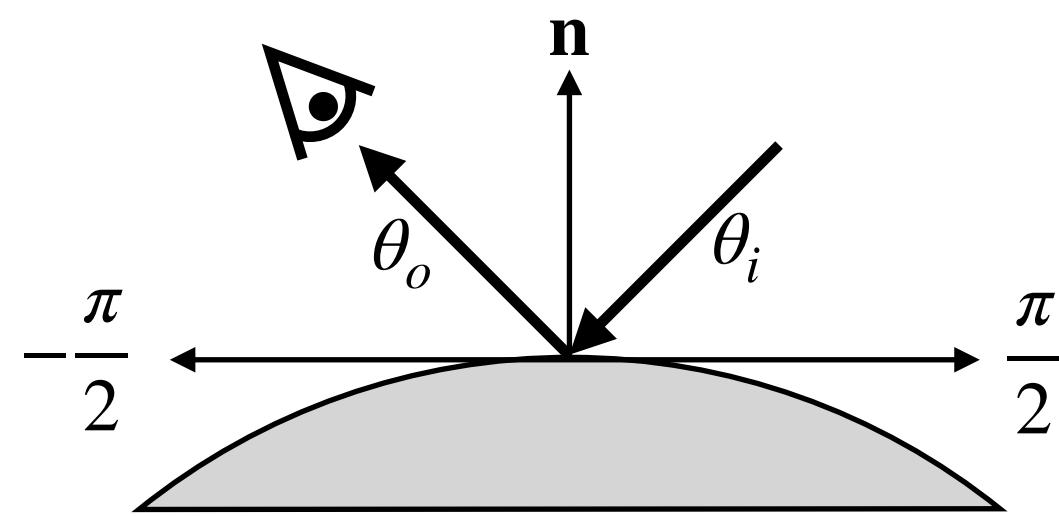
BRDF intuition

perfect mirror



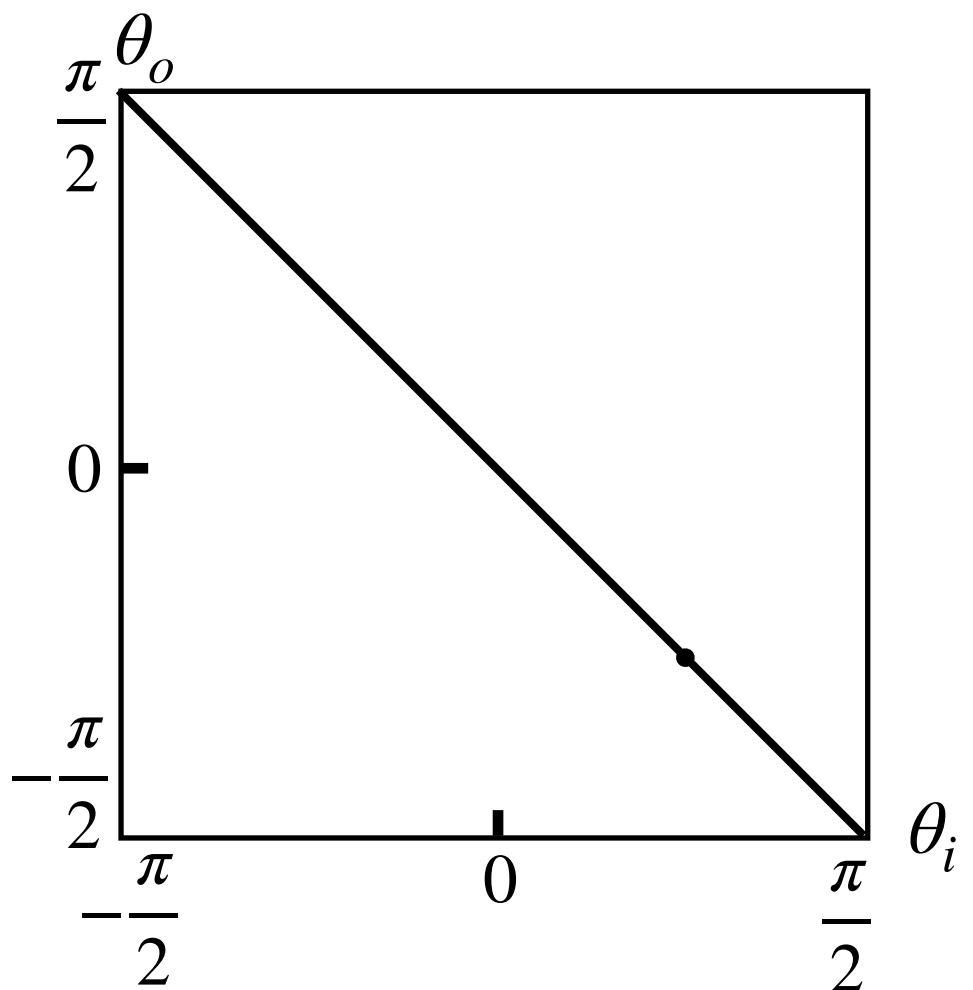
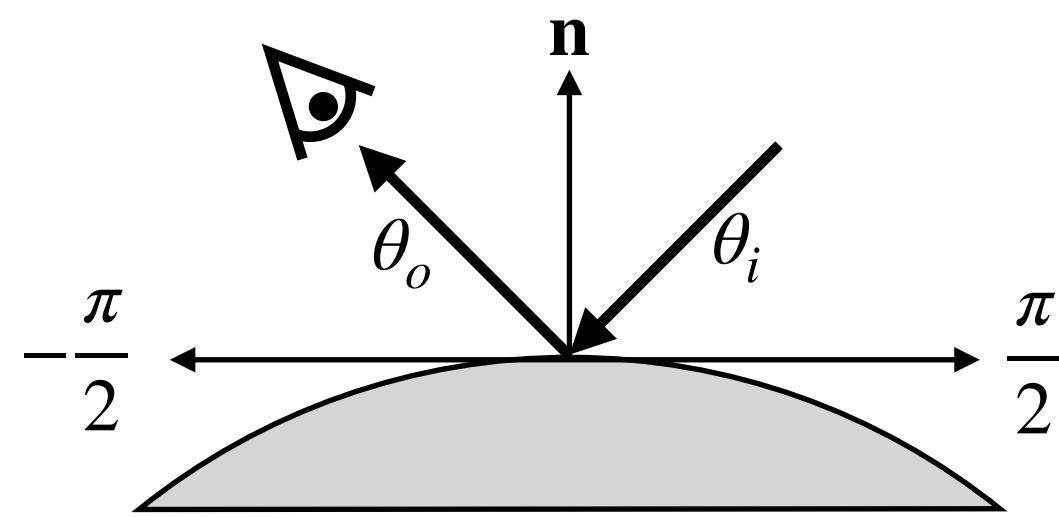
BRDF intuition

perfect mirror



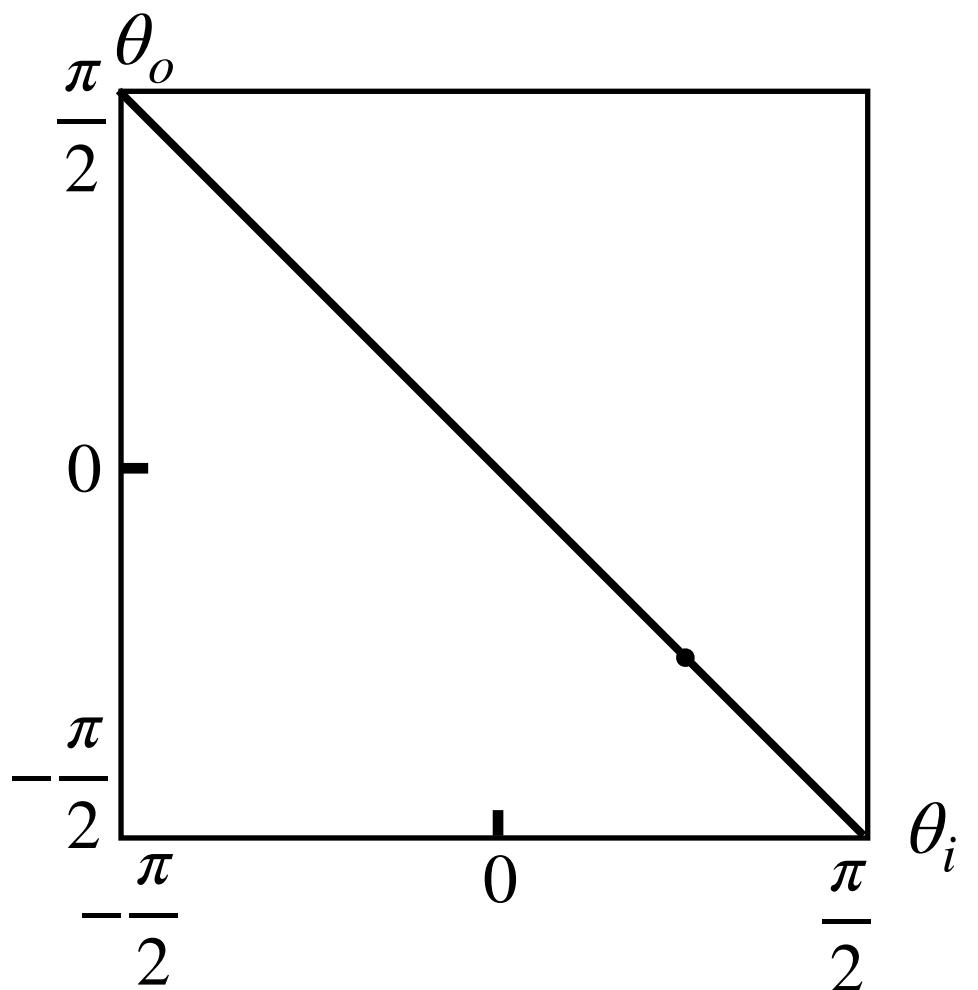
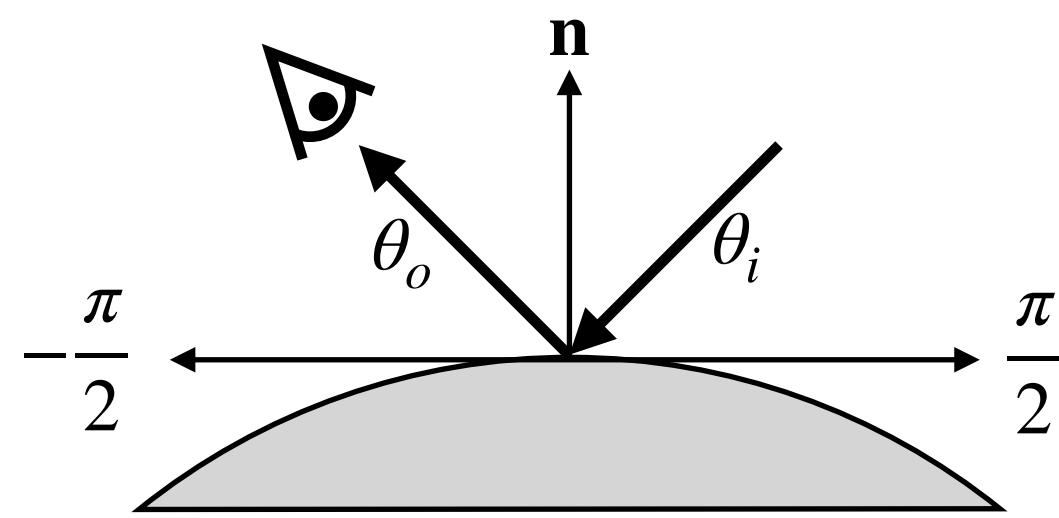
BRDF intuition

perfect mirror

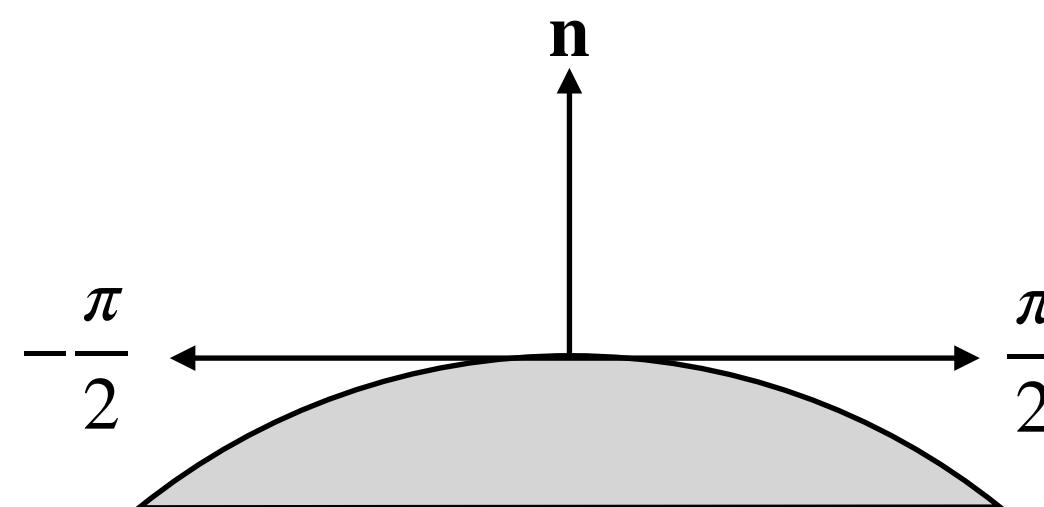


BRDF intuition

perfect mirror

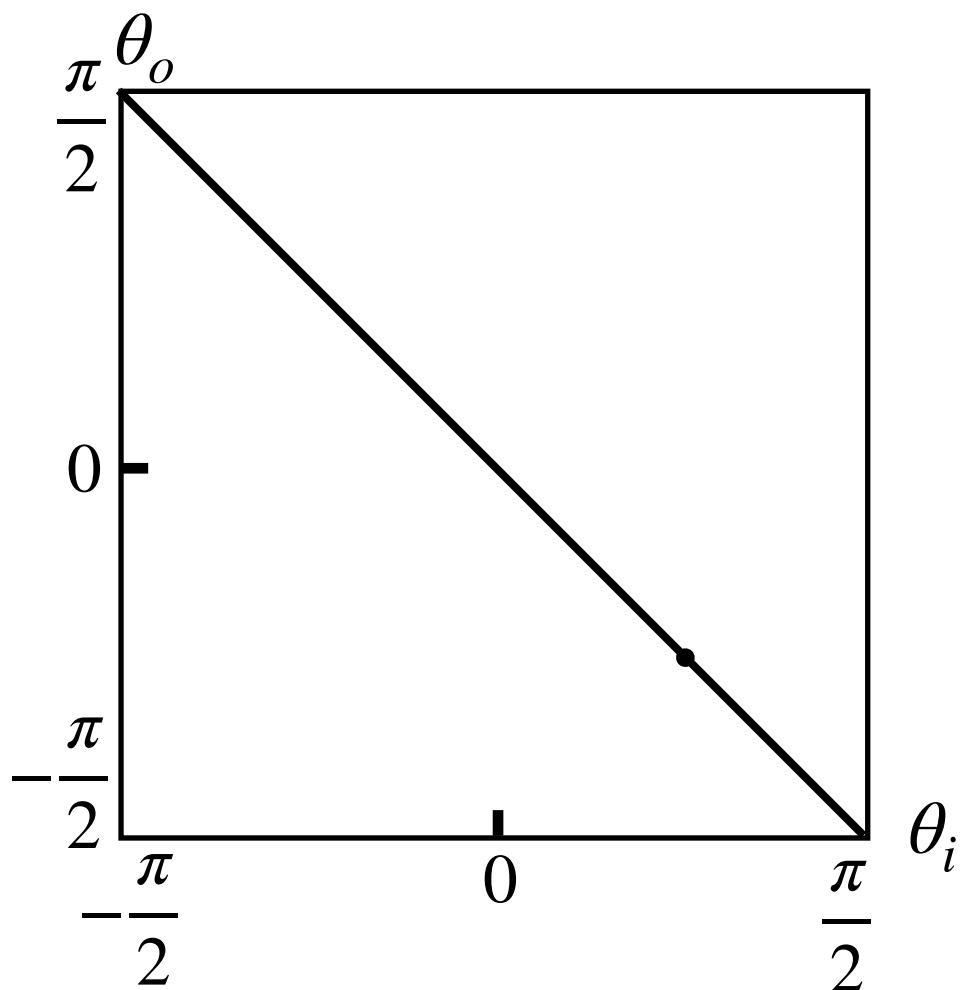
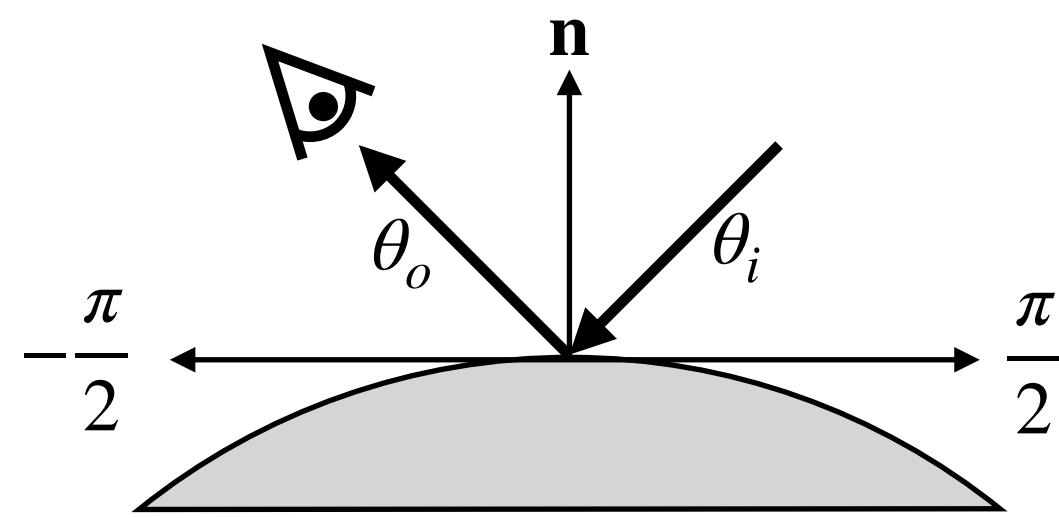


fairly shiny mirror

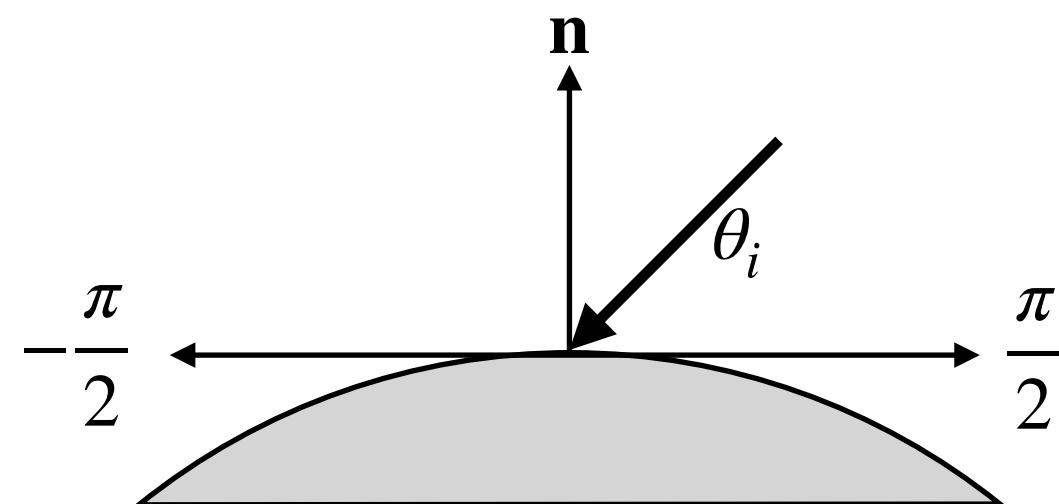


BRDF intuition

perfect mirror

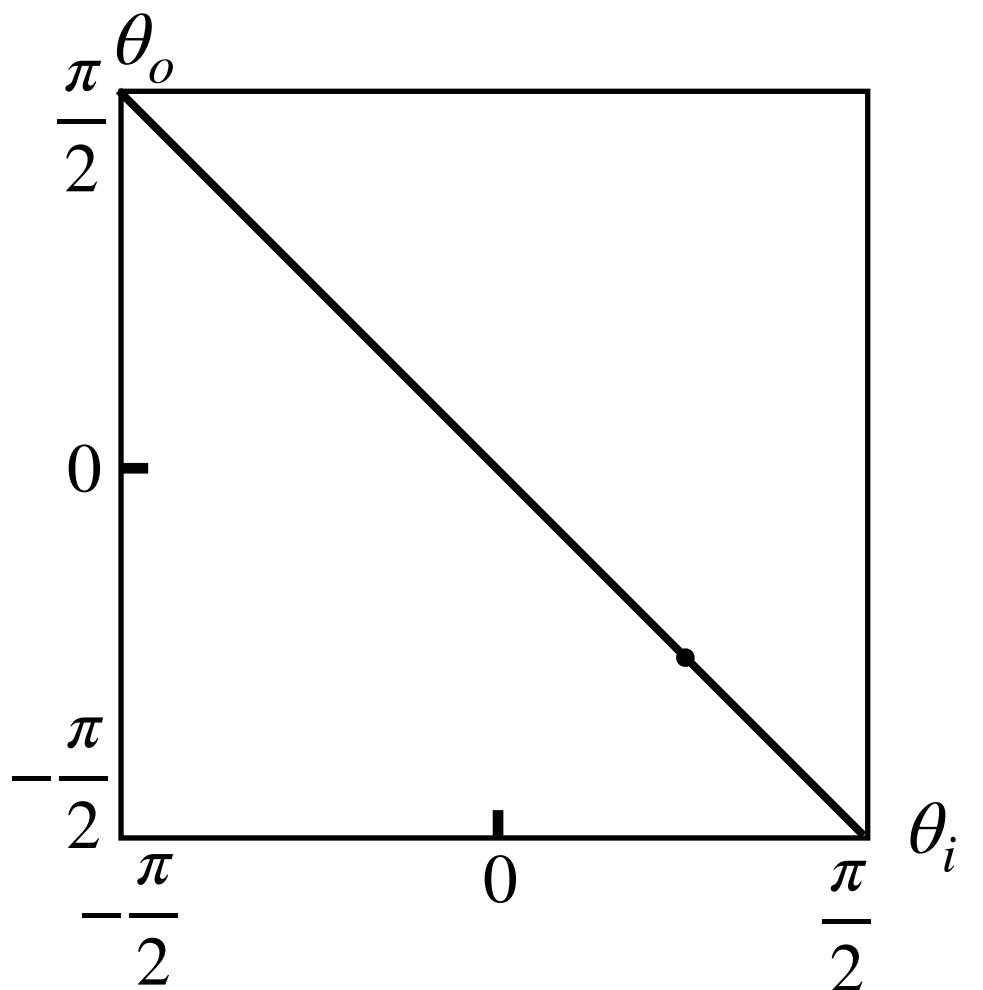
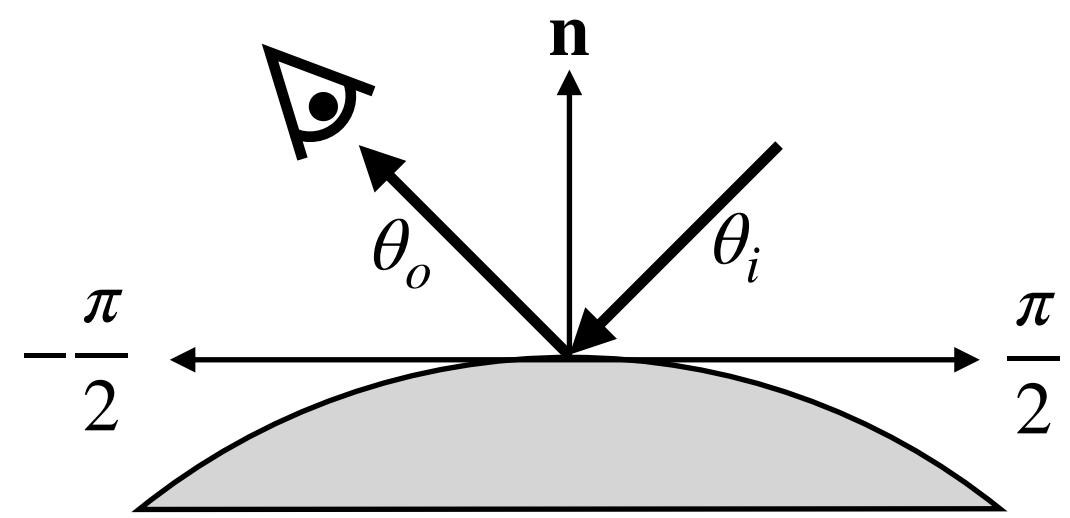


fairly shiny mirror

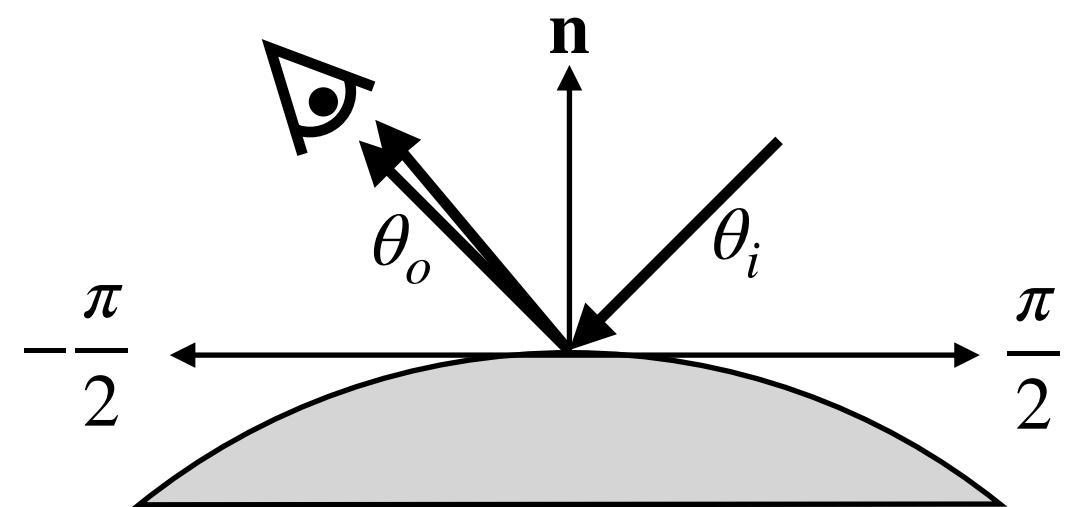


BRDF intuition

perfect mirror

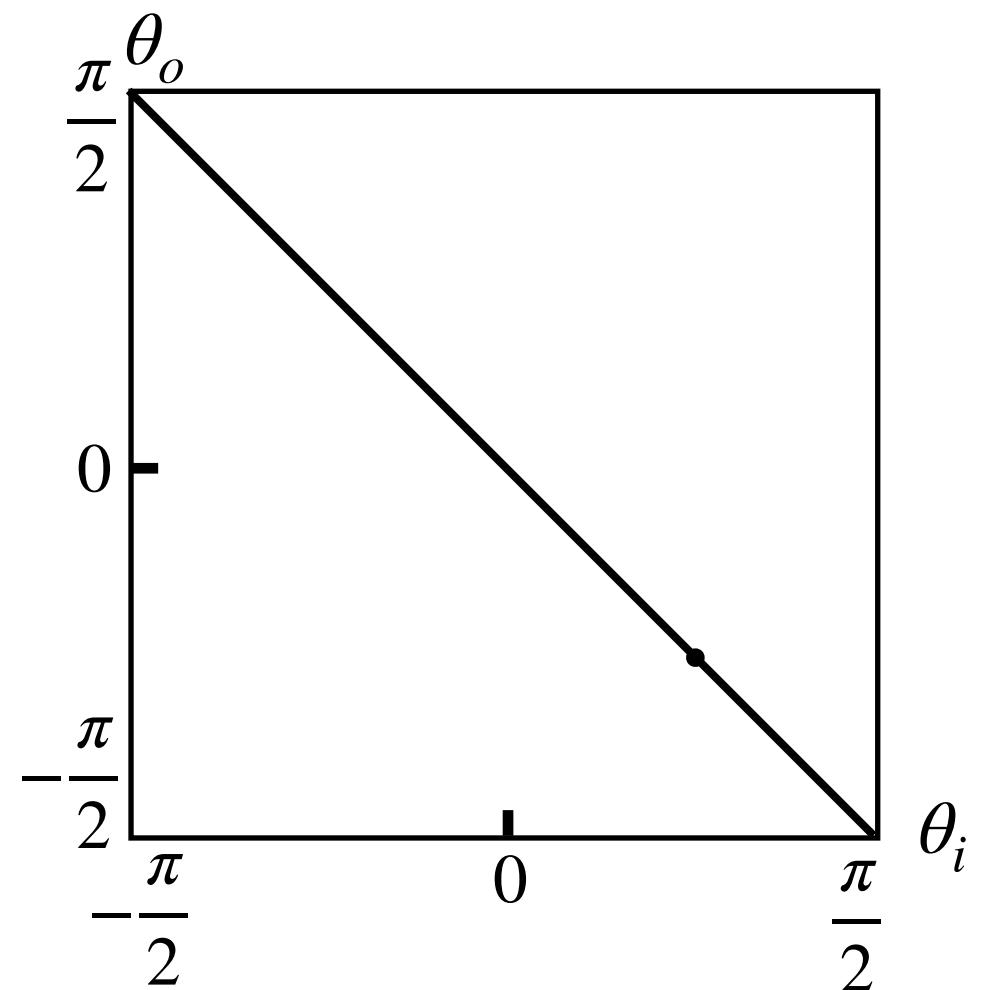
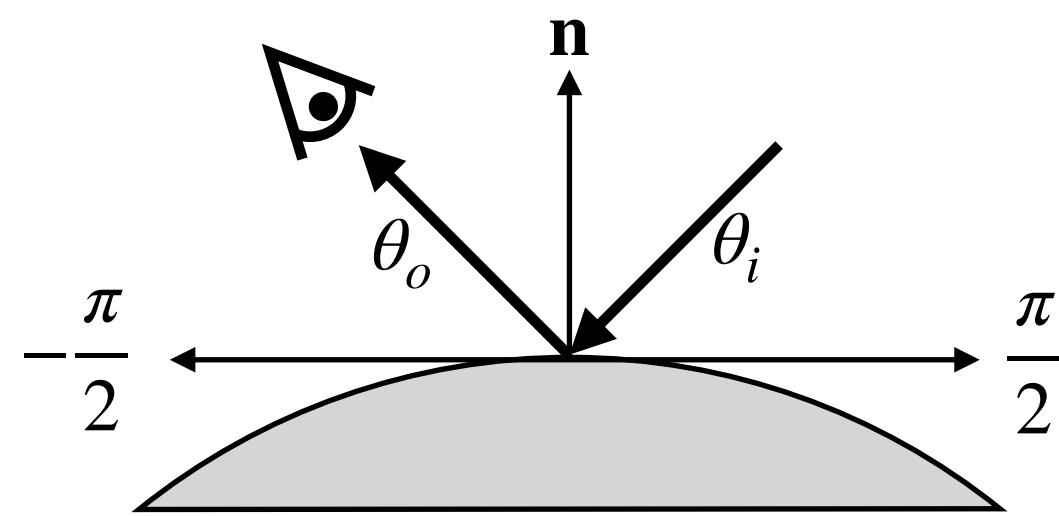


fairly shiny mirror

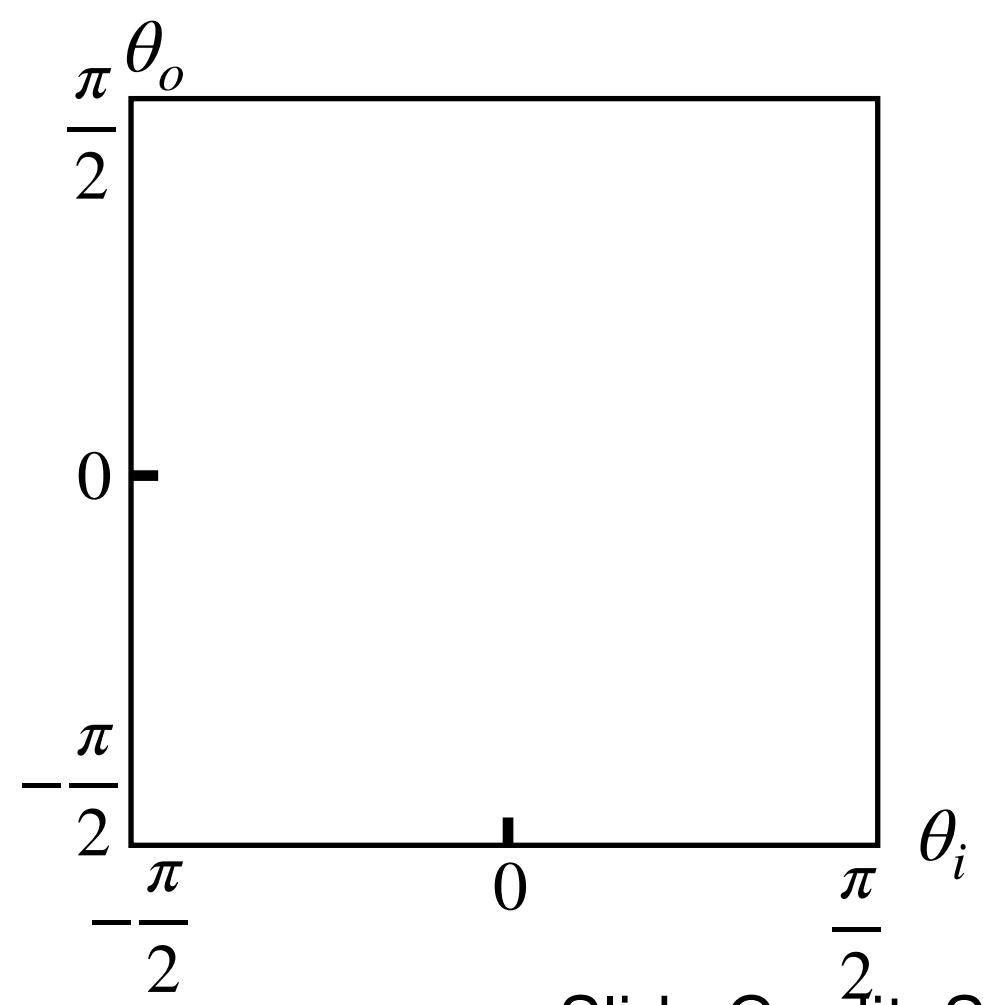
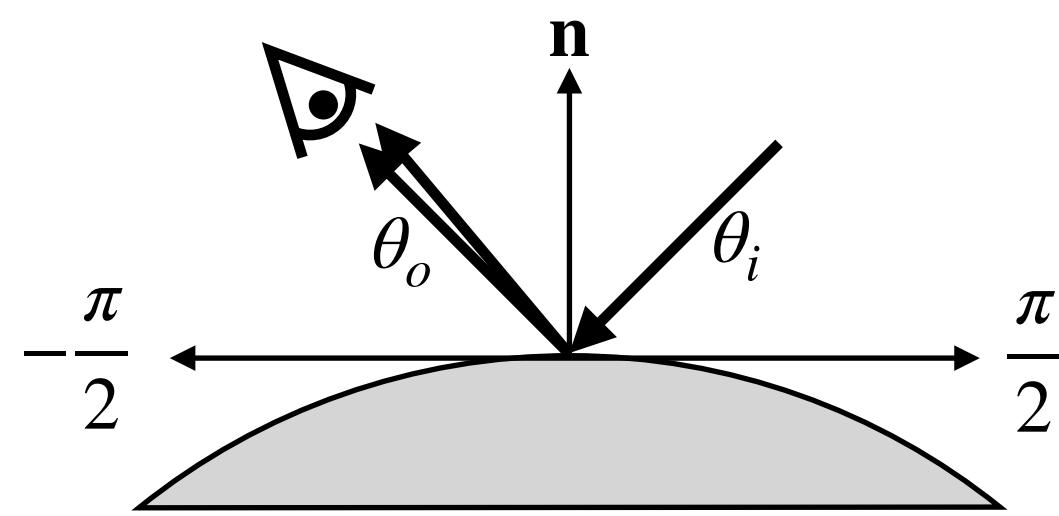


BRDF intuition

perfect mirror

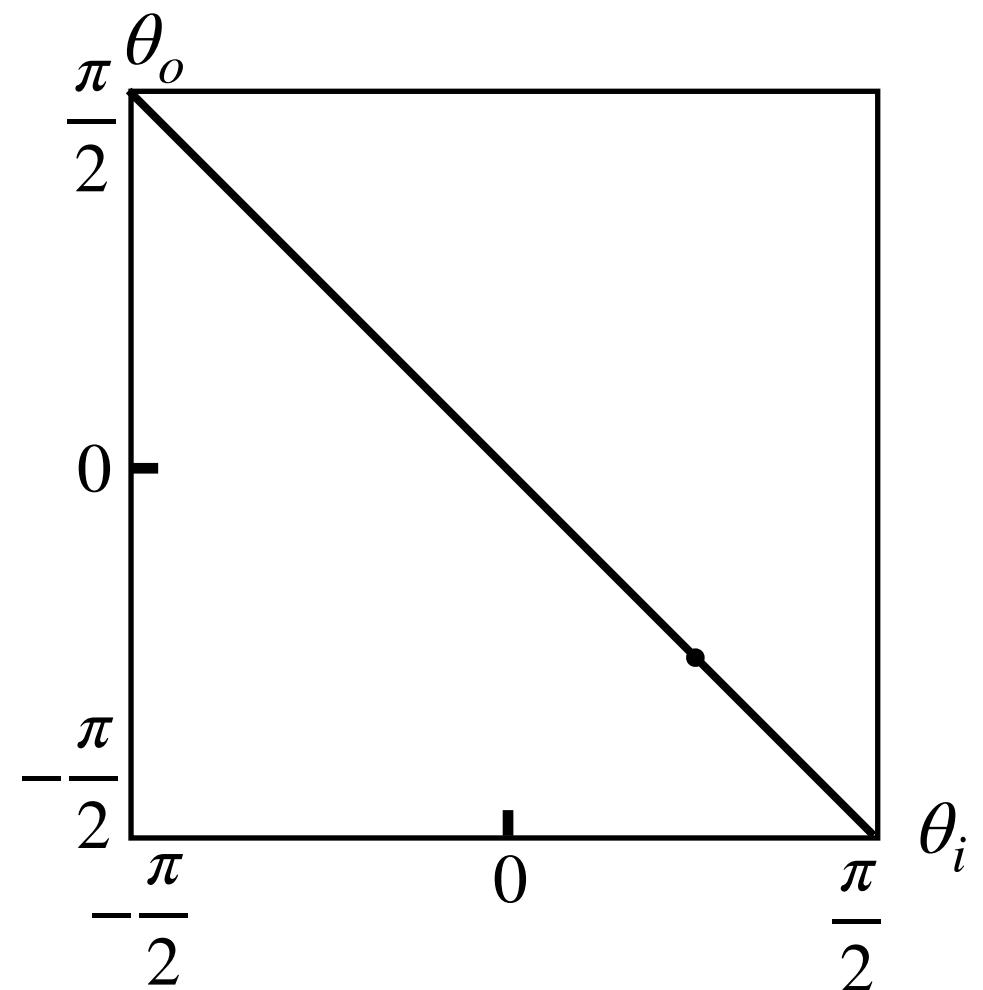
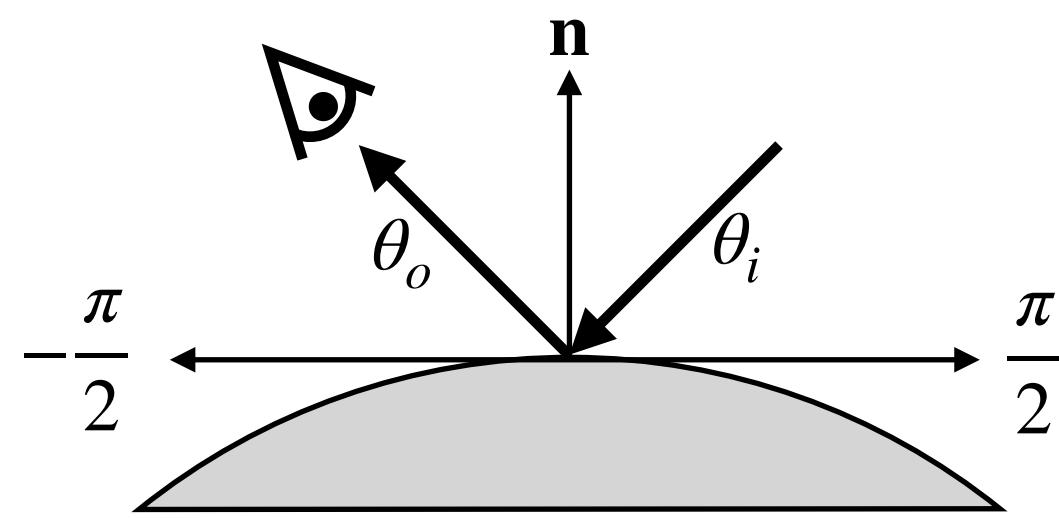


fairly shiny mirror

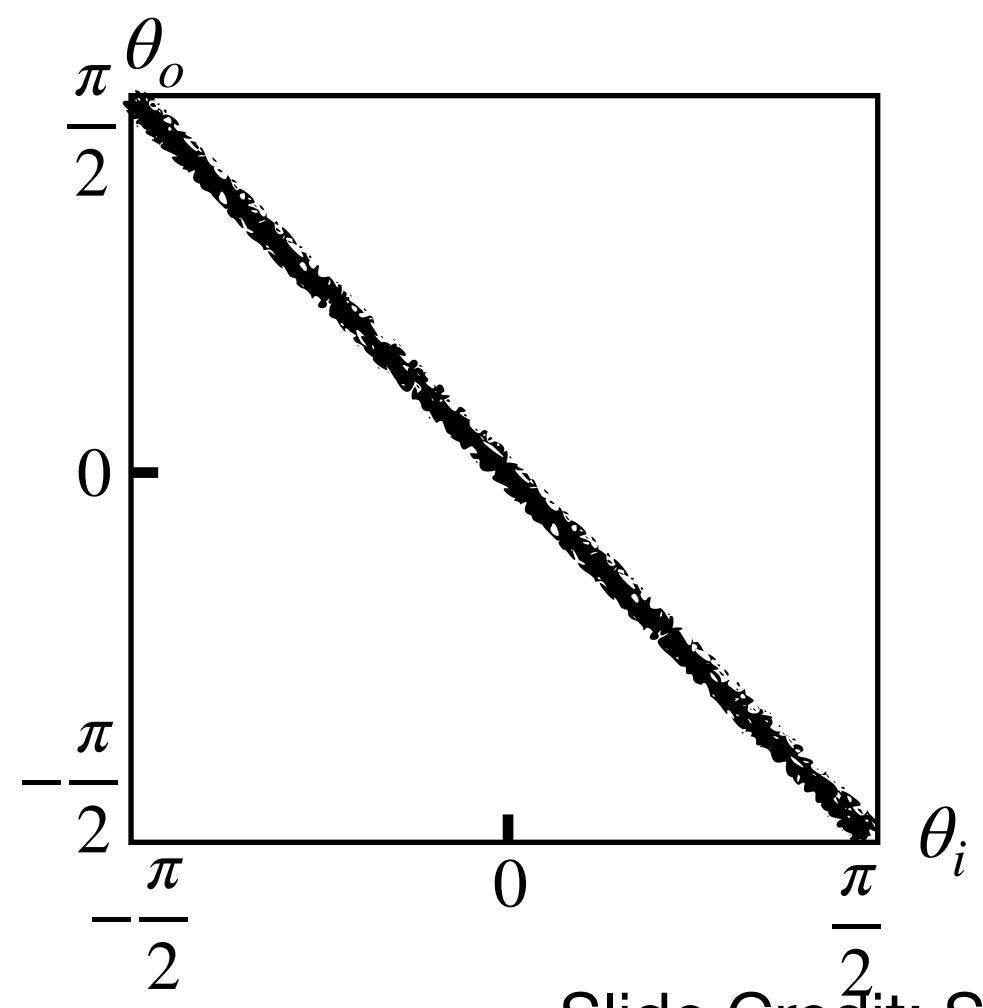
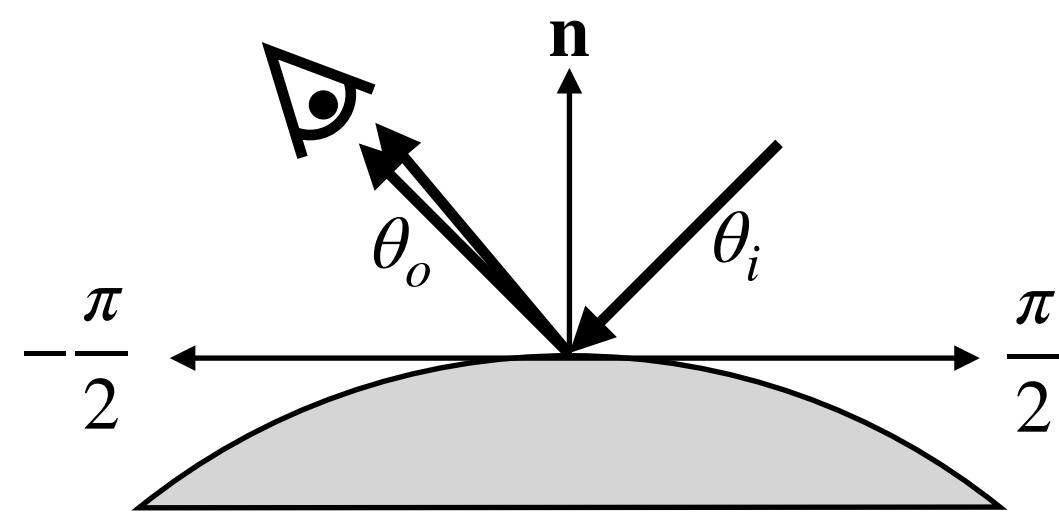


BRDF intuition

perfect mirror

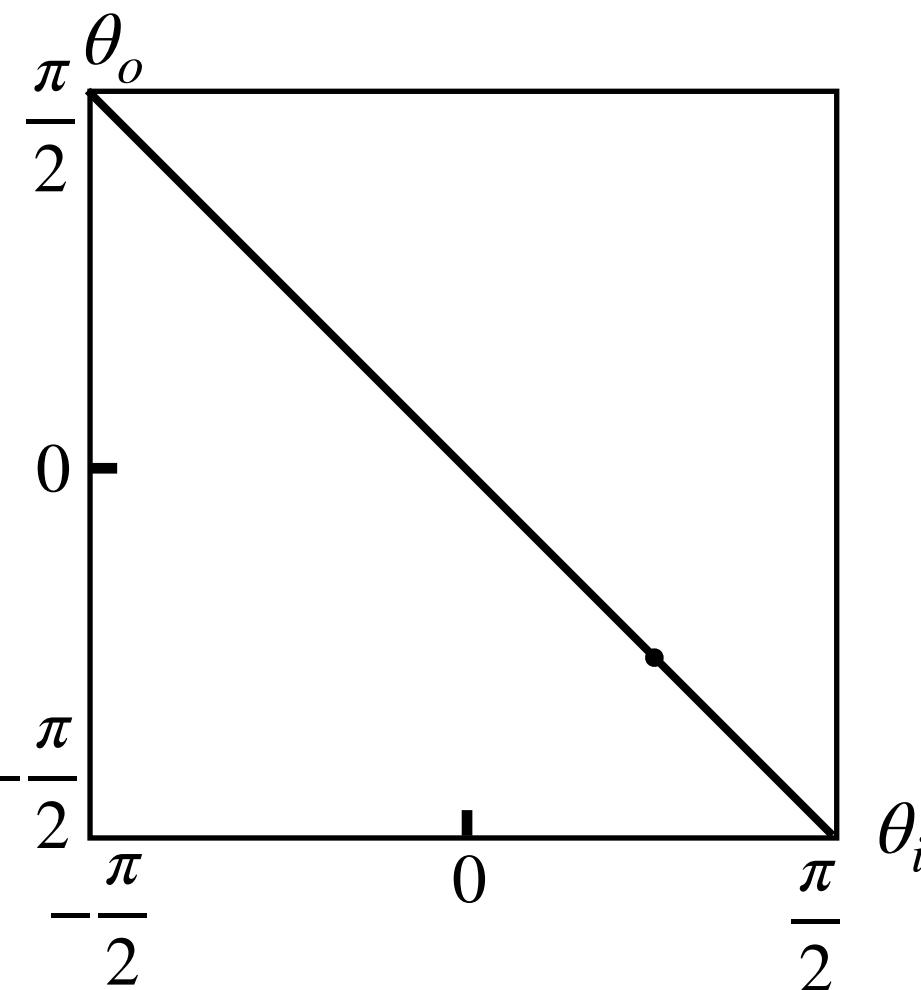
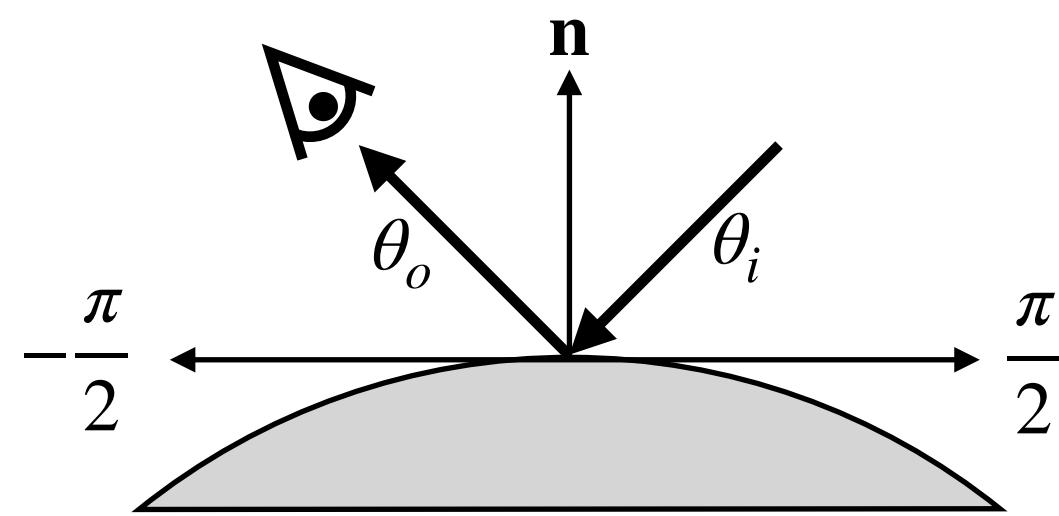


fairly shiny mirror

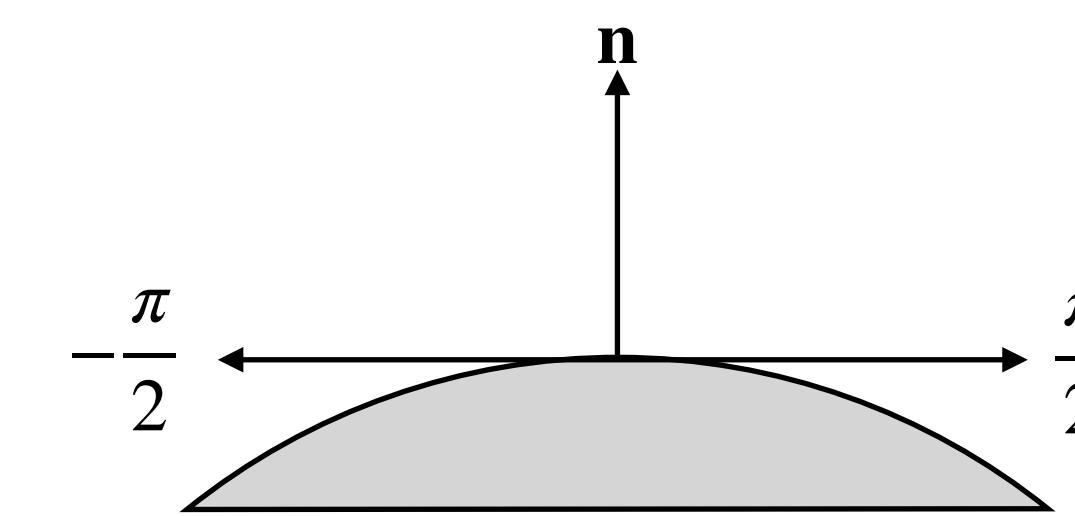


BRDF intuition

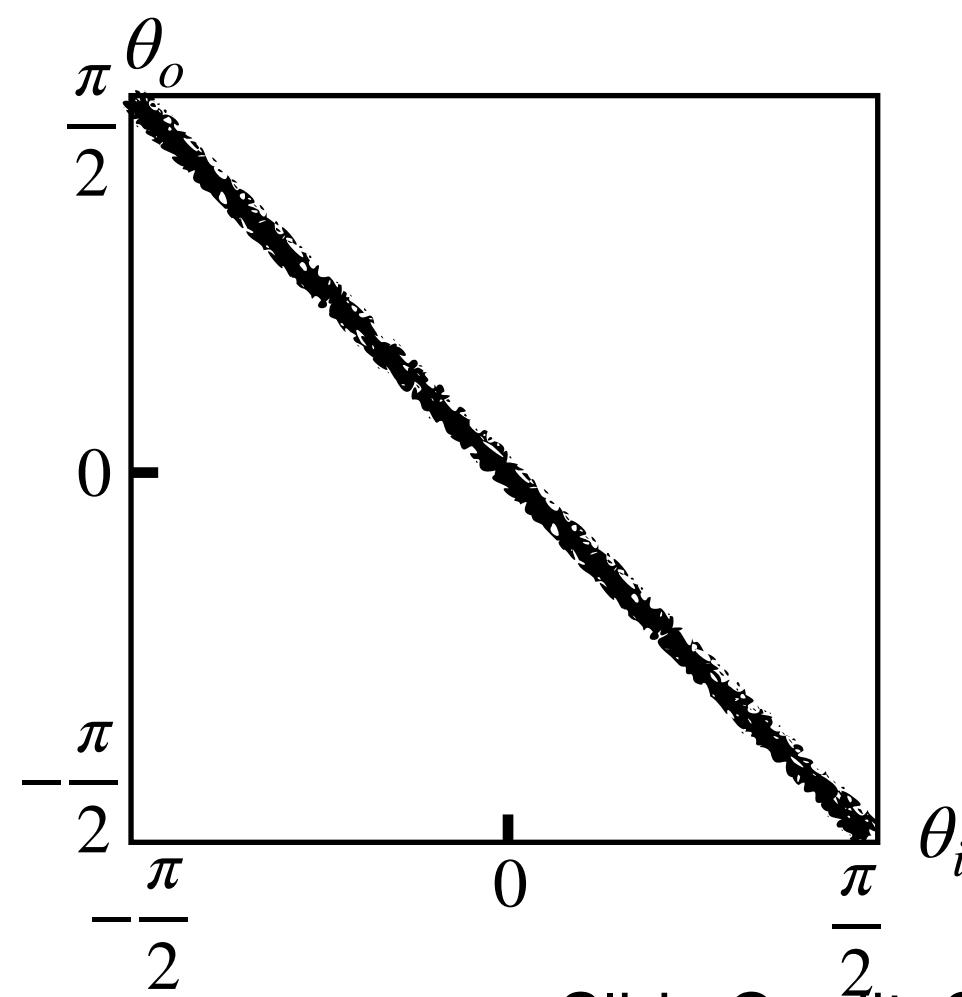
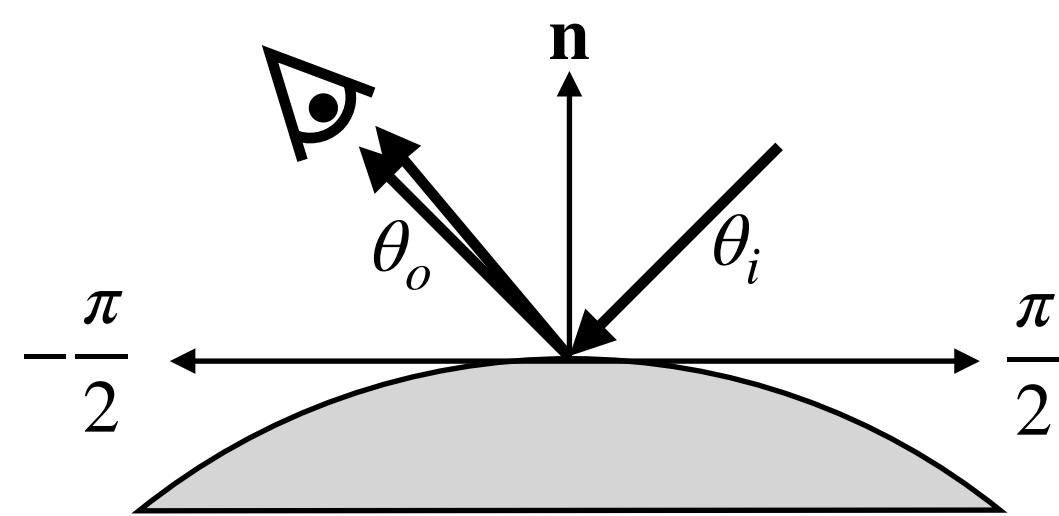
perfect mirror



retro-reflector

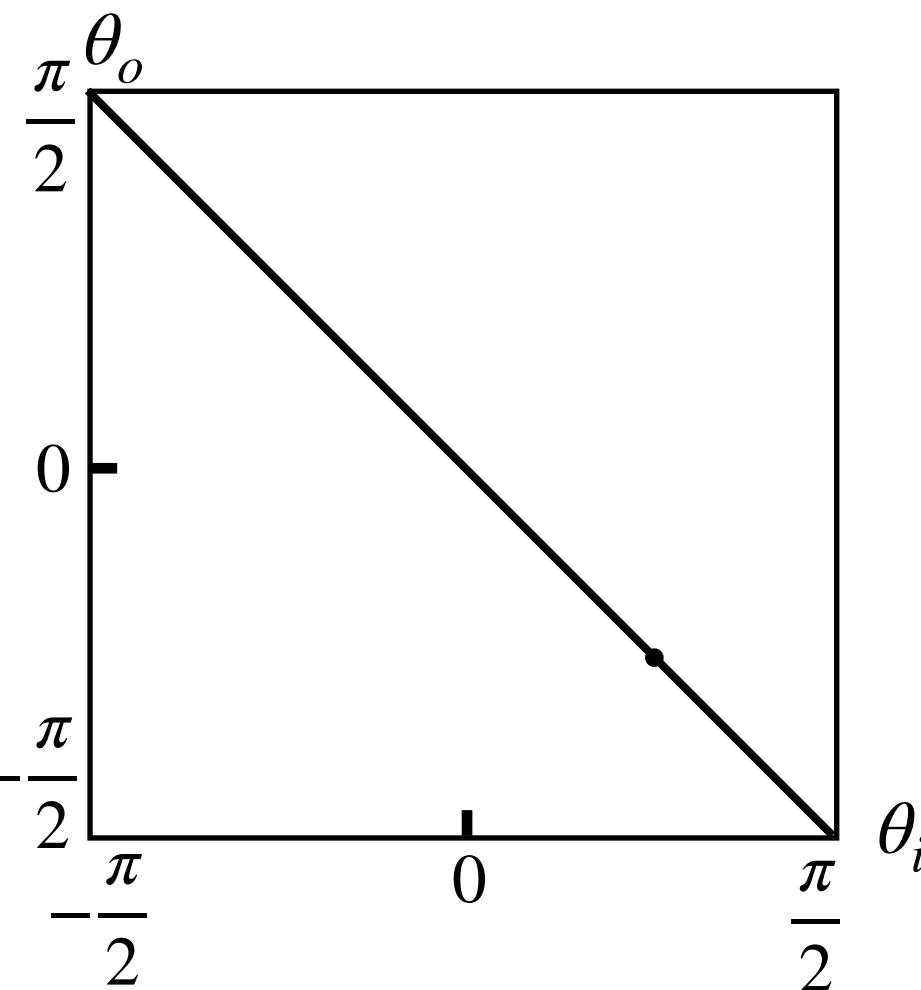
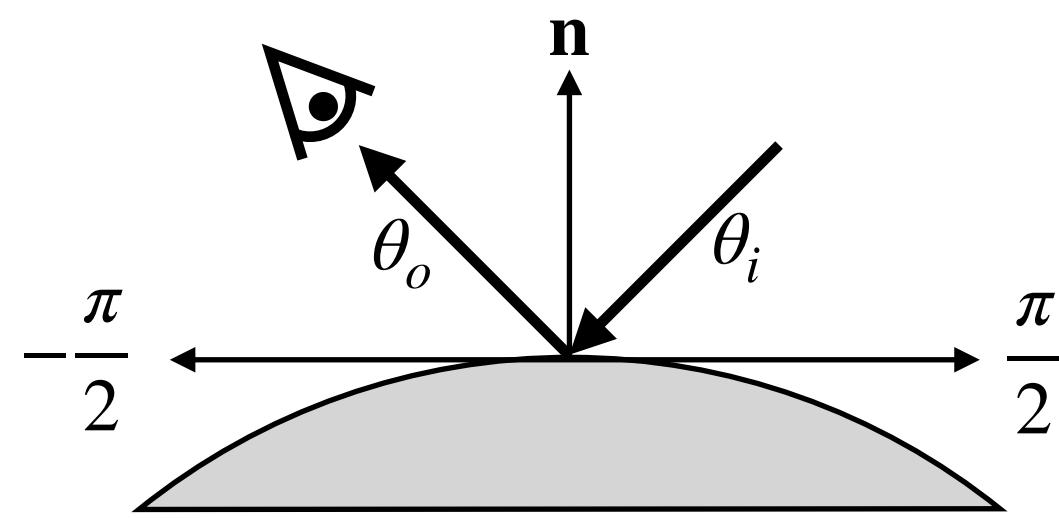


fairly shiny mirror

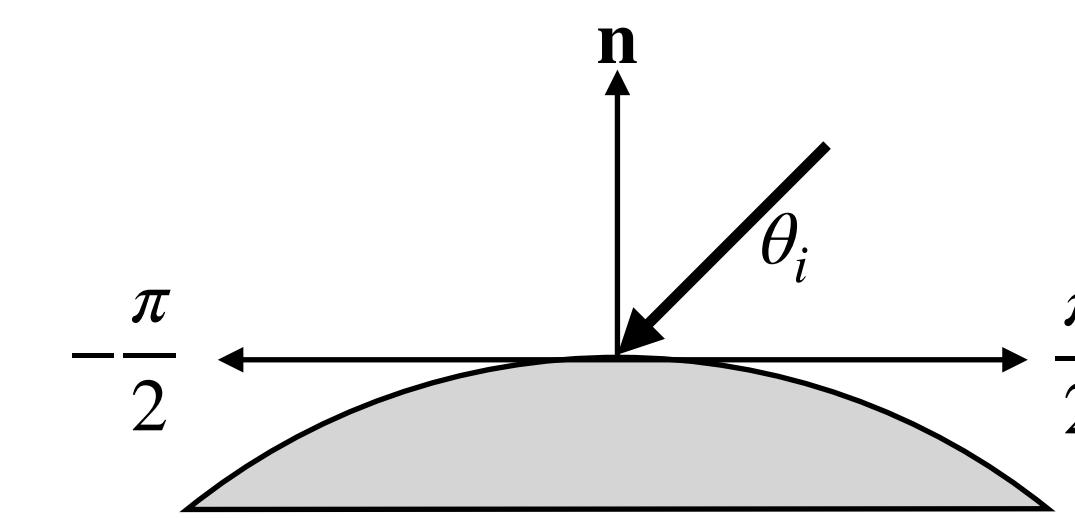


BRDF intuition

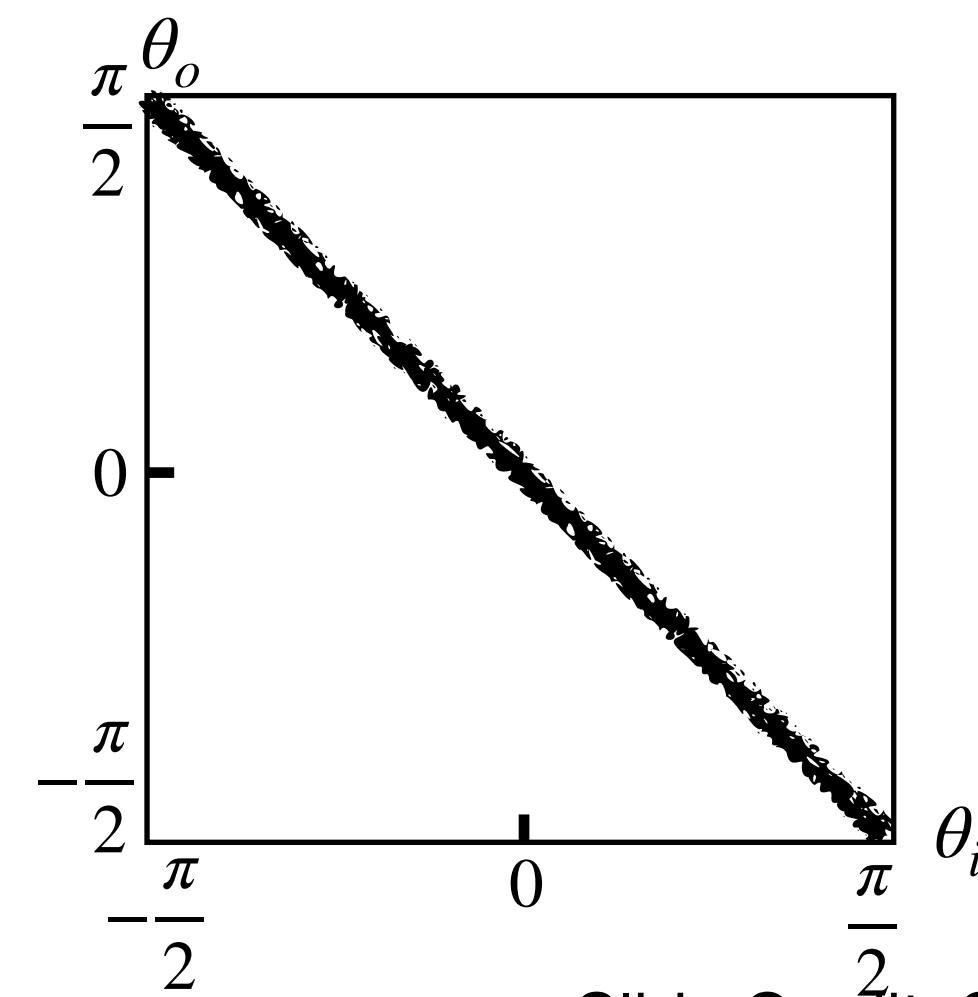
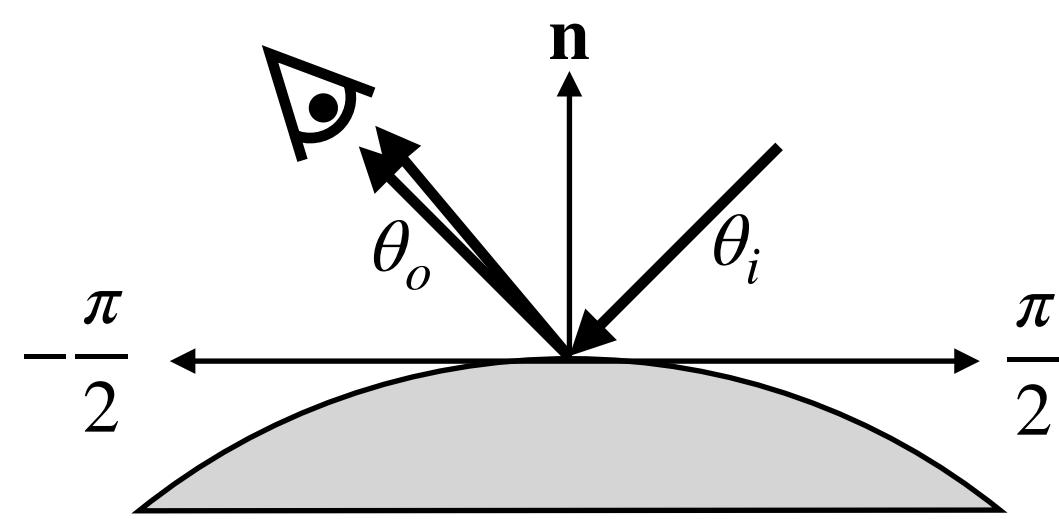
perfect mirror



retro-reflector

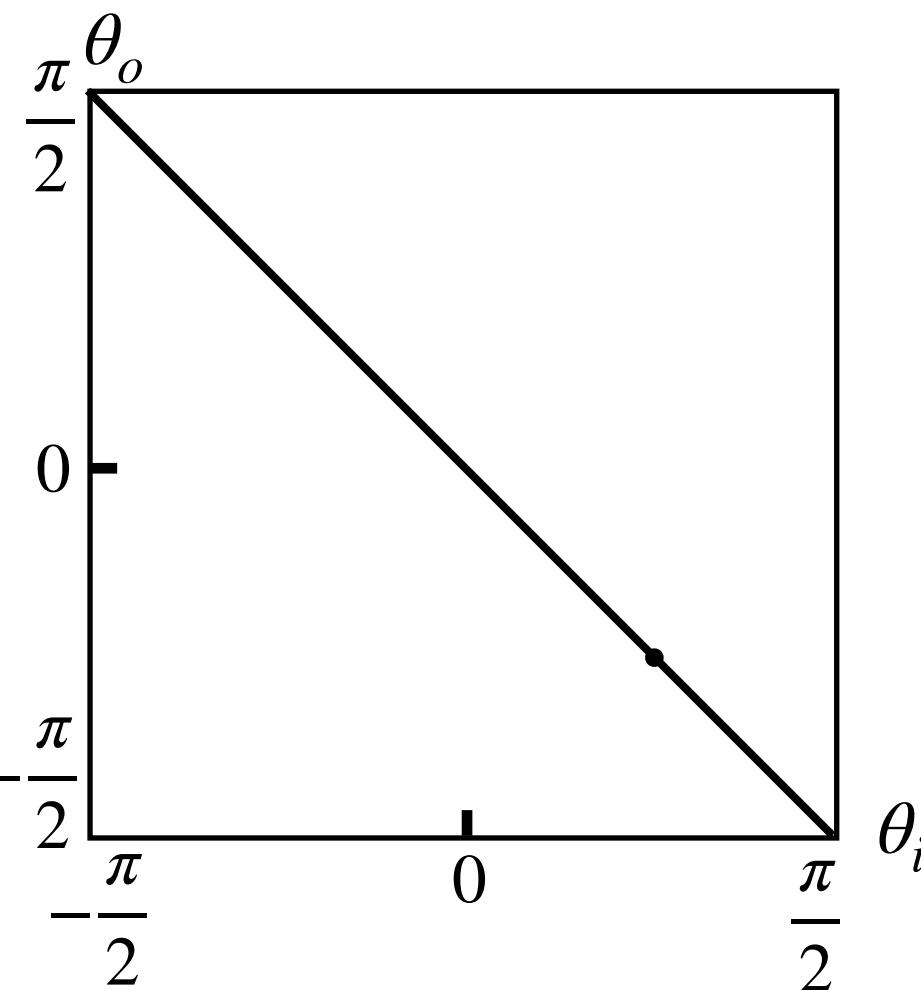
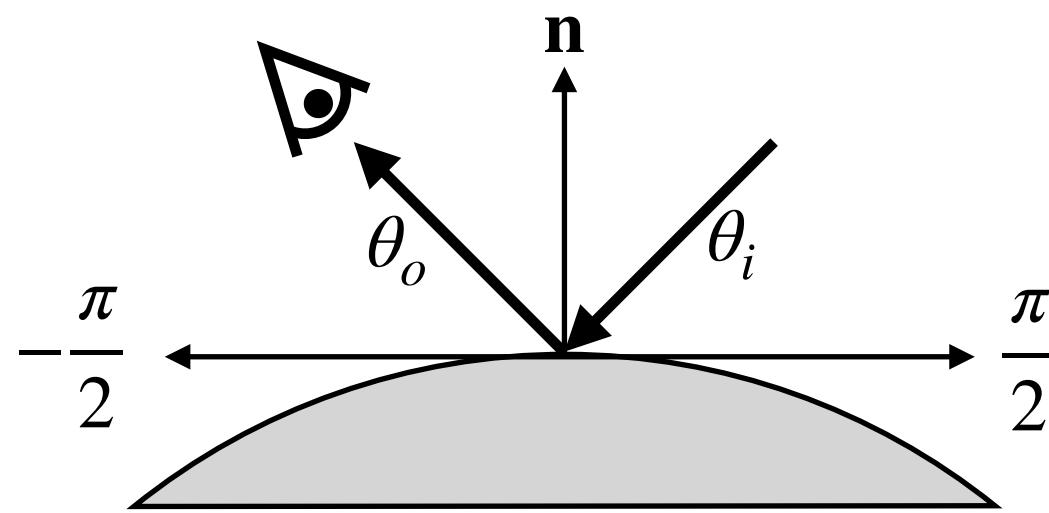


fairly shiny mirror

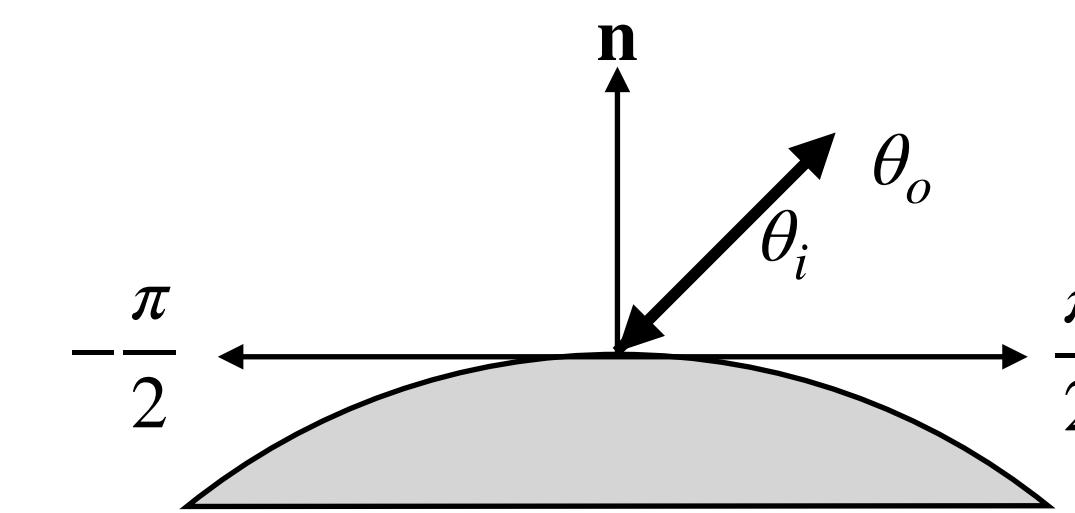


BRDF intuition

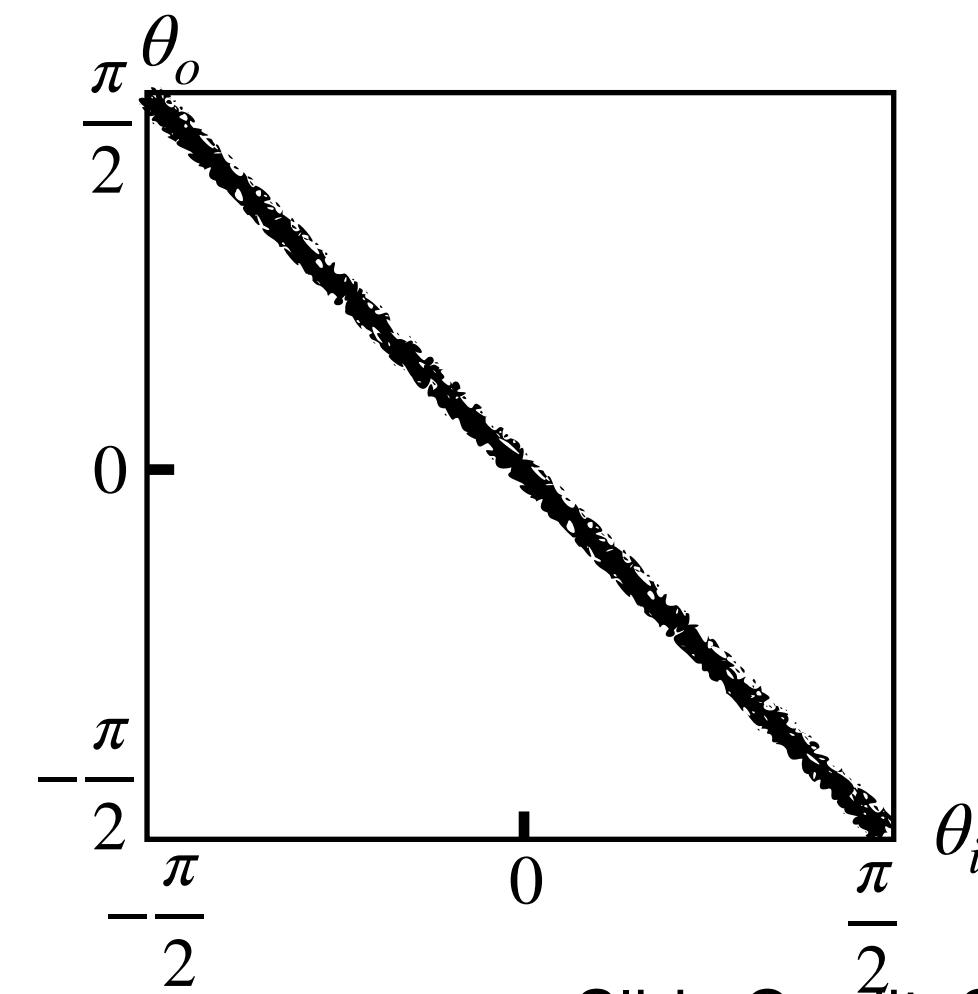
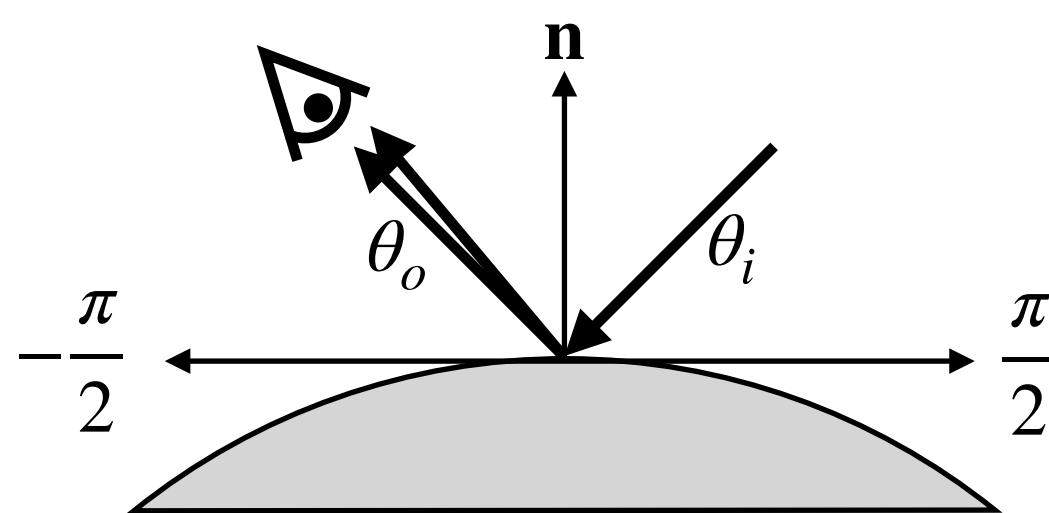
perfect mirror



retro-reflector

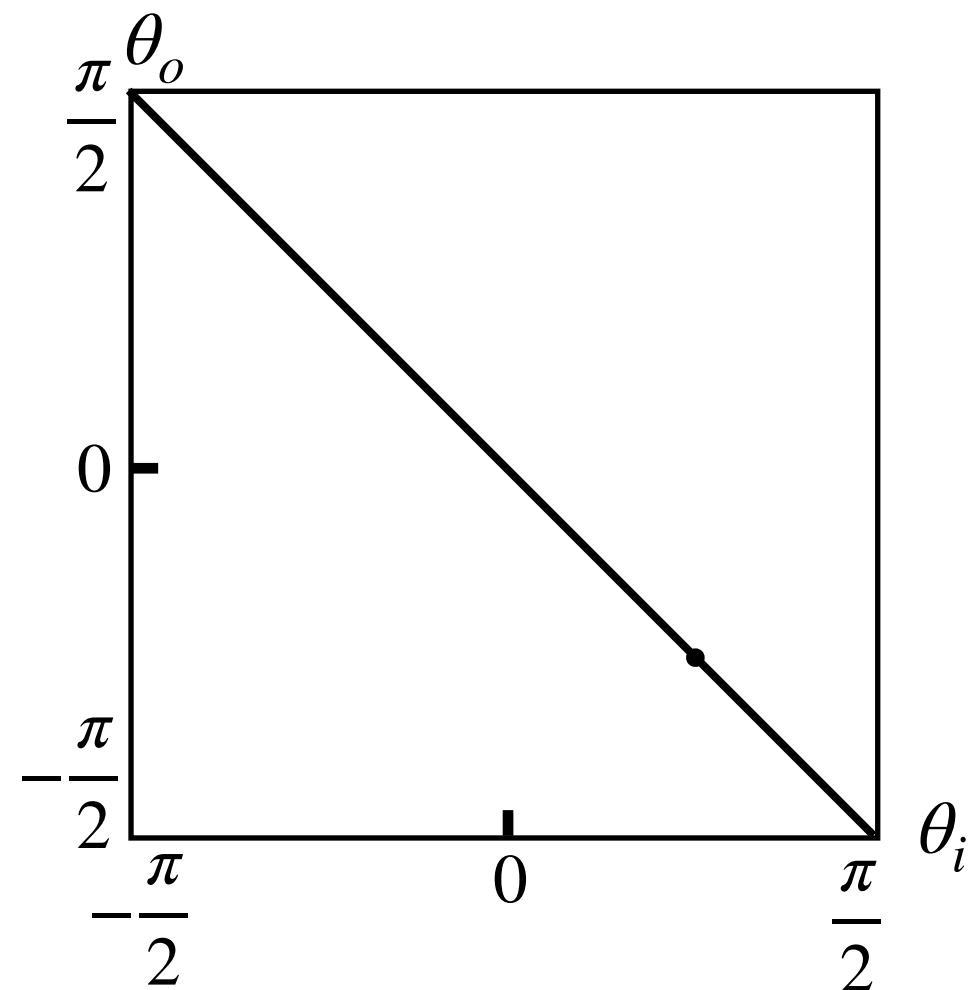
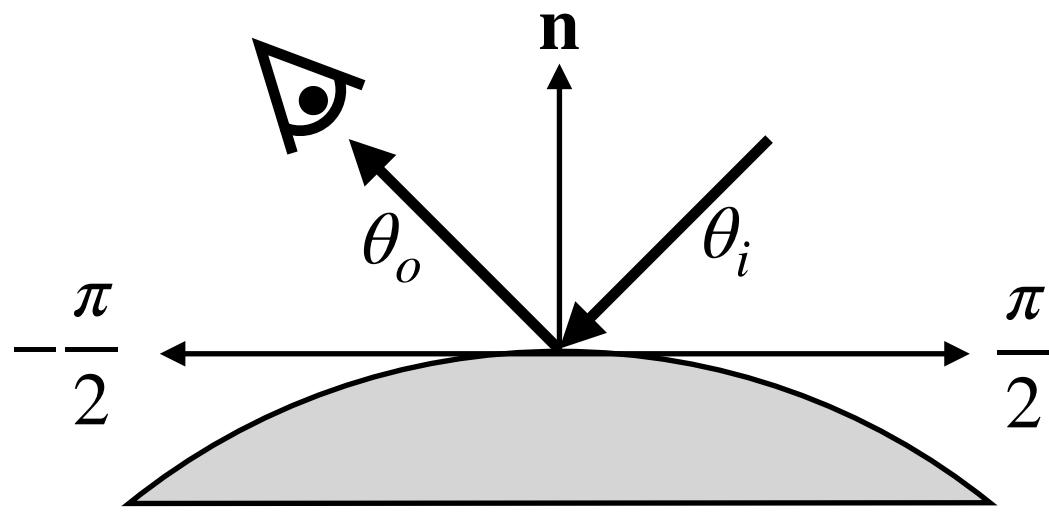


fairly shiny mirror

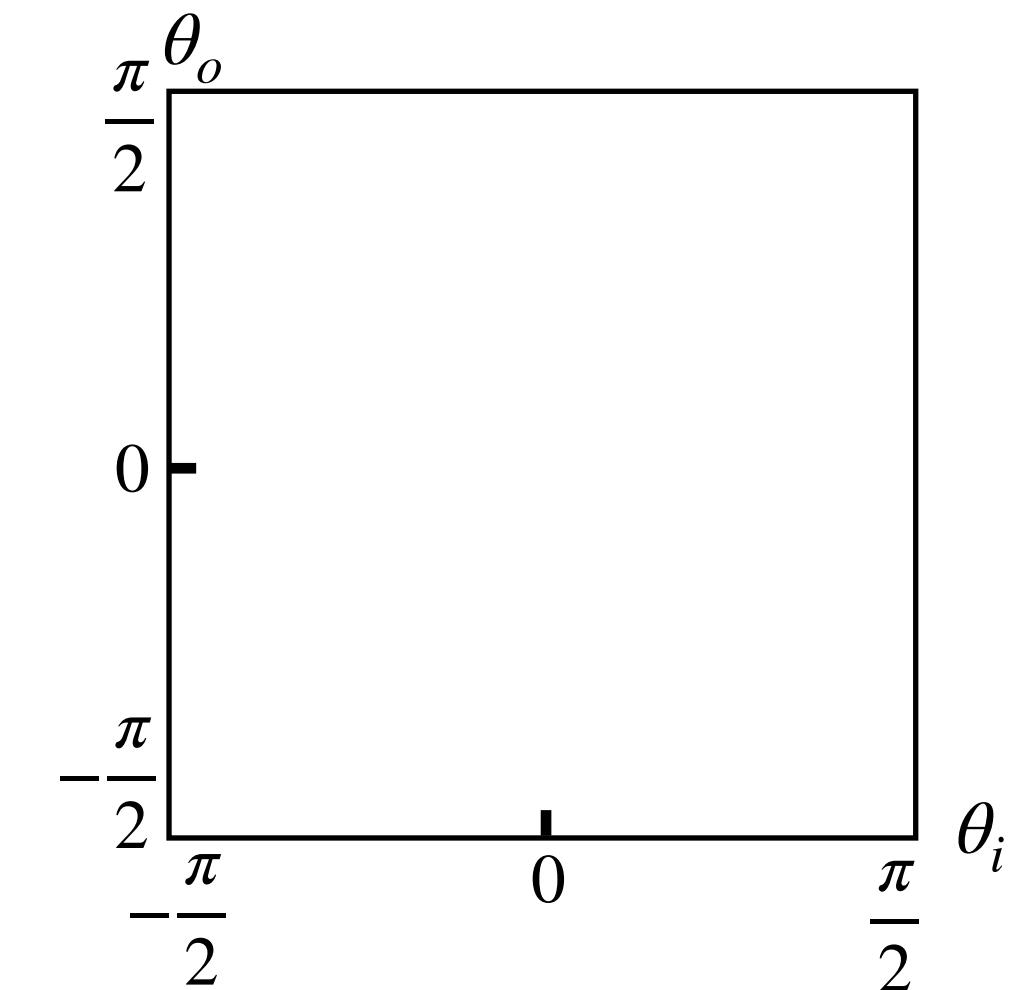
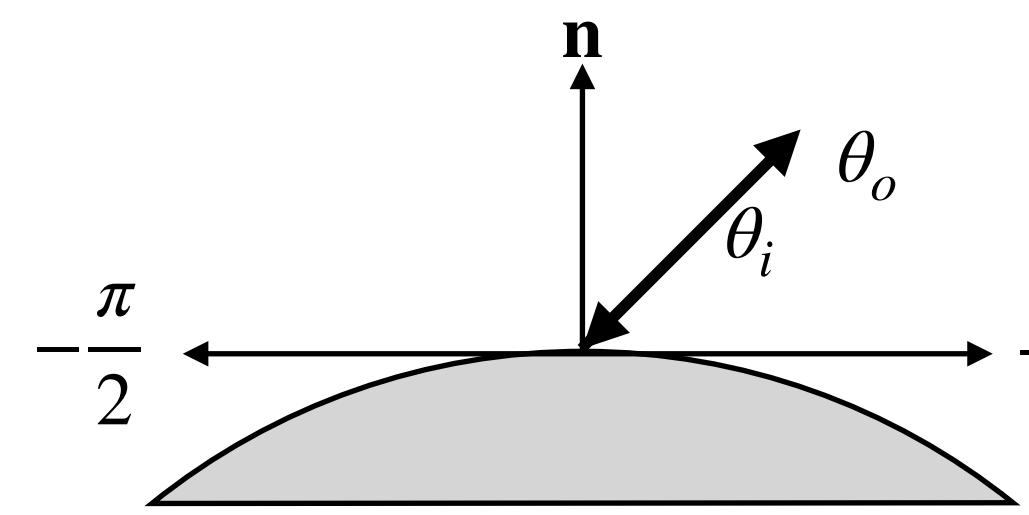


BRDF intuition

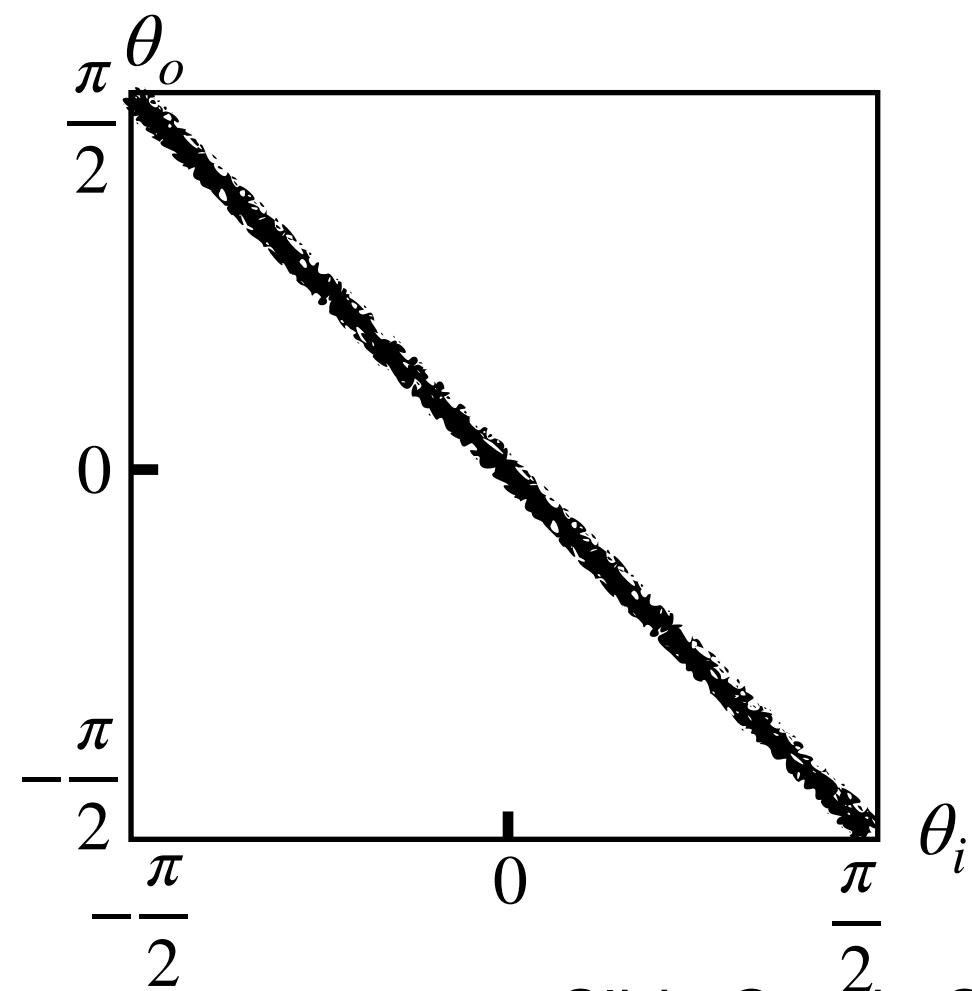
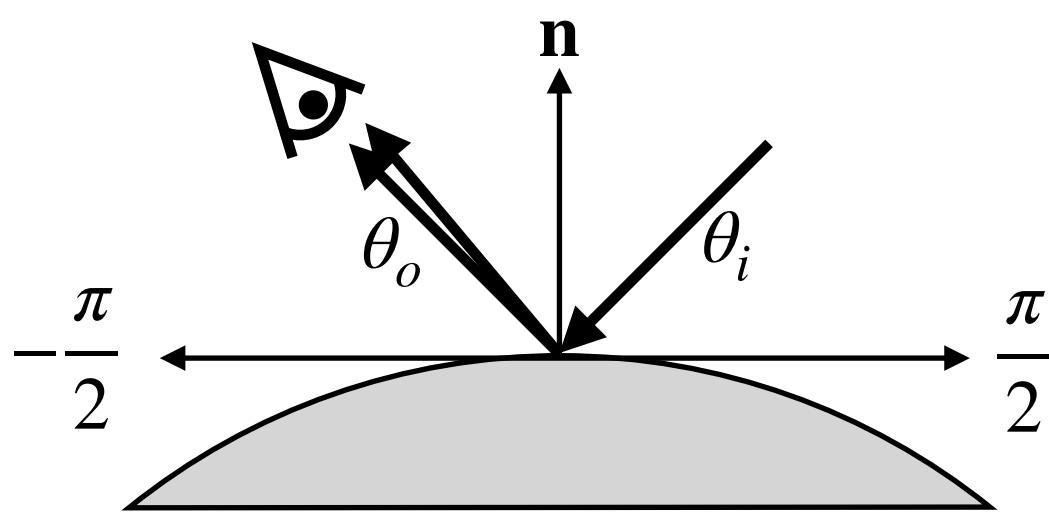
perfect mirror



retro-reflector

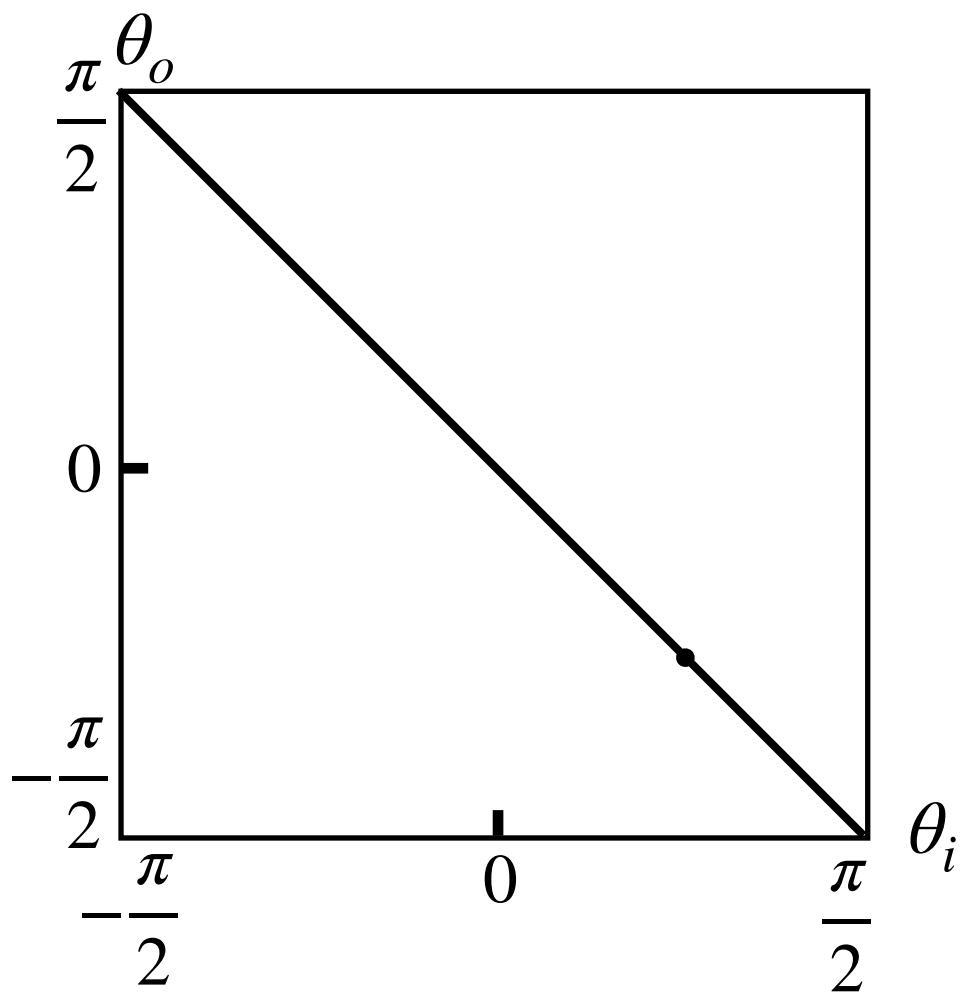
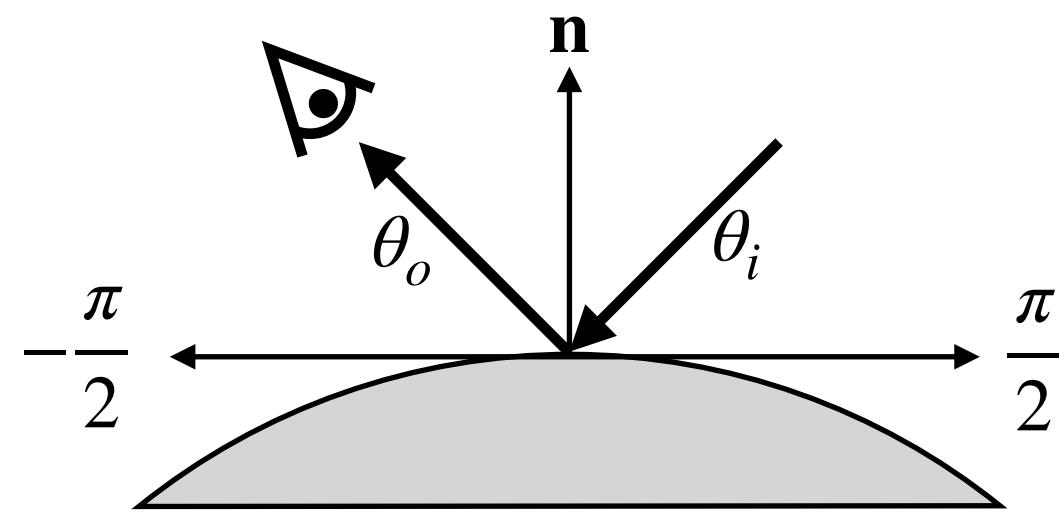


fairly shiny mirror

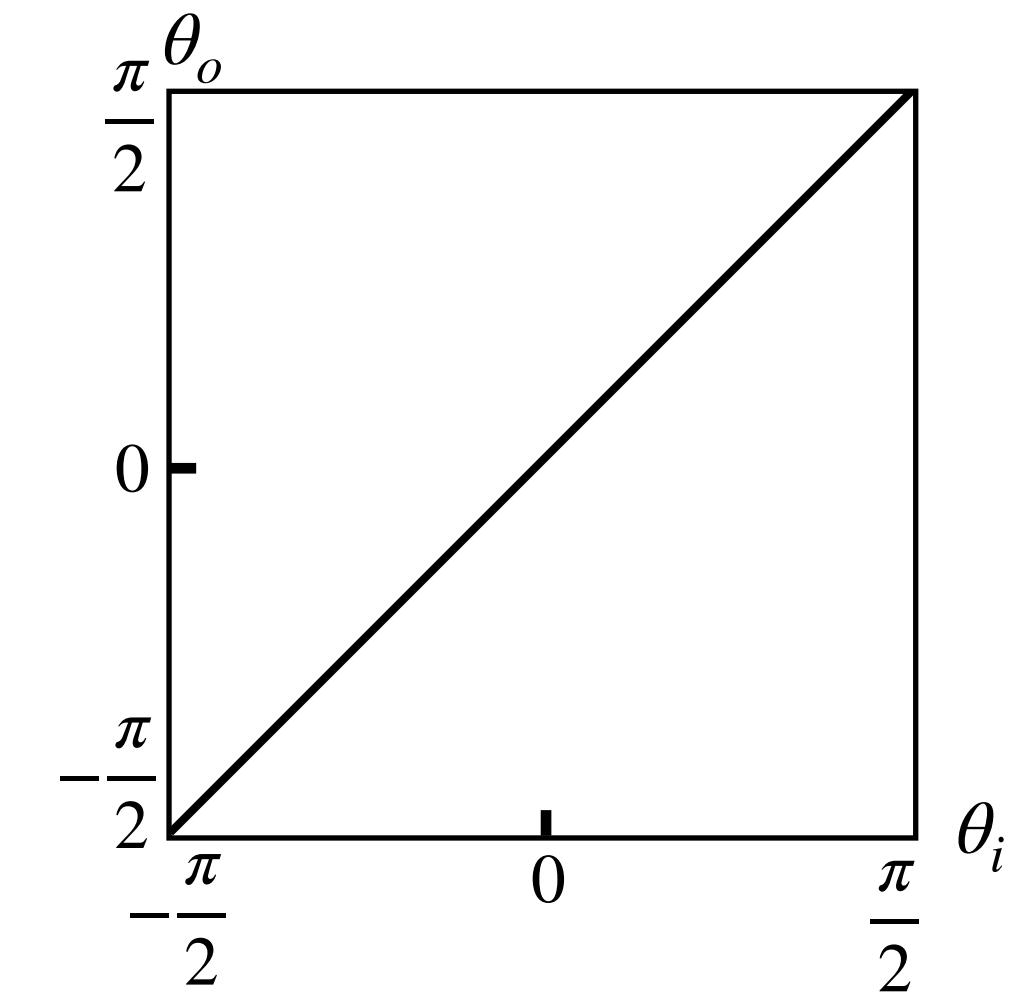
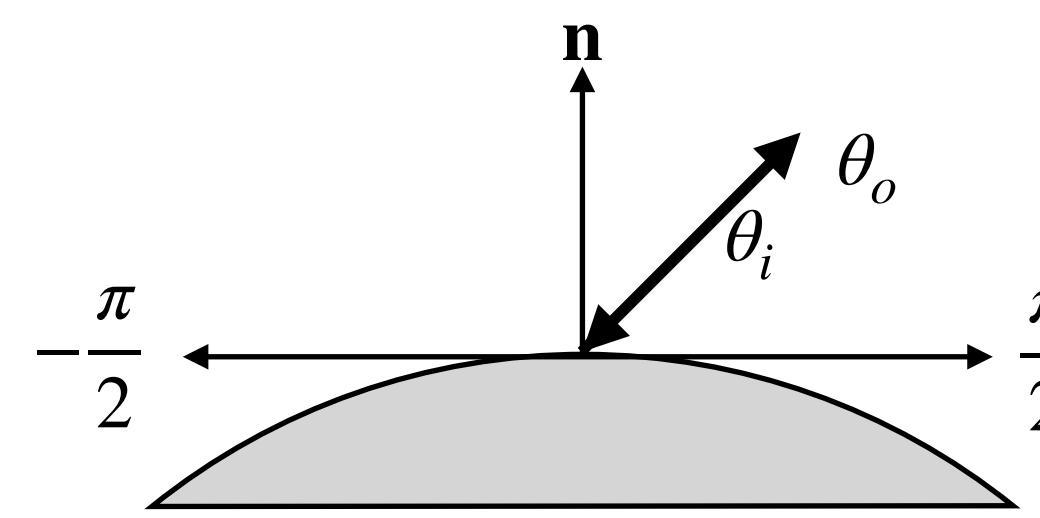


BRDF intuition

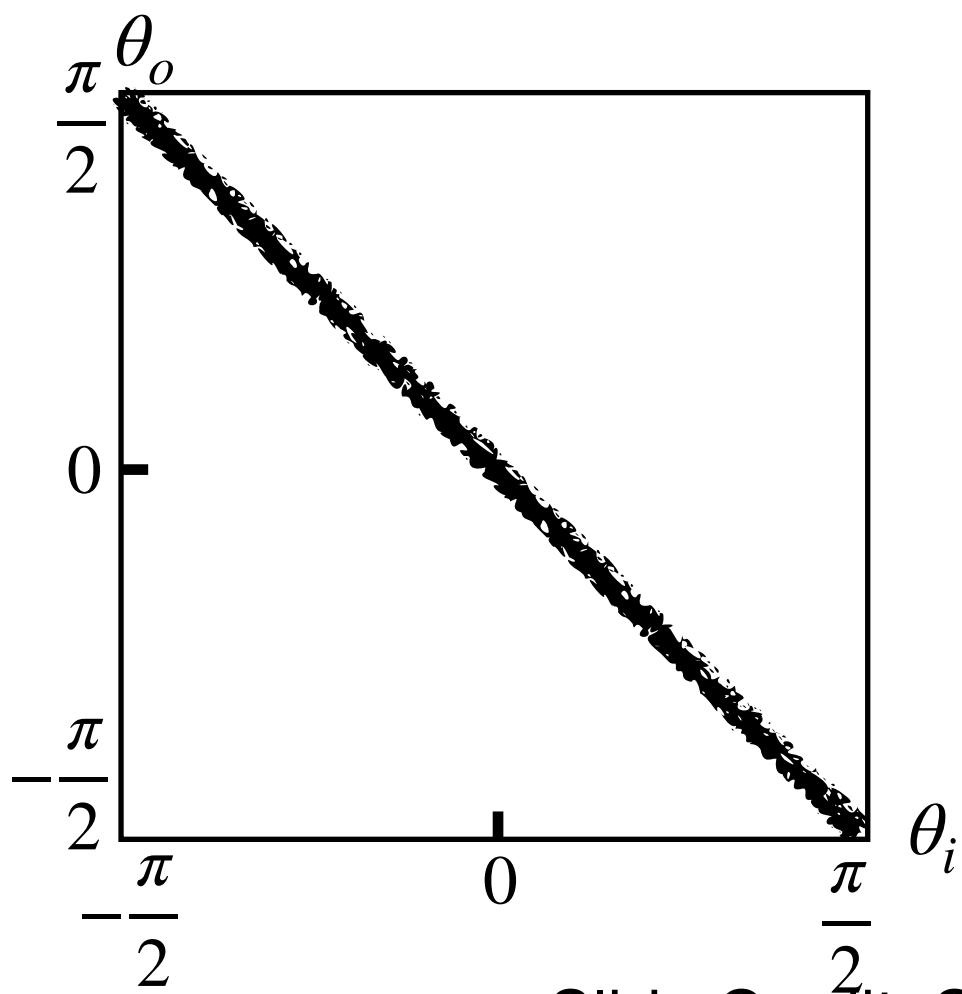
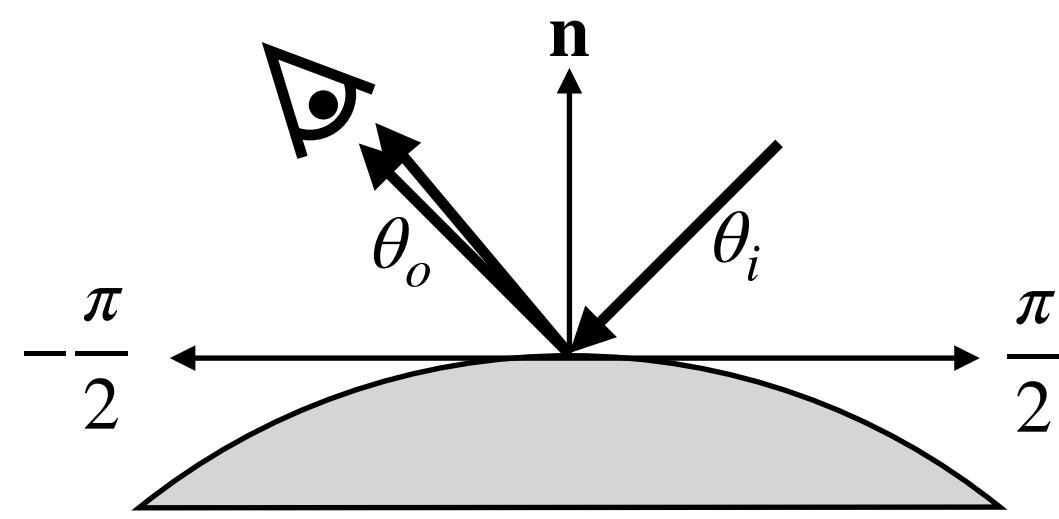
perfect mirror



retro-reflector

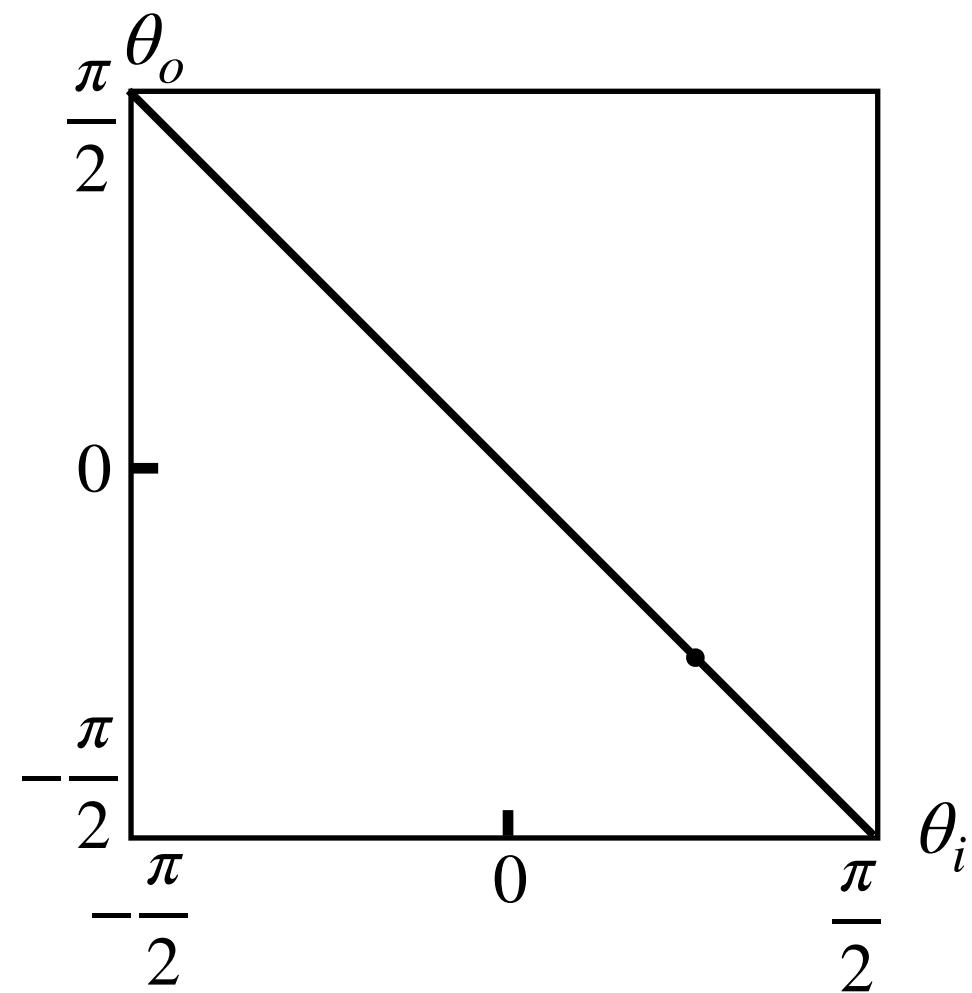
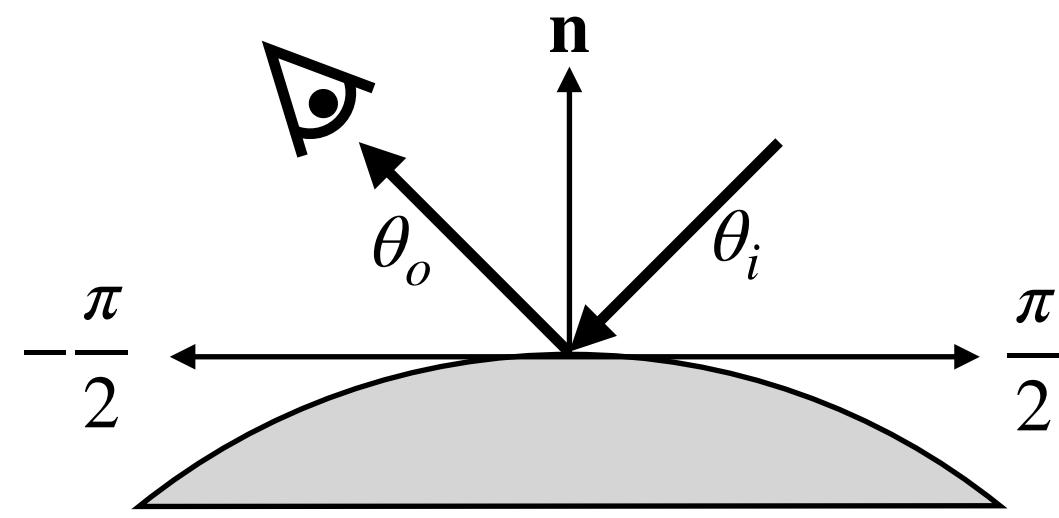


fairly shiny mirror

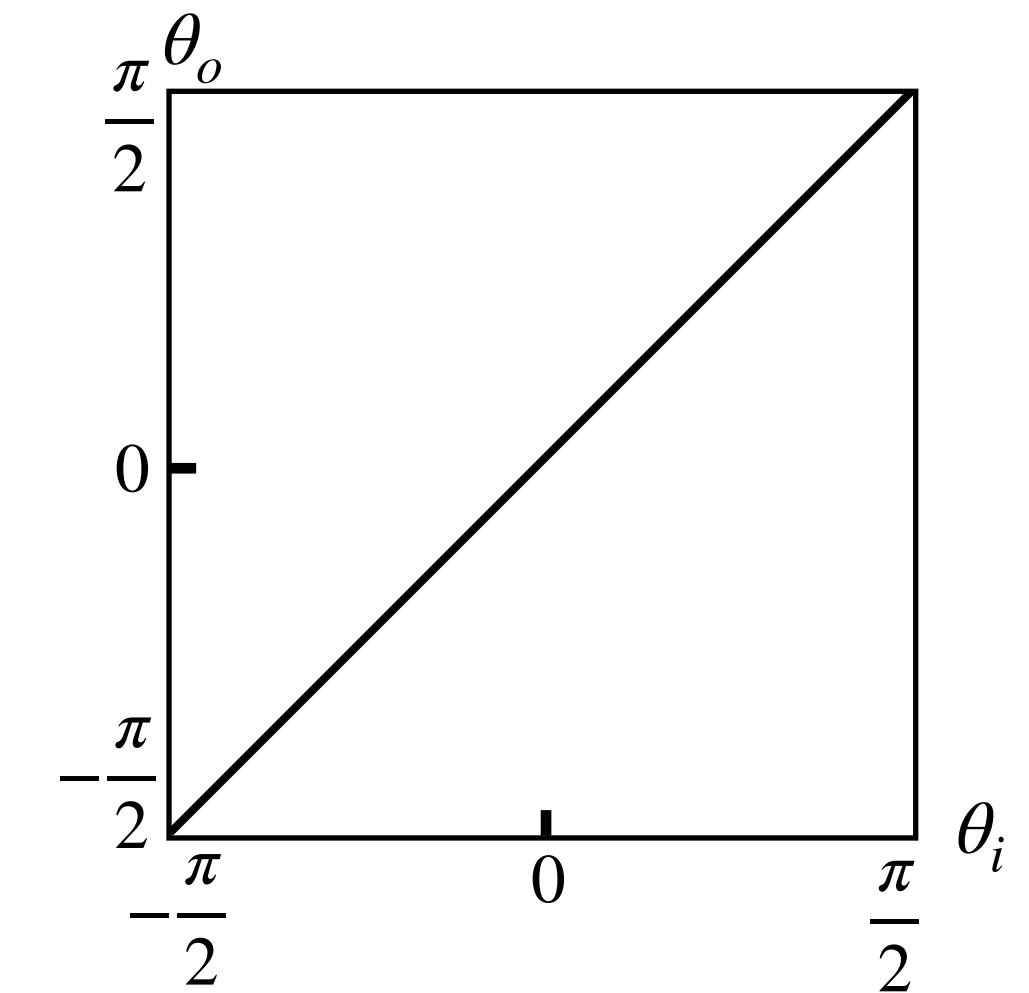
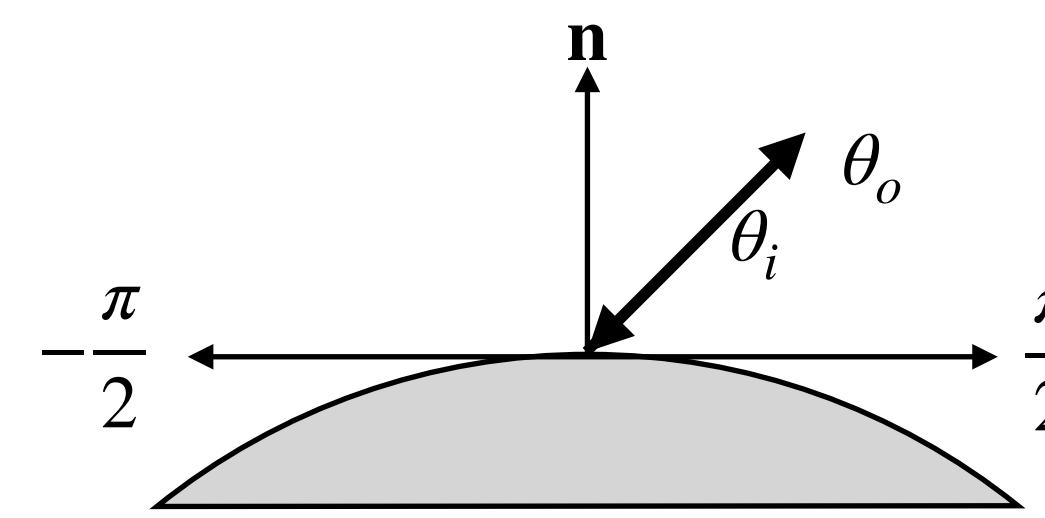


BRDF intuition

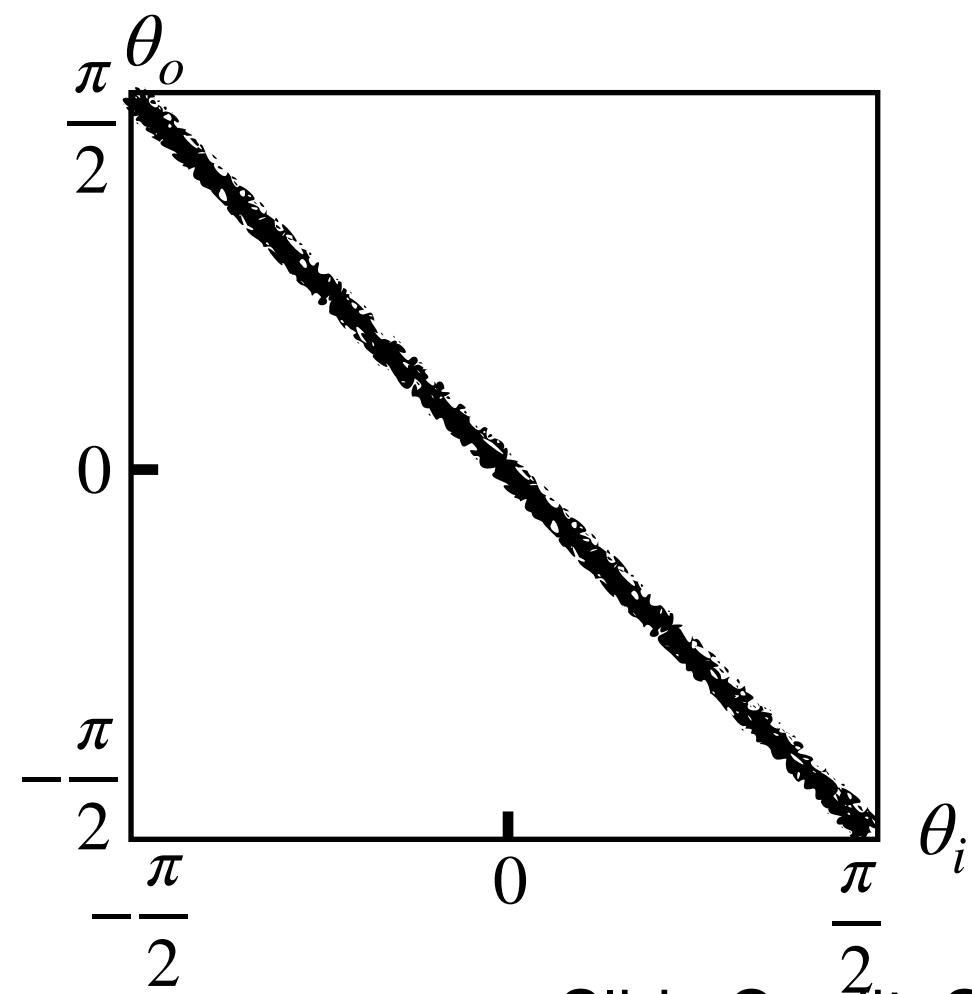
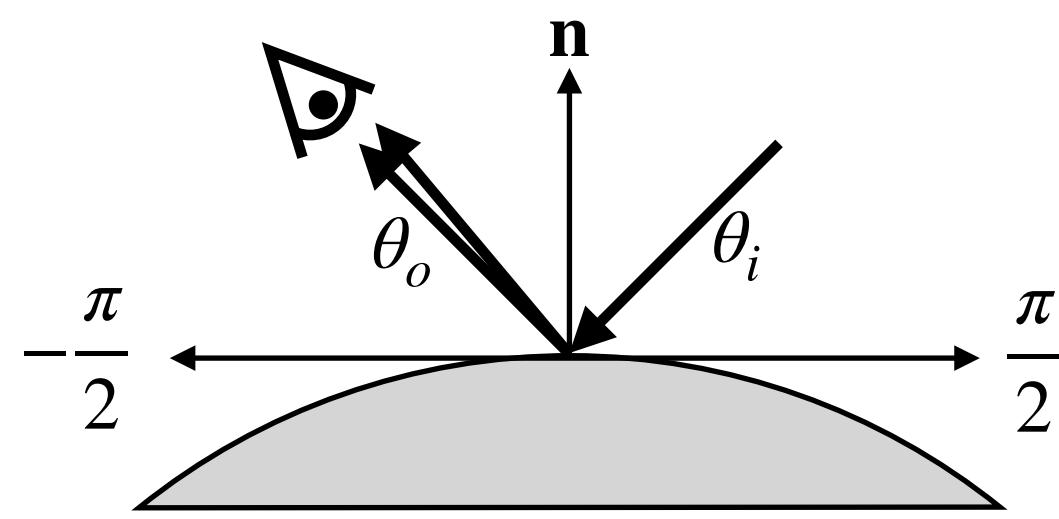
perfect mirror



retro-reflector



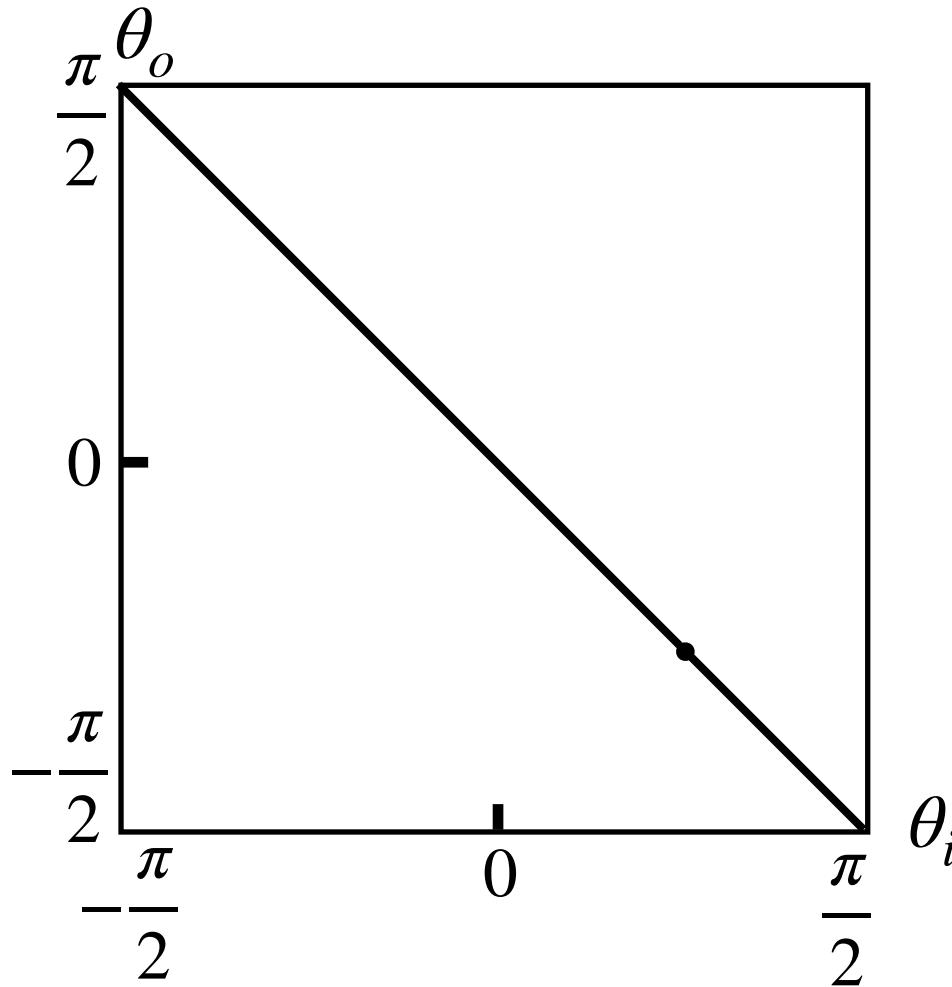
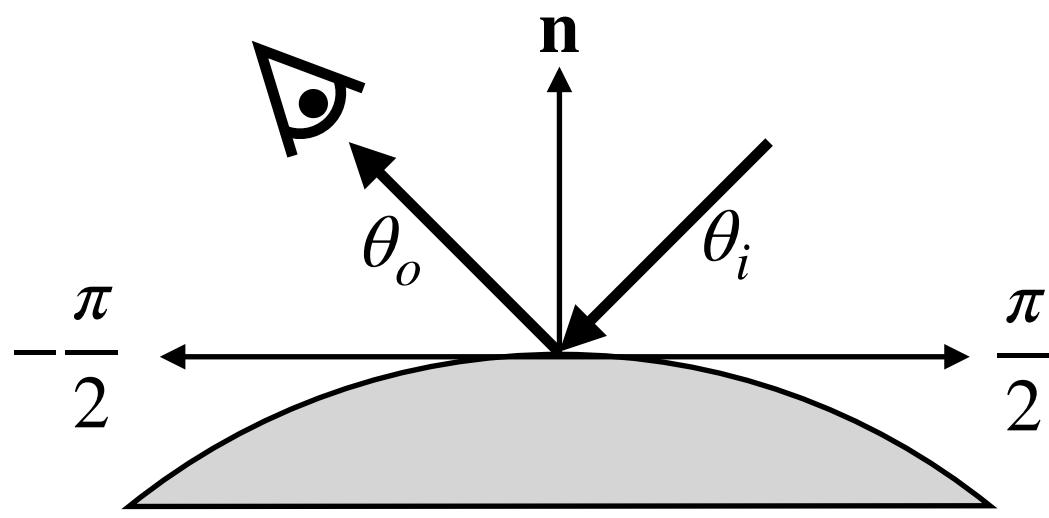
fairly shiny mirror



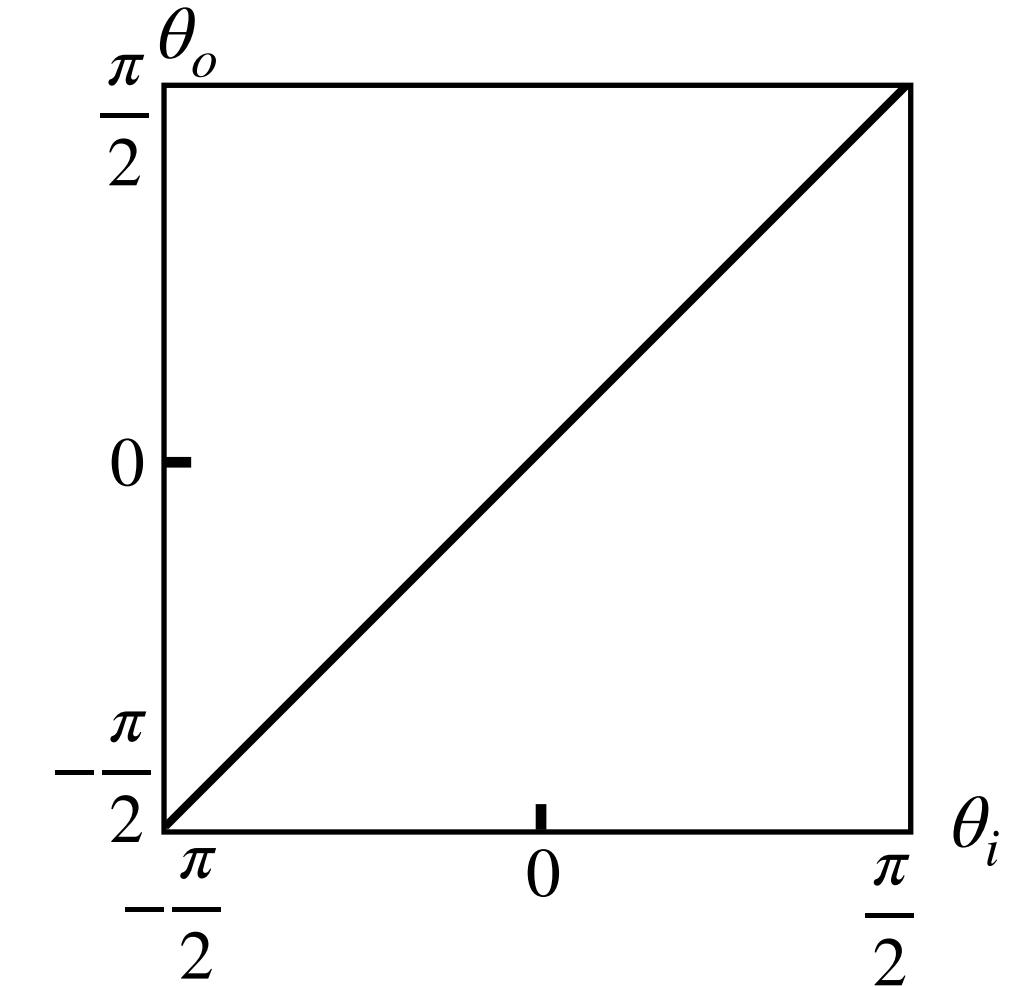
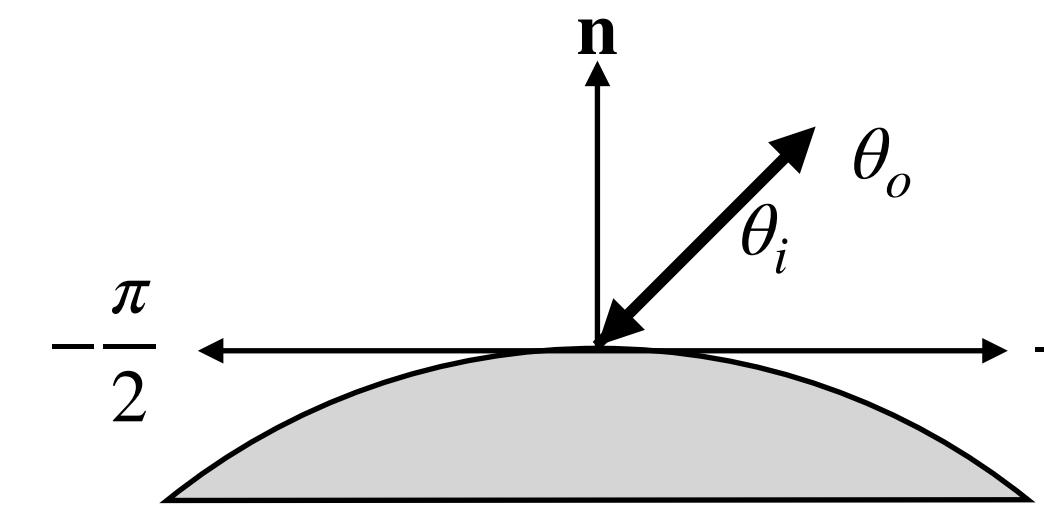
How does BRDF
for paper look like?

BRDF intuition

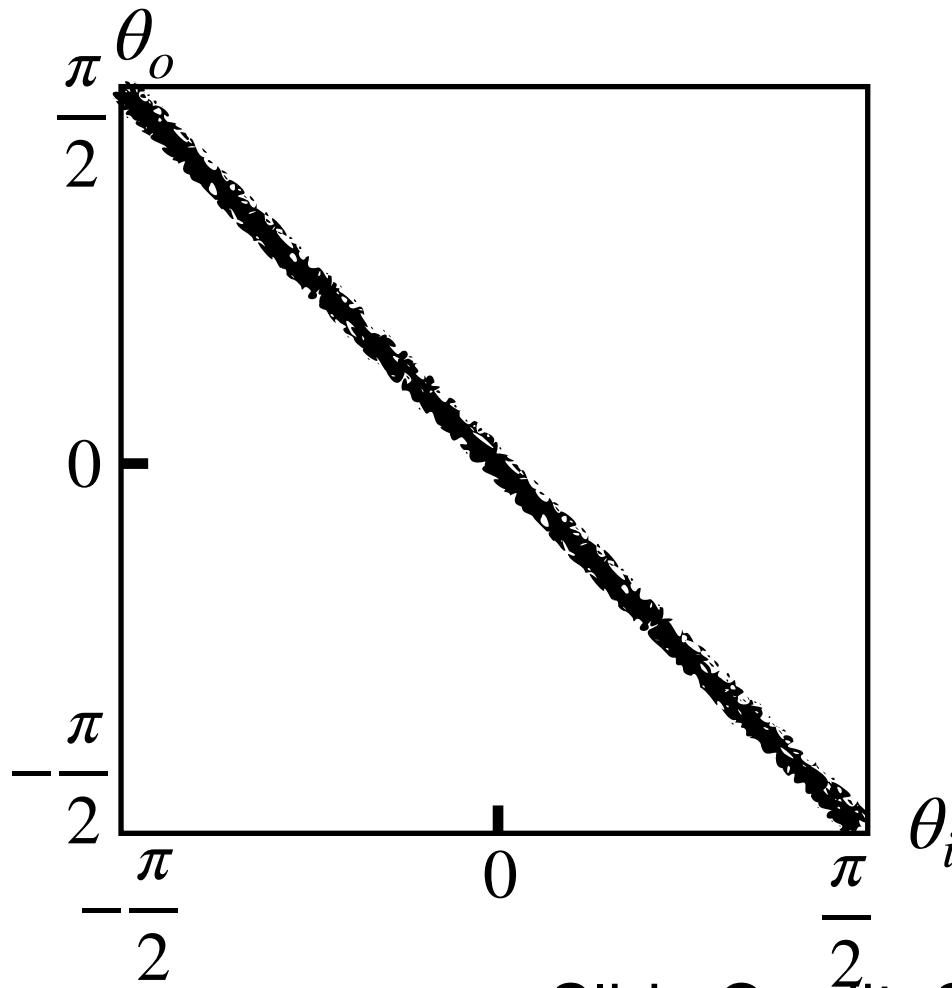
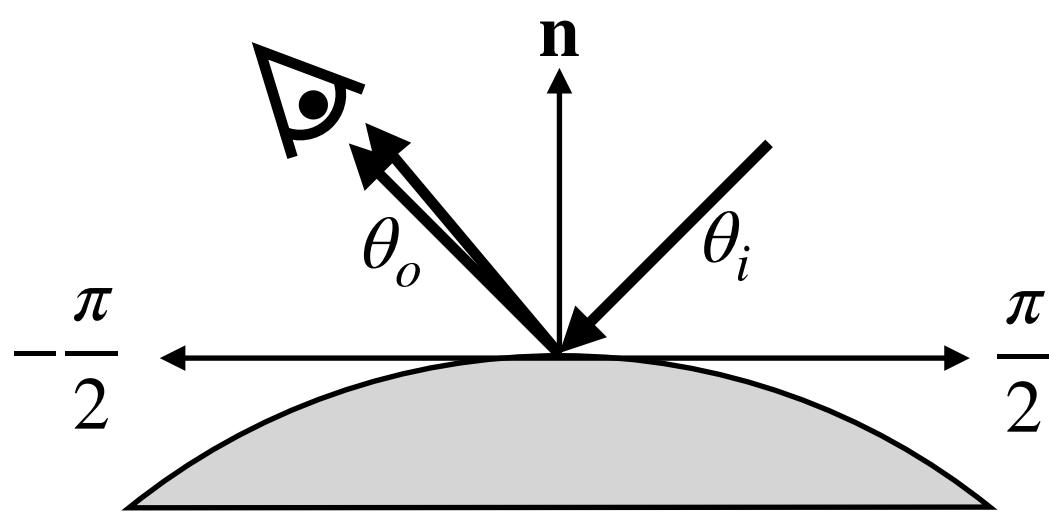
perfect mirror



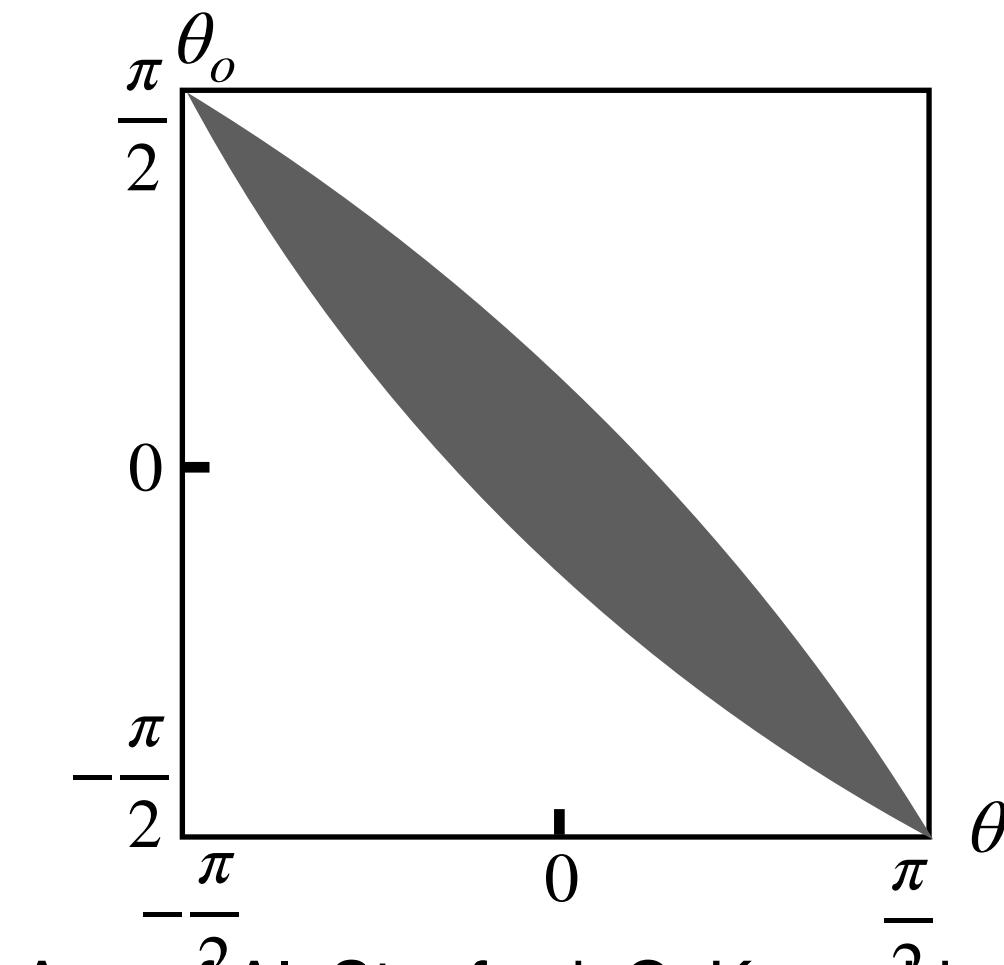
retro-reflector



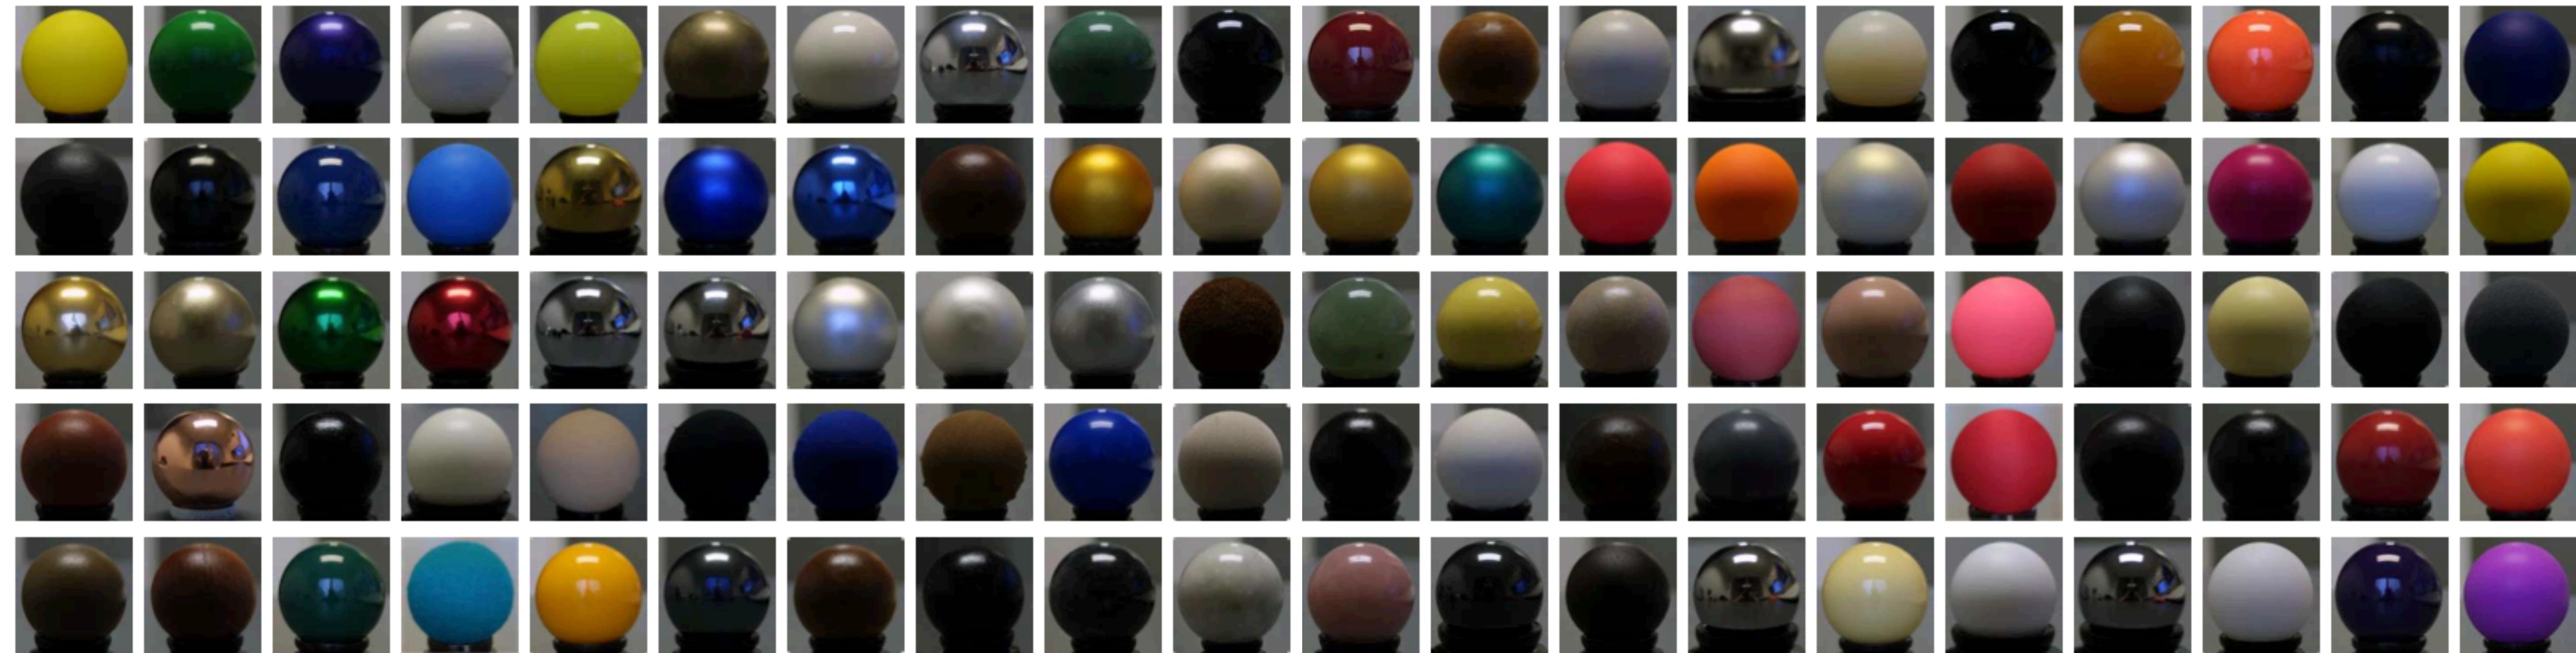
fairly shiny mirror



How does BRDF
for paper look like?

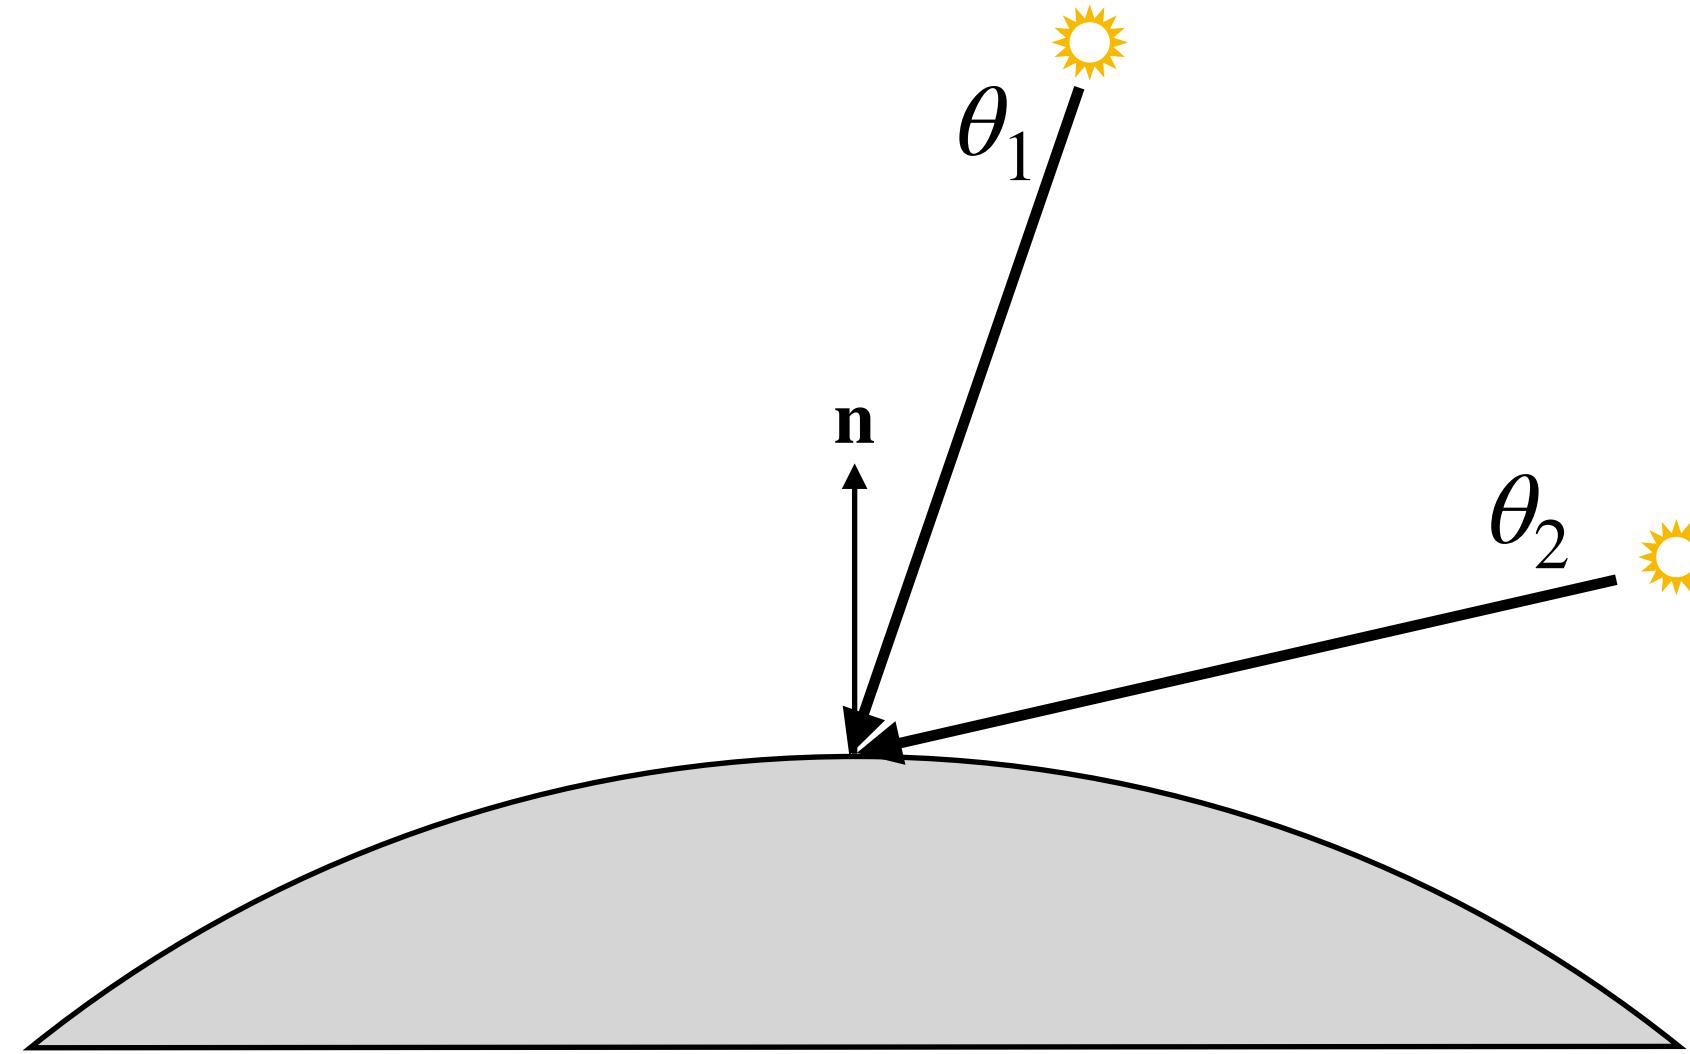


Different materials



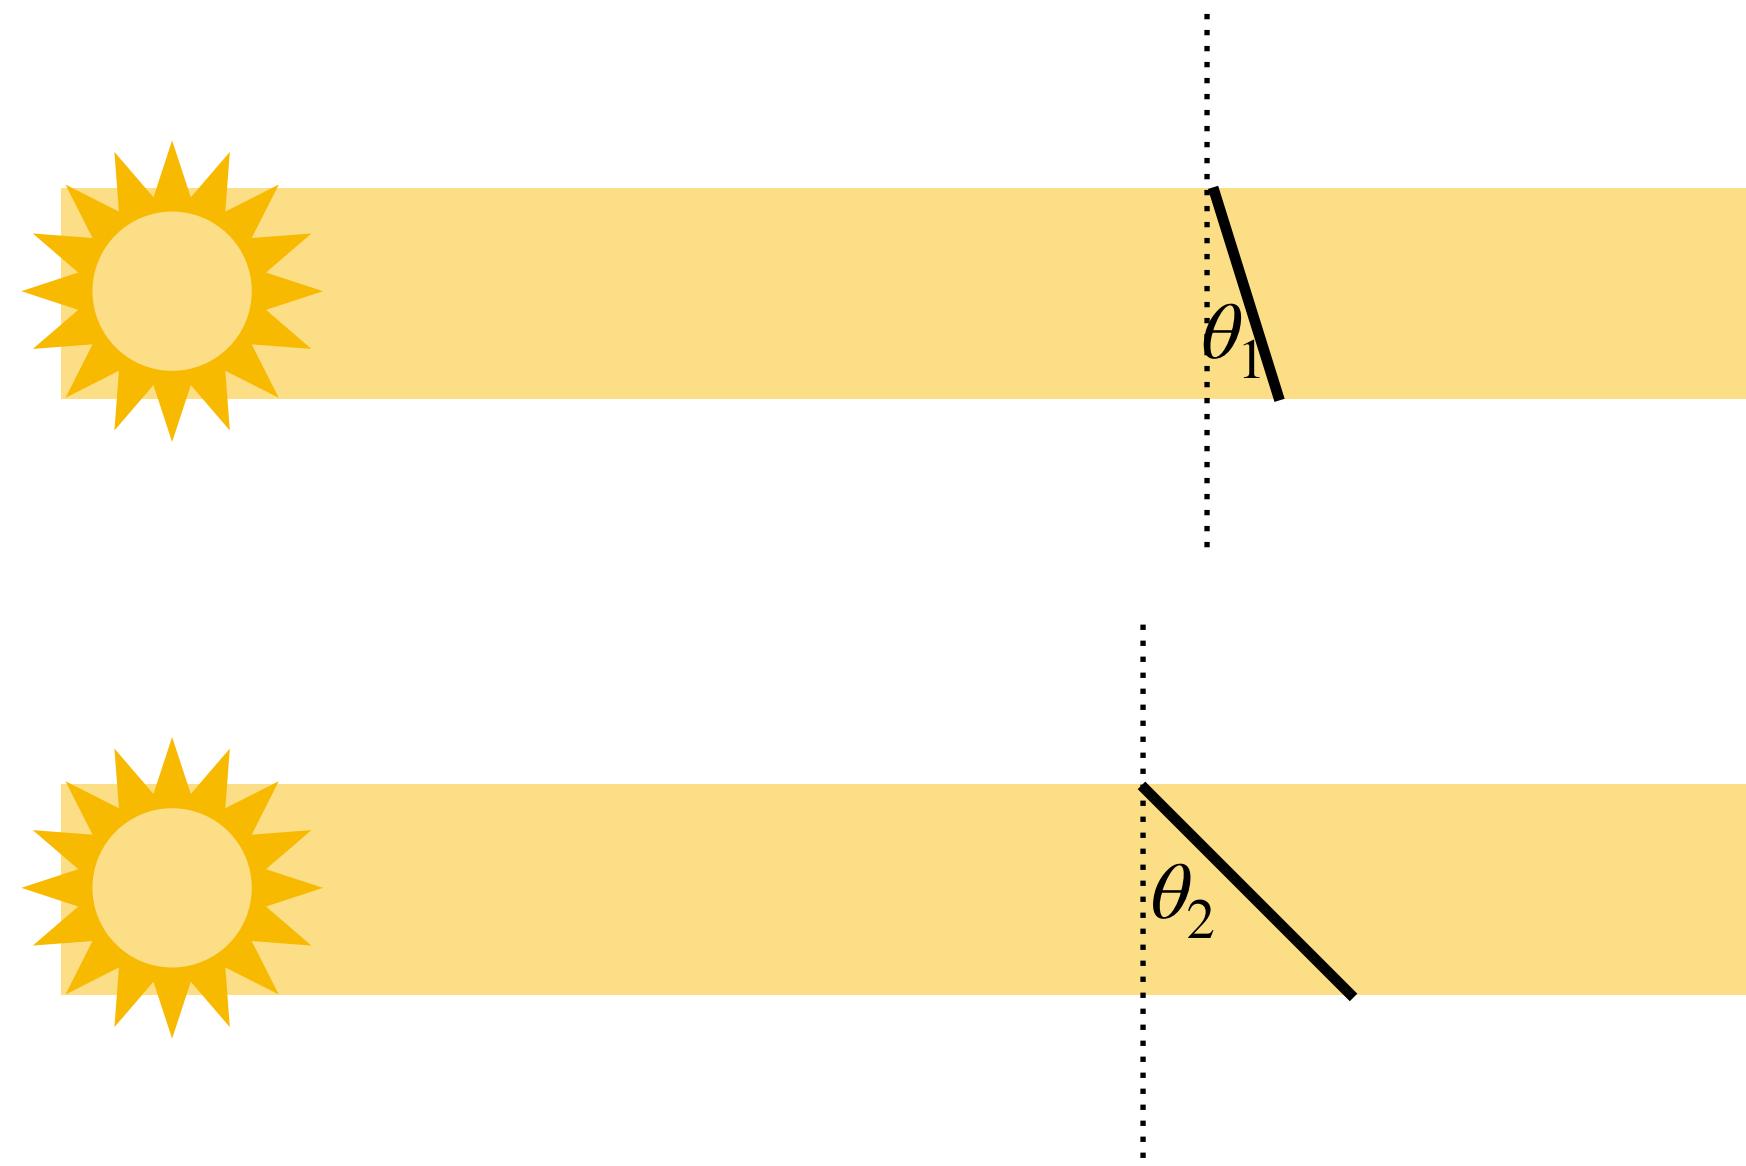
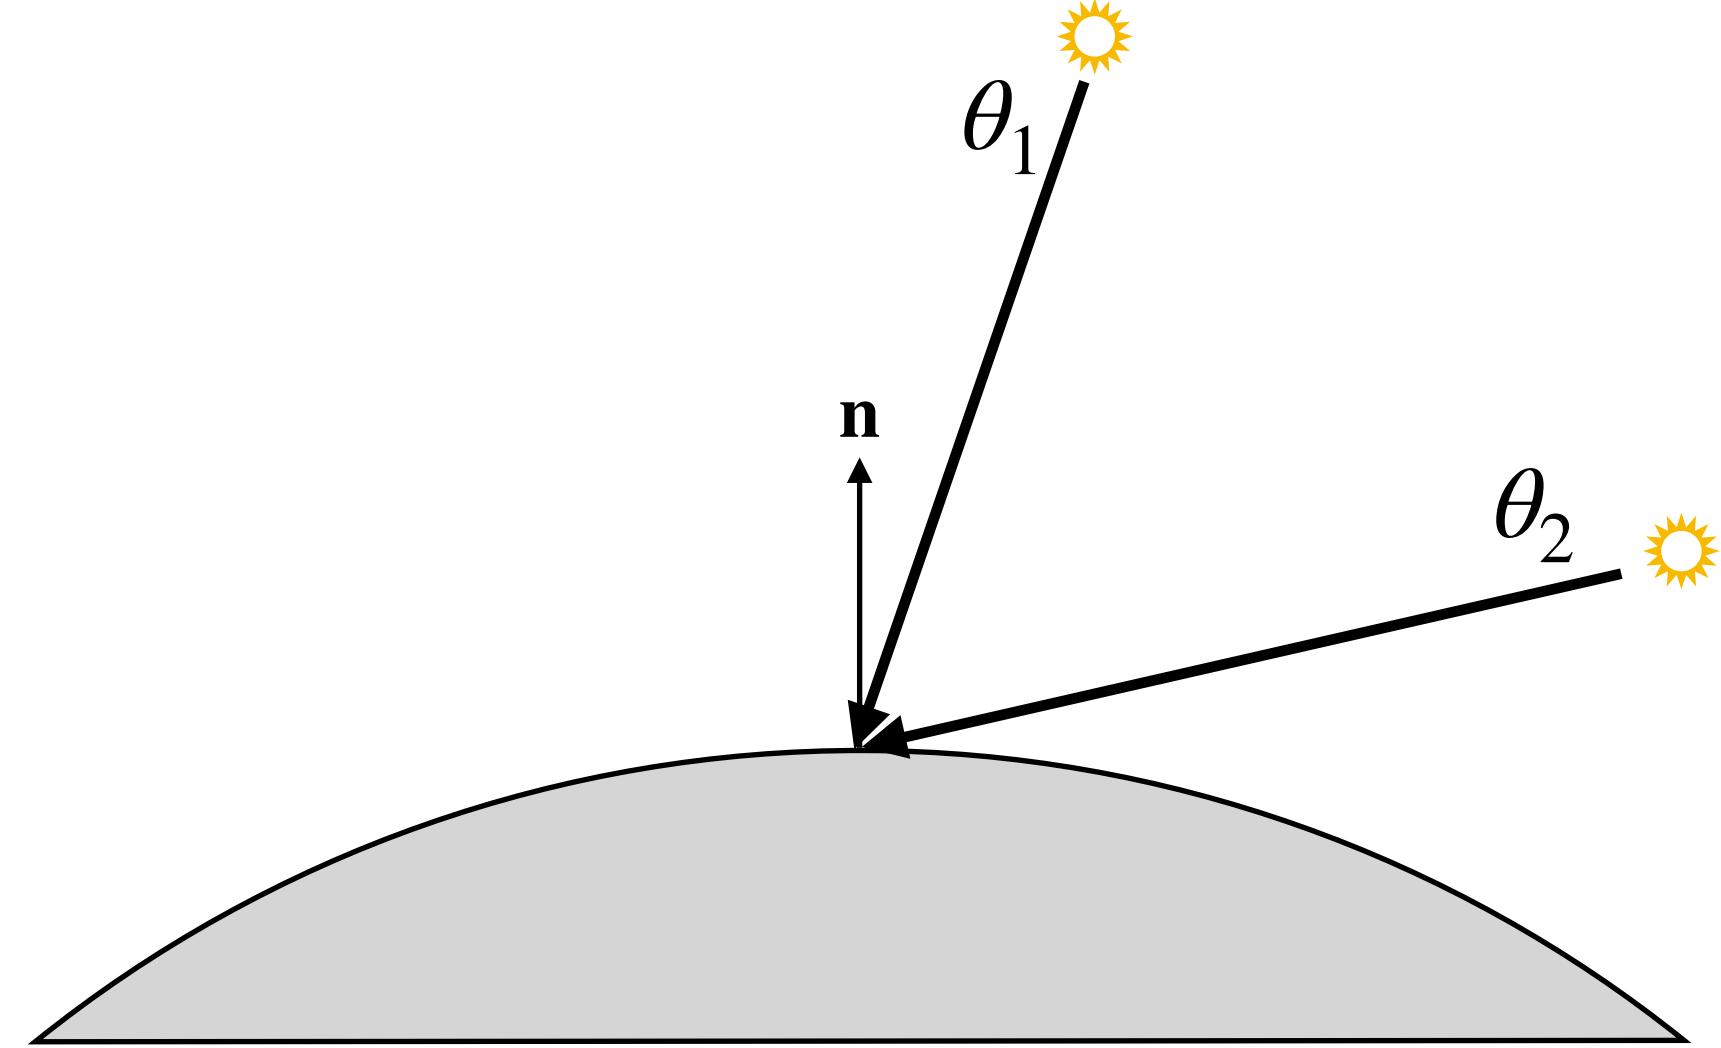
Angle Weight of Photon density

Which incoming light results in more photon density at the surface?



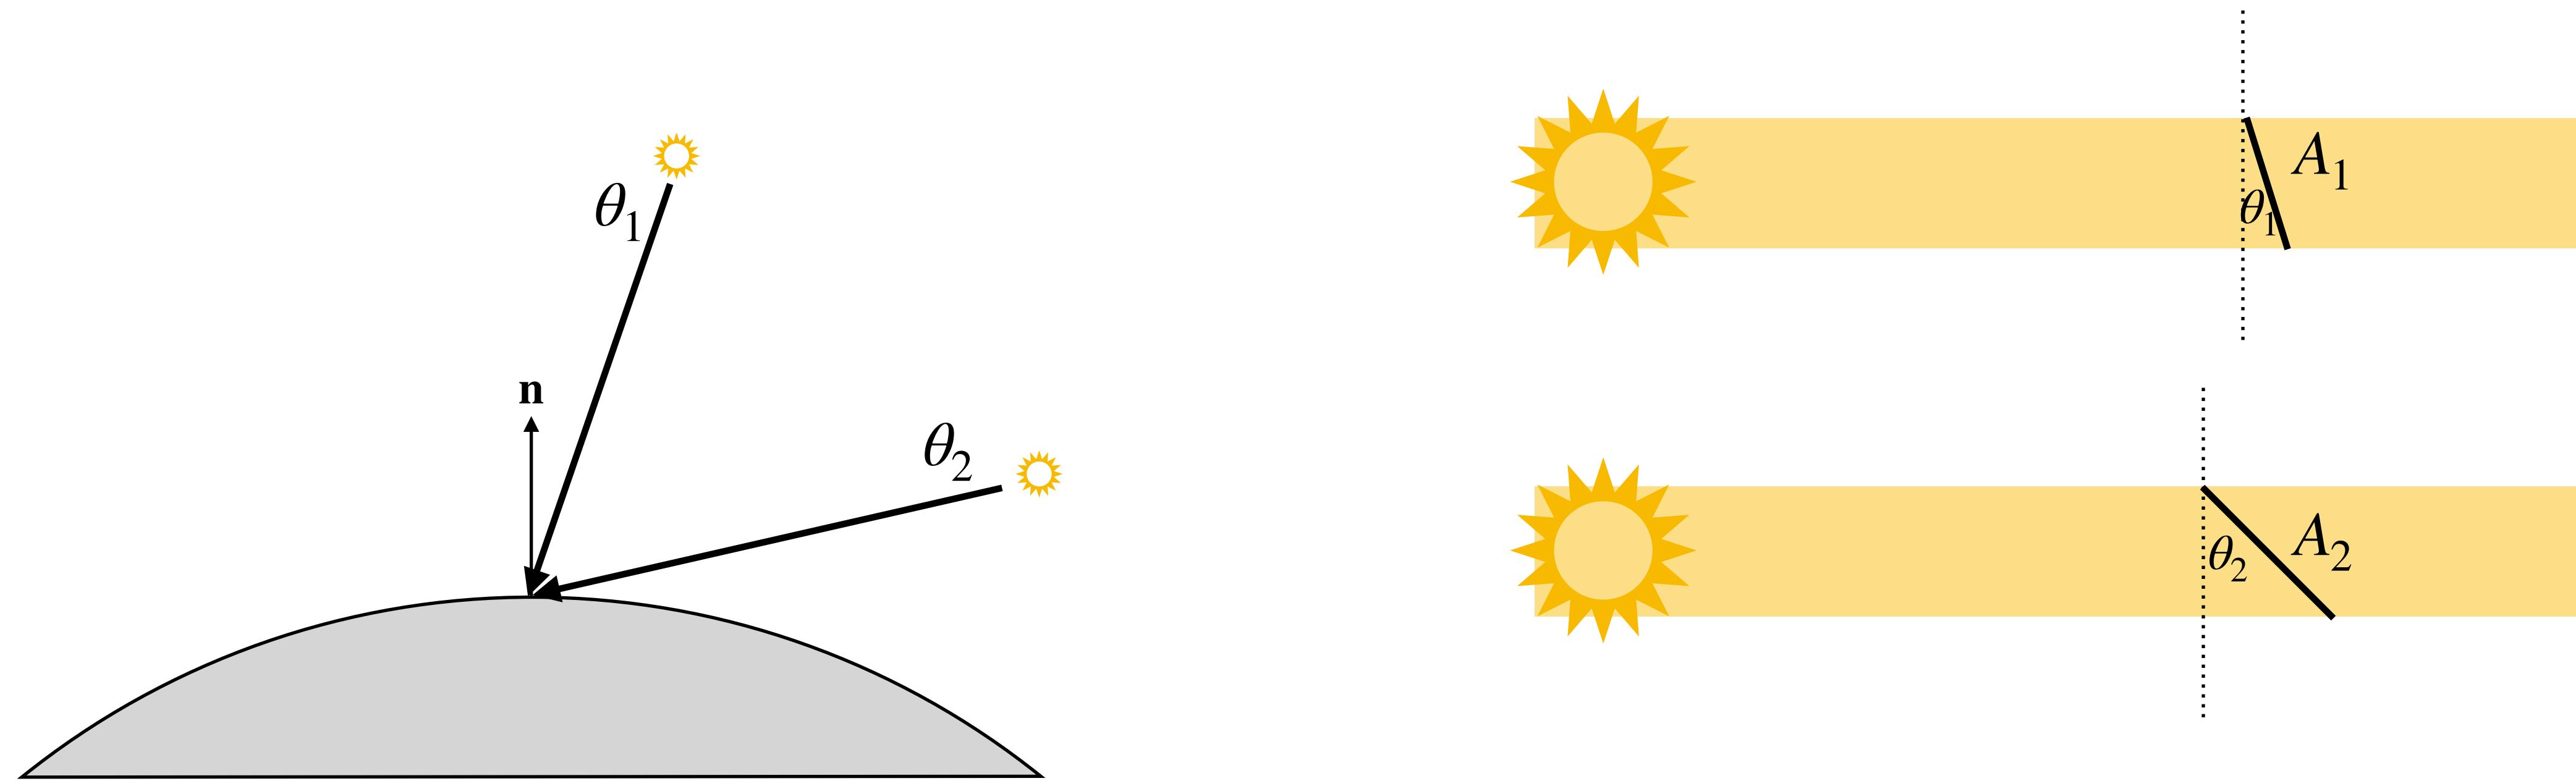
Angle Weight of Photon density

Which incoming light results in more photon density at the surface?



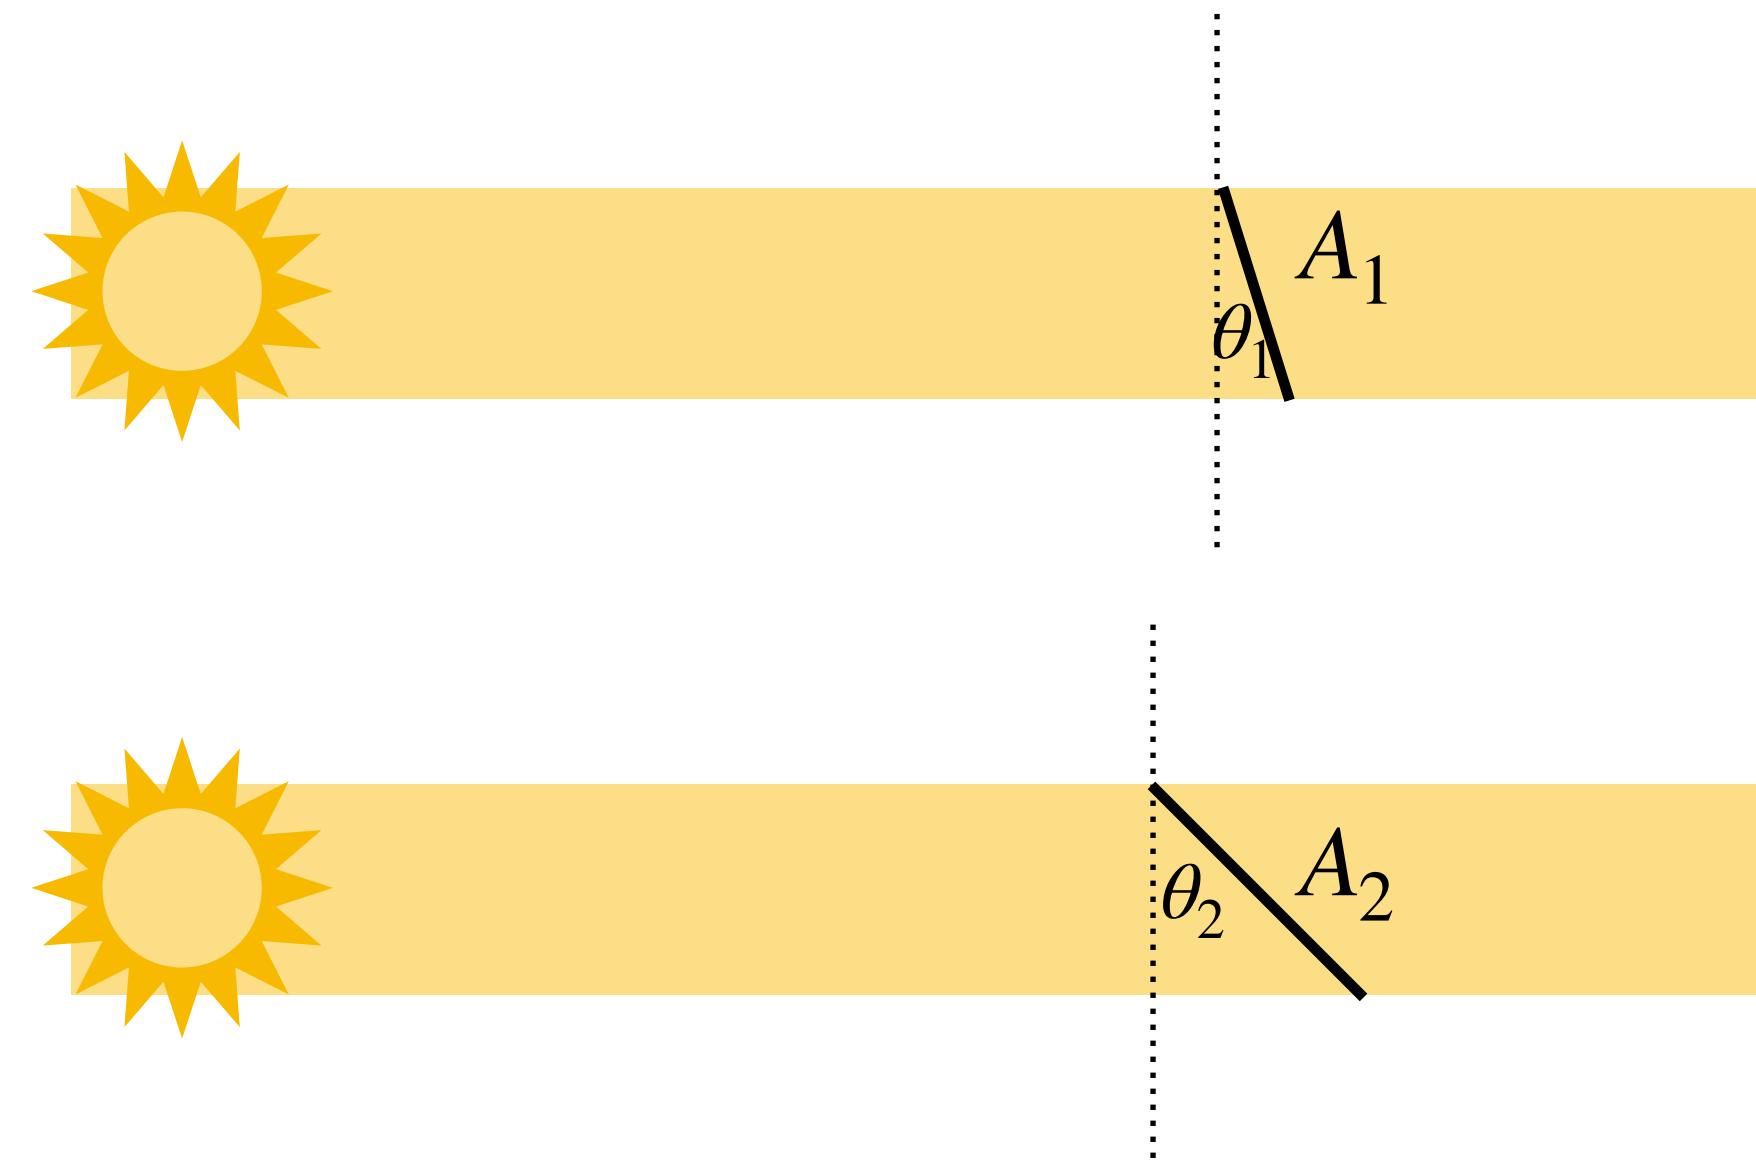
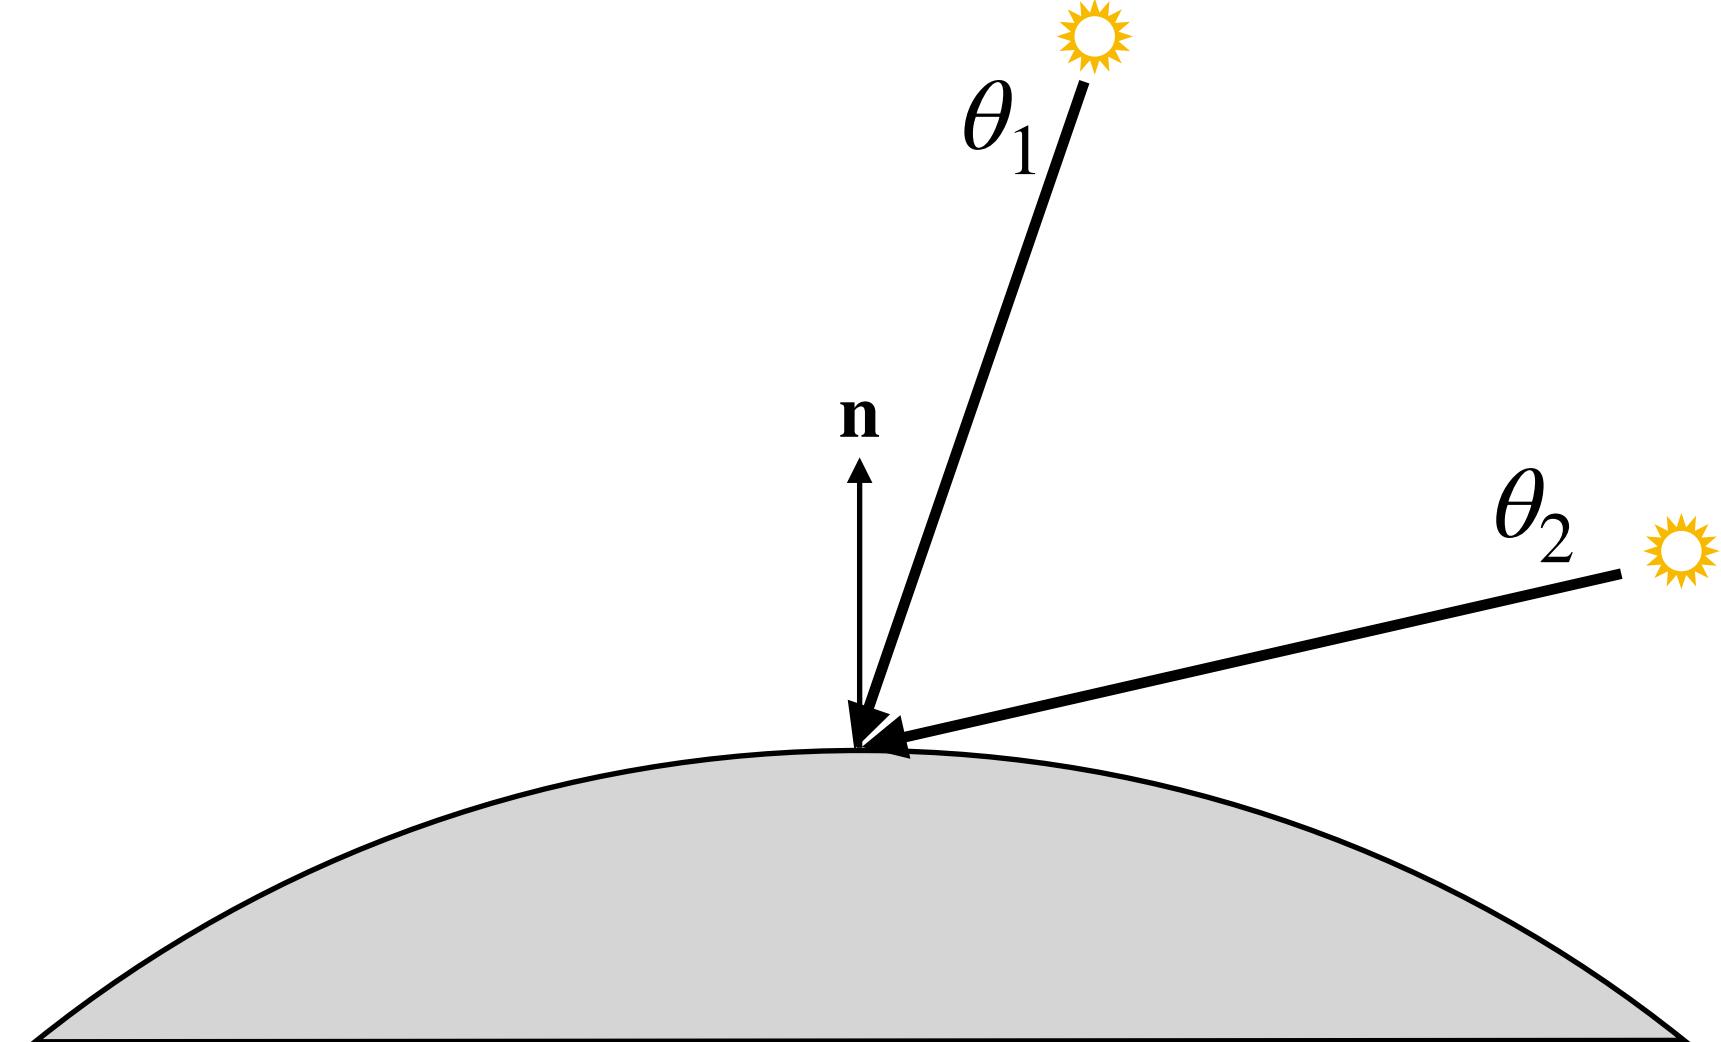
Angle Weight of Photon density

Which incoming light results in more photon density at the surface?



Angle Weight of Photon density

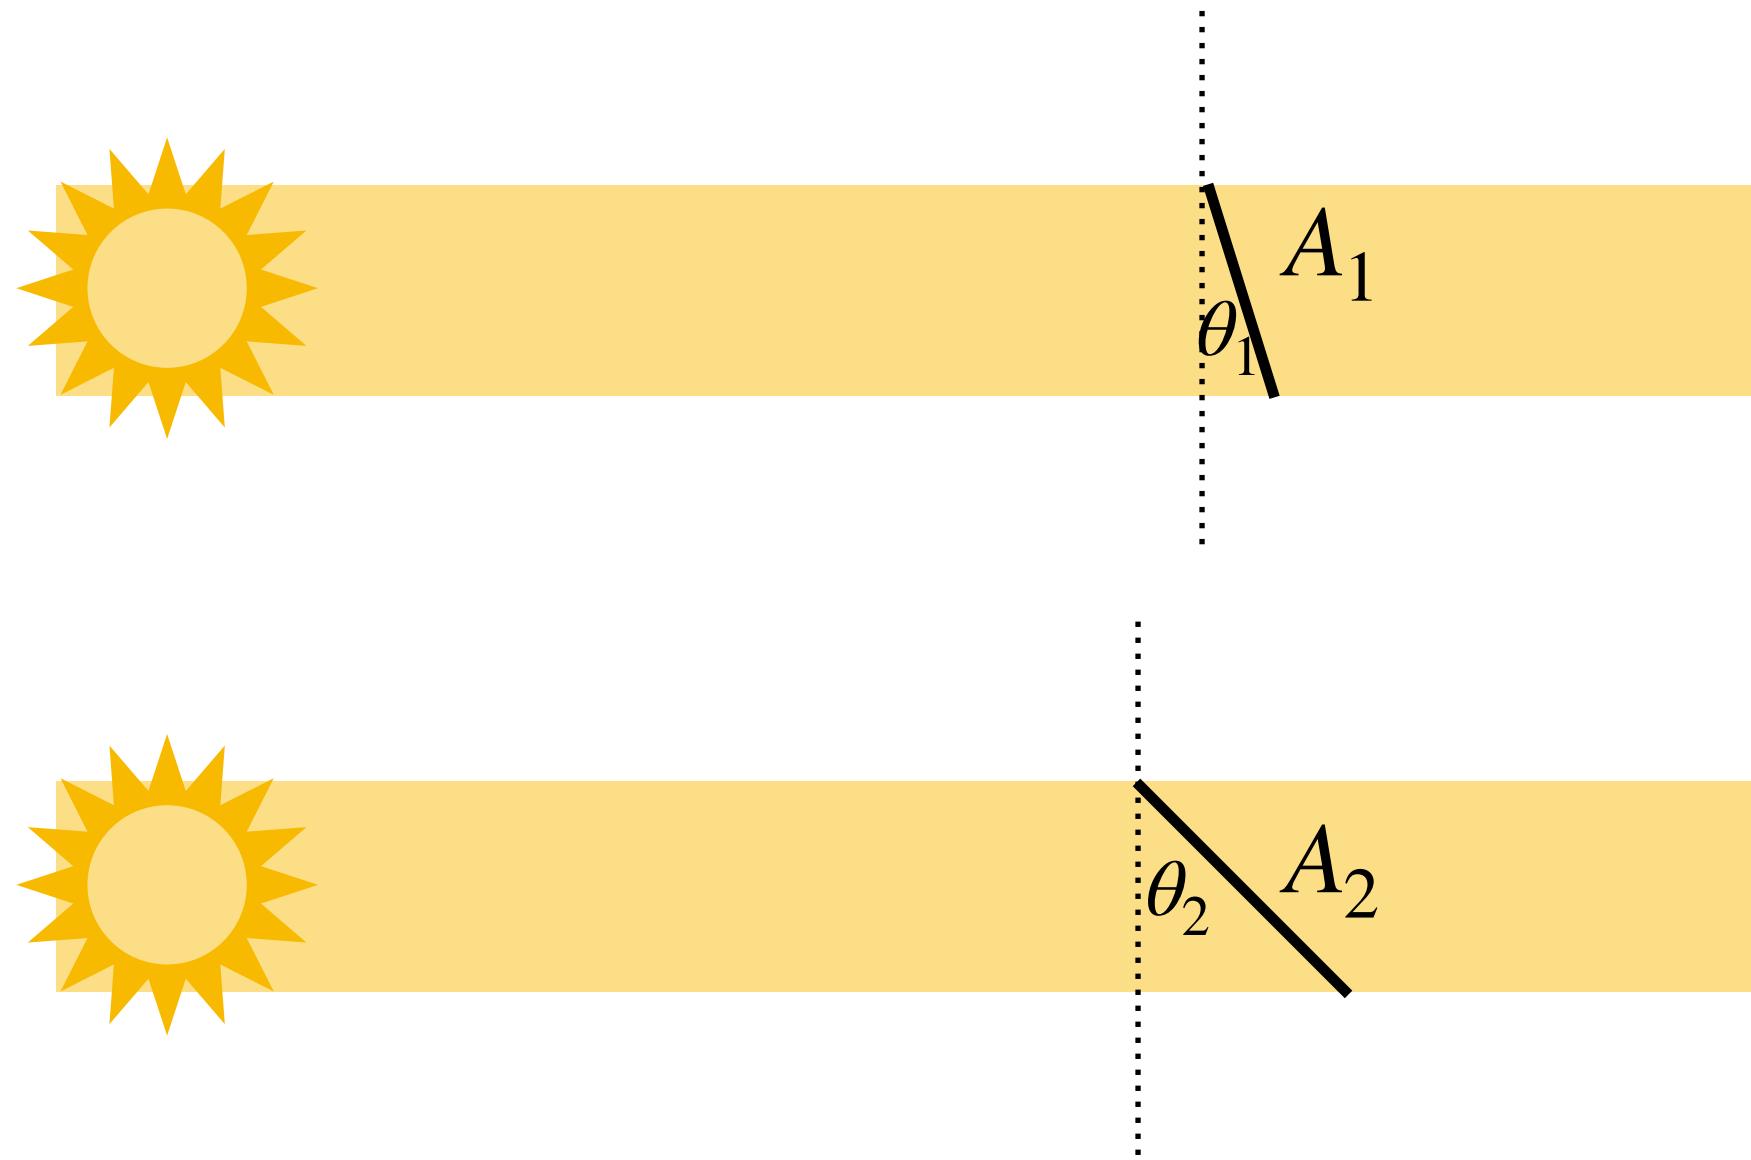
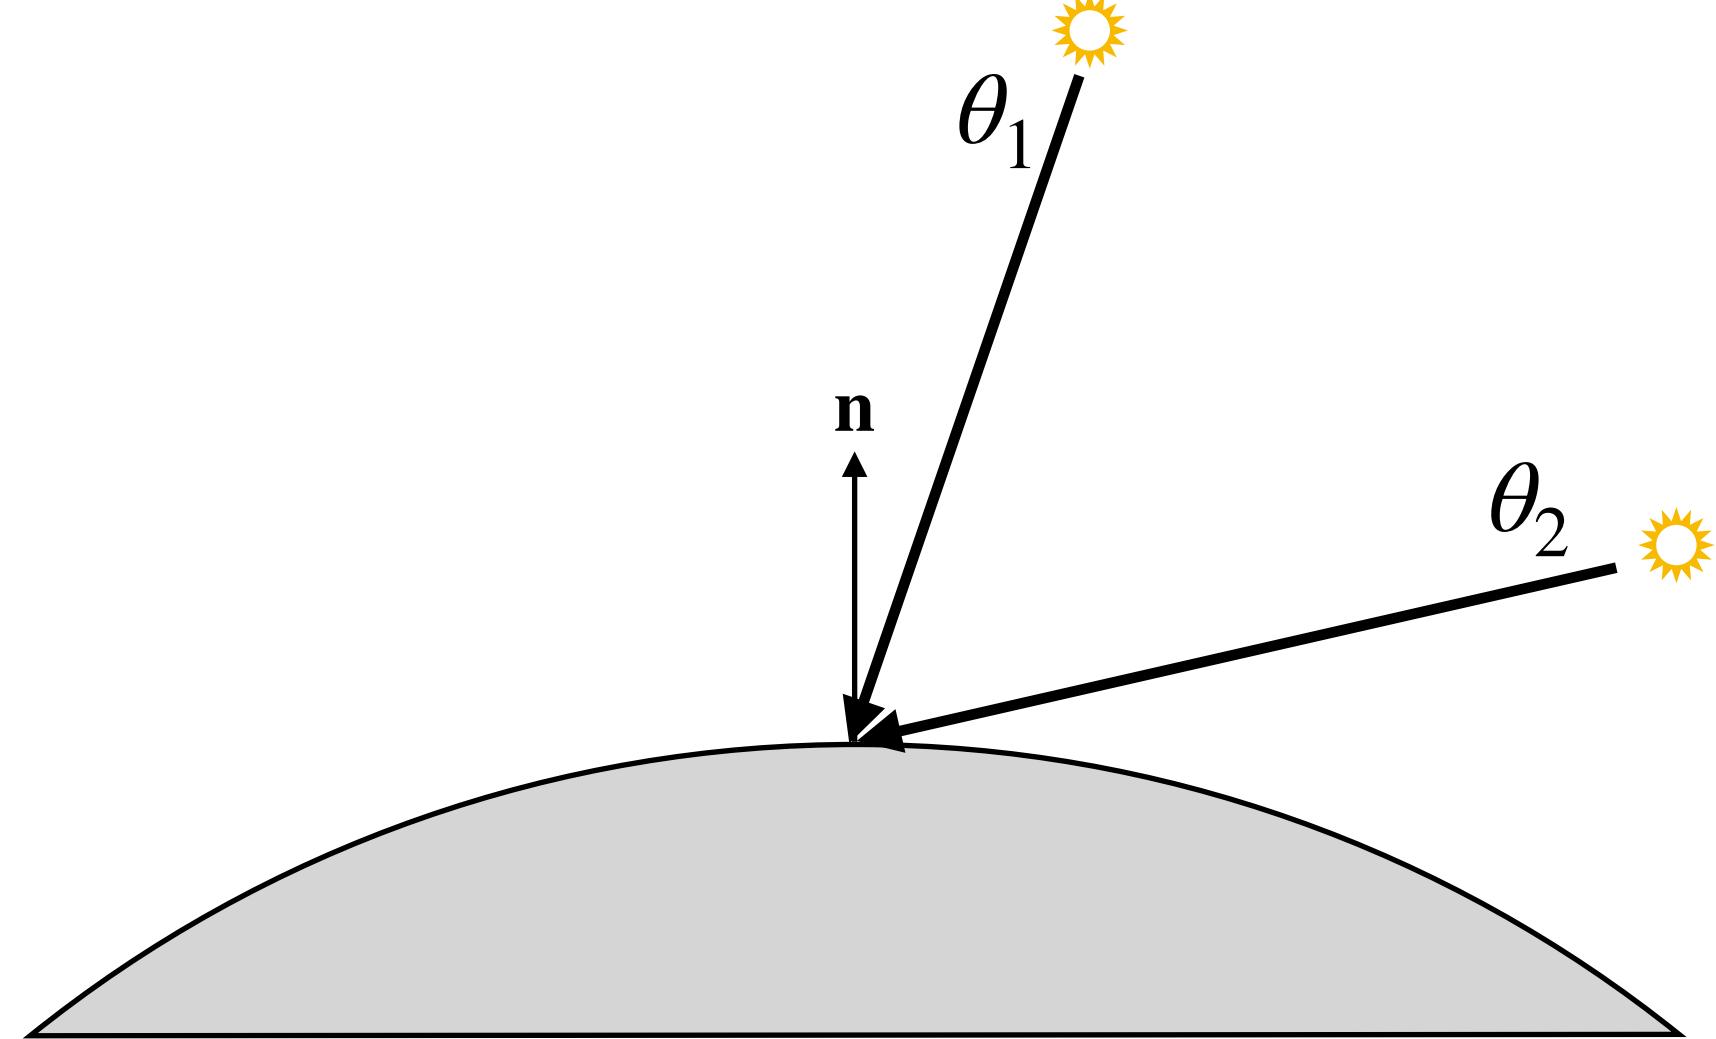
Which incoming light results in more photon density at the surface?



$$\frac{A_1}{A_2} = \frac{\cos \theta_2}{\cos \theta_1}$$

Angle Weight of Photon density

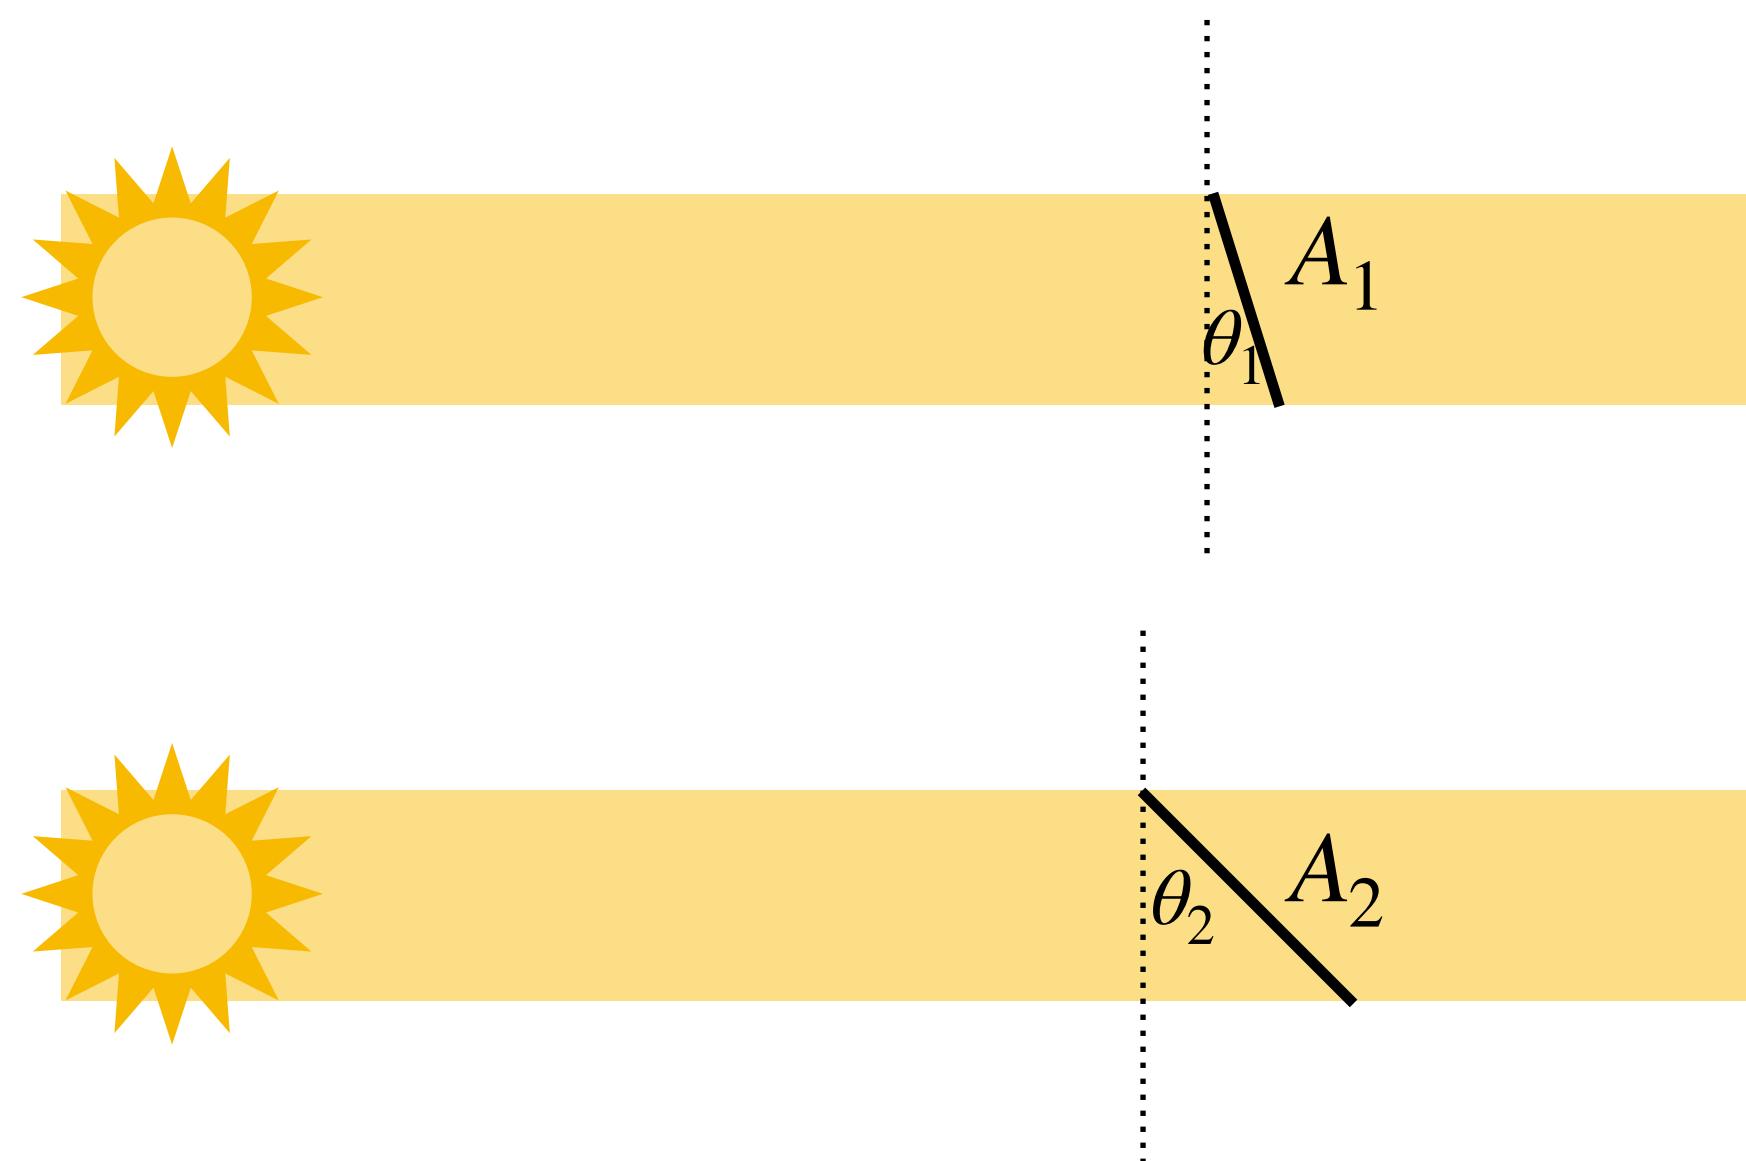
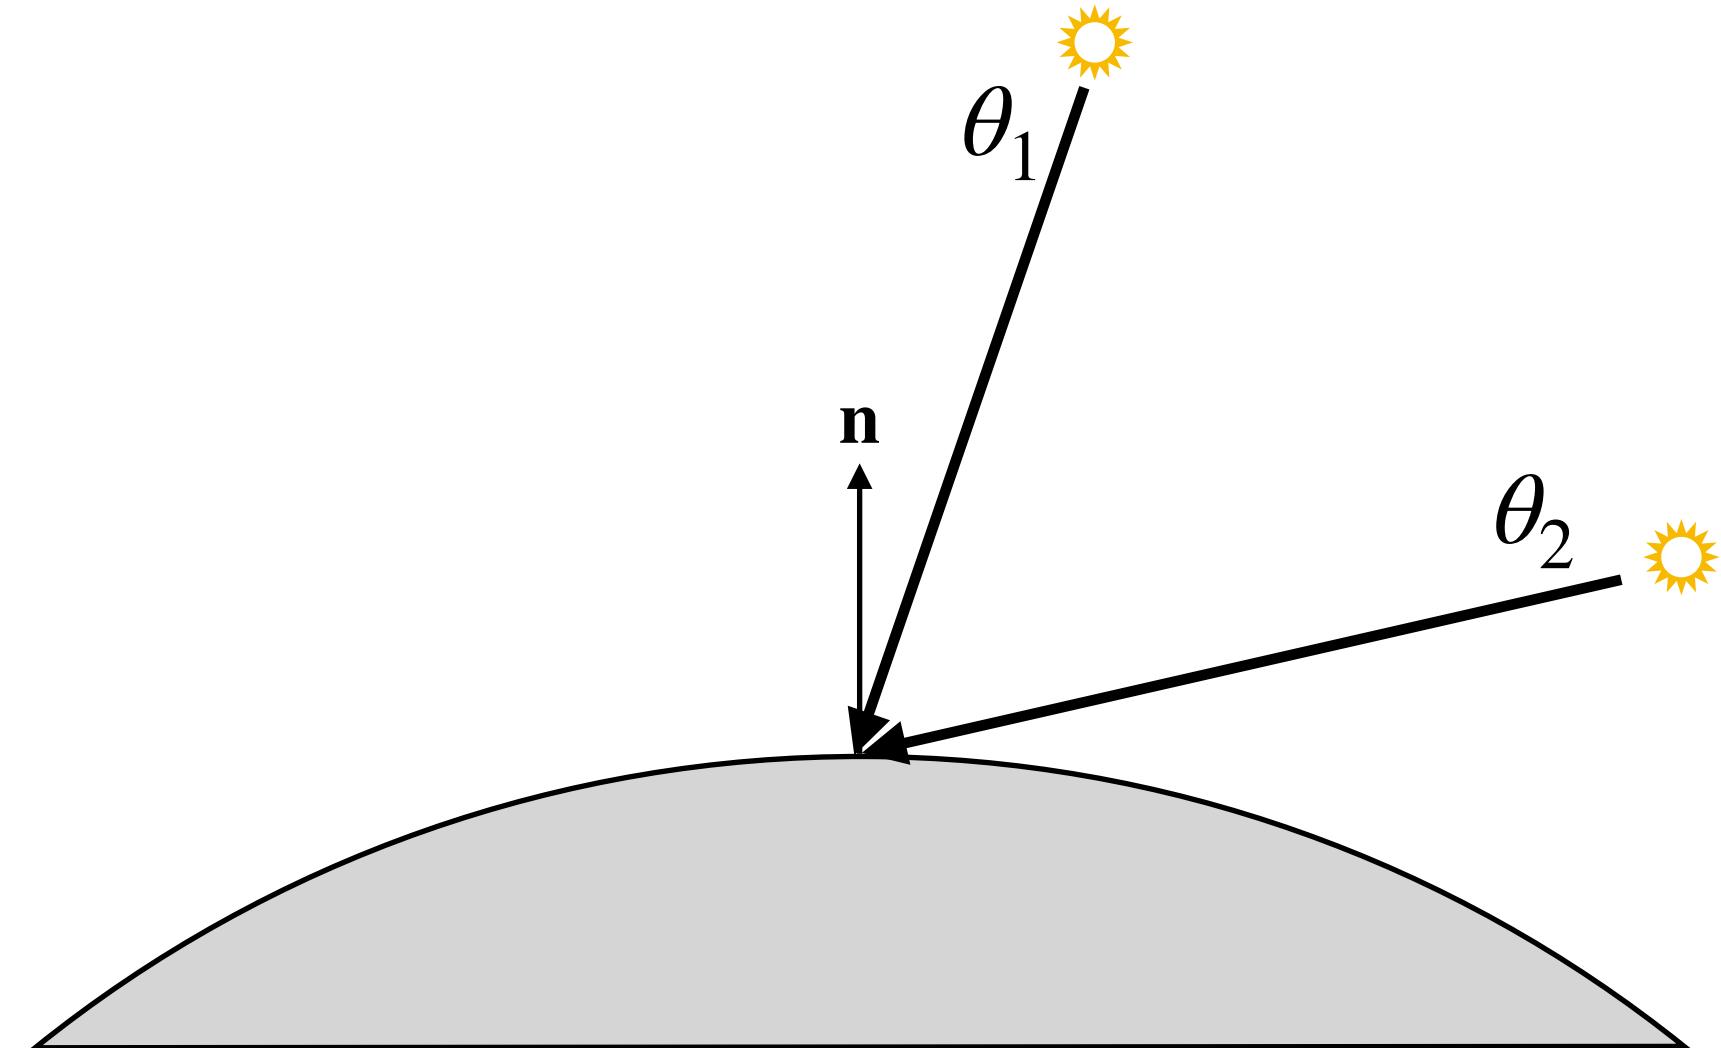
Which incoming light results in more photon density at the surface?



$$\frac{A_1}{A_2} = \frac{\cos \theta_2}{\cos \theta_1} \implies \text{photon density is proportional to } \cos \theta_i$$

Angle Weight of Photon density

Which incoming light results in more photon density at the surface?



$$\frac{A_1}{A_2} = \frac{\cos \theta_2}{\cos \theta_1} \implies \text{photon density is proportional to } \cos \theta_i$$
$$\implies \text{photon density is proportional to } \omega_i \cdot \mathbf{n}$$

Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

↑
outgoing radiance in ω_o

Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

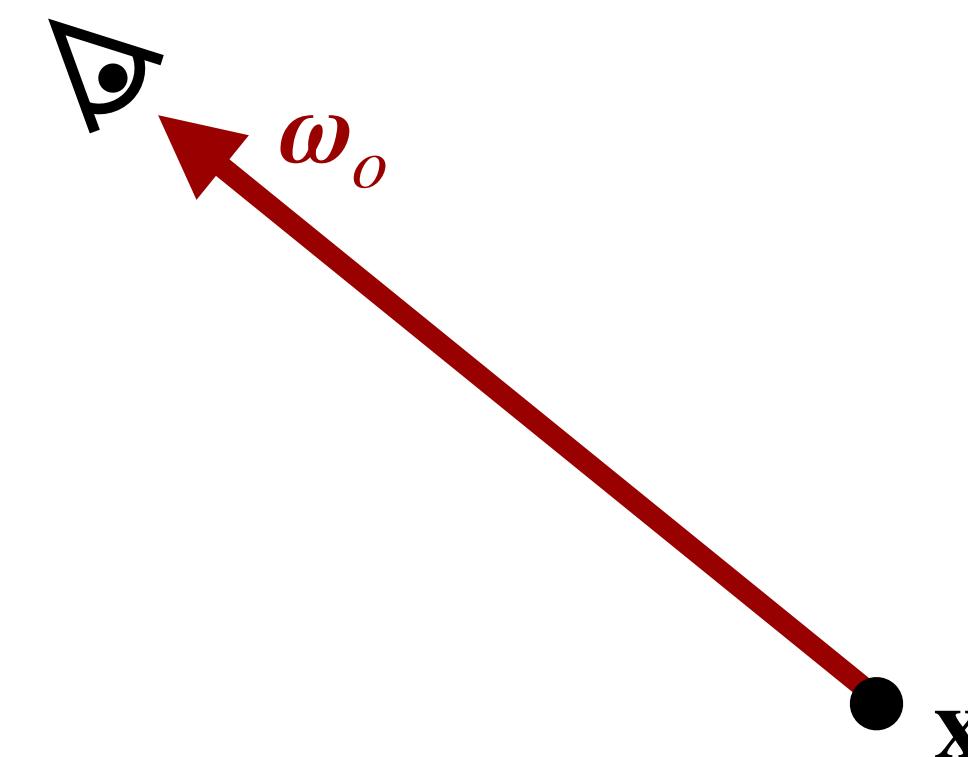
↑
outgoing radiance in ω_o

• \mathbf{x}

Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

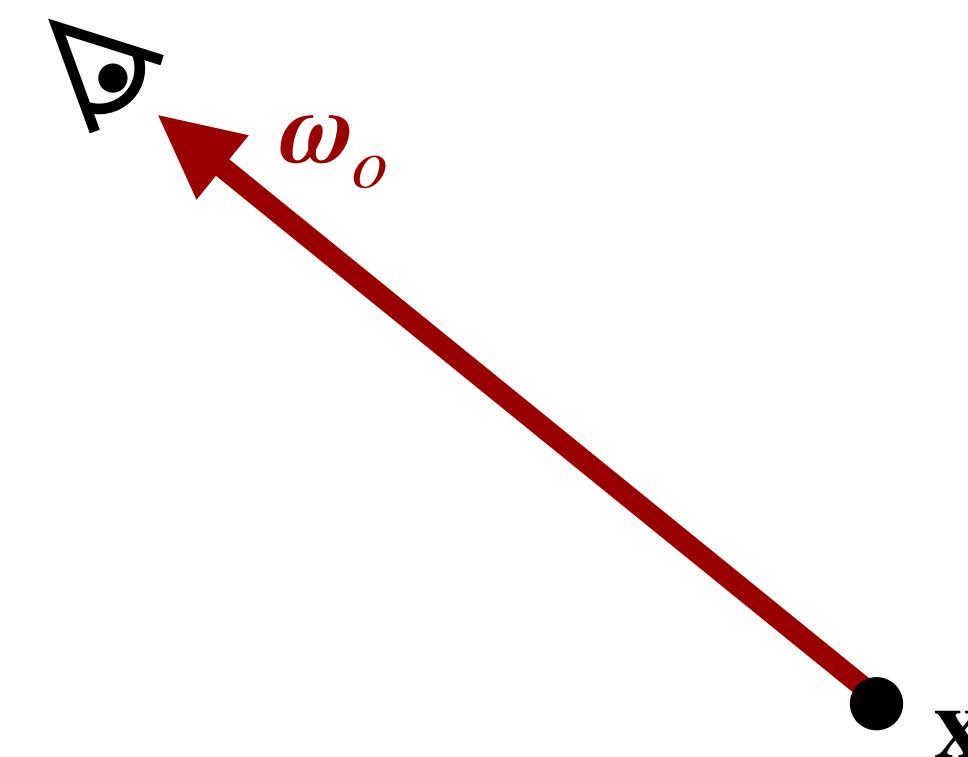
↑
outgoing radiance in ω_o



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance
↓
outgoing radiance in ω_o

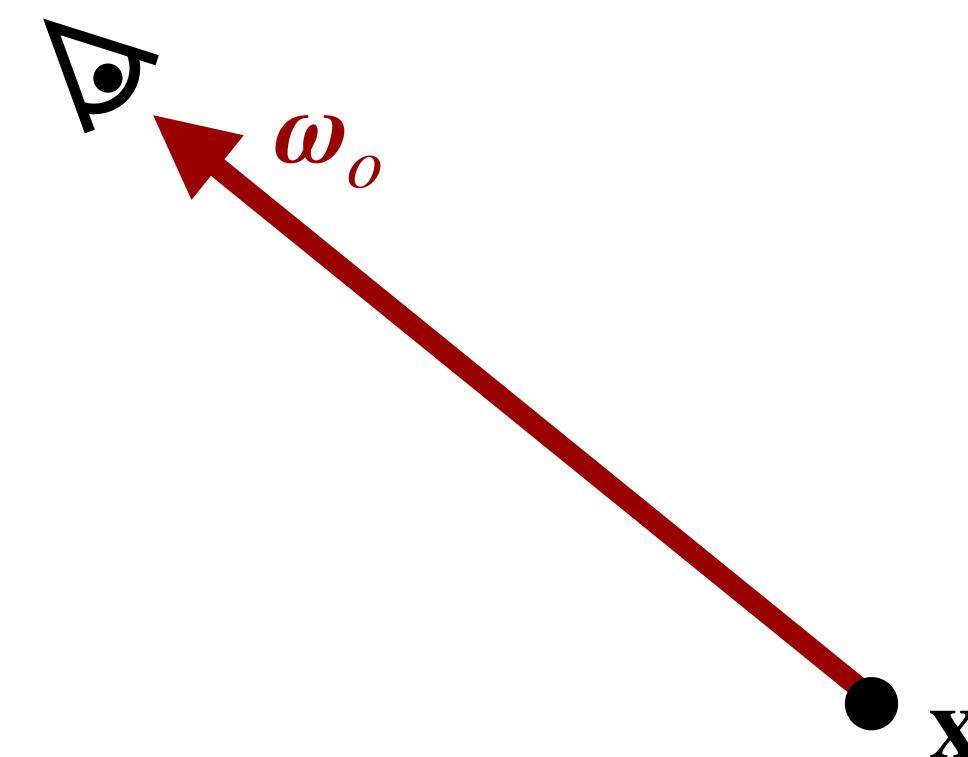


Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance
↓
 $L_o(\mathbf{x}, \omega_o, \lambda)$
↑
outgoing radiance in ω_o

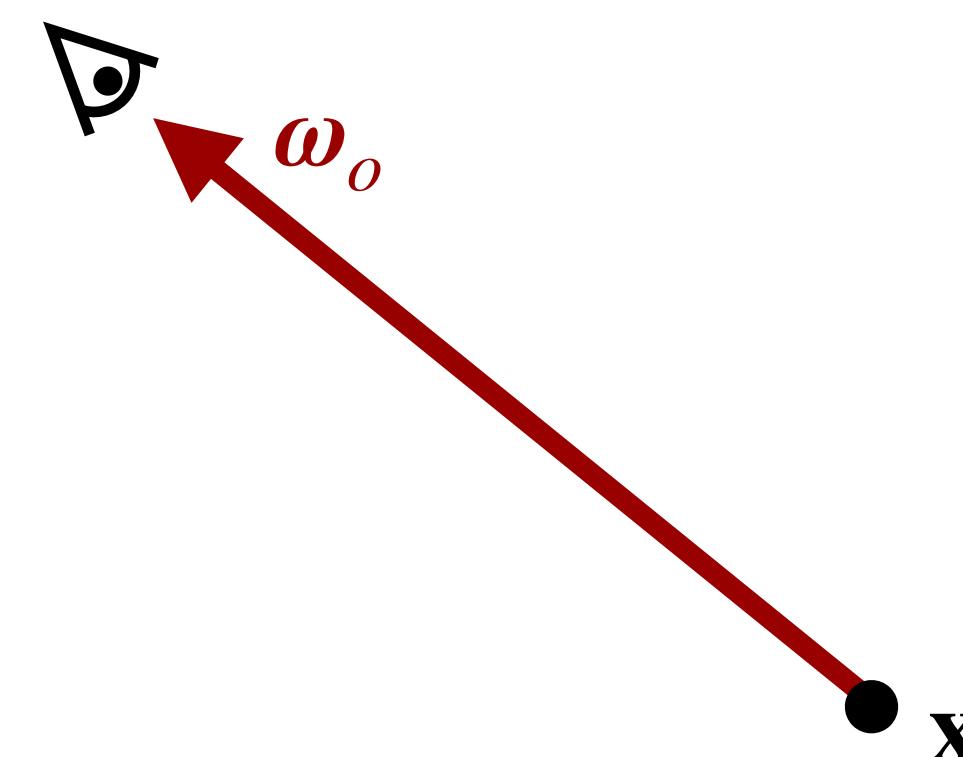
↑
incoming radiance



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

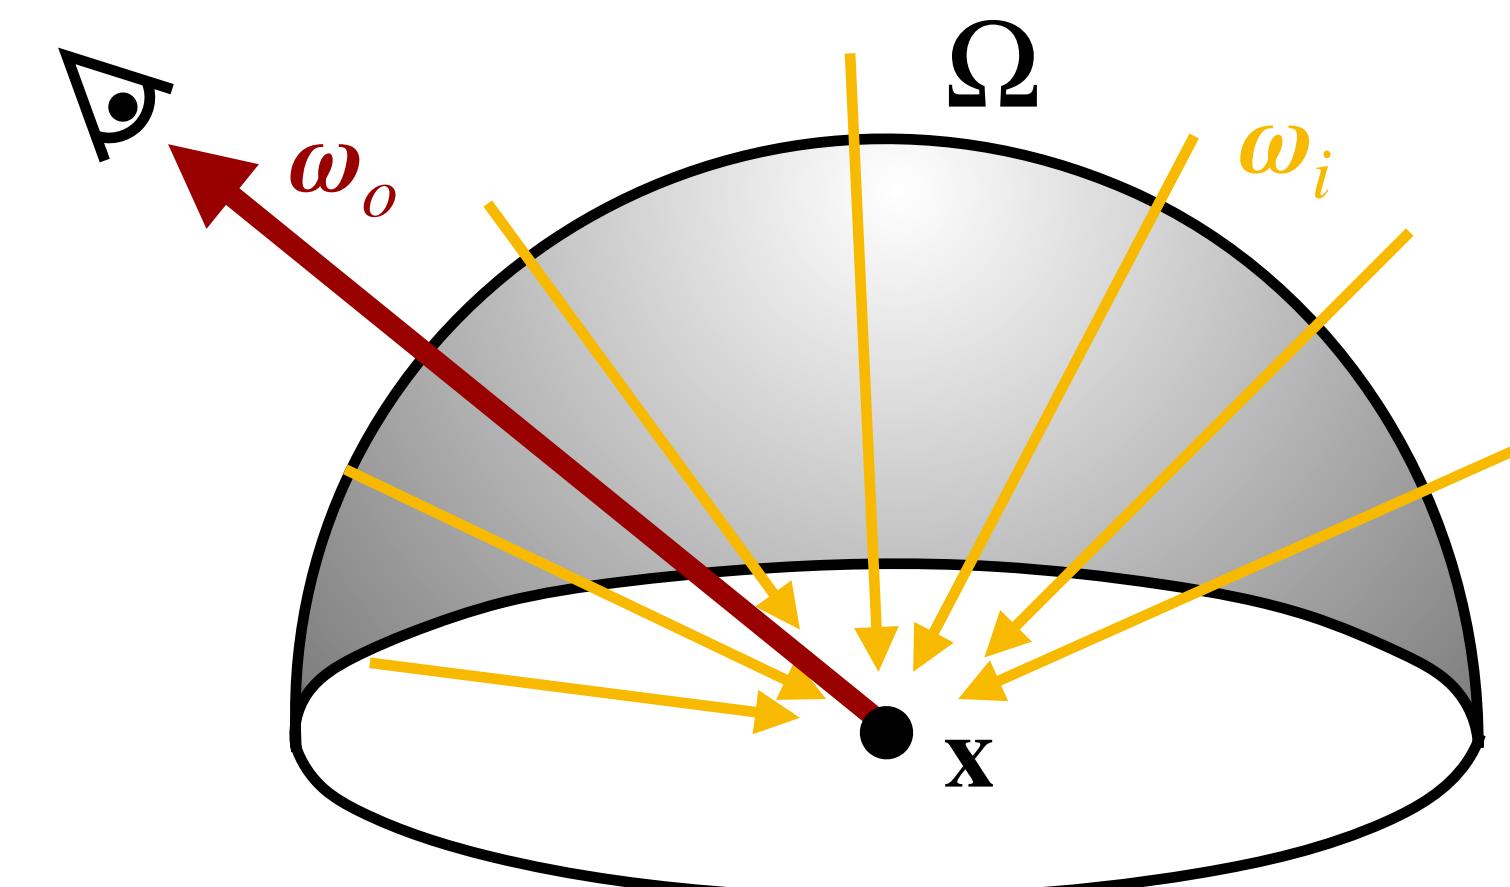
emitted radiance
↓
 $L_o(\mathbf{x}, \omega_o, \lambda)$ = $L_e(\mathbf{x}, \omega_o, \lambda)$ + $\int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$
↑
outgoing radiance in ω_o
↑
incoming radiance



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance
↓
 $L_o(\mathbf{x}, \omega_o, \lambda)$ = $L_e(\mathbf{x}, \omega_o, \lambda)$ + $\int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$
↑
outgoing radiance in ω_o
↑
incoming radiance



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

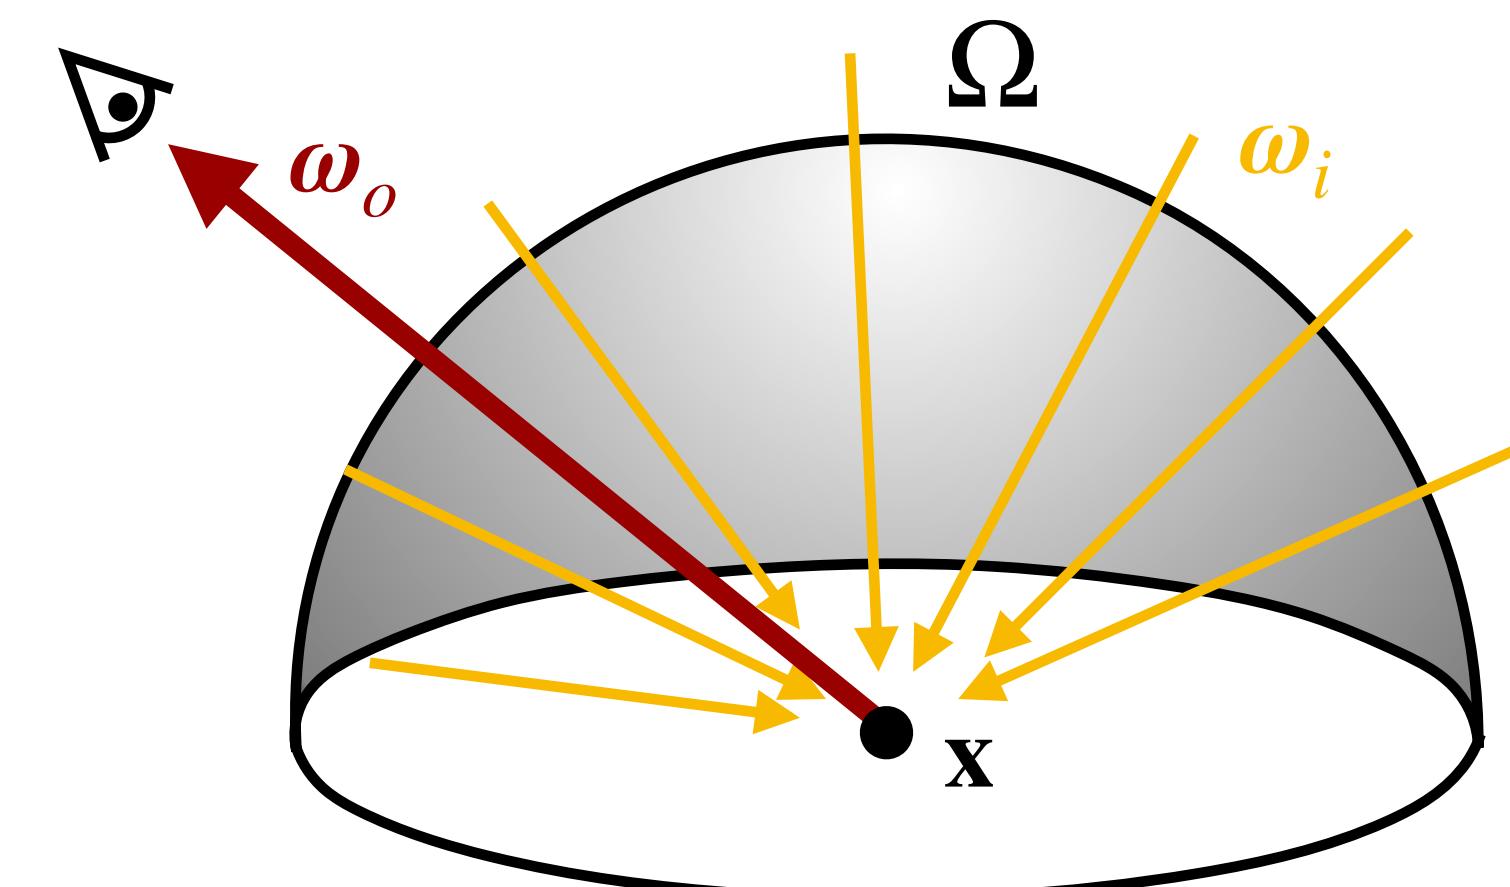
emitted radiance

outgoing radiance in ω_o

solid angle differential

incoming radiance

The diagram illustrates the rendering equation. It shows the total outgoing radiance $L_o(\mathbf{x}, \omega_o, \lambda)$ as the sum of emitted radiance $L_e(\mathbf{x}, \omega_o, \lambda)$ and reflected radiance. The reflected radiance is calculated by integrating over the hemisphere above point \mathbf{x} . The integrand consists of the shading function $f(\mathbf{x}, \omega_i, \omega_o, \lambda)$, the incoming radiance $L_i(\mathbf{x}, \omega_i, \lambda)$, and the cosine of the angle between the incoming direction ω_i and the surface normal \mathbf{n} . The solid angle differential $d\omega_i$ is also shown.



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance

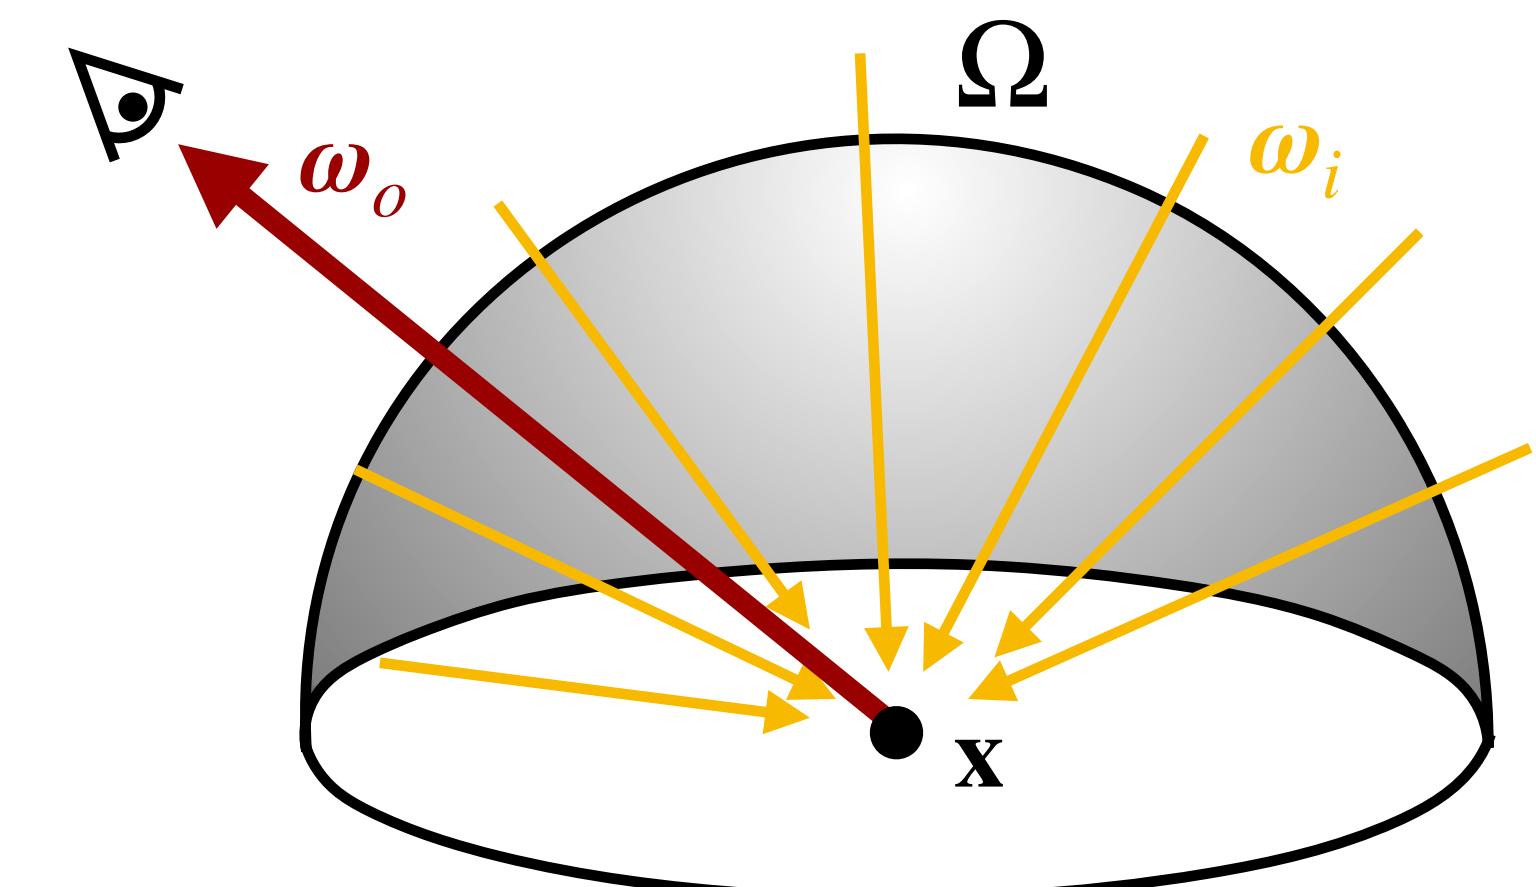
outgoing radiance in ω_o

fraction of incoming light that is reflected into ω_o

incoming radiance

solid angle differential

The diagram illustrates the rendering equation. It shows the total outgoing radiance $L_o(\mathbf{x}, \omega_o, \lambda)$ as the sum of emitted radiance $L_e(\mathbf{x}, \omega_o, \lambda)$ and the integral of reflected radiance. The reflected radiance is calculated by integrating over the solid angle Ω the product of the shading function $f(\mathbf{x}, \omega_i, \omega_o, \lambda)$, the incoming radiance $L_i(\mathbf{x}, \omega_i, \lambda)$, and the cosine of the angle between the incoming direction ω_i and the surface normal \mathbf{n} . Arrows point from each term in the equation to its corresponding label: an arrow points down to the emitted radiance term, an arrow points up to the outgoing radiance term, and three arrows point up to the integral term, each labeled with a different component: the shading function, the incoming radiance, and the solid angle differential.



Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance

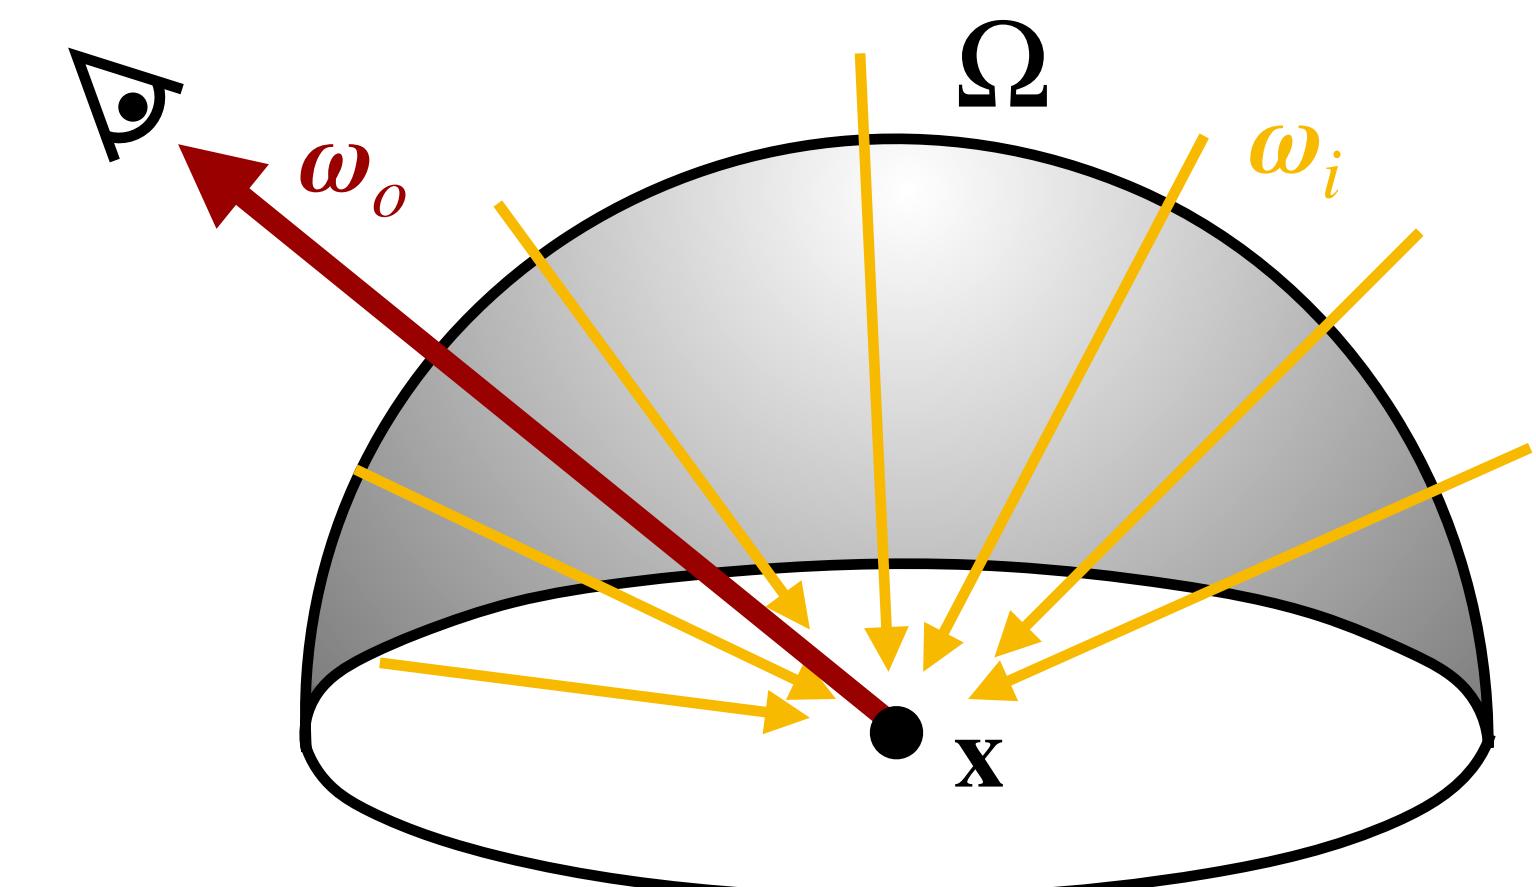
outgoing radiance in ω_o

fraction of incoming light that is reflected into ω_o

incoming radiance

angle weight of photon density

solid angle differential



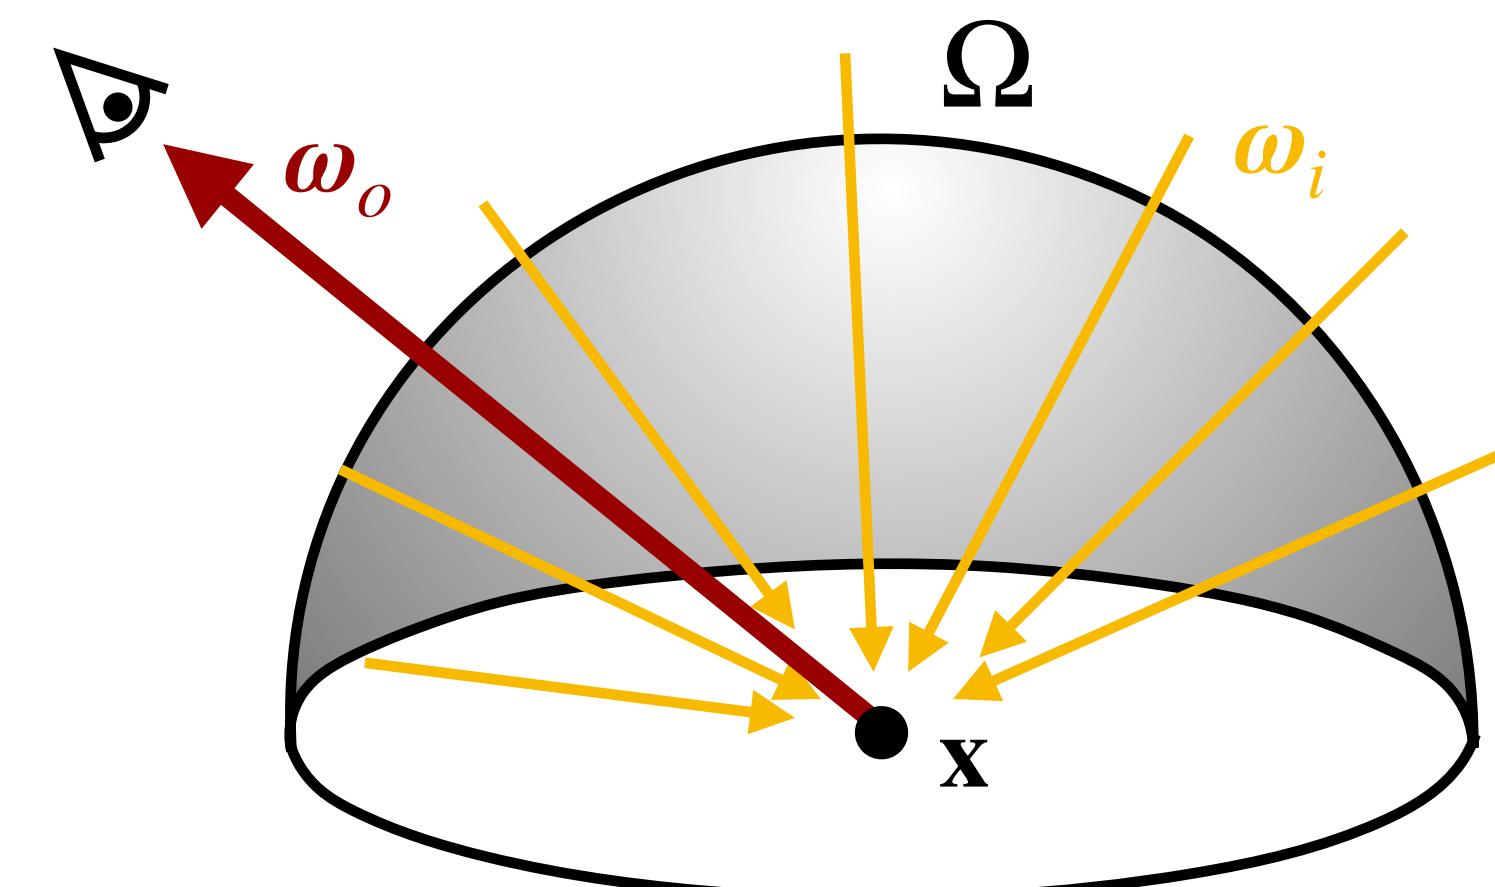
Rendering equation

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

emitted radiance
outgoing radiance in ω_o
fraction of incoming light that is reflected into ω_o
incoming radiance
angle weight of photon density
solid angle differential

Conservation of energy:

$$\int_{\omega_i \in \Omega} f(\omega_i, \omega_o) (\omega_i \cdot \mathbf{n}) d\omega_i \leq 1$$



Colors

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Colors

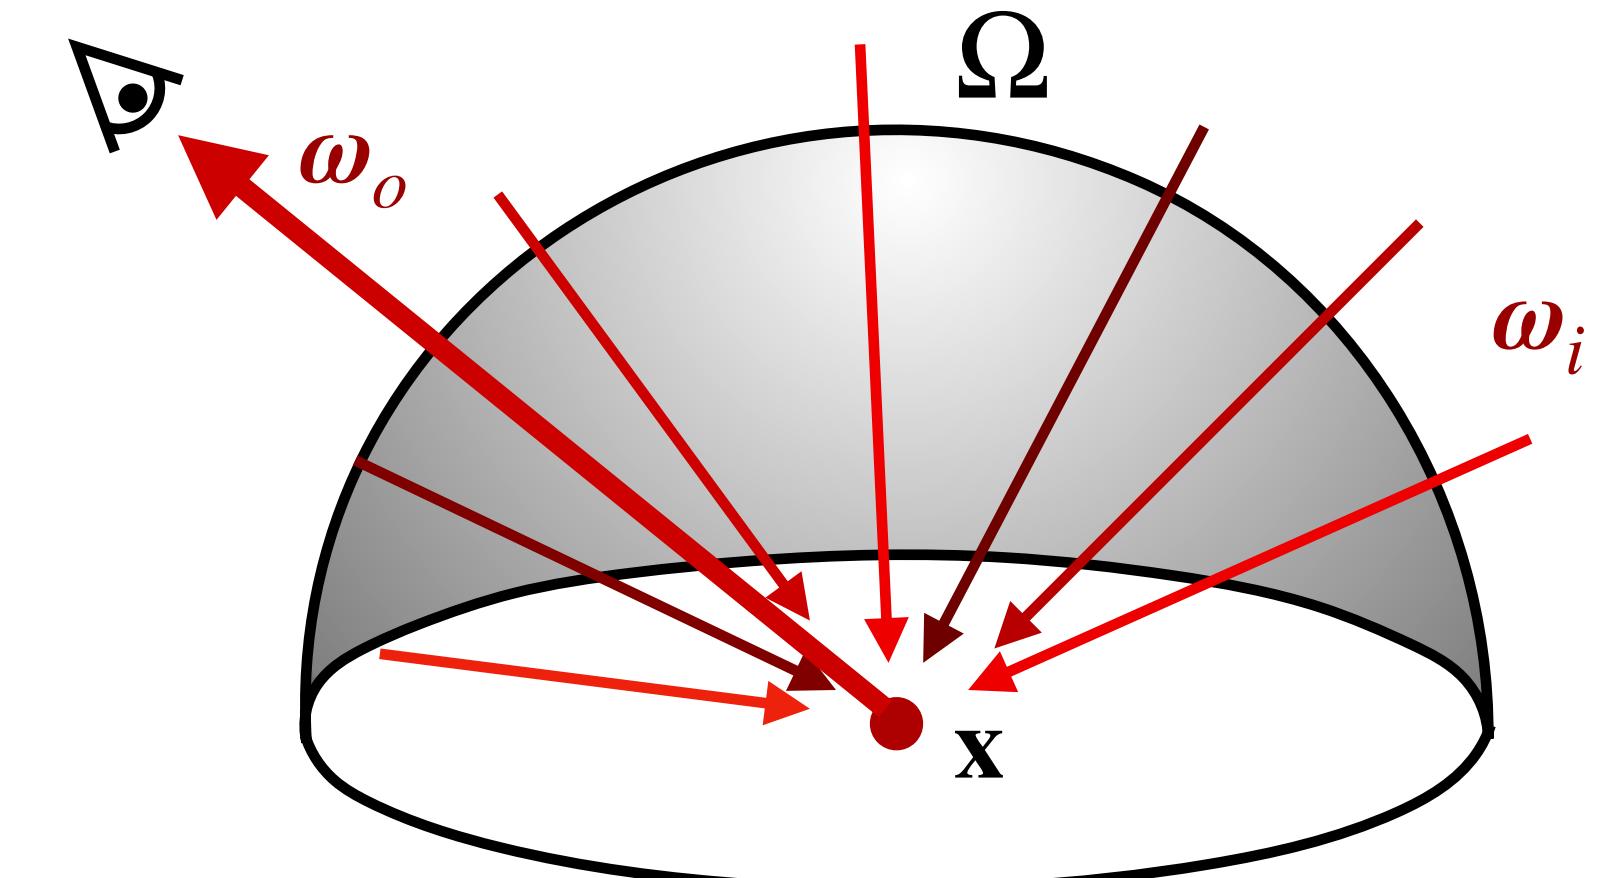
$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^R) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^R) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

Colors

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

$$L_o(\mathbf{x}, \omega_o, \lambda^R) = L_e(\mathbf{x}, \omega_o, \lambda^R) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda^R) L_i(\mathbf{x}, \omega_i, \lambda^R) (\omega_i \cdot \mathbf{n}) d\omega_i$$



Colors

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

Colors

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Colors

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^R) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^R) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

Colors

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^R) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^R) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^G) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^G) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^G) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^G) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

Colors

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^R) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^R) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^R) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^G) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^G) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^G) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^G) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^B) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda^B) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda^B) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda^B) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

Colors

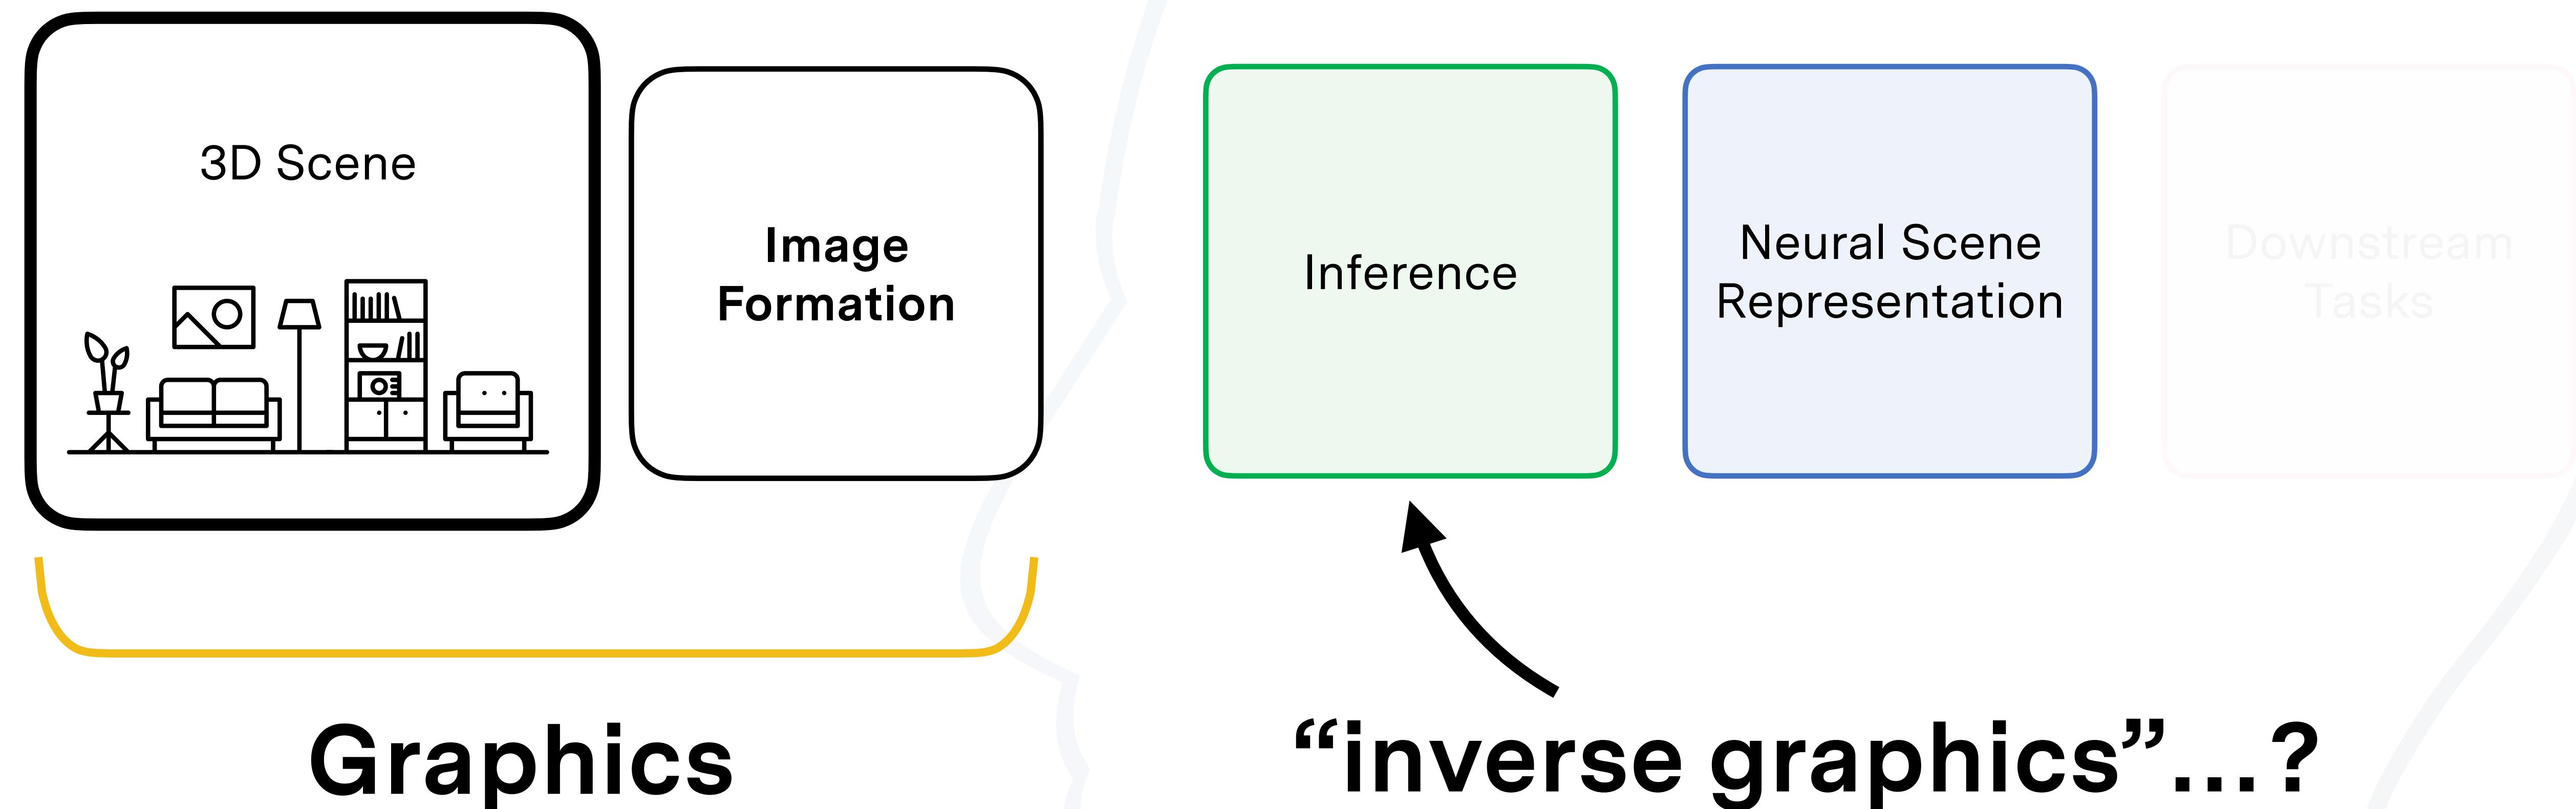
$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Colors

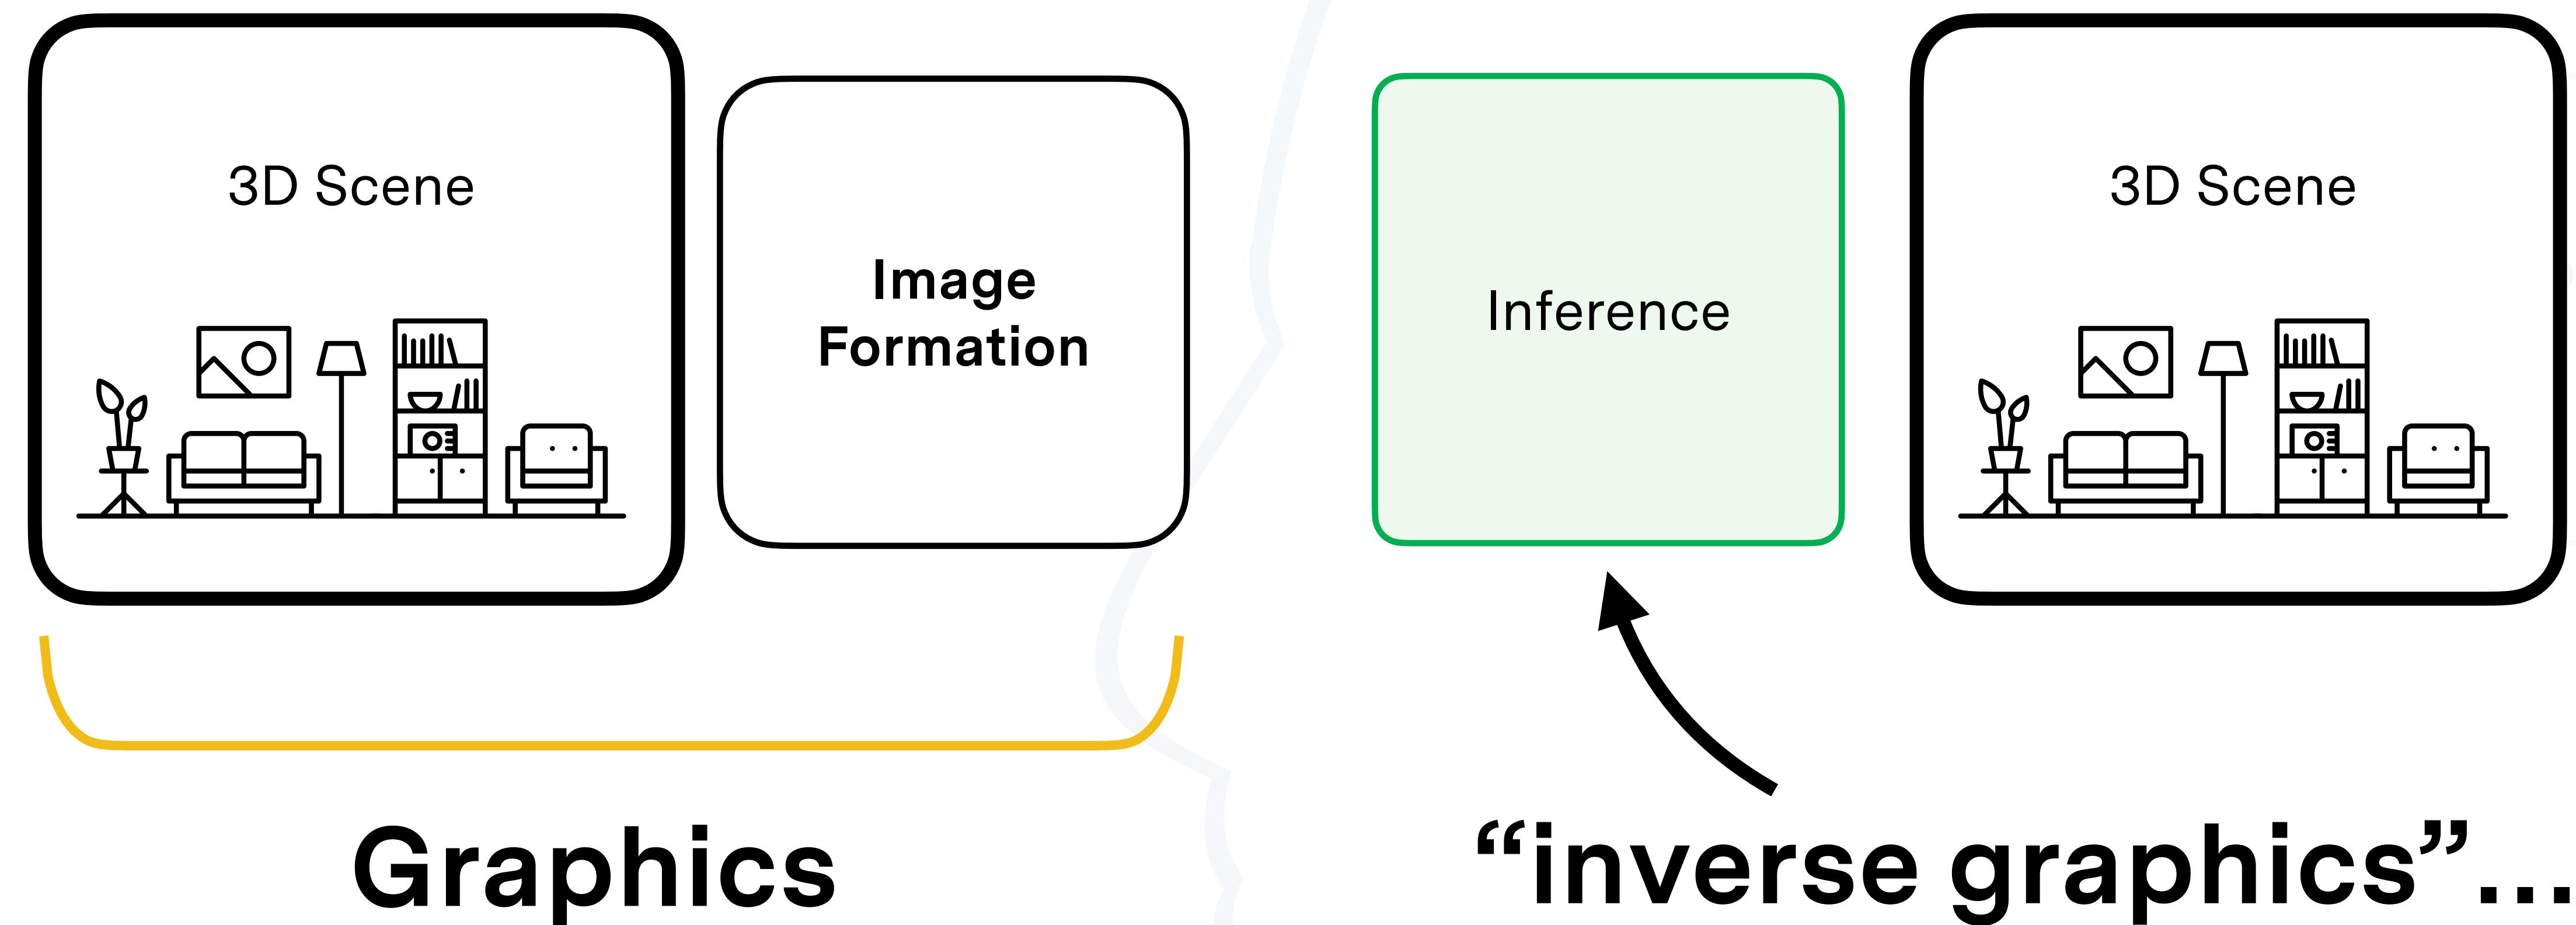
$$L_o(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) = L_e(\mathbf{x}, \boldsymbol{\omega}_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \boldsymbol{\omega}_i, \boldsymbol{\omega}_o, \lambda) L_i(\mathbf{x}, \boldsymbol{\omega}_i, \lambda) (\boldsymbol{\omega}_i \cdot \mathbf{n}) d\boldsymbol{\omega}_i$$

$$L(\mathbf{x}, \boldsymbol{\omega}) : \mathbb{R}^3 \times \mathbb{S}^2 \mapsto \mathbb{R}^3$$

Today: Differentiable Rendering, AKA “Inverse Graphics”



Today: Differentiable Rendering, AKA “Inverse Graphics”



Differentiable Rendering

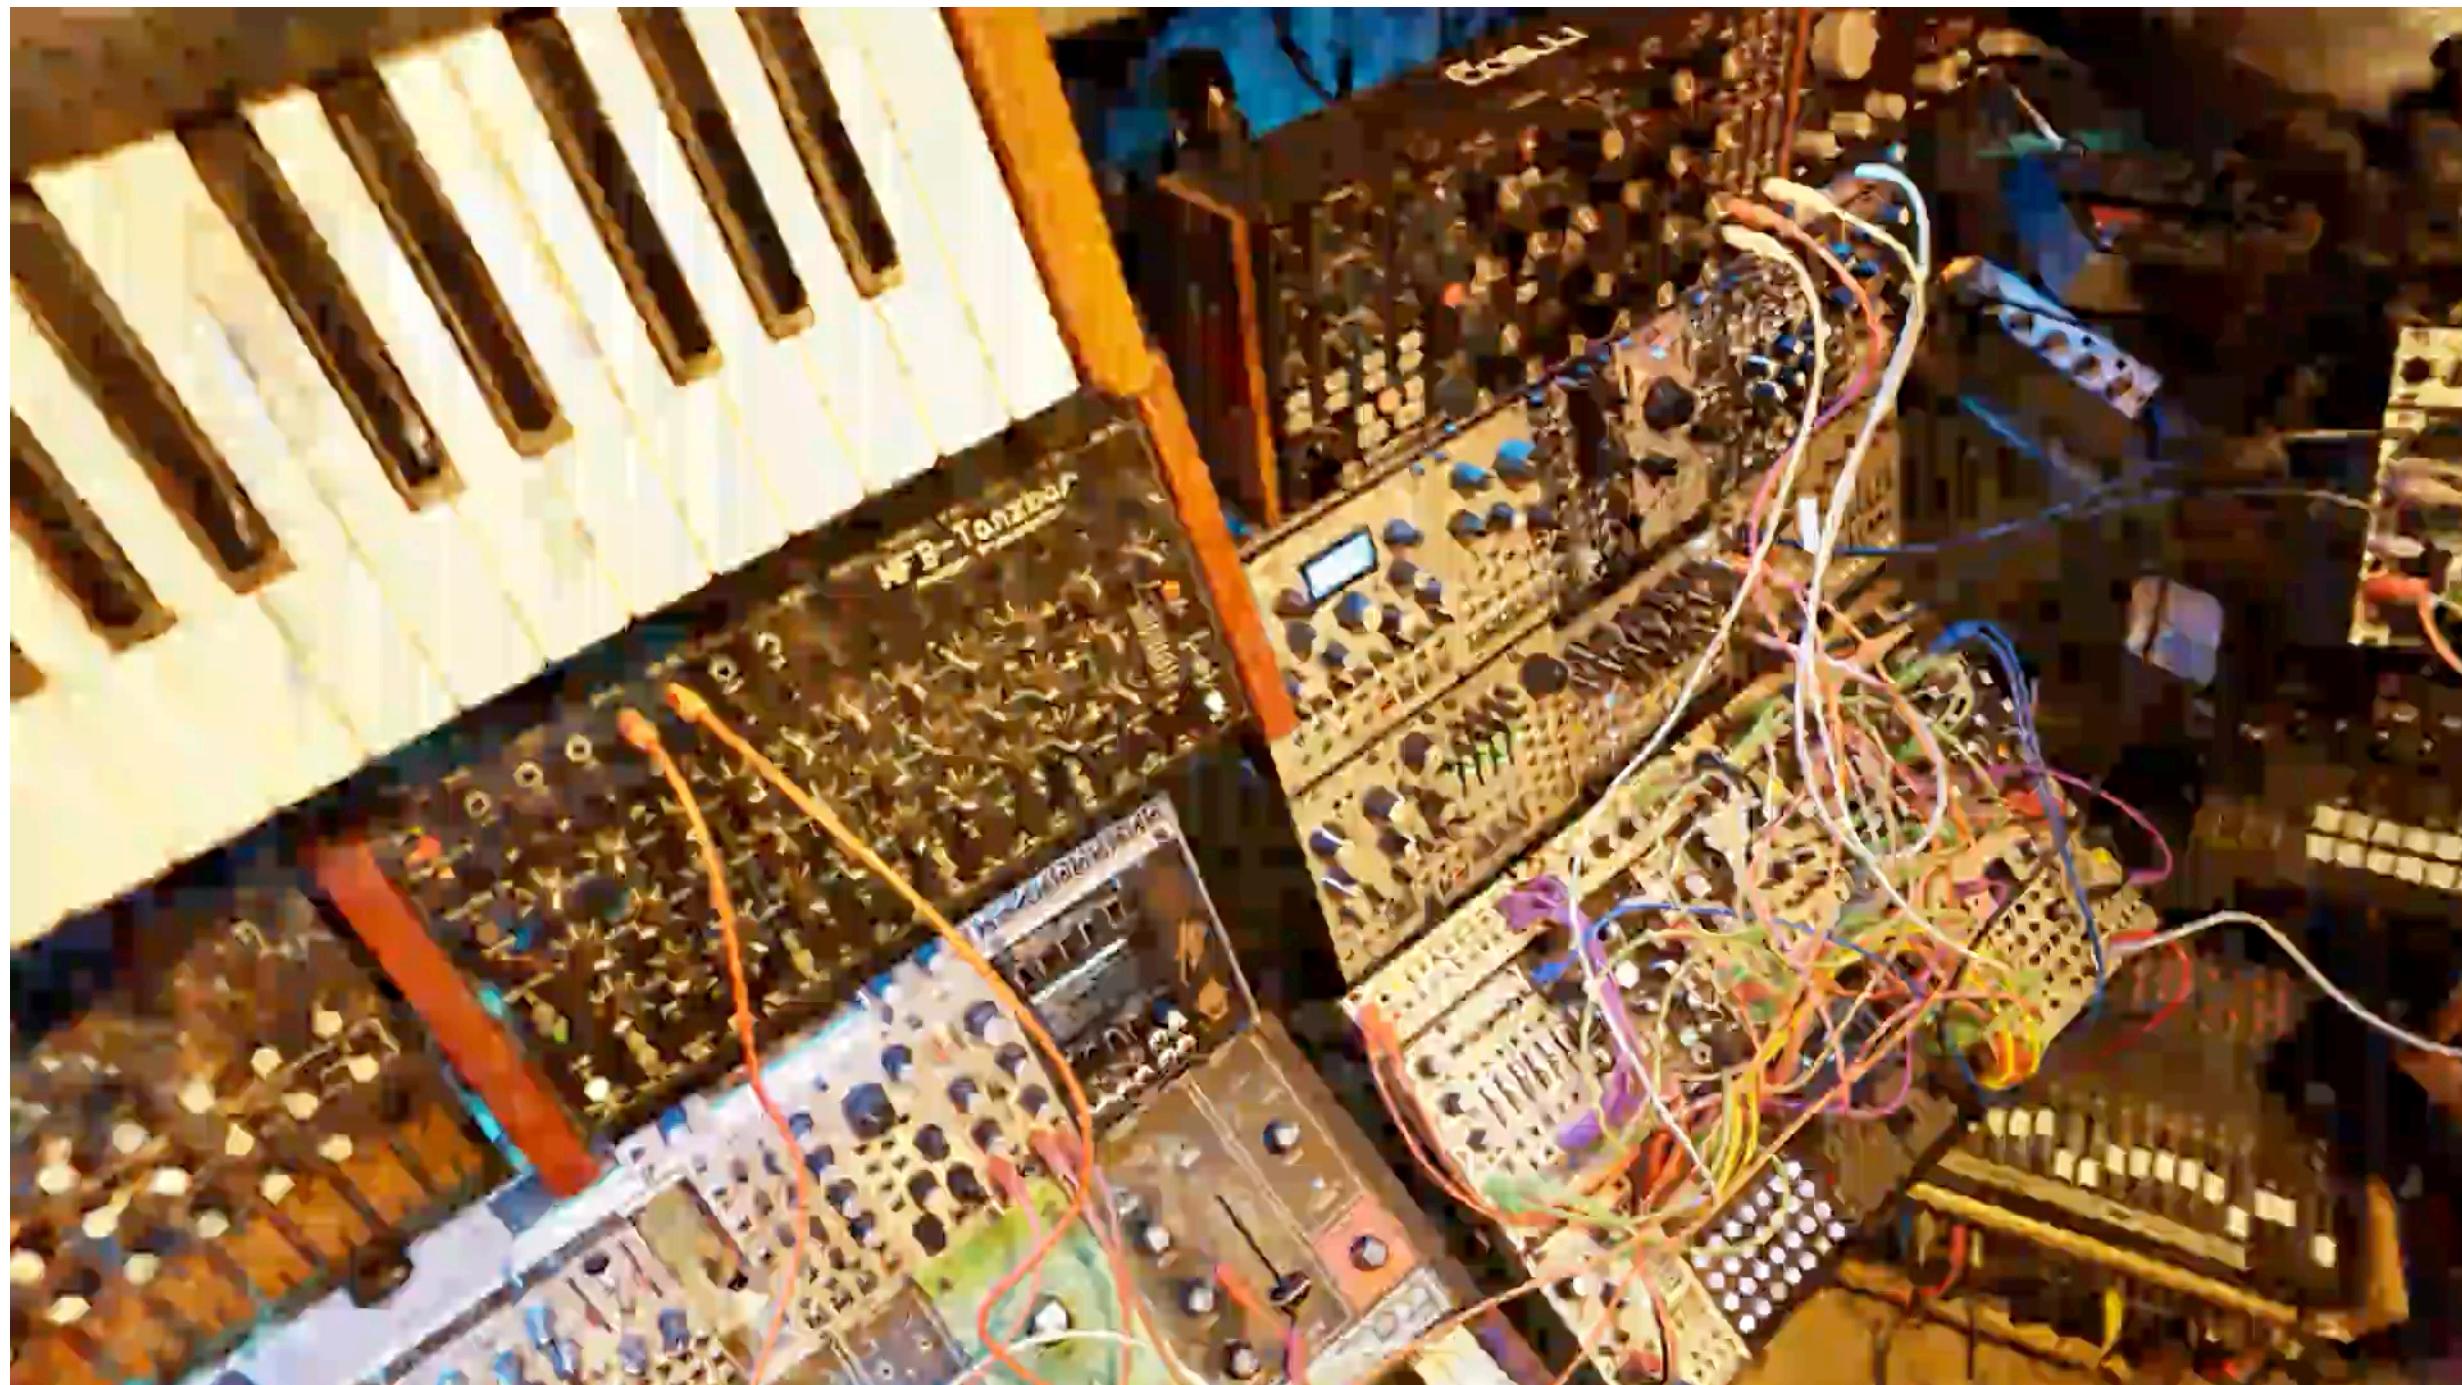
Optimization

Optimizing shape, albedo
& roughness



Iteration 0

Differentiable Signed Distance Function Rendering,
Vicini et al. 2022



InstantNGP, Müller et al. 2022

Why?

Part of our problem relates to **inverting** the rendering process: Given 2D images, we want to reconstruct 3D scenes. Differentiable rendering is one way of exactly inverting the rendering process.

What you'll learn.

The structure of differentiable renderers, pros and cons and assumptions of different algorithms.

Differentiable Rendering

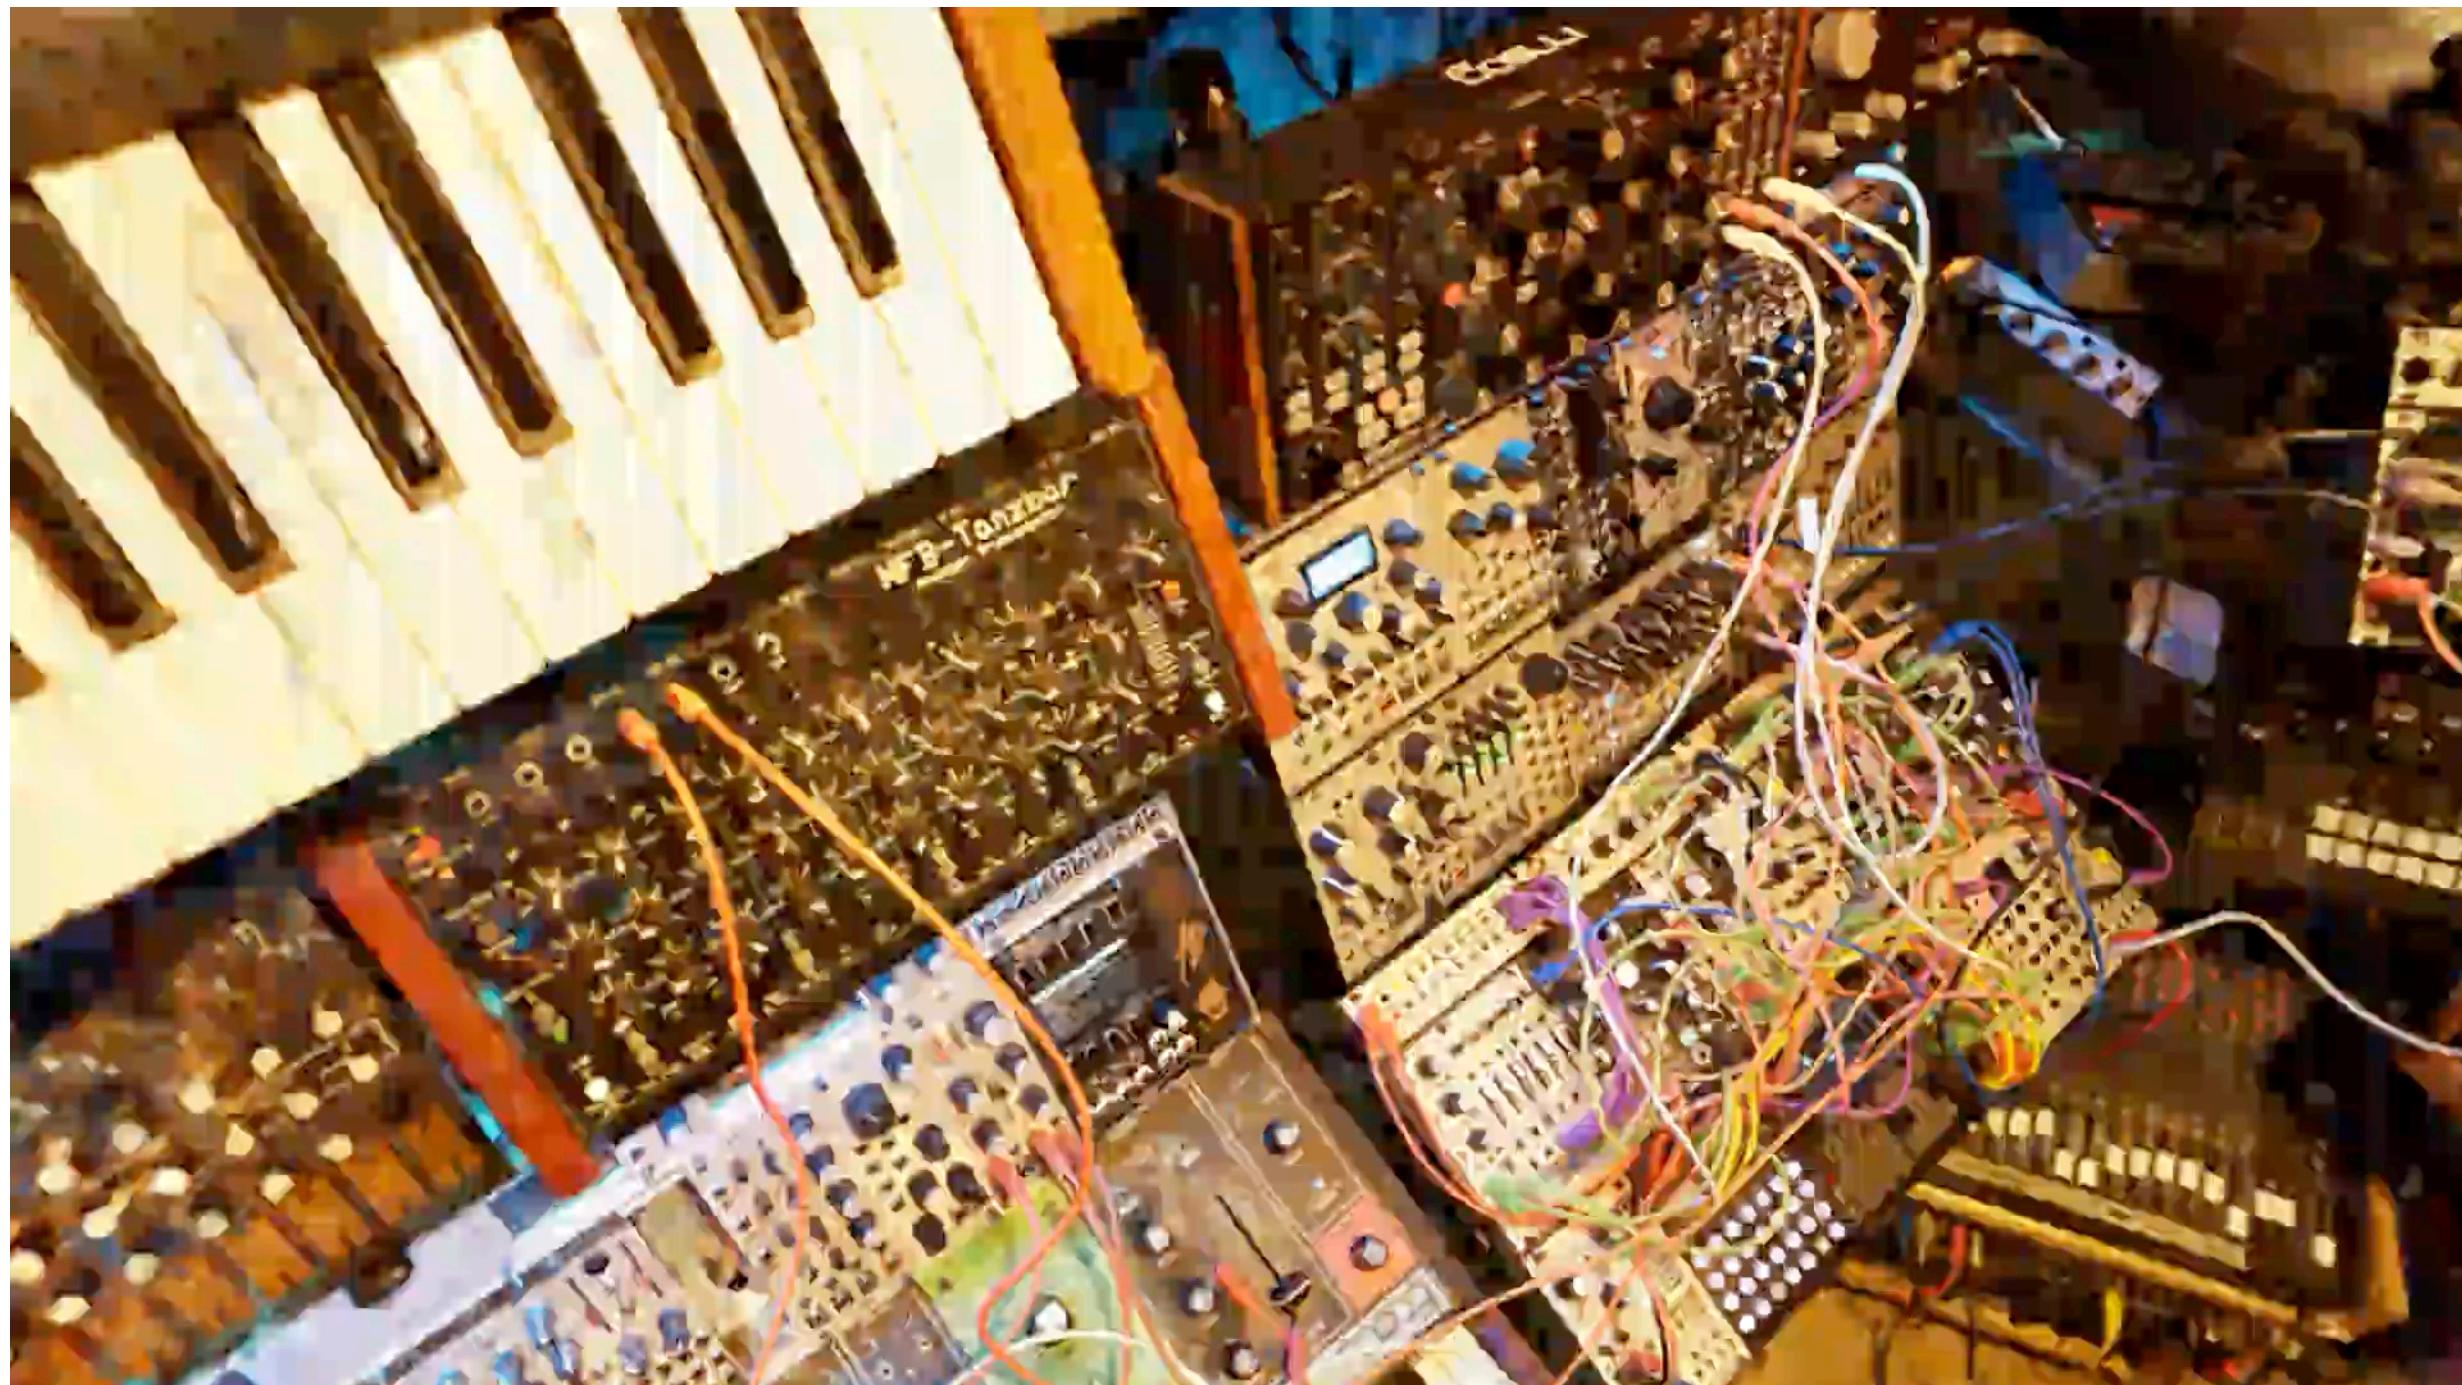
Optimization

Optimizing shape, albedo
& roughness



Iteration 0

Differentiable Signed Distance Function Rendering,
Vicini et al. 2022



InstantNGP, Müller et al. 2022

Why?

Part of our problem relates to **inverting** the rendering process: Given 2D images, we want to reconstruct 3D scenes. Differentiable rendering is one way of exactly inverting the rendering process.

What you'll learn.

The structure of differentiable renderers, pros and cons and assumptions of different algorithms.

Differentiable Rendering

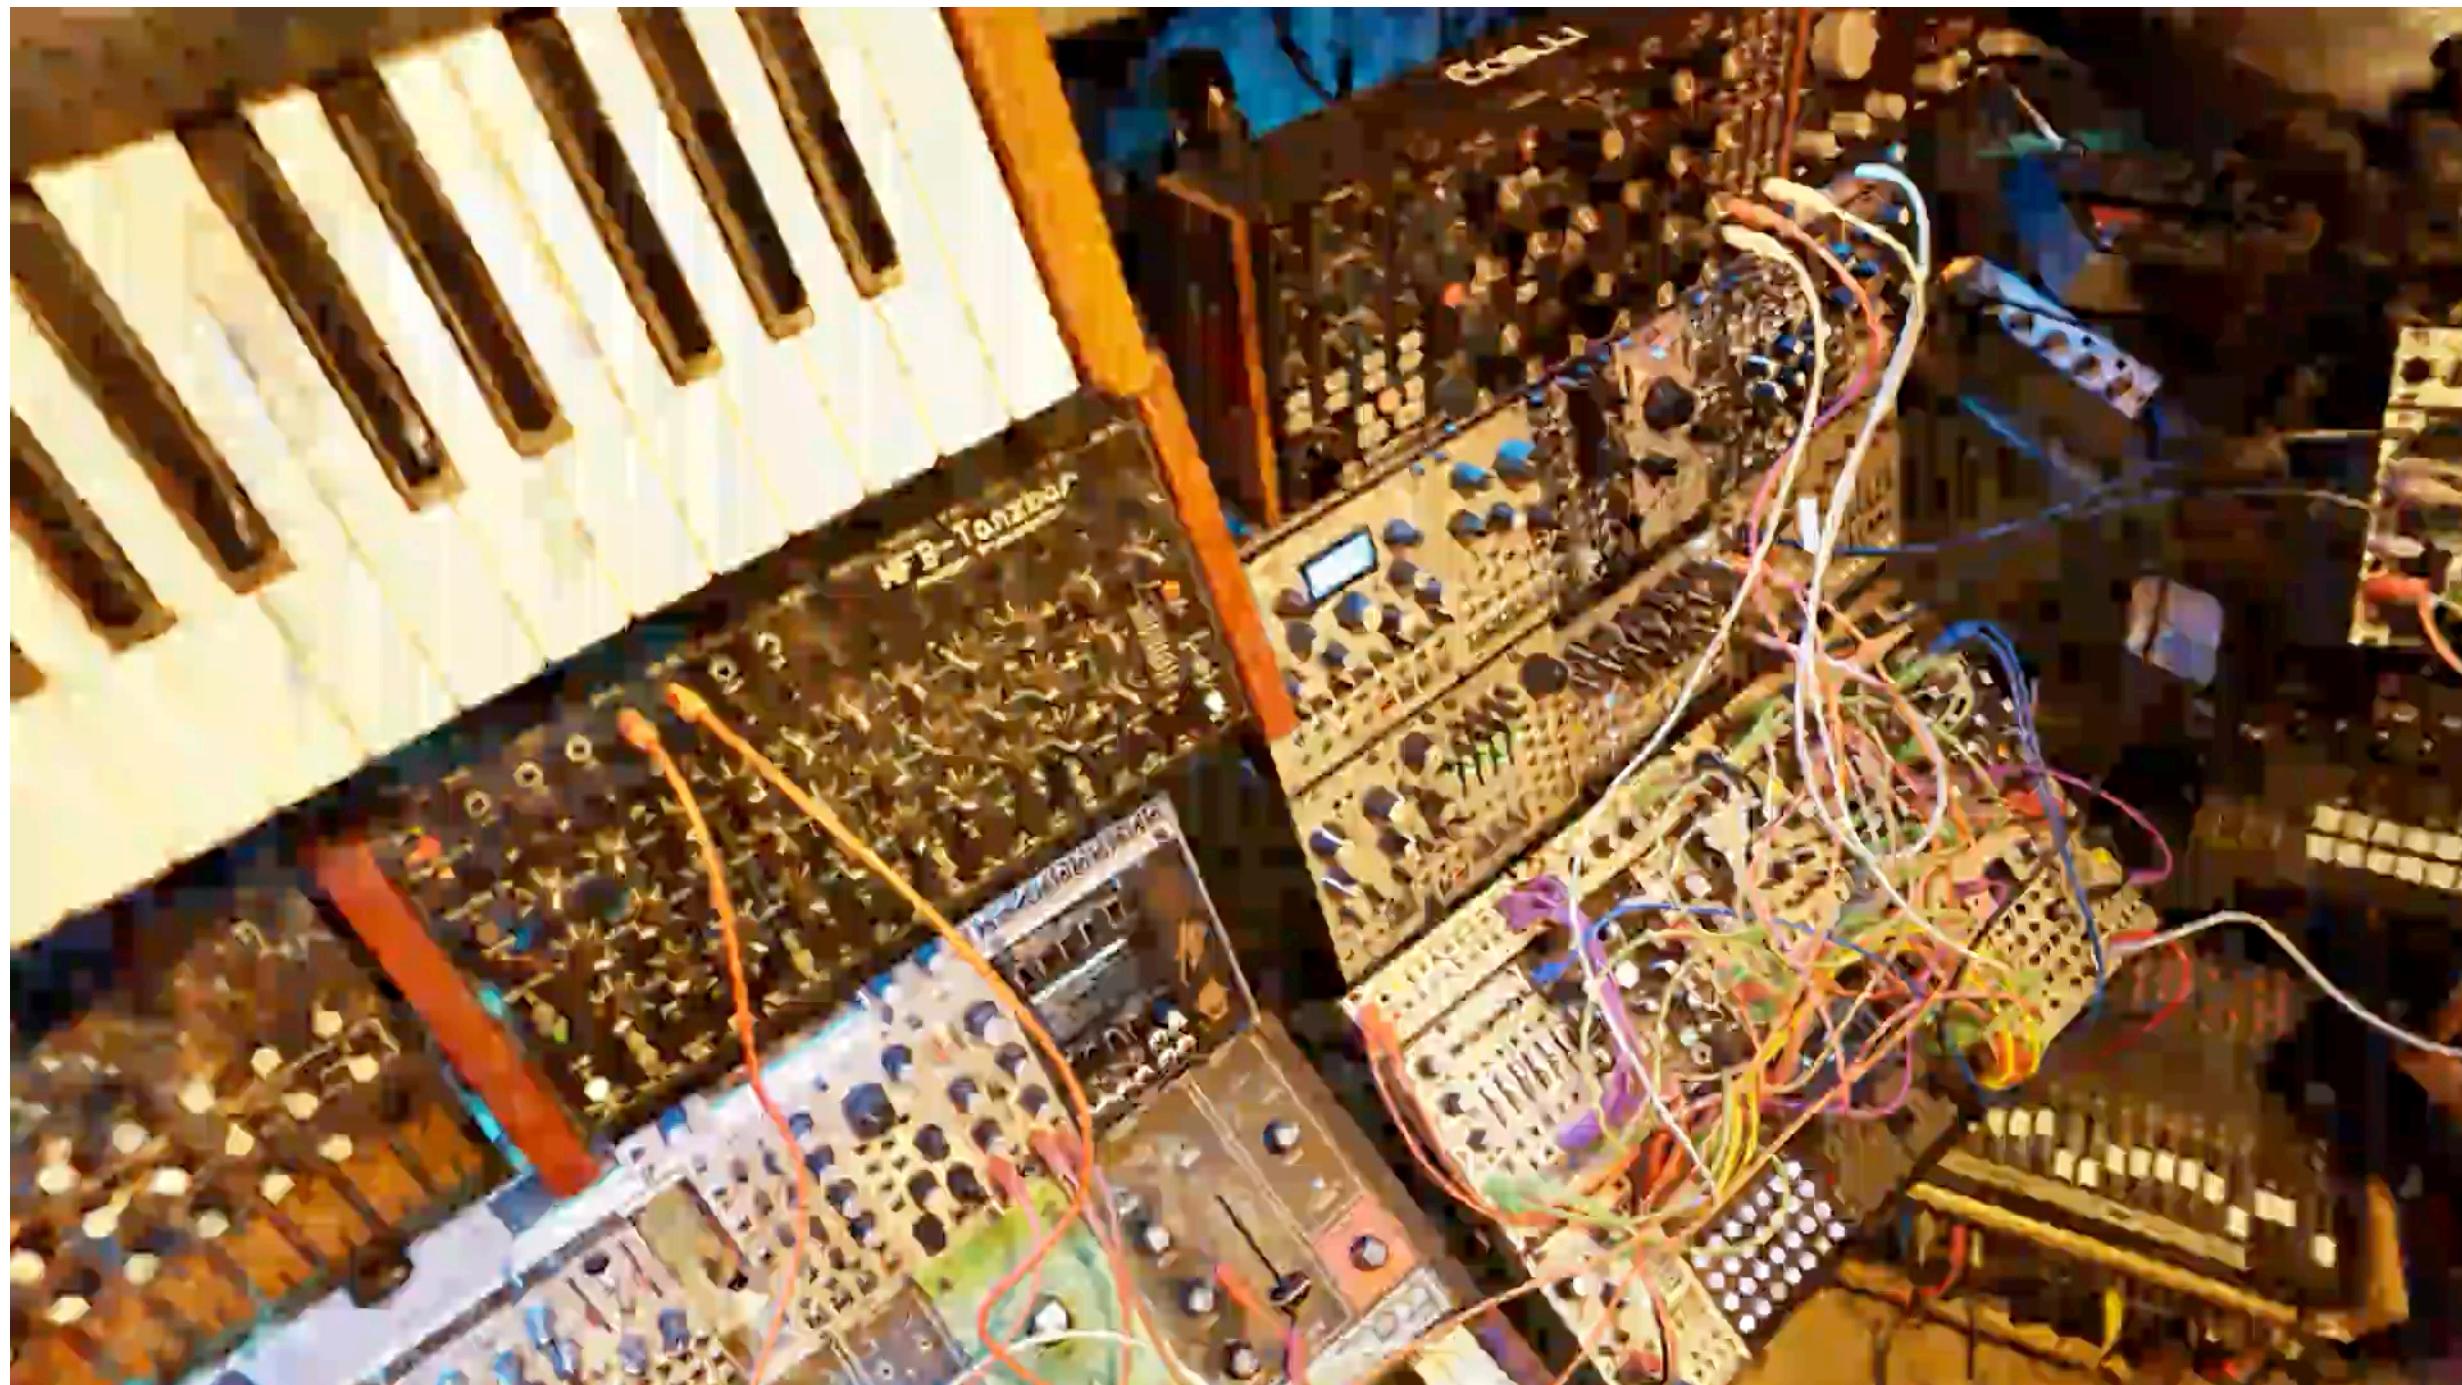
Optimization

Optimizing shape, albedo
& roughness



Iteration 0

Differentiable Signed Distance Function Rendering,
Vicini et al. 2022



InstantNGP, Müller et al. 2022

Why?

Part of our problem relates to **inverting** the rendering process: Given 2D images, we want to reconstruct 3D scenes. Differentiable rendering is one way of exactly inverting the rendering process.

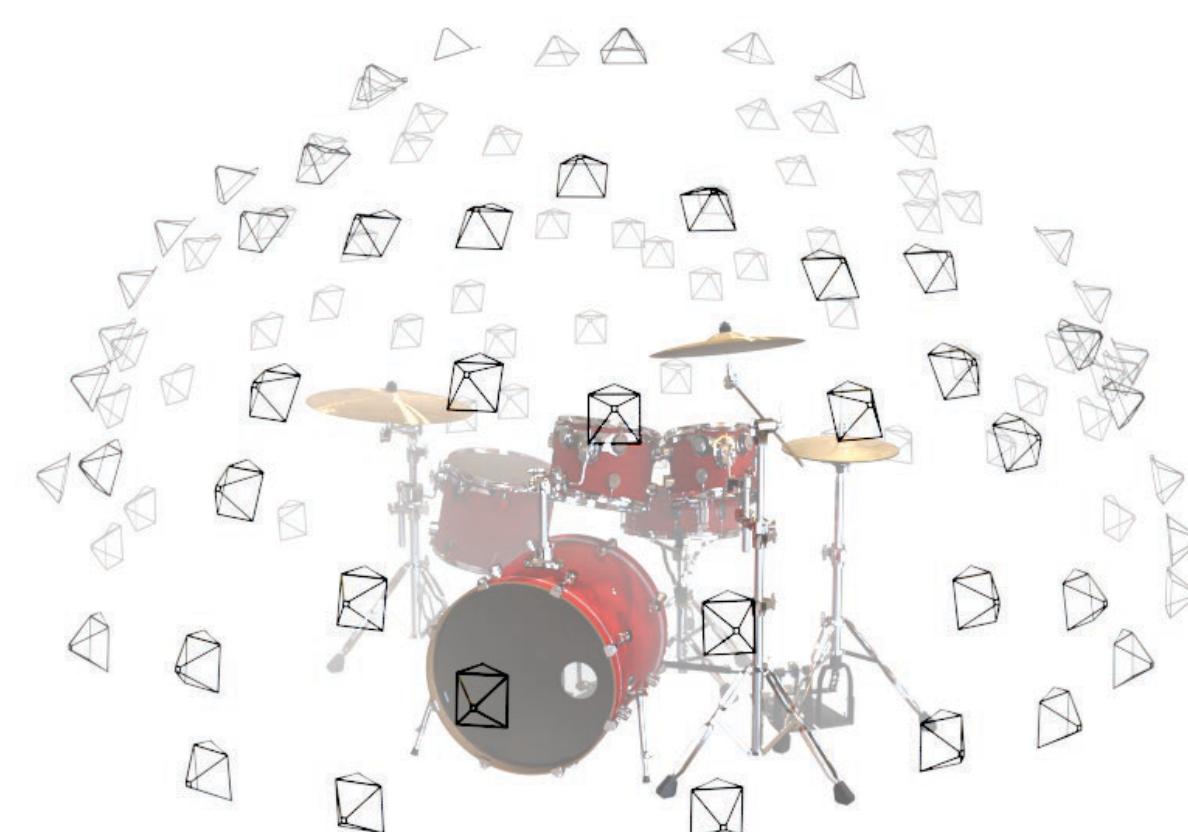
What you'll learn.

The structure of differentiable renderers, pros and cons and assumptions of different algorithms.

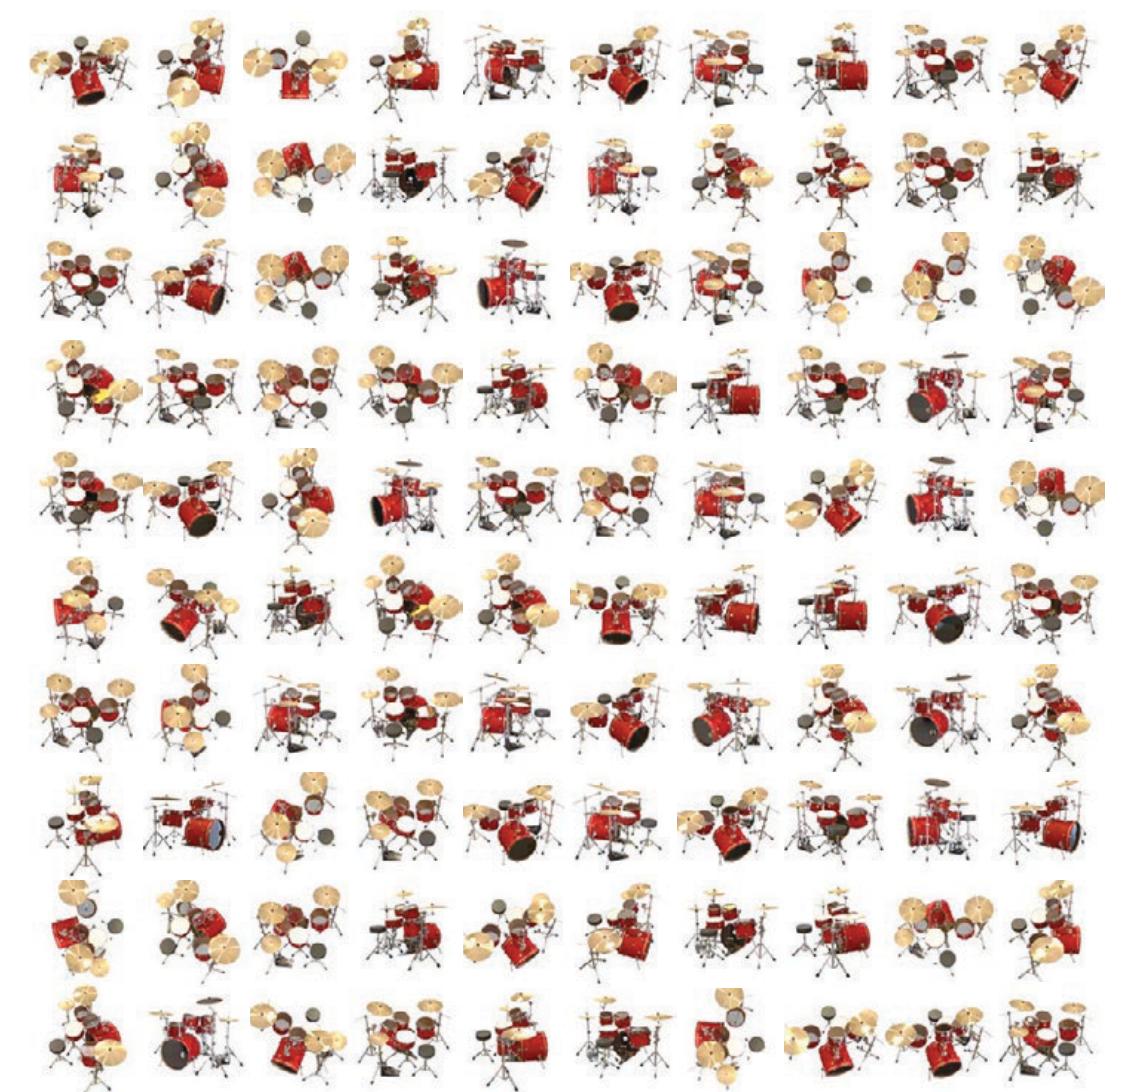
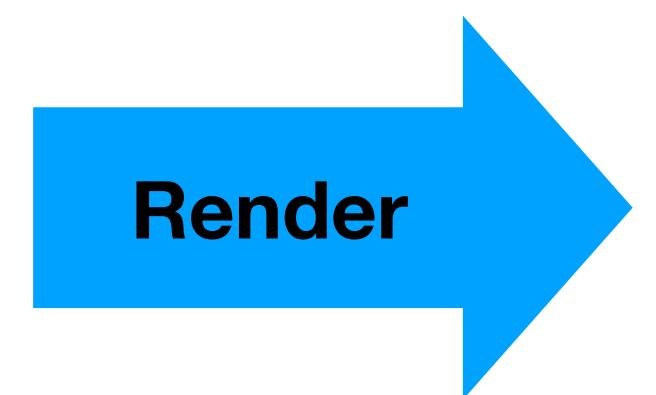
Rendering (Graphics): Given 3D Scene + Camera parameters, yield images



3D Scene



Camera Poses



Images

Inverse Graphics: Given Images, Infer Camera Poses & 3D Scene!



Reconstruct

Images



3D Scene



Camera Poses

How to get camera poses?

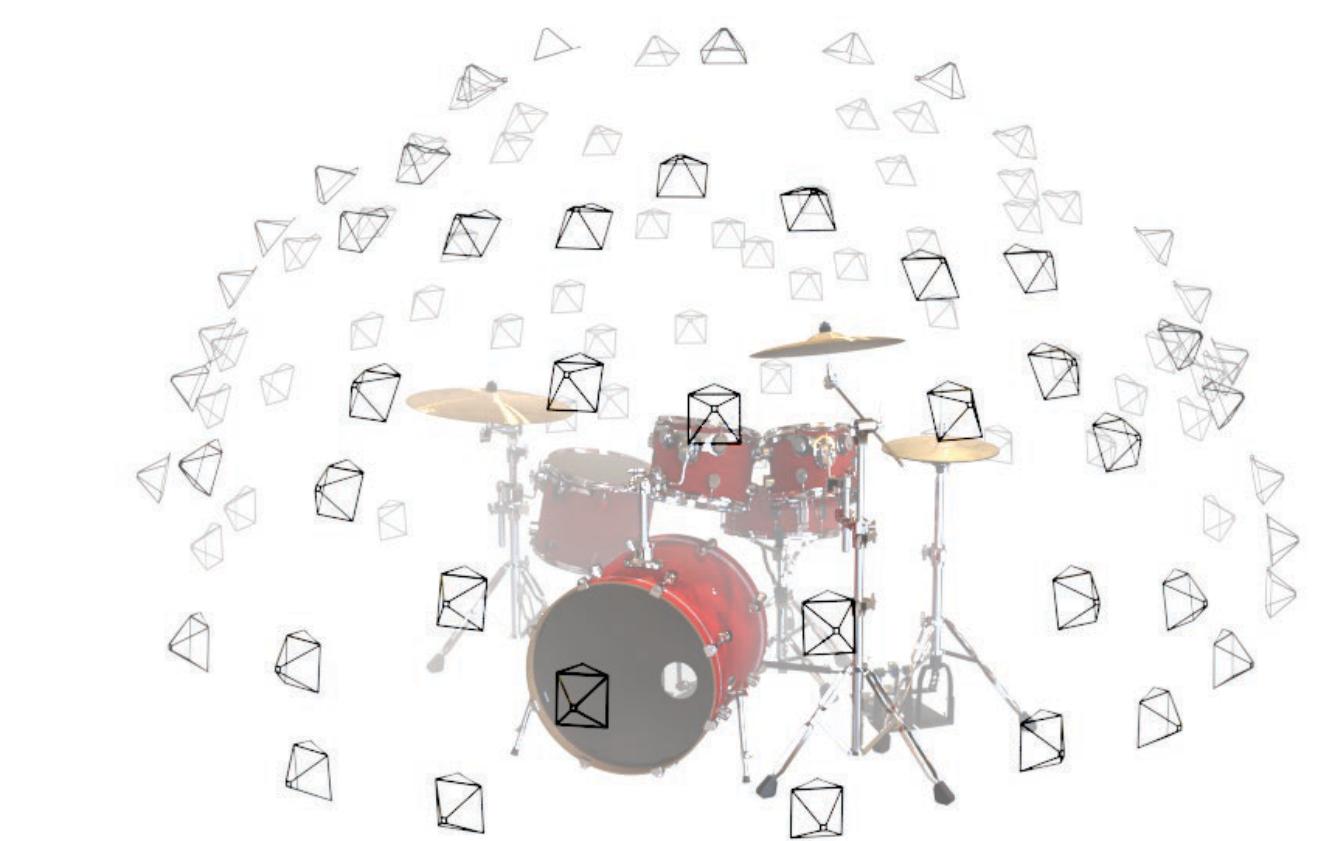


Reconstruct



Images

3D Scene



Camera Poses

Bundle Adjustment (Multi-View Geometry Lecture)!

Bundle adjustment minimizes the reprojection error of all observations:

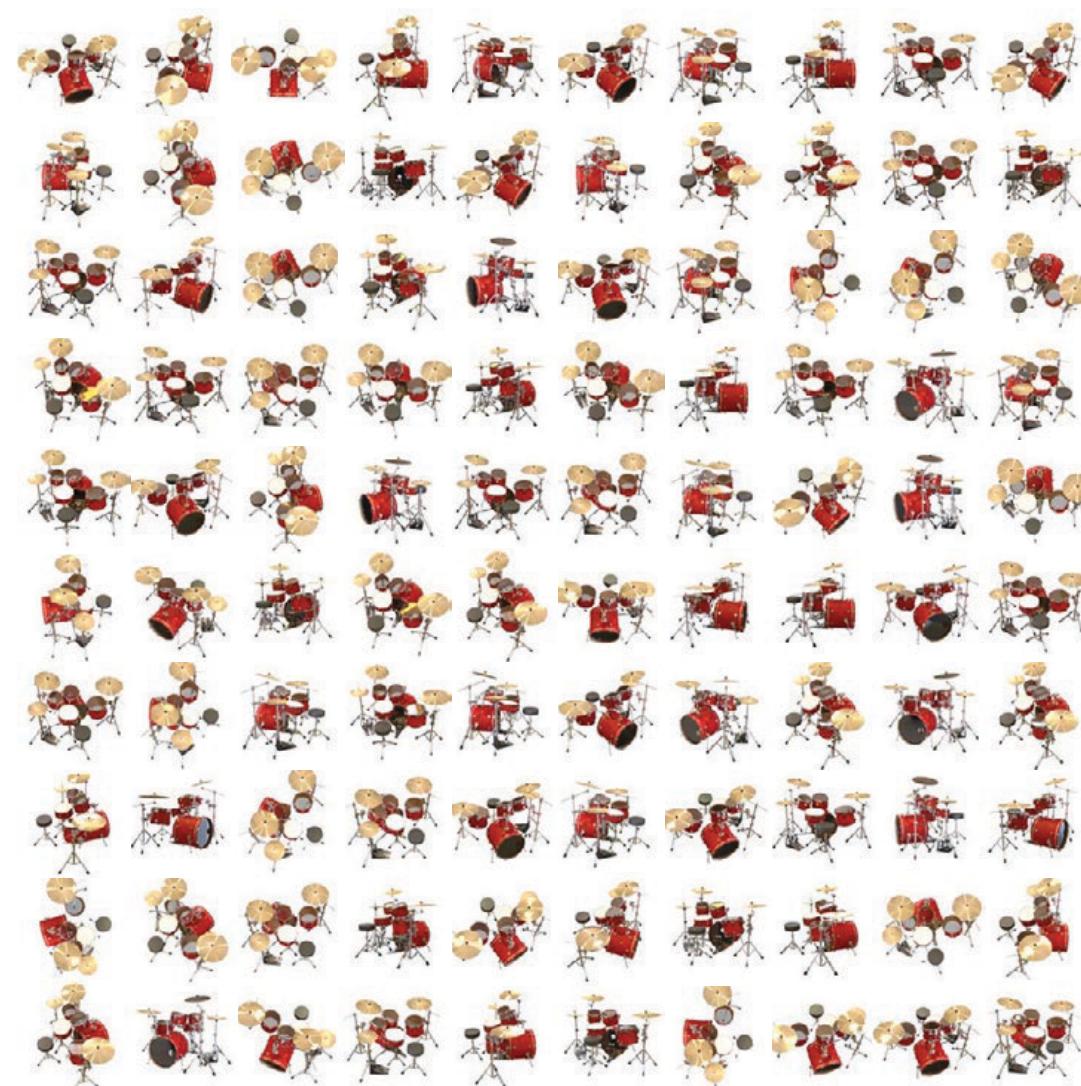
$$\Pi^*, \mathcal{X}_w^* = \operatorname{argmin}_{\Pi, \mathcal{X}_w} \sum_{i=1}^N \sum_{p=1}^P w_{ip} \|\mathbf{x}_{ip}^s - \pi_i(\mathbf{x}_p^w)\|_2^2$$

Images

3D Scene

Camera Poses

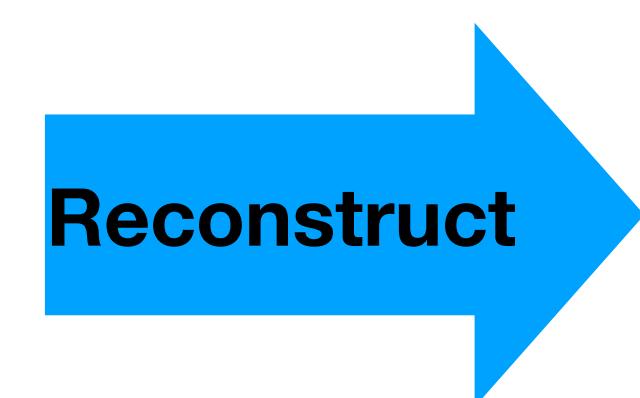
Can assume we know the camera poses.



Images

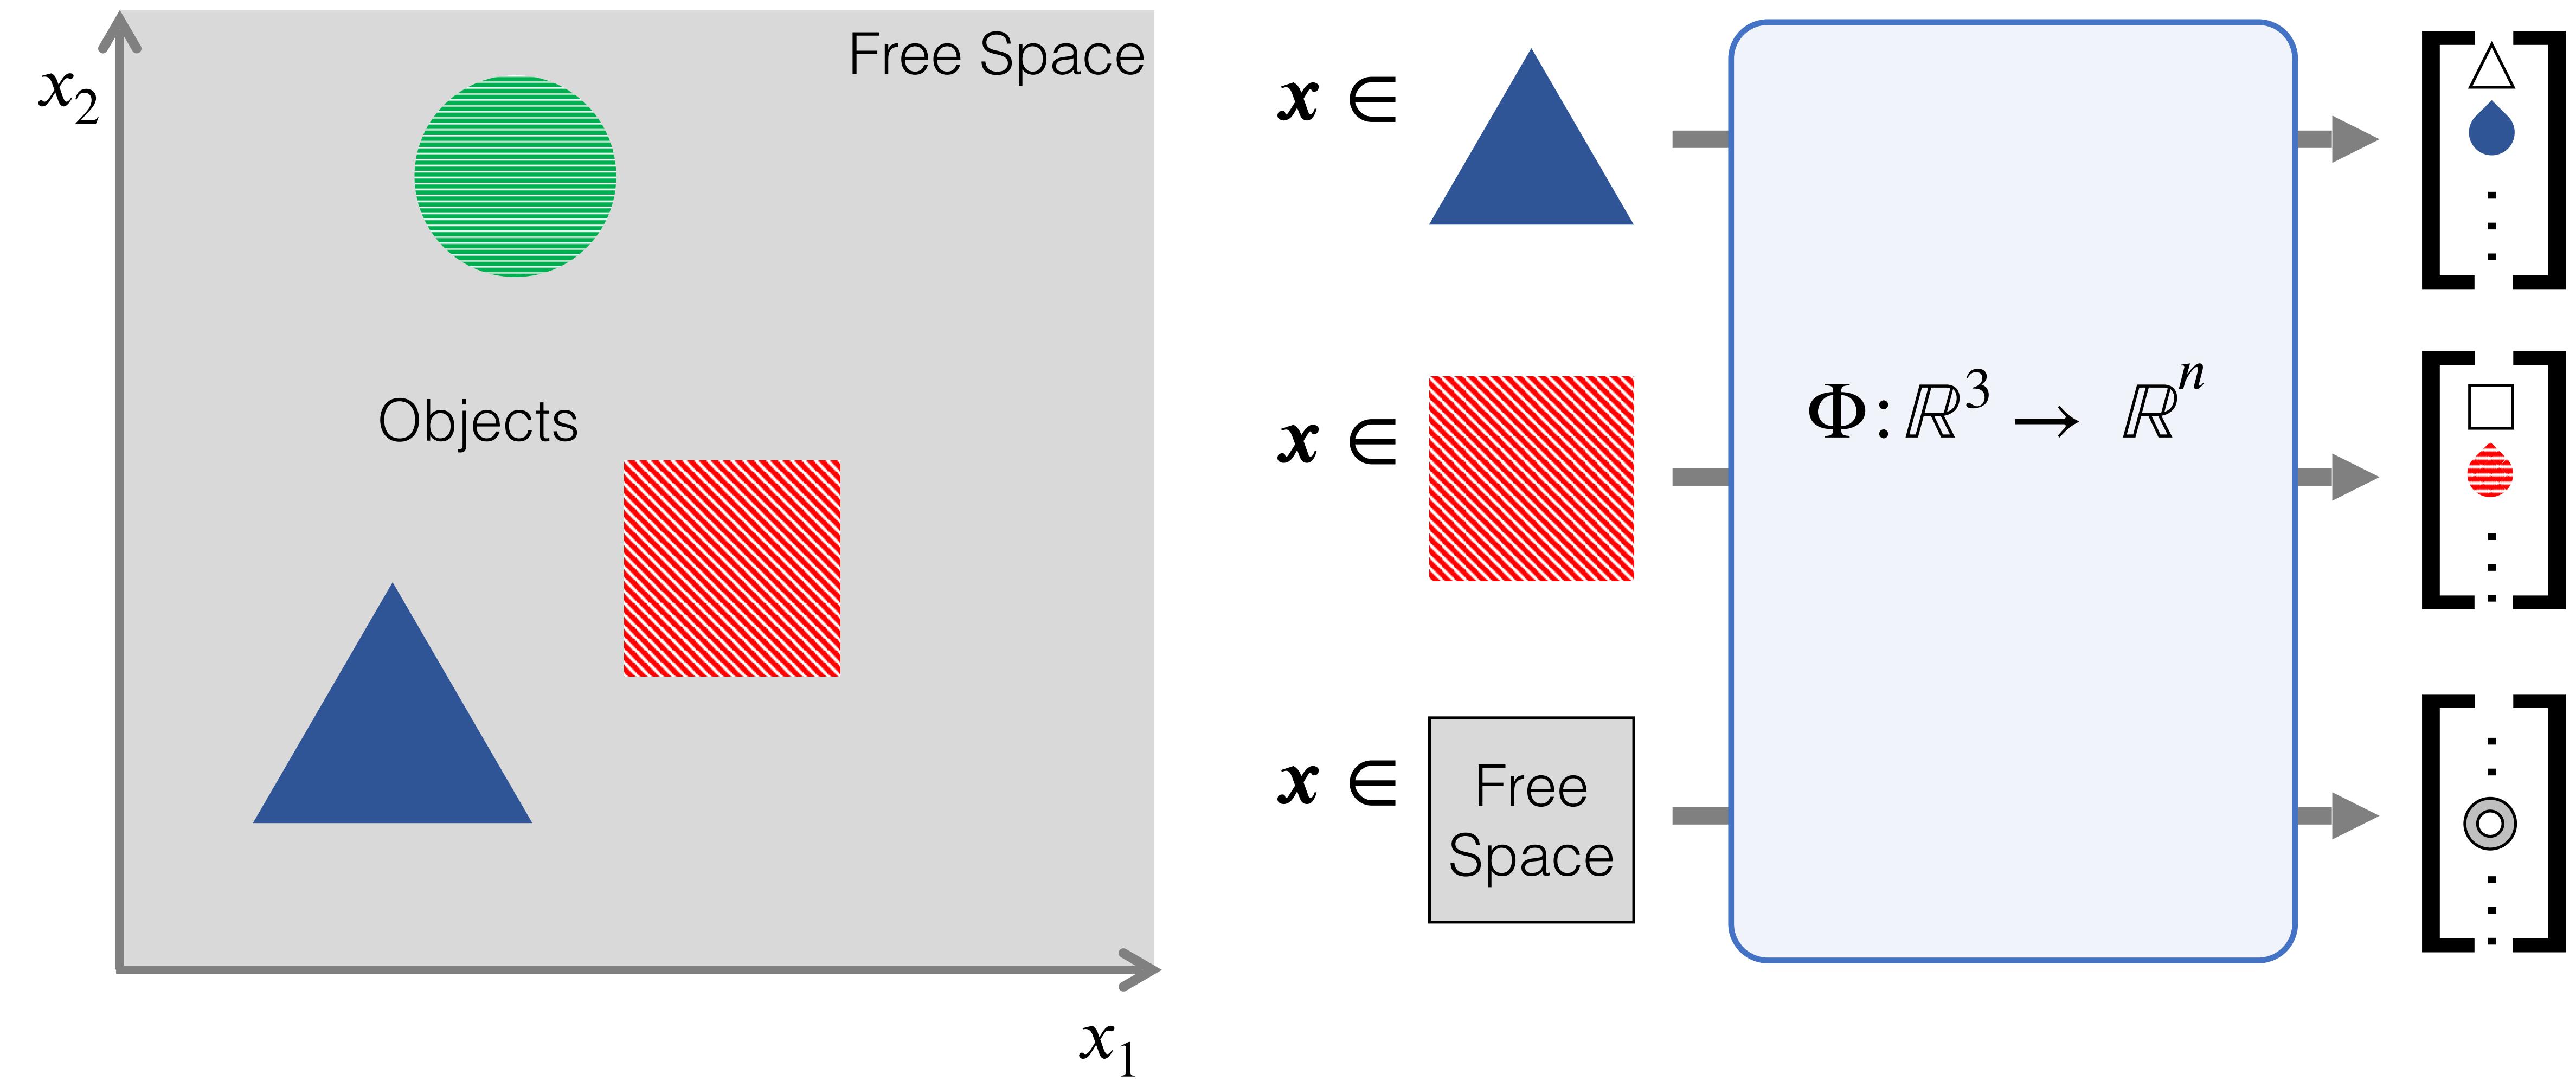


Camera Poses

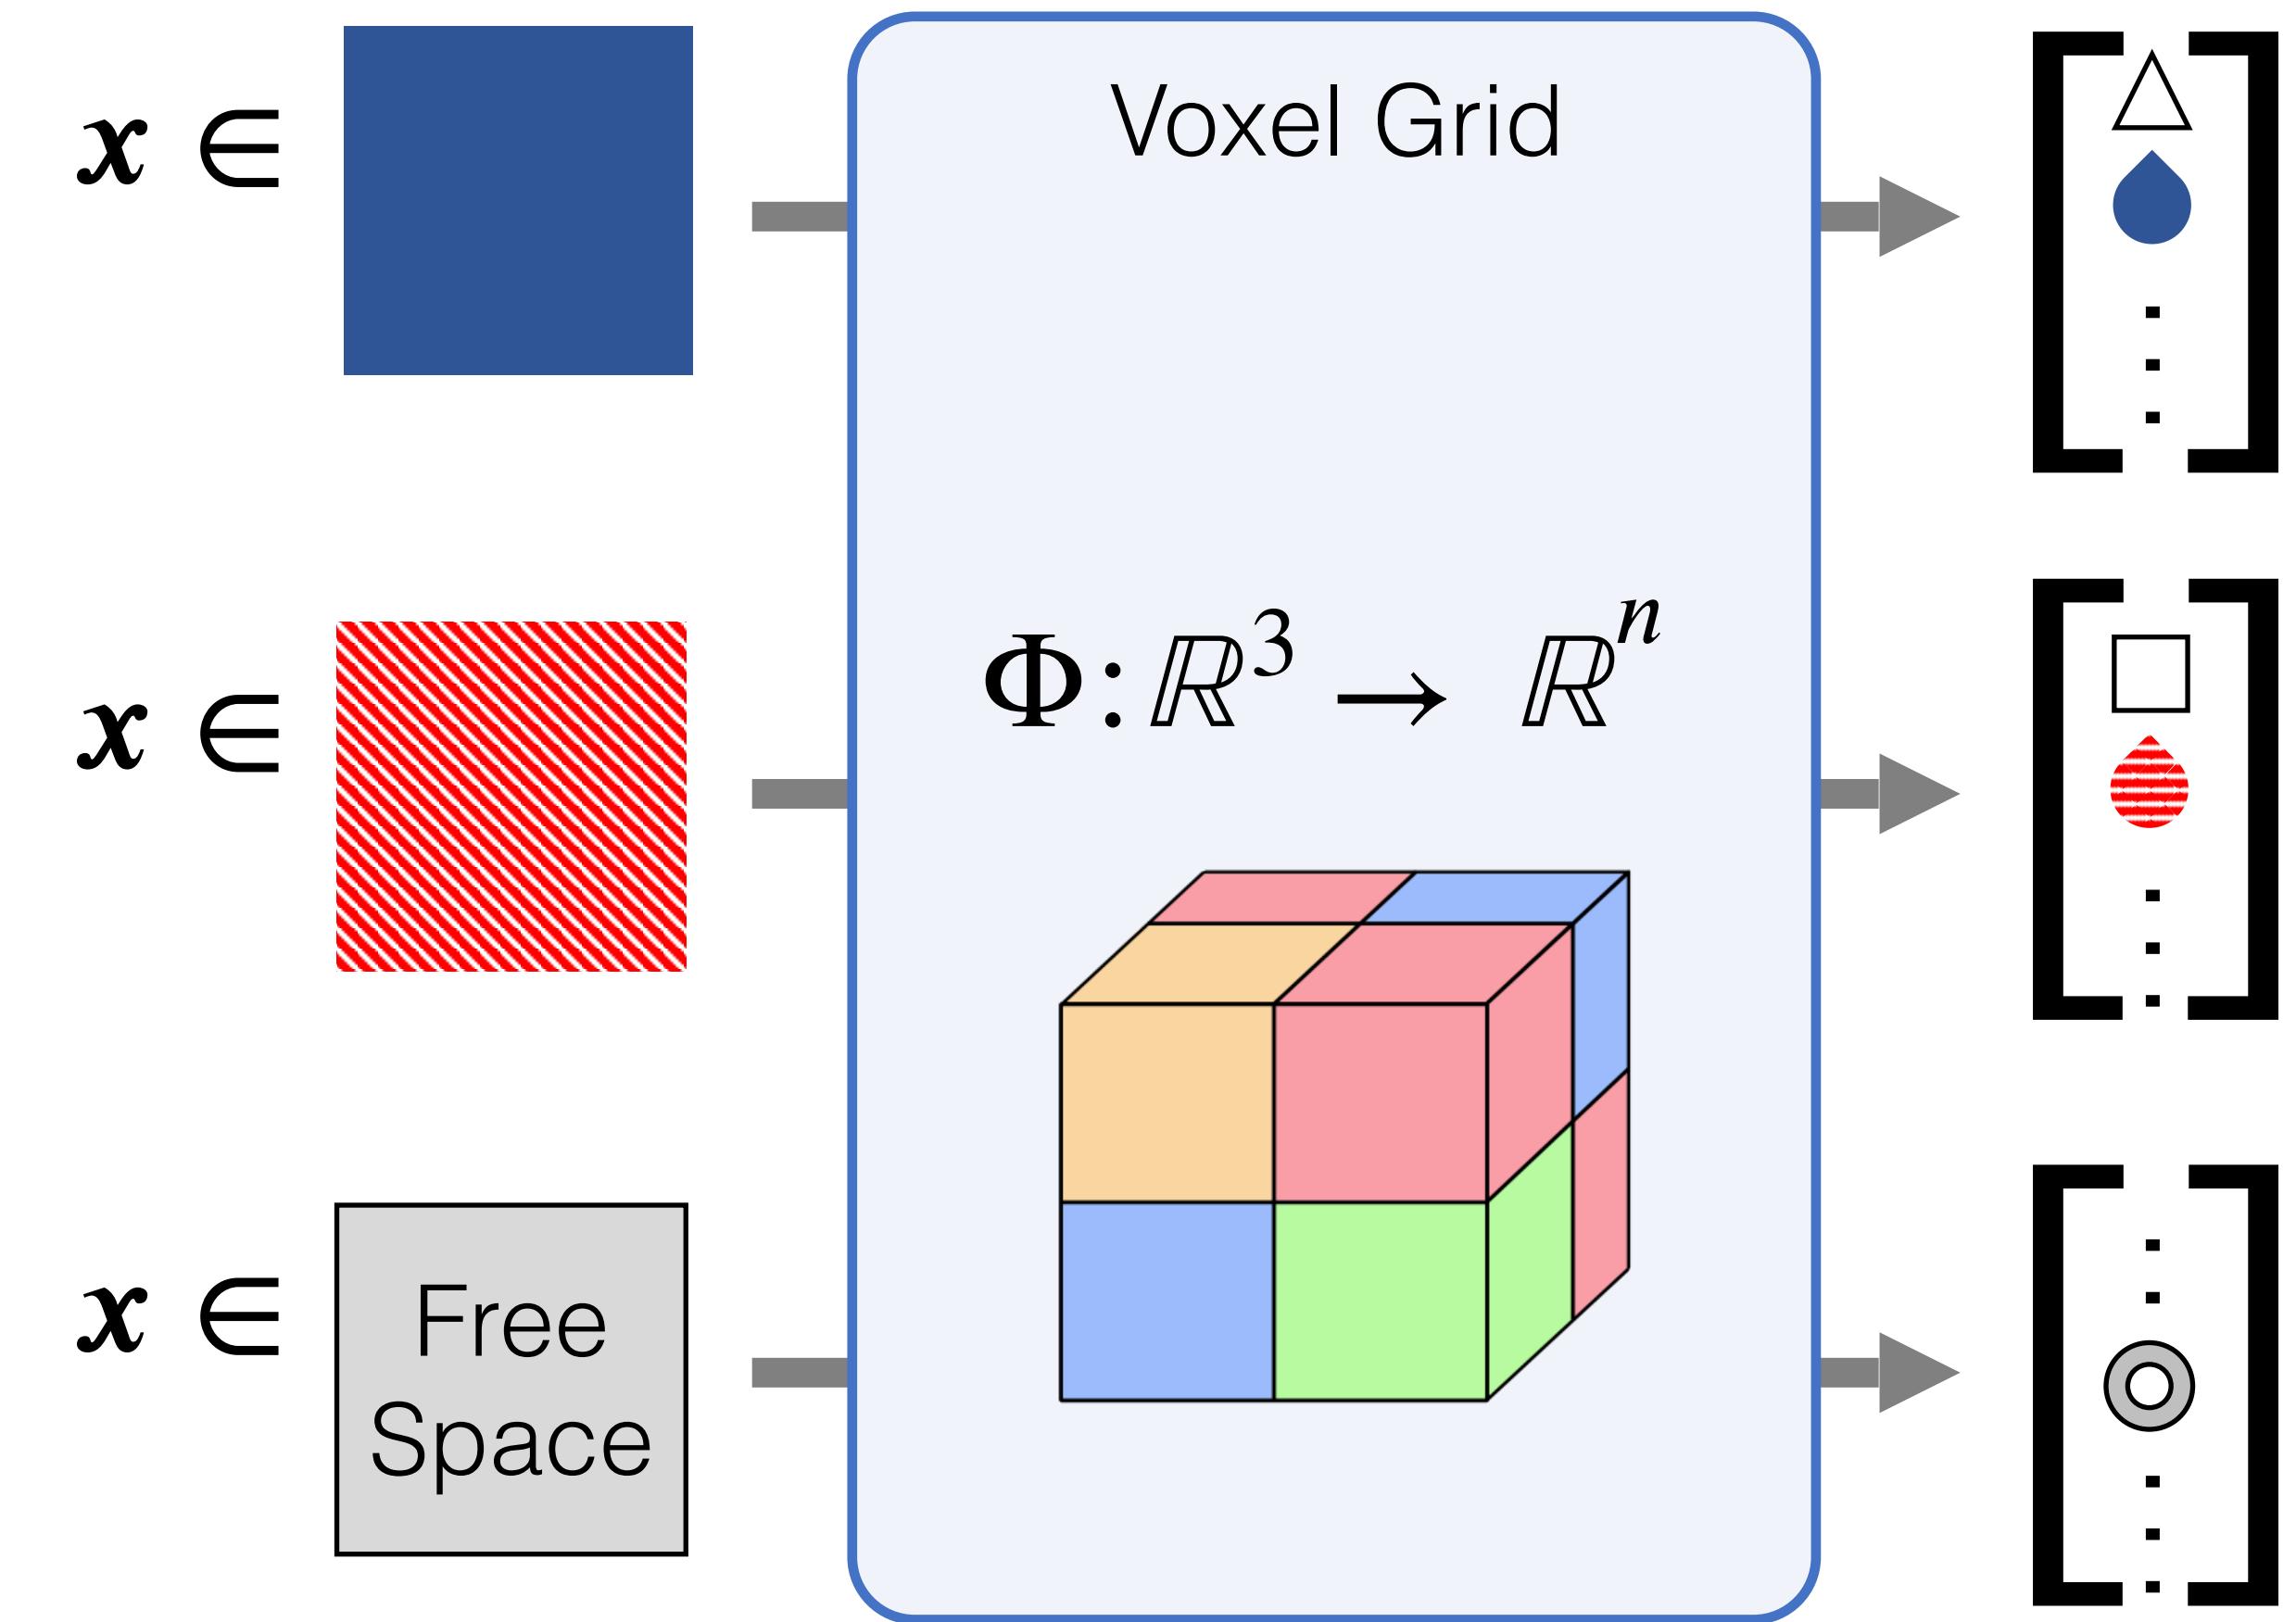
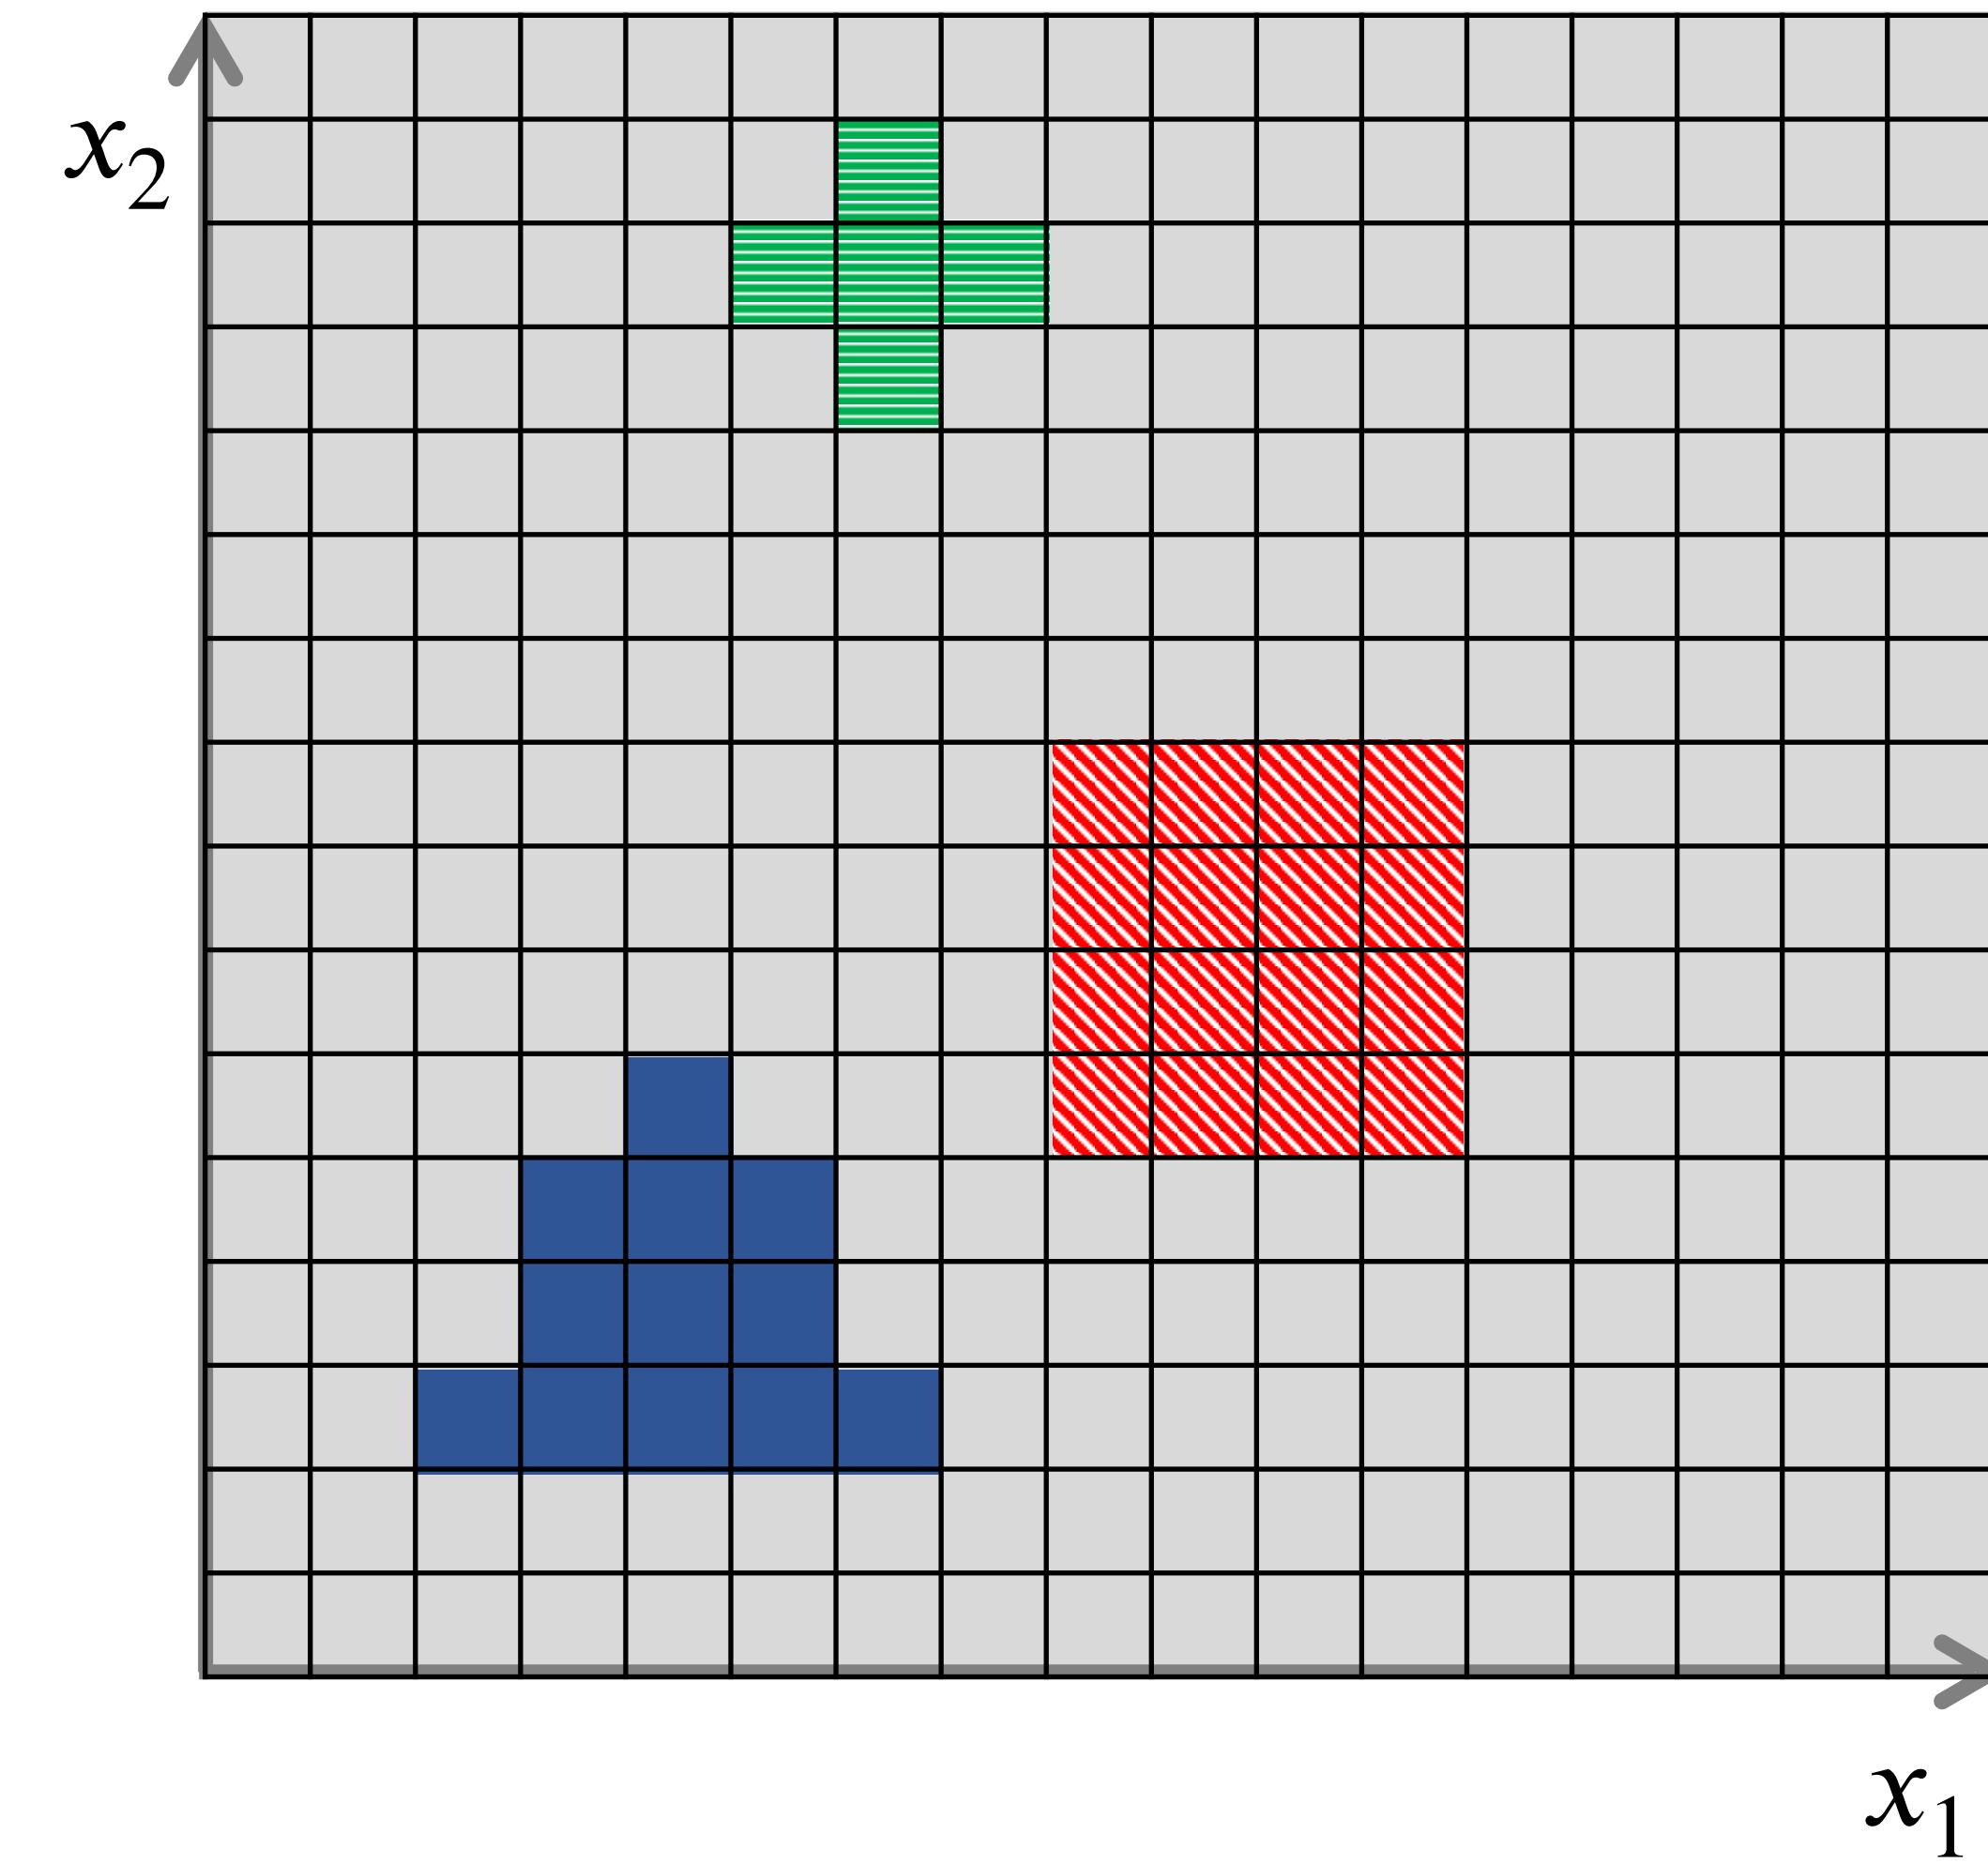


3D Scene

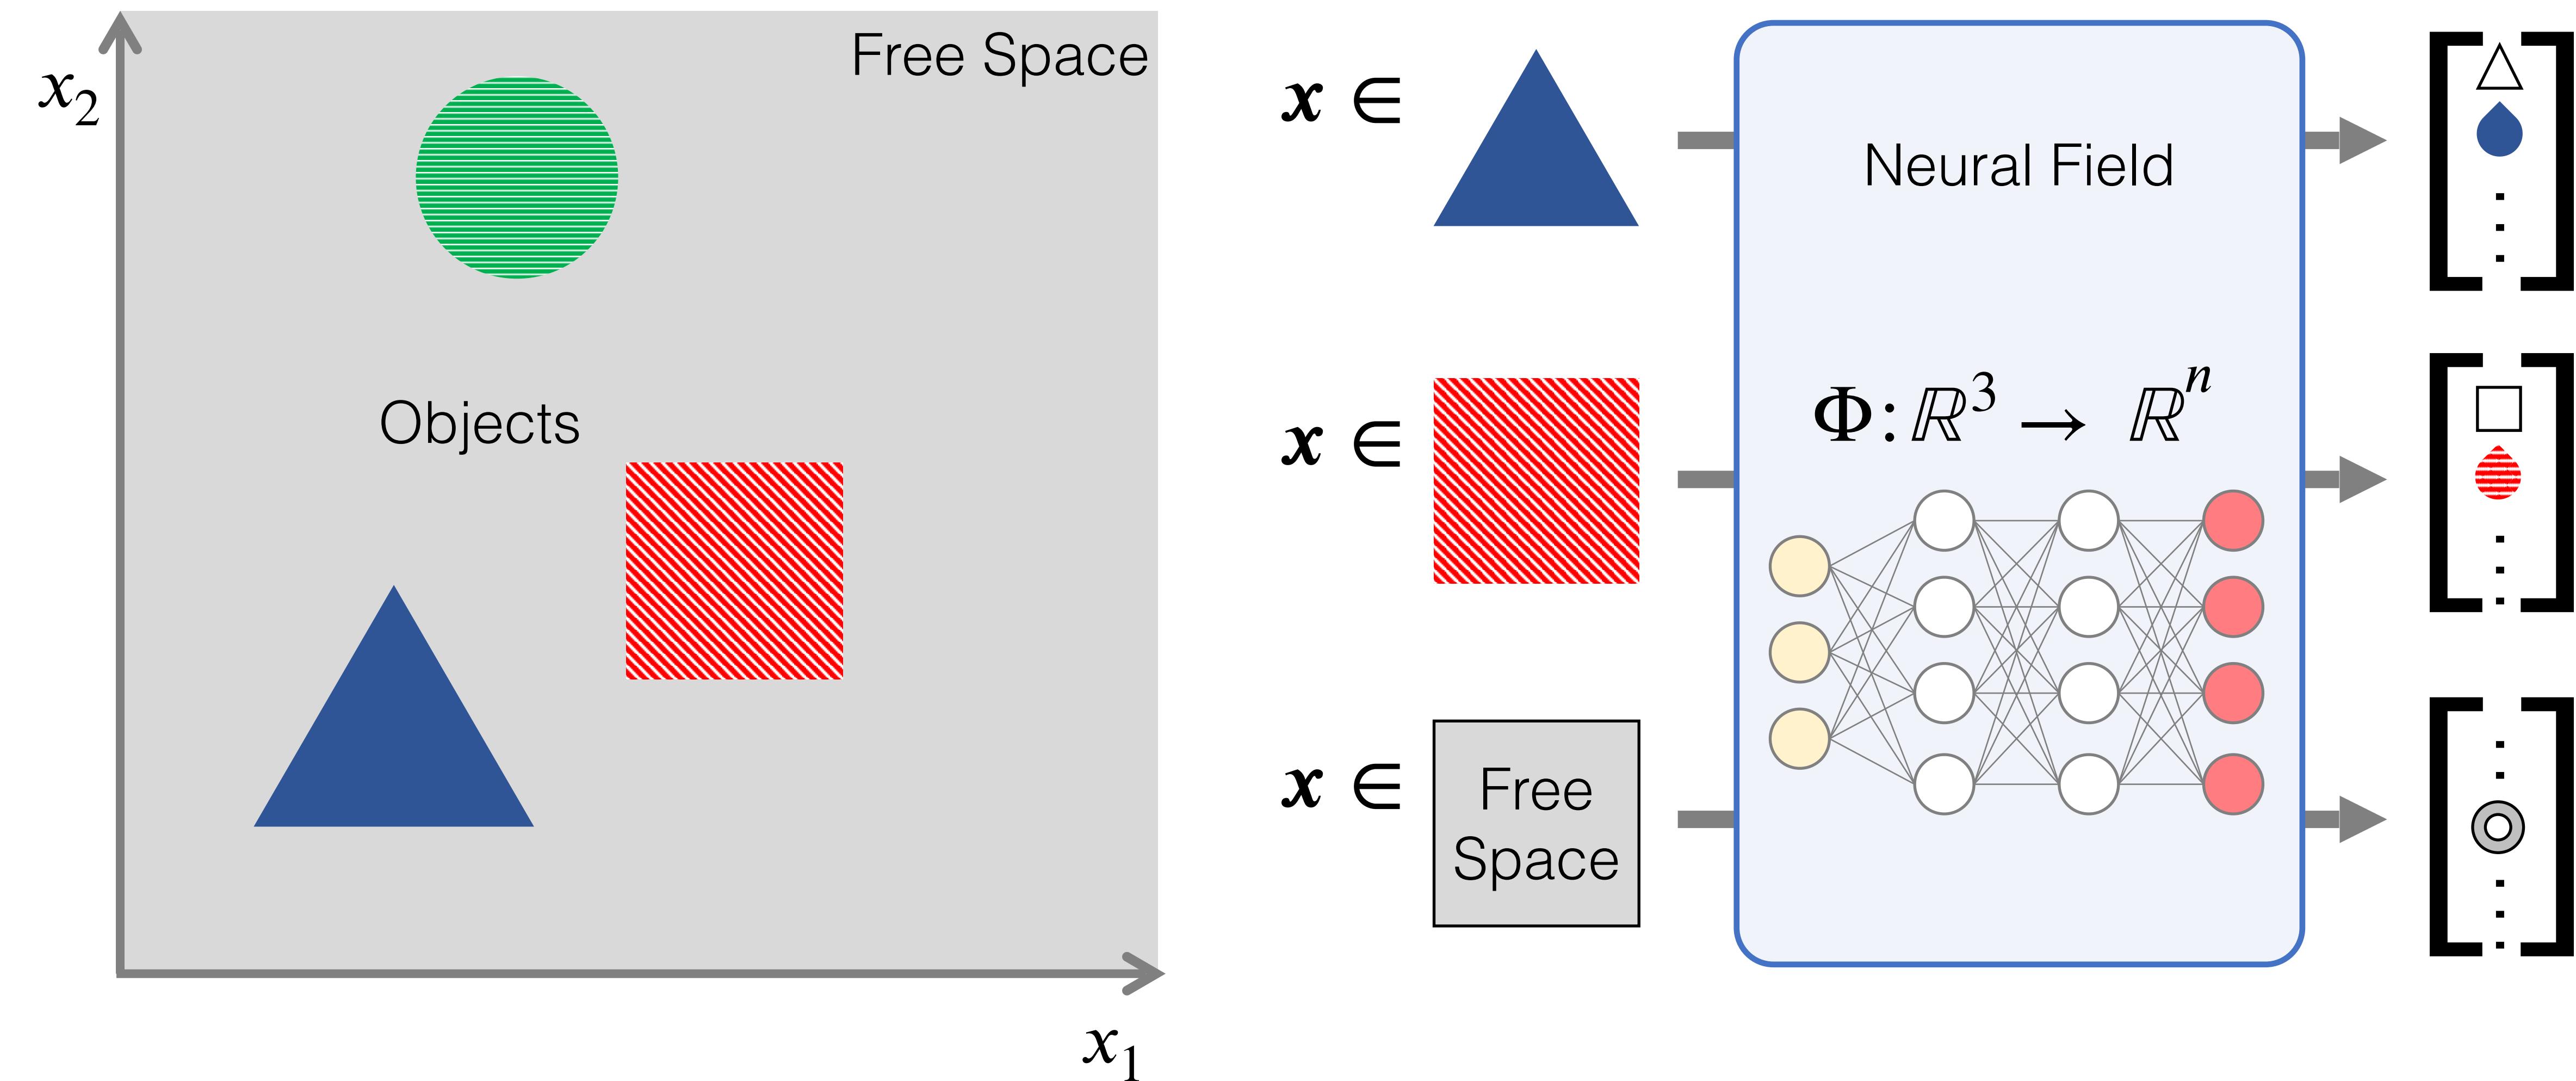
Starting Point: 3D Scene Representation



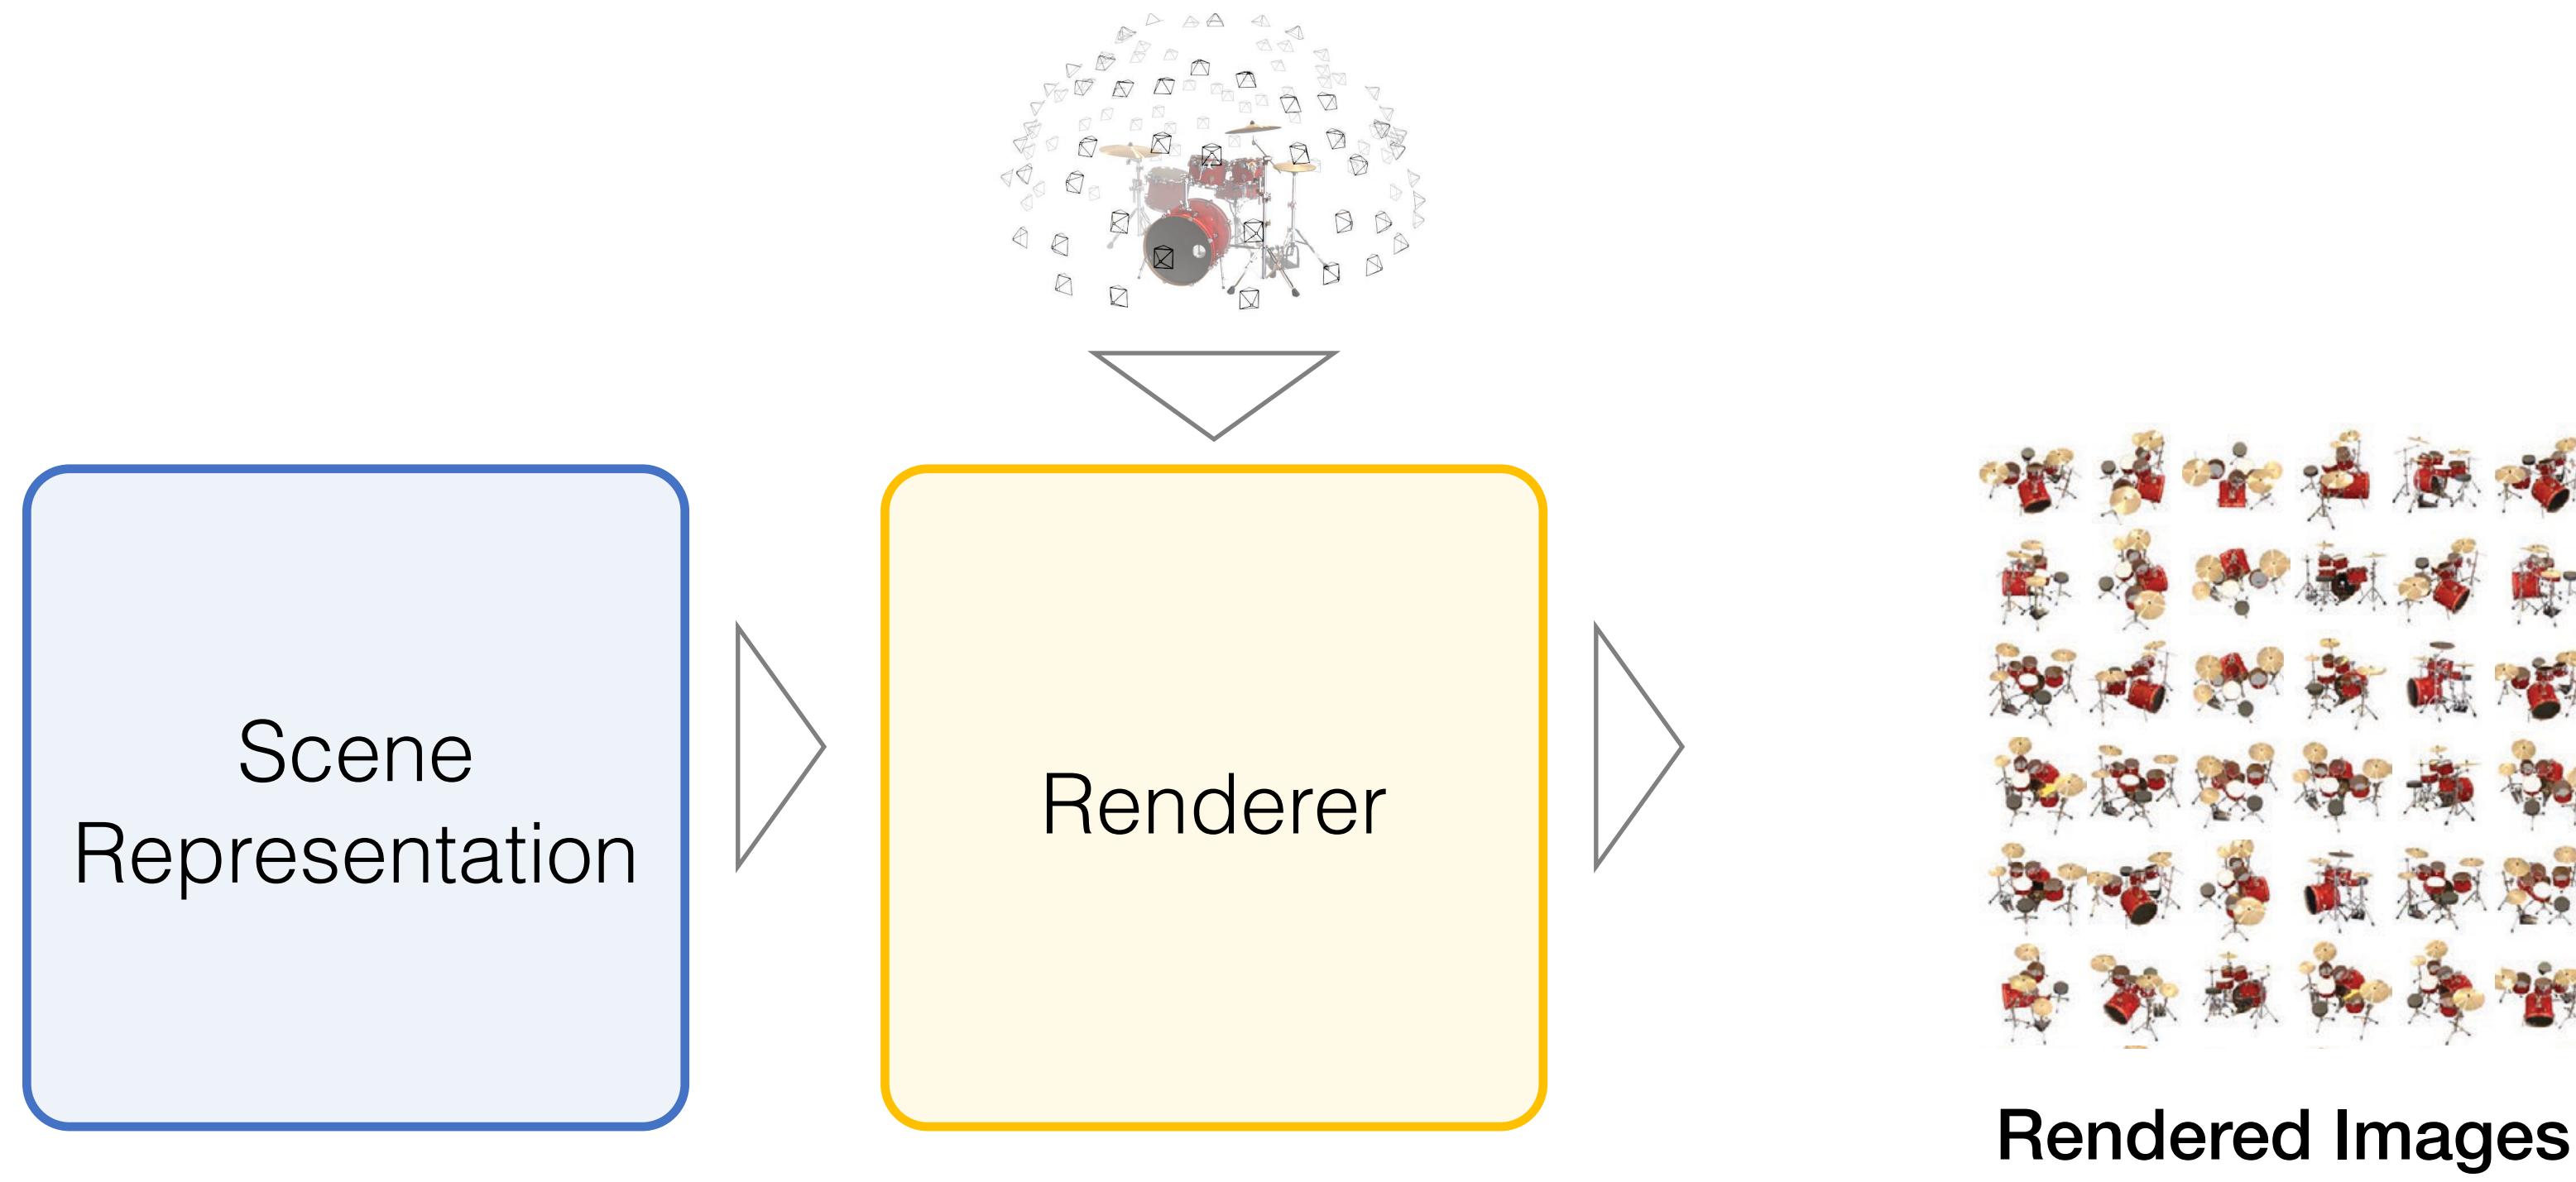
Starting Point: 3D Scene Representation



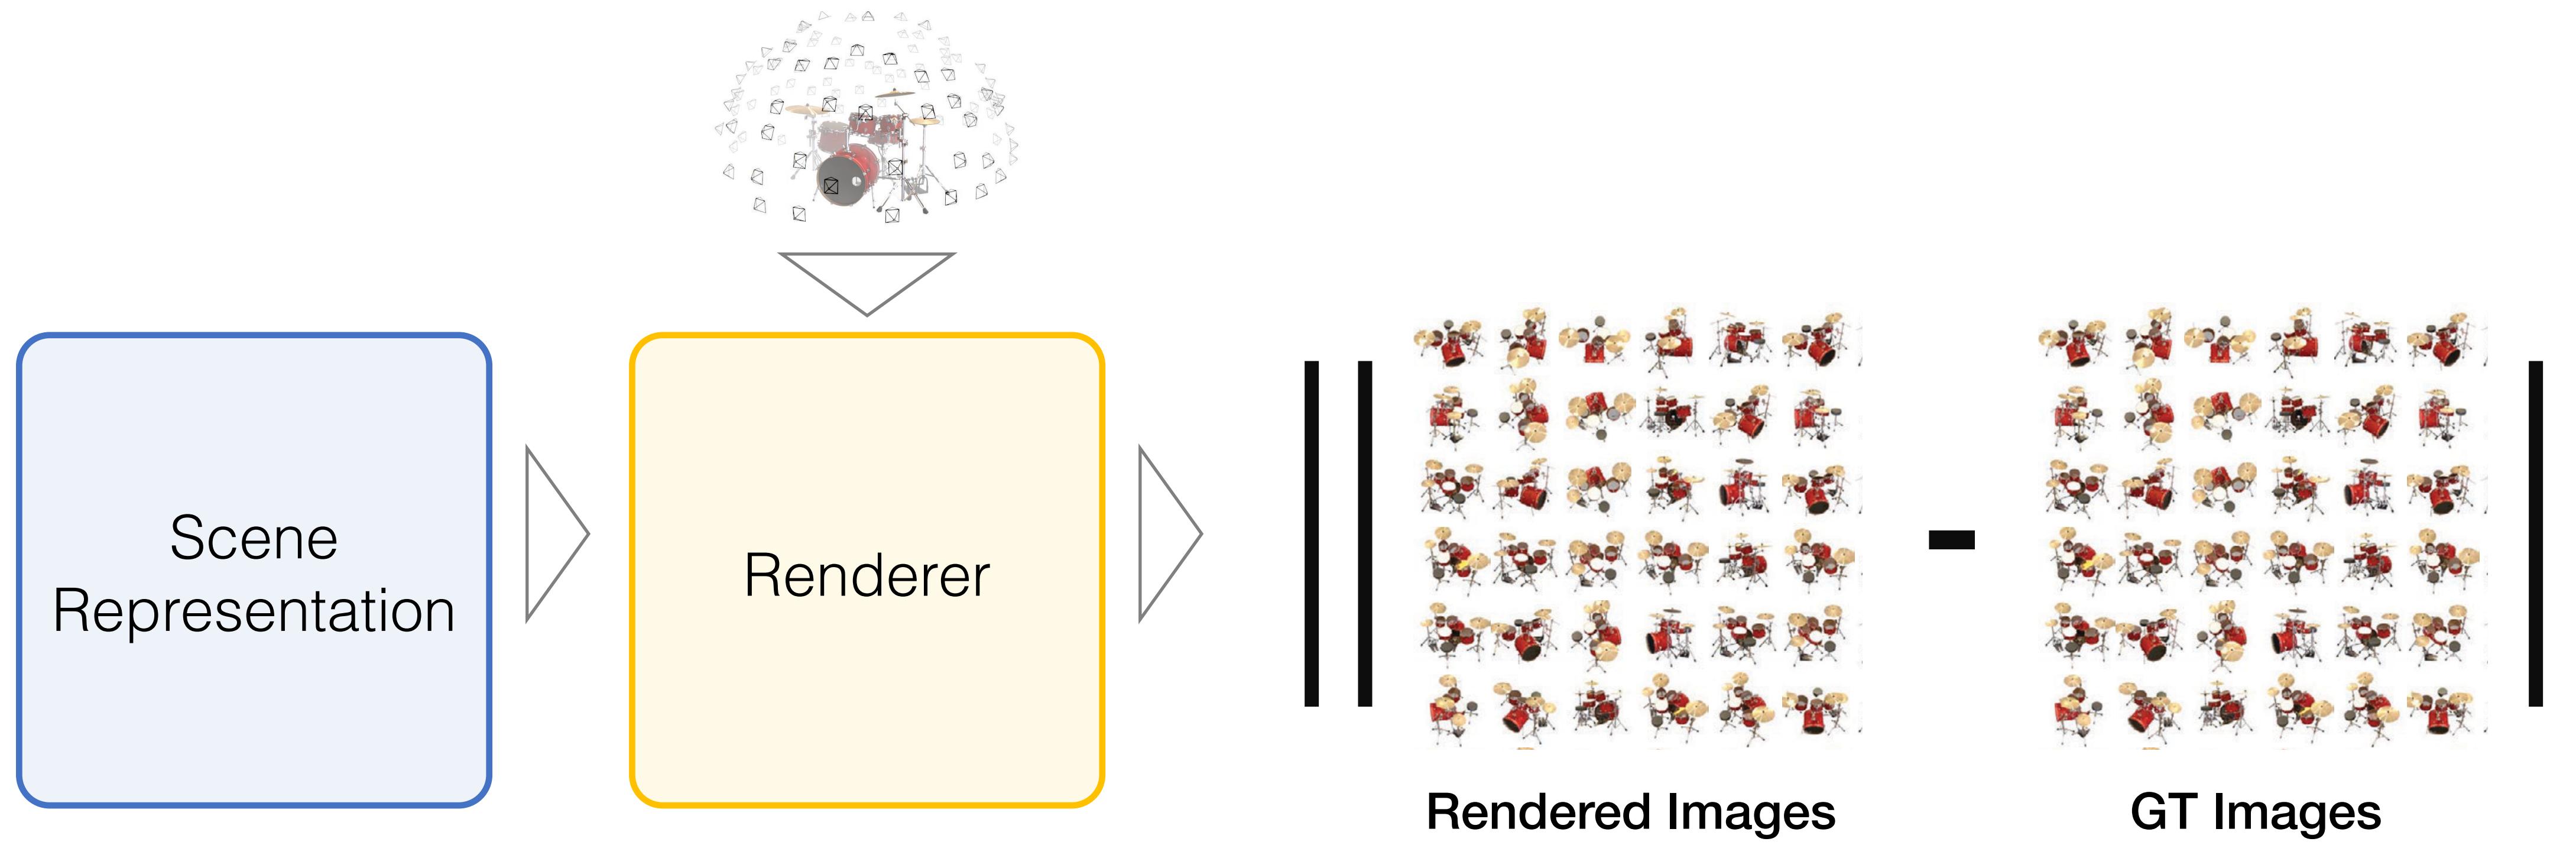
Starting Point: 3D Scene Representation



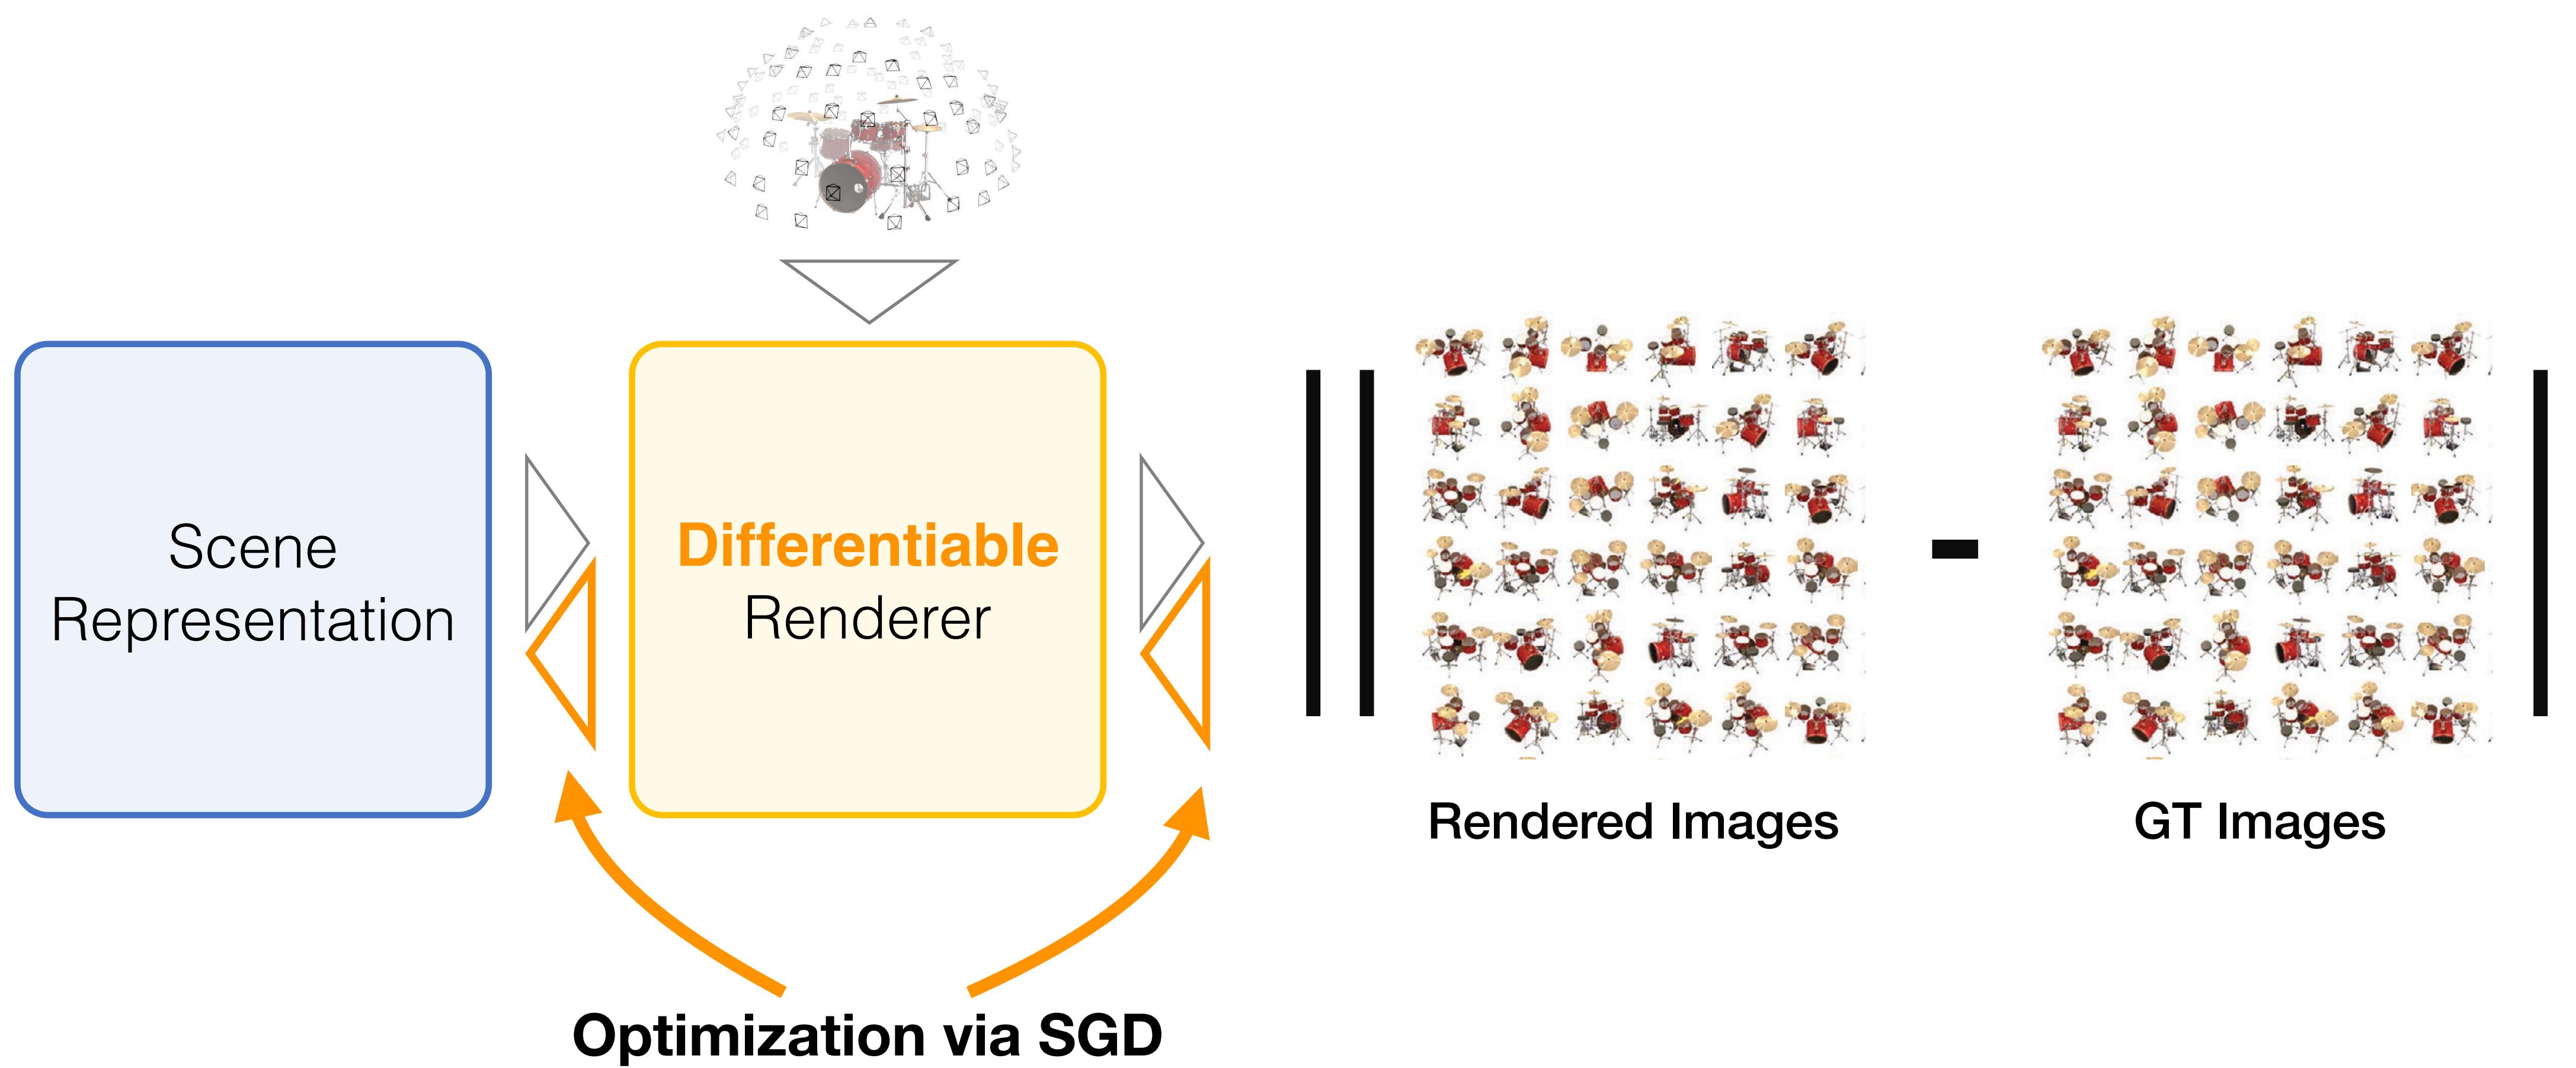
Differentiable Rendering



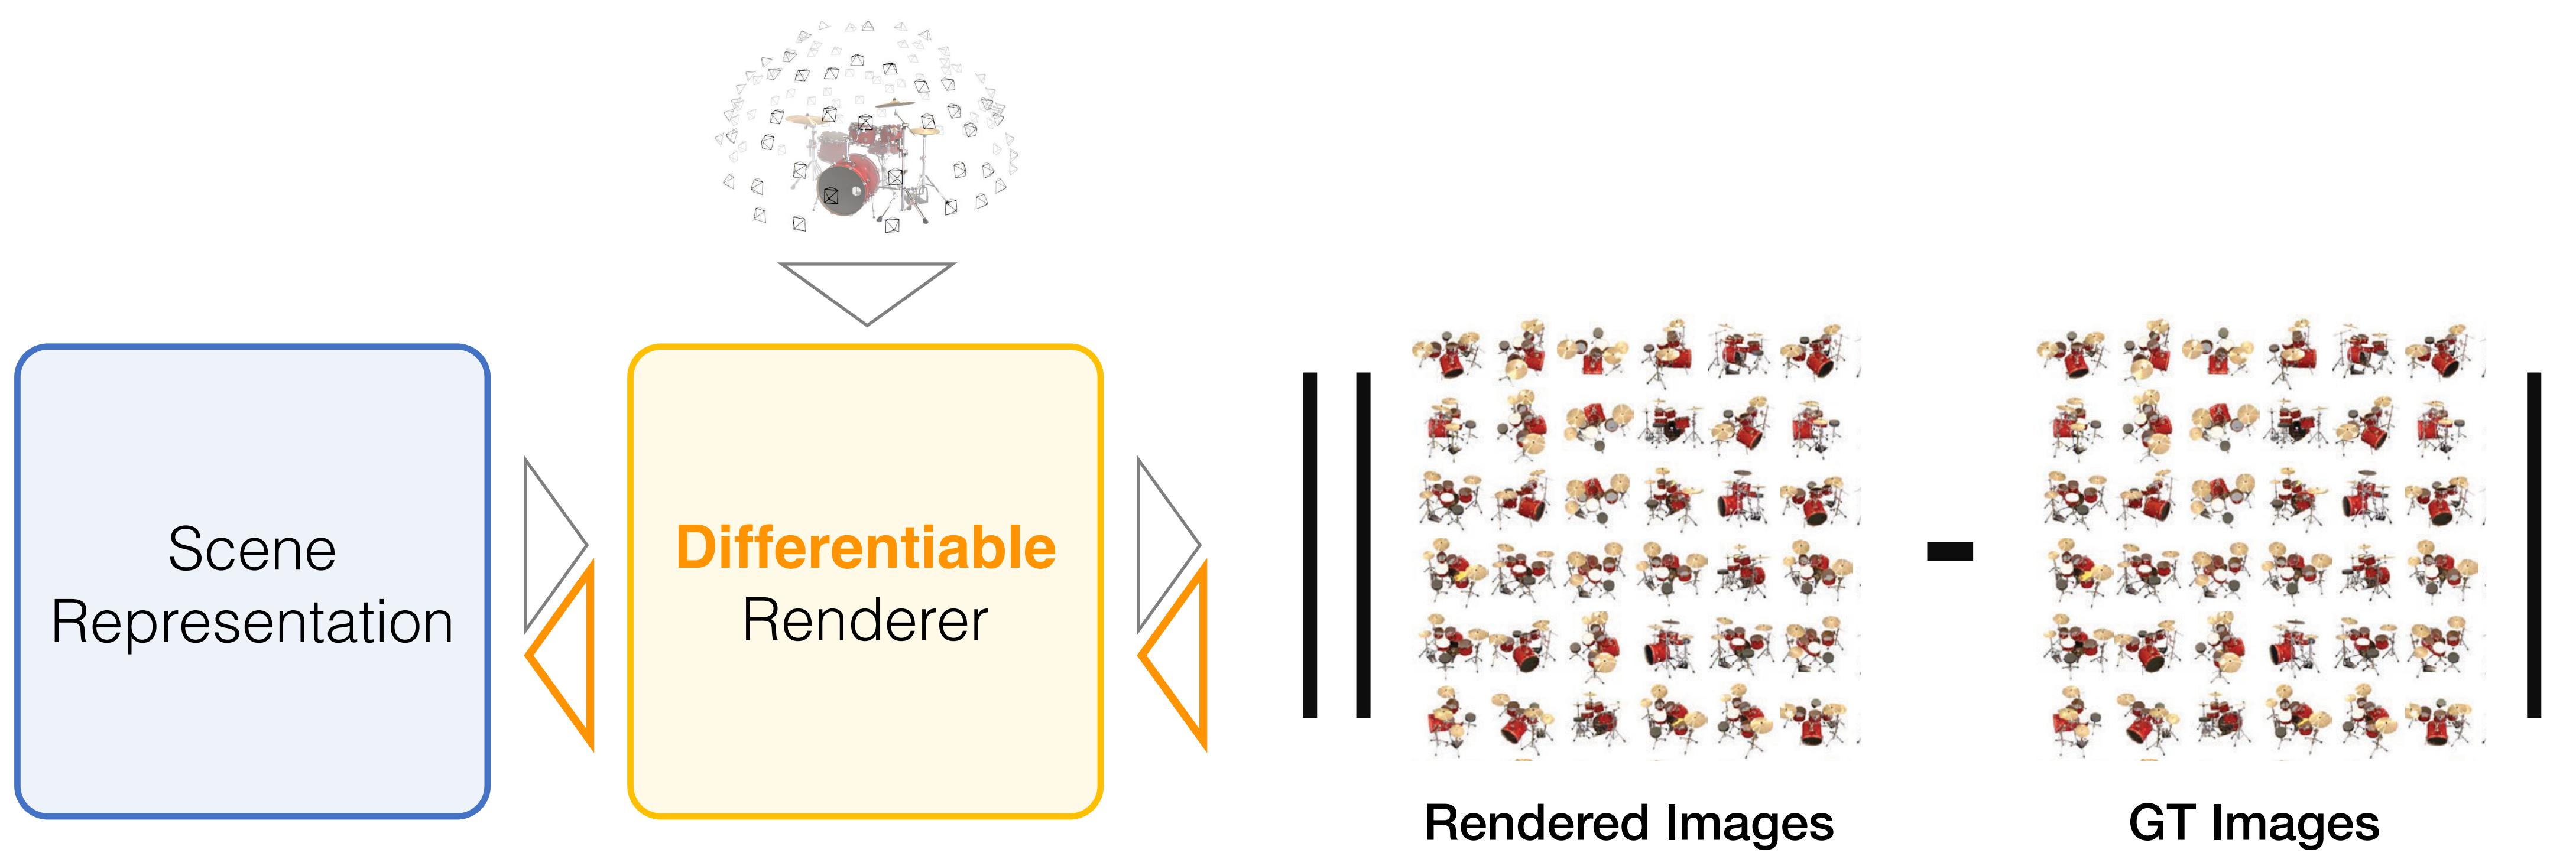
Differentiable Rendering



Differentiable Rendering



Differentiable Rendering



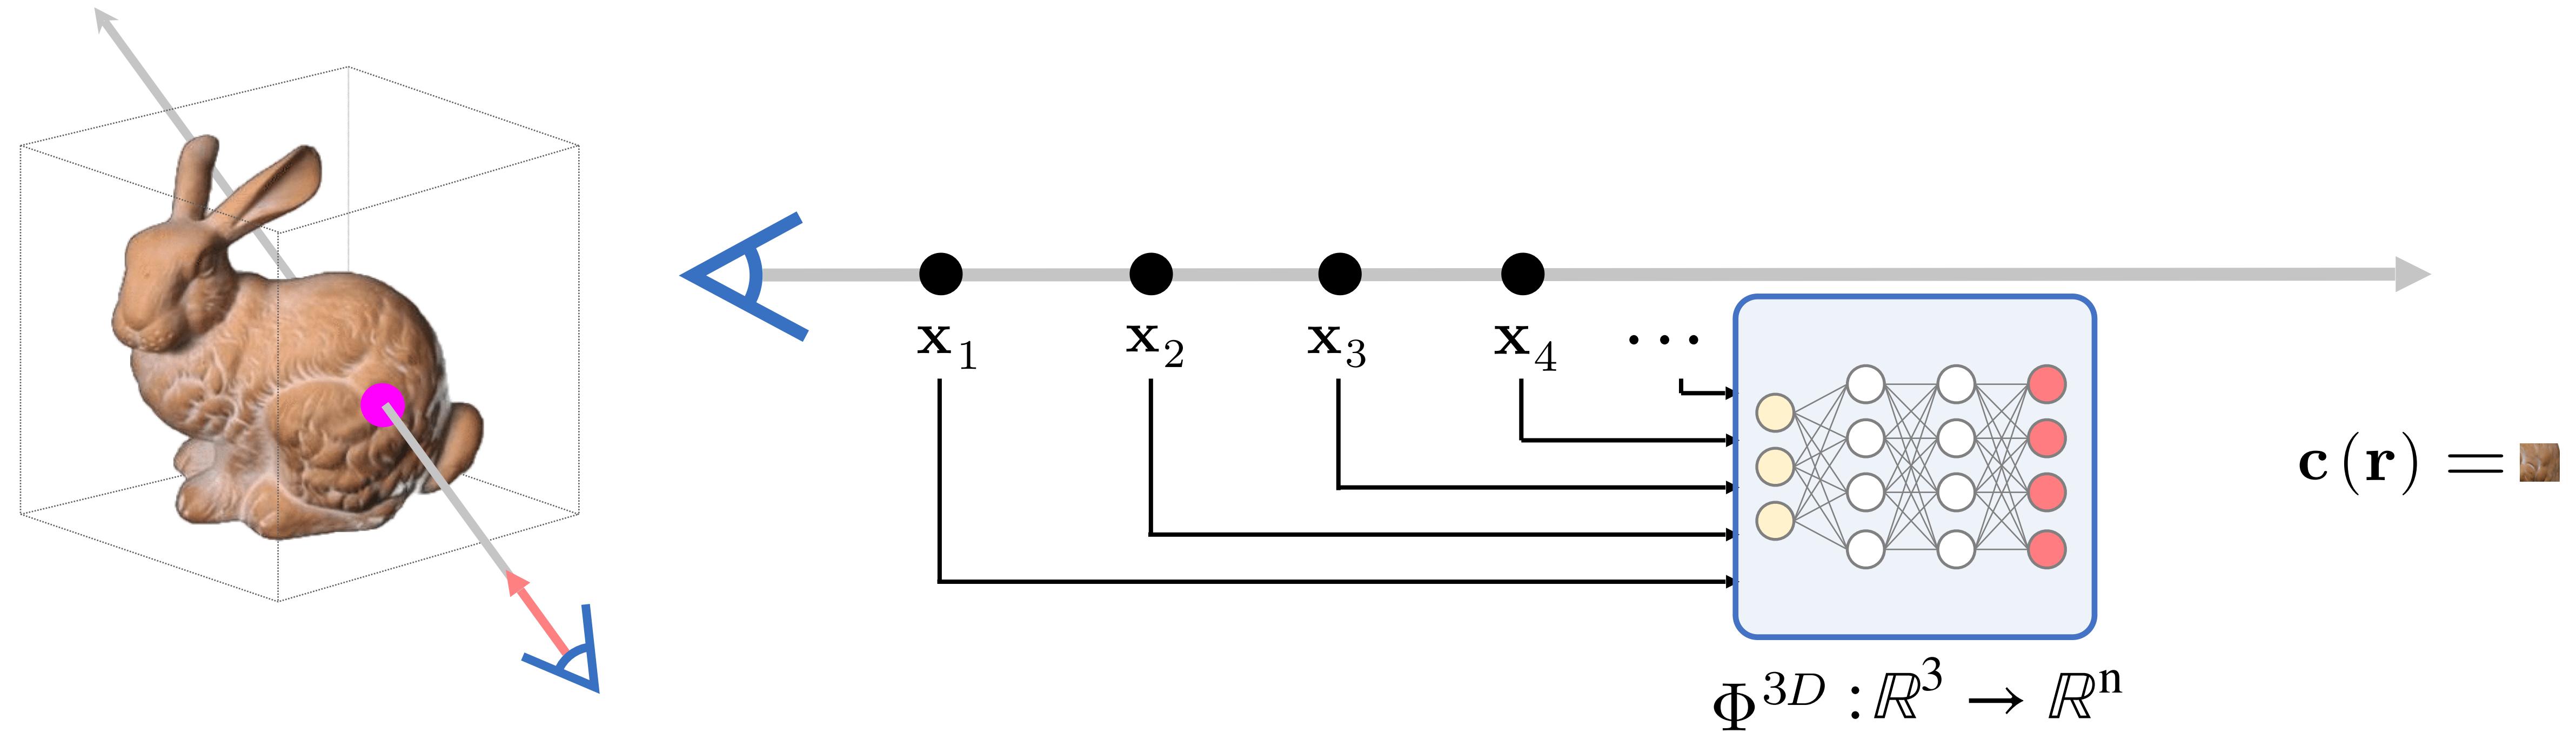
Given an *observable* variable (pixel colors), we will build a differentiable forward model that we then use to estimate *unobserved (latent) variables* (geometry, appearance)!

Questions?

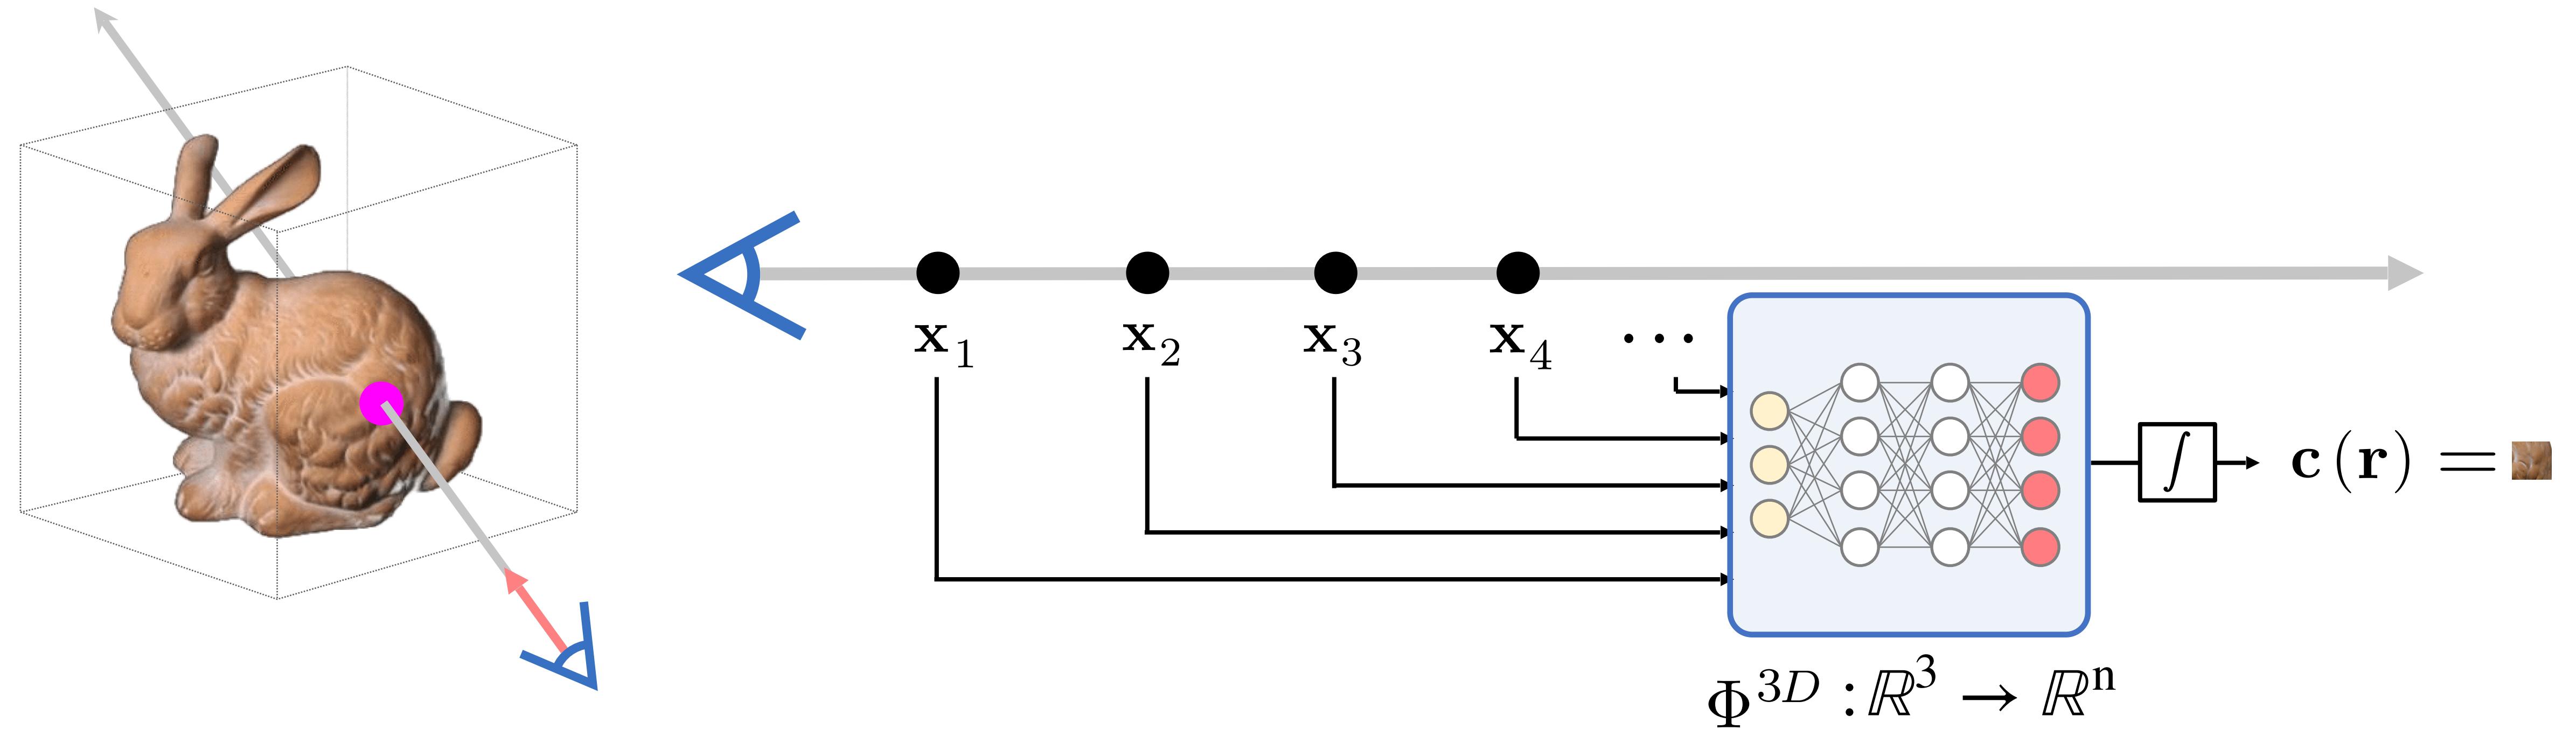
General structure of Neural Renderers for 3D Fields



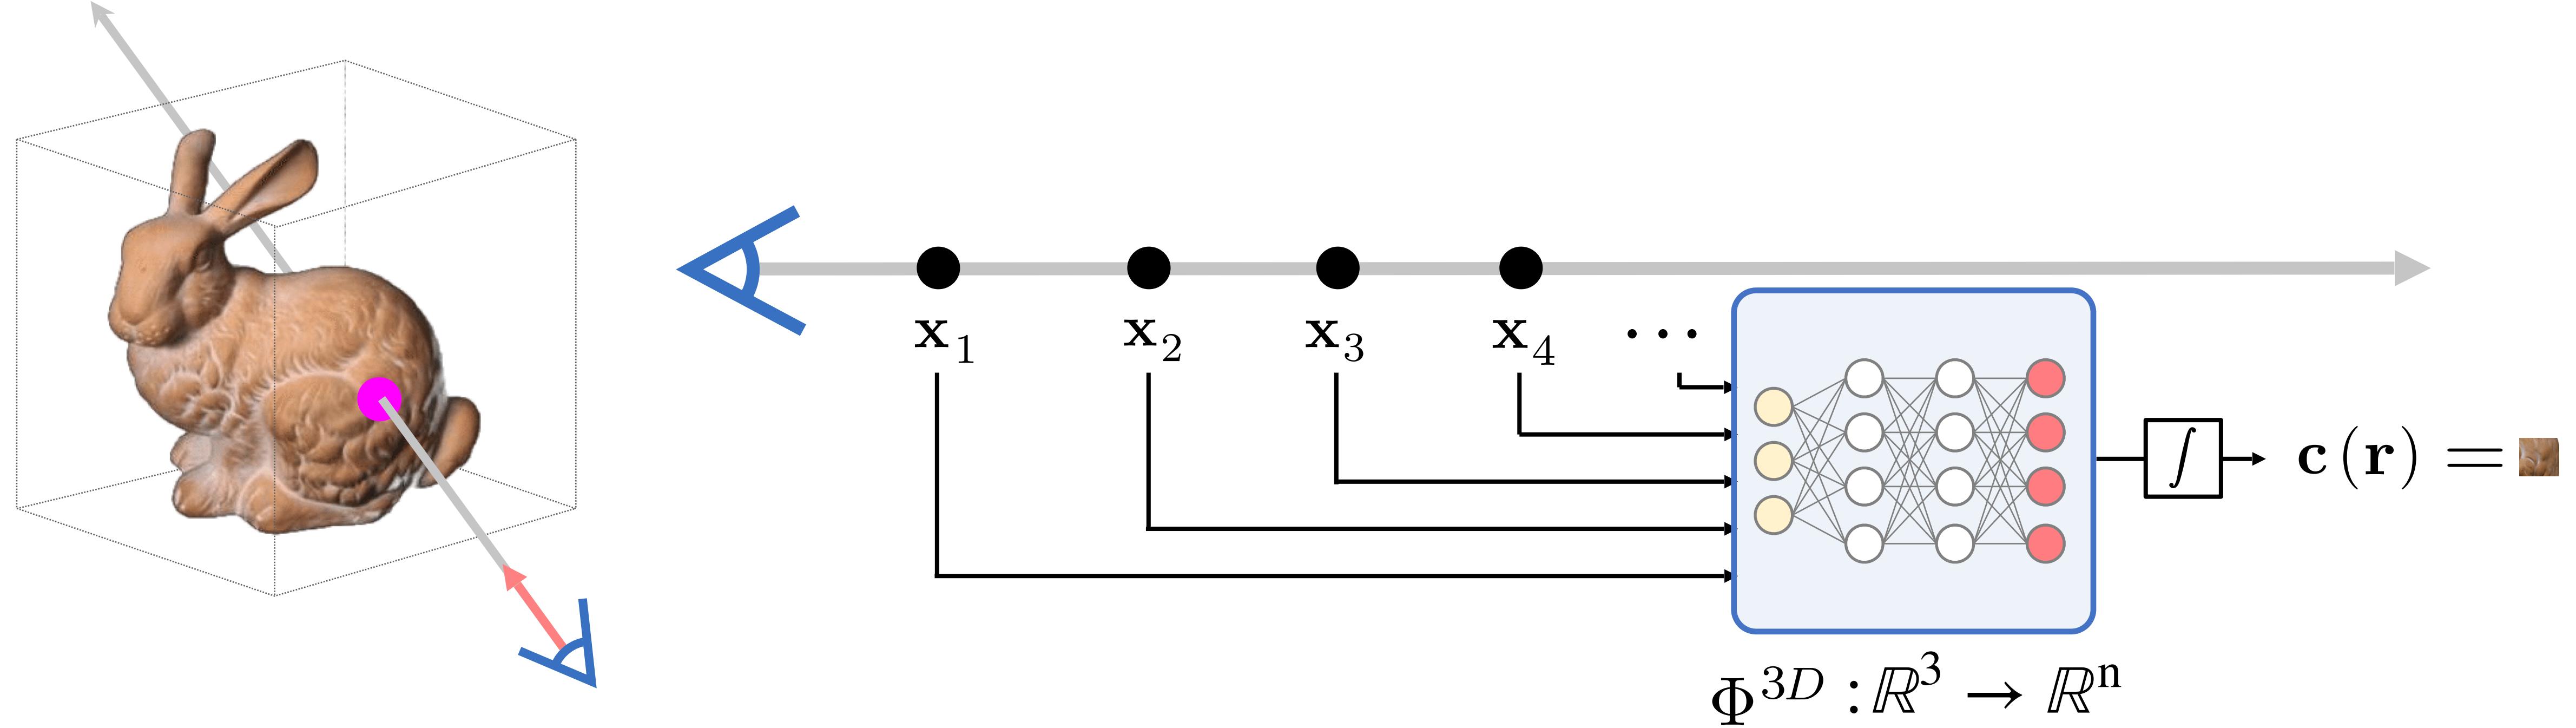
General structure of Neural Renderers for 3D Fields



General structure of Neural Renderers for 3D Fields



General structure of Neural Renderers for 3D Fields



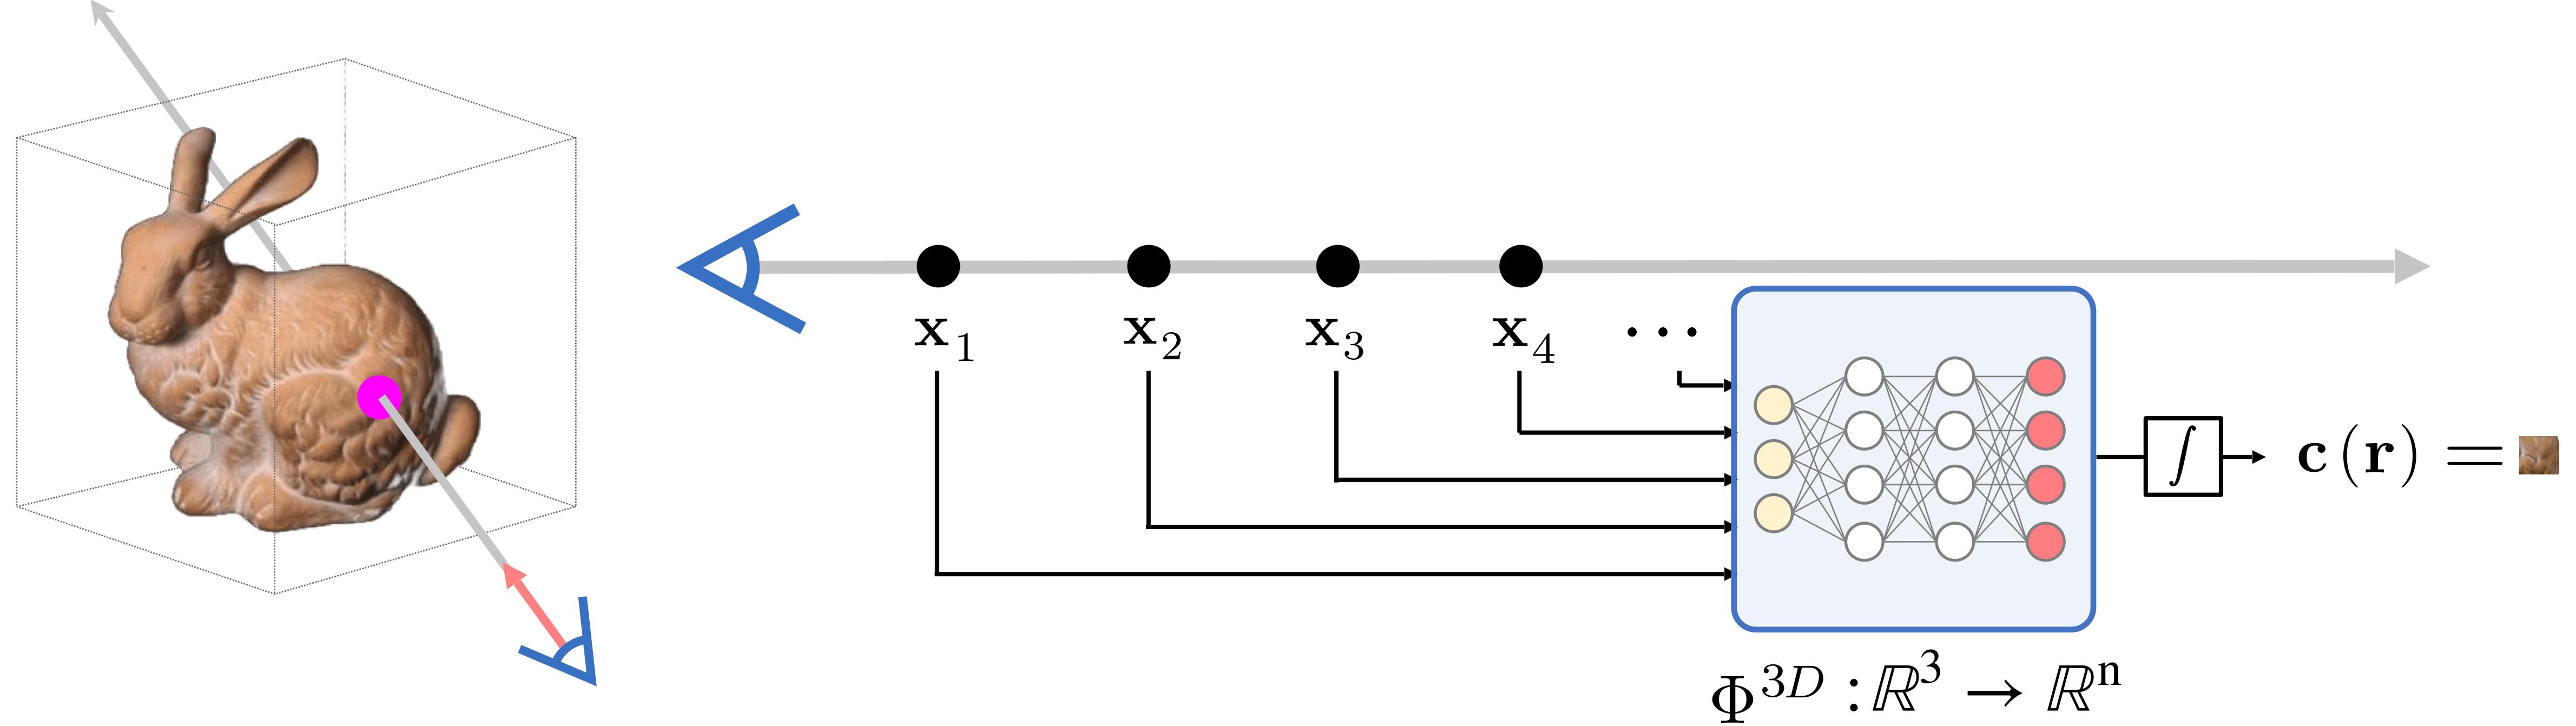
Sphere-Tracing
[JC Hart, 1996]

Volumetric

Hybrid implicit-volumetric

Learned aggregation

General structure of Neural Renderers for 3D Fields



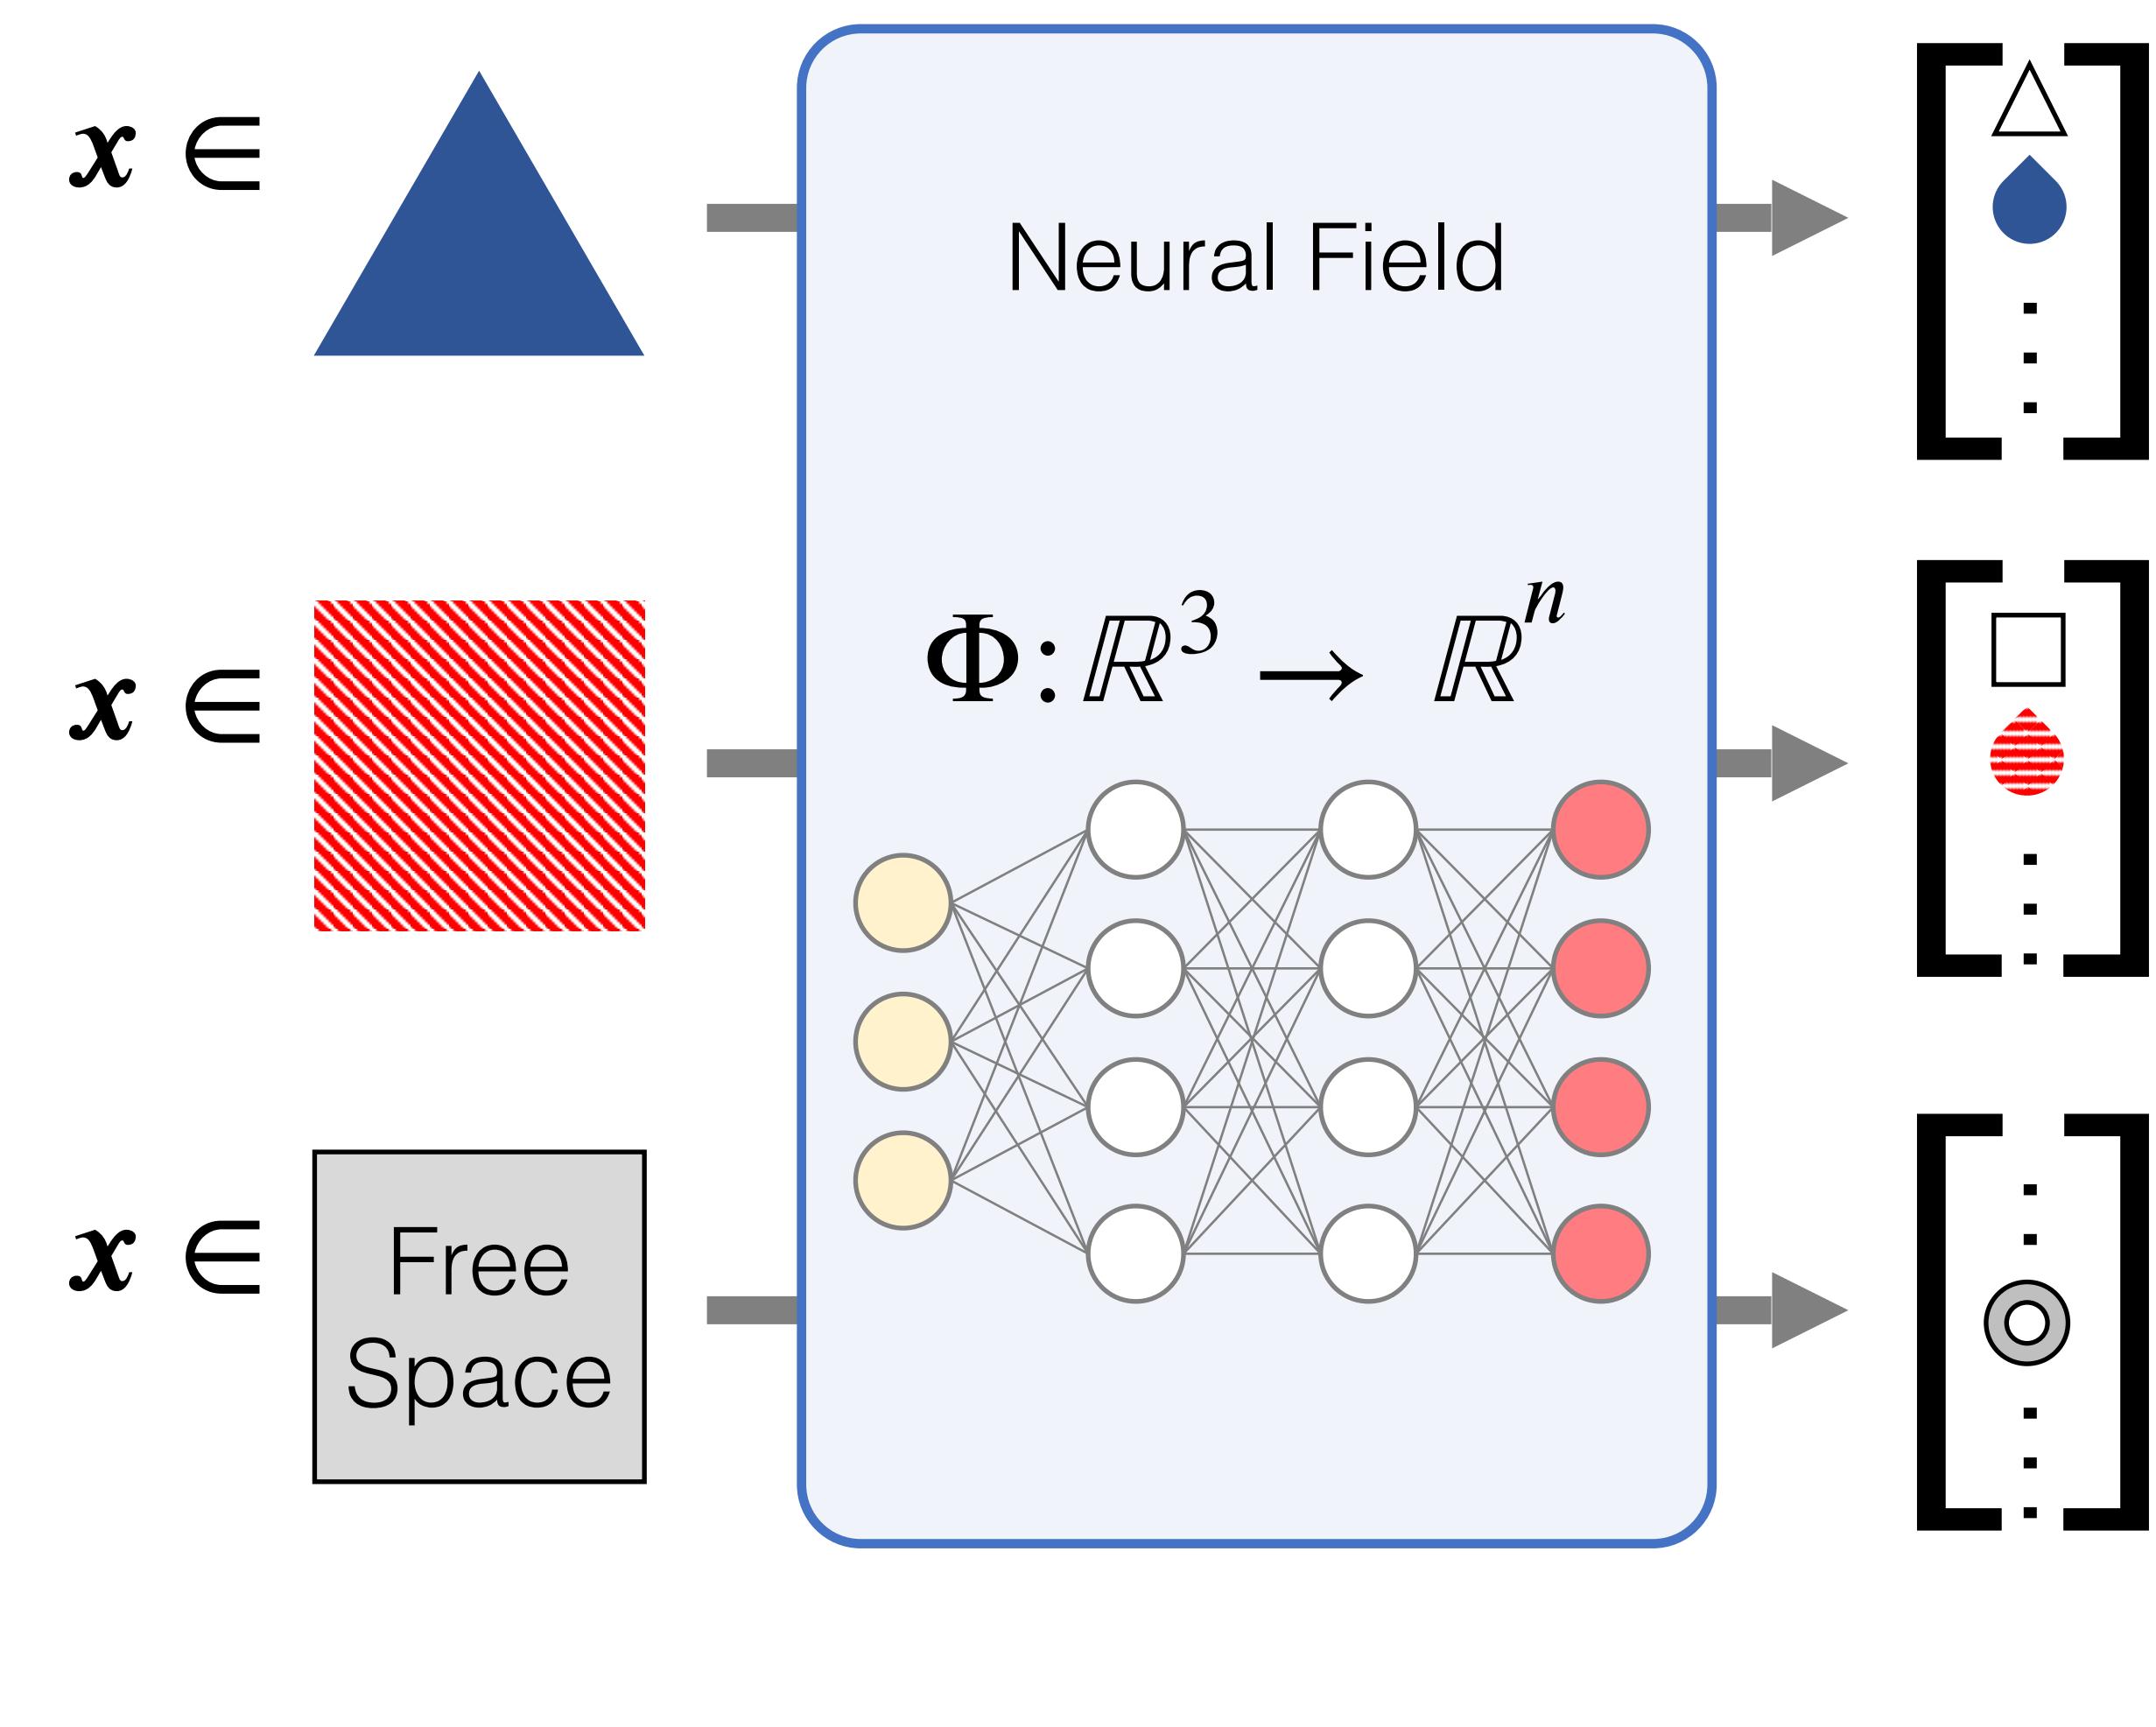
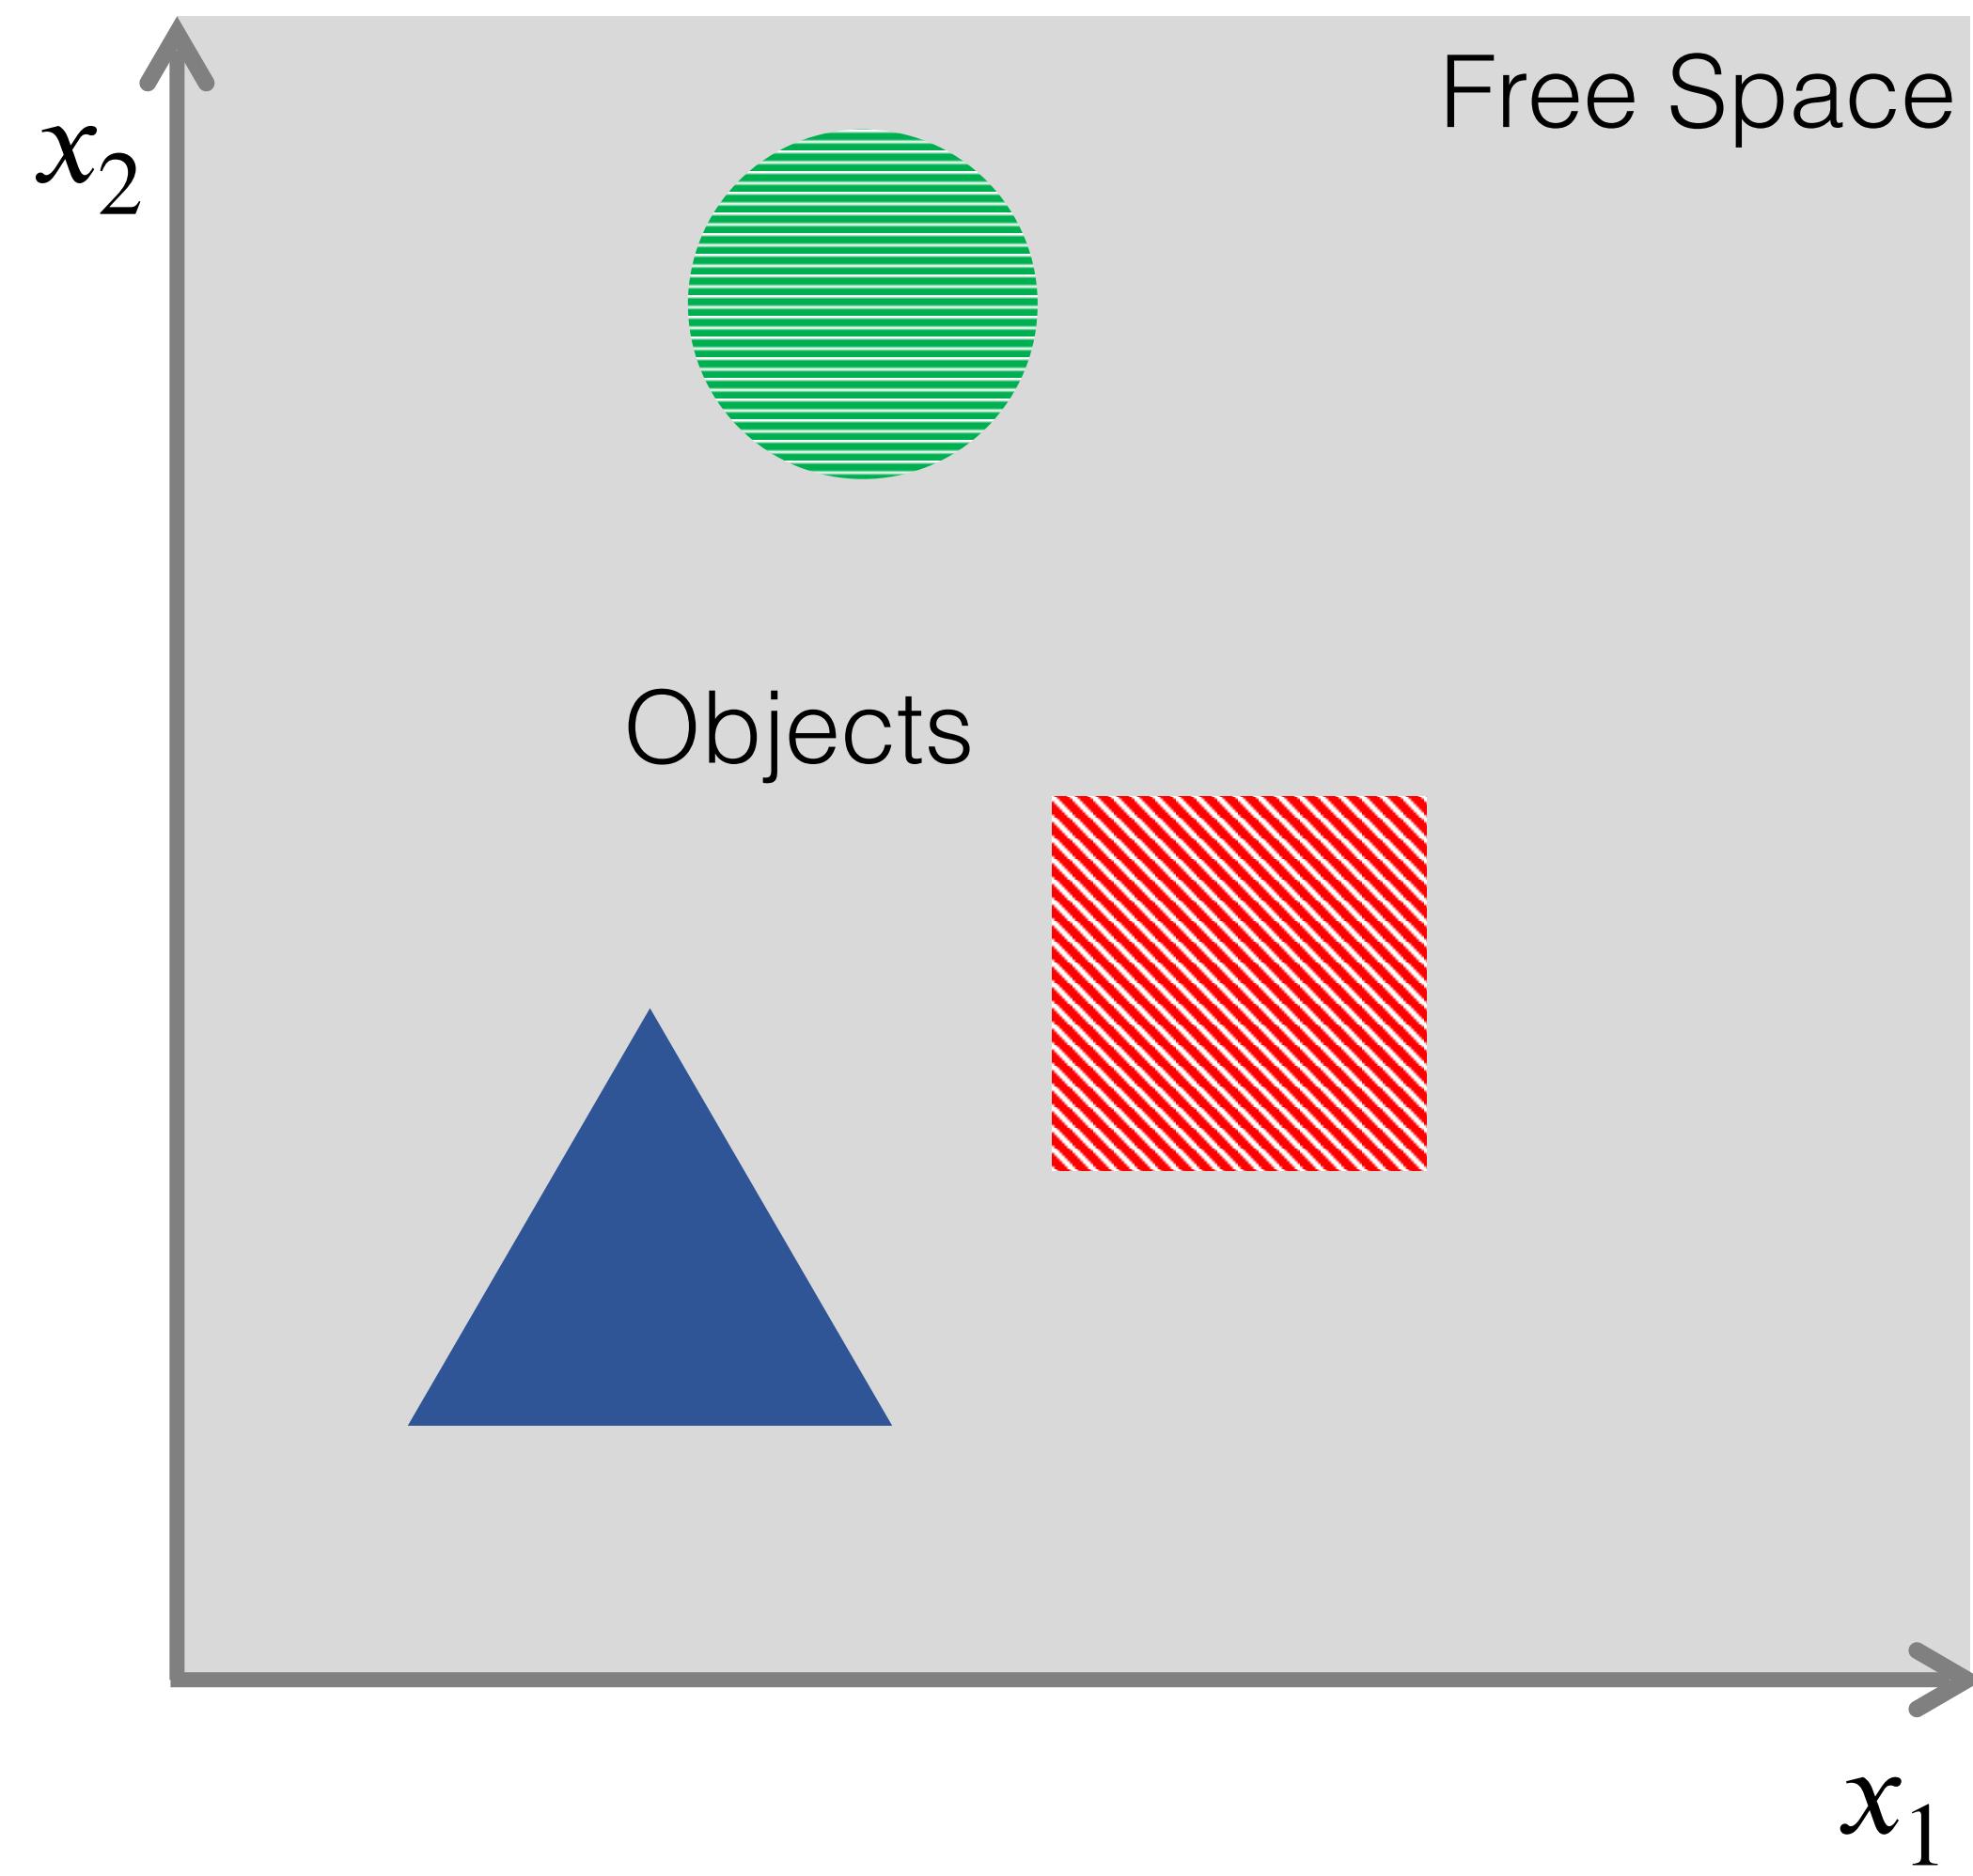
Sphere-Tracing
[JC Hart, 1996]

Volumetric

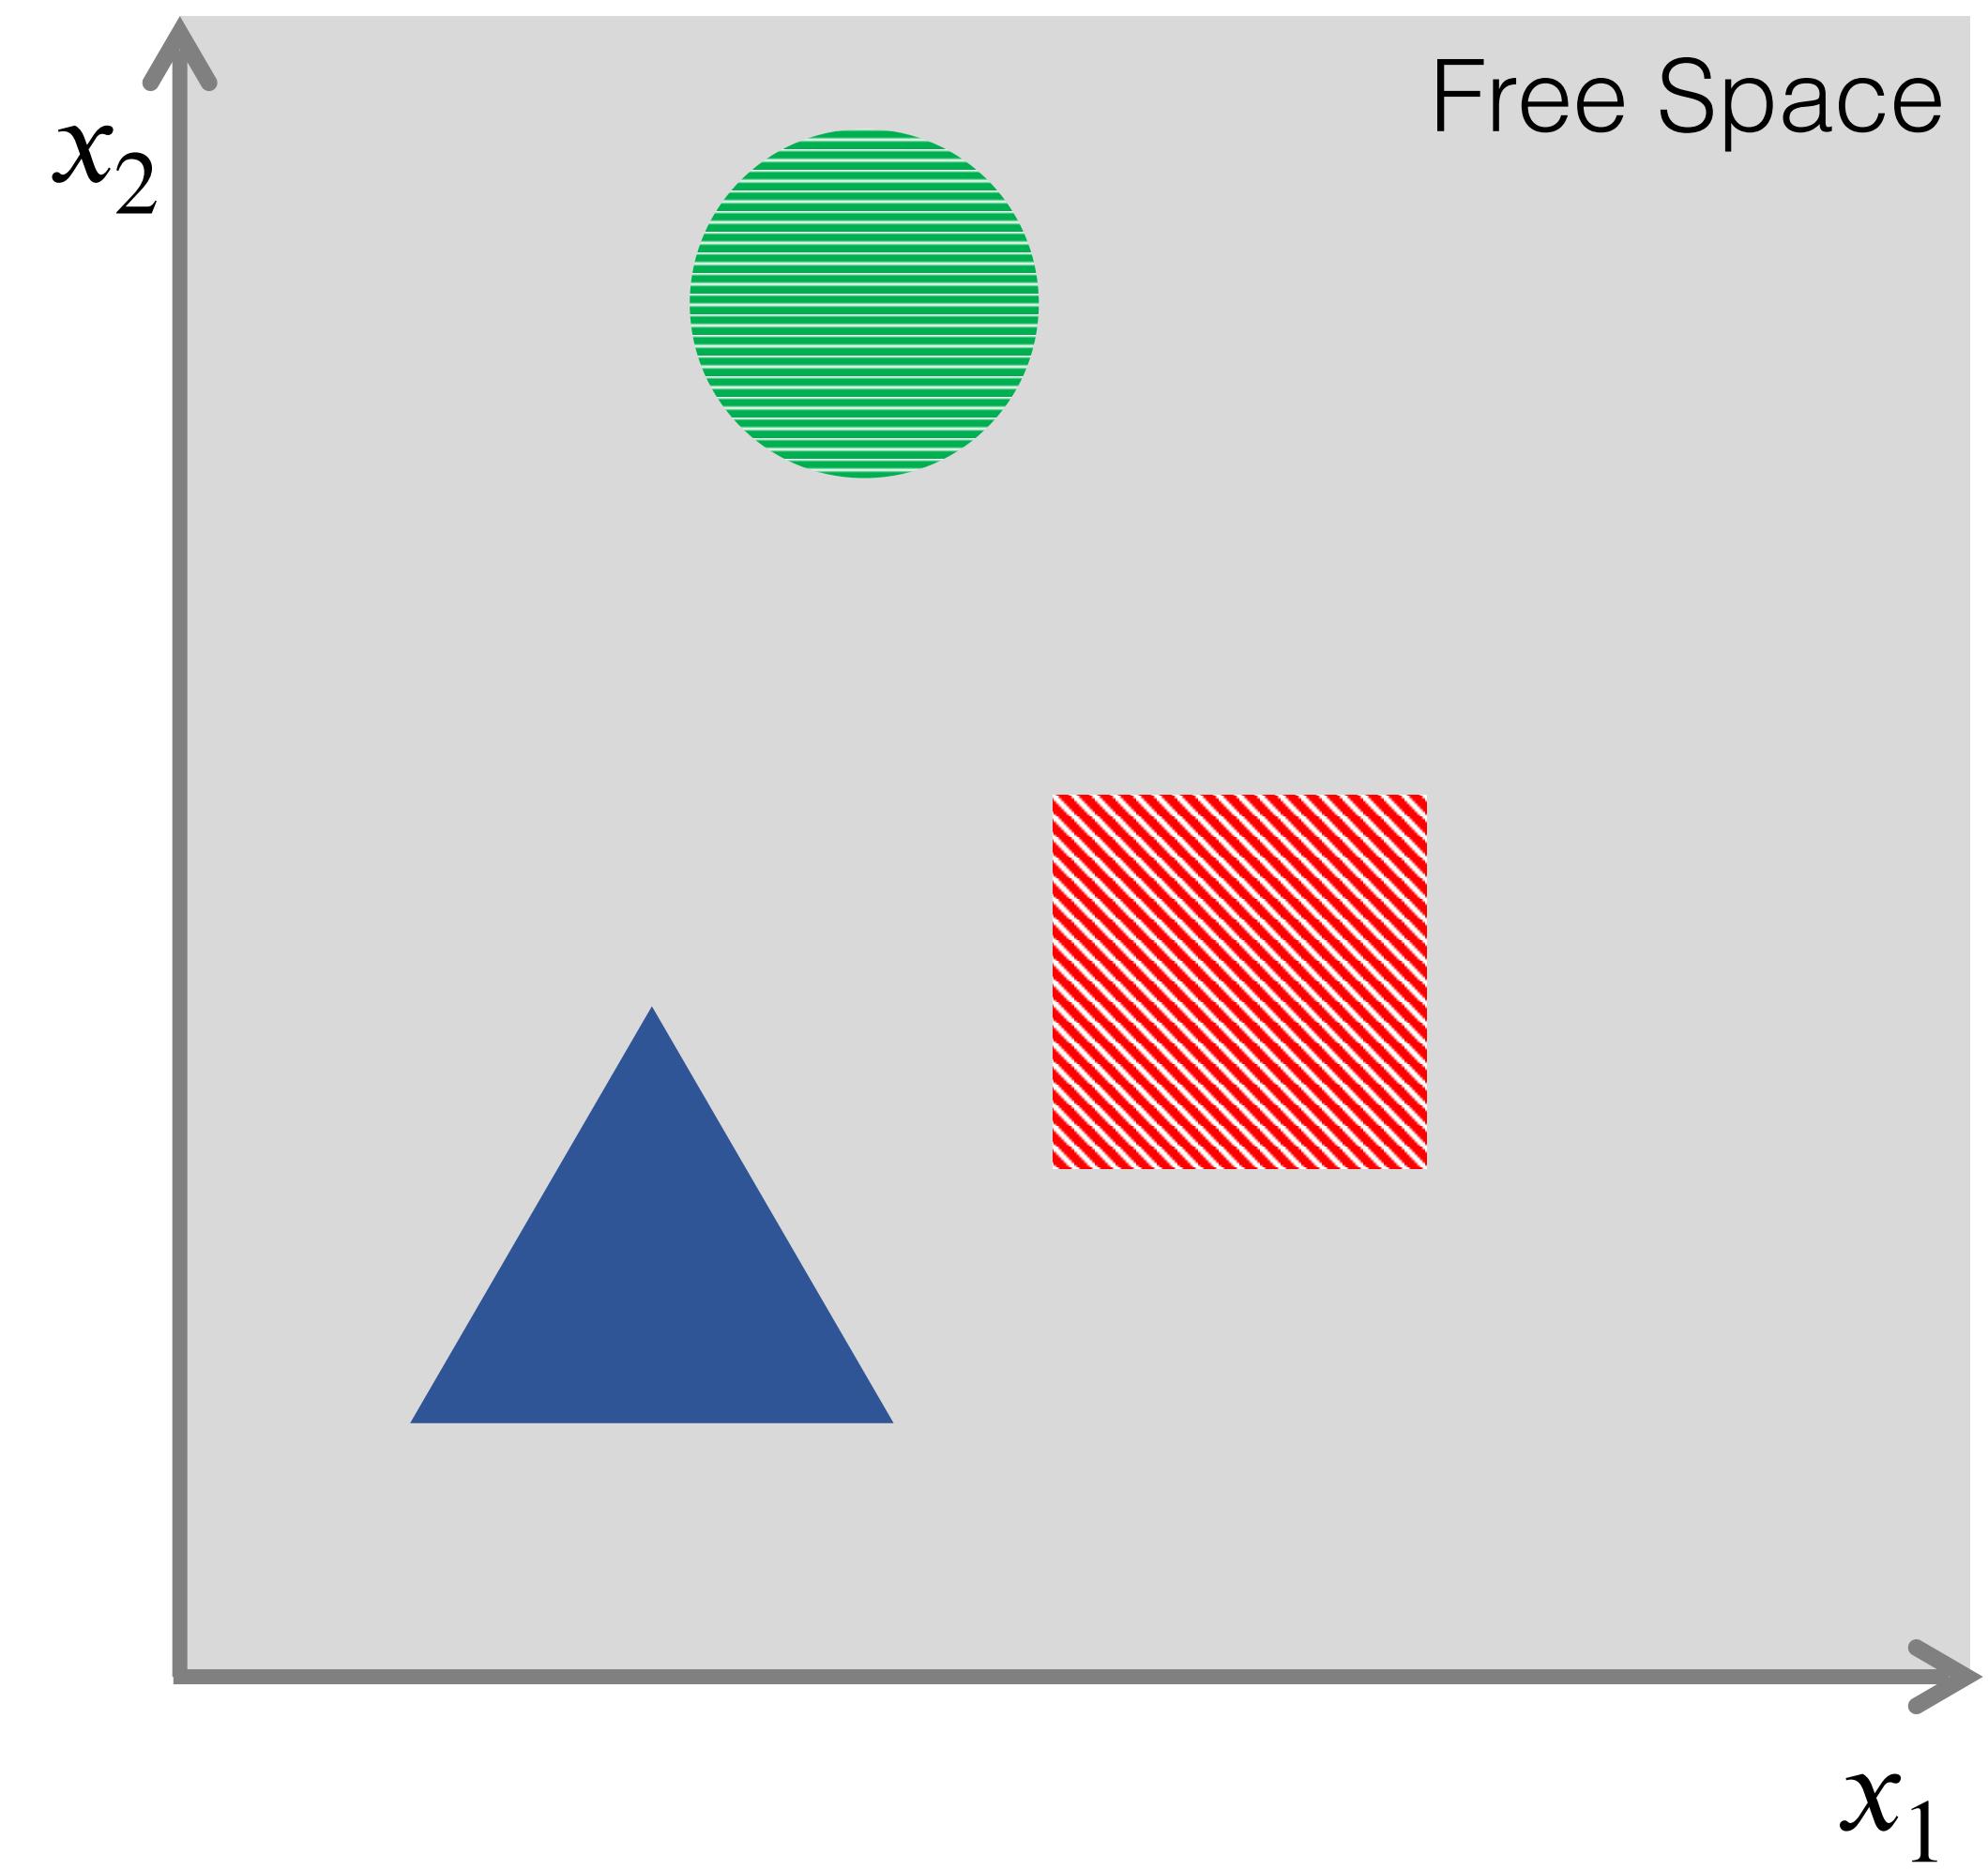
Hybrid implicit-volumetric

Learned aggregation

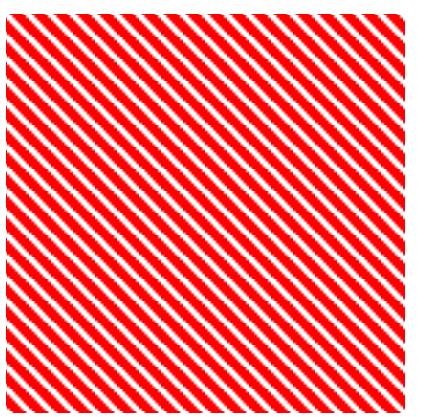
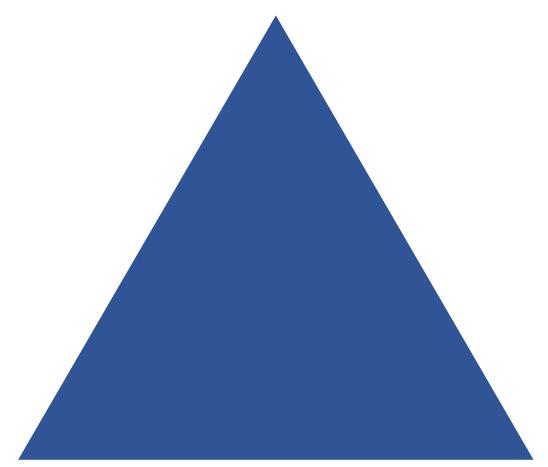
Scene Representation Networks (Sitzmann et al. 2019)



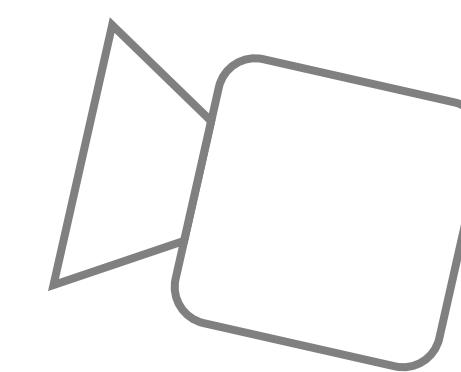
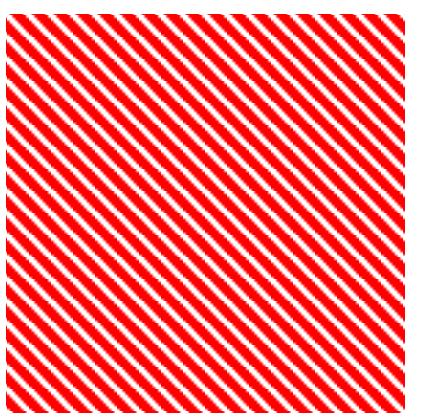
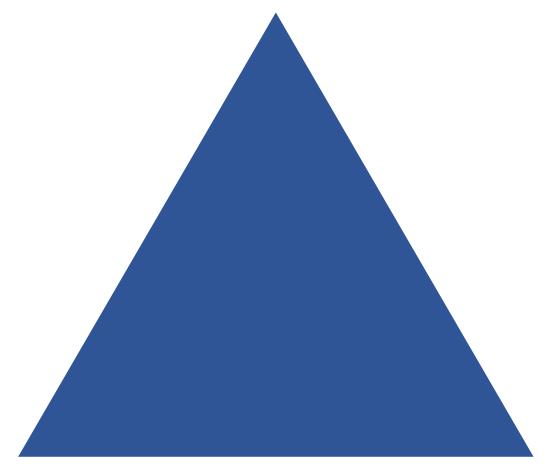
Sphere Tracing



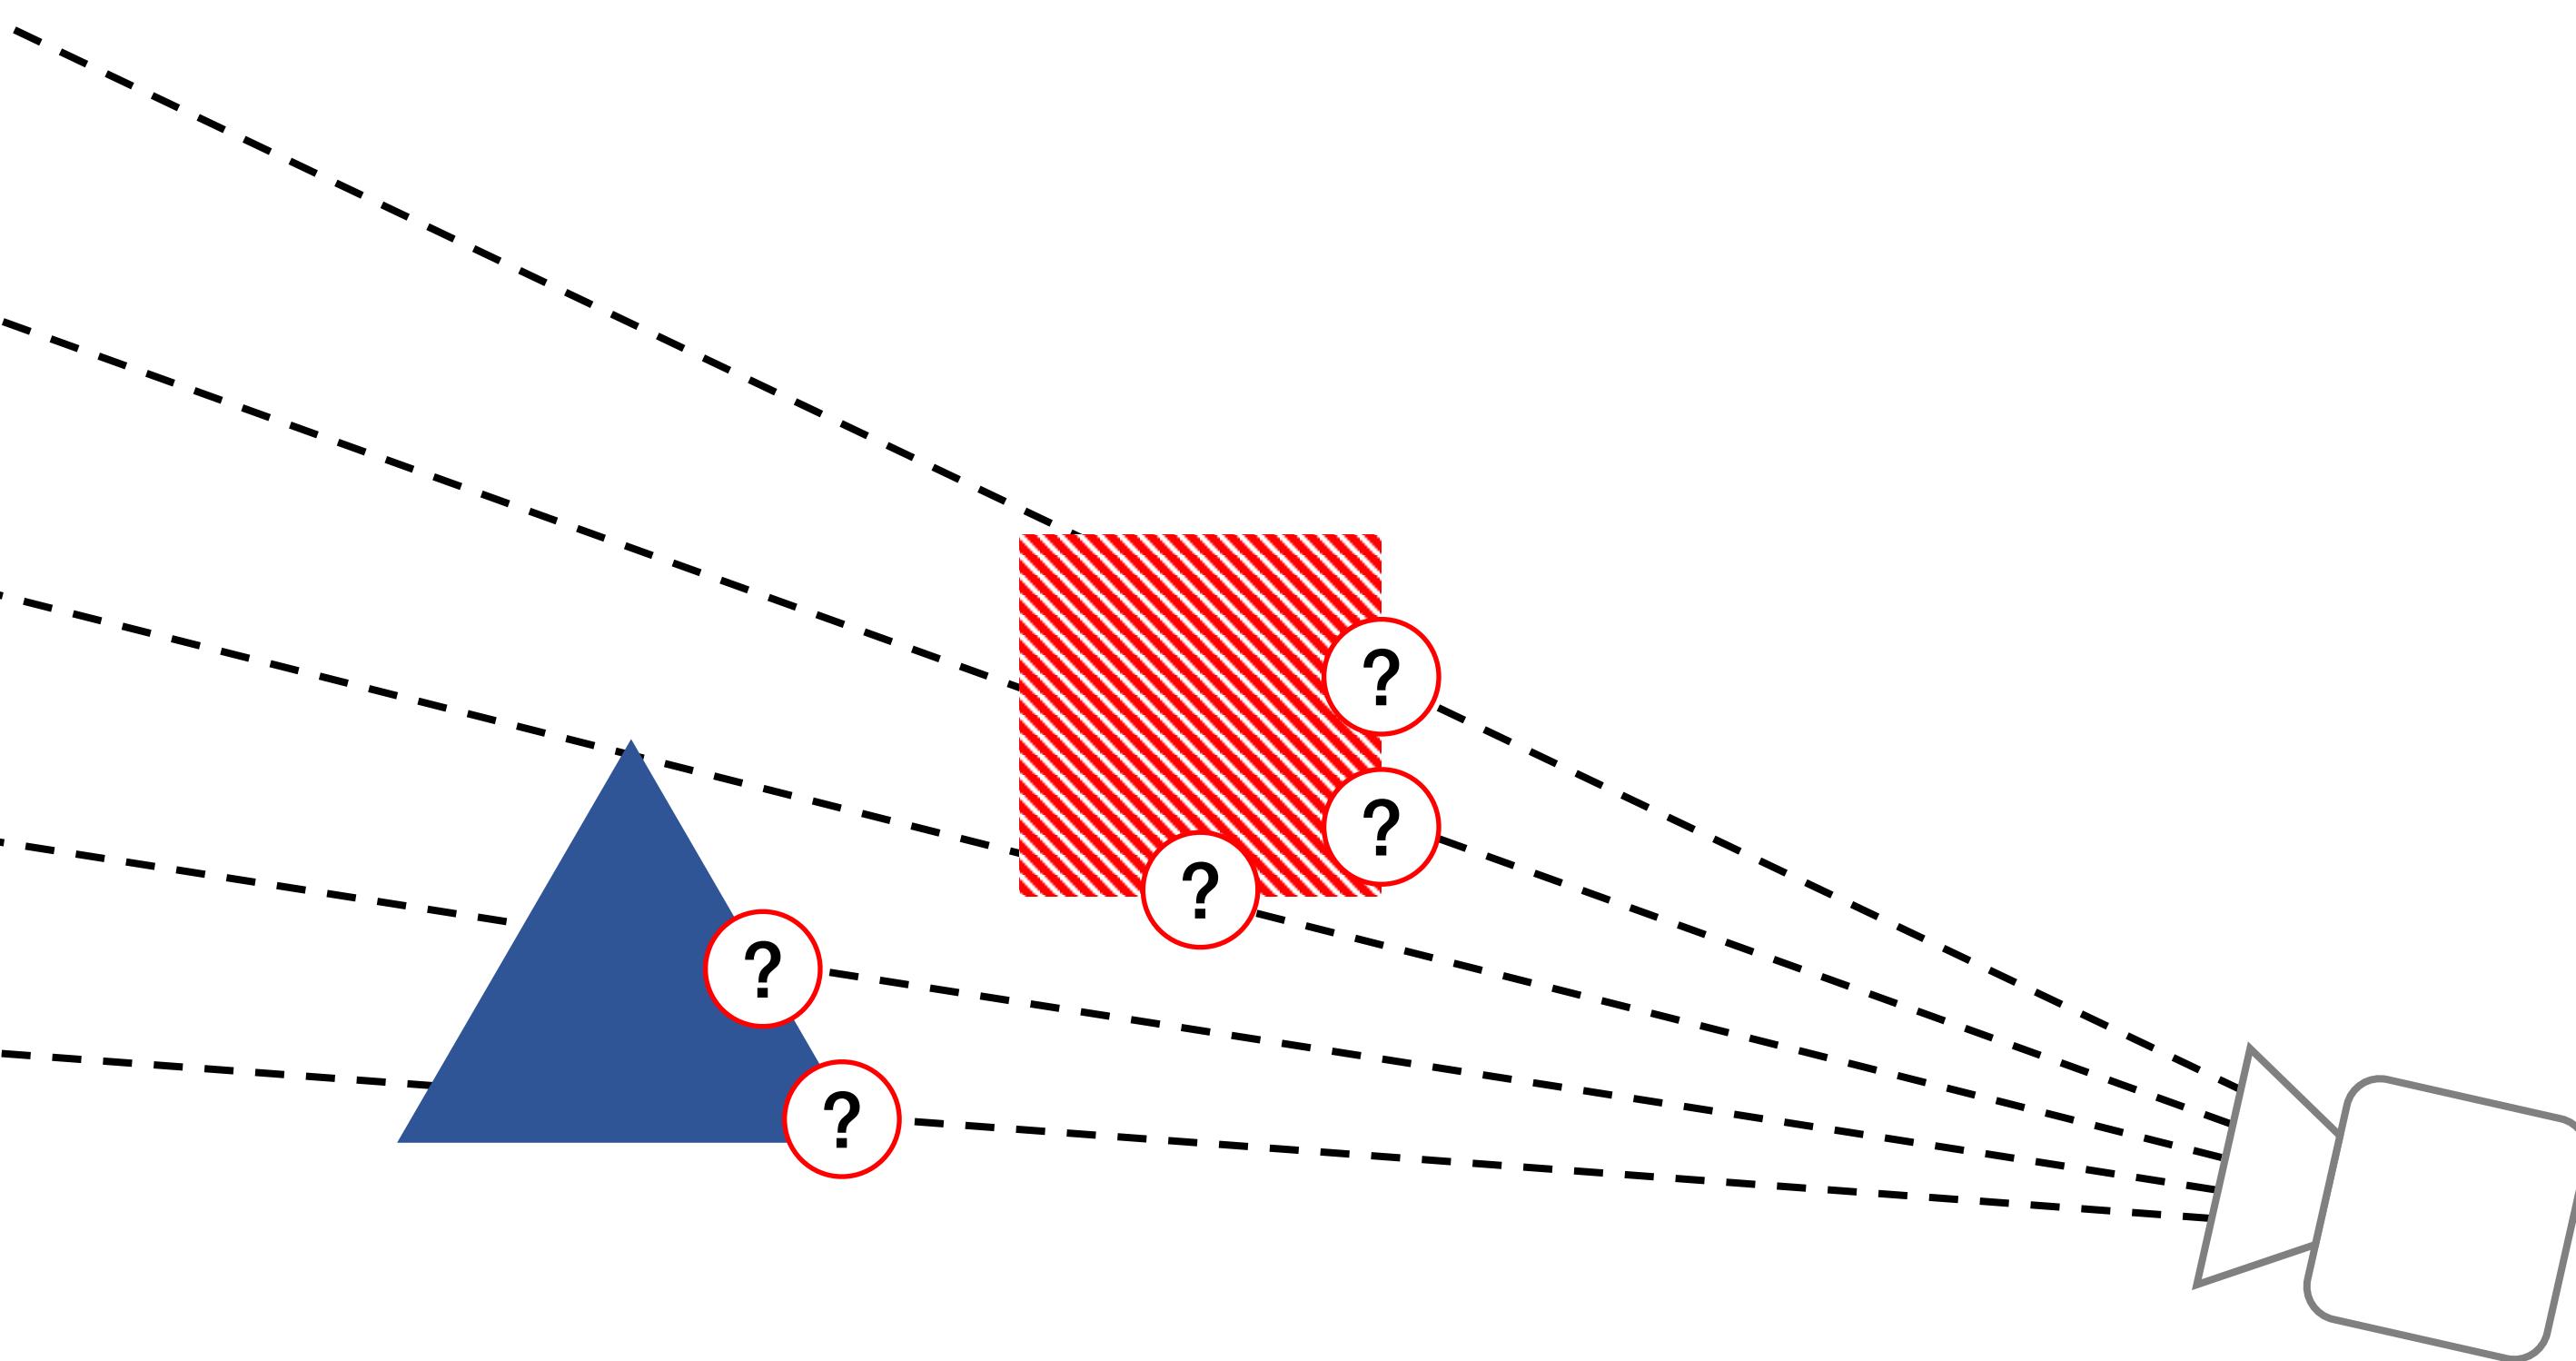
Sphere Tracing



Sphere Tracing

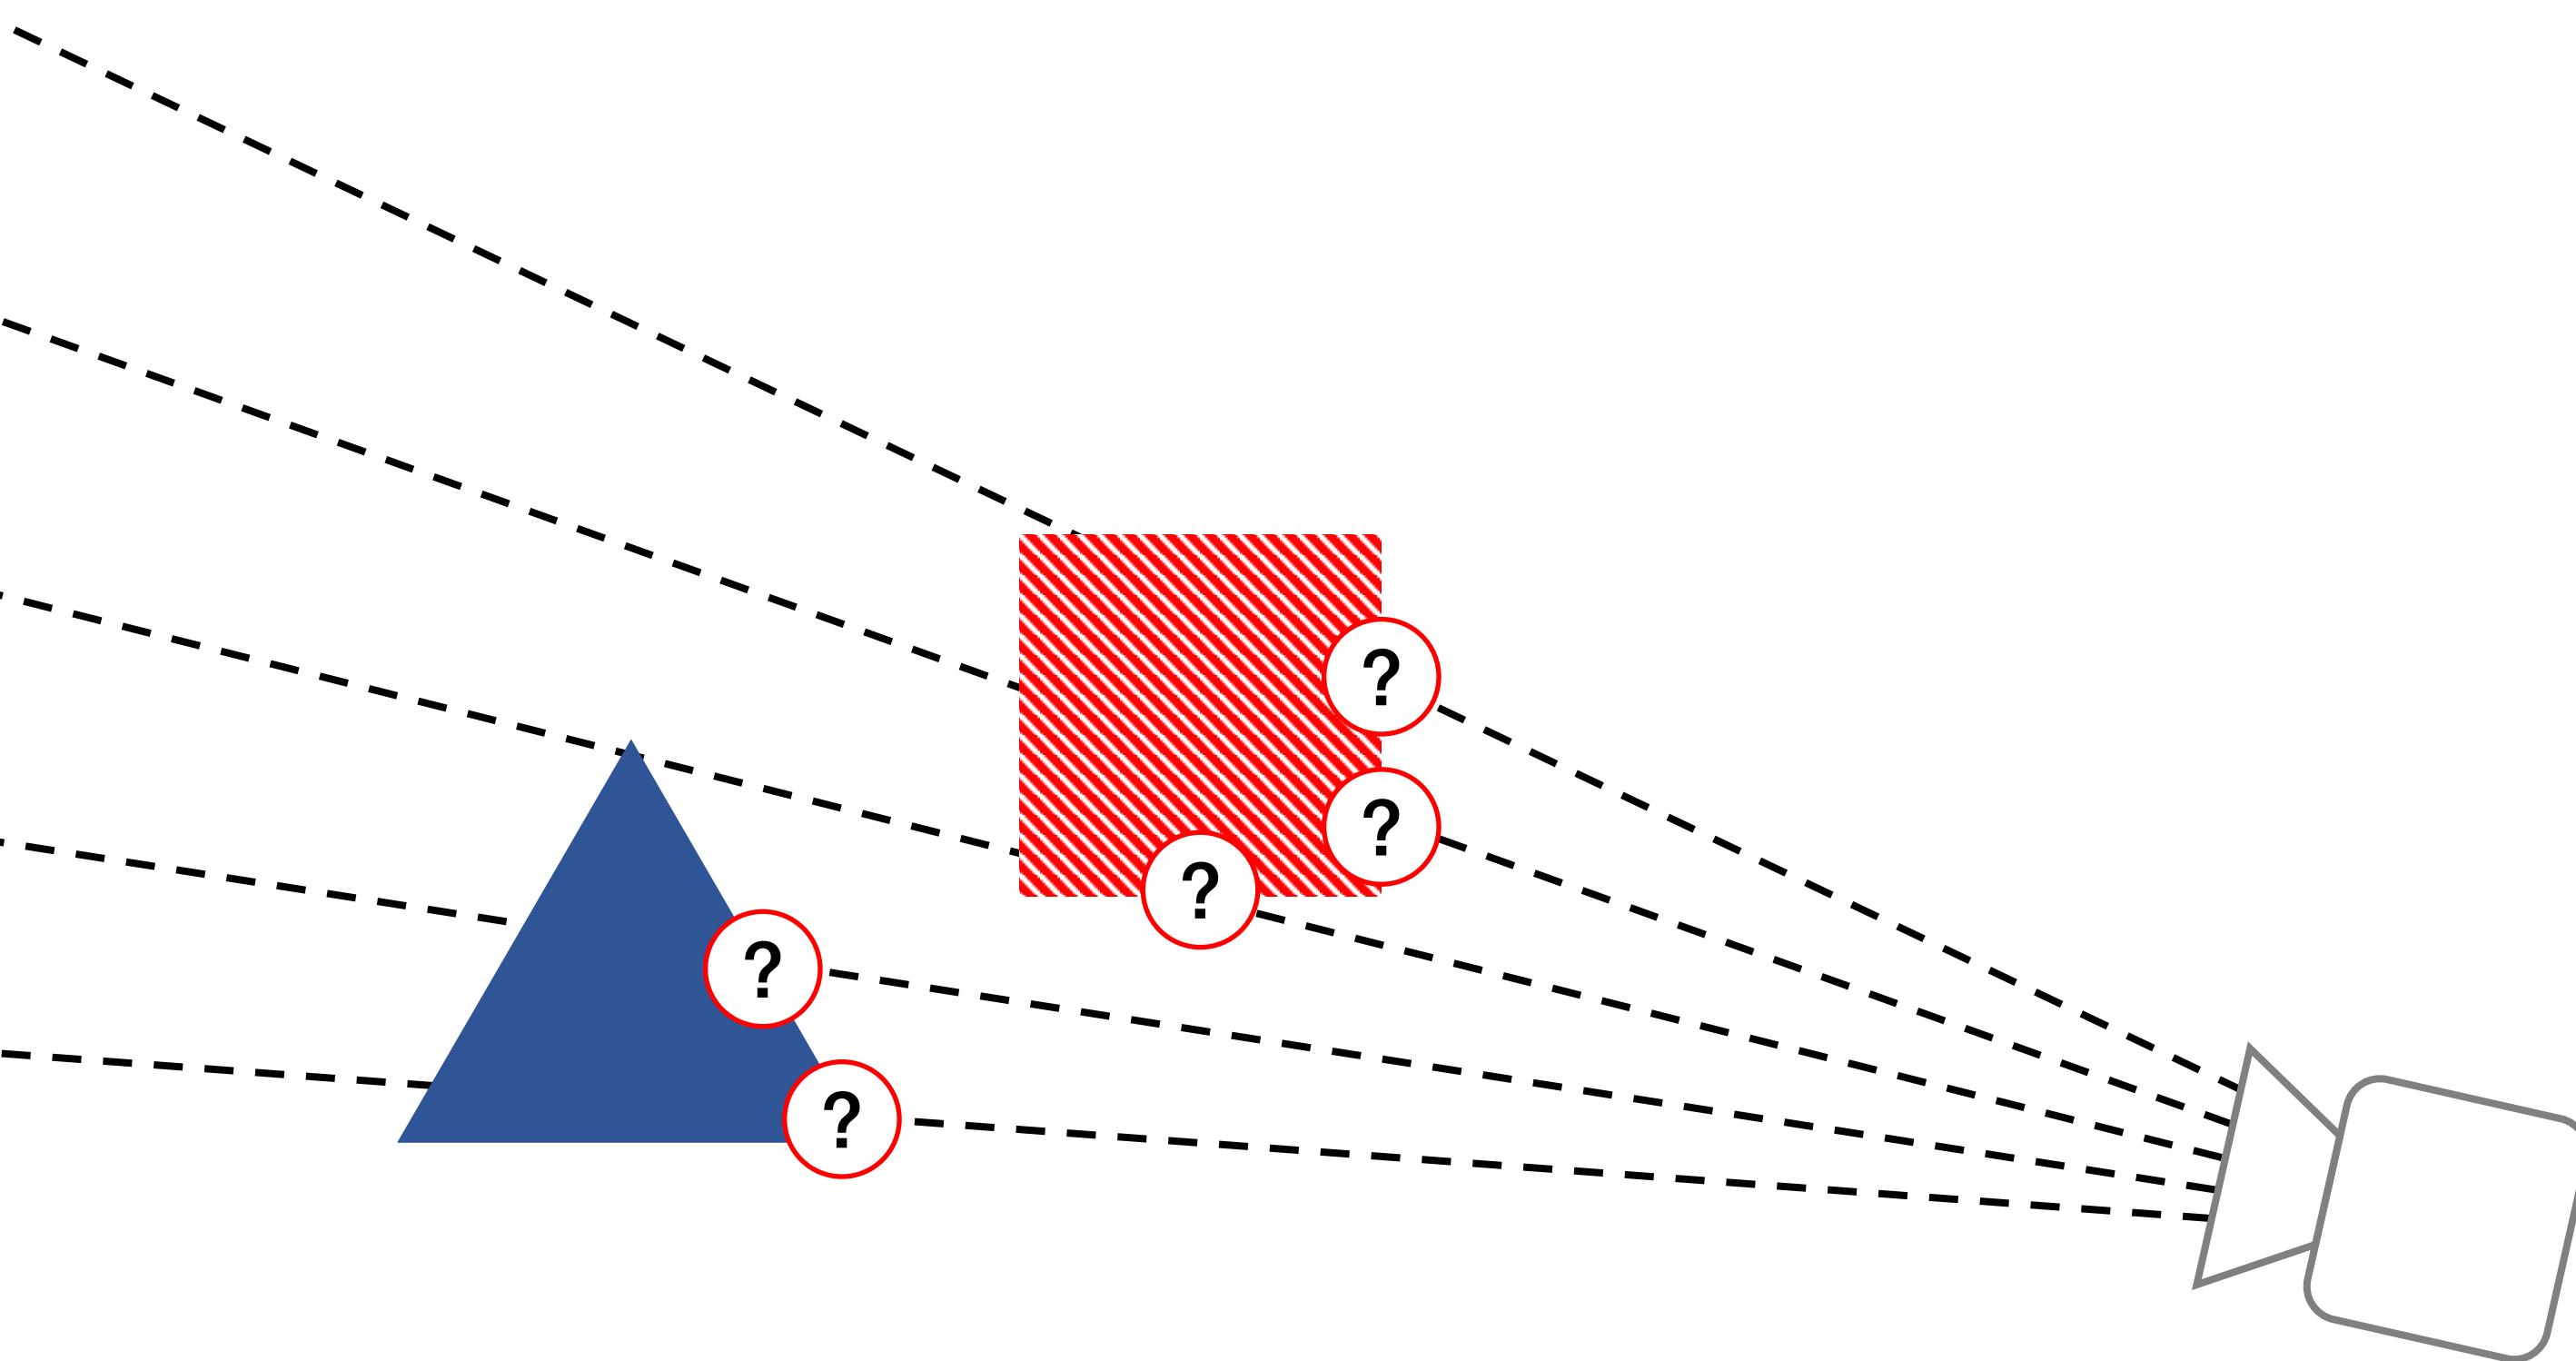


Sphere Tracing Step 1: Intersection Testing.



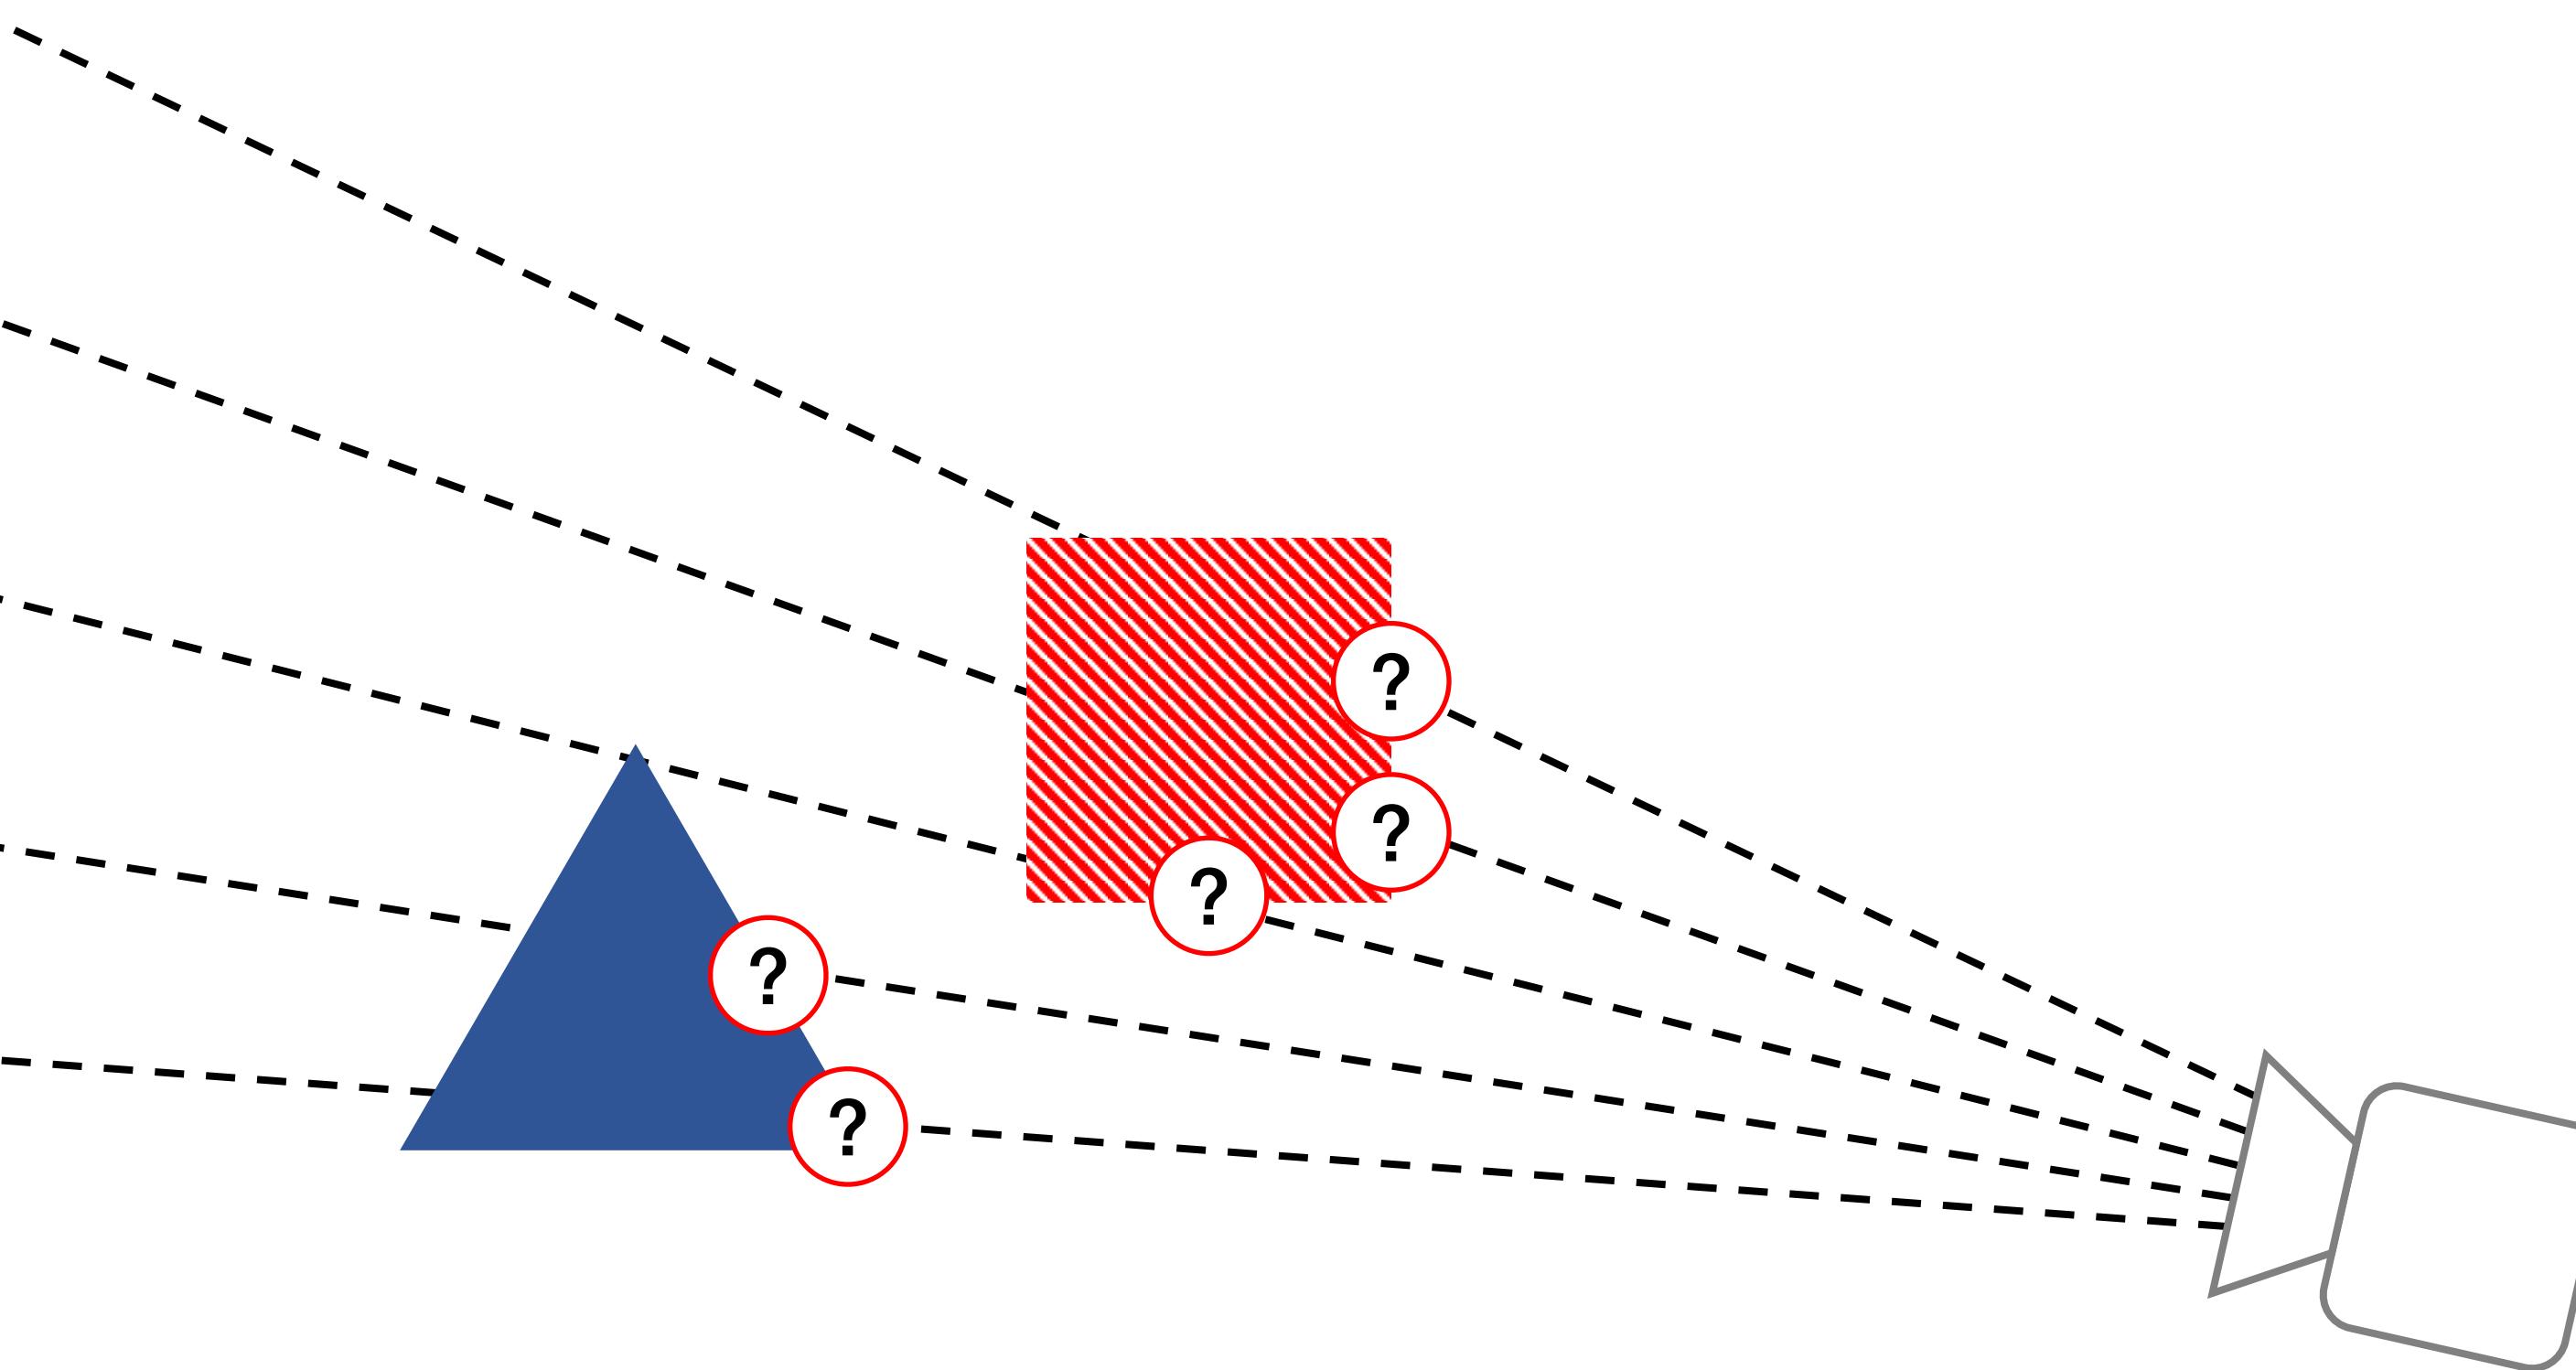
Sphere Tracing Step 1: Intersection Testing.

How to parameterize scene geometry?



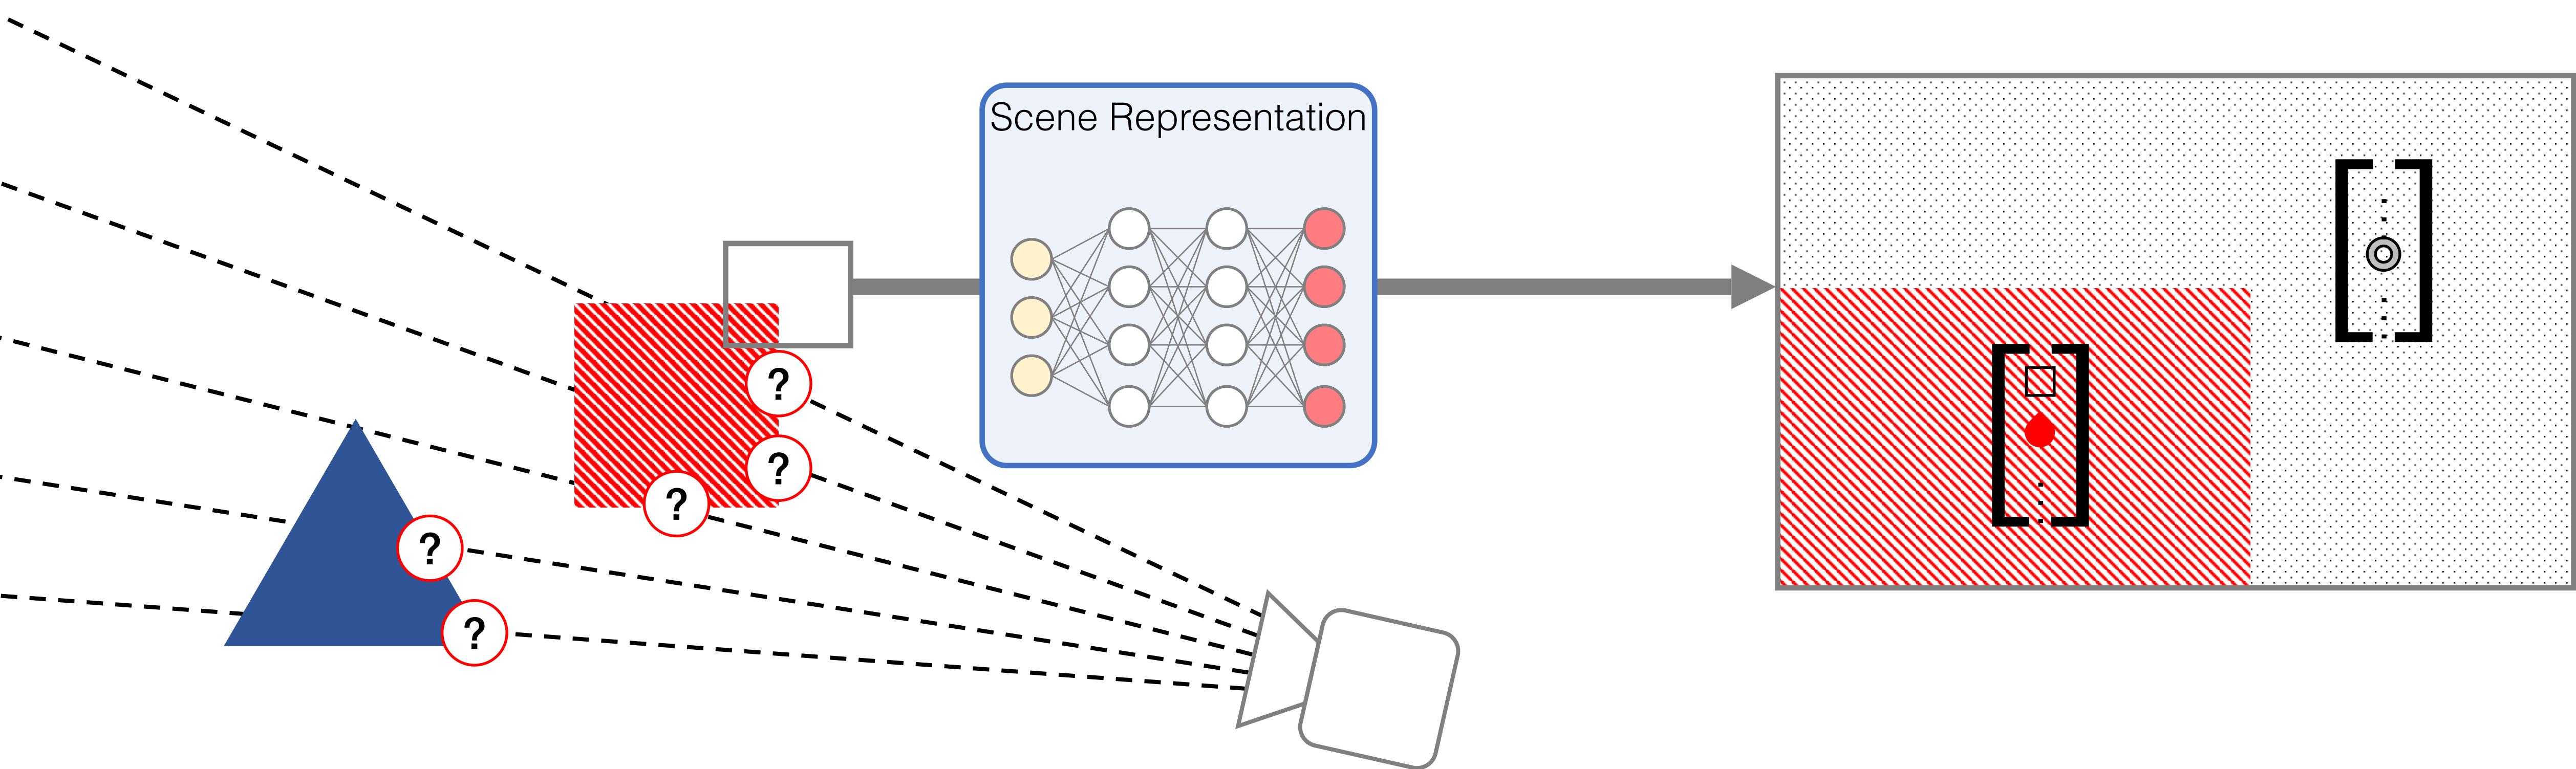
Sphere Tracing Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



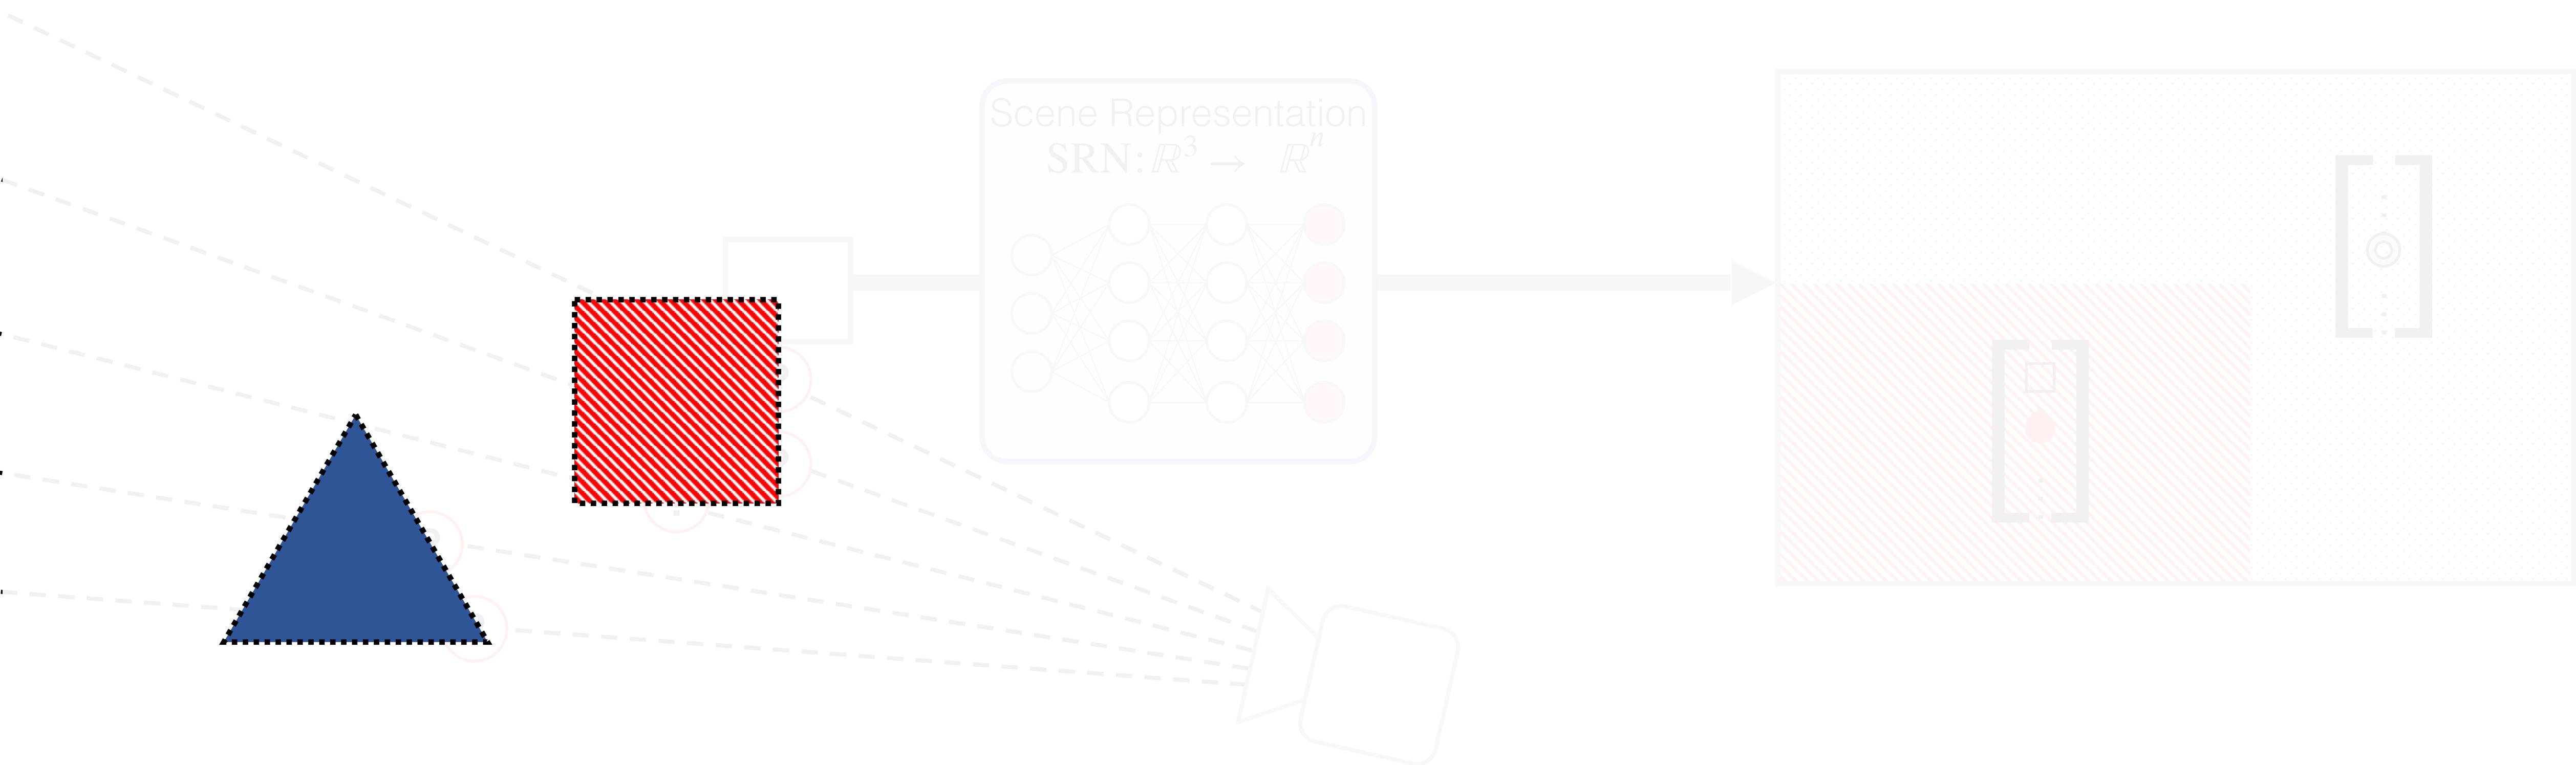
Sphere Tracing Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



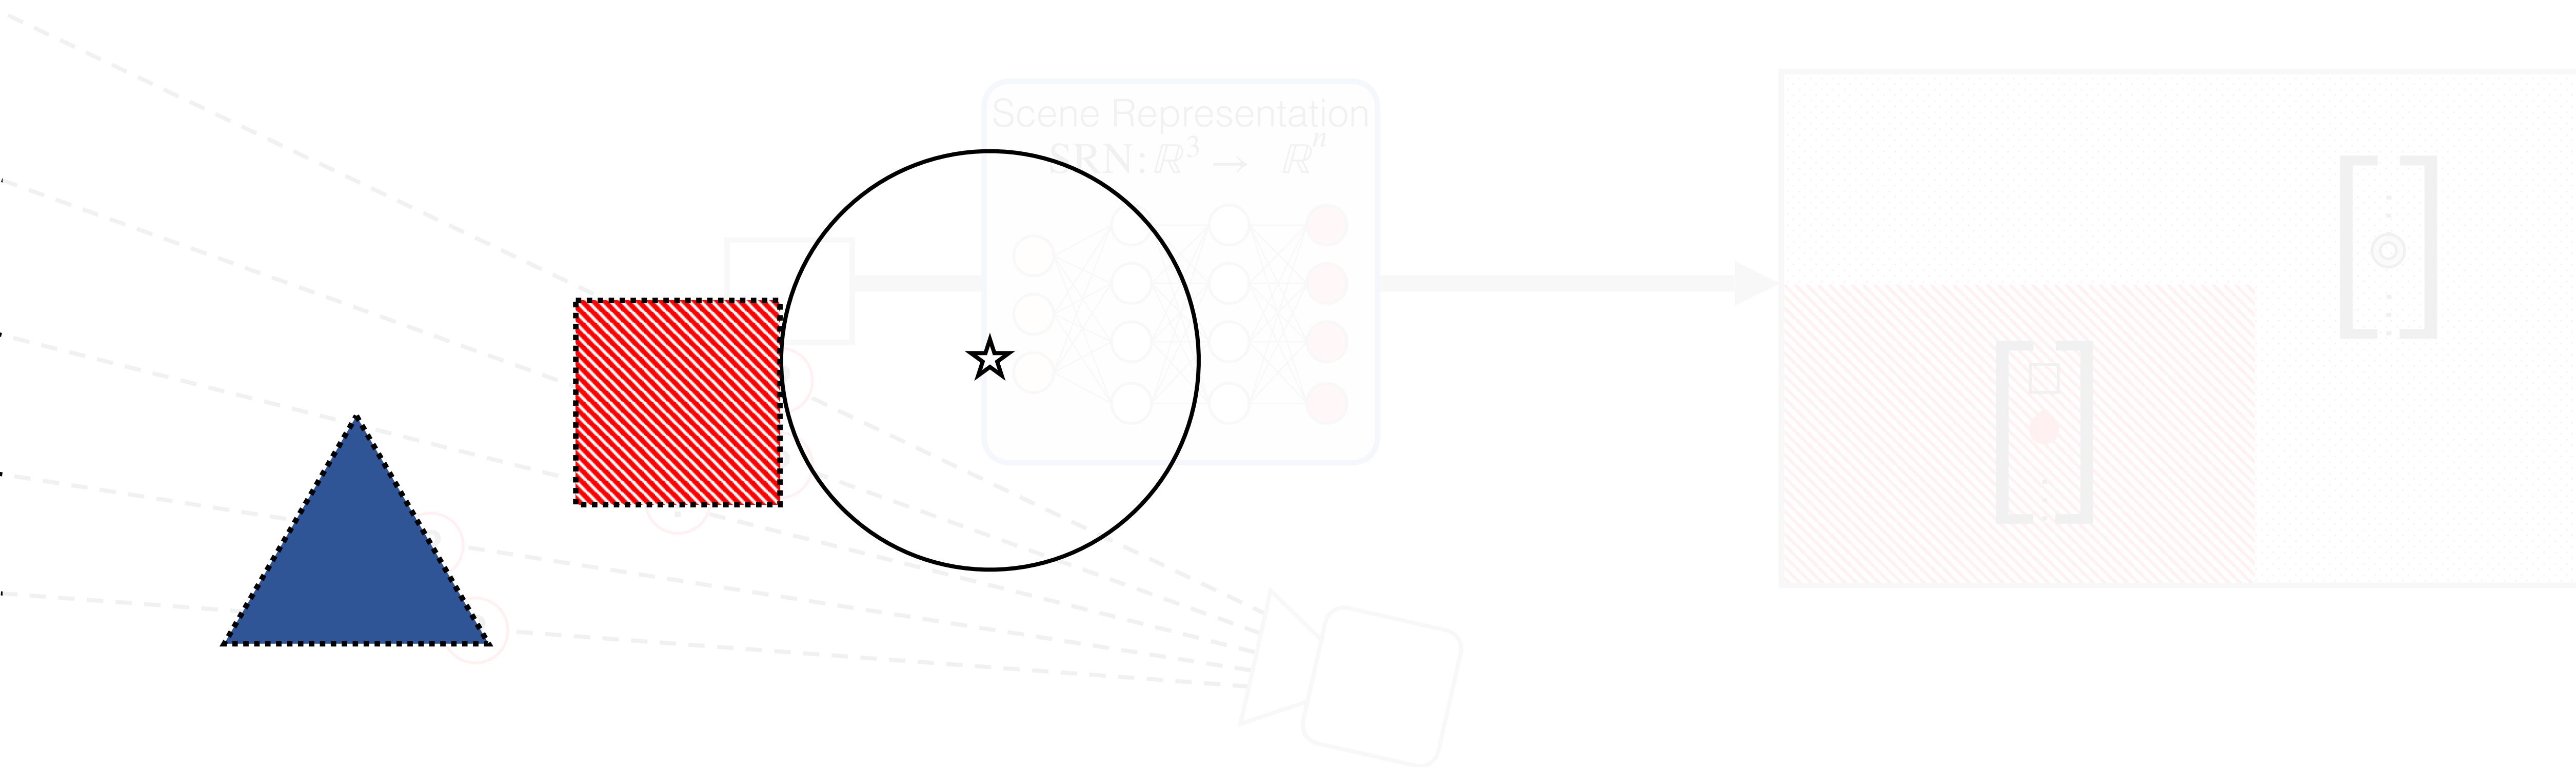
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



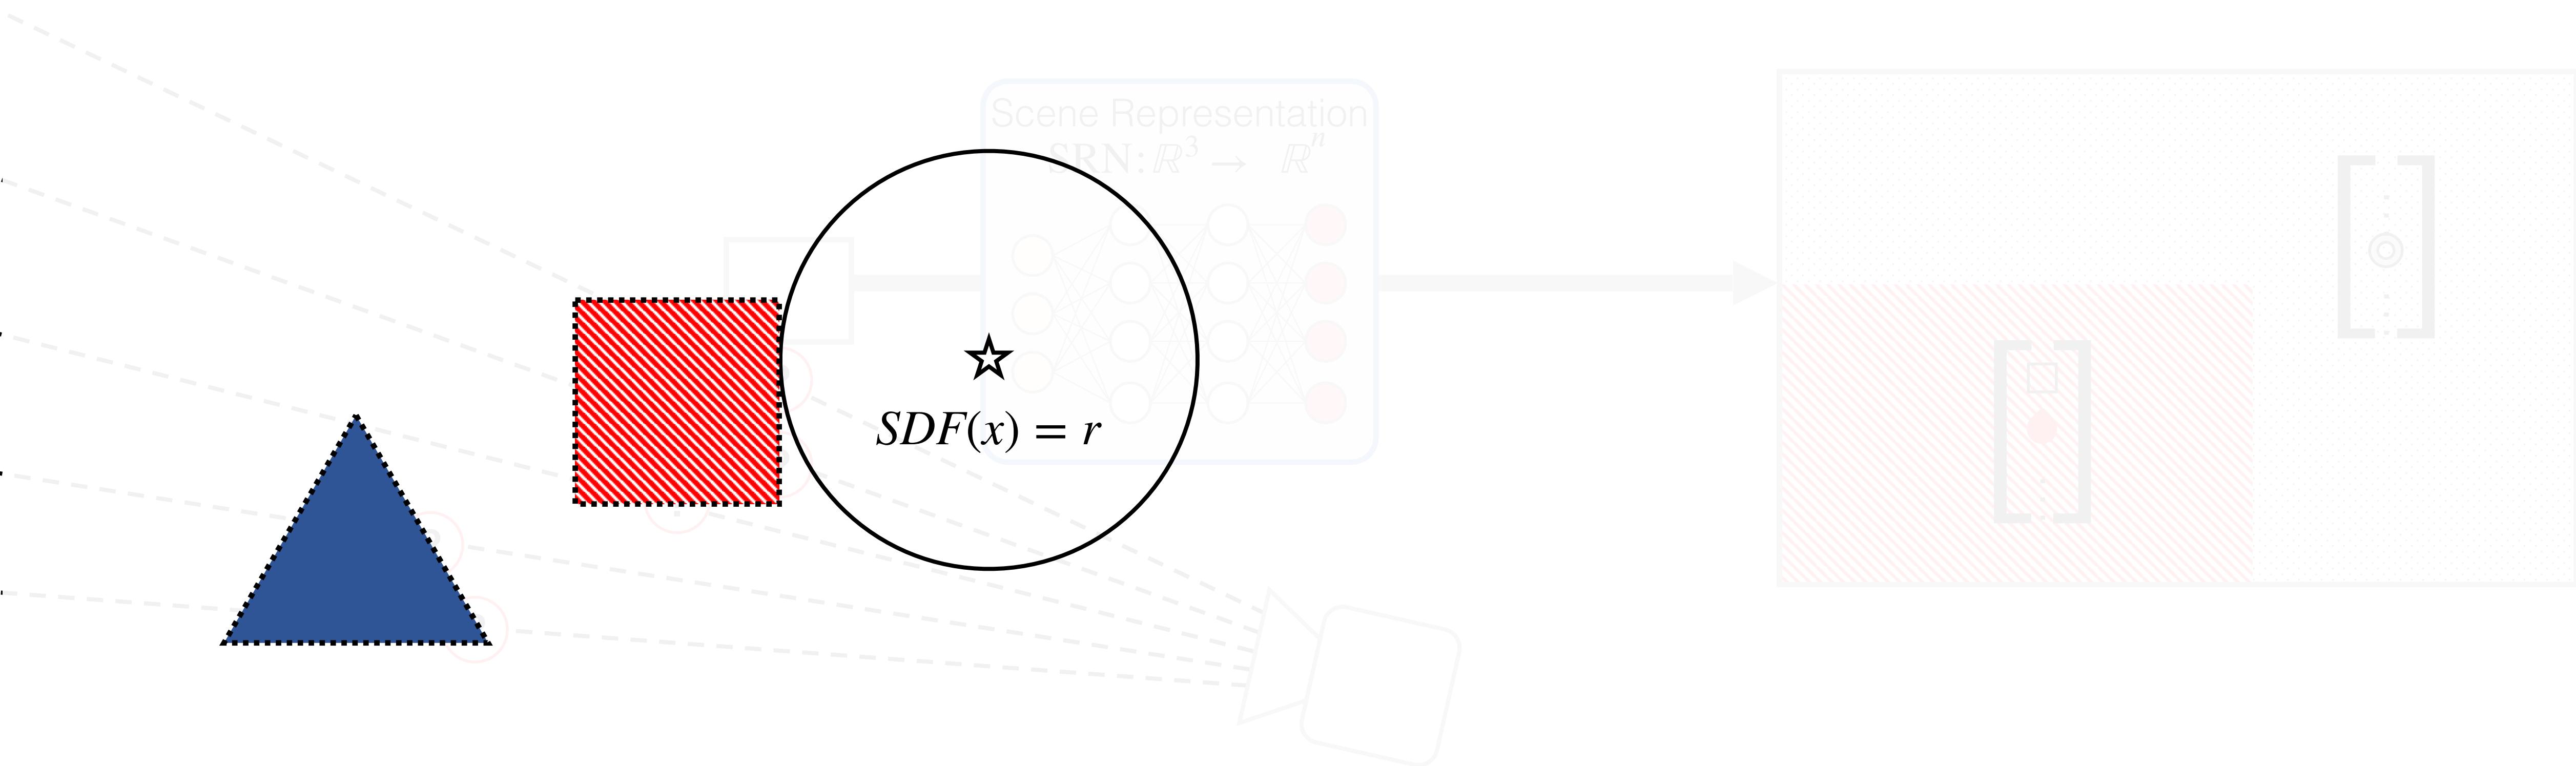
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



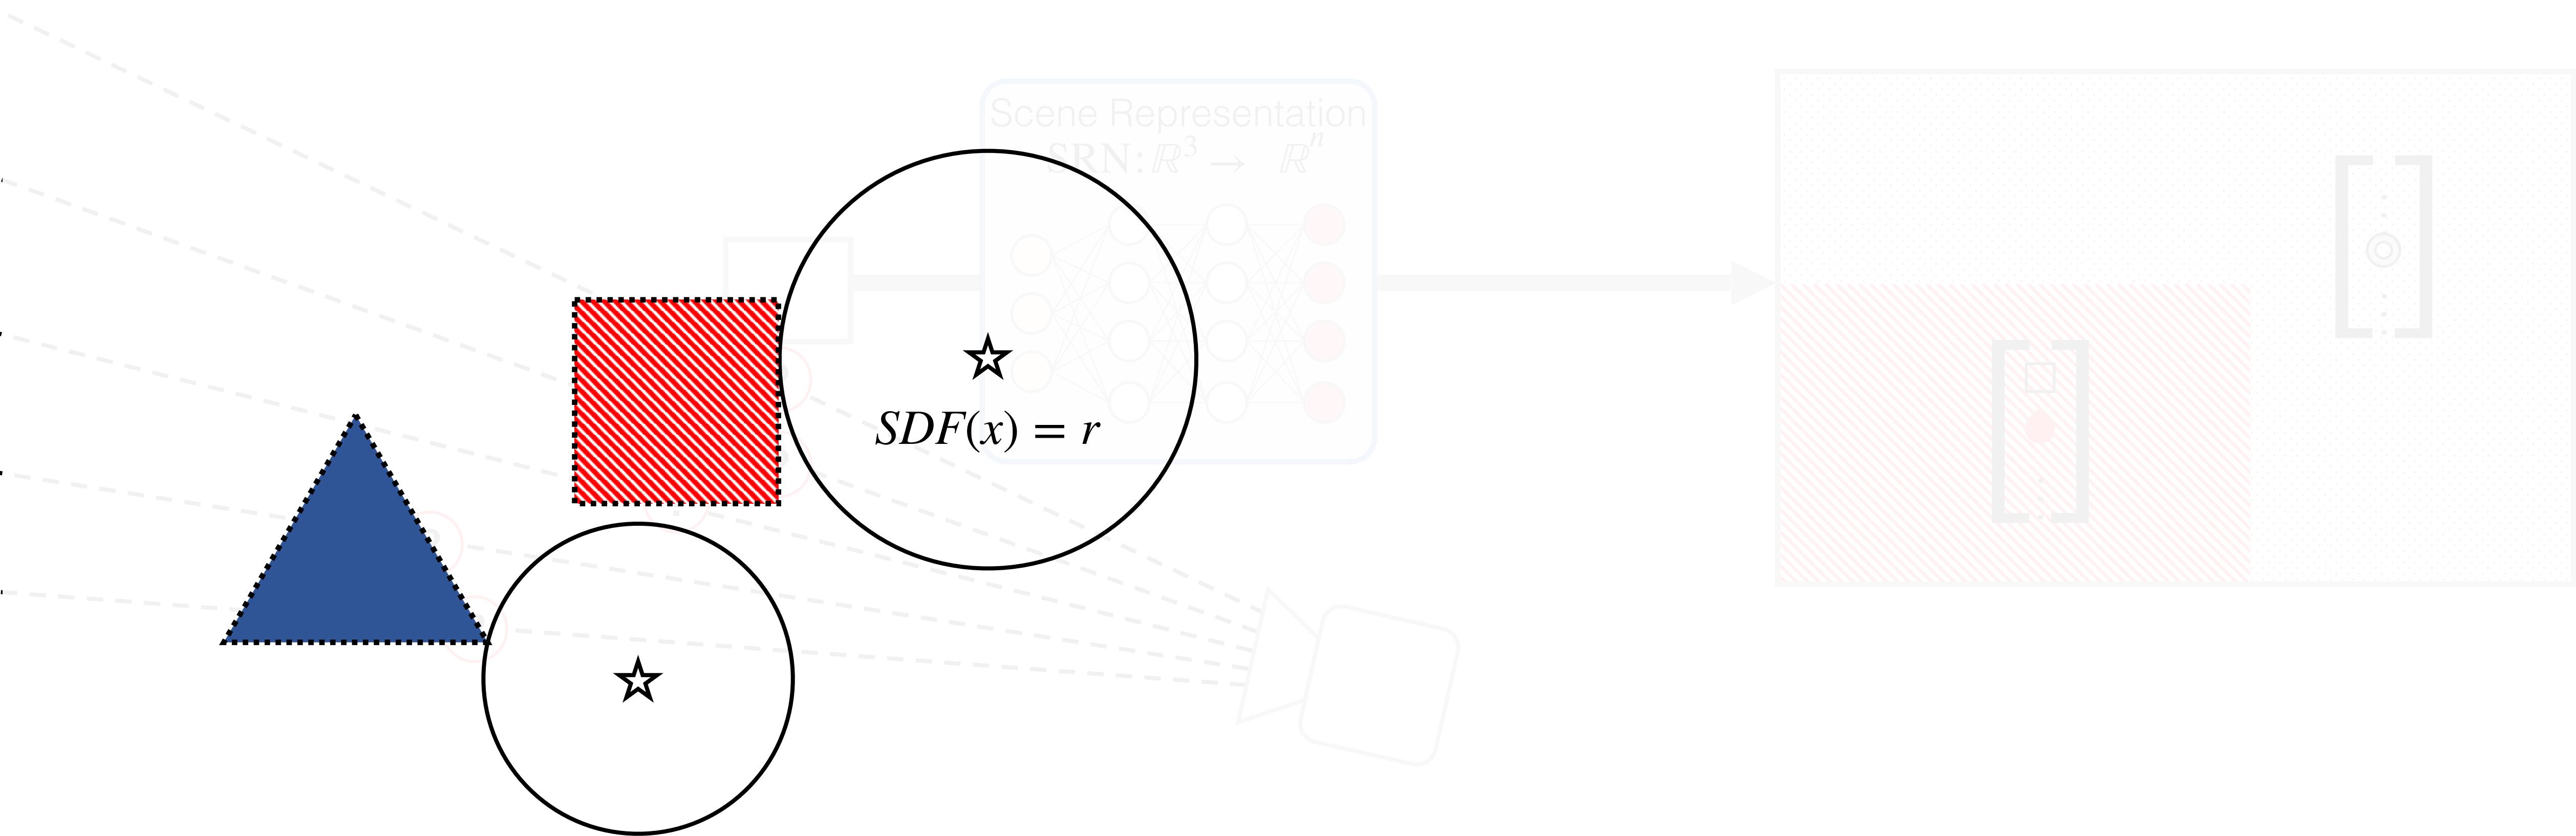
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



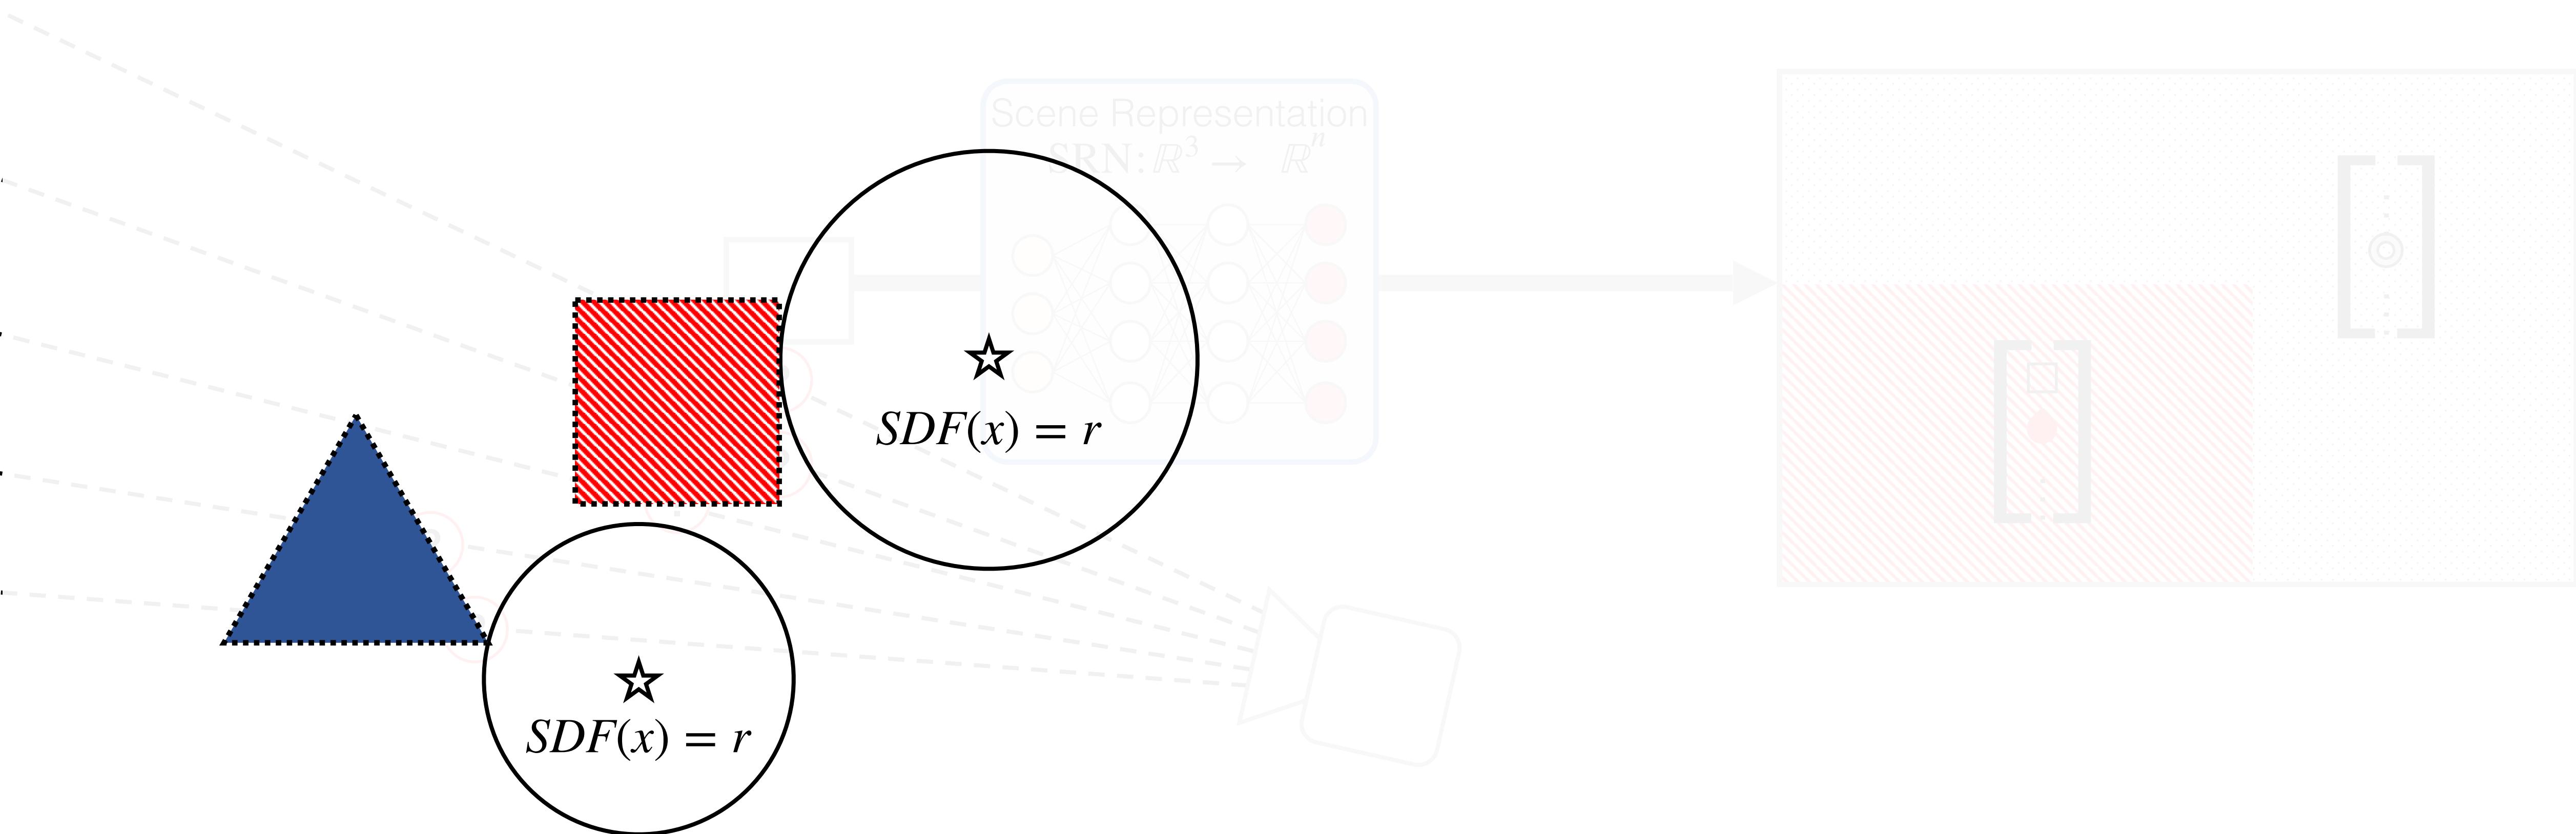
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



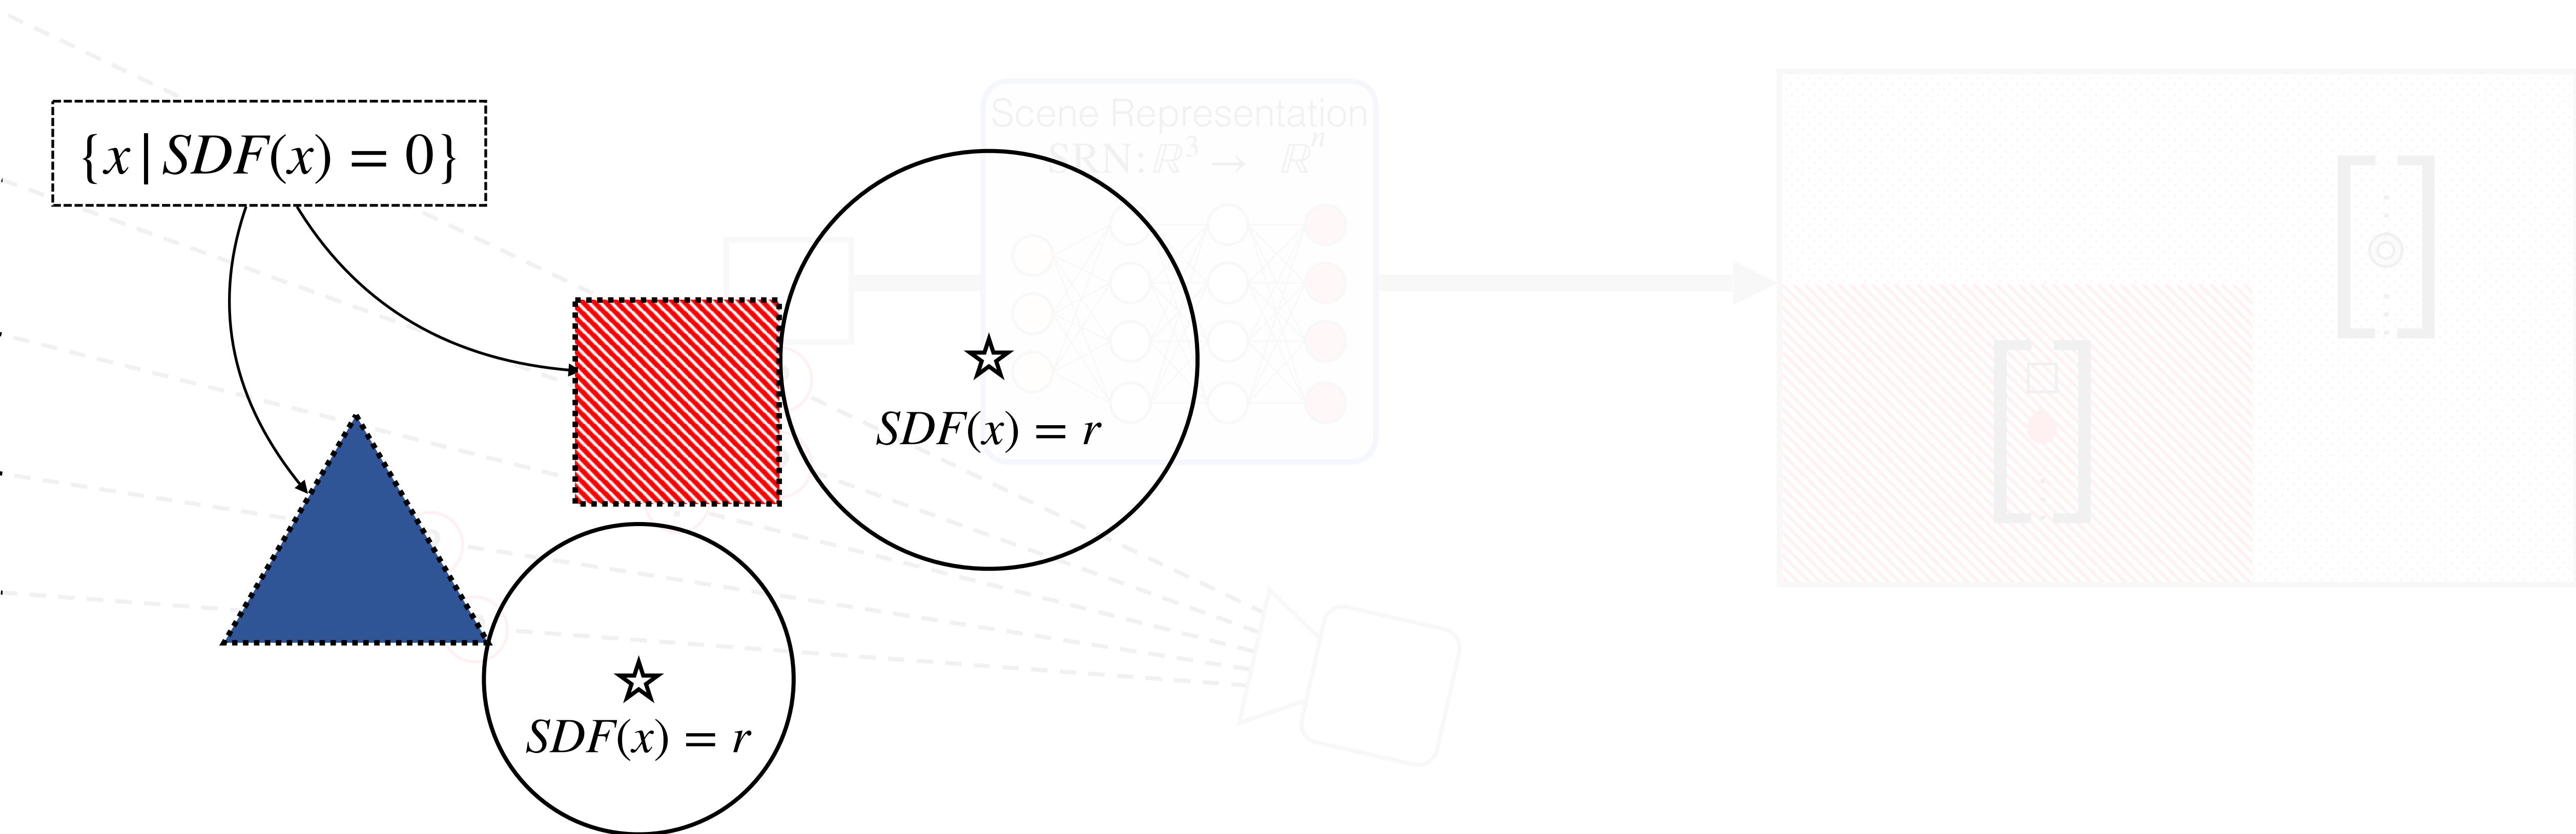
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?



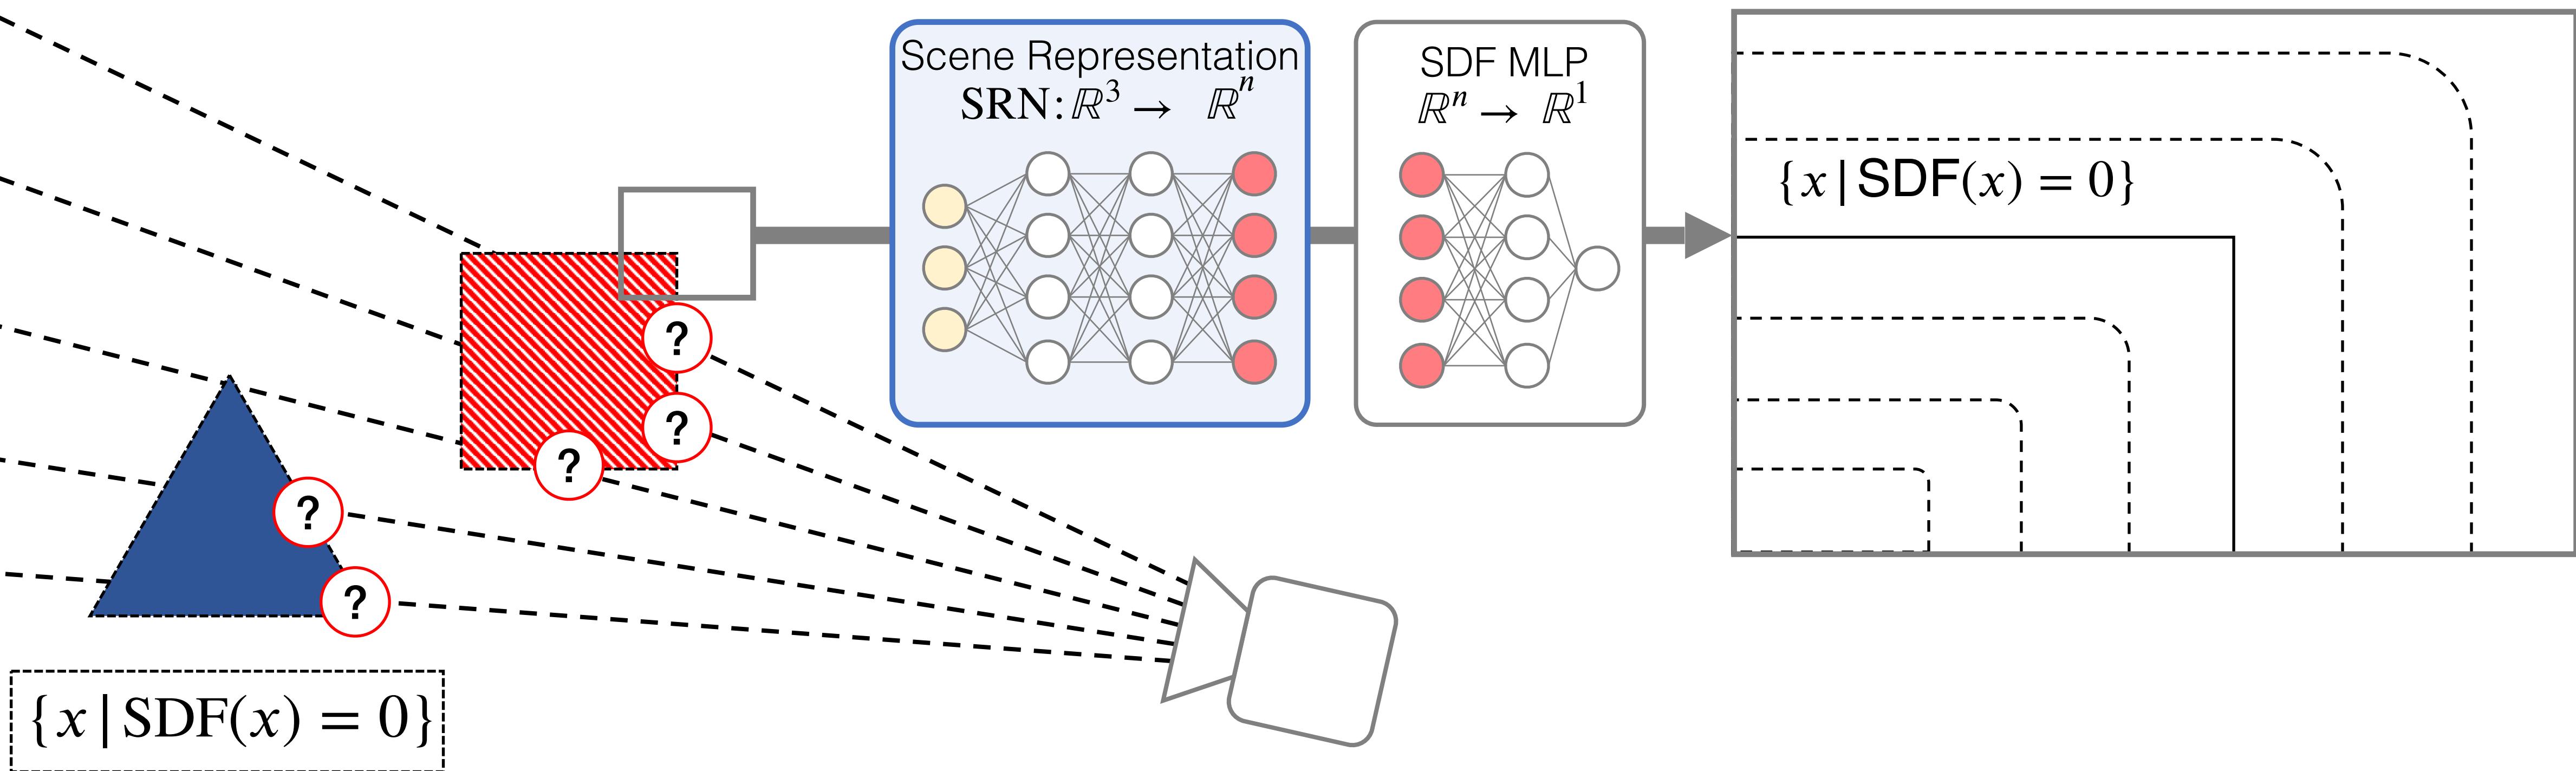
Neural Renderer Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?

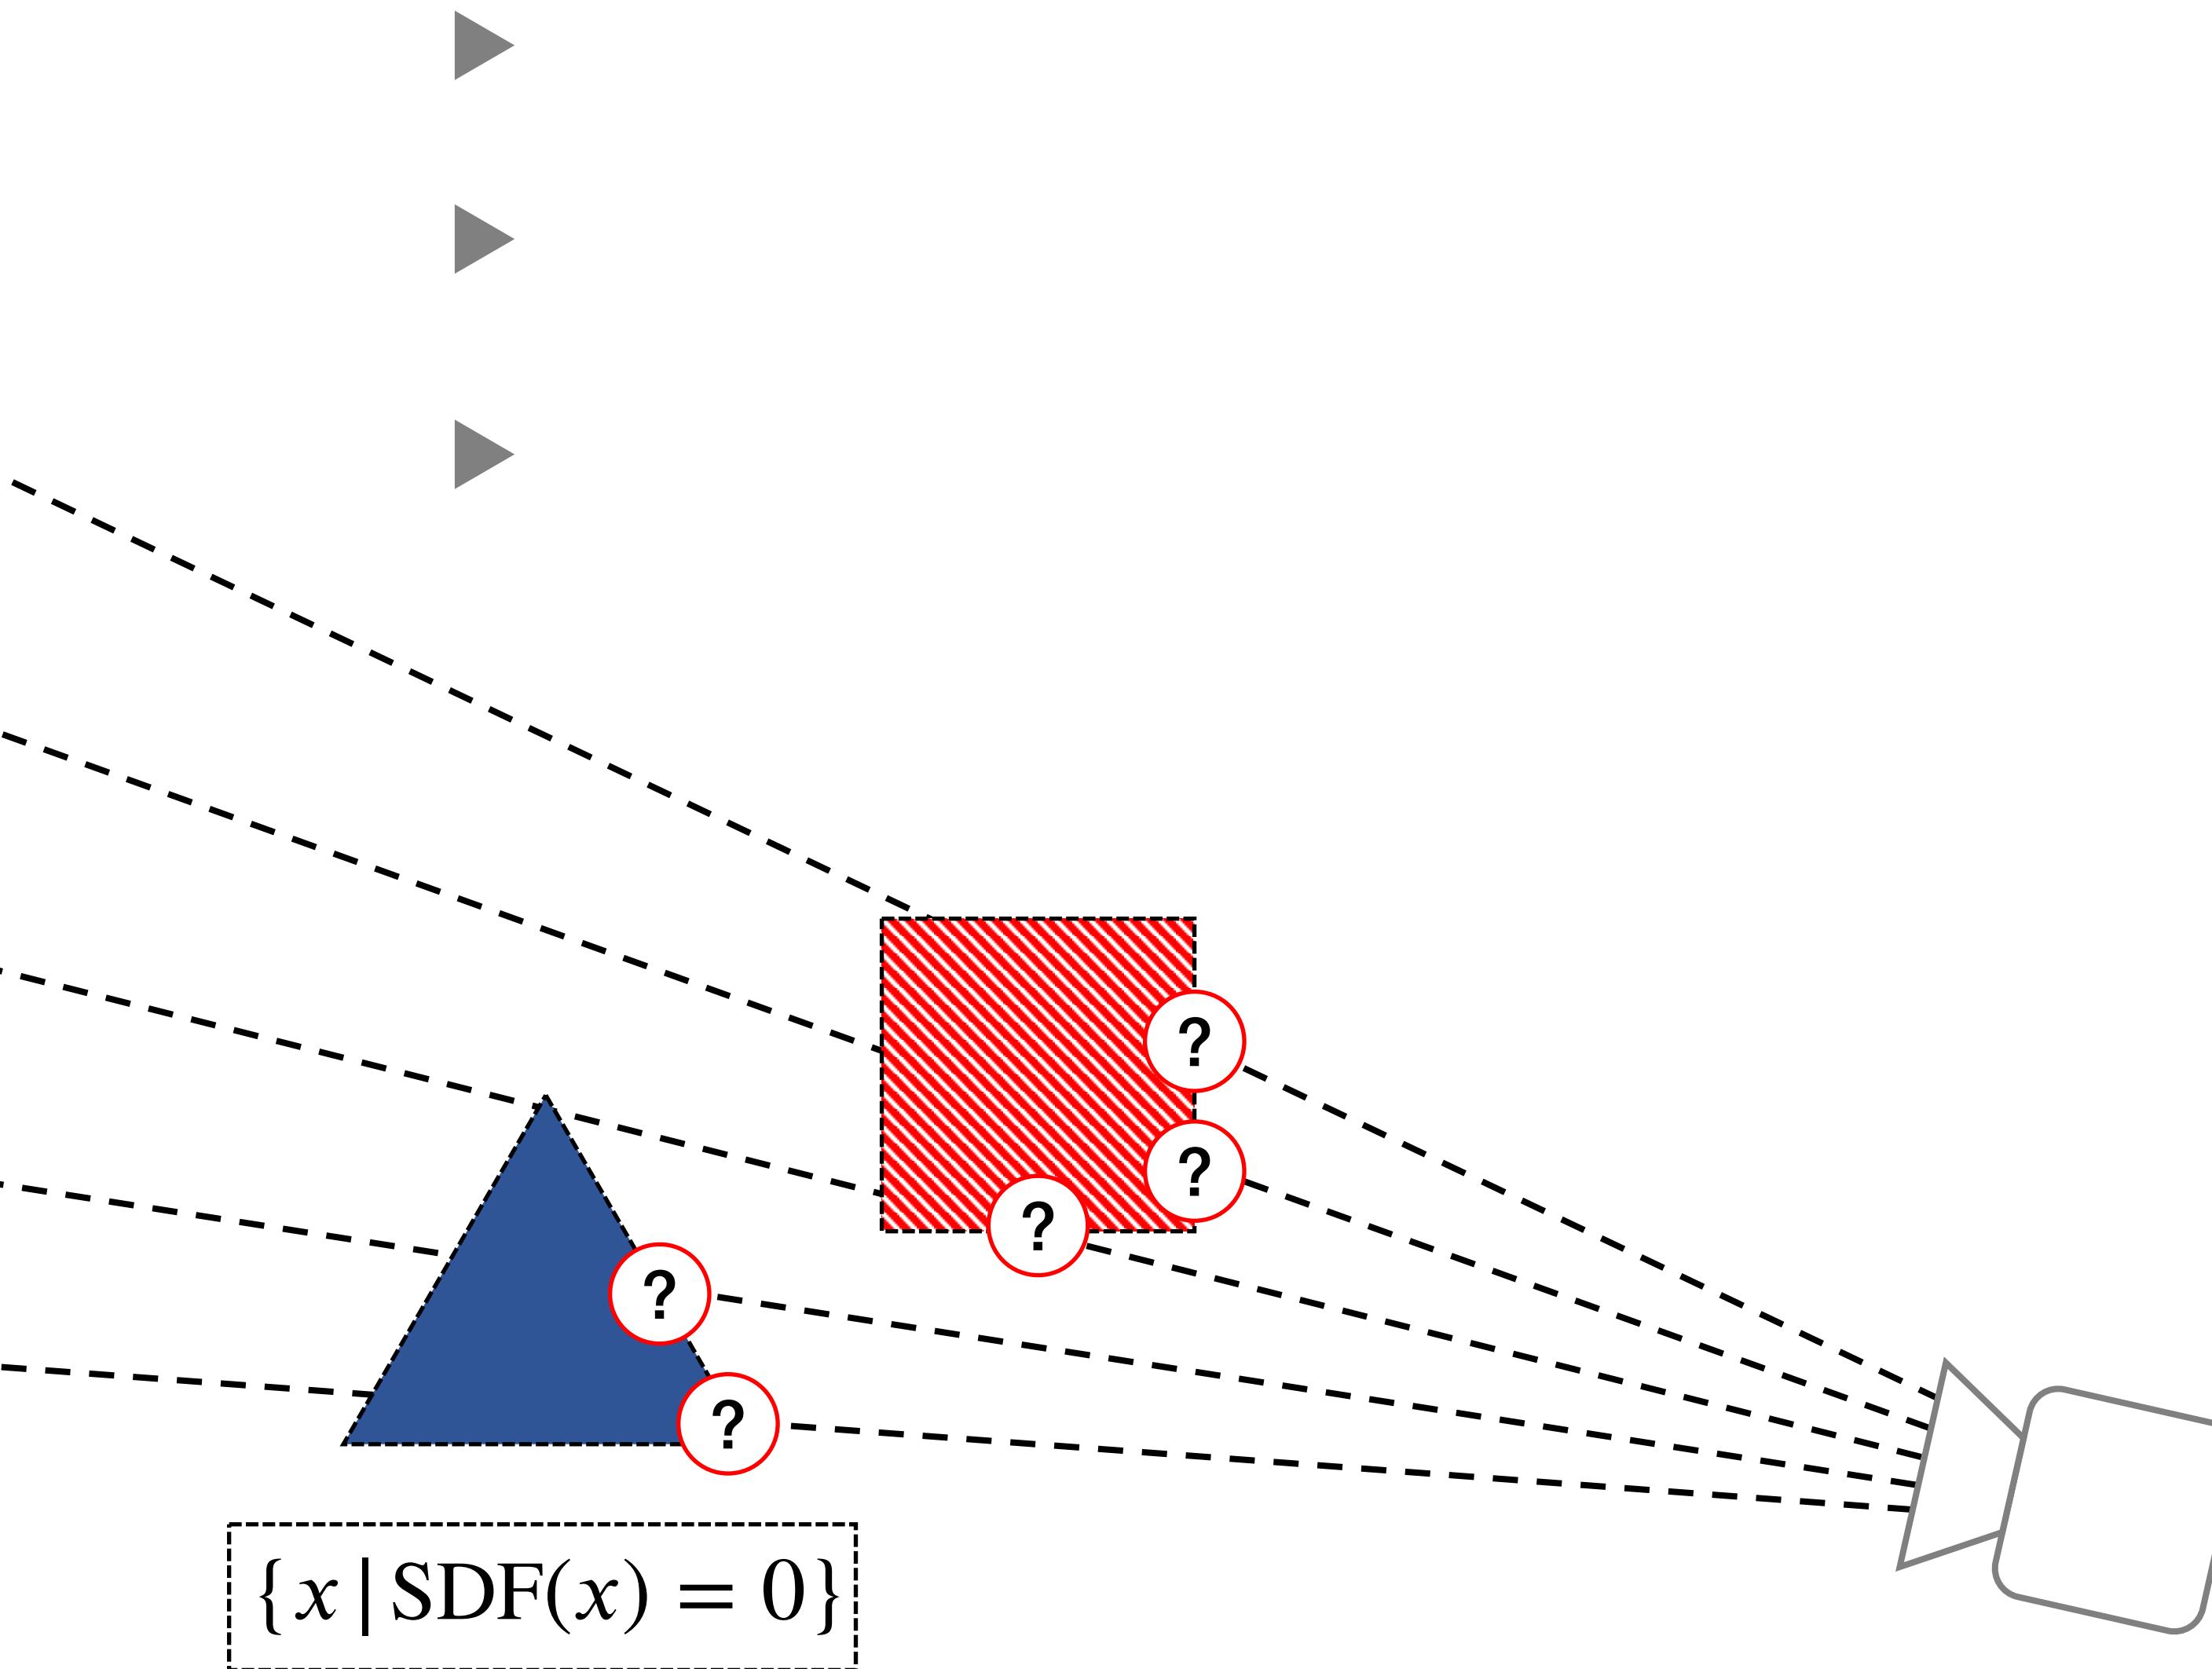


Sphere Tracing Step 1: Intersection Testing.

- ▶ How to parameterize scene geometry?
- ▶ Signed Distance Function maps point to distance to closest surface.

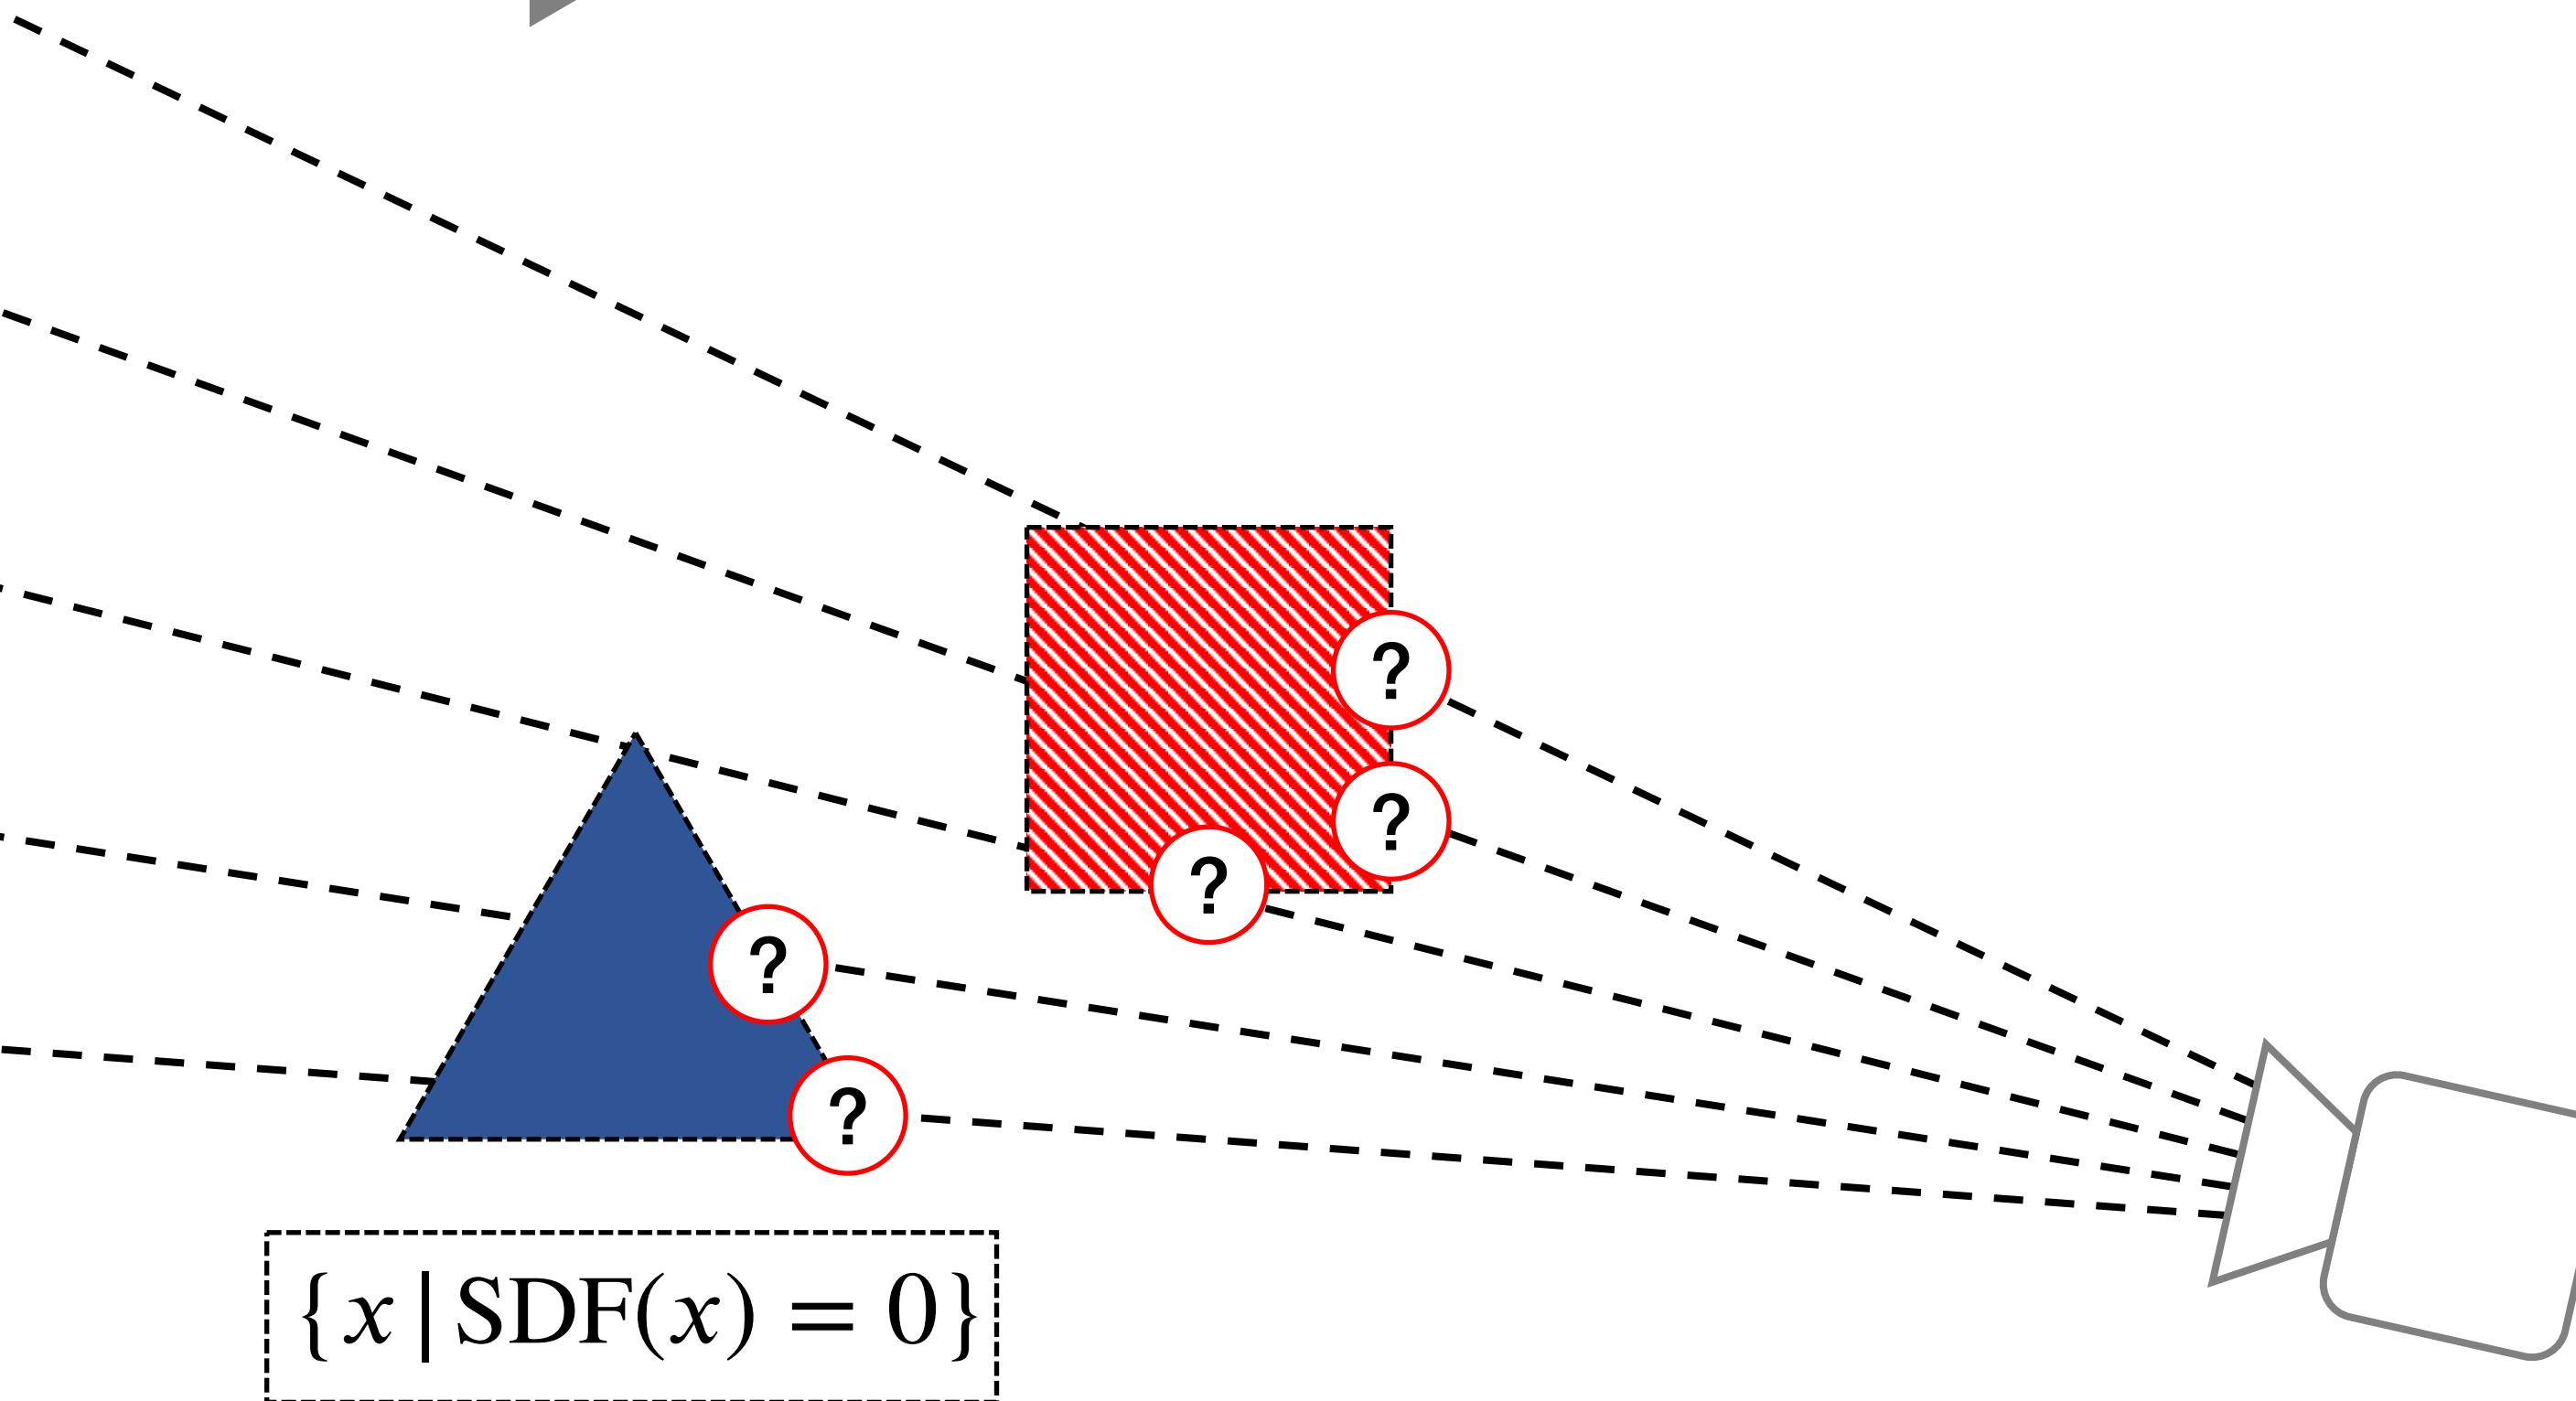


Sphere Tracing Step 1: Intersection Testing.



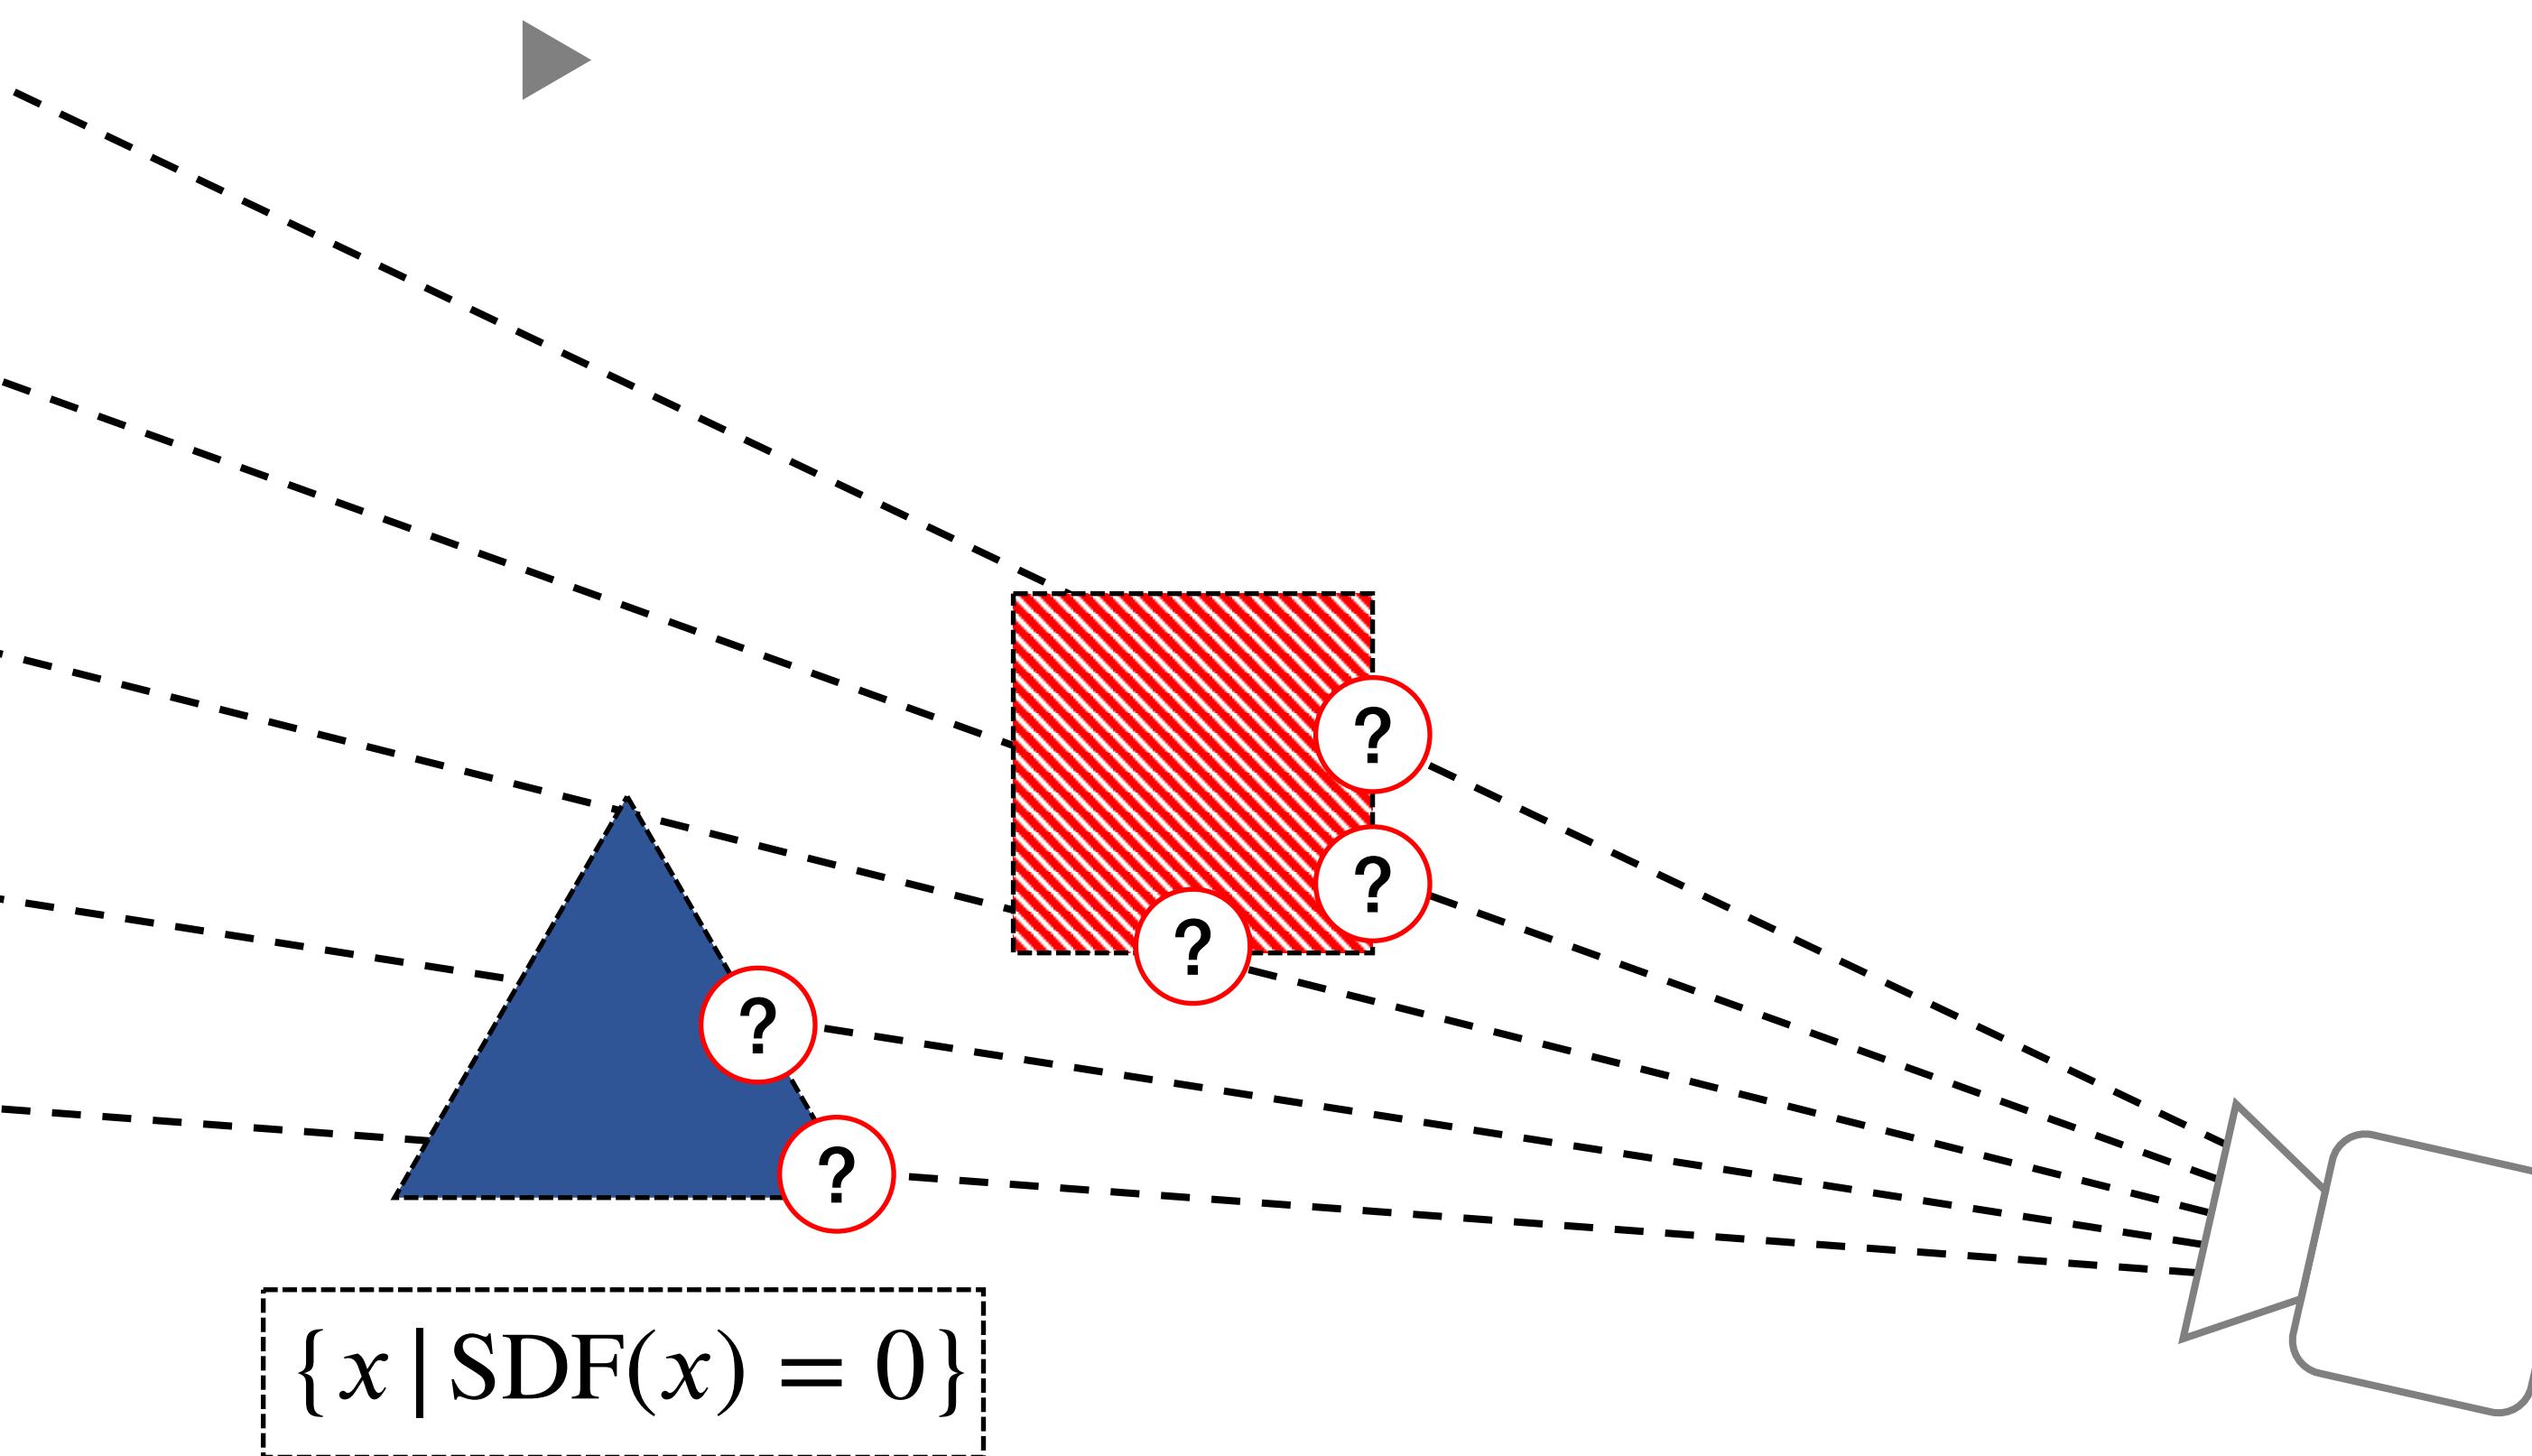
Sphere Tracing Step 1: Intersection Testing.

- ▶ How to find intersections?



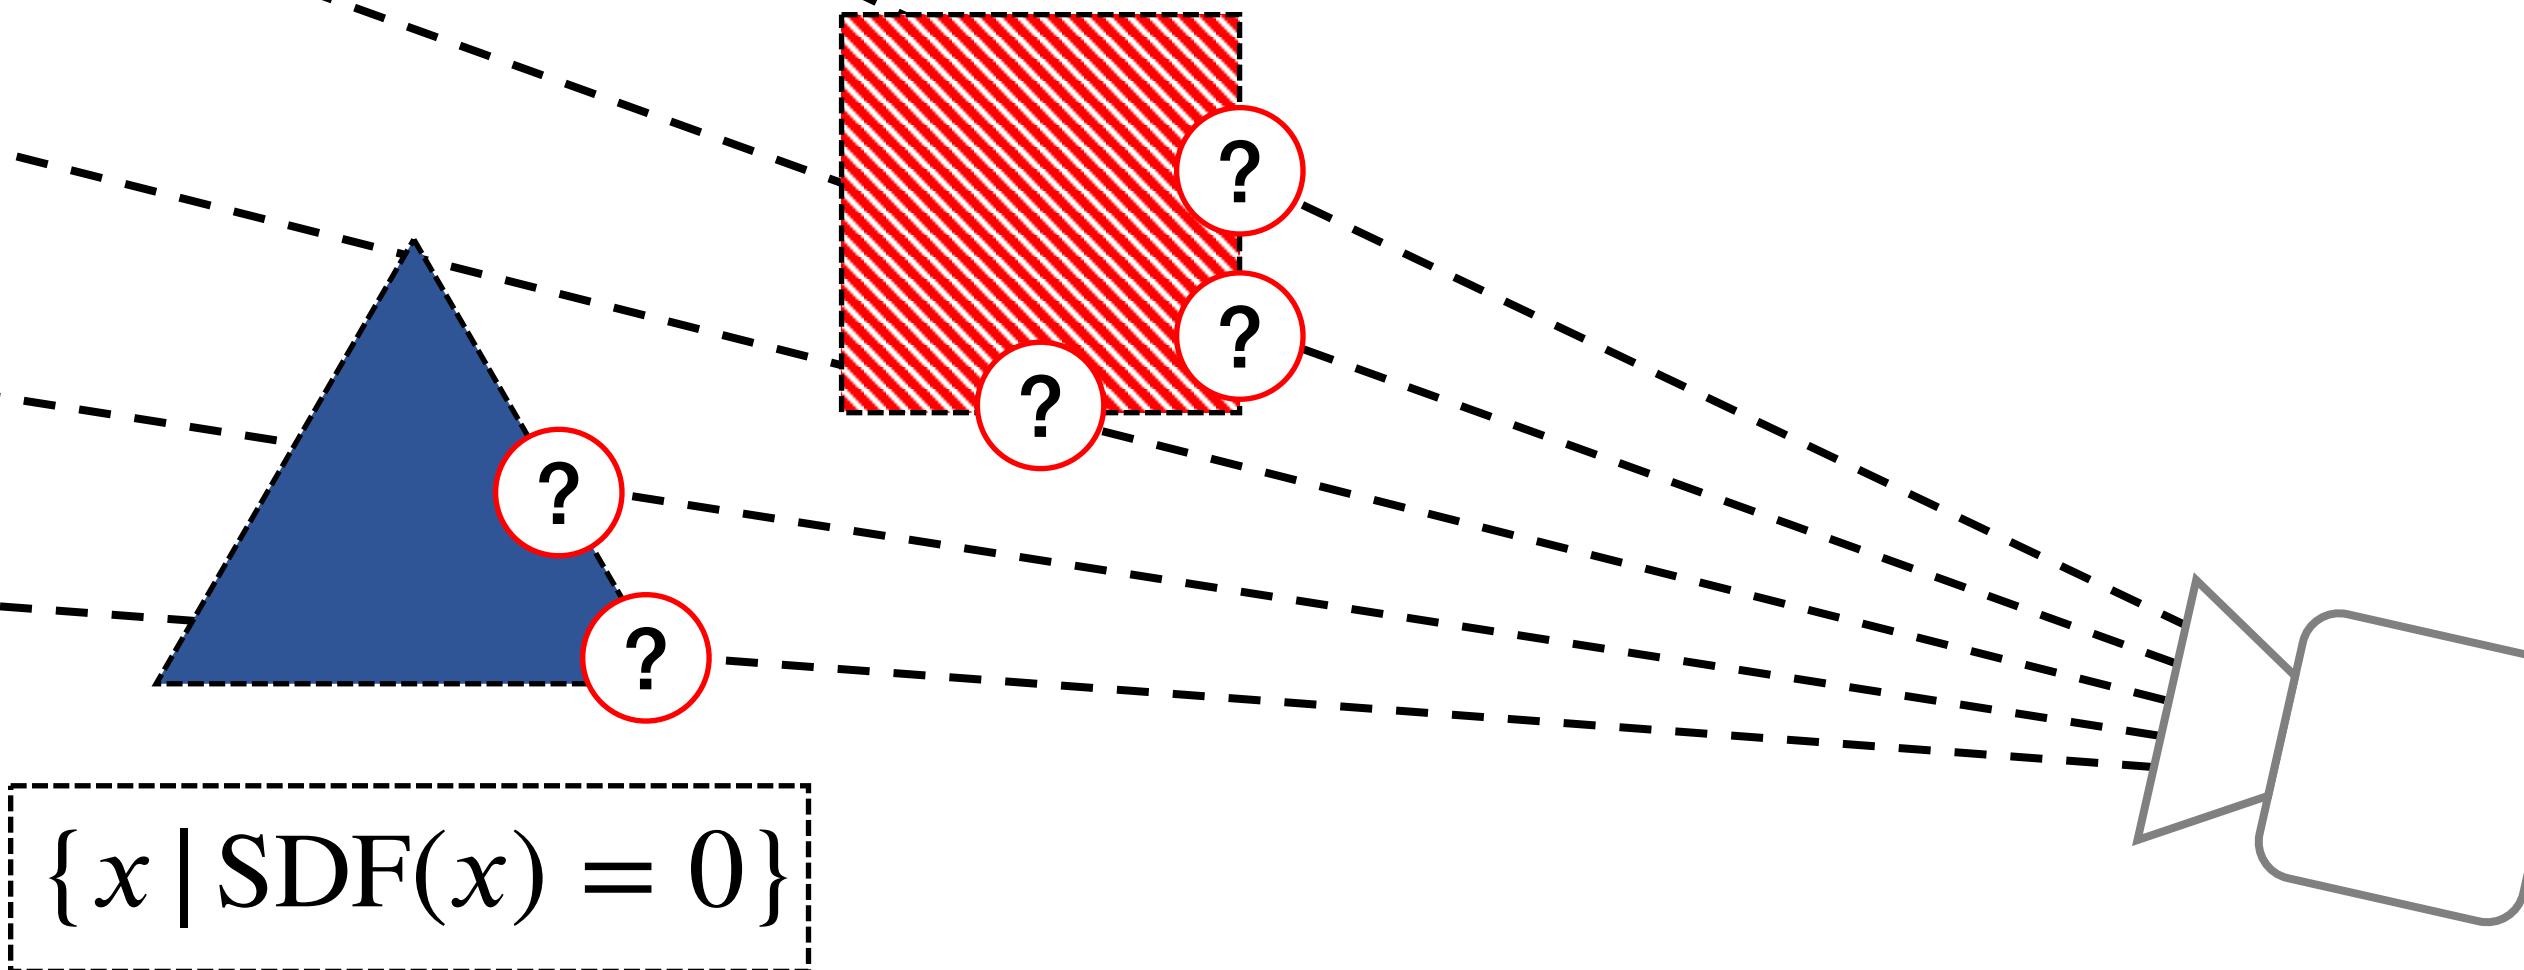
Sphere Tracing Step 1: Intersection Testing.

- ▶ How to find intersections?
- ▶ March along ray until we find zero-levelset.

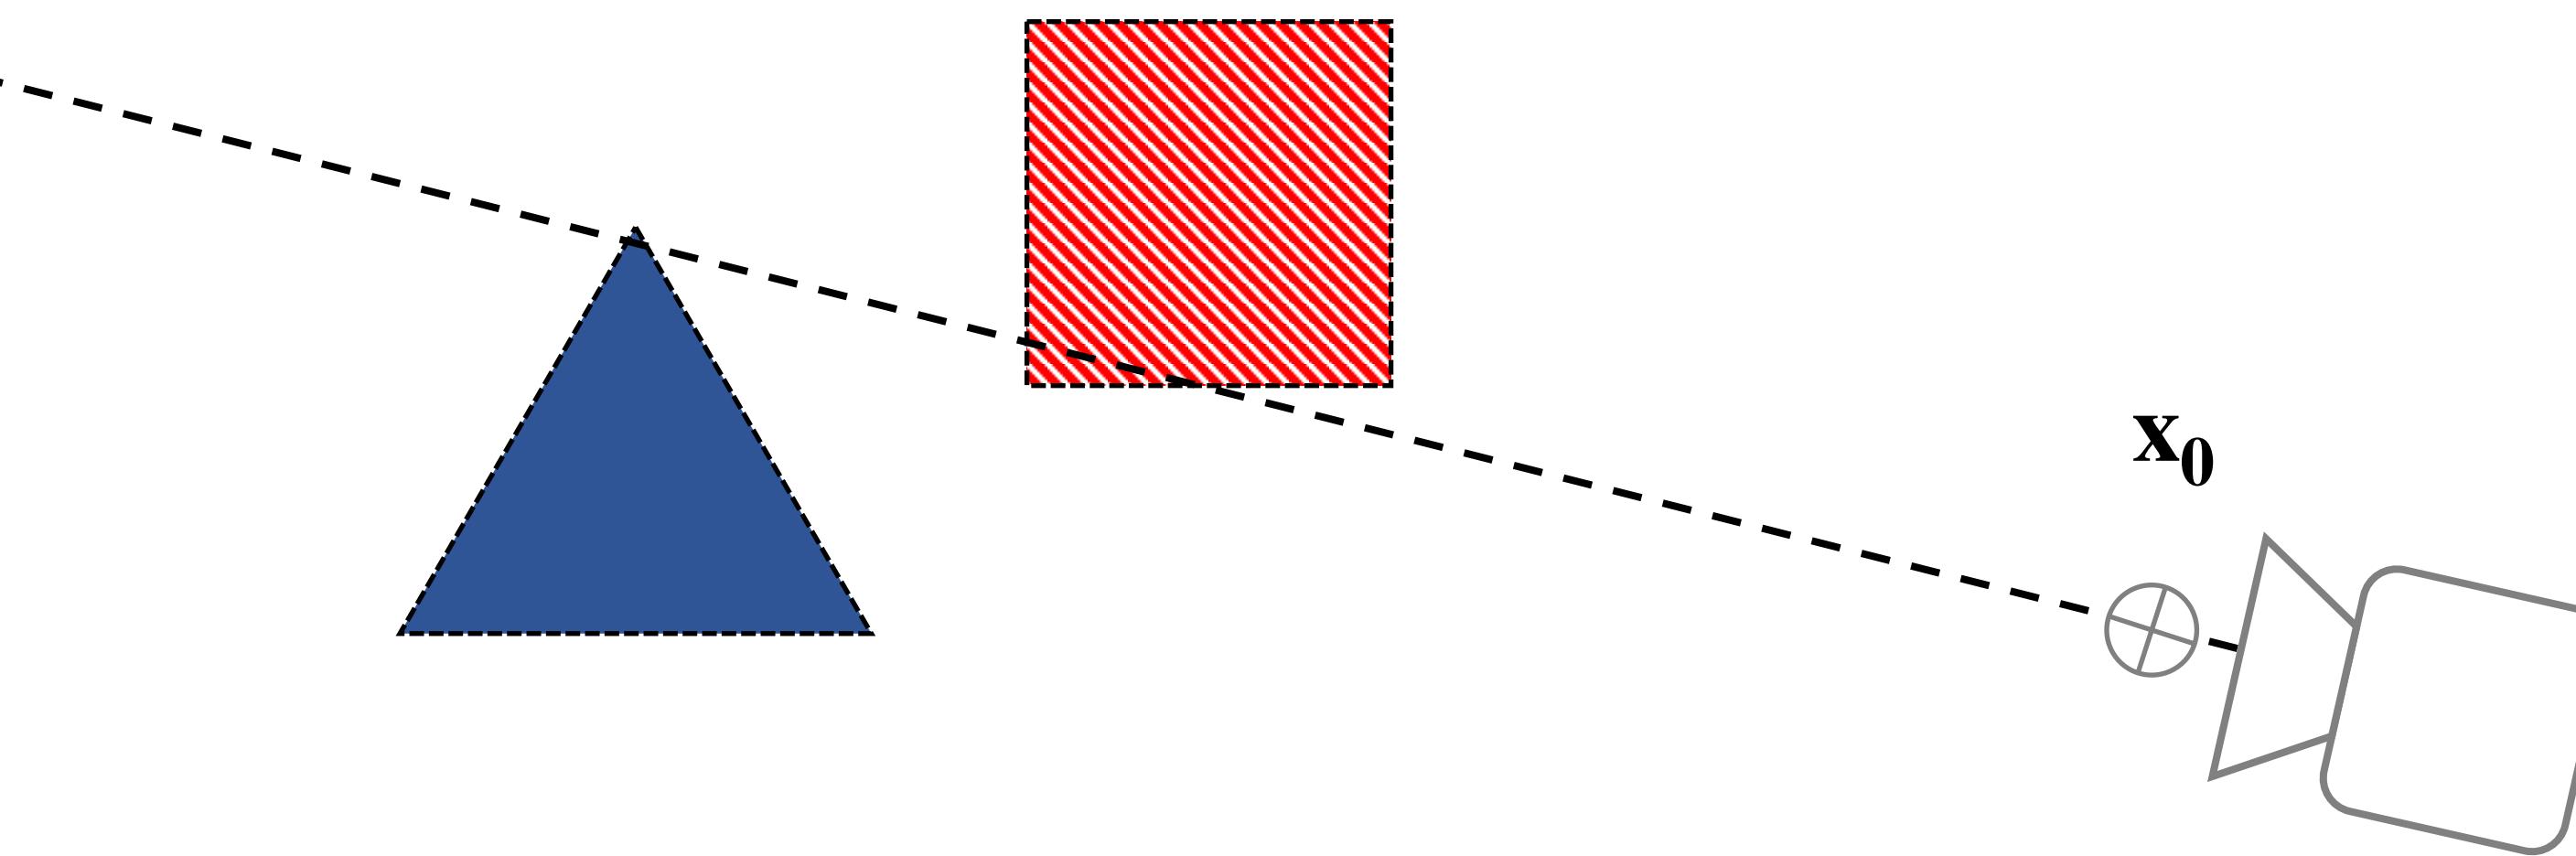


Sphere Tracing Step 1: Intersection Testing.

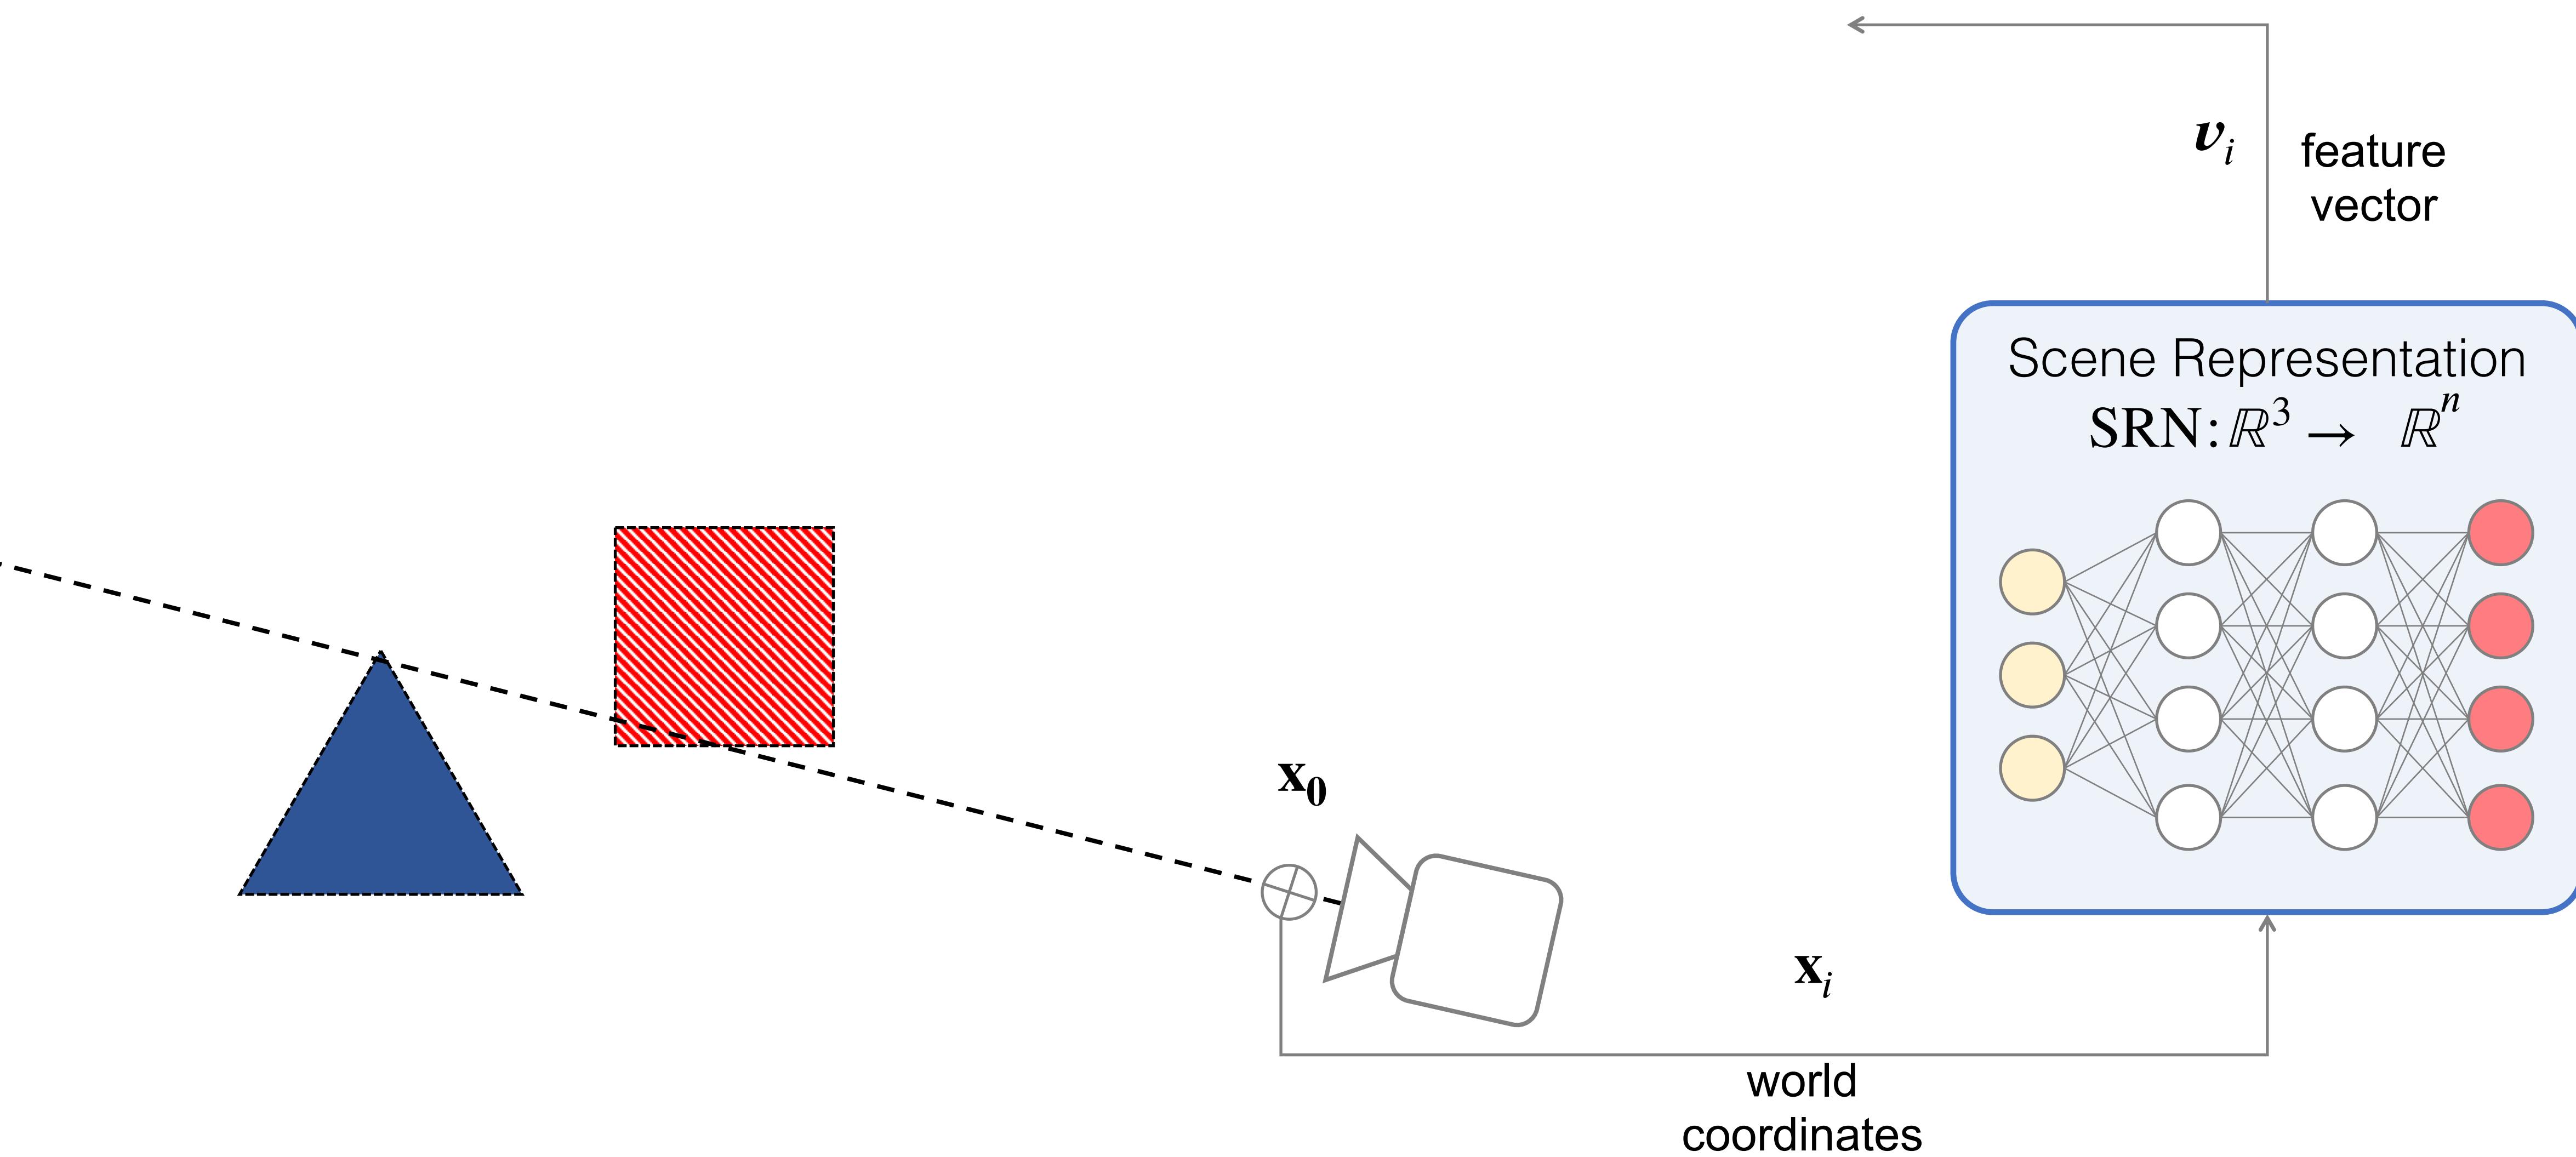
- ▶ How to find intersections?
- ▶ March along ray until we find zero-levelset.
- ▶ Naïve solution: constant steplength. Too expensive!



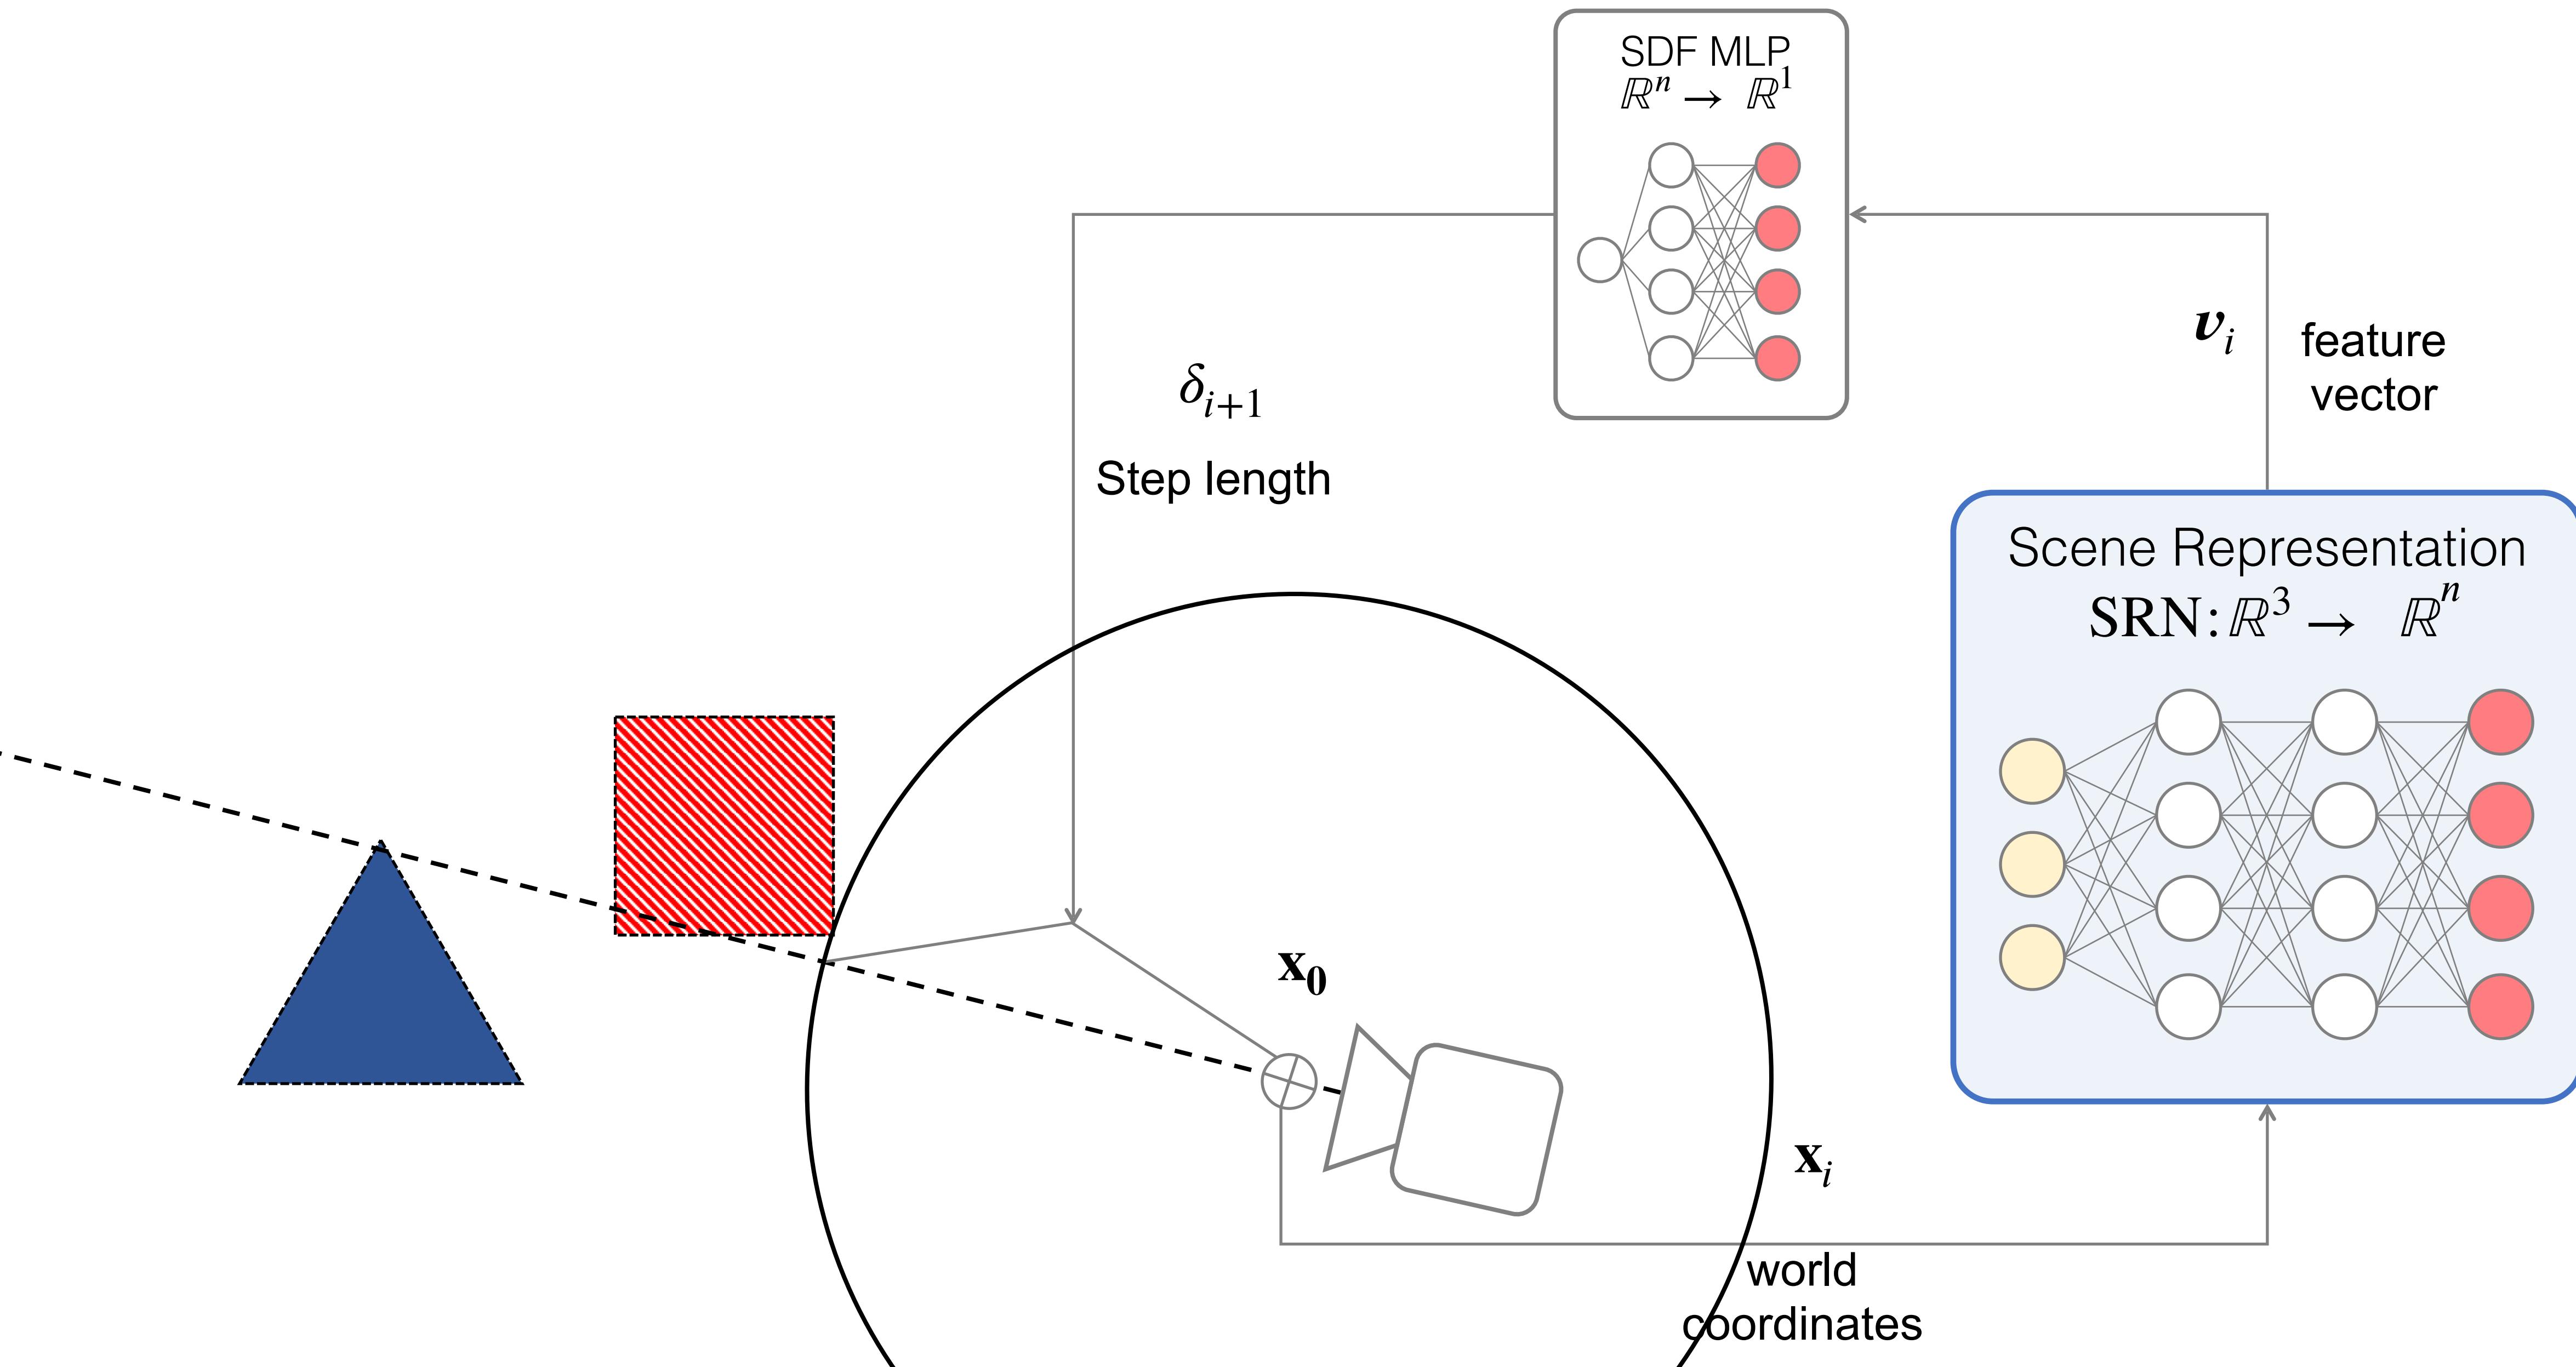
Sphere Tracing Step 1: Intersection Testing.



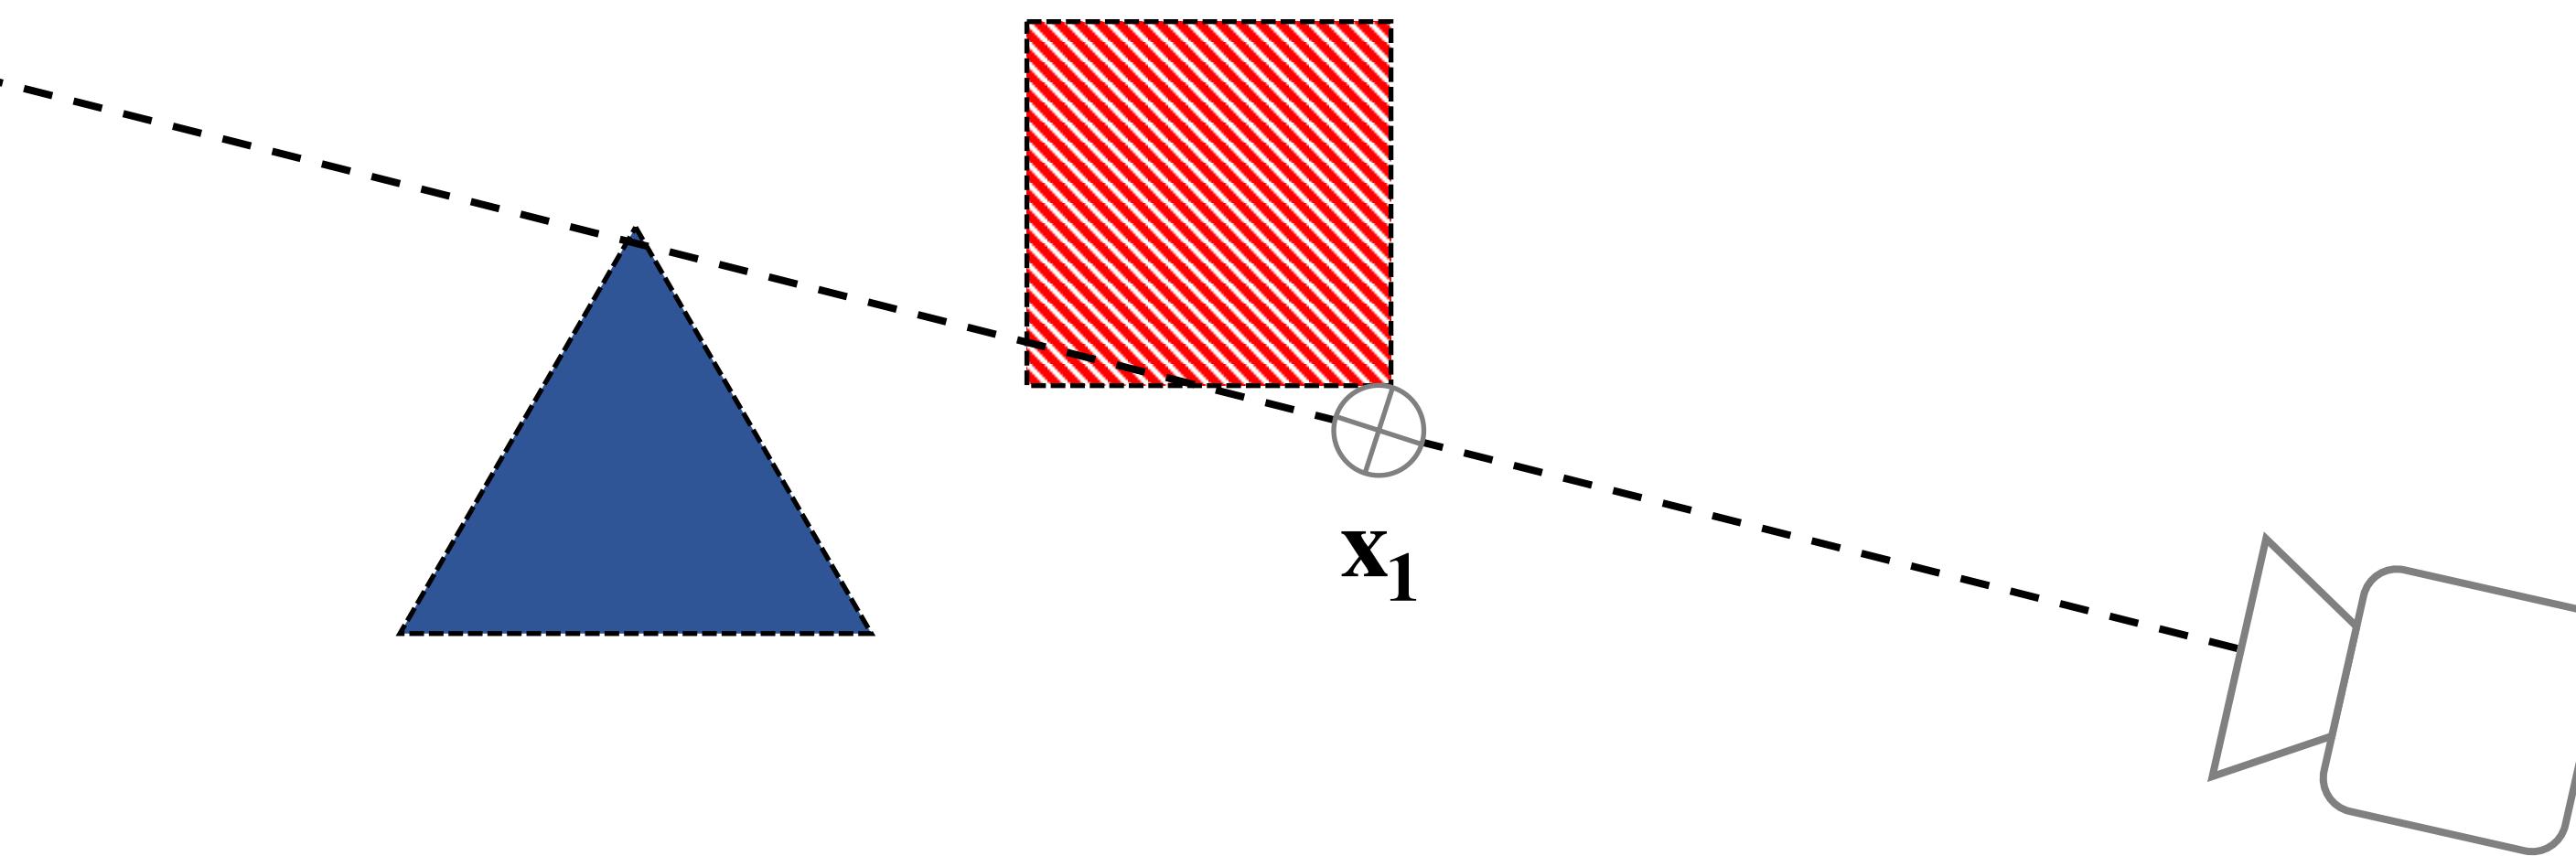
Sphere Tracing Step 1: Intersection Testing.



Sphere Tracing Step 1: Intersection Testing.

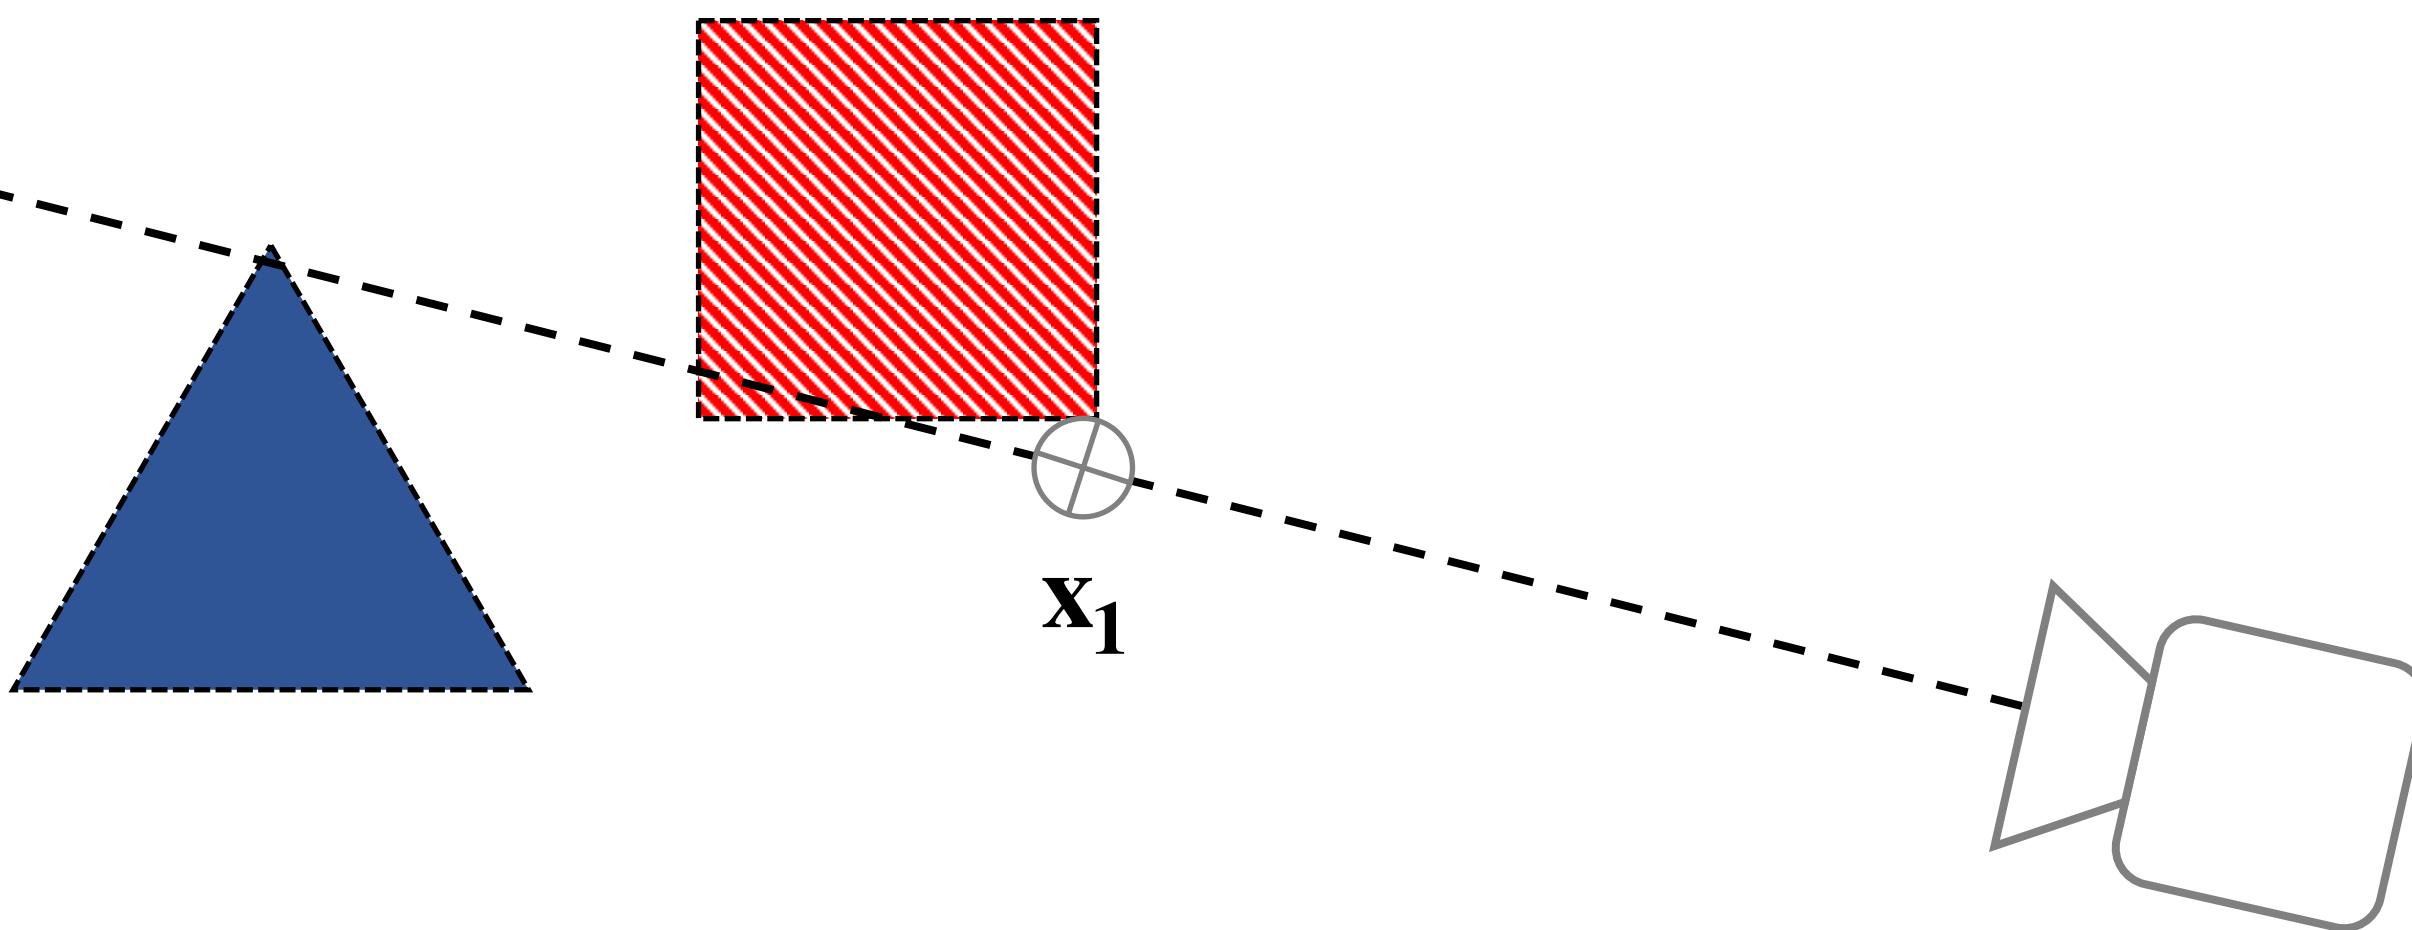


Sphere Tracing Step 1: Intersection Testing.



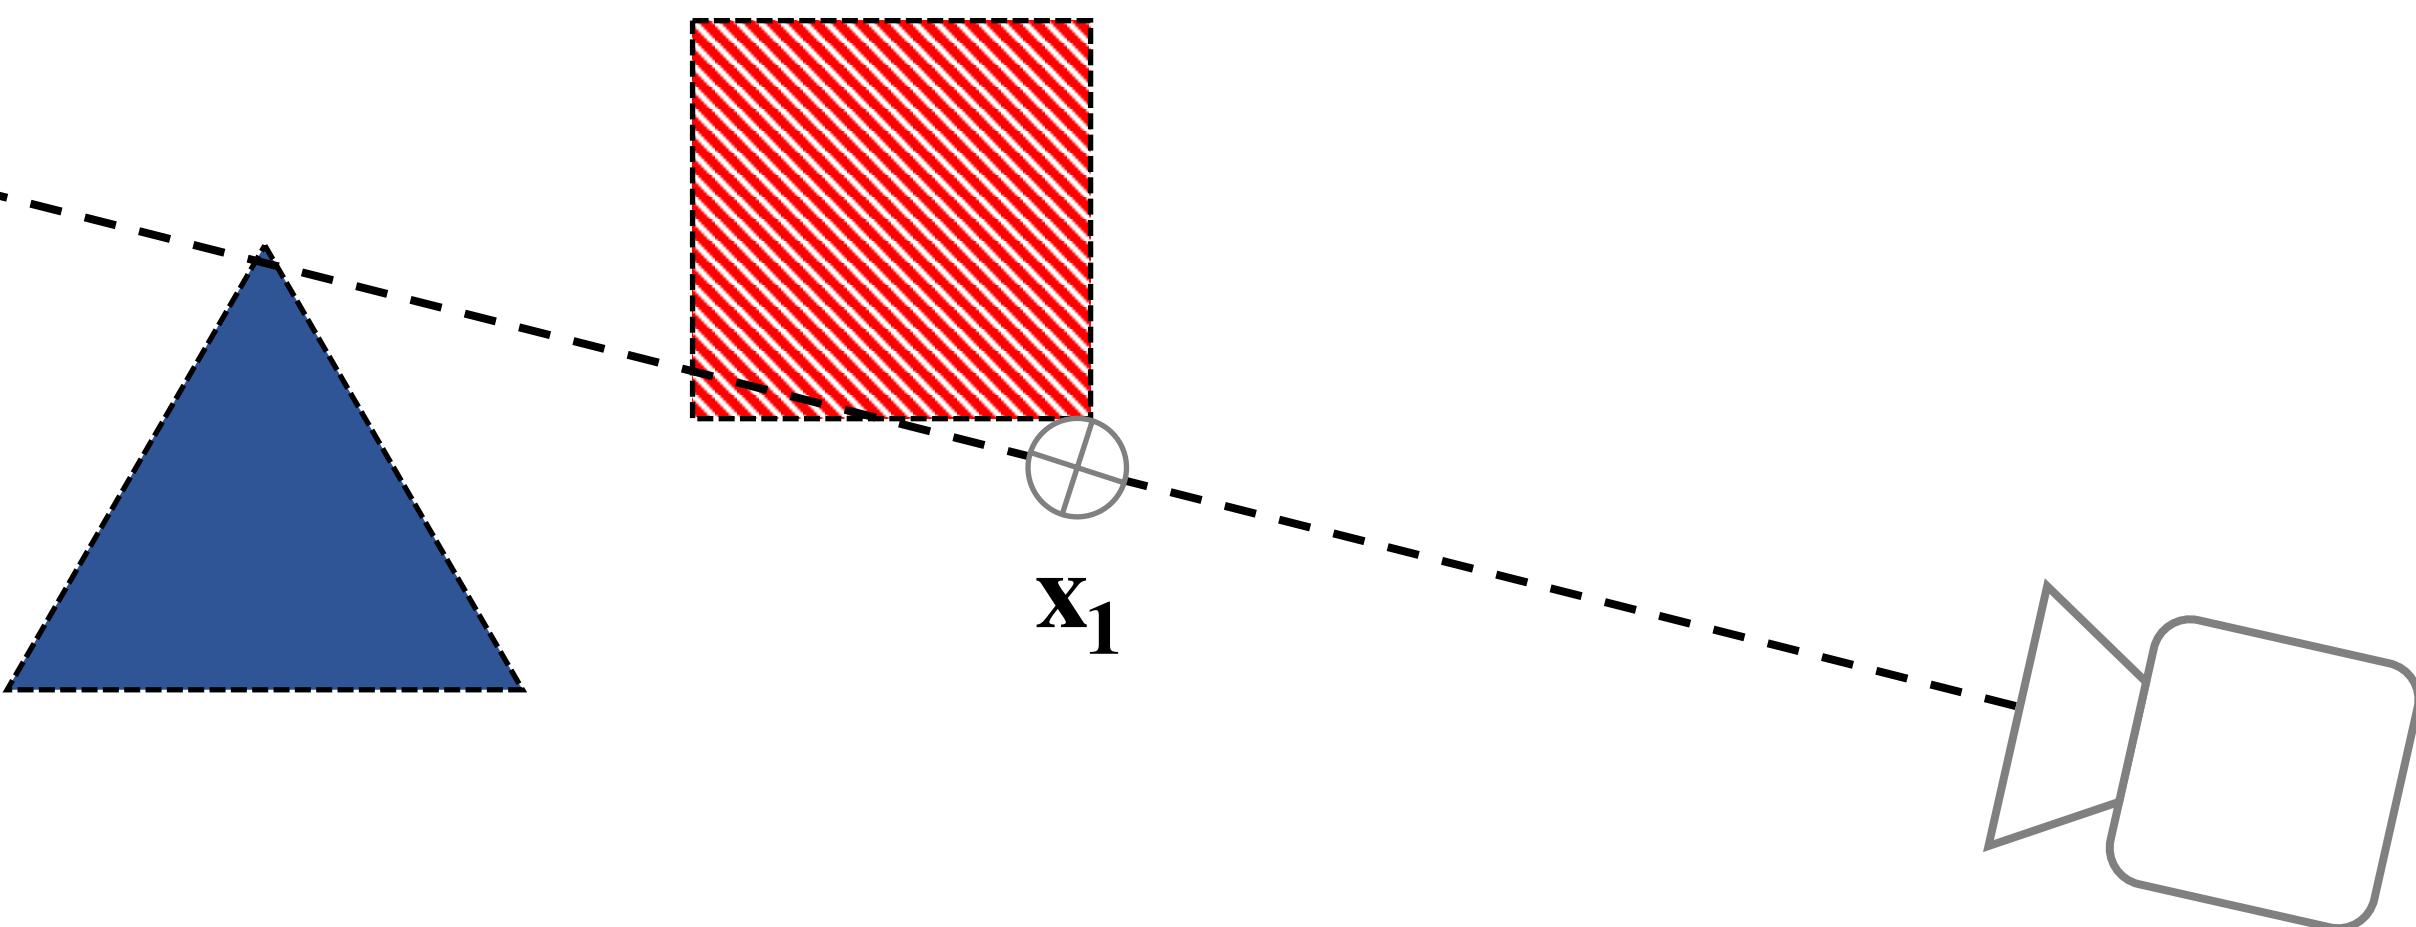
Sphere Tracing Step 1: Intersection Testing.

- Guaranteed to converge to surface.



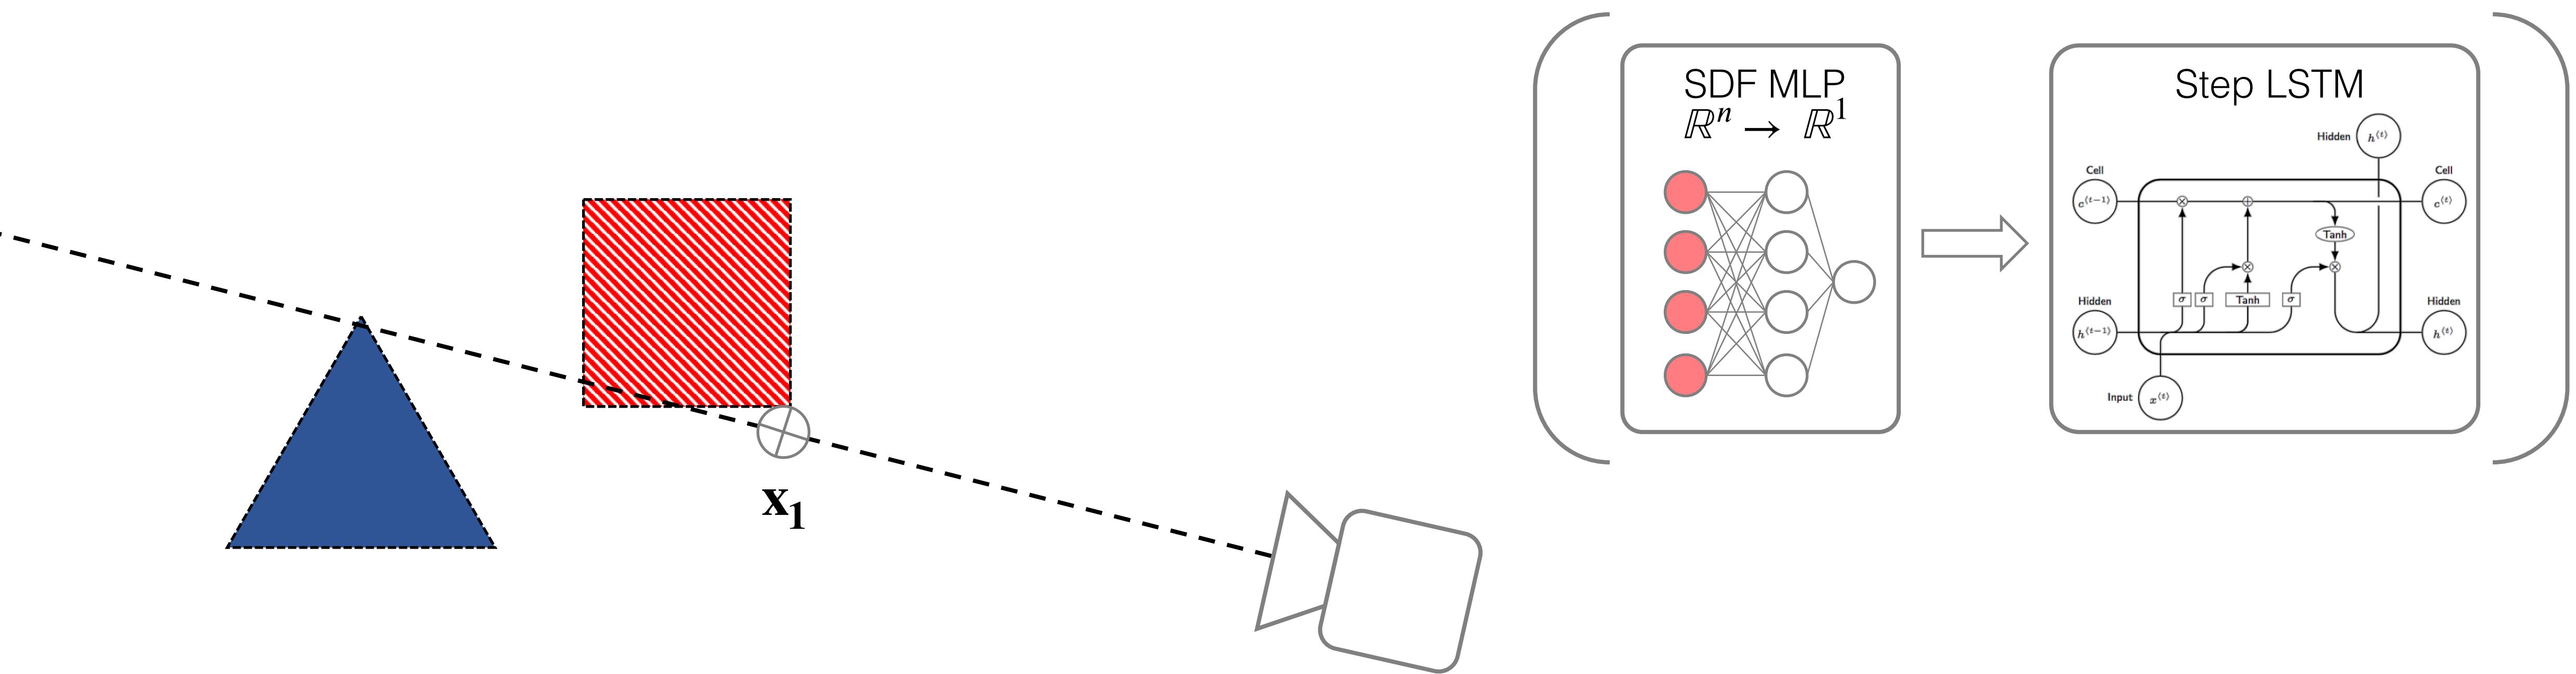
Sphere Tracing Step 1: Intersection Testing.

- Guaranteed to converge to surface.
- Guaranteed to not step *through* a surface.

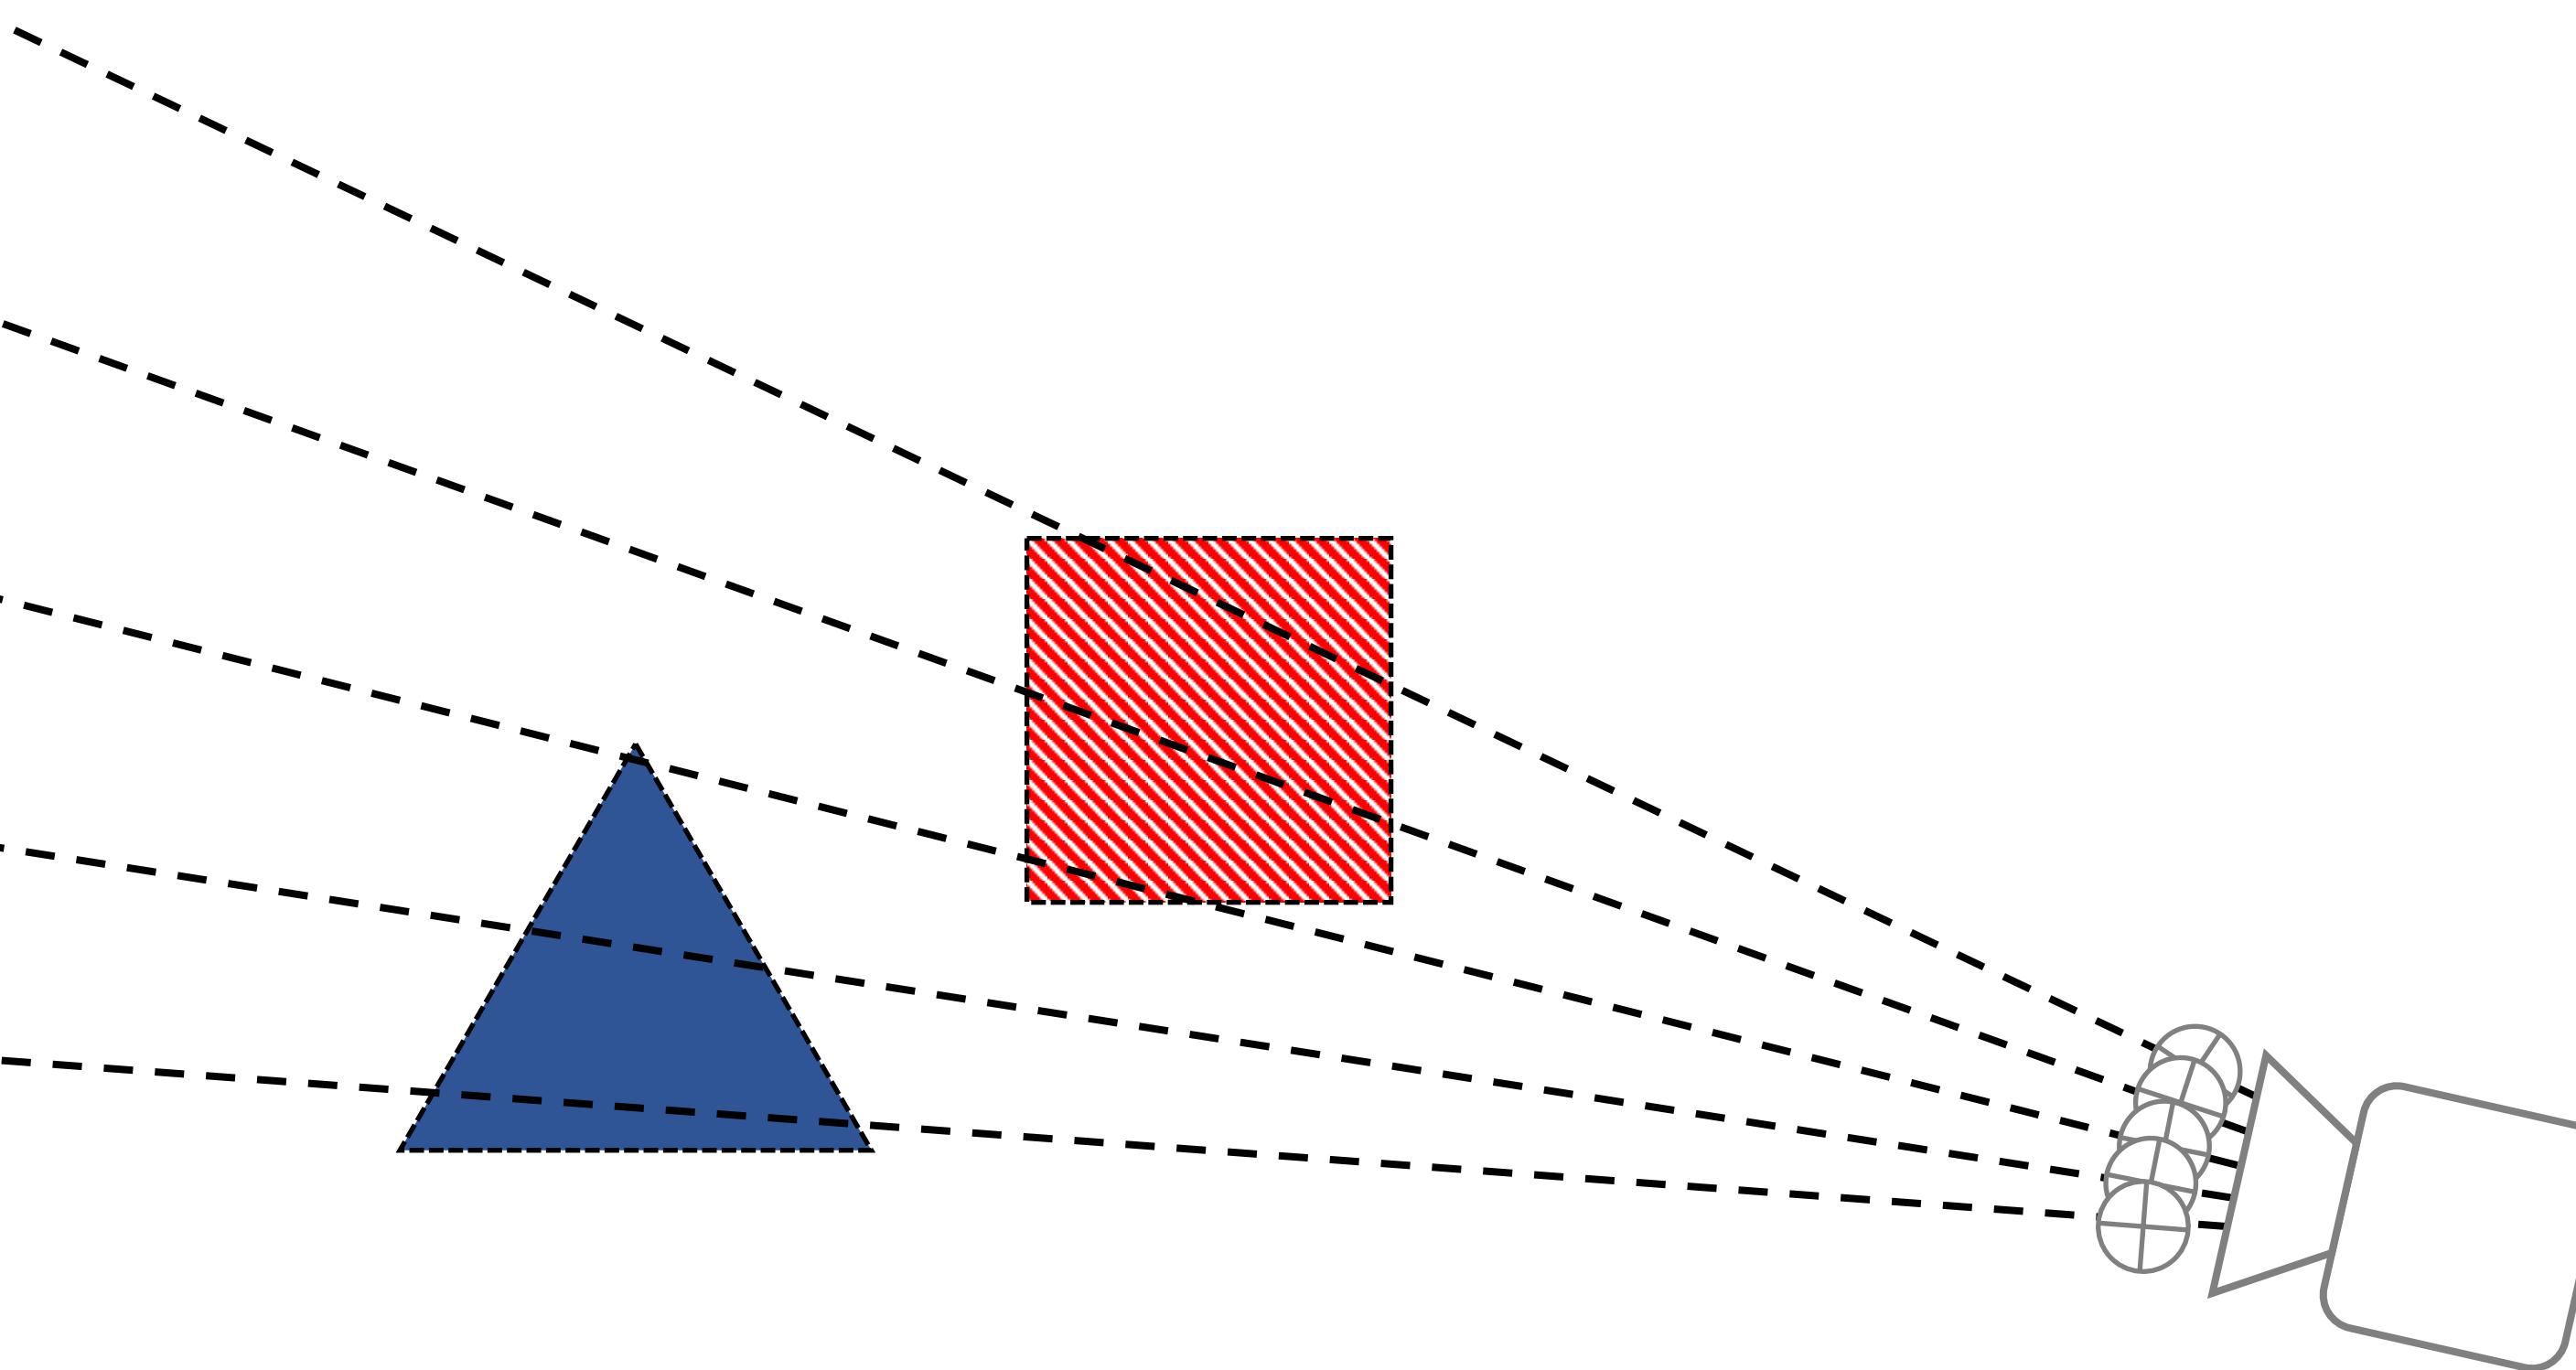


Sphere Tracing Step 1: Intersection Testing.

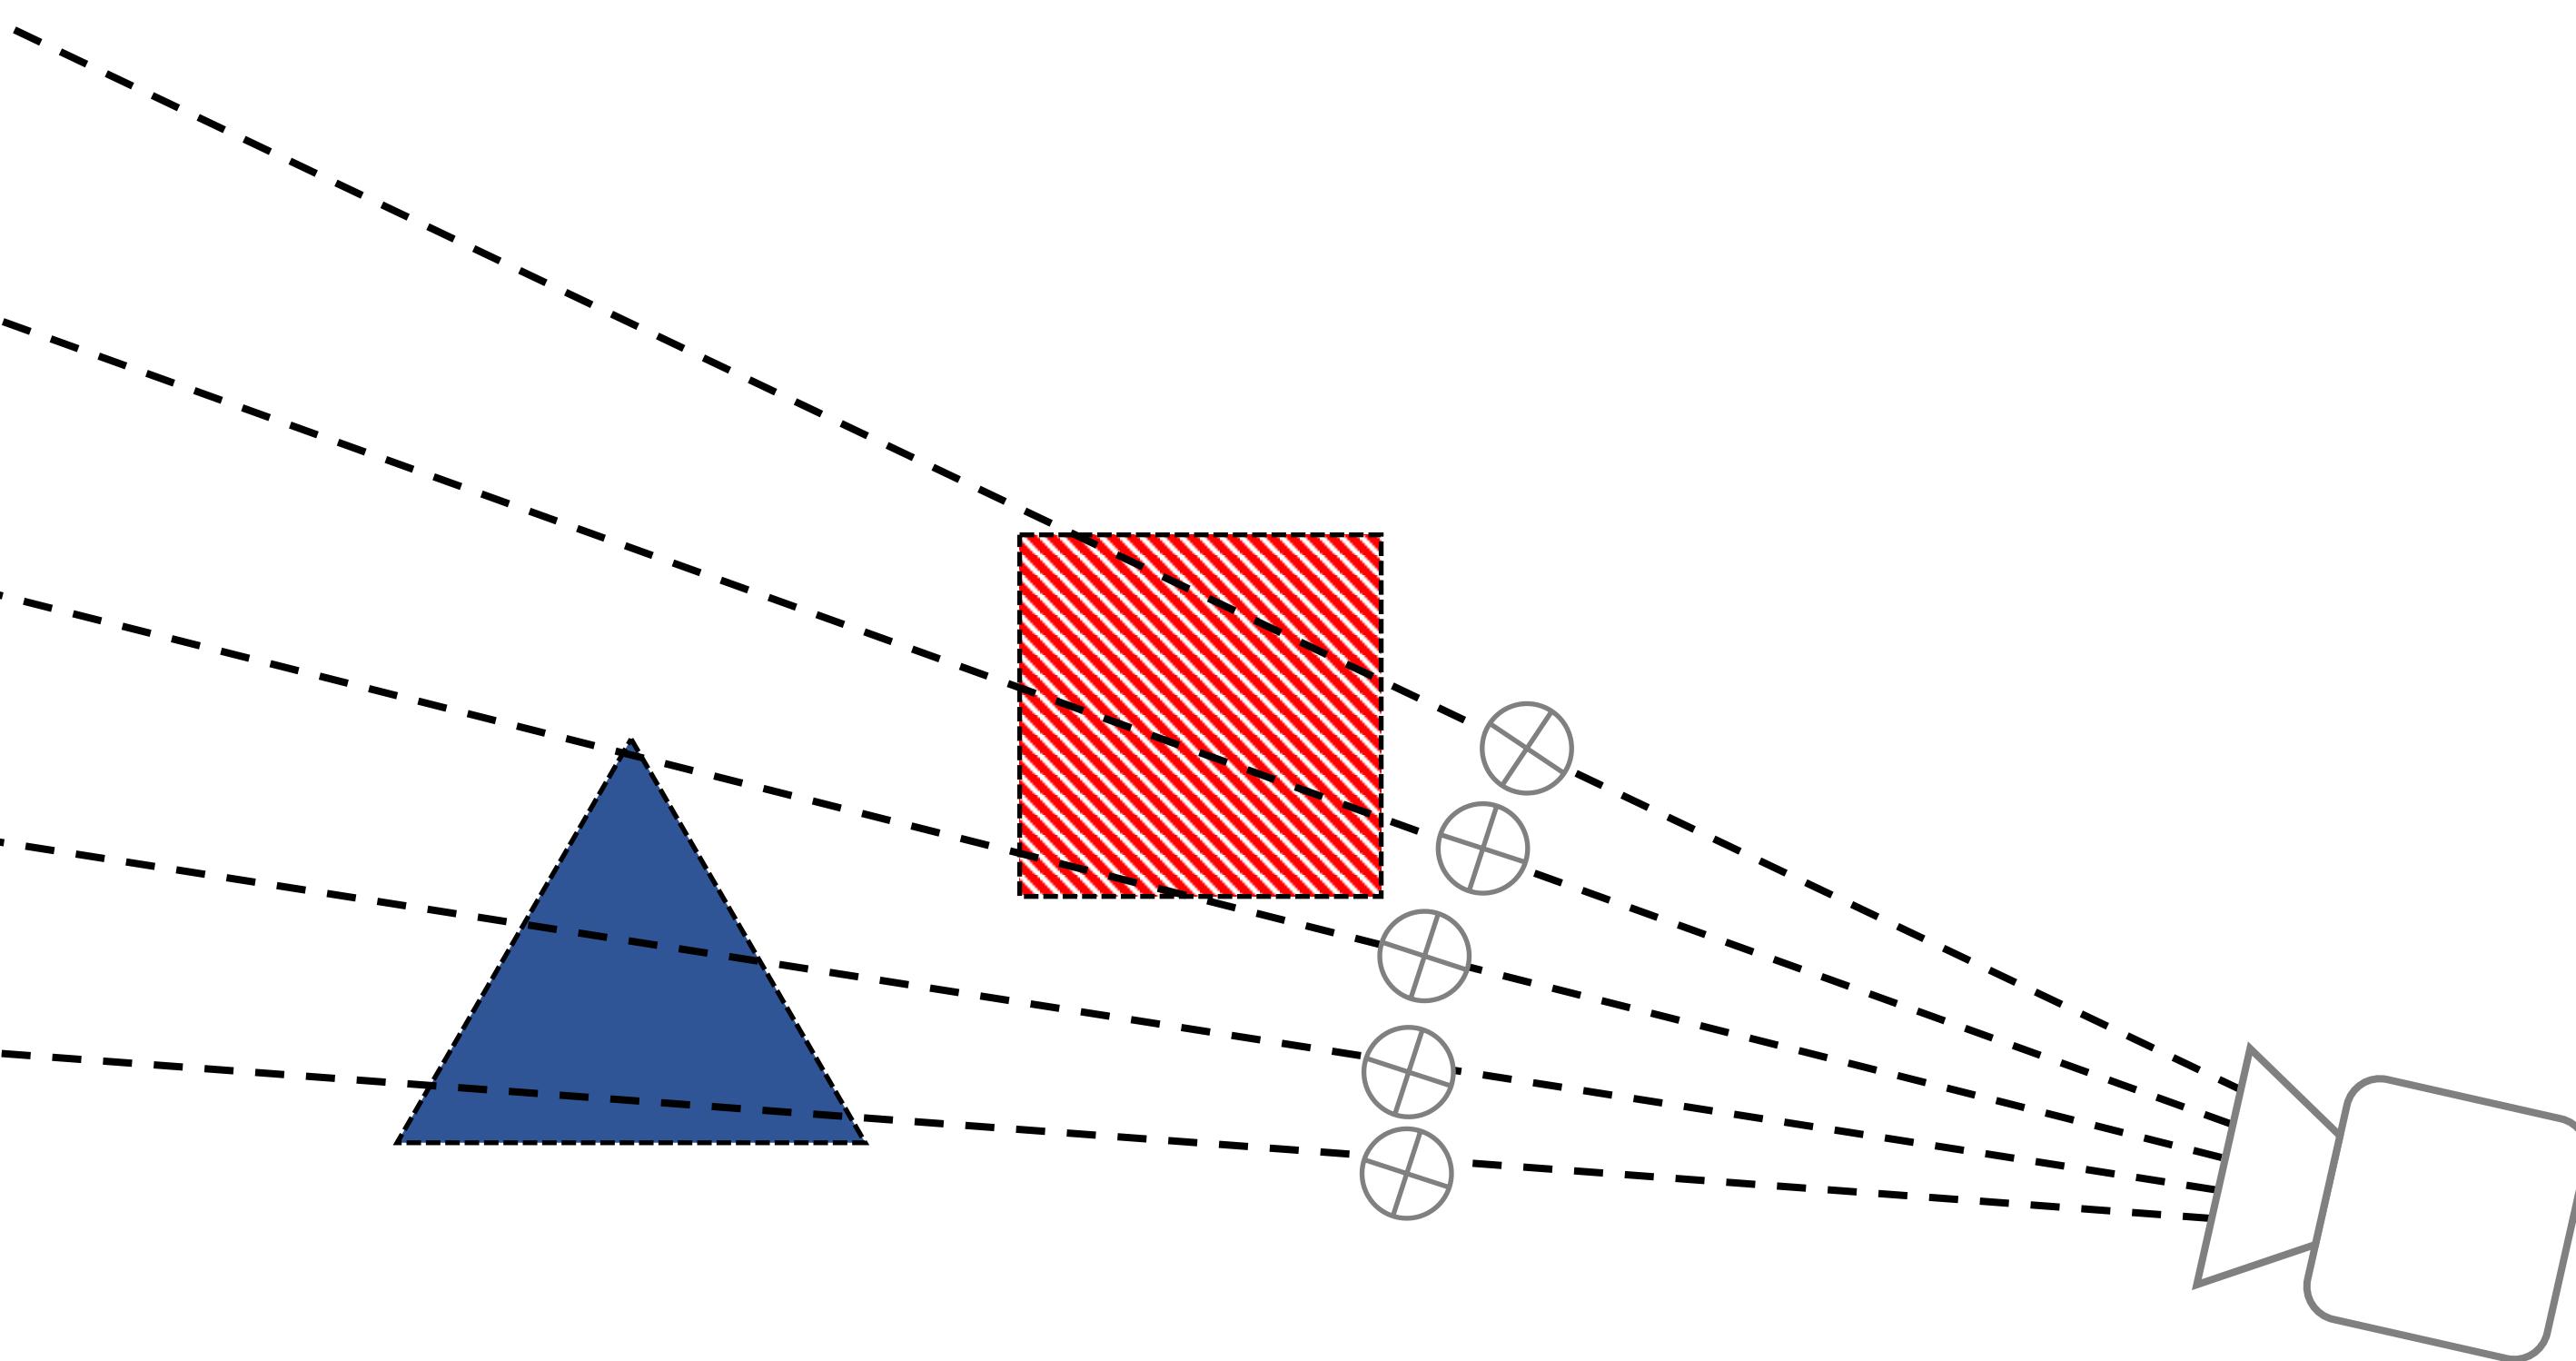
- Guaranteed to converge to surface.
- Guaranteed to not step *through* a surface.



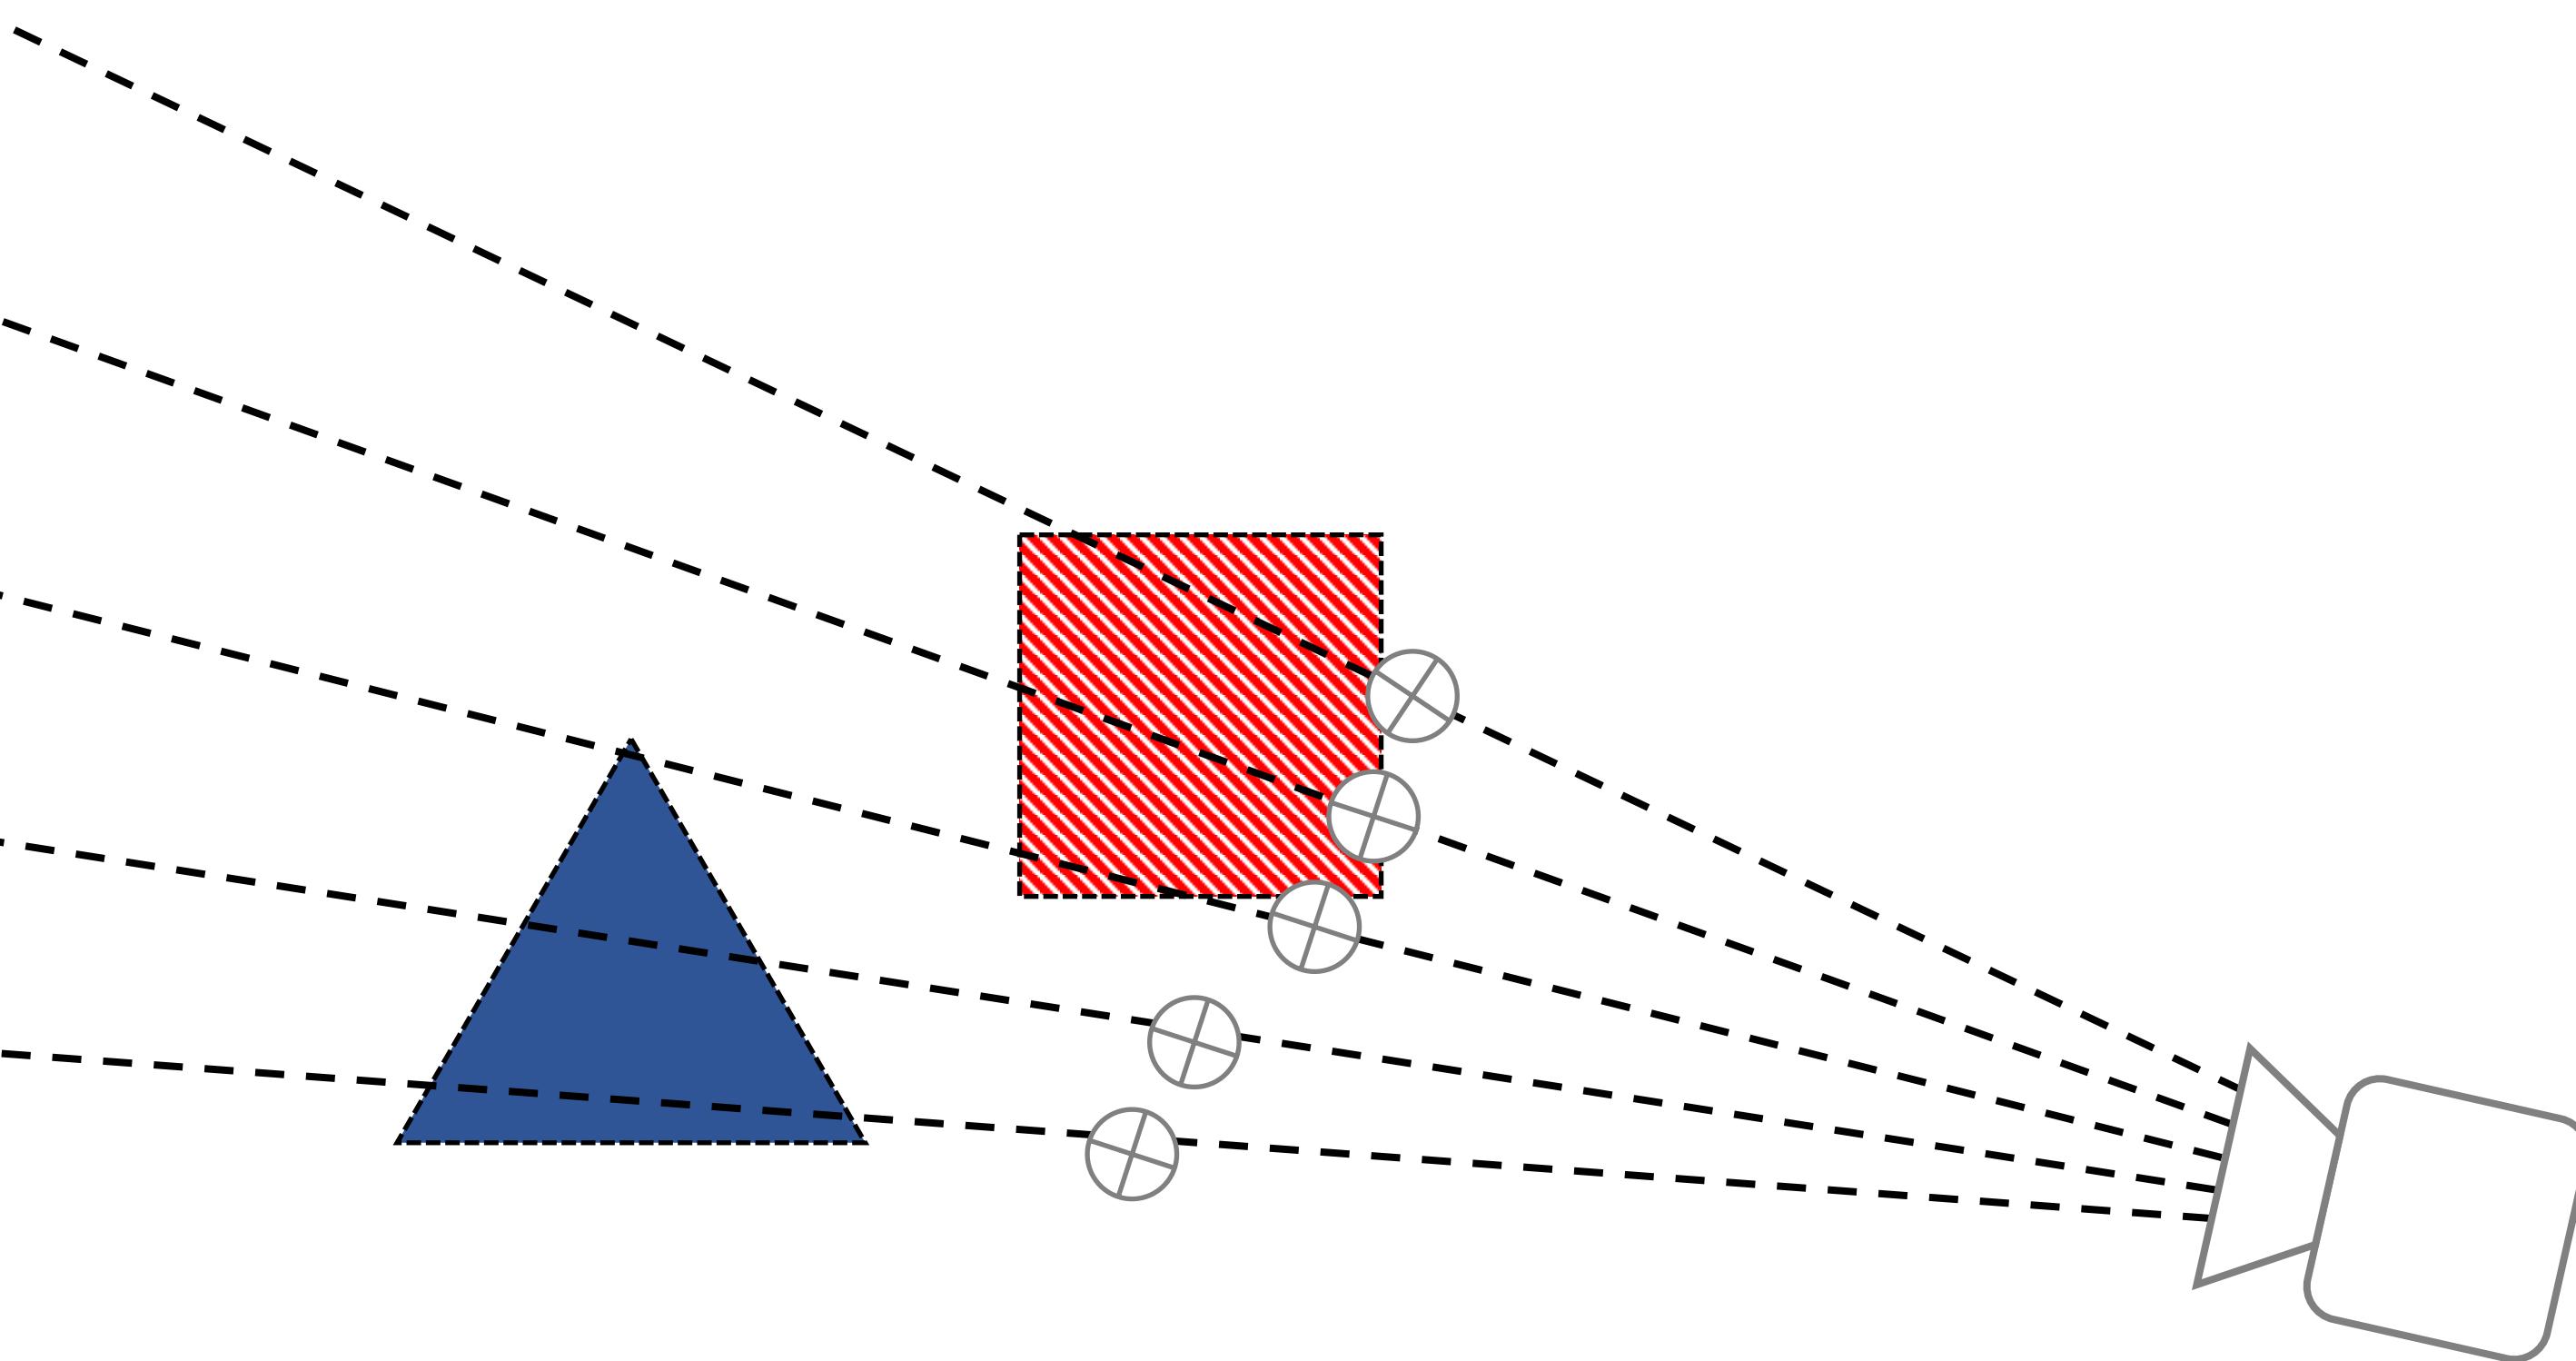
Sphere Tracing Step 1: Intersection Testing.



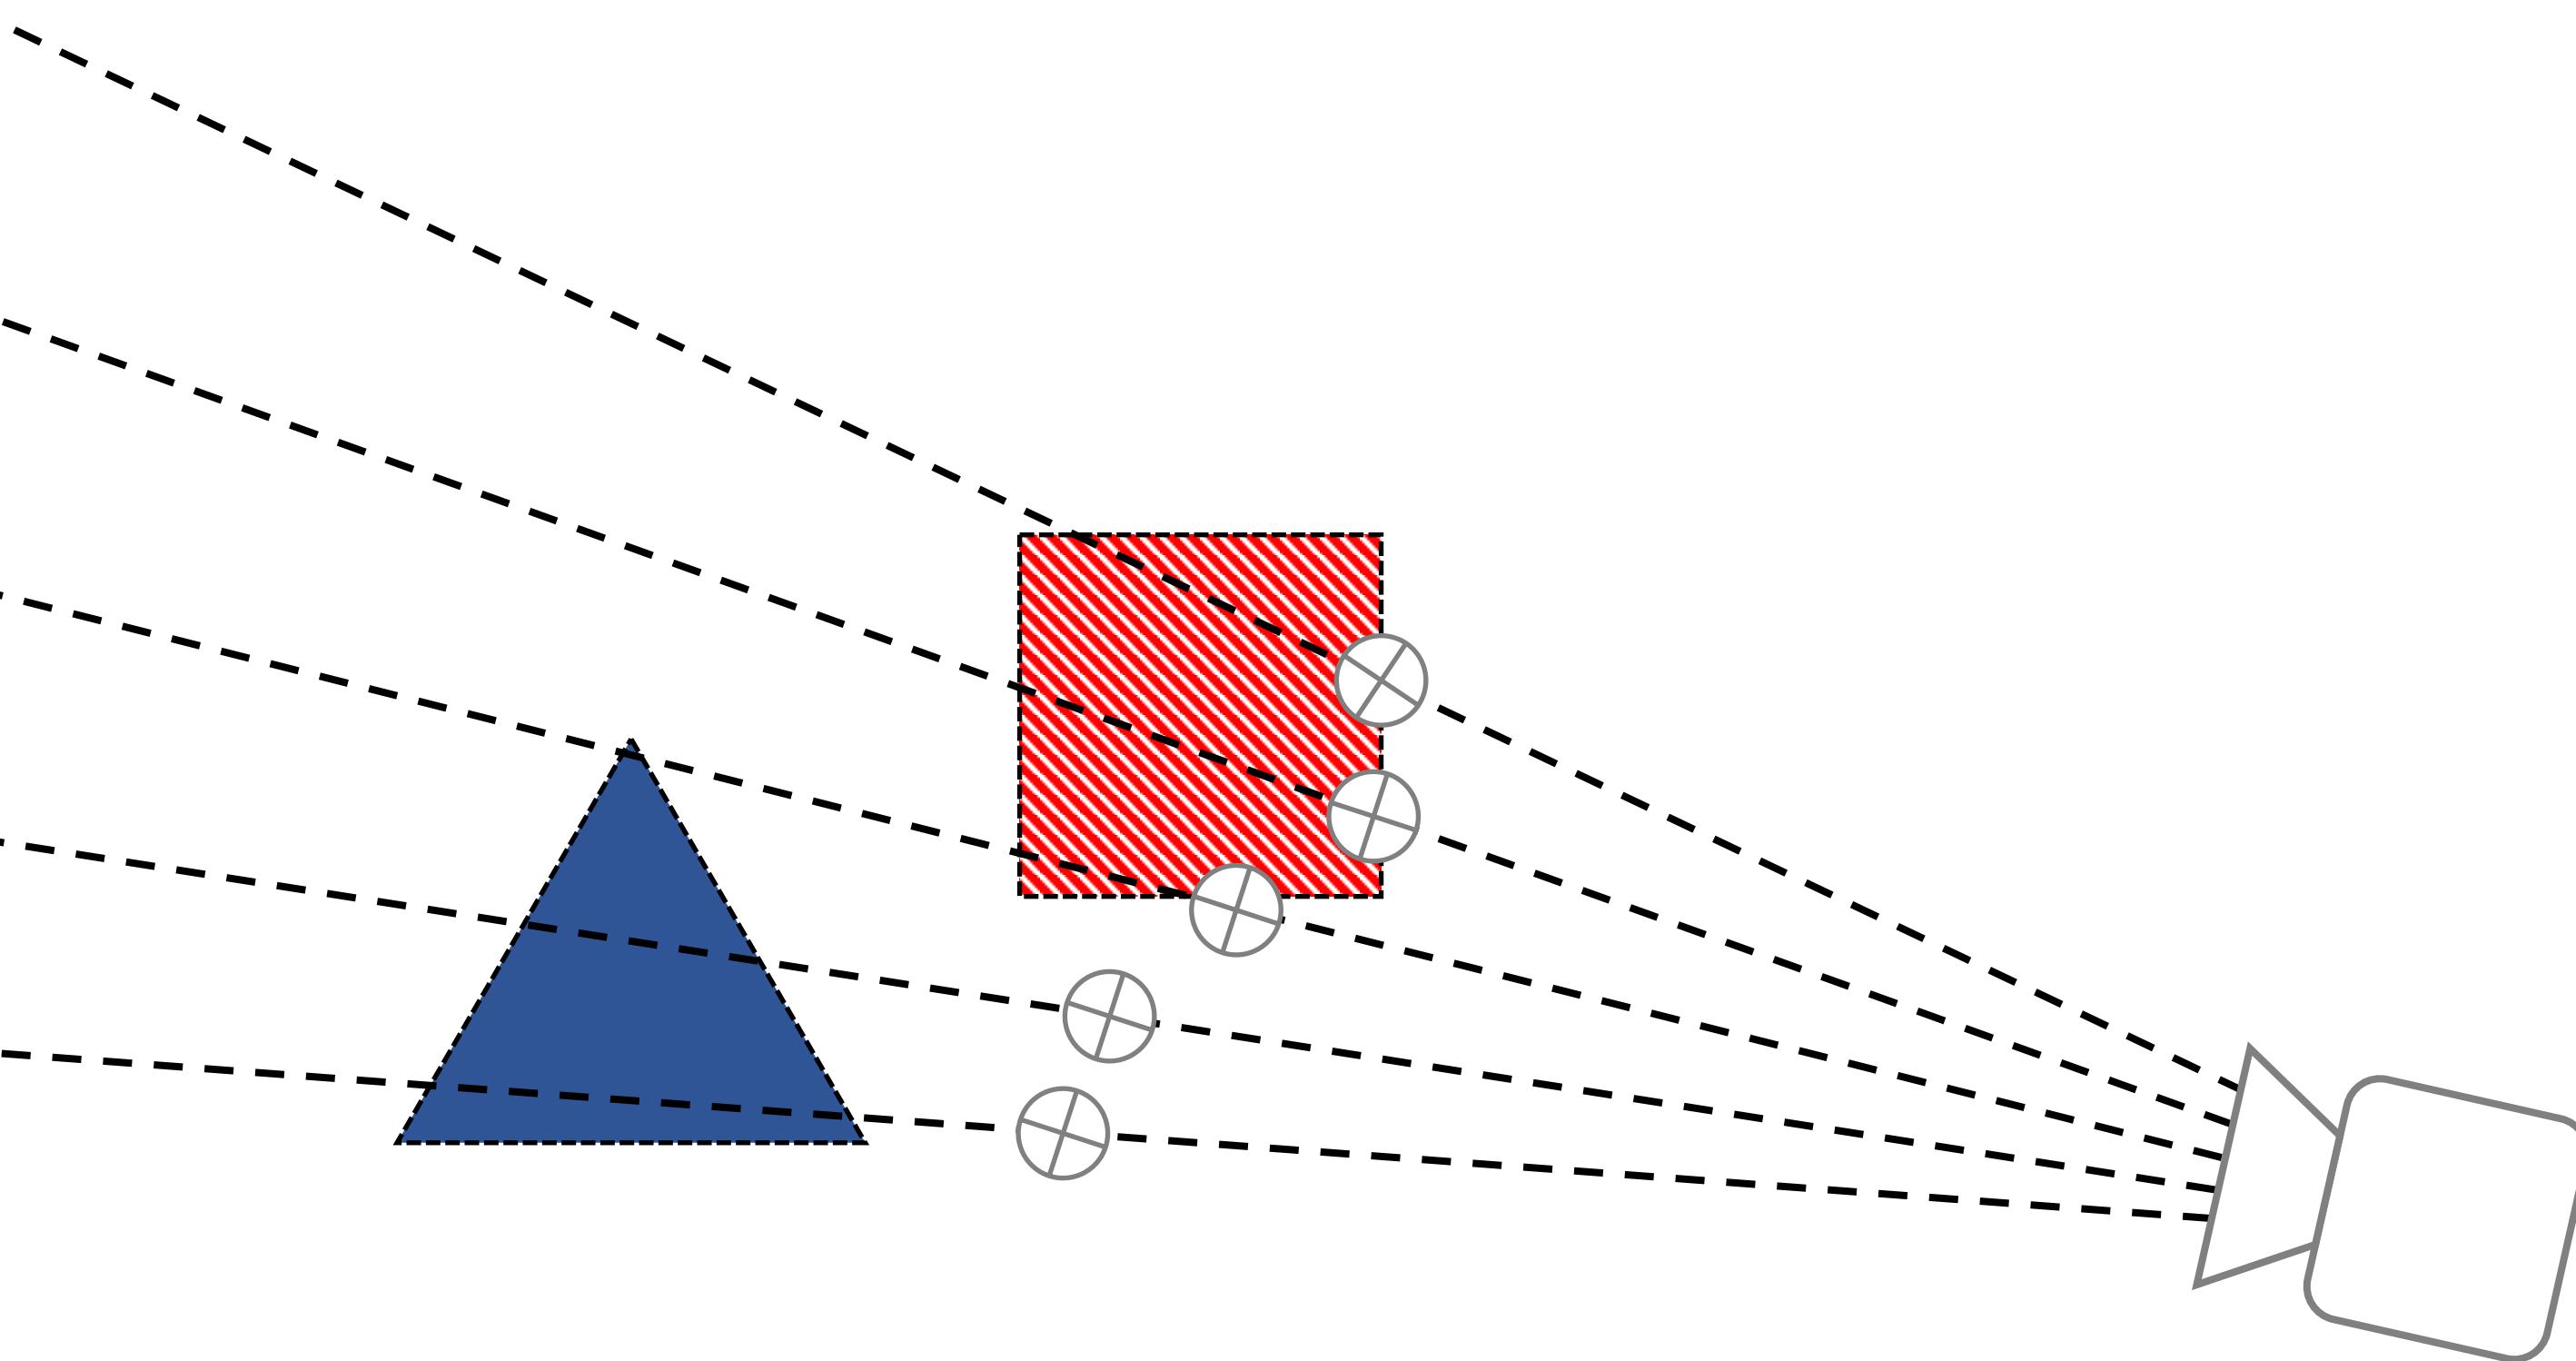
Sphere Tracing Step 1: Intersection Testing.



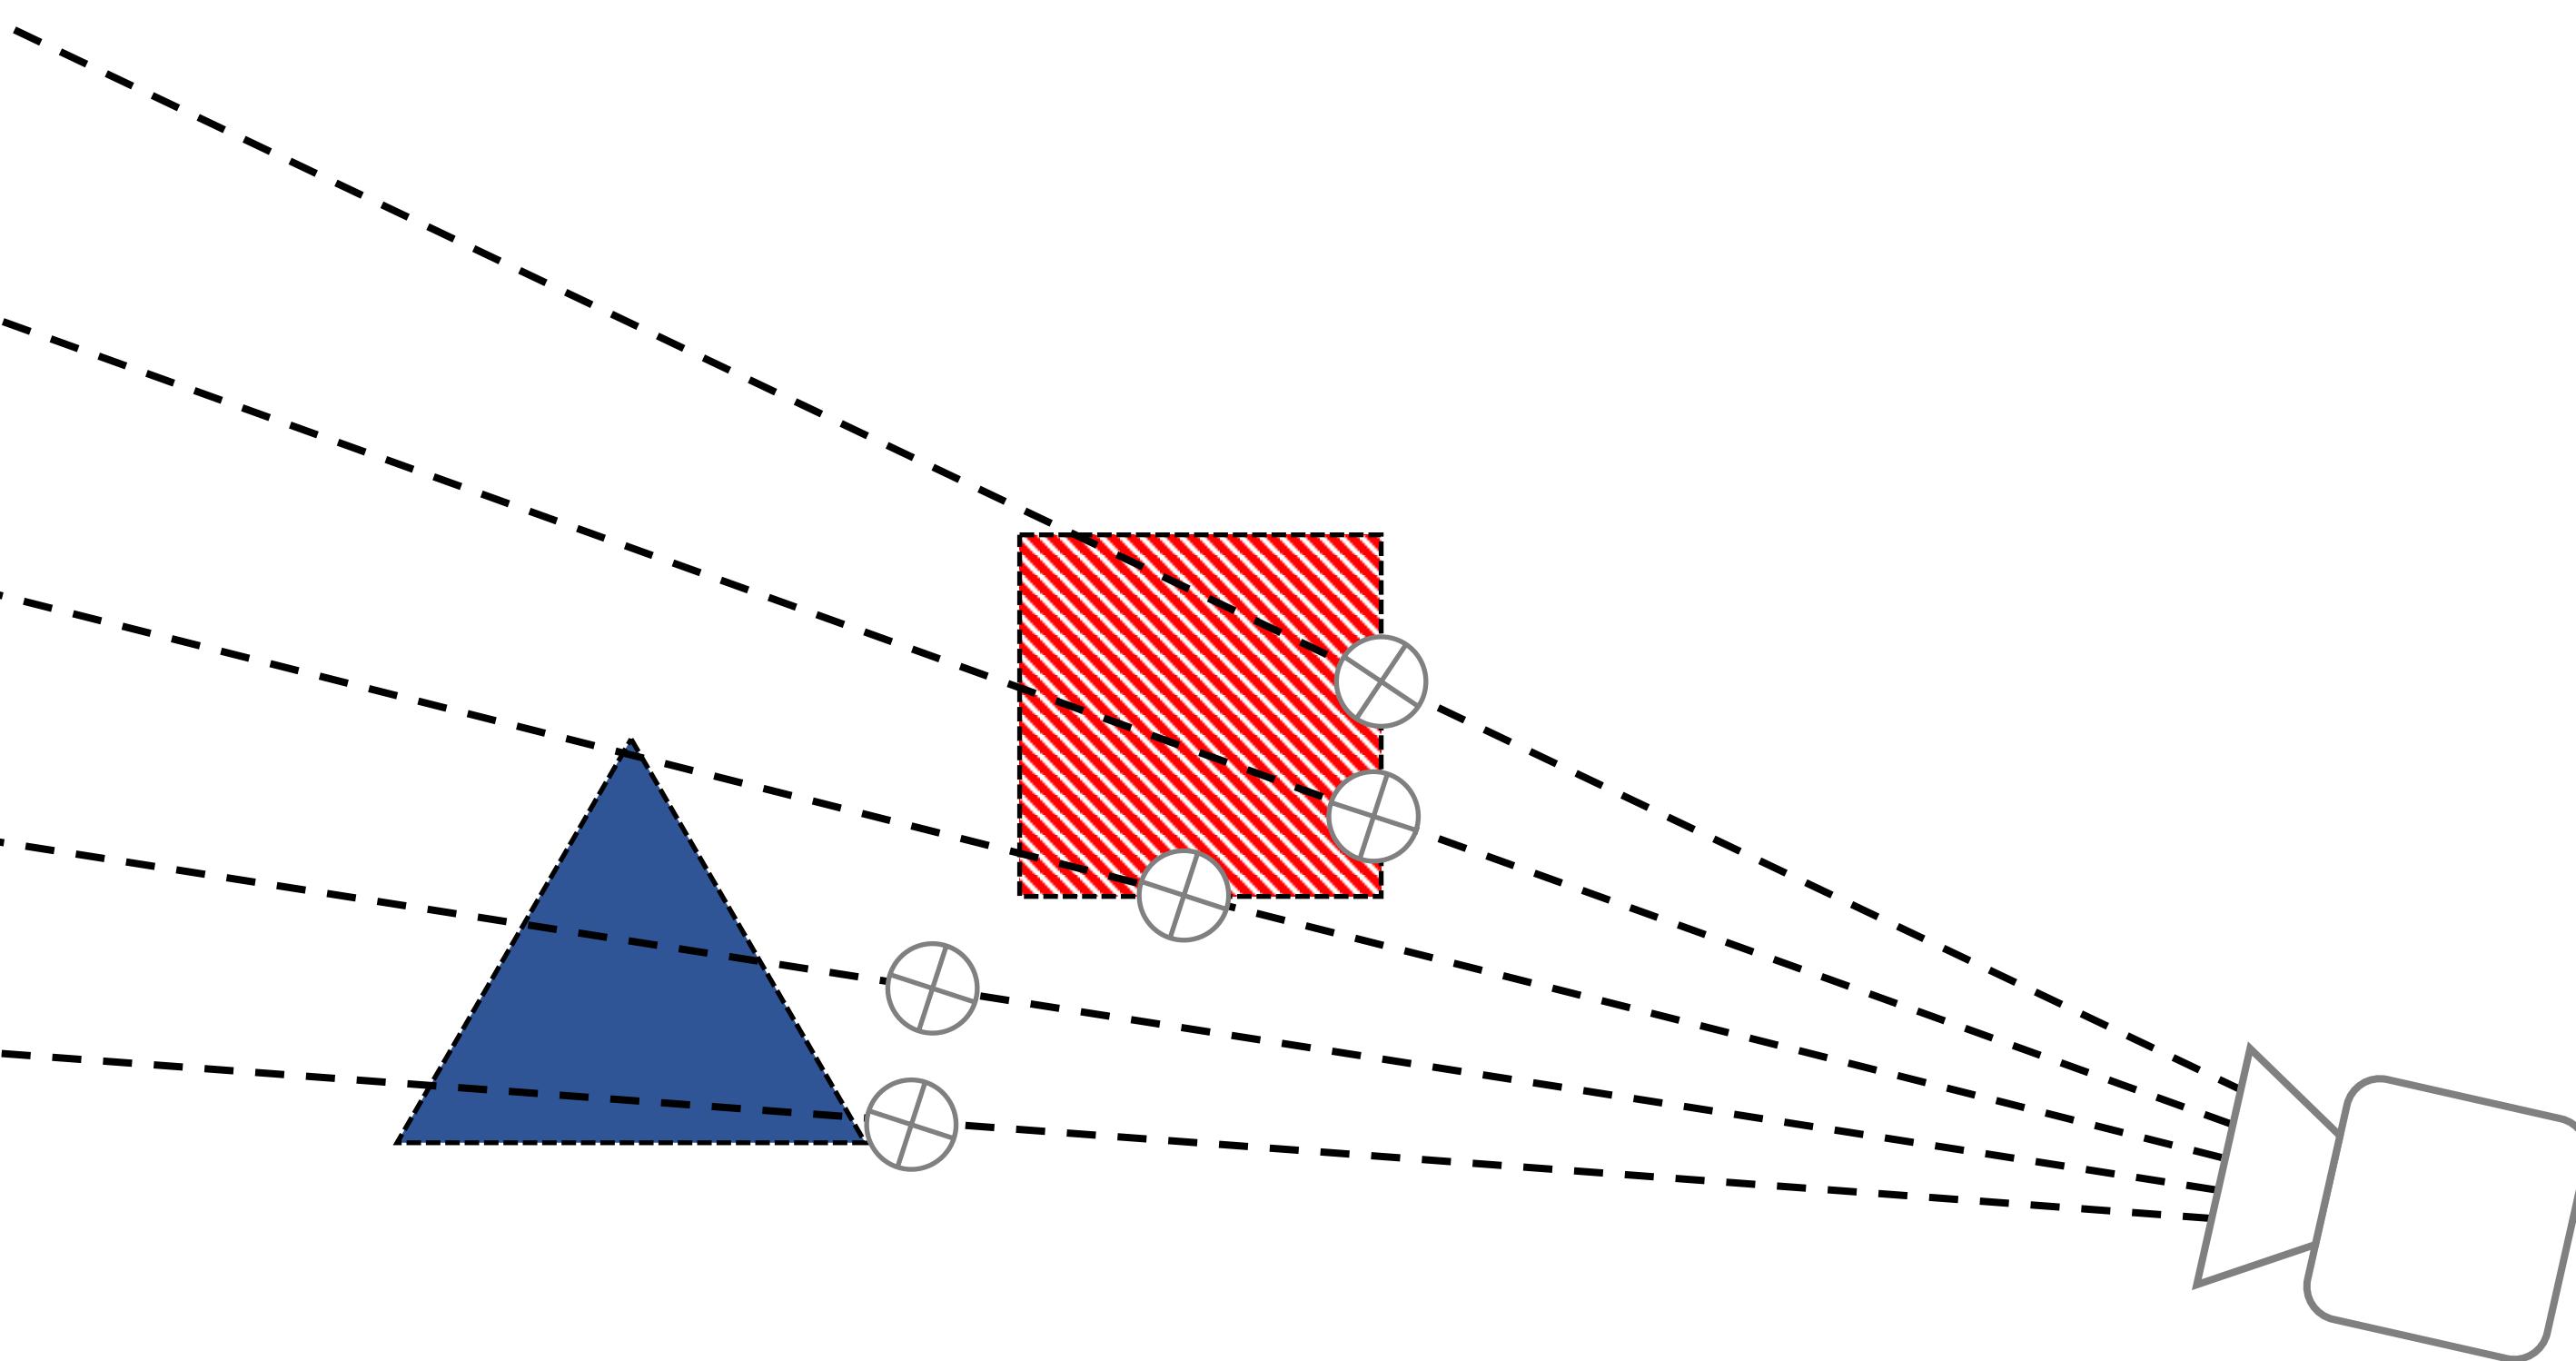
Sphere Tracing Step 1: Intersection Testing.



Sphere Tracing Step 1: Intersection Testing.

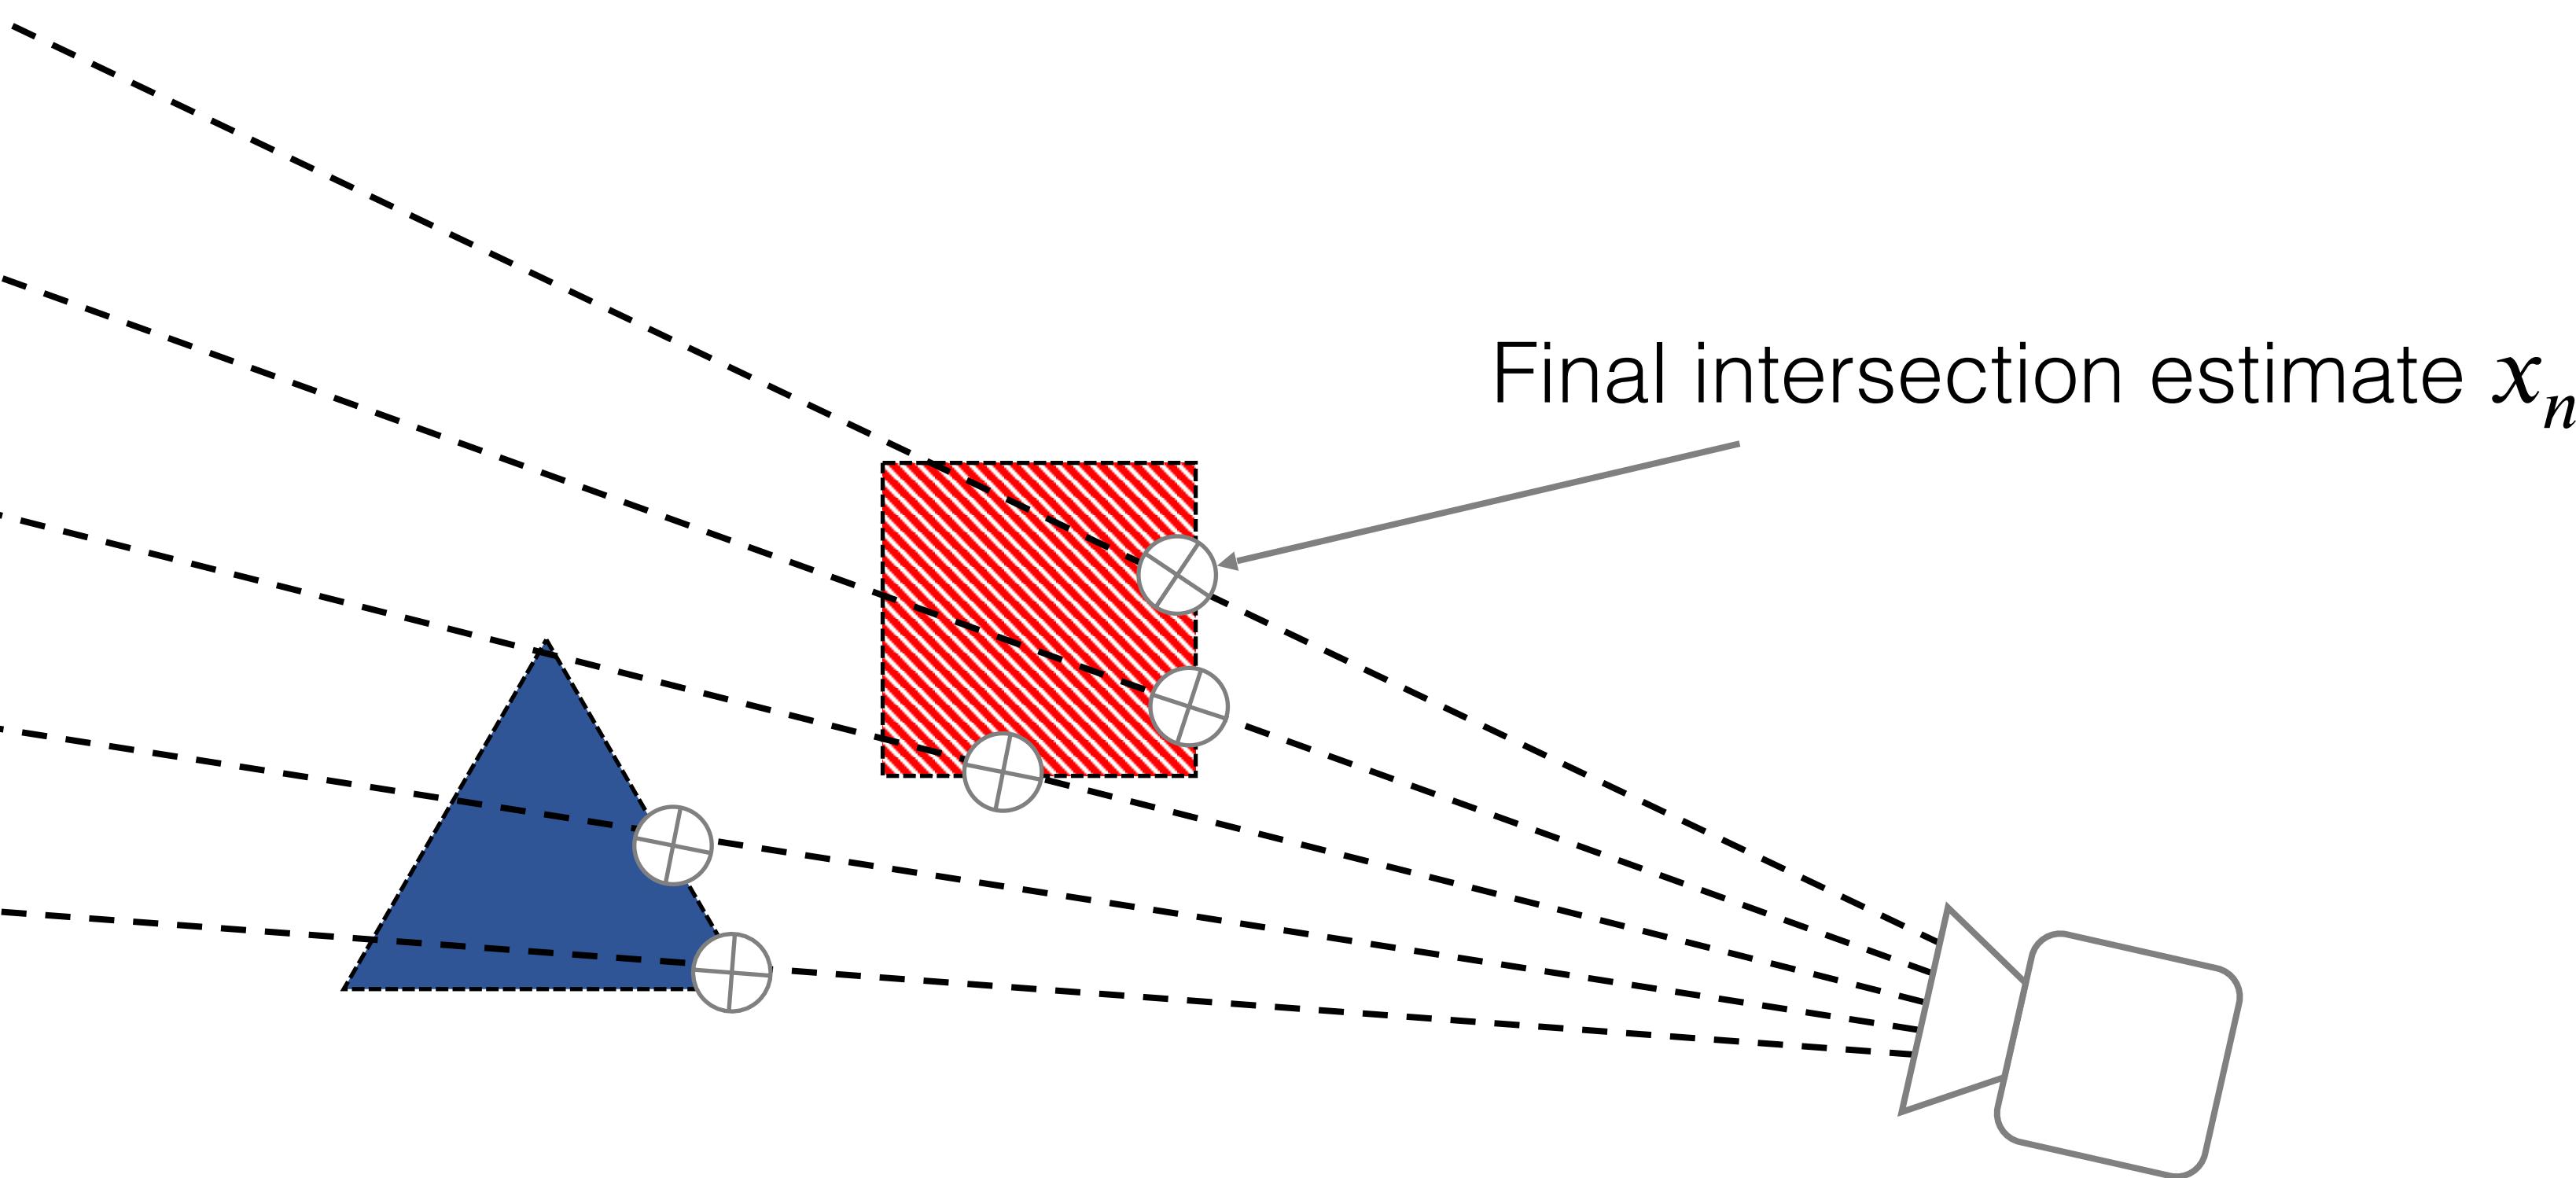


Sphere Tracing Step 1: Intersection Testing.

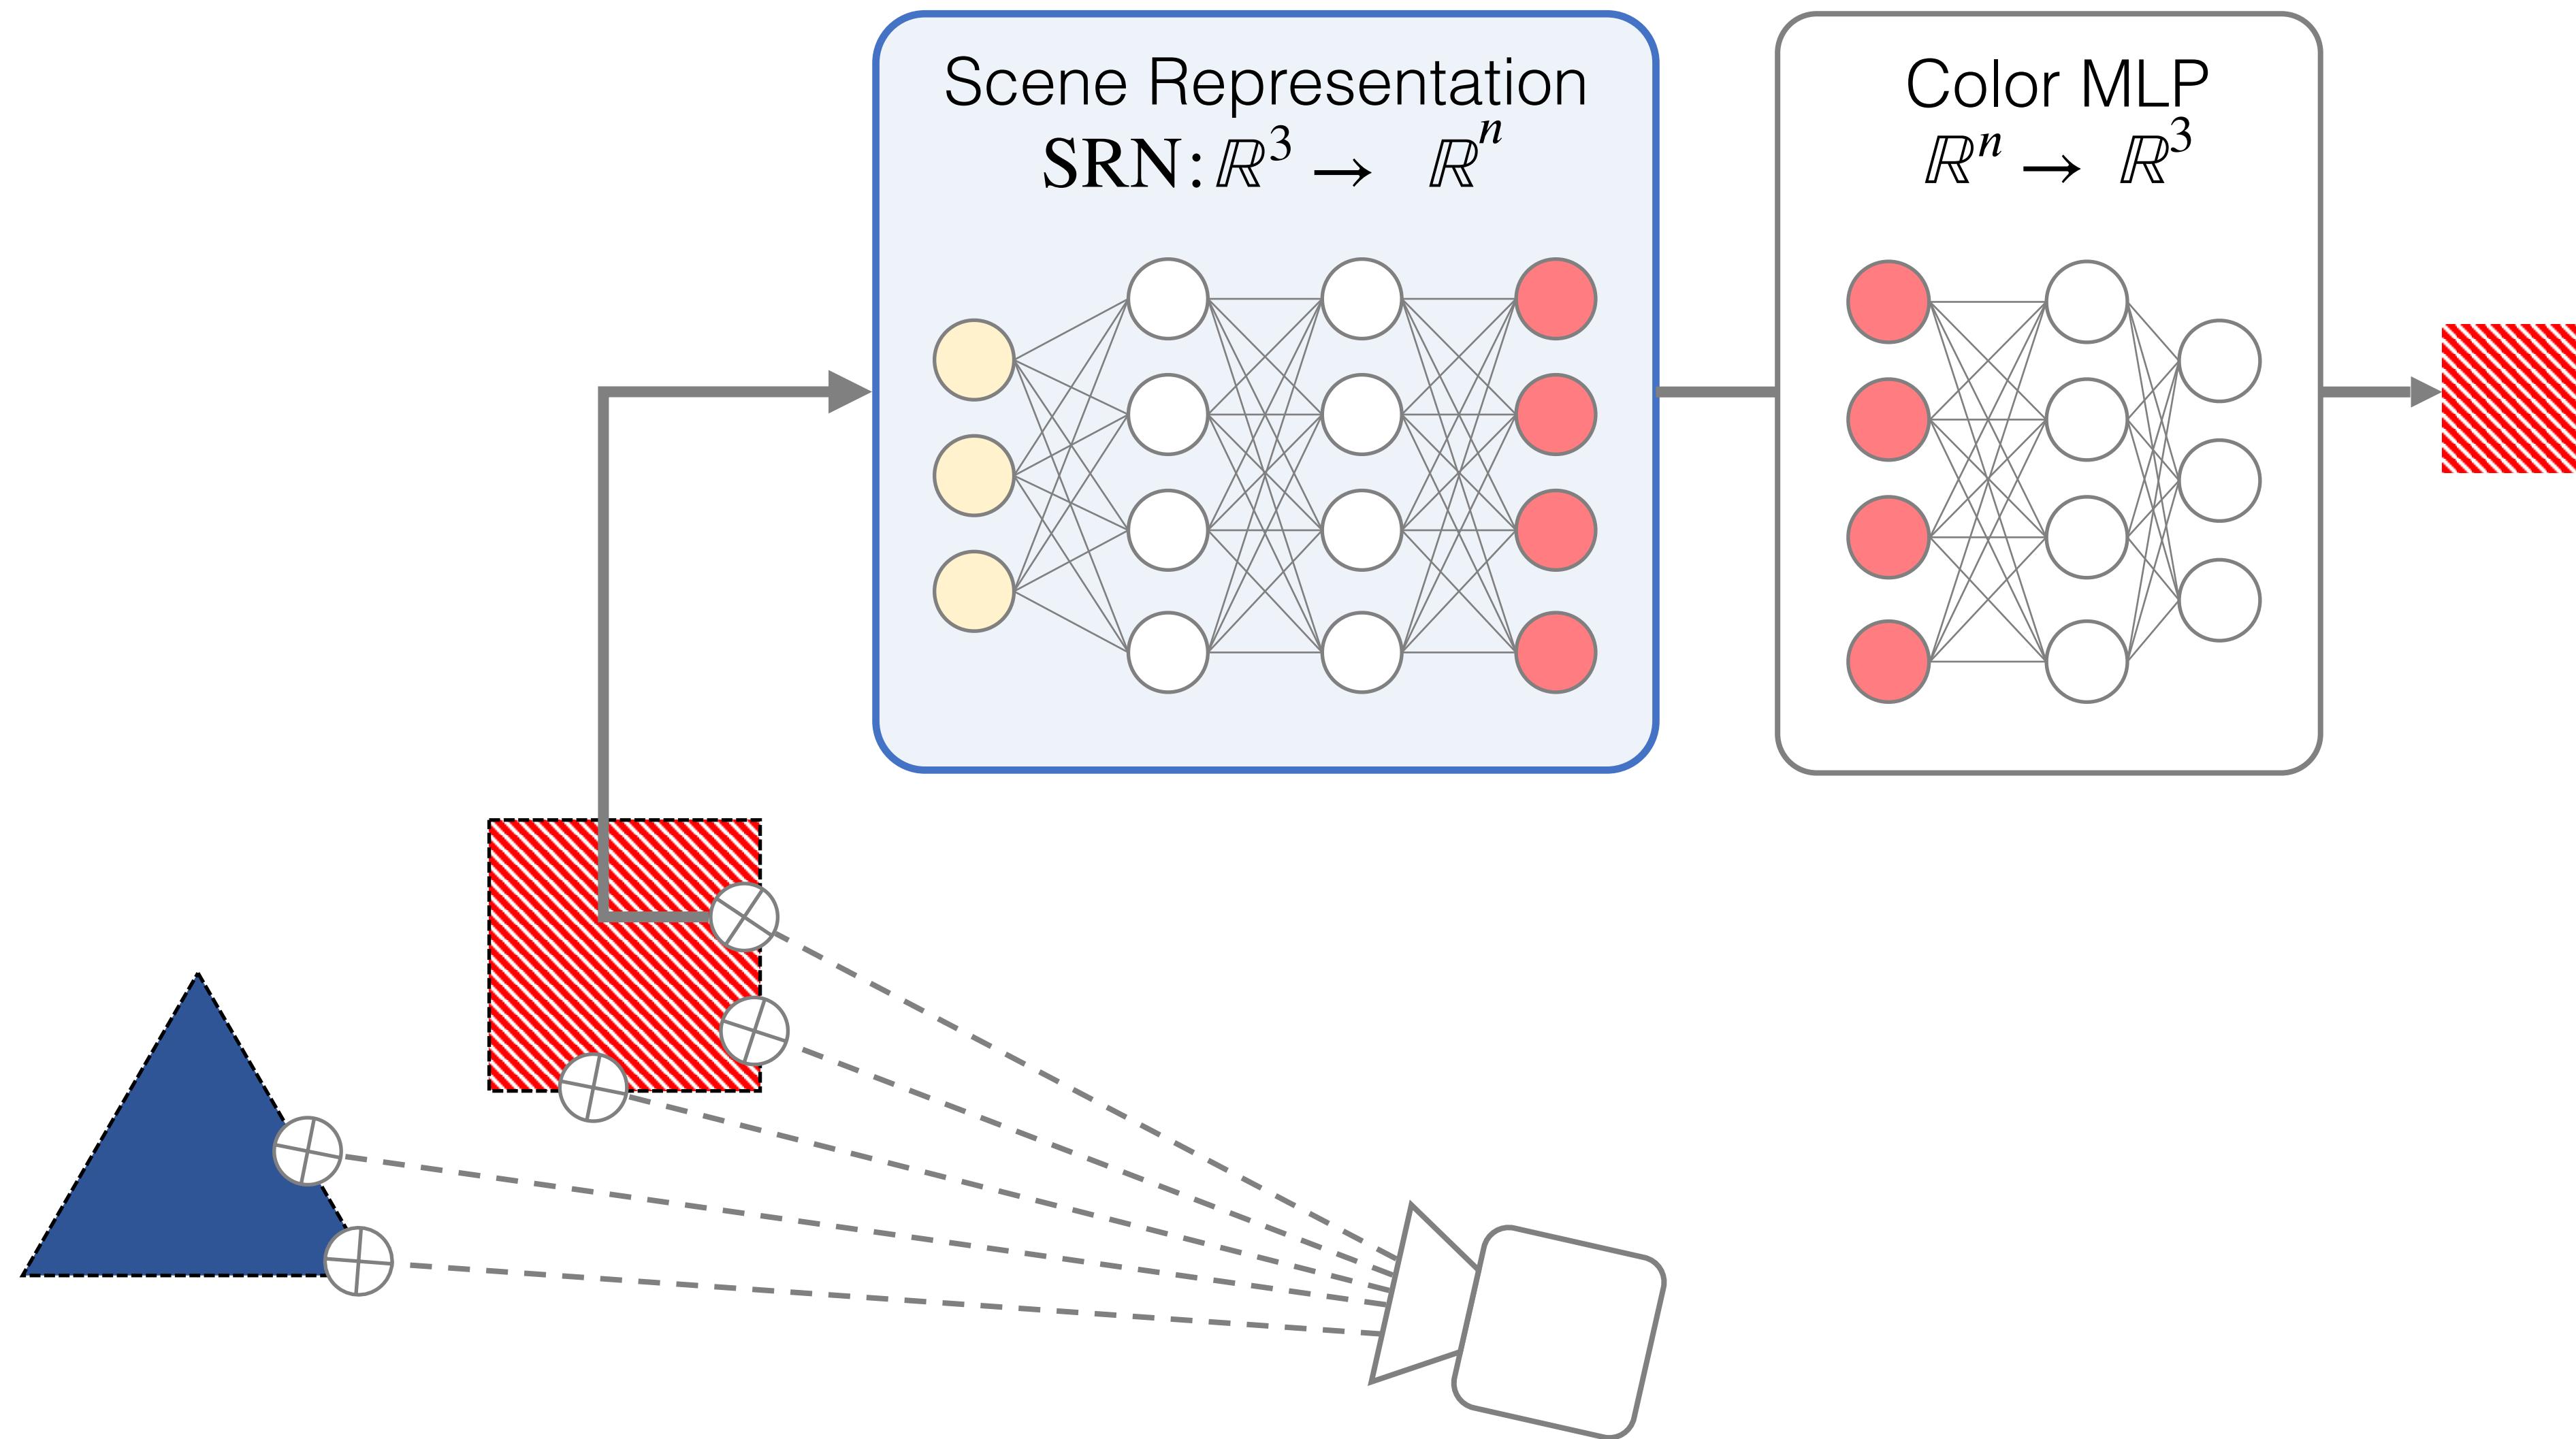


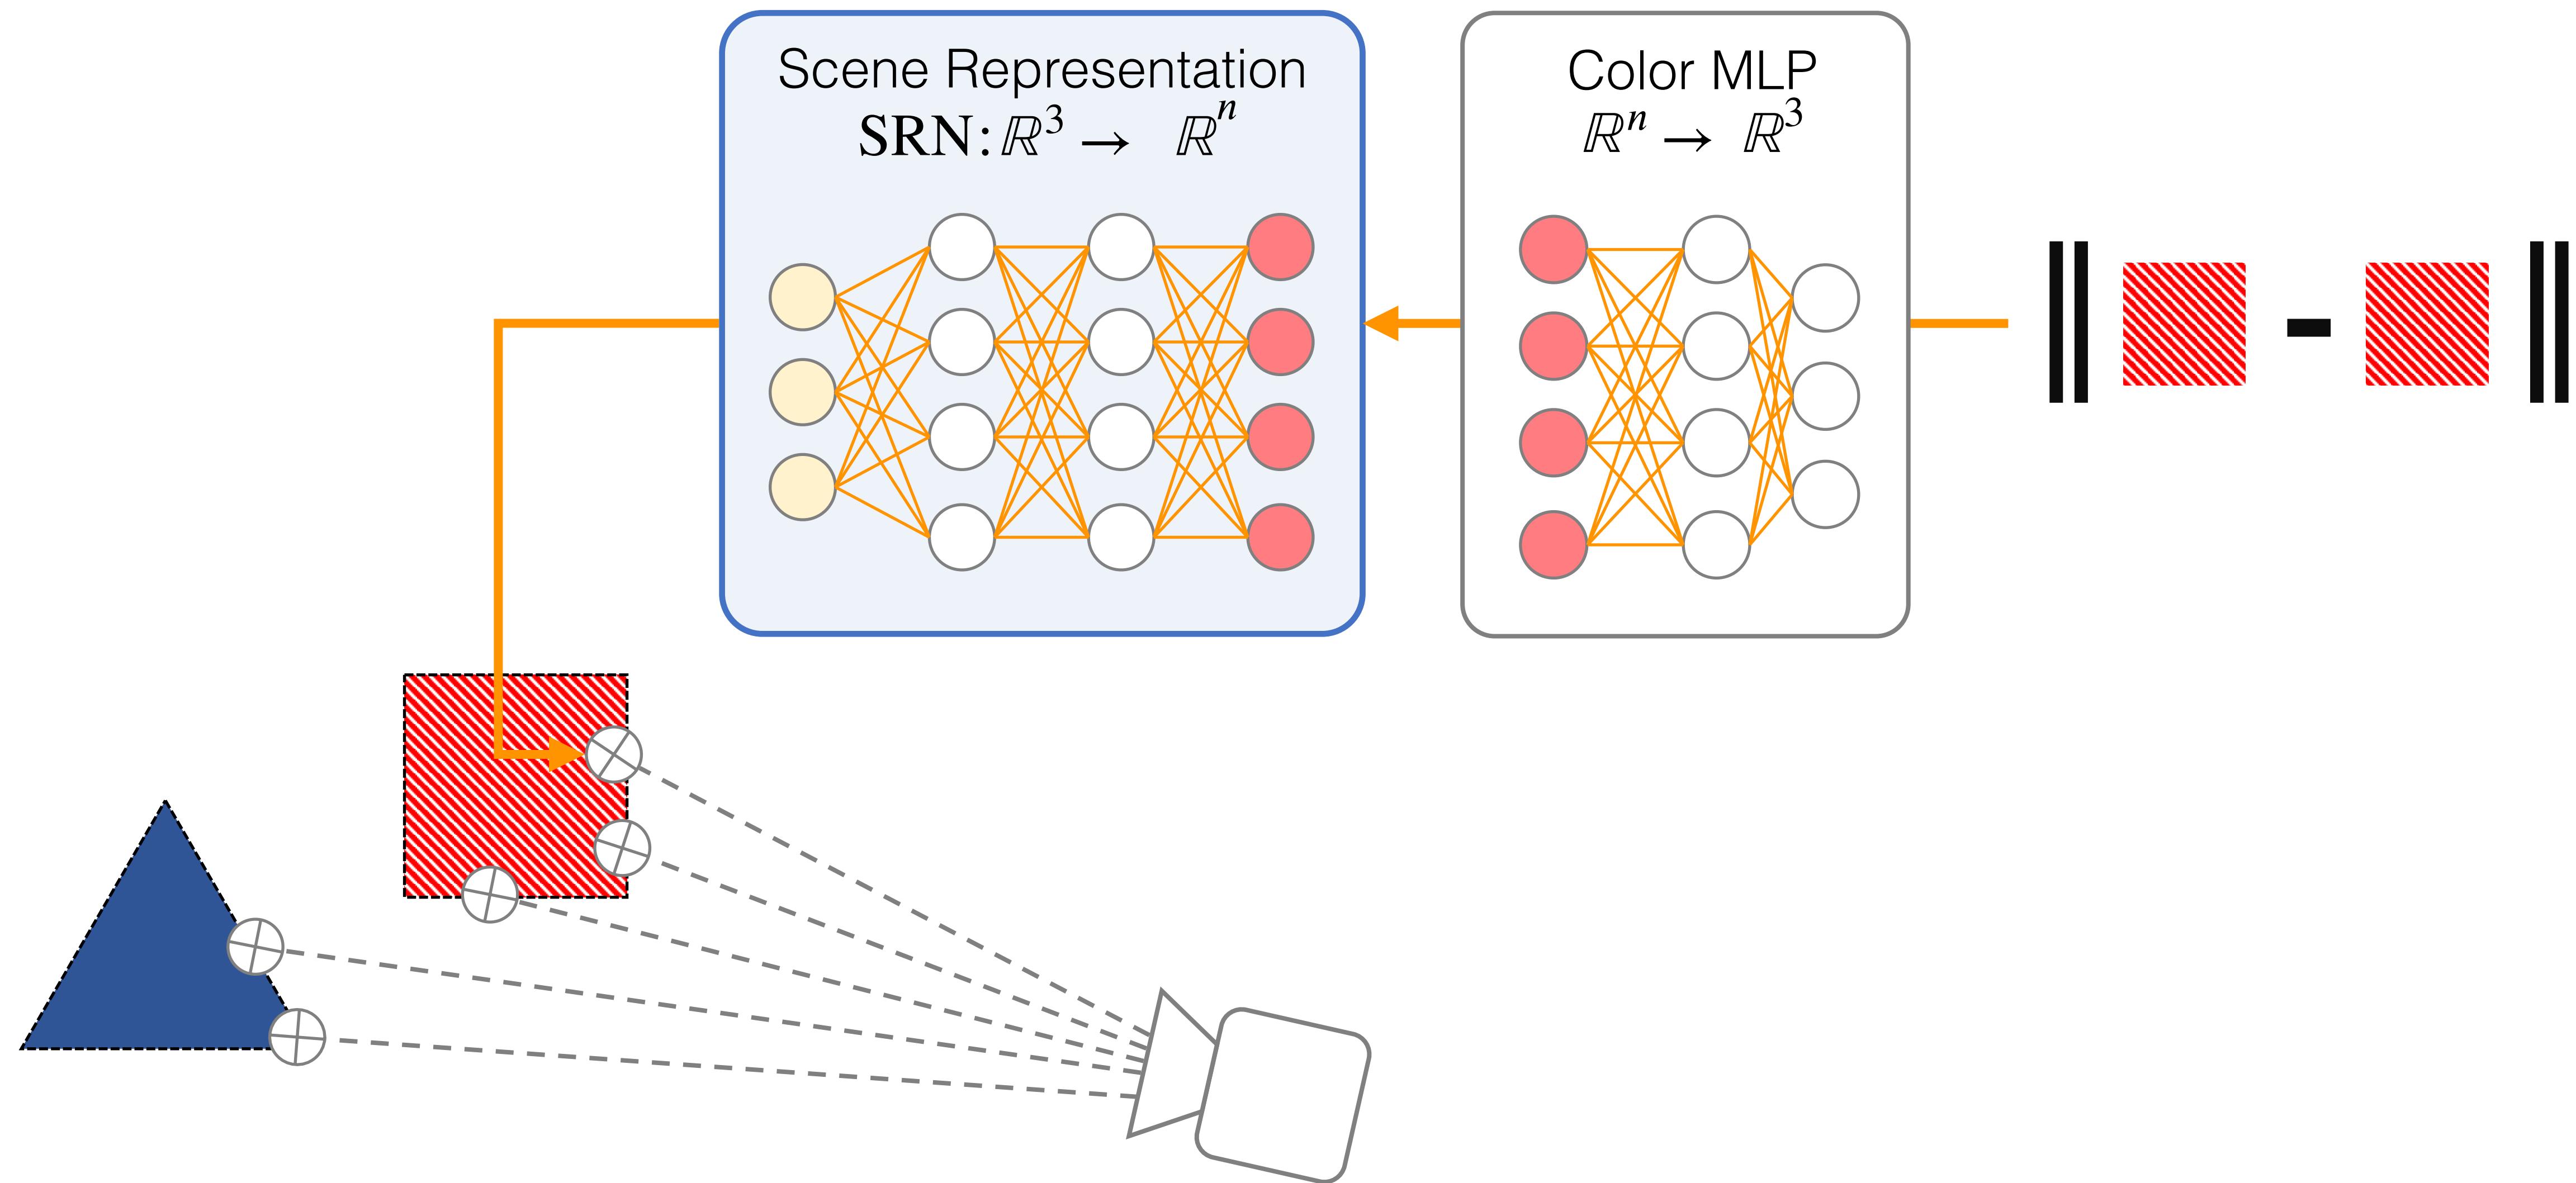
Sphere Tracing Step 1: Intersection Testing.

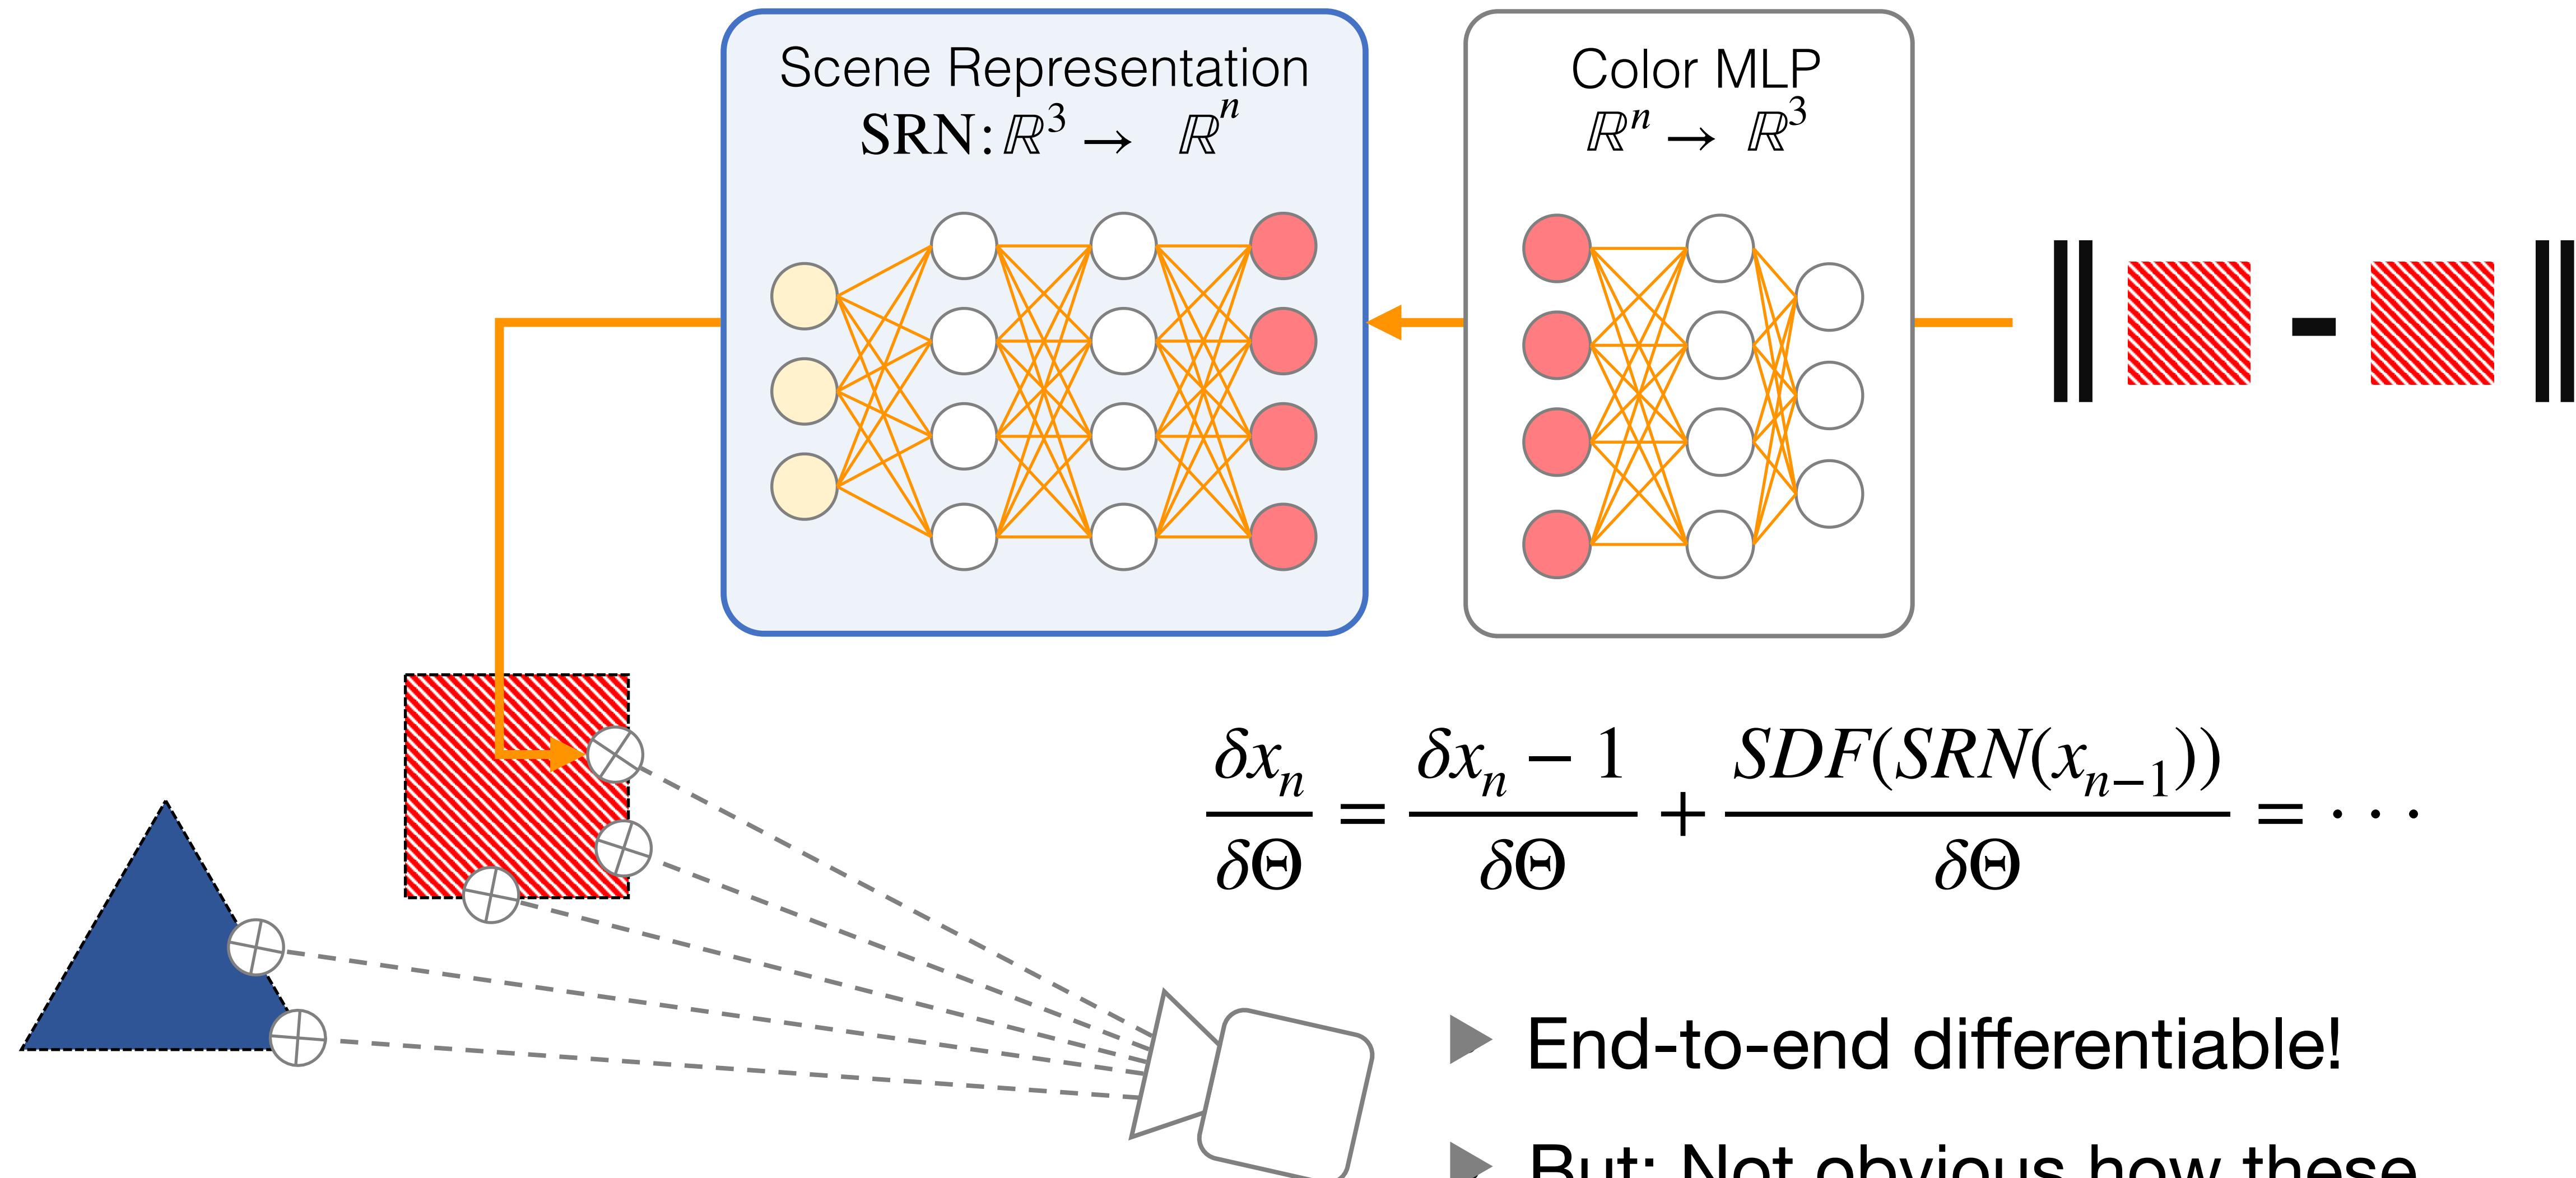
- Length of each ray yields the depth map!



Neural Renderer Step 2: Color Generation







Questions?

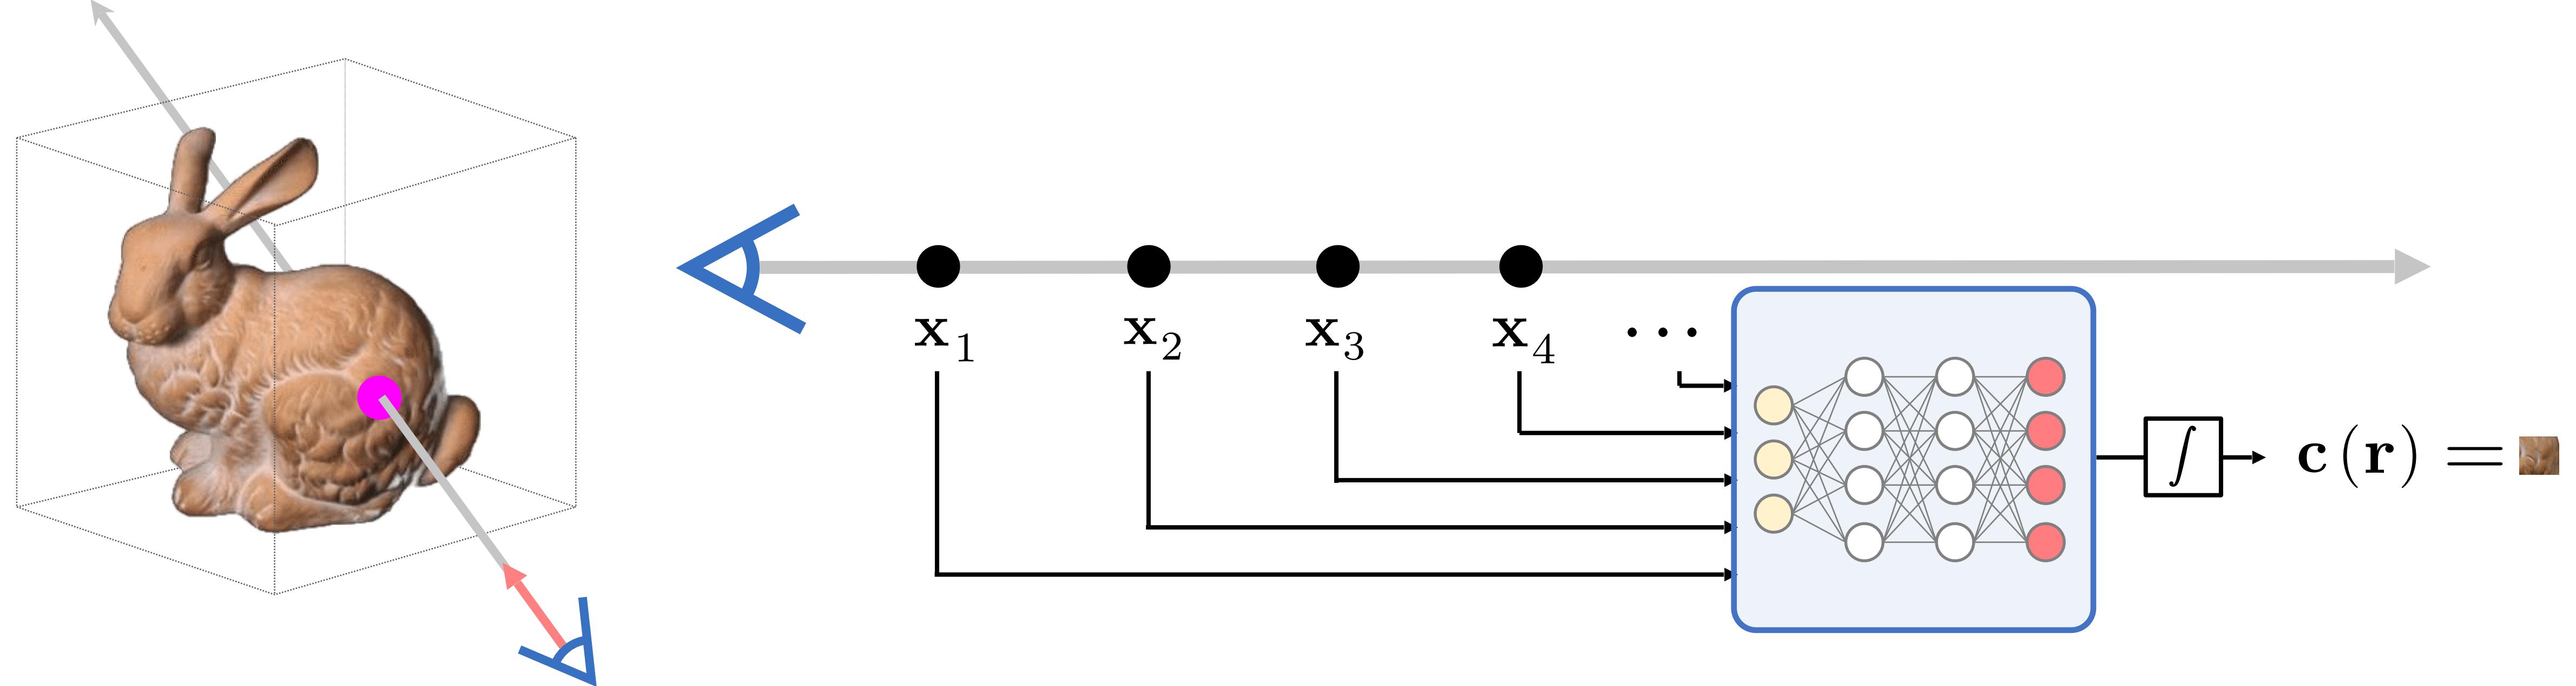
Further reading on computing gradients of sphere-tracing based renderers

Differentiable Rendering of Neural SDFs
through Reparameterization
Bangaru et al. 2022

Differentiable Volumetric Rendering
Niemeyer et al. 2020

Differentiable signed distance function
rendering
Vicini et al. 2022

General structure of Neural Renderers for 3D Fields



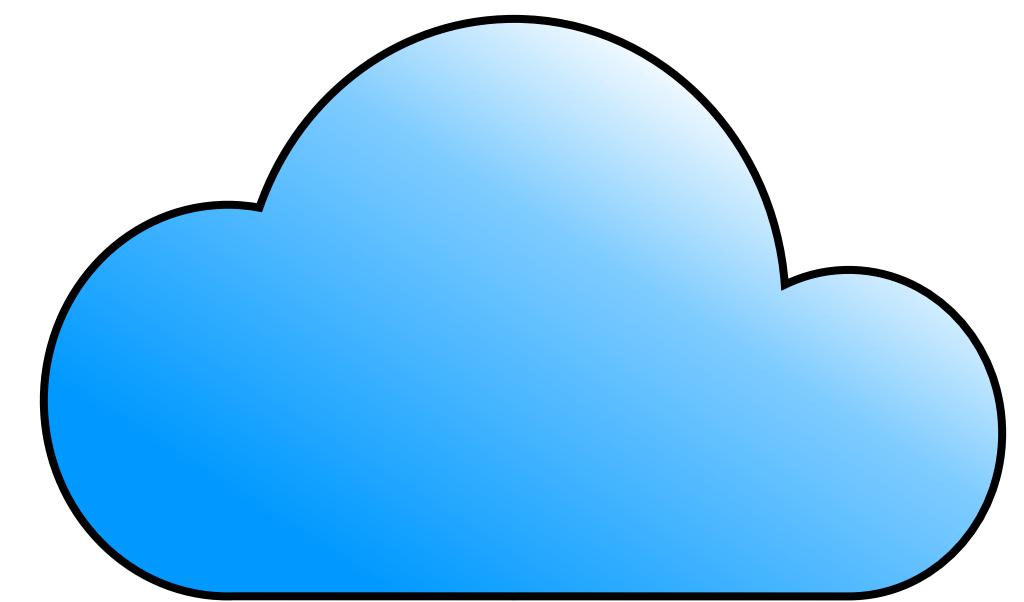
Sphere-Tracing
[JC Hart, 1996]

Volumetric

Hybrid implicit-volumetric

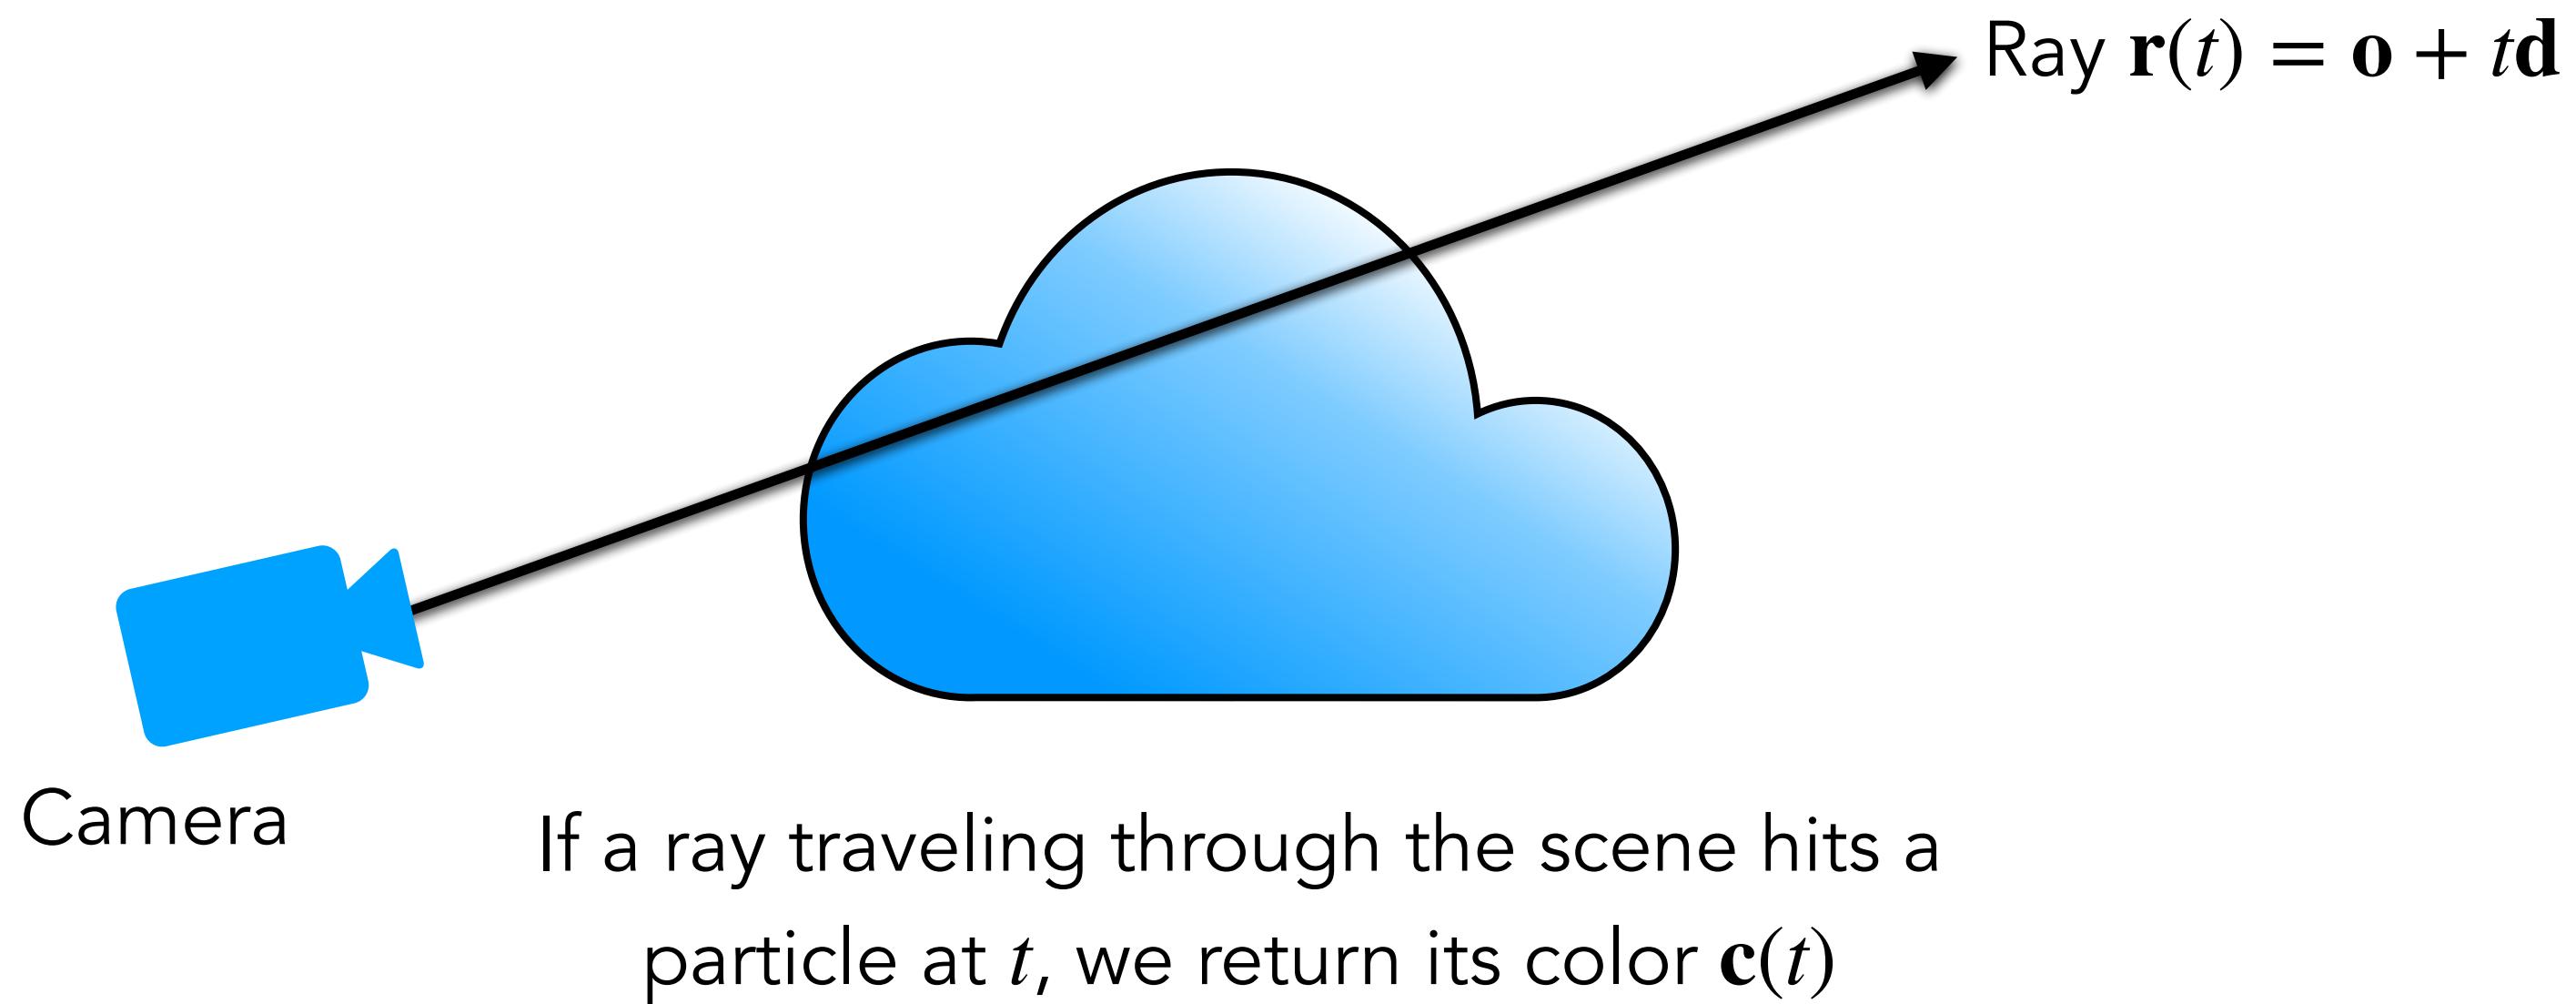
Learned aggregation

Volume Rendering

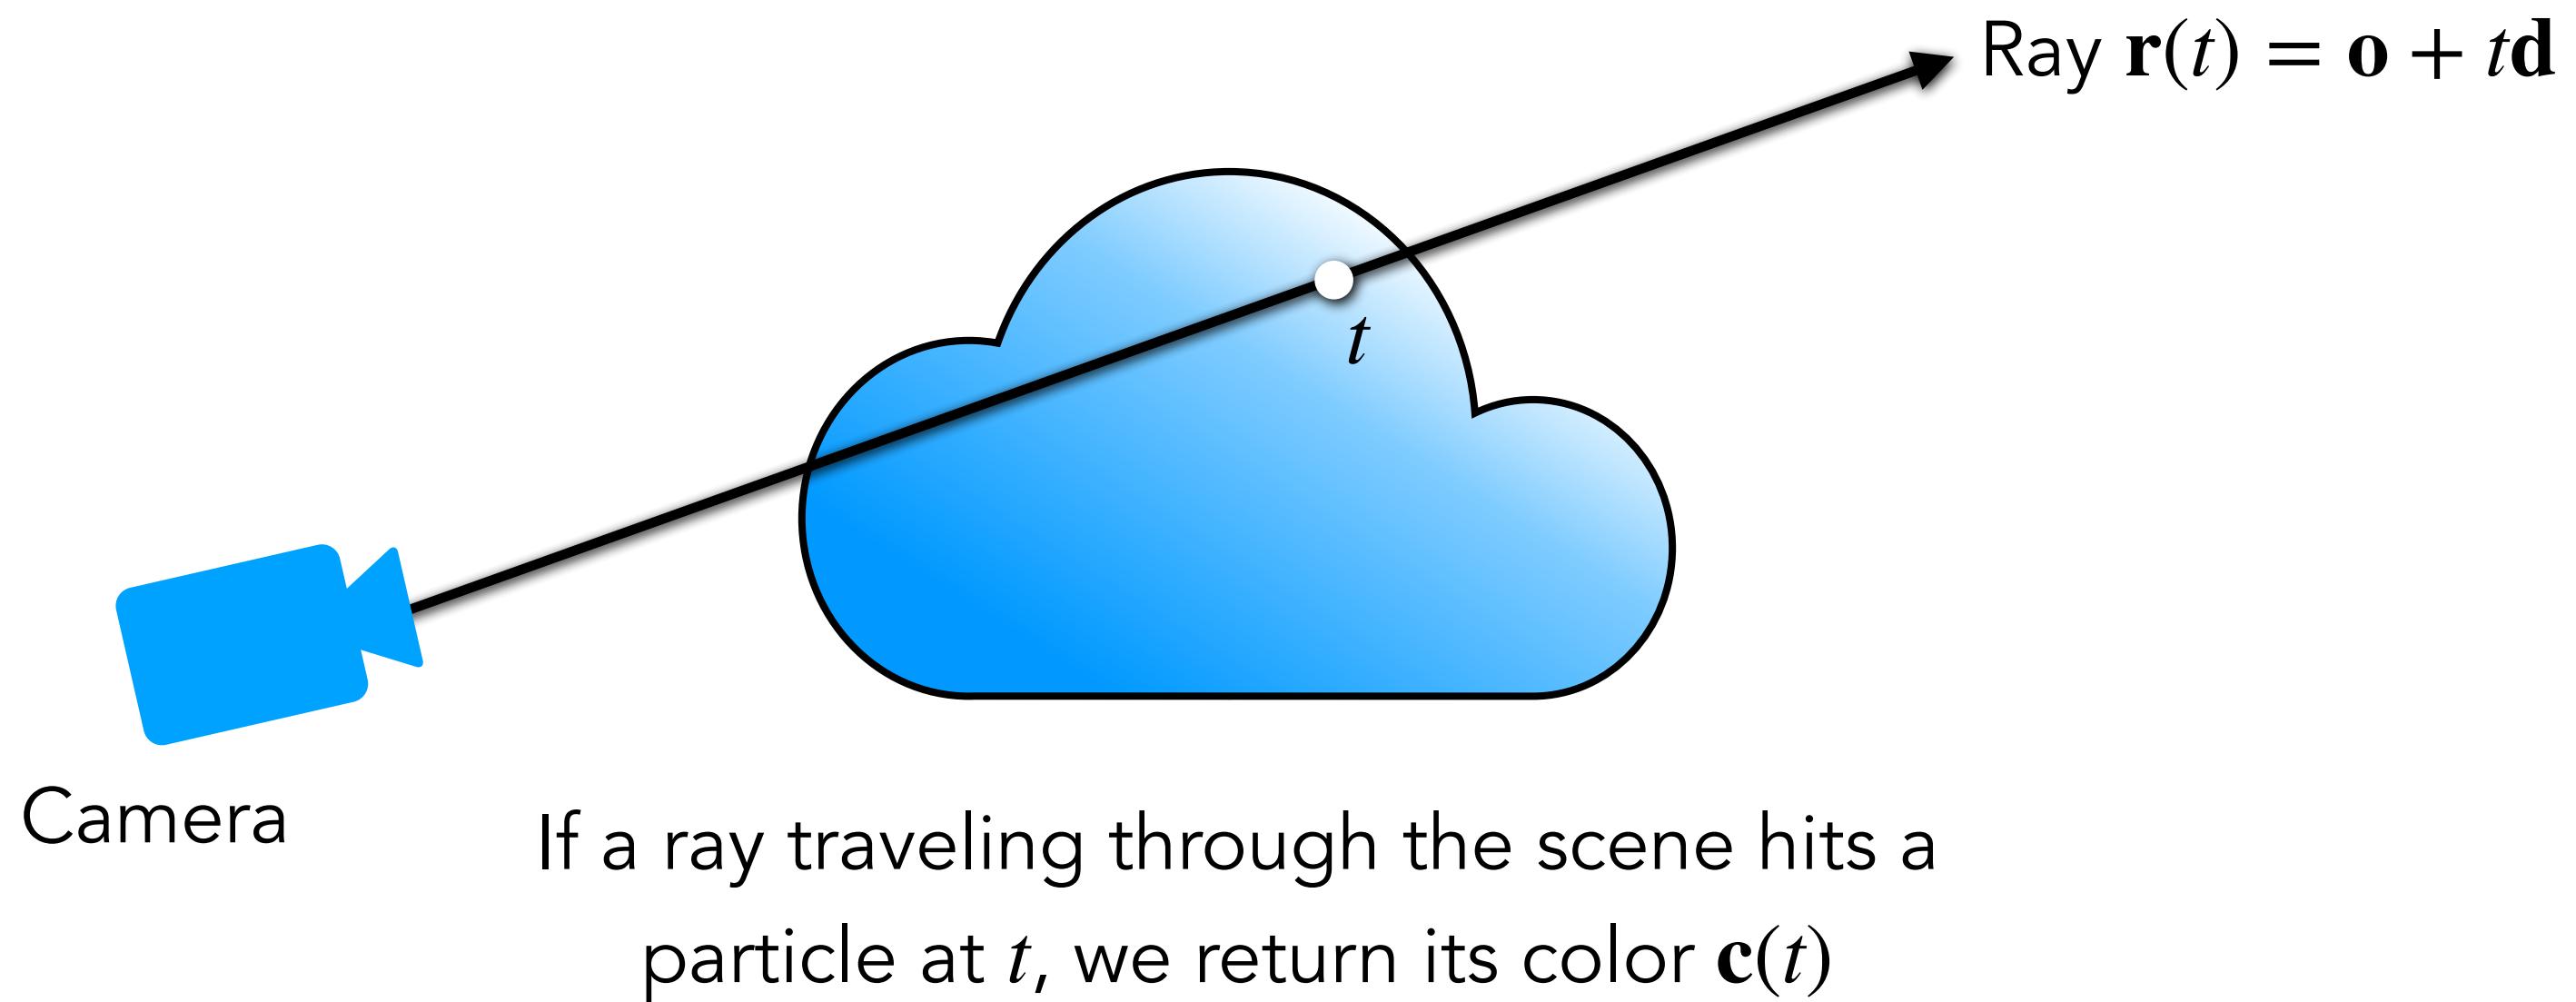


Scene is a cloud of tiny colored particles

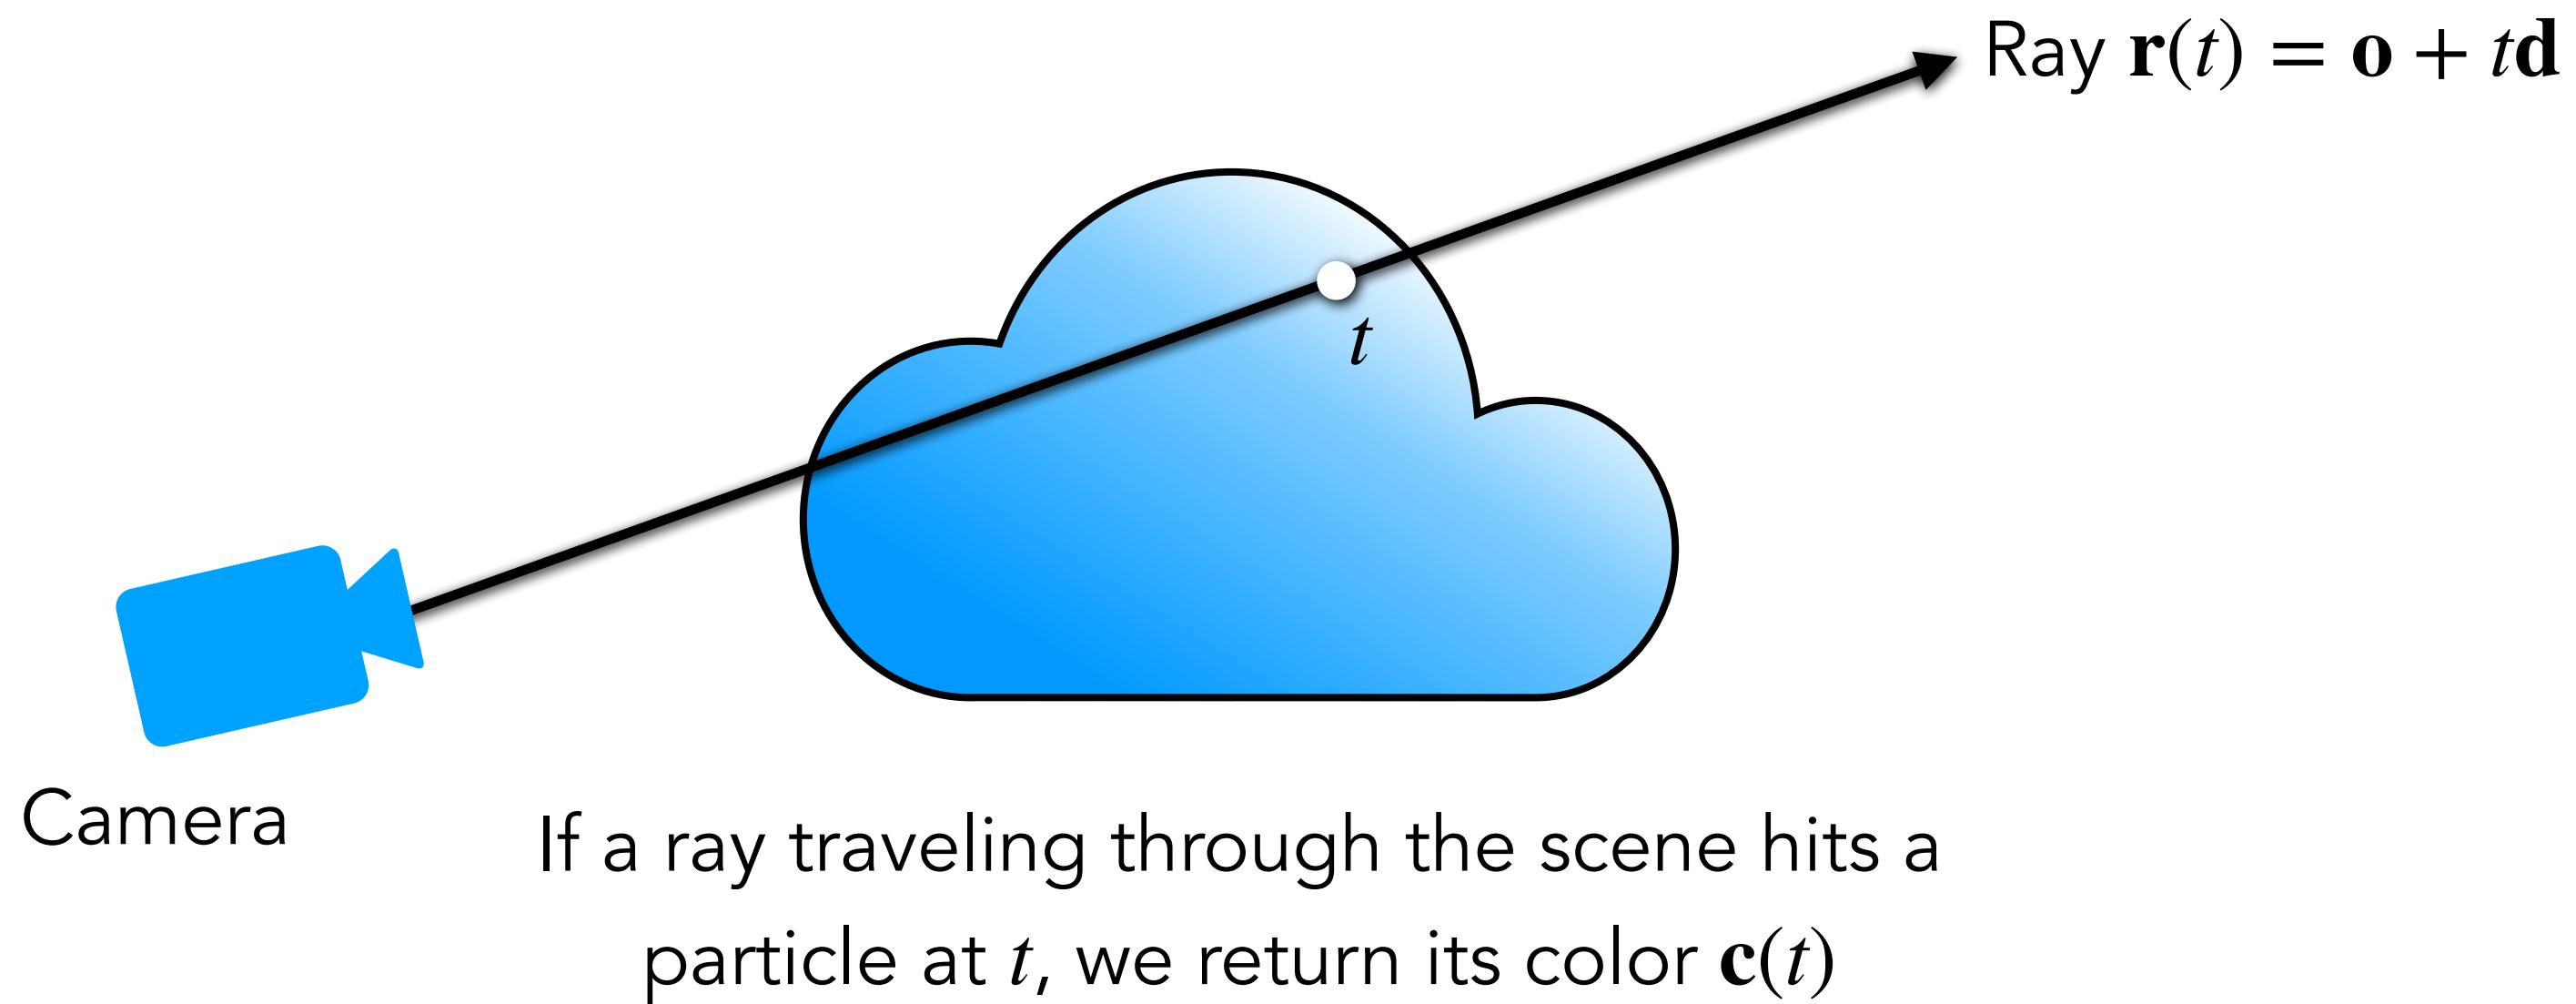
Volume Rendering



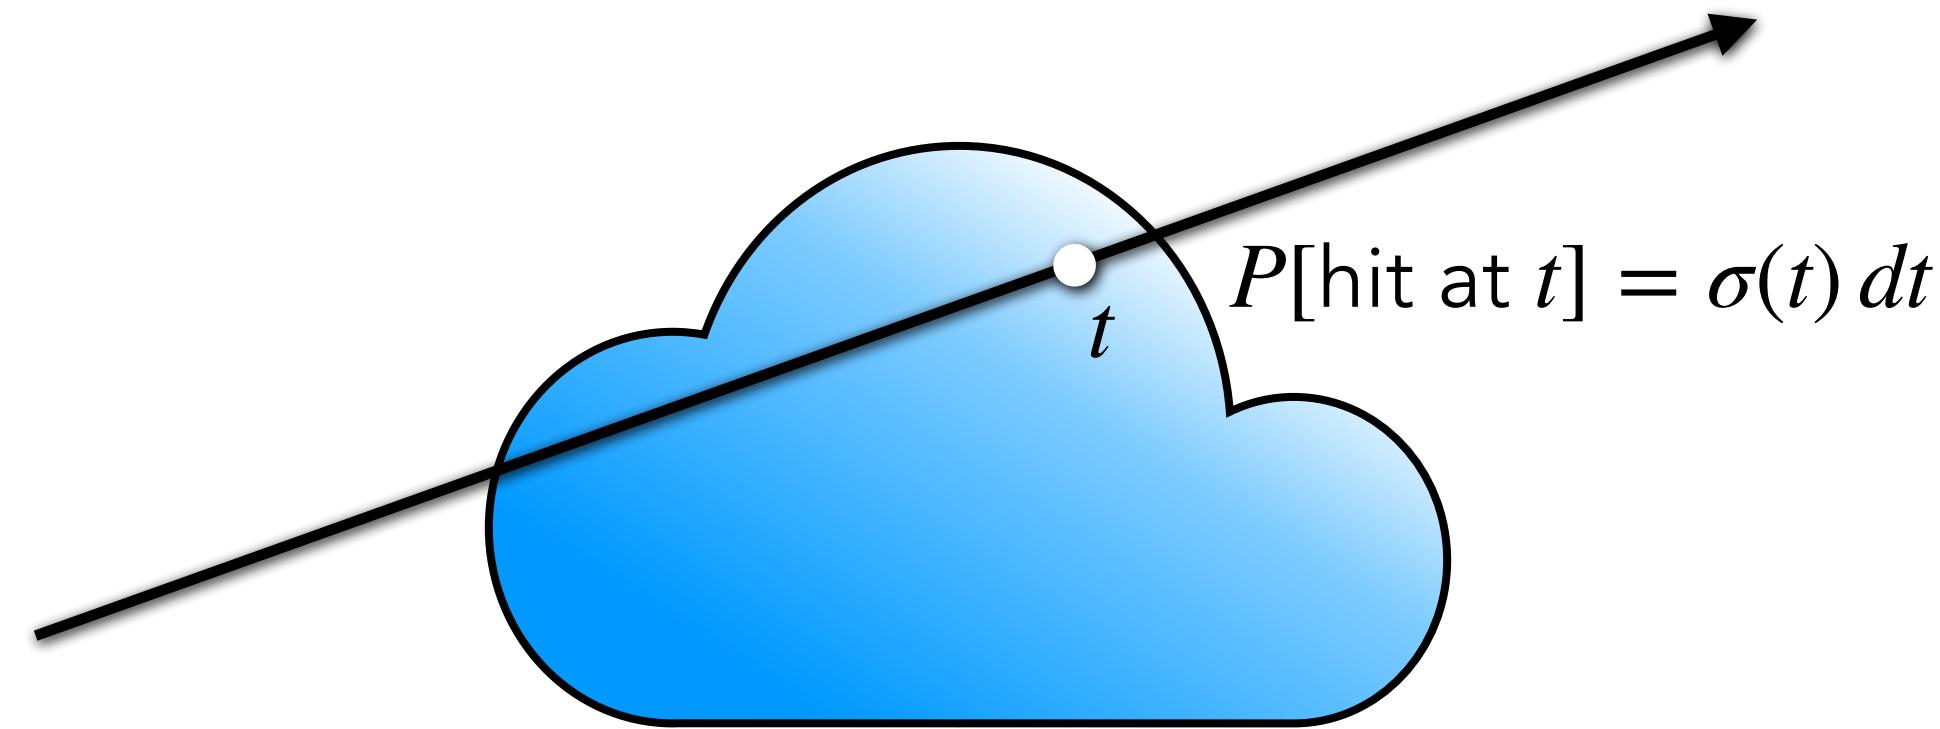
Volume Rendering



Volume Rendering

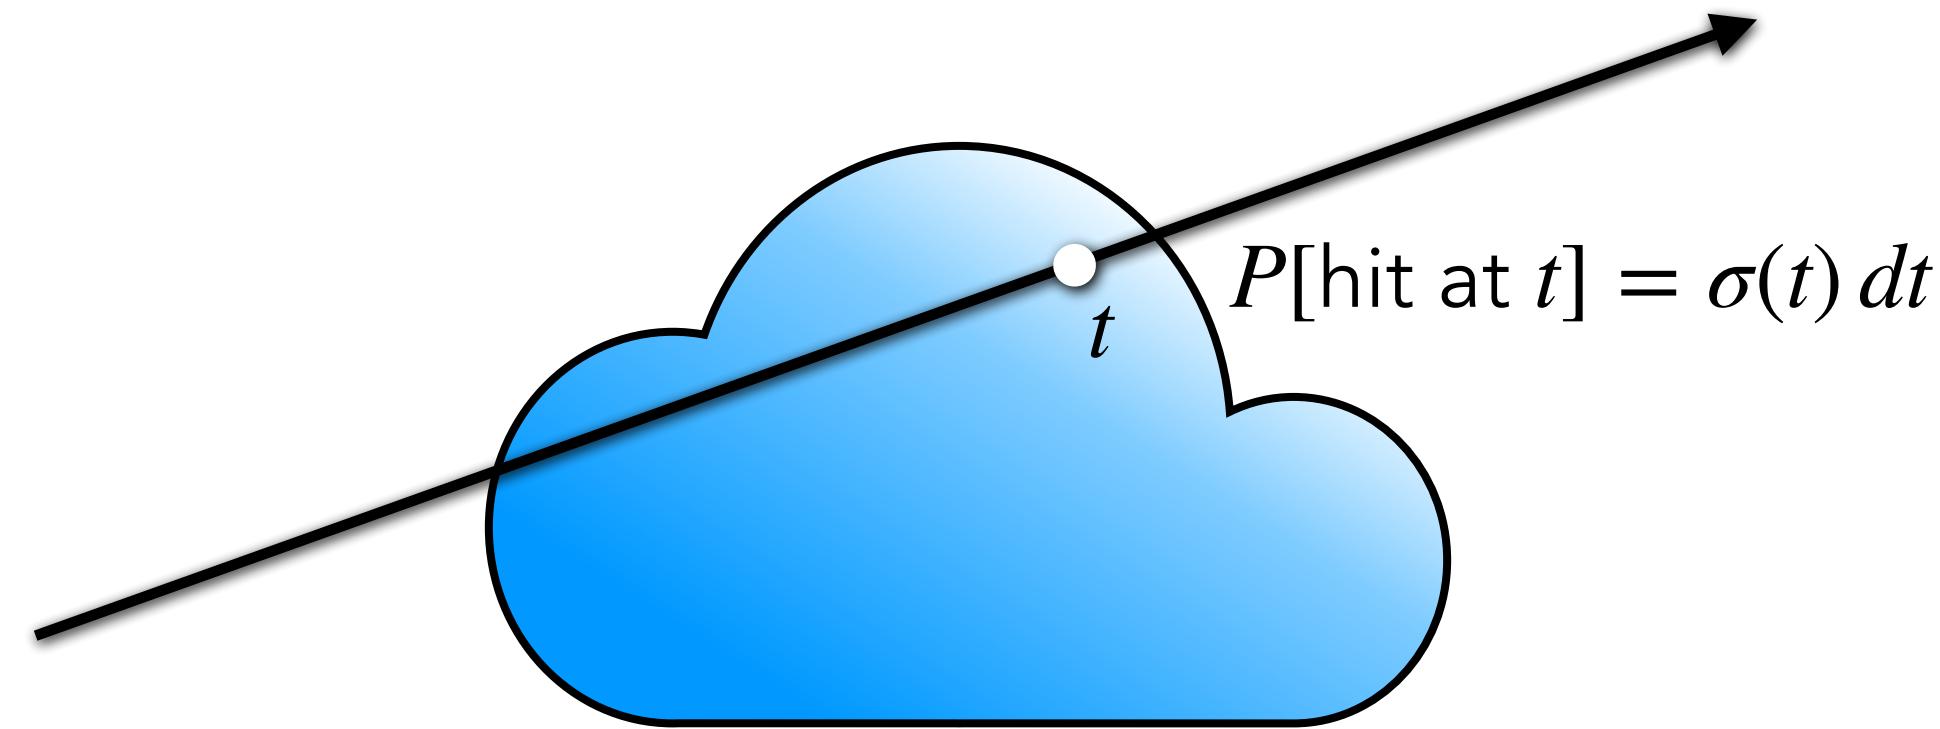


Volume Rendering



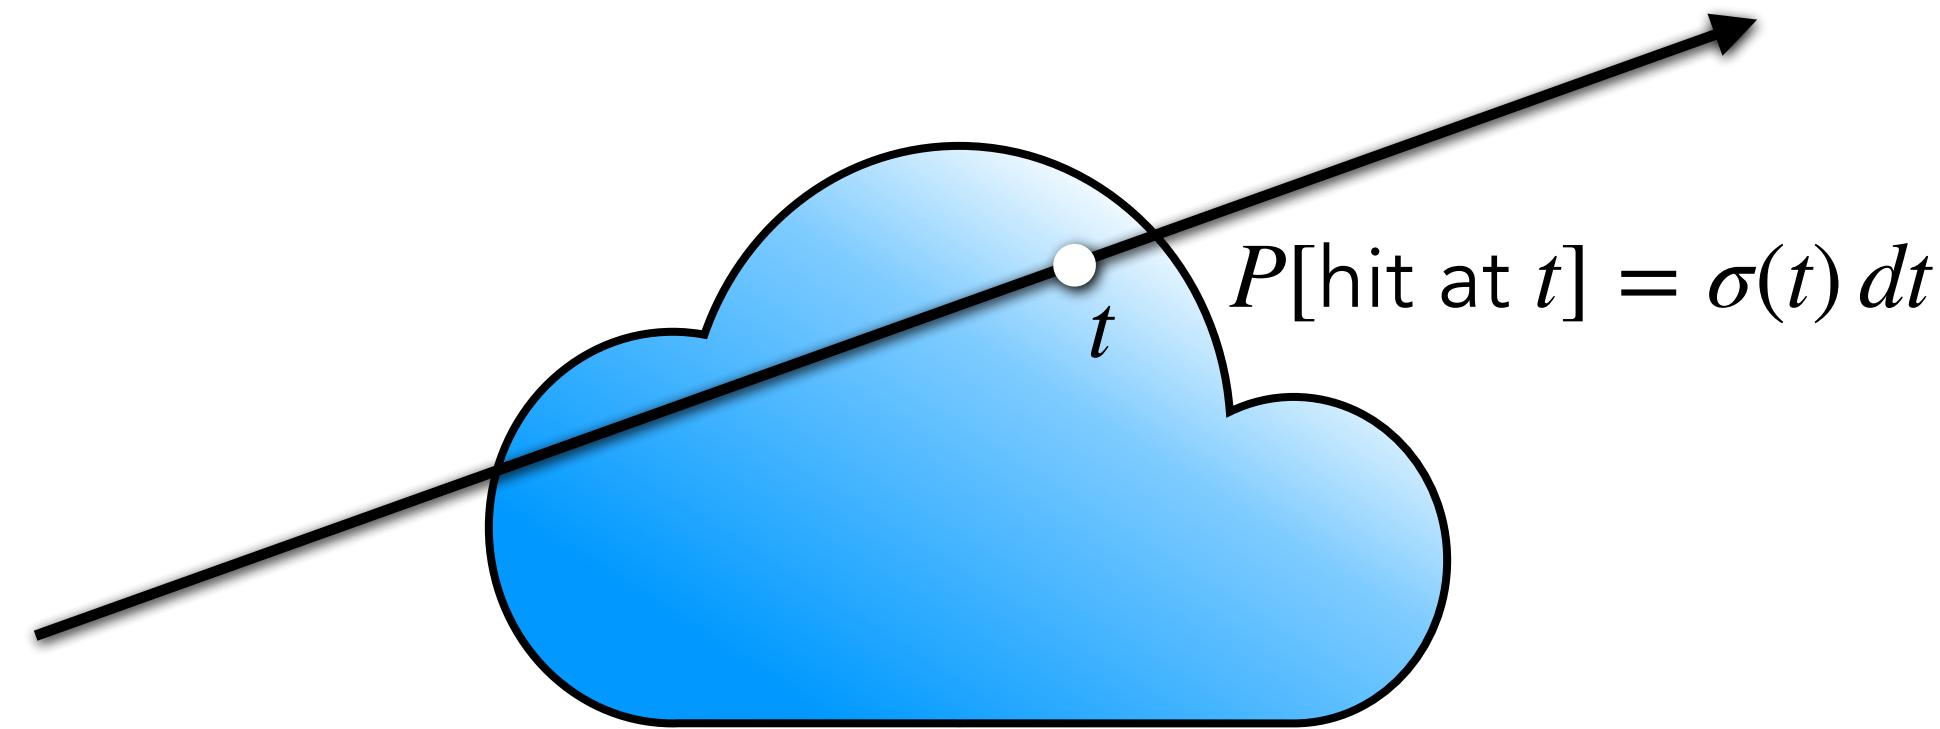
Probability that ray stops in a small interval around t is $\sigma(t) dt$.
 σ is known as the **Volume Density**.

Volume Rendering



Probability that ray stops in a small interval around t is $\sigma(t) dt$.
 σ is known as the **Volume Density**.

Volume Rendering

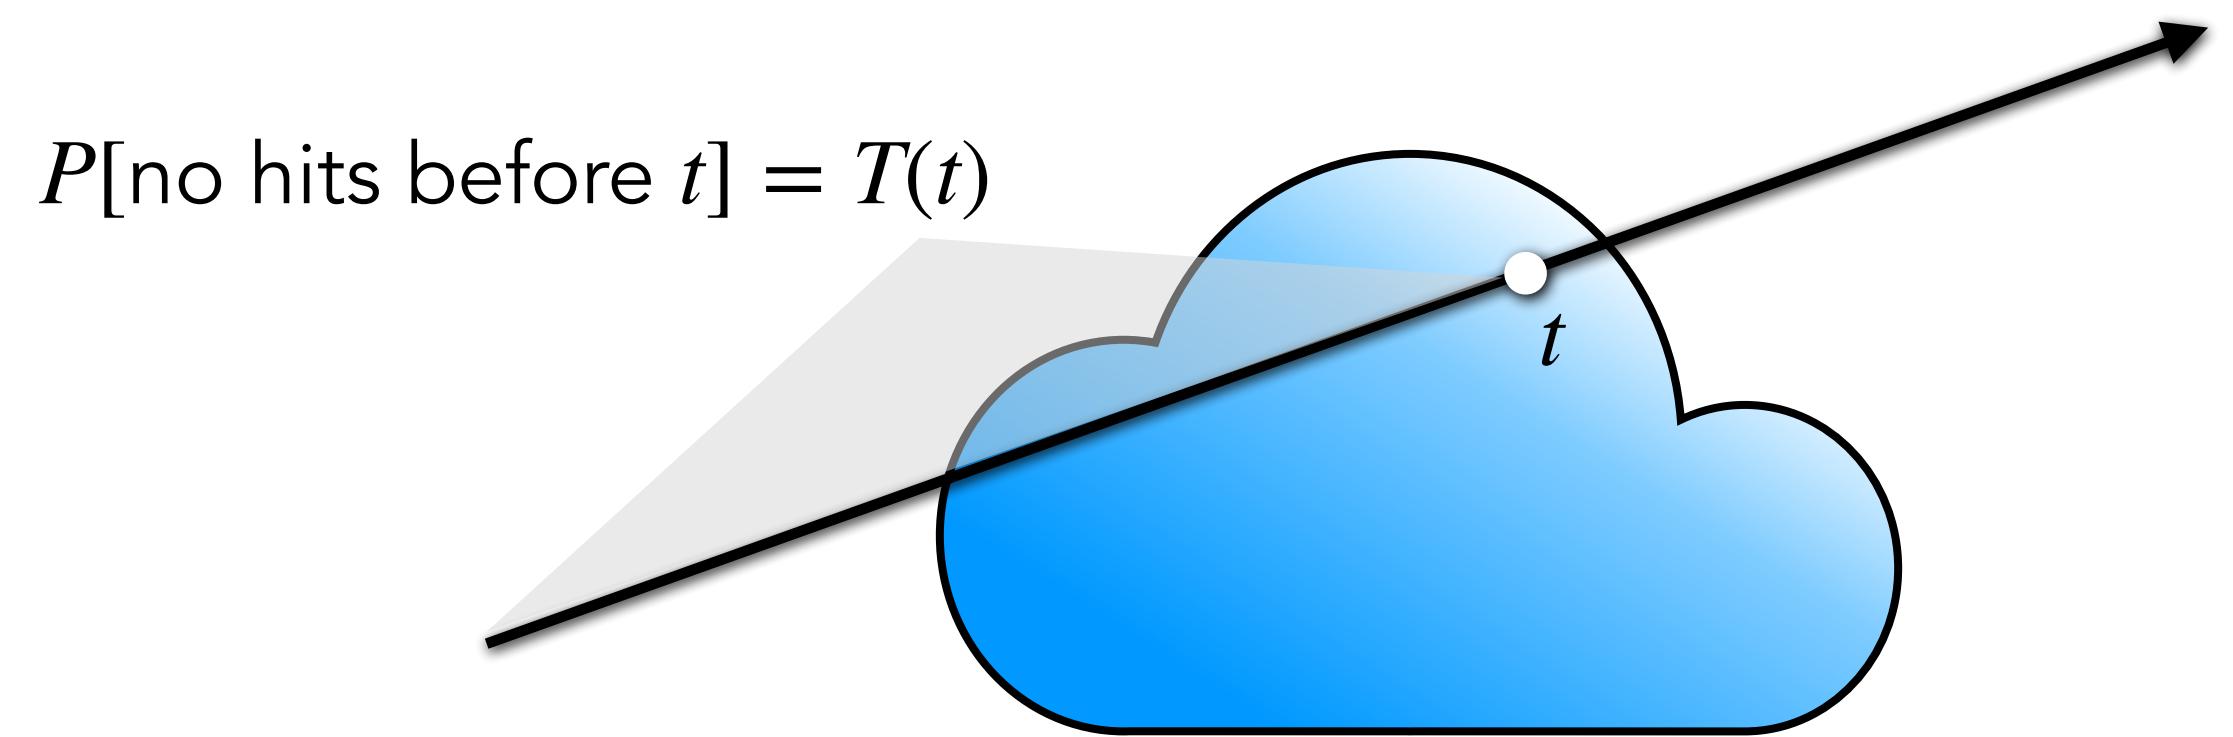


Probability that ray stops in a small interval around t is $\sigma(t) dt$.
 σ is known as the **Volume Density**.

Our field thus has the function signature:

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^3 \times \mathbb{R}^+; \quad \Phi(\mathbf{x}) = (\sigma, \mathbf{c})$$

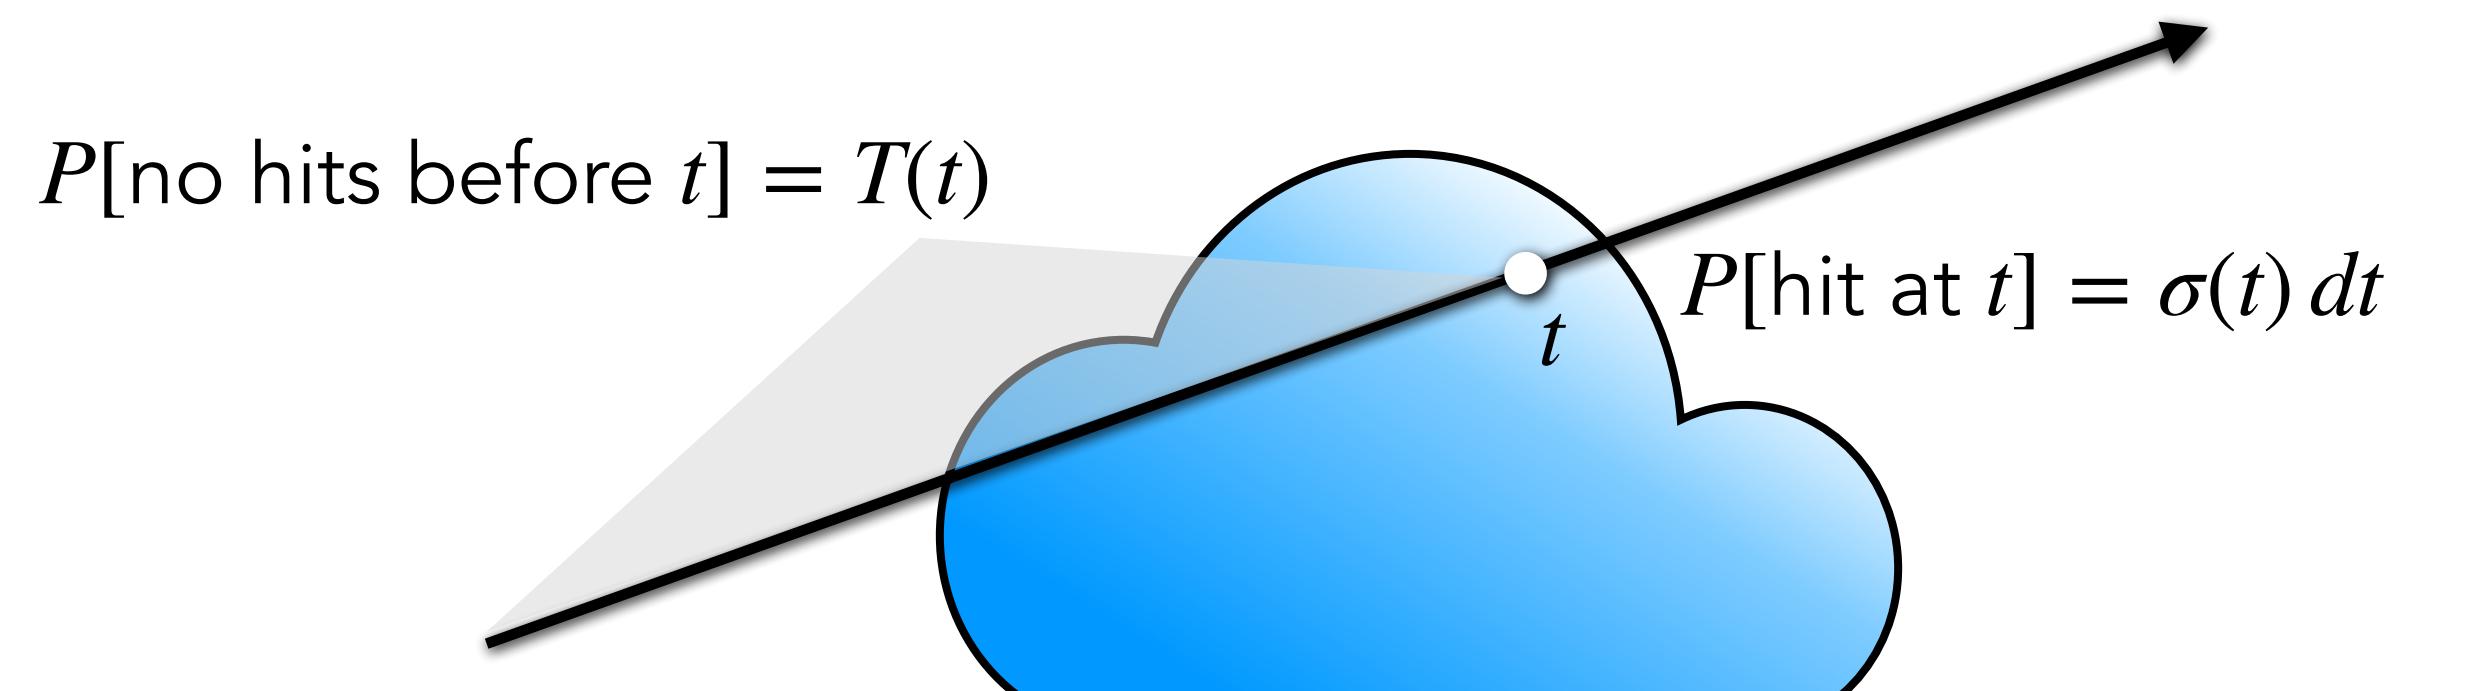
Volume Rendering



To determine if t is the *first* hit, need to know $T(t)$:
probability that the ray didn't hit any particles earlier.

$T(t)$ is called **Transmittance**.

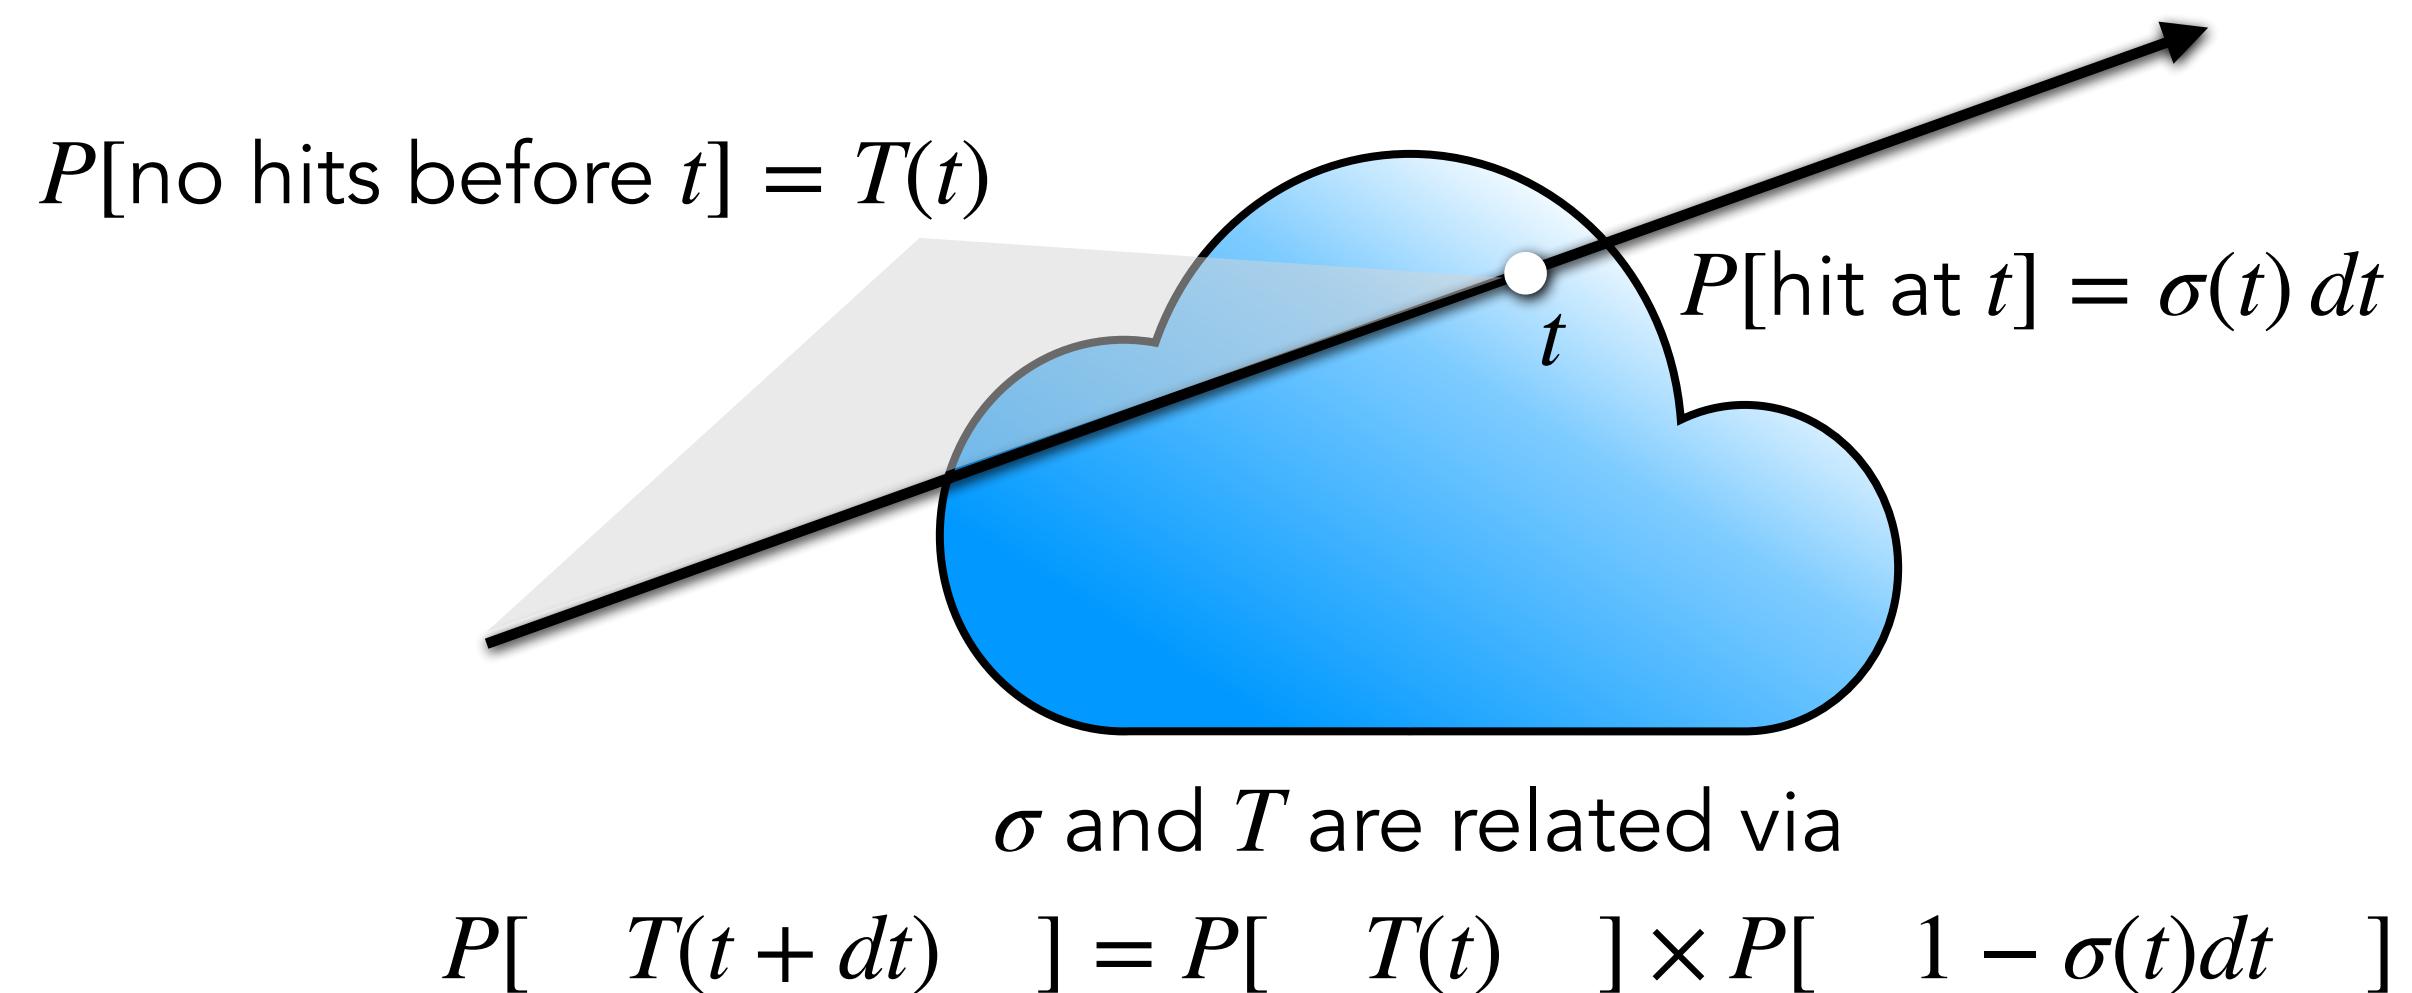
Volume Rendering



σ and T are related via

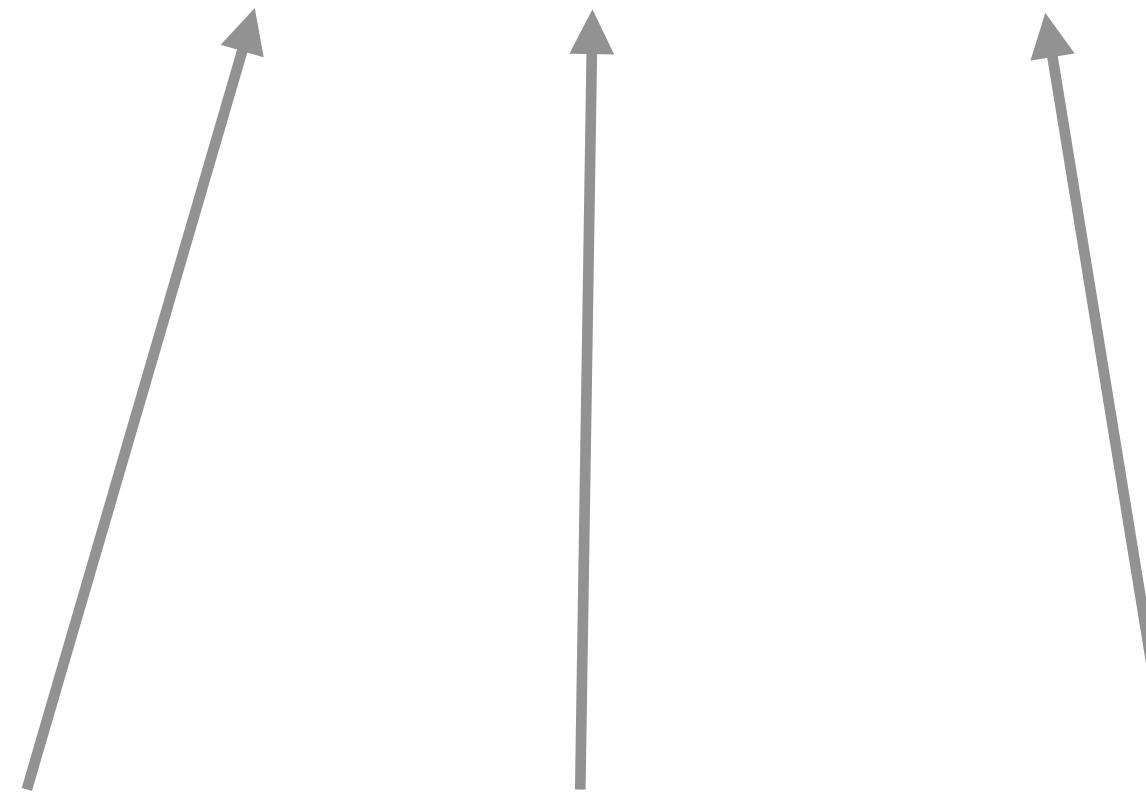
$$P[\text{no hit before } t + dt] = P[\text{no hit before } t] \times P[\text{no hit at } t]$$

Volume Rendering



Volume Rendering

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$



$$P[\quad T(t + dt) \quad] = P[\quad T(t) \quad] \times P[\quad 1 - \sigma(t)dt \quad]$$

Volume Rendering

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

$$\frac{T(t + dt) - T(t)}{dt} = -T(t)\sigma(t)$$

$$T'(t) = -T(t)\sigma(t)$$

$$\frac{T'(t)}{T(t)} = -\sigma(t)$$

$$\int_a^b \frac{T'(t)}{T(t)} dt = - \int_a^b \sigma(t) dt$$

$$\log T(t) \Big|_a^b = - \int_a^b \sigma(t) dt$$

Volume Rendering

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

$$\frac{T(t + dt) - T(t)}{dt} = - T(t)\sigma(t)$$

$$T'(t) = - T(t)\sigma(t)$$

$$\frac{T'(t)}{T(t)} = - \sigma(t)$$

$$\int_a^b \frac{T'(t)}{T(t)} dt = - \int_a^b \sigma(t) dt$$

$$\log T(t) \Big|_a^b = - \int_a^b \sigma(t) dt$$

**Solving simple
differential equation
(for great explanation, see
“Volume Rendering Digest”,
Tagliasacchi et al.)**

Volume Rendering

$$T(t + dt) = T(t)(1 - \sigma(t)dt)$$

$$\frac{T(t + dt) - T(t)}{dt} = -T(t)\sigma(t)$$

$$T'(t) = -T(t)\sigma(t)$$

$$\frac{T'(t)}{T(t)} = -\sigma(t)$$

$$\int_a^b \frac{T'(t)}{T(t)} dt = - \int_a^b \sigma(t) dt$$

$$\log T(t) \Big|_a^b = - \int_a^b \sigma(t) dt$$

$$T(a \rightarrow b) := \frac{T(b)}{T(a)} = \exp \left(- \int_a^b \sigma(t) dt \right)$$

Volume Rendering

$$T(a \rightarrow b) := \frac{T(b)}{T(a)} = \exp\left(-\int_a^b \sigma(t)dt\right)$$

Volume Rendering

$$T(a \rightarrow b) := \frac{T(b)}{T(a)} = \exp\left(-\int_a^b \sigma(t)dt\right)$$

$$T(t) := T(0 \rightarrow t) = \exp\left(-\int_0^t \sigma(a)da\right)$$

Volume Rendering

No hits before t is equal to integral over density up until t .

$$T(t) = \exp\left(-\int_0^t \sigma(a)da\right)$$

Volume Rendering

No hits before t is equal to integral over density up until t .

$$T(t) = \exp\left(-\int_0^t \sigma(a)da\right)$$

$1 - T(t)$ can be seen as cumulative distribution function of probability that ray hits something before reaching t .

Volume Rendering

No hits before t is equal to integral over density up until t .

$$T(t) = \exp\left(-\int_0^t \sigma(a)da\right)$$

$1 - T(t)$ can be seen as cumulative distribution function of probability that ray hits something before reaching t .

Then $T'(t) = T(t)\sigma(t)$ is probability that ray stops *exactly* at t .

Questions?

Volume Rendering

Then $T'(t) = T(t)\sigma(t)$ is probability that ray stops *exactly* at t .

Volume Rendering

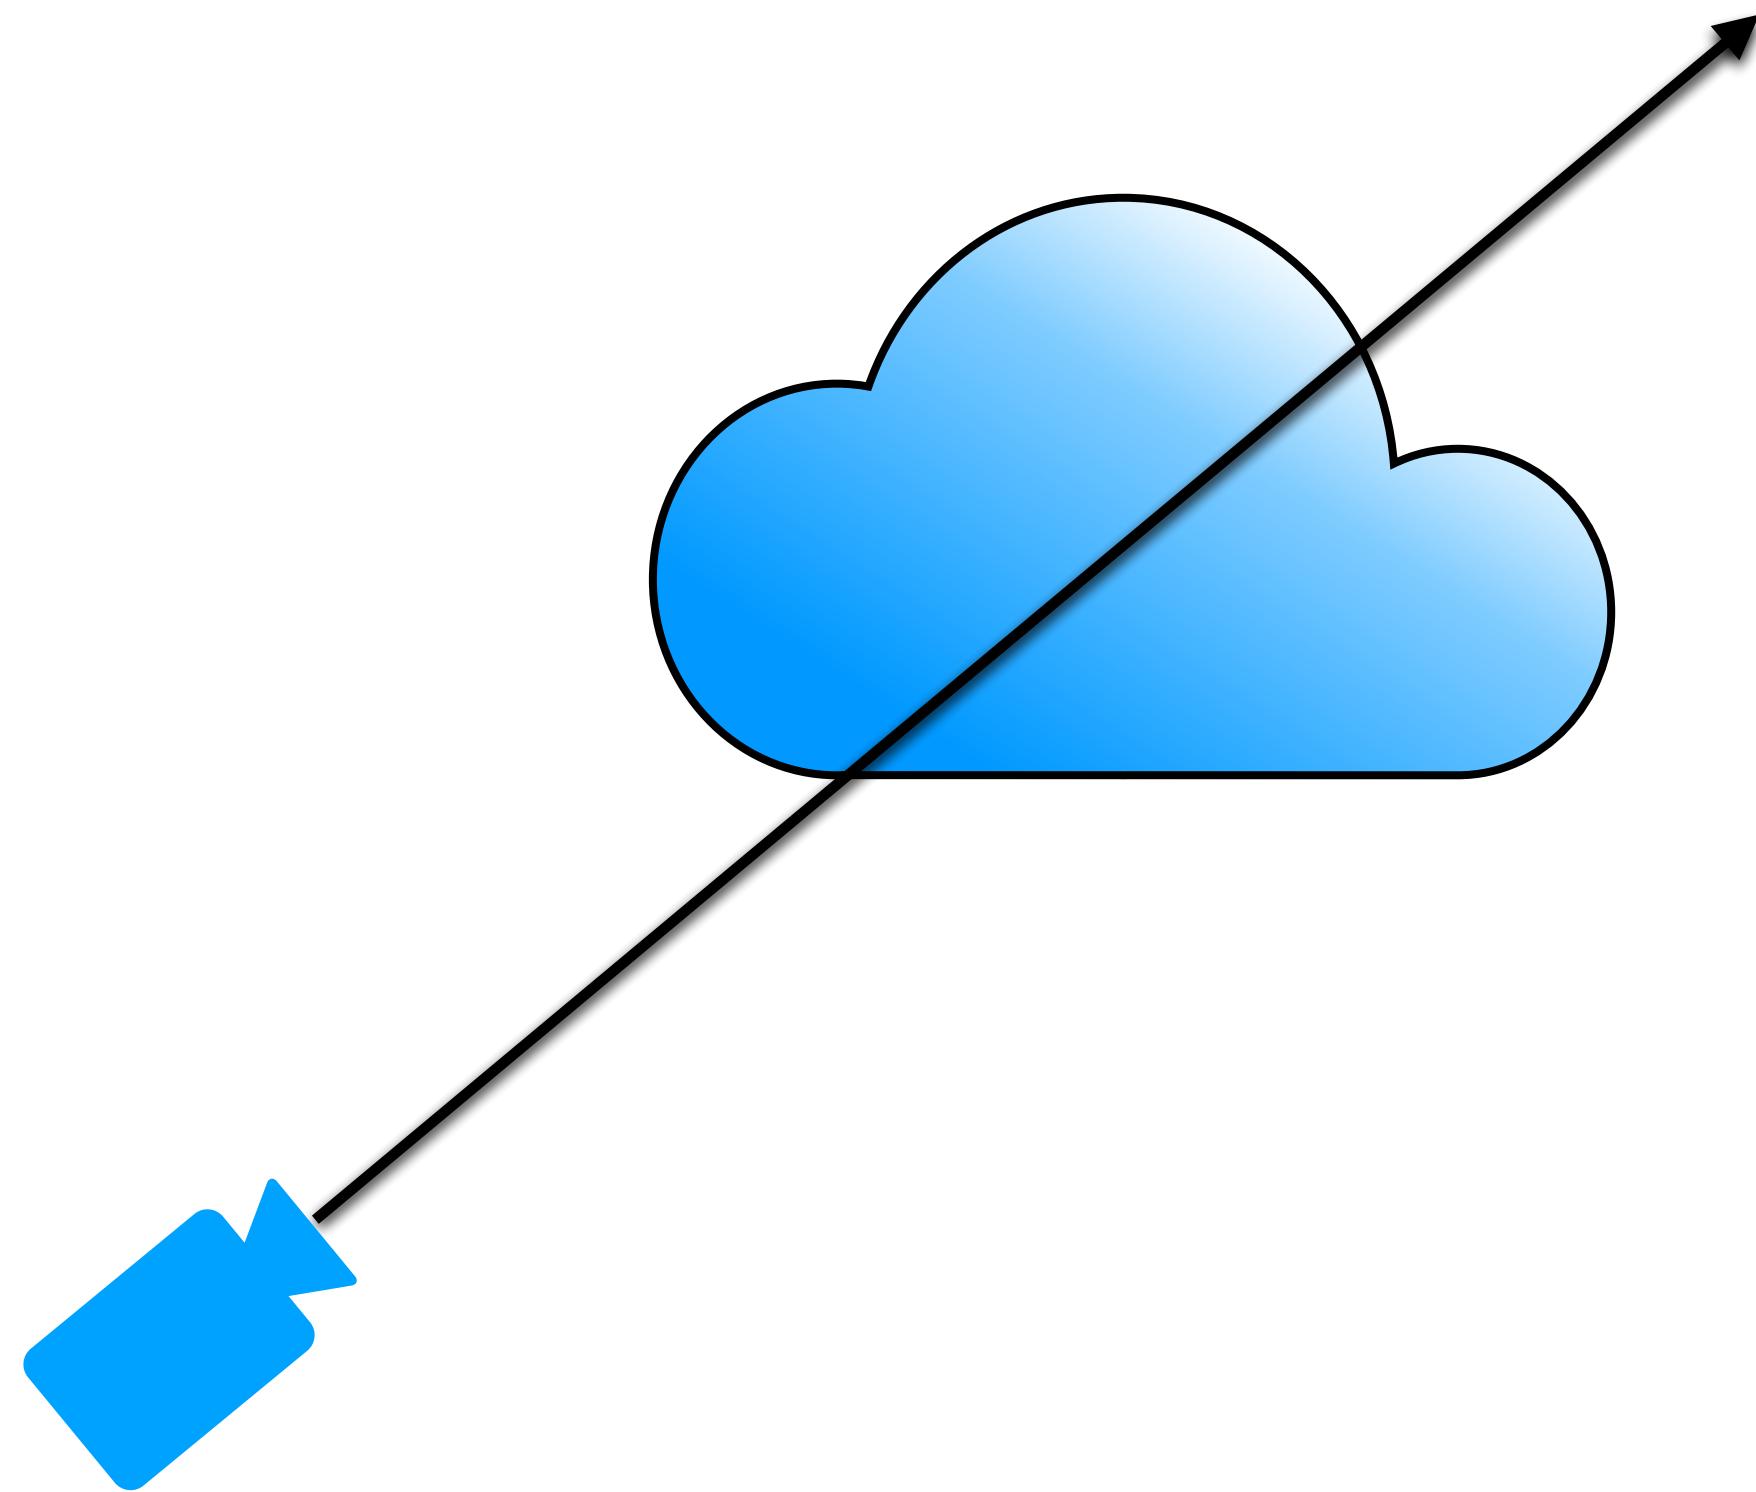
Then $T'(t) = T(t)\sigma(t)$ is probability that ray stops *exactly* at t .

So the expected color returned by the ray will be

$$\int_{t_0}^{t_1} T(t)\sigma(t)\mathbf{c}(t) dt$$

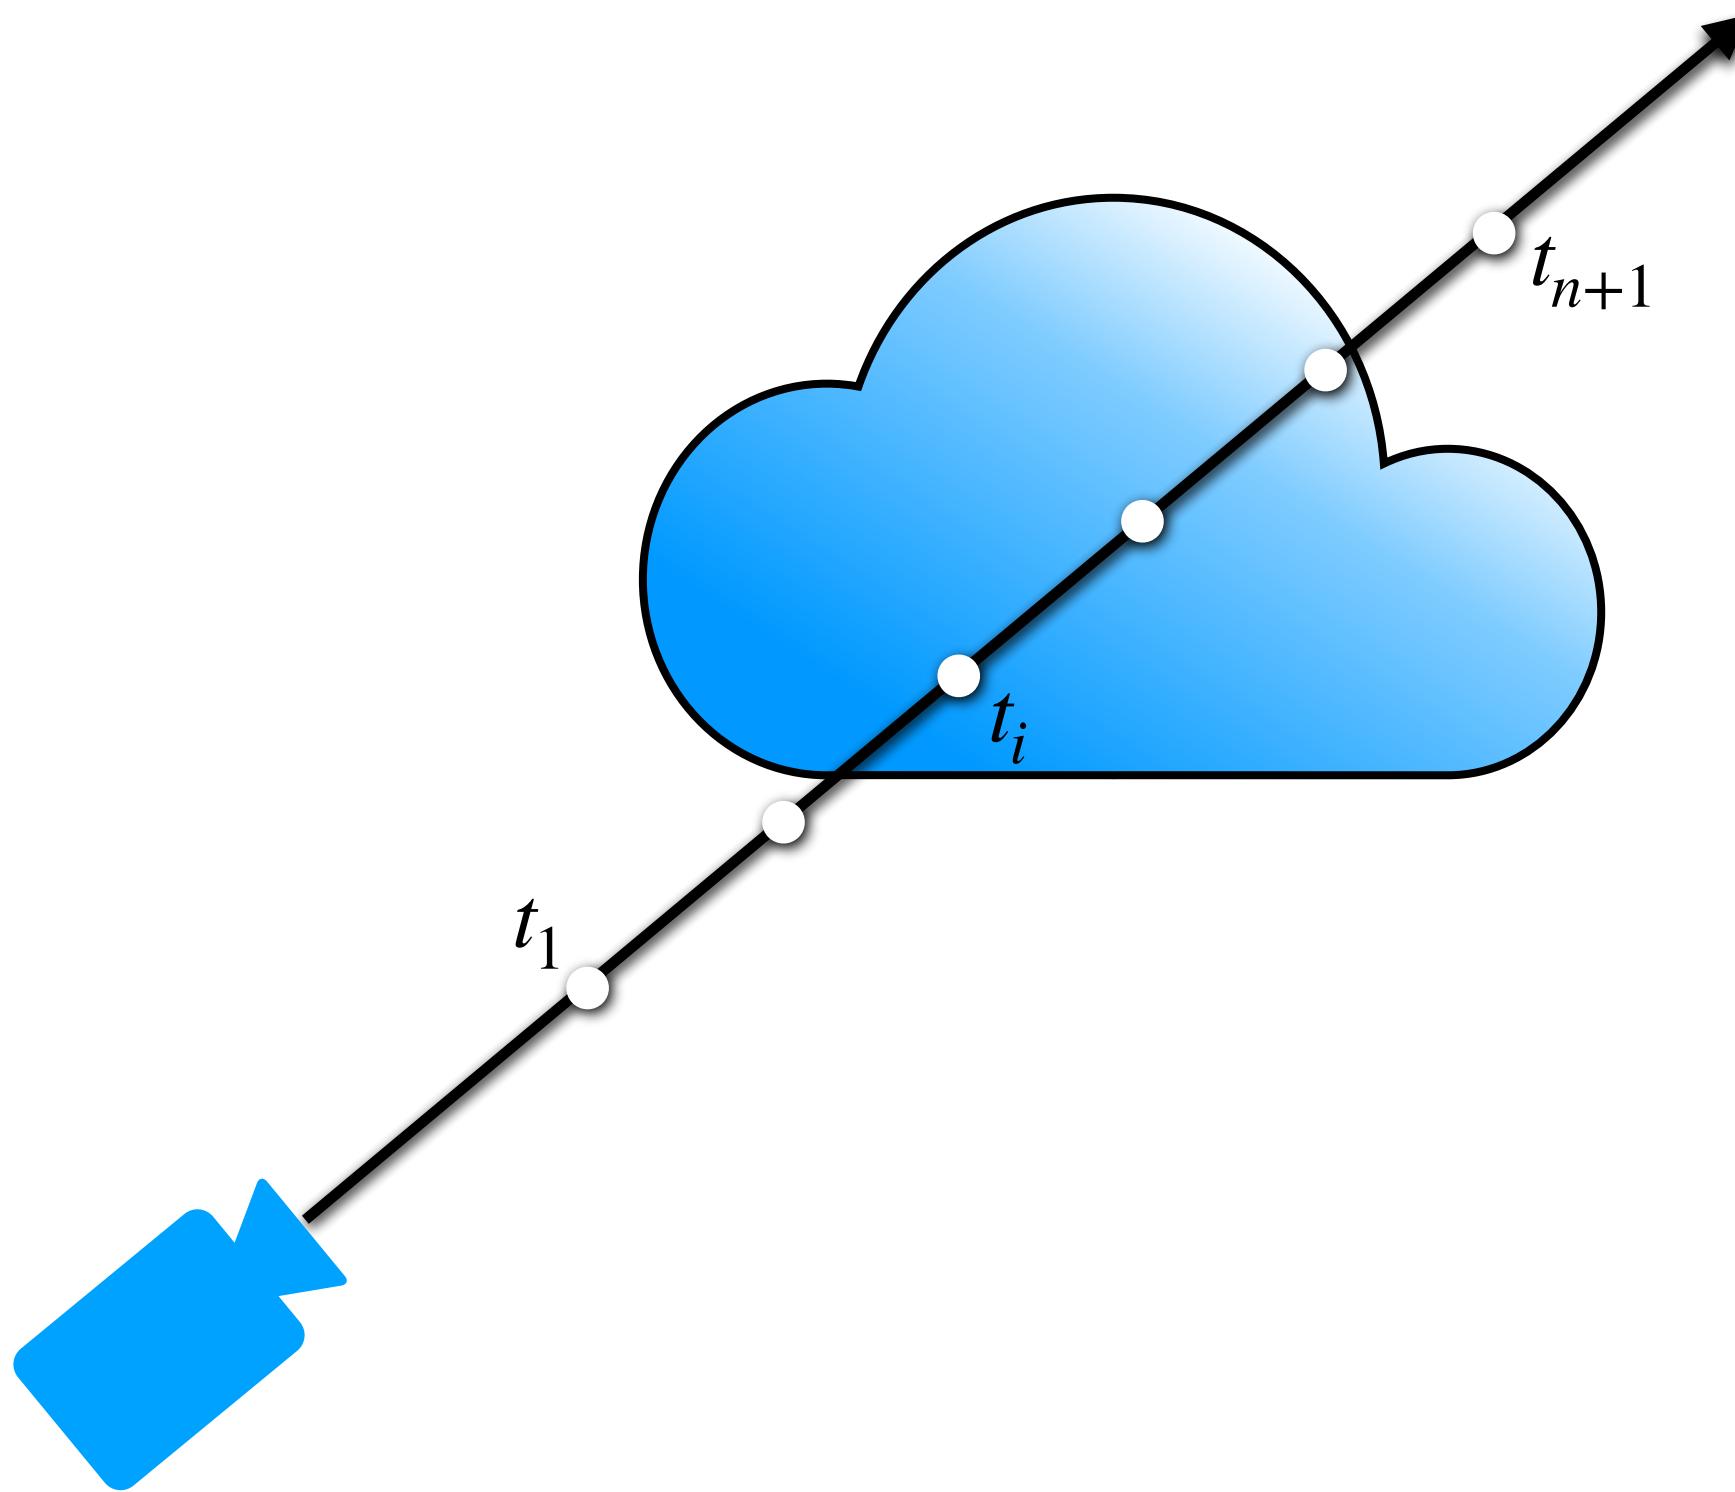
Note the nested integral!

Approximating the nested integral



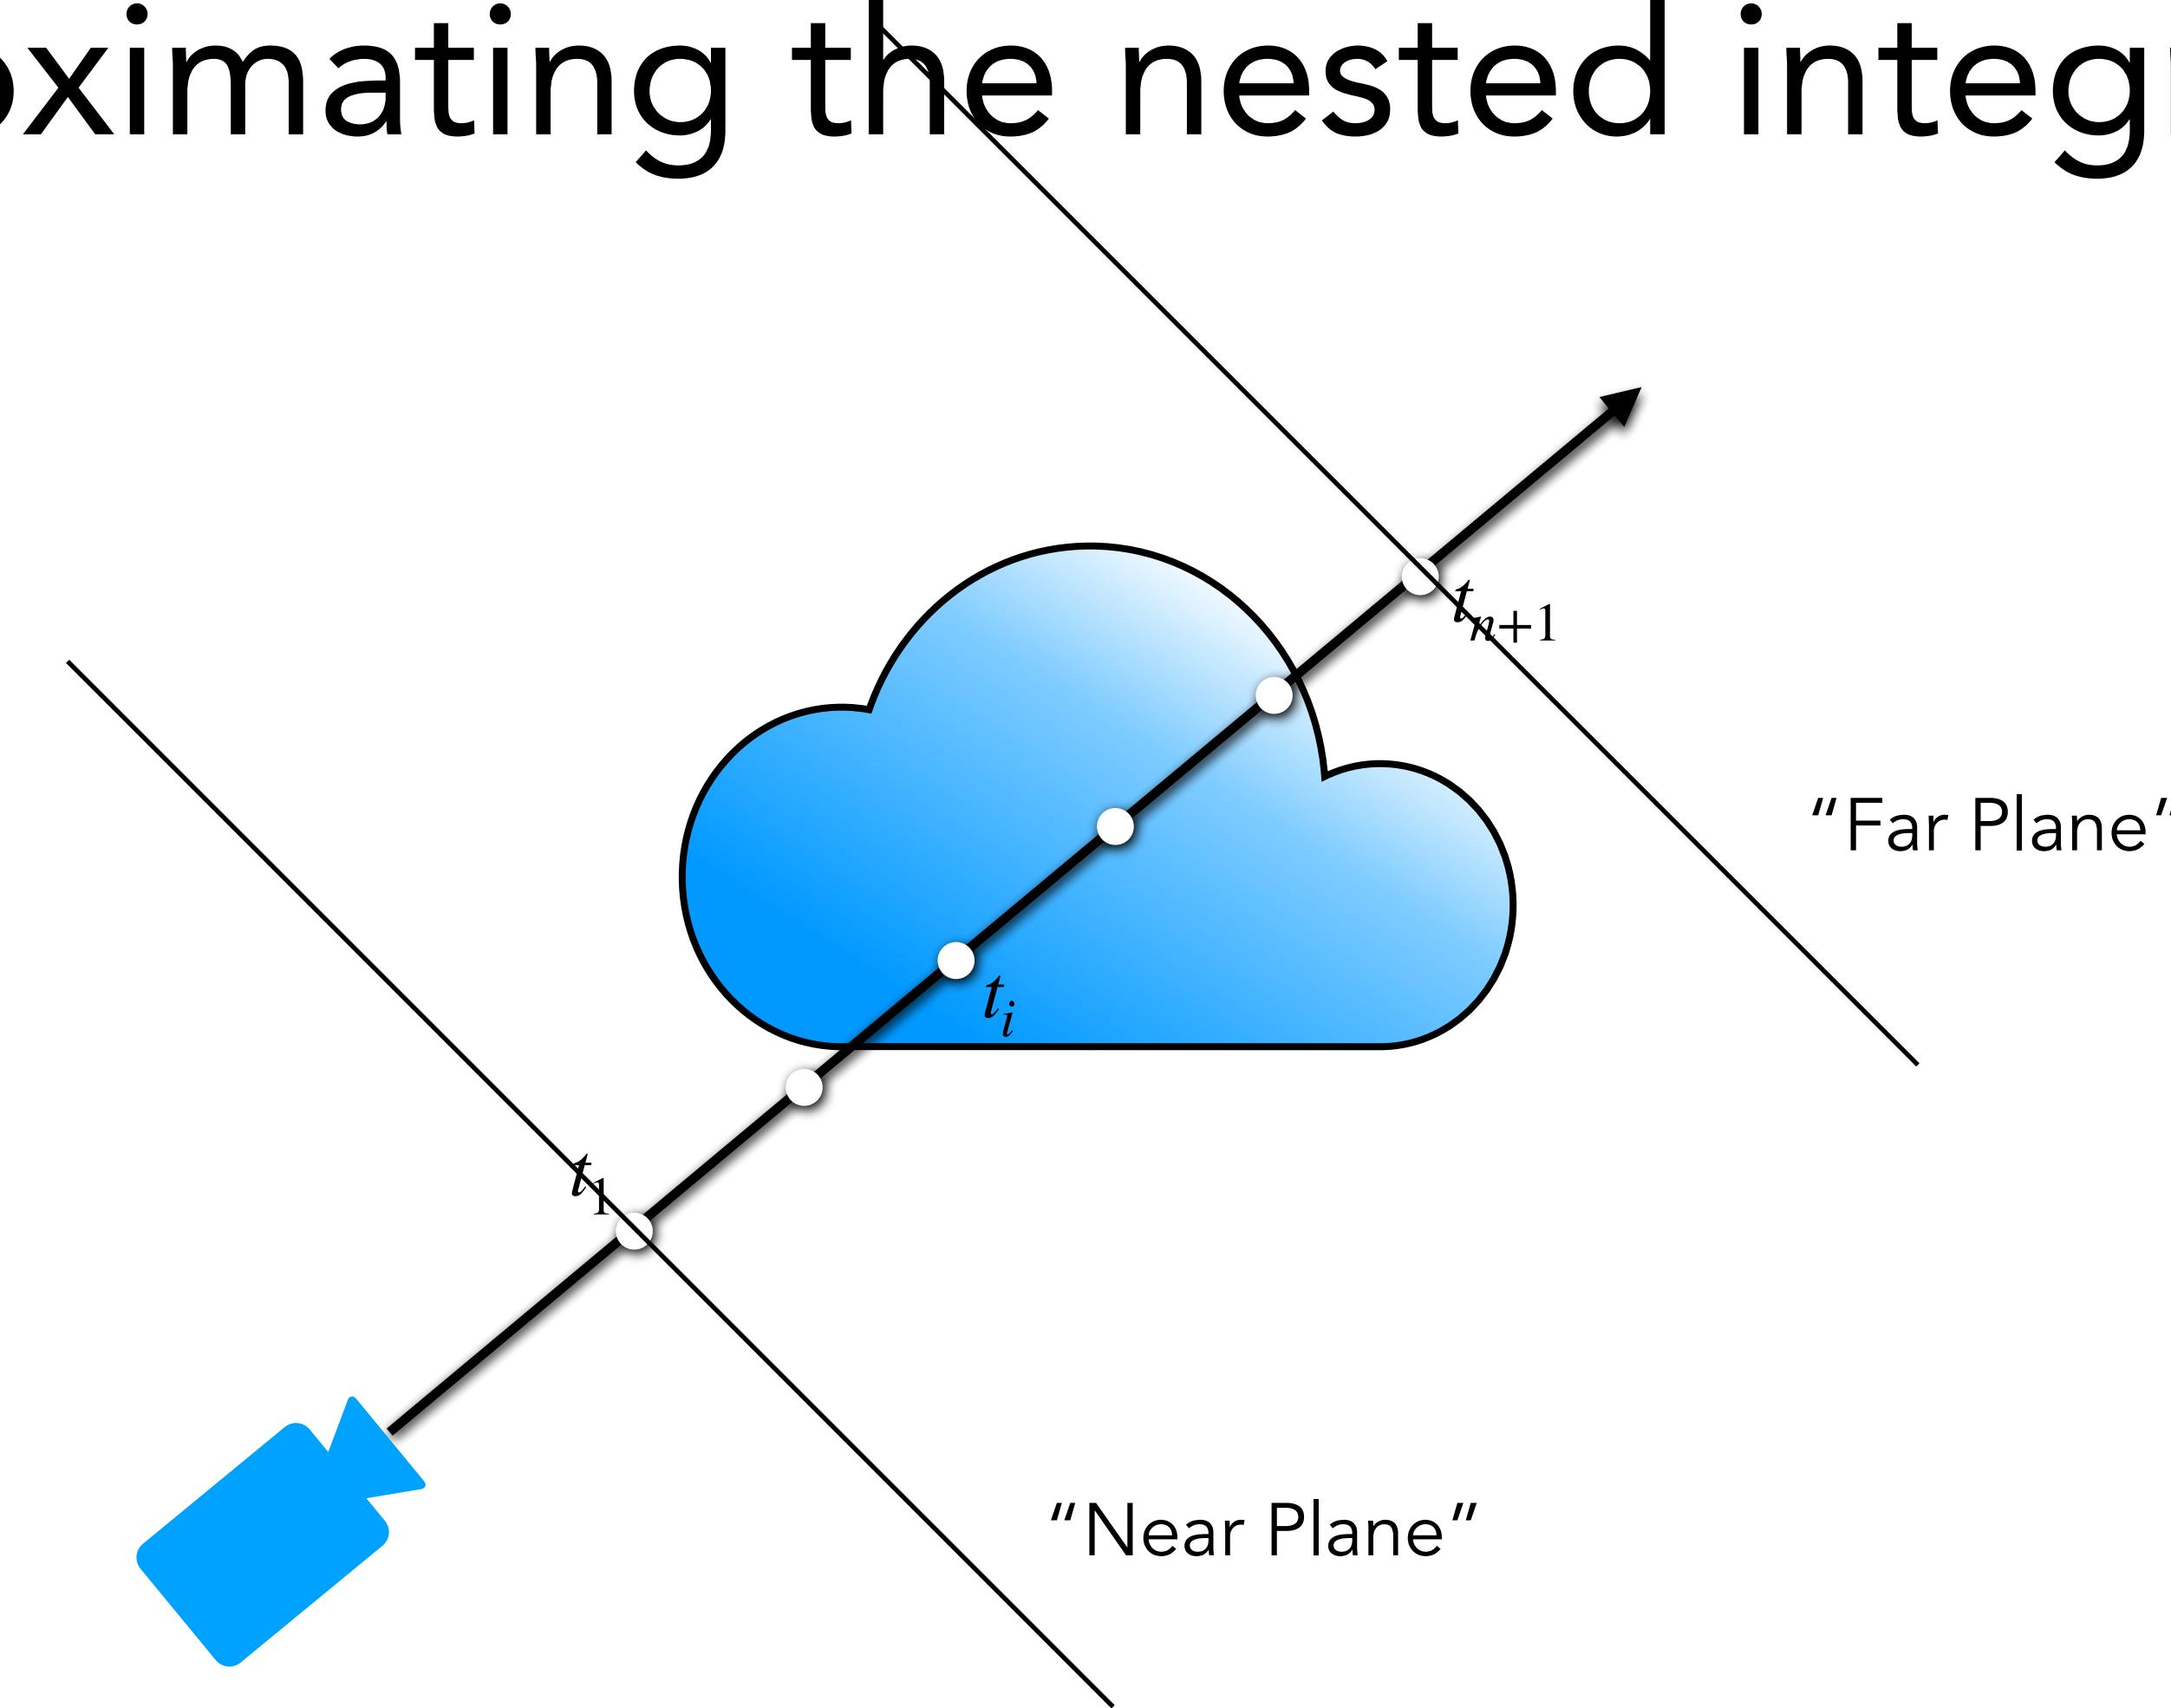
We use quadrature to approximate the nested integral,

Approximating the nested integral



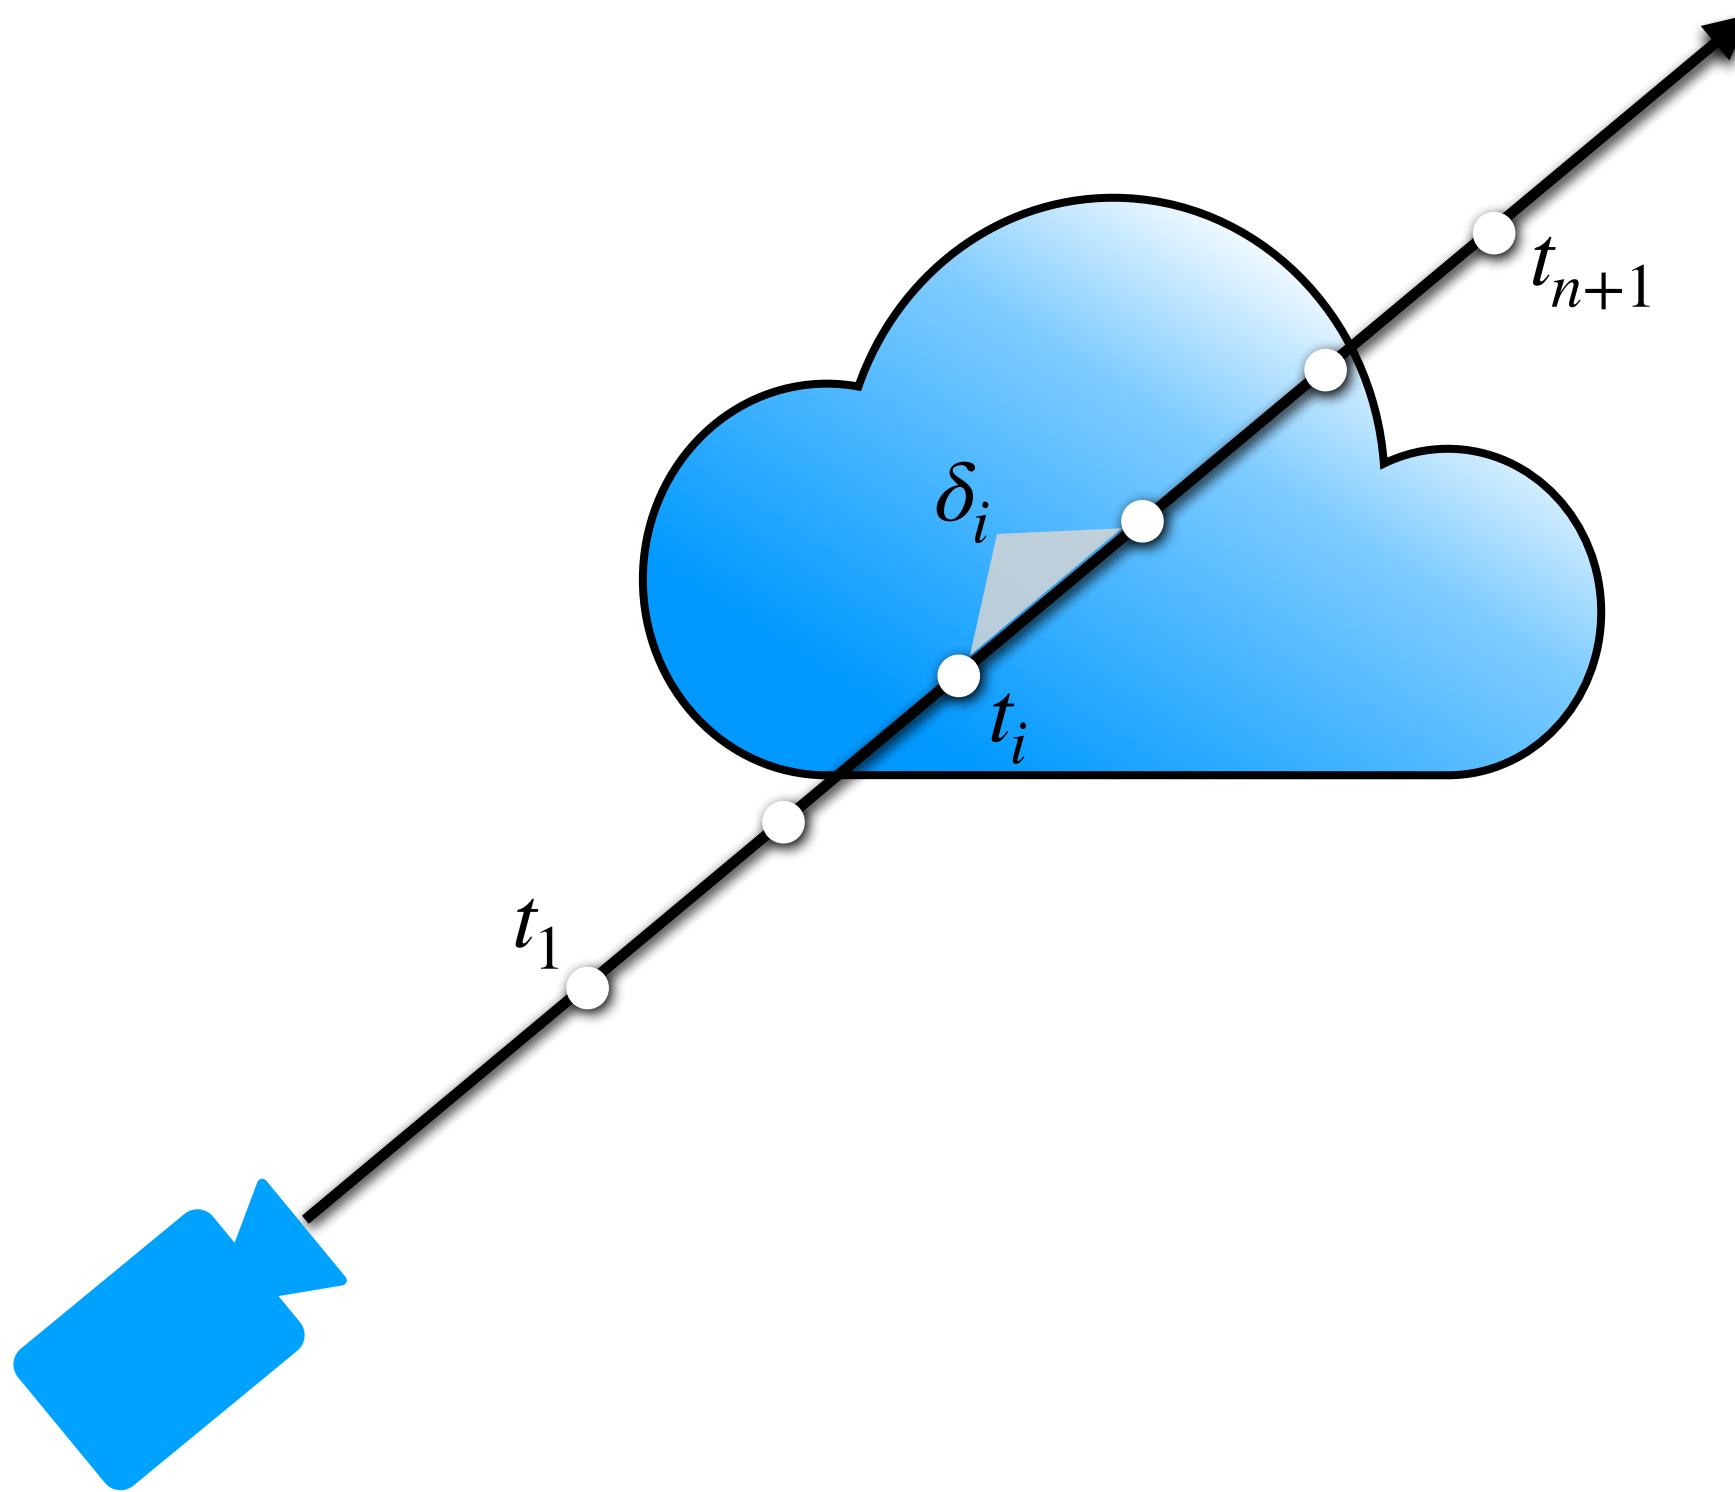
We use quadrature to approximate the nested integral,
splitting the ray up into n segments with endpoints $\{t_1, t_2, \dots, t_{n+1}\}$

Approximating the nested integral



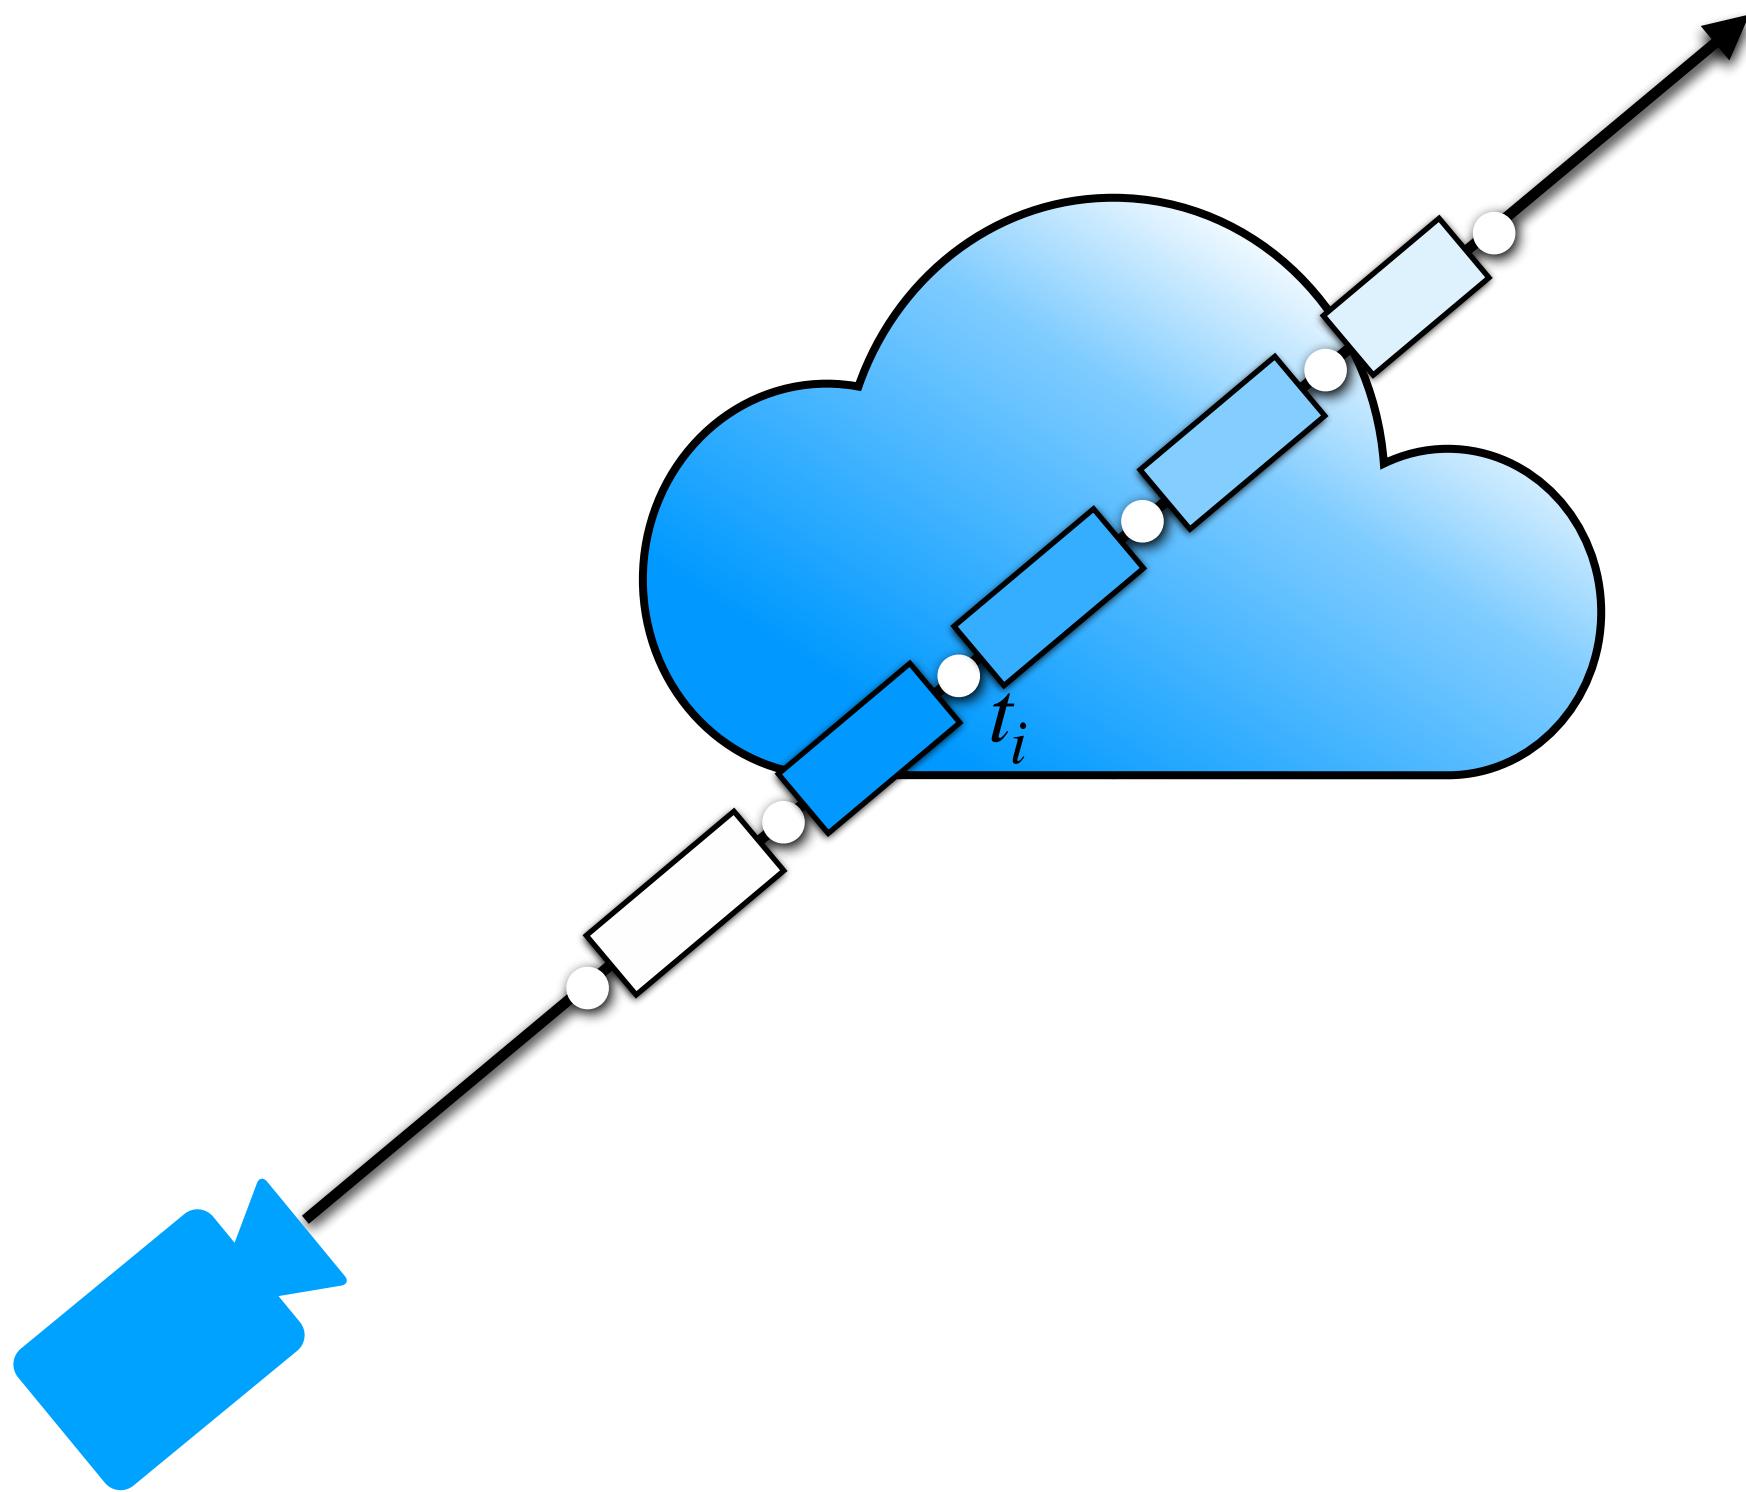
We use quadrature to approximate the nested integral,
splitting the ray up into n segments with endpoints $\{t_1, t_2, \dots, t_{n+1}\}$

Approximating the nested integral



We use quadrature to approximate the nested integral,
splitting the ray up into n segments with endpoints $\{t_1, t_2, \dots, t_{n+1}\}$
with lengths $\delta_i = t_{i+1} - t_i$

Approximating the nested integral



We assume volume density and color are roughly constant within each interval

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx$$

This allows us to break the outer integral

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

This allows us to break the outer integral into a sum of analytically tractable integrals

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

Caveat: piecewise constant density and color
do not imply constant transmittance!

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

Caveat: piecewise constant density and color
do not imply constant transmittance!

Important to account for how early part of a segment blocks later part when σ_i is high

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

For $t \in [t_i, t_{i+1}]$, $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

For $t \in [t_i, t_{i+1}]$, $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$



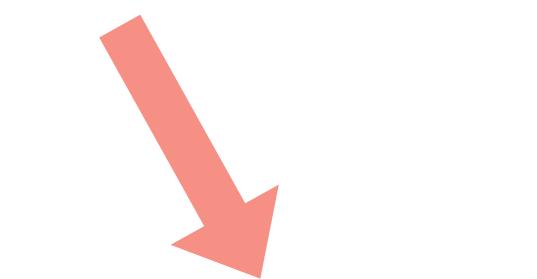
$$\exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right) = T_i \quad \text{"How much is blocked by all previous segments?"}$$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

For $t \in [t_i, t_{i+1}]$, $T(t) = \exp\left(-\int_{t_1}^{t_i} \sigma_i ds\right) \exp\left(-\int_{t_i}^t \sigma_i ds\right)$

"How much is blocked partway through the current segment?"



$$\exp(-\sigma_i(t - t_i))$$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i \mathbf{c}_i dt$$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

Substitute

$$= \sum_{i=1}^n T_i \sigma_i \mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt$$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

$$= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt$$

Integrate $= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i}$

Approximating the nested integral

$$\int T(t)\sigma(t)\mathbf{c}(t) dt \approx \sum_{i=1}^n \int_{t_i}^{t_{i+1}} T(t)\sigma_i\mathbf{c}_i dt$$

$$= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \int_{t_i}^{t_{i+1}} \exp(-\sigma_i(t - t_i)) dt$$

$$= \sum_{i=1}^n T_i\sigma_i\mathbf{c}_i \frac{\exp(-\sigma_i(t_{i+1} - t_i)) - 1}{-\sigma_i}$$

Cancel σ_i

$$= \sum_{i=1}^n T_i\mathbf{c}_i (1 - \exp(-\sigma_i\delta_i))$$

Summary: volume rendering integral estimate

Rendering model for ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$\mathbf{c} \approx \sum_{i=1}^n T_i \alpha_i \mathbf{c}_i$$

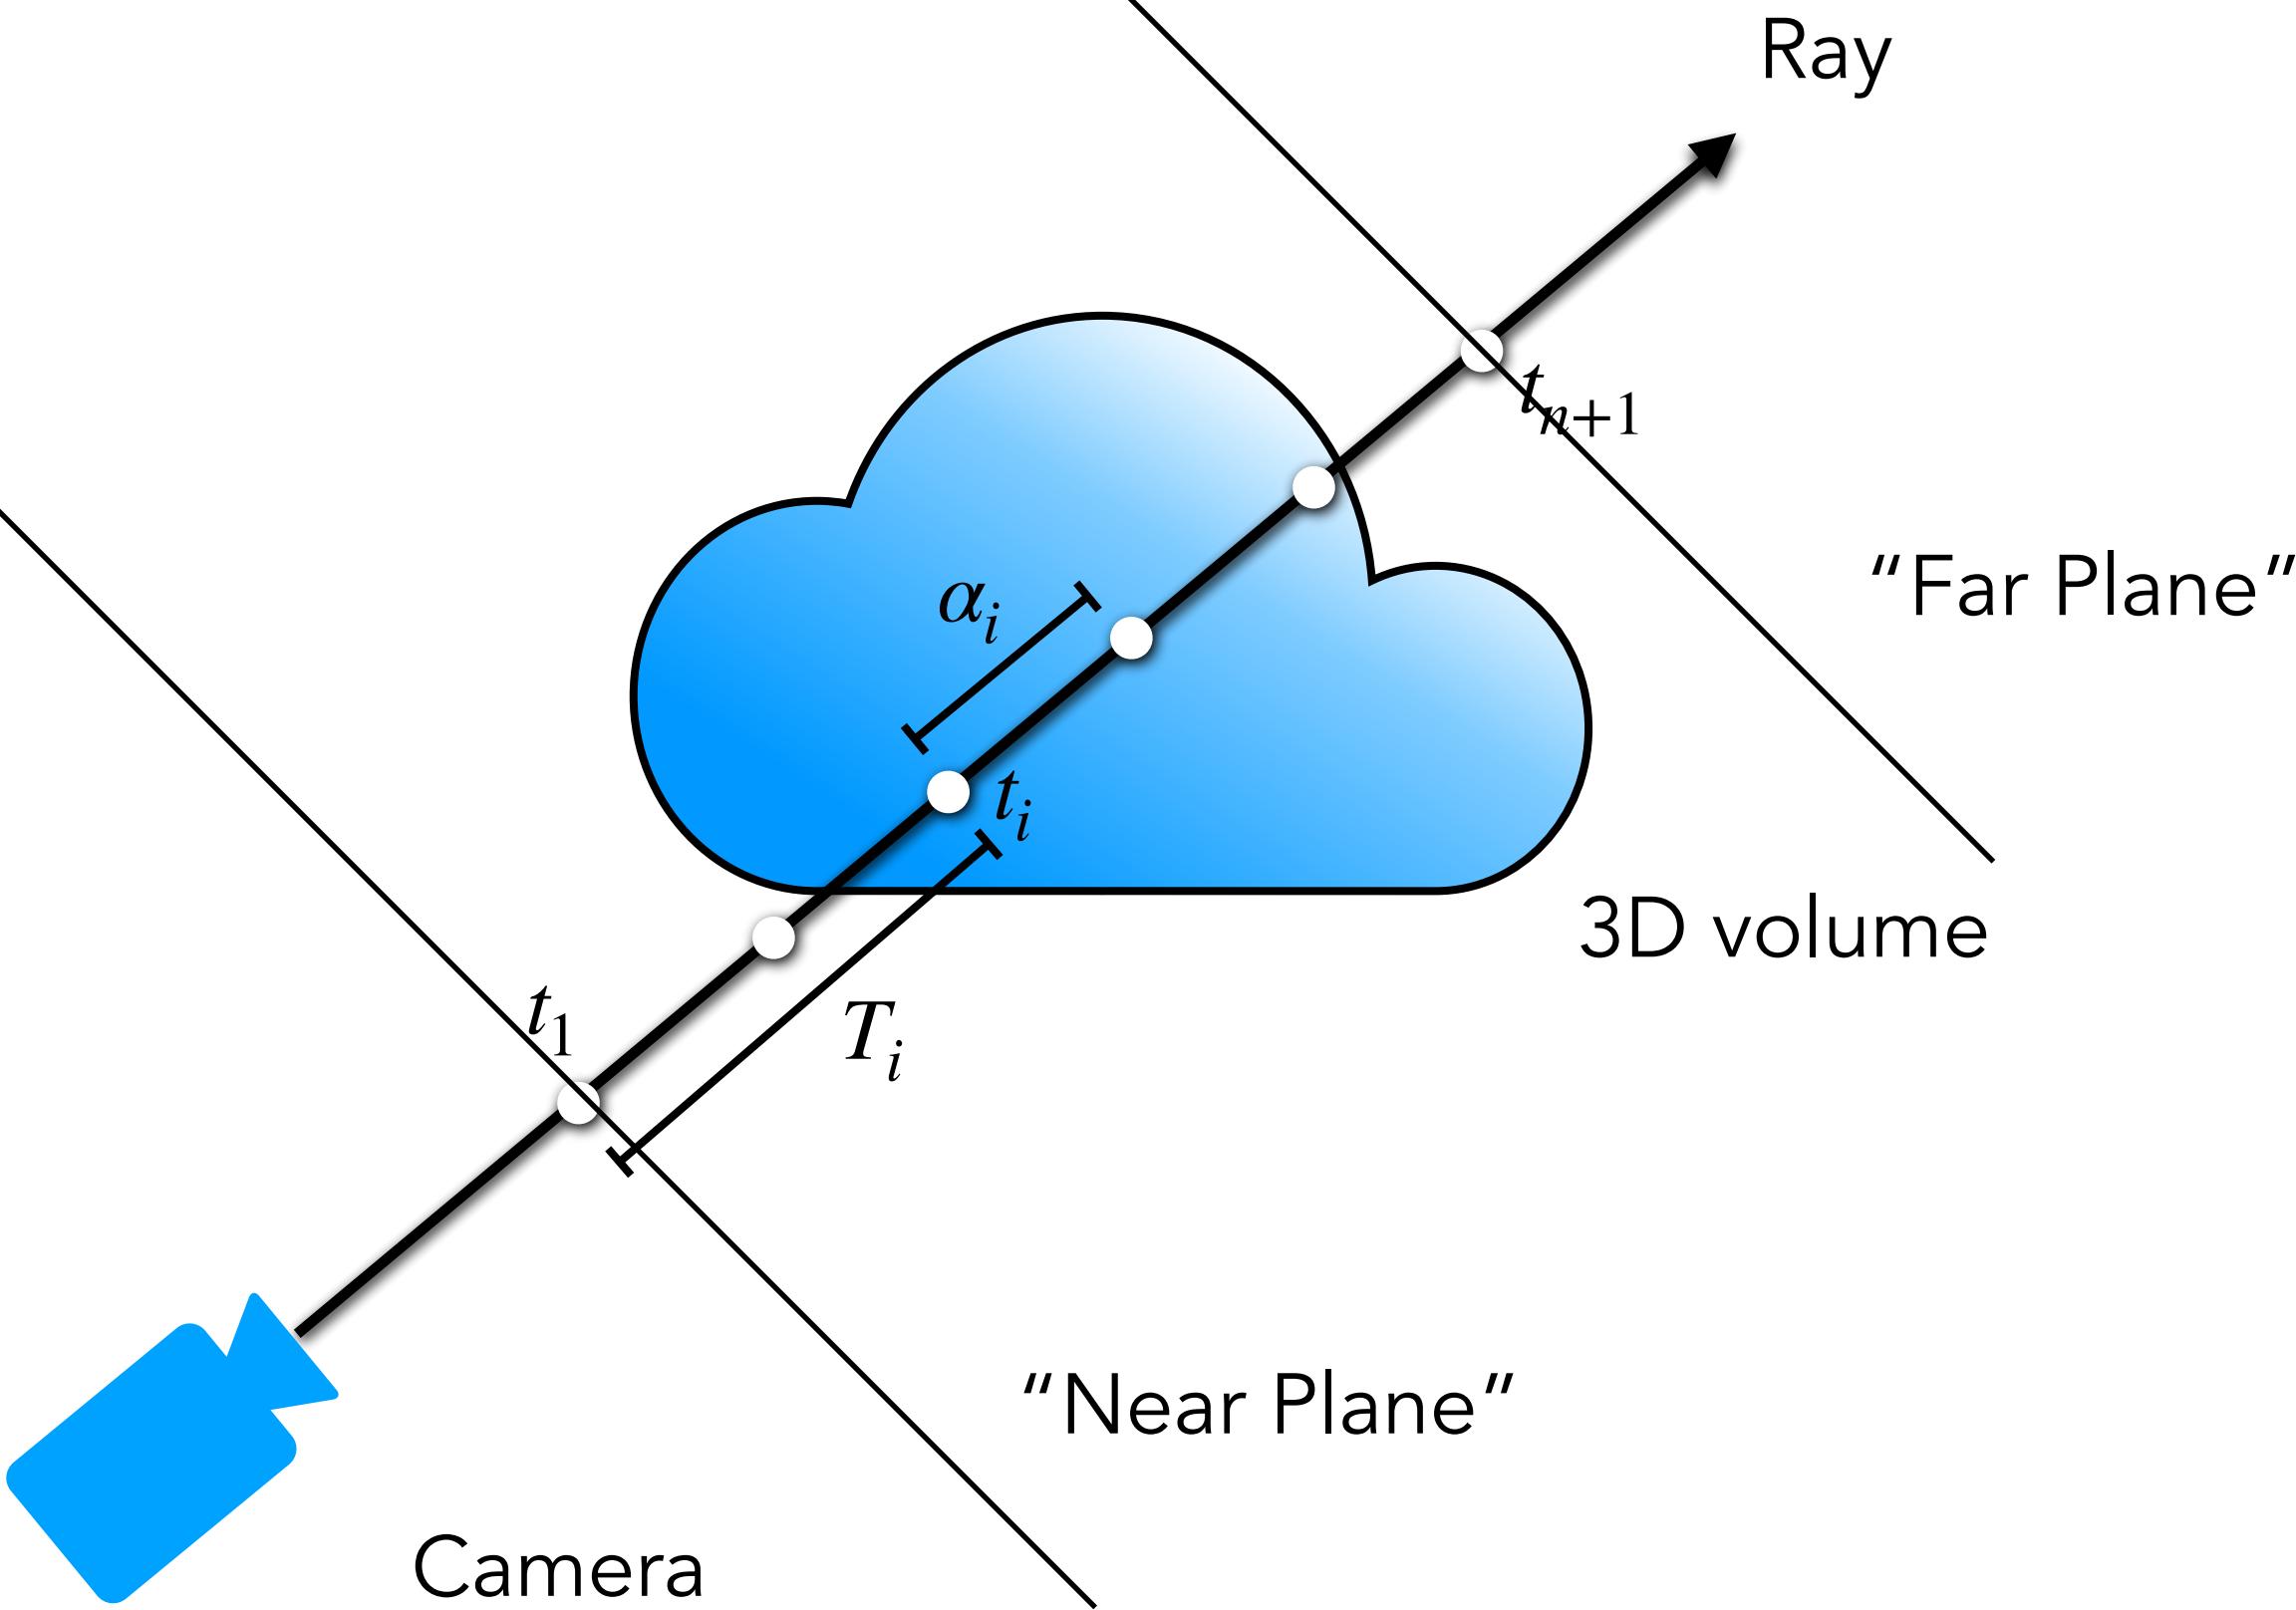
↑
weights colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

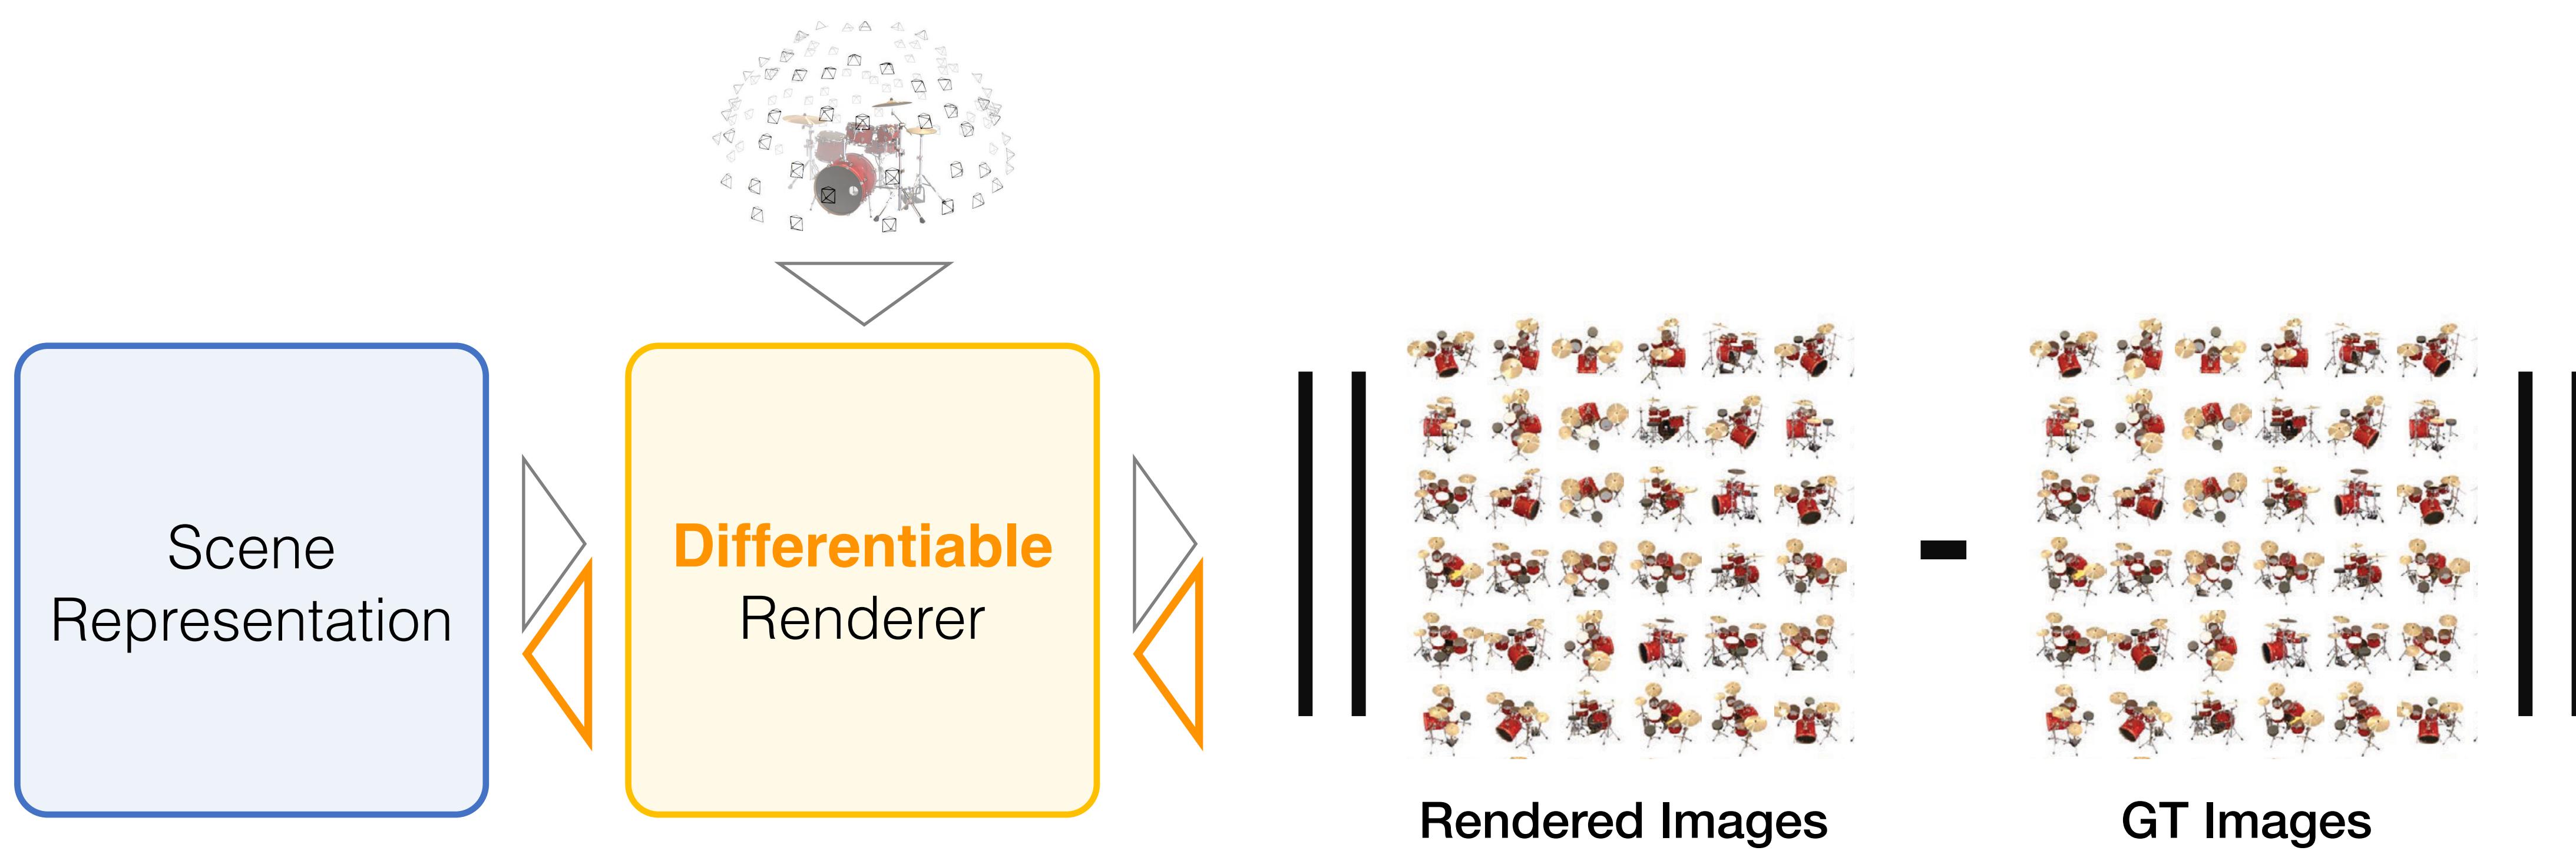
How much light is contributed by ray segment i :

$$\alpha_i = 1 - \exp(-\sigma_i \delta_i)$$



Questions?

Differentiable Renderer is done!



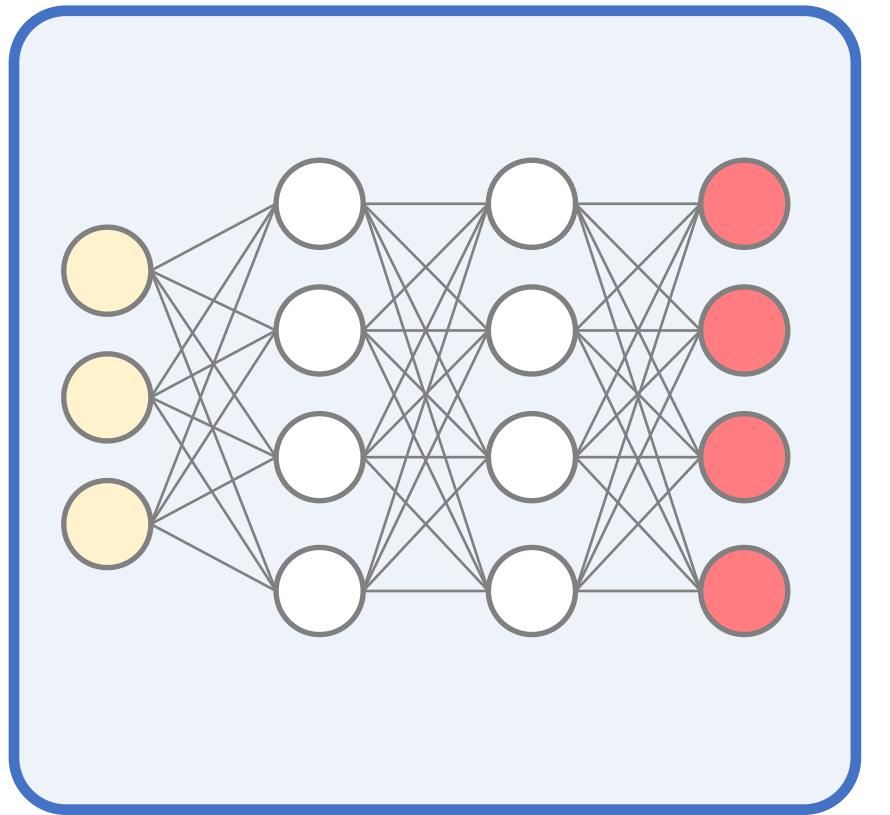
Given an *observable* variable (pixel colors), we will build a differentiable forward model that we then use to estimate *unobserved (latent) variables* (geometry, appearance)!

Test case: Single scene with many observations.

Scene Representation

Neural Renderer

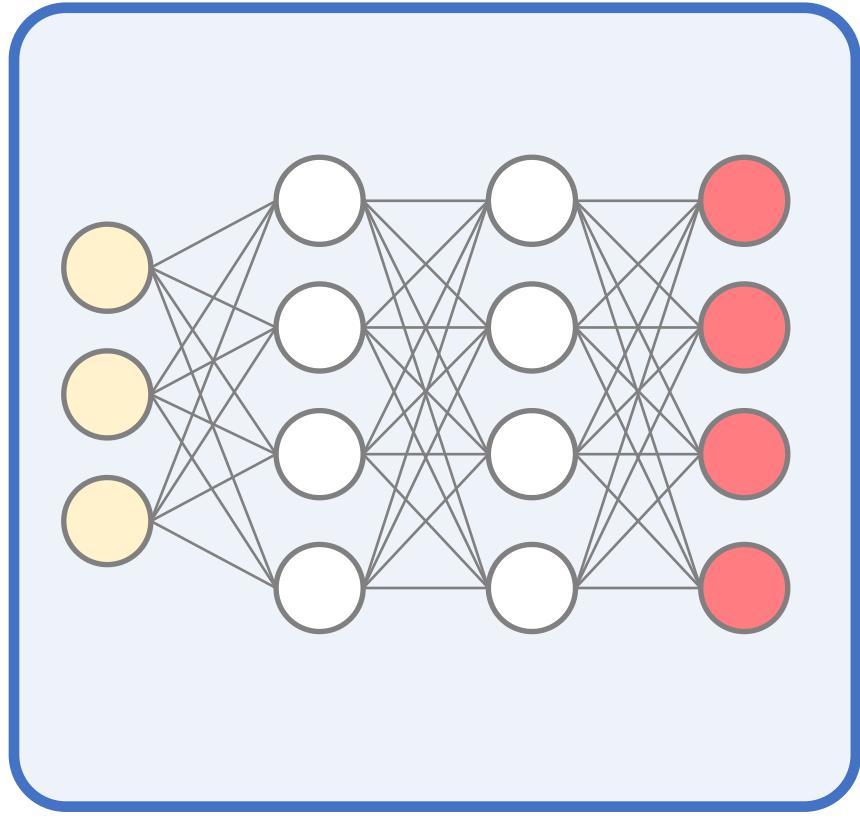
Test case: Single scene with many observations.



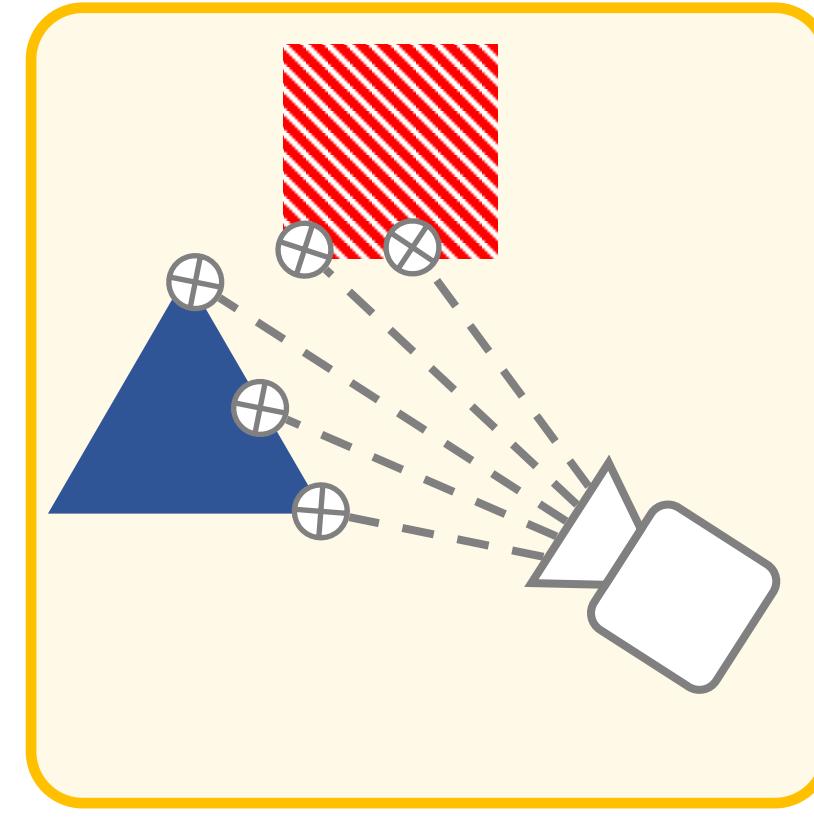
Scene Representation

Neural Renderer

Test case: Single scene with many observations.

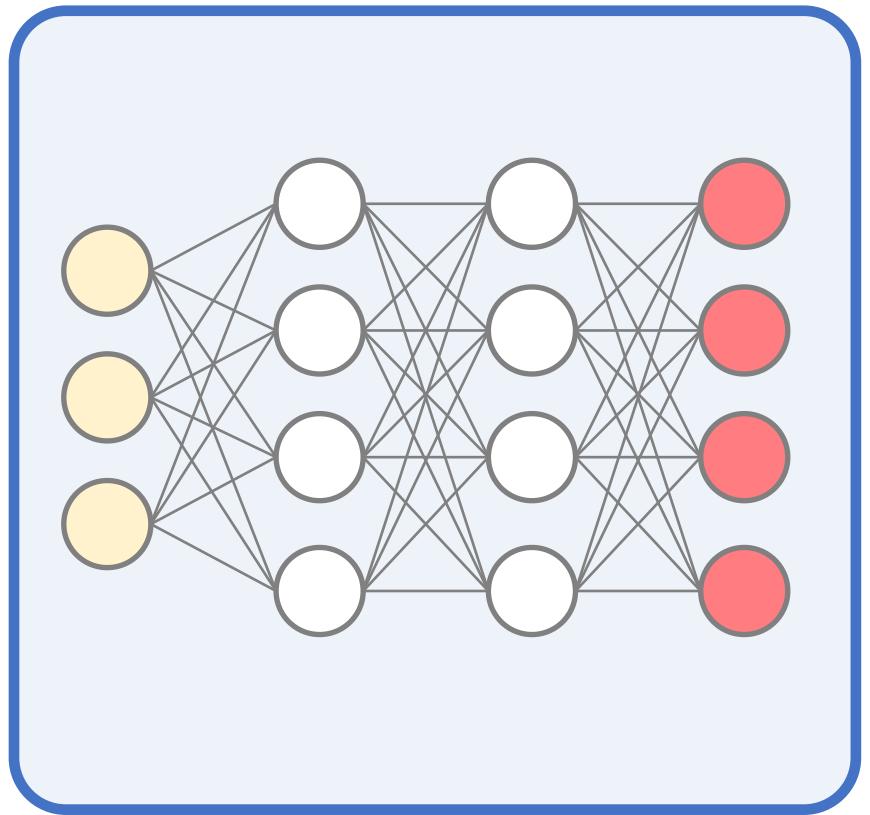


Scene Representation

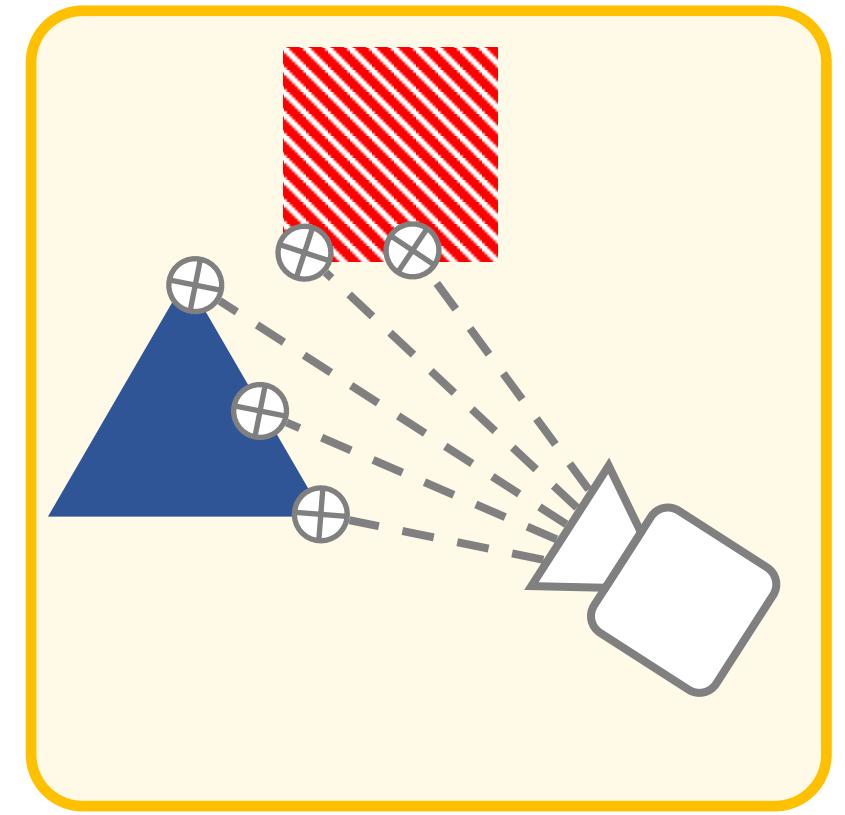


Neural Renderer

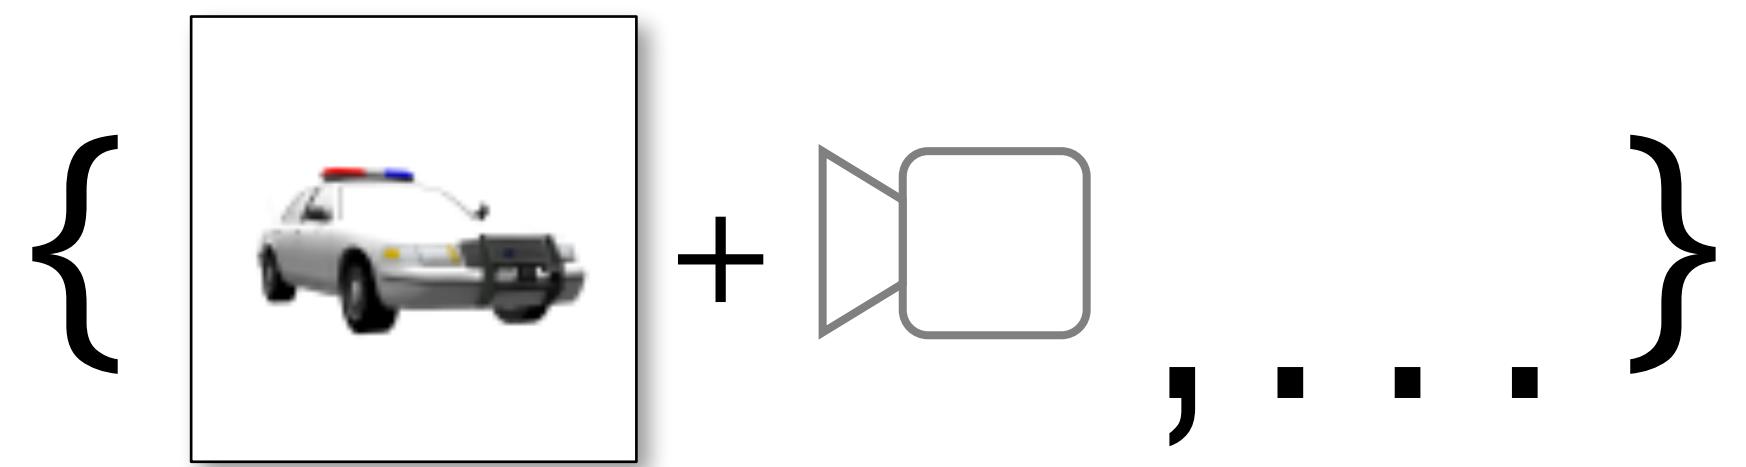
Test case: Single scene with many observations.



Scene Representation

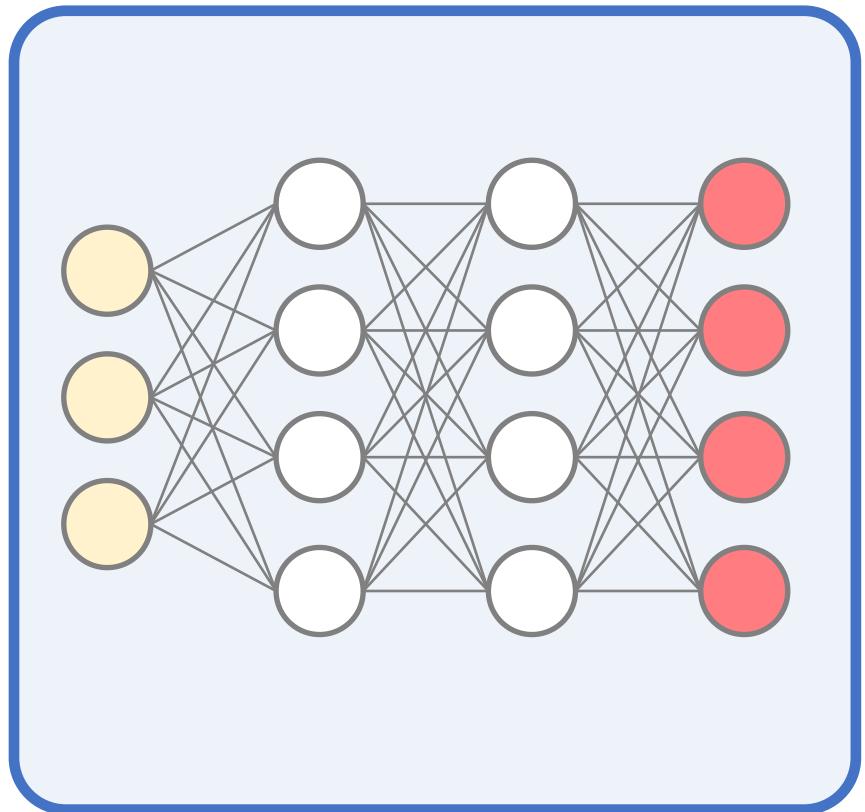


Neural Renderer

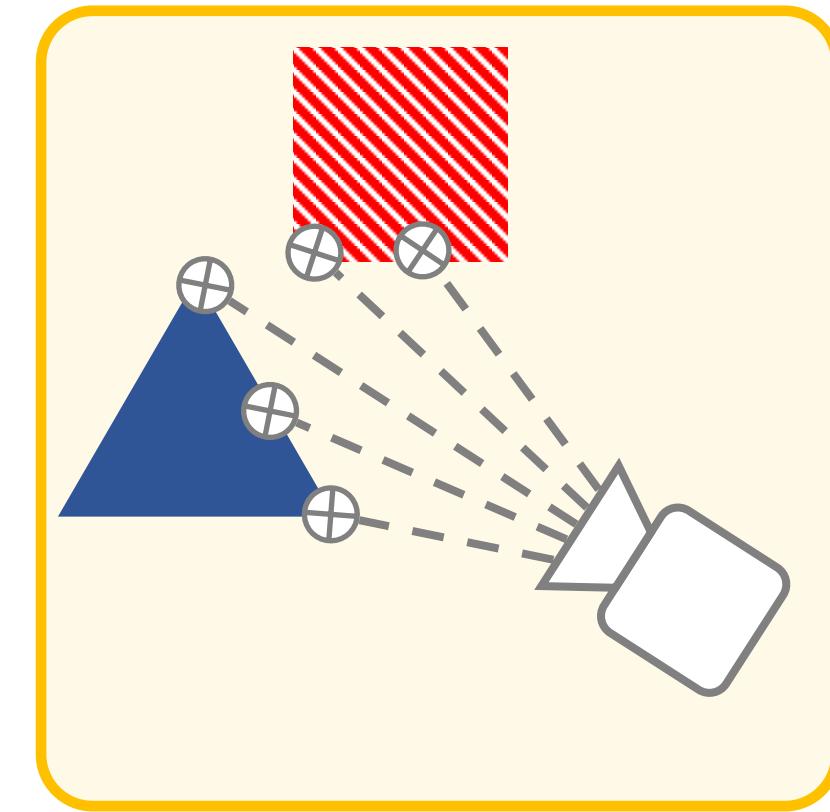


$$\left\{ \left(\mathcal{I}, \xi \right)_i \right\}_{i=1}^N$$

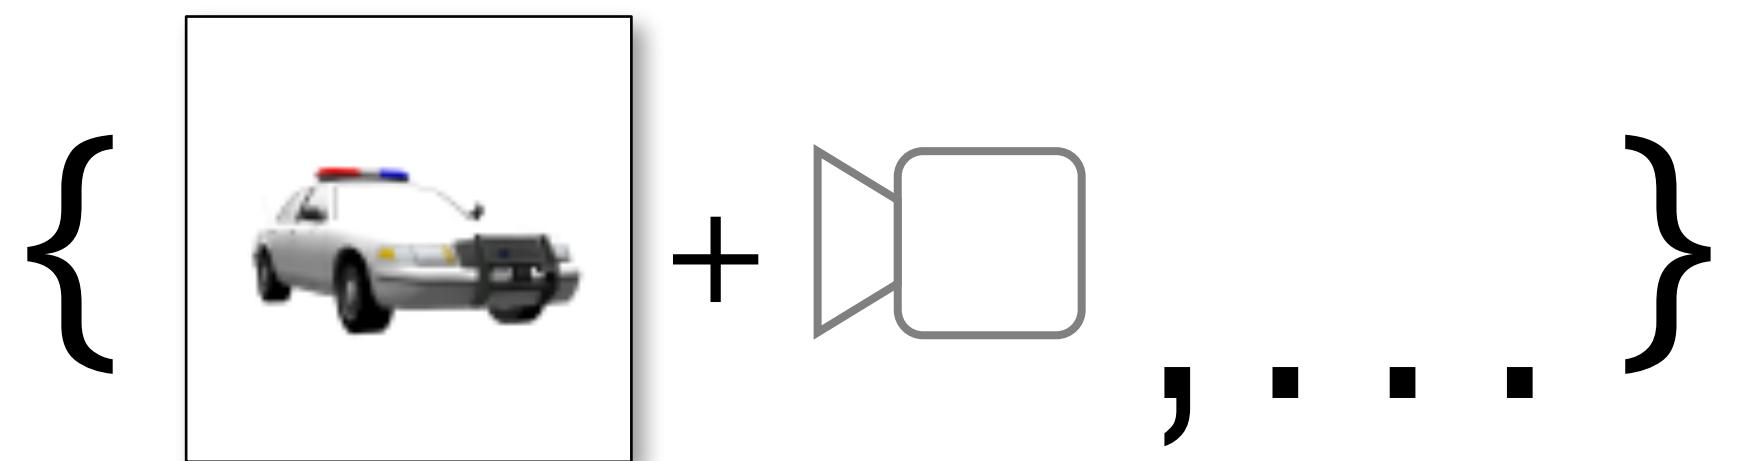
Test case: Single scene with many observations.



Scene Representation



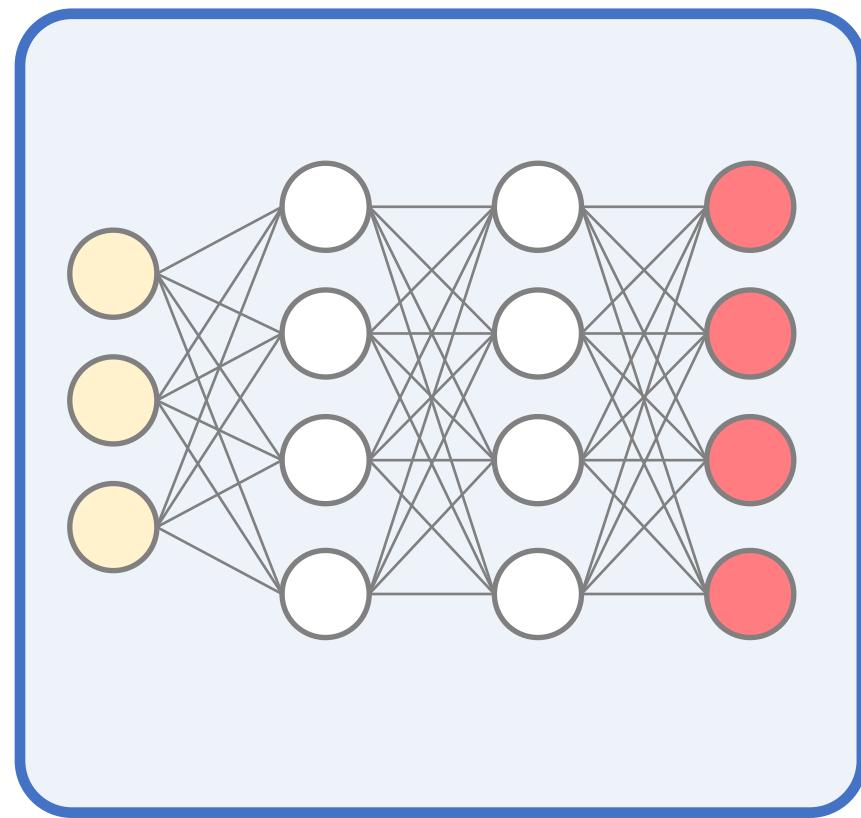
Neural Renderer



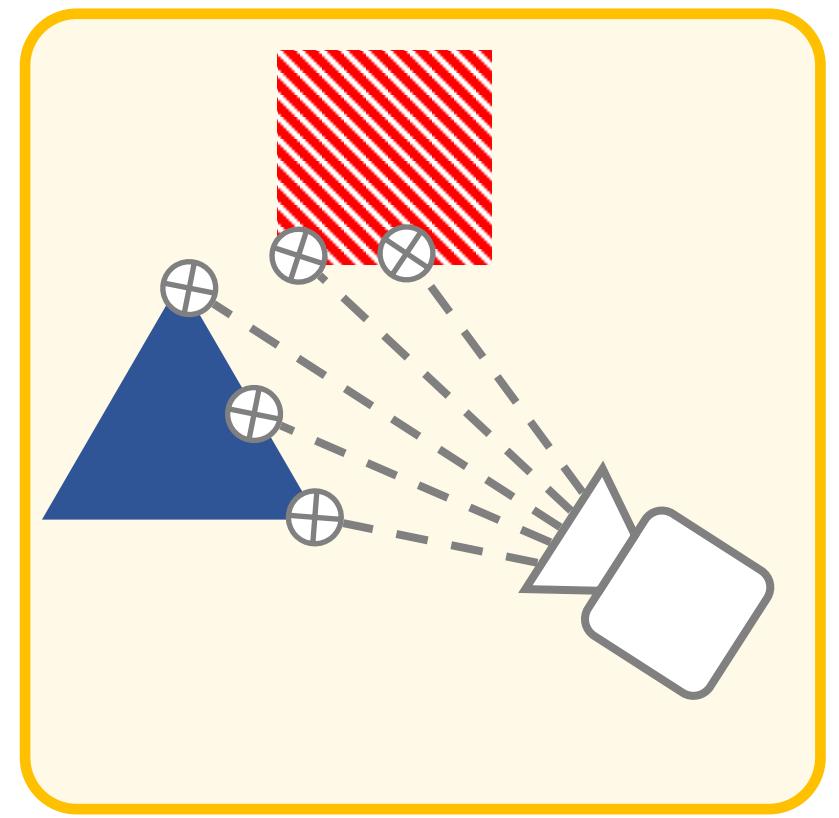
$$\left\{ (\mathcal{I}, \xi)_i \right\}_{i=1}^N$$

$$\operatorname{argmin}_{\phi, \theta} \sum_i \left\| \text{Render}_{\theta}(\text{Srn}_{\phi}, \xi_i) - \mathcal{I}_i \right\|$$

Test case: Single scene with many observations.



Scene Representation



Neural Renderer

Optimization

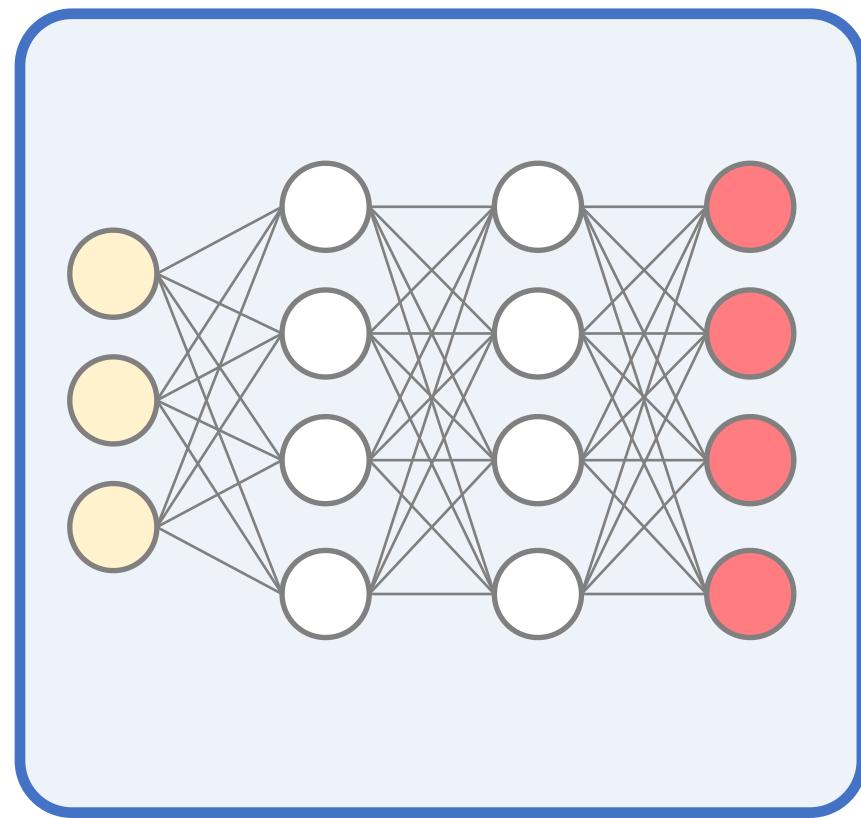


$$\{ \quad \text{[Image of a car]} + \text{[Camera icon]} , \dots \}$$

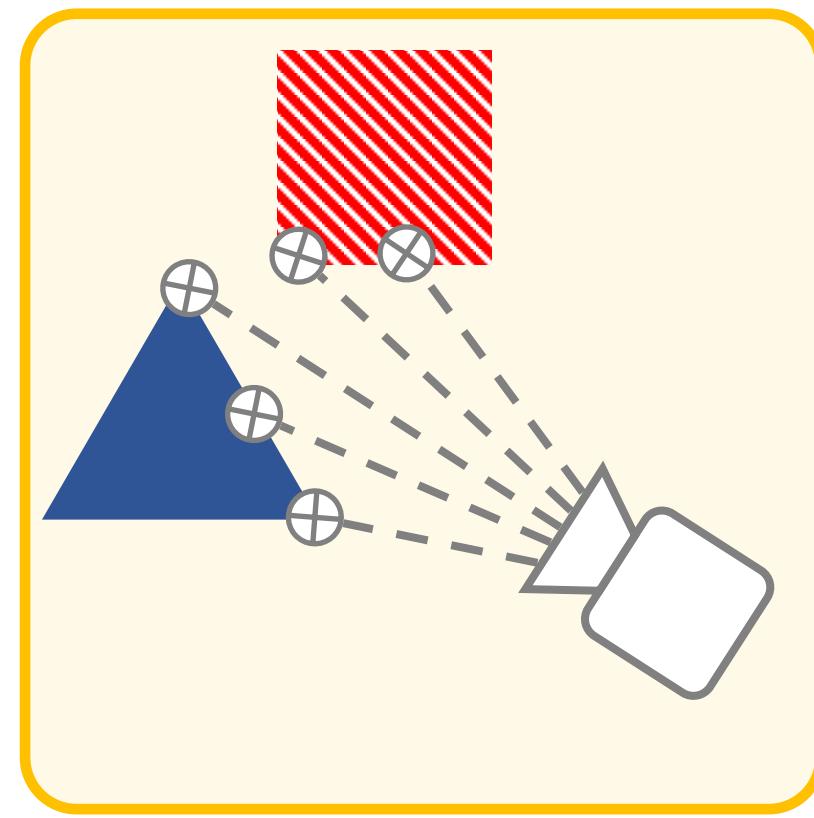
$$\left\{ (\mathcal{I}, \xi)_i \right\}_{i=1}^N$$

$$\operatorname{argmin}_{\phi, \theta} \sum_i \left\| \text{Render}_{\theta}(\text{Srn}_{\phi}, \xi_i) - \mathcal{I}_i \right\|$$

Test case: Single scene with many observations.



Scene Representation



Neural Renderer

Optimization

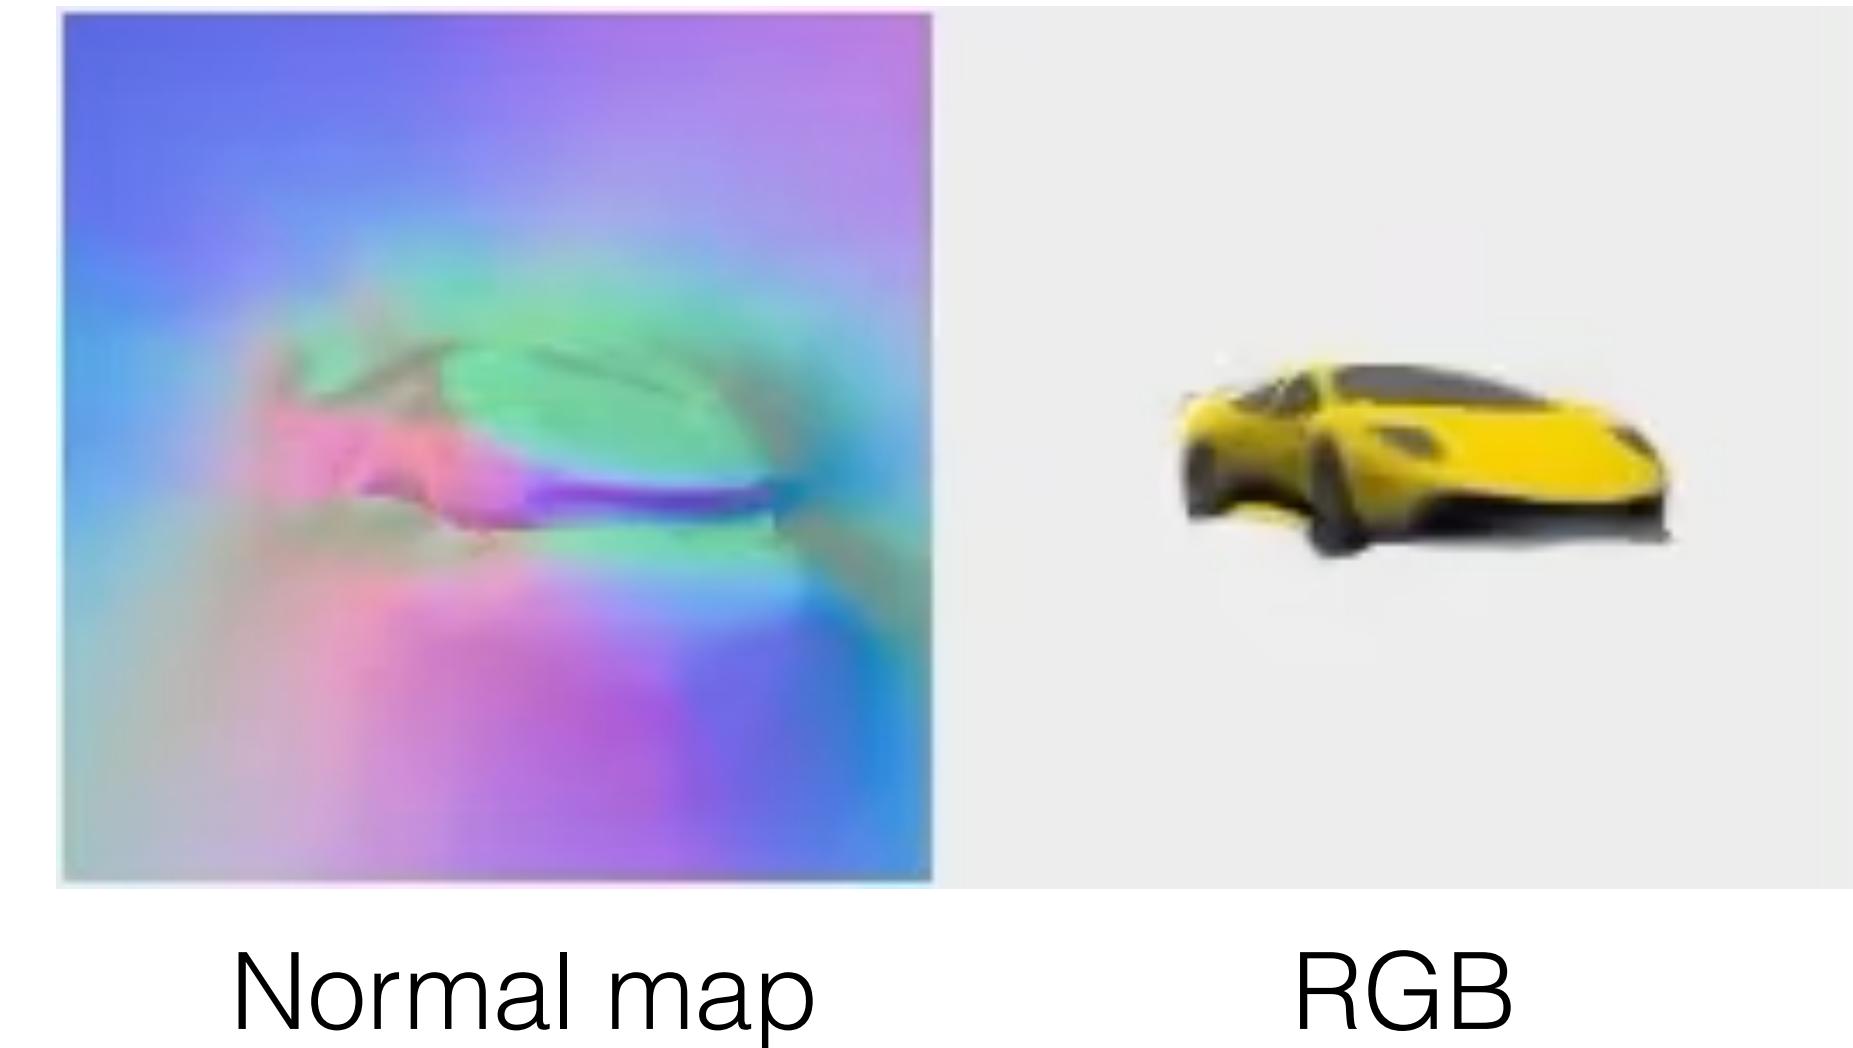


$$\{ \quad \text{[Image of a car]} + \text{[Camera icon]} , \dots \}$$

$$\left\{ (\mathcal{J}, \xi)_i \right\}_{i=1}^N$$

$$\operatorname{argmin}_{\phi, \theta} \sum_i \left\| \text{Render}_{\theta}(\text{Srn}_{\phi}, \xi_i) - \mathcal{J}_i \right\|$$

Novel View Synthesis

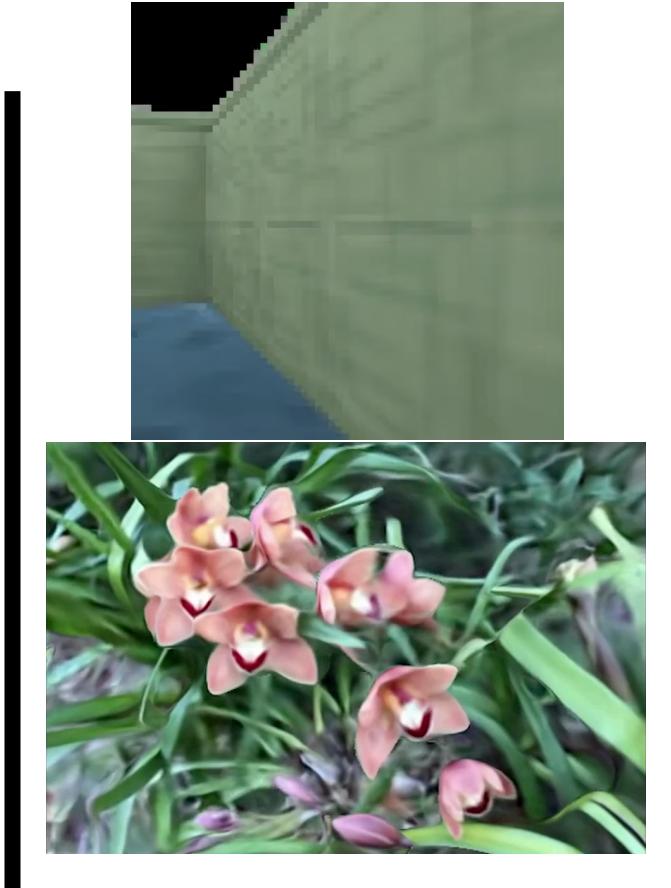


Questions?

The explosion of neural implicit representations in inverse graphics



The explosion of neural implicit representations in inverse graphics



2019

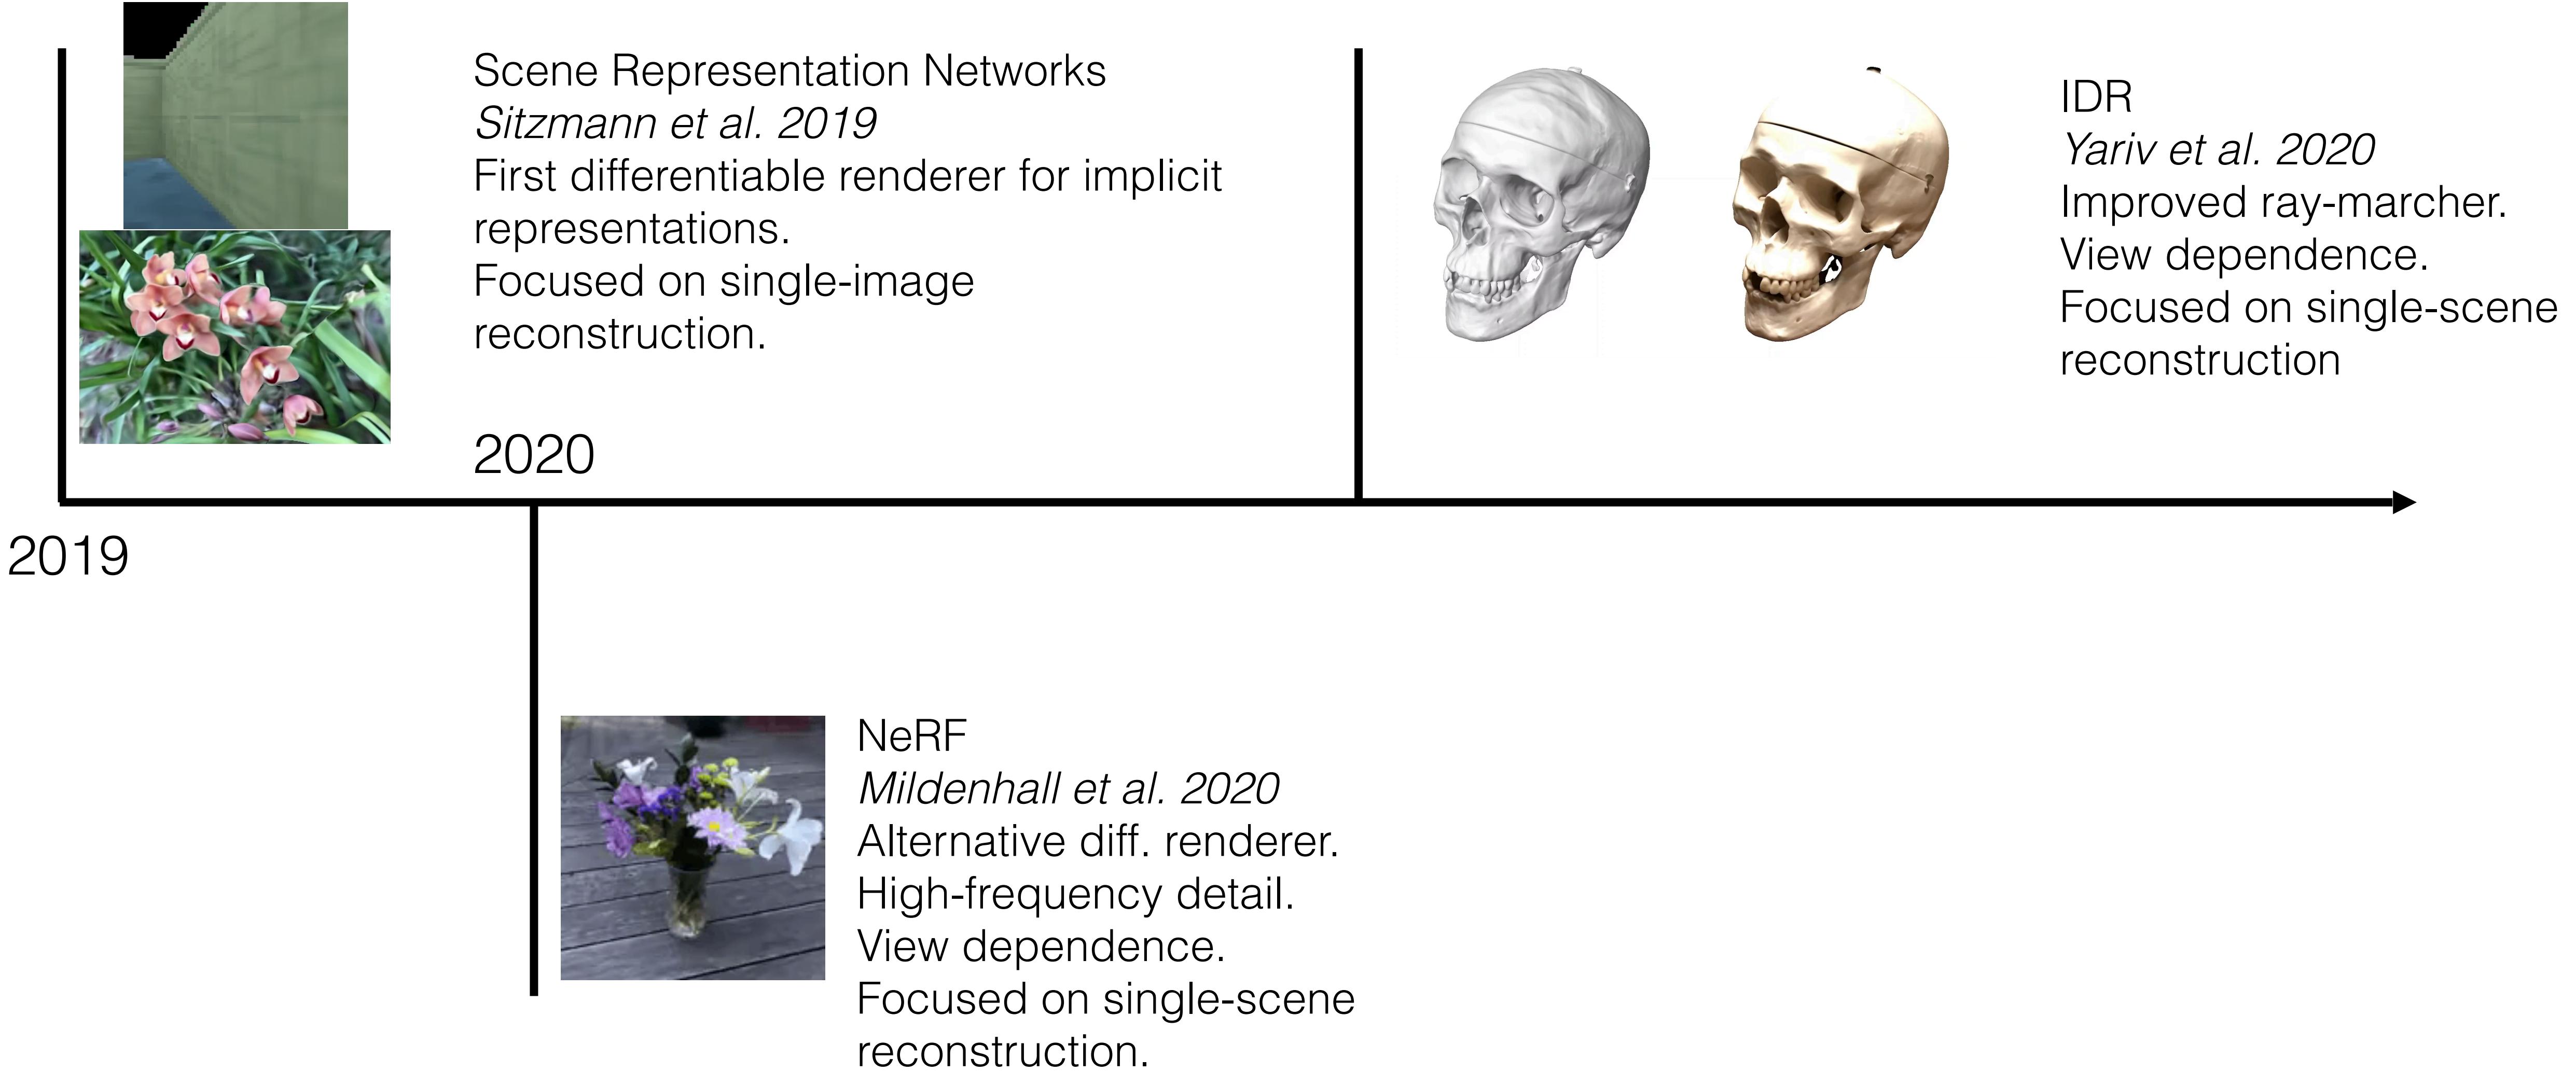
Scene Representation Networks

Sitzmann et al. 2019

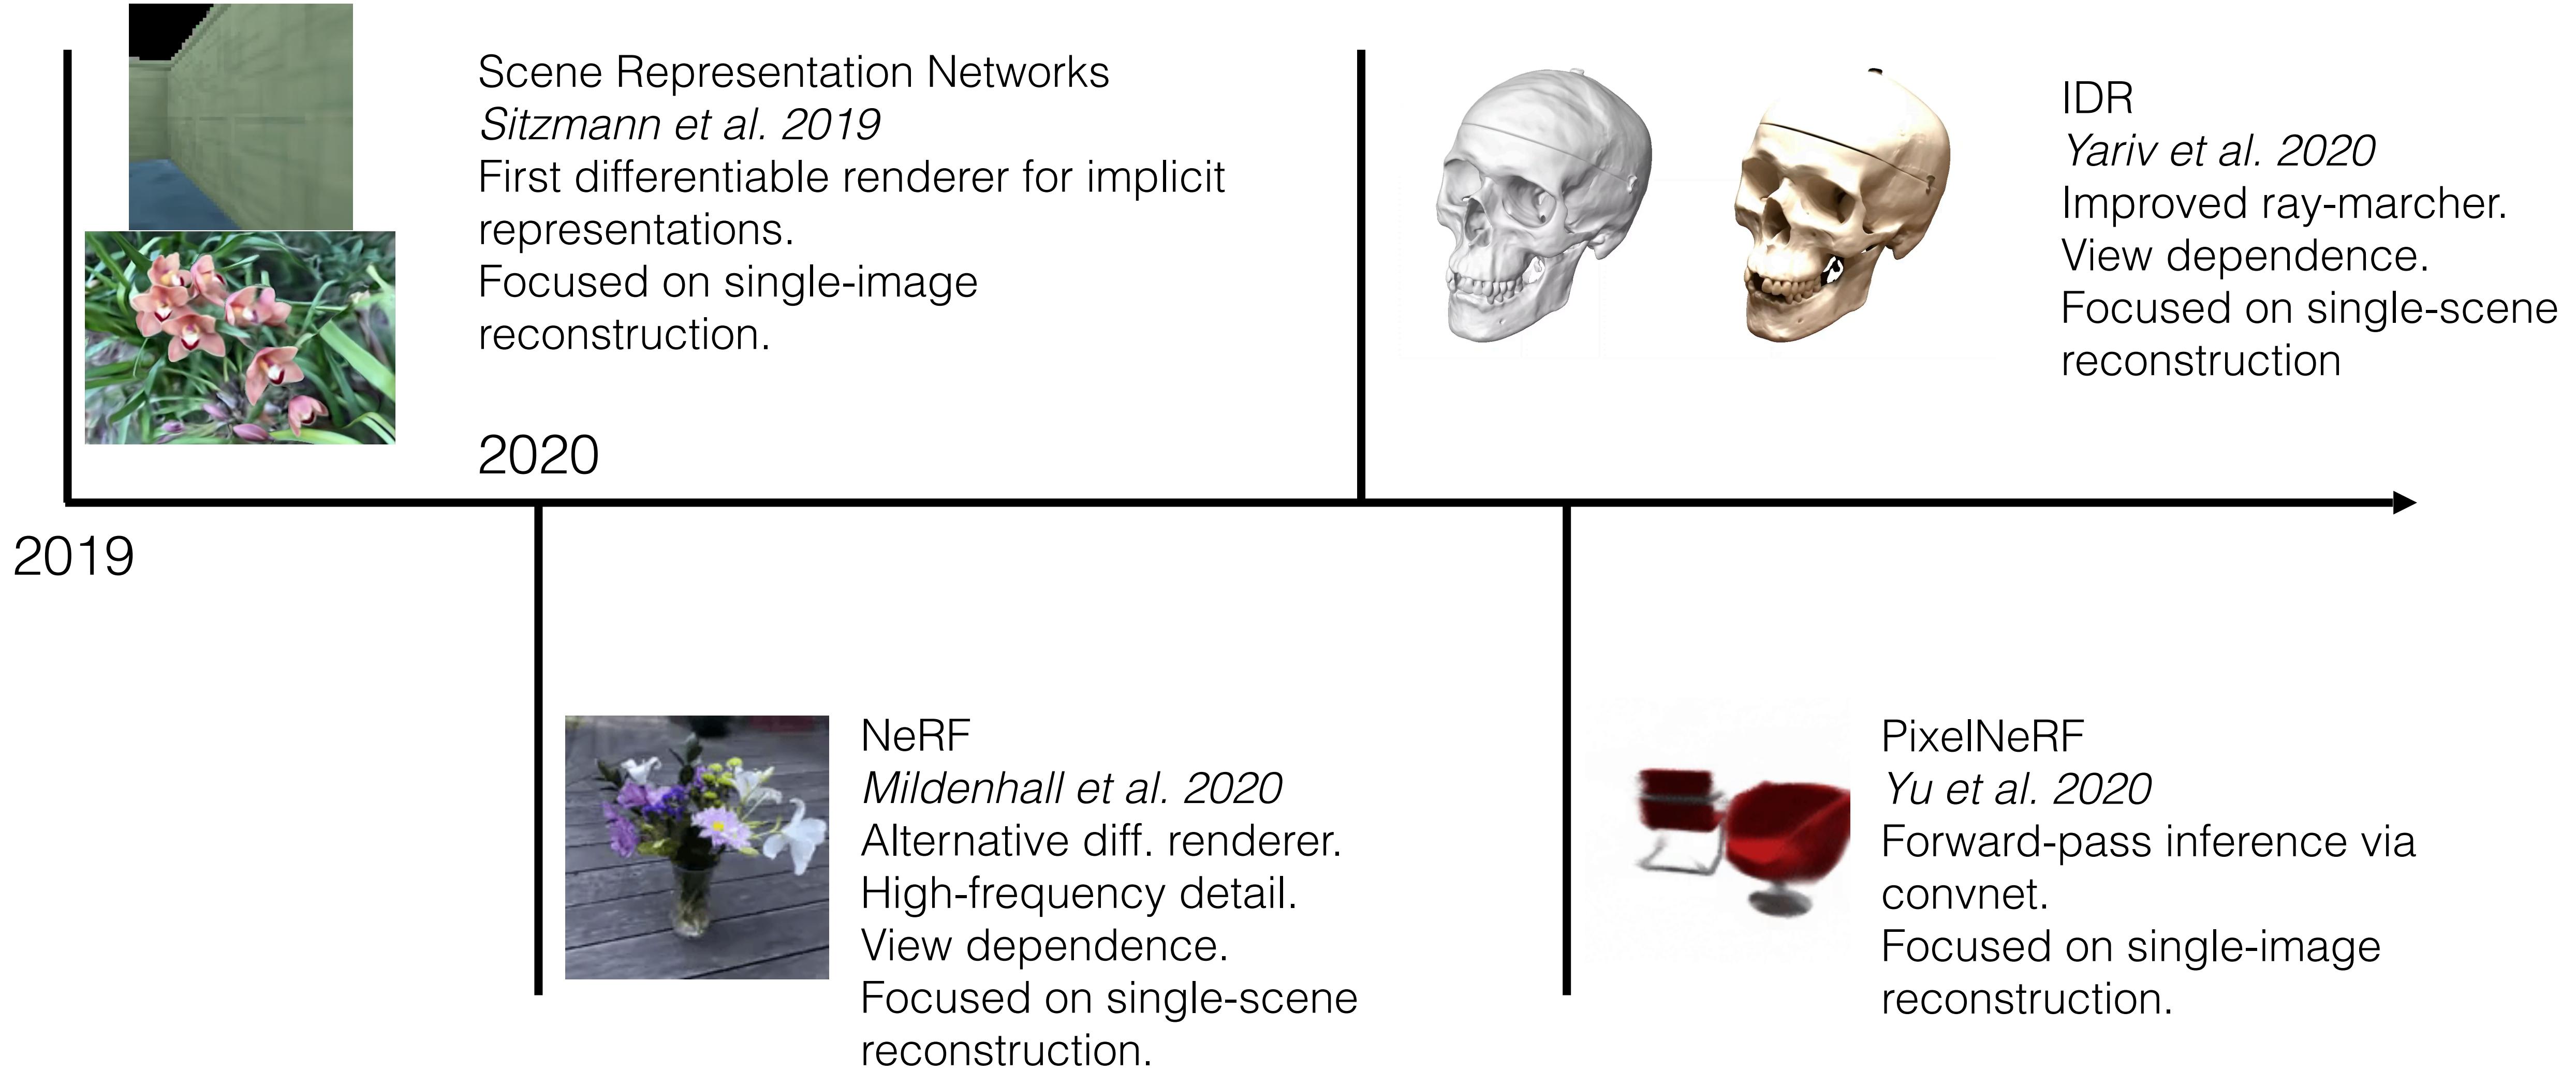
First differentiable renderer for implicit representations.

Focused on single-image reconstruction.

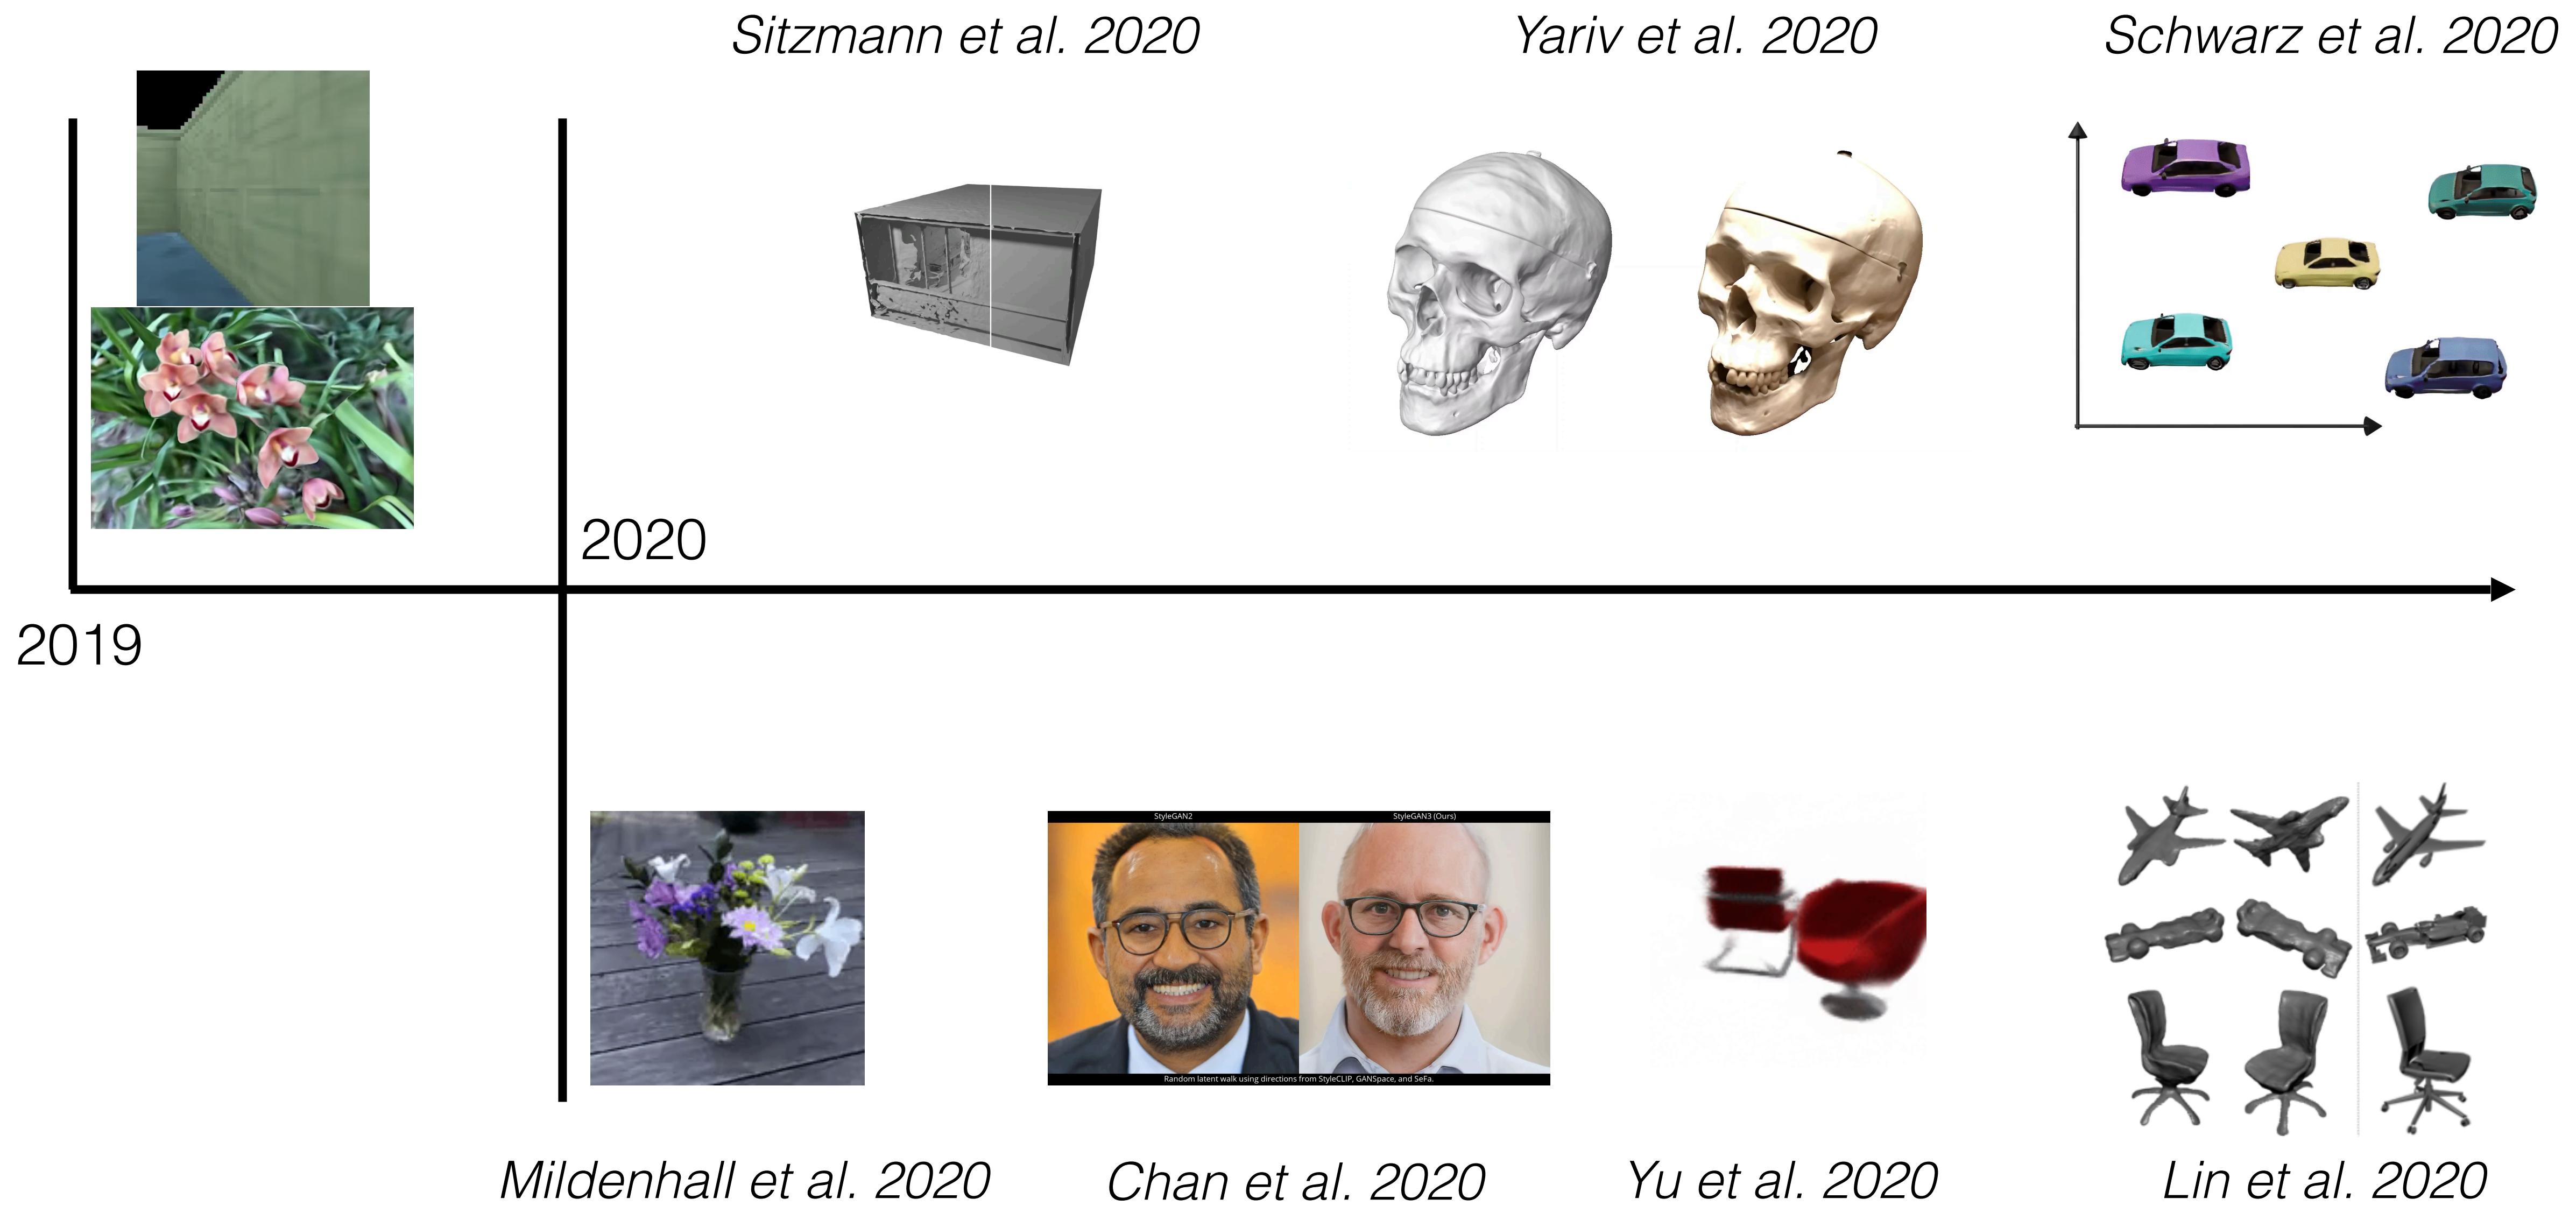
The explosion of neural implicit representations in inverse graphics



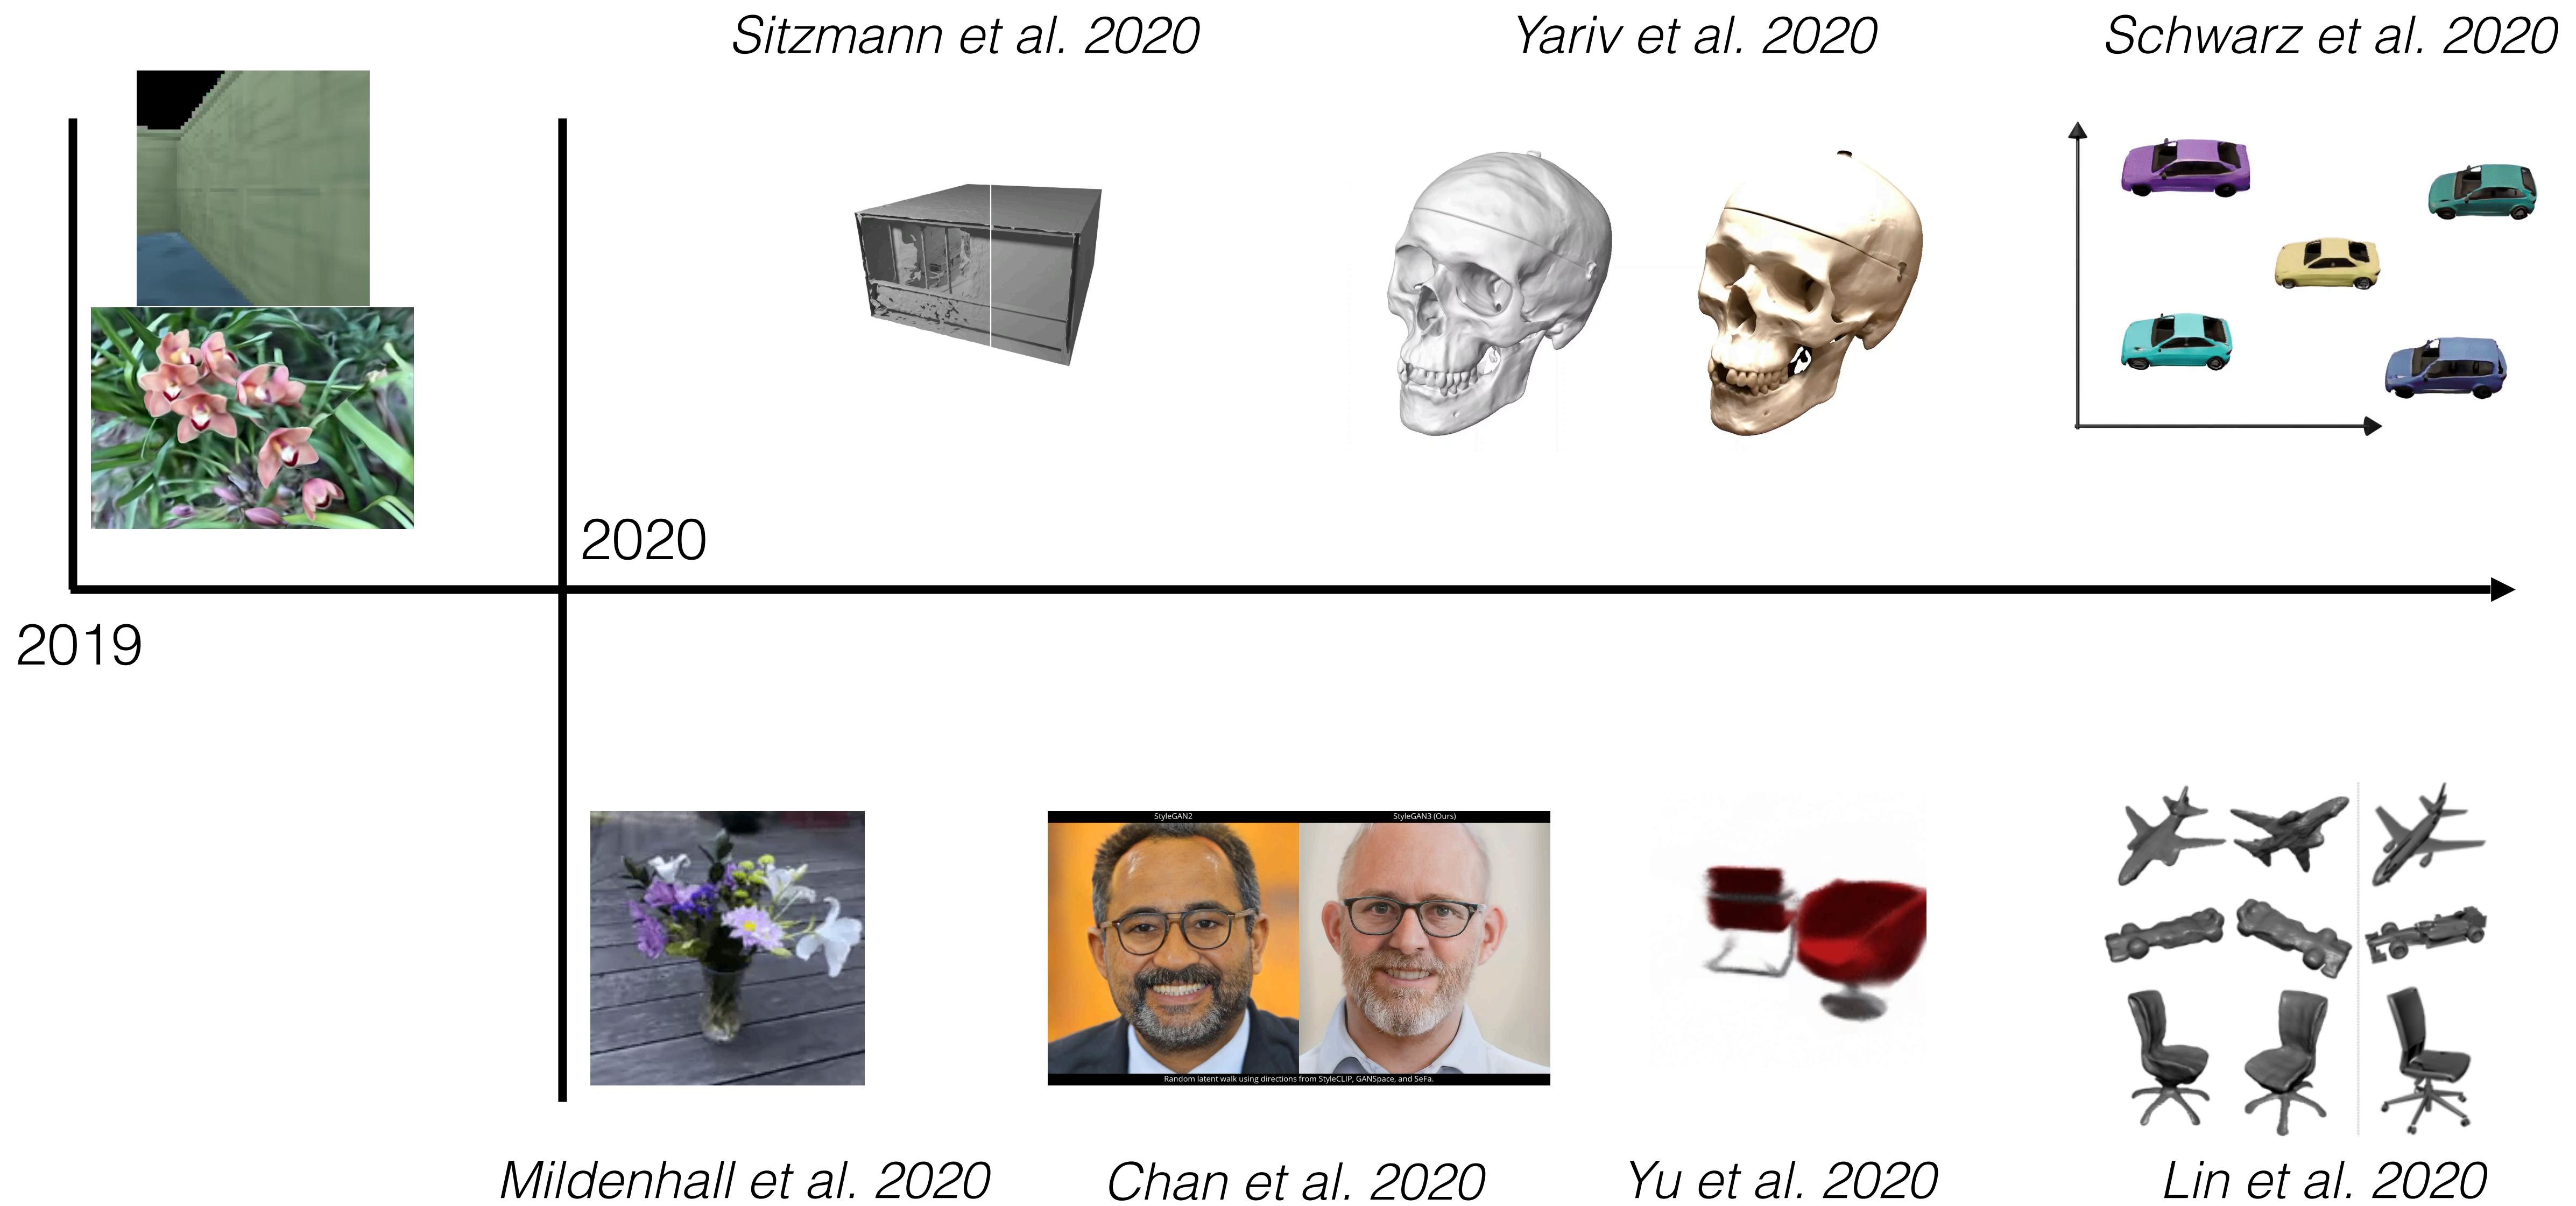
The explosion of neural implicit representations in inverse graphics



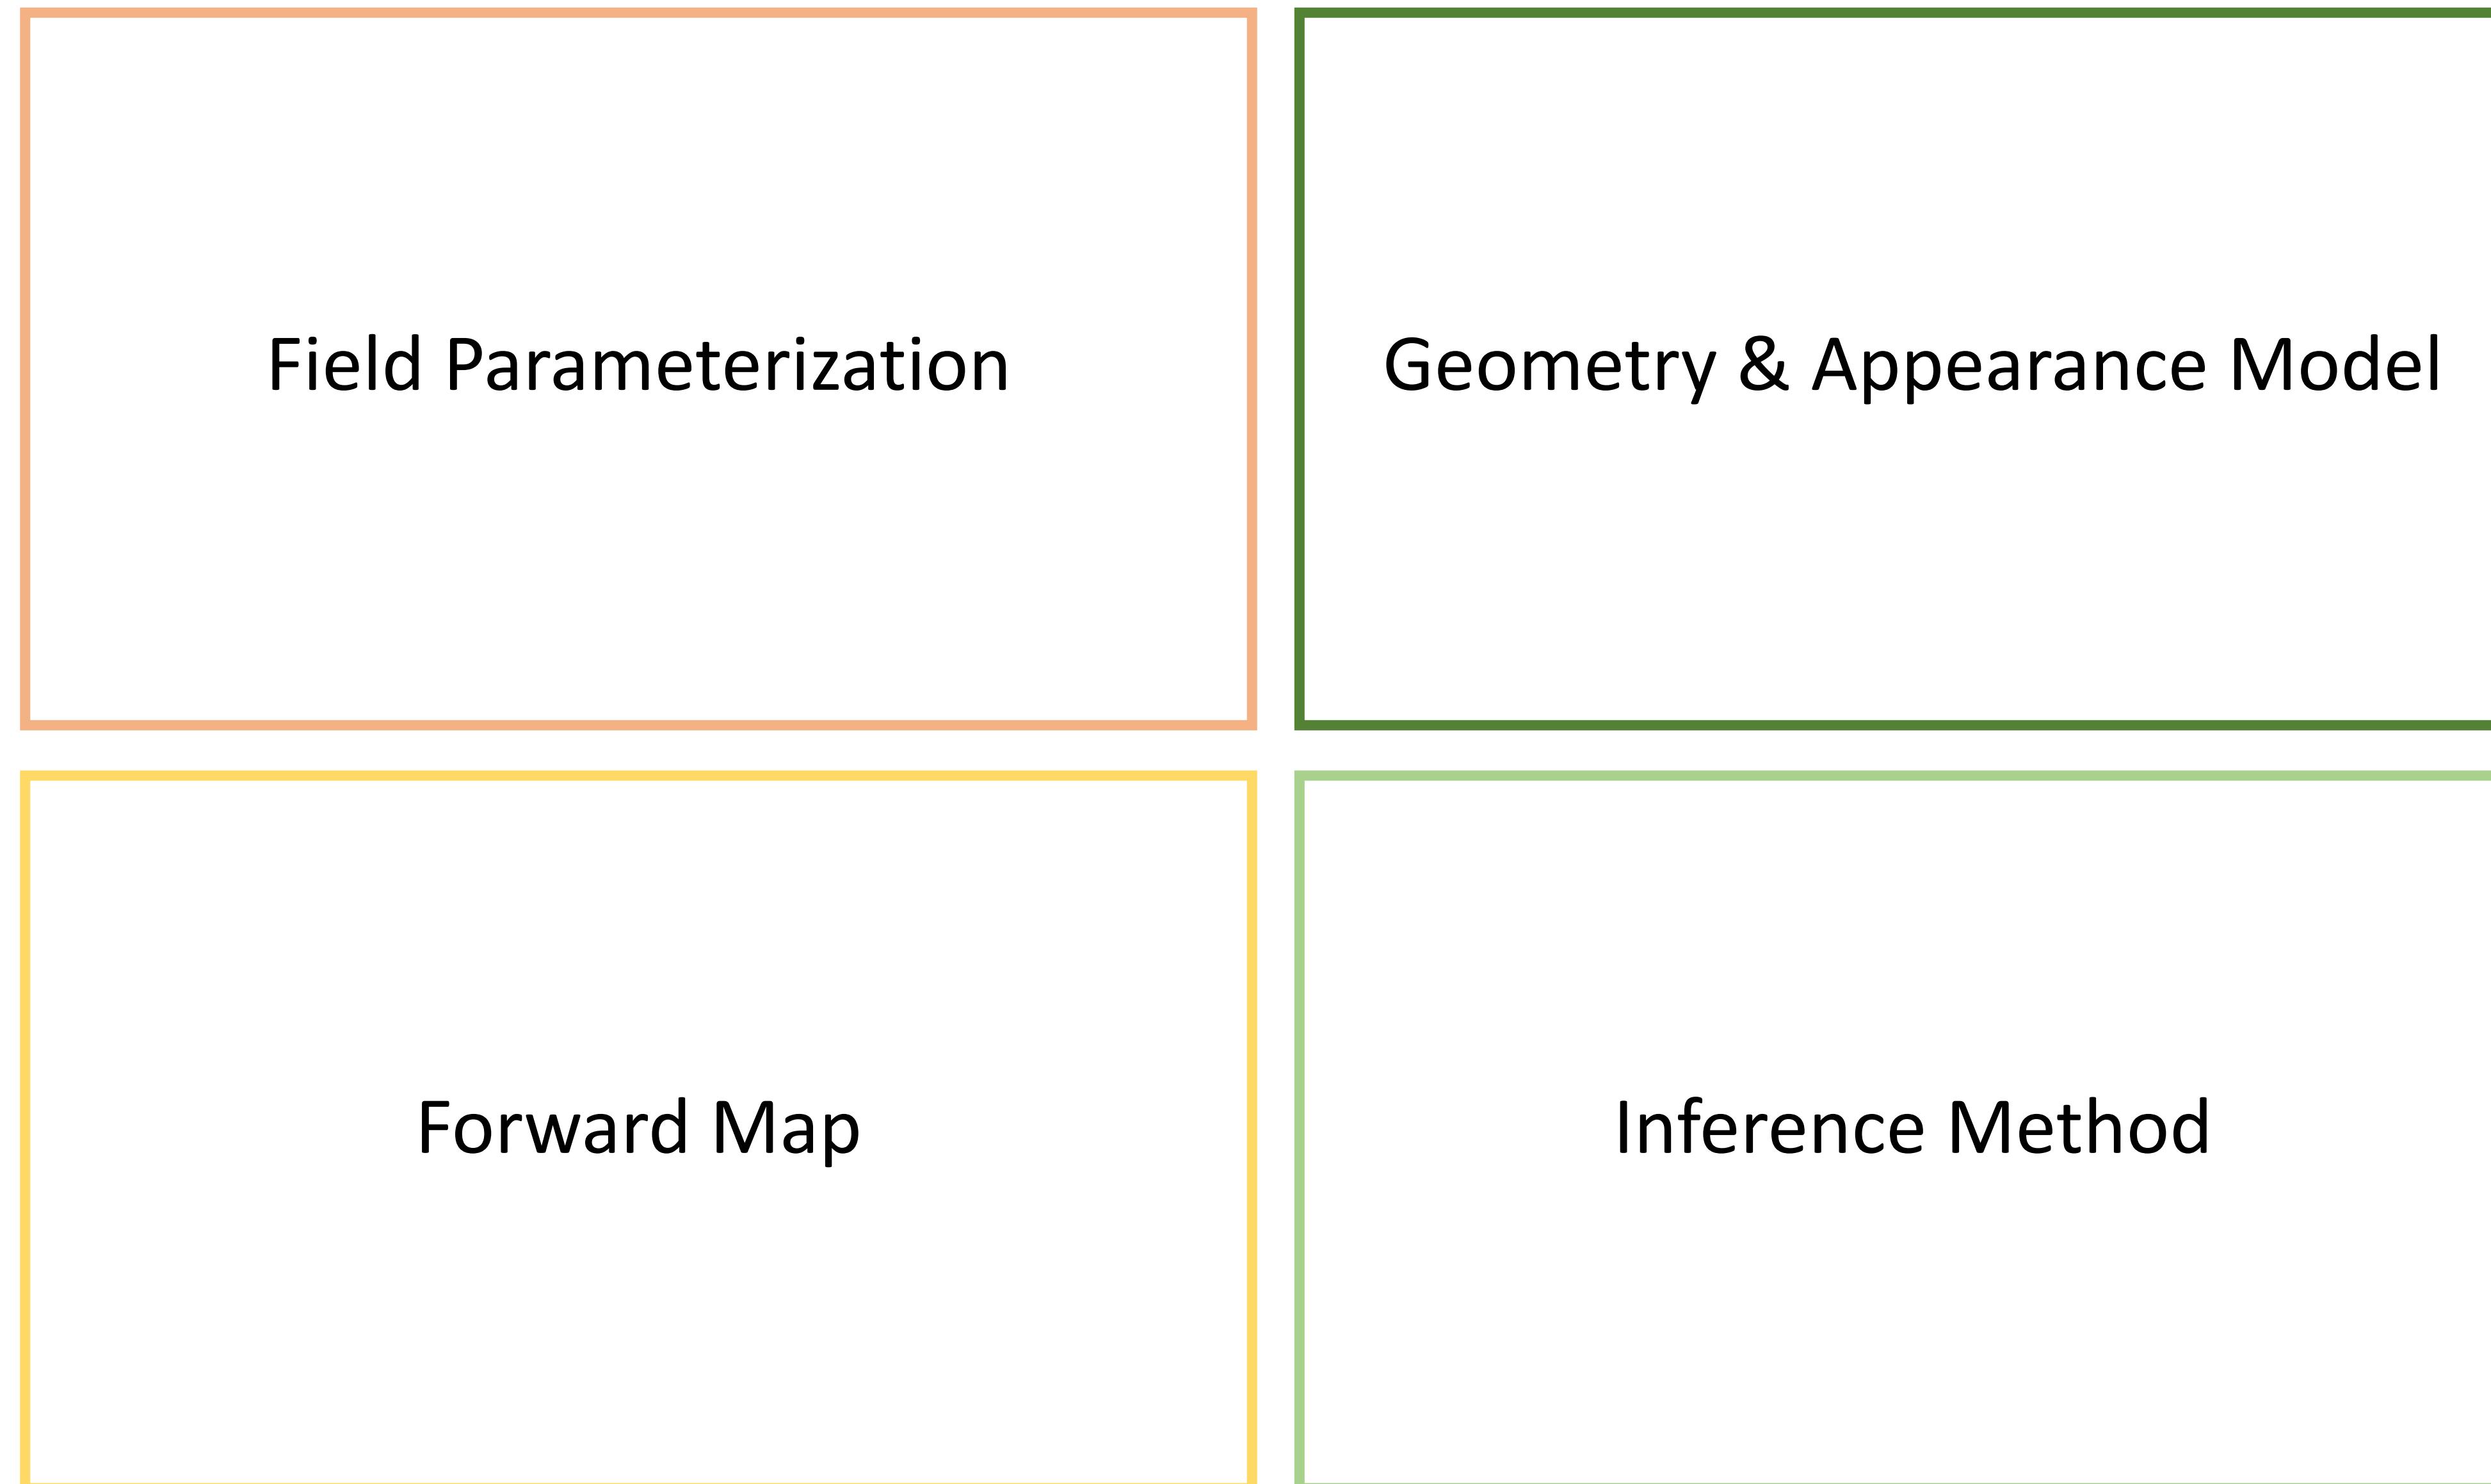
The explosion of neural implicit representations in inverse graphics



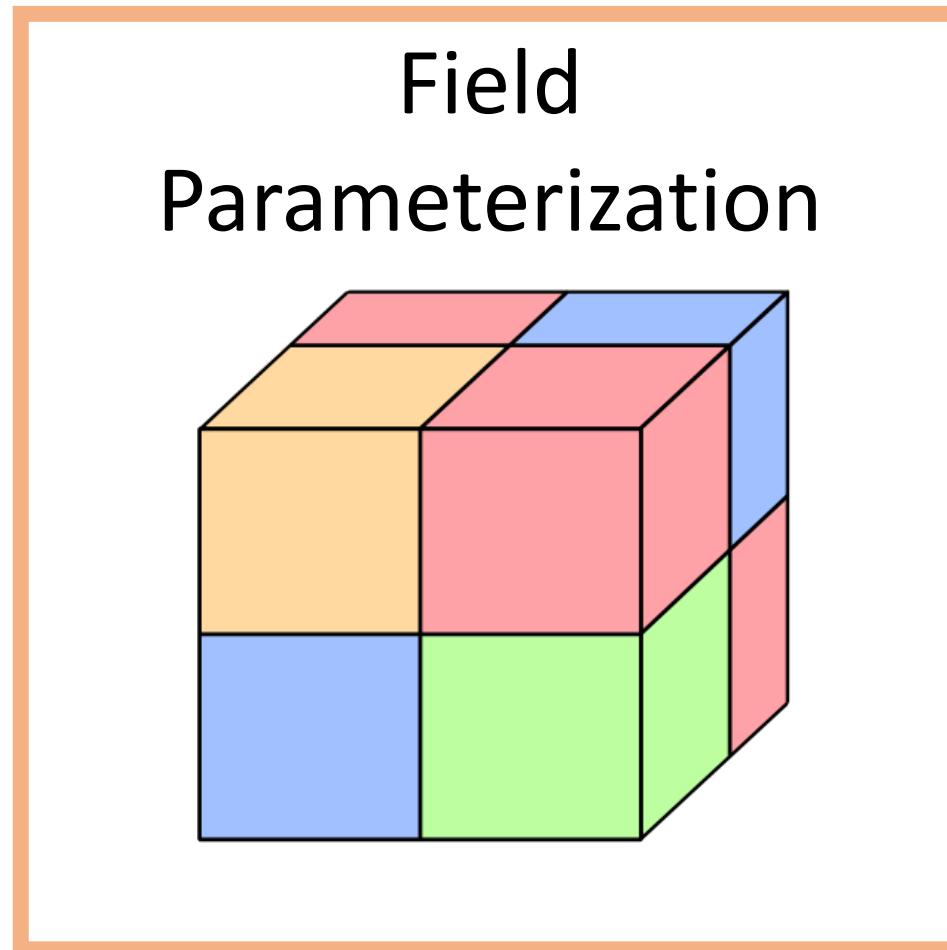
The explosion of neural implicit representations in inverse graphics



3D Inverse Graphics Taxonomy



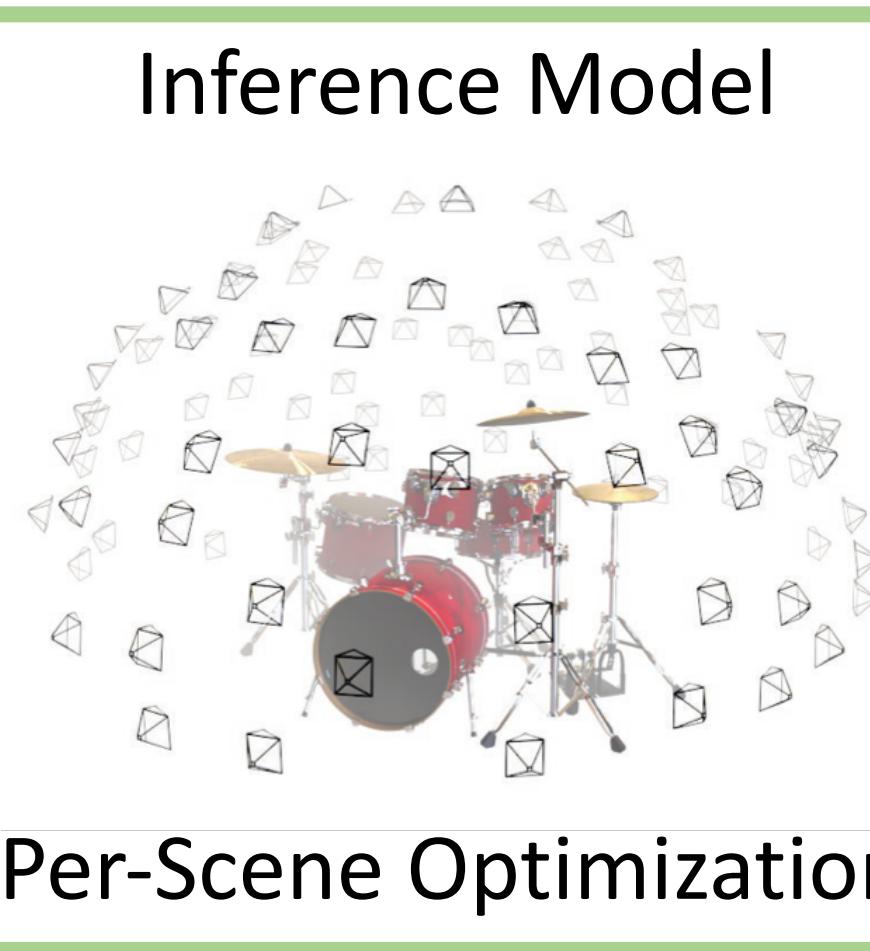
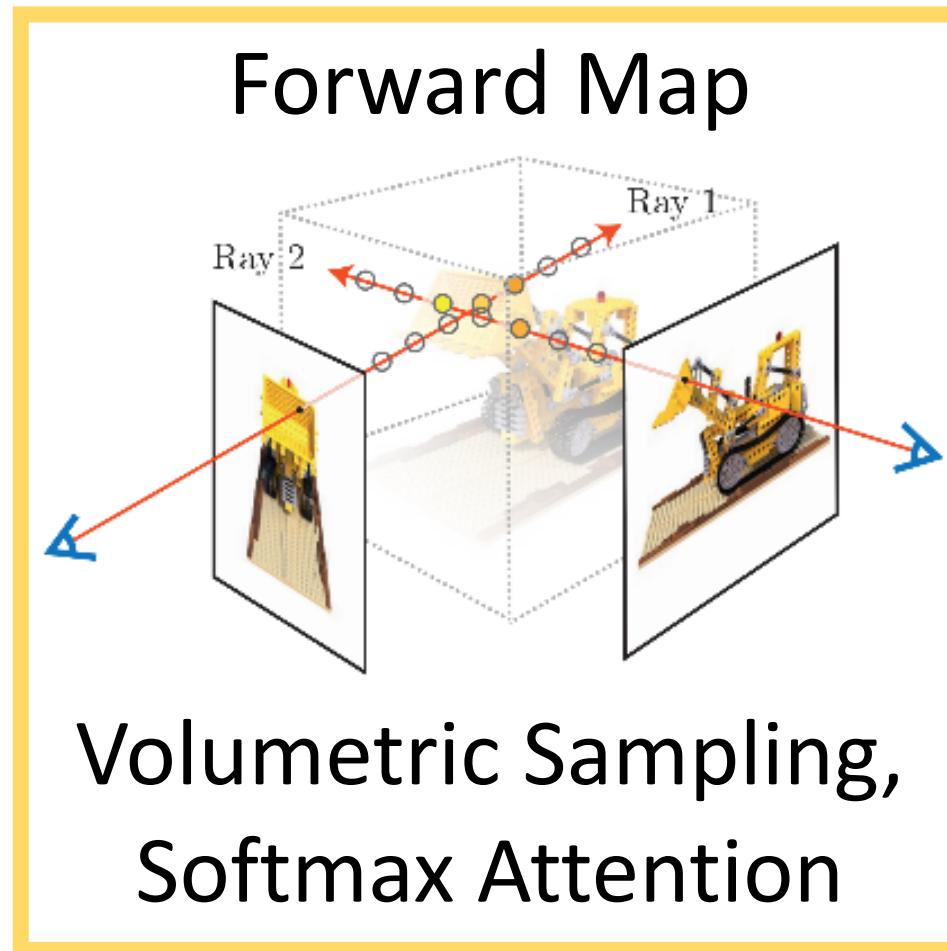
DeepVoxels (CVPR 2018)



Geometry & Appearance Model

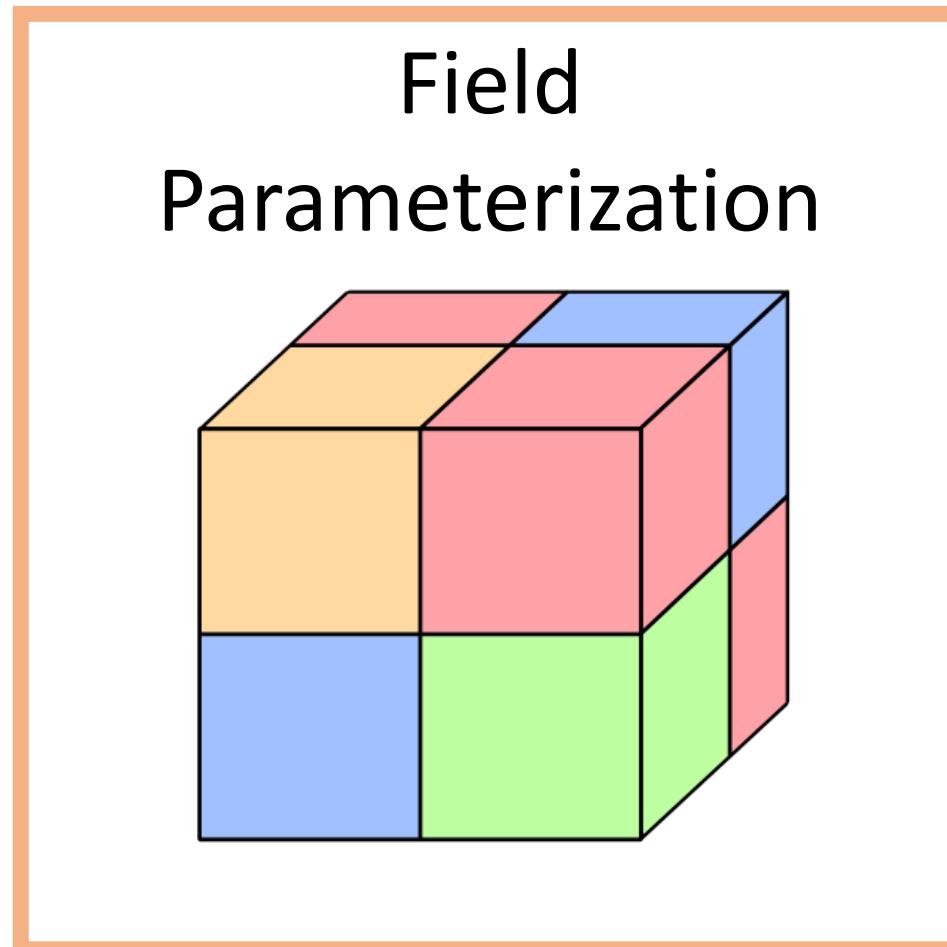
$$\Phi(\mathbf{x}) = \mathbf{f}$$
$$\mathbf{f} \in \mathbb{R}^n$$

Feature-based Representation



- Reasonable quality overfitting
- Fast rendering
- Noisy geometry
- Requires Many Images

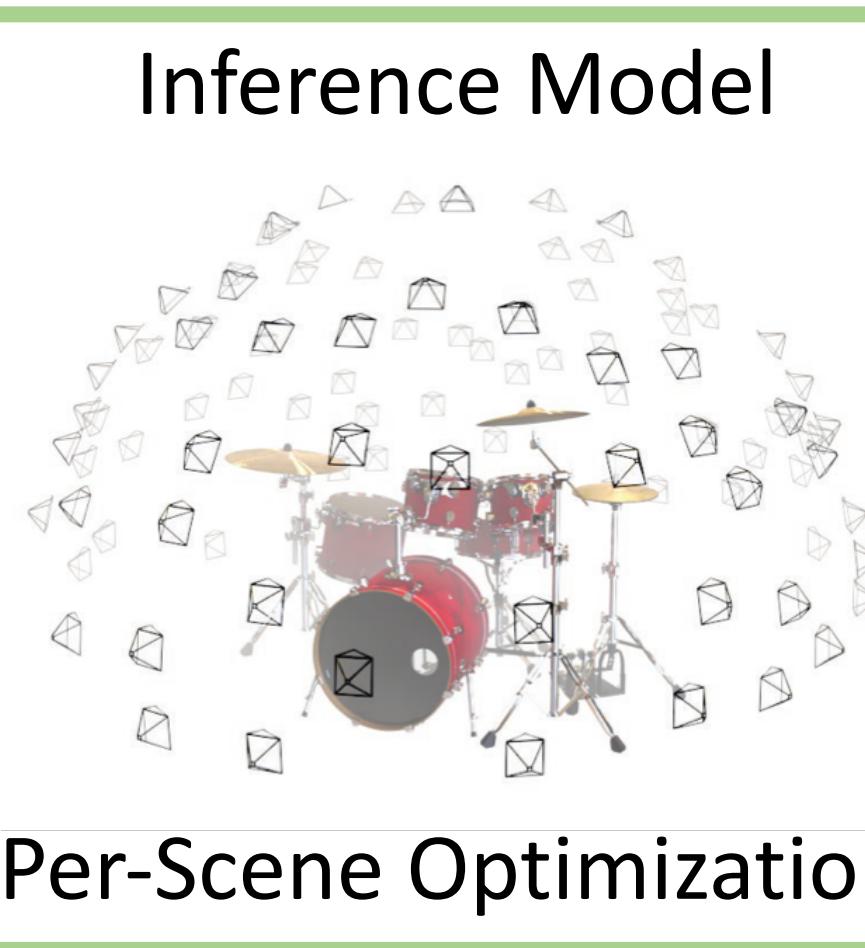
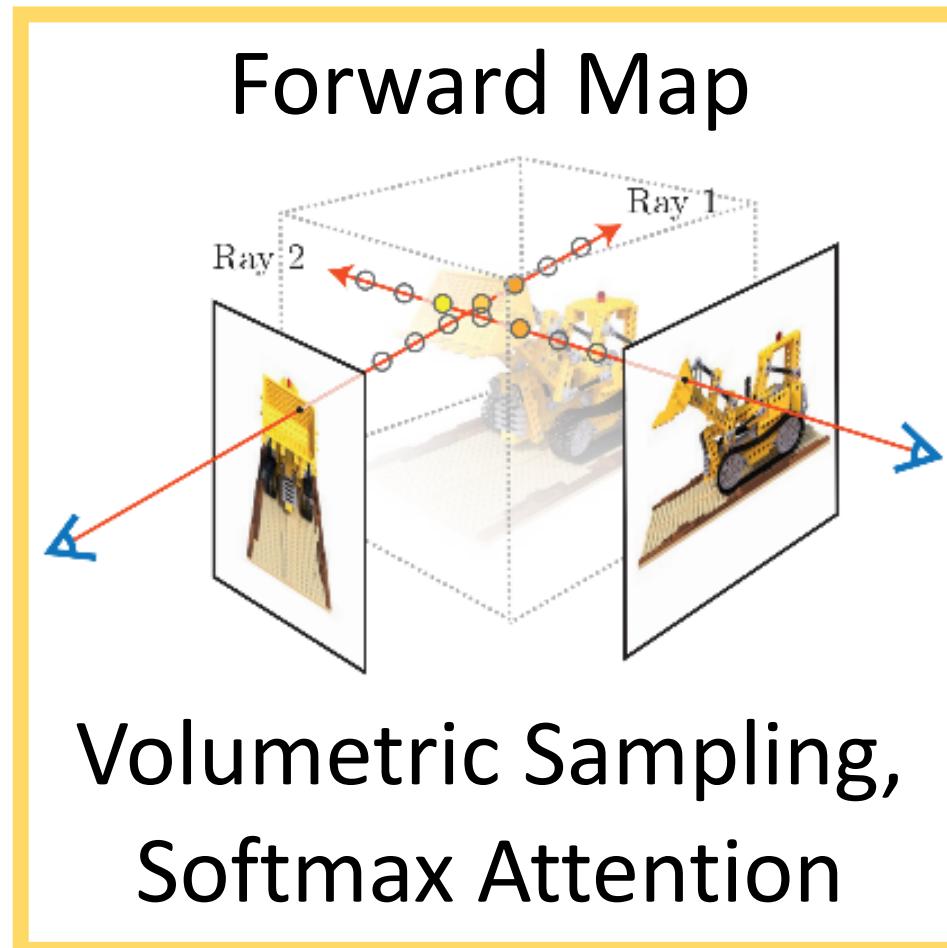
DeepVoxels (CVPR 2018)



Geometry & Appearance Model

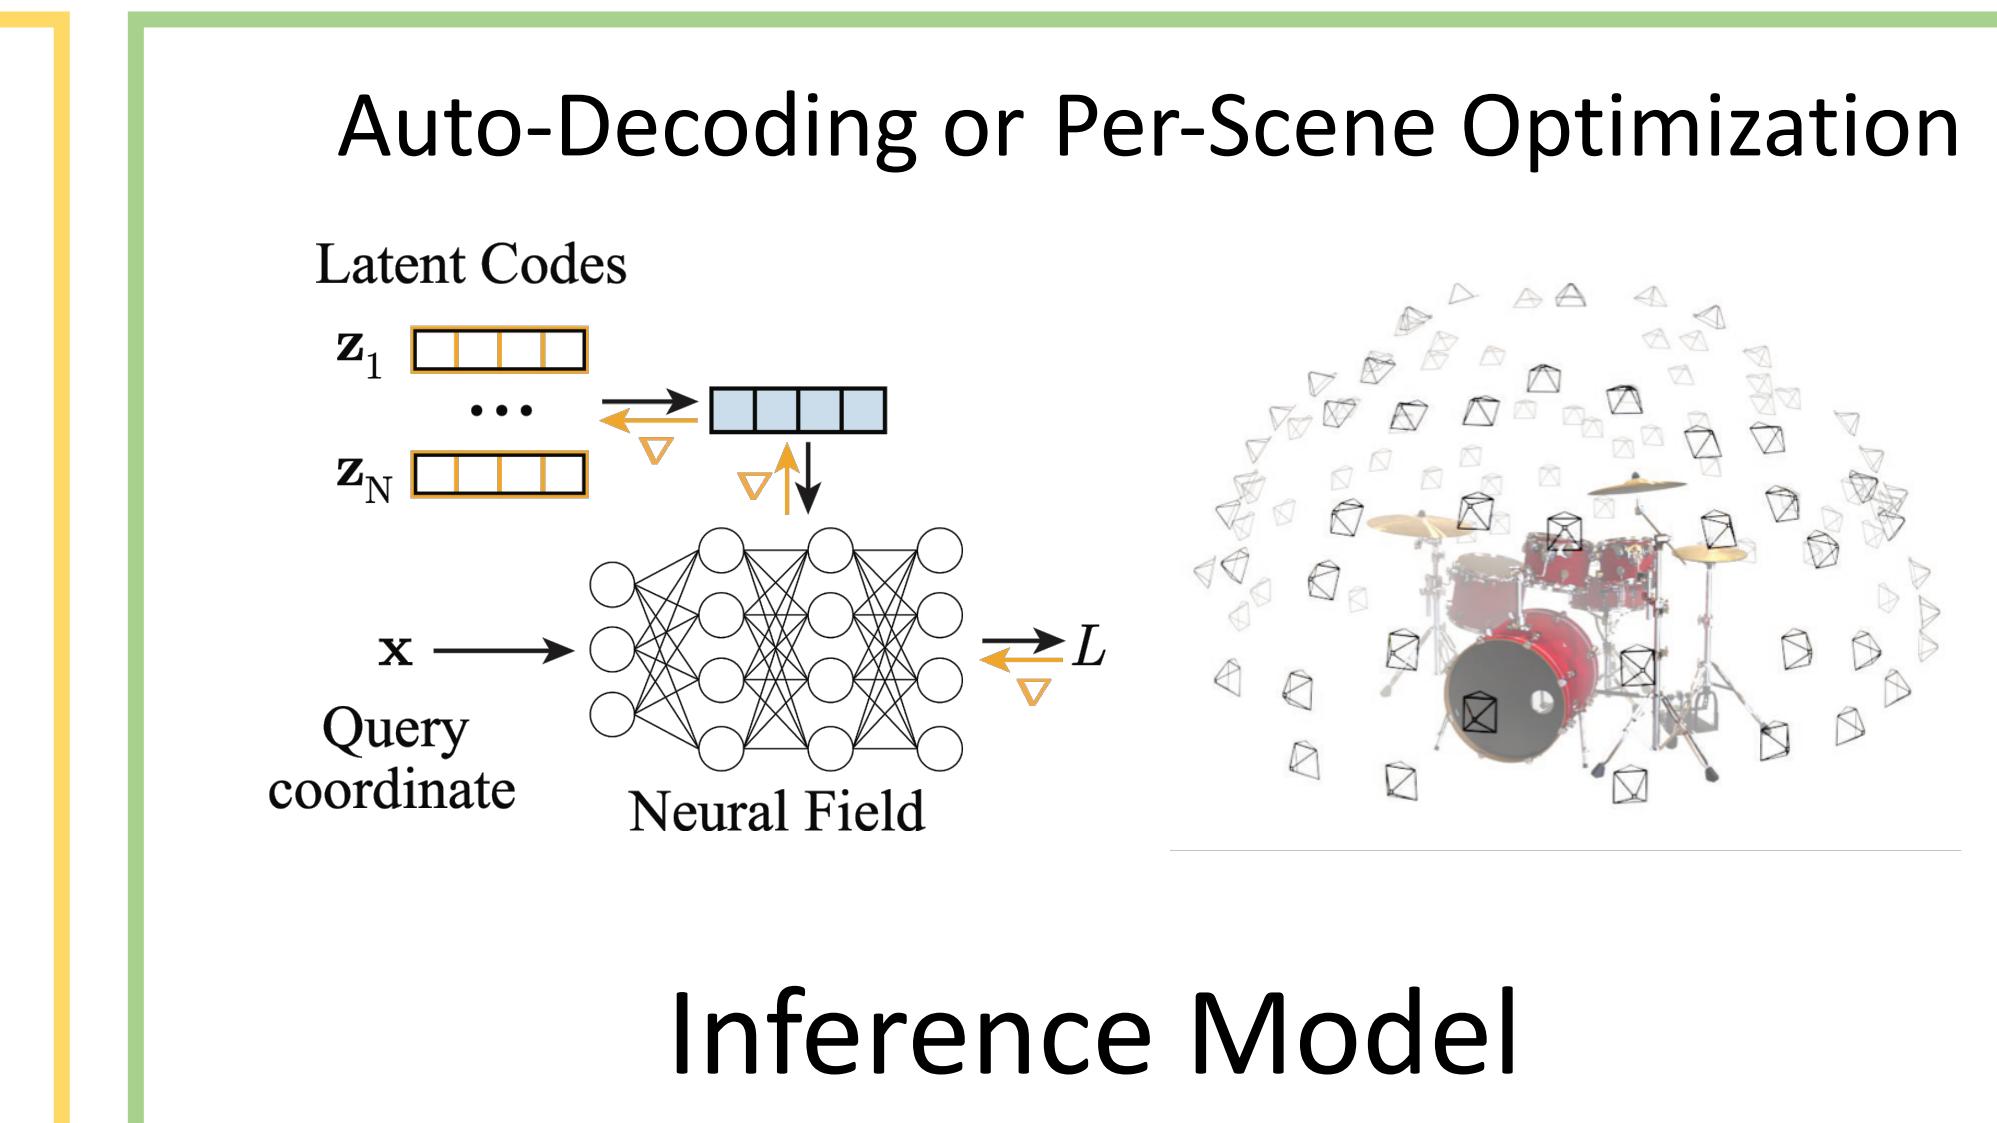
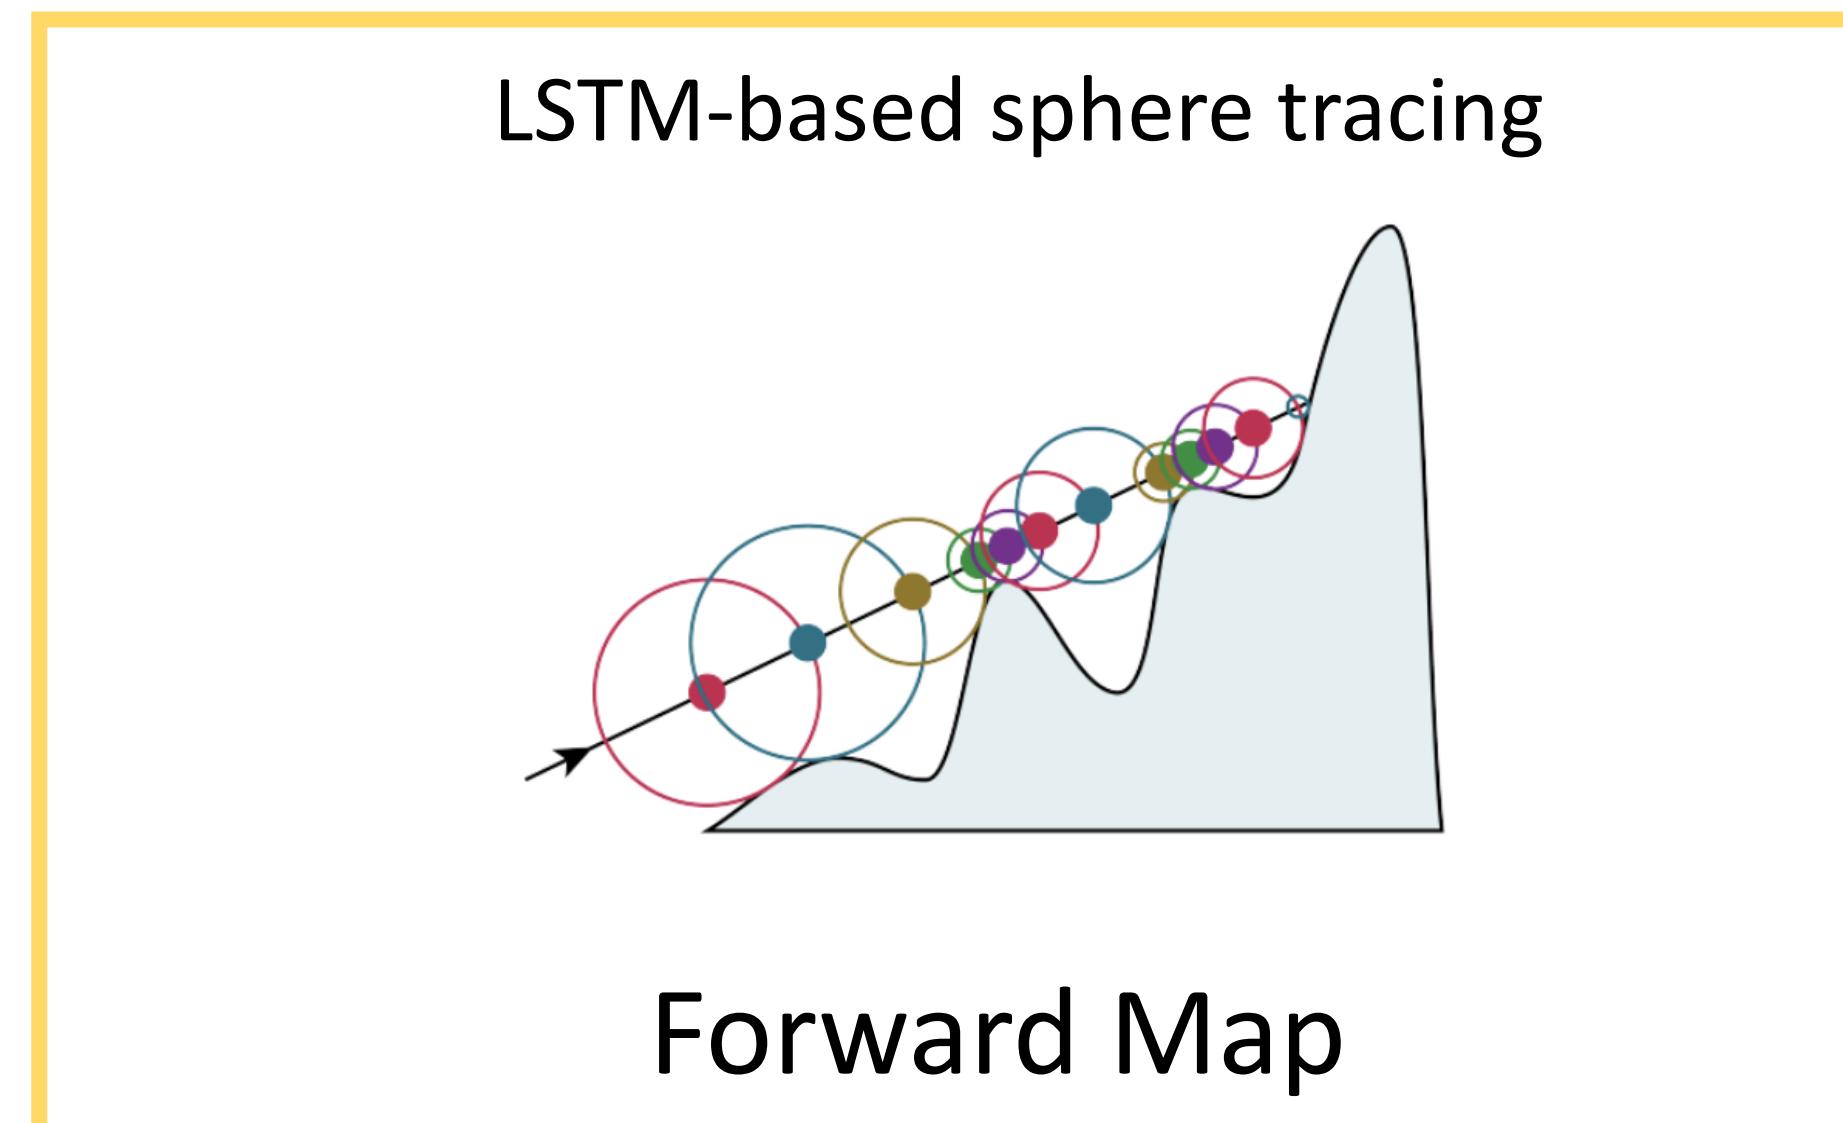
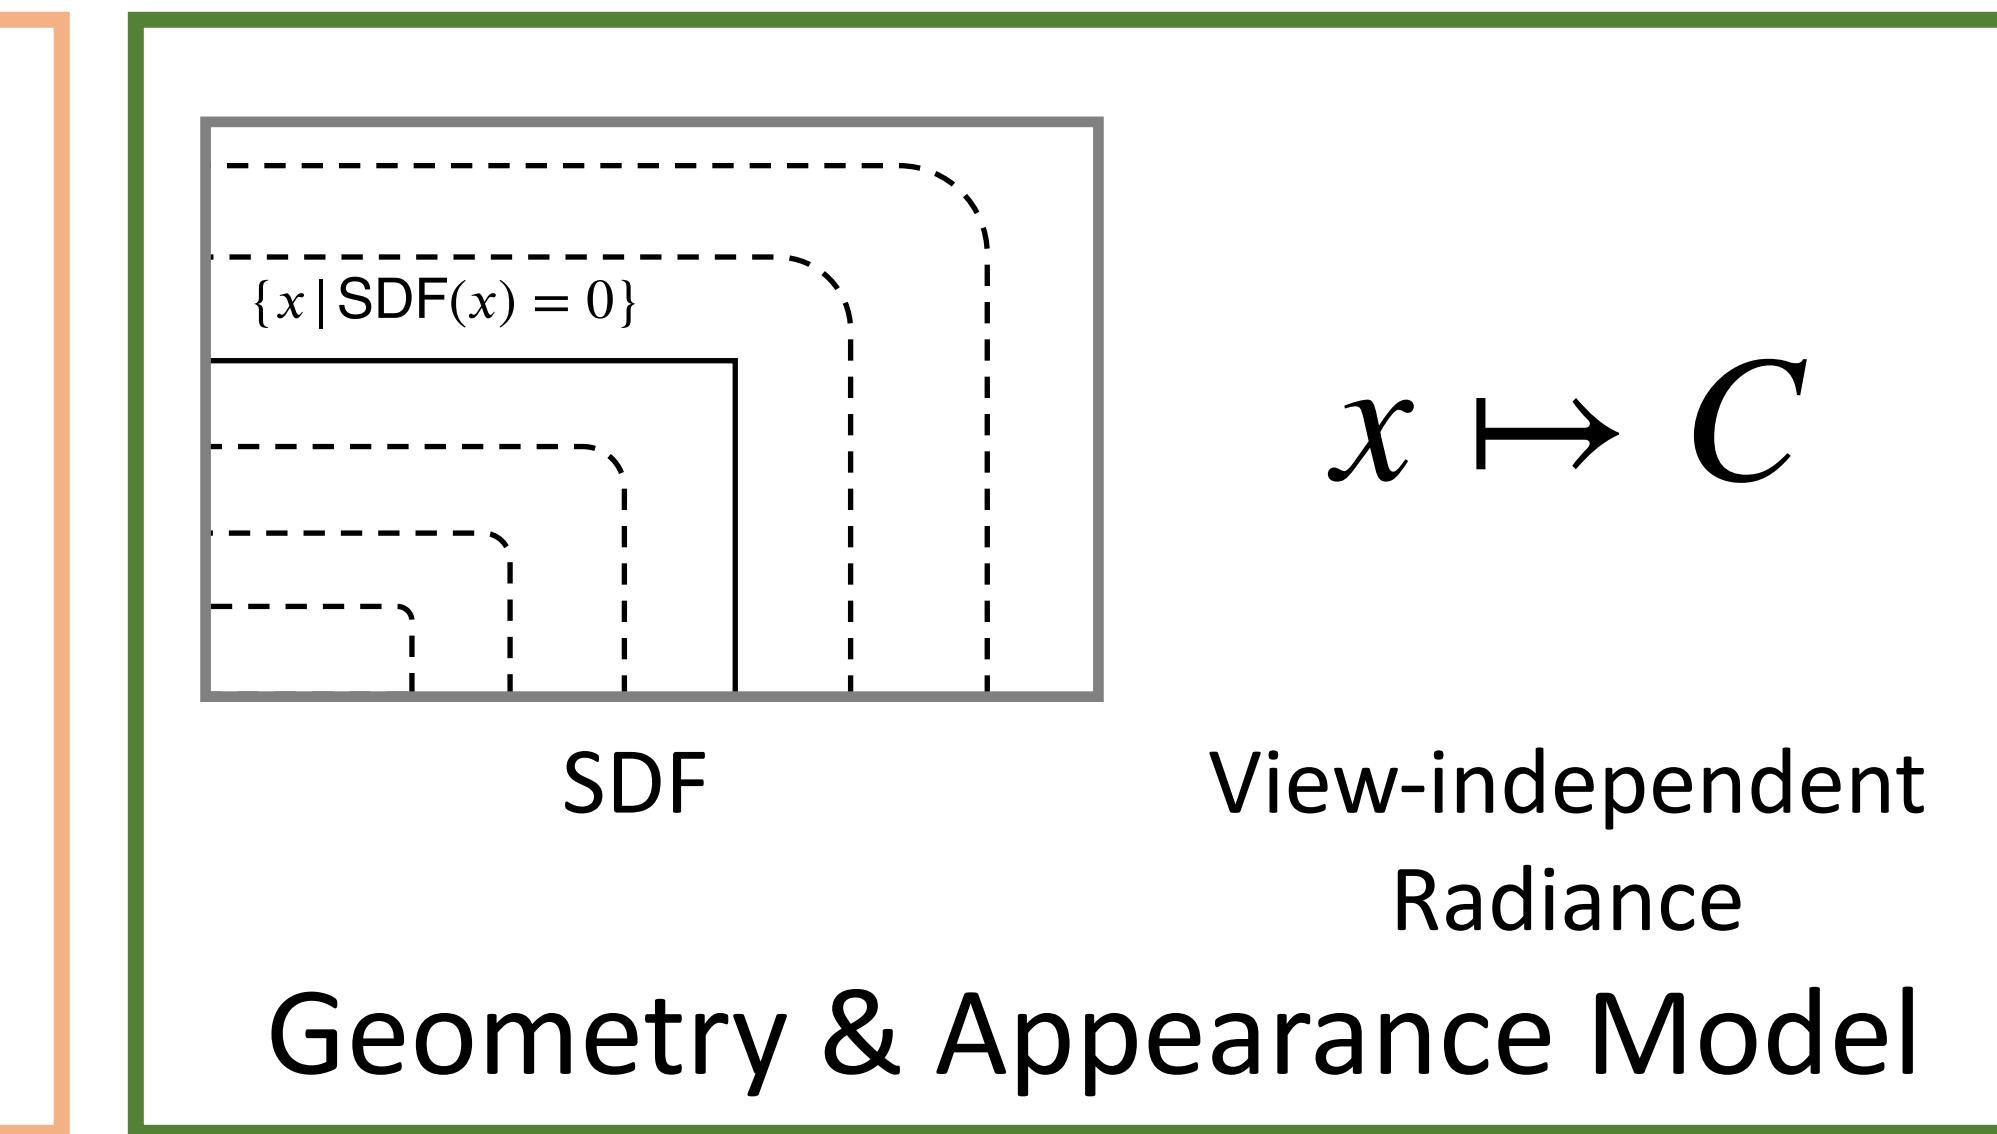
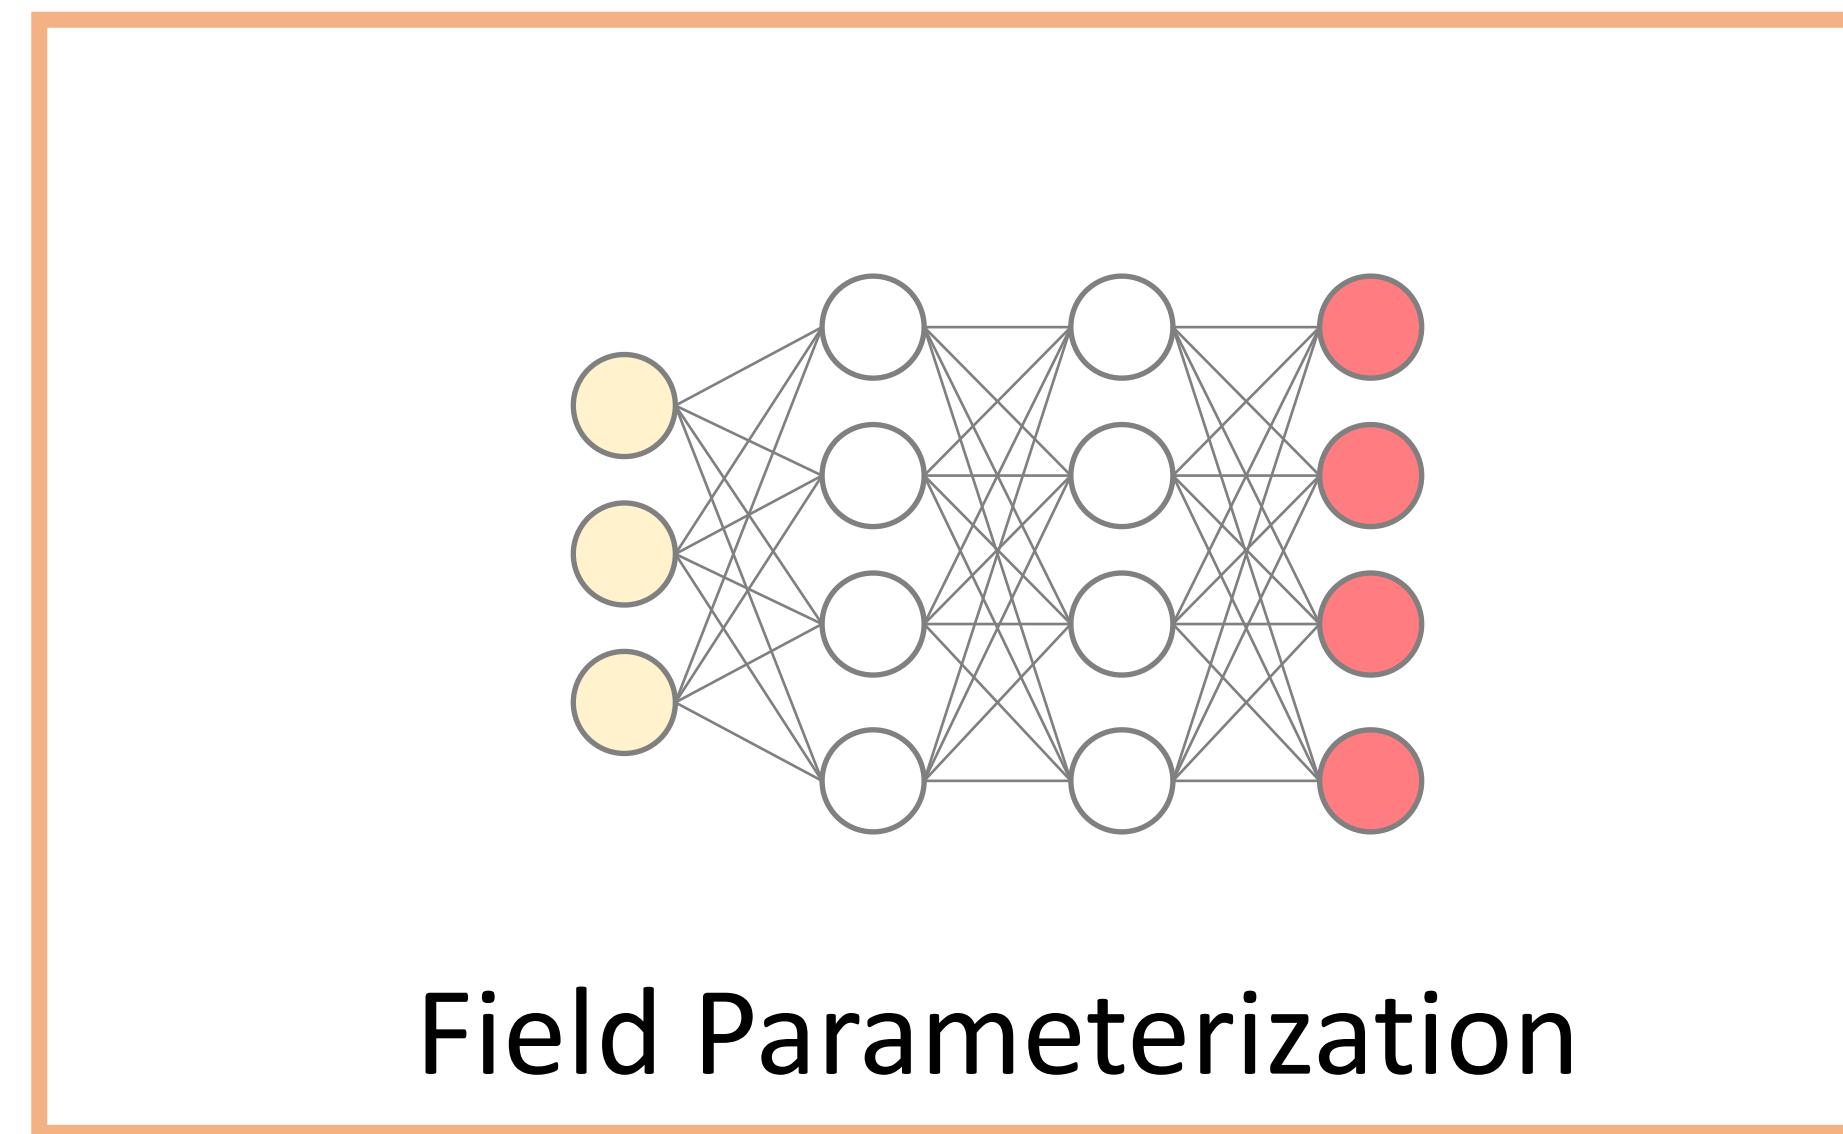
$$\Phi(\mathbf{x}) = \mathbf{f}$$
$$\mathbf{f} \in \mathbb{R}^n$$

Feature-based Representation

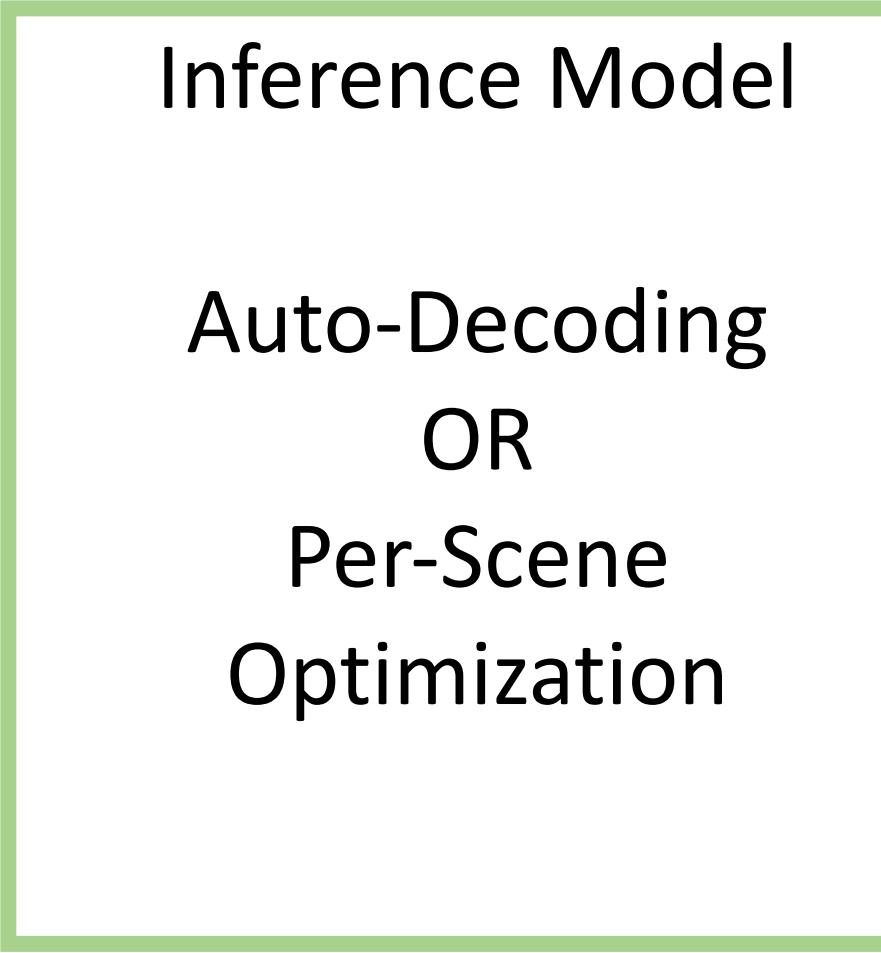
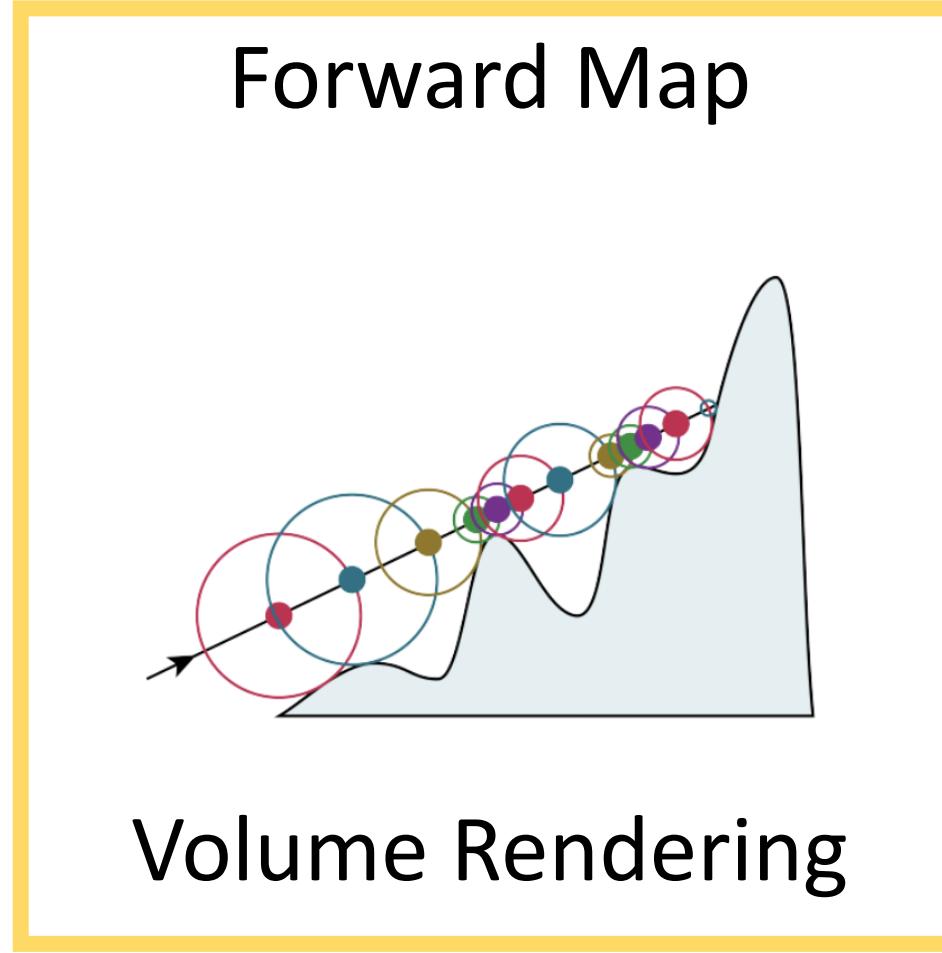
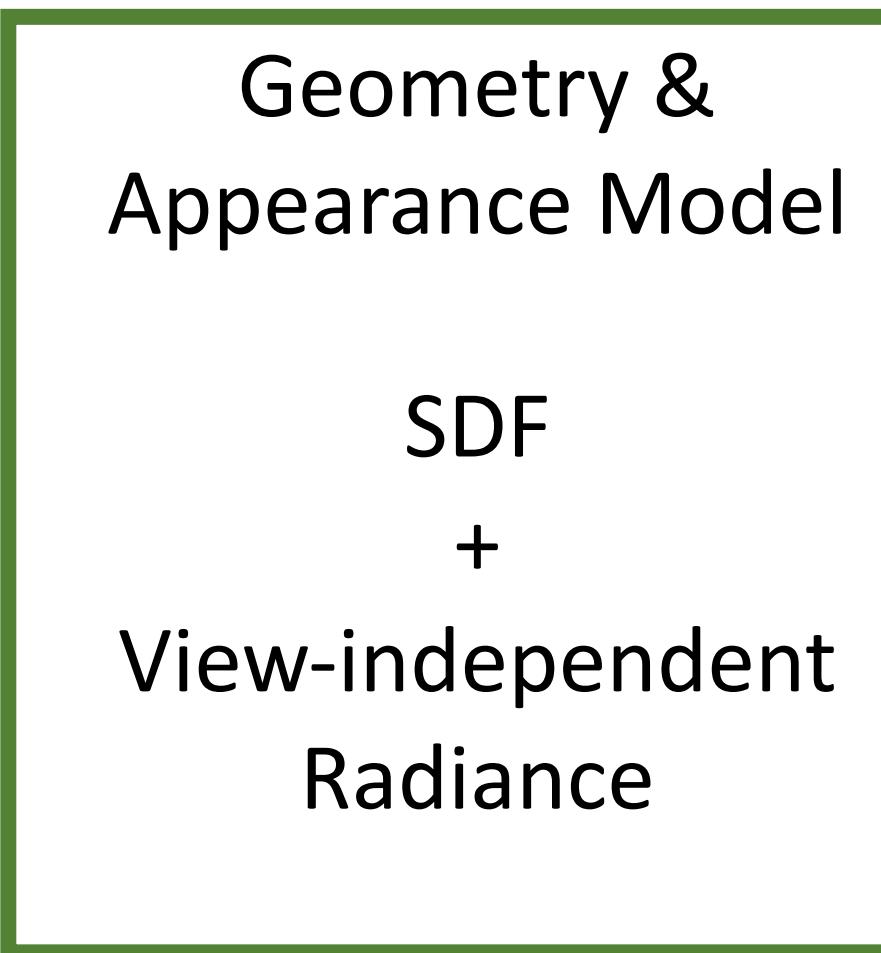
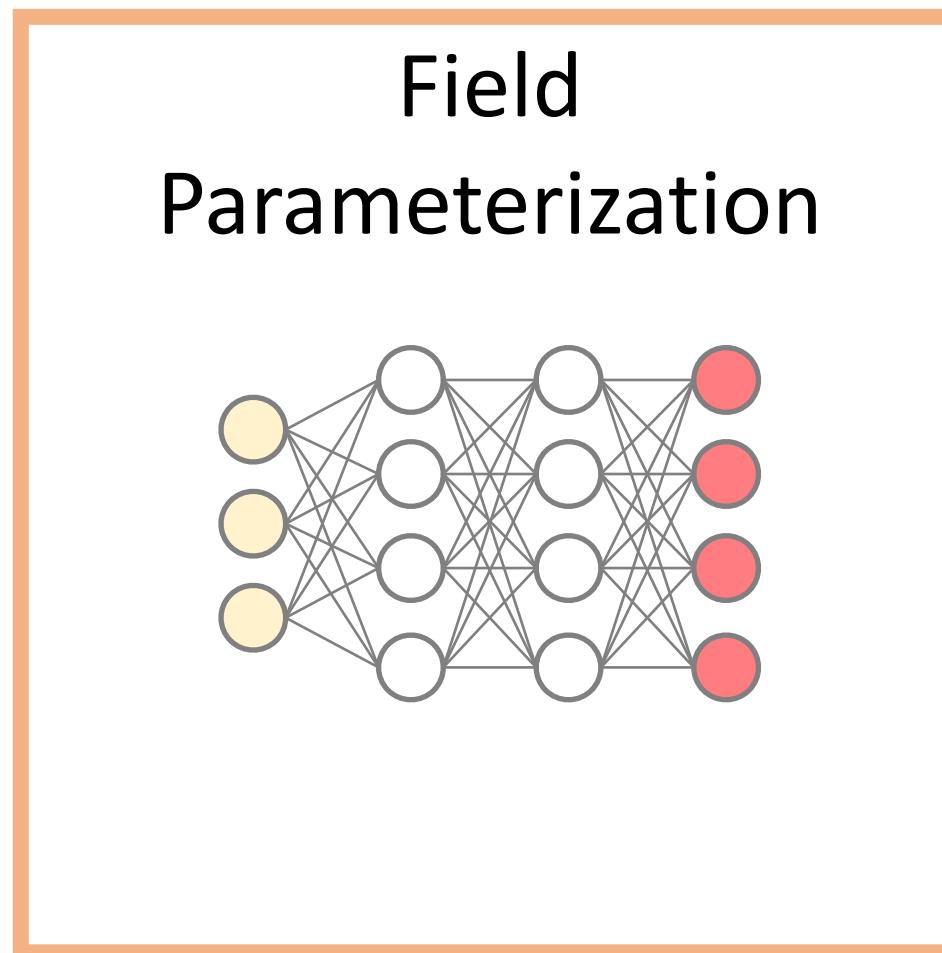


- Reasonable quality overfitting
- Fast rendering
- Noisy geometry
- Requires Many Images

Scene Representation Networks (NeurIPS 2019)



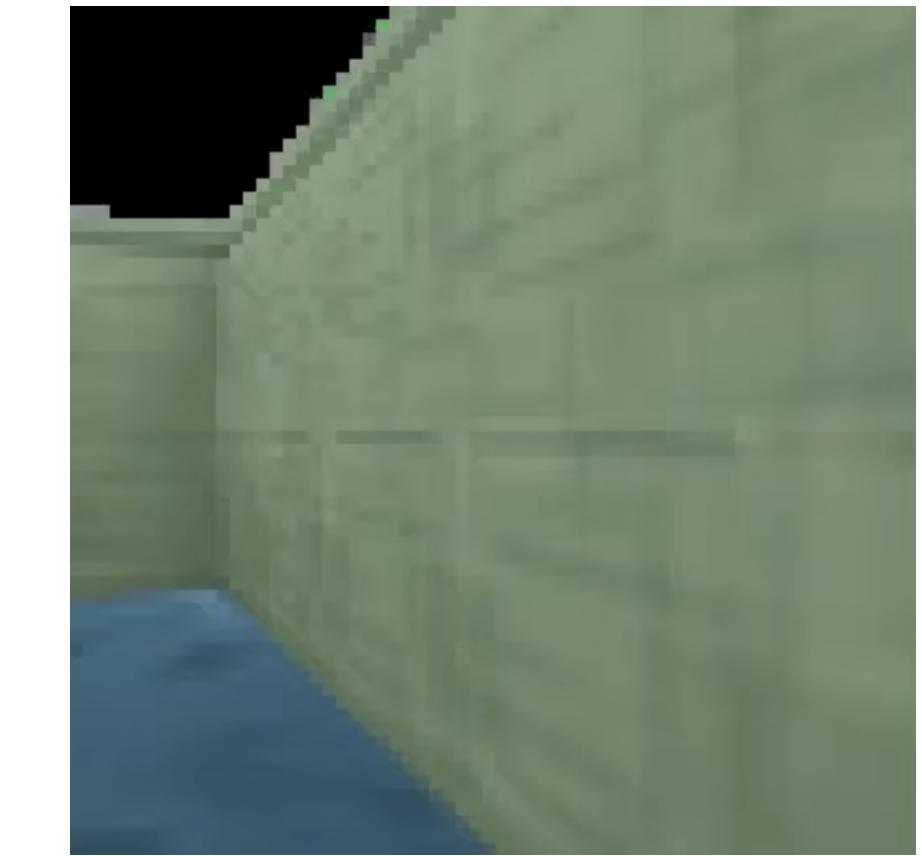
Scene Representation Networks (NeurIPS 2019)



Overfitting

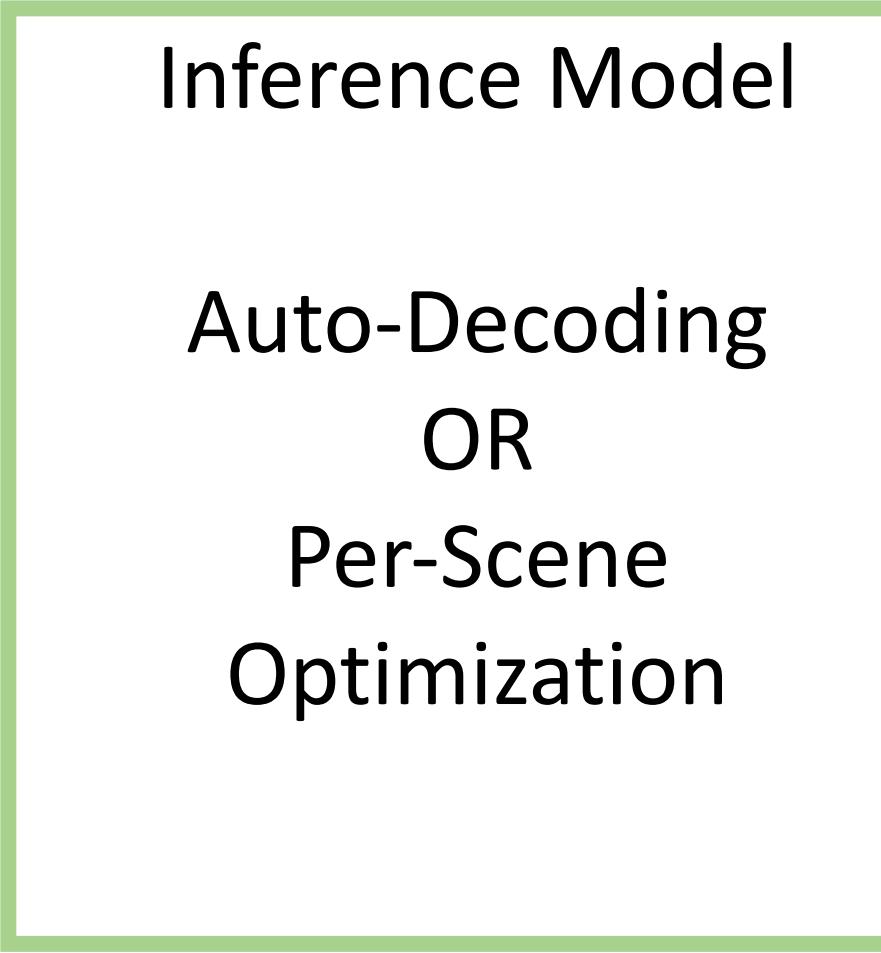
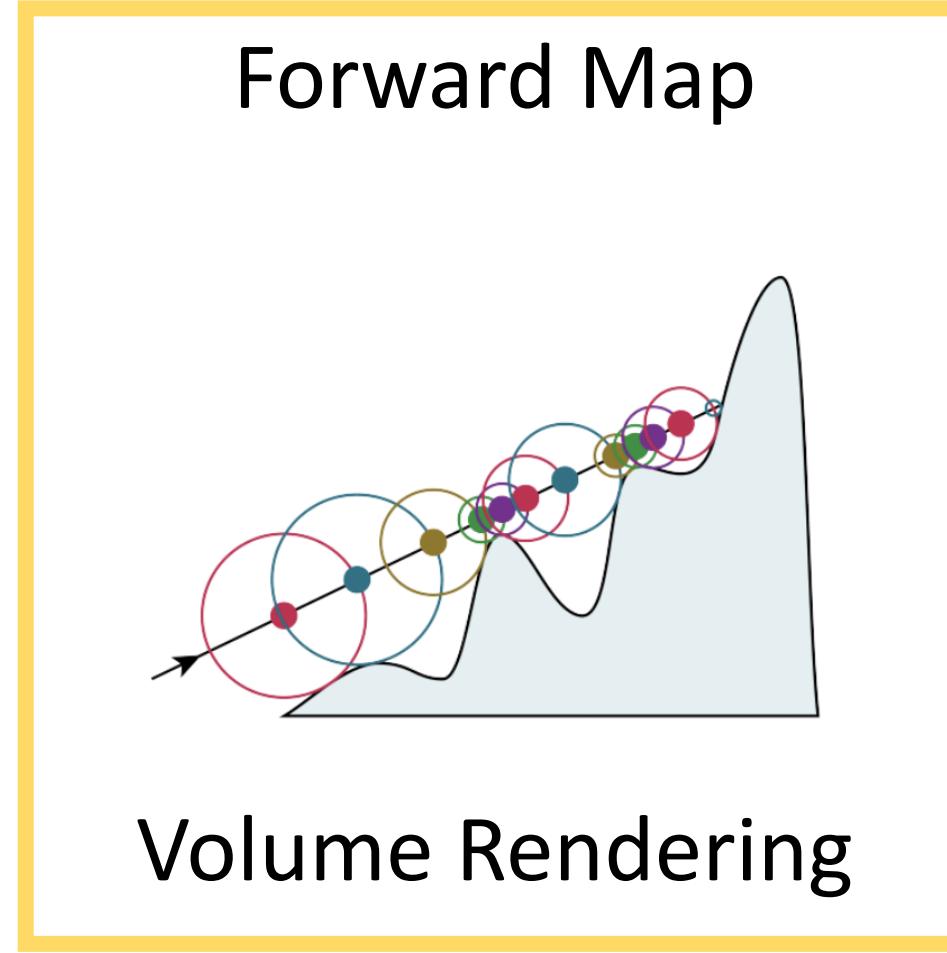
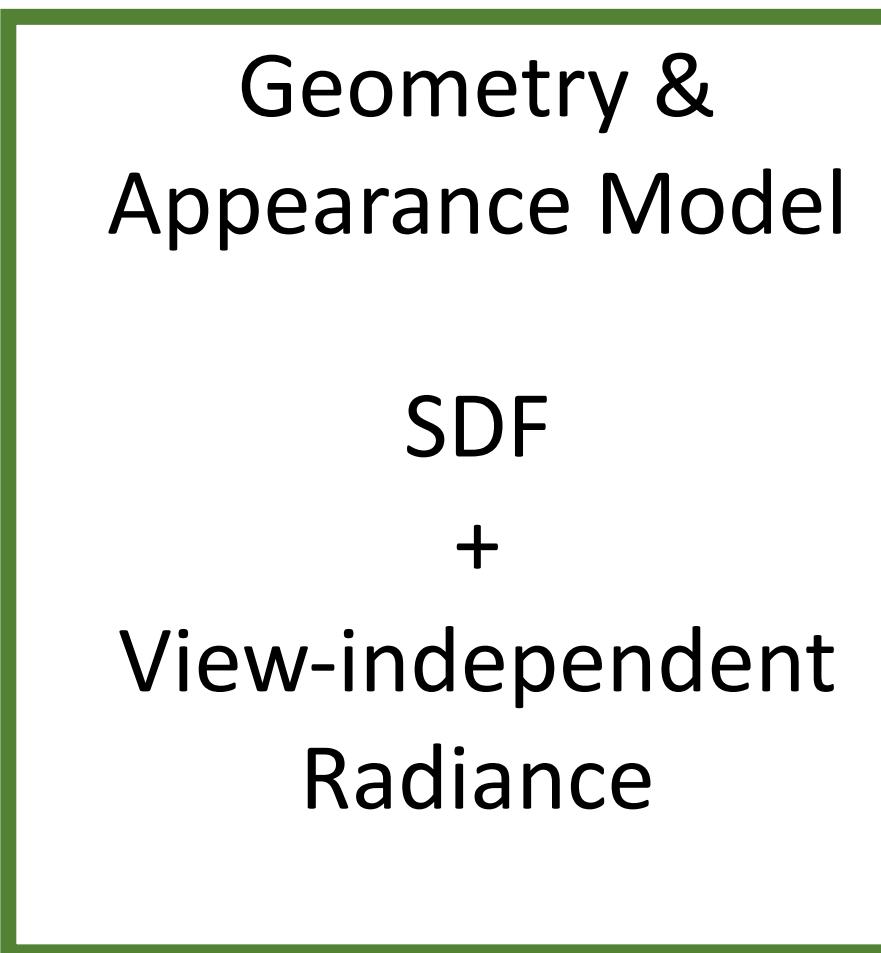
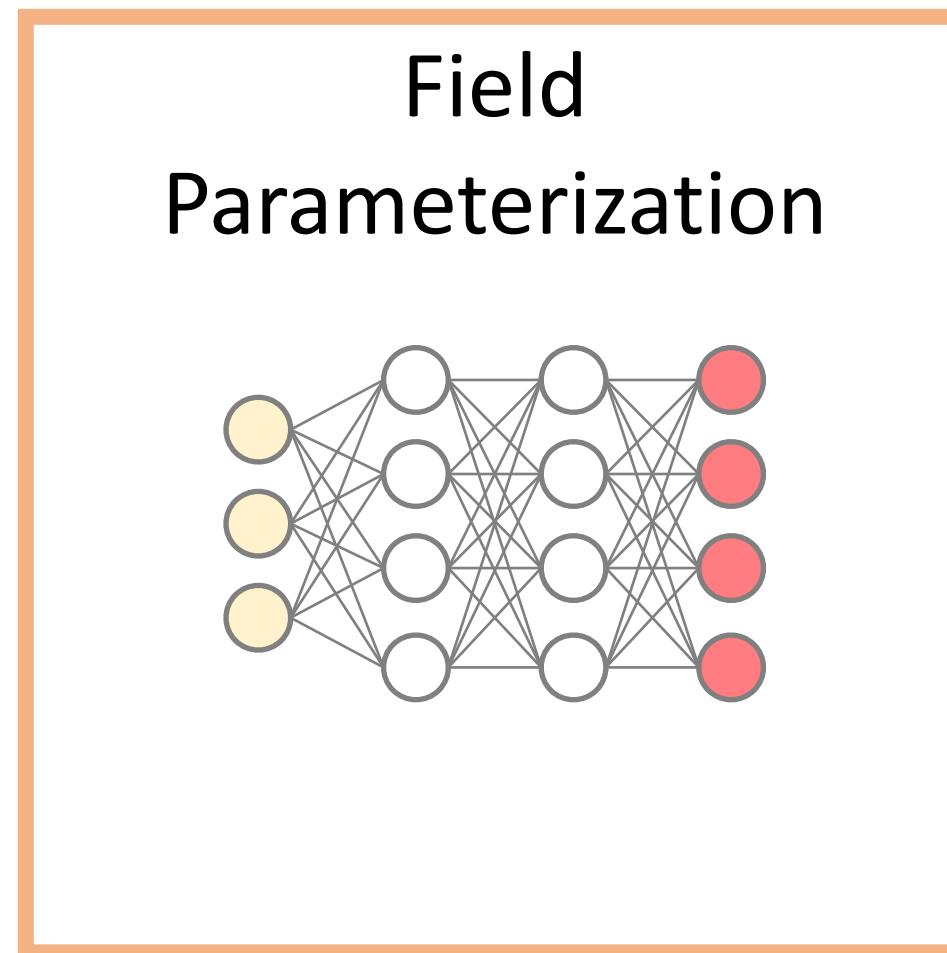


Single Image



- When overfitting: Not photorealistic
- Non-compositional scenes:
Single-image reconstruction
- Well-defined geometry
- Low storage cost (weights)

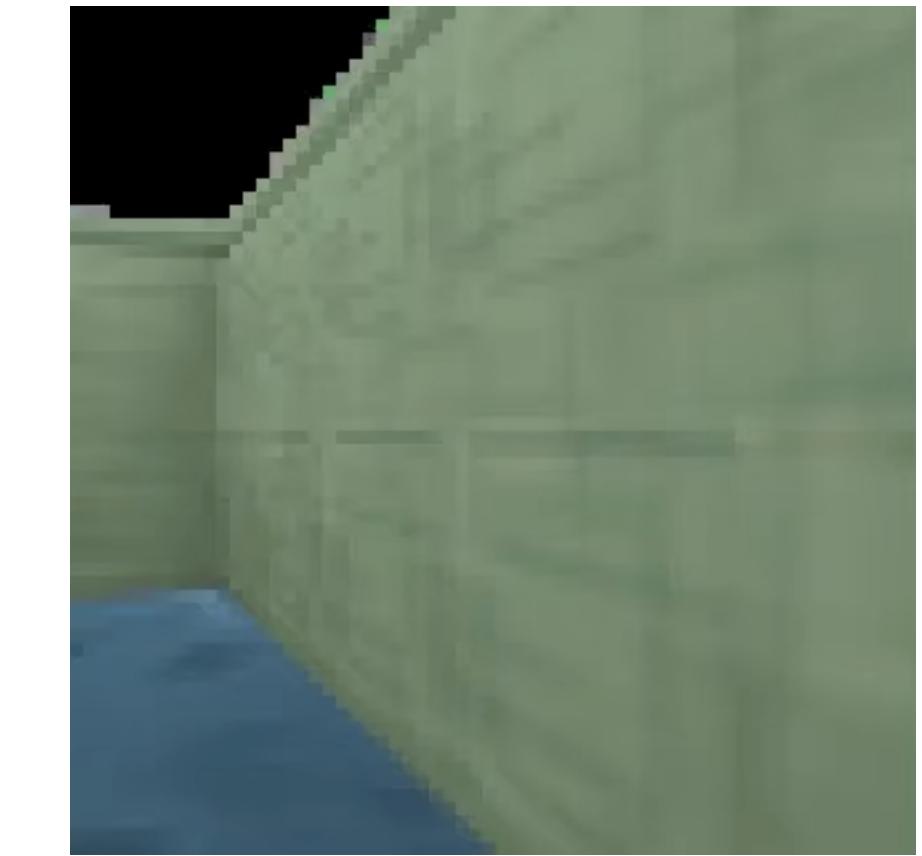
Scene Representation Networks (NeurIPS 2019)



Overfitting



Single Image



- When overfitting: Not photorealistic
- Non-compositional scenes:
Single-image reconstruction
- Well-defined geometry
- Low storage cost (weights)

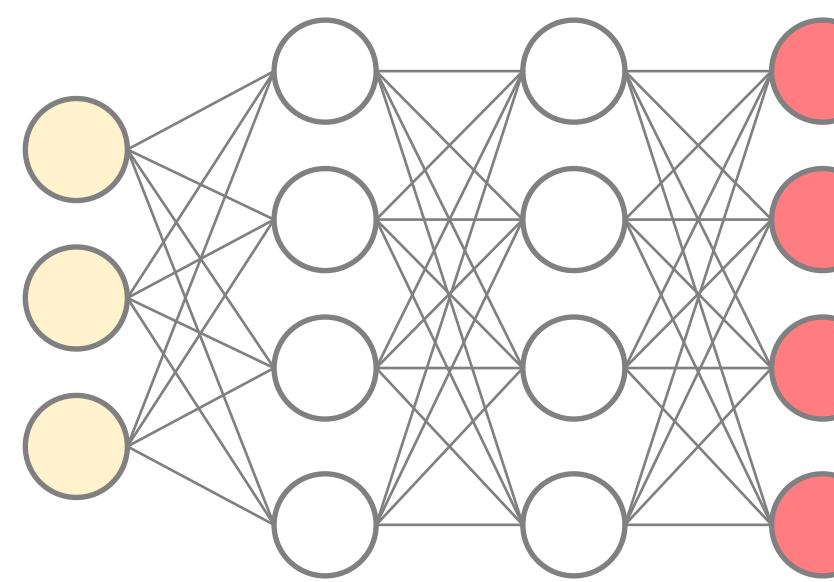
Scene Representation Networks (NeurIPS 2019)



Scene Representation Networks (NeurIPS 2019)



NeRF (Mildenhall et al. ECCV 2020)



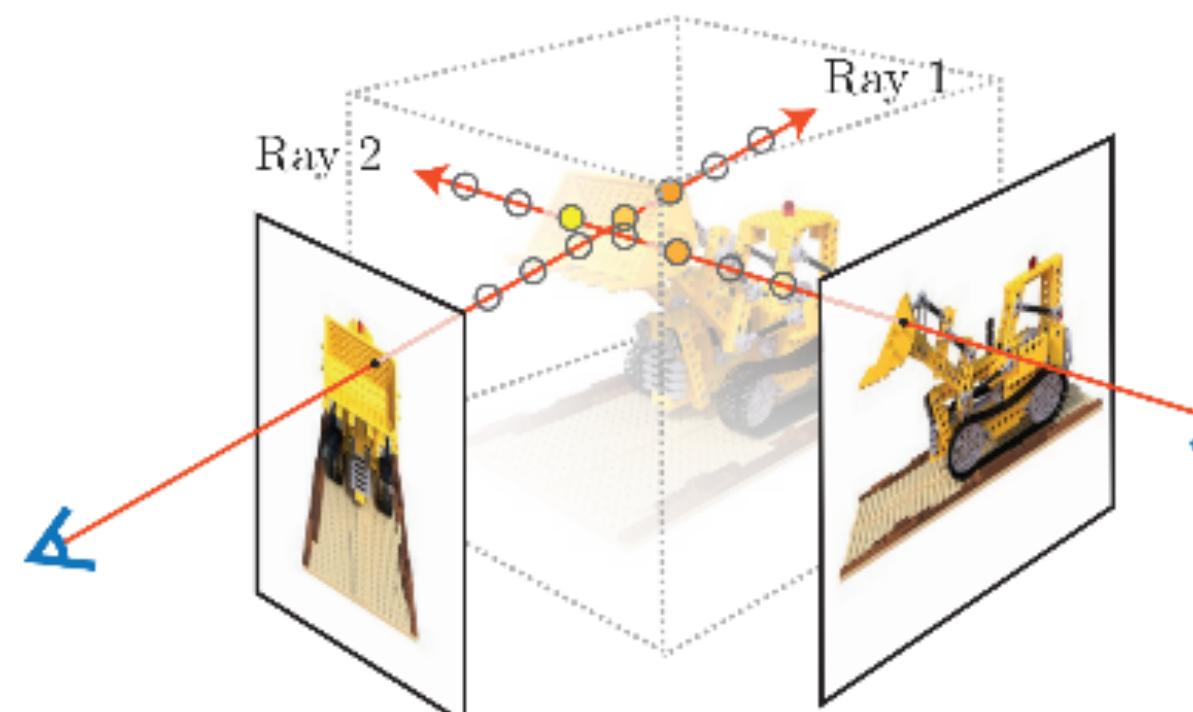
Field Parameterization

$$(x, d) \mapsto (\sigma, C)$$

View-Dependent Density + Radiance

Geometry & Appearance Model

Volume Rendering



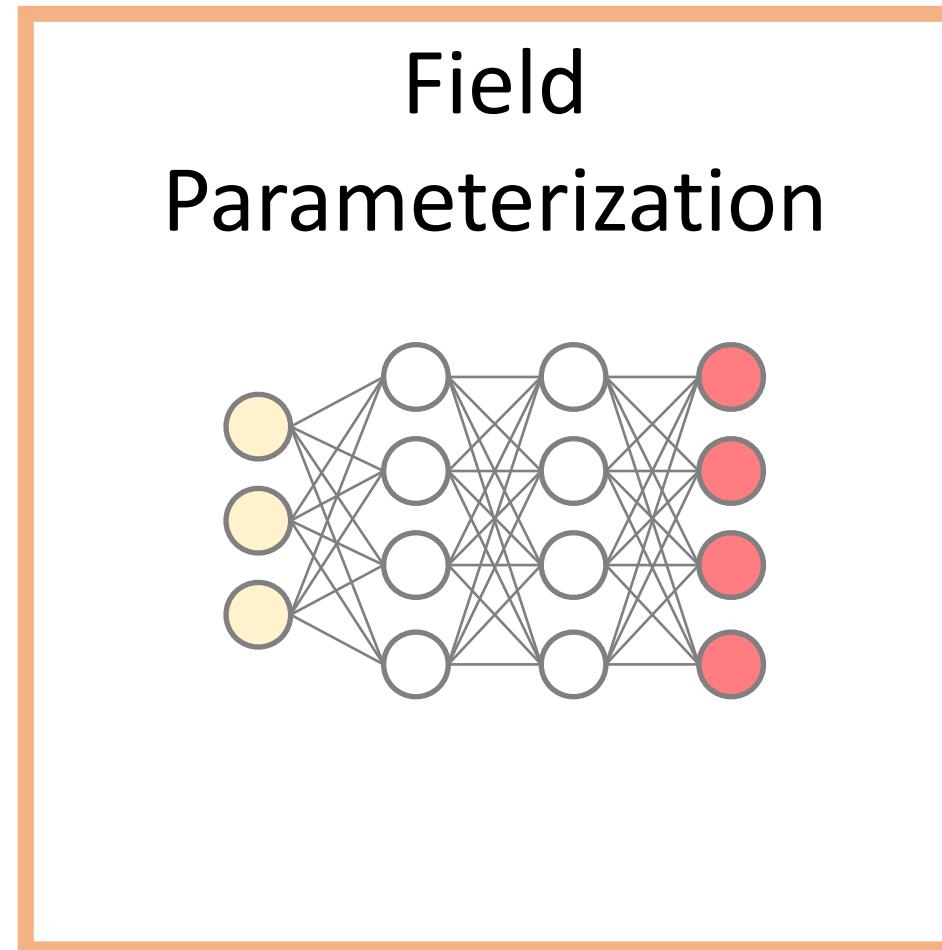
Forward Map

Per-Scene Optimization

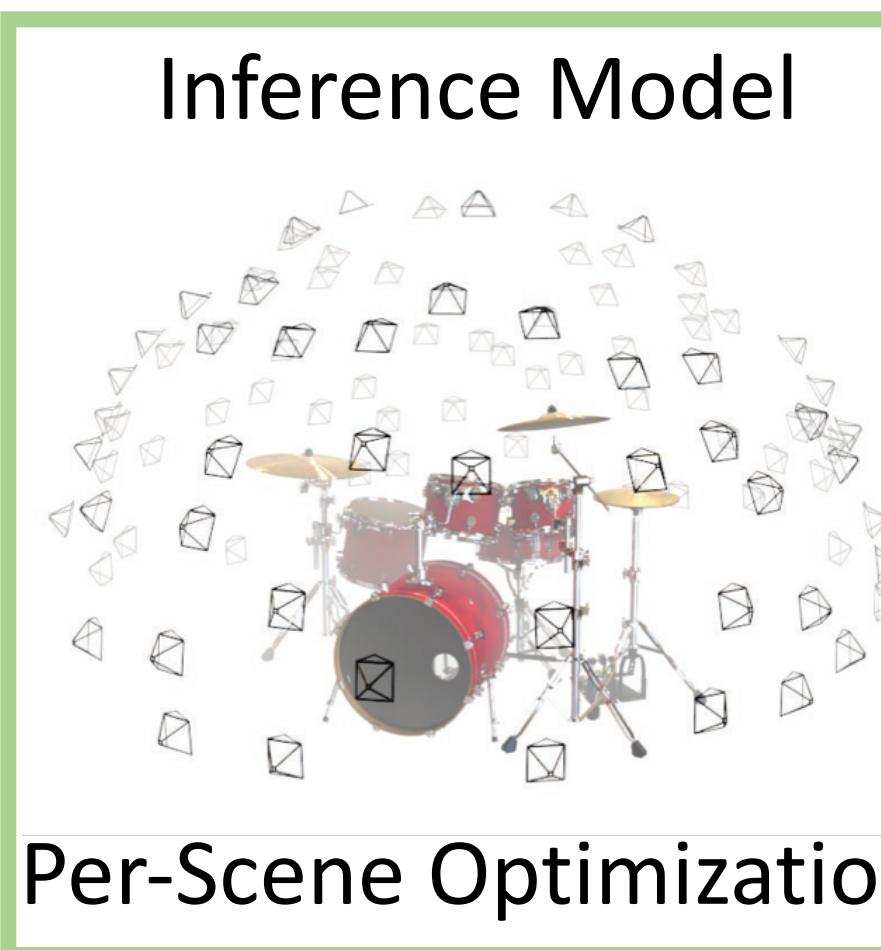
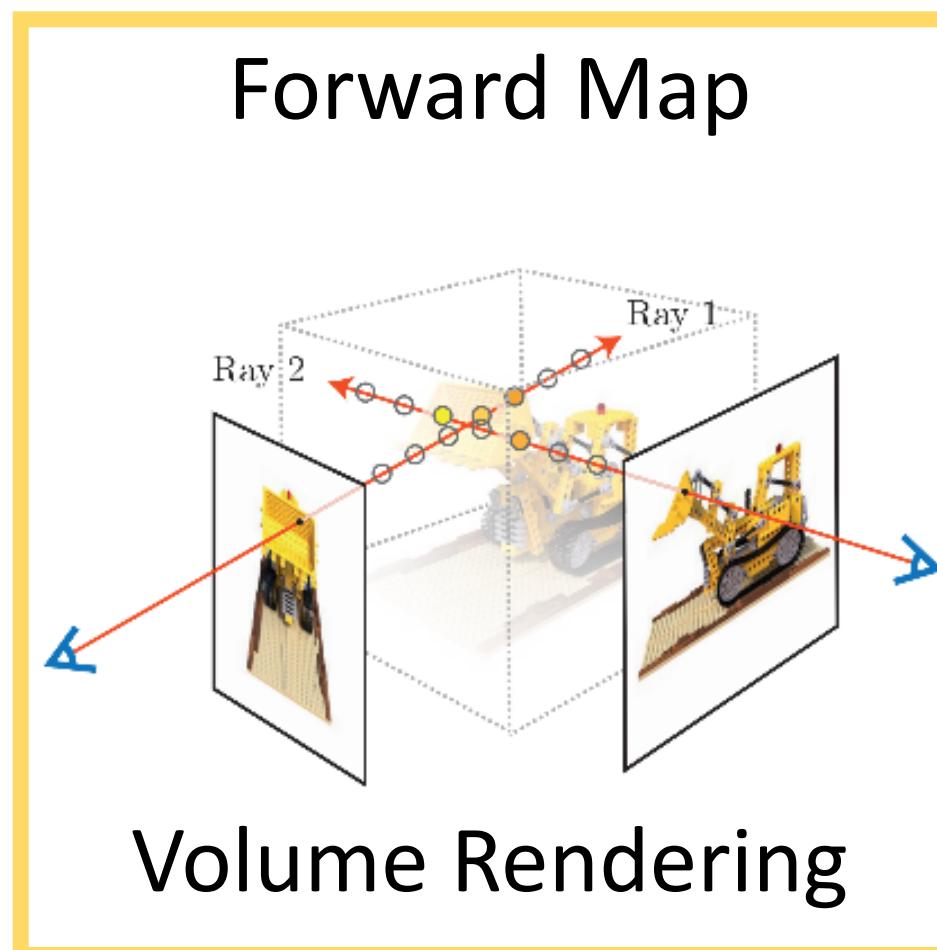


Inference Model

NeRF (Mildenhall et al. ECCV 2020)

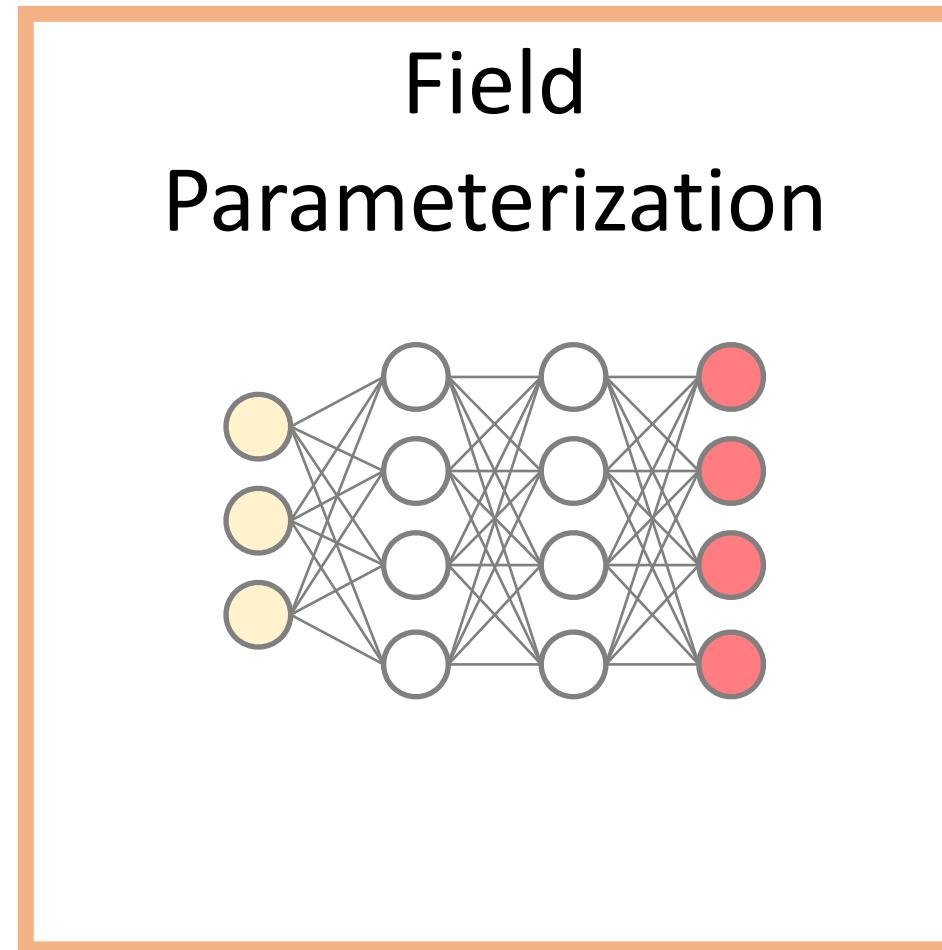


Geometry & Appearance Model
 $(x, d) \mapsto (\sigma, C)$
View-Dependent Density + Radiance

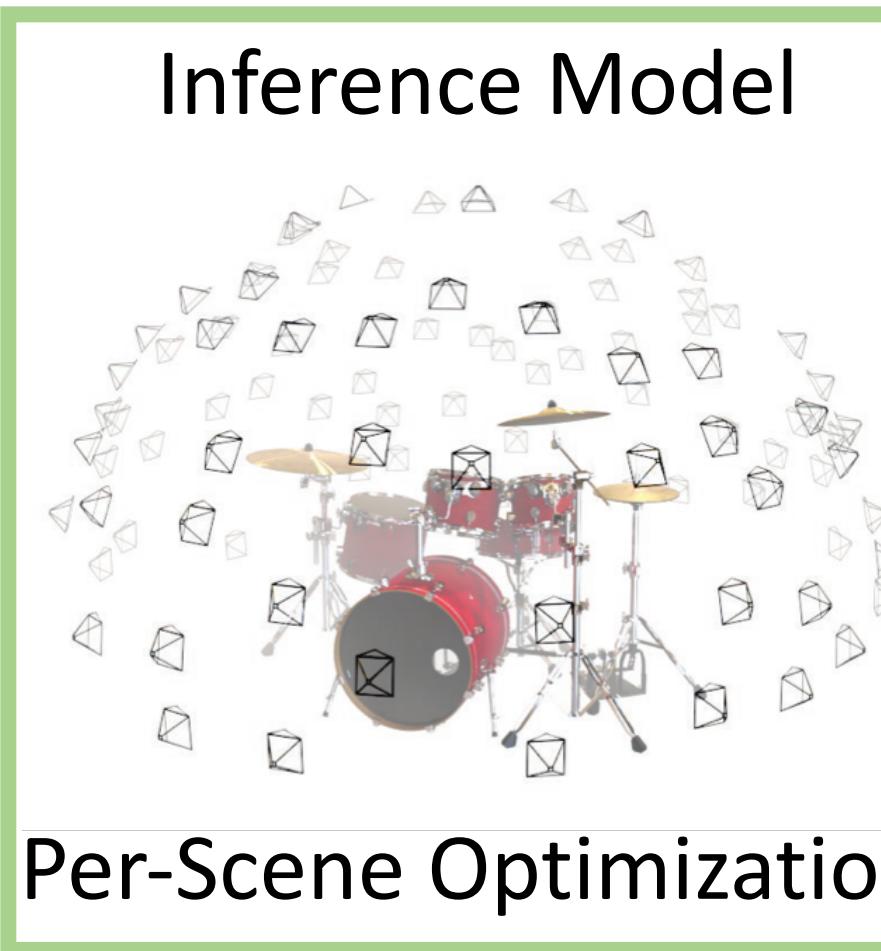
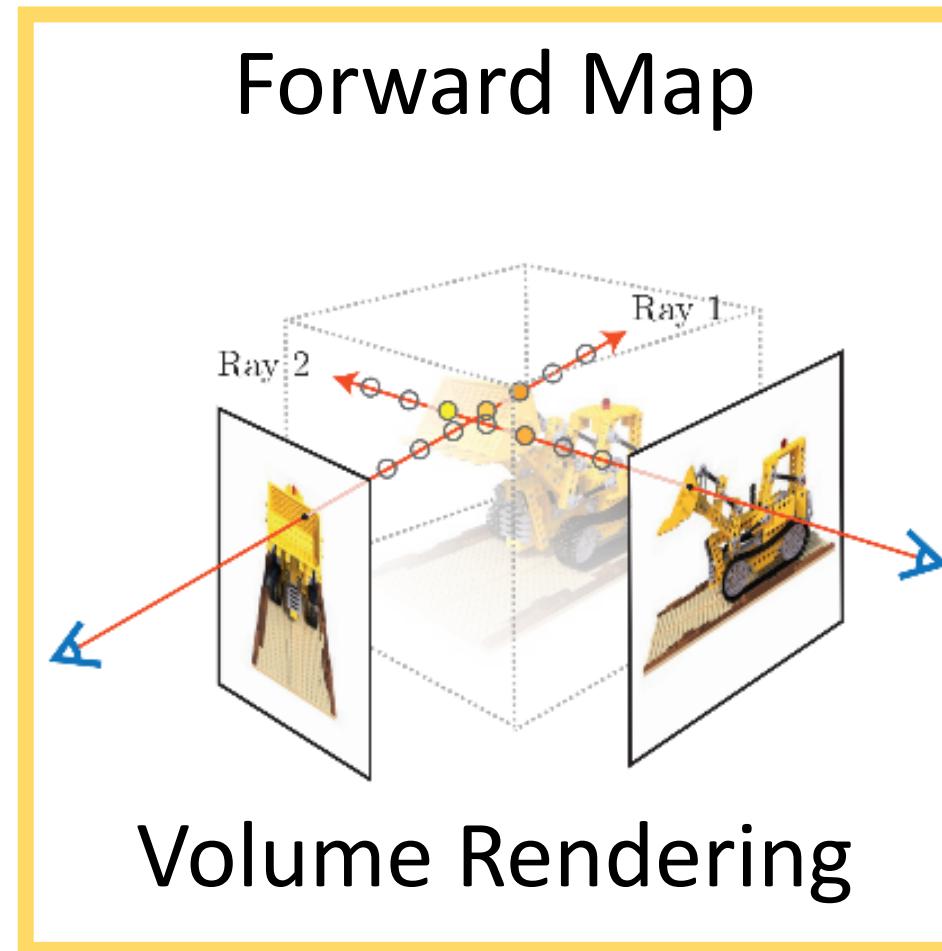


- When overfitting: Almost photorealistic!
- Slow Rendering
- High memory cost in forward pass
- Low Storage Cost (weights)
- Noisy or Cloudy Geometry
- Requires Many Images

NeRF (Mildenhall et al. ECCV 2020)



Geometry & Appearance Model
 $(x, d) \mapsto (\sigma, C)$
View-Dependent Density + Radiance

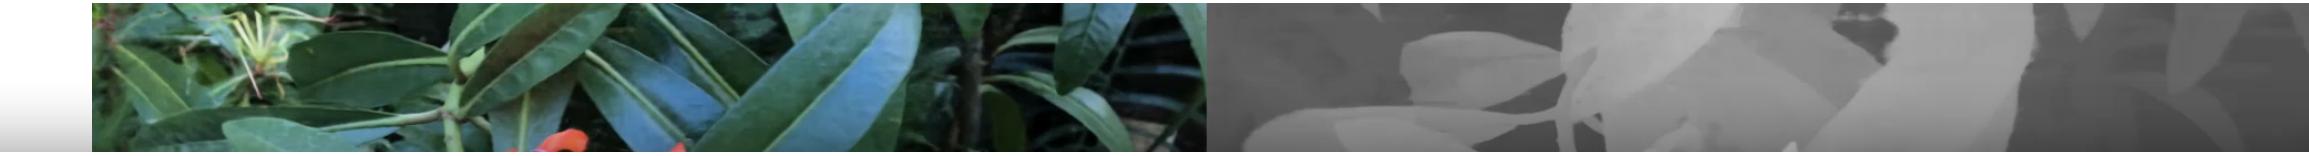


- When overfitting: Almost photorealistic!
- Slow Rendering
- High memory cost in forward pass
- Low Storage Cost (weights)
- Noisy or Cloudy Geometry
- Requires Many Images

NeRF (Mildenhall et al. ECCV 2020)

Field

Geometry &



Great out-of-the-box performance
Great looking results for single-scene reconstruction
Ben donated slides ;)

ic!

**Inspired a boom in 3D reconstruction via differentiable
volumetric rendering!**



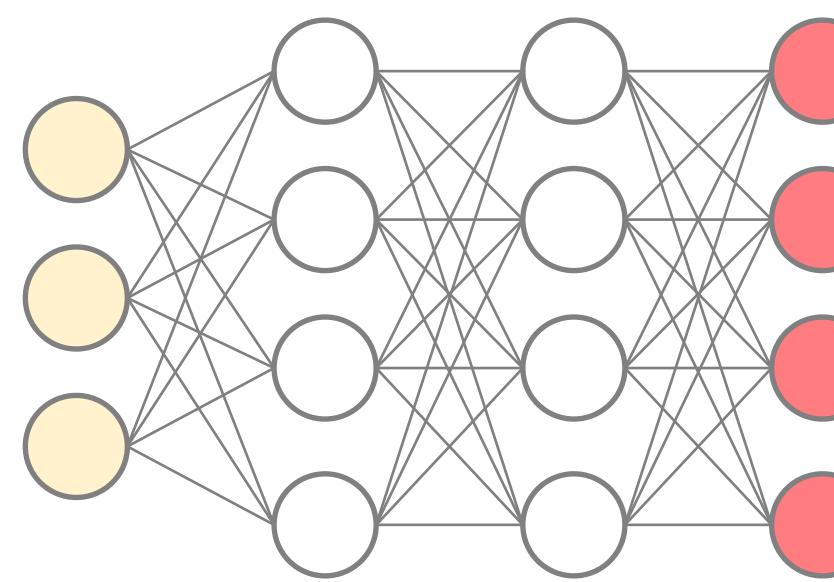
Volume Rendering



Per-Scene Optimization

- Noisy or Cloudy Geometry
- Requires Many Images

NeRF (Mildenhall et al. ECCV 2020)



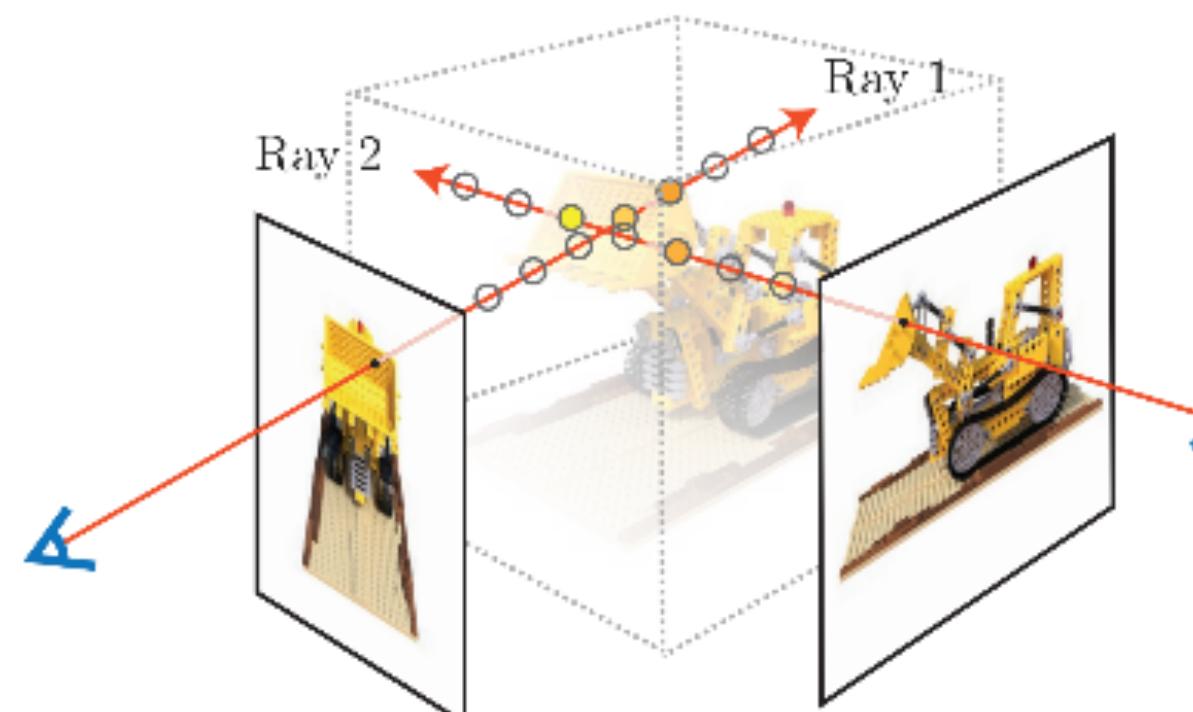
Field Parameterization

$$(x, d) \mapsto (\sigma, C)$$

View-Dependent Density + Radiance

Geometry & Appearance Model

Volume Rendering



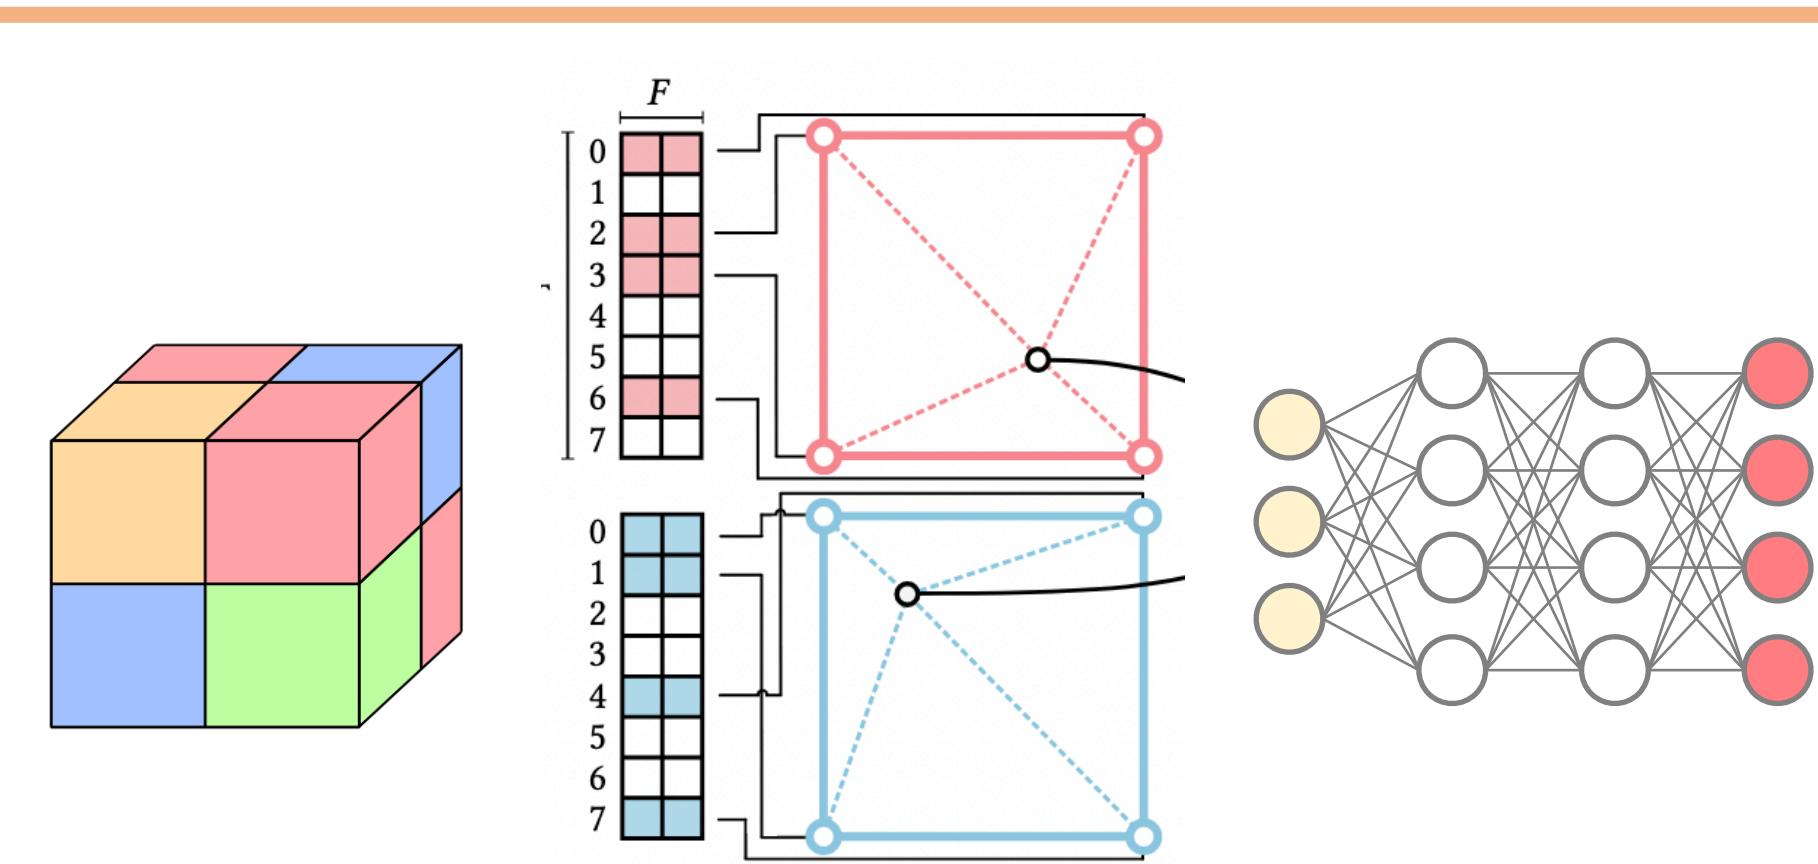
Forward Map

Per-Scene Optimization



Inference Model

Faster Radiance Fields w/ Hybrid Data Structures

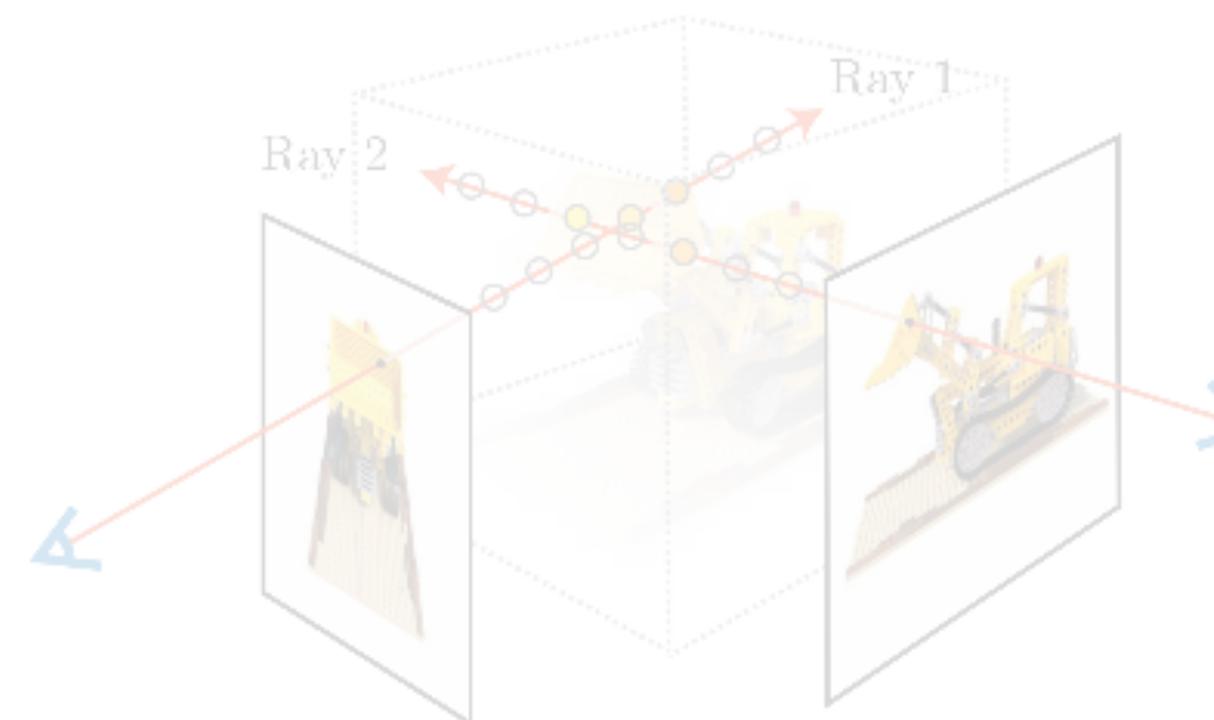


$$(x, d) \mapsto (\sigma, C)$$

View-Dependent Density + Radiance

Geometry & Appearance Model

Volume Rendering

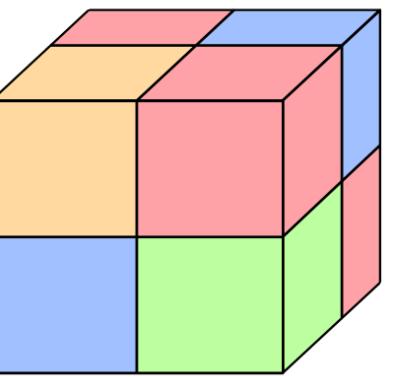


Forward Map

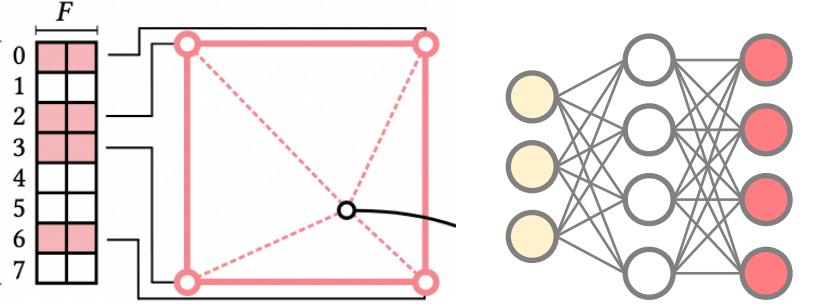
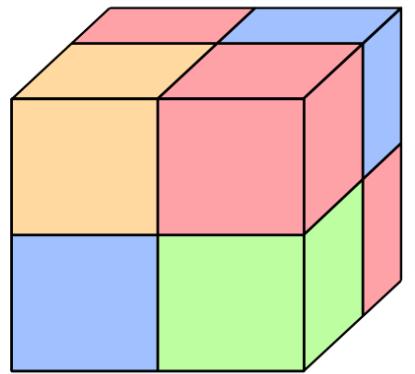
Per-Scene Optimization



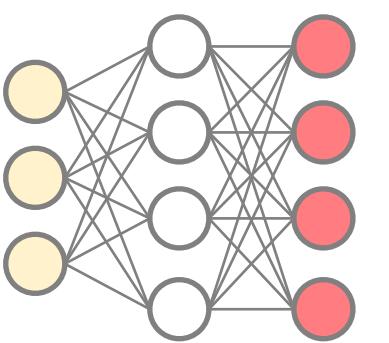
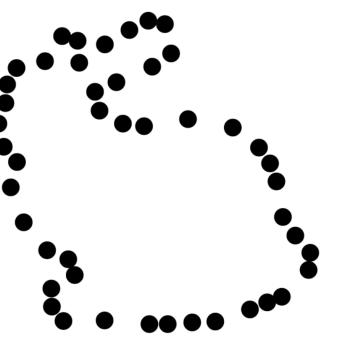
Inference Model



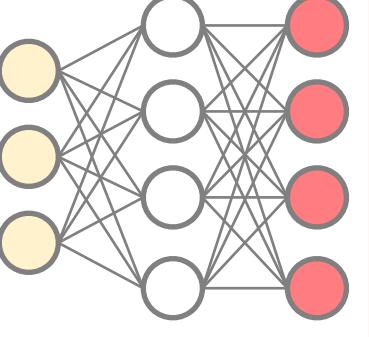
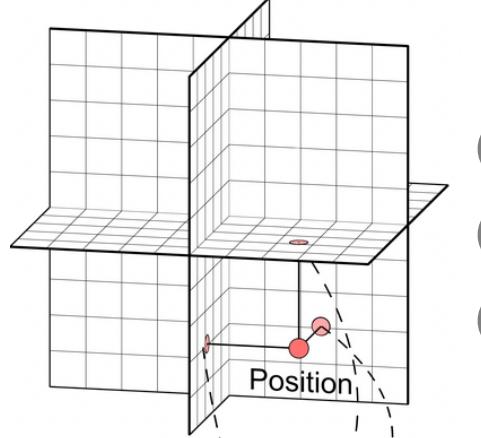
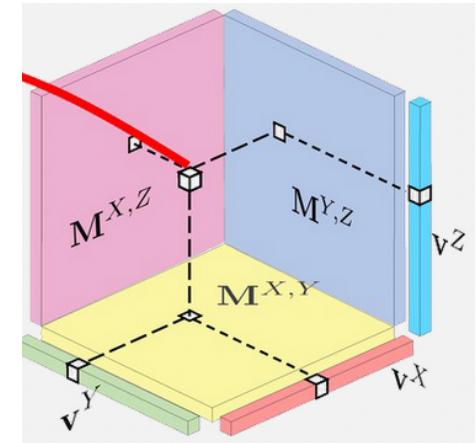
Voxel Grid



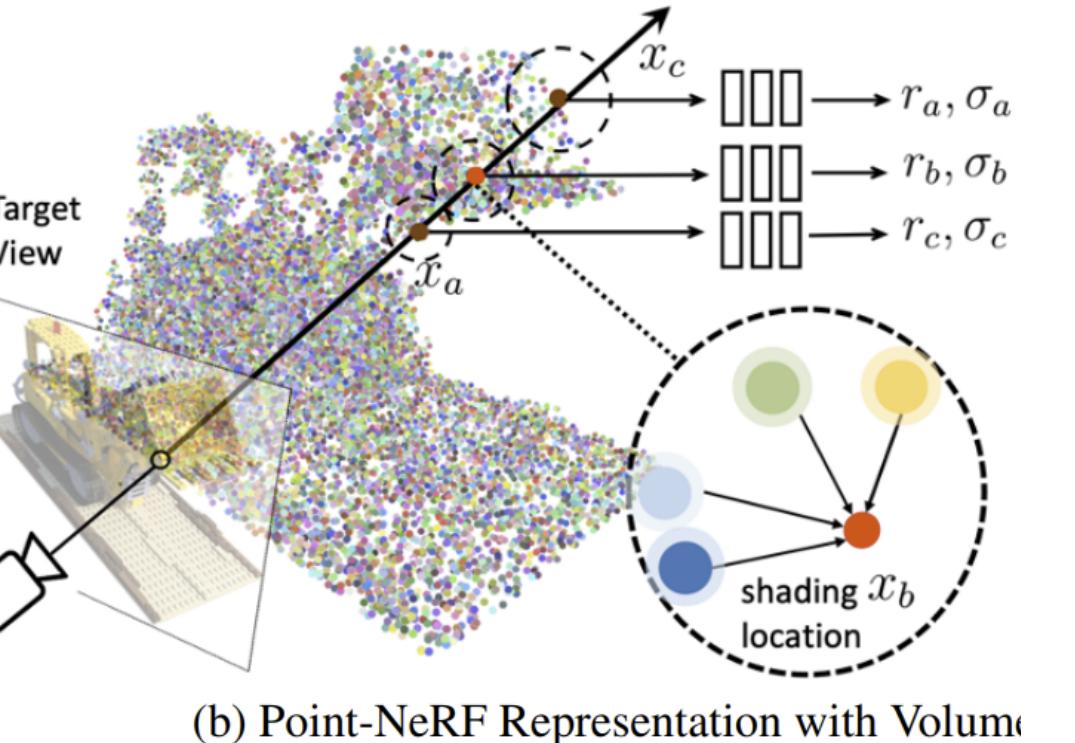
Voxel Grid + Hashmap + MLP



Point Cloud + MLP



Tensor Factorization / Triplane + MLP



(b) Point-NeRF Representation with Volume



Plenoxels: Radiance Fields ...

[Yu et al. 2022]

Direct Voxel Grid Optimization

[Sun et al. 2021]

InstantNGP: Instant Neural ...

[Müller et al. 2022]

PointNeRF: Point-based Neural ...

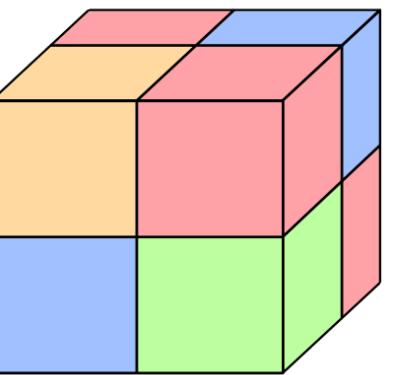
[Xu et al. 2022]

Efficient Geometry-aware 3D...

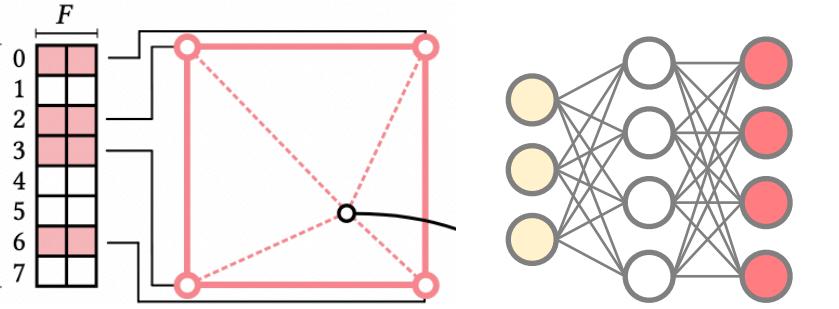
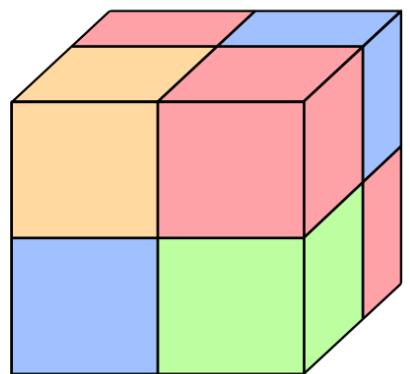
[Chan et al. 2022]

TensorRF: Tensor Radiance Fields

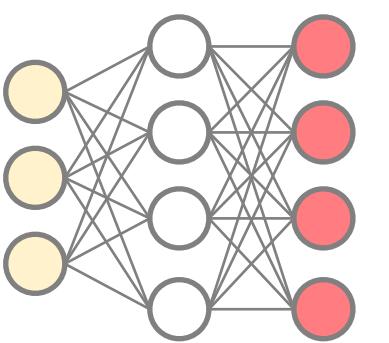
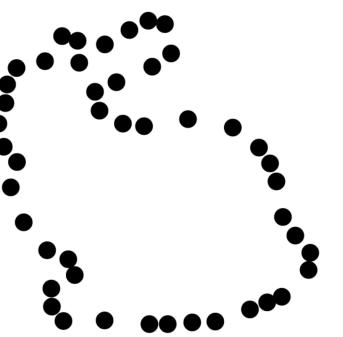
[Chen & Xu et al. 2022]



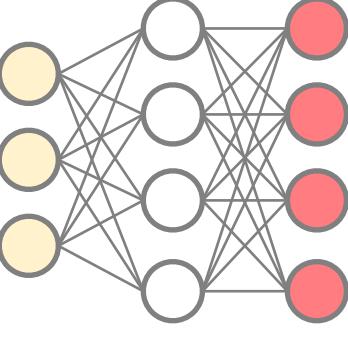
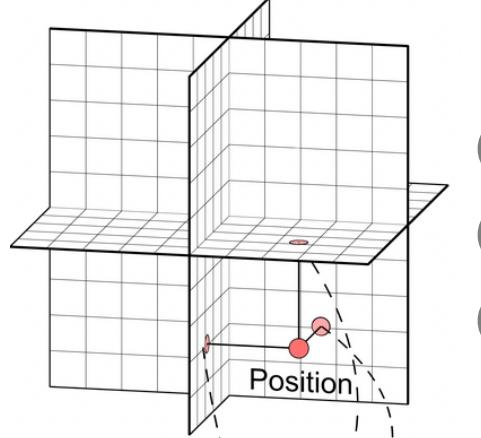
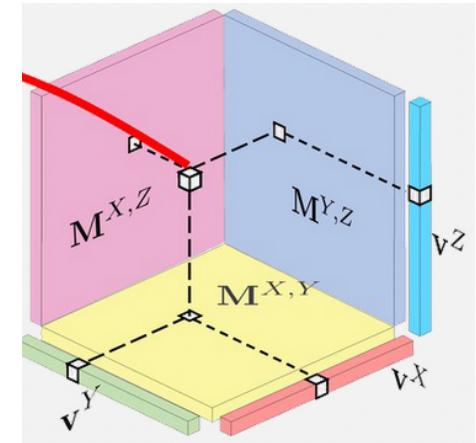
Voxel Grid



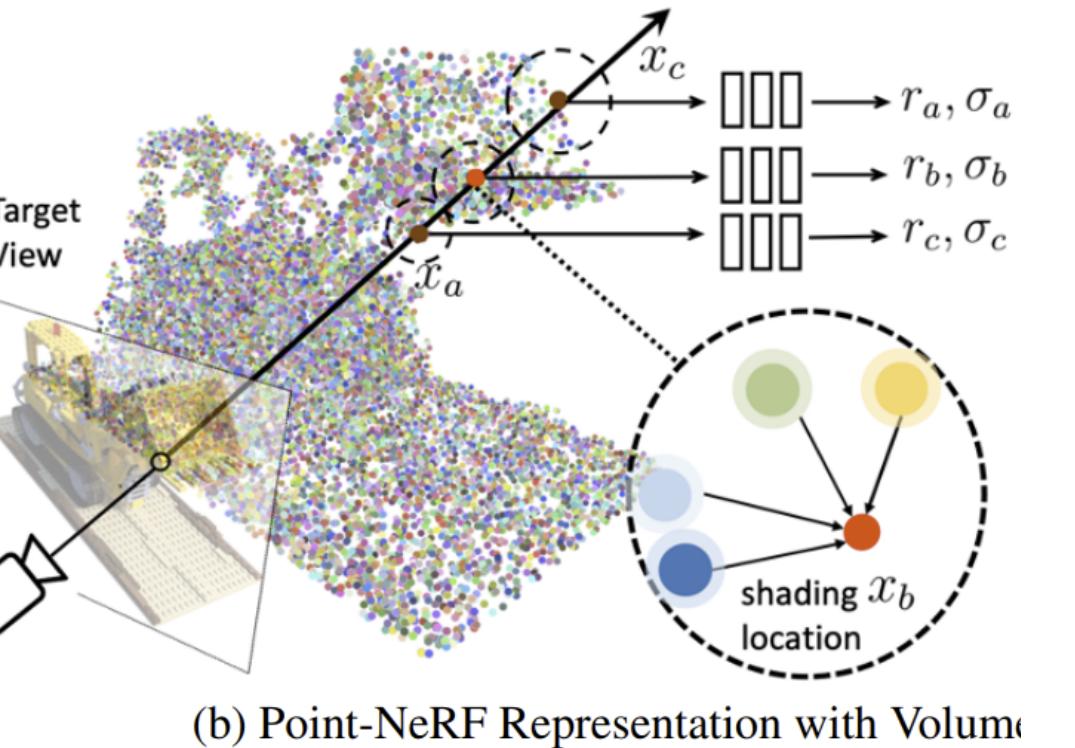
Voxel Grid + Hashmap + MLP



Point Cloud + MLP



Tensor Factorization / Triplane + MLP



(b) Point-NeRF Representation with Volume



Plenoxels: Radiance Fields ...

[Yu et al. 2022]

Direct Voxel Grid Optimization

[Sun et al. 2021]

InstantNGP: Instant Neural ...

[Müller et al. 2022]

PointNeRF: Point-based Neural ...

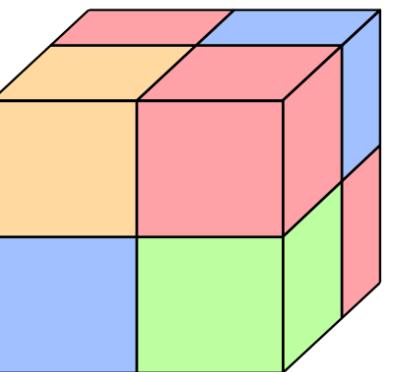
[Xu et al. 2022]

Efficient Geometry-aware 3D...

[Chan et al. 2022]

TensorRF: Tensor Radiance Fields

[Chen & Xu et al. 2022]

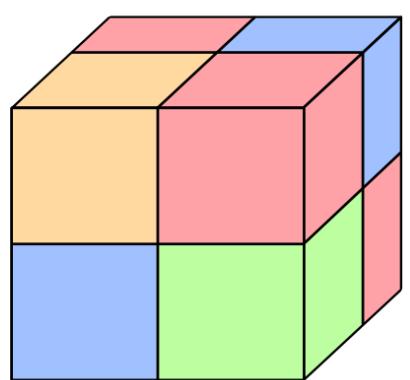


Plenoxels: Radiance Fields ...

[Yu et al. 2022]

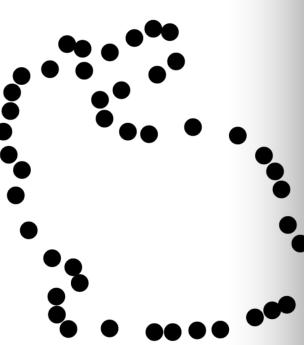
Direct Voxel Grid Optimization

[Sun et al. 2021]

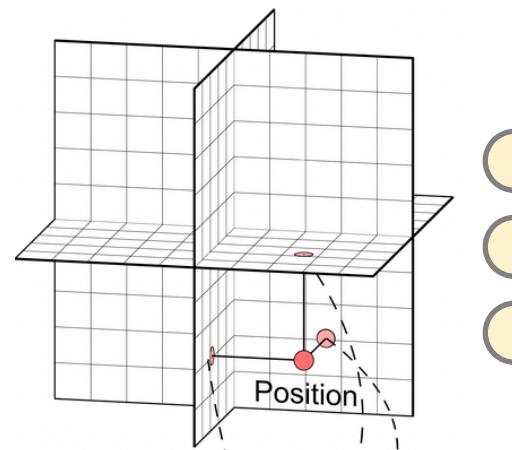
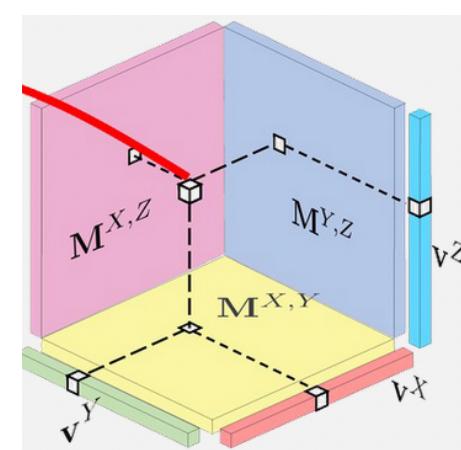


Voxel Grid

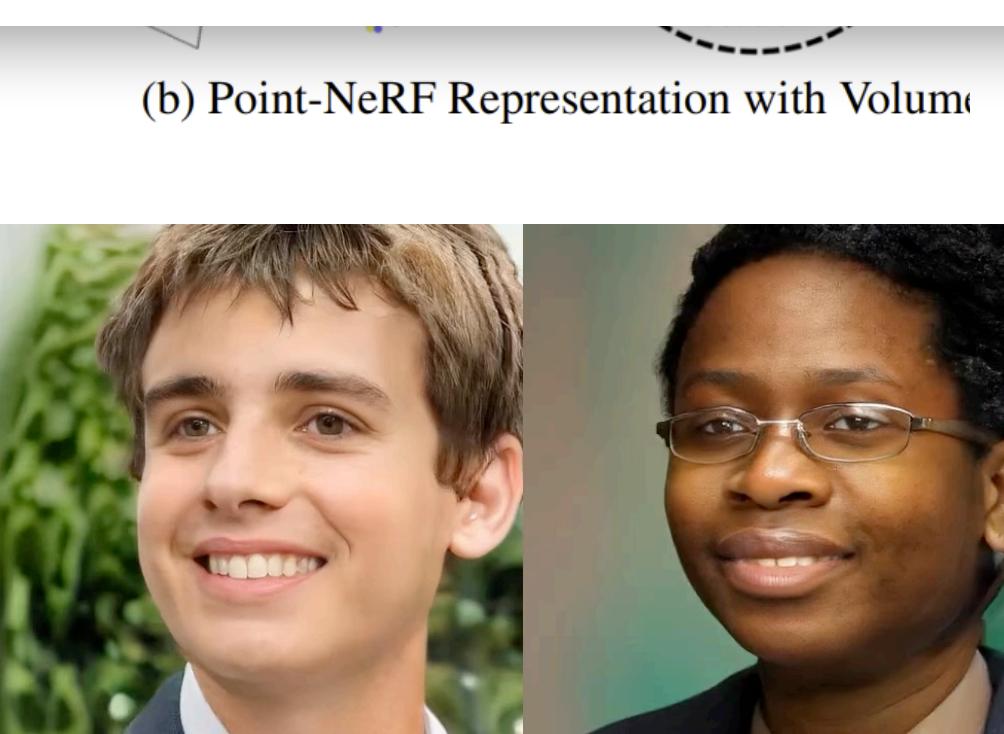
- Faster fitting
- Faster rendering
- Lower forward pass memory footprint
- May achieve higher quality than vanilla NeRF
- Voxels: Memory intense



Point Cloud + MLP



Tensor Factorization / Triplane + MLP



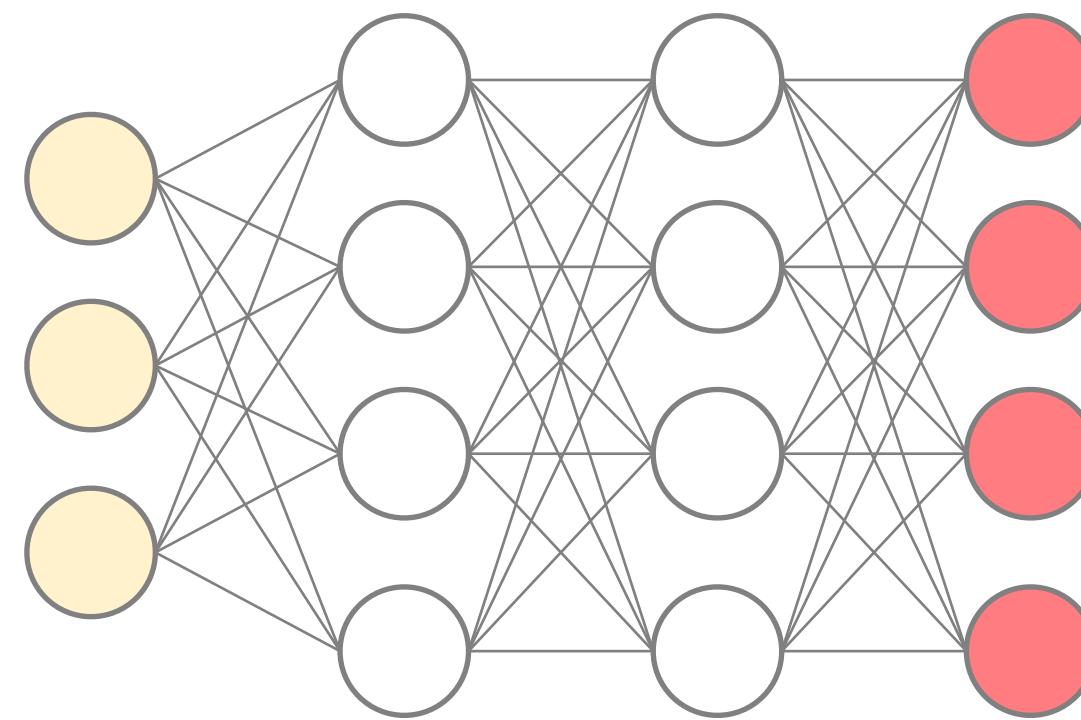
Efficient Geometry-aware 3D...

[Chan et al. 2022]

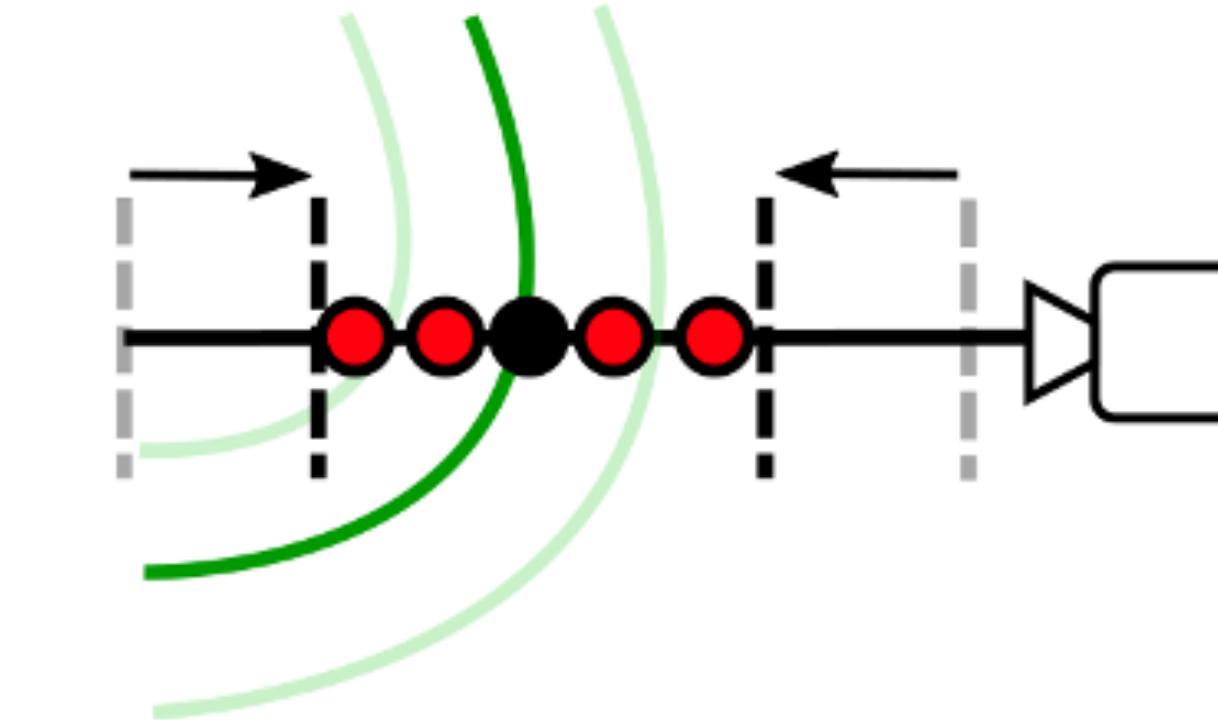
TensorRF: Tensor Radiance Fields

[Chen & Xu et al. 2022]

Hybrid Volumetric / Level Set Representations

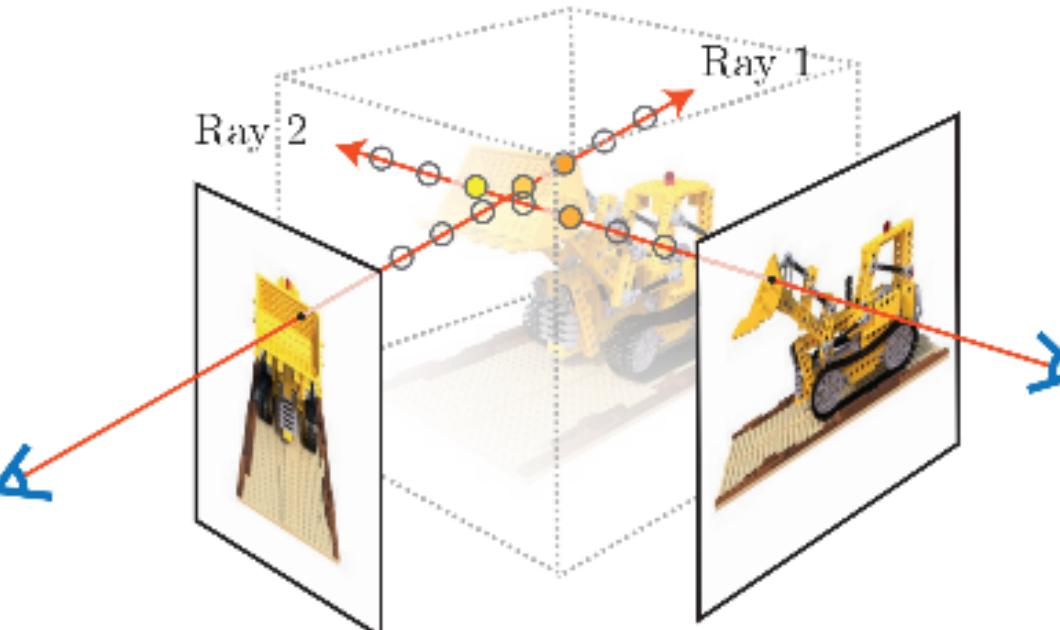


Field Parameterization



Geometry & Appearance Model

Volume Rendering



Forward Map

Per-Scene Optimization

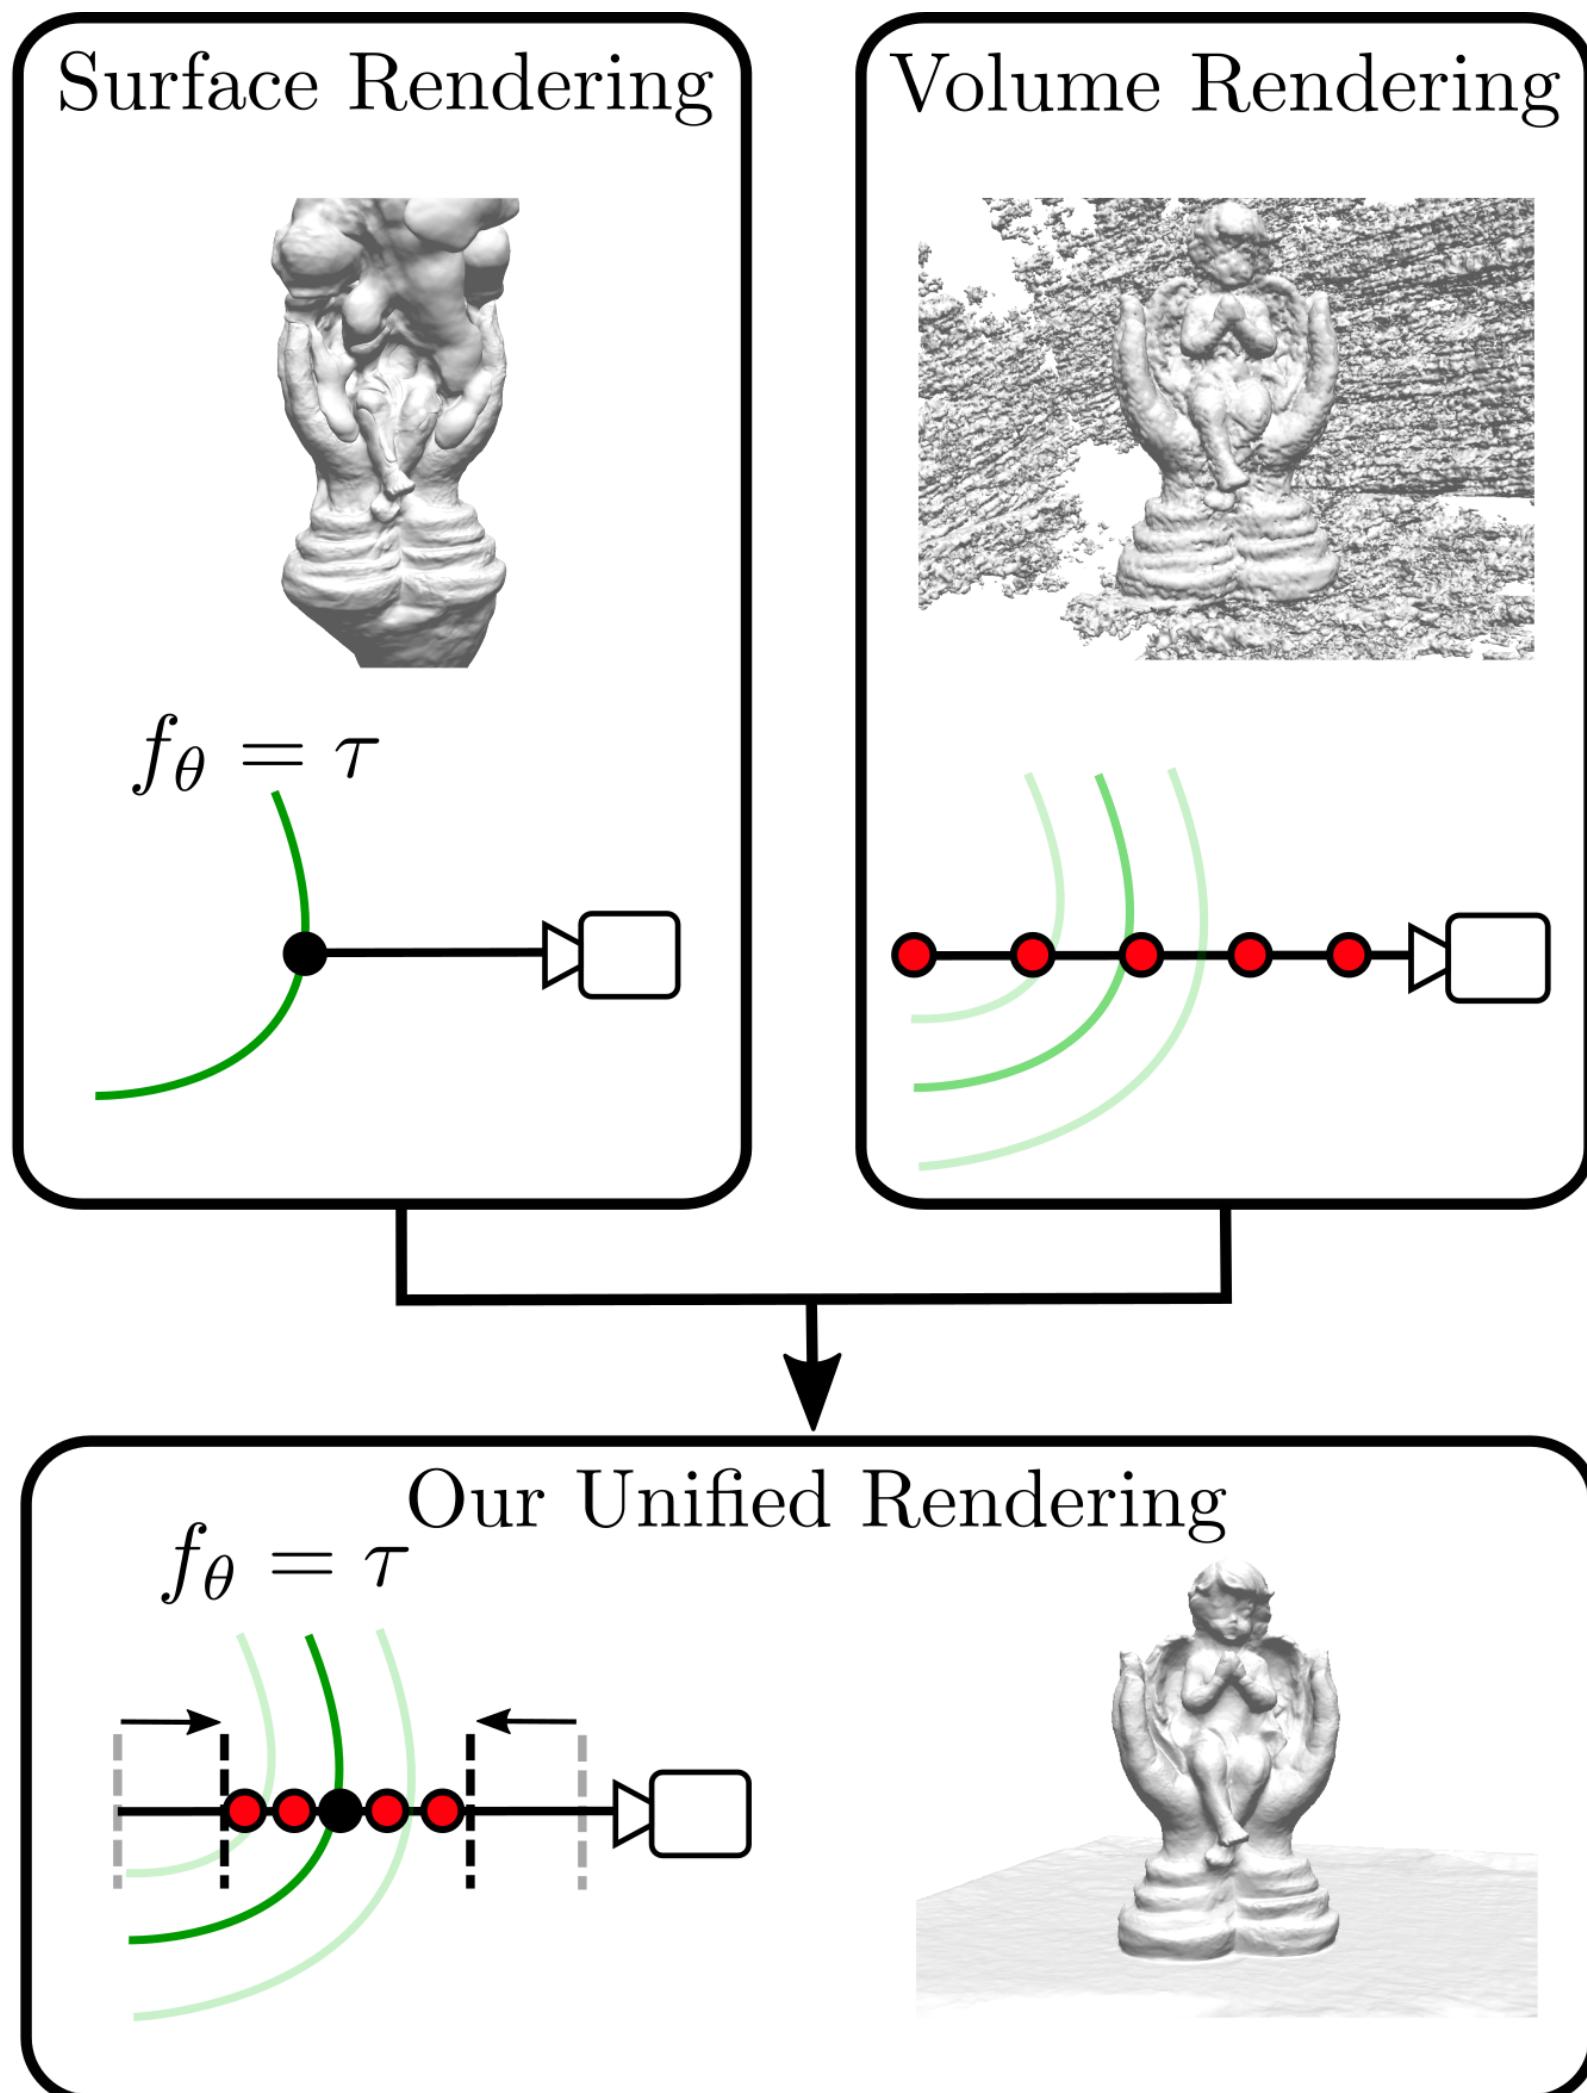


Inference Model

Hybrid Volumetric / Level Set Representations

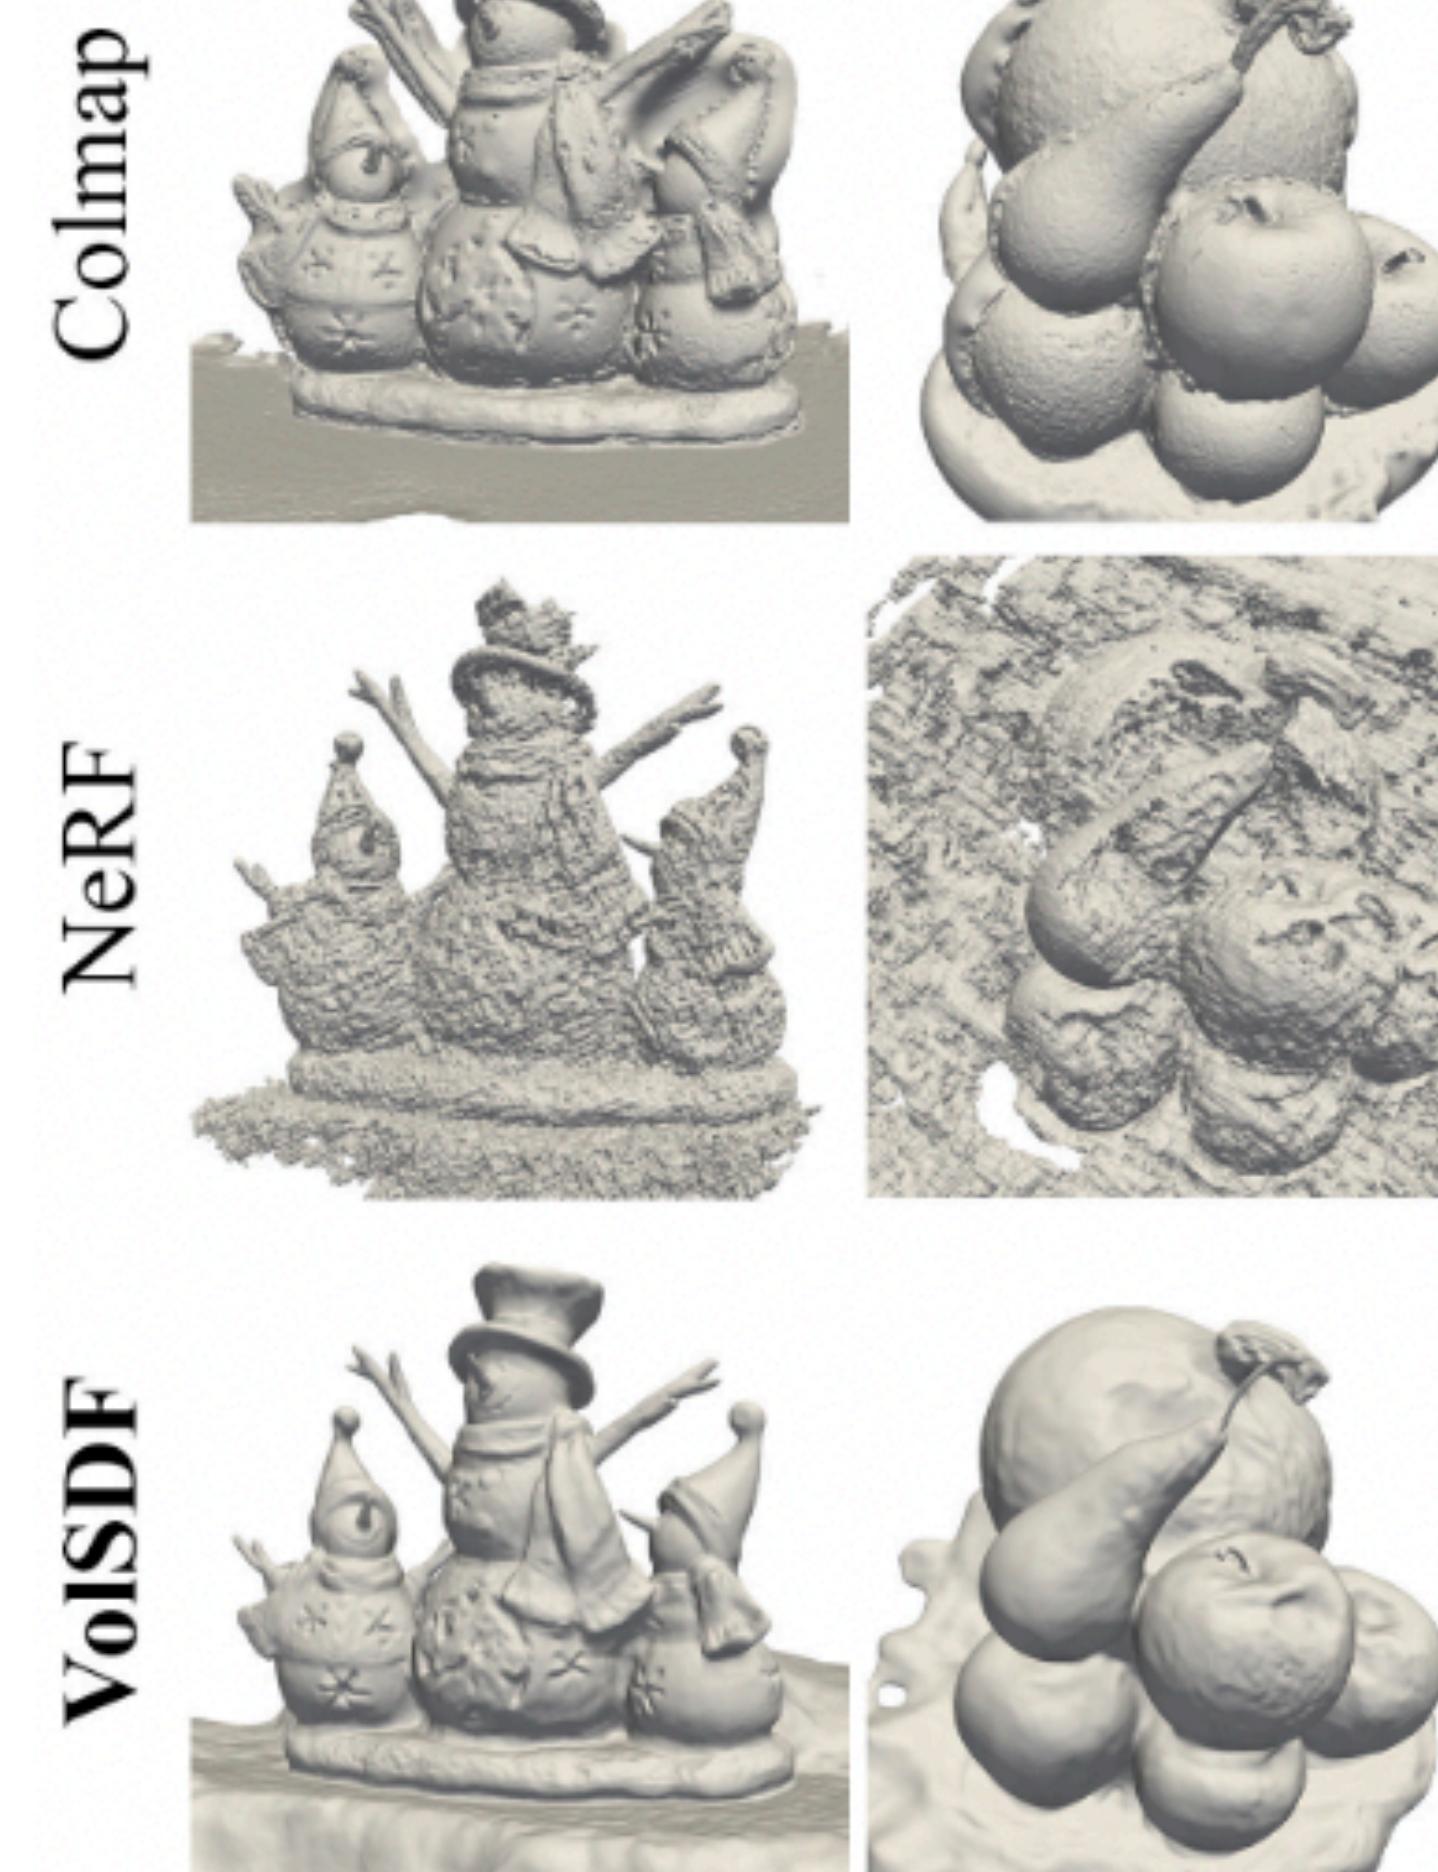
UNISURF

[Oechsle et al., ICCV 2021]



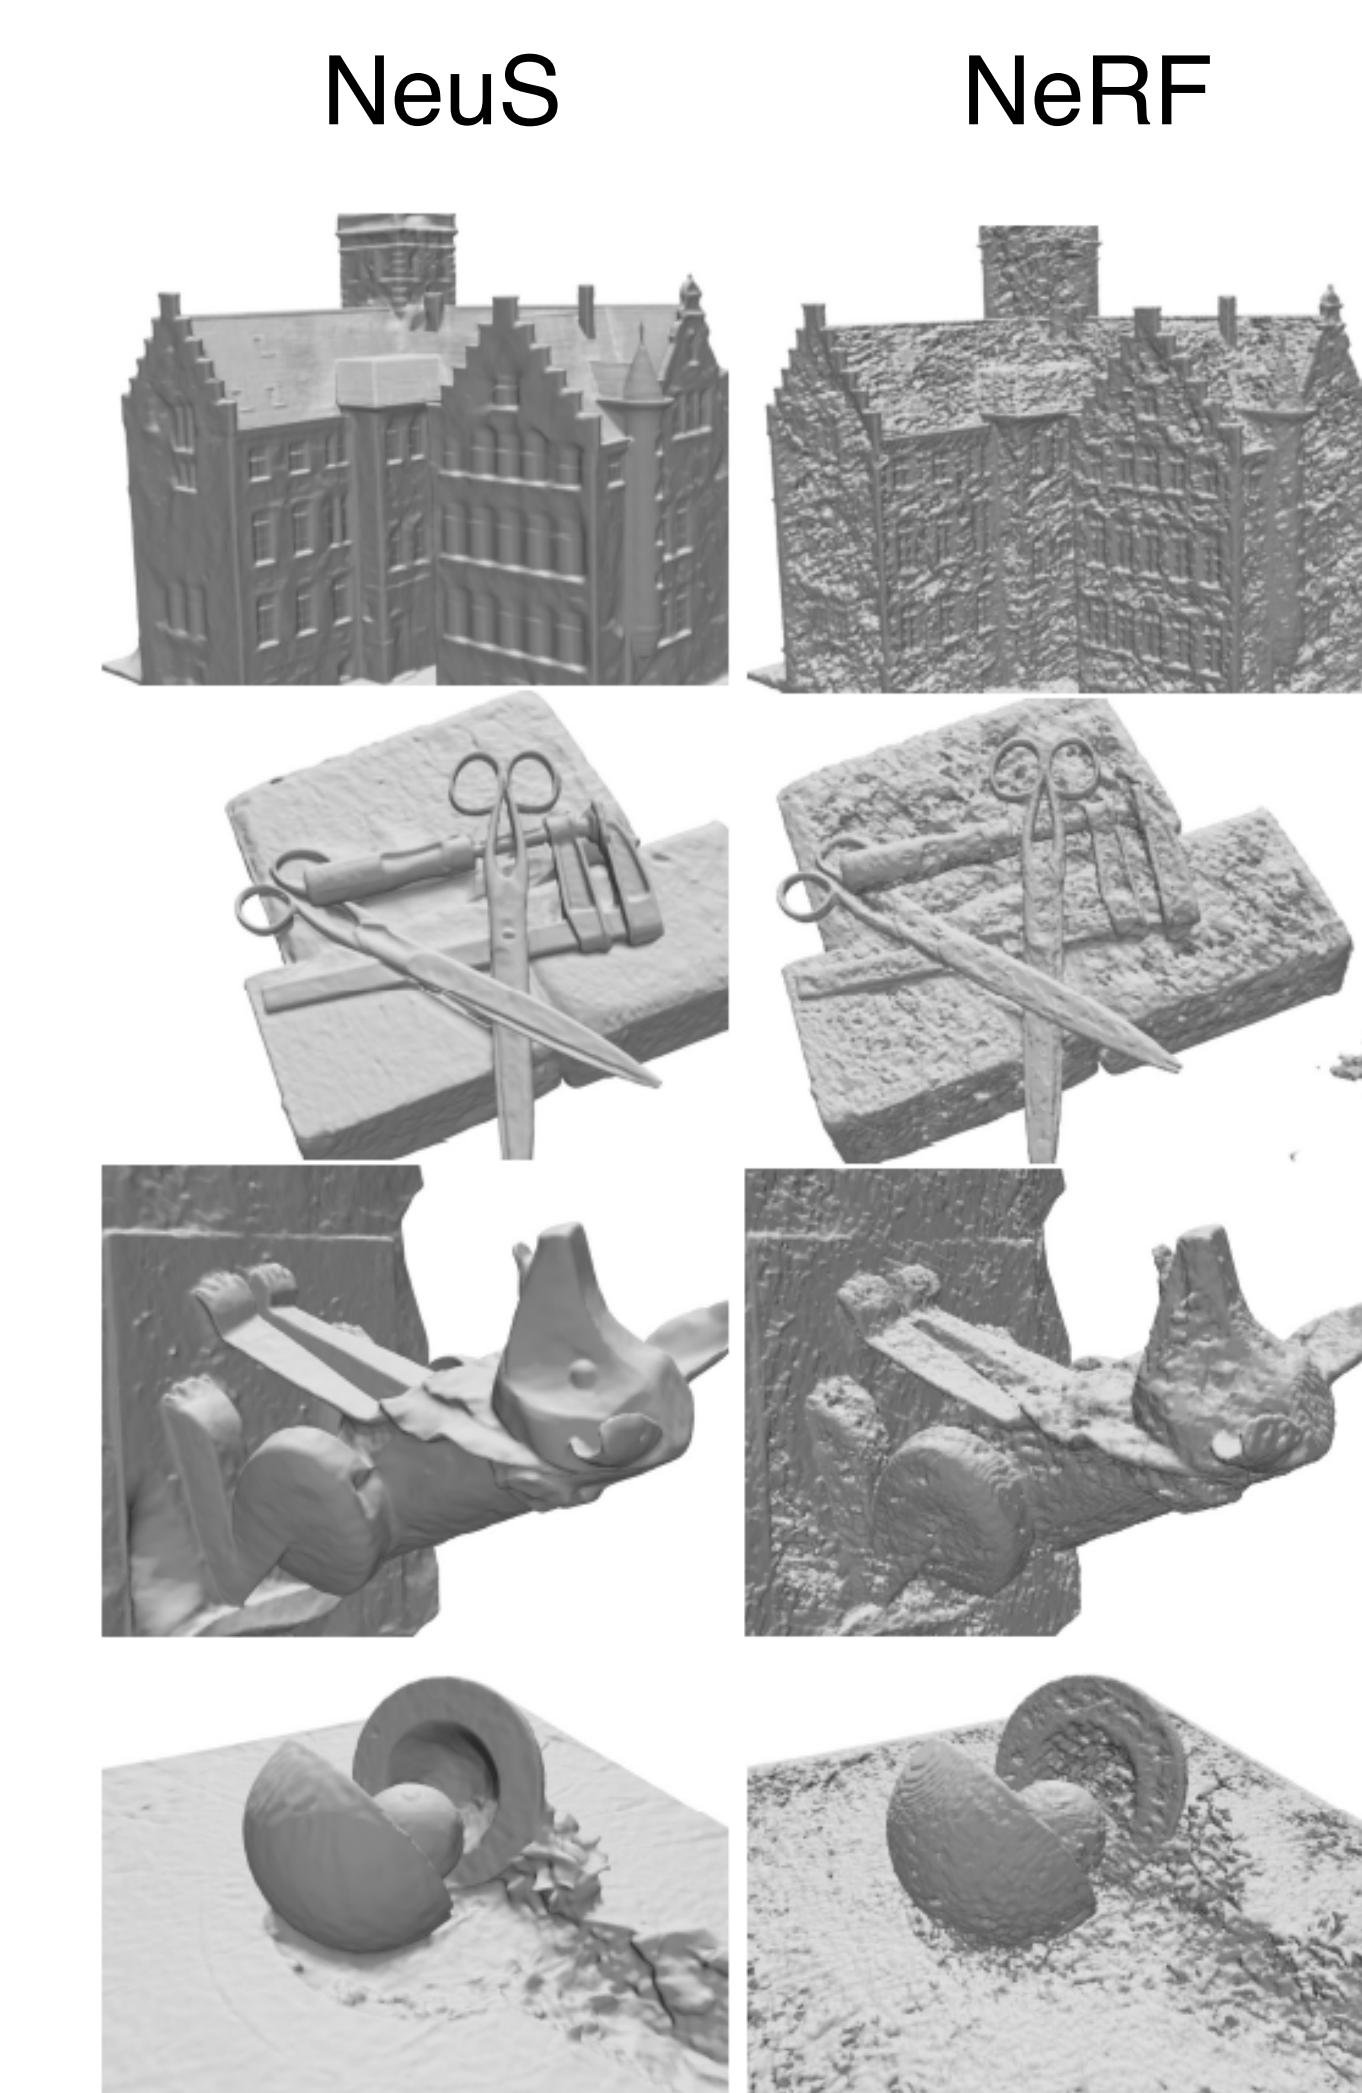
VoISDF

[Yariv et al., NeurIPS 2021]



NeuS

[Wang et al., NeurIPS 2021]

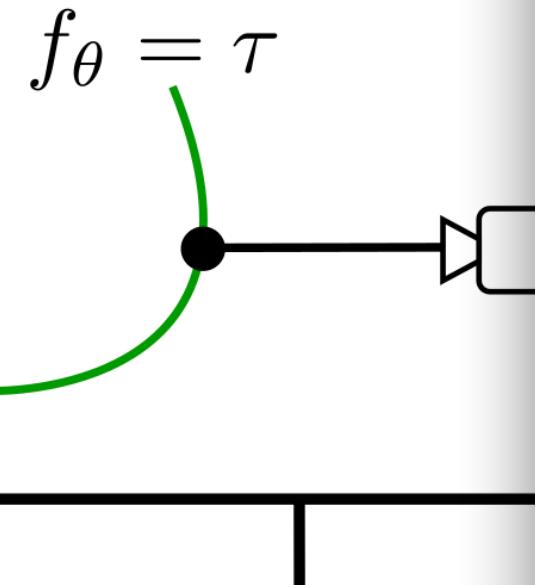


Hybrid Volumetric / Level Set Representations

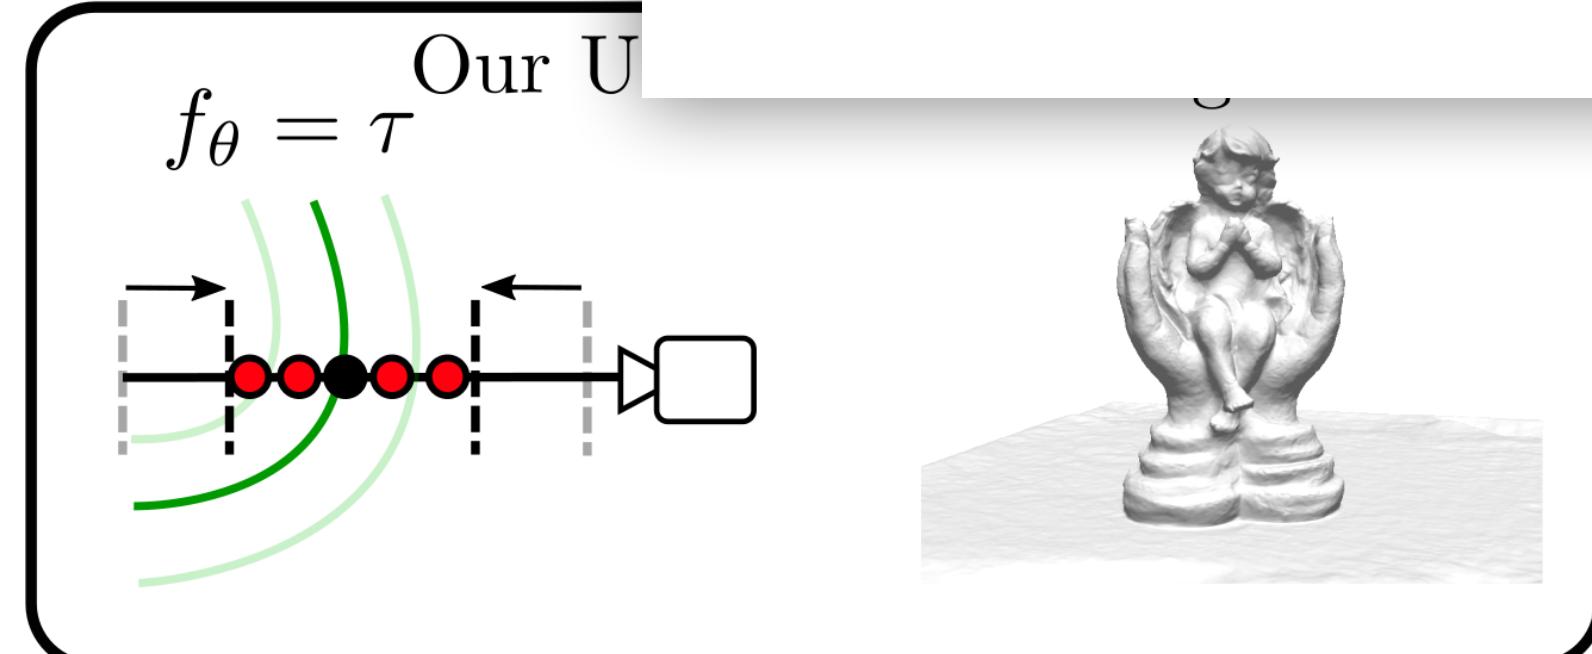
UNISURF

[Oechsle et al., ICCV 2021]

Surface Rendering



- Surface is well defined
- No noisy or cloudy geometry
- Good for solid objects
- Not good for complex scene phenomena (transparency)



VolSDF

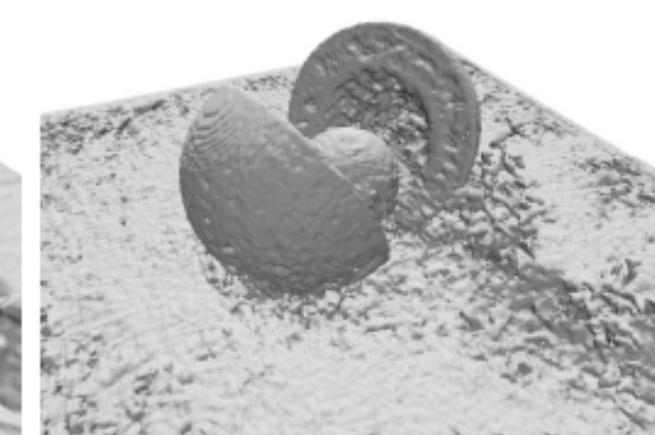
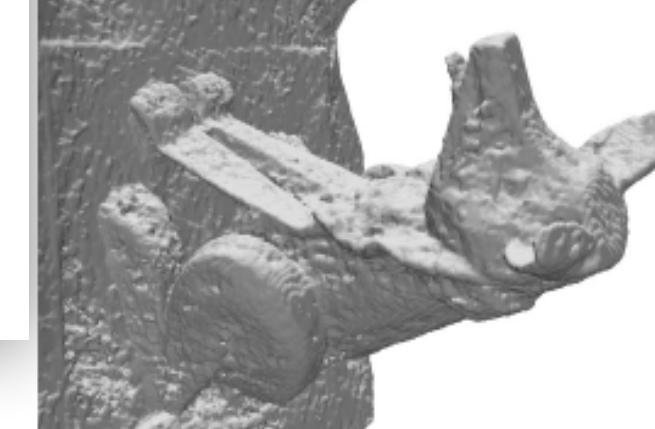
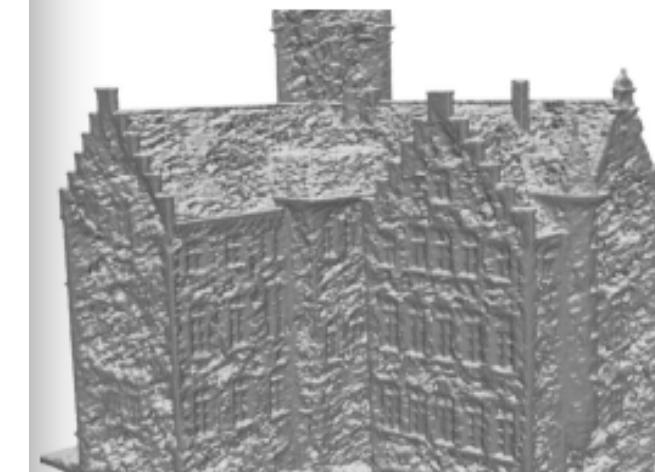
[Yariv et al., NeurIPS 2021]



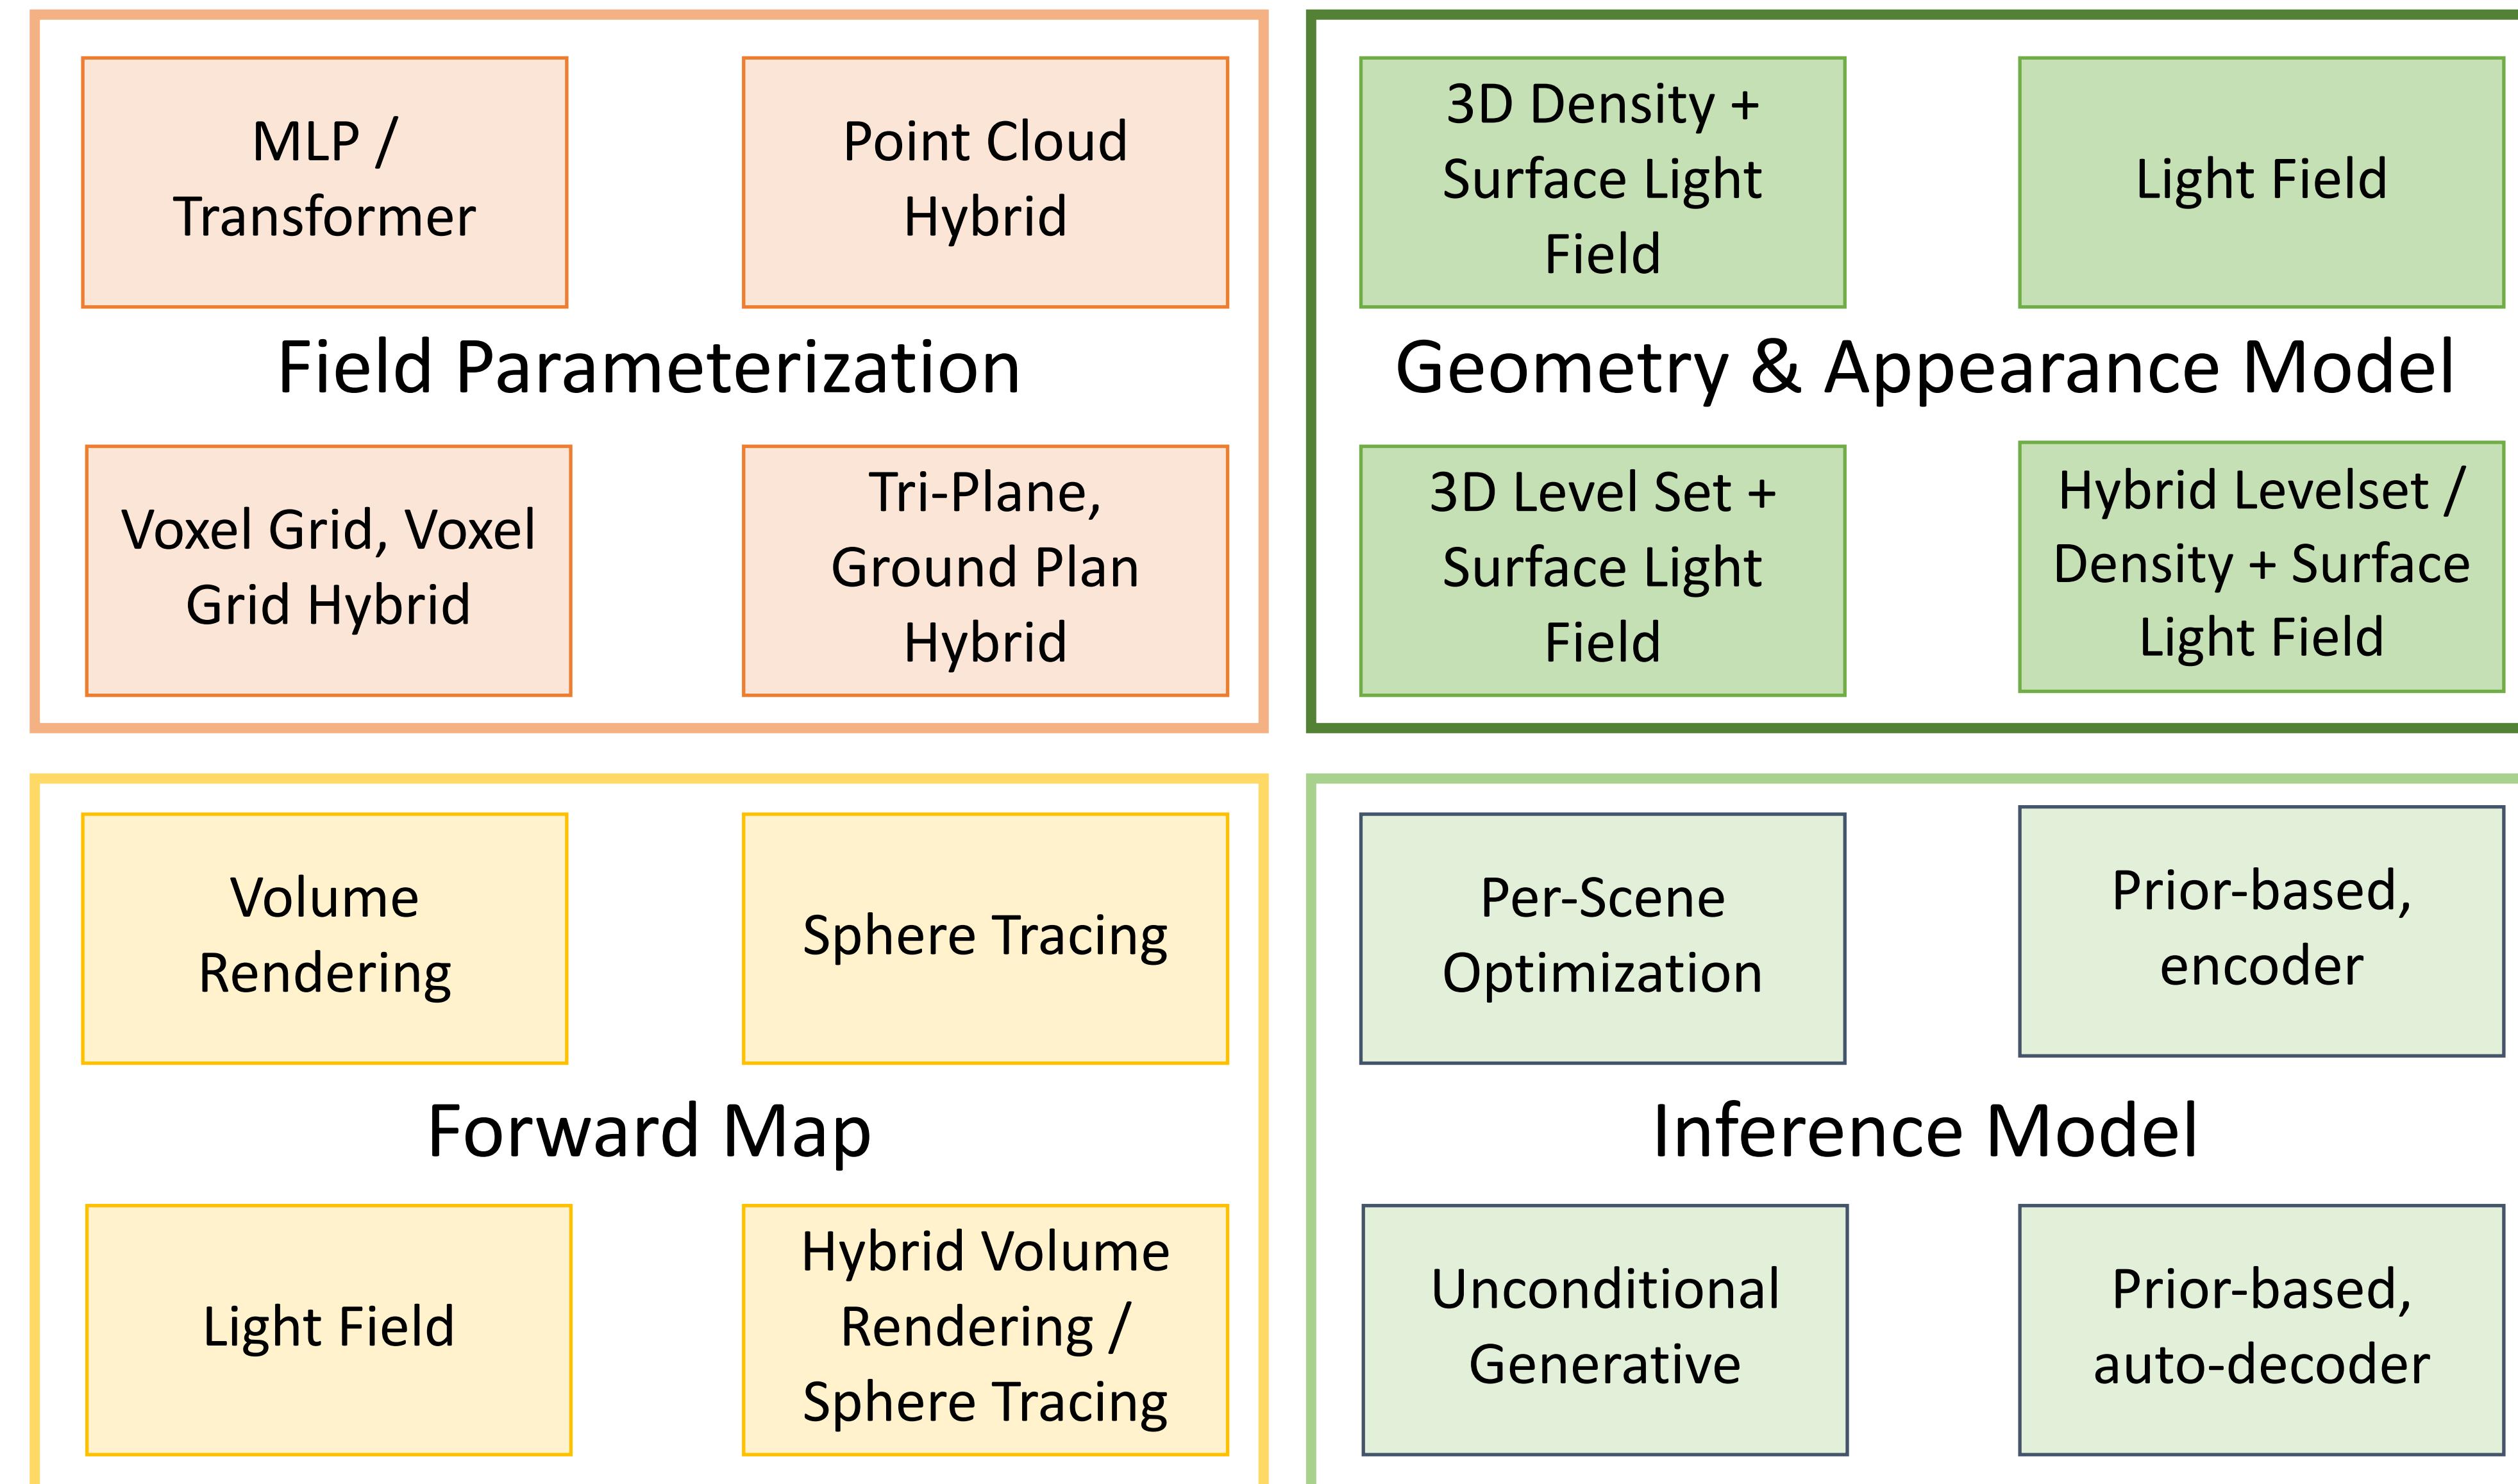
NeuS

[Wang et al., NeurIPS 2021]

NeRF



3D Inverse Graphics Taxonomy



Neural Fields in Visual Computing and Beyond

Yiheng Xie^{1,2} Towaki Takikawa^{3,4} Shunsuke Saito⁵ Or Litany⁴ Shiqin Yan¹ Numair Khan¹ Federico Tombari^{6,7}
 James Tompkin¹ Vincent Sitzmann⁸ Srinath Sridhar^{1†}

¹Brown University ²Unity Technologies ³University of Toronto ⁴NVIDIA ⁵Meta Reality Labs Research ⁶Google ⁷Technical University of Munich
⁸Massachusetts Institute of Technology [†]Equal advising

<https://neuralfields.cs.brown.edu/>

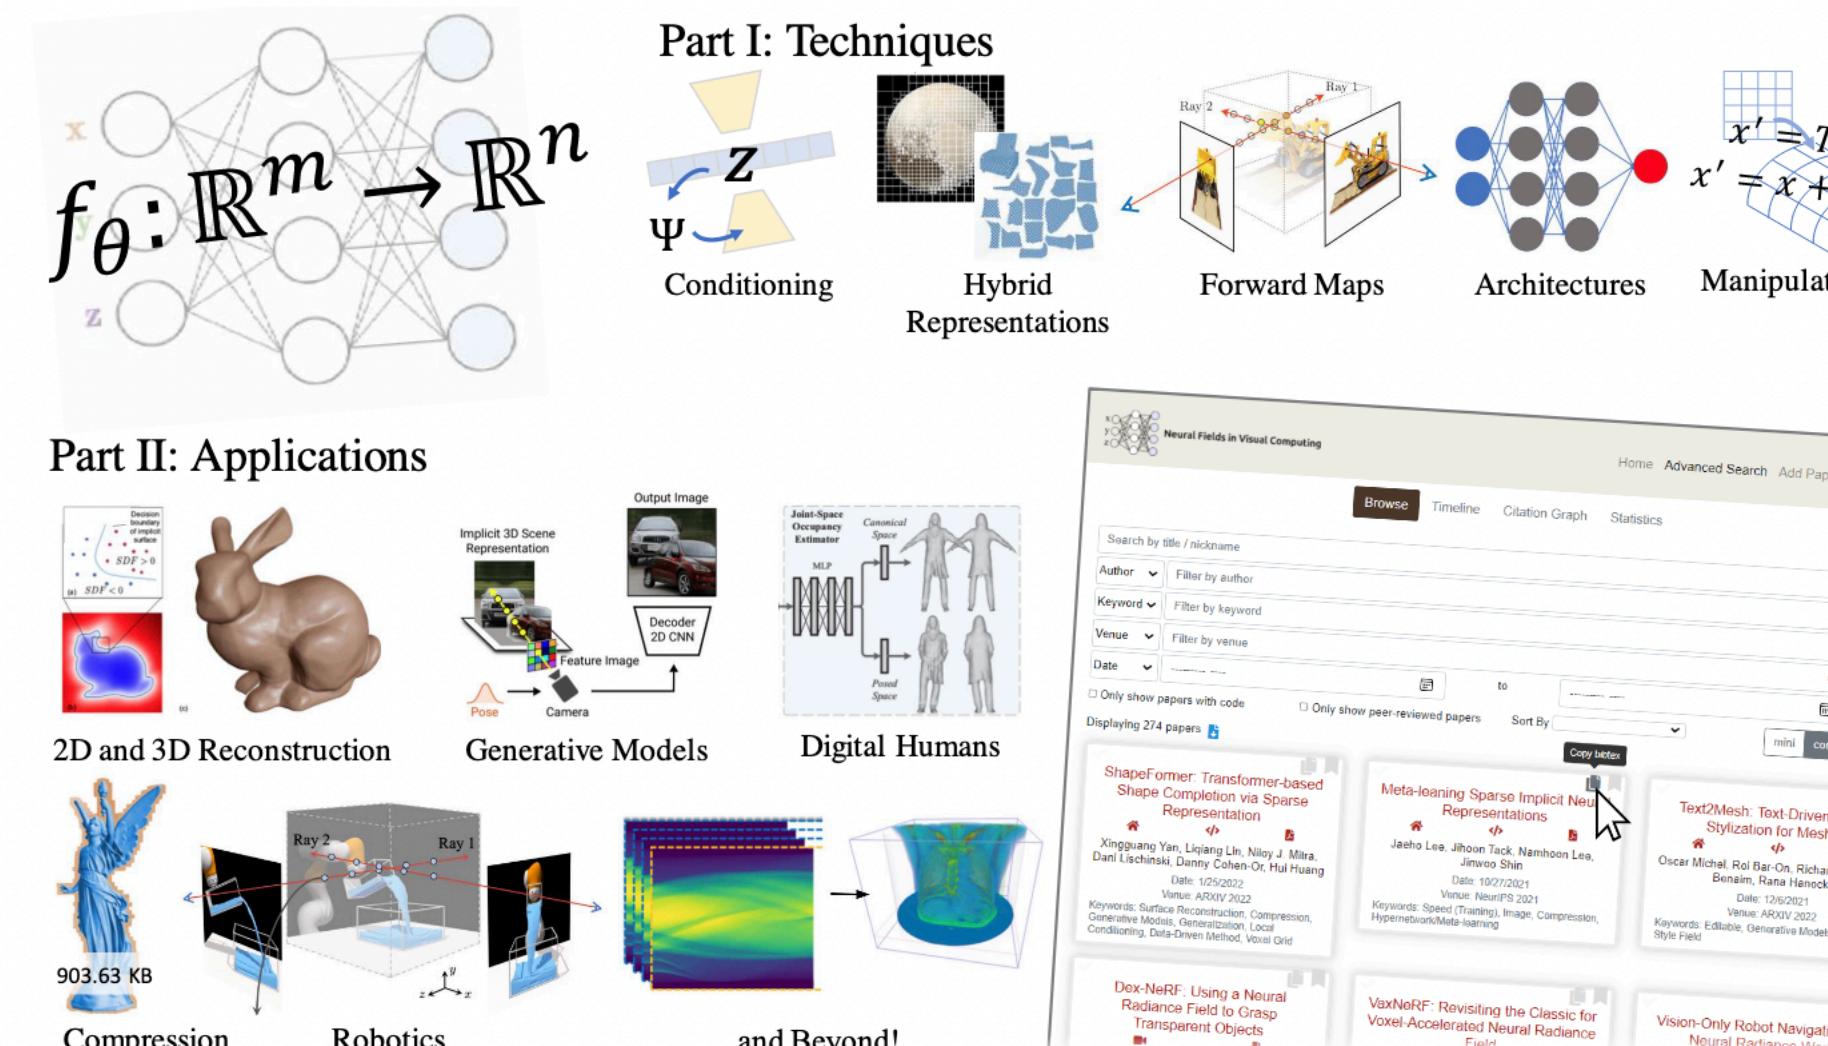


Figure 1: **Contribution of this report.** Following a survey of over 250 papers, we provide a review of (**Part I**) techniques in neural fields such as prior learning and conditioning, representations, forward maps, architectures, and manipulation, and of (**Part II**) applications in visual computing including 2D image processing, 3D scene reconstruction, generative modeling, digital humans, compression, robotics, and beyond. This report is complemented by a [community-driven website](https://neuralfields.cs.brown.edu/) with search, filtering, bibliographic, and visualization features.

How far does this get us?

$$\sigma = \sigma(\mathbf{x})$$

$$\mathbf{c} = \mathbf{c}(\mathbf{x})$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \cancel{\omega_o}, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) L_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \lambda)$$

No “true” reflections, no diffuse inter-reflections, no re-lighting, no materials, no...

How far does this get us?

$$\begin{aligned}\sigma &= \sigma(\mathbf{x}) \\ \mathbf{c} &= \mathbf{c}(\mathbf{x}, \omega_o)\end{aligned}$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) I_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda)$$

Half-way fix: allow direction-dependent color, but keep geometry direction-independent.

How far does this get us?

$$\begin{aligned}\sigma &= \sigma(\mathbf{x}) \\ \mathbf{c} &= \mathbf{c}(\mathbf{x}, \omega_o)\end{aligned}$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda) + \int_{\Omega} f(\mathbf{x}, \omega_i, \omega_o, \lambda) I_i(\mathbf{x}, \omega_i, \lambda) (\omega_i \cdot \mathbf{n}) d\omega_i$$

$$L_o(\mathbf{x}, \omega_o, \lambda) = L_e(\mathbf{x}, \omega_o, \lambda)$$

Can now approximate everything, but materials etc. are not enforced.

Questions?

Direction-dependent color



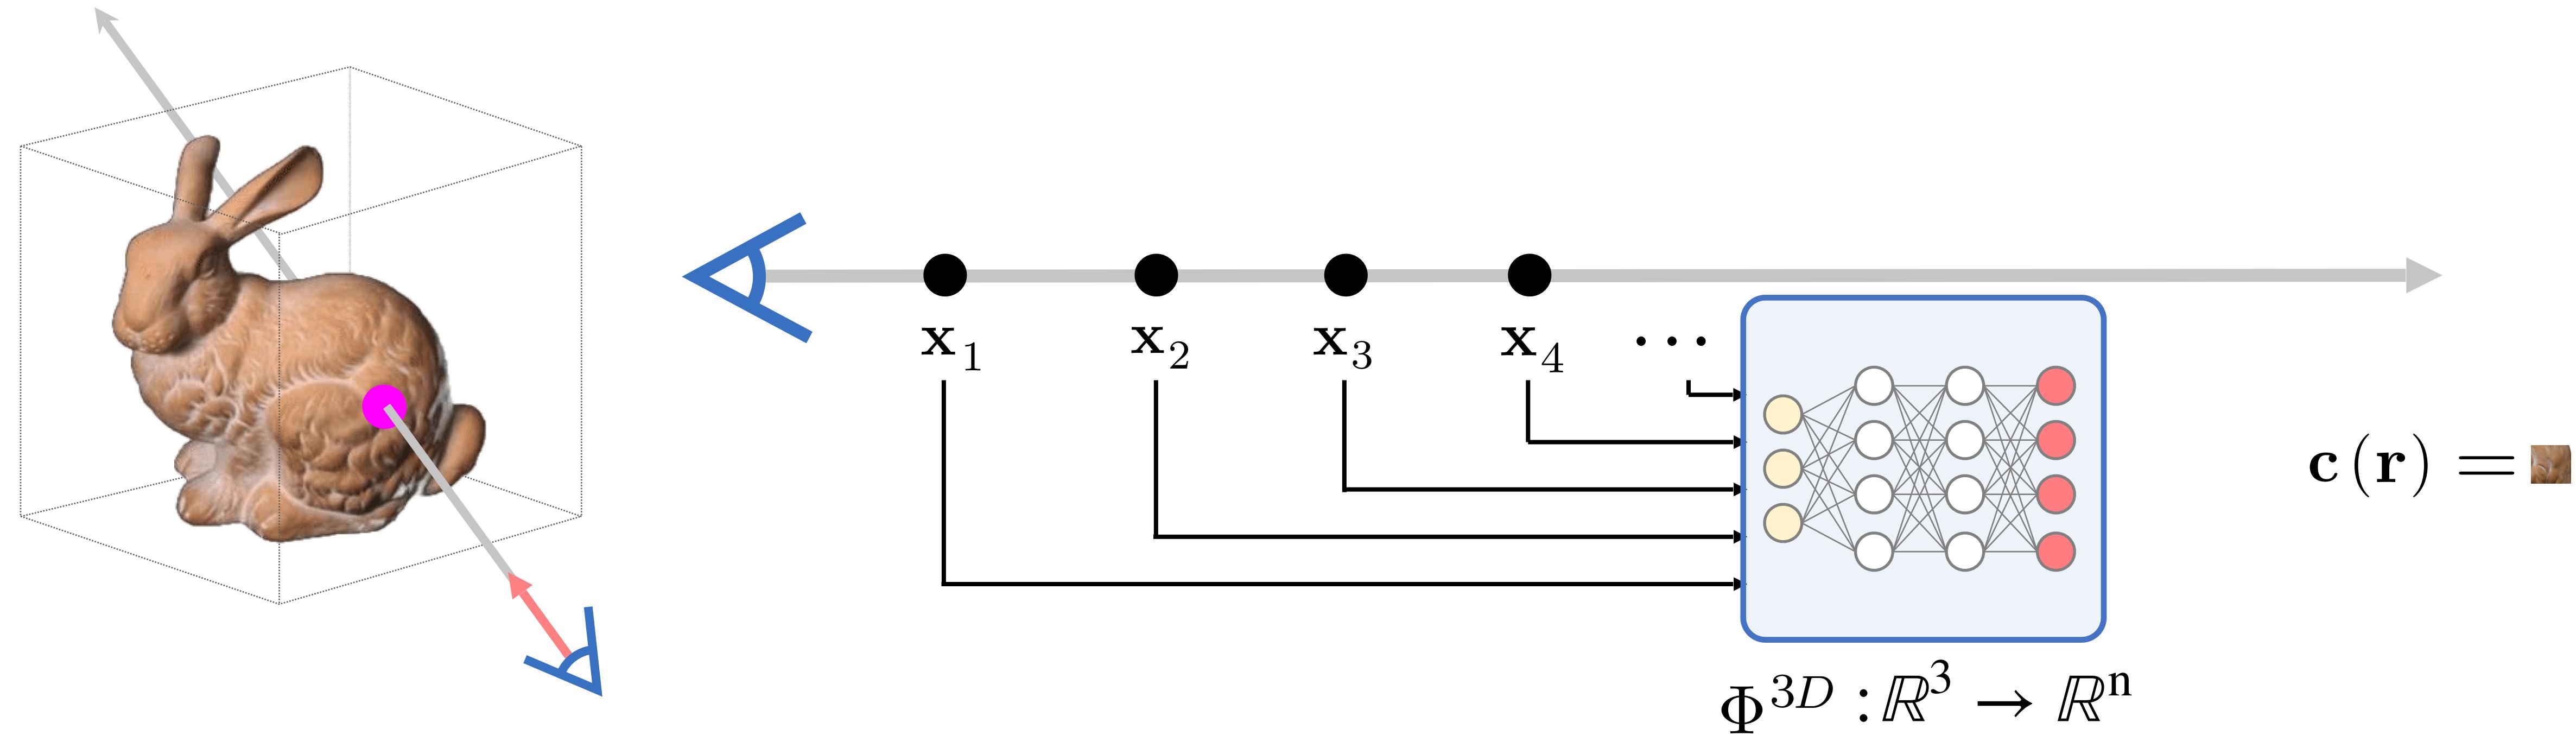
Direction-dependent color



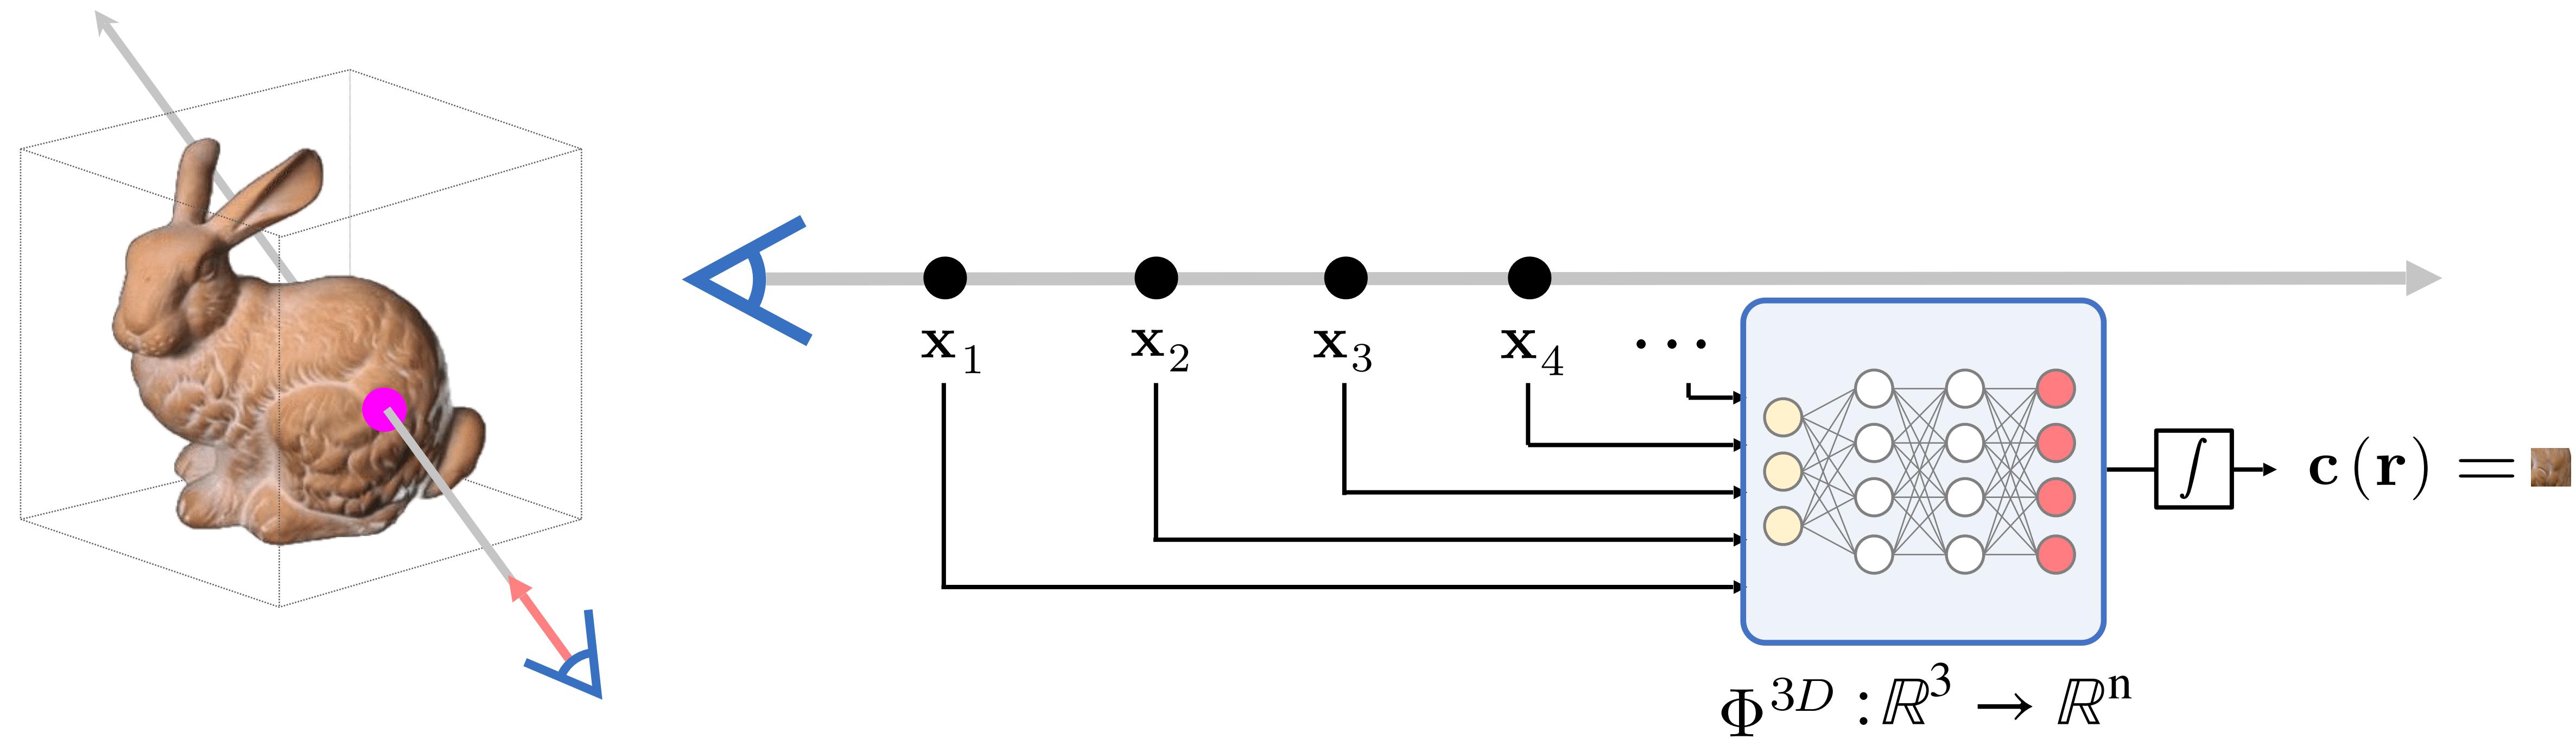
General structure of Neural Renderers for 3D-structured Representations

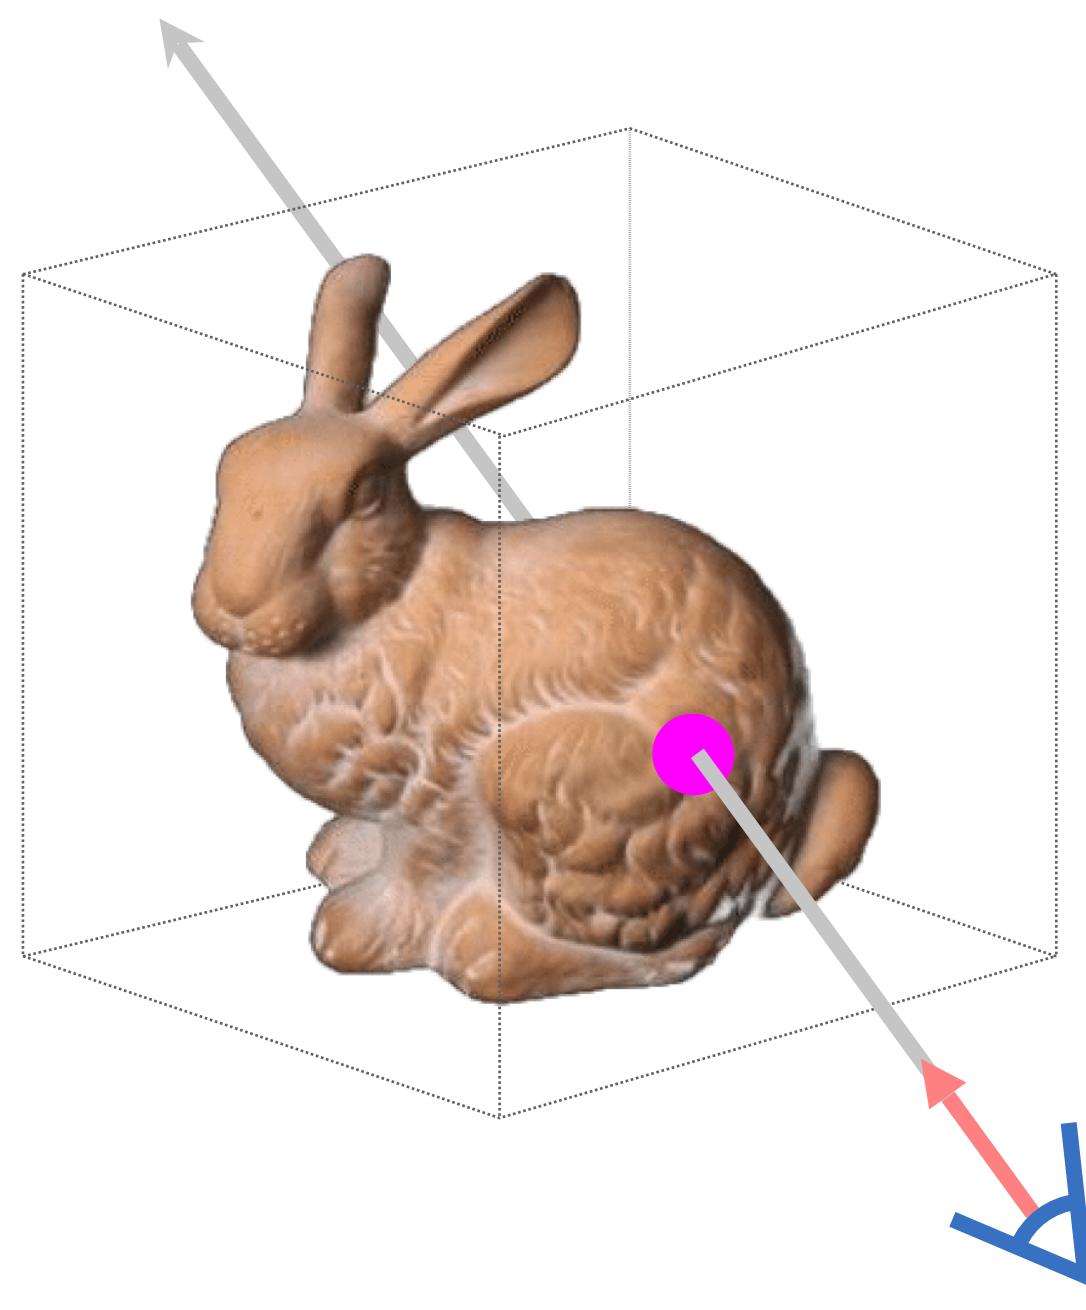


General structure of Neural Renderers for 3D-structured Representations



General structure of Neural Renderers for 3D-structured Representations





A

x_1

x_2

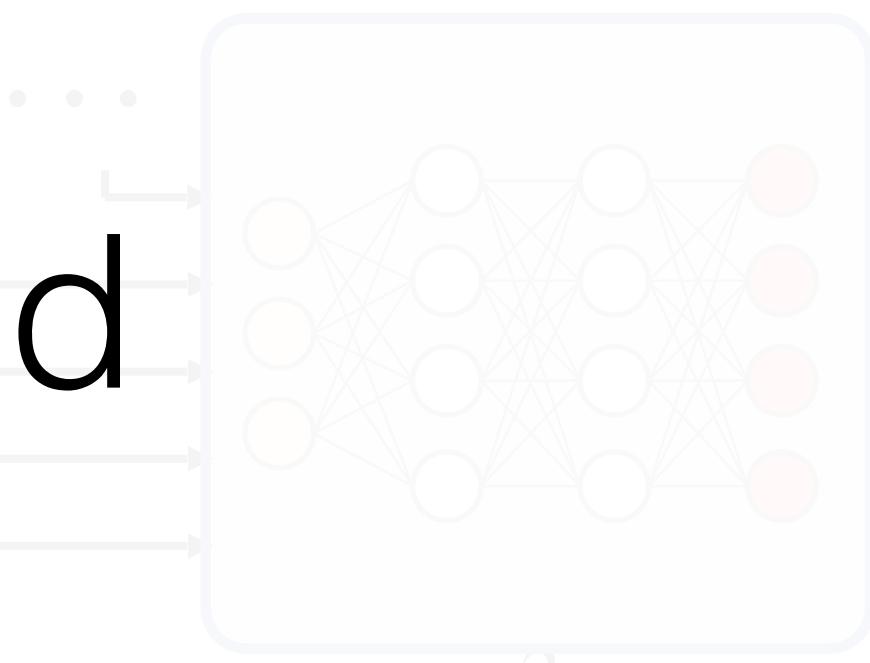
x_3

x_4

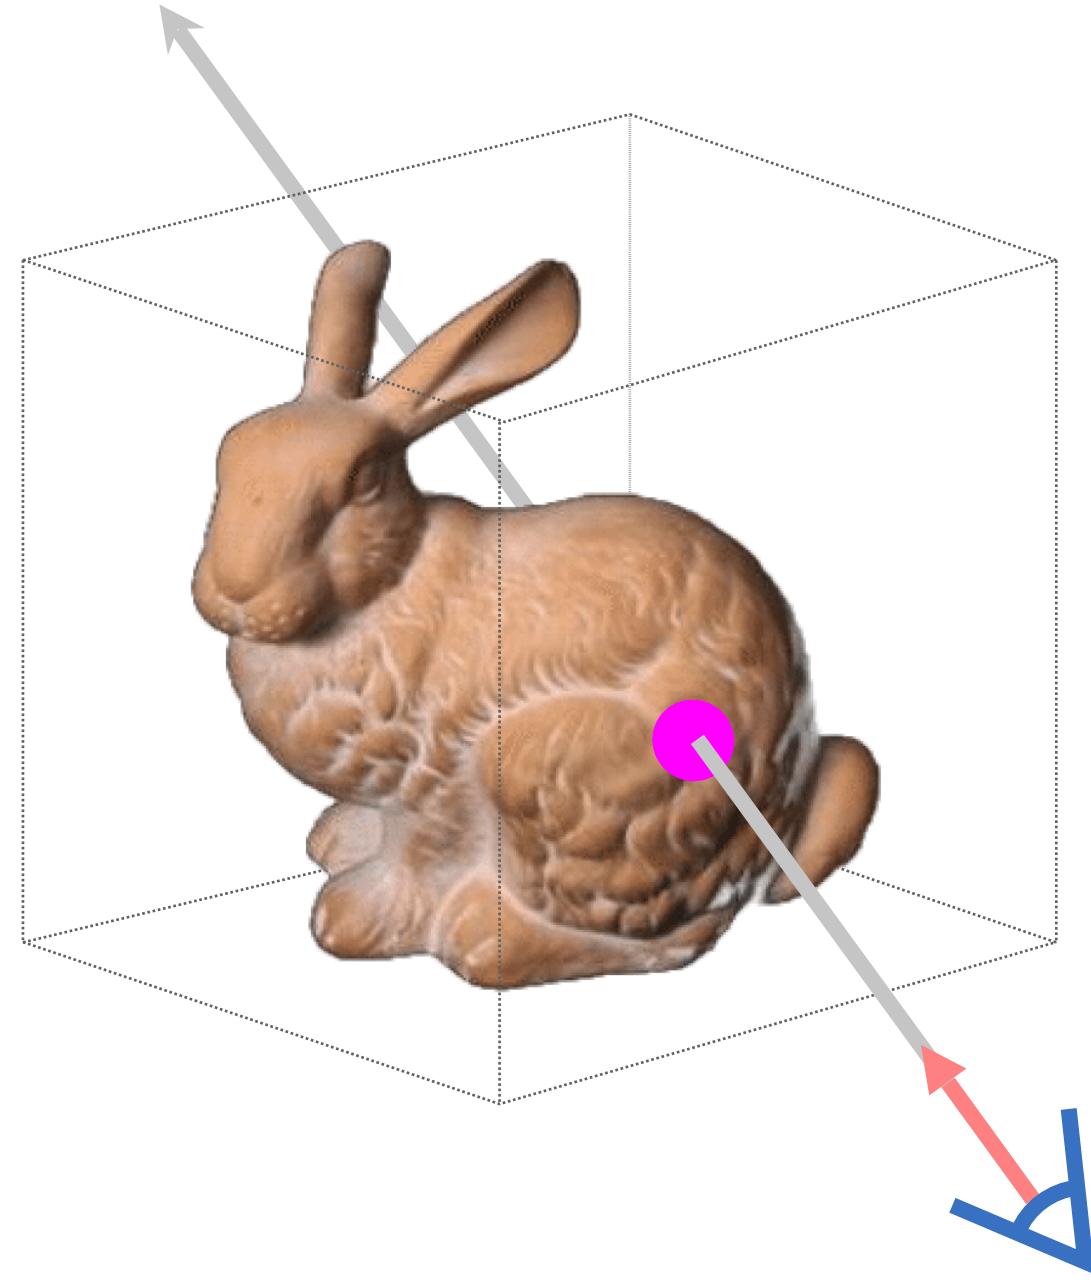
...

Light Field

$$\Phi^{3D} : \mathbb{R}^3 \rightarrow \mathbb{R}^n$$



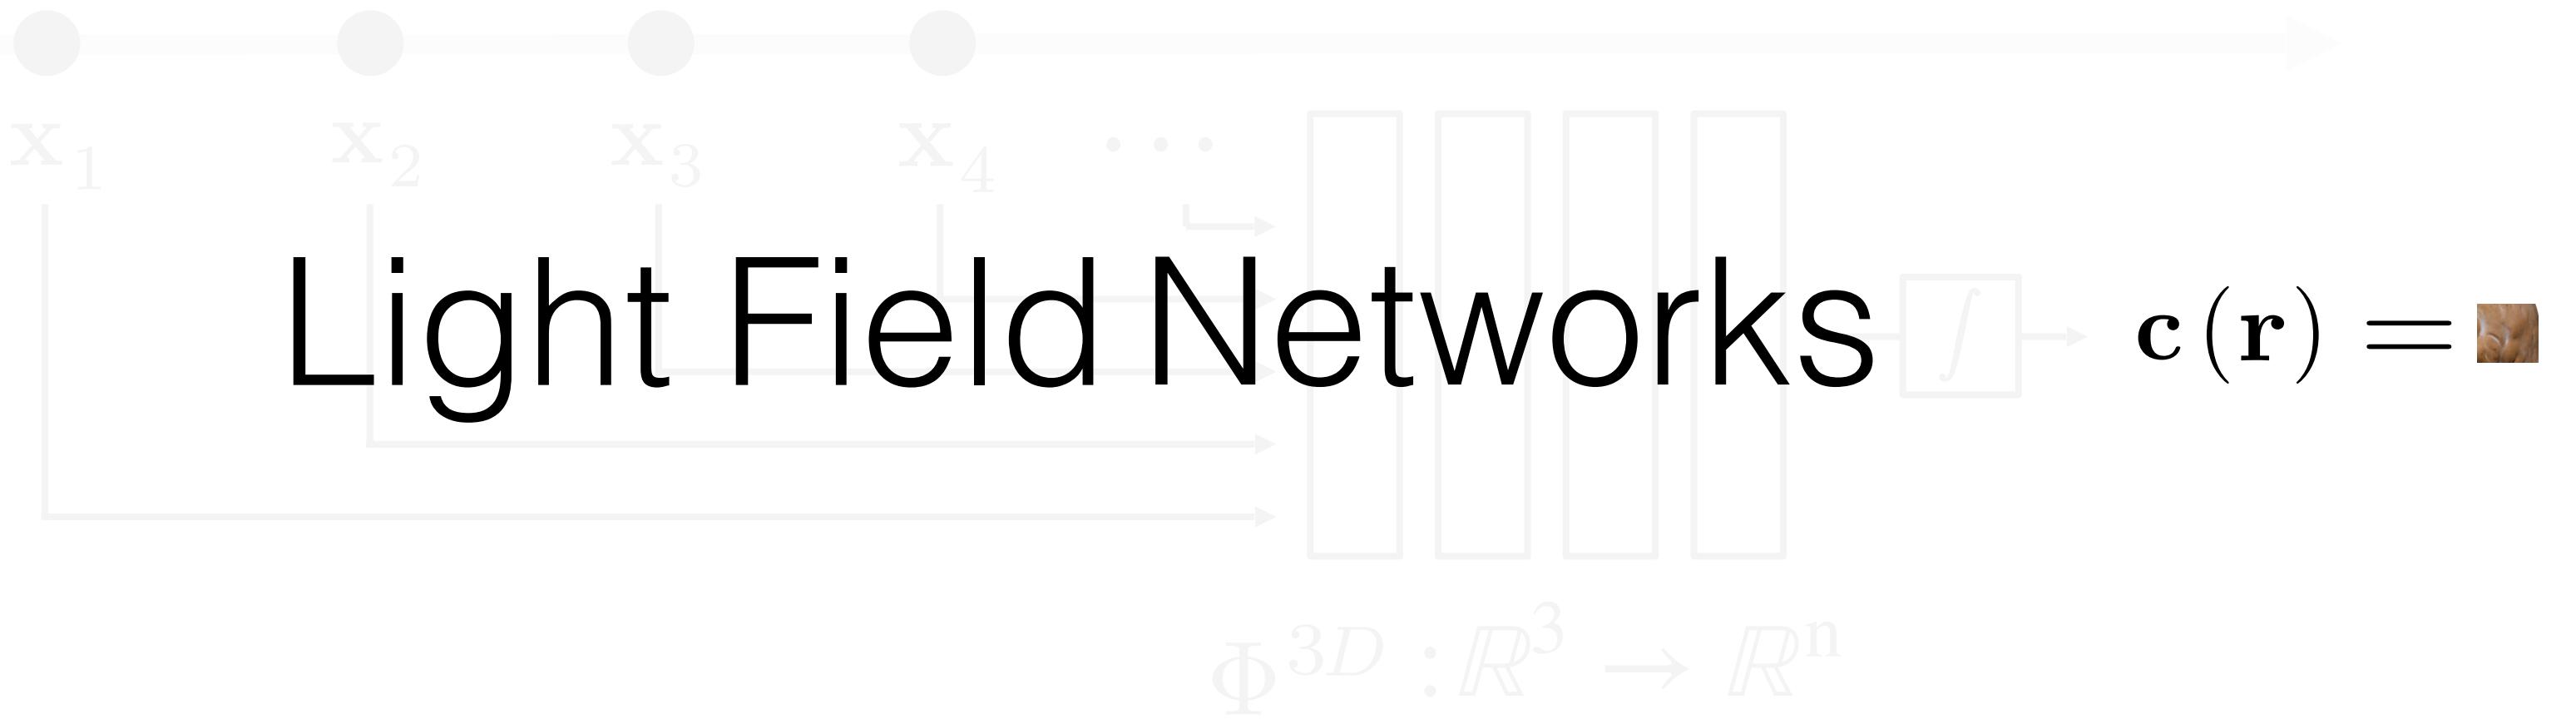
$$c(\mathbf{r}) =$$



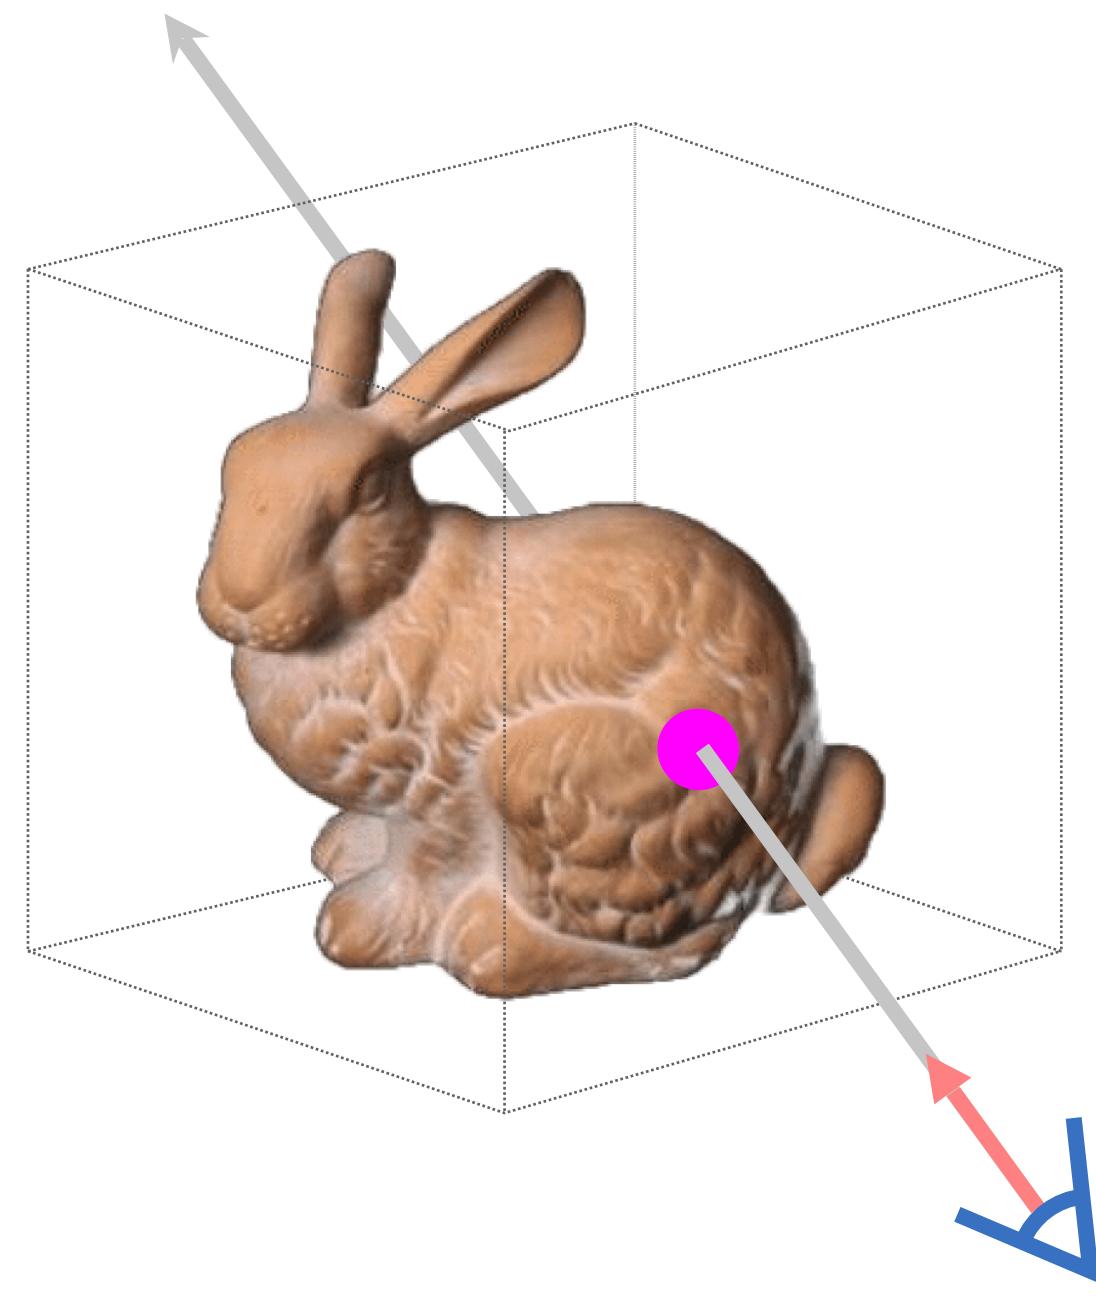
A

x_1

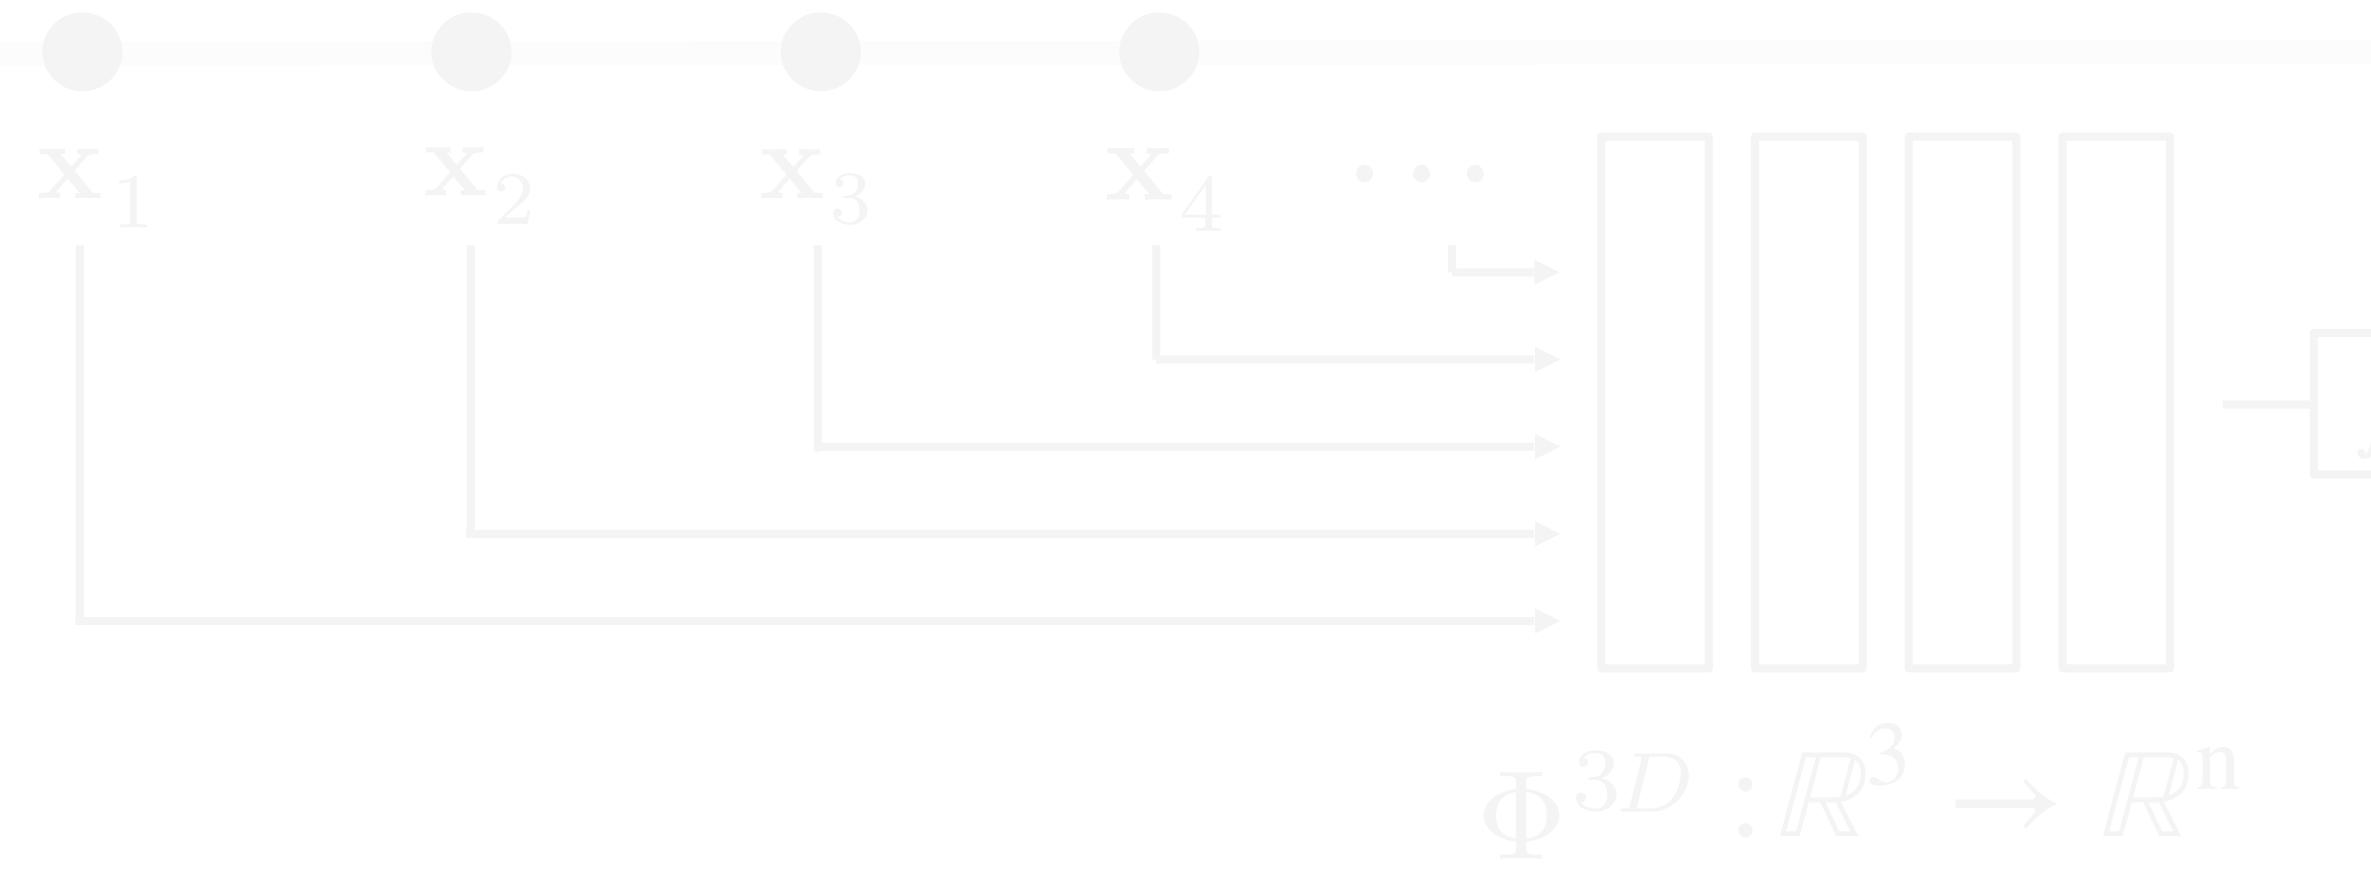
Light Field Networks



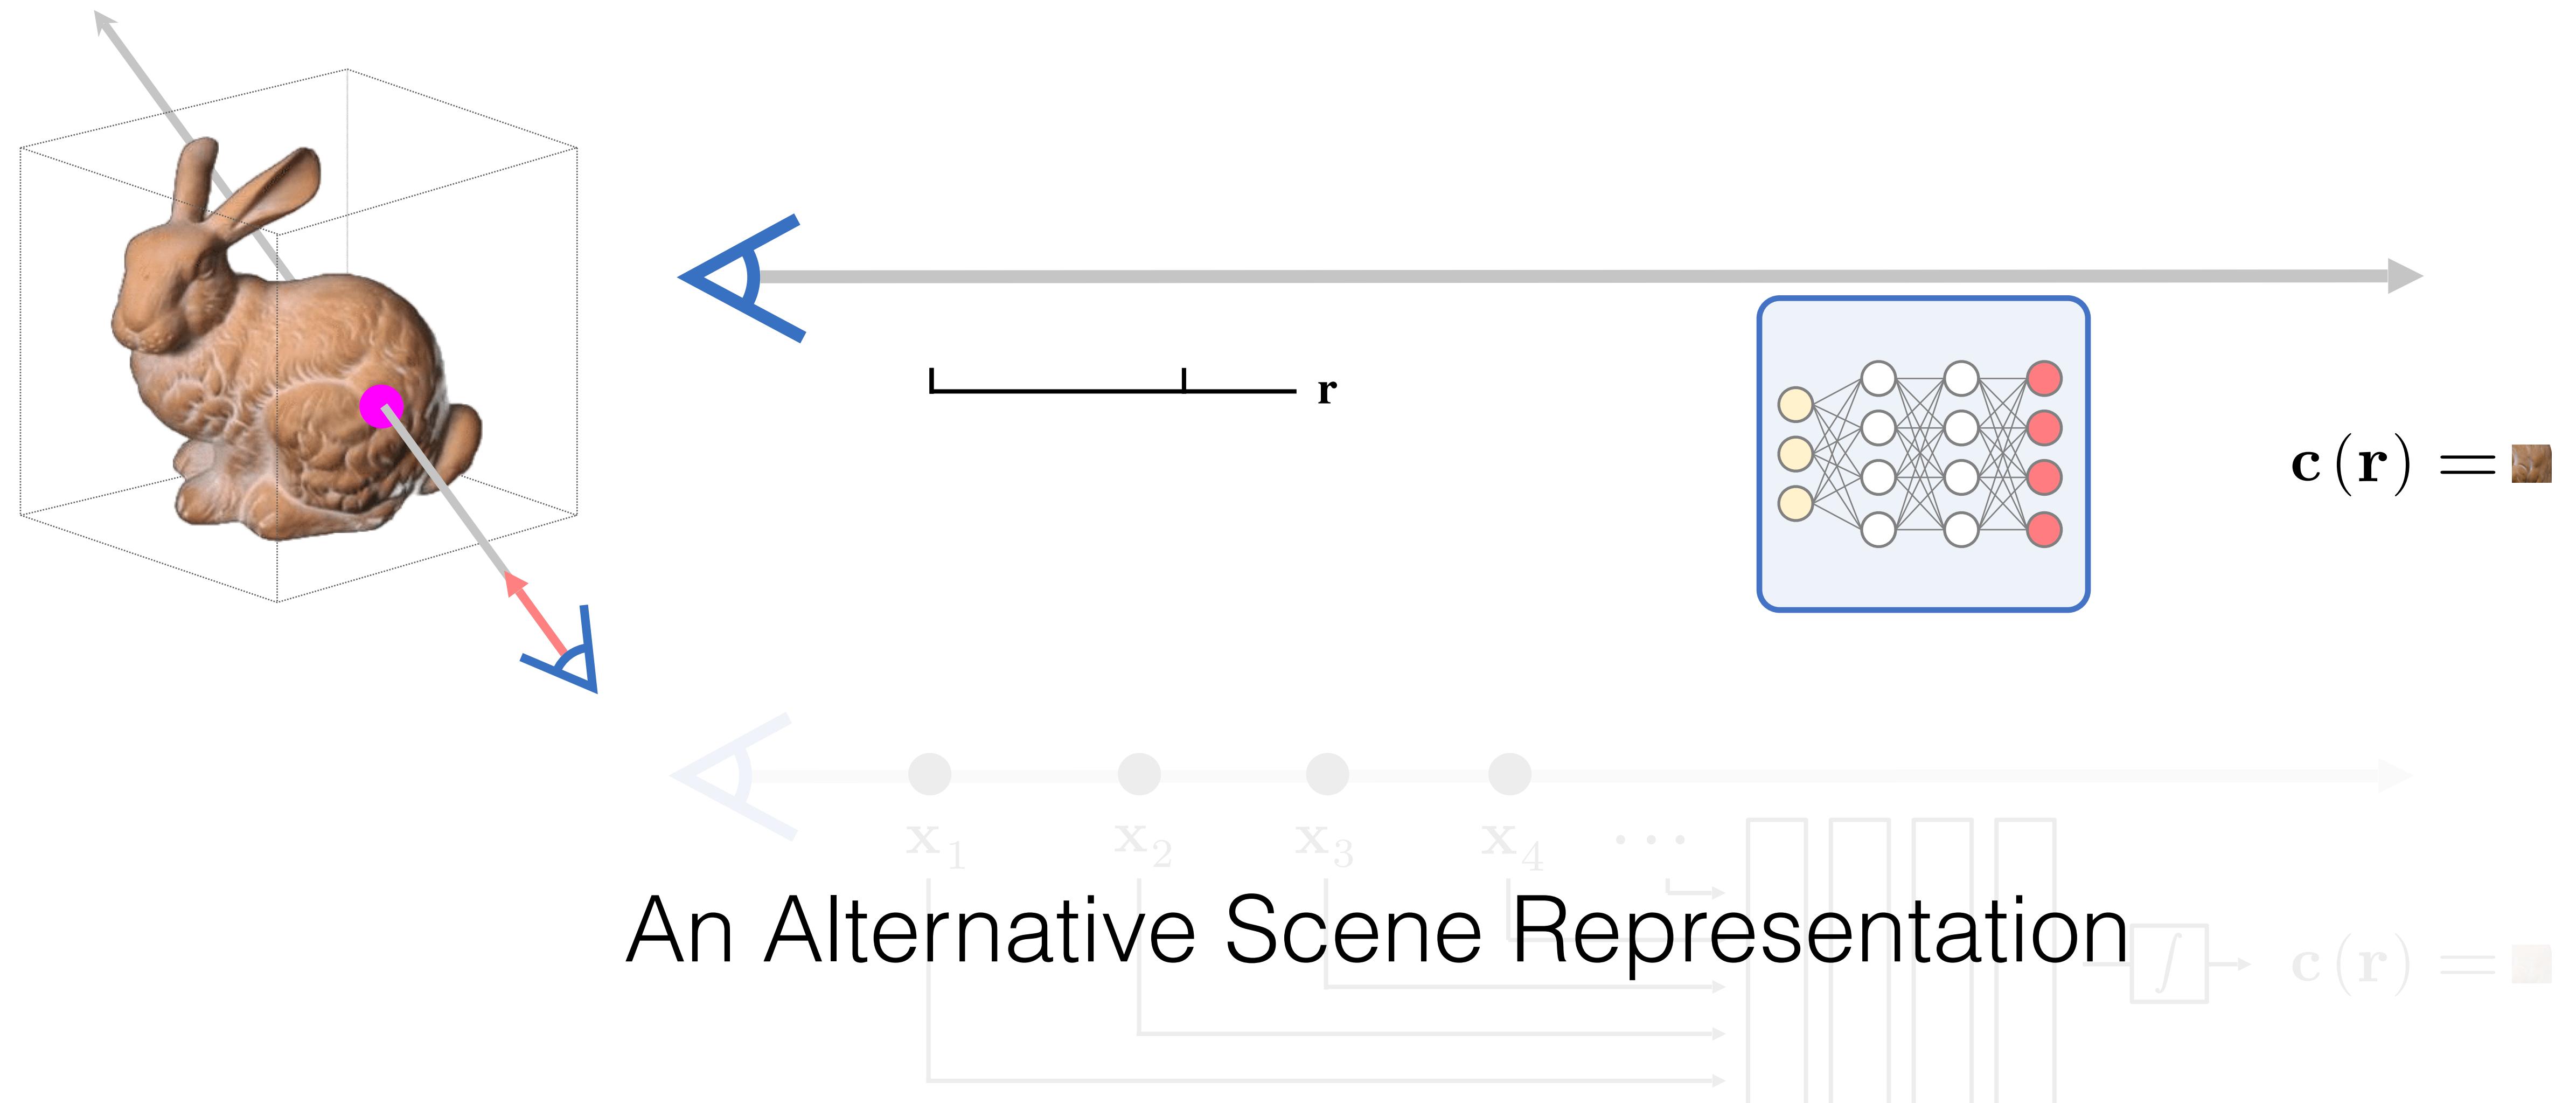
Light Field Networks



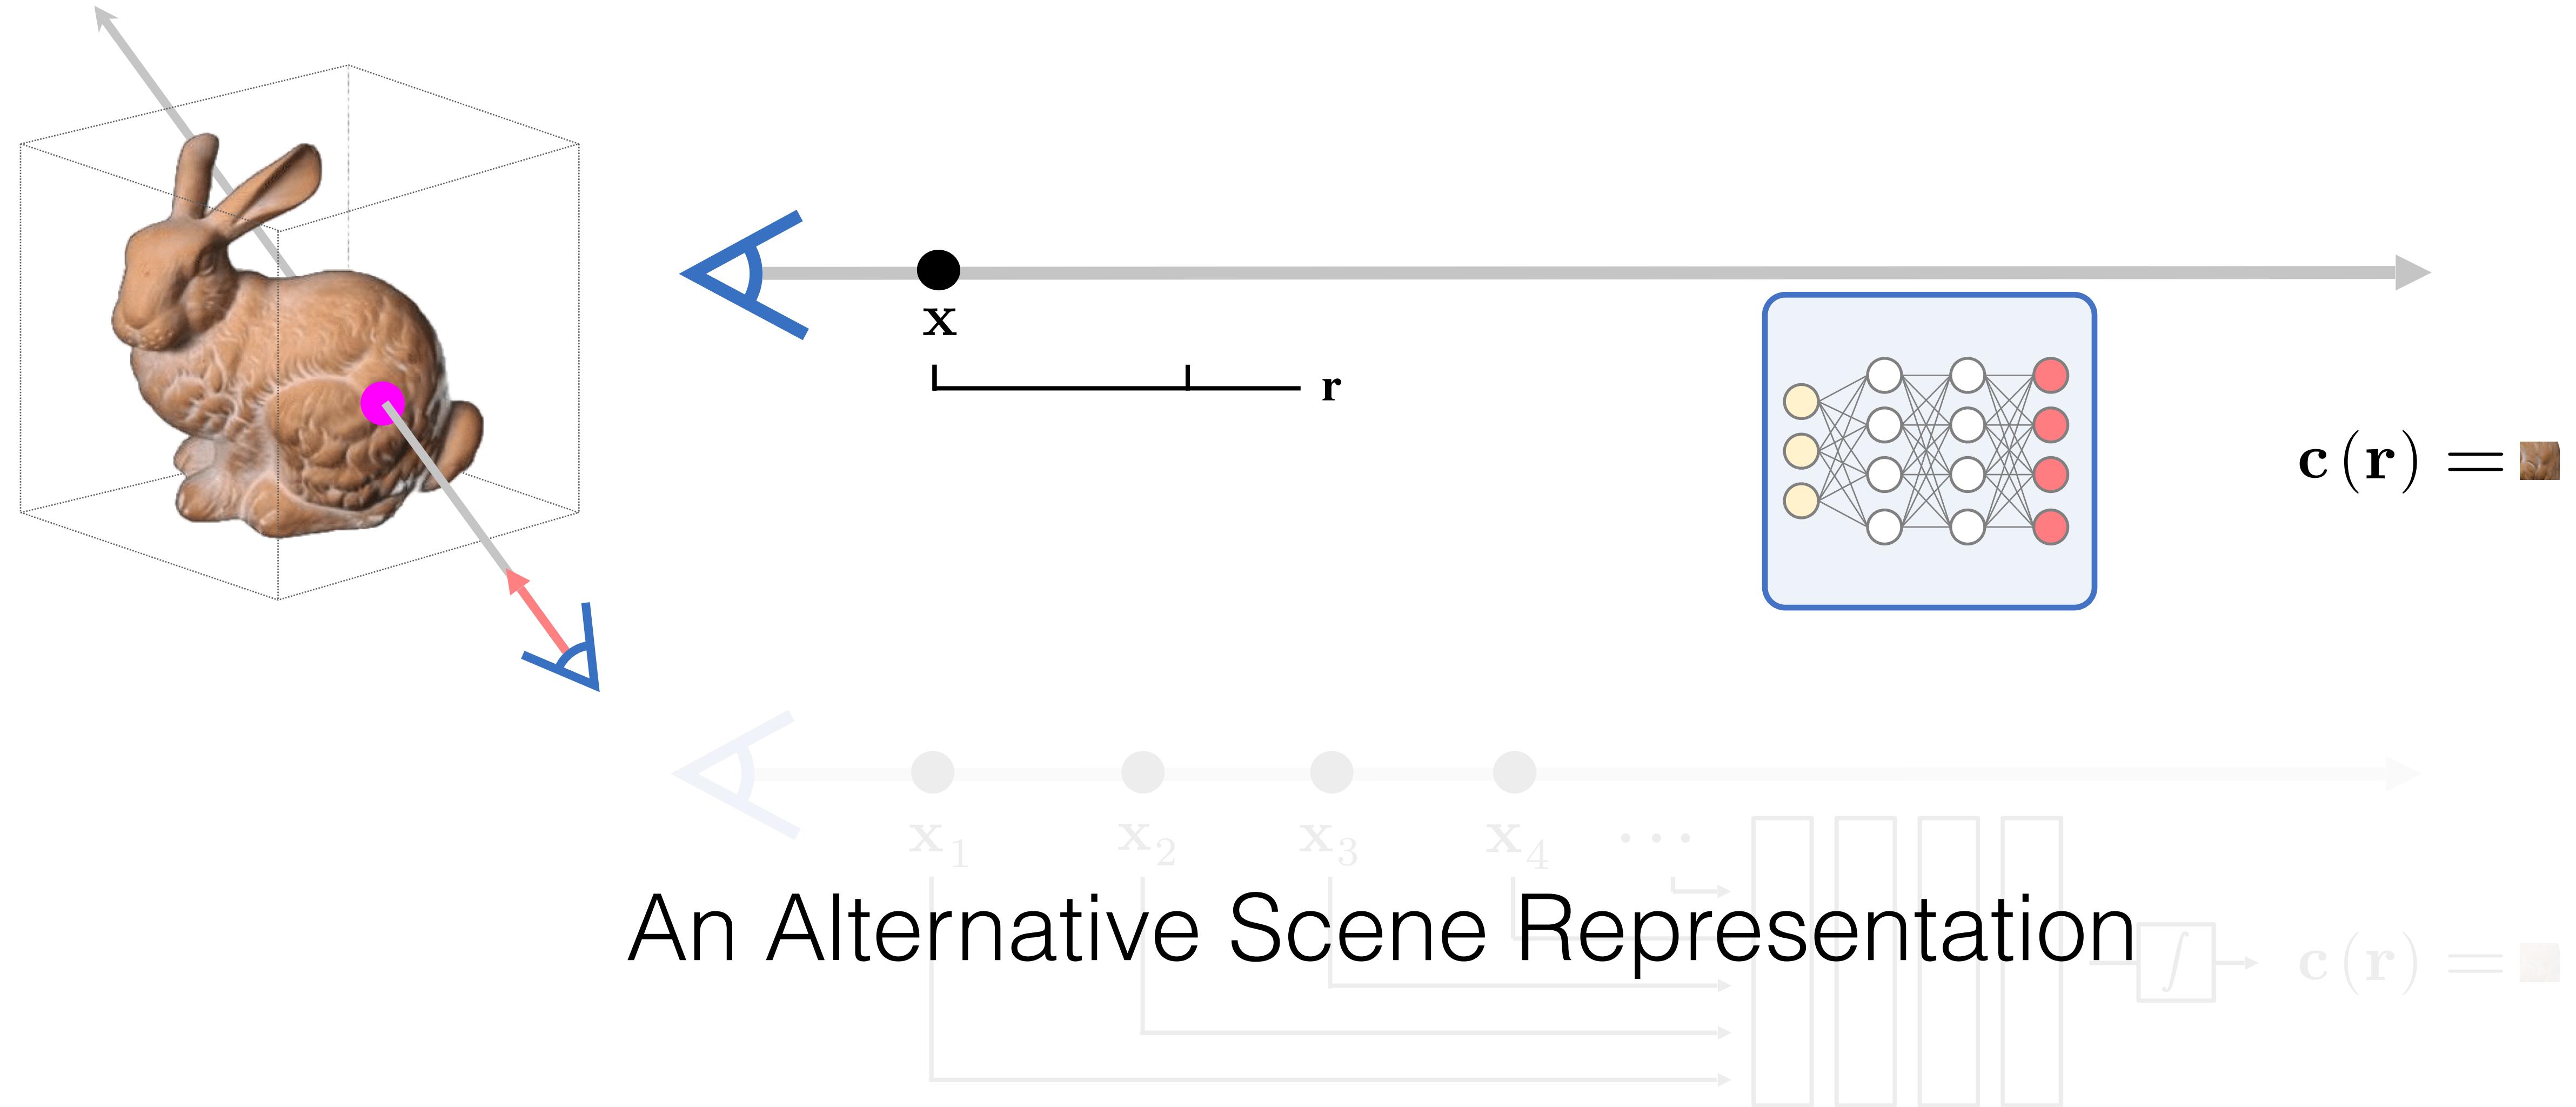
A



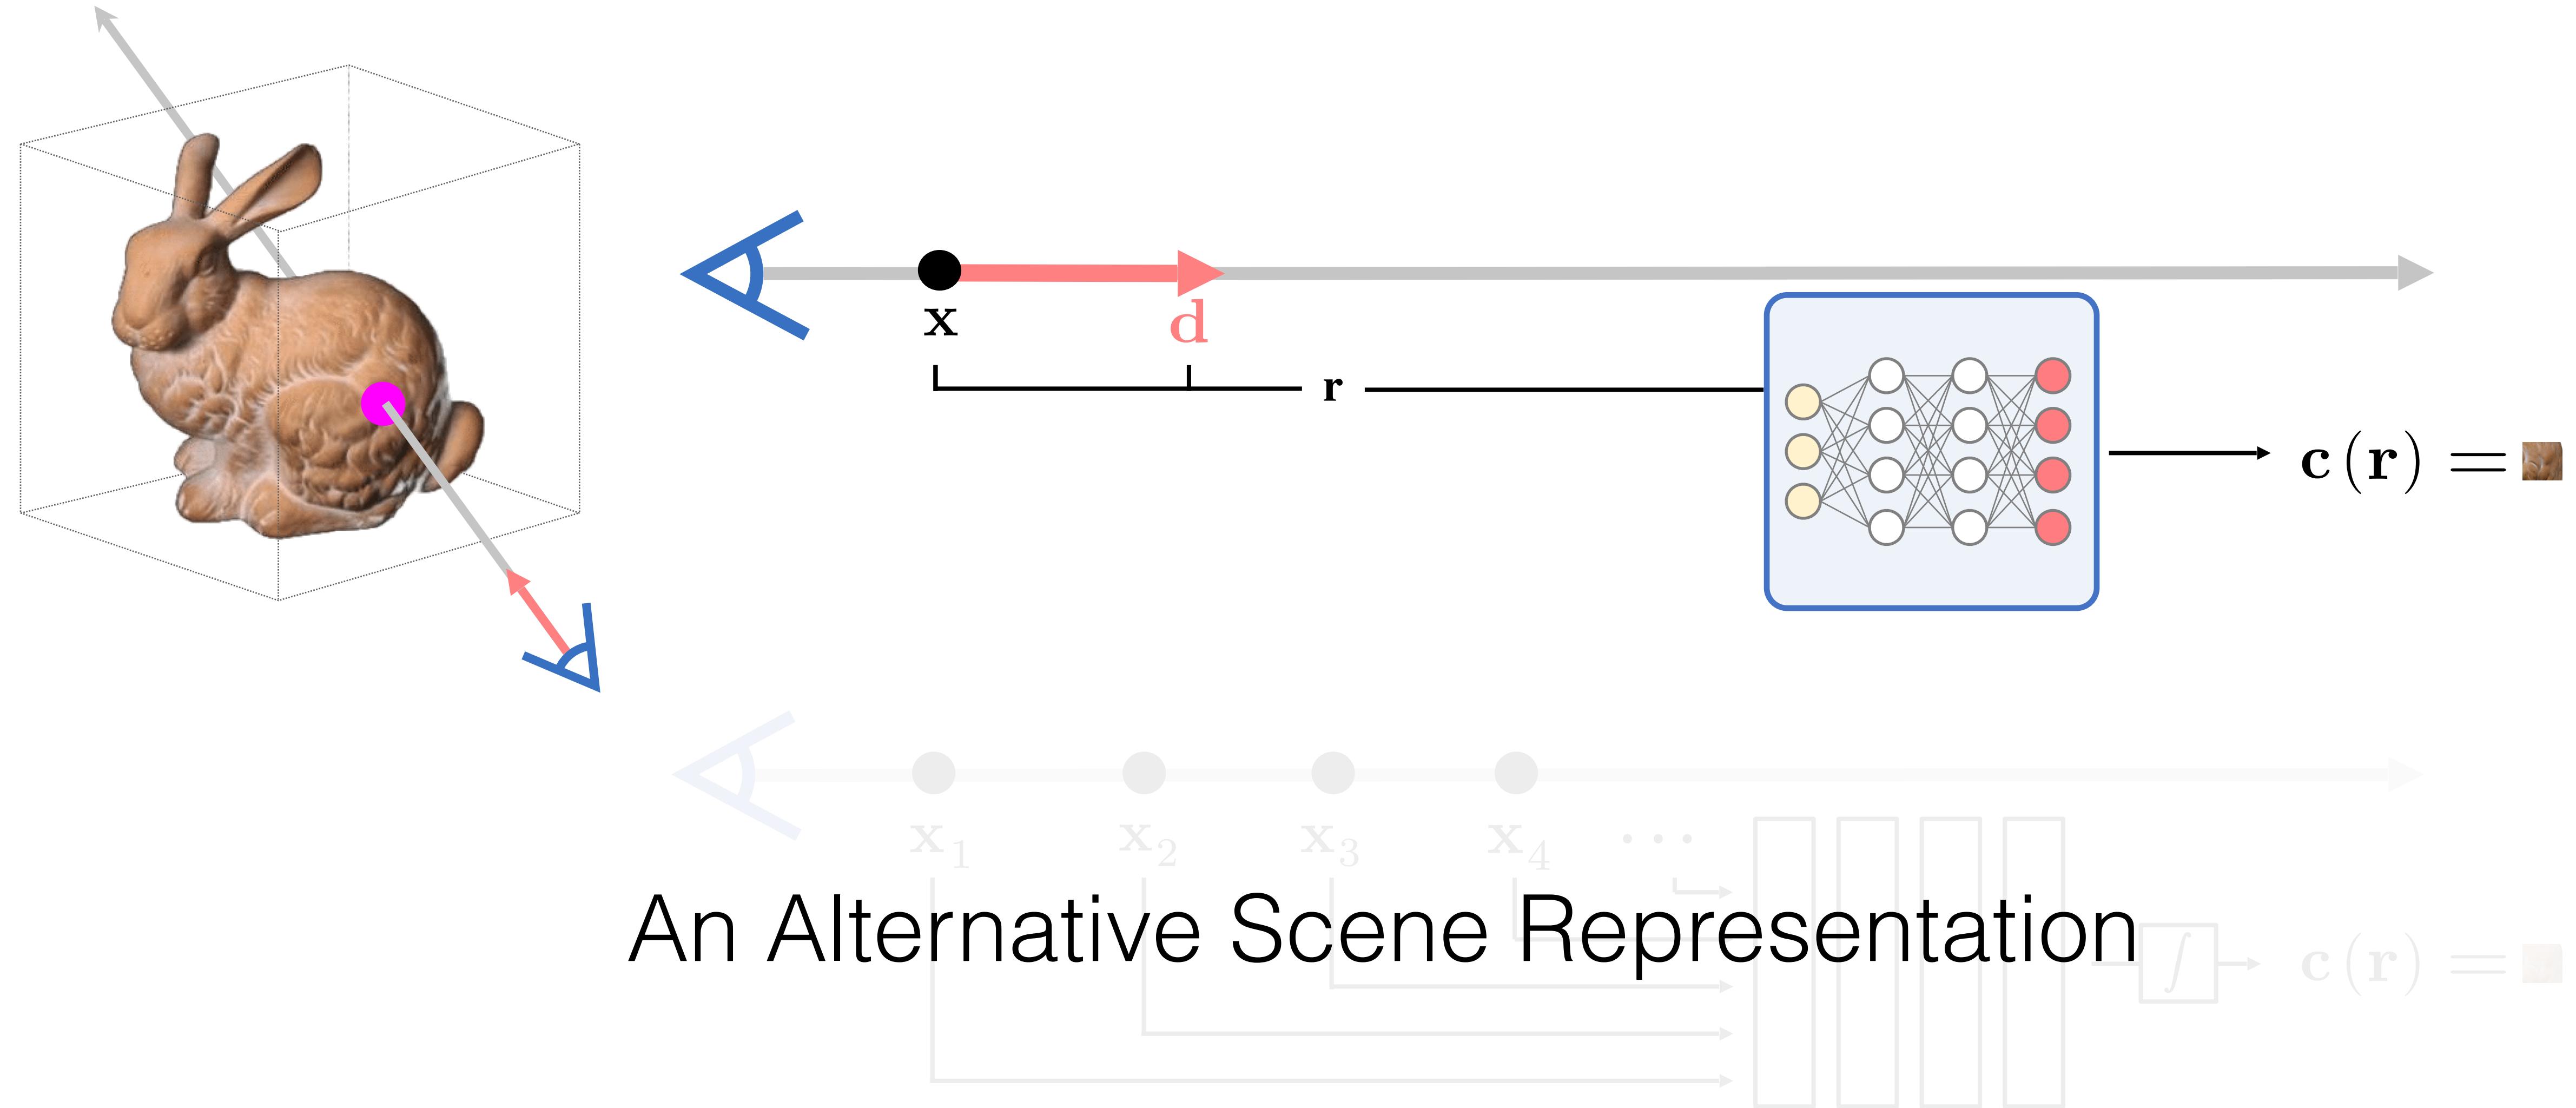
Light Field Networks



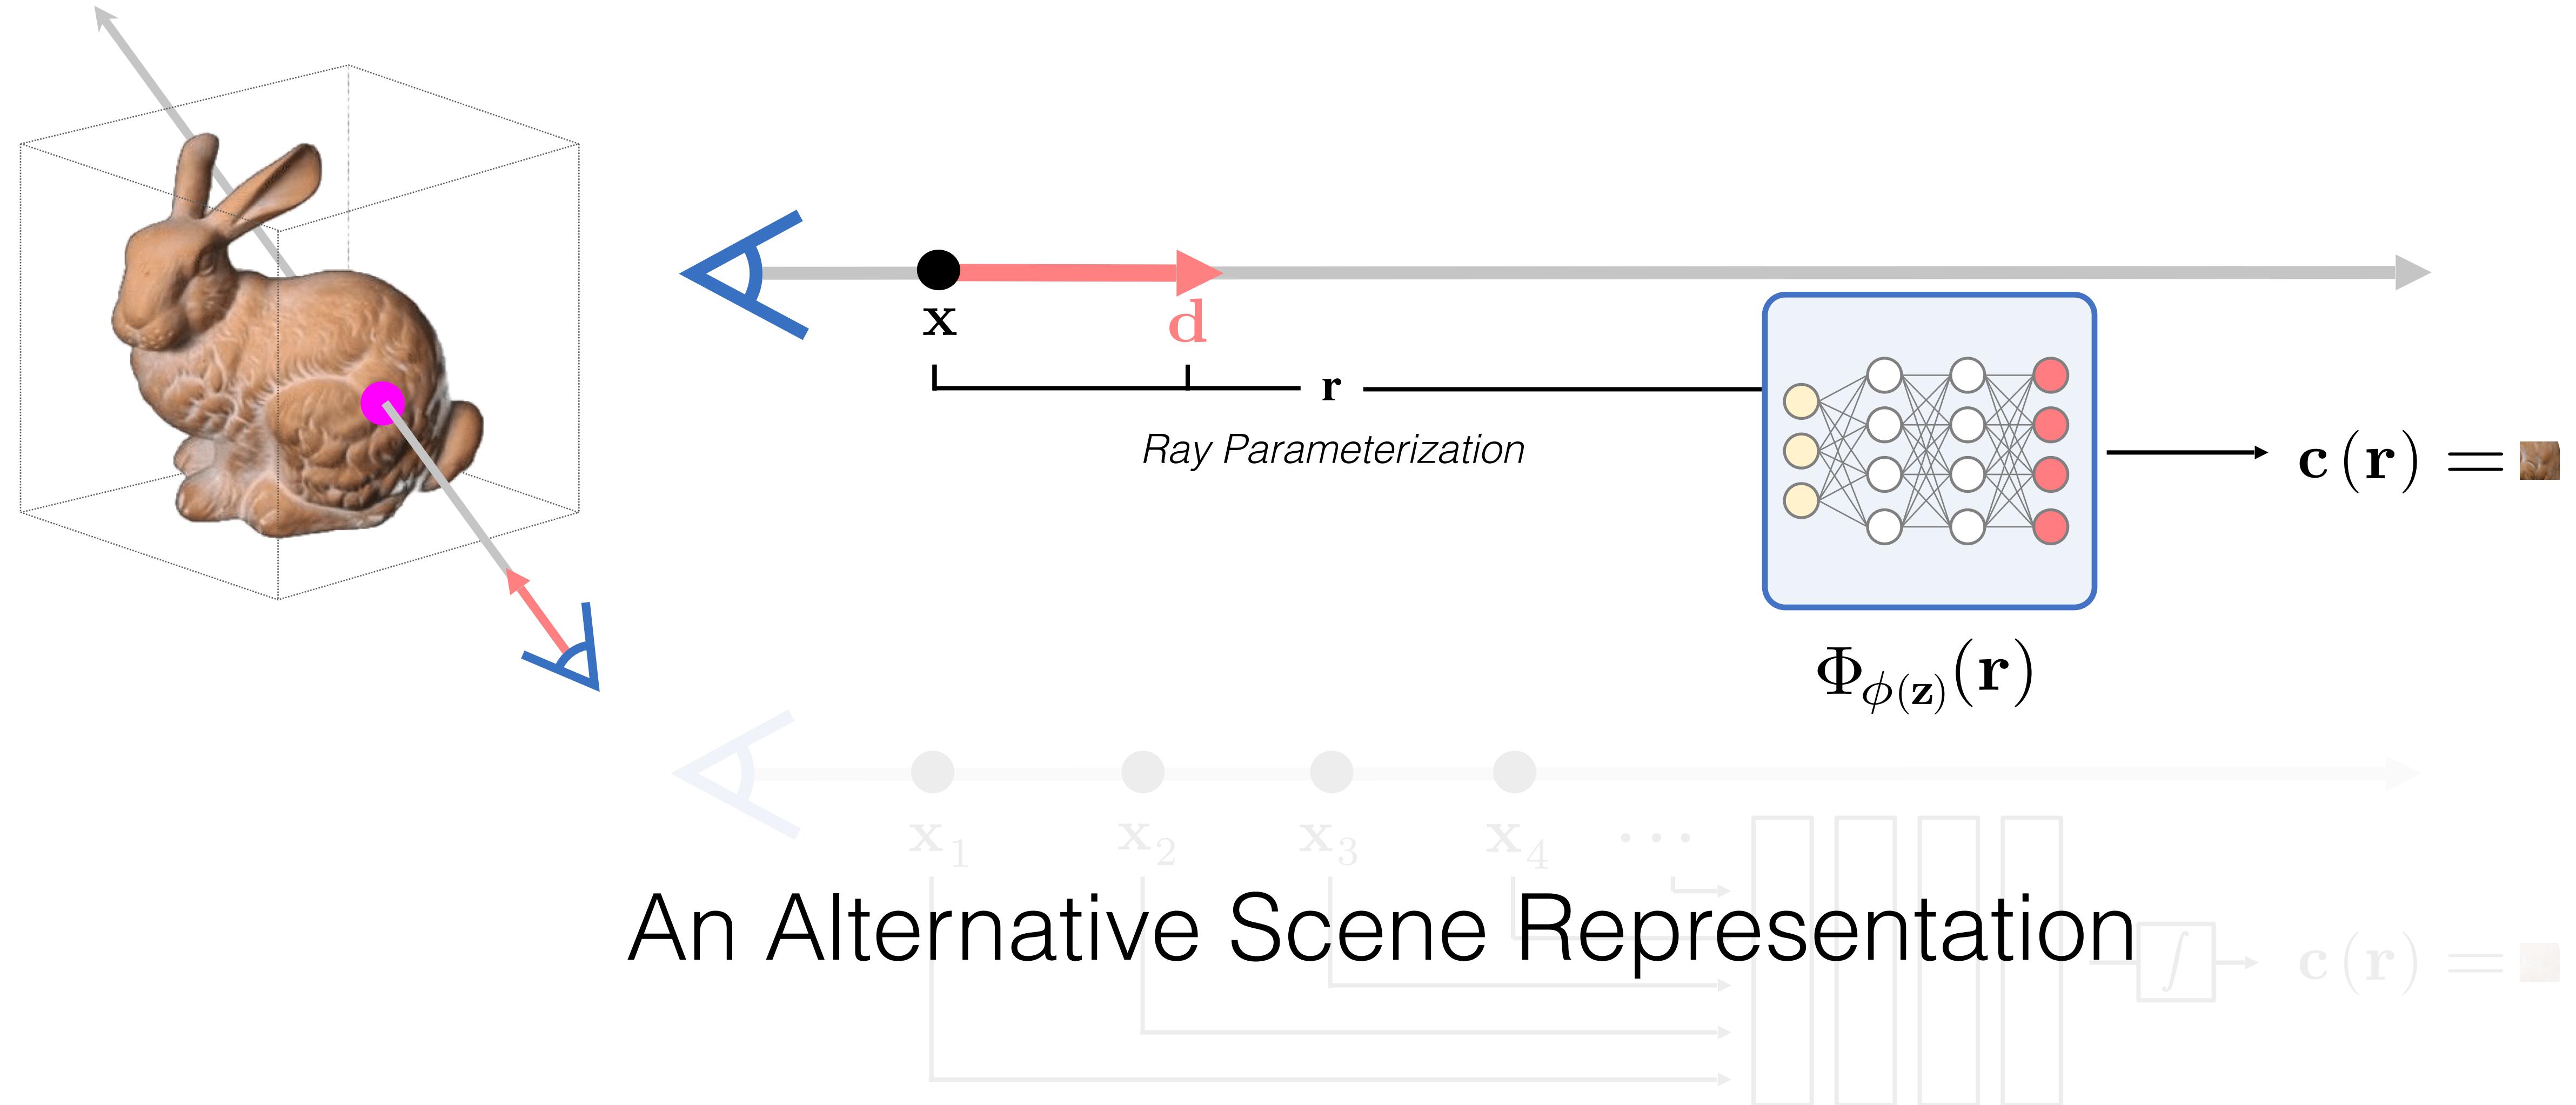
Light Field Networks



Light Field Networks



Light Field Networks



Overfitting doesn't work out-of-the-box: No built-in multi-view consistency!



Context Views

Overfitting doesn't work out-of-the-box: No built-in multi-view consistency!



Context Views

Overfitting doesn't work out-of-the-box: No built-in multi-view consistency!



Context Views



Intermediate Views

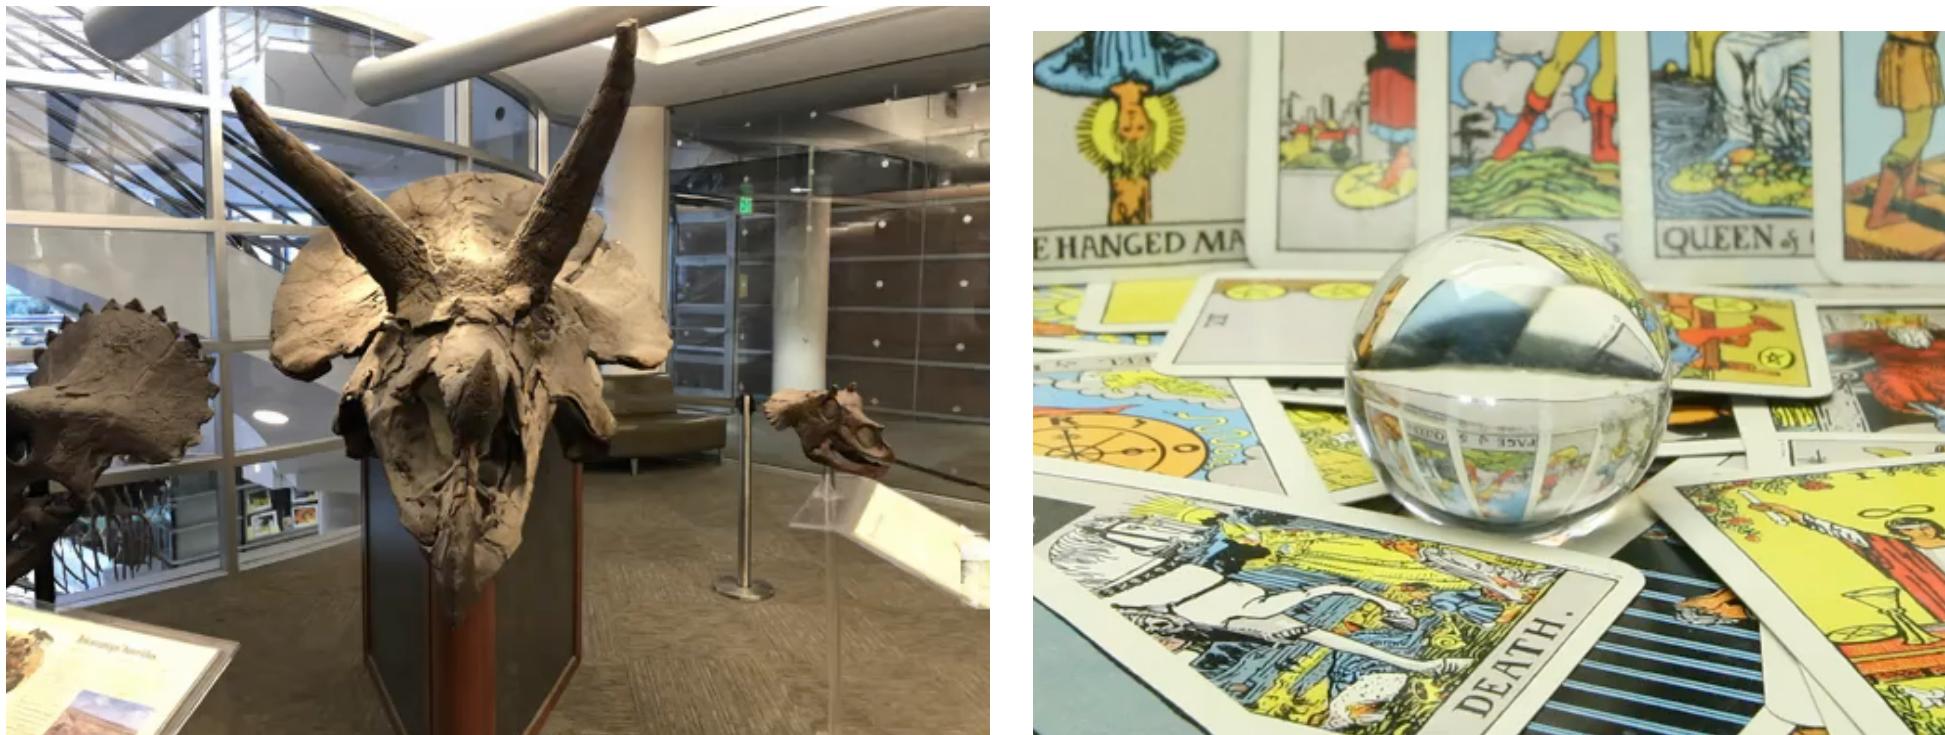
Achieving multi-view consistency in Light Fields

Single-Scene Overfitting

Regularization & Hybrid 3D / Light Field



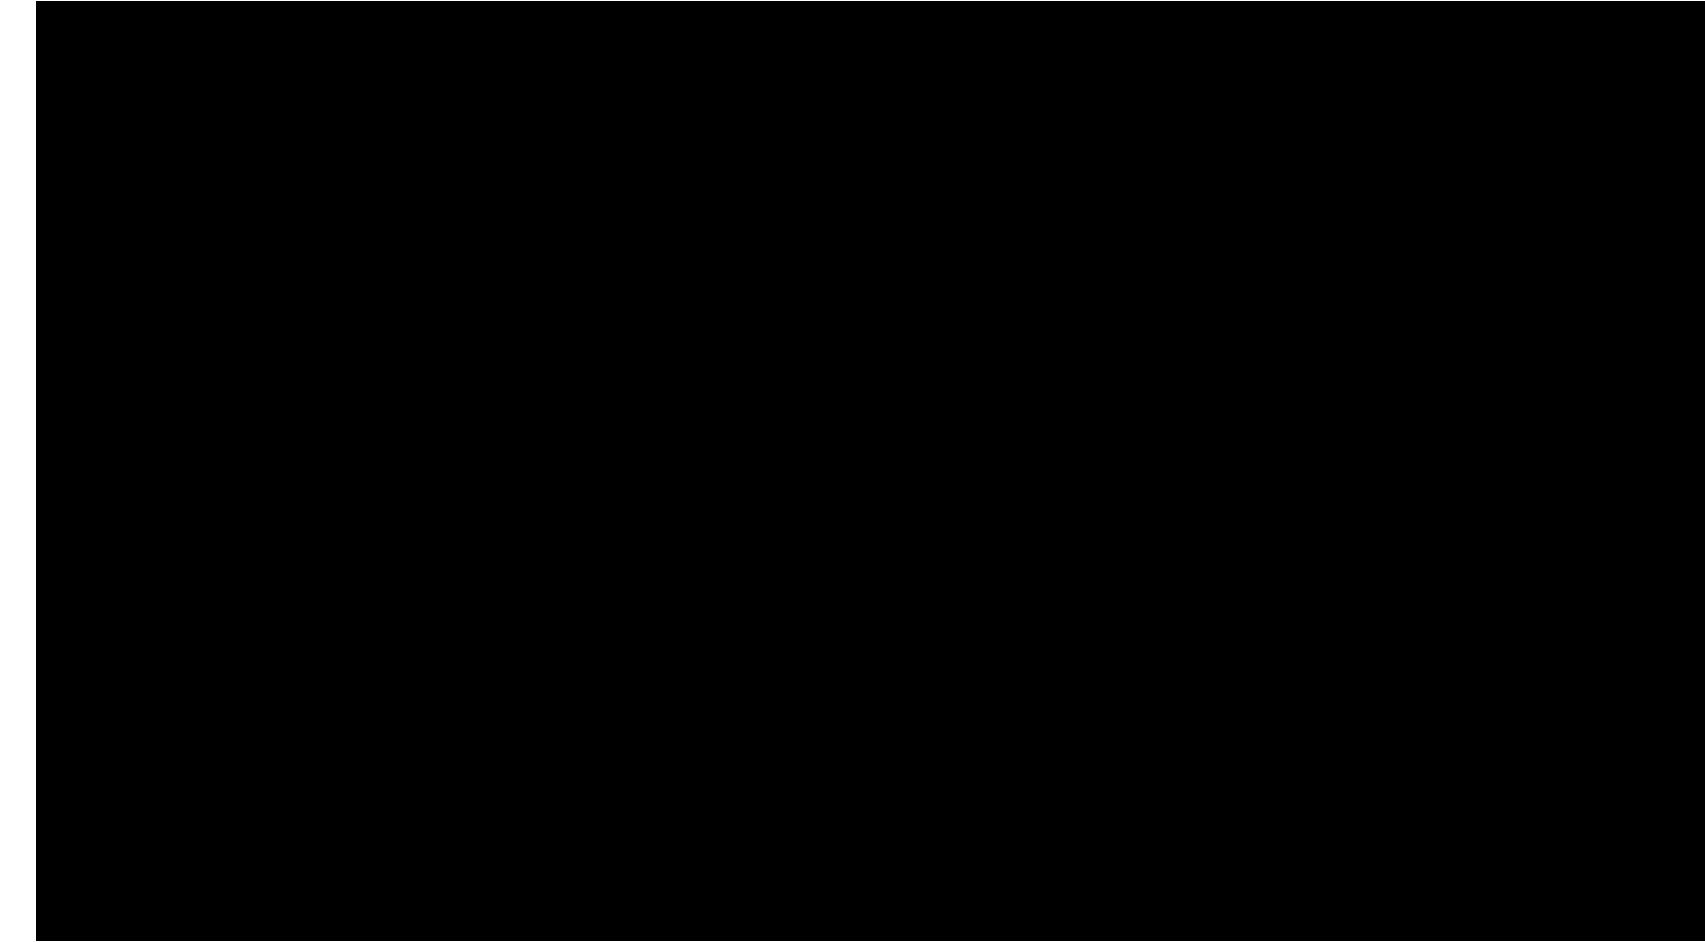
NeuLF: Efficient Novel View Synthesis with Neural 4D Light Field,
Li et al. 2022



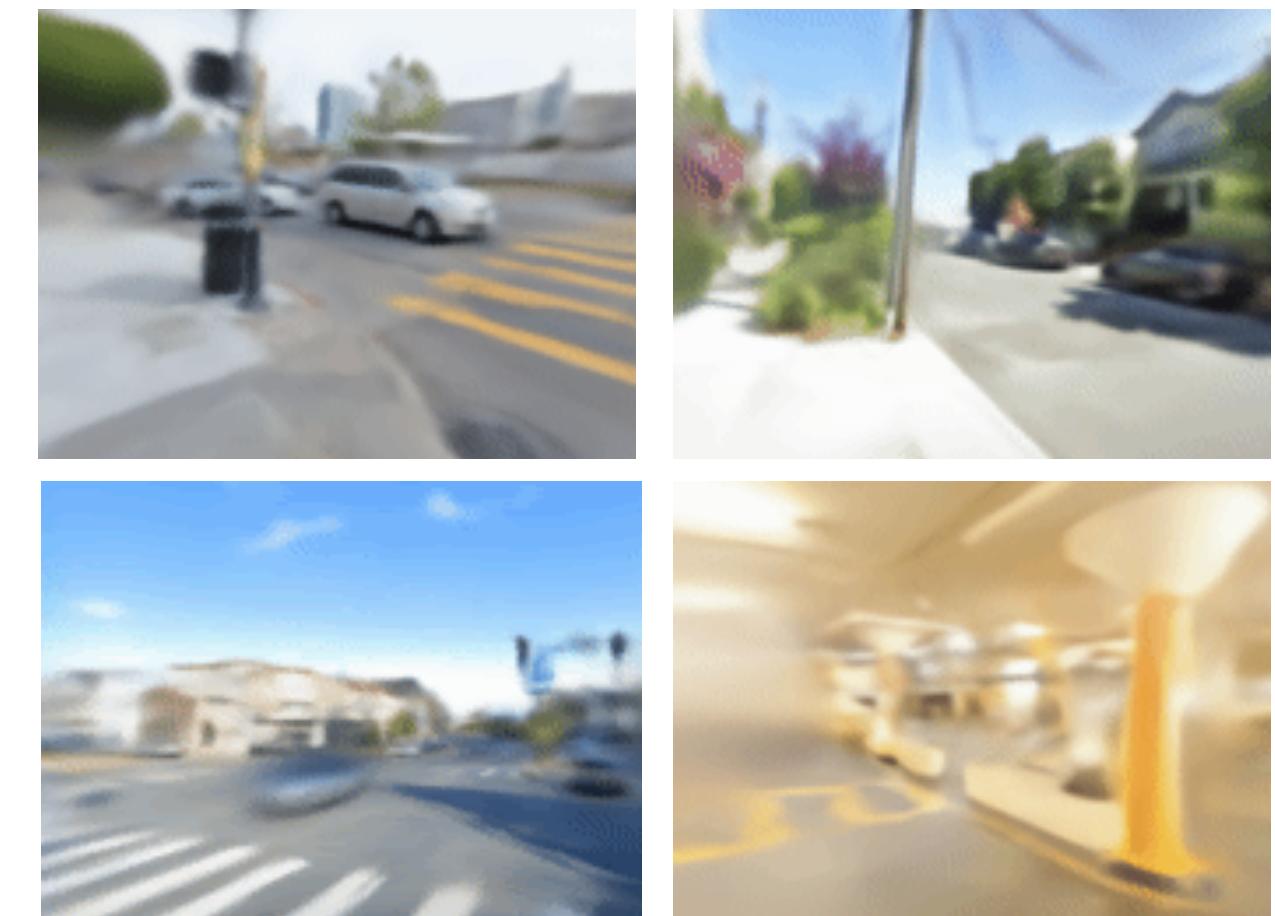
Learning Neural Light Fields with Ray-Space Embedding
Networks, Attal et al. 2022

Machine Learning

Multi-View Consistency Prior



Light Field Networks, Sitzmann et al. 2021



Scene Representation Transformer, Sajjadi et al. 2021

Achieving multi-view consistency in Light Fields

Single-Scene Overfitting

Regularization & Hybrid 3D / Light Field



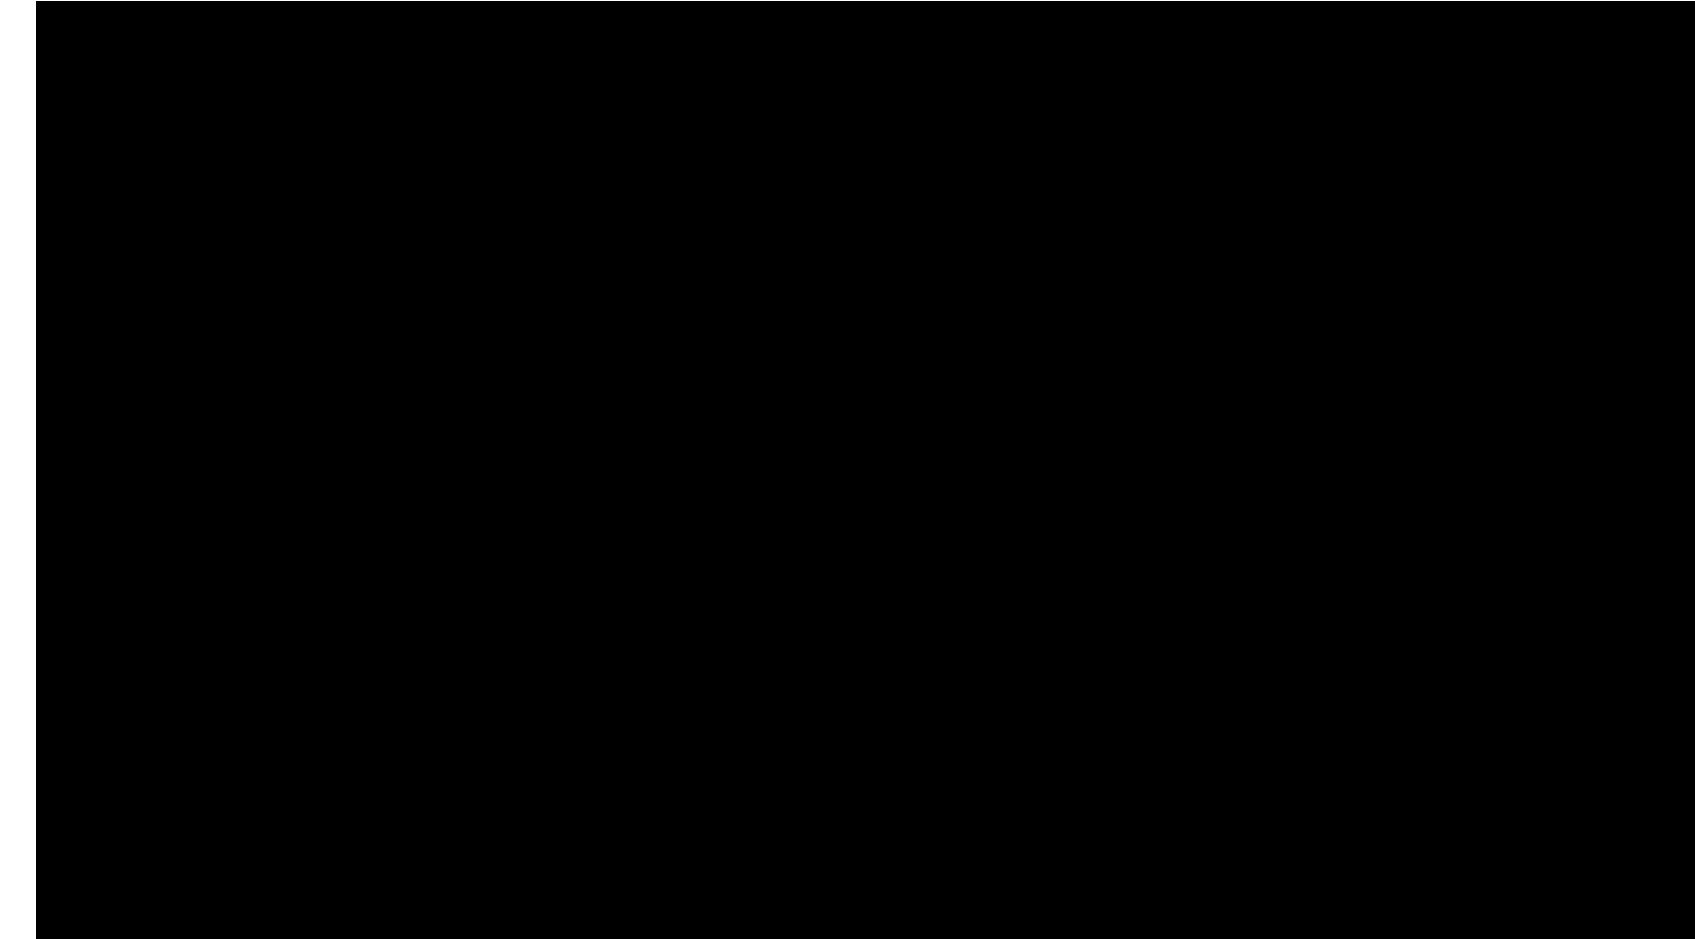
NeuLF: Efficient Novel View Synthesis with Neural 4D Light Field,
Li et al. 2022



Learning Neural Light Fields with Ray-Space Embedding
Networks, Attal et al. 2022

Machine Learning

Multi-View Consistency Prior



Light Field Networks, Sitzmann et al. 2021



Scene Representation Transformer, Sajjadi et al. 2021

Achieving multi-view consistency in Light Fields

Single-Scene Overfitting

Regularization & Hybrid 3D / Light Field



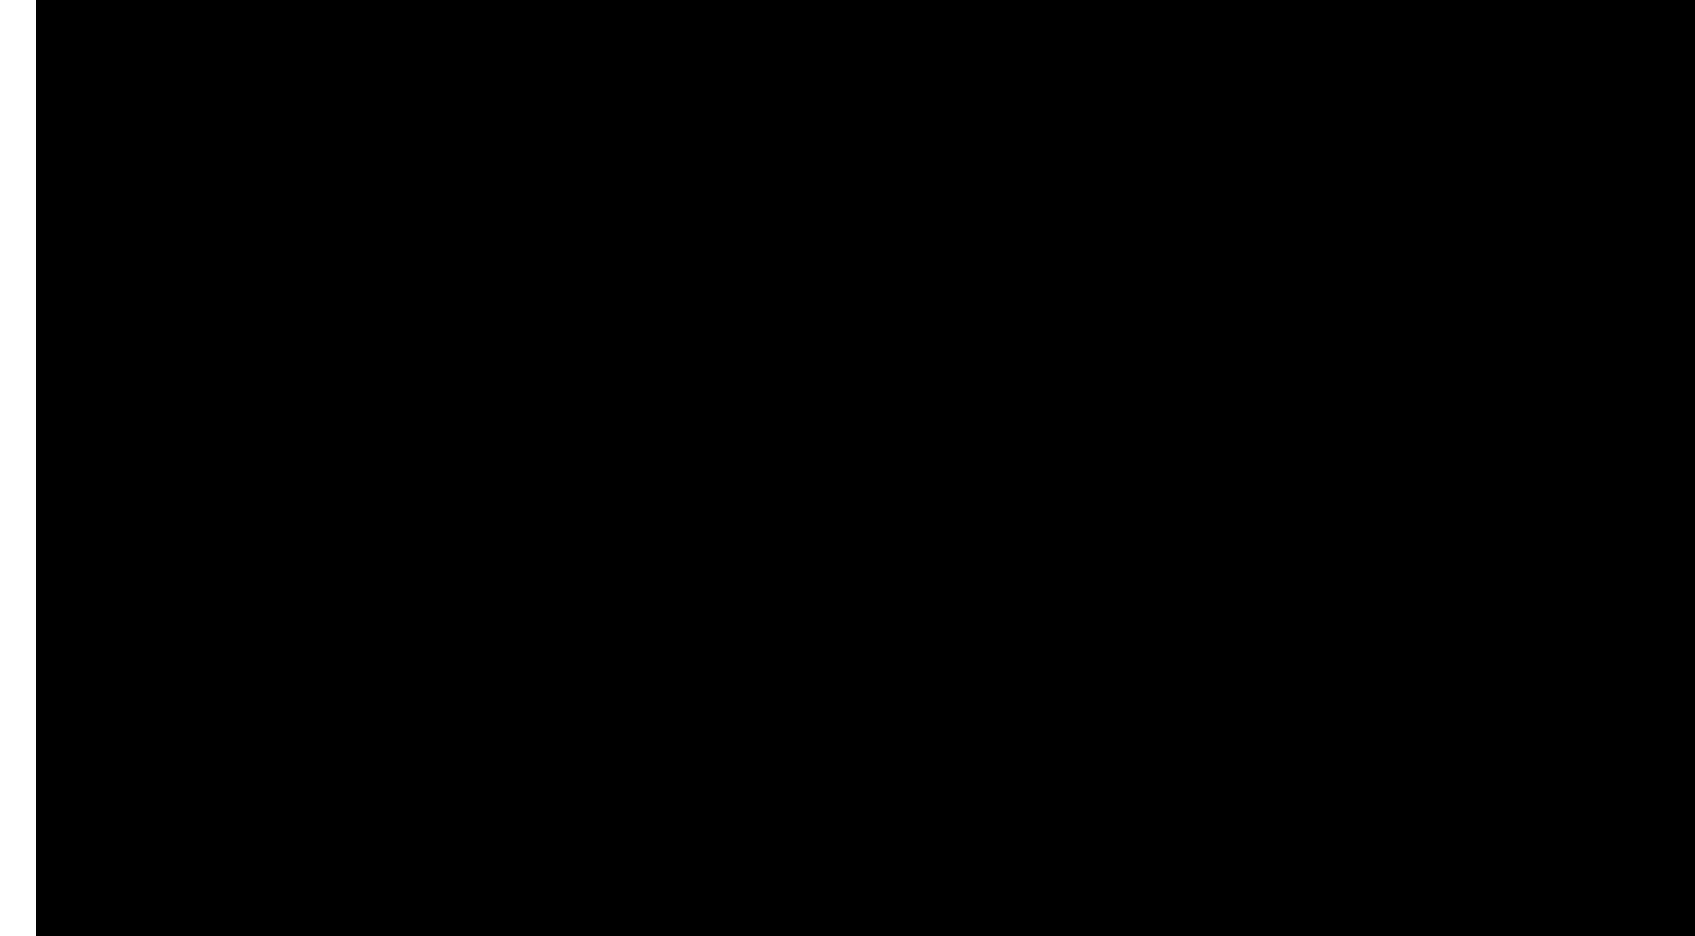
NeuLF: Efficient Novel View Synthesis with Neural 4D Light Field,
Li et al. 2022



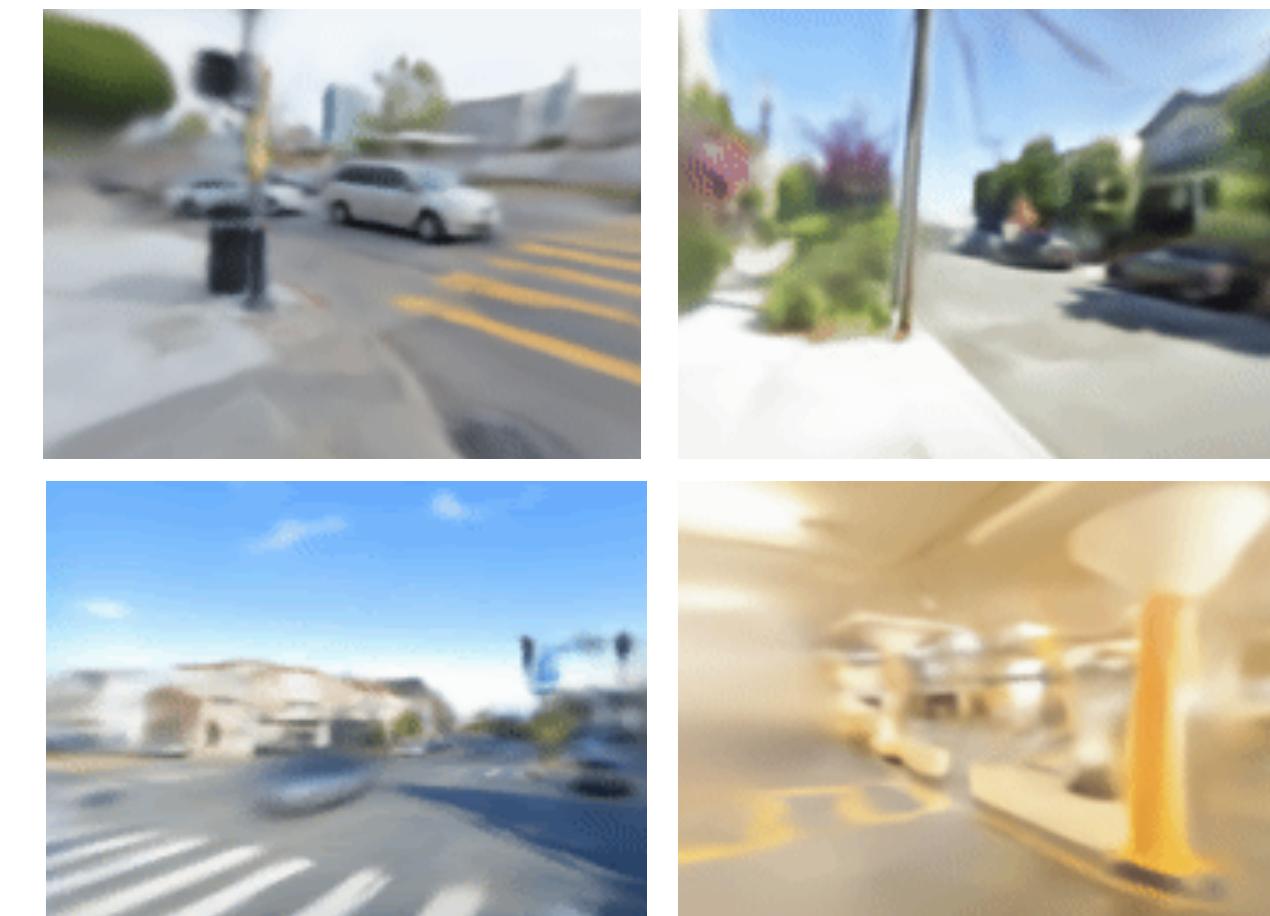
Learning Neural Light Fields with Ray-Space Embedding
Networks, Attal et al. 2022

Machine Learning

Multi-View Consistency Prior

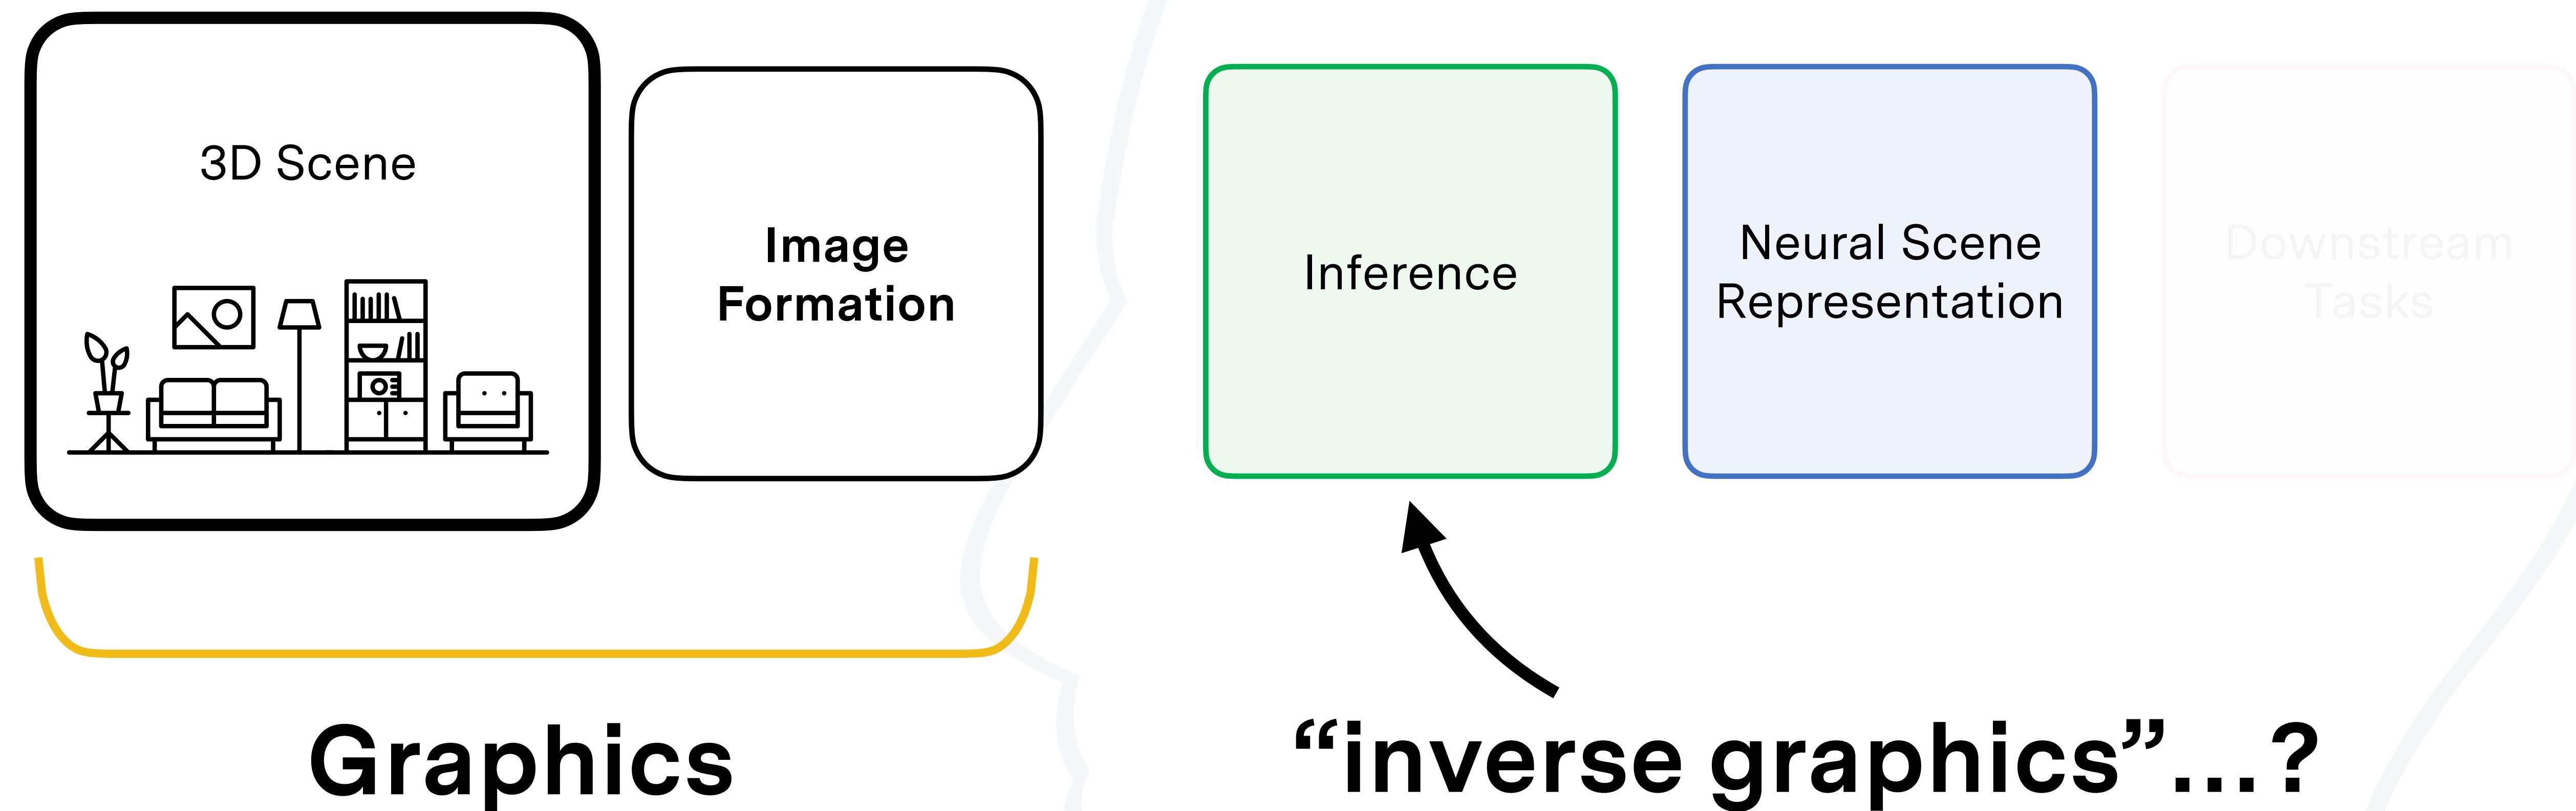


Light Field Networks, Sitzmann et al. 2021



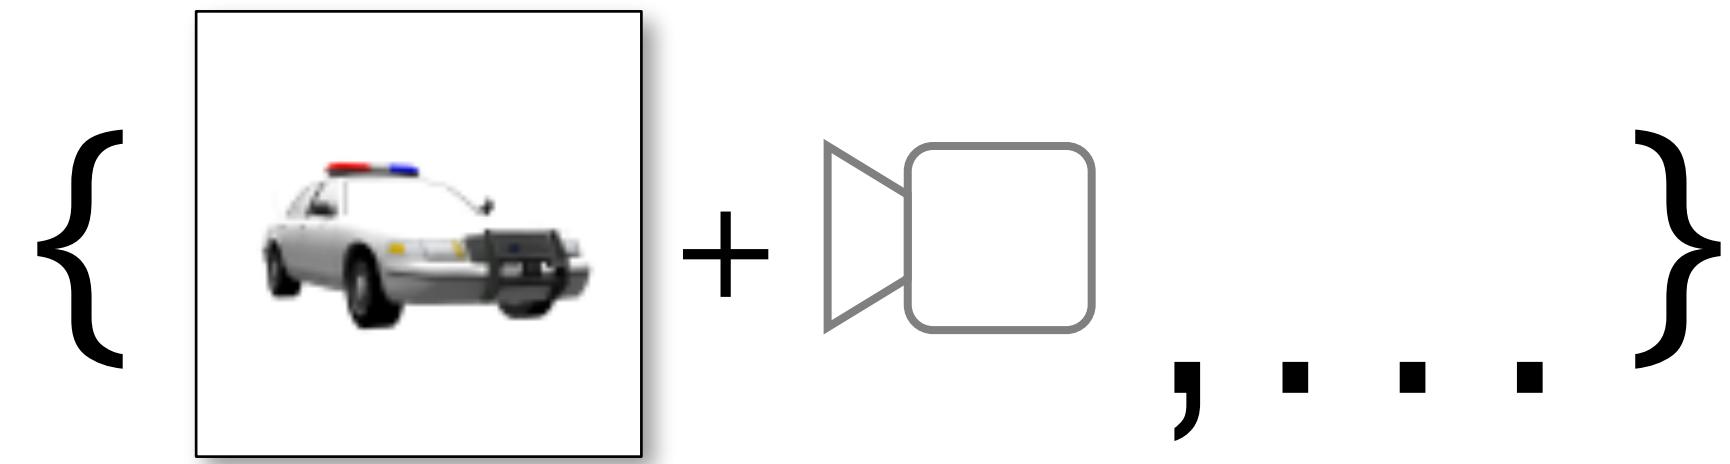
Scene Representation Transformer, Sajjadi et al. 2021

Today: Differentiable Rendering, AKA “Inverse Graphics”

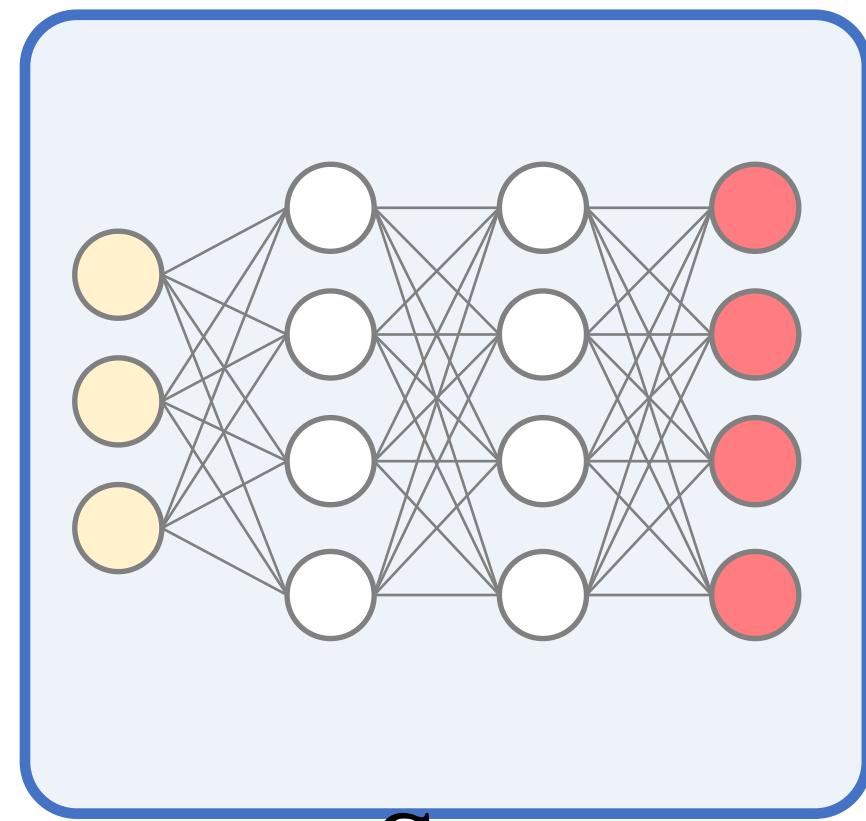
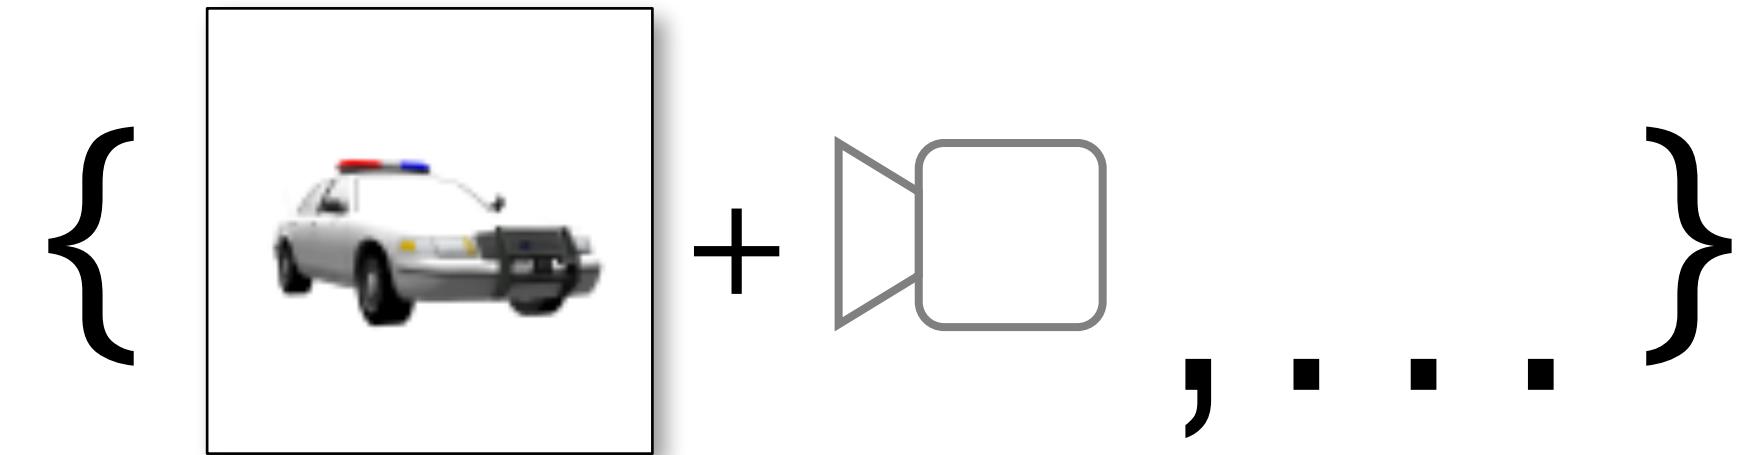


Overfitting case: Inference = Fitting via Gradient Descent

Overfitting case: Inference = Fitting via Gradient Descent

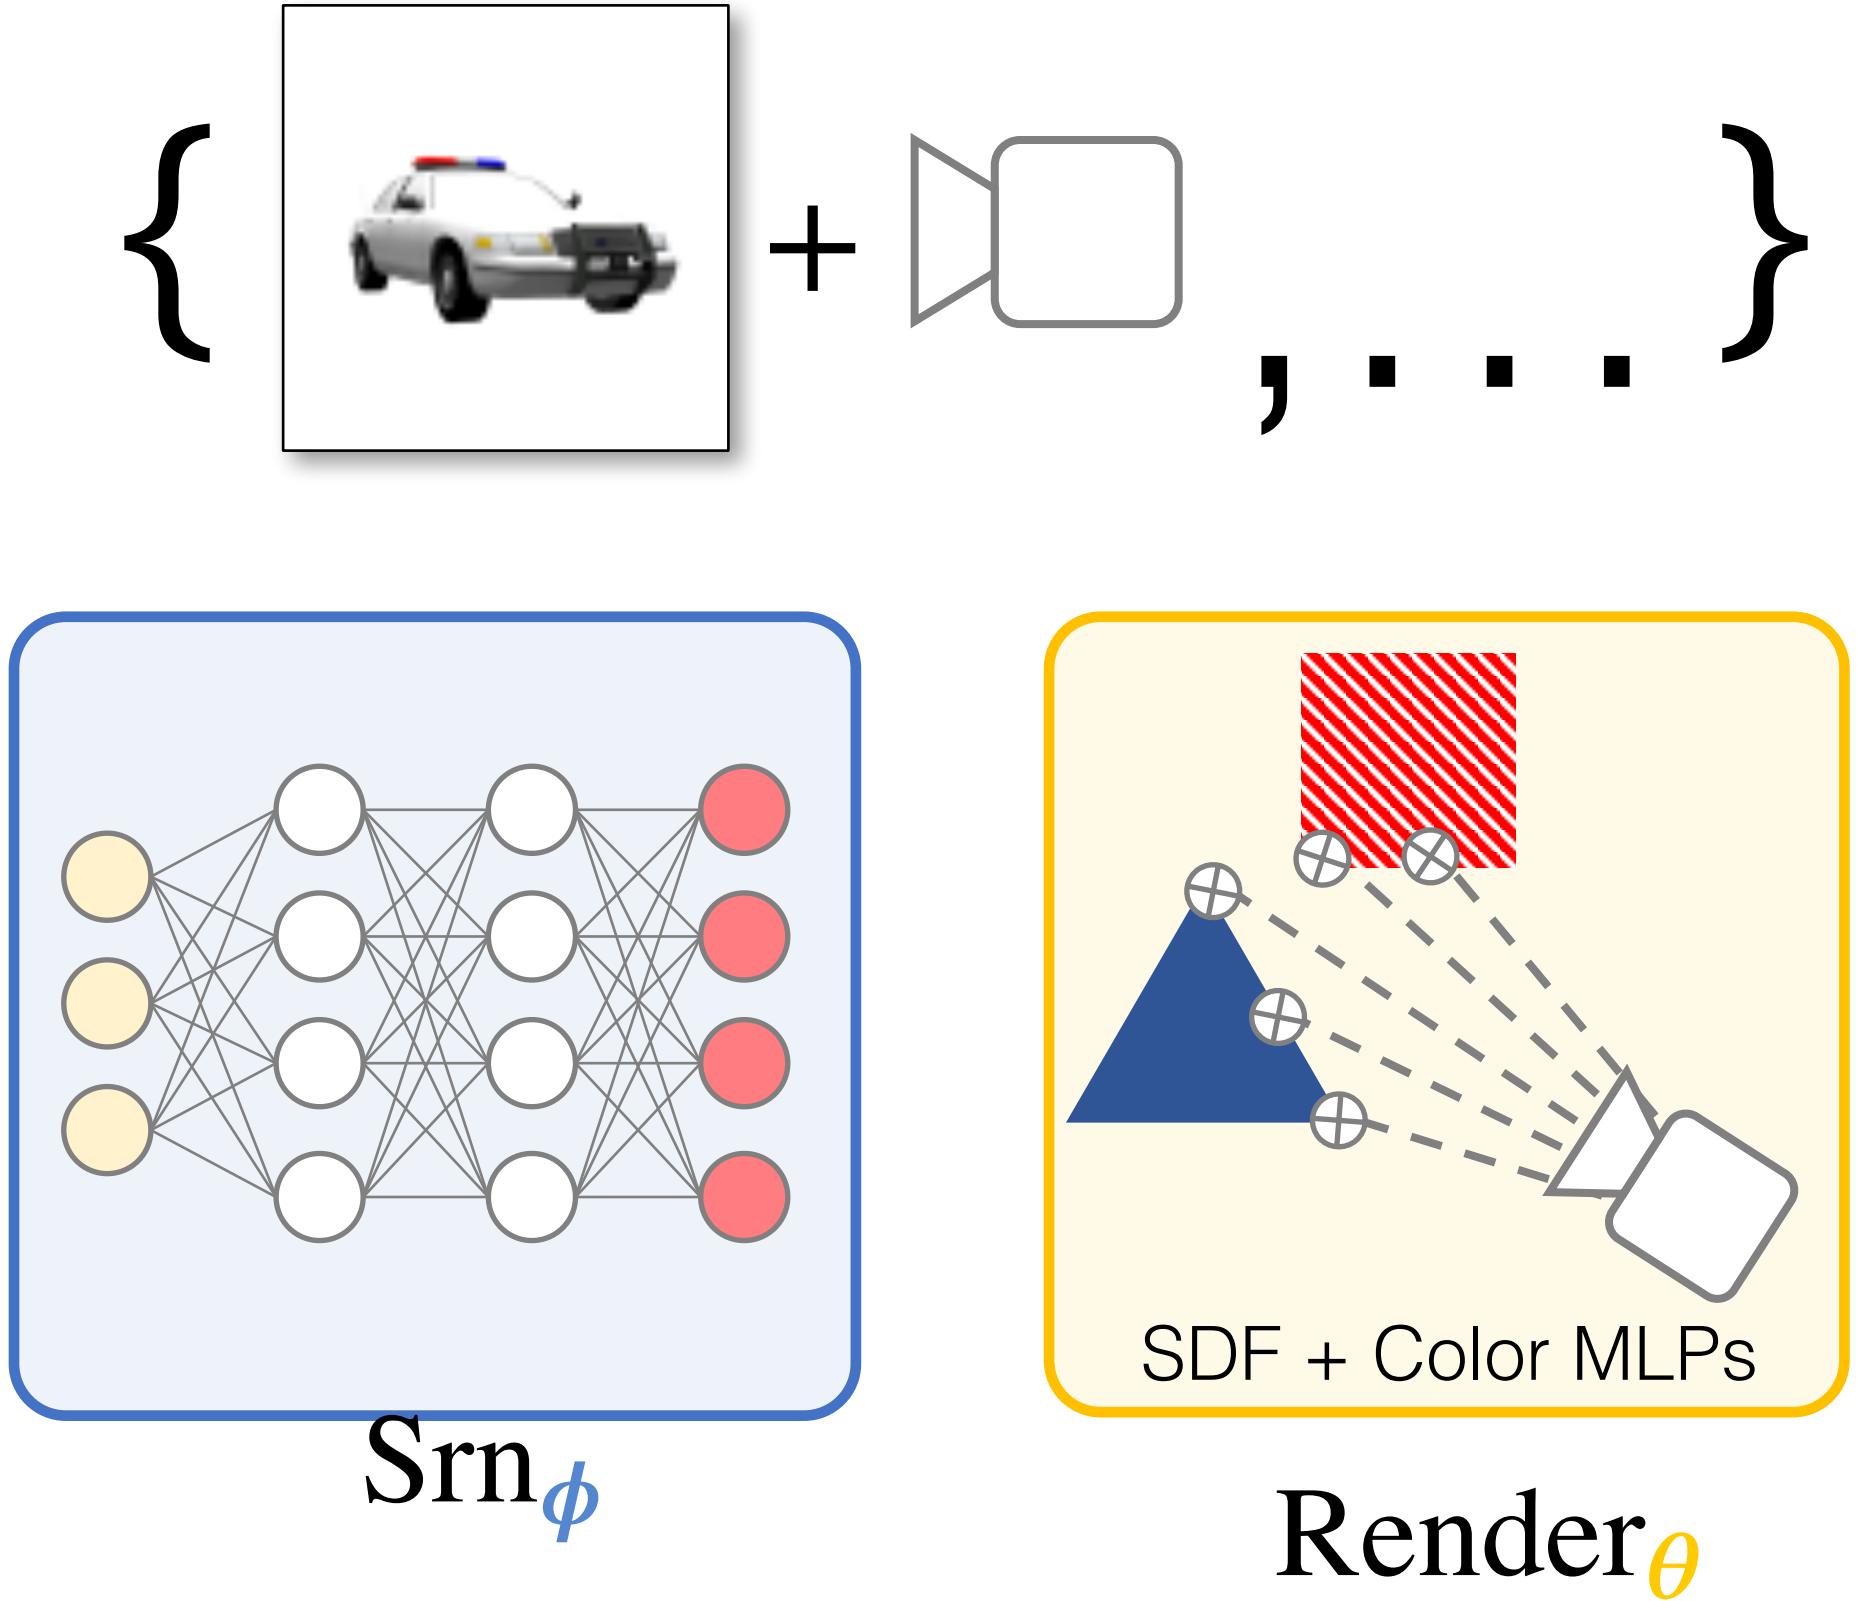


Overfitting case: Inference = Fitting via Gradient Descent

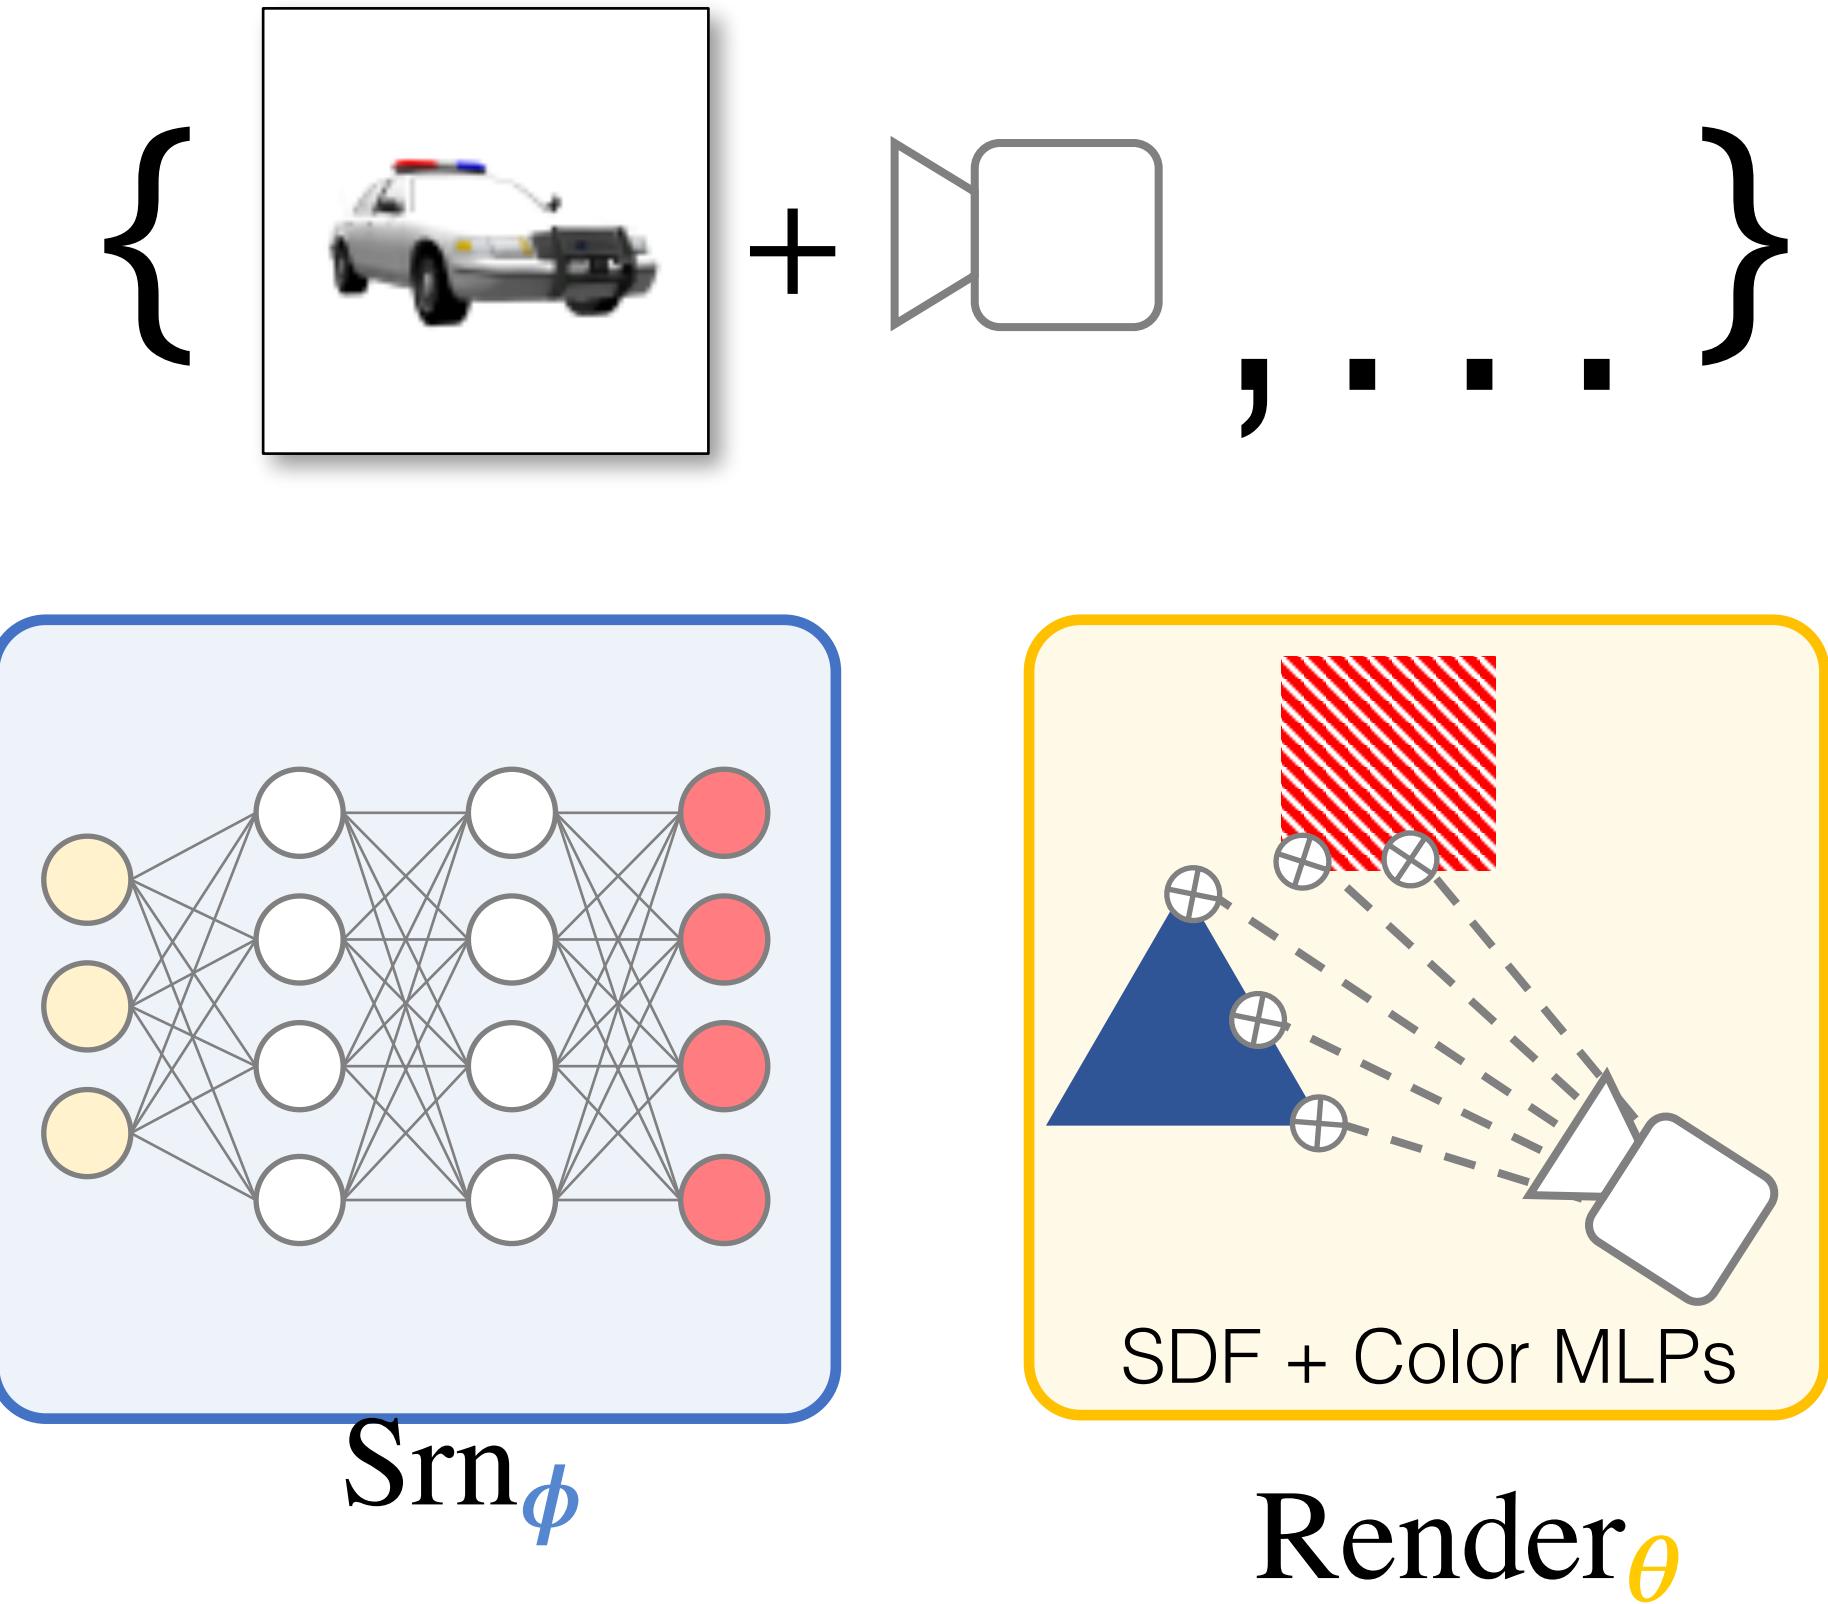


Srn_ϕ

Overfitting case: Inference = Fitting via Gradient Descent

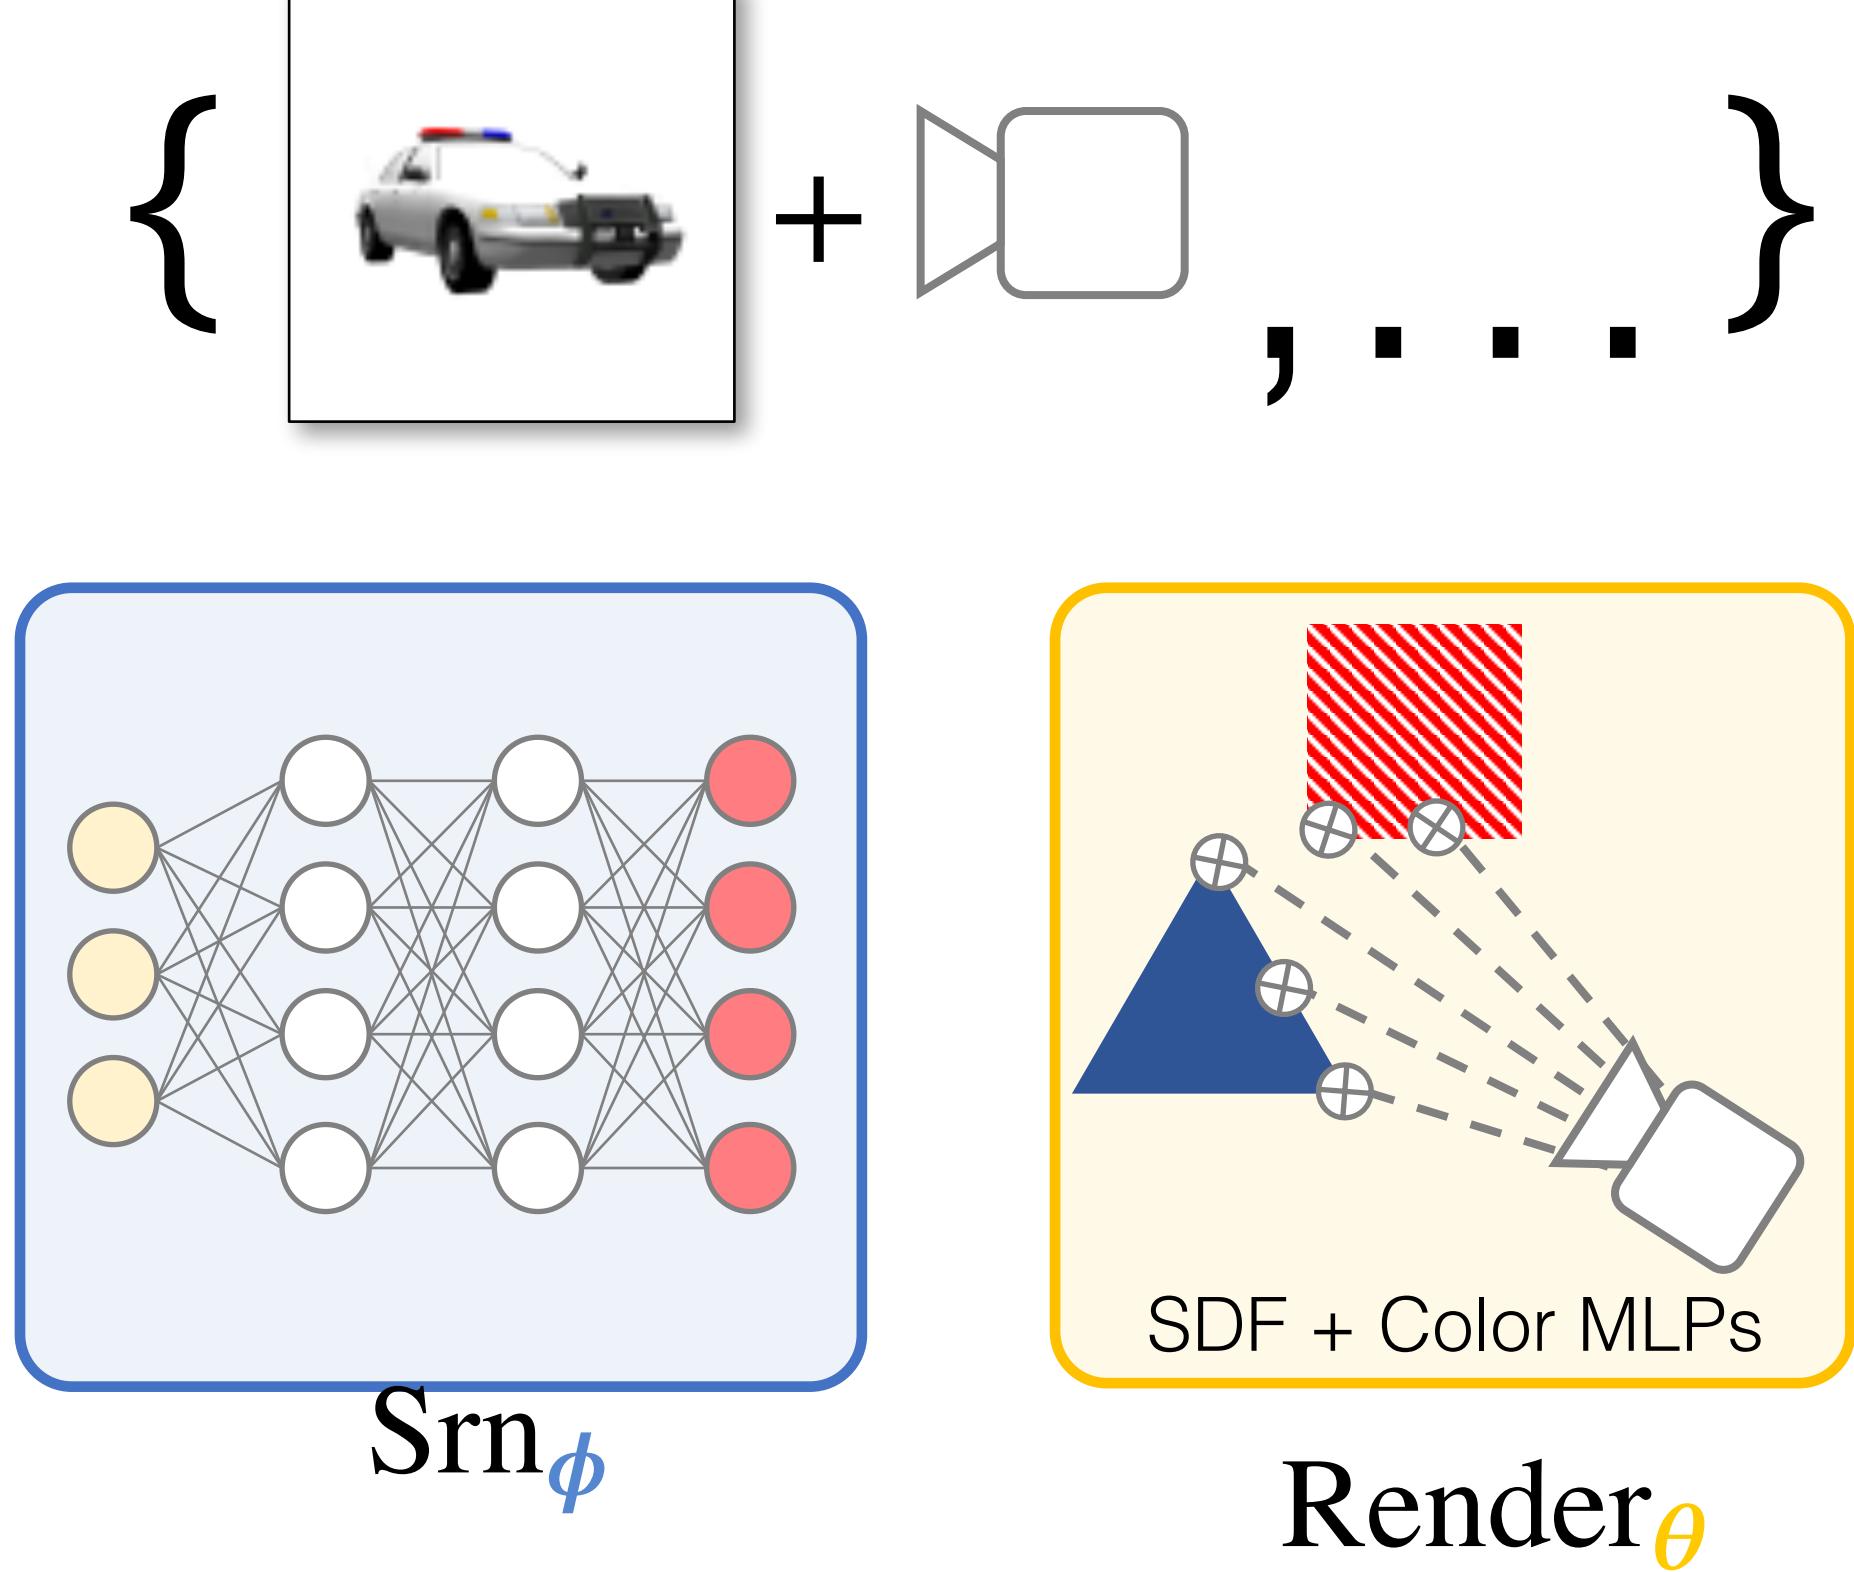


Overfitting case: Inference = Fitting via Gradient Descent



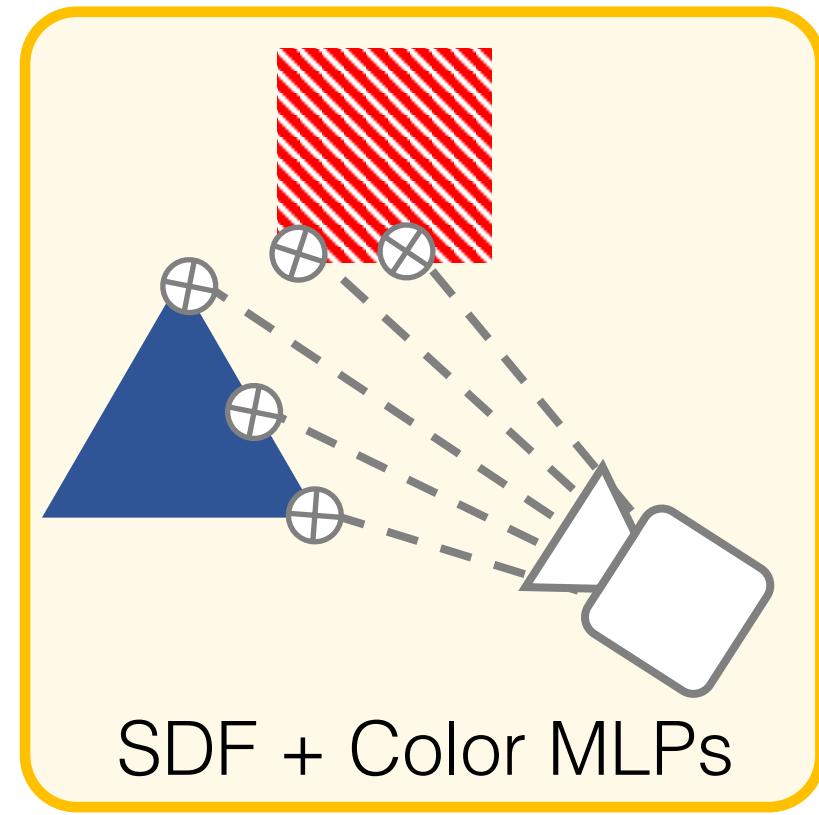
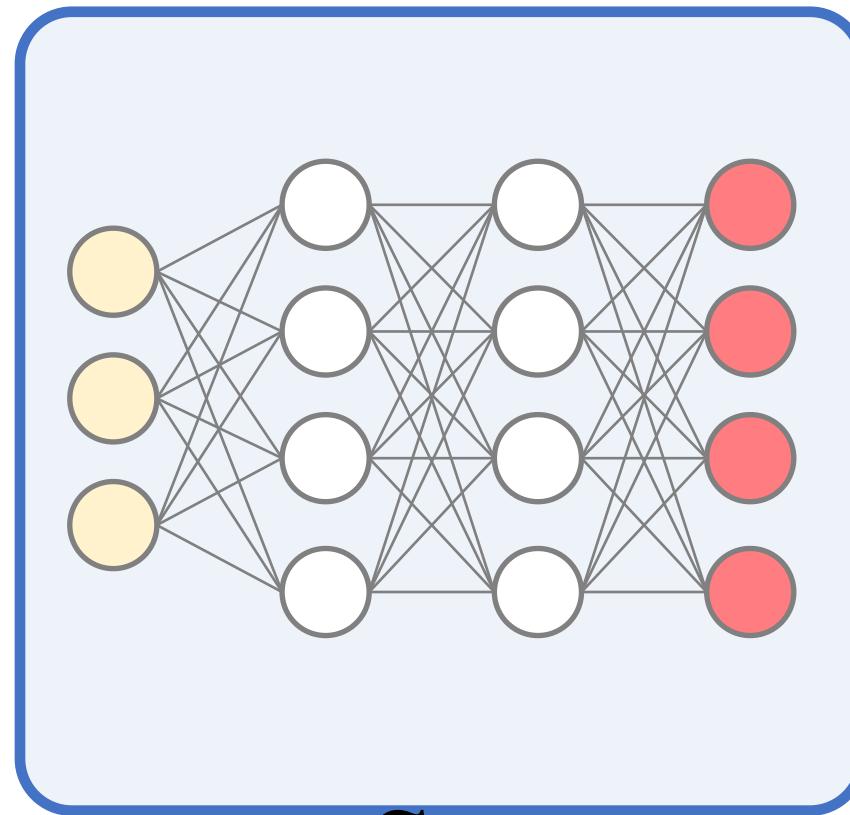
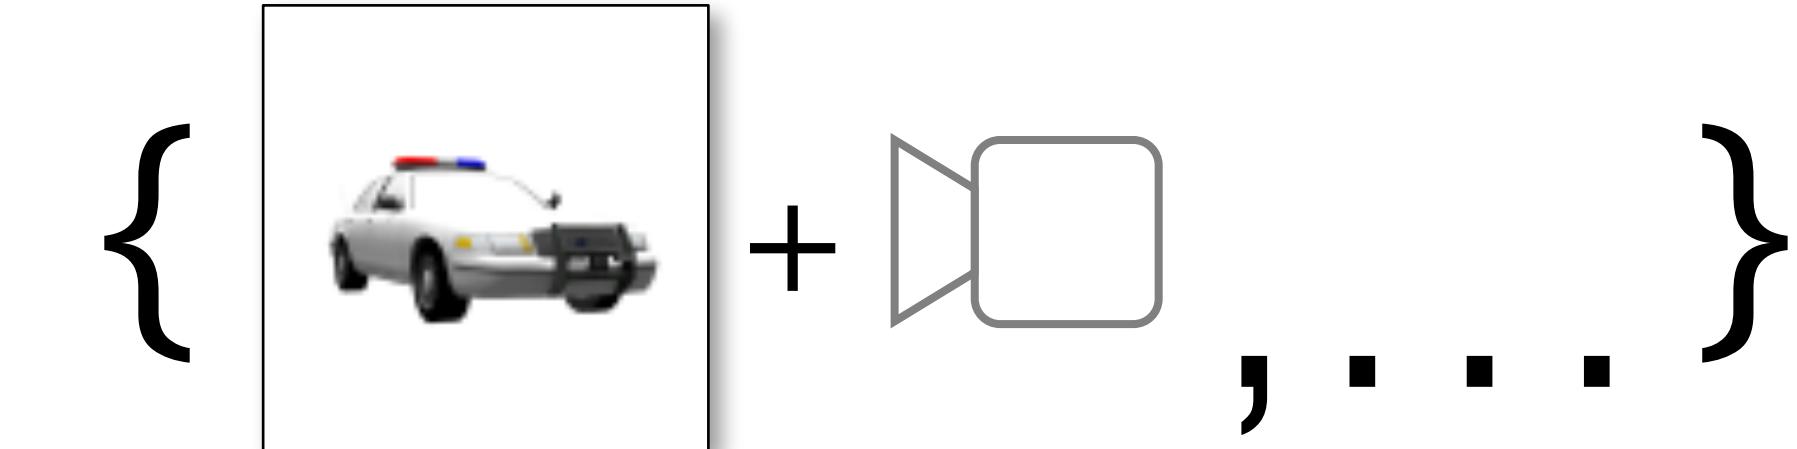
$$\min \left\| \text{Render}_{\theta}(\text{Srn}_{\phi}, \xi_i) - \mathcal{I}_i \right\|$$

Overfitting case: Inference = Fitting via Gradient Descent Optimization



$$\min \left\| \text{Render}_\theta(\text{Srn}_\phi, \xi_i) - \mathcal{I}_i \right\|$$

Overfitting case: Inference = Fitting via Gradient Descent Optimization



Render $_\theta$

$$\min \left\| \text{Render}_\theta(\text{Srn}_\phi, \xi_i) - \mathcal{I}_i \right\|$$



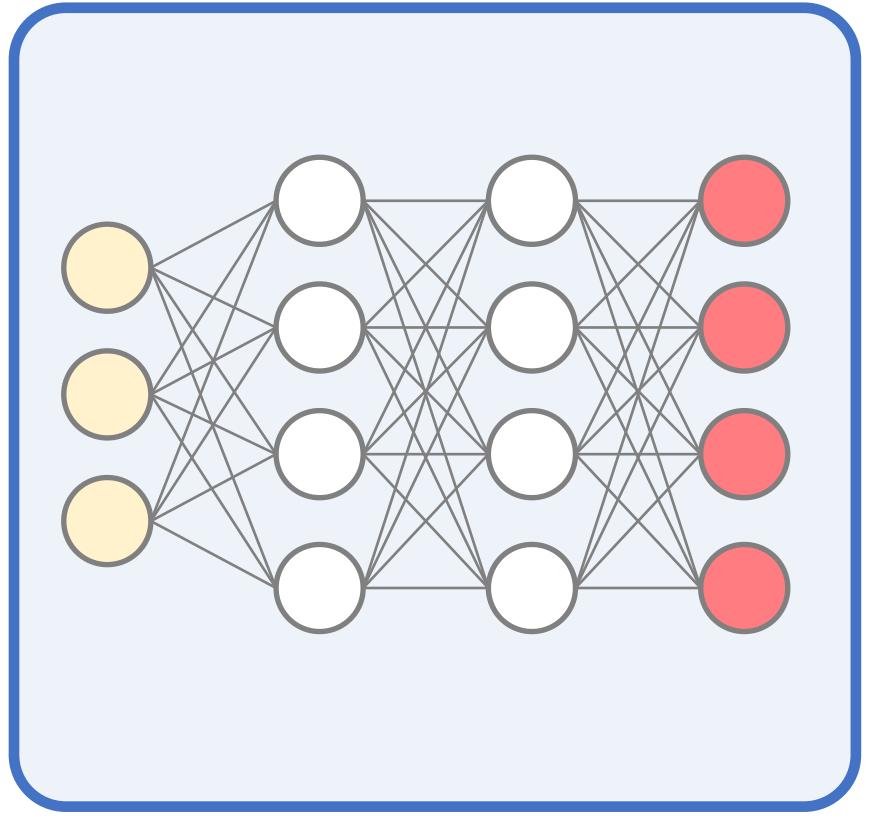
Novel View Synthesis



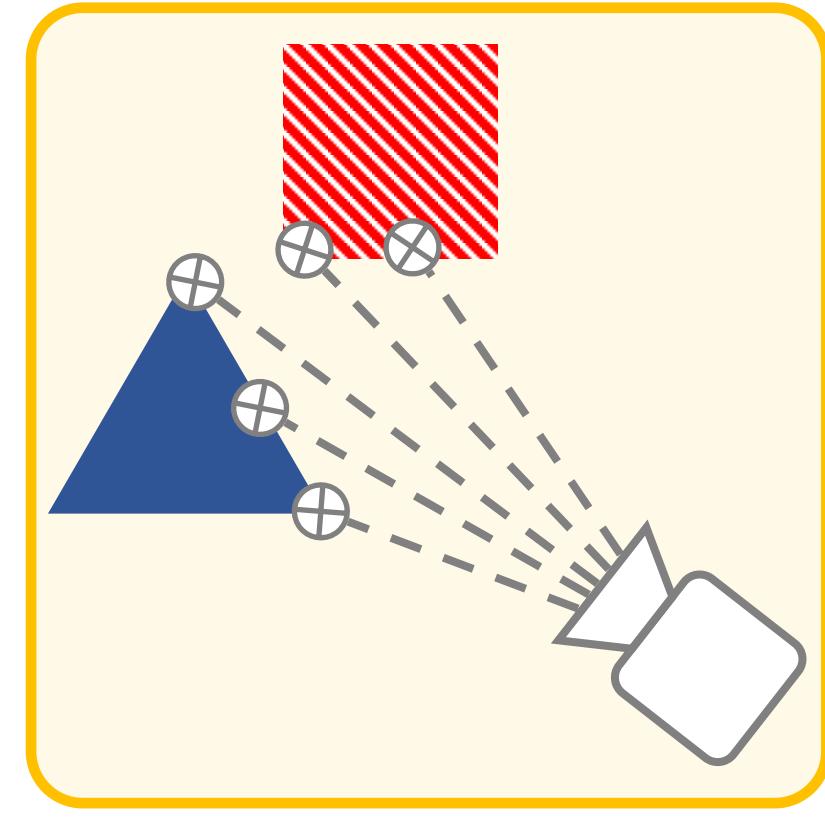
Normal map

RGB

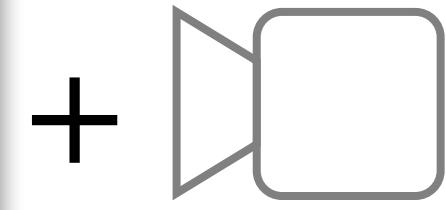
What if observations don't constrain scene representation?



Srn_ϕ



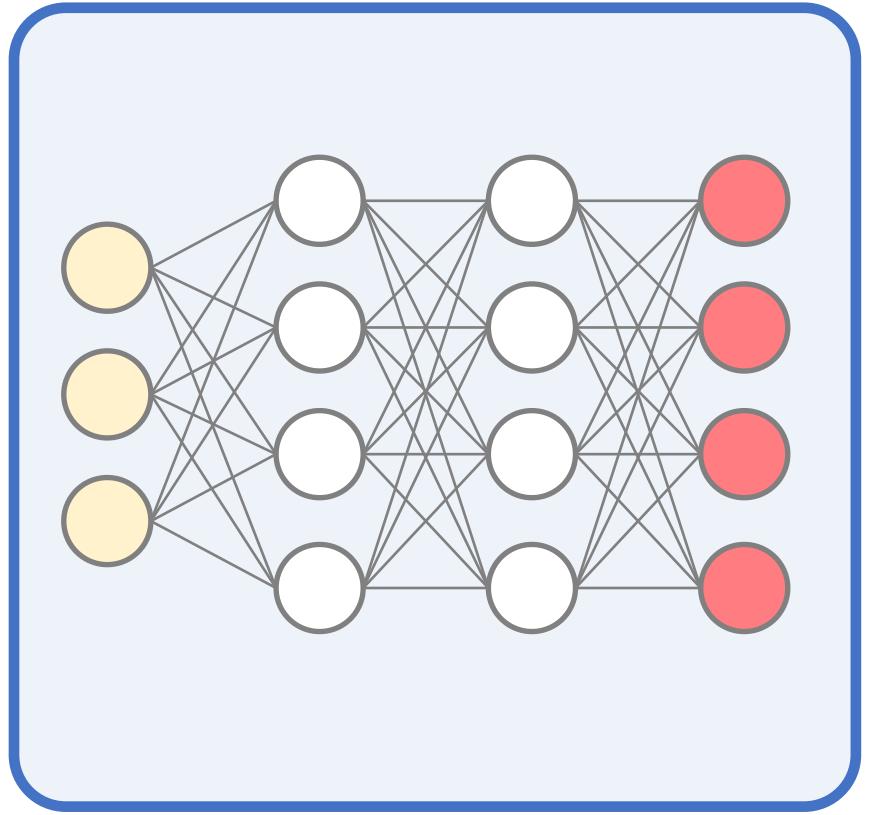
Render_θ



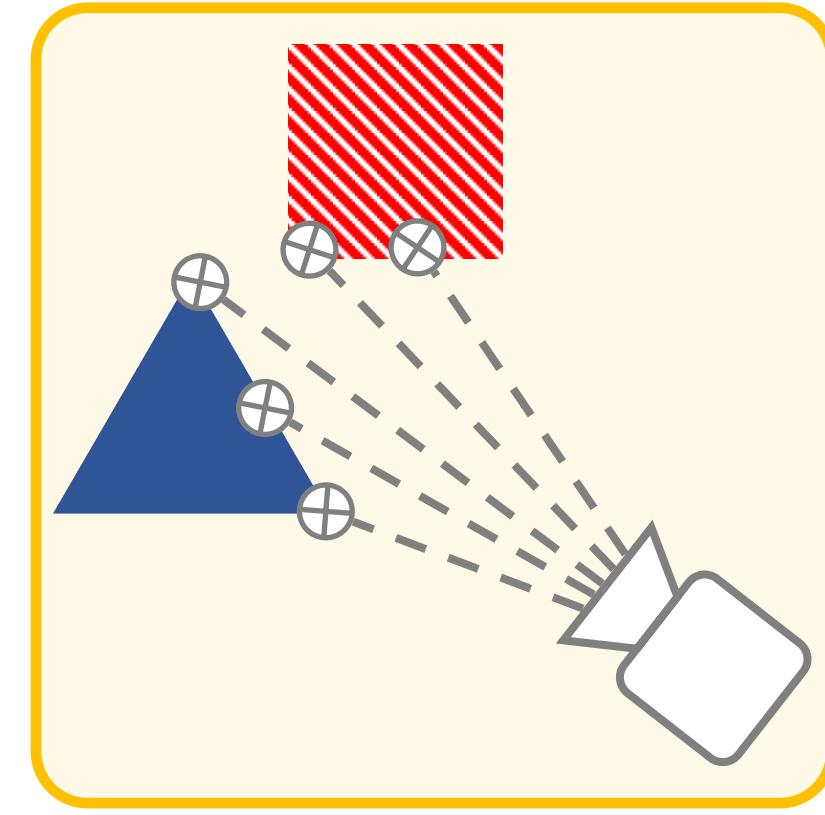
$$(\mathcal{I}, \xi)$$

$$\underset{\phi, \theta}{\operatorname{argmin}} \| \text{Render}_\theta(\text{Srn}_\phi, \xi) - \mathcal{I} \|$$

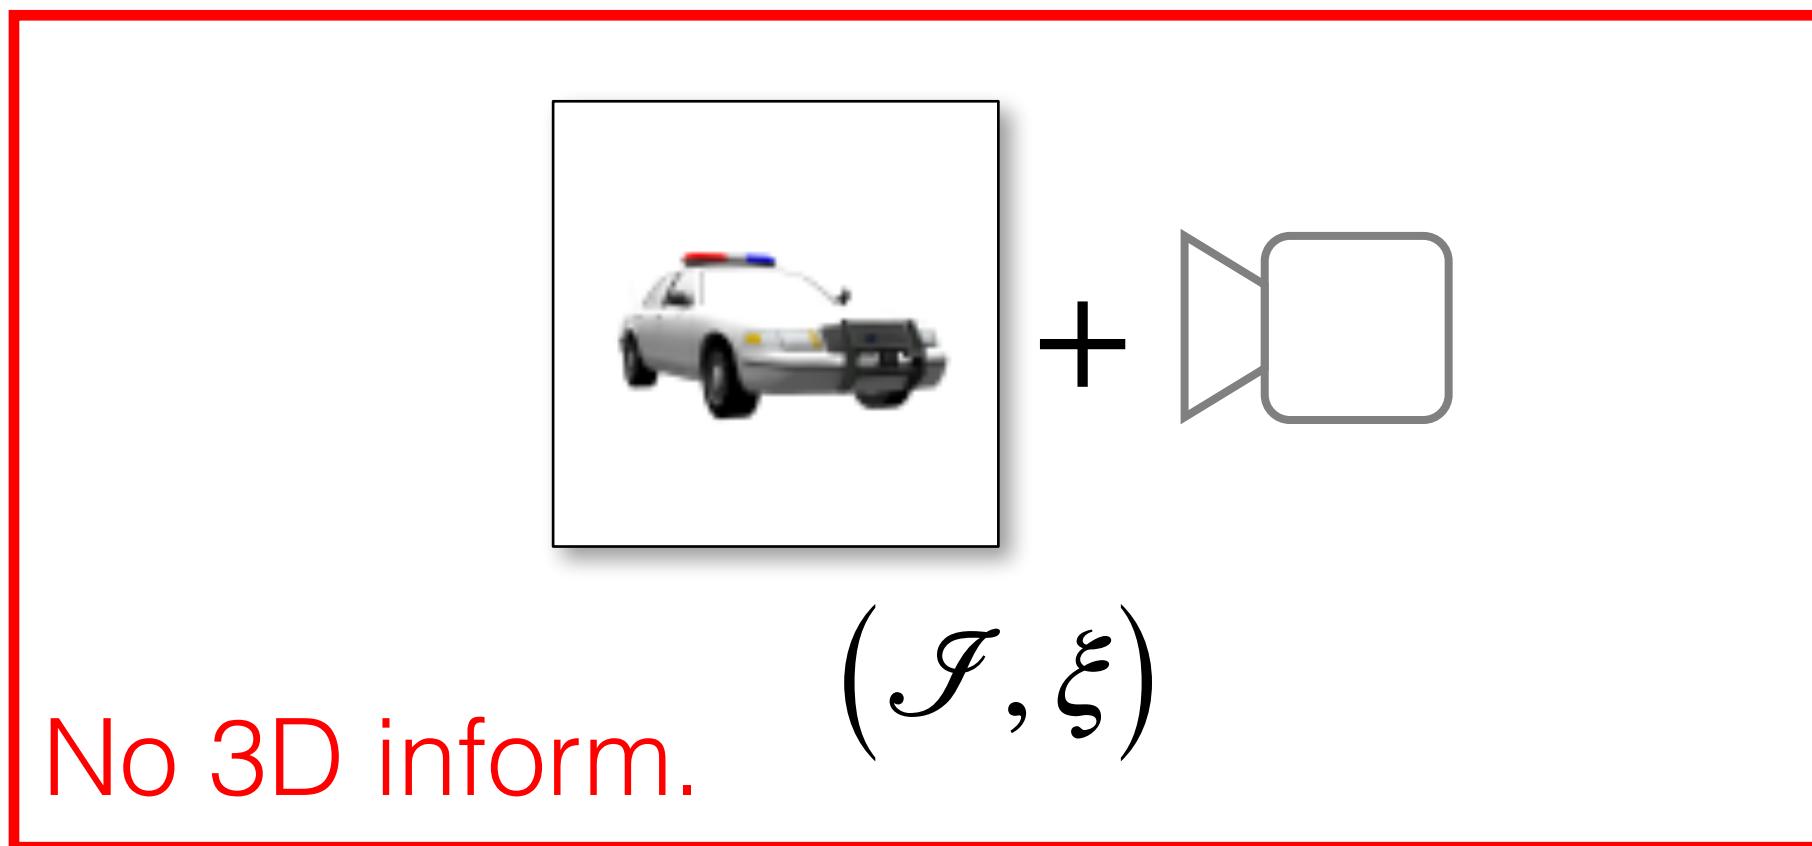
What if observations don't constrain scene representation?



Srn_ϕ

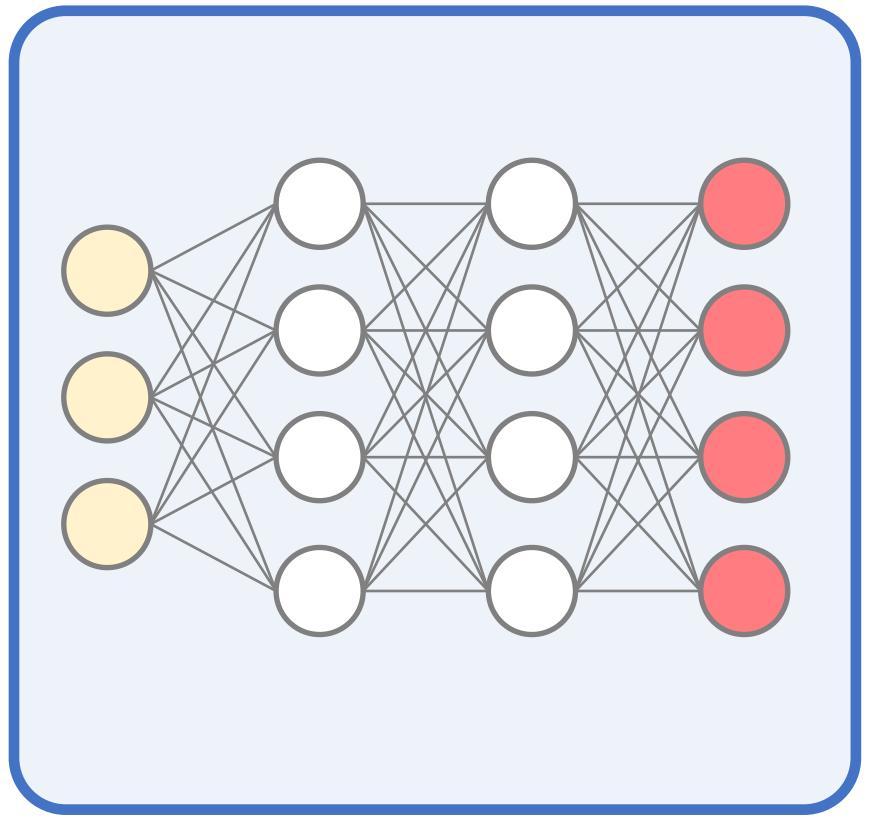


Render_θ

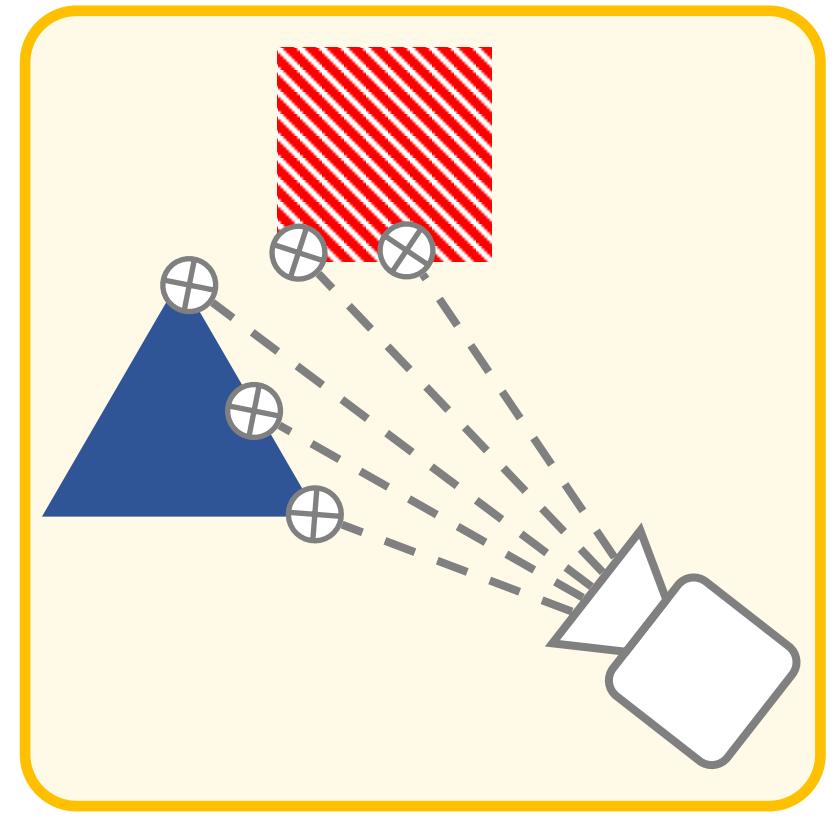


$$\underset{\phi, \theta}{\operatorname{argmin}} \| \text{Render}_\theta(\text{Srn}_\phi, \xi) - \mathcal{I} \|$$

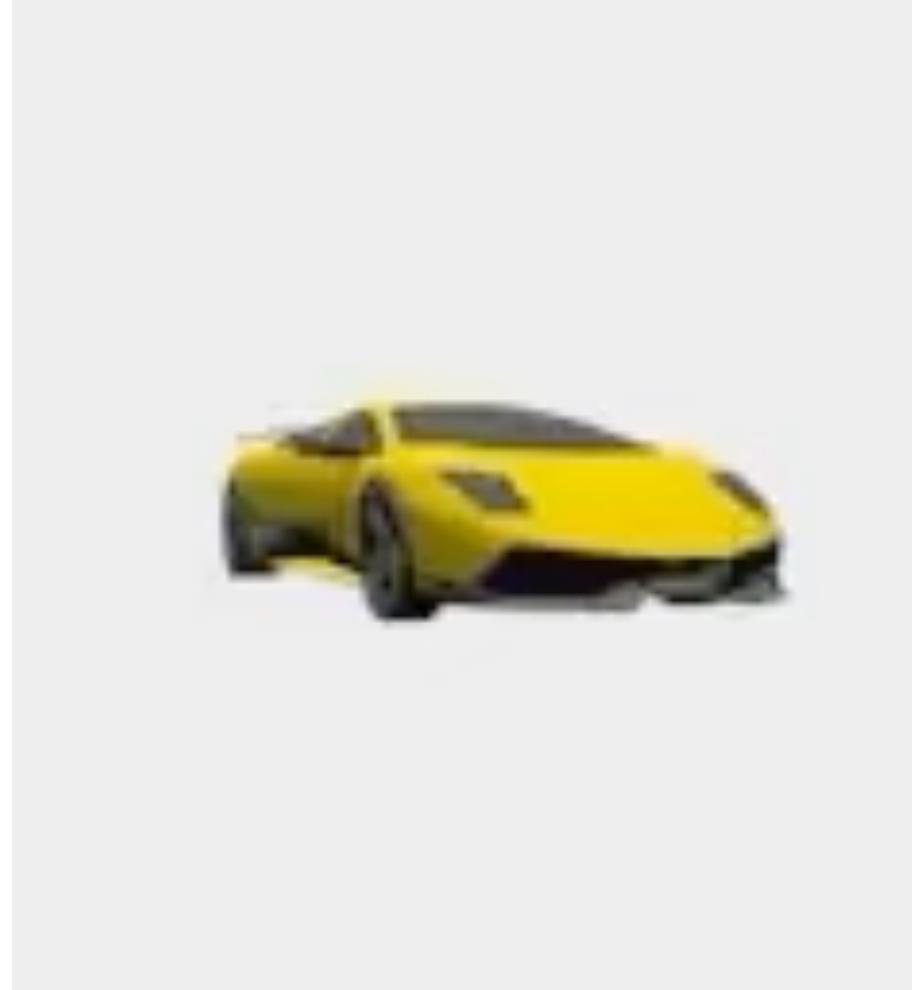
What if observations don't constrain scene representation?



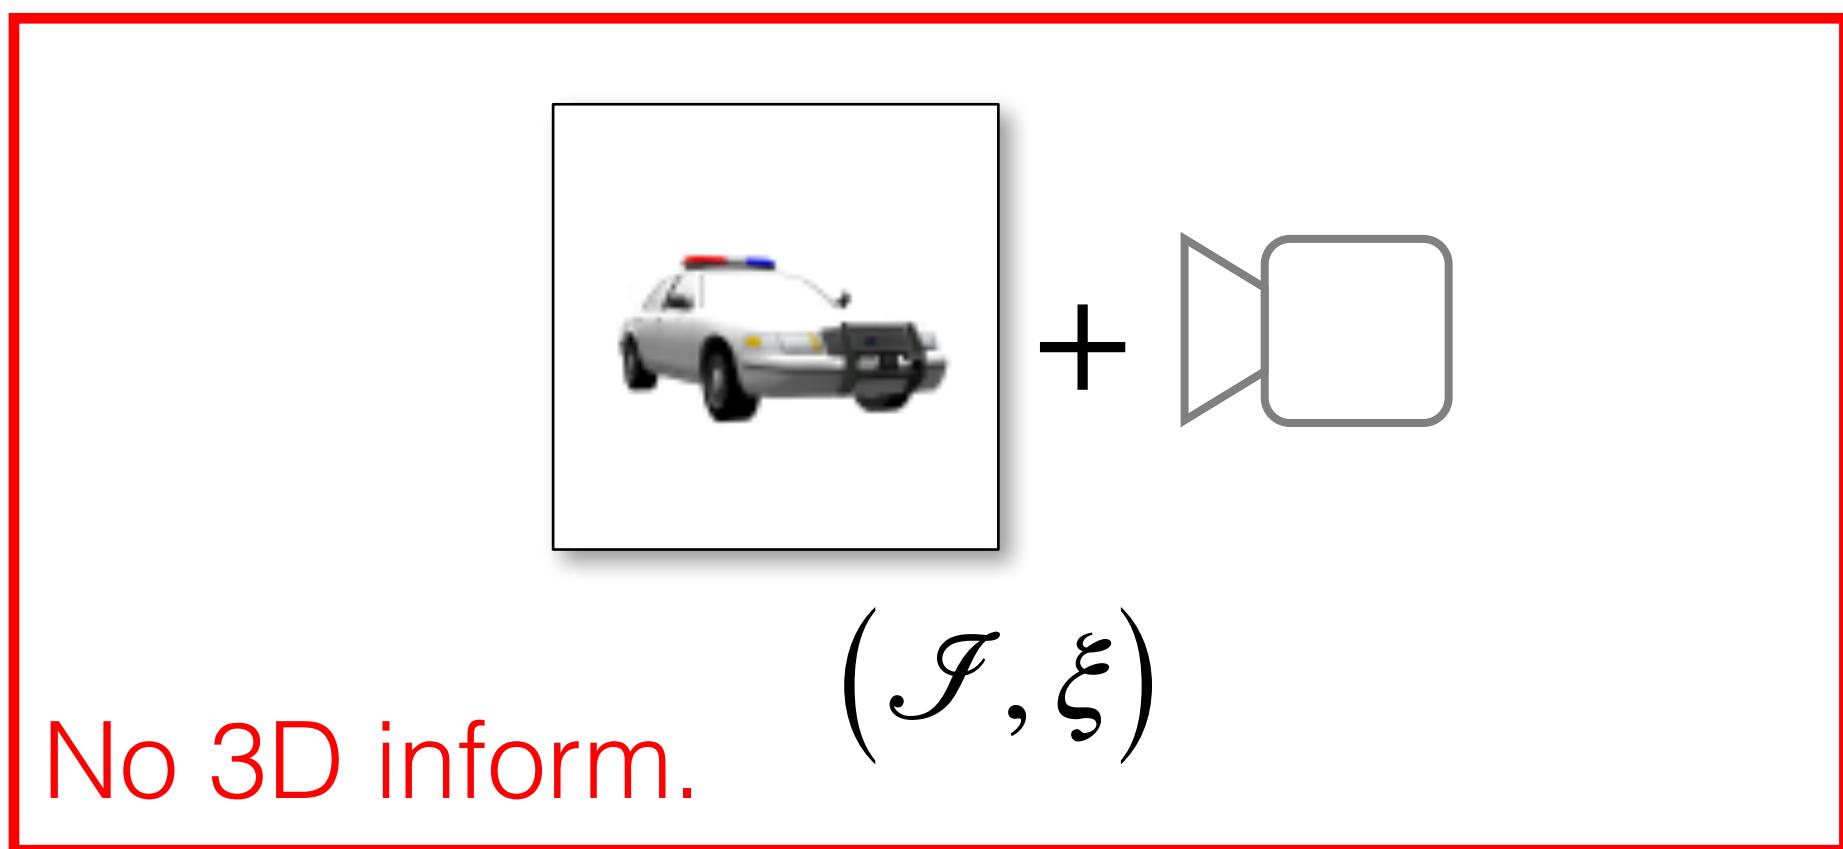
Srn_ϕ



Render_θ

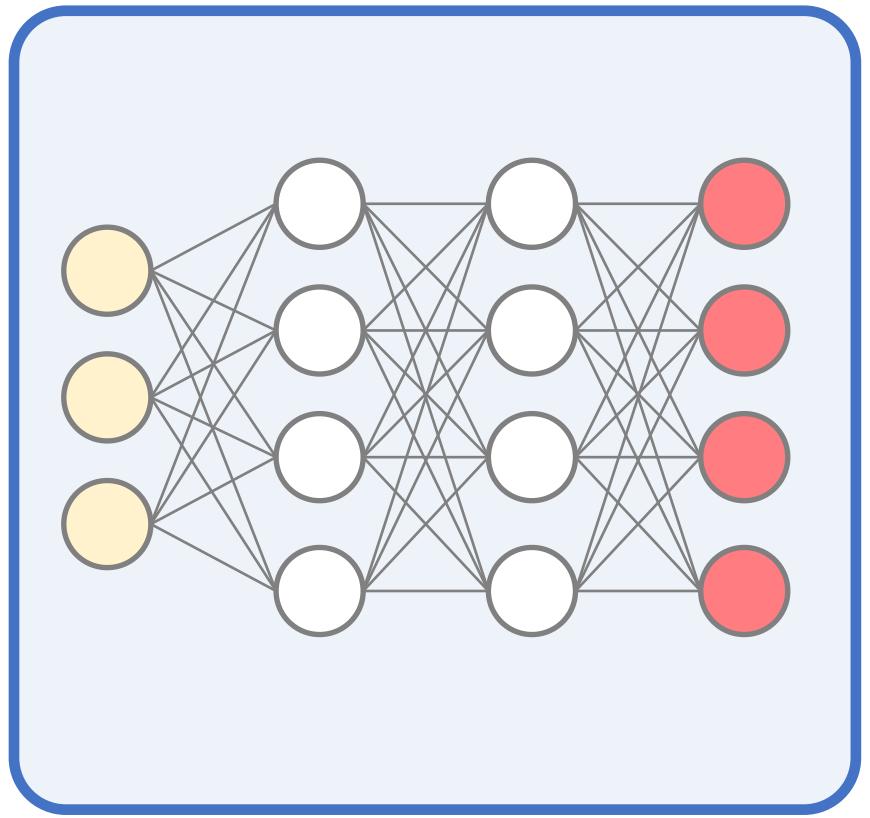


Input

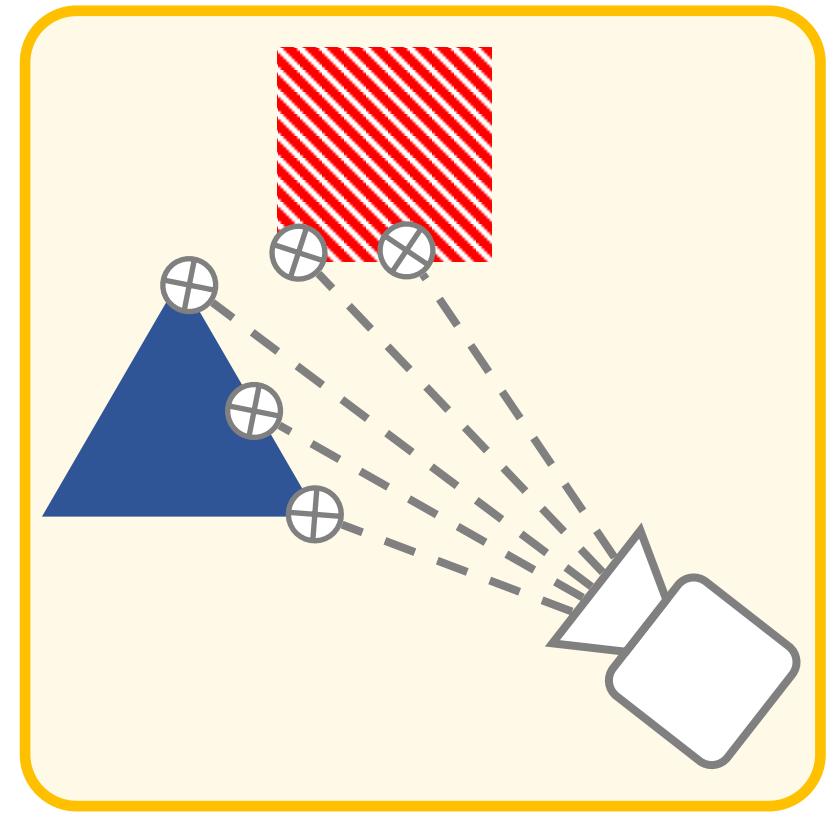


$$\operatorname{argmin}_{\phi, \theta} \| \text{Render}_\theta(\text{Srn}_\phi, \xi) - \mathcal{I} \|$$

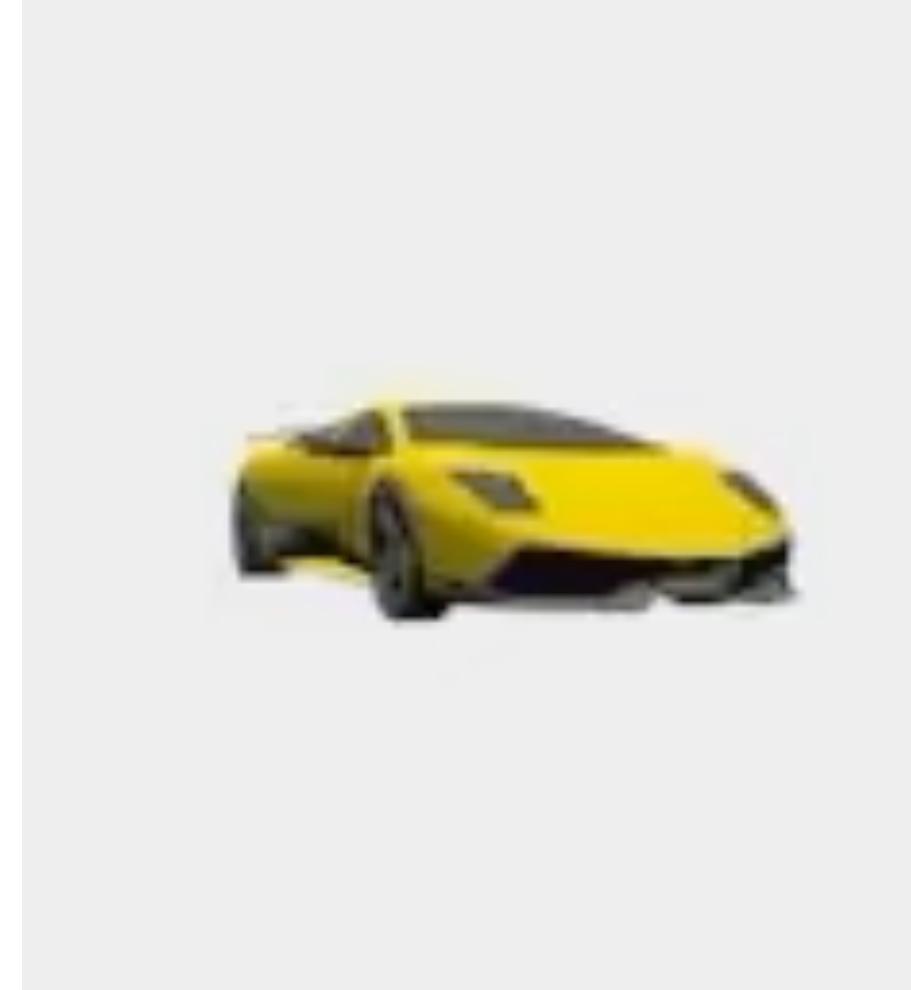
What if observations don't constrain scene representation?



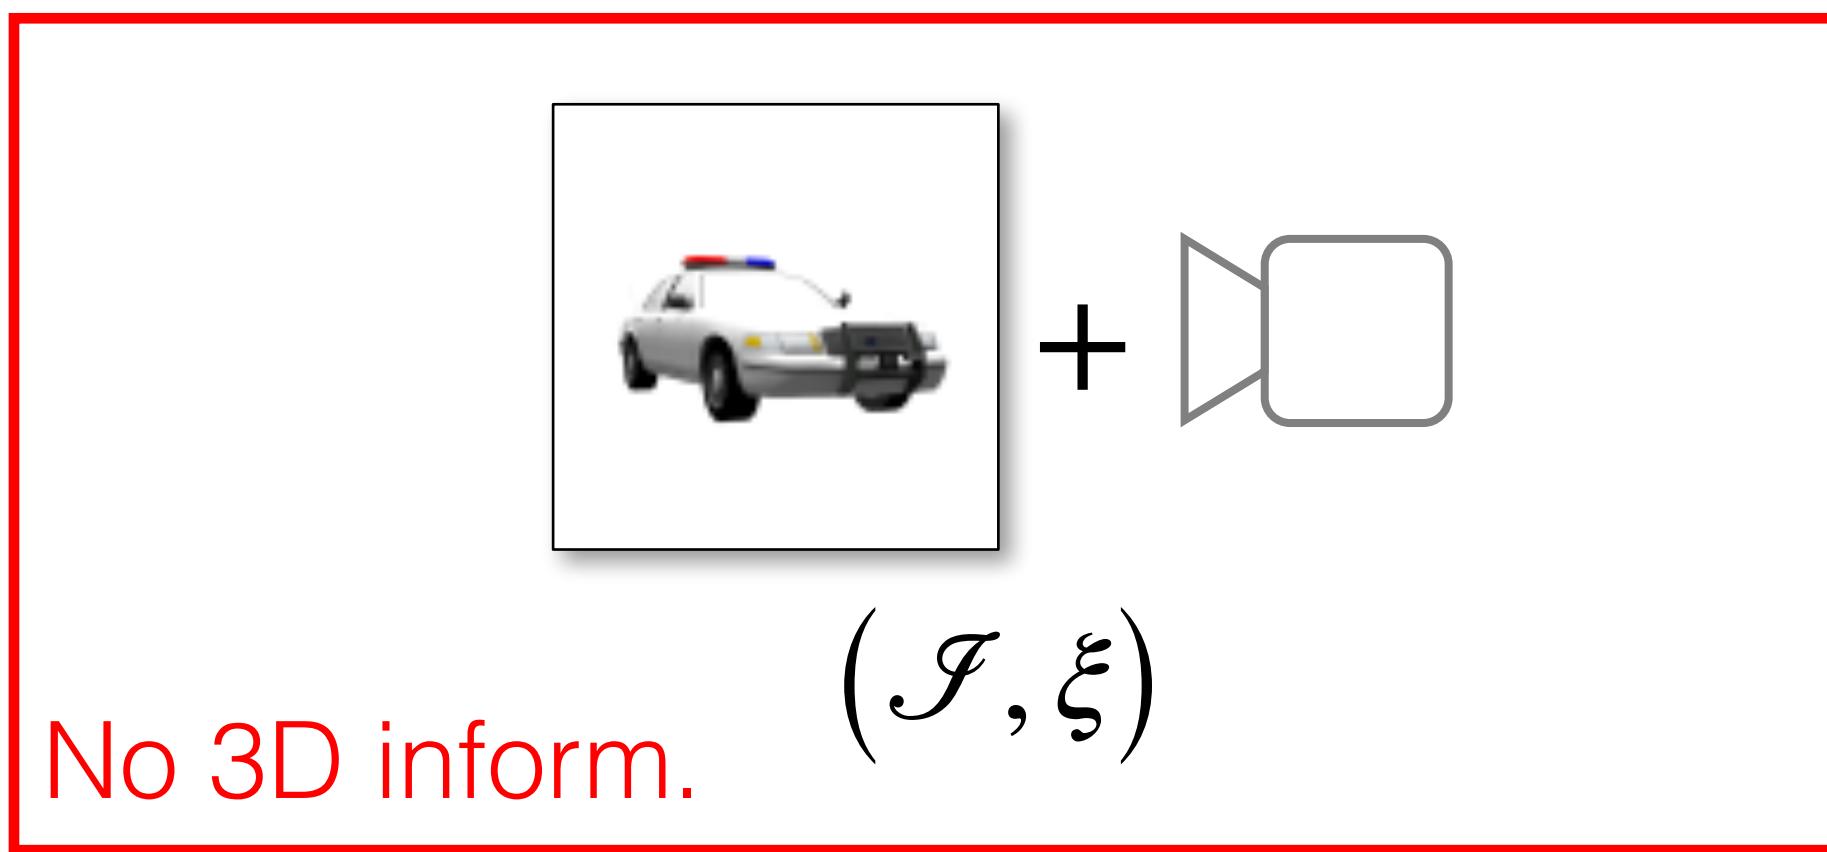
Srn_ϕ



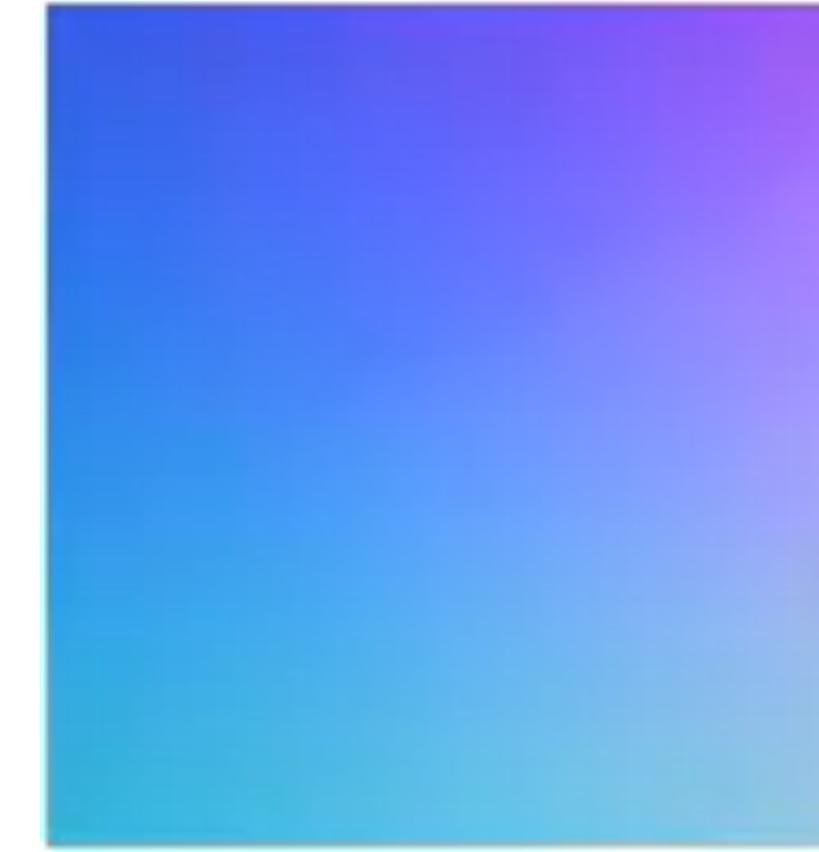
Render_θ



Input



$$\operatorname{argmin}_{\phi, \theta} \|\text{Render}_\theta(\text{Srn}_\phi, \xi) - \mathcal{I}\|$$



Normal map

RGB



GT

Today: Differentiable Rendering, AKA “Inverse Graphics”

